

LeetMastery

Study-Order Edition

Python Only · Priority-Grouped ·
Difficulty-First · 2×1 Landscape

331 questions · 9 rounds · Inline Quick Review ·

Chapter 2 Summary

High Easy → High Med → High Hard → Mid Easy → Mid Med →
Mid Hard → Low Easy → Low Med → Low Hard

p.1

Contents

Round 1 High · E Easy (43)

- ☐ #1 Two Sum
- ☐ #163 Missing Ranges
- ☐ #346 Moving Average from Data Stream
- ☐ #760 Find Anagram Mappings
- ☐ #1426 Counting Elements
- ☐ #1534 Count Good Triplets
- ☐ #383 Ransom Note
- ☐ #734 Sentence Similarity
- ☐ #1165 Single Row Keyboard
- ☐ #88 Merge Sorted Array
- ☐ #125 Valid Palindrome
- ☐ #283 Move Zeroes
- ☐ #408 Valid Word Abbreviation
- ☐ #977 Squares of a Sorted Array
- ☐ #169 Majority Element
- ☐ #217 Contains Duplicate
- ☐ #242 Valid Anagram
- ☐ #252 Meeting Rooms
- ☐ #1086 Largest Unique Number
- ☐ #1122 Relative Sort Array
- ☐ #35 Search Insert Position
- ☐ #69 Sqrt(x)
- ☐ #278 First Bad Version
- ☐ #374 Guess Number Higher or Lower
- ☐ #704 Binary Search
- ☐ #422 Valid Word Square

p.2

- ☐ #566 Reshape the Matrix
- ☐ #766 Toeplitz Matrix
- ☐ #867 Transpose Matrix
- ☐ #2373 Largest Local Values in a Matrix
- ☐ #100 Same Tree
- ☐ #101 Symmetric Tree
- ☐ #104 Maximum Depth of Binary Tree
- ☐ #108 Convert Sorted Array to Binary Search Tree
- ☐ #110 Balanced Binary Tree
- ☐ #112 Path Sum
- ☐ #226 Invert Binary Tree
- ☐ #270 Closest Binary Search Tree Value
- ☐ #501 Find Mode in Binary Search Tree
- ☐ #543 Diameter of Binary Tree
- ☐ #572 Subtree of Another Tree
- ☐ #463 Island Perimeter
- ☐ #733 Flood Fill

Round 2 High - M Medium (116)

- ☐ #57 Insert Interval
- ☐ #238 Product of Array Except Self
- ☐ #281 Zigzag Iterator
- ☐ #307 Range Sum Query – Mutable
- ☐ #454 4Sum II
- ☐ #525 Contiguous Array
- ☐ #560 Subarray Sum Equals K
- ☐ #1010 Pairs of Songs With Total Durations Divisible by 60
- ☐ #1272 Remove Interval
- ☐ #1429 First Unique Number
- ☐ #3159 Find Occurrences of an Element in an Array
- ☐ #8 String to Integer (atoi)

- ☐ #249 Group Shifted Strings
- ☐ #271 Encode and Decode Strings
- ☐ #635 Design Log Storage System
- ☐ #1138 Alphabet Board Path
- ☐ #11 Container With Most Water
- ☐ #15 3Sum
- ☐ #16 3Sum Closest
- ☐ #31 Next Permutation
- ☐ #75 Sort Colors
- ☐ #161 One Edit Distance
- ☐ #167 Two Sum II – Input Array Is Sorted
- ☐ #186 Reverse Words in a String II
- ☐ #189 Rotate Array
- ☐ #532 K-diff Pairs in an Array
- ☐ #923 3Sum With Multiplicity
- ☐ #986 Interval List Intersections
- ☐ #1055 Shortest Way to Form String
- ☐ #2563 Count the Number of Fair Pairs
- ☐ #3 Longest Substring Without Repeating Characters
- ☐ #159 Longest Substring with At Most Two Distinct Characters
- ☐ #340 Longest Substring with At Most K Distinct Characters
- ☐ #424 Longest Repeating Character Replacement
- ☐ #438 Find All Anagrams in a String
- ☐ #487 Max Consecutive Ones II
- ☐ #1100 Find K-Length Substrings with No Repeated Characters
- ☐ #1248 Count Number of Nice Subarrays
- ☐ #49 Group Anagrams
- ☐ #56 Merge Intervals
- ☐ #1152 Analyze User Website Visit Pattern
- ☐ #33 Search in Rotated Sorted Array

- ☐ #34 Find First and Last Position of Element in Sorted Array
- ☐ #153 Find Minimum in Rotated Sorted Array
- ☐ #300 Longest Increasing Subsequence
- ☐ #362 Design Hit Counter
- ☐ #528 Random Pick with Weight
- ☐ #852 Peak Index in a Mountain Array
- ☐ #981 Time Based Key-Value Store
- ☐ #1011 Capacity to Ship Packages Within D Days
- ☐ #1060 Missing Element in Sorted Array
- ☐ #1062 Longest Repeating Substring
- ☐ #1533 Find the Index of the Large Integer
- ☐ #36 Valid Sudoku
- ☐ #48 Rotate Image
- ☐ #54 Spiral Matrix
- ☐ #59 Spiral Matrix II
- ☐ #64 Minimum Path Sum
- ☐ #73 Set Matrix Zeroes
- ☐ #74 Search a 2D Matrix
- ☐ #221 Maximal Square
- ☐ #311 Sparse Matrix Multiplication
- ☐ #498 Diagonal Traverse
- ☐ #531 Lonely Pixel I
- ☐ #723 Candy Crush
- ☐ #835 Image Overlap
- ☐ #1198 Find Smallest Common Element in All Rows
- ☐ #98 Validate Binary Search Tree
- ☐ #102 Binary Tree Level Order Traversal
- ☐ #103 Binary Tree Zigzag Level Order Traversal
- ☐ #105 Construct Binary Tree from Preorder and Inorder Traversal
- ☐ #199 Binary Tree Right Side View

p.5

- ☐ #230 Kth Smallest Element in a BST
- ☐ #235 Lowest Common Ancestor of a Binary Search Tree
- ☐ #236 Lowest Common Ancestor of a Binary Tree
- ☐ #250 Count Univalued Subtrees
- ☐ #285 Inorder Successor in BST
- ☐ #298 Binary Tree Longest Consecutive Sequence
- ☐ #314 Binary Tree Vertical Order Traversal
- ☐ #333 Largest BST Subtree
- ☐ #366 Find Leaves of Binary Tree
- ☐ #437 Path Sum III
- ☐ #545 Boundary of Binary Tree
- ☐ #549 Binary Tree Longest Consecutive Sequence II
- ☐ #582 Kill Process
- ☐ #662 Maximum Width of Binary Tree
- ☐ #863 All Nodes Distance K in Binary Tree
- ☐ #1120 Maximum Average Subtree
- ☐ #1490 Clone N-ary Tree
- ☐ #1522 Diameter of N-Ary Tree
- ☐ #1650 Lowest Common Ancestor of a Binary Tree III
- ☐ #130 Surrounded Regions
- ☐ #133 Clone Graph
- ☐ #200 Number of Islands
- ☐ #207 Course Schedule
- ☐ #210 Course Schedule II
- ☐ #261 Graph Valid Tree
- ☐ #310 Minimum Height Trees
- ☐ #323 Number of Connected Components in an Undirected Graph
- ☐ #417 Pacific Atlantic Water Flow
- ☐ #490 The Maze
- ☐ #694 Number of Distinct Islands

p.6

- ☐ #695 Max Area of Island
- ☐ #721 Accounts Merge
- ☐ #785 Is Graph Bipartite?
- ☐ #1236 Web Crawler
- ☐ #1242 Web Crawler Multithreaded
- ☐ #128 Longest Consecutive Sequence
- ☐ #277 Find the Celebrity
- ☐ #1136 Parallel Courses
- ☐ #286 Walls and Gates
- ☐ #322 Coin Change
- ☐ #542 01 Matrix
- ☐ #994 Rotting Oranges
- ☐ #1197 Minimum Knight Moves
- ☐ #1730 Shortest Path to Get Food

Round 3 High · H Hard (23)

- ☐ #41 First Missing Positive
- ☐ #76 Minimum Window Substring
- ☐ #4 Median of Two Sorted Arrays
- ☐ #315 Count of Smaller Numbers After Self
- ☐ #410 Split Array Largest Sum
- ☐ #668 Kth Smallest Number in Multiplication Table
- ☐ #710 Random Pick with Blacklist
- ☐ #774 Minimize Max Distance to Gas Station
- ☐ #1235 Maximum Profit in Job Scheduling
- ☐ #1074 Number of Submatrices That Sum to Target
- ☐ #124 Binary Tree Maximum Path Sum
- ☐ #297 Serialize and Deserialize Binary Tree
- ☐ #269 Alien Dictionary
- ☐ #329 Longest Increasing Path in a Matrix
- ☐ #685 Redundant Connection II

p.7

- ☐ #1036 Escape a Large Maze
- ☐ #1192 Critical Connections in a Network
- ☐ #305 Number of Islands II
- ☐ #1153 String Transforms Into Another String
- ☐ #127 Word Ladder
- ☐ #317 Shortest Distance from All Buildings
- ☐ #773 Sliding Puzzle
- ☐ #815 Bus Routes

Round 4 Mid · E Easy (12)

- ☐ #21 Merge Two Sorted Lists
- ☐ #141 Linked List Cycle
- ☐ #206 Reverse Linked List
- ☐ #234 Palindrome Linked List
- ☐ #706 Design HashMap
- ☐ #876 Middle of the Linked List
- ☐ #1474 Delete N Nodes After M Nodes of Linked List
- ☐ #20 Valid Parentheses
- ☐ #232 Implement Queue using Stacks
- ☐ #844 Backspace String Compare
- ☐ #14 Longest Common Prefix
- ☐ #409 Longest Palindrome

Round 5 Mid · M Medium (56)

- ☐ #2 Add Two Numbers
- ☐ #19 Remove Nth Node From End of List
- ☐ #24 Swap Nodes in Pairs
- ☐ #61 Rotate List
- ☐ #146 LRU Cache
- ☐ #148 Sort List
- ☐ #328 Odd Even Linked List

p.8

- ☐ #369 Plus One Linked List
- ☐ #430 Flatten a Multilevel Doubly Linked List
- ☐ #622 Design Circular Queue
- ☐ #707 Design Linked List
- ☐ #708 Insert into a Sorted Circular Linked List
- ☐ #1171 Remove Zero Sum Consecutive Nodes from Linked List
- ☐ #143 Reorder List
- ☐ #150 Evaluate Reverse Polish Notation
- ☐ #155 Min Stack
- ☐ #173 Binary Search Tree Iterator
- ☐ #227 Basic Calculator II
- ☐ #255 Verify Preorder Sequence in Binary Search Tree
- ☐ #394 Decode String
- ☐ #439 Ternary Expression Parser
- ☐ #484 Find Permutation
- ☐ #735 Asteroid Collision
- ☐ #739 Daily Temperatures
- ☐ #962 Maximum Width Ramp
- ☐ #1214 Two Sum BSTs
- ☐ #1762 Buildings With an Ocean View
- ☐ #215 Kth Largest Element in an Array
- ☐ #253 Meeting Rooms II
- ☐ #347 Top K Frequent Elements
- ☐ #505 The Maze II
- ☐ #621 Task Scheduler
- ☐ #658 Find K Closest Elements
- ☐ #787 Cheapest Flights Within K Stops
- ☐ #973 K Closest Points to Origin
- ☐ #1057 Campus Bikes
- ☐ #1167 Minimum Cost to Connect Sticks

- ☐ #1424 Diagonal Traverse II
- ☐ #139 Word Break
- ☐ #208 Implement Trie (Prefix Tree)
- ☐ #211 Design Add and Search Words Data Structure
- ☐ #616 Add Bold Tag in String
- ☐ #692 Top K Frequent Words
- ☐ #720 Longest Word in Dictionary
- ☐ #1166 Design File System
- ☐ #17 Letter Combinations of a Phone Number
- ☐ #22 Generate Parentheses
- ☐ #39 Combination Sum
- ☐ #46 Permutations
- ☐ #79 Word Search
- ☐ #113 Path Sum II
- ☐ #134 Gas Station
- ☐ #179 Largest Number
- ☐ #280 Wiggle Sort
- ☐ #954 Array of Doubled Pairs
- ☐ #2131 Longest Palindrome by Concatenating Two Letter Words

Round 6 Mid · H Hard (28)

- ☐ #25 Reverse Nodes in k-Group
- ☐ #432 All O'one Data Structure
- ☐ #460 LFU Cache
- ☐ #32 Longest Valid Parentheses
- ☐ #42 Trapping Rain Water
- ☐ #84 Largest Rectangle in Histogram
- ☐ #224 Basic Calculator
- ☐ #239 Sliding Window Maximum
- ☐ #716 Max Stack
- ☐ #772 Basic Calculator III

- ☐ #895 Maximum Frequency Stack
- ☐ #23 Merge k Sorted Lists
- ☐ #218 The Skyline Problem
- ☐ #272 Closest Binary Search Tree Value II
- ☐ #295 Find Median from Data Stream
- ☐ #358 Rearrange String k Distance Apart
- ☐ #499 The Maze III
- ☐ #632 Smallest Range Covering Elements from K Lists
- ☐ #759 Employee Free Time
- ☐ #1425 Constrained Subsequence Sum
- ☐ #212 Word Search II
- ☐ #336 Palindrome Pairs
- ☐ #588 Design In-Memory File System
- ☐ #642 Design Search Autocomplete System
- ☐ #37 Sudoku Solver
- ☐ #51 N-Queens
- ☐ #126 Word Ladder II
- ☐ #996 Number of Squareful Arrays

Round 7 Low · E Easy (14)

- ☐ #70 Climbing Stairs
- ☐ #121 Best Time to Buy and Sell Stock
- ☐ #67 Add Binary
- ☐ #136 Single Number
- ☐ #190 Reverse Bits
- ☐ #191 Number of 1 Bits
- ☐ #268 Missing Number
- ☐ #338 Counting Bits
- ☐ #9 Palindrome Number
- ☐ #13 Roman to Integer
- ☐ #504 Base 7

p.11

- ☐ #1056 Confusing Number
- ☐ #1228 Missing Number in Arithmetic Progression
- ☐ #1427 Perform String Shifts

Round 8 Low · M Medium (32)

- ☐ #5 Longest Palindromic Substring
- ☐ #53 Maximum Subarray
- ☐ #55 Jump Game
- ☐ #62 Unique Paths
- ☐ #72 Edit Distance
- ☐ #91 Decode Ways
- ☐ #152 Maximum Product Subarray
- ☐ #198 House Robber
- ☐ #256 Paint House
- ☐ #309 Best Time to Buy and Sell Stock with Cooldown
- ☐ #377 Combination Sum IV
- ☐ #416 Partition Equal Subset Sum
- ☐ #435 Non-overlapping Intervals
- ☐ #516 Longest Palindromic Subsequence
- ☐ #576 Out of Boundary Paths
- ☐ #1143 Longest Common Subsequence
- ☐ #78 Subsets
- ☐ #90 Subsets II
- ☐ #287 Find the Duplicate Number
- ☐ #1506 Find Root of N-Ary Tree
- ☐ #7 Reverse Integer
- ☐ #50 Pow(x, n)
- ☐ #380 Insert Delete GetRandom O(1)
- ☐ #1017 Convert to Base -2
- ☐ #3160 Find the Number of Distinct Colors Among the Balls
- ☐ #2627 Debounce

p.12

- ☐ #2628 JSON Deep Equal
- ☐ #2630 Memoize
- ☐ #2633 Convert Object to JSON String
- ☐ #2636 Promise Pool
- ☐ #2675 Array of Objects to Matrix
- ☐ #2700 Differences Between Two Objects

Round 9 Low · H Hard (7)

- ☐ #10 Regular Expression Matching
- ☐ #115 Distinct Subsequences
- ☐ #123 Best Time to Buy and Sell Stock III
- ☐ #132 Palindrome Partitioning II
- ☐ #312 Burst Balloons
- ☐ #1000 Minimum Cost to Merge Stones
- ☐ #68 Text Justification

p.13

Round 1 · High · Easy

Round 1

High Priority · Easy

43 questions · 8 patterns

Arrays & Hashing #1 #163 #346 #760 #1426 #1534

String #383 #734 #1165

Two Pointers #88 #125 #283 #408 #977

Sorting #169 #217 #242 #252 #1086 #1122

Binary Search #35 #69 #278 #374 #704

Matrix #422 #566 #766 #867 #2373

Trees & BST #100 #101 #104 #108 #110 #112 #226 #270
#501 #543 #572

DFS #463 #733

p.14

Arrays & Hashing

Round 1 · High · Easy · 6 questions

Round 1 · High · Easy

p.15

Arrays & Hashing

#1 Two Sum

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: Grind 169

Problem

Description

Given an array of integers `nums` and an integer `target`, return indices of the two numbers such that they add up to `target`.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Round 1 · High · Easy

p.16

Input: nums = [3,3], target = 6
Output: [0,1]

Constraints:

- $2 \leq \text{nums.length} \leq 104$
- $-109 \leq \text{nums}[i] \leq 109$
- $-109 \leq \text{target} \leq 109$
- Only one valid answer exists.

Follow-up: Can you come up with an algorithm that is less than $O(n^2)$ time complexity?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def twoSum(self, nums: list[int], target: int) -> list[int]:
        numToIndex = {}

        for i, num in enumerate(nums):
            if target - num in numToIndex:
                return numToIndex[target - num], i
            numToIndex[num] = i
```

• LeetCode

```
class Solution:
    def twoSum(self, nums: List[int], target: int) -> List[int]:
        d = {}
        for i, x in enumerate(nums):
            if (y := target - x) in d:
                return [d[y], i]
            d[x] = i
```

• SimplyLeet

```
def two_sum(nums, target):
    # Create a dictionary to store the number and its index
    num_to_index = {}
    # Iterate over the list with index
    for index, num in enumerate(nums):
        # Calculate the complement needed to reach the target
        complement = target - num
        # Check if the complement is in the dictionary
        if complement in num_to_index:
            # Return the index of the complement and the current index
```

```

        return [num_to_index[complement],
index]
    # Store the number with its index
    num_to_index[num] = index
    # Since the problem guarantees one solution, we
don't need a return statement here.

# Example usage:
# print(two_sum([2, 7, 11, 15], 9)) # Output: [0,
1]

```

• LC.ca

```

class Solution:
    def twoSum(self, nums: List[int], target: int)
-> List[int]:
        m = {}
        for i, x in enumerate(nums):
            y = target - x
            if y in m:
                return [m[y], i]
            m[x] = i

```

◆ Quick Review

Key Insights

- Use a hash table to store the elements and their indices for fast lookup.
- For each element, calculate its complement: target minus the current element.

- Check if the complement exists in the hash table.
- If found, return the indices of the current element and its complement.
- This approach only requires a single pass through the array.

Space & Time

Time Complexity: $O(n)$ – We traverse the list only once.

Space Complexity: $O(n)$ – In the worst-case, we store all elements in the hash table.

Solution

We use a hash table (dictionary) to map each array element to its index as we iterate. For each number, we compute the complement (target – number) and check if it already exists in our hash table. If the complement is found, we return the indices. This one-pass solution ensures linear time complexity and avoids the need for nested loops.

Arrays & Hashing

#163 Missing Ranges

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array

Lists: Premium 98

Problem

Description

You are given an inclusive range $[lower, upper]$ and a sorted unique integer array `nums`, where all elements are within the inclusive range.

A number x is considered missing if x is in the range $[lower, upper]$ and x is not in `nums`.

Return the shortest sorted list of ranges that exactly covers all the missing numbers. That is, no element of `nums` is included in any of the ranges, and each missing number is covered by one of the ranges.

Example 1:

Input: `nums = [0,1,3,50,75]`, `lower = 0`, `upper = 99`Output: `[[2,2],[4,49],[51,74],[76,99]]`

Explanation: The ranges are:

`[2,2]``[4,49]``[51,74]``[76,99]`

Round 1 · High · Easy

p.21

Example 2:

Input: `nums = [-1]`, `lower = -1`, `upper = -1`Output: `[]`

Explanation: There are no missing ranges since there are no missing numbers.

Constraints:

- $-109 \leq lower \leq upper \leq 109$
- $0 \leq \text{nums.length} \leq 100$
- $lower \leq \text{nums}[i] \leq upper$
- All the values of `nums` are unique.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findMissingRanges(
        self,
        nums: list[int],
        lower: int,
        upper: int,
    ) -> list[list[int]]:
        def getRange(lo: int, hi: int) -> list[int]:
            if lo == hi:
                return [lo, lo]
```

Round 1 · High · Easy

p.22

```

    return [lo, hi]

if not nums:
    return [getRange(lower, upper)]

ans = []

if nums[0] > lower:
    ans.append(getRange(lower, nums[0] - 1))

for prev, curr in zip(nums, nums[1:]):
    if curr > prev + 1:
        ans.append(getRange(prev + 1, curr - 1))

if nums[-1] < upper:
    ans.append(getRange(nums[-1] + 1, upper))

return ans

```

- LeetCode

```

class Solution:
    def findMissingRanges(
        self, nums: List[int], lower: int, upper:
        int
    ) -> List[List[int]]:
        n = len(nums)
        if n == 0:
            return [[lower, upper]]
        ans = []
        if nums[0] > lower:
            ans.append([lower, nums[0] - 1])

        for a, b in pairwise(nums):
            if b - a > 1:
                ans.append([a + 1, b - 1])
        if nums[-1] < upper:
            ans.append([nums[-1] + 1, upper])
        return ans

```

- SimplyLeet

Python solution for Missing Ranges problem

```
def findMissingRanges(nums, lower, upper):
    missing_ranges = [] # List to store the missing
    range intervals
    prev = lower - 1 # Initialize previous value to
    one less than lower

    # Append a sentinel at the end for easier gap
    handling
    nums.append(upper + 1)

    # Loop through each number, including the
    sentinel

    for num in nums:
        # If there is a gap between previous and
        current number
        if num - prev >= 2:
            # If gap is of single number, add that
            number as a range
            if num - prev == 2:
                missing_ranges.append([prev + 1,
                prev + 1])
            else:
                missing_ranges.append([prev + 1,
                num - 1])
            prev = num # Update previous to current
    return missing_ranges
```

Example usage:

```
print(findMissingRanges([0,1,3,50,75], 0, 99))
```

• LC.ca

```
class Solution:
    def findMissingRanges(
        self, nums: List[int], lower: int, upper:
        int
    ) -> List[List[int]]:
        n = len(nums)
        if n == 0:
            return [[lower, upper]]
        ans = []
        if nums[0] > lower:
            ans.append([lower, nums[0] - 1])

        for a, b in pairwise(nums):
            if b - a > 1:
                ans.append([a + 1, b - 1])
        if nums[-1] < upper:
            ans.append([nums[-1] + 1, upper])
        return ans
```

◆ Quick Review

Key Insights

- Use a pointer (or iterate index) to traverse the nums array while keeping track of the previous number considered.

- Check for missing numbers before the first element, between consecutive elements, and after the last element.
- When a gap is found, convert it to a range string that is either a single number (if the gap is exactly one number) or an interval.
- Handle edge cases when the nums array is empty.

Space & Time

Time Complexity: $O(n)$, where n is the length of the nums array, since we iterate through the array once.

Space Complexity: $O(1)$ additional space (ignoring the output list), as we only use a few extra variables.

Solution

The solution leverages a single pass through the input array and calculates the missing ranges by comparing the expected number with the current number in nums. First, initialize a variable "prev" to one less than the lower bound so that missing numbers from the very start of the range can be handled uniformly. For each number in the array, if there is a gap between prev and the current number (i.e., current number is at least 2 more than prev), then add the missing range (prev+1 to current number-1) to the output list. Update "prev" to the current number and finally handle any gap between the last number and the upper bound. The primary data structures used are the list for the output; the approach is iterative.

Arrays & Hashing

#346 Moving Average from Data Stream

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Design · Queue · Data Stream

Lists: Premium 98

Problem

Description

Given a stream of integers and a window size, calculate the moving average of all integers in the sliding window.

Implement the MovingAverage class:

- MovingAverage(int size) Initializes the object with the size of the window size.
- double next(int val) Returns the moving average of the last size values of the stream.

Example 1:

```
Input
["MovingAverage", "next", "next", "next", "next"]
[[3], [1], [10], [3], [5]]
Output
[null, 1.0, 5.5, 4.66667, 6.0]
```

Explanation

```
MovingAverage movingAverage = new MovingAverage(3);
```

Constraints:

- $1 \leq \text{size} \leq 1000$
- $-105 \leq \text{val} \leq 105$
- At most 104 calls will be made to next.

Community Solutions (Python)

• WalkCC

```
class MovingAverage:
    def __init__(self, size: int):
        self.size = size
        self.sum = 0
        self.q = collections.deque()

    def next(self, val: int) -> float:
        if len(self.q) == self.size:
            self.sum -= self.q.popleft()
            self.sum += val

        self.q.append(val)
        return self.sum / len(self.q)
```

• LeetCode

class MovingAverage:

```
    def __init__(self, size: int):
        self.s = 0
        self.data = [0] * size
        self.cnt = 0

    def next(self, val: int) -> float:
        i = self.cnt % len(self.data)
        self.s += val - self.data[i]
```

```
        self.data[i] = val
        self.cnt += 1
        return self.s / min(self.cnt,
                             len(self.data))
```

```
# Your MovingAverage object will be instantiated
and called as such:
# obj = MovingAverage(size)
# param_1 = obj.next(val)
```

```
class MovingAverage:
    def __init__(self, size: int):
        self.n = size
        self.s = 0
        self.q = deque()

    def next(self, val: int) -> float:
        if len(self.q) == self.n:
            self.s -= self.q.popleft()
        self.q.append(val)
```

```
        self.s += val
        return self.s / len(self.q)
```

Your MovingAverage object will be instantiated and called as such:

```
# obj = MovingAverage(size)
# param_1 = obj.next(val)
```

- SimplyLeet

Python solution for Moving Average from Data Stream

from collections import deque # Import deque for efficient queue operations

```
class MovingAverage:
    def __init__(self, size):
        # Initialize with the maximum window size
        self.size = size # Maximum number of
        elements in the window
        self.queue = deque() # Queue to store
        current window values
        self.current_sum = 0 # Running sum of the
        current window
```

```
    def next(self, val):
        # Add new value to the window
        if len(self.queue) == self.size:
            # If the window is full, remove the
            oldest element and update the sum
            removed_val = self.queue.popleft() #
            Remove the front element
            self.current_sum -= removed_val #
            Subtract its value from the running sum
            # Append the new value to the queue and
            update the sum
            self.queue.append(val)
            self.current_sum += val
```

```

        # Calculate and return the moving average
        by dividing the running sum by the current window
        size
        return self.current_sum / len(self.queue)

```

• LC.ca

```

class MovingAverage:
    def __init__(self, size: int):
        self.arr = [0] * size
        self.s = 0
        self.cnt = 0

    def next(self, val: int) -> float:
        idx = self.cnt % len(self.arr) # circular
        array
        self.s += val - self.arr[idx]
        self.arr[idx] = val

```

```

        self.cnt += 1
        return self.s / min(self.cnt,
len(self.arr))

```

Your MovingAverage object will be instantiated and called as such:

```

# obj = MovingAverage(size)
# param_1 = obj.next(val)

```

#####

```

from collections import deque

```

with no sum variable, calculate on the fly

```

class MovingAverage:

```

```

    def __init__(self, size: int):
        self.queue = deque()
        self.size = size
        self.avg = 0

```

```

    def next(self, val: int) -> float:

```

```

        if len(self.queue) < self.size:
            self.queue.append(val)
            self.avg = sum(self.queue) /
len(self.queue)
            return self.avg
        else:
            head = self.queue.popleft()
            self.queue.append(val)
            minus = head / self.size
            add = val / self.size
            self.avg = self.avg + add - minus

```

```

        return self.avg

#####

from collections import deque

class MovingAverage(object):

    def __init__(self, size):
        """

```

```

        Initialize your data structure here.
        :type size: int
        """

        self.windowSize = size
        self.windowSum = 0.0
        self.data = deque([])

    def next(self, val):
        """
        :type val: int

```

```

        :rtype: float
        """

        self.windowSum += val
        data = self.data

        leftTop = 0
        if len(data) >= self.windowSize:
            leftTop = data.popleft()
            data.append(val)

```

```

        self.windowSum -= leftTop
        if len(data) < self.windowSize:
            return self.windowSum / len(data)
        return self.windowSum / self.windowSize

```

```

# Your MovingAverage object will be instantiated
and called as such:
# obj = MovingAverage(size)
# param_1 = obj.next(val)

```

◆ Quick Review

Key Insights

- Use a queue (or circular buffer) to keep track of the current window elements.
- Maintain a running sum to quickly update the window's total when a new element comes in.
- When the window is full, subtract the value of the element that falls out as a new one is added.

- Each call to next() is executed in constant time, $O(1)$.

Space & Time

Time Complexity: $O(1)$ per call to next()

Space Complexity: $O(\text{window size})$

Solution

The solution leverages a queue data structure to store the elements of the sliding window. A running sum variable is maintained for efficient calculation of the moving average.

When a new integer is added:

- If the size of the queue is equal to the maximum window size, remove the oldest element and subtract its value from the running sum.
- Add the new element to both the queue and the running sum.
- Compute and return the moving average as the running sum divided by the number of elements in the queue. This approach ensures that each update is done in constant time without having to sum all elements in the window repeatedly.

Arrays & Hashing

#760 Find Anagram Mappings

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: Premium 98

Problem

Description

You are given two integer arrays `nums1` and `nums2` where `nums2` is an anagram of `nums1`. Both arrays may contain duplicates.

Return an index mapping array mapping from `nums1` to `nums2` where `mapping[i] = j` means the *i*th element in `nums1` appears in `nums2` at index *j*. If there are multiple answers, return any of them.

An array *a* is an anagram of an array *b* means *b* is made by randomizing the order of the elements in *a*.

Example 1:

Input: `nums1 = [12,28,46,32,50]`, `nums2 = [50,12,32,46,28]`

Output: `[1,4,3,2,0]`

Explanation: As `mapping[0] = 1` because the 0th element of `nums1` appears at `nums2[1]`, and `mapping[1] = 4` because the 1st element of `nums1` appears at `nums2[4]`, and so on.

Example 2:

Input: nums1 = [84,46], nums2 = [84,46]
 Output: [0,1]

Constraints:

- 1 <= nums1.length <= 100
- nums2.length == nums1.length
- 0 <= nums1[i], nums2[i] <= 105
- nums2 is an anagram of nums1.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def anagramMappings(self, nums1: list[int],
nums2: list[int]) -> list[int]:
        numToIndices = collections.defaultdict(list)

        for i, num in enumerate(nums2):
            numToIndices[num].append(i)

        return [numToIndices[num].pop() for num in
nums1]
```

• LeetDoocs

class Solution:

```
    def anagramMappings(self, nums1: List[int],
nums2: List[int]) -> List[int]:
        d = {x: i for i, x in enumerate(nums2)}
        return [d[x] for x in nums1]
```

• SimplyLeet

```
# Build a dictionary mapping each number in nums2
to a list of its indices.
def anagramMappings(nums1, nums2):
    # Dictionary to store number : list of indices
    index_dict = {}
    for index, num in enumerate(nums2):
        if num in index_dict:
            index_dict[num].append(index)
        else:
            index_dict[num] = [index]
```



```

mapping = []
# Generate the mapping for each element in
nums1.
for num in nums1:
    # Pop the first available index for num and
    add it to mapping.
    mapping.append(index_dict[num].pop(0))
return mapping

# Example usage:
nums1 = [12,28,46,32,50]
nums2 = [50,12,32,46,28]

print(anagramMappings(nums1, nums2)) # Output:
[1,4,3,2,0] or any valid mapping

```

• LC.ca

```

class Solution:
    def anagramMappings(self, nums1: List[int],
nums2: List[int]) -> List[int]:
        mapper = defaultdict(set)
        for i, num in enumerate(nums2):
            mapper[num].add(i)
        return [mapper[num].pop() for num in
nums1]

```

◆ Quick Review

Key Insights

- Since nums2 is an anagram of nums1, every element in nums1 has a corresponding position in nums2.
- A dictionary (or hash table) can be used to map each number in nums2 to its indices.
- When iterating through nums1, we can assign the next available index from the dictionary for that number.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the arrays (single pass to build the dictionary and another pass to form the mapping).

Space Complexity: $O(n)$ for storing the dictionary that holds indices of elements in nums2.

Solution

The approach consists of two primary steps:

- Build a dictionary to record all indices where each number appears in nums2. This dictionary maps a number to a list of indices.
- Iterate over nums1 and for each element, remove (or pop) one index from the corresponding list in the dictionary, and add that index to the final mapping array. This guarantees that each number from nums1 is correctly mapped to one of its positions in nums2 while efficiently handling duplicates.

Arrays & Hashing

#1426 Counting Elements

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: Premium 98

Problem

Description

Given an integer array `arr`, count how many elements `x` there are, such that `x + 1` is also in `arr`. If there are duplicates in `arr`, count them separately.

Example 1:

Input: `arr = [1,2,3]`
 Output: 2
 Explanation: 1 and 2 are counted cause 2 and 3 are in `arr`.

Example 2:

Input: `arr = [1,1,3,3,5,5,7,7]`
 Output: 0
 Explanation: No numbers are counted, cause there is no 2, 4, 6, or 8 in `arr`.

Constraints:

- `1 <= arr.length <= 1000`
- `0 <= arr[i] <= 1000`

Round 1 · High · Easy

p.43

Community Solutions (Python)

• WalkCC

```
class Solution:
    def countElements(self, arr: list[int]) -> int:
        count = collections.Counter(arr)
        return sum(freq
                    for a, freq in count.items()
                    if count[a + 1] > 0)
```

• LeetDoocs

```
class Solution:
    def countElements(self, arr: List[int]) ->
int:
        cnt = Counter(arr)
        return sum(v for x, v in cnt.items() if
cnt[x + 1])
```

• SimplyLeet

Round 1 · High · Easy

p.44

```
# Function to count elements where (x + 1) is in
the array
def count_elements(arr):
    # Create a set of unique numbers for O(1)
    lookups
    unique_numbers = set(arr)
    count = 0
    # Iterate through each number in the array
    for number in arr:
        # If (number + 1) exists in the unique set,
        increment count
        if number + 1 in unique_numbers:
            count += 1
```

```
return count
```

```
# Example usage:
# arr = [1, 2, 3]
# print(count_elements(arr)) # Expected output: 2
```

• LC.ca

```
class Solution:
    def countElements(self, arr: List[int]) ->
    int:
        cnt = Counter(arr)
        return sum(v for x, v in cnt.items() if
        cnt[x + 1])
```

◆ Quick Review

Round 1 · High · Easy

p.45

Key Insights

- Use a hash table or set to store unique elements from arr for constant time lookups.
- Iterate through each element in arr and check if $x + 1$ exists in the set.
- Count each valid occurrence, including duplicates.
- This approach efficiently handles the problem without needing nested loops.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in arr (set lookups are $O(1)$ on average).

Space Complexity: $O(n)$, due to the use of an auxiliary set to store unique elements.

Solution

The solution leverages a set to achieve constant time membership checks. First, we convert arr to a set of unique numbers. Then, iterate through arr, and for each element, check if (element + 1) exists in the set. If it does, increment the count. This approach correctly handles duplicates by counting each occurrence during iteration.

Round 1 · High · Easy

p.46

Arrays & Hashing

#1534 Count Good Triplets

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Enumeration

Lists: CodeSignal

Problem

Description

Given an array of integers `arr`, and three integers `a`, `b` and `c`. You need to find the number of good triplets.

A triplet $(arr[i], arr[j], arr[k])$ is good if the following conditions are true:

- $0 \leq i < j < k < arr.length$
- $|arr[i] - arr[j]| \leq a$
- $|arr[j] - arr[k]| \leq b$
- $|arr[i] - arr[k]| \leq c$

Where $|x|$ denotes the absolute value of x .

Return the number of good triplets.

Example 1:

Input: `arr = [3,0,1,1,9,7]`, `a = 7`, `b = 2`, `c = 3`

Output: 4

Explanation: There are 4 good triplets: $[(3,0,1), (3,0,1), (3,1,1), (0,1,1)]$.

Example 2:

Input: `arr = [1,1,2,2,3]`, `a = 0`, `b = 0`, `c = 1`

Output: 0

Explanation: No triplet satisfies all conditions.

Constraints:

- $3 \leq arr.length \leq 100$
- $0 \leq arr[i] \leq 1000$
- $0 \leq a, b, c \leq 1000$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def countGoodTriplets(self, arr: list[int], a: int, b: int, c: int) -> int:
        return sum(abs(arr[i] - arr[j]) <= a and
                    abs(arr[j] - arr[k]) <= b and
                    abs(arr[i] - arr[k]) <= c
                    for i, j, k in
                    itertools.combinations(range(len(arr)), 3))
```

- LeetCode

```
class Solution:
    def countGoodTriplets(self, arr: List[int], a:
int, b: int, c: int) -> int:
        ans, n = 0, len(arr)
        for i in range(n):
            for j in range(i + 1, n):
                for k in range(j + 1, n):
                    ans += (
                        abs(arr[i] - arr[j]) <= a
                        and abs(arr[j] - arr[k])
<= b
                        and abs(arr[i] - arr[k])
<= c
                    )
        return ans
```

- SimplyLeet

```
# Python solution for Count Good Triplets
def countGoodTriplets(arr, a, b, c):
    # Initialize counter for good triplets
    good_triplets_count = 0
    n = len(arr)
    # Iterate over all triplet combinations with
    indices i, j, k where i < j < k
    for i in range(n):
        for j in range(i + 1, n):
            for k in range(j + 1, n):
                # Check if absolute differences
                satisfy the conditions
                if abs(arr[i] - arr[j]) <= a and
abs(arr[j] - arr[k]) <= b and abs(arr[i] - arr[k])
<= c:
                    good_triplets_count += 1
    # Return the count of good triplets
    return good_triplets_count

# Example usage:
arr = [3, 0, 1, 1, 9, 7]
a, b, c = 7, 2, 3
print(countGoodTriplets(arr, a, b, c)) # Expected
output: 4
```

- LC.ca

```

class Solution:
    def countGoodTriplets(self, arr: List[int], a:
int, b: int, c: int) -> int:
        ans, n = 0, len(arr)
        for i in range(n):
            for j in range(i + 1, n):
                for k in range(j + 1, n):
                    ans += (
                        abs(arr[i] - arr[j]) <= a
                        and abs(arr[j] - arr[k])
<= b
                        and abs(arr[i] - arr[k])
<= c
                    )
        return ans

```

◆ Quick Review

Key Insights

- Brute force approach using three nested loops is acceptable since the maximum array length is 100.
- Check all possible triplets and count only those meeting the specified conditions.
- Absolute difference comparisons are the key condition validations.

Space & Time

Time Complexity: $O(n^3)$ where n is the length of the array.

Space Complexity: $O(1)$ extra space (ignoring input space).

Solution

We use a brute force algorithm with three nested loops, iterating over all combinations of triplets (i, j, k) ensuring $i < j < k$. For each triplet, we calculate the absolute differences:

– $|arr[i] - arr[j]| - |arr[j] - arr[k]| - |arr[i] - arr[k]|$ The triplet is counted as good if all absolute differences satisfy the conditions ($\leq a$, $\leq b$, and $\leq c$ respectively). This simple approach leverages the small input size, eliminating the need for more complex data structures or optimizations.

String

Round 1 · High · Easy · 3 questions

Round 1 · High · Easy

p.53

String

#383 Ransom Note

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Counting

Lists: Grind 169

Problem

Description

Given two strings ransomNote and magazine, return true if ransomNote can be constructed by using the letters from magazine and false otherwise.

Each letter in magazine can only be used once in ransomNote.

Example 1:

Input: ransomNote = "a", magazine = "b"
Output: false

Example 2:

Input: ransomNote = "aa", magazine = "ab"
Output: false

Example 3:

Input: ransomNote = "aa", magazine = "aab"
Output: true

Constraints:

Round 1 · High · Easy

p.54

- `1 <= ransomNote.length, magazine.length <= 105`
- `ransomNote` and `magazine` consist of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def canConstruct(self, ransomNote: str,
magazine: str) -> bool:
        count1 = collections.Counter(ransomNote)
        count2 = collections.Counter(magazine)
        return all(count1[c] <= count2[c] for c in
string.ascii_lowercase)
```

• LeetCode

```
class Solution:
    def canConstruct(self, ransomNote: str,
magazine: str) -> bool:
        cnt = Counter(magazine)
        for c in ransomNote:
            cnt[c] -= 1
            if cnt[c] < 0:
                return False
        return True
```

• SimplyLeet

```
# Function to check if ransom note can be
constructed from magazine
def canConstruct(ransomNote: str, magazine: str)
-> bool:
    # Create a dictionary to hold the frequency of
each character in magazine
    char_frequency = {}
    for char in magazine:
        char_frequency[char] =
char_frequency.get(char, 0) + 1

    # Check each character in ransomNote against
the frequency map
    for char in ransomNote:
        if char_frequency.get(char, 0) == 0:

            return False # Character not available
or used up
            char_frequency[char] -= 1 # Use the
character

    return True

# Example usage:
if __name__ == "__main__":
    print(canConstruct("aa", "aab")) # Expected
output: True
```

• LC.ca


```
class Solution:
    def canConstruct(self, ransomNote: str,
magazine: str) -> bool:
        cnt = Counter(magazine)
        for c in ransomNote:
            cnt[c] -= 1
            if cnt[c] < 0:
                return False
        return True
```

◆ Quick Review

Key Insights

- Use a hash table (or dictionary) to store the frequency of each character present in magazine.
- Iterate over every character in ransomNote and check if the corresponding count in the hash table is available and sufficient.
- Decrease the character count for every match; if any character count falls below zero, return false.
- If all characters in ransomNote can be accounted for, return true.

Space & Time

Time Complexity: $O(n + m)$, where n is the length of magazine and m is the length of ransomNote.

Space Complexity: $O(1)$, since the extra space required for the hash table is limited to at most 26 entries for lowercase English letters.

Solution

We construct a frequency map for the magazine string using a hash table. This frequency map records how many times each letter appears in magazine. Then we iterate through the ransomNote string, and for every letter, we check the corresponding frequency in the map. If a letter is not found or its frequency is zero, we immediately return false as it means the letter cannot be used to build the note. Otherwise, we decrement the frequency count by one. If we successfully process all characters in ransomNote, then it can be constructed from magazine, and we return true.

String

#734 Sentence Similarity

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String

Lists: Premium 98

Problem

Description

We can represent a sentence as an array of words, for example, the sentence "I am happy with leetcode" can be represented as `arr = ["I","am","happy","with","leetcode"]`.

Given two sentences `sentence1` and `sentence2` each represented as a string array and given an array of string pairs `similarPairs` where `similarPairs[i] = [xi, yi]` indicates that the two words `xi` and `yi` are similar.

Return `true` if `sentence1` and `sentence2` are similar, or `false` if they are not similar.

Two sentences are similar if:

- They have the same length (i.e., the same number of words)
- `sentence1[i]` and `sentence2[i]` are similar.

Notice that a word is always similar to itself, also notice that the similarity relation is not transitive. For example, if the words `a` and `b` are similar, and the words `b` and `c` are similar, `a` and `c` are not necessarily similar.

Round 1 · High · Easy

p.59

Example 1:

```
Input: sentence1 = ["great","acting","skills"],
       sentence2 = ["fine","drama","talent"], similarPairs = [
       ["great","fine"],["drama","acting"],["skills","talent"]
       ]
```

Output: `true`

Explanation: The two sentences have the same length and each word `i` of `sentence1` is also similar to the corresponding word in `sentence2`.

Example 2:

```
Input: sentence1 = ["great"], sentence2 = ["great"],
       similarPairs = []
```

Output: `true`

Explanation: A word is similar to itself.

Example 3:

```
Input: sentence1 = ["great"], sentence2 =
       ["doubleplus","good"], similarPairs =
       [["great","doubleplus"]]
```

Output: `false`

Explanation: As they don't have the same length, we return `false`.

Constraints:

- `1 <= sentence1.length, sentence2.length <= 1000`
- `1 <= sentence1[i].length, sentence2[i].length <= 20`
- `sentence1[i]` and `sentence2[i]` consist of English letters.
- `0 <= similarPairs.length <= 1000`
- `similarPairs[i].length == 2`

Round 1 · High · Easy

p.60

- $1 \leq x_i.length, y_i.length \leq 20$
- x_i and y_i consist of lower-case and upper-case English letters.
- All the pairs (x_i, y_i) are distinct.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def areSentencesSimilar(
        self,
        sentence1: list[str],
        sentence2: list[str],
        similarPairs: list[list[str]],
    ) -> bool:
        if len(sentence1) != len(sentence2):
            return False

        # map[key] := all the similar words of key
        map = collections.defaultdict(set)

        for a, b in similarPairs:
            map[a].add(b)
            map[b].add(a)

        for word1, word2 in zip(sentence1, sentence2):
            if word1 == word2:
                continue
```

```
        if word1 not in map:
            return False
        if word2 not in map[word1]:
            return False

        return True
```

• LeetCode

```
class Solution:
    def areSentencesSimilar(
        self, sentence1: List[str], sentence2:
        List[str], similarPairs: List[List[str]]
    ) -> bool:
        if len(sentence1) != len(sentence2):
            return False
        s = {(x, y) for x, y in similarPairs}
        for x, y in zip(sentence1, sentence2):
            if x != y and (x, y) not in s and (y,
            x) not in s:
                return False

        return True
```

• SimplyLeet

```

# Function to determine if two sentences are
similar
def are_sentences_similar(sentence1, sentence2,
similarPairs):
    # Check if sentences are of the same length
    if len(sentence1) != len(sentence2):
        return False

    # Build dictionary to store similar word pairs
    in both directions
    similar_dict = {}
    for word1, word2 in similarPairs:
        if word1 not in similar_dict:

            similar_dict[word1] = set()
            similar_dict[word1].add(word2)
            if word2 not in similar_dict:
                similar_dict[word2] = set()
                similar_dict[word2].add(word1)

    # Check each pair of words in the sentences
    for w1, w2 in zip(sentence1, sentence2):
        # If words are the same, continue
        if w1 == w2:

```

```

        continue
        # Check if w1 is similar to w2 explicitly
        if (w1 not in similar_dict or w2 not in
similar_dict[w1]) and (w2 not in similar_dict or w1
not in similar_dict[w2]):
            return False

    return True

# Example usage:
# sentence1 = ["great", "acting", "skills"]
# sentence2 = ["fine", "drama", "talent"]

# similarPairs = [["great", "fine"], ["drama",
"acting"], ["skills", "talent"]]
# print(are_sentences_similar(sentence1,
sentence2, similarPairs)) # Output: True

```

- LC.ca

```
class Solution:
    def areSentencesSimilar(
        self, sentence1: List[str], sentence2:
        List[str], similarPairs: List[List[str]]
    ) -> bool:
        if len(sentence1) != len(sentence2):
            return False
        s = {(a, b) for a, b in similarPairs}
        return all(
            a == b or (a, b) in s or (b, a) in s
            for a, b in zip(sentence1, sentence2)
        )
```

◆ Quick Review

Key Insights

- First, check if both sentences have the same length.
- A word is always similar to itself.
- Similarity is defined only by the given pairs; it is not transitive.
- Use a hash table (dictionary) to store bi-directional similar pairs for quick look-ups.

Space & Time

Time Complexity: $O(n + m)$, where n is the number of words in the sentences and m is the number of similar pairs.

Space Complexity: $O(m)$, for storing the similar pairs in a hash table.

Solution

– Check if the two sentences have the same length; if not, return false. – Build a hash table where each word maps to a set of words it is similar to. For each pair $[x, y]$ in `similarPairs`, add y to the set for x and x to the set for y . – Iterate over the words of both sentences simultaneously. For each corresponding pair, if the words are identical, continue; otherwise, check if one word appears in the similar set of the other. – If all pairs meet the similarity condition, return true; otherwise, return false.

String

#1165 Single Row Keyboard

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String

Lists: Premium 98

Problem

Description

There is a special keyboard with all keys in a single row. Given a string keyboard of length 26 indicating the layout of the keyboard (indexed from 0 to 25). Initially, your finger is at index 0. To type a character, you have to move your finger to the index of the desired character. The time taken to move your finger from index i to index j is $|i - j|$.

You want to type a string word. Write a function to calculate how much time it takes to type it with one finger.

Example 1:

Input: keyboard = "abcdefghijklmnopqrstuvwxyz", word = "cba"

Output: 4

Explanation: The index moves from 0 to 2 to write 'c' then to 1 to write 'b' then to 0 again to write 'a'. Total time = 2 + 1 + 1 = 4.

Example 2:

Input: keyboard = "pqrstuvwxyzabcdefghijklmnopqrstuvwxyz", word = "leetcode"
Output: 73

Constraints:

- keyboard.length == 26
- keyboard contains each English lowercase letter exactly once in some order.
- $1 \leq \text{word.length} \leq 104$
- word[i] is an English lowercase letter.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def calculateTime(self, keyboard: str, word: str) -> int:
        letterToIndex = {c: i for i, c in enumerate(keyboard)}
        return (letterToIndex[word[0]] +
                sum(abs(letterToIndex[a] -
                        letterToIndex[b])
                    for a, b in itertools.pairwise(word)))
```

• LeetDoocs

```
class Solution:
    def calculateTime(self, keyboard: str, word: str) -> int:
        pos = {c: i for i, c in enumerate(keyboard)}
        ans = i = 0
        for c in word:
            ans += abs(pos[c] - i)
            i = pos[c]
        return ans
```

• SimplyLeet

```
# Build a dictionary to map each character to its index in the keyboard string.
def calculateTime(keyboard: str, word: str) -> int:
    # Create a mapping for character to index
    char_index = {char: idx for idx, char in enumerate(keyboard)}

    # Initialize the total time and starting position
    total_time = 0
    current_position = 0

    # Iterate over each character in the word
```

```
        for char in word:
            # Find the target index for the character
            target_position = char_index[char]
            # Calculate the time to move to the target position
            total_time += abs(target_position - current_position)
            # Update current position to the target position
            current_position = target_position

        return total_time
```

```
# Example usage:
#
print(calculateTime("abcdefghijklmnopqrstuvwxyz", "cba"))
```

• LC.ca

```
class Solution:
    def calculateTime(self, keyboard: str, word: str) -> int:
        pos = {c: i for i, c in enumerate(keyboard)}
        ans = i = 0
        for c in word:
            ans += abs(pos[c] - i)
            i = pos[c]
        return ans
```

◆ Quick Review

Key Insights

- Create a mapping (hash table/dictionary) from each character to its index on the keyboard for $O(1)$ lookups.
- The cost to type each character is the absolute difference between the current position and the target character's position on the keyboard.
- Start from index 0 and sum up the distances moved for each successive character in the word.

Space & Time

Time Complexity: $O(n)$, where n is the length of the word.

Space Complexity: $O(1)$ (constant space for a mapping of 26 characters).

Solution

We solve the problem by first building a dictionary that maps each character in the keyboard layout to its corresponding index. We then initialize a variable to track the current finger position (starting at 0). For each character in the word, we look up its index, calculate the absolute difference from the current position, add that distance to the total time, and update the current position. This approach uses a single pass over the word with constant time lookups, ensuring an efficient solution.

Two Pointers

Round 1 · High · Easy · 5 questions

Round 1 · High · Easy

p.71

Round 1 · High · Easy

p.72

Two Pointers

#88 Merge Sorted Array

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Sorting

Lists: Denny Zhang

Problem

Description

You are given two integer arrays `nums1` and `nums2`, sorted in non-decreasing order, and two integers `m` and `n`, representing the number of elements in `nums1` and `nums2` respectively.

Merge `nums1` and `nums2` into a single array sorted in non-decreasing order.

The final sorted array should not be returned by the function, but instead be stored inside the array `nums1`. To accommodate this, `nums1` has a length of `m + n`, where the first `m` elements denote the elements that should be merged, and the last `n` elements are set to 0 and should be ignored. `nums2` has a length of `n`.

Example 1:

Input: `nums1 = [1,2,3,0,0,0]`, `m = 3`, `nums2 = [2,5,6]`, `n = 3`

Output: `[1,2,2,3,5,6]`

Explanation: The arrays we are merging are `[1,2,3]` and `[2,5,6]`.

The result of the merge is `[1,2,2,3,5,6]` with the underlined elements coming from `nums1`.

Example 2:

Input: `nums1 = [1]`, `m = 1`, `nums2 = []`, `n = 0`

Output: `[1]`

Explanation: The arrays we are merging are `[1]` and `[]`. The result of the merge is `[1]`.

Example 3:

Input: `nums1 = [0]`, `m = 0`, `nums2 = [1]`, `n = 1`

Output: `[1]`

Explanation: The arrays we are merging are `[]` and `[1]`. The result of the merge is `[1]`.

Note that because `m = 0`, there are no elements in `nums1`. The 0 is only there to ensure the merge result can fit in `nums1`.

Constraints:

- `nums1.length == m + n`
- `nums2.length == n`
- `0 <= m, n <= 200`
- `1 <= m + n <= 200`
- `-109 <= nums1[i], nums2[j] <= 109`

Follow up: Can you come up with an algorithm that runs in $O(m + n)$ time?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def merge(self, nums1: list[int], m: int, nums2:
list[int], n: int) -> None:
        i = m - 1 # nums1's index (the actual nums)
        j = n - 1 # nums2's index
        k = m + n - 1 # nums1's index (the next filled
position)

        while j >= 0:
            if i >= 0 and nums1[i] > nums2[j]:
                nums1[k] = nums1[i]
                k -= 1

            i -= 1
        else:
            nums1[k] = nums2[j]
            k -= 1
            j -= 1
```

• LeetCode

```
class Solution:
    def merge(self, nums1: List[int], m: int,
nums2: List[int], n: int) -> None:
        k = m + n - 1
        i, j = m - 1, n - 1
        while j >= 0:
            if i >= 0 and nums1[i] > nums2[j]:
                nums1[k] = nums1[i]
                i -= 1
            else:
                nums1[k] = nums2[j]

            j -= 1
            k -= 1
```

• SimplyLeet

```
def merge(nums1, m, nums2, n):
    # Initialize pointers for nums1, nums2 and the
tail index
    p1 = m - 1
    p2 = n - 1
    tail = m + n - 1

    # Merge arrays from the end by choosing the
larger element
    while p1 >= 0 and p2 >= 0:
        if nums1[p1] > nums2[p2]:
            nums1[tail] = nums1[p1]
```

```

        p1 -= 1 # Move pointer in nums1
    else:
        nums1[tail] = nums2[p2]
        p2 -= 1 # Move pointer in nums2
        tail -= 1 # Move tail pointer

    # Copy any remaining elements from nums2 into
    nums1
    while p2 >= 0:
        nums1[tail] = nums2[p2]
        p2 -= 1

    tail -= 1

```

• LC.ca

```

class Solution:
    def merge(self, nums1: List[int], m: int,
              nums2: List[int], n: int) -> None:
        """
        Do not return anything, modify nums1
        in-place instead.
        """
        i, j, k = m - 1, n - 1, m + n - 1
        # what if j=-1 already but k is not 0 yet?
        e.g. m=[1,2,0], n=[3]
        # => then no more ops needed, rest of m[]
        already sorted, so this while is good enough
        while j >= 0:
            # i could be -1, eg. [7,8,9, ] and [1]

```

```

        if i >= 0 and nums1[i] > nums2[j]:
            nums1[k] = nums1[i]
            i -= 1
        else:
            nums1[k] = nums2[j]
            j -= 1
        k -= 1

```

♦ Quick Review

Key Insights

- Both arrays are already sorted in non-decreasing order.
- The merge should be done in-place into nums1.
- Start merging from the end to avoid overwriting elements in nums1.
- Use two pointers, one for nums1 and one for nums2, and a tail pointer for placement.

Space & Time

Time Complexity: $O(m + n)$

Space Complexity: $O(1)$

Solution

We use a two-pointer approach starting from the end of both arrays. Initialize three pointers:

- p1 pointing to the last valid element in nums1.
- p2 pointing to the last element in nums2.
- tail pointing to the last index of nums1.

While both pointers are valid, compare the elements pointed by p1 and p2, placing the larger element at the tail position. Decrement the corresponding pointer and the tail pointer. Finally, if any elements remain in nums2 (i.e., p2 $\geq$ 0), copy them over into nums1. This solution leverages the extra space at the end of nums1 and ensures the merge is done in-place in linear time.

Two Pointers

#125 Valid Palindrome

EAS

[LeetDoocs](#) · [SimplyLeet](#) · [WalkCC](#) · [LeetCode](#)[Two Pointers](#) · [String](#)[Lists: Grind 169](#)

Problem

Description

A phrase is a palindrome if, after converting all uppercase letters into lowercase letters and removing all non-alphanumeric characters, it reads the same forward and backward. Alphanumeric characters include letters and numbers.

Given a string s, return true if it is a palindrome, or false otherwise.

Example 1:

Input: s = "A man, a plan, a canal: Panama"

Output: true

Explanation: "amanaplanacanalpanama" is a palindrome.

Example 2:

Input: s = "race a car"

Output: false

Explanation: "raceacar" is not a palindrome.

Example 3:

Input: s = " "

Output: true

Explanation: s is an empty string "" after removing non-alphanumeric characters.

Since an empty string reads the same forward and backward, it is a palindrome.

Constraints:

- $1 \leq \text{s.length} \leq 2 * 10^5$
- s consists only of printable ASCII characters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isPalindrome(self, s: str) -> bool:
        l = 0
        r = len(s) - 1

        while l < r:
            while l < r and not s[l].isalnum():
                l += 1
            while l < r and not s[r].isalnum():
                r -= 1
```

```
        if s[l].lower() != s[r].lower():
            return False
        l += 1
        r -= 1

    return True
```

• LeetCode

```
class Solution:
    def isPalindrome(self, s: str) -> bool:
        i, j = 0, len(s) - 1
        while i < j:
            if not s[i].isalnum():
                i += 1
            elif not s[j].isalnum():
                j -= 1
            elif s[i].lower() != s[j].lower():
                return False
        else:
            i, j = i + 1, j - 1
        return True
```

• SimplyLeet

```

class Solution:
    def isPalindrome(self, s: str) -> bool:
        # Initialize two pointers at the beginning
        # and end of the string
        left, right = 0, len(s) - 1
        while left < right:
            # Skip non-alphanumeric characters
            # from the left
            while left < right and not
s[left].isalnum():
                left += 1
            # Skip non-alphanumeric characters
            # from the right
            while left < right and not
s[right].isalnum():
                right -= 1
            # Compare the characters in lowercase
            if s[left].lower() !=
s[right].lower():
                return False
            left += 1
            right -= 1
        return True

```

• LC.ca

Round 1 · High · Easy

p.83

```

class Solution:
    def isPalindrome(self, s: str) -> bool:
        i, j = 0, len(s) - 1
        while i < j:
            if not s[i].isalnum():
                i += 1
            elif not s[j].isalnum():
                j -= 1
            elif s[i].lower() != s[j].lower():
                return False
        else:
            i, j = i + 1, j - 1
        return True

```

◆ Quick Review

Key Insights

- Use two pointers (one starting at the beginning and one at the end) to compare characters.
- Skip characters that are not alphanumeric.
- Compare characters in a case-insensitive manner.
- Continue until the pointers meet or cross.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, as each character is processed at most once.

Space Complexity: $O(1)$, using only a few extra variables for pointers and temporary comparisons.

Round 1 · High · Easy

p.84

Solution

The solution employs a two-pointer approach:

- Initialize two pointers: one at the start (left) and one at the end (right) of the string.
- Skip any non-alphanumeric characters using these pointers.
- Convert the remaining characters at each pointer to lowercase and compare them.
- If any mismatch is found, return false; otherwise, continue until the pointers cross.
- Return true if all corresponding characters match. This method efficiently checks for a palindrome without needing additional space for a filtered string.

Two Pointers

#283 Move Zeroes

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers

Lists: Grind 169

Problem

Description

Given an integer array `nums`, move all 0's to the end of it while maintaining the relative order of the non-zero elements.

Note that you must do this in-place without making a copy of the array.

Example 1:

Input: `nums = [0,1,0,3,12]`
Output: `[1,3,12,0,0]`

Example 2:

Input: `nums = [0]`
Output: `[0]`

Constraints:

- $1 \leq \text{nums.length} \leq 104$
- $-231 \leq \text{nums}[i] \leq 231 - 1$

Follow up: Could you minimize the total number of operations done?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def moveZeroes(self, nums: list[int]) -> None:
        j = 0
        for num in nums:
            if num != 0:
                nums[j] = num
                j += 1

        for i in range(j, len(nums)):
            nums[i] = 0
```

• LeetDoocs

```
class Solution:
    def moveZeroes(self, nums: List[int]) -> None:
        k = 0
        for i, x in enumerate(nums):
            if x:
                nums[k], nums[i] = nums[i],
nums[k]
                k += 1
```

• SimplyLeet

```
# Moves all zeros in the list 'nums' to the end
in-place.
def moveZeroes(nums):
    # 'j' is the pointer where the next non-zero
    element should be placed.
    j = 0
    # Loop through all the elements with index 'i'
    for i in range(len(nums)):
        # Check if the current element is non-zero.
        if nums[i] != 0:
            # Swap only if i is not equal to j to
            avoid unnecessary operations.
            nums[j], nums[i] = nums[i], nums[j]
```

```
        # Increment j to move to the next
        insertion index.
        j += 1
```

```
# Example usage:
nums = [0, 1, 0, 3, 12]
moveZeroes(nums)
print(nums) # Expected output: [1, 3, 12, 0, 0]
```

• LC.ca


```

class Solution:
    def moveZeroes(self, nums: List[int]) -> None:
        i = -1
        for j, x in enumerate(nums):
            if x:
                i += 1
                nums[i], nums[j] = nums[j],
nums[i]

```

◆ Quick Review

Key Insights

- Use a two-pointer technique: one pointer to iterate through the array and another pointer to track the position for the next non-zero element.
- For each non-zero element encountered, swap it with the element at the tracking pointer.
- This strategy minimizes operations by ensuring each non-zero element is moved at most once.
- After repositioning all non-zero elements, zeros will naturally occupy the remaining positions.
- The in-place requirement ensures space complexity remains $O(1)$.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in the array, since only one pass (or at most two simple passes) is required.

Space Complexity: $O(1)$ extra space, as the reordering is accomplished in-place.

Solution

The solution leverages a two-pointer approach. The left pointer (j) keeps track of the index where the next non-zero element should be placed, while the right pointer (i) iterates through the array. Whenever a non-zero element is encountered, it is swapped with the element at index j , and j is then incremented. This effectively gathers all non-zero elements at the beginning while moving zeros to the end. A key optimization is to perform the swap only when necessary (i.e., when i and j are not the same) to reduce the number of write operations.

Two Pointers

#408 Valid Word Abbreviation

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Two Pointers · String

Lists: Premium 98

Problem

Description

A string can be abbreviated by replacing any number of non-adjacent, non-empty substrings with their lengths. The lengths should not have leading zeros.

For example, a string such as "substitution" could be abbreviated as (but not limited to):

- "s10n" ("s ubstitutio n")
- "sub4u4" ("sub stit u tion")
- "12" ("substitution")
- "su3i1u2on" ("su bst i t u ti on")
- "substitution" (no substrings replaced)

The following are not valid abbreviations:

- "s55n" ("s ubsti tutio n", the replaced substrings are adjacent)
- "s010n" (has leading zeros)
- "s0ubstitution" (replaces an empty substring)

Given a string word and an abbreviation abbr, return whether the string matches the given abbreviation.

A substring is a contiguous non-empty sequence of characters within a string.

Example 1:

Input: word = "internationalization", abbr = "i12iz4n"
 Output: true
 Explanation: The word "internationalization" can be abbreviated as "i12iz4n" ("i nternational iz atio n").

Example 2:

Input: word = "apple", abbr = "a2e"
 Output: false
 Explanation: The word "apple" cannot be abbreviated as "a2e".

Constraints:

- $1 \leq \text{word.length} \leq 20$
- word consists of only lowercase English letters.
- $1 \leq \text{abbr.length} \leq 10$
- abbr consists of lowercase English letters and digits.
- All the integers in abbr will fit in a 32-bit integer.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def validWordAbbreviation(self, word: str, abbr:
str) -> bool:
        i = 0 # word's index
        j = 0 # abbr's index

        while i < len(word) and j < len(abbr):
            if word[i] == abbr[j]:
                i += 1
                j += 1
                continue

            if not abbr[j].isdigit() or abbr[j] == '0':
                return False
            num = 0
            while j < len(abbr) and abbr[j].isdigit():
                num = num * 10 + int(abbr[j])
                j += 1
            i += num

        return i == len(word) and j == len(abbr)

```

- LeetCode

```

class Solution:
    def validWordAbbreviation(self, word: str,
abbr: str) -> bool:
        m, n = len(word), len(abbr)
        i = j = x = 0
        while i < m and j < n:
            if abbr[j].isdigit():
                if abbr[j] == "0" and x == 0:
                    return False
                x = x * 10 + int(abbr[j])
            else:
                i += x
                x = 0
                if i >= m or word[i] != abbr[j]:
                    return False
                i += 1
                j += 1
        return i + x == m and j == n

```

- SimplyLeet

```

# Python solution for valid word abbreviation

def validWordAbbreviation(word: str, abbr: str) -> bool:
    word_index = 0 # pointer for the word
    abbr_index = 0 # pointer for the abbreviation

    while abbr_index < len(abbr):
        # if current abbrev char is a digit
        if abbr[abbr_index].isdigit():
            # leading zero check: if the digit is
            '0', it is invalid

            if abbr[abbr_index] == '0':
                return False
            # build the entire number
            num = 0
            while abbr_index < len(abbr) and
abbr[abbr_index].isdigit():
                num = num * 10 +
int(abbr[abbr_index])
                abbr_index += 1
            # skip num characters in word
            word_index += num
        else:

```

```

        # if letter doesn't match or word index
        is out of bounds, invalid
        if word_index >= len(word) or
word[word_index] != abbr[abbr_index]:
            return False
        word_index += 1
        abbr_index += 1

        # Valid if and only if the entire word was
        covered
        return word_index == len(word)

# Example usage:

#print(validWordAbbreviation("internationalization
", "i12iz4n")) # Expected output: True
#print(validWordAbbreviation("apple", "a2e")) #
Expected output: False

```

- LC.ca

```

class Solution:
    def validWordAbbreviation(self, word: str,
abbr: str) -> bool:
    m, n = len(word), len(abbr)
    i = j = x = 0
    while i < m and j < n:
        if abbr[j].isdigit():
            if abbr[j] == "0" and x == 0:
                return False
            x = x * 10 + int(abbr[j])
        else:
            i += x
            x = 0
            if i >= m or word[i] != abbr[j]:
                return False
            i += 1
            j += 1
    return i + x == m and j == n

```

◆ Quick Review

Key Insights

- Use two pointers: one for looping through the word and one for the abbreviation.
- When encountering a digit in abbr, accumulate the number and skip that many characters in the word.
- Check for invalid cases such as numbers with leading zeros or consecutive digit segments that might imply adjacent skipped substrings.

- Compare letters directly when they are not digits.

Space & Time

Time Complexity: $O(n + m)$ where n is the length of the word and m is the length of the abbreviation.

Space Complexity: $O(1)$ (constant extra space)

Solution

The approach uses two pointers to iterate over the word and abbreviation. For each character in the abbreviation:

- If it is a letter, compare it with the current character in the word.
- If it is a digit, first ensure that it does not start with '0'. Then, convert the sequence of digits into a number which represents how many characters in the word to skip.

After processing, both pointers should have reached the end of their respective strings for the abbreviation to be valid. This approach is efficient due to its linear pass with constant extra space.

Two Pointers

#977 Squares of a Sorted Array

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Sorting

Lists: Grind 169

Problem

Description

Given an integer array `nums` sorted in non-decreasing order, return an array of the squares of each number sorted in non-decreasing order.

Example 1:

Input: `nums = [-4,-1,0,3,10]`Output: `[0,1,9,16,100]`

Explanation: After squaring, the array becomes `[16,1,0,9,100]`.

After sorting, it becomes `[0,1,9,16,100]`.

Example 2:

Input: `nums = [-7,-3,2,3,11]`Output: `[4,9,9,49,121]`

Constraints:

- $1 \leq \text{nums.length} \leq 104$
- $-104 \leq \text{nums}[i] \leq 104$
- `nums` is sorted in non-decreasing order.

Follow up: Squaring each element and sorting the new array is very trivial, could you find an $O(n)$ solution using a different approach?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def sortedSquares(self, nums: list[int]) -> list[int]:
        n = len(nums)
        l = 0
        r = n - 1
        ans = [0] * n

        while n:
            n -= 1
            if abs(nums[l]) > abs(nums[r]):
                ans[n] = nums[l] * nums[l]
                l += 1
            else:
                ans[n] = nums[r] * nums[r]
                r -= 1

        return ans
```

• LeetDoocs

```

class Solution:
    def sortedSquares(self, nums: List[int]) -> List[int]:
        ans = []
        i, j = 0, len(nums) - 1
        while i <= j:
            a = nums[i] * nums[i]
            b = nums[j] * nums[j]
            if a > b:
                ans.append(a)
                i += 1
            else:
                ans.append(b)
                j -= 1
        return ans[::-1]

```

• SimplyLeet

```

# Function to return squares of a sorted array in
# non-decreasing order.
def sortedSquares(nums):
    # Initialize two pointers at the start and end
    # of the array.
    left, right = 0, len(nums) - 1
    # Create a result array of the same length as
    # nums.
    result = [0] * len(nums)
    # Index to insert the largest square at the
    # current position.
    position = len(nums) - 1

    # Loop until the left and right pointers cross.

```

```

while left <= right:
    # Compare absolute values at the left and
    # right pointers.
    if abs(nums[left]) > abs(nums[right]):
        result[position] = nums[left] *
        nums[left]
        left += 1 # Move left pointer
        rightward.
    else:
        result[position] = nums[right] *
        nums[right]
        right -= 1 # Move right pointer
        leftward.
    position -= 1 # Move to the next insertion
    position.
return result

```

```

# Example usage:
print(sortedSquares([-4, -1, 0, 3, 10])) # Output:
[0, 1, 9, 16, 100]

```

• LC.ca

```

class Solution:
    def sortedSquares(self, nums: List[int]) ->
    List[int]:
        n = len(nums)
        res = [0] * n
        i, j, k = 0, n - 1, n - 1
        while i <= j:
            if nums[i] * nums[i] > nums[j] *
            nums[j]:
                res[k] = nums[i] * nums[i]
                i += 1
            else:
                res[k] = nums[j] * nums[j]
                j -= 1
            k -= 1
        return res

```

◆ Quick Review

Key Insights

- Squaring numbers can disrupt the original order since negative numbers become positive.
- A two-pointer approach can efficiently retrieve the largest square from either end of the array.
- By comparing the absolute values at the two pointers, the largest square is inserted from the back of the result array.
- This method achieves an $O(n)$ time complexity without needing to sort the squared values afterward.

Space & Time

Time Complexity: $O(n)$ – We traverse the array only once.

Space Complexity: $O(n)$ – We use an extra array to store the output.

Solution

We use a two-pointer technique to traverse the array from both ends. Initialize pointers (left and right) at the beginning and end of the array, respectively, and an index pointer starting at the end of a new results array. At each step, compare the absolute values of the elements at the left and right pointers. The larger absolute value, when squared, is placed in the result array at the current index, and the corresponding pointer is adjusted. This continues until the left pointer exceeds the right pointer. Filling the result array from the back ensures that the array remains sorted in non-decreasing order.

Sorting

Round 1 · High · Easy · 6 questions

Round 1 · High · Easy

p.105

Round 1 · High · Easy

p.106

Sorting

#169 Majority Element

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Divide and Conquer · Sorting · Counting

Lists: Grind 169

Problem

Description

Given an array `nums` of size `n`, return the majority element.

The majority element is the element that appears more than $n / 2$ times. You may assume that the majority element always exists in the array.

Example 1:

Input: `nums = [3,2,3]`
Output: 3

Example 2:

Input: `nums = [2,2,1,1,1,2,2]`
Output: 2

Constraints:

- $n == \text{nums.length}$
- $1 \leq n \leq 5 * 10^4$
- $-10^9 \leq \text{nums}[i] \leq 10^9$

Round 1 · High · Easy

p.107

- The input is generated such that a majority element will exist in the array.

Follow-up: Could you solve the problem in linear time and in $O(1)$ space?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def majorityElement(self, nums: list[int]) -> int:
        ans = None
        count = 0

        for num in nums:
            if count == 0:
                ans = num
            count += (1 if num == ans else -1)

        return ans
```

• LeetDoocs

Round 1 · High · Easy

p.108

```
class Solution:
    def majorityElement(self, nums: List[int]) ->
    int:
        cnt = m = 0
        for x in nums:
            if cnt == 0:
                m, cnt = x, 1
            else:
                cnt += 1 if m == x else -1
        return m
```

• SimplyLeet

Python implementation of the Boyer-Moore Voting Algorithm

```
def majorityElement(nums):
    # Initialize candidate and counter
    candidate = None
    count = 0

    # Iterate over each number in the input list
    for num in nums:
        # If count is 0, choose the current number
        # as the candidate
```

```
        if count == 0:
            candidate = num
        # Increment count if the current number
        # equals the candidate, otherwise decrement
        count += 1 if num == candidate else -1
        # Return the candidate, which is the majority
        # element
        return candidate

# Example usage
print(majorityElement([2,2,1,1,1,2,2])) # Output:
2
```

• LC.ca

```
class Solution:
    def majorityElement(self, nums: List[int]) ->
    int:
        cnt = m = 0
        for x in nums:
            if cnt == 0:
                m, cnt = x, 1
            else:
                cnt += 1 if m == x else -1
        return m
```

◆ Quick Review

Key Insights

- The majority element appears more than half the length of the array.
- A simple hash table approach can count occurrences, but a more efficient solution exists.
- The Boyer–Moore Voting Algorithm finds the majority in $O(n)$ time and $O(1)$ space.
- The algorithm maintains a candidate and a counter, adjusting the candidate when the count drops to zero.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We use the Boyer–Moore Voting Algorithm to identify the majority element. The algorithm iterates through the array using two variables: one for the candidate and one for the count. Initially, we set the candidate to an arbitrary value and count to 0. For each element, if the count is 0, we update the candidate to the current element. Then, if the current element is the same as the candidate, we increment the count; otherwise, we decrement it. Since the majority element appears more than $n/2$ times, it will remain as the candidate by the end of the iteration.

Sorting

#217 Contains Duplicate

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Sorting

Lists: Grind 169

Problem

Description

Given an integer array `nums`, return `true` if any value appears at least twice in the array, and return `false` if every element is distinct.

Example 1:

Input: `nums = [1,2,3,1]`

Output: `true`

Explanation:

The element 1 occurs at the indices 0 and 3.

Example 2:

Input: `nums = [1,2,3,4]`

Output: `false`

Explanation:

All elements are distinct.

Example 3:

Input: `nums = [1,1,1,3,3,4,3,2,4,2]`

Output: true

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $-109 \leq \text{nums}[i] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def containsDuplicate(self, nums: list[int]) ->
    bool:
        return len(nums) != len(set(nums))
```

• LeetDoocs

```
class Solution:
    def containsDuplicate(self, nums: List[int])
    -> bool:
        return any(a == b for a, b in
        pairwise(sorted(nums)))
```

```
class Solution:
    def containsDuplicate(self, nums: List[int])
    -> bool:
        return len(set(nums)) < len(nums)
```

• SimplyLeet

Function to determine if there are any duplicates in the list.

```
def contains_duplicate(nums):
    # Create a set to keep track of numbers seen so far.
```

```
    seen = set()
```

```
    # Iterate over each number in the list.
```

```
    for num in nums:
```

```
        # If num is already in seen, duplicate found.
```

```
        if num in seen:
```

```
            return True
```

```
        # Add num to the set.
```

```
        seen.add(num)
```

```
    # No duplicates found.
```

```
    return False
```

Example usage:

```
if __name__ == "__main__":
```

```
    nums = [1, 2, 3, 1]
```

```
    print(contains_duplicate(nums)) # Expected
    output: True
```

• LC.ca

```

class Solution:
    def containsDuplicate(self, nums: List[int])
-> bool:
        return len(set(nums)) < len(nums)

#####

class Solution(object):
    def containsDuplicate(self, nums):
        """
        :type nums: List[int]

        :rtype: bool
        """
        nums.sort()
        for i in range(0, len(nums) - 1):
            if nums[i] == nums[i + 1]:
                return True
        return False

```

◆ Quick Review

Key Insights

- A hash set can be used to keep track of seen numbers.
- Iterating once through the array provides an efficient solution.
- If a number is already in the set, a duplicate is found.
- Alternative approaches like sorting have a higher time complexity ($O(n \log n)$).

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution uses a hash set to track the numbers seen while iterating through the array. For each number in the array:

- Check if it exists in the hash set. - If it does, return true immediately as a duplicate is found. - Otherwise, add the number to the hash set. If the iteration completes without finding any duplicates, return false. This approach efficiently determines if any duplicates exist in a single pass.

Sorting

#242 Valid Anagram

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sorting

Lists: Grind 169

Problem

Description

Given two strings *s* and *t*, return true if *t* is an anagram of *s*, and false otherwise.

Example 1:

Input: *s* = "anagram", *t* = "nagaram"

Output: true

Example 2:

Input: *s* = "rat", *t* = "car"

Output: false

Constraints:

- $1 \leq s.length, t.length \leq 5 * 10^4$
- *s* and *t* consist of lowercase English letters.

Follow up: What if the inputs contain Unicode characters? How would you adapt your solution to such a case?

Community Solutions (Python)

Round 1 · High · Easy

p.117

• WalkCC

```
class Solution:
    def isAnagram(self, s: str, t: str) -> bool:
        if len(s) != len(t):
            return False

        count = collections.Counter(s)
        count.subtract(collections.Counter(t))
        return all(freq == 0 for freq in
count.values())
```

• LeetDoocs

```
class Solution:
    def isAnagram(self, s: str, t: str) -> bool:
        if len(s) != len(t):
            return False
        cnt = Counter(s)
        for c in t:
            cnt[c] -= 1
            if cnt[c] < 0:
                return False
        return True
```

```
class Solution:
    def isAnagram(self, s: str, t: str) -> bool:
        return Counter(s) == Counter(t)
```

• SimplyLeet

Round 1 · High · Easy

p.118

Python solution for Valid Anagram

```
def isAnagram(s: str, t: str) -> bool:
    # Check if lengths differ; if so, they can't be
    # anagrams
    if len(s) != len(t):
        return False

    # Create a dictionary to count frequency of
    # each character in s
    char_count = {}
    for char in s:
        # Increase the count of the character
        char_count[char] = char_count.get(char, 0)
        + 1

    # Decrease the count for each character in t
    for char in t:
        if char not in char_count or
        char_count[char] == 0:
            # Character not found or count
            # mismatch, not an anagram
            return False
        char_count[char] -= 1
```

If all counts are zero, it's a valid anagram
return True

Example usage:
print(isAnagram("anagram", "nagaram")) # Output:
True
print(isAnagram("rat", "car")) # Output: False

• LC.ca

```
class Solution:
    def isAnagram(self, s: str, t: str) -> bool:
        if len(s) != len(t):
            return False
        cnt = Counter(s)
        for c in t:
            cnt[c] -= 1
            if cnt[c] < 0:
                return False
        return True
```

◆ Quick Review

Key Insights

- If s and t have different lengths, they cannot be anagrams.
- An anagram requires exactly the same character frequency for both strings.
- Using a hash table (or frequency array when limited to 26 lowercase letters) provides an efficient solution.

- An alternative approach is to sort both strings and compare them.
- For Unicode inputs, a hash table is more adaptable than a fixed-size frequency array.

Space & Time

Time Complexity: $O(n)$, where n is the length of the strings (using a hash table for counting).

Space Complexity: $O(1)$ if we use a fixed-size array (for 26 letters), otherwise $O(n)$ for a hash table with a larger character set.

Solution

We use a hash table (dictionary) to count the frequency of each character in the first string s . Then, we iterate over string t and decrement the count for each character. If at any point a character count goes negative or a character is missing in the hash table, t cannot be an anagram of s . Finally, if all counts return to zero, the strings are anagrams. This approach is efficient and handles the problem in linear time.

Sorting

#252 Meeting Rooms

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Sorting

Lists: Grind 169 | Premium 98

Problem

Description

Given an array of meeting time intervals where $\text{intervals}[i] = [\text{start}_i, \text{end}_i]$, determine if a person could attend all meetings.

Example 1:

Input: $\text{intervals} = [[0,30],[5,10],[15,20]]$
Output: false

Example 2:

Input: $\text{intervals} = [[7,10],[2,4]]$
Output: true

Constraints:

- $0 \leq \text{intervals.length} \leq 104$
- $\text{intervals}[i].\text{length} == 2$
- $0 \leq \text{start}_i < \text{end}_i \leq 106$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def canAttendMeetings(self, intervals:
list[list[int]]) -> bool:
        intervals.sort()

        for i in range(1, len(intervals)):
            if intervals[i - 1][1] > intervals[i][0]:
                return False

        return True
```

- LeetCode

```
class Solution:
    def canAttendMeetings(self, intervals:
List[List[int]]) -> bool:
        intervals.sort()
        return all(a[1] <= b[0] for a, b in
pairwise(intervals))
```

- SimplyLeet

Python solution for checking if one can attend all meetings

```
def canAttendMeetings(intervals):
    # Sort intervals based on start time
    intervals.sort(key=lambda x: x[0])
    # Iterate through intervals to find any overlap
    for i in range(1, len(intervals)):
        # Check if the previous meeting ends after
        the current meeting starts
        if intervals[i-1][1] > intervals[i][0]:
            return False
```

```
    return True
```

```
# Example usage:
print(canAttendMeetings([[0, 30], [5, 10], [15,
20]])) # Expected output: False
print(canAttendMeetings([[7, 10], [2, 4]])) #
Expected output: True
```

- LC.ca

```
class Solution:
    def canAttendMeetings(self, intervals:
List[List[int]]) -> bool:
        intervals.sort()
        return all(a[1] <= b[0] for a, b in
pairwise(intervals))
```

```
#####
```

```
from itertools import pairwise
class Solution:
    def canAttendMeetings(self, intervals:
List[List[int]]) -> bool:
```

```
        intervals.sort(key=lambda x: x[0])
        return not any(a[1] > b[0] for a,b in
pairwise(intervals))
        # any is True, then conflict found
```

```
#####
```

```
class Solution:
    def canAttendMeetings(self, intervals:
List[List[int]]) -> bool:
        intervals.sort(key=lambda x: x[0])
        for i in range(len(intervals) - 1):
```

```
            if intervals[i][1] > intervals[i +
1][0]:
                return False
            return True
```

```
#####
```

```
# Definition for an interval.
# class Interval(object):
# def __init__(self, s=0, e=0):
# self.start = s
```

```
# self.end = e
```

```
class Solution(object):
    def canAttendMeetings(self, intervals):
        """
        :type intervals: List[Interval]
        :rtype: bool
        """
        intervals = sorted(intervals, key=lambda x:
x.start)
        for i in range(1, len(intervals)):
```

```
            if intervals[i].start < intervals[i -
1].end:
                return False
            return True
```

◆ Quick Review

Key Insights

- Sort the intervals based on their starting times.
- Compare each interval's end time with the next interval's start time.
- Detect overlap if the end time of one meeting exceeds the start time of the following meeting.
- If no overlaps are detected, all meetings can be attended.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(1)$ if sorting is done in-place; otherwise $O(n)$.

Solution

The solution starts by sorting the intervals in ascending order by their start times. After sorting, iterate through the intervals and for each consecutive pair, check if the previous meeting's end time exceeds the next meeting's start time. An overlap indicates that the person cannot attend all meetings. The primary data structure used is an array (or list) for holding the intervals. Sorting makes it easier to compare adjacent intervals for potential overlaps.

Sorting

#1086 Largest Unique Number

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Sorting

Lists: Premium 98

Problem

Description

Given an integer array `nums`, return the largest integer that only occurs once. If no integer occurs once, return `-1`.

Example 1:

Input: `nums = [5,7,3,9,4,9,8,3,1]`

Output: 8

Explanation: The maximum integer in the array is 9 but it is repeated. The number 8 occurs only once, so it is the answer.

Example 2:

Input: `nums = [9,9,8,8]`

Output: -1

Explanation: There is no number that occurs only once.

Constraints:

- $1 \leq \text{nums.length} \leq 2000$
- $0 \leq \text{nums}[i] \leq 1000$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def largestUniqueNumber(self, nums: list[int])
    -> int:
        return max([num for num, freq in
collections.Counter(nums).items()
                    if freq == 1], default=-1)
```

• LeetDoocs

```
class Solution:
    def largestUniqueNumber(self, nums: List[int])
    -> int:
        cnt = Counter(nums)
        return max((x for x, v in cnt.items() if v
== 1), default=-1)
```

• SimplyLeet

```
# Define a function to find the largest unique
number
def largestUniqueNumber(nums):
    # Create a dictionary to count the frequency of
each number
    frequency = {}
    for num in nums:
        if num in frequency:
            frequency[num] += 1 # Increment
frequency count if number exists
        else:
            frequency[num] = 1 # Initialize
frequency count for new number
```

```
# Initialize the variable for the maximum
unique number
max_unique = -1
# Iterate over the frequency dictionary
for num, count in frequency.items():
    # Check if the number is unique and greater
than the current maximum unique
    if count == 1 and num > max_unique:
        max_unique = num

return max_unique
```

```
# Example usage:
print(largestUniqueNumber([5,7,3,9,4,9,8,3,1])) #
Output: 8
```

• LC.ca

```
class Solution:
    def largestUniqueNumber(self, nums: List[int])
-> int:
    cnt = Counter(nums)
    return next((x for x in range(1000, -1, -1)
if cnt[x] == 1), -1)
```

◆ Quick Review

Key Insights

- Use a hash table (or dictionary) to count the frequency of each number.
- Iterate over the hash table to find the numbers that appear exactly once.
- Identify the maximum number among the unique numbers.
- Return -1 if no unique number exists.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in `nums`.

Space Complexity: $O(n)$, for the hash table storing frequency counts.

Solution

The approach starts by initializing a hash table to keep track of the frequency of each element in the array. Once the frequency count is complete, we traverse the hash table to collect numbers that appear exactly once. Finally, we determine the maximum number among these unique

elements. If there is no unique element, we return -1. This method is efficient because the frequency count and the search for the maximum unique element both run in linear time relative to the input size.

Sorting

#1122 Relative Sort Array

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Sorting · Counting Sort

Lists: Denny Zhang

Problem

Description

Given two arrays `arr1` and `arr2`, the elements of `arr2` are distinct, and all elements in `arr2` are also in `arr1`.

Sort the elements of `arr1` such that the relative ordering of items in `arr1` are the same as in `arr2`. Elements that do not appear in `arr2` should be placed at the end of `arr1` in ascending order.

Example 1:

Input: `arr1 = [2,3,1,3,2,4,6,7,9,2,19]`, `arr2 = [2,1,4,3,9,6]`
 Output: `[2,2,2,1,4,3,3,9,6,7,19]`

Example 2:

Input: `arr1 = [28,6,22,8,44,17]`, `arr2 = [22,28,8,6]`
 Output: `[22,28,8,6,17,44]`

Constraints:

- $1 \leq \text{arr1.length}, \text{arr2.length} \leq 1000$
- $0 \leq \text{arr1}[i], \text{arr2}[i] \leq 1000$

- All the elements of `arr2` are distinct.
- Each `arr2[i]` is in `arr1`.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def relativeSortArray(self, arr1: list[int],
                        arr2: list[int]) -> list[int]:
        ans = []
        count = [0] * 1001

        for a in arr1:
            count[a] += 1

        for a in arr2:
            while count[a] > 0:
                ans.append(a)
                count[a] -= 1

        for num in range(1001):
            for _ in range(count[num]):
                ans.append(num)

        return ans
```

• LeetDoocs

```
class Solution:
    def relativeSortArray(self, arr1: List[int],
arr2: List[int]) -> List[int]:
        pos = {x: i for i, x in enumerate(arr2)}
        return sorted(arr1, key=lambda x:
pos.get(x, 1000 + x))
```

• SimplyLeet

Python solution using a frequency dictionary

```
def relativeSortArray(arr1, arr2):
    # Count frequency of each element in arr1
    count = {}
    for num in arr1:
        count[num] = count.get(num, 0) + 1

    result = []
    # Process elements defined in arr2: append each
    element count times
```

```
        for num in arr2:
            if num in count:
                result.extend([num] * count[num])
                # Remove the element from the count
                dictionary
                del count[num]
```

Process remaining numbers: sort them in ascending order and append

```
remaining = []
for num, freq in count.items():
    remaining.extend([num] * freq)
```

```
result.extend(sorted(remaining))
```

```
return result
```

Example usage:

```
print(relativeSortArray([2,3,1,3,2,4,6,7,9,2,19],
[2,1,4,3,9,6]))
```

• LC.ca

```
class Solution:
    def relativeSortArray(self, arr1: List[int],
arr2: List[int]) -> List[int]:
        pos = {x: i for i, x in enumerate(arr2)}
        return sorted(arr1, key=lambda x:
pos.get(x, 1000 + x))
```


◆ Quick Review

Key Insights

- Use a hash table (or array for counting given the value constraints) to count the occurrences of each element in arr1.
- Iterate through arr2 and for each element, append it to the result based on its count from the hash table.
- Remove the used elements from the hash table.
- Sort the remaining elements (those not in arr2) in ascending order and append them to the result.

Space & Time

Time Complexity: $O(n + m \log m)$ where n is the length of arr1 and m is the number of leftover elements not in arr2. Given that element values are bounded (0 to 1000), a counting sort can make this nearly $O(n)$.

Space Complexity: $O(n)$ for storing the frequency counts and the result array.

Solution

The solution leverages a frequency counter (hash table) to count how many times each element appears in arr1. We then process arr2 to add each number in its given order according to its frequency. Finally, we sort the remaining numbers that are not in arr2 and append them to the result array. This approach ensures that the relative order specified by arr2 is maintained and that the other elements are sorted in ascending order.

Round 1 · High · Easy

p.137

Binary Search

Round 1 · High · Easy · 5 questions

Round 1 · High · Easy

p.138

Binary Search

#35 Search Insert Position

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Denny Zhang

Problem

Description

Given a sorted array of distinct integers and a target value, return the index if the target is found. If not, return the index where it would be if it were inserted in order.

You must write an algorithm with $O(\log n)$ runtime complexity.

Example 1:

Input: nums = [1,3,5,6], target = 5
Output: 2

Example 2:

Input: nums = [1,3,5,6], target = 2
Output: 1

Example 3:

Input: nums = [1,3,5,6], target = 7
Output: 4

Round 1 · High · Easy

p.139

Constraints:

- $1 \leq \text{nums.length} \leq 104$
- $-104 \leq \text{nums}[i] \leq 104$
- nums contains distinct values sorted in ascending order.
- $-104 \leq \text{target} \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def searchInsert(self, nums: list[int], target:
int) -> int:
        l = 0
        r = len(nums)

        while l < r:
            m = (l + r) // 2
            if nums[m] == target:
                return m
            if nums[m] < target:

                l = m + 1
            else:
                r = m

        return l
```

• LeetDoocs

Round 1 · High · Easy

p.140

```
class Solution:
    def searchInsert(self, nums: List[int],
target: int) -> int:
    l, r = 0, len(nums)
    while l < r:
        mid = (l + r) >> 1
        if nums[mid] >= target:
            r = mid
        else:
            l = mid + 1
    return l
```

```
class Solution:
    def searchInsert(self, nums: List[int],
target: int) -> int:
    return bisect_left(nums, target)
```

- SimplyLeet

```
# Python solution using binary search
def searchInsert(nums, target):
    # Initialize low and high pointers for binary
search
    low, high = 0, len(nums) - 1
    # Perform binary search until low pointer
exceeds high pointer
    while low <= high:
        mid = (low + high) // 2 # Calculate the
middle index
        if nums[mid] == target:
            return mid # Return mid if target is
found
        elif target < nums[mid]:
```

```
            high = mid - 1 # Narrow search on the
left side
        else:
            low = mid + 1 # Narrow search on the
right side
    # Return low pointer as the insertion index
when target is not found
    return low
```

```
# Example usage
print(searchInsert([1, 3, 5, 6], 5)) # Output: 2
```

- LC.ca

```

class Solution:
    def searchInsert(self, nums: List[int],
target: int) -> int:
    left, right = 0, len(nums)
    while left < right:
        mid = (left + right) >> 1
        if nums[mid] >= target:
            right = mid
        else:
            left = mid + 1
    return left

```

◆ Quick Review

Key Insights

- The array is sorted, enabling the use of binary search.
- Binary search achieves a time complexity of $O(\log n)$ which meets the problem constraints.
- When the target is not found during the search, the final low pointer represents the correct insertion index.

Space & Time

Time Complexity: $O(\log n)$ due to binary search.

Space Complexity: $O(1)$ as no extra space is used.

Solution

We use binary search to find the target in the sorted array. Initialize two pointers, low and high. At each iteration, compute the mid index. If the value at mid matches the target, return mid. If the target is less than the mid value,

adjust the high pointer; if greater, adjust the low pointer. When the iterations finish without finding the target, return the low pointer, which represents the correct insertion index.

Binary Search

#69 Sqrt(x)

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Binary Search

Lists: Denny Zhang

Problem

Description

Given a non-negative integer x , return the square root of x rounded down to the nearest integer. The returned integer should be non-negative as well.

You must not use any built-in exponent function or operator.

- For example, do not use `pow(x, 0.5)` in c++ or `x ** 0.5` in python.

Example 1:

Input: $x = 4$
 Output: 2
 Explanation: The square root of 4 is 2, so we return 2.

Example 2:

Input: $x = 8$
 Output: 2
 Explanation: The square root of 8 is 2.82842..., and since we round it down to the nearest integer, 2 is returned.

Constraints:

- $0 \leq x \leq 231 - 1$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def mySqrt(self, x: int) -> int:
        return bisect.bisect_right(range(x + 1), x,
                                     key=lambda m: m *
                                     m) - 1
```

• LeetDoocs

```
class Solution:
    def mySqrt(self, x: int) -> int:
        l, r = 0, x
        while l < r:
            mid = (l + r + 1) >> 1
            if mid > x // mid:
                r = mid - 1
            else:
                l = mid
        return l
```

• SimplyLeet

```

# Binary search solution to find the integer square
root
def mySqrt(x):
    # Edge case for small numbers
    if x < 2:
        return x

    low, high = 0, x
    ans = 0
    # Binary search loop
    while low <= high:

        mid = (low + high) // 2
        # If mid squared equals x, we found the
exact square root
        if mid * mid == x:
            return mid
        # If mid squared is less than x, store mid
and search right half
        elif mid * mid < x:
            ans = mid # candidate answer
            low = mid + 1
        else:
            # If mid squared is more than x, search
left half

```

```

        high = mid - 1
    return ans

```

```

# Example usage:
print(mySqrt(4)) # Output: 2
print(mySqrt(8)) # Output: 2

```

• LC.ca

```

class Solution:
    def mySqrt(self, x: int) -> int:
        l, r = 0, x
        while l < r:
            mid = (l + r + 1) >> 1
            if mid * mid > x:
                r = mid - 1
            else:
                l = mid
        return l

```

◆ Quick Review

Key Insights

- The problem can be solved efficiently using binary search.
- Instead of iterating from 0 to x, binary search narrows down the candidate square roots quickly.
- The algorithm should handle edge cases like $x = 0$ and $x = 1$.

- By checking $mid * mid$ against x , we can determine if we need to search in the lower or upper half.

Space & Time

Time Complexity: $O(\log x)$

Space Complexity: $O(1)$

Solution

We use a binary search approach to find the integer square root. Initialize two pointers, low and high, where low is set to 0 and high is set to x . For each iteration, calculate the mid value. If mid^2 equals x , we have found the square root. If mid^2 is less than x , record mid as a potential answer and move the low pointer to $mid + 1$. Otherwise, move the high pointer to $mid - 1$. Continue the search until low exceeds high, then return the last recorded candidate.

Binary Search

#278 First Bad Version

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Binary Search · Interactive

Lists: Grind 169

Problem

Description

You are a product manager and currently leading a team to develop a new product. Unfortunately, the latest version of your product fails the quality check. Since each version is developed based on the previous version, all the versions after a bad version are also bad. Suppose you have n versions $[1, 2, \dots, n]$ and you want to find out the first bad one, which causes all the following ones to be bad.

You are given an API `bool isBadVersion(version)` which returns whether version is bad. Implement a function to find the first bad version. You should minimize the number of calls to the API.

Example 1:

Input: n = 5, bad = 4
 Output: 4
 Explanation:
 call isBadVersion(3) -> false
 call isBadVersion(5) -> true
 call isBadVersion(4) -> true
 Then 4 is the first bad version.

Example 2:

Input: n = 1, bad = 1
 Output: 1

Constraints:

- $1 \leq \text{bad} \leq n \leq 2^{31} - 1$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def firstBadVersion(self, n: int) -> int:
        l = 1
        r = n

        while l < r:
            m = (l + r) >> 1
            if isBadVersion(m):
                r = m
            else:
```

```
l = m + 1
```

```
return l
```

• LeetCode

The isBadVersion API is already defined for you.
 # def isBadVersion(version: int) -> bool:

```
class Solution:
    def firstBadVersion(self, n: int) -> int:
        l, r = 1, n
        while l < r:
            mid = (l + r) >> 1
            if isBadVersion(mid):
```

```
            r = mid
        else:
            l = mid + 1
        return l
```

• SimplyLeet


```
# Python implementation of First Bad Version
class Solution:
    def firstBadVersion(self, n: int) -> int:
        # Initialize the left pointer to the first
        # version and right pointer to the last version
        left, right = 1, n
        while left < right:
            # Calculate the mid point to avoid
            # potential overflow
            mid = left + (right - left) // 2
            # Call the API isBadVersion to check if
            # mid version is bad
            if isBadVersion(mid):
                # If mid is bad, search the left
                # half including mid
                right = mid
            else:
                # If mid is not bad, search the
                # right half excluding mid
                left = mid + 1
        # When left equals right, we found the
        # first bad version
        return left
```

• LC.ca

Round 1 · High · Easy

p.153

```
# The isBadVersion API is already defined for you.
# @param version, an integer
# @return an integer
# def isBadVersion(version):

class Solution:
    def firstBadVersion(self, n):
        """
        :type n: int

        :rtype: int
        """
        left, right = 1, n
        while left < right:
            mid = (left + right) >> 1
            if isBadVersion(mid):
                right = mid
            else:
                left = mid + 1
        return left
```

◆ Quick Review

Key Insights

- The API returns true for bad versions and all versions after a bad version.
- The problem is essentially about finding a transition point in a sorted sequence which calls for a binary search approach.

Round 1 · High · Easy

p.154

- The aim is to minimize the number of calls to `isBadVersion` by leveraging the ordered nature of the versions.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We use binary search to efficiently locate the first bad version. Initialize two pointers, left and right, at 1 and n respectively. At each step, compute the mid point and call `isBadVersion(mid)`. If mid is bad, it indicates that the first bad version may be mid or to its left, so adjust right to mid. Otherwise, adjust left to mid + 1. Continue the process until the pointers converge. The converged pointer will be the first bad version. Key structures include simple integer variables for left, right, and mid, with binary search being the fundamental algorithm to reduce the search space logarithmically.

Binary Search

#374 Guess Number Higher or Lower

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Binary Search · Interactive

Lists: Denny Zhang

Problem

Description

We are playing the Guess Game. The game is as follows: I pick a number from 1 to n. You have to guess which number I picked (the number I picked stays the same throughout the game).

Every time you guess wrong, I will tell you whether the number I picked is higher or lower than your guess.

You call a pre-defined API `int guess(int num)`, which returns three possible results:

- -1: Your guess is higher than the number I picked (i.e. `num > pick`).
- 1: Your guess is lower than the number I picked (i.e. `num < pick`).
- 0: your guess is equal to the number I picked (i.e. `num == pick`).

Return the number that I picked.

Example 1:

Input: n = 10, pick = 6
Output: 6

Example 2:

Input: n = 1, pick = 1
Output: 1

Example 3:

Input: n = 2, pick = 1
Output: 1

Constraints:

- $1 \leq n \leq 231 - 1$
- $1 \leq \text{pick} \leq n$

Community Solutions (Python)

• WalkCC

```
# The guess API is already defined for you.
# @param num, your guess
# @return -1 if num is higher than the picked
#         number
#         1 if num is lower than the picked number
#         0 otherwise return 0
# def guess(num: int) -> int:

class Solution:
    def guessNumber(self, n: int) -> int:
        l = 1
```

```
r = n

# Find the first guess number that >= the
# target number
while l < r:
    m = (l + r) // 2
    if guess(m) <= 0: # -1, 0
        r = m
    else:
        l = m + 1

return l
```

• LeetCode

```
# The guess API is already defined for you.
# @param num, your guess
# @return -1 if num is higher than the picked
#         number
#         1 if num is lower than the picked number
#         0 otherwise return 0
# def guess(num: int) -> int:

class Solution:
    def guessNumber(self, n: int) -> int:

        return bisect.bisect(range(1, n + 1), 0,
                             key=lambda x: -guess(x))
```

• SimplyLeet

```
# Python solution using binary search
class Solution:
    def guessNumber(self, n: int) -> int:
        # Initialize low and high pointers
        low, high = 1, n
        # Continue binary search until low exceeds
        high
        while low <= high:
            # Calculate midpoint safely to prevent
            overflow
            mid = low + (high - low) // 2
            # Call the guess API with mid

            result = guess(mid)
            # If the guess is correct, return mid
            if result == 0:
                return mid
            # If guess was too high, adjust the
            high pointer
            elif result == -1:
                high = mid - 1
            # If guess was too low, adjust the low
            pointer
            else:
                low = mid + 1

        # Return low (ideally should never reach
        here if pick exists in [1, n])
        return low
```

• LC.ca

Round 1 · High · Easy

p.159

```
# The guess API is already defined for you.
# @param num, your guess
# @return -1 if num is higher than the picked
number
# 1 if num is lower than the picked number
# otherwise return 0
# def guess(num: int) -> int:

class Solution:
    def guessNumber(self, n: int) -> int:

        return bisect.bisect(range(1, n + 1), 0,
            key=lambda x: -guess(x))
```

◆ Quick Review

Key Insights

- The problem can be effectively solved using binary search since the search space is sorted by nature (1 to n).
- Each API call helps to halve the search space, leading to an efficient solution.
- Be cautious with potential overflow when calculating the mid-point by using $low + (high - low) // 2$.

Space & Time

Time Complexity: $O(\log n)$ since we binary search the range.

Space Complexity: $O(1)$ as we use a constant amount of extra space.

Round 1 · High · Easy

p.160

Solution

We use a binary search algorithm to efficiently locate the target number. We initialize two pointers: low (1) and high (n). In each iteration, we compute the mid-point, then call the guess API. Based on the response:

- If guess(mid) returns 0, we have found the number.
- If it returns -1, it indicates that our guess is higher than the pick, so we update the high pointer to mid - 1.
- If it returns 1, it indicates that our guess is lower than the pick, so we update the low pointer to mid + 1.

This approach guarantees a logarithmic time solution.

Binary Search

#704 Binary Search

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Grind 169

Problem

Description

Given an array of integers `nums` which is sorted in ascending order, and an integer `target`, write a function to search `target` in `nums`. If `target` exists, then return its index. Otherwise, return -1.

You must write an algorithm with $O(\log n)$ runtime complexity.

Example 1:

Input: `nums = [-1,0,3,5,9,12]`, `target = 9`

Output: 4

Explanation: 9 exists in `nums` and its index is 4

Example 2:

Input: `nums = [-1,0,3,5,9,12]`, `target = 2`

Output: -1

Explanation: 2 does not exist in `nums` so return -1

Constraints:

- $1 \leq \text{nums.length} \leq 104$

- $-104 < \text{nums}[i], \text{target} < 104$
- All the integers in nums are unique.
- nums is sorted in ascending order.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def search(self, nums: list[int], target: int) -> int:
        i = bisect.bisect_left(nums, target)
        return -1 if i == len(nums) or nums[i] != target else i
```

• LeetCode

```
class Solution:
    def search(self, nums: List[int], target: int) -> int:
        l, r = 0, len(nums) - 1
        while l < r:
            mid = (l + r) >> 1
            if nums[mid] >= target:
                r = mid
            else:
                l = mid + 1
        return l if nums[l] == target else -1
```

• SimplyLeet

Python code for binary search

```
def search(nums, target):
    # Set initial boundaries for the search
    left, right = 0, len(nums) - 1

    # Loop until the search space is empty
    while left <= right:
        # Calculate middle index to avoid overflow
        mid = left + (right - left) // 2

        # Check if the mid element matches the target
        if nums[mid] == target:
            return mid
        # If target is smaller than mid element, narrow search to left half
        elif target < nums[mid]:
            right = mid - 1
        # Else, narrow search to right half
        else:
            left = mid + 1
```

```
# Target not found, return -1
return -1
```

```
# Example usage:
print(search([-1,0,3,5,9,12], 9)) # Expected
output: 4
```

• LC.ca

```
class Solution:
    def search(self, nums: List[int], target: int)
-> int:
    left, right = 0, len(nums) - 1
    while left < right:
        mid = (left + right) >> 1
        if nums[mid] >= target:
            right = mid
        else:
            left = mid + 1
    return left if nums[left] == target else -1
```

◆ Quick Review

Key Insights

- Utilize binary search to achieve $O(\log n)$ runtime on a sorted array.
- Use two pointers to maintain the current search region.
- Narrow down the search range by comparing the middle element with the target.

- Return the index when the target is found; otherwise, return -1 if not found.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We use an iterative binary search algorithm. We initialize two pointers, left and right, to represent the current segment of the array that could contain the target. At each iteration, we compute the middle index and compare the element there with the target. If they match, we return the index. Otherwise, based on whether the target is smaller or larger than the middle element, we adjust our pointers to narrow the search to either the left or right half of the array. This process continues until the target is found or until the pointers cross, indicating that the target is not present. No extra data structures are used, making the space complexity $O(1)$.

Matrix

Round 1 · High · Easy · 5 questions

Round 1 · High · Easy

p.167

Matrix

#422 Valid Word Square

EAS

[LeetDoocs](#) · [SimplyLeet](#) · [WalkCC](#) · [LeetCode](#)

[Array](#) · [Matrix](#)

Lists: Premium 98

Problem

Description

Given an array of strings `words`, return true if it forms a valid word square.

A sequence of strings forms a valid word square if the k th row and column read the same string, where $0 \leq k < \max(\text{numRows}, \text{numColumns})$.

Example 1:

Round 1 · High · Easy

p.168

a	b	c	d
b	n	r	t
c	r	m	y
d	t	y	e

Input: words = ["abcd","bnrt","crmy","dtye"]
Output: true
Explanation:
The 1st row and 1st column both read "abcd".
The 2nd row and 2nd column both read "bnrt".
The 3rd row and 3rd column both read "crmy".
The 4th row and 4th column both read "dtye".
Therefore, it is a valid word square.

Example 2:

a	b	c	d
b	n	r	t
c	r	m	
d	t		

Input: words = ["abcd","bnrt","crm","dt"]
Output: true
Explanation:
The 1st row and 1st column both read "abcd".
The 2nd row and 2nd column both read "bnrt".
The 3rd row and 3rd column both read "crm".
The 4th row and 4th column both read "dt".
Therefore, it is a valid word square.

Example 3:

b	a	l	l
a	r	e	a
r	e	a	d
l	a	d	y

Input: words = ["ball","area","read","lady"]

Output: false

Explanation:

The 3rd row reads "read" while the 3rd column reads "lead".

Therefore, it is NOT a valid word square.

Constraints:

- $1 \leq \text{words.length} \leq 500$
- $1 \leq \text{words}[i].\text{length} \leq 500$
- $\text{words}[i]$ consists of only lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def validWordSquare(self, words: list[str]) -> bool:
        for i, word in enumerate(words):
            for j, c in enumerate(word):
                if len(words) <= j or len(words[j]) <= i: # out-of-bounds
                    return False
                if c != words[j][i]:
                    return False
        return True
```

• LeetDoocs

```
class Solution:
    def validWordSquare(self, words: List[str]) -> bool:
        m = len(words)
        for i, w in enumerate(words):
            for j, c in enumerate(w):
                if j >= m or i >= len(words[j]) or c != words[j][i]:
                    return False
        return True
```

• SimplyLeet

```
# Function to check if the provided words form a
valid word square
def validWordSquare(words):
    # Iterate over each row with its index
    for i in range(len(words)):
        # Iterate over each character in the
        current row
        for j in range(len(words[i])):
            # Check if the corresponding column
            exists
            if j >= len(words) or i >=
            len(words[j]) or words[i][j] != words[j][i]:
                return False
    return True
```

```
# Example test case
words = ["abcd","bnrt","crmy","dtye"]
print(validWordSquare(words)) # Expected output:
True
```

- LC.ca

```
class Solution:
    def validWordSquare(self, words: List[str]) ->
    bool:
        m = len(words)
        for i, w in enumerate(words):
            for j, c in enumerate(w):
                if j >= m or i >= len(words[j]) or
                c != words[j][i]:
                    return False
        return True
```

◆ Quick Review

Key Insights

- The k-th row should be equal to the k-th column.
- Not every word is the same length; checks must be performed within bounds.
- Iterating over each character and cross-verifying with the corresponding column character is sufficient.

Space & Time

Time Complexity: $O(N * L)$ where N is the number of words and L is the average length of the words.

Space Complexity: $O(1)$ additional space is used for the checking procedure.

Solution

We iterate through each word (row) and each character (column) in the word. For every character at position (i, j), we confirm that there is a corresponding word at index j

and that its length is greater than i . Then we check that the character at position (j, i) in that word is the same as the current character. If any of these checks fail, we return false. If all checks pass, the words form a valid word square.

Round 1 · High · Easy

p.175

Matrix

#566 Reshape the Matrix

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix · Simulation

Lists: CodeSignal

Problem

Description

In MATLAB, there is a handy function called reshape which can reshape an $m \times n$ matrix into a new one with a different size $r \times c$ keeping its original data.

You are given an $m \times n$ matrix `mat` and two integers `r` and `c` representing the number of rows and the number of columns of the wanted reshaped matrix.


The reshaped matrix should be filled with all the elements of the original matrix in the same row-traversing order as they were.

If the reshape operation with given parameters is possible and legal, output the new reshaped matrix; Otherwise, output the original matrix.

Example 1:


Round 1 · High · Easy

p.176



Input: mat = [[1,2],[3,4]], r = 1, c = 4
 Output: [[1,2,3,4]]

Example 2:



Input: mat = [[1,2],[3,4]], r = 2, c = 4
 Output: [[1,2],[3,4]]

Constraints:

- m == mat.length
- n == mat[i].length
- 1 <= m, n <= 100
- -1000 <= mat[i][j] <= 1000
- 1 <= r, c <= 300

Community Solutions (Python)

• WalkCC

```
class Solution:
    def matrixReshape(self, nums: list[list[int]],
                      r: int, c: int) ->
        list[list[int]]:
        if nums == [] or r * c != len(nums) *
            len(nums[0]):
            return nums

        ans = [[0 for j in range(c)] for i in range(r)]
        k = 0

        for row in nums:

            for num in row:
                ans[k // c][k % c] = num
                k += 1

        return ans
```

• LeetDoocs

```
class Solution:
    def matrixReshape(self, mat: List[List[int]],
r: int, c: int) -> List[List[int]]:
    m, n = len(mat), len(mat[0])
    if m * n != r * c:
        return mat
    ans = [[0] * c for _ in range(r)]
    for i in range(m * n):
        ans[i // c][i % c] = mat[i // n][i % n]
    return ans
```

• SimplyLeet

```
# Function to reshape the matrix
def matrixReshape(mat, r, c):
    # Calculate the dimensions of the original
matrix
    m = len(mat)
    n = len(mat[0])

    # Check if reshape is possible
    if m * n != r * c:
        return mat
```

```
# Flatten the original matrix into a list
flat_list = []
for row in mat:
    for num in row:
        flat_list.append(num)

# Create the new matrix by slicing the flat
list
new_matrix = []
for i in range(0, len(flat_list), c):
    new_matrix.append(flat_list[i:i+c])

return new_matrix

# Example call:
# print(matrixReshape([[1,2],[3,4]], 1, 4))
```

• LC.ca

```
class Solution:
    def matrixReshape(self, mat: List[List[int]],
r: int, c: int) -> List[List[int]]:
    m, n = len(mat), len(mat[0])
    if m * n != r * c:
        return mat
    ans = [[0] * c for _ in range(r)]
    for i in range(m * n):
        ans[i // c][i % c] = mat[i // n][i % n]
    return ans
```

◆ Quick Review

Key Insights

- Check if the reshape operation is possible by comparing the total elements ($m * n$ with $r * c$).
- Flatten the original matrix while preserving the row-traversing order.
- Use the flattened list to build the new matrix row by row.
- If the total number of elements does not match, simply return the original matrix.

Space & Time

Time Complexity: $O(mn)$ as we iterate through all elements of the matrix once.

Space Complexity: $O(mn)$ for the additional list used to store flattened elements (or the new matrix).

Solution

The solution involves first checking whether the reshape is possible by ensuring that the product of the rows and columns of the original matrix equals that of the desired matrix ($r * c$). If not, the original matrix is returned. Otherwise, the matrix is flattened into a one-dimensional list, preserving the row-order traversal. Finally, the flattened list is segmented by the new column size c to form the reshaped matrix. The primary data structures used include a list (or array) for flattening the matrix and then another list (or vector/array) to hold the final reshaped matrix.

Matrix

#766 Toeplitz Matrix

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Matrix

Lists: CodeSignal


Problem

Description

Given an $m \times n$ matrix, return true if the matrix is Toeplitz. Otherwise, return false.

A matrix is Toeplitz if every diagonal from top-left to bottom-right has the same elements.

Example 1:



Input: matrix = [[1,2,3,4],[5,1,2,3],[9,5,1,2]]

Output: true


Explanation:

In the above grid, the diagonals are:

"[9]", "[5, 5]", "[1, 1, 1]", "[2, 2, 2]", "[3, 3]", "[4]".

In each diagonal all elements are the same, so the answer is True.

Example 2:



Input: matrix = [[1,2],[2,2]]

Output: false

Explanation:

The diagonal "[1, 2]" has different elements.

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[i].\text{length}$
- $1 \leq m, n \leq 20$

Round 1 · High · Easy

p.183

- $0 \leq \text{matrix}[i][j] \leq 99$

Follow up:

- What if the matrix is stored on disk, and the memory is limited such that you can only load at most one row of the matrix into the memory at once?
- What if the matrix is so large that you can only load up a partial row into the memory at once?

Community Solutions (Python)

- WalkCC

```

class Solution:
    def isToeplitzMatrix(self, matrix:
list[list[int]]) -> bool:
        for i in range(len(matrix) - 1):
            for j in range(len(matrix[0]) - 1):
                if matrix[i][j] != matrix[i + 1][j + 1]:
                    return False

        return True
  
```

- LeetCode

Round 1 · High · Easy

p.184


```

class Solution:
    def isToeplitzMatrix(self, matrix:
List[List[int]]) -> bool:
        m, n = len(matrix), len(matrix[0])
        for i in range(1, m):
            for j in range(1, n):
                if matrix[i][j] != matrix[i - 1][j
- 1]:
                    return False
        return True

```

• SimplyLeet

```

# Function to determine if the matrix is Toeplitz
def isToeplitzMatrix(matrix):
    # Get the number of rows and columns in the
    matrix
    rows = len(matrix)
    cols = len(matrix[0])

    # Iterate starting from row 1 and column 1
    because first row and column are trivially Toeplitz
    for i in range(1, rows):
        for j in range(1, cols):
            # Compare current element with its
            top-left neighbor

```

```

                if matrix[i][j] != matrix[i - 1][j -
1]:
                    return False
            return True

# Example usage:
matrix1 = [[1,2,3,4],[5,1,2,3],[9,5,1,2]]
print(isToeplitzMatrix(matrix1)) # Expected
output: True

```

• LC.ca

```

class Solution:
    def isToeplitzMatrix(self, matrix:
List[List[int]]) -> bool:
        m, n = len(matrix), len(matrix[0])
        return all(
            matrix[i][j] == matrix[i - 1][j - 1]
            for i in range(1, m)
            for j in range(1, n)
        )

```

◆ Quick Review

Key Insights

- Each element (except those in the first row and first column) must be equal to its top-left neighbor.
- You can validate the matrix by iterating from the second row/column and comparing each element with `matrix[i-1][j-1]`.

- This simple check allows you to determine if every diagonal has the same elements without extra space.
- For streaming data (loading one row at a time), maintain the previous row for comparison.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(1)$

Solution

The problem is solved by iterating through the matrix starting from the element at position (1, 1). For every element at `matrix[i][j]`, compare it with the element at `matrix[i-1][j-1]`. If any two elements in this pair do not match, the matrix is not Toeplitz. This check ensures that every diagonal from the top-left to bottom-right has identical elements. In terms of auxiliary space, only constant extra space is used, which makes this approach efficient both in time and space.

Matrix

#867 Transpose Matrix

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix · Simulation


Lists: CodeSignal

Problem

Description

Given a 2D integer array `matrix`, return the transpose of `matrix`.

The transpose of a matrix is the matrix flipped over its main diagonal, switching the matrix's row and column indices.



Example 1:

Input: `matrix = [[1,2,3],[4,5,6],[7,8,9]]`
 Output: `[[1,4,7],[2,5,8],[3,6,9]]`

Example 2:

Input: matrix = [[1,2,3],[4,5,6]]

Output: [[1,4],[2,5],[3,6]]

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[i].\text{length}$
- $1 \leq m, n \leq 1000$
- $1 \leq m * n \leq 105$
- $-109 \leq \text{matrix}[i][j] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def transpose(self, A: list[list[int]]) -> list[list[int]]:
        ans = [[0] * len(A) for _ in range(len(A[0]))]

        for i in range(len(A)):
            for j in range(len(A[0])):
                ans[j][i] = A[i][j]

        return ans
```

• LeetCode

```
class Solution:
    def transpose(self, matrix: List[List[int]])
-> List[List[int]]:
        return list(zip(*matrix))
```

• SimplyLeet

Python solution with clear comments

```
def transpose(matrix):
    # Number of rows in the original matrix
    num_rows = len(matrix)
    # Number of columns in the original matrix
    num_cols = len(matrix[0])

    # Create a new matrix with dimensions swapped:
    num_cols x num_rows
    transposed = [[0] * num_rows for _ in
range(num_cols)]
```

```

    # Iterate through each element in the original
    matrix
    for i in range(num_rows):
        for j in range(num_cols):
            # Place the element at the swapped
            index in the transposed matrix
            transposed[j][i] = matrix[i][j]

    return transposed

```

Example usage:

```

matrix = [[1, 2, 3], [4, 5, 6]]
print(transpose(matrix))

```

• LC.ca

```

class Solution:
    def transpose(self, matrix: List[List[int]])
    -> List[List[int]]:
        return list(zip(*matrix))

```

◆ Quick Review

Key Insights

- The value at position (i, j) in the original matrix moves to (j, i) in the transposed matrix.
- The transposed matrix dimensions are the reverse of the original: if the original is m x n, then the transposed is n x m.

- A simple double loop is sufficient to swap the indices, ensuring $O(m \cdot n)$ time complexity.
- This problem handles both square and rectangular matrices.

Space & Time

Time Complexity: $O(m \cdot n)$

Space Complexity: $O(m \cdot n)$ for the new matrix holding the transposed values

Solution

The solution involves iterating over each element of the original matrix and assigning it to the corresponding position in the new matrix where the indexes are swapped. We create the transposed matrix with dimensions based on the number of columns and rows of the original matrix, respectively. The approach leverages basic array manipulation with nested loops to fill in the new matrix.

Matrix

#2373 Largest Local Values in a Matrix

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix

Lists: CodeSignal

Problem

Description

You are given an $n \times n$ integer matrix grid.

Generate an integer matrix maxLocal of size $(n - 2) \times (n - 2)$ such that:

- $\text{maxLocal}[i][j]$ is equal to the largest value of the 3×3 matrix in grid centered around row $i + 1$ and column $j + 1$.

In other words, we want to find the largest value in every contiguous 3×3 matrix in grid.

Return the generated matrix.

Example 1:

9	9	8	1
5	6	2	6
8	2	6	4
6	2	2	2

9	9
8	6

Input: grid = [[9,9,8,1],[5,6,2,6],[8,2,6,4],[6,2,2,2]]

Output: [[9,9],[8,6]]

Explanation: The diagram above shows the original matrix and the generated matrix.

Notice that each value in the generated matrix corresponds to the largest value of a contiguous 3×3 matrix in grid.

Example 2:

1	1	1	1	1
1	1	1	1	1
1	1	2	1	1
1	1	1	1	1
1	1	1	1	1

2	2	2
2	2	2
2	2	2

Input: grid = [[1,1,1,1,1],[1,1,1,1,1],[1,1,2,1,1],[1,1,1,1,1],[1,1,1,1,1]]
 Output: [[2,2,2],[2,2,2],[2,2,2]]
 Explanation: Notice that the 2 is contained within every contiguous 3 x 3 matrix in grid.

Constraints:

- $n == \text{grid.length} == \text{grid}[i].\text{length}$
- $3 \leq n \leq 100$
- $1 \leq \text{grid}[i][j] \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def largestLocal(self, grid: list[list[int]]) -> list[list[int]]:
        n = len(grid)
        ans = [[0] * (n - 2) for _ in range(n - 2)]

        for i in range(n - 2):
            for j in range(n - 2):
                for x in range(i, i + 3):
                    for y in range(j, j + 3):
                        ans[i][j] = max(ans[i][j], grid[x][y])

        return ans
```

- LeetCode

```
class Solution:
    def largestLocal(self, grid: List[List[int]]) -> List[List[int]]:
        n = len(grid)
        ans = [[0] * (n - 2) for _ in range(n - 2)]
        for i in range(n - 2):
            for j in range(n - 2):
                ans[i][j] = max(
                    grid[x][y] for x in range(i, i + 3) for y in range(j, j + 3)
                )
        return ans
```

- SimplyLeet

Python solution for finding the largest local values in a matrix

```
def largestLocal(grid):
    n = len(grid) # dimension of the grid
    # Initialize the result matrix with zeros, of
    size (n-2) x (n-2)
    result = [[0] * (n - 2) for _ in range(n - 2)]

    # Loop through each possible top-left index of
    a 3x3 submatrix
    for i in range(n - 2):
        for j in range(n - 2):

            # Compute the maximum in the current
            3x3 submatrix
            local_max = 0 # since grid elements are
            >= 1, starting with 0 is safe
            for k in range(3):
                for l in range(3):
                    local_max = max(local_max,
                    grid[i + k][j + l])
            # Store the local maximum in the result
            matrix
            result[i][j] = local_max
    return result

# Example usage:
```

```
grid = [[9,9,8,1],[5,6,2,6],[8,2,6,4],[6,2,2,2]]
print(largestLocal(grid))
```

• LC.ca

```
class Solution:
    def largestLocal(self, grid: List[List[int]])
    -> List[List[int]]:
        n = len(grid)
        ans = [[0] * (n - 2) for _ in range(n - 2)]
        for i in range(n - 2):
            for j in range(n - 2):
                ans[i][j] = max(
                    grid[x][y] for x in range(i, i
                    + 3) for y in range(j, j + 3)
                )
        return ans
```

◆ Quick Review

Key Insights

- The size of the output matrix is $(n - 2) \times (n - 2)$.
- For each valid index (i, j) in the output, examine the 3×3 block in grid starting at $\text{grid}[i][j]$ to $\text{grid}[i+2][j+2]$.
- Since each 3×3 block contains exactly 9 elements, finding the maximum in a block takes constant time.
- Use nested loops to slide over the grid and efficiently compute each local maximum.

Space & Time

Time Complexity: $O(n^2)$ – We iterate over $(n-2) \times (n-2)$ positions, and for each position, we look at 9 elements.

Space Complexity: $O(n^2)$ – The additional space is required for the output matrix.

Solution

The solution involves iterating over each possible 3×3 submatrix in grid. For each (i, j) starting index of the submatrix:

- Examine the nine elements from $grid[i][j]$ to $grid[i+2][j+2]$.
- Compute the maximum value from these elements.
- Store the maximum in the corresponding position of the result matrix.

Data structures used include the result matrix for storing the maximum values. The algorithmic approach involves two nested loops (one for rows and one for columns) to traverse each possible starting position of a 3×3 window. The simplicity of the solution comes from the fixed 3×3 block size, which enables constant-time maximum computation for each block.

Trees & BST

Round 1 · High · Easy · 11 questions

Trees & BST

#100 Same Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS

Lists: Grind 169


Problem

Description

Given the roots of two binary trees p and q , write a function to check if they are the same or not.

Two binary trees are considered the same if they are structurally identical, and the nodes have the same value.

Example 1:




Input: $p = [1,2,3]$, $q = [1,2,3]$
Output: true

Example 2:


Round 1 · High · Easy

p.201



Input: $p = [1,2]$, $q = [1,null,2]$
Output: false

Example 3:



Input: $p = [1,2,1]$, $q = [1,1,2]$
Output: false

Constraints:

- The number of nodes in both trees is in the range $[0, 100]$.
- $-104 \leq \text{Node.val} \leq 104$

Round 1 · High · Easy

p.202

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isSameTree(self, p: TreeNode | None, q:
TreeNode | None) -> bool:
        if not p or not q:
            return p == q
        return (p.val == q.val and
                self.isSameTree(p.left, q.left) and
                self.isSameTree(p.right, q.right))
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def isSameTree(self, p: Optional[TreeNode], q:
Optional[TreeNode]) -> bool:
        if p == q:
            return True
```

```
        if p is None or q is None or p.val !=
q.val:
            return False
        return self.isSameTree(p.left, q.left) and
self.isSameTree(p.right, q.right)
```

• SimplyLeet

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

def isSameTree(p, q):
    # If both nodes are None, trees are identical
    at this branch.
    if not p and not q:
```

```

        return True
    # If one of the nodes is None, trees are not
    identical.
    if not p or not q:
        return False
    # Check if current nodes have the same value.
    if p.val != q.val:
        return False
    # Recursively compare the left subtrees and the
    right subtrees.
    return isSameTree(p.left, q.left) and
    isSameTree(p.right, q.right)

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x

```

```

# self.left = None
# self.right = None

class Solution(object):
    def isSameTree(self, p, q):
        """
        :type p: TreeNode
        :type q: TreeNode
        :rtype: bool
        """

```

```

        if not p or not q:
            return p == q # covers: p=None and q!=None,
            q=None and p!=None, both None
            return p.val == q.val and
            self.isSameTree(p.left, q.left) and
            self.isSameTree(p.right, q.right)

```

```

# iteration
class Solution:
    def isSameTree(self, p: TreeNode, q: TreeNode)
    -> bool:
        if p is None:
            return q is None

```

```

if q is None:
    return p is None

```

```

stack1 = [p]
stack2 = [q]

```

```

while stack1 and stack2:
    current1 = stack1.pop()
    current2 = stack2.pop()

```

```

None:
    if current1 is None and current2 is
        continue
    elif current1 is None or current2 is
None:
        return False
    elif current1.val != current2.val:
        return False

    stack1.append(current1.left)
    stack2.append(current2.left)

```

```

    stack1.append(current1.right)
    stack2.append(current2.right)

return not stack1 and not stack2

```

◆ Quick Review

Key Insights

- Use recursion to compare nodes in both trees.
- If both nodes being compared are null, they are the same.
- If one node is null while the other is not, the trees differ.
- Check if current nodes' values match, then recursively verify the left and right subtrees.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the trees, since every node is visited once.

Space Complexity: $O(n)$ in the worst-case (unbalanced tree), due to recursion call stack usage. In a balanced tree, the space would be $O(\log n)$.

Solution

The problem can be solved using a recursive approach (Depth-First Search) where we compare the current nodes from both trees. If the current nodes are both null, they match; if one is null, then the trees are not identical. For non-null nodes, check if the values are equal. Then, recursively apply the same logic on the left and right children of the nodes. This approach ensures that both the structure and the node values are identical across the two trees.

Trees & BST

#101 Symmetric Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS


Lists: Grind 169

Problem

Description

Given the root of a binary tree, check whether it is a mirror of itself (i.e., symmetric around its center).


Example 1:



Input: root = [1,2,2,3,4,4,3]

Output: true

Example 2:



Input: root = [1,2,2,null,3,null,3]
Output: false

Constraints:

- The number of nodes in the tree is in the range [1, 1000].
- $-100 \leq \text{Node.val} \leq 100$

Follow up: Could you solve it both recursively and iteratively?

Community Solutions (Python)

• WalkCC

```

class Solution:
    def isSymmetric(self, root: TreeNode | None) -> bool:
        def isSymmetric(p: TreeNode | None, q: TreeNode | None) -> bool:
            if not p or not q:
                return p == q
            return (p.val == q.val and
                    isSymmetric(p.left, q.right) and
                    isSymmetric(p.right, q.left))

        return isSymmetric(root, root)
  
```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isSymmetric(self, root: Optional[TreeNode]) -> bool:
        def dfs(root1: Optional[TreeNode], root2: Optional[TreeNode]) -> bool:
            if root1 == root2:
  
```

```

        return True
    if root1 is None or root2 is None or
    root1.val != root2.val:
        return False
    return dfs(root1.left, root2.right)
    and dfs(root1.right, root2.left)

    return dfs(root.left, root.right)

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
    right=None):
        self.val = val # current node's value
        self.left = left # pointer to left child
        self.right = right # pointer to right child

class Solution:
    def isSymmetric(self, root: TreeNode) -> bool:
        # Helper function to check mirror symmetry
        between two trees.

```

```

def isMirror(t1: TreeNode, t2: TreeNode)
-> bool:
    # Both nodes are null, so they are
    mirrors.
    if t1 is None and t2 is None:
        return True
    # If one is null or the values do not
    match, symmetry fails.
    if t1 is None or t2 is None or t1.val
    != t2.val:
        return False
    # Recursively check mirror symmetry
    for children:
        # t1's left with t2's right and t1's
        right with t2's left.
        return isMirror(t1.left, t2.right) and
        isMirror(t1.right, t2.left)

    # Start recursion comparing the tree with
    itself.
    return isMirror(root, root)

```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isSymmetric(self, root:
Optional[TreeNode]) -> bool:
        def dfs(root1, root2):
            if root1 is None and root2 is None:

                return True
            if root1 is None or root2 is None or
root1.val != root2.val:
                return False
            return dfs(root1.left, root2.right)
and dfs(root1.right, root2.left)

        return dfs(root, root)
```

◆ Quick Review

Key Insights

- Use recursion (or iteration) to compare pairs of nodes.
- For two trees (or two nodes) to be mirror images:
- Their root values must be the same.

- The left subtree of one must match the right subtree of the other.
- The right subtree of one must match the left subtree of the other.
- Alternatively, iterative approaches using a queue can be applied to compare the nodes in pairs.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree, since we traverse each node once.

Space Complexity: $O(h)$ for recursive call stack (worst-case $O(n)$ for a skewed tree) or $O(n)$ for the iterative approach using a queue.

Solution

We solve the problem using recursion by defining a helper function that takes two nodes and checks if they are mirrors of each other. The helper function:

- Returns true if both nodes are null.
- Returns false if one node is null or if their values differ.
- Recursively compares the left child of the first node with the right child of the second node, and vice-versa. This method ensures each comparison verifies symmetry at every level. A similar logic can be applied iteratively by using a queue to compare pairs of nodes.

Trees & BST

#104 Maximum Depth of Binary Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS

Lists: Grind 169


Problem

Description

Given the root of a binary tree, return its maximum depth.

A binary tree's maximum depth is the number of nodes along the longest path from the root node down to the farthest leaf node.

Example 1:



Input: root = [3,9,20,null,null,15,7]

Output: 3

Example 2:

Input: root = [1,null,2]

Output: 2

Constraints:

- The number of nodes in the tree is in the range [0, 104].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def maxDepth(self, root: TreeNode | None) -> int:
        if not root:
            return 0
        return 1 + max(self.maxDepth(root.left),
self.maxDepth(root.right))
```

- LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def maxDepth(self, root: TreeNode) -> int:
        if root is None:
            return 0

        l, r = self.maxDepth(root.left),
self.maxDepth(root.right)
        return 1 + max(l, r)
```

- SimplyLeet

```
# Python implementation using DFS

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

class Solution:
```

```
    def maxDepth(self, root: Optional[TreeNode])
-> int:
        # Base case: if the node is null, return 0.
        if not root:
            return 0

        # Recursively compute the maximum depth of
the left subtree.
        left_depth = self.maxDepth(root.left)
        # Recursively compute the maximum depth of
the right subtree.
        right_depth = self.maxDepth(root.right)

        # Return the greater depth plus one for the
current node.
        return 1 + max(left_depth, right_depth)
```

- LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
# right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def maxDepth(self, root: TreeNode) -> int:
        if root is None:
            return 0

        l, r = self.maxDepth(root.left),
self.maxDepth(root.right)
        return 1 + max(l, r)
```

◆ Quick Review

Key Insights

- Use a recursive strategy (DFS) to traverse the binary tree.
- At each node, compute the maximum depth of its left and right subtree.
- The maximum depth at the current node is 1 (current node) plus the maximum of the left and right subtree depths.
- Alternatively, an iterative approach using BFS (level-order traversal) can be used.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since every node is visited once.

Space Complexity: $O(h)$ for recursion stack in DFS, where h is the height of the tree (worst-case $O(n)$ for a skewed tree); $O(n)$ for BFS in the worst-case scenario if the tree is balanced.

Solution

The problem is approached using Depth-First Search (DFS) with recursion. The idea is to visit every node and calculate the maximum depth of the left and right subtrees, then add 1 for the current node. This method naturally handles the base case when a null node (empty tree) is reached by returning 0. Alternatively, a Breadth-First Search (BFS) approach can determine the tree level by level, with the total number of levels equaling the maximum depth. Key data structures include recursion call stack for DFS or a queue for BFS.

Trees & BST

#108 Convert Sorted Array to Binary Search Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Divide and Conquer · Tree


Lists: Grind 169

Problem

Description

Given an integer array `nums` where the elements are sorted in ascending order, convert it to a height-balanced binary search tree.

Example 1:




Input: `nums = [-10,-3,0,5,9]`

Output: `[0,-3,9,-10,null,5]`

Explanation: `[0,-10,5,null,-3,null,9]` is also accepted:

Example 2:



Input: nums = [1,3]

Output: [3,1]

Explanation: [1,null,3] and [3,1] are both height-balanced BSTs.

Constraints:

- $1 \leq \text{nums.length} \leq 104$
- $-104 \leq \text{nums}[i] \leq 104$
- nums is sorted in a strictly increasing order.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def sortedArrayToBST(self, nums: list[int]) ->
    TreeNode | None:
        def build(l: int, r: int) -> TreeNode | None:
            if l > r:
                return None
            m = (l + r) // 2
            return TreeNode(nums[m],
                            build(l, m - 1),
                            build(m + 1, r))

        return build(0, len(nums) - 1)
  
```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
#     right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def sortedArrayToBST(self, nums: List[int]) ->
    Optional[TreeNode]:
        def dfs(l: int, r: int) ->
        Optional[TreeNode]:
            if l > r:
  
```

```

        return None
    mid = (l + r) >> 1
    return TreeNode(nums[mid], dfs(l, mid
- 1), dfs(mid + 1, r))

    return dfs(0, len(nums) - 1)

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    # Initializes a tree node with a value and
    optional left/right children.
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def sortedArrayToBST(self, nums):

```

```

        # Helper function to build BST from
        subarray indices.
        def buildBST(left, right):
            # Base condition: if no elements to
            process, return None.
            if left > right:
                return None
            # Choose the middle element as the root
            for balance.
            mid = (left + right) // 2
            node = TreeNode(nums[mid])
            # Recursively build the left subtree.
            node.left = buildBST(left, mid - 1)

            # Recursively build the right subtree.
            node.right = buildBST(mid + 1, right)
            return node
        # Build and return the BST from the entire
        array.
        return buildBST(0, len(nums) - 1)

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def sortedArrayToBST(self, nums: List[int]) ->
Optional[TreeNode]:
        def dfs(l, r):
            if l > r:

                return None
            mid = (l + r) >> 1
            left = dfs(l, mid - 1)
            right = dfs(mid + 1, r)
            return TreeNode(nums[mid], left,
right)

        return dfs(0, len(nums) - 1)

#####

```

```

# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

class Solution(object):
    def sortedArrayToBST(self, nums):
        """

        :type nums: List[int]
        :rtype: TreeNode
        """
        if nums:
            midPos = len(nums) / 2
            mid = nums[midPos]
            root = TreeNode(mid)
            root.left =
self.sortedArrayToBST(nums[:midPos])
            root.right =
self.sortedArrayToBST(nums[midPos + 1:])
            return root

```

◆ Quick Review

Key Insights

- The array is sorted in ascending order, which allows the use of a divide and conquer strategy.

- Choosing the middle element of the array as the root ensures that the left and right subtrees are roughly equal in size, yielding a balanced BST.
- A recursive approach is natural for building the tree: the base case is when the subarray is empty.
- The algorithm splits the array into two halves to recursively build the left and right subtrees.

Space & Time

Time Complexity: $O(n)$ – Every element is processed exactly once.

Space Complexity: $O(n)$ – $O(n)$ space is used to store the BST, and additional $O(\log n)$ space may be used for the recursion stack.

Solution

We use a divide and conquer approach with recursion. The main idea is to select the middle element of the current subarray as the root node so that the left subarray elements form the left subtree and the right subarray elements form the right subtree, ensuring the tree is balanced. The recursion continues until the subarray is empty, which serves as the base case. This strategy naturally maintains the BST property since all elements in the left subarray are smaller than the root, and all in the right subarray are larger.

Trees & BST

#110 Balanced Binary Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS


Lists: Grind 169

Problem

Description

Given a binary tree, determine if it is height-balanced.


Example 1:



Input: root = [3,9,20,null,null,15,7]

Output: true

Example 2:



Input: root = [1,2,2,3,3,null,null,4,4]

Output: false

Example 3:

Input: root = []

Output: true

Constraints:

- The number of nodes in the tree is in the range [0, 5000].
- $-104 \leq \text{Node.val} \leq 104$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def isBalanced(self, root: TreeNode | None) -> bool:
        if not root:
            return True

        def maxDepth(root: TreeNode | None) -> int:
            if not root:
                return 0
            return 1 + max(maxDepth(root.left),
                           maxDepth(root.right))

        return (abs(maxDepth(root.left) -
                    maxDepth(root.right)) <= 1 and
                self.isBalanced(root.left) and
                self.isBalanced(root.right))
  
```

```

class Solution:
    def isBalanced(self, root: TreeNode | None) ->
bool:
    ans = True

    def maxDepth(root: TreeNode | None) -> int:
        """Returns the height of root and sets ans to
false if root unbalanced."""
        nonlocal ans
        if not root or not ans:
            return 0
        left = maxDepth(root.left)

        right = maxDepth(root.right)
        if abs(left - right) > 1:
            ans = False
        return max(left, right) + 1

    maxDepth(root)
    return ans

```

```

class Solution:
    def isBalanced(self, root: TreeNode | None) ->
bool:
    def maxDepth(root: TreeNode | None) -> int:
        """Returns the height of root if root is
balanced; otherwise, returns -1."""
        if not root:
            return 0

        left = maxDepth(root.left)
        if left == -1:
            return -1

        right = maxDepth(root.right)
        if right == -1:
            return -1
        if abs(left - right) > 1:
            return -1

        return 1 + max(maxDepth(root.left),
maxDepth(root.right))

    return maxDepth(root) != -1

```

- LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isBalanced(self, root: Optional[TreeNode])
-> bool:
        def height(root):
            if root is None:
                return 0
            l, r = height(root.left),
height(root.right)
            if l == -1 or r == -1 or abs(l - r) >
1:
                return -1
            return 1 + max(l, r)

        return height(root) >= 0

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isBalanced(self, root: TreeNode) -> bool:
        # Helper function to return the height of
the subtree or -1 if unbalanced.

        def checkHeight(node: TreeNode) -> int:
            # An empty tree has height 0 and is
balanced.
            if not node:
                return 0

            # Recursively get the height of the
left subtree.
            leftHeight = checkHeight(node.left)
            # If left subtree is unbalanced, return
-1.
            if leftHeight == -1:
                return -1

```

```

        # Recursively get the height of the
right subtree.
        rightHeight = checkHeight(node.right)
        # If right subtree is unbalanced,
return -1.
        if rightHeight == -1:
            return -1

        # If the current node's subtrees differ
in height by more than 1, it's unbalanced.
        if abs(leftHeight - rightHeight) > 1:
            return -1

    # Return the height of the current
node.
    return max(leftHeight, rightHeight) +
1

    # The tree is balanced if checkHeight does
not return -1.
    return checkHeight(root) != -1

```

• LC.ca

Round 1 · High · Easy

p.239

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isBalanced(self, root: Optional[TreeNode])
-> bool:
        def height(root):
            if root is None:

                return 0
            l, r = height(root.left),
height(root.right)
            if l == -1 or r == -1 or abs(l - r) >
1:

                return -1
            return 1 + max(l, r)

        return height(root) >= 0

```

◆ Quick Review

Key Insights

- Use a bottom-up DFS (Depth-First Search) traversal to compute the height of each subtree.
- If any subtree is found unbalanced, propagate an error indicator (e.g., -1) upward, allowing early termination.

Round 1 · High · Easy

p.240

- An empty tree is considered balanced.
- The key check is to ensure that for each node, the absolute difference between the heights of its left and right subtrees is no more than 1.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes in the tree, since we visit each node once.

Space Complexity: $O(h)$ where h is the height of the tree, which corresponds to the recursion stack ($O(n)$ in the worst-case of a skewed tree).

Solution

The solution uses a recursive DFS approach. A helper function recursively computes the height of a subtree. If a subtree is found to be unbalanced (the height difference between its left and right subtrees exceeds 1), the helper returns a special value (e.g., -1) to indicate that the subtree is unbalanced. This value is then propagated up the recursion to terminate further unnecessary computations. The main function checks the result of the helper function and returns true if the tree is balanced (i.e., the helper did not return -1) and false otherwise.

Trees & BST

#112 Path Sum

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS

Lists: Denny Zhang


Problem

Description

Given the root of a binary tree and an integer `targetSum`, return true if the tree has a root-to-leaf path such that adding up all the values along the path equals `targetSum`.


A leaf is a node with no children.

Example 1:



Input: root =
 [5,4,8,11,null,13,4,7,2,null,null,null,1], targetSum =
 22
 Output: true
 Explanation: The root-to-leaf path with the target sum
 is shown.

Example 2:



Input: root = [1,2,3], targetSum = 5
 Output: false
 Explanation: There are two root-to-leaf paths in the
 tree:
 (1 --> 2): The sum is 3.
 (1 --> 3): The sum is 4.
 There is no root-to-leaf path with sum = 5.

Example 3:

Input: root = [], targetSum = 0

Output: false

Explanation: Since the tree is empty, there are no root-to-leaf paths.

Constraints:

- The number of nodes in the tree is in the range [0, 5000].
- $-1000 \leq \text{Node.val} \leq 1000$
- $-1000 \leq \text{targetSum} \leq 1000$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def hasPathSum(self, root: TreeNode, summ: int)
    -> bool:
        if not root:
            return False
        if root.val == summ and not root.left and not
        root.right:
            return True
        return (self.hasPathSum(root.left, summ -
        root.val) or
                self.hasPathSum(root.right, summ -
        root.val))
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
    right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def hasPathSum(self, root: Optional[TreeNode],
    targetSum: int) -> bool:
        def dfs(root, s):
            if root is None:
                return False
            s += root.val
            if root.left is None and root.right is
            None and s == targetSum:
                return True
            return dfs(root.left, s) or
            dfs(root.right, s)
        return dfs(root, 0)
```

• SimplyLeet

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # Node's value
        self.left = left # Pointer to left child
        self.right = right # Pointer to right child

def hasPathSum(root, targetSum):
    # If the current node is None, no valid path
    exists
    if not root:

        return False
    # Check if the current node is a leaf node
    if not root.left and not root.right:
        # Return true if the remaining targetSum
        equals the leaf value
        return targetSum == root.val
    # Recursively check the left and right subtree
    with the updated targetSum
    return hasPathSum(root.left, targetSum -
root.val) or hasPathSum(root.right, targetSum -
root.val)

# Example usage:
# root = TreeNode(5, TreeNode(4, TreeNode(11,
TreeNode(7), TreeNode(2))), TreeNode(8,
TreeNode(13), TreeNode(4, None, TreeNode(1))))
```

```
# print(hasPathSum(root, 22)) # Output: True
```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def hasPathSum(self, root: Optional[TreeNode],
targetSum: int) -> bool:
        def dfs(root, s):
            if root is None:

                return False
            s += root.val
            if root.left is None and root.right is
None and s == targetSum:
                return True
            return dfs(root.left, s) or
dfs(root.right, s)

        return dfs(root, 0)
```

◆ Quick Review

Key Insights

- Use a depth-first search (DFS) strategy to traverse the tree from the root to the leaves.
- At each node, subtract the node's value from the targetSum.
- Check if the node is a leaf; if yes and the remaining targetSum equals the node's value then a valid path is found.
- If the tree is empty, return false as no path exists.

Space & Time

Time Complexity: $O(N)$ where N is the number of nodes, since every node may need to be visited.

Space Complexity: $O(H)$ where H is the height of the tree, accounting for the recursion stack (worst-case $O(N)$ for a skewed tree).

Solution

The approach uses a recursive depth-first search. The algorithm subtracts the current node's value from targetSum as it moves down the tree. When it reaches a leaf, it checks if the modified targetSum is equal to the leaf's value. If yes, it returns true. The recursion continues until either a valid path is found or all paths are exhausted. The primary data structure used here is the call stack for recursion.

Trees & BST

#226 Invert Binary Tree

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS


Lists: Grind 169

Problem

Description

Given the root of a binary tree, invert the tree, and return its root.


Example 1:



Input: root = [4,2,7,1,3,6,9]

Output: [4,7,2,9,6,3,1]

Example 2:



Input: root = [2,1,3]
Output: [2,3,1]

Example 3:

Input: root = []
Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 100].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

- WalkCC

```

class Solution:
    def invertTree(self, root: TreeNode | None) ->
    TreeNode | None:
        if not root:
            return None

        left = root.left
        right = root.right
        root.left = self.invertTree(right)
        root.right = self.invertTree(left)
        return root
  
```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def invertTree(self, root: Optional[TreeNode])
    -> Optional[TreeNode]:
        if root is None:
            return None

        l, r = self.invertTree(root.left),
        self.invertTree(root.right)
        root.left, root.right = r, l
        return root
  
```

• SimplyLeet

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # Initialize node value
        self.left = left # Pointer to left child
        self.right = right # Pointer to right child

class Solution:
    def invertTree(self, root: TreeNode) ->
TreeNode:
    # Base case: if the node is None, return
None.

    if not root:
        return None
    # Recursively invert the left and right
subtrees.
    left_inverted = self.invertTree(root.left)
    right_inverted =
self.invertTree(root.right)
    # Swap the left and right children.
    root.left, root.right = right_inverted,
left_inverted
    return root
```

• LC.ca

Round 1 · High · Easy

p.253

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def invertTree(self, root: Optional[TreeNode])
-> Optional[TreeNode]:
        def dfs(root):
            if root is None:

                return
            root.left, root.right = root.right,
root.left
            dfs(root.left)
            dfs(root.right)

        dfs(root)
        return root
```

◆ Quick Review

Key Insights

- The problem can be solved using recursion (depth-first search) by swapping the children of every node.
- Alternatively, an iterative approach (using a queue for breadth-first search) can be used.

Round 1 · High · Easy

p.254

- The recursion terminates when a null node is encountered.
- The operation is performed in-place, modifying the original tree structure.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$ where h is the height of the tree, which is the recursion stack space in the recursive approach. In worst-case (skewed tree), this could be $O(n)$.

Solution

We solve the problem using recursion (depth-first search). The approach is to:

- Check the base case: if the node is null, simply return null.
- Recursively invert the left subtree and the right subtree.
- Swap the left and right children of the current node.
- Return the current node after the swap. The recursive approach elegantly handles all nodes and ensures the inversion process covers the complete tree.

Trees & BST

#270 Closest Binary Search Tree Value

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Binary Search · Tree · DFS · BST


Lists: Premium 98

Problem

Description

Given the root of a binary search tree and a target value, return the value in the BST that is closest to the target. If there are multiple answers, print the smallest.

Example 1:



Input: root = [4,2,5,1,3], target = 3.714286
Output: 4

Example 2:

Input: root = [1], target = 4.428571
Output: 1

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- $0 \leq \text{Node.val} \leq 109$
- $-109 \leq \text{target} \leq 109$

Round 1 · High · Easy

p.257

Community Solutions (Python)

• WalkCC

```

class Solution:
    def closestValue(self, root: TreeNode | None,
target: float) -> int:
        # If target < root.val, search the left
        subtree.
        if target < root.val and root.left:
            left = self.closestValue(root.left, target)
            if abs(left - target) <= abs(root.val -
target):
                return left

        # If target > root.val, search the right
        subtree.
        if target > root.val and root.right:

            right = self.closestValue(root.right,
target)
            if abs(right - target) < abs(root.val -
target):
                return right

        return root.val
  
```

• LeetCode

Round 1 · High · Easy

p.258

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def closestValue(self, root:
Optional[TreeNode], target: float) -> int:
        def dfs(node: Optional[TreeNode]):
            if node is None:
```

```
                return
                nxt = abs(target - node.val)
                nonlocal ans, diff
                if nxt < diff or (nxt == diff and
node.val < ans):
                    diff = nxt
                    ans = node.val
                    node = node.left if target < node.val
else node.right
                    dfs(node)
```

```
        ans = 0
```

```
        diff = inf
        dfs(root)
        return ans
```

• SimplyLeet

Round 1 · High · Easy

p.259

Python solution for Closest Binary Search Tree Value

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
```

```
class Solution:
```

```
    def closestValue(self, root: TreeNode, target:
float) -> int:
        # Initialize the best value with the value
of the root
        best_value = root.val

        # Traverse the BST iteratively
        current = root
        while current:
            # Calculate the difference between
current node's value and target
            current_diff = abs(current.val -
target)

            best_diff = abs(best_value - target)
```

Round 1 · High · Easy

p.260

```

        # Update best_value if a closer
        difference is found, or if equal and current value
        is smaller
        if current_diff < best_diff or
        (current_diff == best_diff and current.val <
        best_value):
            best_value = current.val

        # Decide the traversal direction based
        on the target comparison
        if target < current.val:
            current = current.left
        else:
            current = current.right

    return best_value

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode(object):
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution(object):
    def closestValue(self, p, target):
        """

```

```

:type root: TreeNode
:type target: float
:rtype: int
"""

ans = p.val
while p:
    if abs(target - p.val) < abs(ans - target):
        ans = p.val
        if target < p.val:
            p = p.left

    else:
        p = p.right
return ans

#####

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
#         right=None):

```

```

# self.val = val
# self.left = left
# self.right = right
class Solution:
    def closestValue(self, root:
Optional[TreeNode], target: float) -> int:
        ans, mi = root.val, inf
        while root:
            t = abs(root.val - target)
            if t < mi:
                mi = t

            ans = root.val
            if root.val > target:
                root = root.left
            else:
                root = root.right
        return ans

```

◆ Quick Review

Key Insights

- BST property allows us to traverse the tree efficiently by comparing the target with node values.
- We can iteratively traverse the tree, moving left if the target is smaller than the current node's value and moving right if greater.
- At each node, calculate the absolute difference between node's value and the target, and update the answer if the difference is smaller—or equal with a smaller candidate

value.

- This approach minimizes the search area by leveraging the inherent ordering in BSTs.

Space & Time

Time Complexity: $O(h)$ on average, where h is the height of the tree ($O(\log n)$ for a balanced tree, and $O(n)$ in the worst-case scenario for a skewed tree).

Space Complexity: $O(1)$ for an iterative solution (or $O(h)$ if using recursion due to the call stack).

Solution

We solve the problem by performing an iterative traversal of the BST. Starting at the root, we maintain a variable to store the value closest to the target. For every node encountered, we:

- Compute the absolute difference between the node's value and the target.
- Update the stored closest value if the computed difference is less than the current smallest difference; if the difference is the same, choose the smaller node value.
- Depending on whether the target is less than or greater than the current node's value, move to the left or right child accordingly. This approach efficiently defines a search path that is consistent with the properties of a BST.

Trees & BST

#501 Find Mode in Binary Search Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BST

Lists: Denny Zhang

Problem

Description


Given the root of a binary search tree (BST) with duplicates, return all the mode(s) (i.e., the most frequently occurred element) in it.

If the tree has more than one mode, return them in any order.

Assume a BST is defined as follows:

- The left subtree of a node contains only nodes with keys less than or equal to the node's key.
- The right subtree of a node contains only nodes with keys greater than or equal to the node's key.
- Both the left and right subtrees must also be binary search trees.

Example 1:



Input: root = [1,null,2,2]
Output: [2]

Example 2:

Input: root = [0]
Output: [0]

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- $-105 \leq \text{Node.val} \leq 105$

Follow up: Could you do that without using any extra space? (Assume that the implicit stack space incurred due to recursion does not count).

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findMode(self, root: TreeNode | None) -> list[int]:
        self.ans = []
        self.pred = None
        self.count = 0
        self.maxCount = 0

        def updateCount(root: TreeNode | None) -> None:
            if self.pred and self.pred.val == root.val:
                self.count += 1

        else:
            self.count = 1

        if self.count > self.maxCount:
            self.maxCount = self.count
            self.ans = [root.val]
        elif self.count == self.maxCount:
            self.ans.append(root.val)

        self.pred = root
```

```
def inorder(root: TreeNode | None) -> None:
    if not root:
        return

    inorder(root.left)
    updateCount(root)
    inorder(root.right)

inorder(root)

return self.ans
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def findMode(self, root: TreeNode) -> List[int]:
        def dfs(root):
            if root is None:
```

```

        return
    nonlocal mx, prev, ans, cnt
    dfs(root.left)
    cnt = cnt + 1 if prev == root.val else
1
    if cnt > mx:
        ans = [root.val]
        mx = cnt
    elif cnt == mx:
        ans.append(root.val)
    prev = root.val

    dfs(root.right)

    prev = None
    mx = cnt = 0
    ans = []
    dfs(root)
    return ans

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right

class Solution:
    def findMode(self, root):
        # Initialize state variables: prev value,
        # current frequency count, maximum frequency, and
        # modes list.

        self.prev = None
        self.currentCount = 0
        self.maxCount = 0
        self.modes = []

        # First in-order traversal to find the
        # maximum frequency (maxCount).
        def inorder_first(node):
            if not node:
                return
            inorder_first(node.left)

```

```

        # Update the count for the current
node.
        if self.prev is None or node.val !=
self.prev:
            self.currentCount = 1
        else:
            self.currentCount += 1
        self.prev = node.val
        # Update maxCount if currentCount
exceeds it.
        self.maxCount = max(self.maxCount,
self.currentCount)
        inorder_first(node.right)

```

```
inorder_first(root)
```

```

# Second in-order traversal to collect all
node values with frequency equal to maxCount.
self.prev = None # Reset previous value.
self.currentCount = 0 # Reset count.

```

```

def inorder_second(node):
    if not node:
        return
    inorder_second(node.left)

```

```

        if self.prev is None or node.val !=
self.prev:
            self.currentCount = 1
        else:
            self.currentCount += 1
        self.prev = node.val
        if self.currentCount == self.maxCount:
            self.modes.append(node.val)
            inorder_second(node.right)

        inorder_second(root)

    return self.modes

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def findMode(self, root: TreeNode) ->
List[int]:
        def dfs(root):
            if root is None:

```

```

        return
    nonlocal mx, prev, ans, cnt
    dfs(root.left)
    cnt = cnt + 1 if prev == root.val else
1
    if cnt > mx:
        ans = [root.val]
        mx = cnt
    elif cnt == mx:
        ans.append(root.val)
    prev = root.val

    dfs(root.right)

    prev = None
    mx = cnt = 0
    ans = []
    dfs(root)
    return ans

```

◆ Quick Review

Key Insights

- The BST in-order traversal yields a sorted order, which helps in counting consecutive duplicates.
- A two-pass in-order traversal can be used: one pass to determine the maximum frequency (maxCount) and a second pass to collect all values that match this frequency.
- The solution can be implemented without extra space (apart from recursion stack) by leveraging constant space

counters and state variables.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is processed twice (once in each traversal).

Space Complexity: $O(h)$ where h is the height of the tree due to the recursion stack. $O(1)$ extra space is used aside from this.

Solution

We perform an in-order traversal of the BST twice. The first traversal determines the maximum frequency of any value in the BST by comparing the current node's value with the previously visited node. Since the in-order traversal processes nodes in sorted order, duplicate values are consecutive, making the counting simple. After calculating maxCount, we reset our state and perform a second in-order traversal to collect all node values with frequency equal to maxCount. The key trick is to leverage the BST property (sorted in-order traversal) to count frequencies in one pass without using additional data structures like hash maps.

Trees & BST

#543 Diameter of Binary Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Grind 169

Problem


Description

Given the root of a binary tree, return the length of the diameter of the tree.

The diameter of a binary tree is the length of the longest path between any two nodes in a tree. This path may or may not pass through the root.

The length of a path between two nodes is represented by the number of edges between them.

Example 1:



Input: root = [1,2,3,4,5]

Output: 3

Explanation: 3 is the length of the path [4,2,1,3] or [5,2,1,3].

Example 2:

Input: root = [1,2]

Output: 1

Constraints:

- The number of nodes in the tree is in the range [1, 104].

- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def diameterOfBinaryTree(self, root: TreeNode | None) -> int:
        ans = 0

        def maxDepth(root: TreeNode | None) -> int:
            nonlocal ans
            if not root:
                return 0

            l = maxDepth(root.left)

            r = maxDepth(root.right)
            ans = max(ans, l + r)
            return 1 + max(l, r)

        maxDepth(root)
        return ans
```

- LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def diameterOfBinaryTree(self, root: TreeNode) -> int:
        def dfs(root):
            if root is None:
                return 0
            nonlocal ans
            left, right = dfs(root.left), dfs(root.right)
            ans = max(ans, left + right)
            return 1 + max(left, right)

        ans = 0
        dfs(root)
        return ans
```

- SimplyLeet

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

class Solution:
    def diameterOfBinaryTree(self, root: TreeNode)
-> int:
        # Global variable to track the maximum
diameter found
```

```
        self.max_diameter = 0

    def dfs(node):
        # Base case: if node is None, depth is
0.
        if not node:
            return 0
        # Compute the depth of left and right
subtrees.
        left_depth = dfs(node.left)
        right_depth = dfs(node.right)
        # Update the global diameter: the path
through this node
```

```
        self.max_diameter =
max(self.max_diameter, left_depth + right_depth)
        # Return the depth of the tree rooted
at this node.
        return max(left_depth, right_depth) +
1

    def dfs(root):
        return self.max_diameter
```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def diameterOfBinaryTree(self, root: TreeNode)
-> int:
        def dfs(root):
            if root is None:
```



```

        return 0
    nonlocal ans
    left, right = dfs(root.left),
dfs(root.right)
    ans = max(ans, left + right)
    return 1 + max(left, right)

ans = 0
dfs(root)
return ans

```

◆ Quick Review

Key Insights

- Use depth-first search (DFS) to explore the tree.
- The diameter at a node is the sum of the depths of its left and right subtrees.
- Track the maximum diameter found during the DFS traversal.
- Recursively calculating the depth of subtrees provides an efficient solution.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree.

Space Complexity: $O(n)$, in the worst-case scenario due to the recursion stack (skewed tree).

Solution

The solution employs a DFS traversal of the binary tree. At every node in the tree, compute the depth of the left and right subtrees. The sum of these depths gives the potential diameter (longest path through that node). Update a global maximum diameter variable if this sum exceeds the current maximum. Finally, return the maximum diameter after completing the DFS.

Trees & BST

#572 Subtree of Another Tree

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · String Matching

Lists: Grind 169


Problem

Description

Given the roots of two binary trees `root` and `subRoot`, return `true` if there is a subtree of `root` with the same structure and node values of `subRoot` and `false` otherwise.

A subtree of a binary tree `tree` is a tree that consists of a node in `tree` and all of this node's descendants. The tree `tree` could also be considered as a subtree of itself.


Example 1:



Input: `root = [3,4,5,1,2]`, `subRoot = [4,1,2]`

Output: `true`

Example 2:



Input: root = [3,4,5,1,2,null,null,null,null,0],
 subRoot = [4,1,2]
 Output: false

Constraints:

- The number of nodes in the root tree is in the range [1, 2000].
- The number of nodes in the subRoot tree is in the range [1, 1000].
- $-104 \leq \text{root.val} \leq 104$
- $-104 \leq \text{subRoot.val} \leq 104$

Community Solutions (Python)

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isSubtree(self, root: Optional[TreeNode],
                  subRoot: Optional[TreeNode]) -> bool:
        def same(p: Optional[TreeNode], q:
                  Optional[TreeNode]) -> bool:
            if p is None or q is None:
                return p is q
            return p.val == q.val and same(p.left,
                                             q.left) and same(p.right, q.right)

        if root is None:
            return False
        return (
            same(root, subRoot)
            or self.isSubtree(root.left, subRoot)
            or self.isSubtree(root.right, subRoot)
        )
  
```

• SimplyLeet

Python solution for checking if subRoot is a subtree of root.

Definition for a binary tree node.

class TreeNode:

def __init__(self, val=0, left=None, right=None):

self.val = val

self.left = left

self.right = right

class Solution:

def isSubtree(self, root: TreeNode, subRoot: TreeNode) -> bool:

Helper function to check if two trees are identical.

def isSameTree(s, t):

Both nodes are None, trees match.

if not s and not t:

return True

One node is None or values differ.

if not s or not t or s.val != t.val:

return False

Recurse on left and right subtrees.

return isSameTree(s.left, t.left) and
isSameTree(s.right, t.right)

Base case: if root is None, then no
subtree exists.

if not root:

return False

Check if current subtree is identical to
subRoot.

if isSameTree(root, subRoot):

return True

Recursively check left and right
subtrees.

return self.isSubtree(root.left, subRoot)
or self.isSubtree(root.right, subRoot)

• LC.ca

Definition for a binary tree node.

class TreeNode:

def __init__(self, val=0, left=None, right=None):

self.val = val

self.left = left

self.right = right

class Solution:

def isSubtree(self, root: TreeNode, subRoot: TreeNode) -> bool:

def dfs(root1, root2):

if root1 is None and root2 is None:

```

        return True
    if root1 is None or root2 is None:
        return False
    return (
        root1.val == root2.val
        and dfs(root1.left, root2.left)
        and dfs(root1.right, root2.right)
    )

if root is None:

    return False
return (
    dfs(root, subRoot)
    or self.isSubtree(root.left, subRoot)
    or self.isSubtree(root.right, subRoot)
)

```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to visit each node in the main tree.
- At each node, check if the subtree starting from that node is identical to subRoot.
- A helper function is used to compare the structure and node values of two trees.
- Consider recursion as the primary approach to traverse and compare trees.

- Early exit when a mismatch is found can optimize average-case performance.

Space & Time

Time Complexity: $O(m * n)$ in the worst case, where m is the number of nodes in the main tree and n is the number of nodes in the subRoot tree.

Space Complexity: $O(\max(\text{depth of main tree, depth of subRoot}))$ due to recursive call stack.

Solution

We solve the problem by performing a DFS traversal on the main tree. For each visited node, a helper function checks whether the subtree starting at that node is identical to subRoot by comparing node values and recursively checking both left and right subtrees. This method leverages recursion effectively and provides early termination if mismatches occur, allowing for efficient traversal and comparison.

DFS

Round 1 · High · Easy · 2 questions

Round 1 · High · Easy

p.291

DFS

#463 Island Perimeter

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Matrix

Lists: Denny Zhang

Problem

Description

You are given row x col grid representing a map where $\text{grid}[i][j] = 1$ represents land and $\text{grid}[i][j] = 0$ represents water.

Grid cells are connected horizontally/vertically (not diagonally). The grid is completely surrounded by water, and there is exactly one island (i.e., one or more connected land cells).

The island doesn't have "lakes", meaning the water inside isn't connected to the water around the island.


One cell is a square with side length 1. The grid is rectangular, width and height don't exceed 100.

Determine the perimeter of the island.

Example 1:

Round 1 · High · Easy

p.292



Input: grid = [[0,1,0,0],[1,1,1,0],[0,1,0,0],[1,1,0,0]]
 Output: 16
 Explanation: The perimeter is the 16 yellow stripes in the image above.

Example 2:

Input: grid = [[1]]
 Output: 4

Example 3:

Input: grid = [[1,0]]
 Output: 4

Constraints:

- row == grid.length
- col == grid[i].length
- 1 <= row, col <= 100
- grid[i][j] is 0 or 1.
- There is exactly one island in grid.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def islandPerimeter(self, grid: list[list[int]])
    -> int:
        m = len(grid)
        n = len(grid[0])

        islands = 0
        neighbors = 0

        for i in range(m):
            for j in range(n):
```

```

if grid[i][j] == 1:
    islands += 1
    if i + 1 < m and grid[i + 1][j] == 1:
        neighbors += 1
    if j + 1 < n and grid[i][j + 1] == 1:
        neighbors += 1

return islands * 4 - neighbors * 2

```

• LeetCode

```

class Solution:
    def islandPerimeter(self, grid:
List[List[int]]) -> int:
    m, n = len(grid), len(grid[0])
    ans = 0
    for i in range(m):
        for j in range(n):
            if grid[i][j] == 1:
                ans += 4
                if i < m - 1 and grid[i + 1][j]
== 1:
                    ans -= 2

                if j < n - 1 and grid[i][j + 1]
== 1:
                    ans -= 2

    return ans

```

• SimplyLeet

Python solution for Island Perimeter

```

def islandPerimeter(grid):
    # Initialize perimeter to 0
    perimeter = 0

    # Get number of rows and columns
    rows = len(grid)
    cols = len(grid[0])

    # Iterate through each cell in the grid

    for i in range(rows):
        for j in range(cols):
            # Only process if current cell is land
            if grid[i][j] == 1:
                # Each land cell contributes 4
                edges initially
                cell_perimeter = 4

                # Check if the upper neighbor is
                also land
                if i > 0 and grid[i-1][j] == 1:
                    cell_perimeter -= 1

```



```

        # Check if the lower neighbor is
also land    if i < rows - 1 and grid[i+1][j] ==
1:
        cell_perimeter -= 1
        # Check if the left neighbor is
also land    if j > 0 and grid[i][j-1] == 1:
        cell_perimeter -= 1
        # Check if the right neighbor is
also land    if j < cols - 1 and grid[i][j+1] ==
1:
        cell_perimeter -= 1

```

```

        # Add the effective perimeter of
the current cell to the overall perimeter
        perimeter += cell_perimeter

```

```

    return perimeter

```

```

# Example usage:
# grid = [[0,1,0,0],[1,1,1,0],[0,1,0,0],[1,1,0,0]]
# print(islandPerimeter(grid)) # Output should be
16

```

• [LC.ca](#)

Round 1 · High · Easy

p.297

```

class Solution:
    def islandPerimeter(self, grid:
List[List[int]]) -> int:
        m, n = len(grid), len(grid[0])
        ans = 0
        for i in range(m):
            for j in range(n):
                if grid[i][j] == 1:
                    ans += 4
                    if i < m - 1 and grid[i + 1][j]
== 1:
                        ans -= 2
                    if j < n - 1 and grid[i][j + 1]
== 1:
                        ans -= 2
        return ans

```

◆ Quick Review

Key Insights

- Each land cell contributes 4 edges if isolated.
- For every adjacent pair of land cells (neighbors horizontally or vertically), the shared edge should be subtracted from the total perimeter.
- Instead of DFS/BFS, we can iterate the grid and check the four neighbors for each land cell.
- Alternatively, one may use DFS/BFS to visit all connected cells if needed, but an iterative solution is straightforward

Round 1 · High · Easy

p.298

with $O(\text{rows} * \text{cols})$ time complexity.

Space & Time

Time Complexity: $O(n * m)$, where n is the number of rows and m is the number of columns.

Space Complexity: $O(1)$ for the iterative solution (excluding the input storage).

Solution

We iterate through every cell in the grid. For each land cell, we initially add 4 to the perimeter. Then we check its four neighbors (up, down, left, right). For every neighbor that is also land, we subtract 1 from the current cell's contribution since that side is not exposed. This simple approach ensures that shared borders between adjacent land cells are not over-counted.

DFS

#733 Flood Fill

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Matrix

Lists: Grind 169

Problem

Description

You are given an image represented by an $m \times n$ grid of integers `image`, where `image[i][j]` represents the pixel value of the image. You are also given three integers `sr`, `sc`, and `color`. Your task is to perform a flood fill on the image starting from the pixel `image[sr][sc]`.

To perform a flood fill:

1. Begin with the starting pixel and change its color to color.
2. Perform the same process for each pixel that is directly adjacent (pixels that share a side with the original pixel, either horizontally or vertically) and shares the same color as the starting pixel.
3. Keep repeating this process by checking neighboring pixels of the updated pixels and modifying their color if it matches the original color of the starting pixel.
4. The process stops when there are no more adjacent pixels of the original color to update.


Return the modified image after performing the flood fill.

Example 1:

Input: image = [[1,1,1],[1,1,0],[1,0,1]], sr = 1, sc = 1, color = 2

Output: [[2,2,2],[2,2,0],[2,0,1]]

Explanation:



From the center of the image with position (sr, sc) = (1, 1) (i.e., the red pixel), all pixels connected by a path of the same color as the starting pixel (i.e., the blue pixels) are colored with the new color.

Note the bottom corner is not colored 2, because it is not horizontally or vertically connected to the starting pixel.

Example 2:

Input: image = [[0,0,0],[0,0,0]], sr = 0, sc = 0, color = 0

Output: [[0,0,0],[0,0,0]]

Explanation:

The starting pixel is already colored with 0, which is the same as the target color. Therefore, no changes are made to the image.

Constraints:

- m == image.length
- n == image[i].length
- 1 <= m, n <= 50
- 0 <= image[i][j], color < 216
- 0 <= sr < m
- 0 <= sc < n

Community Solutions (Python)

• WalkCC

```
class Solution:
    def floodFill(self, image: list[list[int]],
                  sr: int, sc: int, newColor: int) ->
list[list[int]]:
    startColor = image[sr][sc]
    seen = set()

    def dfs(i: int, j: int) -> None:
        if i < 0 or i == len(image) or j < 0 or j ==
len(image[0]):
            return
        if image[i][j] != startColor or (i, j) in
seen:
```

```

    return

    image[i][j] = newColor
    seen.add((i, j))

    dfs(i + 1, j)
    dfs(i - 1, j)
    dfs(i, j + 1)
    dfs(i, j - 1)

```

```

dfs(sr, sc)
return image

```

• LeetCode

```

class Solution:
    def floodFill(
        self, image: List[List[int]], sr: int, sc:
        int, color: int
    ) -> List[List[int]]:
        def dfs(i: int, j: int):
            image[i][j] = color
            for a, b in pairwise(dirs):
                x, y = i + a, j + b
                if 0 <= x < len(image) and 0 <= y <
                len(image[0]) and image[x][y] == oc:
                    dfs(x, y)

```

```

    oc = image[sr][sc]
    if oc != color:
        dirs = (-1, 0, 1, 0, -1)
        dfs(sr, sc)
    return image

```

• SimplyLeet

```

def floodFill(image, sr, sc, color):
    # Get the original color of the starting pixel
    orig_color = image[sr][sc]
    # If the new color is the same as the original,
    return the image as is
    if orig_color == color:
        return image
    rows, cols = len(image), len(image[0])

    # Recursive DFS helper function to fill the
    image
    def dfs(r, c):

```

```

        # Check boundaries and if the current pixel
        is the same as the original color
        if r < 0 or r >= rows or c < 0 or c >= cols
        or image[r][c] != orig_color:
            return
        # Update the current pixel's color
        image[r][c] = color
        # Recursively call DFS in all four
        directions
        dfs(r - 1, c) # up
        dfs(r + 1, c) # down
        dfs(r, c - 1) # left
        dfs(r, c + 1) # right

```

```

dfs(sr, sc)
return image

```

• LC.ca

```

class Solution:
    def floodFill(
        self, image: List[List[int]], sr: int, sc:
        int, color: int
    ) -> List[List[int]]:
        def dfs(i, j):
            if (
                not 0 <= i < m
                or not 0 <= j < n
                or image[i][j] != oc
                or image[i][j] == color

```

```

):
    return
    image[i][j] = color
    for a, b in pairwise(dirs):
        dfs(i + a, j + b)

```

```

dirs = (-1, 0, 1, 0, -1)
m, n = len(image), len(image[0])
oc = image[sr][sc]
dfs(sr, sc)

```

```

return image

```

◆ Quick Review

Key Insights

- The flood fill operation can be executed using either DFS (Depth First Search) or BFS (Breadth First Search).
- A check is needed at the beginning to determine if the new color is the same as the original; if so, no changes are needed.
- The algorithm recursively or iteratively visits each pixel that is directly adjacent and matches the original color, then updates its color.
- The problem has a worst-case scenario where all pixels are the same color, leading to a full traversal of the grid.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the dimensions of the image grid, because in the worst-case all

cells are visited.

Space Complexity: $O(m * n)$ in the worst-case due to the recursion stack (or BFS queue) if the entire image needs to be filled.

Solution

We use a DFS approach to perform the flood fill. The algorithm starts at the given starting pixel and uses recursion to fill connected pixels with the same original color. The following steps are key:

- Check if the starting pixel's color is already the target color; if so, return the image immediately.
- Define a recursive helper function that: Checks boundary conditions. Updates the current pixel's color. Recursively calls itself on the four adjacent directions (up, down, left, right).
- Return the modified image after the recursion completes.

Data structures used:

- Recursion call stack for the DFS implementation.
- Variables to store the image dimensions and original color.

Round 2 · High · Medium

Round 2

High Priority · Medium

116 questions · 11 patterns

Arrays & Hashing #57 #238 #281 #307 #454 #525 #560
#1010 #1272 #1429 #3159

String #8 #249 #271 #635 #1138

Two Pointers #11 #15 #16 #31 #75 #161 #167 #186 #189
#532 #923 #986 #1055 #2563

Sliding Window #3 #159 #340 #424 #438 #487 #1100 #1248

Sorting #49 #56 #1152

Binary Search #33 #34 #153 #300 #362 #528 #852 #981
#1011 #1060 #1062 #1533

Matrix #36 #48 #54 #59 #64 #73 #74 #221 #311 #498 #531
#723 #835 #1198

Trees & BST #98 #102 #103 #105 #199 #230 #235 #236
#250 #285 #298 #314 #333 #366 #437 #545 #549 #582
#662 #863 #1120 #1490 #1522 #1650

Round 1 · High · Easy

p.307

Round 1 · High · Easy

p.308

**DFS #130 #133 #200 #207 #210 #261 #310 #323 #417 #490
#694 #695 #721 #785 #1236 #1242**

Graphs #128 #277 #1136

BFS #286 #322 #542 #994 #1197 #1730

Arrays & Hashing

Round 2 · High · Medium · 11 questions

Round 2 · High · Medium

p.309

Round 2 · High · Medium

p.310

Arrays & Hashing

#57 Insert Interval

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array

Lists: Grind 169

Problem

Description

You are given an array of non-overlapping intervals `intervals` where `intervals[i] = [starti, endi]` represent the start and the end of the *i*th interval and `intervals` is sorted in ascending order by `starti`. You are also given an interval `newInterval = [start, end]` that represents the start and end of another interval.

Insert `newInterval` into `intervals` such that `intervals` is still sorted in ascending order by `starti` and `intervals` still does not have any overlapping intervals (merge overlapping intervals if necessary).

Return `intervals` after the insertion.

Note that you don't need to modify `intervals` in-place. You can make a new array and return it.

Example 1:

Input: `intervals = [[1,3],[6,9]]`, `newInterval = [2,5]`
 Output: `[[1,5],[6,9]]`

Example 2:

Input: `intervals = [[1,2],[3,5],[6,7],[8,10],[12,16]]`,
`newInterval = [4,8]`
 Output: `[[1,2],[3,10],[12,16]]`
 Explanation: Because the new interval `[4,8]` overlaps with `[3,5]`, `[6,7]`, `[8,10]`.

Constraints:

- `0 <= intervals.length <= 104`
- `intervals[i].length == 2`
- `0 <= starti <= endi <= 105`
- `intervals` is sorted by `starti` in ascending order.
- `newInterval.length == 2`
- `0 <= start <= end <= 105`

Community Solutions (Python)

- WalkCC


```

class Solution:
    def insert(self, intervals: list[list[int]],
               newInterval: list[int]) ->
list[list[int]]:
    n = len(intervals)
    ans = []
    i = 0

    while i < n and intervals[i][1] <
newInterval[0]:
        ans.append(intervals[i])
        i += 1

```

```

    # Merge overlapping intervals.
    while i < n and intervals[i][0] <=
newInterval[1]:
        newInterval[0] = min(newInterval[0],
intervals[i][0])
        newInterval[1] = max(newInterval[1],
intervals[i][1])
        i += 1

    ans.append(newInterval)

    while i < n:

```

```

        ans.append(intervals[i])
        i += 1

```

```

    return ans

```

• LeetCode

```

class Solution:
    def insert(
        self, intervals: List[List[int]],
        newInterval: List[int]
    ) -> List[List[int]]:
        def merge(intervals: List[List[int]]) ->
List[List[int]]:
            intervals.sort()
            ans = [intervals[0]]
            for s, e in intervals[1:]:
                if ans[-1][1] < s:
                    ans.append([s, e])

            else:
                ans[-1][1] = max(ans[-1][1],
e)

            return ans

        intervals.append(newInterval)
        return merge(intervals)

```

• SimplyLeet

Python solution for Insert Interval

```
class Solution:
    def insert(self, intervals, newInterval):
        # List to store the resulting intervals
        merged_intervals = []
        i = 0
        n = len(intervals)

        # Add all intervals that end before
        newInterval starts
        while i < n and intervals[i][1] <
        newInterval[0]:
            merged_intervals.append(intervals[i])
            i += 1

        # Merge overlapping intervals with
        newInterval
        while i < n and intervals[i][0] <=
        newInterval[1]:
            # Update the start and end of
            newInterval by merging with overlapping intervals
            newInterval[0] = min(newInterval[0],
            intervals[i][0])
            newInterval[1] = max(newInterval[1],
            intervals[i][1])
            i += 1
```

```
# Add the merged newInterval
merged_intervals.append(newInterval)
```

```
# Add any remaining intervals that come
after newInterval
while i < n:
    merged_intervals.append(intervals[i])
    i += 1

return merged_intervals
```

• LC.ca

```
...
>>> a = [1,2,3,4,5]
>>> a[1]
2
>>> a[~1]
4
...
class Solution:
    def insert(self, intervals: List[List[int]],
    newInterval: List[int]) -> List[List[int]]:
        start = newInterval[0]
```

```

        end = newInterval[1]

        left = list(filter(lambda x: x[1] < start,
            intervals)) # cast via list()
        right = list(filter(lambda x: x[0] > end,
            intervals))

        if left + right != intervals:
            start = min(start,
                intervals[len(left)][0]) # note, left not -1,
            # because index starts at 0
            end = max(end,
                intervals[~len(right)][1]) # same, starting at
            # right with index=0

        return left + [[start, end]] + right

```

```

#####

class Solution: # re-use Leetcode-56's merge
    solution
    def insert(self, intervals: List[List[int]],
        newInterval: List[int]) -> List[List[int]]:
        intervals.append(newInterval)
        return self.merge(intervals)

    def merge(self, intervals: List[List[int]]) ->
        List[List[int]]:
        ans = []

```

```

        for intv in sorted(intervals, key=lambda
            x: x[0]):
            if ans and ans[-1][1] >= intv[0]:
                ans[-1][1] = max(ans[-1][1],
                    intv[1])
            else:
                ans.append(intv)
        return ans

```

#####

class Solution:

```

    def insert(
        self, intervals: List[List[int]],
        newInterval: List[int]
    ) -> List[List[int]]:
        def merge(intervals: List[List[int]]) ->
            List[List[int]]:
            intervals.sort()
            ans = [intervals[0]]
            for s, e in intervals[1:]:
                if ans[-1][1] < s:
                    ans.append([s, e])
            else:

```

```

e)
        ans[-1][1] = max(ans[-1][1],
            return ans

        intervals.append(newInterval)
        return merge(intervals)

```

◆ Quick Review

Key Insights

- The intervals array is sorted by start times.
- Three main steps are used:
- Add intervals that come completely before the new interval.
- Merge overlapping intervals with the new interval.
- Append intervals that start after the new (merged) interval.
- Iterating through the list once allows an $O(n)$ solution.

Space & Time

Time Complexity: $O(n)$ since we scan the list once.

Space Complexity: $O(n)$ to store the resulting list of intervals.

Solution

Traverse the intervals list and perform the following:

- Append all intervals that end before the new interval starts.
- For intervals that overlap with newInterval, update newInterval to merge them (using min for start and max for

end). – Append the new merged interval. – Append all remaining intervals. This ensures that all intervals remain non-overlapping and sorted with a single pass.

Arrays & Hashing

#238 Product of Array Except Self

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Prefix Sum

Lists: Grind 169

Problem

Description

Given an integer array `nums`, return an array `answer` such that `answer[i]` is equal to the product of all the elements of `nums` except `nums[i]`.

The product of any prefix or suffix of `nums` is guaranteed to fit in a 32-bit integer.

You must write an algorithm that runs in $O(n)$ time and without using the division operation.

Example 1:

Input: `nums = [1,2,3,4]`
 Output: `[24,12,8,6]`

Example 2:

Input: `nums = [-1,1,0,-3,3]`
 Output: `[0,0,9,0,0]`

Constraints:

- $2 \leq \text{nums.length} \leq 105$

- $-30 \leq \text{nums}[i] \leq 30$
- The input is generated such that `answer[i]` is guaranteed to fit in a 32-bit integer.

Follow up: Can you solve the problem in $O(1)$ extra space complexity? (The output array does not count as extra space for space complexity analysis.)

Community Solutions (Python)

• WalkCC

```
class Solution:
    def productExceptSelf(self, nums: list[int]) -> list[int]:
        n = len(nums)
        prefix = [1] * n # prefix product
        suffix = [1] * n # suffix product

        for i in range(1, n):
            prefix[i] = prefix[i - 1] * nums[i - 1]

        for i in reversed(range(n - 1)):
            suffix[i] = suffix[i + 1] * nums[i + 1]

        return [prefix[i] * suffix[i] for i in range(n)]
```

```

class Solution:
    def productExceptSelf(self, nums: list[int]) -> list[int]:
        n = len(nums)
        ans = [1] * n

        # Use ans as the prefix product array.
        for i in range(1, n):
            ans[i] = ans[i - 1] * nums[i - 1]

        suffix = 1 # suffix product

        for i, num in reversed(list(enumerate(nums))):
            ans[i] *= suffix
            suffix *= num

        return ans

```

• LeetCode

```

class Solution:
    def productExceptSelf(self, nums: List[int]) -> List[int]:
        n = len(nums)
        ans = [0] * n
        left = right = 1
        for i, x in enumerate(nums):
            ans[i] = left
            left *= x
        for i in range(n - 1, -1, -1):
            ans[i] *= right

```

Round 2 · High · Medium

p.323

```

        right *= nums[i]
        return ans

```

• SimplyLeet

```

# Function to compute product of array except self
def productExceptSelf(nums):
    n = len(nums)
    # Initialize the output array with 1s for
    # prefix multiplication
    answer = [1] * n

    # Compute prefix products: for each element
    # answer[i] holds the product of all numbers left of
    # i
    prefix_product = 1
    for i in range(n):
        answer[i] = prefix_product # store prefix
        # product so far
        prefix_product *= nums[i]

```

Round 2 · High · Medium

p.324

```

        prefix_product *= nums[i] # update the
        prefix product

    # Compute suffix products and multiply with
    prefix products in answer array
    suffix_product = 1
    for i in range(n - 1, -1, -1):
        answer[i] *= suffix_product # multiply
        current answer with suffix product
        suffix_product *= nums[i] # update the
        suffix product

    return answer

```

```

# Example usage
nums = [1, 2, 3, 4]
print(productExceptSelf(nums)) # Output: [24, 12,
8, 6]

```

• LC.ca

```

class Solution: # 2 extra arrays
    def productExceptSelf(self, nums):
        if nums is None or len(nums) == 0:
            return nums

        # Product from index=0 to index=i-1
        from_left = [1] * len(nums)
        for i in range(1, len(nums)):
            from_left[i] = nums[i - 1] *
            from_left[i - 1]

```

```

        # Product from index=n-1 to current
        position
        from_right = [1] * len(nums)
        for i in range(len(nums) - 2, -1, -1):
            from_right[i] = nums[i + 1] *
            from_right[i + 1]

        # Calculate result
        result = [0] * len(nums)
        for i in range(len(nums)):
            result[i] = from_left[i] *
            from_right[i]

```

```

    return result

```

```

#####

```

```

class Solution: # follow up, no extra space
    def productExceptSelf(self, nums: List[int])
    -> List[int]:
        n = len(nums)
        ans = [0] * n
        left = right = 1
        for i, v in enumerate(nums):

```

```

        ans[i] = left
        left *= v
    for i in range(n - 1, -1, -1):
        ans[i] *= right
        right *= nums[i]
    return ans

```

```
#####
```

```

class Solution(object):
    def productExceptSelf(self, nums):
        """
        :type nums: List[int]
        :rtype: List[int]
        """
        dp = [1] * len(nums)
        for i in range(1, len(nums)):
            dp[i] = dp[i - 1] * nums[i - 1]

```

```

        prod = 1
        for i in reversed(range(0, len(nums))):
            dp[i] = dp[i] * prod
            prod *= nums[i]
        return dp

```

◆ Quick Review

Key Insights

- The product for each index can be computed by multiplying the product of all numbers to the left and the product of all numbers to the right of that index.
- Division is not allowed, so we must compute the cumulative products explicitly.
- Create two passes: one forward pass for prefix products and one backward pass for suffix products.
- Use the output array to store the cumulative product values, achieving $O(1)$ extra space by not using additional arrays.

Space & Time

Time Complexity: $O(n)$ because we iterate through the array twice.

Space Complexity: $O(1)$ extra space (excluding the output array).

Solution

The approach uses two passes over the input array:

- Initialize the answer array and fill it with prefix products. For each index i , $answer[i]$ contains the product of all elements before i .
- Traverse the array from the end using a running suffix product. Update the answer array by multiplying the current suffix product with the prefix product already stored. This method leverages an output array to store intermediate results and uses a temporary variable for the suffix product, ensuring constant extra space.

Arrays & Hashing

#281 Zigzag Iterator

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Design · Queue · Iterator

Lists: Denny Zhang

Problem

Description

Given two vectors of integers `v1` and `v2`, implement an iterator to return their elements alternately.

Implement the `ZigzagIterator` class:

- `ZigzagIterator(List<int> v1, List<int> v2)` initializes the object with the two vectors `v1` and `v2`.
- boolean `hasNext()` returns true if the iterator still has elements, and false otherwise.
- int `next()` returns the current element of the iterator and moves the iterator to the next element.

Example 1:

Input: `v1 = [1,2]`, `v2 = [3,4,5,6]`

Output: `[1,3,2,4,5,6]`

Explanation: By calling `next` repeatedly until `hasNext` returns false, the order of elements returned by `next` should be: `[1,3,2,4,5,6]`.

Example 2:

Input: `v1 = [1]`, `v2 = []`

Output: `[1]`

Example 3:

Input: `v1 = []`, `v2 = [1]`

Output: `[1]`

Constraints:

- $0 \leq v1.length, v2.length \leq 1000$
- $1 \leq v1.length + v2.length \leq 2000$
- $-231 \leq v1[i], v2[i] \leq 231 - 1$

Follow up: What if you are given `k` vectors? How well can your code be extended to such cases?

Clarification for the follow-up question:

The "Zigzag" order is not clearly defined and is ambiguous for $k > 2$ cases. If "Zigzag" does not look right to you, replace "Zigzag" with "Cyclic".

Follow-up Example:

Input: `v1 = [1,2,3]`, `v2 = [4,5,6,7]`, `v3 = [8,9]`

Output: `[1,4,8,2,5,9,3,6,7]`

Community Solutions (Python)

- WalkCC

```
class ZigzagIterator:
    def __init__(self, v1: list[int], v2:
list[int]):
        def vals():
            for i in itertools.count():
                for v in v1, v2:
                    if i < len(v):
                        yield v[i]
        self.vals = vals()
        self.n = len(v1) + len(v2)
```

```
def next(self):
    self.n -= 1
    return next(self.vals)
```

```
def hasNext(self):
    return self.n > 0
```

• LeetCode

```
class ZigzagIterator:
    def __init__(self, v1: List[int], v2:
List[int]):
        self.cur = 0
        self.size = 2
        self.indexes = [0] * self.size
        self.vectors = [v1, v2]

    def next(self) -> int:
        vector = self.vectors[self.cur]
        index = self.indexes[self.cur]
```

```
        res = vector[index]
        self.indexes[self.cur] = index + 1
        self.cur = (self.cur + 1) % self.size
        return res

    def hasNext(self) -> bool:
        start = self.cur
        while self.indexes[self.cur] ==
len(self.vectors[self.cur]):
            self.cur = (self.cur + 1) % self.size
            if self.cur == start:
```

```
                return False
            return True
```

```
# Your ZigzagIterator object will be instantiated
and called as such:
# i, v = ZigzagIterator(v1, v2), []
# while i.hasNext(): v.append(i.next())
```

• SimplyLeet

```

class ZigzagIterator:
    def __init__(self, v1, v2):
        # Initialize the queue with non-empty
        # vectors paired with starting index 0
        from collections import deque
        self.queue = deque()
        if v1:
            self.queue.append((v1, 0))
        if v2:
            self.queue.append((v2, 0))

    def next(self):
        # Pop from the left
        vec, idx = self.queue.popleft()
        result = vec[idx]
        # If there are more elements, put it back
        # with the next index
        if idx + 1 < len(vec):
            self.queue.append((vec, idx + 1))
        return result

    def hasNext(self):
        return len(self.queue) > 0

# Example usage:
# v1 = [1,2]; v2 = [3,4,5,6]
# iterator = ZigzagIterator(v1, v2)
# while iterator.hasNext():
#     print(iterator.next(), end=" ")

```

Round 2 · High · Medium

p.333

• LC.ca

```

class ZigzagIteratorK:

    # int[0] is index of vector of vectors list,
    # int[1] is start of a vector, int[2] is end
    q = []

    # no map needed, if move 'vectors' a class
    # variable
    vectors = []

    def __init__(self, vectors):
        self.vectors = vectors

    for i in range(len(vectors)):
        self.q.append([i, 0, len(vectors[i])])

    def next(self):
        current = self.q.pop(0)

        i = current[0]
        start = current[1]
        end = current[2]

```

Round 2 · High · Medium

p.334

```

        if start + 1 < end:
            self.q.append([i, start + 1, end])

        return self.vectors[i][start]

    def hasNext(self):
        return len(self.q) != 0

#####

```

```

class ZigzagIterator:
    def __init__(self, v1: List[int], v2:
List[int]):
        self.cur = 0
        self.size = 2
        self.indexes = [0] * self.size
        self.vectors = [v1, v2]

    def next(self) -> int:
        vector = self.vectors[self.cur]
        index = self.indexes[self.cur]

```

```

        res = vector[index]
        self.indexes[self.cur] = index + 1
        self.cur = (self.cur + 1) % self.size
        return res

    def hasNext(self) -> bool:
        start = self.cur
        while self.indexes[self.cur] ==
len(self.vectors[self.cur]):
            self.cur = (self.cur + 1) % self.size
            if self.cur == start:

                return False
            return True

# Your ZigzagIterator object will be instantiated
and called as such:
# i, v = ZigzagIterator(v1, v2), []
# while i.hasNext(): v.append(i.next())

#####

```

```
'''
iter() - builtin function, return an iterator
object
https://docs.python.org/3/library/functions.ht
ml#iter
```

```
>>> a = map(iter, [[1,2],[3,4]])
>>>
>>> b = next(a)
>>> b
<list_iterator object at 0x101630280>
```

```
>>> next(b)
1
>>> next(b)
2
>>> next(b)
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
StopIteration
```

```
>>> b = deque(map(iter, [[1,2],[3,4]]))
>>>
>>> c = b.popleft()
>>> c
<list_iterator object at 0x1016304f0>
>>> c.next()
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
AttributeError: 'list_iterator' object has no
attribute 'next'
>>> next(c)
```

```
1
>>> next(c)
2
>>> next(c)
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
StopIteration
```

```
len(s) - Return the length (the number of items) of
an object
https://docs.python.org/3/library/functions.ht
ml#len
```

```
'''

from collections import deque

class ZigzagIterator(object):

    def __init__(self, v1, v2):
        """
        Initialize your data structure here.
        :type v1: List[int]

        :type v2: List[int]
        """
        self.iters = deque(map(iter, [v1, v2]))
        self.total = sum(map(len, [v1, v2]))

    def next(self):
        """
        :rtype: int
        """
        top = self.iters.popleft()
```

```
try:
    value = top.next()
except StopIteration:
    return self.next()
self.total -= 1
if self.total != 0:
    self.iters.append(top)
return value

def hasNext(self):
```

```
    """
    :rtype: bool
    """
    return self.total > 0

# Your ZigzagIterator object will be instantiated
# and called as such:
# i, v = ZigzagIterator(v1, v2), []
# while i.hasNext(): v.append(i.next())
```

◆ Quick Review

Key Insights

- Use a FIFO (first-in-first-out) structure to maintain the order of vectors which still have remaining elements.
- For each call to next(), select the next element from the current vector and then move that vector to the end of the queue if it still contains elements.

- The approach ensures proper zigzag (or cyclic) ordering among the vectors.
- When extending to k vectors, the same approach works: initialize the queue with each non-empty vector.

Space & Time

Time Complexity: $O(1)$ per `next()` call on average, with an overall $O(n)$ for n total elements.

Space Complexity: $O(k)$ for the queue structure (where k is the number of vectors), plus the storage for the input vectors.

Solution

We can solve the Zigzag Iterator problem by using a queue (or deque) to hold pairs of the vector's current index and the vector itself. For each call to `next()`, we remove the front element from the queue, return the current item from its associated vector, and if the vector has more elements, reinsert the pair (with an updated index) into the back of the queue. This cyclic process naturally interleaves the elements from the provided vectors.

Arrays & Hashing

#307 Range Sum Query – Mutable

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Design · Binary Indexed Tree · Segment Tree

Lists: Denny Zhang

Problem

Description

Given an integer array `nums`, handle multiple queries of the following types:

1. Update the value of an element in `nums`.
2. Calculate the sum of the elements of `nums` between indices `left` and `right` inclusive where `left <= right`.

Implement the `NumArray` class:

- `NumArray(int[] nums)` Initializes the object with the integer array `nums`.
- `void update(int index, int val)` Updates the value of `nums[index]` to be `val`.
- `int sumRange(int left, int right)` Returns the sum of the elements of `nums` between indices `left` and `right` inclusive (i.e. `nums[left] + nums[left + 1] + ... + nums[right]`).

Example 1:

Input

```
["NumArray", "sumRange", "update", "sumRange"]
```

```
[[[1, 3, 5]], [0, 2], [1, 2], [0, 2]]
```

Output

```
[null, 9, null, 8]
```

Explanation

```
NumArray numArray = new NumArray([1, 3, 5]);
```

Constraints:

- $1 \leq \text{nums.length} \leq 3 * 10^4$
- $-100 \leq \text{nums}[i] \leq 100$
- $0 \leq \text{index} < \text{nums.length}$
- $-100 \leq \text{val} \leq 100$
- $0 \leq \text{left} \leq \text{right} < \text{nums.length}$
- At most $3 * 10^4$ calls will be made to update and sumRange.

Community Solutions (Python)

- WalkCC

class FenwickTree:

```
def __init__(self, n: int):
```

```
    self.sums = [0] * (n + 1)
```

```
def add(self, i: int, delta: int) -> None:
```

```
    while i < len(self.sums):
```

```
        self.sums[i] += delta
```

```
        i += FenwickTree.lowbit(i)
```

```
def get(self, i: int) -> int:
```

```
    summ = 0
```

```
    while i > 0:
```

```
        summ += self.sums[i]
```

```
        i -= FenwickTree.lowbit(i)
```

```
    return summ
```

```
@staticmethod
```

```
def lowbit(i: int) -> int:
```

```
    return i & -i
```



```
class NumArray:
    def __init__(self, nums: list[int]):
        self.nums = nums
        self.tree = FenwickTree(len(nums))
        for i, num in enumerate(nums):
            self.tree.add(i + 1, num)

    def update(self, index: int, val: int) -> None:
        self.tree.add(index + 1, val -
self.nums[index])
```

```
        self.nums[index] = val
```

```
    def sumRange(self, left: int, right: int) -> int:
        return self.tree.get(right + 1) -
self.tree.get(left)
```

• LeetCode

```
class BinaryIndexedTree:
    __slots__ = ["n", "c"]

    def __init__(self, n):
        self.n = n
        self.c = [0] * (n + 1)

    def update(self, x: int, delta: int):
        while x <= self.n:
            self.c[x] += delta
```

```
        x += x & -x
```

```
    def query(self, x: int) -> int:
        s = 0
        while x > 0:
            s += self.c[x]
            x -= x & -x
        return s
```

```
class NumArray:
    __slots__ = ["tree"]

    def __init__(self, nums: List[int]):
        self.tree = BinaryIndexedTree(len(nums))
        for i, v in enumerate(nums, 1):
            self.tree.update(i, v)

    def update(self, index: int, val: int) -> None:
        prev = self.sumRange(index, index)
```

```

        self.tree.update(index + 1, val - prev)

    def sumRange(self, left: int, right: int) ->
int:
        return self.tree.query(right + 1) -
self.tree.query(left)

```

Your NumArray object will be instantiated and called as such:

```

# obj = NumArray(nums)
# obj.update(index,val)
# param_2 = obj.sumRange(left,right)

```

```

class Node:
    __slots__ = ["l", "r", "v"]

    def __init__(self):
        self.l = self.r = self.v = 0

```

```

class SegmentTree:
    __slots__ = ["nums", "tr"]

```

```

def __init__(self, nums):
    self.nums = nums
    n = len(nums)
    self.tr = [Node() for _ in range(n << 2)]
    self.build(1, 1, n)

```

```

def build(self, u, l, r):
    self.tr[u].l, self.tr[u].r = l, r
    if l == r:
        self.tr[u].v = self.nums[l - 1]

```

```

        return
        mid = (l + r) >> 1
        self.build(u << 1, l, mid)
        self.build(u << 1 | 1, mid + 1, r)
        self.pushup(u)

```

```

def modify(self, u, x, v):
    if self.tr[u].l == x and self.tr[u].r == x:
        self.tr[u].v = v
    return

```

```

mid = (self.tr[u].l + self.tr[u].r) >> 1
if x <= mid:
    self.modify(u << 1, x, v)
else:
    self.modify(u << 1 | 1, x, v)
self.pushup(u)

def query(self, u, l, r):
    if self.tr[u].l >= l and self.tr[u].r <= r:
        return self.tr[u].v

```

```

mid = (self.tr[u].l + self.tr[u].r) >> 1
result = 0
if l <= mid:
    result += self.query(u << 1, l, r)
if r > mid:
    result += self.query(u << 1 | 1, l, r)
return result

def pushup(self, u):
    self.tr[u].v = self.tr[u << 1].v +
self.tr[u << 1 | 1].v

```

```

class NumArray:
    __slots__ = ["tree"]

    def __init__(self, nums: List[int]):
        self.tree = SegmentTree(nums)

    def update(self, index: int, val: int) -> None:
        self.tree.modify(1, index + 1, val)

    def sumRange(self, left: int, right: int) ->
int:
        return self.tree.query(1, left + 1, right +
1)

# Your NumArray object will be instantiated and
called as such:
# obj = NumArray(nums)
# obj.update(index, val)
# param_2 = obj.sumRange(left, right)

```

• SimplyLeet

```
# Python implementation using a Binary Indexed Tree
(Fenwick Tree)
class NumArray:
    def __init__(self, nums):
        # Store the original array
        self.nums = nums[:]
        self.n = len(nums)
        # Initialize the BIT with 0s, size is (n+1)
        # because BIT is 1-indexed
        self.tree = [0] * (self.n + 1)
        # Build the BIT: update tree for each value
        # in the nums array
        for i, num in enumerate(nums):
```

```
            self._update_tree(i, num)

    def _update_tree(self, index, delta):
        # Update the BIT for index with delta
        # change
        i = index + 1 # Convert to 1-indexed for
        # BIT operations
        while i <= self.n:
            self.tree[i] += delta # Add delta to
            # current tree node
            i += i & -i # Move to the next
            # responsible node

    def update(self, index, val):
```

```
        # Calculate the difference between new
        # value and current value
        delta = val - self.nums[index]
        self.nums[index] = val # Update the
        # original array
        self._update_tree(index, delta) # Update
        # the BIT to reflect this change

    def _prefix_sum(self, index):
        # Compute the prefix sum from start up to
        # and including index
        result = 0
        i = index + 1 # Convert to 1-indexed
        while i:
```

```
            result += self.tree[i]
            i -= i & -i # Move to parent node
        return result

    def sumRange(self, left, right):
        # Return the sum of the range [left, right]
        return self._prefix_sum(right) -
        self._prefix_sum(left - 1) if left > 0 else
        self._prefix_sum(right)

# Example usage:
# numArray = NumArray([1, 3, 5])
```

```
# print(numArray.sumRange(0, 2))
# numArray.update(1, 2)
# print(numArray.sumRange(0, 2))
```

• LC.ca

```
class BinaryIndexedTree:
    def __init__(self, n):
        self.n = n
        self.c = [0] * (n + 1)

    @staticmethod
    def lowbit(x):
        return x & -x

    def update(self, x, delta):
        while x <= self.n:
            self.c[x] += delta
            x += BinaryIndexedTree.lowbit(x)

    def query(self, x):
        s = 0
        while x > 0:
            s += self.c[x]
            x -= BinaryIndexedTree.lowbit(x)
        return s
```

```
class NumArray:
    def __init__(self, nums: List[int]):
        self.tree = BinaryIndexedTree(len(nums))
        for i, v in enumerate(nums, 1):
            self.tree.update(i, v)

    def update(self, index: int, val: int) -> None:
        prev = self.sumRange(index, index)

        self.tree.update(index + 1, val - prev)

    def sumRange(self, left: int, right: int) -> int:
        return self.tree.query(right + 1) - self.tree.query(left)

# Your NumArray object will be instantiated and
# called as such:
# obj = NumArray(nums)
# obj.update(index, val)
# param_2 = obj.sumRange(left, right)
```

```
#####
```

```
# Segment tree node
```

```
class STNode(object):
```

```
    def __init__(self, start, end):
```

```
        self.start = start
```

```
        self.end = end
```

```
        self.total = 0
```

```
        self.left = None
```

```
        self.right = None
```

```
class SegmentedTree(object):
```

```
    def __init__(self, nums, start, end):
```

```
        self.root = self.buildTree(nums, start, end)
```

```
    def buildTree(self, nums, start, end):
```

```
        if start > end:
```

```
            return None
```

```
    if start == end:
```

```
        node = STNode(start, end)
```

```
        node.total = nums[start]
```

```
        return node
```

```
    mid = start + (end - start) / 2
```

```
    root = STNode(start, end)
```

```
    root.left = self.buildTree(nums, start, mid)
```

```
    root.right = self.buildTree(nums, mid + 1, end)
```

```
    root.total = root.left.total +
```

```
    root.right.total
```

```
    return root
```

```
    def updateVal(self, i, val):
```

```
        def updateVal(root, i, val):
```

```
            if root.start == root.end:
```

```
                root.total = val
```

```
                return val
```

```
                mid = root.start + (root.end - root.start) /
```

```
                2
```

```

    if i <= mid:
        updateVal(root.left, i, val)
    else:
        updateVal(root.right, i, val)

    root.total = root.left.total +
root.right.total
    return root.total

return updateVal(self.root, i, val)

def sumRange(self, i, j):
    def rangeSum(root, start, end):
        if root.start == start and root.end == end:
            return root.total

        mid = root.start + (root.end - root.start) /
2
        if j <= mid:
            return rangeSum(root.left, start, end)
        elif i >= mid + 1:
            return rangeSum(root.right, start, end)

```

```

    else:
        return rangeSum(root.left, start, mid) +
rangeSum(root.right, mid + 1, end)

    return rangeSum(self.root, i, j)

class NumArray(object):
    def __init__(self, nums):
        """
        initialize your data structure here.

        :type nums: List[int]
        """
        self.stTree = SegmentedTree(nums, 0, len(nums)
- 1)

    def update(self, i, val):
        """
        :type i: int
        :type val: int
        :rtype: int
        """

```

```

    return self.stTree.updateVal(i, val)

def sumRange(self, i, j):
    """
    sum of elements nums[i..j], inclusive.
    :type i: int
    :type j: int
    :rtype: int
    """
    return self.stTree.sumRange(i, j)

```

```

# Your NumArray object will be instantiated and
called as such:
# numArray = NumArray(nums)
# numArray.sumRange(0, 1)
# numArray.update(1, 10)
# numArray.sumRange(1, 2)

```

◆ Quick Review

Key Insights

- To efficiently support both point updates and range sum queries, a specialized data structure like a Binary Indexed Tree (Fenwick Tree) or a Segment Tree is ideal.
- Both data structures provide update and sum query operations in $O(\log n)$ time.
- The challenge is to maintain a mutable array while enabling fast sum computation over any specified range.

Space & Time

Time Complexity: $O(\log n)$ for both update and sumRange operations

Space Complexity: $O(n)$ to store the additional tree structure

Solution

We can use a Binary Indexed Tree (Fenwick Tree) to solve this problem. The Fenwick Tree allows for efficient prefix sum calculations and point updates, both in $O(\log n)$ time. The approach involves:

- Building the Fenwick Tree using the initial array.
- On an update, computing the difference between the new value and the current value, then propagating the change appropriately in the tree.
- For the sum range query, computing the prefix sum up to right and subtracting the prefix sum up to left-1.

Arrays & Hashing

#454 4Sum II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: CodeSignal

Problem

Description

Given four integer arrays `nums1`, `nums2`, `nums3`, and `nums4` all of length `n`, return the number of tuples (i, j, k, l) such that:

- $0 \leq i, j, k, l < n$
- $nums1[i] + nums2[j] + nums3[k] + nums4[l] == 0$

Example 1:

Input: `nums1 = [1,2]`, `nums2 = [-2,-1]`, `nums3 = [-1,2]`, `nums4 = [0,2]`

Output: 2

Explanation:

The two tuples are:

1. $(0, 0, 0, 1) \rightarrow nums1[0] + nums2[0] + nums3[0] + nums4[1] = 1 + (-2) + (-1) + 2 = 0$
2. $(1, 1, 0, 0) \rightarrow nums1[1] + nums2[1] + nums3[0] + nums4[0] = 2 + (-1) + (-1) + 0 = 0$

Example 2:

Input: `nums1 = [0]`, `nums2 = [0]`, `nums3 = [0]`, `nums4 = [0]`

Output: 1

Constraints:

- $n == nums1.length$
- $n == nums2.length$
- $n == nums3.length$
- $n == nums4.length$
- $1 \leq n \leq 200$
- $-228 \leq nums1[i], nums2[i], nums3[i], nums4[i] \leq 228$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def fourSumCount(self, nums1: list[int], nums2: list[int],
                    nums3: list[int], nums4: list[int]) -> int:
        count = collections.Counter(a + b for a in nums1 for b in nums2)
        return sum(count[-c - d] for c in nums3 for d in nums4)
```

• LeetDoocs

```
class Solution:
    def fourSumCount(
        self, nums1: List[int], nums2: List[int],
        nums3: List[int], nums4: List[int]
    ) -> int:
        cnt = Counter(a + b for a in nums1 for b in
            nums2)
        return sum(cnt[-(c + d)] for c in nums3 for
            d in nums4)
```

- SimplyLeet

```
# Function to count quadruplets that sum to zero
using a hash map approach
def four_sum_count(nums1, nums2, nums3, nums4):
    # Dictionary to store the frequency of sums of
    elements from nums1 and nums2
    sum_count = {}
    for num1 in nums1:
        for num2 in nums2:
            current_sum = num1 + num2
            sum_count[current_sum] =
                sum_count.get(current_sum, 0) + 1

    count = 0
```

```
# For each pair from nums3 and nums4, find the
complement that makes the total sum zero
for num3 in nums3:
    for num4 in nums4:
        complement = -(num3 + num4)
        count += sum_count.get(complement, 0)
    return count
```

```
# Example usage
nums1 = [1,2]
nums2 = [-2,-1]
```

```
nums3 = [-1,2]
nums4 = [0,2]
print(four_sum_count(nums1, nums2, nums3, nums4))
# Expected Output: 2
```

- LC.ca

```

from collections import Counter

class Solution:
    def fourSumCount(
        self, nums1: List[int], nums2: List[int],
        nums3: List[int], nums4: List[int]
    ) -> int:
        cnt = Counter(a + b for a in nums1 for b in
            nums2)
        return sum(cnt[-(c + d)] for c in nums3 for
            d in nums4)

#####

```

```

class Solution:
    def fourSumCount(
        self, nums1: List[int], nums2: List[int],
        nums3: List[int], nums4: List[int]
    ) -> int:
        counter = Counter()
        for a in nums1:
            for b in nums2:
                counter[a + b] += 1
        ans = 0

```

```

        for c in nums3:
            for d in nums4:
                # adding for every time -(c+d)
                ans += counter[-(c + d)]
        return ans

#####

class Solution(object):
    def fourSumCount(self, A, B, C, D):

```

```

        """
        :type A: List[int]
        :type B: List[int]
        :type C: List[int]
        :type D: List[int]
        :rtype: int
        """
        ans = 0
        abDict = {}
        for i in range(len(A)):

```

```

for j in range(len(B)):
    if A[i] + B[j] not in abDict:
        abDict[A[i] + B[j]] = 1
    else:
        abDict[A[i] + B[j]] += 1

for i in range(len(C)):
    for j in range(len(D)):
        if -C[i] - D[j] in abDict:
            ans += abDict[-C[i] - D[j]]

```

```

return ans

```

◆ Quick Review

Key Insights

- Use a hash table to store the frequency of all possible sums from nums1 and nums2.
- For each pair from nums3 and nums4, compute the complement (negative sum) and look it up in the hash table.
- This reduces the time complexity from $O(n^4)$ to $O(n^2)$.

Space & Time

Time Complexity: $O(n^2)$ where n is the length of each array.

Space Complexity: $O(n^2)$ due to the storage of pair sums in the hash table.

Solution

The solution employs a hash table to precompute and store the frequency of sums of all pairs from the first two arrays. Then, by iterating over all pairs from the third and fourth arrays, we calculate the required complement that would result in a total sum of zero. This complement is checked in the hash table, and the frequency (number of valid pairs) is added to the result. This method efficiently finds the number of valid quadruplets without resorting to a brute-force $O(n^4)$ approach.

Arrays & Hashing

#525 Contiguous Array

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Prefix Sum

Lists: Grind 169

Problem

Description

Given a binary array `nums`, return the maximum length of a contiguous subarray with an equal number of 0 and 1.

Example 1:

Input: `nums = [0,1]`
 Output: 2
 Explanation: `[0, 1]` is the longest contiguous subarray with an equal number of 0 and 1.

Example 2:

Input: `nums = [0,1,0]`
 Output: 2
 Explanation: `[0, 1]` (or `[1, 0]`) is a longest contiguous subarray with equal number of 0 and 1.

Example 3:

Input: `nums = [0,1,1,1,1,1,0,0,0]`
 Output: 6
 Explanation: `[1,1,1,0,0,0]` is the longest contiguous subarray with equal number of 0 and 1.

Constraints:

- `1 <= nums.length <= 105`
- `nums[i]` is either 0 or 1.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findMaxLength(self, nums: list[int]) -> int:
        ans = 0
        prefix = 0
        prefixToIndex = {0: -1}

        for i, num in enumerate(nums):
            prefix += 1 if num else -1
            ans = max(ans, i -
prefixToIndex.setdefault(prefix, i))

        return ans
```

• LeetCode

```
class Solution:
    def findMaxLength(self, nums: List[int]) ->
int:
    d = {0: -1}
    ans = s = 0
    for i, x in enumerate(nums):
        s += 1 if x else -1
        if s in d:
            ans = max(ans, i - d[s])
        else:
            d[s] = i

    return ans
```

• SimplyLeet

```
# Python solution for Contiguous Array

def findMaxLength(nums):
    # Dictionary to store the earliest occurrence
    of each prefix sum
    prefix_indices = {0: -1}
    prefix_sum = 0
    max_length = 0

    # Iterate over the array
    for index, num in enumerate(nums):
```

```
        # Convert 0 to -1, keep 1 as is
        if num == 0:
            prefix_sum -= 1
        else:
            prefix_sum += 1

        # If the prefix sum has been seen before, a
        balanced subarray exists
        if prefix_sum in prefix_indices:
            # Calculate the length of the subarray
            current_length = index -
            prefix_indices[prefix_sum]

            max_length = max(max_length,
            current_length)
        else:
            # Store the index for this prefix sum
            if not already seen
            prefix_indices[prefix_sum] = index

        # Return the maximum length of balanced
        subarray
        return max_length

# Example usage:
# print(findMaxLength([0,1,0])) # Expected output:
2
```

• LC.ca

```

class Solution:
    def findMaxLength(self, nums: List[int]) ->
    int:
        s = ans = 0
        mp = {0: -1}
        for i, v in enumerate(nums):
            s += 1 if v == 1 else -1
            if s in mp:
                ans = max(ans, i - mp[s])
            else:
                mp[s] = i

        return ans

```

◆ Quick Review

Key Insights

- Convert 0s to -1 so that a balanced subarray (equal number of 0s and 1s) sums to zero.
- Use a prefix sum to track the cumulative sum as you iterate through the array.
- Use a hash table (or dictionary) to store the first occurrence of each prefix sum.
- When the same prefix sum is encountered again, it indicates that the subarray between these indices is balanced.

Space & Time

Time Complexity: $O(n)$ – We iterate through the array once.

Space Complexity: $O(n)$ – In the worst case, we store an entry for each prefix sum.

Solution

We use a technique where we convert each 0 in the array to -1. Then, we compute a running prefix sum as we iterate through the array. A prefix sum of 0 at any point indicates that the subarray from the beginning to the current index is balanced. For each prefix sum, if we have seen it before, the subarray between the previous index (where the prefix sum was first seen) and the current index sums to 0. We update the maximum length accordingly. We use a hash table to store the first index where each prefix sum occurs.

Arrays & Hashing

#560 Subarray Sum Equals K

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Prefix Sum

Lists: Grind 169

Problem

Description

Given an array of integers `nums` and an integer `k`, return the total number of subarrays whose sum equals to `k`.

A subarray is a contiguous non-empty sequence of elements within an array.

Example 1:

Input: `nums = [1,1,1]`, `k = 2`
 Output: 2

Example 2:

Input: `nums = [1,2,3]`, `k = 3`
 Output: 2

Constraints:

- $1 \leq \text{nums.length} \leq 2 \times 10^4$
- $-1000 \leq \text{nums}[i] \leq 1000$
- $-10^7 \leq k \leq 10^7$

Round 2 · High · Medium

p.375

Community Solutions (Python)

• WalkCC

```
class Solution:
    def subarraySum(self, nums: list[int], k: int) -> int:
        ans = 0
        prefix = 0
        count = collections.Counter({0: 1})

        for num in nums:
            prefix += num
            ans += count[prefix - k]
            count[prefix] += 1

        return ans
```

• LeetDoocs

```
class Solution:
    def subarraySum(self, nums: List[int], k: int) -> int:
        cnt = Counter({0: 1})
        ans = s = 0
        for x in nums:
            s += x
            ans += cnt[s - k]
            cnt[s] += 1
        return ans
```

Round 2 · High · Medium

p.376

- SimplyLeet

```
# Function to calculate the number of subarrays
that sum to k
def subarraySum(nums, k):
    # Dictionary to store frequency of prefix sums
    prefix_counts = {0: 1} # Initialize with prefix
    sum 0
    current_sum = 0 # Running sum initialization
    count = 0 # To count the valid subarrays
    # Iterate over the array
    for num in nums:
        current_sum += num # Update the running
(prefix) sum
        # Check if there's a prefix sum that
differs from current_sum by k

        count += prefix_counts.get(current_sum -
k, 0)
        # Update the frequency count for the
current sum
        prefix_counts[current_sum] =
prefix_counts.get(current_sum, 0) + 1
    return count

# Example usage:
print(subarraySum([1,1,1], 2)) # Expected output:
2
```

- LC.ca

```
class Solution:
    def subarraySum(self, nums: List[int], k: int)
-> int:
        dp = Counter({0: 1})
        ans = s = 0
        for num in nums:
            s += num
            ans += dp[s - k] # not matter if the
same num like [1,1]
            dp[s] += 1
        return ans
```

#####

```
class Solution: # over time limit, not fast enough
    def subarraySum(self, nums: List[int], k: int)
-> int:
        count = 0

        # sum value from 1 to i-th element
        sum_arr = [0] * (len(nums) + 1)

        # cached sum
```

```

        for i in range(1, len(nums) + 1):
            sum_arr[i] = sum_arr[i - 1] + nums[i -
1]

        for start in range(len(nums)):
            for end in range(start + 1, len(nums) +
1):
                if sum_arr[end] - sum_arr[start]
== k:
                    count += 1

        return count

```

◆ Quick Review

Key Insights

- Utilize the prefix sum technique to simplify sum calculation for subarrays.
- Use a hash table (or dictionary/map) to store the frequency of prefix sums encountered.
- For each element, check if (current prefix sum - k) exists in the hash table; if it does, it indicates a valid subarray.
- Initialize the hash table with {0: 1} to manage cases where a subarray's sum is exactly k.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

We solve the problem by maintaining a running sum (prefix sum) and using a hash table to record the number of times each sum has occurred. As we iterate through the array, we update the running sum. For each new sum, we check if (current sum - k) exists in our hash table. If it does, the value associated with (current sum - k) gives the number of subarrays ending at the current index that sum to k. This method efficiently computes the result in a single pass through the array.

Arrays & Hashing

#1010 Pairs of Songs With Total Durations Divisible by 60

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Counting

Lists: CodeSignal

Problem

Description

You are given a list of songs where the i th song has a duration of $\text{time}[i]$ seconds.

Return the number of pairs of songs for which their total duration in seconds is divisible by 60. Formally, we want the number of indices i, j such that $i < j$ with $(\text{time}[i] + \text{time}[j]) \% 60 == 0$.

Example 1:

Input: `time = [30,20,150,100,40]`

Output: 3

Explanation: Three pairs have a total duration divisible by 60:

($\text{time}[0] = 30$, $\text{time}[2] = 150$): total duration 180

($\text{time}[1] = 20$, $\text{time}[3] = 100$): total duration 120

($\text{time}[1] = 20$, $\text{time}[4] = 40$): total duration 60

Example 2:

Input: `time = [60,60,60]`

Output: 3

Explanation: All three pairs have a total duration of 120, which is divisible by 60.

Constraints:

- $1 \leq \text{time.length} \leq 6 * 10^4$
- $1 \leq \text{time}[i] \leq 500$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numPairsDivisibleBy60(self, time: list[int])
-> int:
    ans = 0
    count = [0] * 60

    for t in time:
        t %= 60
        ans += count[(60 - t) % 60]
        count[t] += 1

    return ans
```

• LeetDoocs

```
class Solution:
    def numPairsDivisibleBy60(self, time:
List[int]) -> int:
        cnt = Counter()
        ans = 0
        for x in time:
            x %= 60
            y = (60 - x) % 60
            ans += cnt[y]
            cnt[x] += 1
        return ans
```

• SimplyLeet

```
# Python solution for counting pairs of songs whose
duration sums are divisible by 60
def numPairsDivisibleBy60(time):
    # Frequency array for remainders 0-59
    remainder_freq = [0] * 60
    pair_count = 0

    # Loop over each song duration
    for t in time:
        # Calculate the remainder when dividing by
60
        remainder = t % 60
```

```
        # Calculate the complementary remainder
needed for the sum to be divisible by 60
        complement = (60 - remainder) % 60
        # Increase pair count by frequency of
complement seen so far
        pair_count += remainder_freq[complement]
        # Update the frequency counter for the
current remainder
        remainder_freq[remainder] += 1

    return pair_count
```

Example usage

```
# print(numPairsDivisibleBy60([30,20,150,100,40]))
# Expected output: 3
```

• LC.ca

```
class Solution:
    def numPairsDivisibleBy60(self, time:
List[int]) -> int:
        cnt = Counter(t % 60 for t in time)
        ans = sum(cnt[x] * cnt[60 - x] for x in
range(1, 30))
        ans += cnt[0] * (cnt[0] - 1) // 2
        ans += cnt[30] * (cnt[30] - 1) // 2
        return ans
```

♦ Quick Review

Key Insights

- Each song duration's effect is determined by its remainder when divided by 60.
- For any given module value r (where $0 \leq r < 60$), the complementary module needed to complete 60 is $(60 - r) \% 60$.
- Special handling is needed when r is 0 (or 30, since $30 + 30 = 60$) to avoid double-counting.
- A frequency counter for remainders can be efficiently used to count valid pairs in one pass through the list.

Space & Time

Time Complexity: $O(n)$, where n is the number of songs.

Space Complexity: $O(60) \rightarrow O(1)$, since we only need to keep track of 60 possible remainders.

Solution

We iterate over the list of song durations and calculate the remainder of each duration modulo 60. A frequency array (or hash table) is used to record how many times each remainder has been seen so far. For each song, we look up the frequency of its complementary remainder (i.e., $(60 - \text{current_remainder}) \% 60$) already encountered. This frequency indicates how many valid pairs can be formed with the current song. After counting pairs with the current song, we update the frequency of its remainder for future pairings.

Arrays & Hashing

#1272 Remove Interval

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array

Lists: Premium 98

Problem


Description

A set of real numbers can be represented as the union of several disjoint intervals, where each interval is in the form $[a, b)$. A real number x is in the set if one of its intervals $[a, b)$ contains x (i.e. $a \leq x < b$).

You are given a sorted list of disjoint intervals representing a set of real numbers as described above, where $\text{intervals}[i] = [a_i, b_i]$ represents the interval $[a_i, b_i)$. You are also given another interval toBeRemoved .


Return the set of real numbers with the interval toBeRemoved removed from intervals. In other words, return the set of real numbers such that every x in the set is in intervals but not in toBeRemoved . Your answer should be a sorted list of disjoint intervals as described above.

Example 1:



Input: intervals = [[0,2],[3,4],[5,7]], toBeRemoved = [1,6]
 Output: [[0,1],[6,7]]

Example 2:



Input: intervals = [[0,5]], toBeRemoved = [2,3]
 Output: [[0,2],[3,5]]

Example 3:

Input: intervals = [[-5,-4],[-3,-2],[1,2],[3,5],[8,9]], toBeRemoved = [-1,4]
 Output: [[-5,-4],[-3,-2],[4,5],[8,9]]

Constraints:

- $1 \leq \text{intervals.length} \leq 104$
- $-109 \leq a_i < b_i \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def removeInterval(self, intervals:
list[list[int]],
                        toBeRemoved: list[int]) ->
list[list[int]]:
    ans = []

    for a, b in intervals:
        if a >= toBeRemoved[1] or b <=
toBeRemoved[0]:
            ans.append([a, b])
        else: # a < toBeRemoved[1] and b >
toBeRemoved[0]
            if a < toBeRemoved[0]:
                ans.append([a, toBeRemoved[0]])
            if b > toBeRemoved[1]:
                ans.append([toBeRemoved[1], b])

    return ans
```

• LeetDoocs

```
class Solution:
    def removeInterval(
        self, intervals: List[List[int]],
        toBeRemoved: List[int]
    ) -> List[List[int]]:
        x, y = toBeRemoved
        ans = []
        for a, b in intervals:
            if a >= y or b <= x:
                ans.append([a, b])
            else:
                if a < x:
                    ans.append([a, x])
                if b > y:
                    ans.append([y, b])

        return ans
```

• SimplyLeet

```
def removeInterval(intervals, toBeRemoved):
    # Unpack the removal interval boundaries.
    remove_start, remove_end = toBeRemoved
    result = []

    # Iterate over each interval.
    for interval in intervals:
        start, end = interval

        # If the interval is completely outside
        toBeRemoved, add as is.
```

```

        if end <= remove_start or start >=
remove_end:
            result.append([start, end])
        else:
            # If there is a non-overlapping left
part, add it.
            if start < remove_start:
                result.append([start,
remove_start])
            # If there is a non-overlapping right
part, add it.
            if end > remove_end:
                result.append([remove_end, end])

return result

# Example usage:
intervals = [[0,2],[3,4],[5,7]]
toBeRemoved = [1,6]
print(removeInterval(intervals, toBeRemoved)) #
Expected output: [[0,1],[6,7]]

```

• LC.ca

```

...
>>> a
[1, 2, 3, 4, 5]
>>>
>>> x, y = a
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: too many values to unpack
>>> x,y,z,p,q = a
>>> x

1
>>> y
2
>>> z
3
...
class Solution:
    def removeInterval(
        self, intervals: List[List[int]],
        toBeRemoved: List[int]
    ) -> List[List[int]]:

```



```

x, y = toBeRemoved
ans = []
for a, b in intervals:
    if a >= y or b <= x:
        ans.append([a, b])
    else:
        if a < x:
            ans.append([a, x]) # in this
else block, meaning b>x, so no need to append
[a,min(b,x)]

        if b > y:
            ans.append([y, b])

```

```

        # if x < a < b < y, then delete
(a,b), do nothing and skip here
return ans

```

```

#####

```

```

class Solution:
    def removeInterval(
        self, intervals: List[List[int]],
        toBeRemoved: List[int]
    ) -> List[List[int]]:
        result = []

```

```

        for ind, intv in enumerate(intervals):
            if intv[1] < toBeRemoved[0] or
toBeRemoved[1] < intv[0]:
                result.append(intv)
            else:
                if intv[0] < toBeRemoved[0]:
                    result.append([intv[0],
min(intv[1], toBeRemoved[0])]) # min() is not
needed

                    if intv[1] > toBeRemoved[1]:
                        result.append([max(intv[0],
toBeRemoved[1]), intv[1]])

        return result

```

◆ Quick Review

Key Insights

- Iterate over each interval and check if it overlaps with the toBeRemoved interval.
- If the current interval is completely outside of toBeRemoved, add it to the result as is.
- For overlapping intervals, compute the non-overlapping portions:
 - If the left side of the interval lies before toBeRemoved, include [start, min(end, toBeRemoved.start)].
 - If the right side of the interval lies after toBeRemoved, include [max(start, toBeRemoved.end), end].

Space & Time

Time Complexity: $O(n)$, where n is the number of intervals.

Space Complexity: $O(n)$, for storing the resulting intervals.

Solution

The approach involves iterating through the list of sorted disjoint intervals and checking for overlap with the `toBeRemoved` interval. For each interval:

- If it is completely before or after the removal interval, add it directly to the result.
- If it overlaps with `toBeRemoved`, determine the remaining portion(s): The left side of the interval that does not intersect (if any). The right side of the interval that does not intersect (if any). Use simple conditional checks and list manipulation to form the final set of intervals.

Arrays & Hashing

#1429 First Unique Number

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Design · Queue · Data Stream

Lists: Premium 98

Problem

Description

You have a queue of integers, you need to retrieve the first unique integer in the queue.

Implement the `FirstUnique` class:

- `FirstUnique(int[] nums)` Initializes the object with the numbers in the queue.
- `int showFirstUnique()` returns the value of the first unique integer of the queue, and returns `-1` if there is no such integer.
- `void add(int value)` insert value to the queue.

Example 1:

```

Input:
["FirstUnique","showFirstUnique","add","showFirstUnique",
,"add","showFirstUnique","add","showFirstUnique"]
[[[2,3,5]],[],[5],[],[2],[],[3],[]]
Output:
[null,2,null,2,null,3,null,-1]
Explanation:
FirstUnique firstUnique = new FirstUnique([2,3,5]);
firstUnique.showFirstUnique(); // return 2

```

Example 2:

```

Input:
["FirstUnique","showFirstUnique","add","add","add","add",
,"add","showFirstUnique"]
[[[7,7,7,7,7,7]],[],[7],[3],[3],[7],[17],[]]
Output:
[null,-1,null,null,null,null,null,17]
Explanation:
FirstUnique firstUnique = new
FirstUnique([7,7,7,7,7,7]);
firstUnique.showFirstUnique(); // return -1

```

Example 3:

```

Input:
["FirstUnique","showFirstUnique","add","showFirstUnique",
"]
[[[809]],[],[809],[]]
Output:
[null,809,null,-1]
Explanation:
FirstUnique firstUnique = new FirstUnique([809]);
firstUnique.showFirstUnique(); // return 809

```

Constraints:

- $1 \leq \text{nums.length} \leq 10^5$
- $1 \leq \text{nums}[i] \leq 10^8$
- $1 \leq \text{value} \leq 10^8$
- At most 50000 calls will be made to showFirstUnique and add.

Community Solutions (Python)

• WalkCC

```

class FirstUnique:
    def __init__(self, nums: list[int]):
        self.seen = set()
        self.unique = {}
        for num in nums:
            self.add(num)

    def showFirstUnique(self) -> int:
        return next(iter(self.unique), -1)

```

```
def add(self, value: int) -> None:
    if value not in self.seen:
        self.seen.add(value)
        self.unique[value] = 1
    elif value in self.unique:
        # We have added this value before, and this
        # is the second time we're
        # adding it. So, erase the value from
        # `unique`.
        self.unique.pop(value)
```

• LeetCode

```
class FirstUnique:
    def __init__(self, nums: List[int]):
        self.cnt = Counter(nums)
        self.unique = OrderedDict({v: 1 for v in
nums if self.cnt[v] == 1})

    def showFirstUnique(self) -> int:
        return -1 if not self.unique else next(v
for v in self.unique.keys())

    def add(self, value: int) -> None:
        self.cnt[value] += 1
```

```
if self.cnt[value] == 1:
    self.unique[value] = 1
elif value in self.unique:
    self.unique.pop(value)
```

```
# Your FirstUnique object will be instantiated and
# called as such:
# obj = FirstUnique(nums)
# param_1 = obj.showFirstUnique()
# obj.add(value)
```

• SimplyLeet

```
import collections

class FirstUnique:
    # Initialize the object with the numbers in the
    # queue.
    def __init__(self, nums):
        # Using deque for efficient popleft
        # operations.
        self.queue = collections.deque()
        # Dictionary to store frequency count.
        self.count = collections.defaultdict(int)
        # Process the initial list of numbers.
```

```

        for num in nums:
            self.add(num)

    # Returns the value of the first unique integer
    # of the queue.
    def showFirstUnique(self):
        # Remove numbers from the front if they are
        # no longer unique.
        while self.queue and
self.count[self.queue[0]] > 1:
            self.queue.popleft()
        # Return the first unique number if
        # available.
        return self.queue[0] if self.queue else -1

    # Inserts value to the queue.
    def add(self, value):
        self.count[value] += 1
        # If it is the first occurrence, append it
        # to the queue.
        if self.count[value] == 1:
            self.queue.append(value)

```

• LC.ca

```

class FirstUnique:
    def __init__(self, nums: List[int]):
        self.cnt = Counter(nums)
        self.unique = OrderedDict({v: 1 for v in
nums if self.cnt[v] == 1})

    def showFirstUnique(self) -> int:
        return -1 if not self.unique else next(v
for v in self.unique.keys())

    def add(self, value: int) -> None:
        self.cnt[value] += 1

        if self.cnt[value] == 1:
            self.unique[value] = 1
        elif value in self.unique:
            self.unique.pop(value)

    # Your FirstUnique object will be instantiated and
    # called as such:
    # obj = FirstUnique(nums)
    # param_1 = obj.showFirstUnique()
    # obj.add(value)

```

◆ Quick Review

Key Insights

- Use a queue to maintain the order of potential unique numbers.

- Use a hash table (or map) to count the occurrences of each number.
- When showing the first unique number, remove numbers from the front of the queue if their frequency is more than one.
- All operations (add and showFirstUnique) can be done in $O(1)$ amortized time.

Space & Time

Time Complexity: $O(1)$ amortized for each add and showFirstUnique call.

Space Complexity: $O(n)$, where n is the number of distinct integers encountered.

Solution

We design the FirstUnique class using two main data structures: a queue and a hash map (or dictionary). The queue preserves the order of insertion for unique candidates, while the hash map holds the frequency count for each integer.

When a new number is added:

– Increment its count in the hash map. – If the count is one, append it to the queue.

To retrieve the first unique number:

– While the queue is not empty and the count for the element at the front is greater than one, remove it from the queue. – Return the front element of the queue if it exists, otherwise return -1 .

This approach ensures that we are always up-to-date with the current first unique number.

Arrays & Hashing

#3159 Find Occurrences of an Element in an Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: CodeSignal

Problem

Description

You are given an integer array `nums`, an integer array `queries`, and an integer `x`.

For each `queries[i]`, you need to find the index of the `queries[i]`th occurrence of `x` in the `nums` array. If there are fewer than `queries[i]` occurrences of `x`, the answer should be -1 for that query.

Return an integer array `answer` containing the answers to all queries.

Example 1:

Input: `nums = [1,3,1,7]`, `queries = [1,3,2,4]`, `x = 1`

Output: `[0,-1,2,-1]`

Explanation:

- For the 1st query, the first occurrence of 1 is at index 0.
- For the 2nd query, there are only two occurrences of 1 in `nums`, so the answer is -1 .

- For the 3rd query, the second occurrence of 1 is at index 2.
- For the 4th query, there are only two occurrences of 1 in nums, so the answer is -1.

Example 2:

Input: nums = [1,2,3], queries = [10], x = 5

Output: [-1]

Explanation:

- For the 1st query, 5 doesn't exist in nums, so the answer is -1.

Constraints:

- $1 \leq \text{nums.length}, \text{queries.length} \leq 105$
- $1 \leq \text{queries}[i] \leq 105$
- $1 \leq \text{nums}[i], x \leq 104$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def occurrencesOfElement(
        self,
        nums: list[int],
        queries: list[int],
        x: int,
    ) -> list[int]:
        indices = [i for i, num in enumerate(nums) if
            num == x]
        return [indices[query - 1] if query <=
            len(indices) else -1
                for query in queries]
```

- LeetCode

```
class Solution:
    def occurrencesOfElement(
        self, nums: List[int], queries: List[int],
        x: int
    ) -> List[int]:
        ids = [i for i, v in enumerate(nums) if v
            == x]
        return [ids[i - 1] if i - 1 < len(ids) else
            -1 for i in queries]
```

- SimplyLeet

```
# Function to find occurrences of x at specified
query indices
def find_occurrences(nums, queries, x):
    # Precompute the indices where x occurs in nums
    positions = []
    for index, num in enumerate(nums):
        if num == x:
            positions.append(index)

    # Process each query
    results = []
```

```
    for occurrence in queries:
        # Check if the requested occurrence exists
        if occurrence <= len(positions):
            results.append(positions[occurrence -
1]) # Convert to zero-based index
        else:
            results.append(-1)
    return results
```

```
# Example usage:
nums = [1, 3, 1, 7]
```

```
queries = [1, 3, 2, 4]
x = 1
print(find_occurrences(nums, queries, x)) #
Output: [0, -1, 2, -1]
```

• LC.ca

Round 2 · High · Medium

p.407

```
class Solution:
    def occurrencesOfElement(
        self, nums: List[int], queries: List[int],
x: int
    ) -> List[int]:
        ids = [i for i, v in enumerate(nums) if v
== x]
        return [ids[i - 1] if i - 1 < len(ids) else
-1 for i in queries]
```

♦ Quick Review

Key Insights

- Scan the array once and record every index where `nums[i] == x`.
- Store those indices in order inside a positions array.
- For each query `k`, check whether the `k`-th occurrence exists.
- If it exists, return `positions[k - 1]`.
- Otherwise, return `-1`.
- This lets every query be answered in constant time after the preprocessing pass.

Space & Time

Time Complexity: $O(n + q)$, where n is the length of `nums` and q is the number of queries.

Space Complexity: $O(m)$, where m is the number of occurrences of `x`.

Round 2 · High · Medium

p.408

Solution

We solve the problem by first precomputing all occurrences of `x`. As we iterate through `nums`, every time we see `x`, we append its index to an array called `positions`. Then, for each query value `k`, we check whether `k` is within the bounds of `positions`. If it is, the answer is `positions[k - 1]`; otherwise, the required occurrence does not exist, so the answer is `-1`. This avoids rescanning the array for every query.

String

Round 2 · High · Medium · 5 questions

Round 2 · High · Medium

p.409

Round 2 · High · Medium

p.410

String

#8 String to Integer (atoi)

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String

Lists: Grind 169

Problem

Description

Implement the `myAtoi(string s)` function, which converts a string to a 32-bit signed integer.

The algorithm for `myAtoi(string s)` is as follows:

1. **Whitespace:** Ignore any leading whitespace (" ").
2. **Signedness:** Determine the sign by checking if the next character is '-' or '+', assuming positivity if neither present.
3. **Conversion:** Read the integer by skipping leading zeros until a non-digit character is encountered or the end of the string is reached. If no digits were read, then the result is 0.
4. **Rounding:** If the integer is out of the 32-bit signed integer range $[-2^{31}, 2^{31} - 1]$, then round the integer to remain in the range. Specifically, integers less than -2^{31} should be rounded to -2^{31} , and integers greater than $2^{31} - 1$ should be rounded to $2^{31} - 1$.

Return the integer as the final result.

Example 1:

Input: `s = "42"`

Output: 42

Explanation:

The underlined characters are what is read in and the caret is the current reader position.

Step 1: "42" (no characters read because there is no leading whitespace)

^

Step 2: "42" (no characters read because there is neither a '-' nor '+')

^

Step 3: "42" ("42" is read in)

^

Example 2:

Input: `s = " -042"`

Output: -42

Explanation:

Step 1: " -042" (leading whitespace is read and ignored)

^

Step 2: " -042" ('-' is read, so the result should be negative)

^

Step 3: " -042" ("042" is read in, leading zeros ignored in the result)

^

Example 3:

Input: `s = "1337c0d3"`

Output: 1337

Explanation:

Step 1: "1337c0d3" (no characters read because there is no leading whitespace)

^

Step 2: "1337c0d3" (no characters read because there is neither a '-' nor '+')

^

Step 3: "1337c0d3" ("1337" is read in; reading stops because the next character is a non-digit)

^

Example 4:

Input: `s = "0-1"`

Output: 0

Explanation:

Step 1: "0-1" (no characters read because there is no leading whitespace)

^

Step 2: "0-1" (no characters read because there is neither a '-' nor '+')

^

Step 3: "0-1" ("0" is read in; reading stops because the next character is a non-digit)

^

Example 5:

Input: `s = "words and 987"`

Output: 0

Explanation:

Reading stops at the first non-digit character 'w'.

Constraints:

- $0 \leq s.length \leq 200$
- `s` consists of English letters (lower-case and upper-case), digits (0-9), '-', '+', and '.'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def myAtoi(self, s: str) -> int:
        s = s.strip()
        if not s:
            return 0

        sign = -1 if s[0] == '-' else 1
        if s[0] in {'-', '+'}:
            s = s[1:]
```

```

num = 0

for c in s:
    if not c.isdigit():
        break
    num = num * 10 + int(c)
    if sign * num <= -2**31:
        return -2**31
    if sign * num >= 2**31 - 1:
        return 2**31 - 1

```

```

return sign * num

```

• LeetCode

```

class Solution:
    def myAtoi(self, s: str) -> int:
        if not s:
            return 0
        n = len(s)
        if n == 0:
            return 0
        i = 0
        while s[i] == ' ':
            i += 1

```

```

        if i == n:
            return 0
        sign = -1 if s[i] == '-' else 1
        if s[i] in ['-', '+']:
            i += 1
        res, flag = 0, (2**31 - 1) // 10
        while i < n:
            if not s[i].isdigit():
                break
            c = int(s[i])

```

```

        if res > flag or (res == flag and c >
7):
            return 2**31 - 1 if sign > 0 else
            -(2**31)
            res = res * 10 + c
            i += 1
        return sign * res

```

• SimplyLeet

```

# Python solution for the myAtoi problem.
def myAtoi(s: str) -> int:
    INT_MAX = 2**31 - 1
    INT_MIN = -2**31

    i, n = 0, len(s)
    # Skip leading whitespace.
    while i < n and s[i] == ' ':
        i += 1

```

```
# Handle sign.
sign = 1
if i < n and (s[i] == '+' or s[i] == '-'):
    if s[i] == '-':
        sign = -1
    i += 1

result = 0
# Convert digits and build the integer.
while i < n and s[i].isdigit():
```

```
    digit = int(s[i])
    # Check for overflow or underflow.
    if result > (INT_MAX - digit) // 10:
        return INT_MAX if sign == 1 else
INT_MIN
    result = result * 10 + digit
    i += 1

return sign * result
```

Example usage:

```
print(myAtoi("42"))
print(myAtoi(" -042"))
print(myAtoi("1337c0d3"))
print(myAtoi("0-1"))
print(myAtoi("words and 987"))
```

• LC.ca

Round 2 · High · Medium

p.417

```
class Solution:
    def myAtoi(self, s: str) -> int:
        if not s:
            return 0
        n = len(s)
        if n == 0:
            return 0
        i = 0
        while s[i] == ' ':
            i += 1
```

```
        # only contains blank space
        if i == n:
            return 0
        sign = -1 if s[i] == '-' else 1
        if s[i] in ['-', '+']:
            i += 1
        res, flag = 0, (2**31 - 1) // 10
        while i < n:
            # not a number, exit the loop
            if not s[i].isdigit():
```

Round 2 · High · Medium

p.418

```

        break
    c = int(s[i])
    # if overflows
    if res > flag or (res == flag and c >
7):
        return 2**31 - 1 if sign > 0 else
-(2**31)
    res = res * 10 + c
    i += 1
    return sign * res

```

◆ Quick Review

Key Insights

- Skip leading spaces in the string.
- Detect and handle the optional sign indicator ('+' or '-').
- Read and convert digit characters until a non-digit is encountered.
- Handle cases with no digits or leading zeros.
- Clamp the result to fit within the 32-bit signed integer range.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string.

Space Complexity: $O(1)$ as only a constant amount of extra space is used.

Solution

We iterate through the string to first eliminate any leading whitespace. Next, we check for a sign character to determine if the resulting integer should be negative. We then convert each consecutive digit into an integer, all the while updating the result and checking for potential overflow. If an overflow is detected, we immediately return the appropriate maximum or minimum 32-bit signed integer boundary value. This approach ensures that we strictly follow the conversion rules and efficiently handle edge cases.

String

#249 Group Shifted Strings

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String

Lists: Premium 98

Problem

Description

Perform the following shift operations on a string:

- Right shift: Replace every letter with the successive letter of the English alphabet, where 'z' is replaced by 'a'. For example, "abc" can be right-shifted to "bcd" or "xyz" can be right-shifted to "yza".
- Left shift: Replace every letter with the preceding letter of the English alphabet, where 'a' is replaced by 'z'. For example, "bcd" can be left-shifted to "abc" or "yza" can be left-shifted to "xyz".

We can keep shifting the string in both directions to form an endless shifting sequence.

- For example, shift "abc" to form the sequence: ... <-> "abc" <-> "bcd" <-> ... <-> "xyz" <-> "yza" <-> ... <-> "zab" <-> "abc" <-> ...

You are given an array of strings `strings`, group together all `strings[i]` that belong to the same shifting sequence. You may return the answer in any order.

Example 1:

Input: `strings = ["abc","bcd","acef","xyz","az","ba","a","z"]`

Output: `[["acef"],["a","z"],["abc","bcd","xyz"],["az","ba"]]`

Example 2:

Input: `strings = ["a"]`

Output: `[["a"]]`

Constraints:

- `1 <= strings.length <= 200`
- `1 <= strings[i].length <= 50`
- `strings[i]` consists of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def groupStrings(self, strings: list[str]) -> list[list[str]]:
        keyToStrings = collections.defaultdict(list)

        def getKey(s: str) -> str:
            """
            Returns the key of 's' by pairwise
            calculation of differences.
            e.g. getKey("abc") -> "1,1" because diff(a,
            b) = 1 and diff(b, c) = 1.
            """
            diffs = []
```

```

    for i in range(1, len(s)):
        diff = (ord(s[i]) - ord(s[i - 1]) + 26) %
26        diffs.append(str(diff))

    return ','.join(diffs)

for s in strings:
    keyToStrings[getKey(s)].append(s)

return keyToStrings.values()

```

• LeetCode

```

class Solution:
    def groupStrings(self, strings: List[str]) ->
List[List[str]]:
    g = defaultdict(list)
    for s in strings:
        diff = ord(s[0]) - ord("a")
        t = []
        for c in s:
            c = ord(c) - diff
            if c < ord("a"):
                c += 26

        t.append(chr(c))
        g["".join(t)].append(s)
    return list(g.values())

```

Round 2 · High · Medium

p.423

• SimplyLeet

```

# Python solution for grouping shifted strings

def group_shifted_strings(strings):
    # Helper function to generate a unique key for
    # each string based on differences
    def get_key(s):
        # Single character strings return a fixed
        # key
        if len(s) == 1:
            return "0"
        # Compute differences modulo 26 for each
        # adjacent characters
        key = []

        for i in range(1, len(s)):
            diff = (ord(s[i]) - ord(s[i-1]) + 26) %
26            key.append(str(diff))

        # Use tuple interleaving or join to form a
        # string key
        return ','.join(key)

    groups = {}
    # Group strings by their computed key
    for s in strings:
        key = get_key(s)

```

Round 2 · High · Medium

p.424


```

    if key not in groups:
        groups[key] = []
    groups[key].append(s)
# Return grouped strings as a list of lists
return list(groups.values())

```

Example usage:

```

print(group_shifted_strings(["abc", "bcd", "acef", "xyz", "az", "ba", "a", "z"]))

```

• LC.ca

```

...
>>> a = defaultdict(list)
>>> b = defaultdict([])
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: first argument must be callable or None

```

```

>>> ord('c')
99

```

```

>>> ord('c') - ord('a')
2
>>> ord('d') - 2
98
>>> chr(98)
'b'

```

```

>>> a = {"s":1, "t":2, "x":3}
>>> a

```

```

{'s': 1, 't': 2, 'x': 3}
>>> a.items()
dict_items([('s', 1), ('t', 2), ('x', 3)])
>>> a.values()
dict_values([1, 2, 3])
>>> list(a.values())
[1, 2, 3]
...

```

```

from collections import defaultdict

```

```

class Solution:
    def groupStrings(self, strings: List[str]) -> List[List[str]]:
        mp = defaultdict(list)
        for s in strings:
            t = []
            diff = ord(s[0]) - ord('a')
            for c in s:
                d = ord(c) - diff
                if d < ord('a'):
                    d += 26
                t.append(chr(d))
            k = ''.join(t)
            mp[k].append(s)
        return list(mp.values())

```

◆ Quick Review

Key Insights

- All strings in the same group have the same relative differences between adjacent characters.
- Normalize each string's pattern by converting it to a key representation based on the differences between consecutive letters.
- A single character string forms its own group as there are no differences to compute.

Space & Time

Time Complexity: $O(n * m)$ where n is the number of strings and m is the average length of the strings (to compute each string's key).

Space Complexity: $O(n * m)$ to store the keys and grouping in the hash table.

Solution

The solution involves iterating over each string and generating a unique key that represents the relative differences between adjacent characters. This key is calculated by computing the difference (mod 26) between every consecutive pair in the string. Strings that share the same key can be grouped together since they can be shifted into one another. A hash table (or dictionary) is used to store and group the strings by their computed key. This approach handles edge cases like single-character strings separately.

String

#271 Encode and Decode Strings

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · Design

Lists: Grind 169 | Premium 98

Problem

Description

Design an algorithm to encode a list of strings to a string. The encoded string is then sent over the network and is decoded back to the original list of strings.

Machine 1 (sender) has the function:

```
string encode(vector<string> strs) {
// ... your code
return encoded_string;
}
```

Machine 2 (receiver) has the function:

```
vector<string> decode(string s) {
//... your code
return strs;
}
```

So Machine 1 does:

```
string encoded_string = encode(strs);
```

and Machine 2 does:

```
vector<string> strs2 = decode(encoded_string);
```

strs2 in Machine 2 should be the same as strs in Machine 1.

Implement the encode and decode methods.

You are not allowed to solve the problem using any serialize methods (such as eval).

Example 1:

```
Input: dummy_input = ["Hello","World"]
Output: ["Hello","World"]
Explanation:
Machine 1:
Codec encoder = new Codec();
String msg = encoder.encode(strs);
Machine 1 ---msg---> Machine 2
```

Example 2:

```
Input: dummy_input = [""]
Output: [""]
```

Constraints:

- $1 \leq \text{strs.length} \leq 200$
- $0 \leq \text{strs}[i].\text{length} \leq 200$
- $\text{strs}[i]$ contains any possible characters out of 256 valid ASCII characters.

Follow up: Could you write a generalized algorithm to work on any possible set of characters?

Community Solutions (Python)

• WalkCC

```
class Codec:
    def encode(self, strs: list[str]) -> str:
        """Encodes a list of strings to a single
        string."""
        return ''.join(str(len(s)) + '/' + s for s in
        strs)

    def decode(self, s: str) -> list[str]:
        """Decodes a single string to a list of
        strings."""
        decoded = []

        i = 0

        while i < len(s):
            slash = s.find('/', i)
            length = int(s[i:slash])
            i = slash + length + 1
            decoded.append(s[slash + 1:i])

        return decoded
```

• LeetCode

```
class Codec:
    def encode(self, strs: List[str]) -> str:
        """Encodes a list of strings to a single
        string."""
        ans = []
        for s in strs:
            ans.append('{:4}'.format(len(s)) + s)
        return ''.join(ans)

    def decode(self, s: str) -> List[str]:
        """Decodes a single string to a list of
        strings."""
```

```
        ans = []
        i, n = 0, len(s)
        while i < n:
            size = int(s[i : i + 4])
            i += 4
            ans.append(s[i : i + size])
            i += size
        return ans
```

```
# Your Codec object will be instantiated and called
as such:
# codec = Codec()
# codec.decode(codec.encode(strs))
```

• SimplyLeet

```

class Codec:
    # Encodes a list of strings to a single string.
    def encode(self, strs):
        encoded = []
        # Iterate through each string in the list.
        for s in strs:
            # Append the length of the string, a
            '#' as delimiter, and the string.
            encoded.append(str(len(s)) + '#' + s)
        # Join all parts into one encoded string.
        return ''.join(encoded)

```

```

    # Decodes a single string back to a list of
    strings.
    def decode(self, s):
        decoded = []
        i = 0 # Pointer to traverse the encoded
        string.
        while i < len(s):
            j = i
            # Find the '#' delimiter to determine
            the length.
            while s[j] != '#':
                j += 1

```

```

        # Convert the substring from i to j
        into an integer length.
        length = int(s[i:j])
        # Extract the actual string using the
        obtained length.
        string = s[j+1:j+1+length]
        decoded.append(string)
        # Move the pointer to the start of the
        next encoded segment.
        i = j + 1 + length
        return decoded

```

Example usage:

```

# codec = Codec()
# encoded_str = codec.encode(["Hello", "World"])
# decoded_list = codec.decode(encoded_str)

```

• LC.ca

```

'''
>>> s = 'abcde'
>>> '{:4}'.format(len(s)) + s
' 5abcde'
>>>
>>> '{:4}'.format(333)
' 333'
>>>
>>> '{:4}'.format(1234)
'1234'
'''

```

```
>>> '{:4}'.format(12345)
'12345'

>>> '{0:4}'.format(12345)
'12345'
>>> '{1:4}'.format(12345)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
IndexError: Replacement index 1 out of range for
positional args tuple
>>> '{1:4}'.format(12345, 678910)
'678910'
>>> '{:4}'.format(12345, 678910)
'12345'
>>> '{0:4}'.format(12345, 678910)
'12345'
'''

class Codec: # no splitter
    def encode(self, strs: List[str]) -> str:
        """Encodes a list of strings to a single
        string."""
        ans = []
```

```
        for s in strs:
            ans.append('{:4}'.format(len(s)) + s)
        return ''.join(ans)

    def decode(self, s: str) -> List[str]:
        """Decodes a single string to a list of
        strings."""
        ans = []
        i, n = 0, len(s)
        while i < n:
            size = int(s[i : i + 4]) # so here
            # potential overflow, if larger than int-max, need to
            # clarify assumption with interviewer

            i += 4
            ans.append(s[i : i + size]) # note:
            # exclusive for right index
            i += size
        return ans

# Your Codec object will be instantiated and called
# as such:
# codec = Codec()
# codec.decode(codec.encode(strs))
```

```
#####
```

```
class Codec:
    def encode(self, strs: List[str]) -> str:
        encoded = ""
        for s in strs:
            # cover for input: ["abcd", "####"]
            encoded += str(len(s)) + "#" +
s.replace("#", "$$") + "#"
```

```
return encoded
```

```
def decode(self, s: str) -> List[str]:
    decoded = []
    i = 0
    while i < len(s):
        j = s.find("#", i)
        length = int(s[i:j])
        decoded.append(s[j+1:j+1+length].repla
ce("$$", "#"))
        i = j + 1 + length + 1
```

```
return decoded
```

```
#####
```

```
class Codec: # wrong compared with above, for
input: ["abcd", "####"]
    def encode(self, strs):
        encoded = ""
        for s in strs:
            encoded += str(len(s)) + "#" + s
```

```
return encoded
```

```
def decode(self, s):
    decoded = []
    i = 0
    while i < len(s):
        delimiter = s.index("#", i)
        length = int(s[i:delimiter])
        decoded.append(s[delimiter+1
:delimiter+1+length])
        i = delimiter + 1 + length
```

```
return decoded
```

◆ Quick Review

Key Insights

- Use length-prefix encoding for each string to avoid ambiguity.
- Append a delimiter (e.g., "#") after the length to separate it from the string content.
- This method ensures that even if the original strings contain special characters, including the delimiter, the decoding remains unambiguous.
- The approach handles edge cases, such as empty strings.

Space & Time

Time Complexity: $O(n)$, where n is the total number of characters across all strings.

Space Complexity: $O(n)$, required for the encoded string and the resulting decoded list.

Solution

We encode by iterating over each string and building a new string that contains the length of the string, a delimiter (e.g., "#"), and then the string itself. This way, during decoding, we can safely determine how many characters to extract by first reading until the delimiter to get the length, and then reading that number of characters. This length-prefix approach avoids conflicts with potential characters in the strings.

String

#635 Design Log Storage System

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Design · Ordered Set

Lists: Denny Zhang

Problem

Description

You are given several logs, where each log contains a unique ID and timestamp. Timestamp is a string that has the following format:

Year:Month:Day:Hour:Minute:Second, for example, 2017:01:01:23:59:59. All domains are zero-padded decimal numbers.

Implement the LogSystem class:

- `LogSystem()` Initializes the LogSystem object.
- `void put(int id, string timestamp)` Stores the given log (id, timestamp) in your storage system.
- `int[] retrieve(string start, string end, string granularity)` Returns the IDs of the logs whose timestamps are within the range from start to end inclusive. start and end all have the same format as timestamp, and granularity means how precise the range should be (i.e. to the exact Day, Minute, etc.). For example, start = "2017:01:01:23:59:59", end =

"2017:01:02:23:59:59", and granularity = "Day" means that we need to find the logs within the inclusive range from Jan. 1st 2017 to Jan. 2nd 2017, and the Hour, Minute, and Second for each log entry can be ignored.

Example 1:

```
Input
["LogSystem", "put", "put", "put", "retrieve",
"retrieve"]
[[], [1, "2017:01:01:23:59:59"], [2,
"2017:01:01:22:59:59"], [3, "2016:01:01:00:00:00"],
["2016:01:01:01:01:01", "2017:01:01:23:00:00", "Year"],
["2016:01:01:01:01:01", "2017:01:01:23:00:00", "Hour"]]
Output
[null, null, null, null, [3, 2, 1], [2, 1]]
```

Explanation

```
LogSystem logSystem = new LogSystem();
```

Constraints:

- $1 \leq id \leq 500$
- $2000 \leq Year \leq 2017$
- $1 \leq Month \leq 12$
- $1 \leq Day \leq 31$
- $0 \leq Hour \leq 23$
- $0 \leq Minute, Second \leq 59$
- granularity is one of the values ["Year", "Month", "Day", "Hour", "Minute", "Second"].
- At most 500 calls will be made to put and retrieve.

Community Solutions (Python)

• WalkCC

```
class LogSystem:
    def __init__(self):
        self.granularityToIndices = {'Year': 4,
'Month': 7, 'Day': 10,
'Hour': 13,
'Minute': 16, 'Second': 19}
        self.idAndTimestamps = []

    def put(self, id: int, timestamp: str) -> None:
        self.idAndTimestamps.append((id, timestamp))

    def retrieve(self, start: str, end: str,
granularity: str) -> list[int]:

        index = self.granularityToIndices[granularity]
        s = start[:index]
        e = end[:index]
        return [id for id, timestamp in
self.idAndTimestamps
                if s <= timestamp[:index] <= e]
```

• LeetDoocs

```

class LogSystem:
    def __init__(self):
        self.logs = []
        self.d = {
            "Year": 4,
            "Month": 7,
            "Day": 10,
            "Hour": 13,
            "Minute": 16,
            "Second": 19,
        }

    def put(self, id: int, timestamp: str) -> None:
        self.logs.append((id, timestamp))

    def retrieve(self, start: str, end: str,
granularity: str) -> List[int]:
        i = self.d[granularity]
        return [id for id, ts in self.logs if
start[:i] <= ts[:i] <= end[:i]]

# Your LogSystem object will be instantiated and
called as such:
# obj = LogSystem()
# obj.put(id,timestamp)
# param_2 = obj.retrieve(start,end,granularity)

```

• SimplyLeet

Round 2 · High · Medium

p.443

```

class LogSystem:
    def __init__(self):
        # List to store logs as (timestamp, id)
        tuples
        self.logs = []
        # Mapping granularity to index for the
substring (fixed format:
Year:Month:Day:Hour:Minute:Second)
        self.granularity_to_index = {
            "Year": 4,
            "Month": 7,
            "Day": 10,
            "Hour": 13,
            "Minute": 16,
            "Second": 19
        }

    def put(self, id, timestamp):
        # Append the log as a tuple
        self.logs.append((timestamp, id))

    def retrieve(self, start, end, granularity):
        # Determine the slicing index based on
granularity

```

Round 2 · High · Medium

p.444

```

        idx =
self.granularity_to_index[granularity]
        # Slice the given start and end timestamps
        start_slice = start[:idx]
        end_slice = end[:idx]
        result = []
        # Iterate over all logs and check if the
timestamp slice is within the range
        for ts, log_id in self.logs:
            ts_slice = ts[:idx]
            if start_slice <= ts_slice <=
end_slice:
                result.append(log_id)

```

```

return result

```

```

# Example usage:
# logSystem = LogSystem()
# logSystem.put(1, "2017:01:01:23:59:59")
# logSystem.put(2, "2017:01:01:22:59:59")
# logSystem.put(3, "2016:01:01:00:00:00")
# print(logSystem.retrieve("2016:01:01:01:01:01",
"2017:01:01:23:00:00", "Year")) # Output: [3, 2, 1]
# print(logSystem.retrieve("2016:01:01:01:01:01",
"2017:01:01:23:00:00", "Hour")) # Output: [2, 1]

```

• LC.ca

```

class LogSystem:
    def __init__(self):
        self.logs = []
        self.d = {
            "Year": 4,
            "Month": 7,
            "Day": 10,
            "Hour": 13,
            "Minute": 16,
            "Second": 19,

        }

    def put(self, id: int, timestamp: str) -> None:
        self.logs.append((id, timestamp))

    def retrieve(self, start: str, end: str,
granularity: str) -> List[int]:
        i = self.d[granularity]
        return [id for id, ts in self.logs if
start[:i] <= ts[:i] <= end[:i]]

```

```
# Your LogSystem object will be instantiated and
called as such:
# obj = LogSystem()
# obj.put(id,timestamp)
# param_2 = obj.retrieve(start,end,granularity)
```

```
#####
```

```
# pip install sortedcontainers
from sortedcontainers import SortedDict
```

```
# optimized, sorted
class LogSystem:
    def __init__(self):
        self.logs = SortedDict()

    def put(self, id: int, timestamp: str) -> None:
        if timestamp not in self.logs:
            self.logs[timestamp] = []
        self.logs[timestamp].append(id)
```

```
def retrieve(self, start: str, end: str,
granularity: str) -> List[int]:
    index = {
        "Year": 4,
        "Month": 7,
        "Day": 10,
        "Hour": 13,
        "Minute": 16,
        "Second": 19,
    }
```

```
start = start[:index[granularity]]
end = end[:index[granularity]]

# Adjust end to ensure it captures the
correct range for the granularity
end = end[:index[granularity]] + "9" * (19
- index[granularity])

result = []
# Using irange to efficiently iterate over
the range of timestamps
# By default, both are set to True for
inclusive{}
# eg: sd.irange('banana', 'elderberry',
inclusive=(True, True))
```

```

        for timestamp in self.logs.irange(start,
end):
            if timestamp[:index[granularity]] >
end:
                break
            result.extend(self.logs[timestamp])
        return result

```

◆ Quick Review

Key Insights

- Store logs in a simple data structure such as a list or array.
- Convert the timestamps into strings; due to fixed-length and zero padding, lexicographical string comparisons are valid.
- Use a mapping to determine the slice length for each granularity, e.g. "Year" corresponds to the first 4 characters, "Month" corresponds to the first 7 characters, etc.
- During retrieval, compare the relevant substring of each stored timestamp against the start and end query timestamp slices.
- The inclusive nature of the query simplifies checking conditions.

Space & Time

Time Complexity: $O(N)$ per retrieval query, where N is the number of logs stored (since each log is checked). In the

worst case, all log entries are examined.

Space Complexity: $O(N)$ where N is the number of logs stored.

Solution

We implement a `LogSystem` class that supports putting log entries and retrieving them using a specified granularity.

– Data Structures: Use a list/array to store log entries as tuples (or objects) containing the timestamp and log ID. A dictionary/map is used to store the mapping from granularity string to the substring index (e.g., "Year" => 4, "Month" => 7, etc.). – For the `put()` method, simply append the new log entry into the list. – For the `retrieve()` method, compute the substring length based on the provided granularity. Then iterate over all stored logs and compare the sliced timestamp against the sliced start and end query values. If it falls within the bounds, include the log ID in the result. – Since the timestamps have fixed width with leading zeros, slicing and lexicographical comparisons are both valid and simple.

String

#1138 Alphabet Board Path

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String


Lists: Denny Zhang

Problem

Description

On an alphabet board, we start at position (0, 0), corresponding to character board[0][0].

Here, board = ["abcde", "fghij", "klmno", "pqrst", "uvwxy", "z"], as shown in the diagram below.



We may make the following moves:

- 'U' moves our position up one row, if the position exists on the board;
- 'D' moves our position down one row, if the position exists on the board;
- 'L' moves our position left one column, if the position exists on the board;
- 'R' moves our position right one column, if the position exists on the board;
- '!' adds the character board[r][c] at our current position (r, c) to the answer.

(Here, the only positions that exist on the board are positions with letters on them.)

Return a sequence of moves that makes our answer equal to target in the minimum number of moves. You may return any path that does so.

Example 1:

Input: target = "leet"

Output: "DDR!UURRR!!DDD!"

Example 2:

Input: target = "code"

Output: "RR!DDRR!UUL!R!"

Constraints:

- $1 \leq \text{target.length} \leq 100$
- target consists only of English lowercase letters.

Community Solutions (Python)

- LeetDoocs

```
class Solution:
    def alphabetBoardPath(self, target: str) -> str:
        i = j = 0
        ans = []
        for c in target:
            v = ord(c) - ord("a")
            x, y = v // 5, v % 5
            while j > y:
                j -= 1
                ans.append("L")

            while i > x:
                i -= 1
                ans.append("U")
            while j < y:
                j += 1
                ans.append("R")
            while i < x:
                i += 1
                ans.append("D")
            ans.append("!")

        return "".join(ans)
```

- SimplyLeet

```

# Python solution with clear variable names and
inline comments.
class Solution:
    def alphabetBoardPath(self, target: str) ->
str:
    # Starting position at the top-left corner
    (0,0)
    current_row, current_col = 0, 0
    result = []

    # Helper function to convert letter to its
    (row, col) coordinates
    def getPosition(letter):
        index = ord(letter) - ord('a')

        return index // 5, index % 5

    for char in target:
        target_row, target_col =
getPosition(char)
        # If target character is 'z', move
        horizontally first then vertically
        if char == 'z':
            # Move left if needed
            while current_col > target_col:
                result.append("L")
                current_col -= 1

```

```

        # Move right if needed
        while current_col < target_col:
            result.append("R")
            current_col += 1
        # Then move down if needed
        while current_row < target_row:
            result.append("D")
            current_row += 1
        # Then move up if needed
        while current_row > target_row:
            result.append("U")
            current_row -= 1
        else:
            # For all other characters, move
            vertically first
            while current_row > target_row:
                result.append("U")
                current_row -= 1
            while current_row < target_row:
                result.append("D")
                current_row += 1

```



```

        # Then move horizontally
        while current_col > target_col:
            result.append("L")
            current_col -= 1
        while current_col < target_col:
            result.append("R")
            current_col += 1
        # Append pick-up action for the letter
        result.append("!")
    return "".join(result)

```

• LC.ca

```

class Solution:
    def alphabetBoardPath(self, target: str) -> str:
        i = j = 0
        ans = []
        for c in target:
            v = ord(c) - ord("a")
            x, y = v // 5, v % 5
            while j > y:
                j -= 1
                ans.append("L")

```

```

        while i > x:
            i -= 1
            ans.append("U")
        while j < y:
            j += 1
            ans.append("R")
        while i < x:
            i += 1
            ans.append("D")
        ans.append("!")

```

```

    return "".join(ans)

```

◆ Quick Review

Key Insights

- Determine the row and column for each letter (using integer division and modulus).
- Because "z" is in a special row that only contains one column, the move order is critical to avoid invalid positions.
- For target letter "z", perform horizontal moves before vertical moves; for all other letters, do vertical moves before horizontal moves.
- Build the answer step-by-step by moving from your current position to each target letter.

Space & Time

Time Complexity: $O(N)$ where N is the length of target (each move is processed in constant time)

Space Complexity: $O(N)$ for the resulting string of moves

Solution

The approach calculates the destination coordinates for each character. For each letter in the target:

- Convert the letter to its board coordinates (row and col). For example, for letter c, $\text{row} = (\text{ord}(\text{c}) - \text{ord}(\text{'a'})) // 5$ and $\text{col} = (\text{ord}(\text{c}) - \text{ord}(\text{'a'})) \% 5$. Note that "z" (the 26th letter) always maps to (5, 0).
- To move from the current position (curRow, curCol) to the target's (targetRow, targetCol), decide on the order of moves: If the target letter is "z", first move horizontally (left or right) then vertically (up or down) to avoid stepping into invalid regions from the extra row. For any other letter, move vertically (up or down) first, then horizontally (left or right).
- Append "!" to the moves after reaching each letter.
- Update the current position to the letter's coordinates.

This method ensures you only traverse valid cells on the board and minimizes moves by taking the optimal coordinate difference each step.

Two Pointers

Round 2 · High · Medium · 14 questions

Round 2 · High · Medium

p.459

Round 2 · High · Medium

p.460

Two Pointers

#11 Container With Most Water

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Greedy

Lists: Grind 169

Problem

Description


You are given an integer array `height` of length `n`. There are `n` vertical lines drawn such that the two endpoints of the `i`th line are $(i, 0)$ and $(i, \text{height}[i])$.

Find two lines that together with the x-axis form a container, such that the container contains the most water.

Return the maximum amount of water a container can store.

Notice that you may not slant the container.

Example 1:



Input: `height = [1,8,6,2,5,4,8,3,7]`

Output: 49

Explanation: The above vertical lines are represented by array `[1,8,6,2,5,4,8,3,7]`. In this case, the max area of water (blue section) the container can contain is 49.

Example 2:

Input: `height = [1,1]`

Output: 1

Constraints:

- `n == height.length`
- `2 <= n <= 105`
- `0 <= height[i] <= 104`

Community Solutions (Python)

- WalkCC

```

class Solution:
    def maxArea(self, height: list[int]) -> int:
        ans = 0
        l = 0
        r = len(height) - 1

        while l < r:
            minHeight = min(height[l], height[r])
            ans = max(ans, minHeight * (r - l))
            if height[l] < height[r]:
                l += 1
            else:
                r -= 1

        return ans

```

- LeetCode

```

class Solution:
    def maxArea(self, height: List[int]) -> int:
        l, r = 0, len(height) - 1
        ans = 0
        while l < r:
            t = min(height[l], height[r]) * (r - l)
            ans = max(ans, t)
            if height[l] < height[r]:
                l += 1
            else:

```

```

        r -= 1
    return ans

```

- SimplyLeet

```

# Function to calculate the maximum water container area
def maxArea(height):
    left, right = 0, len(height) - 1 # Initialize two pointers
    max_water = 0 # Variable to store the maximum area

    # Loop until pointers meet
    while left < right:
        # Calculate current area based on smaller height and distance between pointers
        current_area = min(height[left], height[right]) * (right - left)
        max_water = max(max_water, current_area) # Update maximum area if current is larger

```

```

        # Move the pointer corresponding to the
        smaller height
        if height[left] < height[right]:
            left += 1
        else:
            right -= 1
    return max_water

```

```

# Test the function with an example
if __name__ == "__main__":

```

```

    input_heights = [1,8,6,2,5,4,8,3,7]
    print(maxArea(input_heights)) # Expected
    output: 49

```

• LC.ca

```

class Solution:
    def maxArea(self, height: List[int]) -> int:
        i, j = 0, len(height) - 1
        ans = 0
        while i < j:
            t = (j - i) * min(height[i], height[j])
            ans = max(ans, t)
            if height[i] < height[j]:
                i += 1
            else:
                j -= 1
        return ans

```

Round 2 · High · Medium

p.465

◆ Quick Review

Key Insights

- The container's area is determined by the distance between two lines multiplied by the shorter line's height.
- A brute-force approach checking all possible pairs would result in $O(n^2)$ time complexity, which is inefficient for large inputs.
- A two-pointer approach allows us to solve the problem in $O(n)$ time by initializing pointers at both ends and gradually moving them inward.
- By always moving the pointer corresponding to the shorter line, we are more likely to find a taller line that could potentially form a container with a larger capacity.

Space & Time

Time Complexity: $O(n)$ – The algorithm makes a single pass through the array with two pointers.

Space Complexity: $O(1)$ – Only a few extra variables are used, regardless of the input size.

Solution

The optimal approach uses a two-pointer strategy:

- Start with two pointers at each end of the array.
- Calculate the area formed between the two lines at these pointers.
- Update the maximum area if the calculated value is larger.
- Move the pointer pointing to the shorter line inward, as this is the only possible way to potentially increase the height of the shorter line and thus the area.

Round 2 · High · Medium

p.466

Continue until the two pointers meet.

This method effectively reduces the number of comparisons and ensures that every possible pair that can yield a larger area is considered.

Two Pointers

#15 3Sum

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Sorting

Lists: Grind 169

Problem

Description

Given an integer array `nums`, return all the triplets `[nums[i], nums[j], nums[k]]` such that $i \neq j$, $i \neq k$, and $j \neq k$, and $nums[i] + nums[j] + nums[k] == 0$.

Notice that the solution set must not contain duplicate triplets.

Example 1:

Input: `nums = [-1,0,1,2,-1,-4]`

Output: `[[-1,-1,2],[-1,0,1]]`

Explanation:

`nums[0] + nums[1] + nums[2] = (-1) + 0 + 1 = 0.`

`nums[1] + nums[2] + nums[4] = 0 + 1 + (-1) = 0.`

`nums[0] + nums[3] + nums[4] = (-1) + 2 + (-1) = 0.`

The distinct triplets are `[-1,0,1]` and `[-1,-1,2]`.

Notice that the order of the output and the order of the triplets does not matter.

Example 2:

Input: nums = [0,1,1]

Output: []

Explanation: The only possible triplet does not sum up to 0.

Example 3:

Input: nums = [0,0,0]

Output: [[0,0,0]]

Explanation: The only possible triplet sums up to 0.

Constraints:

- $3 \leq \text{nums.length} \leq 3000$
- $-105 \leq \text{nums}[i] \leq 105$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def threeSum(self, nums: list[int]) ->
list[list[int]]:
    if len(nums) < 3:
        return []

    ans = []

    nums.sort()

    for i in range(len(nums) - 2):
```

```
        if i > 0 and nums[i] == nums[i - 1]:
            continue
        # Choose nums[i] as the first number in the
        triplet, then search the
        # remaining numbers in [i + 1, n - 1].
        l = i + 1
        r = len(nums) - 1
        while l < r:
            summ = nums[i] + nums[l] + nums[r]
            if summ == 0:
                ans.append((nums[i], nums[l], nums[r]))

            l += 1
            r -= 1
            while nums[l] == nums[l - 1] and l < r:
                l += 1
            while nums[r] == nums[r + 1] and l < r:
                r -= 1
            elif summ < 0:
                l += 1
            else:
                r -= 1

    return ans
```

• LeetCode

```

class Solution:
    def threeSum(self, nums: List[int]) ->
List[List[int]]:
    nums.sort()
    n = len(nums)
    ans = []
    for i in range(n - 2):
        if nums[i] > 0:
            break
        if i and nums[i] == nums[i - 1]:
            continue

        j, k = i + 1, n - 1
        while j < k:
            x = nums[i] + nums[j] + nums[k]
            if x < 0:
                j += 1
            elif x > 0:
                k -= 1
            else:
                ans.append([nums[i], nums[j],
nums[k]])
                j, k = j + 1, k - 1

```

```

                                while j < k and nums[j] ==
nums[j - 1]:
                                    j += 1
                                while j < k and nums[k] ==
nums[k + 1]:
                                    k -= 1

                                return ans

```

• SimplyLeet

```

# Function to find all unique triplets that sum to
zero.
def threeSum(nums):
    # Sort the array to facilitate finding
duplicates and using two pointers.
    nums.sort()
    result = []
    n = len(nums)

    # Iterate through the array, fixing one element
at a time.
    for i in range(n - 2):
        # Skip duplicate elements for the first
element.

```



```

    if i > 0 and nums[i] == nums[i - 1]:
        continue

    left, right = i + 1, n - 1
    while left < right:
        current_sum = nums[i] + nums[left] +
nums[right]
        # Check if the triplet adds up to zero.
        if current_sum == 0:
            result.append([nums[i],
nums[left], nums[right]])
            # Skip duplicates for the second
element.

```

```

        while left < right and nums[left]
== nums[left + 1]:
            left += 1
            # Skip duplicates for the third
element.
        while left < right and nums[right]
== nums[right - 1]:
            right -= 1
            left += 1
            right -= 1
        elif current_sum < 0:
            left += 1
        else:

```

```

        right -= 1

    return result

# Example usage:
#print(threeSum([-1,0,1,2,-1,-4]))

```

• LC.ca

```

class Solution:
    def threeSum(self, nums: List[int]) ->
List[List[int]]:
        nums.sort()
        n = len(nums)
        ans = []
        for i in range(n - 2):
            if nums[i] > 0:
                break
            if i and nums[i] == nums[i - 1]:
                continue

```

```

        j, k = i + 1, n - 1
        while j < k:
            x = nums[i] + nums[j] + nums[k]
            if x < 0:
                j += 1
            elif x > 0:
                k -= 1
            else:
                ans.append([nums[i], nums[j],
nums[k]])
                j, k = j + 1, k - 1

        while j < k and nums[j] ==
nums[j - 1]:
            j += 1
        while j < k and nums[k] ==
nums[k + 1]:
            k -= 1

    return ans

```

◆ Quick Review

Key Insights

- Sort the array to efficiently skip duplicates and facilitate the two pointers approach.
- Use a fixed pointer for the first element, and then use two additional pointers (left and right) to find pairs that sum to the negation of the fixed element.
- Skip duplicate elements to ensure unique triplets.

Space & Time

Time Complexity: $O(n^2)$ where n is the number of elements in the array.

Space Complexity: $O(1)$ extra space (ignoring the space required for the output).

Solution

We first sort the array to bring duplicates together and allow the two-pointer technique to work efficiently. Then, iterate over the array and for each element, use two pointers (one at the beginning right after the current element and one at the end of the array) to find a pair that sums with the current element to zero. Skip duplicates for the current element and within the inner loop to avoid duplicate triplets. This ensures that each valid triplet is unique and the overall runtime remains efficient.

Two Pointers

#16 3Sum Closest

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Sorting

Lists: Grind 169

Problem

Description

Given an integer array `nums` of length `n` and an integer `target`, find three integers at distinct indices in `nums` such that the sum is closest to `target`.

Return the sum of the three integers.

You may assume that each input would have exactly one solution.

Example 1:

Input: `nums = [-1,2,1,-4]`, `target = 1`

Output: 2

Explanation: The sum that is closest to the target is 2. ($-1 + 2 + 1 = 2$).

Example 2:

Input: `nums = [0,0,0]`, `target = 1`

Output: 0

Explanation: The sum that is closest to the target is 0. ($0 + 0 + 0 = 0$).

Round 2 · High · Medium

p.477

Constraints:

- $3 \leq \text{nums.length} \leq 500$
- $-1000 \leq \text{nums}[i] \leq 1000$
- $-104 \leq \text{target} \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def threeSumClosest(self, nums: list[int],
                        target: int) -> int:
        ans = nums[0] + nums[1] + nums[2]

        nums.sort()

        for i in range(len(nums) - 2):
            if i > 0 and nums[i] == nums[i - 1]:
                continue
            # Choose nums[i] as the first number in the
            triplet, then search the
```

Round 2 · High · Medium

p.478

```
# remaining numbers in [i + 1, n - 1].
l = i + 1
r = len(nums) - 1
while l < r:
    summ = nums[i] + nums[l] + nums[r]
    if summ == target:
        return summ
    if abs(summ - target) < abs(ans - target):
        ans = summ
    if summ < target:
```

```
        l += 1
    else:
        r -= 1
```

```
return ans
```

• LeetCode

```
class Solution:
    def threeSumClosest(self, nums: List[int],
target: int) -> int:
    nums.sort()
    n = len(nums)
    ans = inf
    for i, v in enumerate(nums):
        j, k = i + 1, n - 1
        while j < k:
            t = v + nums[j] + nums[k]
            if t == target:
```

```
                return t
            if abs(t - target) < abs(ans -
target):
                ans = t
            if t > target:
                k -= 1
            else:
                j += 1
        return ans
```

• SimplyLeet

```
# Python solution for 3Sum Closest
def threeSumClosest(nums, target):
    # Sort the array to allow two pointer technique
    nums.sort()
    # Initialize best_sum with the sum of the first
three elements
    best_sum = sum(nums[:3])

    # Iterate through each number as the fixed
element
    for i in range(len(nums) - 2):
        left, right = i + 1, len(nums) - 1
```

```

    # Use two pointers to check pairs
    while left < right:
        current_sum = nums[i] + nums[left] +
nums[right]
        # Update best_sum if current_sum is
closer to target
        if abs(target - current_sum) <
abs(target - best_sum):
            best_sum = current_sum
        # Move left pointer if current_sum is
less than target to increase sum
        if current_sum < target:
            left += 1
        # Move right pointer if current_sum is
greater than target to decrease sum

    elif current_sum > target:
        right -= 1
    else:
        # If the current_sum exactly
equals target, return the target sum
        return current_sum
    # Return the closest sum found
    return best_sum

```

```

# Example usage:
# print(threeSumClosest([-1,2,1,-4], 1))

```

• LC.ca

Round 2 · High · Medium

p.481

```

class Solution:
    def threeSumClosest(self, nums: List[int],
target: int) -> int:
        nums.sort()
        n = len(nums)
        ans = inf
        for i, v in enumerate(nums):
            j, k = i + 1, n - 1
            while j < k:
                t = v + nums[j] + nums[k]
                if t == target:

                    return t
                if abs(t - target) < abs(ans -
target):

                    ans = t
                    if t > target:
                        k -= 1
                    else:
                        j += 1
        return ans

```

◆ Quick Review

Key Insights

- Sort the array to make it easier to use the two pointers technique.
- Fix one element and then use two pointers (left and right) to find the best pair that, along with the fixed element, produces a sum closest to the target.

Round 2 · High · Medium

p.482

- Update the best sum whenever a closer sum is found.
- The sorted order helps eliminate unnecessary iterations and allows efficient movement of pointers.

Space & Time

Time Complexity: $O(n^2)$

Space Complexity: $O(\log n)$ to $O(n)$ depending on the sorting algorithm used (typically in-place sort, so extra space might be minimal)

Solution

The solution starts by sorting the array. For each index in the array (considered as the first element of the triple), we use two pointers: one starting just after the fixed element and the other at the end of the array. We calculate the sum of the fixed element and the two pointers. If this sum is closer to the target than the current best, we update our best sum. Depending on whether the current sum is less than or greater than the target, we adjust the left or right pointer to try and get closer to the target. This two-pointer method leverages the sorted order of the array for efficient summing.

Two Pointers

#31 Next Permutation

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers

Lists: Grind 169

Problem

Description

A permutation of an array of integers is an arrangement of its members into a sequence or linear order.

- For example, for $arr = [1,2,3]$, the following are all the permutations of arr : $[1,2,3]$, $[1,3,2]$, $[2, 1, 3]$, $[2, 3, 1]$, $[3,1,2]$, $[3,2,1]$.

The next permutation of an array of integers is the next lexicographically greater permutation of its integer. More formally, if all the permutations of the array are sorted in one container according to their lexicographical order, then the next permutation of that array is the permutation that follows it in the sorted container. If such arrangement is not possible, the array must be rearranged as the lowest possible order (i.e., sorted in ascending order).

- For example, the next permutation of $arr = [1,2,3]$ is $[1,3,2]$.

- Similarly, the next permutation of $\text{arr} = [2,3,1]$ is $[3,1,2]$.
- While the next permutation of $\text{arr} = [3,2,1]$ is $[1,2,3]$ because $[3,2,1]$ does not have a lexicographical larger rearrangement.

Given an array of integers `nums`, find the next permutation of `nums`.

The replacement must be in place and use only constant extra memory.

Example 1:

Input: `nums = [1,2,3]`
Output: `[1,3,2]`

Example 2:

Input: `nums = [3,2,1]`
Output: `[1,2,3]`

Example 3:

Input: `nums = [1,1,5]`
Output: `[1,5,1]`

Constraints:

- $1 \leq \text{nums.length} \leq 100$
- $0 \leq \text{nums}[i] \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def nextPermutation(self, nums: list[int]) ->
    None:
        n = len(nums)

        # From back to front, find the first number <
        nums[i + 1].
        i = n - 2
        while i >= 0:
            if nums[i] < nums[i + 1]:
                break
            i -= 1
```

```
        # From back to front, find the first number >
        nums[i], swap it with nums[i].
        if i >= 0:
            for j in range(n - 1, i, -1):
                if nums[j] > nums[i]:
                    nums[i], nums[j] = nums[j], nums[i]
                    break

        def reverse(nums: list[int], l: int, r: int) ->
        None:
            while l < r:
```

```

    nums[l], nums[r] = nums[r], nums[l]
    l += 1
    r -= 1

# Reverse nums[i + 1..n - 1].
reverse(nums, i + 1, len(nums) - 1)

```

- LeetCode

```

class Solution:
    def nextPermutation(self, nums: List[int]) ->
    None:
        n = len(nums)
        i = next((i for i in range(n - 2, -1, -1)
        if nums[i] < nums[i + 1]), -1)
        if ~i:
            j = next((j for j in range(n - 1, i,
            -1) if nums[j] > nums[i]))
            nums[i], nums[j] = nums[j], nums[i]
            nums[i + 1 :] = nums[i + 1 :][::-1]

```

- SimplyLeet

```

def nextPermutation(nums):
    # Length of the array
    n = len(nums)
    # Find the first index from the right where the
    current number is less than the next number
    i = n - 2
    while i >= 0 and nums[i] >= nums[i + 1]:
        i -= 1
    # If such an index is found, find the number
    just larger than nums[i] from the right side to
    swap
    if i >= 0:
        j = n - 1

```

```

        while nums[j] <= nums[i]:
            j -= 1
        nums[i], nums[j] = nums[j], nums[i] # Swap
        the two numbers
        # Reverse the subarray from i+1 to the end to
        get the smallest sequence
        left, right = i + 1, n - 1
        while left < right:
            nums[left], nums[right] = nums[right],
            nums[left]
            left += 1
            right -= 1

```



```
# Example usage:
nums = [1, 2, 3]
nextPermutation(nums)
print(nums) # Output: [1, 3, 2]
```

• LC.ca

```
...
>>> i = 3
>>> ~i
-4
>>> bool(~i)
True
#####
```

```
>>> j = -1
>>> ~j
```

```
0
>>> bool(~j)
False
#####

>>> a = (i for i in range(10, -1, -1) if i < 6)
>>> a
<generator object <genexpr> at 0x10a17eeb0>
>>> next(a)
5
```

```
>>>
>>>
>>> b = (i for i in range(10, -1, -1) if i < 0)
>>> next(b)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
StopIteration
>>> next(b, -1)
-1
>>>
```

```
...
class Solution:
    def nextPermutation(self, nums: List[int]) ->
None:
    """
    Do not return anything, modify nums
    in-place instead.
    """
    n = len(nums)
    # next(func, -1) => default value -1
    i = next((i for i in range(n - 2, -1, -1)
if nums[i] < nums[i + 1]), -1)
    if i != -1:
```

```

        j = next((j for j in range(n - 1, i,
-1) if nums[j] > nums[i]))
        nums[i], nums[j] = nums[j], nums[i]
        nums[i + 1:] = nums[i + 1:][:-1]
        # wrong reverse: nums[i + 1:] = nums[i +
1::-1]

```

```
#####
```

```

...
>>> a=[1,2,3]
>>> reversed(a)

```

```

<list_reverseiterator object at 0x108458be0>
>>> list(reversed(a))
[3, 2, 1]

```

```

# but, use reversed(a) to directly assign values is
ok
>>> b=[4,5,6,7,8]
>>> b[:3] = reversed(a)
>>> b
[3, 2, 1, 7, 8]
...

```

```

class Solution:
    def nextPermutation(self, nums: List[int]) ->
None:
        if not nums:
            return

        # next
        for i in range(len(nums)-2, -1, -1): # 3,
4, 5, 2, 1
            if nums[i] < nums[i+1]: # ==> 4
                j = next(j for j in
range(len(nums)-1, i, -1) if nums[j] > nums[i]) #
nums[i+1] ==> 5
                nums[i], nums[j] = nums[j],
nums[i] # 3,5,4,2,1

```

```

        # reverse
        #
        nums[i+1:] = reversed(nums[i+1:])
        # [4,2,1] reverse to [1,2,4] ==> 3, 5, 1, 2, 4
        return

```

```

        nums.reverse() # for return

```

```
#####
```

```

class Solution:
    def nextPermutation(self, nums: List[int]) ->
None:
        n = len(nums)
        to_left = [i for i in range(n - 2, -1, -1)
if nums[i] < nums[i + 1]]

        if not to_left:
            nums = nums[::-1]
            return

        i = max(to_left)
        if i >= 0:
            to_right = [ j for j in range(n - 1, i,
-1) if nums[j] > nums[i] ]
            j = to_right[0]
            nums[i], nums[j] = nums[j], nums[i]
            nums[i + 1:] = nums[i + 1::-1]

```

◆ Quick Review

Key Insights

- The problem is solved by identifying the first pair from the right where the left element is less than the right element.
- Once this pivot is found, find the rightmost element that is greater than this pivot.
- Swap these two elements.

- Reverse the subarray that comes after the pivot to ensure the next permutation is the next lexicographically larger sequence.
- If no such pivot exists, the array is in descending order and should be reversed to get the smallest permutation.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We start by scanning the array from the right to find the first index (i) such that `nums[i]` is less than `nums[i+1]`. This index 'i' represents the pivot point where the next permutation can be formed. If we find such an index, we then scan from the right again to locate the first number that is greater than `nums[i]` (let this index be j), and swap these two numbers. Finally, we reverse the subarray starting from index i+1 to the end of the array, which gives us the next permutation in lexicographical order. This approach leverages in-place modifications using basic operations like swapping and reversing, thus meeting the constant space requirement.

Two Pointers

#75 Sort Colors

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Sorting

Lists: Grind 169

Problem

Description

Given an array `nums` with `n` objects colored red, white, or blue, sort them in-place so that objects of the same color are adjacent, with the colors in the order red, white, and blue.

We will use the integers 0, 1, and 2 to represent the color red, white, and blue, respectively.

You must solve this problem without using the library's sort function.

Example 1:

Input: `nums = [2,0,2,1,1,0]`Output: `[0,0,1,1,2,2]`

Example 2:

Input: `nums = [2,0,1]`Output: `[0,1,2]`

Constraints:

- `n == nums.length`
- `1 <= n <= 300`
- `nums[i]` is either 0, 1, or 2.

Follow up: Could you come up with a one-pass algorithm using only constant extra space?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def sortColors(self, nums: list[int]) -> None:
        zero = -1
        one = -1
        two = -1

        for num in nums:
            if num == 0:
                two += 1
                one += 1
```

```

    zero += 1
    nums[two] = 2
    nums[one] = 1
    nums[zero] = 0
elif num == 1:
    two += 1
    one += 1
    nums[two] = 2
    nums[one] = 1
else:

```

```

    two += 1
    nums[two] = 2

```

```

class Solution:
    def sortColors(self, nums: list[int]) -> None:
        l = 0 # The next 0 should be placed in l.
        r = len(nums) - 1 # The next 2 should be placed
        in r.

        i = 0
        while i <= r:
            if nums[i] == 0:
                nums[i], nums[l] = nums[l], nums[i]
                i += 1

```

```

        l += 1
    elif nums[i] == 1:
        i += 1
    else:
        # We may swap a 0 to index i, but we're
        still not sure whether this 0
        # is placed in the correct index, so we
        can't move pointer i.
        nums[i], nums[r] = nums[r], nums[i]
        r -= 1

```

• LeetCode

```

class Solution:
    def sortColors(self, nums: List[int]) -> None:
        i, j, k = -1, len(nums), 0
        while k < j:
            if nums[k] == 0:
                i += 1
                nums[i], nums[k] = nums[k],
                nums[i]

                k += 1
            elif nums[k] == 2:
                j -= 1
                nums[j], nums[k] = nums[k],
                nums[j]

            else:
                k += 1

```

• SimplyLeet

```
# Function to sort the colors using Dutch National
Flag algorithm
def sortColors(nums):
    # Initialize pointers for the start, current,
    and end positions
    low, mid, high = 0, 0, len(nums) - 1

    # Process elements until mid pointer exceeds
    high pointer
    while mid <= high:
        # If the current element is 0 (red)
        if nums[mid] == 0:
            # Swap current element with the element
            at low pointer
```

```
        nums[low], nums[mid] = nums[mid],
nums[low]
        # Increment both low and mid pointers
        low += 1
        mid += 1
        # If the current element is 1 (white)
        elif nums[mid] == 1:
            # Move mid pointer forward since it's
            already in the correct place
            mid += 1
        # If the current element is 2 (blue)
        else:
```

```
            # Swap current element with the element
            at high pointer
            nums[mid], nums[high] = nums[high],
nums[mid]
            # Decrement high pointer
            high -= 1

# Example usage:
nums = [2,0,2,1,1,0]
sortColors(nums)
print(nums) # Output: [0,0,1,1,2,2]
```

• LC.ca

```
class Solution:
    def sortColors(self, nums: List[int]) -> None:
        i, j, k = -1, len(nums), 0
        while k < j:
            if nums[k] == 0:
                i += 1
                nums[i], nums[k] = nums[k],
nums[i]
                k += 1
            elif nums[k] == 2:
                j -= 1
                nums[j], nums[k] = nums[k],
nums[j]
            else:
                k += 1
```

◆ Quick Review

Key Insights

- Use the Dutch National Flag algorithm to partition the array into three sections.
- Maintain three pointers: one for the next position for 0 (low), one for the next position for 2 (high), and one to traverse the array (mid).
- Swap elements appropriately to ensure that all 0's are at the beginning, all 2's are at the end, and 1's remain in the middle.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution uses a three-pointer approach (Dutch National Flag algorithm). Initialize three pointers: low, mid, and high. As you iterate through the array:

- If the element is 0, swap it with the element at the low pointer and increment both low and mid.
 - If the element is 1, simply move mid forward.
 - If the element is 2, swap it with the element at the high pointer and decrement high.
- This single pass method ensures that the array is sorted in-place with constant extra space. Key details include correctly handling swaps without losing any data and ensuring that the pointers are updated in the proper order.

Two Pointers

#161 One Edit Distance

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · String

Lists: Premium 98

Problem

Description

Given two strings s and t , return true if they are both one edit distance apart, otherwise return false.

A string s is said to be one distance apart from a string t if you can:

- Insert exactly one character into s to get t .
- Delete exactly one character from s to get t .
- Replace exactly one character of s with a different character to get t .

Example 1:

Input: $s = "ab"$, $t = "acb"$

Output: true

Explanation: We can insert 'c' into s to get t .

Example 2:

Input: $s = ""$, $t = ""$

Output: false

Explanation: We cannot get t from s by only one step.

Constraints:

- $0 \leq s.length, t.length \leq 104$
- s and t consist of lowercase letters, uppercase letters, and digits.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isOneEditDistance(self, s: str, t: str) -> bool:
        m = len(s)
        n = len(t)
        if m > n: # Make sure that |s| <= |t|.
            return self.isOneEditDistance(t, s)

        for i in range(m):
            if s[i] != t[i]:
                if m == n:

                    return s[i + 1:] == t[i + 1:] # Replace
s[i] with t[i].
                    return s[i:] == t[i + 1:] # Delete t[i].

                return m + 1 == n # Delete t[-1].
```

• LeetCode

```
class Solution:
    def isOneEditDistance(self, s: str, t: str) -> bool:
        if len(s) < len(t):
            return self.isOneEditDistance(t, s)
        m, n = len(s), len(t)
        if m - n > 1:
            return False
        for i, c in enumerate(t):
            if c != s[i]:
                return s[i + 1:] == t[i + 1:] if
m == n else s[i + 1:] == t[i:]

        return m == n + 1
```

• SimplyLeet

```
# Function to check if two strings are one edit
distance apart
def isOneEditDistance(s: str, t: str) -> bool:
    # Get lengths of both strings
    m, n = len(s), len(t)

    # The difference in length must not be greater
than 1
    if abs(m - n) > 1:
        return False

    # Ensure s is the shorter string, or they are
equal in length
```



```

if m > n:
    return isOneEditDistance(t, s)

# Initialize variables; found difference flag
and indices for both strings
found_difference = False
i = j = 0

while i < m and j < n:
    if s[i] != t[j]:
        # If a difference has already been
        found, return False

        if found_difference:
            return False
        found_difference = True
        # If lengths are equal, move both
        pointers
        if m == n:
            i += 1
        else:
            # If characters are the same, move the
            pointer in the shorter string
            i += 1
        # Always move pointer for longer string

```

```

j += 1

# If no difference was found during iteration,
then one extra character at the end is the edit
return found_difference or (j - i == 1)

# Example usage:
#print(isOneEditDistance("ab", "acb")) # Output:
True

```

• LC.ca

```

class Solution:
    def isOneEditDistance(self, s: str, t: str) ->
    bool:
        if len(s) < len(t): # goal is to ensure t
        is always the shorter one
            return self.isOneEditDistance(t, s)
        m, n = len(s), len(t)
        if m - n > 1:
            # case: abc, abcdefg, return false
            return False
        for i, c in enumerate(t):
            if c != s[i]:

```

```

        return s[i + 1 :] == t[i + 1 :] if
m == n else s[i + 1 :] == t[i:]
        # case: abc, abcd. return true
        # case: abc, abc. the same 2 strings should
return false
        # case: abc, abcdefg, return false
        return m == n + 1

```

```
#####
```

```

class Solution(object):
    def isOneEditDistance(self, s, t):

```

```

        """
        :type s: str
        :type t: str
        :rtype: bool
        """
        if len(s) >= len(t):
            i = 0
            while i < len(t) and s[i] == t[i]:
                i += 1
            return s != t and (s[i + 1:] == t[i:] if
len(s) != len(t) else s[i + 1:] == t[i + 1:])

```

```

        else:
            return self.isOneEditDistance(t, s)

```

◆ Quick Review

Round 2 · High · Medium

p.507

Key Insights

- The difference in length between s and t must be at most 1.
- When lengths are equal, the only possibility is that exactly one character is different (replacement).
- When one string is longer by one character, the difference can be fixed by one insertion (or deletion in the other direction).
- Use two pointers to compare corresponding characters and carefully handle the first difference.

Space & Time

Time Complexity: $O(n)$, where n is the length of the shorter string, as we traverse both strings at most once.

Space Complexity: $O(1)$, using only constant extra space.

Solution

The solution uses a two-pointer approach. First, we check if the length difference is greater than 1, in which case they cannot be one edit apart. For the two primary cases:

- When strings are of equal length, iterate through both strings and count the number of differing characters. There should be exactly one difference.
- When lengths differ by one, use pointers for both strings. Upon encountering a difference, increment the pointer for the longer string only and ensure that no previous difference has been encountered. If a second discrepancy is found, return false. The key is to handle the edge case where the difference is at the end of the string.

Round 2 · High · Medium

p.508

Two Pointers

#167 Two Sum II – Input Array Is Sorted

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Binary Search

Lists: Denny Zhang

Problem

Description

Given a 1-indexed array of integers `numbers` that is already sorted in non-decreasing order, find two numbers such that they add up to a specific target number. Let these two numbers be `numbers[index1]` and `numbers[index2]` where $1 \leq \text{index1} < \text{index2} \leq \text{numbers.length}$.

Return the indices of the two numbers `index1` and `index2`, each incremented by one, as an integer array `[index1, index2]` of length 2.

The tests are generated such that there is exactly one solution. You may not use the same element twice. Your solution must use only constant extra space.

Example 1:

Input: `numbers = [2,7,11,15]`, `target = 9`

Output: `[1,2]`

Explanation: The sum of 2 and 7 is 9. Therefore, `index1 = 1`, `index2 = 2`. We return `[1, 2]`.

Example 2:

Input: `numbers = [2,3,4]`, `target = 6`

Output: `[1,3]`

Explanation: The sum of 2 and 4 is 6. Therefore `index1 = 1`, `index2 = 3`. We return `[1, 3]`.

Example 3:

Input: `numbers = [-1,0]`, `target = -1`

Output: `[1,2]`

Explanation: The sum of `-1` and `0` is `-1`. Therefore `index1 = 1`, `index2 = 2`. We return `[1, 2]`.

Constraints:

- $2 \leq \text{numbers.length} \leq 3 \times 10^4$
- $-1000 \leq \text{numbers}[i] \leq 1000$
- `numbers` is sorted in non-decreasing order.
- $-1000 \leq \text{target} \leq 1000$
- The tests are generated such that there is exactly one solution.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def twoSum(self, numbers: list[int], target:
int) -> list[int]:
        l = 0
        r = len(numbers) - 1

        while l < r:
            summ = numbers[l] + numbers[r]
            if summ == target:
                return [l + 1, r + 1]
            if summ < target:
                l += 1
            else:
                r -= 1

```

• LeetCode

```

class Solution:
    def twoSum(self, numbers: List[int], target:
int) -> List[int]:
        n = len(numbers)
        for i in range(n - 1):
            x = target - numbers[i]
            j = bisect_left(numbers, x, lo=i + 1)
            if j < n and numbers[j] == x:
                return [i + 1, j + 1]

```

```

class Solution:
    def twoSum(self, numbers: List[int], target:
int) -> List[int]:
        i, j = 0, len(numbers) - 1
        while i < j:
            x = numbers[i] + numbers[j]
            if x == target:
                return [i + 1, j + 1]
            if x < target:
                i += 1
            else:
                j -= 1

```

• SimplyLeet

```

# Python solution for Two Sum II - Input Array Is
Sorted
def twoSum(numbers, target):
    # Initialize two pointers at the beginning and
    end of the array
    left = 0
    right = len(numbers) - 1

    # Continue until the two pointers meet
    while left < right:
        # Calculate the sum of the two pointed
        values
        current_sum = numbers[left] +
        numbers[right]

```

```

    # If the sum matches the target, return the
    1-indexed results
    if current_sum == target:
        return [left + 1, right + 1]
    # If the current sum is less than the
    target, move the left pointer to the right
    elif current_sum < target:
        left += 1
    # Otherwise, move the right pointer to the
    left since the current sum is too high
    else:
        right -= 1

```

```

# Example call
print(twoSum([2, 7, 11, 15], 9)) # Output: [1, 2]

```

• LC.ca

```

class Solution:
    def twoSum(self, numbers: List[int], target:
    int) -> List[int]:
        i, j = 0, len(numbers) - 1
        while i < j:
            x = numbers[i] + numbers[j]
            if x == target:
                return [i + 1, j + 1]
            if x < target:
                i += 1
            else:
                j -= 1

```

◆ Quick Review

Key Insights

- The array is sorted in non-decreasing order, which makes the two pointers technique very effective.
- Using two pointers (one at the start and one at the end) allows checking pairs and adjusting the pointers based on the sum.
- Since the problem guarantees exactly one solution and requires constant space, using two pointers satisfies both constraints.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in the array, because in the worst-case scenario each element is visited at most once.

Space Complexity: $O(1)$, as no additional data structure is used that scales with the input size.

Solution

We use the two pointers approach:

- Initialize two pointers, left at the beginning (index 0) and right at the end (last index) of the array.
- While left is less than right, compute the sum of the elements at the two pointers.
- If the sum equals the target, return the indices as [left+1, right+1] (to adjust for the 1-indexed array).
- If the sum is less than the target, move the left pointer to the right to increase the sum.
- If the sum is greater than the target, move the right pointer to the left to decrease the

sum. This approach leverages the sorted property and uses constant extra space.

Two Pointers

#186 Reverse Words in a String II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · String

Lists: Premium 98

Problem

Description

Given a character array *s*, reverse the order of the words. A word is defined as a sequence of non-space characters. The words in *s* will be separated by a single space.

Your code must solve the problem in-place, i.e. without allocating extra space.

Example 1:

```
Input: s = ["t","h","e"," ","s","k","y"," ","i","s"," ","b","l","u","e"]
Output: ["b","l","u","e"," ","i","s"," ","s","k","y"," ","t","h","e"]
```

Example 2:

```
Input: s = ["a"]
Output: ["a"]
```

Constraints:

- $1 \leq s.length \leq 105$

- `s[i]` is an English letter (uppercase or lowercase), digit, or space ' '.
- There is at least one word in `s`.
- `s` does not contain leading or trailing spaces.
- All the words in `s` are guaranteed to be separated by a single space.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def reverseWords(self, s: list[str]) -> None:
        def reverse(l: int, r: int) -> None:
            while l < r:
                s[l], s[r] = s[r], s[l]
                l += 1
                r -= 1

        def reverseWords(n: int) -> None:
            i = 0
```

```
        j = 0

        while i < n:
            while i < j or (i < n and s[i] == ' '): #
                Skip the spaces.
                i += 1
            while j < i or (j < n and s[j] != ' '): #
                Skip the spaces.
                j += 1
            reverse(i, j - 1) # Reverse the word.

            reverse(0, len(s) - 1) # Reverse the whole
            string.

        reverseWords(len(s)) # Reverse each word.
```

• LeetCode

```
class Solution:
    def reverseWords(self, s: List[str]) -> None:
        def reverse(i: int, j: int):
            while i < j:
                s[i], s[j] = s[j], s[i]
                i, j = i + 1, j - 1

        i, n = 0, len(s)
        for j, c in enumerate(s):
            if c == " ":
```

```

        reverse(i, j - 1)
        i = j + 1
    elif j == n - 1:
        reverse(i, j)
    reverse(0, n - 1)

```

• SimplyLeet

```

def reverse_range(s, left, right):
    # Reverse the elements in s from index left to
    # right.
    while left < right:
        s[left], s[right] = s[right], s[left]
        left += 1
        right -= 1

```

```

def reverseWords(s):
    # Step 1: Reverse the entire array.
    reverse_range(s, 0, len(s) - 1)

```

```

    n = len(s)
    start = 0
    # Step 2: Reverse each word to correct their
    # order.
    for i in range(n + 1):
        if i == n or s[i] == ' ':
            reverse_range(s, start, i - 1)
            start = i + 1

```

• LC.ca

Round 2 · High · Medium

p.519

```

class Solution:
    def reverseWords(self, s: List[str]) -> None:
        """
        Do not return anything, modify s in-place
        instead.
        """

```

```

    def reverse(s, i, j):
        while i < j:
            s[i], s[j] = s[j], s[i]
            i += 1

```

```

        j -= 1

```

```

    i, j, n = 0, 0, len(s)
    while j < n:
        if s[j] == ' ':
            reverse(s, i, j - 1)
            i = j + 1
        elif j == n - 1:
            reverse(s, i, j)
        j += 1

```

```

    reverse(s, 0, n - 1)

```

◆ Quick Review

Key Insights

- Reverse the entire array first to obtain the reversed order of characters.

Round 2 · High · Medium

p.520

- Then, reverse each individual word to restore the correct letter order.
- Doing so leverages the two-pointer technique for in-place reversals.
- This method accomplishes the goal in $O(n)$ time while using $O(1)$ extra space.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution employs a two-step reversal process using the two-pointer technique. First, reverse the entire character array to reverse the overall order of characters. Although this reverses both the words and the letters in each word, it sets up the scenario where each word's letters are in reverse order. The next step is to iterate through the array and locate each word (using the space character as a delimiter) and then reverse each word individually to correct its letter order. This process is done in-place without any additional data structures, conserving space.

Two Pointers

#189 Rotate Array

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Math · Two Pointers

Lists: Grind 169

Problem

Description

Given an integer array `nums`, rotate the array to the right by `k` steps, where `k` is non-negative.

Example 1:

Input: `nums = [1,2,3,4,5,6,7]`, `k = 3`
 Output: `[5,6,7,1,2,3,4]`
 Explanation:
 rotate 1 steps to the right: `[7,1,2,3,4,5,6]`
 rotate 2 steps to the right: `[6,7,1,2,3,4,5]`
 rotate 3 steps to the right: `[5,6,7,1,2,3,4]`

Example 2:

Input: `nums = [-1,-100,3,99]`, `k = 2`
 Output: `[3,99,-1,-100]`
 Explanation:
 rotate 1 steps to the right: `[99,-1,-100,3]`
 rotate 2 steps to the right: `[3,99,-1,-100]`

Constraints:

- $1 \leq \text{nums.length} \leq 105$

- $-231 \leq \text{nums}[i] \leq 231 - 1$
- $0 \leq k \leq 105$

Follow up:

- Try to come up with as many solutions as you can. There are at least three different ways to solve this problem.
- Could you do it in-place with $O(1)$ extra space?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def rotate(self, nums: list[int], k: int) ->
    None:
        k %= len(nums)
        self.reverse(nums, 0, len(nums) - 1)
        self.reverse(nums, 0, k - 1)
        self.reverse(nums, k, len(nums) - 1)

    def reverse(self, nums: list[int], l: int, r:
    int) -> None:
        while l < r:
            nums[l], nums[r] = nums[r], nums[l]

            l += 1
            r -= 1
```

• LeetCode

```
class Solution:
    def rotate(self, nums: List[int], k: int) ->
    None:
        def reverse(i: int, j: int):
            while i < j:
                nums[i], nums[j] = nums[j],
                nums[i]
                i, j = i + 1, j - 1

        n = len(nums)
        k %= n
        reverse(0, n - 1)

        reverse(0, k - 1)
        reverse(k, n - 1)
```

• SimplyLeet

```
# Python solution using the reverse method

def rotate(nums, k):
    # Calculate the effective rotations needed
    n = len(nums)
    k = k % n # in case k is greater than the
    length of the array

    # Helper function to reverse a portion of the
    list in place
    def reverse(left, right):
        while left < right:
```

```

        # Swap the elements at indices left and
right      nums[left], nums[right] = nums[right],
nums[left]
        left += 1
        right -= 1

```

```

    # Reverse the entire array
    reverse(0, n - 1)
    # Reverse the first k elements to restore their
order      reverse(0, k - 1)
    # Reverse the remaining n - k elements

```

```
reverse(k, n - 1)
```

```

# Example usage:
# nums = [1,2,3,4,5,6,7]
# rotate(nums, 3)
# print(nums) # Output: [5,6,7,1,2,3,4]

```

- LC.ca

```

class Solution:
    def rotate(self, nums: List[int], k: int) ->
None:
        k %= len(nums)
        nums[:] = nums[-k:] + nums[:-k]

```

```
#####
```

```

class Solution:
    def rotate(self, nums: List[int], k: int) ->
None:
        """"

```

Do not return anything, modify nums
in-place instead.

```

        """"
        n = len(nums)
        k %= n
        if n < 2 or k == 0:
            return
        nums[:] = nums[::-1]
        nums[:k] = nums[:k][::-1]
        nums[k:] = nums[k:][::-1]

```

```
#####
```

```
class Solution(object):
    def rotate(self, nums, k):
        """
        :type nums: List[int]
        :type k: int
        :rtype: void Do not return anything, modify
        nums in-place instead.
        """
        if len(nums) == 0 or k == 0:

            return

        def reverse(start, end, s):
            while start < end:
                s[start], s[end] = s[end], s[start]
                start += 1
                end -= 1

        n = len(nums) - 1
        k = k % len(nums)

        reverse(0, n - k, nums)
        reverse(n - k + 1, n, nums)
        reverse(0, n, nums)
```

◆ Quick Review

Key Insights

- The rotation amount k might be greater than the length of the array, so use $k \% n$ (where n is the array length) to compute the effective rotation.
- A smart in-place solution involves reversing sections of the array:
 - Reverse the entire array.
 - Reverse the first k elements.
 - Reverse the remaining $n-k$ elements.
- Alternative approaches include using extra space or performing k single-step rotations (which is less optimal).

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the array.

Space Complexity: $O(1)$ if done in-place.

Solution

We solve the problem using the "reverse" technique to perform the rotation in-place. First, we adjust k by taking k modulo the length of the array to handle cases where k exceeds the array length. Then, we reverse the entire array. Next, we reverse the first k elements to restore their correct order. Finally, we reverse the remaining $n-k$ elements. This sequence yields the array rotated to the right by k positions while strictly using $O(1)$ extra space.

Two Pointers

#532 K-diff Pairs in an Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Two Pointers · Binary Search · Sorting
Lists: CodeSignal

Problem

Description

Given an array of integers `nums` and an integer `k`, return the number of unique `k`-diff pairs in the array.

A `k`-diff pair is an integer pair $(\text{nums}[i], \text{nums}[j])$, where the following are true:

- $0 \leq i, j < \text{nums.length}$
- $i \neq j$
- $|\text{nums}[i] - \text{nums}[j]| == k$

Notice that $|\text{val}|$ denotes the absolute value of `val`.

Example 1:

Input: `nums = [3,1,4,1,5]`, `k = 2`
 Output: 2
 Explanation: There are two 2-diff pairs in the array, (1, 3) and (3, 5).
 Although we have two 1s in the input, we should only return the number of unique pairs.

Example 2:

Input: `nums = [1,2,3,4,5]`, `k = 1`
 Output: 4
 Explanation: There are four 1-diff pairs in the array, (1, 2), (2, 3), (3, 4) and (4, 5).

Example 3:

Input: `nums = [1,3,1,5,4]`, `k = 0`
 Output: 1
 Explanation: There is one 0-diff pair in the array, (1, 1).

Constraints:

- $1 \leq \text{nums.length} \leq 104$
- $-107 \leq \text{nums}[i] \leq 107$
- $0 \leq k \leq 107$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def findPairs(self, nums: list[int], k: int) ->
    int:
        ans = 0
        numToIndex = {num: i for i, num in
        enumerate(nums)}

        for i, num in enumerate(nums):
            target = num + k
            if target in numToIndex and
            numToIndex[target] != i:
                ans += 1
                del numToIndex[target]

        return ans
```

• LeetCode

```
class Solution:
    def findPairs(self, nums: List[int], k: int) ->
    int:
        ans = set()
        vis = set()
        for x in nums:
            if x - k in vis:
                ans.add(x - k)
            if x + k in vis:
                ans.add(x)
            vis.add(x)
```

```
return len(ans)
```

• SimplyLeet

```
# Define a function to count k-diff pairs in the
array.
def find_pairs(nums, k):
    # If k is negative, return 0 as absolute
    difference can't be negative.
    if k < 0:
        return 0

    count = 0

    # k == 0: Need to find numbers that appear more
    than once.
    if k == 0:

        # Create a frequency map for nums.
        freq = {}
        for num in nums:
            freq[num] = freq.get(num, 0) + 1
        # Count numbers with frequency greater than
        1.
        for val in freq.values():
            if val > 1:
                count += 1
        else:
            # k > 0: Use a set to check for the
            existence of each number's complement.
```

```

unique_nums = set(nums)
for num in unique_nums:
    if num + k in unique_nums:
        count += 1
return count

```

Example usage:
 # print(find_pairs([3,1,4,1,5], 2)) # Expected
 output: 2

• LC.ca

```

class Solution:
    def findPairs(self, nums: List[int], k: int) ->
    int:
        vis, ans = set(), set()
        for v in nums:
            if v - k in vis:
                ans.add(v - k)
            if v + k in vis:
                ans.add(v)
            vis.add(v)
        return len(ans)

```

◆ Quick Review

Key Insights

- For $k < 0$, the answer is always 0 because the absolute difference cannot be negative.

- When k is 0, you are looking for numbers that appear at least twice.
- For $k > 0$, it is efficient to use a hash set or a hash map to check if each number's complement (number + k) exists.
- Uniqueness is ensured by processing each number only once from the set of unique numbers.

Space & Time

Time Complexity: $O(n)$, where n is the length of the array since we iterate over the array and then over the unique numbers.

Space Complexity: $O(n)$, for storing the frequency count or the unique elements in a hash map or hash set.

Solution

The solution first checks if k is negative, returning 0 immediately if true. For k equals 0, we count numbers that appear more than once using a frequency map. For k greater than 0, a set of unique numbers is created and for each number, we check if (number + k) is also present in the set. Each valid case contributes to the result count, ensuring that only unique pairs are considered.

Two Pointers

#923 3Sum With Multiplicity

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Two Pointers · Sorting · Counting
Lists: CodeSignal

Problem

Description

Given an integer array `arr`, and an integer `target`, return the number of tuples `i, j, k` such that $i < j < k$ and $arr[i] + arr[j] + arr[k] == target$.

As the answer can be very large, return it modulo $10^9 + 7$.

Example 1:

Input: `arr = [1,1,2,2,3,3,4,4,5,5]`, `target = 8`
 Output: 20
 Explanation:
 Enumerating by the values (`arr[i]`, `arr[j]`, `arr[k]`):
 (1, 2, 5) occurs 8 times;
 (1, 3, 4) occurs 8 times;
 (2, 2, 4) occurs 2 times;
 (2, 3, 3) occurs 2 times.

Example 2:

Input: `arr = [1,1,2,2,2,2]`, `target = 5`
 Output: 12
 Explanation:
`arr[i] = 1`, `arr[j] = arr[k] = 2` occurs 12 times:
 We choose one 1 from `[1,1]` in 2 ways,
 and two 2s from `[2,2,2,2]` in 6 ways.

Example 3:

Input: `arr = [2,1,3]`, `target = 6`
 Output: 1
 Explanation: (1, 2, 3) occurred one time in the array so we return 1.

Constraints:

- $3 \leq arr.length \leq 3000$
- $0 \leq arr[i] \leq 100$
- $0 \leq target \leq 300$

Community Solutions (Python)

- WalkCC


```

class Solution:
    def threeSumMulti(self, arr: list[int], target:
int) -> int:
        MOD = 1_000_000_007
        ans = 0
        count = collections.Counter(arr)

        for i, x in count.items():
            for j, y in count.items():
                k = target - i - j
                if k not in count:
                    continue
                if i == j and j == k:
                    ans = (ans + x * (x - 1) * (x - 2) // 6)
% MOD
                elif i == j and j != k:
                    ans = (ans + x * (x - 1) // 2 * count[k])
% MOD
                elif i < j and j < k:
                    ans = (ans + x * y * count[k]) % MOD

        return ans % MOD

```

• LeetCode

Round 2 · High · Medium

p.537

```

class Solution:
    def threeSumMulti(self, arr: List[int],
target: int) -> int:
        mod = 10**9 + 7
        cnt = Counter(arr)
        ans = 0
        for j, b in enumerate(arr):
            cnt[b] -= 1
            for a in arr[:j]:
                c = target - a - b
                ans = (ans + cnt[c]) % mod

        return ans

```

• SimplyLeet

```

MOD = 10**9 + 7

def threeSumMulti(arr, target):
    # Create a frequency array for numbers 0 to 100
    freq = [0] * 101
    for num in arr:
        freq[num] += 1

    count = 0
    # Iterate through all possible values for x, y,
    and implicitly z

```

Round 2 · High · Medium

p.538

```
for x in range(101):
    if freq[x] == 0:
        continue
    for y in range(x, 101):
        if freq[y] == 0:
            continue
        z = target - x - y
        if z < 0 or z > 100 or freq[z] == 0:
            continue
        if x == y and y == z:
```

```
            # All numbers are the same: nC3 =
n*(n-1)*(n-2)/6
            count += freq[x] * (freq[x] - 1) *
(freq[x] - 2) // 6
            elif x == y and y != z:
                # x and y are the same, z is
different: nC2 for x multiplied by count[z]
                count += (freq[x] * (freq[x] - 1)
// 2) * freq[z]
            elif x < y and y < z:
                # All distinct: product of counts
                count += freq[x] * freq[y] *
freq[z]
            elif x < y and y == z:
                # x different, y and z are the
same: count[x] multiplied by nC2 for y
```

```
        count += freq[x] * (freq[y] *
(freq[y] - 1) // 2)
        count %= MOD
    return count
```

```
# Example usage:
print(threeSumMulti([1,1,2,2,3,3,4,4,5,5], 8))
```

• LC.ca

```
class Solution:
    def threeSumMulti(self, arr: List[int],
target: int) -> int:
        cnt = Counter(arr)
        ans = 0
        mod = 10**9 + 7
        for j, b in enumerate(arr):
            cnt[b] -= 1
            for i in range(j):
                a = arr[i]
                c = target - a - b

        ans = (ans + cnt[c]) % mod
    return ans
```

◆ Quick Review

Key Insights

- Use a frequency counter for elements since `arr[i]` is in the range `[0, 100]`.

- Iterate through all possible combinations of numbers x , y , and z with $x \leq y \leq z$.
- For each valid triplet (x, y, z) that sums to target, count the number of ways using combinatorial formulas:
 - All distinct: $\text{count}[x] * \text{count}[y] * \text{count}[z]$.
 - Two numbers the same: use $nC2$ for the repeated element.
 - All three the same: use $nC3$.
- Use modulo arithmetic ($\text{MOD} = 10^9 + 7$) at every step to avoid overflow.

Space & Time

Time Complexity: $O(M^2)$, where $M = 101$ (the range of possible numbers), which is efficient given the constraints.

Space Complexity: $O(M)$ for the frequency array.

Solution

The solution first counts the frequency of each number in the array. Then it iterates through all valid triplets (x, y, z) (with $x \leq y \leq z$) and checks if they sum up to the target. Depending on whether x , y , and z are distinct or have duplicates, the combination count is computed accordingly:

- If x , y , and z are all the same, compute the combination as $nC3$.
- If only two are the same, compute as $nC2$ multiplied by the count of the third number.
- If all are distinct, use the direct product of the counts.

This covers all cases while applying the modulo operation to the result.

Two Pointers

#986 Interval List Intersections

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers

Lists: Denny Zhang

Problem

Description


You are given two lists of closed intervals, `firstList` and `secondList`, where `firstList[i] = [starti, endi]` and `secondList[j] = [startj, endj]`. Each list of intervals is pairwise disjoint and in sorted order.

Return the intersection of these two interval lists.

A closed interval $[a, b]$ (with $a \leq b$) denotes the set of real numbers x with $a \leq x \leq b$.

The intersection of two closed intervals is a set of real numbers that are either empty or represented as a closed interval. For example, the intersection of $[1, 3]$ and $[2, 4]$ is $[2, 3]$.

Example 1:



Input: firstList = [[0,2],[5,10],[13,23],[24,25]],
 secondList = [[1,5],[8,12],[15,24],[25,26]]
 Output: [[1,2],[5,5],[8,10],[15,23],[24,24],[25,25]]

Example 2:

Input: firstList = [[1,3],[5,9]], secondList = []
 Output: []

Constraints:

- $0 \leq \text{firstList.length}, \text{secondList.length} \leq 1000$
- $\text{firstList.length} + \text{secondList.length} \geq 1$
- $0 \leq \text{start}_i < \text{end}_i \leq 109$
- $\text{end}_i < \text{start}_{i+1}$
- $0 \leq \text{start}_j < \text{end}_j \leq 109$
- $\text{end}_j < \text{start}_{j+1}$

Community Solutions (Python)

- [WalkCC](#)

```
class Solution:
    def intervalIntersection(self, firstList:
        list[list[int]],
                               secondList:
        list[list[int]]) -> list[list[int]]:
        ans = []
        i = 0
        j = 0

        while i < len(firstList) and j <
            len(secondList):
            # lo := the start of the intersection
            # hi := the end of the intersection

            lo = max(firstList[i][0], secondList[j][0])
            hi = min(firstList[i][1], secondList[j][1])
            if lo <= hi:
                ans.append([lo, hi])
            if firstList[i][1] < secondList[j][1]:
                i += 1
            else:
                j += 1

        return ans
```

- [LeetDoocs](#)

```

class Solution:
    def intervalIntersection(
        self, firstList: List[List[int]],
        secondList: List[List[int]]
    ) -> List[List[int]]:
        i = j = 0
        ans = []
        while i < len(firstList) and j <
len(secondList):
            s1, e1, s2, e2 = *firstList[i],
*secondList[j]
            l, r = max(s1, s2), min(e1, e2)
            if l <= r:

                ans.append([l, r])
            if e1 < e2:
                i += 1
            else:
                j += 1
        return ans

```

• SimplyLeet

```

# Function to compute the interval intersections
using two pointers
def intervalIntersection(firstList, secondList):
    intersections = [] # List to store the
    resulting intersections
    i, j = 0, 0 # Pointers for firstList and
    secondList

    # Loop until either list is fully traversed
    while i < len(firstList) and j <
len(secondList):
        start1, end1 = firstList[i]
        start2, end2 = secondList[j]

        # Find the intersection boundaries
        start_overlap = max(start1, start2)
        end_overlap = min(end1, end2)

        # Check if there is an intersection
        if start_overlap <= end_overlap:
            intersections.append([start_overlap,
end_overlap])

        # Move the pointer that has the smaller
        interval ending
        if end1 < end2:

```

```

        i += 1
    else:
        j += 1

    return intersections

```

Example usage:

```

firstList = [[0,2],[5,10],[13,23],[24,25]]
secondList = [[1,5],[8,12],[15,24],[25,26]]
print(intervalIntersection(firstList, secondList))
# Output:
[[1,2],[5,5],[8,10],[15,23],[24,24],[25,25]]

```

• LC.ca

```

class Solution:
    def intervalIntersection(
        self, firstList: List[List[int]],
        secondList: List[List[int]]
    ) -> List[List[int]]:
        i = j = 0
        ans = []
        while i < len(firstList) and j <
len(secondList):
            s1, e1, s2, e2 = *firstList[i],
*secondList[j]
            l, r = max(s1, s2), min(e1, e2)
            if l <= r:

```

```

        ans.append([l, r])
        if e1 < e2:
            i += 1
        else:
            j += 1
    return ans

```

◆ Quick Review

Key Insights

- Use two pointers to traverse both lists simultaneously.
- Determine the overlapping interval by taking the maximum of the start points and the minimum of the end points.
- If the overlapping condition (start <= end) is met, add the intersection to the result.
- Move the pointer that has the smaller end value to explore further potential intersections efficiently.
- Achieve an overall time complexity of $O(m+n)$, where m and n are the lengths of the two interval lists.

Space & Time

Time Complexity: $O(m + n)$ where m and n are the lengths of firstList and secondList respectively.

Space Complexity: $O(1)$ extra space (excluding the space used for the output).

Solution

The solution uses a two pointers approach. Initialize two indices, one for each list, and iterate through both lists. For each pair of intervals, the potential intersection is computed by:

- Calculating the maximum of the two start values.
- Calculating the minimum of the two end values. If the maximum start is less than or equal to the minimum end, there is an intersection, and it is added to the result list. The pointer corresponding to the interval with the smaller end value is then incremented since that interval cannot intersect with any further intervals from the other list. This process repeats until one of the lists is fully traversed.

Two Pointers

#1055 Shortest Way to Form String

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · String · Binary Search

Lists: Premium 98

Problem

Description

A subsequence of a string is a new string that is formed from the original string by deleting some (can be none) of the characters without disturbing the relative positions of the remaining characters. (i.e., "ace" is a subsequence of "abcde" while "aec" is not).

Given two strings source and target, return the minimum number of subsequences of source such that their concatenation equals target. If the task is impossible, return -1.

Example 1:

Input: source = "abc", target = "abcbc"

Output: 2

Explanation: The target "abcbc" can be formed by "abc" and "bc", which are subsequences of source "abc".

Example 2:

Input: source = "abc", target = "acdbc"

Output: -1

Explanation: The target string cannot be constructed from the subsequences of source string due to the character "d" in target string.

Example 3:

Input: source = "xyz", target = "xzyxz"

Output: 3

Explanation: The target string can be constructed as follows "xz" + "y" + "xz".

Constraints:

- $1 \leq \text{source.length}, \text{target.length} \leq 1000$
- source and target consist of lowercase English letters.

Community Solutions (Python)

• LeetCode

```
class Solution:
    def shortestWay(self, source: str, target: str) -> int:
        def f(i, j):
            while i < m and j < n:
                if source[i] == target[j]:
                    j += 1
                i += 1
            return j

        m, n = len(source), len(target)
```

```
ans = j = 0
while j < n:
    k = f(0, j)
    if k == j:
        return -1
    j = k
    ans += 1
return ans
```

• SimplyLeet

Python solution using greedy two-pointer approach

```
def shortestWay(source, target):
    count = 0 # Count of subsequences used
    t_index = 0
    t_len = len(target)

    # Pre-check: if any character in target is not
    in source, return -1
    if set(target) - set(source):
        return -1
```



```

    # While there are still characters in target to
    match
    while t_index < t_len:
        s_index = 0 # Pointer for source
        progress = False # Flag to check if we made
        progress this pass

        # Iterate through source and try to match
        target characters
        while s_index < len(source) and t_index <
        t_len:
            if source[s_index] == target[t_index]:
                t_index += 1 # Move to next
                character in target

```

```

                progress = True
                s_index += 1

            count += 1 # One subsequence used for this
            pass

            # If no progress was made, then it is
            impossible to form the target
            if not progress:
                return -1

        return count

```

```

# Example usage:
# print(shortestWay("abc", "abcbc")) # Output: 2

```

• LC.ca

```

class Solution:
    def shortestWay(self, source: str, target:
    str) -> int:
        def f(i, j):
            while i < m and j < n:
                if source[i] == target[j]:
                    j += 1
                i += 1
            return j

        m, n = len(source), len(target)

        ans = j = 0
        while j < n:
            k = f(0, j)
            if k == j:
                return -1
            j = k
            ans += 1
        return ans

```

◆ Quick Review

Key Insights

- Check if every character in target exists in source; if not, the answer is -1.
- Use a greedy two-pointer strategy: iterate over target while scanning source for matching characters.
- During each pass over source, match as many characters as possible from target.
- Count the number of passes (subsequences) required until the target is completely matched.
- An optimized approach may precompute indices for binary search; however, the basic greedy simulation is efficient for the given constraints.

Space & Time

Time Complexity: $O(n * m)$ in the worst-case, where n = length of source and m = length of target.

Space Complexity: $O(1)$ for the greedy simulation ($O(n)$ if additional data structures for binary search are used).

Solution

The solution uses a greedy two-pointer approach. We simulate the process of reading the target by repeatedly scanning through the source string. Each scan through source is treated as one subsequence concatenated to form the target. A pointer in the target moves forward whenever a matching character is found in source. If a complete pass through source does not advance the target pointer, then it is impossible to form the target and we return -1. The algorithm counts the number of passes required until the target is fully matched.

Two Pointers

#2563 Count the Number of Fair Pairs

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Binary Search · Sorting

Lists: CodeSignal

Problem

Description

Given a 0-indexed integer array `nums` of size `n` and two integers `lower` and `upper`, return the number of fair pairs.

A pair (i, j) is fair if:

- $0 \leq i < j < n$, and
- $\text{lower} \leq \text{nums}[i] + \text{nums}[j] \leq \text{upper}$

Example 1:

Input: `nums = [0,1,7,4,4,5]`, `lower = 3`, `upper = 6`

Output: 6

Explanation: There are 6 fair pairs: $(0,3)$, $(0,4)$, $(0,5)$, $(1,3)$, $(1,4)$, and $(1,5)$.

Example 2:

Input: `nums = [1,7,9,2,5]`, `lower = 11`, `upper = 11`

Output: 1

Explanation: There is a single fair pair: $(2,3)$.

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $\text{nums.length} == n$
- $-109 \leq \text{nums}[i] \leq 109$
- $-109 \leq \text{lower} \leq \text{upper} \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def countFairPairs(self, nums: list[int], lower:
int, upper: int) -> int:
        # nums[i] + nums[j] == nums[j] + nums[i], so
the condition that i < j
        # degrades to i != j and we can sort the array.
        nums.sort()

        def countLess(summ: int) -> int:
            res = 0
            i = 0
            j = len(nums) - 1
```

```
        while i < j:
            while i < j and nums[i] + nums[j] > summ:
                j -= 1
            res += j - i
            i += 1
        return res

return countLess(upper) - countLess(lower - 1)
```

• LeetCode

```
class Solution:
    def countFairPairs(self, nums: List[int],
lower: int, upper: int) -> int:
        nums.sort()
        ans = 0
        for i, x in enumerate(nums):
            j = bisect_left(nums, lower - x, lo=i +
1)
            k = bisect_left(nums, upper - x + 1,
lo=i + 1)
            ans += k - j
        return ans
```

• SimplyLeet

```

import bisect

def countFairPairs(nums, lower, upper):
    # Sort the array to allow binary search
    nums.sort()
    fair_pairs = 0
    n = len(nums)

    for i in range(n):
        # Find lower bound index for the required
        # complement so that nums[i] + nums[j] >= lower
        left = bisect.bisect_left(nums, lower -
                                   nums[i], i + 1, n)
        # Find upper bound index for the required
        # complement so that nums[i] + nums[j] <= upper
        right = bisect.bisect_right(nums, upper -
                                     nums[i], i + 1, n)
        fair_pairs += (right - left)

    return fair_pairs

# Example usage:
print(countFairPairs([0,1,7,4,4,5], 3, 6)) #
Expected output: 6

```

- LC.ca

```

class Solution:
    def countFairPairs(self, nums: List[int],
                      lower: int, upper: int) -> int:
        nums.sort()
        ans = 0
        for i, x in enumerate(nums):
            j = bisect_left(nums, lower - x, lo=i +
                           1)
            k = bisect_left(nums, upper - x + 1,
                           lo=i + 1)
            ans += k - j
        return ans

```

♦ Quick Review

Key Insights

- Sorting the array simplifies searching for pairs with required sum limits.
- For each element, binary search is used to locate:
 - The first index where the complement makes the sum $\geq$ lower.
 - The last index where the complement makes the sum $\leq$ upper.
- This avoids the need for a nested loop through the entire array, yielding a more efficient solution.

Space & Time

Time Complexity: $O(n \log n)$, primarily due to sorting and (for each element) binary search operations.

Space Complexity: $O(n)$ if sorting creates a copy; $O(1)$ additional space if sorting in-place.

Solution

The solution starts by sorting the input array. For each element at index i in the sorted array, we search for indices j (where $j > i$) such that the sum of $\text{nums}[i]$ and $\text{nums}[j]$ falls between lower and upper. Two binary search operations are performed:

- To find the smallest index j where $\text{nums}[i] + \text{nums}[j] \geq \text{lower}$.
- To find the largest index j where $\text{nums}[i] + \text{nums}[j] \leq \text{upper}$.

Subtracting the two positions (and summing the counts for each i) gives the total number of fair pairs. Edge cases (like no valid pair) are automatically handled by the binary search approach.

Sliding Window

Round 2 · High · Medium · 8 questions

Round 2 · High · Medium

p.561

Round 2 · High · Medium

p.562

Sliding Window

#3 Longest Substring Without Repeating Characters

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Grind 169

Problem

Description

Given a string *s*, find the length of the longest substring without duplicate characters.

Example 1:

Input: *s* = "abcabcbb"
 Output: 3
 Explanation: The answer is "abc", with the length of 3.
 Note that "bca" and "cab" are also correct answers.

Example 2:

Input: *s* = "bbbbbb"
 Output: 1
 Explanation: The answer is "b", with the length of 1.

Example 3:

Input: *s* = "pwwkew"

Output: 3

Explanation: The answer is "wke", with the length of 3.
 Notice that the answer must be a substring, "pwke" is a subsequence and not a substring.

Constraints:

- $0 \leq s.length \leq 5 * 10^4$
- *s* consists of English letters, digits, symbols and spaces.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def lengthOfLongestSubstring(self, s: str) -> int:
        ans = 0
        count = collections.Counter()

        l = 0
        for r, c in enumerate(s):
            count[c] += 1
            while count[c] > 1:
                count[s[l]] -= 1
                l += 1
            ans = max(ans, r - l + 1)

        return ans
```

```
class Solution:
    def lengthOfLongestSubstring(self, s: str) ->
    int:
        ans = 0
        # The substring s[j + 1..i] has no repeating
        characters.
        j = -1
        # lastSeen[c] := the index of the last time c
        appeared
        lastSeen = {}

        for i, c in enumerate(s):
            # Update j to lastSeen[c], so the window must
            start from j + 1.
```

```
        j = max(j, lastSeen.get(c, -1))
        ans = max(ans, i - j)
        lastSeen[c] = i
```

```
    return ans
```

• LeetCode

Round 2 · High · Medium

p.565

```
class Solution:
    def lengthOfLongestSubstring(self, s: str) ->
    int:
        cnt = Counter()
        ans = l = 0
        for r, c in enumerate(s):
            cnt[c] += 1
            while cnt[c] > 1:
                cnt[s[l]] -= 1
                l += 1
            ans = max(ans, r - l + 1)

        return ans
```

• SimplyLeet

```
# Function to find the length of the longest
# substring without repeating characters
def length_of_longest_substring(s):
    # Dictionary to store the most recent index of
    # each character
    char_index = {}
    # Initialize starting index of current window
    # and max length
    left = 0
    max_length = 0

    # Iterate over the string with the right
    # pointer
    for right, char in enumerate(s):
```

Round 2 · High · Medium

p.566

```

        # If char is in dictionary and its index is
        within the current window
        if char in char_index and char_index[char]
        >= left:
            # Move the left pointer to one position
            after the last occurrence of char
            left = char_index[char] + 1
            # Update the latest index of the char
            char_index[char] = right
            # Update max_length with the current window
            size if larger
            max_length = max(max_length, right - left +
            1)

        return max_length

```

```

# Example usage:
# print(length_of_longest_substring("abcabcbb")) #
Expected output: 3

```

- LC.ca

```

class Solution:
    def lengthOfLongestSubstring(self, s): # map for
    index
        d = {} # value => its index
        i = 0
        ans = 0
        for j, c in enumerate(s):
            if c in d:
                # must max check i, example "abba"
                i = max(i, d[c] + 1)
                d[c] = j

            ans = max(ans, j - i + 1)
        return ans

```

```

class Solution:
    def lengthOfLongestSubstring(self, s: str) ->
    int:

        if not s:
            return 0

```



```

l = 0 # left
r = 0 # right

result = 0
isFoundInWindow = [False] * 256

while r < len(s):

    while isFoundInWindow[ord(s[r])]:
        isFoundInWindow[ord(s[l])] = False
        l += 1

    isFoundInWindow[ord(s[r])] = True

    # check before shrink
    result = max(result, r - l + 1)

    r += 1

return result

```

```

class Solution: # extra space for set()
    def lengthOfLongestSubstring(self, s: str) ->
    int:

        d = collections.defaultdict(int)
        start = 0
        ans = 0
        for i,c in enumerate(s):
            # shrink
            while d[c] > 0:

                d[s[start]] -= 1 # not s[start],
                not start
                start += 1
            ans = max(ans, i - start + 1)
            d[c] = 1
        return ans

#####

class Solution(object):
    def _lengthOfLongestSubstring(self, s): # no
    extra data structure

```

```

"""
:type s: str
:rtype: int
"""
d = collections.defaultdict(int)
l = ans = 0
for i, c in enumerate(s):
    while l > 0 and d[c] > 0:
        d[s[i - l]] -= 1
        l -= 1

```

```

    d[c] += 1
    l += 1
    ans = max(ans, l)
return ans

```

```
#####
```

```

class Solution:
    def lengthOfLongestSubstring(self, s: str) ->
    int:
        ss = set()

```

```

i = ans = 0
for j, c in enumerate(s):
    while c in ss:
        ss.remove(s[i])
        i += 1
    ss.add(c)
    ans = max(ans, j - i + 1)
return ans

```

◆ Quick Review

Key Insights

- Use the sliding window technique to maintain a window of unique characters.
- Use a hash table (or dictionary) to quickly check for duplicate characters and update window boundaries.
- The window is expanded by moving a pointer (right) and contracted by moving another pointer (left) when duplicates are encountered.
- Time complexity can be maintained at $O(n)$ by ensuring each character is processed a limited number of times.

Space & Time

Time Complexity: $O(n)$, where n is the length of the input string.

Space Complexity: $O(\min(n, m))$, where m is the size of the character set (constant for fixed character sets).

Solution

The solution uses a sliding window technique with two pointers, left and right, to keep track of the current substring without repeating characters. A hash table (dictionary or map) is used to store the most recent index of each character encountered. As the right pointer moves forward, if a duplicate character is found within the current window, the left pointer is advanced to the position right after where the duplicate was last seen. This ensures that the window always contains unique characters. The maximum length is updated throughout the process.

Sliding Window

#159 Longest Substring with At Most Two Distinct Characters

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Premium 98

Problem

Description

Given a string s , return the length of the longest substring that contains at most two distinct characters.

Example 1:

Input: $s = \text{"eceba"}$
 Output: 3
 Explanation: The substring is "ece" which its length is 3.

Example 2:

Input: $s = \text{"ccaabbb"}$
 Output: 5
 Explanation: The substring is "aabbb" which its length is 5.

Constraints:

- $1 \leq s.length \leq 105$
- s consists of English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def lengthOfLongestSubstringTwoDistinct(self, s:
str) -> int:
    ans = 0
    distinct = 0
    count = [0] * 128

    l = 0
    for r, c in enumerate(s):
        count[ord(c)] += 1
        if count[ord(c)] == 1:

            distinct += 1
            while distinct == 3:
                count[ord(s[l])] -= 1
                if count[ord(s[l])] == 0:
                    distinct -= 1
                l += 1
            ans = max(ans, r - l + 1)

    return ans
```

• LeetCode

```
class Solution:
    def lengthOfLongestSubstringTwoDistinct(self,
s: str) -> int:
        cnt = Counter()
        ans = j = 0
        for i, c in enumerate(s):
            cnt[c] += 1
            while len(cnt) > 2:
                cnt[s[j]] -= 1
                if cnt[s[j]] == 0:
                    cnt.pop(s[j])

                j += 1
            ans = max(ans, i - j + 1)
        return ans
```

• SimplyLeet

```
# Python solution for Longest Substring with At
Most Two Distinct Characters
class Solution:
    def lengthOfLongestSubstringTwoDistinct(self,
s: str) -> int:
        left = 0 # Initialize left pointer of the
window
        max_length = 0 # Variable to store the
maximum window size found
        char_count = {} # Dictionary to count
characters in the current window

        for right, char in enumerate(s):
            # Add the current character to the
window (increase count)
            char_count[char] =
char_count.get(char, 0) + 1
```

```
        # If there are more than two distinct
characters, shrink the window from the left
        while len(char_count) > 2:
            left_char = s[left]
            char_count[left_char] -= 1
            if char_count[left_char] == 0:
                del char_count[left_char]
            left += 1

        # Update max_length if current window
is larger
```

```
        max_length = max(max_length, right -
left + 1)

        return max_length
```

- LC.ca

```
'''
>>> from collections import Counter
>>> c = Counter()
>>> c
Counter()
>>> c['a'] += 1
>>> c
Counter({'a': 1})

>>> c
Counter({'a': 1, 'c': 1, 'b': 1})
>>> c.pop(2)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 2
>>>
>>> c.pop('c')
1
>>> c
```

```

Counter({'a': 1, 'b': 1})
...

from collections import Counter

class Solution:
    def lengthOfLongestSubstringTwoDistinct(self,
s: str) -> int:
        cnt = Counter() # or, cnt =
defaultdict(int)
        ans = i = 0
        for j, c in enumerate(s):

            cnt[c] += 1
            while len(cnt) > 2: # shrink
                cnt[s[i]] -= 1
                if cnt[s[i]] == 0:
                    cnt.pop(s[i])
                i += 1
            ans = max(ans, j - i + 1)
        return ans

```

◆ Quick Review

Key Insights

- Use a sliding window to maintain a contiguous substring.
- Keep track of character frequencies within the current window using a hash map or dictionary.
- When the window contains more than two distinct characters, shrink the window from the left until only two

remain.

- Update the maximum length each time the window satisfies the condition.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, since each character is processed at most twice.

Space Complexity: $O(1)$ because the hash map holds at most 3 entries (usually up to 2 distinct characters).

Solution

We use the sliding window technique with two pointers (left and right) to represent the current window. A hash table/dictionary records the count of each character inside the window. As we extend the right pointer, we update the count. If the hash table contains more than two distinct characters, we start moving the left pointer to reduce the number of distinct characters until we have at most two distinct keys. The length of the current valid window is then used to update our maximum length answer.

Sliding Window

#340 Longest Substring with At Most K Distinct Characters

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Premium 98

Problem

Description

Given a string s and an integer k , return the length of the longest substring of s that contains at most k distinct characters.

Example 1:

Input: $s = \text{"eceba"}, k = 2$

Output: 3

Explanation: The substring is "ece" with length 3.

Example 2:

Input: $s = \text{"aa"}, k = 1$

Output: 2

Explanation: The substring is "aa" with length 2.

Constraints:

- $1 \leq s.length \leq 5 * 10^4$
- $0 \leq k \leq 50$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def lengthOfLongestSubstringKDistinct(self, s: str, k: int) -> int:
        ans = 0
        distinct = 0
        count = collections.Counter()

        l = 0
        for r, c in enumerate(s):
            count[c] += 1
            if count[c] == 1:
                distinct += 1
            while distinct == k + 1:
                count[s[l]] -= 1
                if count[s[l]] == 0:
                    distinct -= 1
                l += 1
            ans = max(ans, r - l + 1)

        return ans
```

• LeetDoocs

```

class Solution:
    def lengthOfLongestSubstringKDistinct(self, s:
str, k: int) -> int:
        l = 0
        cnt = Counter()
        for c in s:
            cnt[c] += 1
            if len(cnt) > k:
                cnt[s[l]] -= 1
                if cnt[s[l]] == 0:
                    del cnt[s[l]]

                l += 1
        return len(s) - l

```

• SimplyLeet

Python solution for Longest Substring with At Most K Distinct Characters

```

def longest_substring_with_k_distinct(s: str, k:
int) -> int:

```

Handle edge case: if k is 0, no valid substring exists

```

    if k == 0:
        return 0

```

Dictionary to store frequency of characters in the current window

```

    char_count = {}
    left = 0 # Left pointer for the window

```

```

        max_length = 0 # Variable to track the maximum
length found

```

Iterate through the string using the right pointer

```

        for right in range(len(s)):
            # Add/update the count of the current
character in the dictionary
            current_char = s[right]
            char_count[current_char] =
char_count.get(current_char, 0) + 1

```

If window contains more than k distinct characters, shrink from the left

```

        while len(char_count) > k:

```

```

            left_char = s[left]
            char_count[left_char] -= 1
            # Remove the character count if it
becomes 0

```

```

            if char_count[left_char] == 0:
                del char_count[left_char]
            left += 1

```

Update the maximum length of the valid window found so far

```

        max_length = max(max_length, right - left +
1)

```



```
return max_length
```

```
# Example usage:
# s = "eceba", k = 2
# print(longest_substring_with_k_distinct(s, 2)) #
Output: 3
```

• LC.ca

```
class Solution:
    def lengthOfLongestSubstringTwoDistinct(self,
s: str) -> int:
    cnt = Counter() # or, cnt =
defaultdict(int)
    ans = i = 0
    for j, c in enumerate(s):
        cnt[c] += 1
        while len(cnt) > k:
            cnt[s[i]] -= 1
            if cnt[s[i]] == 0:
                cnt.pop(s[i])

        i += 1
        ans = max(ans, j - i + 1)
    return ans
```

◆ Quick Review

Key Insights

- Use the sliding window technique to efficiently scan through the string.

Round 2 · High · Medium

p.585

- Maintain a hash table (or dictionary) that keeps track of the frequency of characters in the current window.
- When the number of distinct characters exceeds k, increment the left pointer to shrink the window until the condition is met.
- Update the maximum length of valid substrings throughout the process.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, as each character is processed at most twice (once when expanding the window and once when contracting it).

Space Complexity: $O(\min(m, k))$, where m is the size of the charset (or number of unique characters in s), since at most $k+1$ keys are stored.

Solution

This problem is solved using the sliding window technique. The algorithm keeps track of a window defined by two indices (left and right pointers). A hash map is used to count the frequency of each character in the current window. As the right pointer extends the window, the count of the character is incremented. When the count of distinct keys in the map exceeds k , the left pointer is moved forward while decrementing the count of the leftmost character, until the condition is restored. The maximum window length encountered is recorded and returned as the result.

Round 2 · High · Medium

p.586

Sliding Window

#424 Longest Repeating Character Replacement

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Grind 169

Problem

Description

You are given a string s and an integer k . You can choose any character of the string and change it to any other uppercase English character. You can perform this operation at most k times.

Return the length of the longest substring containing the same letter you can get after performing the above operations.

Example 1:

Input: $s = \text{"ABAB"}, k = 2$

Output: 4

Explanation: Replace the two 'A's with two 'B's or vice versa.

Example 2:

Input: $s = \text{"AABABBA"}, k = 1$

Output: 4

Explanation: Replace the one 'A' in the middle with 'B' and form "AABBBBA".

The substring "BBBB" has the longest repeating letters, which is 4.

There may exist other ways to achieve this answer too.

Constraints:

- $1 \leq s.length \leq 105$
- s consists of only uppercase English letters.
- $0 \leq k \leq s.length$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def characterReplacement(self, s: str, k: int) -> int:
        ans = 0
        maxCount = 0
        count = collections.Counter()

        l = 0
        for r, c in enumerate(s):
            count[c] += 1
            maxCount = max(maxCount, count[c])
```

```

while maxCount + k < r - l + 1:
    count[s[l]] -= 1
    l += 1
ans = max(ans, r - l + 1)

return ans

```

```

class Solution:
    def characterReplacement(self, s: str, k: int) -> int:
        maxCount = 0
        count = collections.Counter()

        # l and r track the maximum window instead of
        # the valid window.
        l = 0
        for r, c in enumerate(s):
            count[c] += 1
            maxCount = max(maxCount, count[c])

```

```

while maxCount + k < r - l + 1:
    count[s[l]] -= 1
    l += 1

return r - l + 1

```

• LeetCode

```

class Solution:
    def characterReplacement(self, s: str, k: int) -> int:
        cnt = Counter()
        l = mx = 0
        for r, c in enumerate(s):
            cnt[c] += 1
            mx = max(mx, cnt[c])
            if r - l + 1 - mx > k:
                cnt[s[l]] -= 1
                l += 1

        return len(s) - l

```

• SimplyLeet

```

# Python solution using the sliding window
# technique
def characterReplacement(s: str, k: int) -> int:
    # Dictionary to store frequency of characters
    # in the current window
    char_counts = {}
    max_freq = 0 # Track the maximum frequency of a
    # single character in the window
    max_length = 0
    left = 0 # Left pointer of the sliding window

    # Expand the sliding window with the right
    # pointer
    for right in range(len(s)):

```

```

        # Count the current character
        char_counts[s[right]] =
char_counts.get(s[right], 0) + 1
        # Update the max frequency in the current
window
        max_freq = max(max_freq,
char_counts[s[right]])

        # Check if the number of replacements
needed exceeds k
        # (current window size) - (count of most
frequent character) > k
        if (right - left + 1) - max_freq > k:
            # If condition fails, shrink the window
by moving left pointer
            char_counts[s[left]] -= 1

```

```

        left += 1

        # Update the maximum length found so far
max_length = max(max_length, right - left +
1)

        return max_length

# Example usage:
# print(characterReplacement("AABABBA", 1)) #
Expected output: 4

```

• LC.ca

Round 2 · High · Medium

p.591

```

class Solution:
    def characterReplacement(self, s: str, k: int)
-> int:
        counter = [0] * 26
        i = j = maxCnt = 0
        while i < len(s):
            counter[ord(s[i]) - ord('A')] += 1
            maxCnt = max(maxCnt, counter[ord(s[i])
- ord('A')])
            if i - j + 1 > maxCnt + k:
                counter[ord(s[j]) - ord('A')] -= 1
                j += 1

        i += 1
        return i - j

```

◆ Quick Review

Key Insights

- Use a Sliding Window algorithm to keep track of a contiguous substring.
- Maintain a frequency count of letters in the current window.
- Keep track of the count of the most frequent character in the window.
- If the current window size minus the maximum frequency is greater than k, shrink the window.
- The maximum window size seen during the process is the answer.

Round 2 · High · Medium

p.592

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, since the window traverses the string in linear time.

Space Complexity: $O(1)$ because the frequency count for letters (A–Z) has a constant size.

Solution

The problem is solved using the sliding window technique. We initialize two pointers representing the left and right bounds of the current window. As we expand the window by moving the right pointer, we update the frequency count of the letter at the right pointer and track the letter that appears most frequently. The key observation is that the number of characters that need to be replaced in the current window is the window's size minus the count of the most frequently occurring letter. If this number exceeds k , we increment the left pointer to shrink the window until the condition is satisfied again. The maximum size of the window during the process gives the answer.

Sliding Window

#438 Find All Anagrams in a String

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Grind 169

Problem

Description

Given two strings s and p , return an array of all the start indices of p 's anagrams in s . You may return the answer in any order.

Example 1:

Input: $s = \text{"cbaebabacd"}, p = \text{"abc"}$

Output: $[0,6]$

Explanation:

The substring with start index = 0 is "cba", which is an anagram of "abc".

The substring with start index = 6 is "bac", which is an anagram of "abc".

Example 2:

Input: s = "abab", p = "ab"

Output: [0,1,2]

Explanation:

The substring with start index = 0 is "ab", which is an anagram of "ab".

The substring with start index = 1 is "ba", which is an anagram of "ab".

The substring with start index = 2 is "ab", which is an anagram of "ab".

Constraints:

- $1 \leq s.length, p.length \leq 3 * 10^4$
- s and p consist of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findAnagrams(self, s: str, p: str) ->
    list[int]:
        ans = []
        count = collections.Counter(p)
        required = len(p)

        for r, c in enumerate(s):
            count[c] -= 1
            if count[c] >= 0:
                required -= 1
```

```
if r >= len(p):
    count[s[r - len(p)]] += 1
    if count[s[r - len(p)]] > 0:
        required += 1
    if required == 0:
        ans.append(r - len(p) + 1)

return ans
```

• LeetCode

```
class Solution:
    def findAnagrams(self, s: str, p: str) ->
    List[int]:
        m, n = len(s), len(p)
        ans = []
        if m < n:
            return ans
        cnt1 = Counter(p)
        cnt2 = Counter(s[: n - 1])
        for i in range(n - 1, m):
            cnt2[s[i]] += 1

            if cnt1 == cnt2:
                ans.append(i - n + 1)
            cnt2[s[i - n + 1]] -= 1
        return ans
```

• SimplyLeet

```

# Define the function to find all anagram indices
in the string
def findAnagrams(s: str, p: str):
    # Import Counter for frequency counting
    from collections import Counter
    # List to store the starting indices of all
    anagrams
    result = []
    # Length of string p
    p_length = len(p)
    # Frequency count for characters in p
    p_count = Counter(p)

    # Frequency count for the current window in s
    s_count = Counter()

    # Iterate over the string s
    for i in range(len(s)):
        # Add current character to the window count
        s_count[s[i]] += 1

        # Remove the leftmost character from window
        if window size exceeds p_length
        if i >= p_length:

```

```

        if s_count[s[i - p_length]] == 1:
            del s_count[s[i - p_length]] #
Remove key completely if count becomes 0
        else:
            s_count[s[i - p_length]] -= 1

        # Compare window count with p_count when
        window size is equal to p_length
        if s_count == p_count:
            result.append(i - p_length + 1)

    return result

```

```

# Example usage:
print(findAnagrams("cbaebabacd", "abc")) #
Expected output: [0, 6]

```

• LC.ca

```

class Solution:
    def findAnagrams(self, s: str, p: str) ->
    List[int]:
        m, n = len(s), len(p)
        ans = []
        if m < n:
            return ans
        cnt1 = Counter(p)
        cnt2 = Counter(s[: n - 1])
        for i in range(n - 1, m):
            cnt2[s[i]] += 1

```

```

        if cnt1 == cnt2:
            ans.append(i - n + 1)
            cnt2[s[i - n + 1]] -= 1
        return ans

```

◆ Quick Review

Key Insights

- Use a sliding window of length equal to p over s .
- Maintain frequency counts of characters in p and within the current window in s .
- Compare frequency counts to determine if the current window is an anagram.
- Update the window efficiently by adding one character and removing the character that goes out of the window.

Space & Time

Time Complexity: $O(n)$ where n is the length of s .

Space Complexity: $O(1)$ since we use fixed-size frequency arrays or hash maps.

Solution

We use the sliding window technique to iterate over s with a window size equal to the length of p . First, count the frequencies of characters in p . Then, iterate over s and update a frequency count for the current window. When the window size equals the length of p , compare the two frequency counts. If they match, record the start index. To slide the window efficiently, increment the count for the incoming character and decrement the count for the

outgoing character. This way, we achieve an $O(n)$ solution that is optimal for the problem constraints.

Sliding Window

#487 Max Consecutive Ones II

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Sliding Window

Lists: Premium 98

Problem

Description

Given a binary array `nums`, return the maximum number of consecutive 1's in the array if you can flip at most one 0.

Example 1:

Input: `nums = [1,0,1,1,0]`

Output: 4

Explanation:

- If we flip the first zero, `nums` becomes `[1,1,1,1,0]` and we have 4 consecutive ones.
 - If we flip the second zero, `nums` becomes `[1,0,1,1,1]` and we have 3 consecutive ones.
- The max number of consecutive ones is 4.

Example 2:

Input: `nums = [1,0,1,1,0,1]`

Output: 4

Explanation:

- If we flip the first zero, `nums` becomes `[1,1,1,1,0,1]` and we have 4 consecutive ones.
 - If we flip the second zero, `nums` becomes `[1,0,1,1,1,1]` and we have 4 consecutive ones.
- The max number of consecutive ones is 4.

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- `nums[i]` is either 0 or 1.

Follow up: What if the input numbers come in one by one as an infinite stream? In other words, you can't store all numbers coming from the stream as it's too large to hold in memory. Could you solve it efficiently?

Community Solutions (Python)

- WalkCC

```

class Solution:
    def findMaxConsecutiveOnes(self, nums:
list[int]) -> int:
    ans = 0
    zeros = 0

    l = 0
    for r, num in enumerate(nums):
        if num == 0:
            zeros += 1
            while zeros == 2:

                if nums[l] == 0:
                    zeros -= 1
                l += 1
            ans = max(ans, r - l + 1)

    return ans

```

Round 2 · High · Medium

p.603

```

class Solution:
    def findMaxConsecutiveOnes(self, nums:
list[int]) -> int:
    maxZeros = 1
    ans = 0
    q = collections.deque() # Store indices of
zero.

    l = 0
    for r, num in enumerate(nums):
        if num == 0:
            q.append(r)

            if len(q) > maxZeros:
                l = q.popleft() + 1
            ans = max(ans, r - l + 1)

    return ans

```

Round 2 · High · Medium

p.604

```

class Solution:
    def findMaxConsecutiveOnes(self, nums:
list[int]) -> int:
    ans = 0
    lastZeroIndex = -1

    l = 0
    for r, num in enumerate(nums):
        if num == 0:
            l = lastZeroIndex + 1
            lastZeroIndex = r

```

```
ans = max(ans, r - l + 1)

return ans
```

- LeetCode

```
class Solution:
    def findMaxConsecutiveOnes(self, nums:
List[int]) -> int:
    l = cnt = 0
    for x in nums:
        cnt += x ^ 1
        if cnt > 1:
            cnt -= nums[l] ^ 1
            l += 1
    return len(nums) - l
```

- SimplyLeet

```
# Function to find the maximum consecutive ones
with at most one zero flip
def findMaxConsecutiveOnes(nums):
    max_length = 0 # Tracks the maximum length of
the window
    left = 0 # Left pointer for the sliding window
    zero_count = 0 # Counter for zeros in the
current window

    # Iterate with right pointer over the entire
array
    for right in range(len(nums)):
        # If current element is zero, increase the
zero_count
        if nums[right] == 0:

            zero_count += 1

        # If window invalid (more than one zero),
shrink from left
        while zero_count > 1:
            if nums[left] == 0:
                zero_count -= 1
            left += 1

        # Update max_length if current window is
larger
        max_length = max(max_length, right - left +
1)
```

```
return max_length
```

```
# Example usage:
# nums = [1,0,1,1,0]
# print(findMaxConsecutiveOnes(nums)) # Expected
output: 4
```

• LC.ca

```
from collections import deque

class Solution:
    def findMaxConsecutiveOnes(self, nums:
List[int]) -> int:

        res, zero, left, k = 0, 0, 0, 1

        for right in range(len(nums)):
            if nums[right] == 0:
                zero += 1
```

```
        while zero > k:
            if nums[left] == 0:
                zero -= 1
            left += 1
            res = max(res, right - left + 1) #
max-check every for iteration

        return res
```

```
class Solution: # followup, extra space
    def findMaxConsecutiveOnes(self, nums:
List[int]) -> int:
```

```
        res, left, k = 0, 0, 1
        q = deque()

        for right in range(len(nums)):
            if nums[right] == 0:
                q.append(right)
            if len(q) > k:
                left = q.popleft() + 1
            res = max(res, right - left + 1) #
max-check every for iteration
```

```
return res
```

```
#####
```

```
# maintain a sliding window of size 2, that keeps  
track of the counts.
```

```
class Solution: # follow up, optimized
```

```
    def findMaxConsecutiveOnes(self, nums:  
List[int]) -> int:  
        count = 0 # Count of consecutive ones  
        prev_count = 0 # Count of consecutive ones  
        before a zero
```

```
        max_count = 0 # Maximum count of  
consecutive ones
```

```
        for num in nums:  
            if num == 1:  
                count += 1  
                prev_count += 1  
            else:  
                count = prev_count + 1  
                prev_count = 0
```

```
max_count = max(max_count, count)
```

```
# Reset the count and prev_count if we  
encounter two zeros in a row
```

```
    if num == 0 and prev_count == 0:  
        count = 0  
        prev_count = 0
```

```
return max_count
```

```
#####
```

```
class Solution:
```

```
    def findMaxConsecutiveOnes(self, nums:  
List[int]) -> int:  
        l = r = 0  
        k = 1  
        while r < len(nums):  
            if nums[r] == 0:  
                k -= 1  
            if k < 0:  
                if nums[l] == 0:
```

```

        k += 1
        l += 1
        r += 1
    return r - l

```

```
#####
```

```

class Solution(object):
    def __init__(self):
        self.ans = 0

```

```

        self.count = 0
        self.lastCount = 0

```

```

    def findMaxConsecutiveOnes(self, nums):
        """
        :type nums: List[int]
        :rtype: int
        """
        for num in nums:
            self.readNum(num) # stream the input

```

```

    return self.ans

```

```

    def readNum(self, num):

```

```

        """

```

```

        :type nums: int

```

```

        """

```

```

        if num == 1:

```

```

            self.count += 1

```

```

        else:

```

```

            self.count = self.count - self.lastCount + 1

```

```

        self.lastCount = self.count

```

```

        self.ans = max(self.ans, self.count)

```

◆ Quick Review

Key Insights

- Utilize a sliding window to maintain the current segment.
- The window can contain at most one 0.
- Expand the window with the right pointer and contract it with the left pointer when the constraint is violated.
- The maximum window size observed during this process is the desired result.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the array.

Space Complexity: $O(1)$.

Solution

We use a sliding window approach to efficiently find the longest contiguous subarray that contains at most one 0. Two pointers (left and right) define the window. As we iterate with the right pointer, counting zeros in the window, we adjust the left pointer when the count of zeros exceeds one. The length of the valid window gives us the number of consecutive 1's (considering one flipped 0), and we update the maximum length accordingly.

Sliding Window

#1100 Find K-Length Substrings with No Repeated Characters

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Premium 98

Problem

Description

Given a string s and an integer k , return the number of substrings in s of length k with no repeated characters.

Example 1:

Input: $s = \text{"havefunonleetcode"}, k = 5$

Output: 6

Explanation: There are 6 substrings they are:

'havef', 'avefu', 'vefun', 'efuno', 'etcod', 'tcode'.

Example 2:

Input: $s = \text{"home"}, k = 5$

Output: 0

Explanation: Notice k can be larger than the length of s . In this case, it is not possible to find any substring.

Constraints:

- $1 \leq s.length \leq 104$

- s consists of lowercase English letters.
- $1 \leq k \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numKLenSubstrNoRepeats(self, s: str, k: int)
    -> int:
        ans = 0
        unique = 0
        count = collections.Counter()

        for i, c in enumerate(s):
            count[c] += 1
            if count[c] == 1:
                unique += 1

            if i >= k:
                count[s[i - k]] -= 1
                if count[s[i - k]] == 0:
                    unique -= 1
                if unique == k:
                    ans += 1

        return ans
```

• LeetDoocs

```
class Solution:
    def numKLenSubstrNoRepeats(self, s: str, k:
    int) -> int:
        cnt = Counter(s[:k])
        ans = int(len(cnt) == k)
        for i in range(k, len(s)):
            cnt[s[i]] += 1
            cnt[s[i - k]] -= 1
            if cnt[s[i - k]] == 0:
                cnt.pop(s[i - k])
            ans += int(len(cnt) == k)

        return ans
```

• SimplyLeet

```
def numKLenSubstrNoRepeats(s: str, k: int) -> int:
    # If k is larger than the string length, no
    valid substring exists.
    if k > len(s):
        return 0

    count = 0
    left = 0
    # Dictionary to store the frequency of
    characters in the current window.
    char_count = {}
```



```
# Iterate over the string using right pointer.
for right in range(len(s)):
    # Update frequency count for the current
    character.
    char_count[s[right]] =
    char_count.get(s[right], 0) + 1

    # If current character frequency > 1, slide
    window from the left to remove duplicates.
    while char_count[s[right]] > 1:
        char_count[s[left]] -= 1
        left += 1

    # When window size equals k, we have a
    valid substring.
    if right - left + 1 == k:
        count += 1
        # Slide the window by removing the
        leftmost character.
        char_count[s[left]] -= 1
        left += 1
    return count

# Example usage:
#
print(numKLenSubstrNoRepeats("havefunonleetcode",
5)) -> 6
```

• LC.ca

Round 2 · High · Medium

p.617

```
class Solution:
    def numKLenSubstrNoRepeats(self, s: str, k:
int) -> int:
        n = len(s)
        if k > n or k > 26:
            return 0
        ans = j = 0
        cnt = Counter()
        for i, c in enumerate(s):
            cnt[c] += 1
            while cnt[c] > 1 or i - j + 1 > k:
                cnt[s[j]] -= 1
                j += 1
            ans += i - j + 1 == k
        return ans
```

◆ Quick Review

Key Insights

- Use a sliding window of fixed size k.
- Track character counts within the current window using a hash table (or frequency array).
- When a duplicate is found in the window, slide the window from the left to remove duplicates.
- Count a valid window when its size equals k and all characters are unique.
- Optimize by adjusting the window immediately after counting a valid substring.

Round 2 · High · Medium

p.618

Space & Time

Time Complexity: $O(n)$ where n is the length of s , as each character is visited at most twice.

Space Complexity: $O(1)$ since the frequency tracking is over a fixed set (26 lowercase letters).

Solution

We use a sliding window approach with two pointers: left and right. A hash table (or frequency array) tracks the occurrences of characters in the window. As we move the right pointer, we update the frequency and check for duplicates. If a duplicate is encountered, we increment the left pointer until the duplicate is removed. Once the window size exactly equals k , we have a valid substring, so we increment our count and then slide the window forward. This ensures that every possible k -length substring is examined once.

Sliding Window

#1248 Count Number of Nice Subarrays

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Math · Sliding Window · Prefix Sum
Lists: CodeSignal

Problem

Description

Given an array of integers `nums` and an integer k . A continuous subarray is called nice if there are k odd numbers on it.

Return the number of nice sub-arrays.

Example 1:

Input: `nums = [1,1,2,1,1]`, $k = 3$
Output: 2
Explanation: The only sub-arrays with 3 odd numbers are `[1,1,2,1]` and `[1,2,1,1]`.

Example 2:

Input: `nums = [2,4,6]`, $k = 1$
Output: 0
Explanation: There are no odd numbers in the array.

Example 3:

Input: nums = [2,2,2,1,2,2,1,2,2,2], k = 2
Output: 16

Constraints:

- $1 \leq \text{nums.length} \leq 50000$
- $1 \leq \text{nums}[i] \leq 10^5$
- $1 \leq k \leq \text{nums.length}$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numberOfSubarrays(self, nums: list[int], k:
int) -> int:
        def numberOfSubarraysAtMost(k: int) -> int:
            ans = 0
            l = 0
            r = 0

            while r <= len(nums):
                if k >= 0:
                    ans += r - l
```

```
            if r == len(nums):
                break
            if nums[r] & 1:
                k -= 1
                r += 1
            else:
                if nums[l] & 1:
                    k += 1
                    l += 1
            return ans
```

```
        return numberOfSubarraysAtMost(k) -
        numberOfSubarraysAtMost(k - 1)
```

• LeetCode

```
class Solution:
    def numberOfSubarrays(self, nums: List[int],
k: int) -> int:
        cnt = Counter({0: 1})
        ans = t = 0
        for v in nums:
            t += v & 1
            ans += cnt[t - k]
            cnt[t] += 1
        return ans
```

• SimplyLeet

```
# Python solution with line-by-line comments

def numberOfSubarrays(nums, k):
    # Initialize dictionary to store frequency of
    # prefix sums
    prefix_count = {0: 1} # There is one way to
    # have zero odd numbers initially.
    count = 0 # This will store the final count of
    # nice subarrays
    odd_sum = 0 # Prefix sum representing the count
    # of odd numbers so far

    # Iterate through each number in the array
    for num in nums:

        # If the number is odd, increment the
        # odd_count prefix sum
        if num % 2 == 1:
            odd_sum += 1

        # If there is a prefix such that (current
        # odd_sum - prefix) equals k,
        # then we found subarrays ending here with
        # k odd numbers
        count += prefix_count.get(odd_sum - k, 0)

        # Update the frequency for the current
        # odd_sum in the hash table
        prefix_count[odd_sum] =
        prefix_count.get(odd_sum, 0) + 1
```

Round 2 · High · Medium

p.623

```
# Return the total number of nice subarrays
return count

# Example usage:
# nums = [1,1,2,1,1]
# k = 3
# print(numberOfSubarrays(nums, k)) # Expected
# output is 2
```

• LC.ca

```
class Solution:
    def numberOfSubarrays(self, nums: List[int],
    k: int) -> int:
        cnt = Counter({0: 1})
        ans = t = 0
        for v in nums:
            t += v & 1
            ans += cnt[t - k]
            cnt[t] += 1
        return ans
```

◆ Quick Review

Key Insights

- Convert the problem to counting subarrays with a specific "prefix sum" value where the sum represents the count of odd numbers.

Round 2 · High · Medium

p.624

- Use a hash table (or dictionary) to store the frequency of prefix sums encountered.
- The number of nice subarrays ending at the current index equals the frequency of (current prefix sum – k) seen so far.
- This approach allows you to find the answer in one pass over the array.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in `nums`, because we traverse the array once.

Space Complexity: $O(n)$, in the worst case where each prefix sum is unique and stored in the hash table.

Solution

We solve the problem using a prefix sum approach. First, we iterate through the array and convert it into a prefix sum representing the cumulative count of odd numbers encountered. As we compute the prefix sums, we use a hash table to record how many times each prefix sum appears. For every current prefix sum value, we check how many times we have seen a prefix sum equal to (current prefix sum – k). The idea is that if `prefix[j] – prefix[i]` equals `k`, then between indices `i+1` and `j`, there are exactly `k` odd numbers. This gives us the count of valid subarrays ending at the current index.

This method works efficiently as it requires only a single pass through the array and maintains a running frequency count using the hash table.

Round 2 · High · Medium

p.625

Sorting

Round 2 · High · Medium · 3 questions

Round 2 · High · Medium

p.626

Sorting

#49 Group Anagrams

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Sorting

Lists: Grind 169

Problem

Description

Given an array of strings `strs`, group the anagrams together. You can return the answer in any order.

Example 1:

Input: `strs = ["eat","tea","tan","ate","nat","bat"]`

Output: `[["bat"],["nat","tan"],["ate","eat","tea"]]`

Explanation:

- There is no string in `strs` that can be rearranged to form "bat".
- The strings "nat" and "tan" are anagrams as they can be rearranged to form each other.
- The strings "ate", "eat", and "tea" are anagrams as they can be rearranged to form each other.

Example 2:

Input: `strs = [""]`

Output: `[[""]]`

Example 3:

Input: `strs = ["a"]`

Output: `[["a"]]`

Constraints:

- `1 <= strs.length <= 104`
- `0 <= strs[i].length <= 100`
- `strs[i]` consists of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def groupAnagrams(self, strs: list[str]) ->
list[list[str]]:
    dict = collections.defaultdict(list)

    for str in strs:
        key = ''.join(sorted(str))
        dict[key].append(str)

    return dict.values()
```

• LeetDoocs

```
class Solution:
    def groupAnagrams(self, strs: List[str]) ->
List[List[str]]:
    d = defaultdict(list)
    for s in strs:
        k = ''.join(sorted(s))
        d[k].append(s)
    return list(d.values())
```

• SimplyLeet

Python solution for grouping anagrams
from collections import defaultdict

```
def groupAnagrams(strs):
    # Create a dictionary to hold lists of
    anagrams; default to an empty list for new keys
    anagrams = defaultdict(list)
    # Iterate over each string in the input list
    for s in strs:
        # Sort the string to use as a key; sorted
        returns a list so join back to string
        key = ''.join(sorted(s))
```

```
        # Append the original string under the
        sorted key
        anagrams[key].append(s)
    # Return the grouped anagrams as a list of
    lists
    return list(anagrams.values())
```

Example usage:

```
if __name__ == "__main__":
    strs = ["eat", "tea", "tan", "ate", "nat", "bat"]
    result = groupAnagrams(strs)
    print(result)
```

• LC.ca

```
'''
>>> from collections import defaultdict
>>> d = defaultdict(list)
>>> d["a"]=1
>>> d["b"]=2
>>> d
defaultdict(<class 'list'>, {'a': 1, 'b': 2})
>>> d.values()
dict_values([1, 2])
>>> list(d.values())
```

```
[1, 2]
```

```
### sorted(str) will return a list of chars
```

```
>>> a = "sdfddxyz"
```

```
>>> sorted(a)
```

```
['d', 'd', 'd', 'f', 's', 'x', 'y', 'z']
```

```
>>> "".join(sorted(a))
```

```
'dddfsxyz'
```

```
...
```

```
class Solution:
```

```
    def groupAnagrams(self, strs: List[str]) ->  
    List[List[str]]:
```

```
        d = defaultdict(list)
```

```
        for s in strs:
```

```
            k = "".join(sorted(s))
```

```
            d[k].append(s)
```

```
        return list(d.values())
```

◆ Quick Review

Key Insights

- Anagrams have the same characters when sorted alphabetically.
- Sorting each string provides a unique key for grouping.
- Use a hash table (dictionary) to map each sorted key to a list of its original strings.
- Finally, gather all the lists from the dictionary to form the result.

Space & Time

Time Complexity: $O(n * k \log k)$, where n is the number of strings and k is the maximum length of a string (due to sorting each string).

Space Complexity: $O(n * k)$ for storing the grouped strings in the hash table.

Solution

The primary idea is to iterate through each string, sort it to obtain a key that represents its character composition, and then group the original string under this key in a hash table. After processing all strings, the values of the hash table will be the groups of anagrams. This approach efficiently groups anagrams by leveraging sorting and a dictionary for constant-time lookups.

Sorting

#56 Merge Intervals

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Sorting

Lists: Grind 169

Problem

Description

Given an array of intervals where $\text{intervals}[i] = [\text{start}_i, \text{end}_i]$, merge all overlapping intervals, and return an array of the non-overlapping intervals that cover all the intervals in the input.

Example 1:

Input: $\text{intervals} = [[1,3],[2,6],[8,10],[15,18]]$
 Output: $[[1,6],[8,10],[15,18]]$
 Explanation: Since intervals $[1,3]$ and $[2,6]$ overlap, merge them into $[1,6]$.

Example 2:

Input: $\text{intervals} = [[1,4],[4,5]]$
 Output: $[[1,5]]$
 Explanation: Intervals $[1,4]$ and $[4,5]$ are considered overlapping.

Example 3:

Input: $\text{intervals} = [[4,7],[1,4]]$
 Output: $[[1,7]]$
 Explanation: Intervals $[1,4]$ and $[4,7]$ are considered overlapping.

Constraints:

- $1 \leq \text{intervals.length} \leq 104$
- $\text{intervals}[i].\text{length} == 2$
- $0 \leq \text{start}_i \leq \text{end}_i \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def merge(self, intervals: list[list[int]]) -> list[list[int]]:
        ans = []

        for interval in sorted(intervals):
            if not ans or ans[-1][1] < interval[0]:
                ans.append(interval)
            else:
                ans[-1][1] = max(ans[-1][1], interval[1])

        return ans
```

• LeetDoocs

```

class Solution:
    def merge(self, intervals: List[List[int]]) -> List[List[int]]:
        intervals.sort()
        ans = []
        st, ed = intervals[0]
        for s, e in intervals[1:]:
            if ed < s:
                ans.append([st, ed])
                st, ed = s, e
            else:
                ed = max(ed, e)
        ans.append([st, ed])
        return ans

```

• SimplyLeet

```

def merge(intervals):
    # Edge case: if the list is empty, return an empty list
    if not intervals:
        return []
    # Sort intervals based on the starting time
    intervals.sort(key=lambda x: x[0])
    merged = [] # List to store merged intervals

    # Iterate over each interval in the sorted list
    for interval in intervals:

```

```

        # If merged list is not empty and the current interval
        # overlaps with the last one in merged, merge them
        if merged and interval[0] <= merged[-1][1]:
            # Update the end of the last interval in merged if needed
            merged[-1][1] = max(merged[-1][1], interval[1])
        else:
            # No overlap, so add the current interval to merged
            merged.append(interval)
    return merged

```

```

# Example usage:
print(merge([[1,3],[2,6],[8,10],[15,18]]))

```

• LC.ca

```

# Definition for an interval.
# class Interval(object):
# def __init__(self, s=0, e=0):
# self.start = s
# self.end = e

class Solution(object):
    def merge(self, intervals):
        """
        :type intervals: List[Interval]

```

```

:rtype: List[Interval]
"""
ans = []
for intv in sorted(intervals, key=lambda x:
x.start):
    if ans and ans[-1].end >= intv.start:
        ans[-1].end = max(ans[-1].end, intv.end)
    else:
        ans.append(intv)
return ans

```

#####

```

...
>>> intervals =
[[111,222],[1,3],[2,6],[8,10],[15,18]]
>>> intervals.sort()
>>> intervals
[[1, 3], [2, 6], [8, 10], [15, 18], [111, 222]]
...
class Solution:
    def merge(self, intervals: List[List[int]]) ->
List[List[int]]:

```

```

intervals.sort() # default sort also ok
ans = [intervals[0]]
for s, e in intervals[1:]:
    if ans[-1][1] < s:
        ans.append([s, e])
    else:
        ans[-1][1] = max(ans[-1][1], e)
return ans

```

◆ Quick Review

Key Insights

- Sort the intervals based on their start values.
- Use a pointer or iteration to compare the current interval with the last merged interval.
- If the current interval overlaps with the previous one, merge them by updating the end time.
- Otherwise, add the current interval to the results as a new merged interval.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(n)$ to store the merged intervals.

Solution

First, sort the array of intervals based on the starting time. Then, iterate through the sorted list and check if the current interval overlaps with the last interval in the merged list. If an overlap is detected (i.e., the start of the current interval is less than or equal to the end of the last

merged interval), update the last merged interval's end to be the maximum of both ends. If there is no overlap, simply add the current interval to the merged list. This approach efficiently merges overlapping intervals and ensures that all intervals are processed in a single pass after sorting.

Round 2 · High · Medium

p.639

Sorting

#1152 Analyze User Website Visit Pattern

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Sorting

Lists: Premium 98

Problem

Description

You are given two string arrays `username` and `website` and an integer array `timestamp`. All the given arrays are of the same length and the tuple `[username[i], website[i], timestamp[i]]` indicates that the user `username[i]` visited the website `website[i]` at time `timestamp[i]`.

A pattern is a list of three websites (not necessarily distinct).

- For example, `["home", "away", "love"]`, `["leetcode", "love", "leetcode"]`, and `["luffy", "luffy", "luffy"]` are all patterns.

The score of a pattern is the number of users that visited all the websites in the pattern in the same order they appeared in the pattern.

- For example, if the pattern is `["home", "away", "love"]`, the score is the number of users `x` such that `x` visited

Round 2 · High · Medium

p.640

"home" then visited "away" and visited "love" after that.

- Similarly, if the pattern is ["leetcode", "love", "leetcode"], the score is the number of users x such that x visited "leetcode" then visited "love" and visited "leetcode" one more time after that.

- Also, if the pattern is ["luffy", "luffy", "luffy"], the score is the number of users x such that x visited "luffy" three different times at different timestamps.

Return the pattern with the largest score. If there is more than one pattern with the same largest score, return the lexicographically smallest such pattern.

Note that the websites in a pattern do not need to be visited contiguously, they only need to be visited in the order they appeared in the pattern.

Example 1:

Input: username = ["joe","joe","joe","james","james","james","james","mary","mary","mary"], timestamp = [1,2,3,4,5,6,7,8,9,10], website = ["home","about","career","home","cart","maps","home","home","about","career"]

Output: ["home","about","career"]

Explanation: The tuples in this example are:

["joe","home",1],["joe","about",2],["joe","career",3],["james","home",4],["james","cart",5],["james","maps",6],["james","home",7],["mary","home",8],["mary","about",9], and ["mary","career",10].

The pattern ("home", "about", "career") has score 2 (joe and mary).

The pattern ("home", "cart", "maps") has score 1 (james).

The pattern ("home", "cart", "home") has score 1 (james).

The pattern ("home", "maps", "home") has score 1 (james).

Example 2:

Input: username = ["ua","ua","ua","ub","ub","ub"], timestamp = [1,2,3,4,5,6], website = ["a","b","a","a","b","c"]

Output: ["a","b","a"]

Constraints:

- $3 \leq \text{username.length} \leq 50$
- $1 \leq \text{username}[i].\text{length} \leq 10$
- $\text{timestamp.length} == \text{username.length}$
- $1 \leq \text{timestamp}[i] \leq 109$

- `website.length == username.length`
- `1 <= website[i].length <= 10`
- `username[i]` and `website[i]` consist of lowercase English letters.
- It is guaranteed that there is at least one user who visited at least three websites.
- All the tuples `[username[i], timestamp[i], website[i]]` are unique.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def mostVisitedPattern(
        self,
        username: list[str],
        timestamp: list[int],
        website: list[str],
    ) -> list[str]:
        userToSites = collections.defaultdict(list)

        # Sort websites of each user by timestamp.
```

```
        for user, _, site in sorted(
            zip(username, timestamp, website),
            key=lambda x: x[1]):
            userToSites[user].append(site)

        # For each of three websites, count its
        frequency.
        patternCount = collections.Counter()

        for user, sites in userToSites.items():
            patternCount.update(Counter(set(itertools.co
mbinations(sites, 3))))

        return max(sorted(patternCount),
key=patternCount.get)
```

• LeetCode

```
class Solution:
    def mostVisitedPattern(
        self, username: List[str], timestamp:
List[int], website: List[str]
    ) -> List[str]:
        d = defaultdict(list)
        for user, _, site in sorted(
            zip(username, timestamp, website),
            key=lambda x: x[1]
        ):
            d[user].append(site)
```

```

cnt = Counter()
for sites in d.values():
    m = len(sites)
    s = set()
    if m > 2:
        for i in range(m - 2):
            for j in range(i + 1, m - 1):
                for k in range(j + 1, m):
                    s.add((sites[i],
sites[j], sites[k]))
                    for t in s:
                        cnt[t] += 1
return sorted(cnt.items(), key=lambda x:
(-x[1], x[0]))[0][0]

```

- SimplyLeet

Python code solution for Analyze User Website Visit Pattern

```

from collections import defaultdict
from itertools import combinations

def mostVisitedPattern(username, timestamp,
website):
    # Combine the logs into a list of tuples and
    sort by timestamp.
    logs = sorted(zip(timestamp, username,
website))

    # Dictionary to hold each user's list of
    website visits

```

```

user_visits = defaultdict(list)
for time, user, site in logs:
    user_visits[user].append(site)

    # Dictionary to map a pattern (tuple of 3
    websites) to set of users who visited that pattern
    pattern_users = defaultdict(set)

    # For each user, generate all unique
    3-sequences from their visit list and update
    pattern_users
    for user, sites in user_visits.items():
        # Use set to avoid counting duplicate
        patterns from the same user multiple times

```

```

seen_patterns = set()
if len(sites) < 3:
    continue
# Generate all 3-sequence combinations
(order matters since visits are in time order)
for pattern in combinations(sites, 3):
    if pattern not in seen_patterns:
        seen_patterns.add(pattern)
        pattern_users[pattern].add(user)

```

```

# Find the pattern with the highest score
(i.e., most users)

```

```

# In case of a tie, lexicographically smallest
pattern wins.
max_count = 0
result = None
for pattern, users in pattern_users.items():
    count = len(users)
    if count > max_count or (count == max_count
and (result is None or pattern < result)):
        max_count = count
        result = pattern

return list(result)

```

```

# Example usage:
print(mostVisitedPattern(
    ["joe","joe","joe","james","james","james","ja
mes","mary","mary","mary"],
    [1,2,3,4,5,6,7,8,9,10],
    ["home","about","career","home","cart","maps",
"home","home","about","career"]
))

```

- LC.ca

```

class Solution:
    def mostVisitedPattern(
        self, username: List[str], timestamp:
List[int], website: List[str]
    ) -> List[str]:
        d = defaultdict(list)
        for user, _, site in sorted(
            zip(username, timestamp, website),
            key=lambda x: x[1]
        ):
            d[user].append(site)

```



```

cnt = Counter()
for sites in d.values():
    m = len(sites)
    s = set()
    if m > 2:
        for i in range(m - 2):
            for j in range(i + 1, m - 1):
                for k in range(j + 1, m):
                    s.add((sites[i],
sites[j], sites[k]))
                    for t in s:
                        cnt[t] += 1
return sorted(cnt.items(), key=lambda x:
(-x[1], x[0]))[0][0]

```

◆ Quick Review

Key Insights

- Combine the inputs into a single list of events and sort them by timestamp.
- Group the visits by each user, preserving the sorted order.
- Generate all unique combinations of 3-sequences for each user (order matters and non-contiguous visits are allowed).
- Use a dictionary (or hash map) to map each 3-sequence pattern to the set of users who visited that pattern.

- The problem constraints are small, so generating combinations per user (even using triple nested loops) is acceptable.
- When multiple patterns have the same frequency, choose the lexicographically smallest one.

Space & Time

Time Complexity: $O(U * L^3)$ where U is the number of users and L is the maximum length of visits per user (since we generate combinations of 3 visits per user). Sorting all events costs $O(N \log N)$ for N total events.

Space Complexity: $O(P)$ where P is the number of unique 3-sequence patterns stored in the dictionary, plus $O(N)$ for grouping the events.

Solution

- Aggregate the events by user while sorting them based on timestamp.
- For each user, generate all sets of 3-website patterns (to avoid duplicate patterns per user).
- Use a hash map to count the number of unique users for each pattern.
- Traverse the map to find the pattern with the highest count. In case of a tie, compare patterns lexicographically.
- Return the pattern that meets the criteria.

Binary Search

Round 2 · High · Medium · 12 questions

Round 2 · High · Medium

p.651

Binary Search

#33 Search in Rotated Sorted Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Grind 169

Problem

Description

There is an integer array `nums` sorted in ascending order (with distinct values).

Prior to being passed to your function, `nums` is possibly left rotated at an unknown index k ($1 \leq k < \text{nums.length}$) such that the resulting array is `[nums[k], nums[k+1], ..., nums[n-1], nums[0], nums[1], ..., nums[k-1]]` (0-indexed). For example, `[0,1,2,4,5,6,7]` might be left rotated by 3 indices and become `[4,5,6,7,0,1,2]`.

Given the array `nums` after the possible rotation and an integer `target`, return the index of `target` if it is in `nums`, or `-1` if it is not in `nums`.

You must write an algorithm with $O(\log n)$ runtime complexity.

Example 1:

Round 2 · High · Medium

p.652

Input: nums = [4,5,6,7,0,1,2], target = 0
Output: 4

Example 2:

Input: nums = [4,5,6,7,0,1,2], target = 3
Output: -1

Example 3:

Input: nums = [1], target = 0
Output: -1

Constraints:

- $1 \leq \text{nums.length} \leq 5000$
- $-104 \leq \text{nums}[i] \leq 104$
- All values of nums are unique.
- nums is an ascending array that is possibly rotated.
- $-104 \leq \text{target} \leq 104$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def search(self, nums: list[int], target: int) -> int:
        l = 0
        r = len(nums) - 1

        while l <= r:
            m = (l + r) // 2
            if nums[m] == target:
                return m
            if nums[l] <= nums[m]: # nums[l..m] are sorted.
```

```
                if nums[l] <= target < nums[m]:
                    r = m - 1
                else:
                    l = m + 1
            else: # nums[m..n - 1] are sorted.
                if nums[m] < target <= nums[r]:
                    l = m + 1
                else:
                    r = m - 1
```

```
        return -1
```

- LeetCode

```

class Solution:
    def search(self, nums: List[int], target: int)
-> int:
    n = len(nums)
    left, right = 0, n - 1
    while left < right:
        mid = (left + right) >> 1
        if nums[0] <= nums[mid]:
            if nums[0] <= target <= nums[mid]:
                right = mid
            else:
                left = mid + 1
        else:
            if nums[mid] < target <= nums[n -
1]:
                left = mid + 1
            else:
                right = mid
    return left if nums[left] == target else -1

```

• SimplyLeet

Round 2 · High · Medium

p.655

```

# Define a function to search for the target in a
rotated sorted array.
def search(nums, target):
    # Initialize pointers for binary search.
    left, right = 0, len(nums) - 1

    # Loop until the search space is exhausted.
    while left <= right:
        # Compute the middle index.
        mid = (left + right) // 2

        # If the mid element is the target, return
its index.
        if nums[mid] == target:
            return mid

        # Check if the left part of the array is
sorted.
        if nums[left] <= nums[mid]:
            # If the target is within the sorted
left half, adjust the right pointer.
            if nums[left] <= target < nums[mid]:
                right = mid - 1
            else:

```

Round 2 · High · Medium

p.656

```

        # Otherwise, search the right
half.
        left = mid + 1
    else:
        # The right half must be sorted.
        if nums[mid] < target <= nums[right]:
            left = mid + 1
        else:
            right = mid - 1

    # If the target is not found, return -1.

    return -1

# Example usage:
# result = search([4,5,6,7,0,1,2], 0)
# print(result) # Expected output: 4

```

• LC.ca

Round 2 · High · Medium

p.657

```

# below 2 solutions, diff is while condition: left
('<' or '<=') right

# I like this better
class Solution:
    def search(self, nums: List[int], target: int)
-> int:
        n = len(nums)
        left, right = 0, n - 1
        while left <= right:
            mid = (left + right) >> 1
            if nums[mid] == target:
                return mid
            elif nums[0] <= nums[mid]: # left half
sorted
                if nums[0] <= target <= nums[mid]:
                    right = mid - 1
                else:
                    left = mid + 1
            else: # right half sorted
                if nums[mid] < target <= nums[n -
1]:
                    left = mid + 1
                else:

```

Round 2 · High · Medium

p.658

```

        right = mid - 1
    return -1

#####

class Solution:
    def search(self, nums: List[int], target: int)
-> int:
    n = len(nums)
    left, right = 0, n - 1
    while left < right:

        mid = (left + right) >> 1
        if nums[0] <= nums[mid]: # left half
sorted
            if nums[0] <= target <= nums[mid]:
                right = mid
            else:
                left = mid + 1
        else: # right half sorted
            if nums[mid] < target <= nums[n -
1]:
                left = mid + 1
            else:

                right = mid
    return left if nums[left] == target else -1

```

◆ Quick Review

Round 2 · High · Medium

p.659

Key Insights

- The array is originally sorted and then rotated at an unknown pivot.
- Use a modified binary search to accommodate the rotation.
- At every step, determine which half of the array is properly sorted.
- Narrow down the search based on the sorted half and the target's value.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We perform a binary search while accounting for the possibility that the array is rotated. At each iteration, we check if the left-to-mid portion is sorted. If it is, and the target falls within that range, we search the left half; otherwise, we search the right half. If the left half is not sorted, then the right half must be sorted, and similarly, we check if the target is in that sorted half. This approach ensures logarithmic time complexity while handling the rotation correctly.

Round 2 · High · Medium

p.660

Binary Search

#34 Find First and Last Position of Element in Sorted Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Denny Zhang

Problem

Description

Given an array of integers `nums` sorted in non-decreasing order, find the starting and ending position of a given target value.

If target is not found in the array, return `[-1, -1]`.

You must write an algorithm with $O(\log n)$ runtime complexity.

Example 1:

Input: `nums = [5,7,7,8,8,10]`, `target = 8`
Output: `[3,4]`

Example 2:

Input: `nums = [5,7,7,8,8,10]`, `target = 6`
Output: `[-1,-1]`

Example 3:

Input: `nums = []`, `target = 0`
Output: `[-1,-1]`

Constraints:

- $0 \leq \text{nums.length} \leq 105$
- $-109 \leq \text{nums}[i] \leq 109$
- `nums` is a non-decreasing array.
- $-109 \leq \text{target} \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def searchRange(self, nums: List[int], target:
int) -> List[int]:
        l = bisect_left(nums, target)
        if l == len(nums) or nums[l] != target:
            return -1, -1
        r = bisect_right(nums, target) - 1
        return l, r
```

• LeetDoocs

```
class Solution:
    def searchRange(self, nums: List[int], target:
int) -> List[int]:
        l = bisect_left(nums, target)
        r = bisect_left(nums, target + 1)
        return [-1, -1] if l == r else [l, r - 1]
```

• SimplyLeet

```
def searchRange(nums, target):
    # Helper function to find the boundary index
    (leftmost or rightmost) for target.
    def findBoundary(isLeft):
        left, right = 0, len(nums) - 1
        boundary = -1
        while left <= right:
            mid = (left + right) // 2
            # If the target is found, update
            boundary and continue to search on one side.
            if nums[mid] == target:
                boundary = mid

        if isLeft:
            right = mid - 1 # Continue
            search on left half
        else:
            left = mid + 1 # Continue
            search on right half
        elif nums[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
        return boundary
```

```
# Find the leftmost index
leftIndex = findBoundary(True)
if leftIndex == -1:
    return [-1, -1]
# Find the rightmost index
rightIndex = findBoundary(False)
return [leftIndex, rightIndex]
```

```
# Example usage:
# nums = [5,7,7,8,8,10]
```

```
# target = 8
# print(searchRange(nums, target)) # Output: [3,4]
```

• LC.ca

```
...
>>> bisect.bisect_left([1,1,1,2,2,2,7,7,7], 2)
3
>>> bisect.bisect_left([1,1,1,2,2,2,7,7,7], 2+1)
6

>>> bisect.bisect_right([1,1,1,2,2,2,7,7,7], 2)
6
>>> bisect.bisect_right([1,1,1,2,2,2,7,7,7], 2+1)
```


6

```
>>> bisect.bisect_left([1,1,1,2,2,2,7,7,7], 0)
```

0

```
>>> bisect.bisect_left([1,1,1,2,2,2,7,7,7], 1)
```

0

```
# below, even single value for 2, after +1 the
index will be different
```

```
>>> bisect.bisect_left([1,2,3,4,5], 2)
```

1

```
>>> bisect.bisect_left([1,2,3,4,5], 2+1)
```

2

...

```
import bisect
```

```
class Solution:
```

```
    def searchRange(self, nums: List[int], target:
```

```
int) -> List[int]:
```

```
        l = bisect_left(nums, target)
```

```
        r = bisect_left(nums, target + 1)
        return [-1, -1] if l == r else [l, r - 1]
```

```
#####
```

```
class Solution:
```

```
    def searchRange(self, nums: List[int], target:
int) -> List[int]:
```

```
        def findRange(to_left):
```

```
            l, r = 0, len(nums) - 1
```

```
            m = 0
```

```
        while l <= r:
```

```
            m = l + (r - l) // 2
```

```
            if nums[m] < target:
```

```
                l = m + 1
```

```
            elif nums[m] > target:
```

```
                r = m - 1
```

```
            elif to_left and m > 0 and nums[m -
1] == nums[m]: # so now nums[m] == target, but
maybe not the leftmost or rightmost
```

```
                r = m - 1
```

```
            elif not to_left and m + 1 <
len(nums) and nums[m + 1] == nums[m]: # so now
nums[m] == target, but maybe not the leftmost or
rightmost
```

```
                l = m + 1
```

```

        else:
            return m
        if m < len(nums) and nums[m] == target:
            # 1. cover not-found,
            # or 2. cover input=[3] and target=3,
            first and last both 0, return [0,0]
            return m
        else:
            return -1
    return [findRange(True), findRange(False)]

```

◆ Quick Review

Key Insights

- The array is sorted, allowing the use of binary search.
- Two binary searches can efficiently find the leftmost (first) and rightmost (last) indices of the target.
- Handling edge cases where the target is not present by checking bounds.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

The approach is to use binary search twice:

- The first binary search finds the leftmost occurrence of the target by continuing to search in the left half even after finding the target.
- The second binary search finds the rightmost occurrence by searching in the right half. The

algorithm maintains pointers for the current search bounds and narrows the search space using comparisons. If the target is not found during the first binary search, the solution immediately returns $[-1, -1]$. This method leverages the sorted property of the array, ensuring $O(\log n)$ time complexity with constant extra space.

Binary Search

#153 Find Minimum in Rotated Sorted Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Grind 169

Problem

Description

Suppose an array of length n sorted in ascending order is rotated between 1 and n times. For example, the array `nums = [0,1,2,4,5,6,7]` might become:

- `[4,5,6,7,0,1,2]` if it was rotated 4 times.
- `[0,1,2,4,5,6,7]` if it was rotated 7 times.

Notice that rotating an array `[a[0], a[1], a[2], ..., a[n-1]]` 1 time results in the array `[a[n-1], a[0], a[1], a[2], ..., a[n-2]]`.

Given the sorted rotated array `nums` of unique elements, return the minimum element of this array.

You must write an algorithm that runs in $O(\log n)$ time.

Example 1:

Input: `nums = [3,4,5,1,2]`

Output: 1

Explanation: The original array was `[1,2,3,4,5]` rotated 3 times.

Example 2:

Input: `nums = [4,5,6,7,0,1,2]`

Output: 0

Explanation: The original array was `[0,1,2,4,5,6,7]` and it was rotated 4 times.

Example 3:

Input: `nums = [11,13,15,17]`

Output: 11

Explanation: The original array was `[11,13,15,17]` and it was rotated 4 times.

Constraints:

- $n == \text{nums.length}$
- $1 \leq n \leq 5000$
- $-5000 \leq \text{nums}[i] \leq 5000$
- All the integers of `nums` are unique.
- `nums` is sorted and rotated between 1 and n times.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def findMin(self, nums: list[int]) -> int:
        l = 0
        r = len(nums) - 1

        while l < r:
            m = (l + r) // 2
            if nums[m] < nums[r]:
                r = m
            else:
                l = m + 1

        return nums[l]

```

- LeetCode

```

class Solution:
    def findMin(self, nums: List[int]) -> int:
        l, r = 0, len(nums) - 1
        while l < r:
            mid = (l + r) >> 1
            if nums[mid] > nums[-1]:
                l = mid + 1
            else:
                r = mid
        return nums[l]

```

- SimplyLeet

```

def findMin(nums):
    # Initialize left and right pointers
    left, right = 0, len(nums) - 1
    # Continue while the search space is valid
    while left < right:
        mid = (left + right) // 2 # Compute middle
        index
        # If middle element is greater than the
        rightmost element,
        # the minimum must be in the right half
        if nums[mid] > nums[right]:
            left = mid + 1
        else:
            # Otherwise, the minimum is in the left
            half (including mid)
            right = mid
    # When loop ends, left == right, pointing to
    the minimum element
    return nums[left]

```

- LC.ca

```

class Solution:
    def findMin(self, nums: List[int]) -> int:
        if nums[0] <= nums[-1]:
            return nums[0]
        left, right = 0, len(nums) - 1
        while left < right:
            mid = (left + right) >> 1
            if nums[0] <= nums[mid]:
                left = mid + 1
            else:
                right = mid
        return nums[left]

```

◆ Quick Review

Key Insights

- The rotated sorted array consists of two sorted subarrays.
- A modified binary search can be used to locate the pivot point where the rotation occurs.
- By comparing the middle element to the rightmost element, we can decide which half of the array contains the minimum.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

Use a binary search approach. Initialize two pointers, left and right, at the start and end of the array. While left is less than right, calculate the middle index. If the middle element is greater than the right element, then the minimum is in the right half, so move the left pointer to mid + 1; otherwise, the minimum must be in the left half including mid, so adjust the right pointer to mid. When the loop terminates, left equals right and points to the minimum element. This approach leverages the properties of the rotated sorted array to achieve logarithmic time complexity.

Binary Search

#300 Longest Increasing Subsequence

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode
 Array · Binary Search · Dynamic Programming
 Lists: Grind 169

Problem

Description

Given an integer array `nums`, return the length of the longest strictly increasing subsequence.

Example 1:

Input: `nums = [10,9,2,5,3,7,101,18]`
 Output: 4
 Explanation: The longest increasing subsequence is `[2,3,7,101]`, therefore the length is 4.

Example 2:

Input: `nums = [0,1,0,3,2,3]`
 Output: 4

Example 3:

Input: `nums = [7,7,7,7,7,7,7]`
 Output: 1

Constraints:

- $1 \leq \text{nums.length} \leq 2500$

Round 2 · High · Medium

p.675

- $-104 \leq \text{nums}[i] \leq 104$

Follow up: Can you come up with an algorithm that runs in $O(n \log(n))$ time complexity?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def lengthOfLIS(self, nums: list[int]) -> int:
        if not nums:
            return 0

        # dp[i] := the length of LIS ending in nums[i]
        dp = [1] * len(nums)

        for i in range(1, len(nums)):
            for j in range(i):

                if nums[j] < nums[i]:
                    dp[i] = max(dp[i], dp[j] + 1)

        return max(dp)
```

Round 2 · High · Medium

p.676

```

class Solution:
    def lengthOfLIS(self, nums: list[int]) -> int:
        # tails[i] := the minimum tails of all the
        # increasing subsequences having
        # length i + 1
        tails = []

        for num in nums:
            if not tails or num > tails[-1]:
                tails.append(num)
            else:
                tails[bisect.bisect_left(tails, num)] = num

        return len(tails)

```

• LeetCode

```

class Solution:
    def lengthOfLIS(self, nums: List[int]) -> int:
        n = len(nums)
        f = [1] * n
        for i in range(1, n):
            for j in range(i):
                if nums[j] < nums[i]:
                    f[i] = max(f[i], f[j] + 1)
        return max(f)

```

```

class BinaryIndexedTree:
    def __init__(self, n: int):
        self.n = n
        self.c = [0] * (n + 1)

    def update(self, x: int, v: int):
        while x <= self.n:
            self.c[x] = max(self.c[x], v)
            x += x & -x

    def query(self, x: int) -> int:
        mx = 0
        while x:
            mx = max(mx, self.c[x])
            x -= x & -x
        return mx

```

```

class Solution:
    def lengthOfLIS(self, nums: List[int]) -> int:
        s = sorted(set(nums))
        m = len(s)
        tree = BinaryIndexedTree(m)
        for x in nums:
            x = bisect_left(s, x) + 1
            t = tree.query(x - 1) + 1
            tree.update(x, t)
        return tree.query(m)

```

• SimplyLeet

```
import bisect

def lengthOfLIS(nums):
    dp = [] # dp array stores smallest tail for
    each subsequence length
    for num in nums:
        # Find the leftmost insertion point using
        binary search
        idx = bisect.bisect_left(dp, num)
        if idx == len(dp):
            dp.append(num) # Extend dp if no larger
            tail exists
        else:
```

```
            dp[idx] = num # Replace to maintain the
            smallest tail
    return len(dp)
```

```
# Example usage:
# print(lengthOfLIS([10,9,2,5,3,7,101,18])) #
Output: 4
```

• LC.ca

```
# find the largest end element in tails that is
smaller than nums[i]
# and then replace it with nums[i] and discard the
list in the same length
# which is implemented by `tail[idx] = num`
```

```
class Solution(object):
    def lengthOfLIS(self, nums):
        """
        :type nums: List[int]
        :rtype: int
        """
```

```
        tail = []
        for num in nums:
            # if using bisect_right(tail, num),
            # then input=[7,7,7,7,7,7,7] will
            output 7 but expected result is 1
            idx = bisect.bisect_left(tail, num)
            if idx == len(tail): # same as in java:
            if (i == len) len++;
                tail.append(num)
            else:
                tail[idx] = num
        return len(tail)
```



```

# implementation of bisect.bisect_left()
# similar to Leetcode-302, find
left/right/top/bottom callable
def bisect_left(a, x, lo=0, hi=None):
    if hi is None:
        hi = len(a)

    while lo < hi:
        mid = (lo + hi) // 2

        if a[mid] < x:
            lo = mid + 1
        else:
            hi = mid

    return lo

#####

class BinaryIndexedTree:

```

```

def __init__(self, n: int):
    self.n = n
    self.c = [0] * (n + 1)

def update(self, x: int, v: int):
    while x <= self.n:
        self.c[x] = max(self.c[x], v)
        x += x & -x

def query(self, x: int) -> int:

```

```

mx = 0
while x:
    mx = max(mx, self.c[x])
    x -= x & -x
return mx

```

```

class Solution:
    def lengthOfLIS(self, nums: List[int]) -> int:
        s = sorted(set(nums))

```

```

m = len(s)
tree = BinaryIndexedTree(m)
for x in nums:
    x = bisect_left(s, x) + 1
    t = tree.query(x - 1) + 1
    tree.update(x, t)
return tree.query(m)

```

◆ Quick Review

Key Insights

- A dynamic programming solution with $O(n^2)$ time complexity exists, but we focus on an optimized approach.
- Using binary search and a dp array, we can achieve $O(n \log n)$ time.
- The dp array maintains the smallest possible tail for increasing subsequences of various lengths.
- For each number, perform binary search to decide if it can extend or replace a value in dp.

Space & Time

Time Complexity: $O(n \log n)$ where n is the length of nums

Space Complexity: $O(n)$ for storing the dp array

Solution

We use a dynamic programming approach combined with binary search. The idea is to maintain a dp array where $dp[i]$ holds the smallest tail value for any increasing subsequence with length $i+1$. For each number in the array, we perform a binary search on dp to find its correct insertion position:

- If the number is larger than all elements, we append it, effectively extending the current longest subsequence.
- Otherwise, we replace the first element in dp that is greater than or equal to the current number to keep the tail as small as possible. This process ensures that dp remains sorted and its length gives the length of the longest increasing subsequence.

Binary Search

#362 Design Hit Counter

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Binary Search · Design · Queue

Lists: Grind 169

Problem

Description

Design a hit counter which counts the number of hits received in the past 5 minutes (i.e., the past 300 seconds).

Your system should accept a timestamp parameter (in seconds granularity), and you may assume that calls are being made to the system in chronological order (i.e., timestamp is monotonically increasing). Several hits may arrive roughly at the same time.

Implement the HitCounter class:

- **HitCounter()** Initializes the object of the hit counter system.
- **void hit(int timestamp)** Records a hit that happened at timestamp (in seconds). Several hits may happen at the same timestamp.
- **int getHits(int timestamp)** Returns the number of hits in the past 5 minutes from timestamp (i.e., the past 300 seconds).

Example 1:

```

Input
["HitCounter", "hit", "hit", "hit", "getHits", "hit",
"getHits", "getHits"]
[[], [1], [2], [3], [4], [300], [300], [301]]
Output
[null, null, null, null, 3, null, 4, 3]

Explanation
HitCounter hitCounter = new HitCounter();

```

Constraints:

- $1 \leq \text{timestamp} \leq 2 * 10^9$
- All the calls are being made to the system in chronological order (i.e., timestamp is monotonically increasing).
- At most 300 calls will be made to hit and getHits.

Follow up: What if the number of hits per second could be huge? Does your design scale?

Community Solutions (Python)

- WalkCC

```

class HitCounter:
    def __init__(self):
        self.timestamps = [0] * 300
        self.hits = [0] * 300

    def hit(self, timestamp: int) -> None:
        i = timestamp % 300
        if self.timestamps[i] == timestamp:
            self.hits[i] += 1
        else:
            self.timestamps[i] = timestamp
            self.hits[i] = 1 # Reset the hit count to 1.

    def getHits(self, timestamp: int) -> int:
        return sum(h for t, h in zip(self.timestamps,
                                     self.hits)
                  if timestamp - t < 300)

```

- LeetCode

```
class HitCounter:

    def __init__(self):
        self.ts = []

    def hit(self, timestamp: int) -> None:
        self.ts.append(timestamp)

    def getHits(self, timestamp: int) -> int:
        return len(self.ts) - bisect_left(self.ts,
timestamp - 300 + 1)
```

```
# Your HitCounter object will be instantiated and
called as such:
# obj = HitCounter()
# obj.hit(timestamp)
# param_2 = obj.getHits(timestamp)
```

- SimplyLeet

Python implementation of HitCounter

```
from collections import deque
```

```
class HitCounter:
    def __init__(self):
        # Initialize the deque to store (timestamp,
count) pairs.
        self.hits_queue = deque()

    def hit(self, timestamp: int) -> None:
```

```
        # If the last recorded hit is at the same
timestamp, increment its count.
        if self.hits_queue and
self.hits_queue[-1][0] == timestamp:
            self.hits_queue[-1][1] += 1
        else:
            # Append new (timestamp, count) pair.
            self.hits_queue.append([timestamp, 1])
```

```
    def getHits(self, timestamp: int) -> int:
        # Remove outdated hits from the front of
the queue.
        while self.hits_queue and
self.hits_queue[0][0] <= timestamp - 300:
```

```

        self.hits_queue.popleft()
        # Sum counts of the remaining hit entries.
        return sum(count for _, count in
self.hits_queue)

```

• LC.ca

```

from collections import deque

```

```

class HitCounter: # queue

```

```

    def __init__(self):
        """

```

```

        Initialize your data structure here.
        """

```

```

        self.hits = deque()

```

```

    def hit(self, timestamp: int) -> None:
        """

```

```

        Record a hit.

```

```

        @param timestamp - The current timestamp
        (in seconds granularity).
        """

```

```

        self.hits.append(timestamp)

```

```

    def getHits(self, timestamp: int) -> int:
        """

```

```

        Return the number of hits in the past 5
        minutes.

```

```

        @param timestamp - The current timestamp
        (in seconds granularity).
        """

```

```

        while self.hits and self.hits[0] <=
timestamp - 300:

```

```

            self.hits.popleft()
            return len(self.hits)

```

```

class HitCounter: # follow-up

```

```

    def __init__(self):
        self.times = [0] * 300

```

```

        self.hits = [0] * 300

```

```

    def hit(self, timestamp: int) -> None:

```

```

        idx = timestamp % 300

```

```

        if self.times[idx] != timestamp:

```

```

            self.times[idx] = timestamp

```

```

            self.hits[idx] = 1

```

```

        else:

```

```

            self.hits[idx] += 1

```

```
def getHits(self, timestamp: int) -> int:
    res = 0
    for i in range(300):
        if timestamp - self.times[i] < 300:
            res += self.hits[i]
    return res
```

Your HitCounter object will be instantiated and called as such:

```
# obj = HitCounter()
# obj.hit(timestamp)
```

```
# param_2 = obj.getHits(timestamp)
```

```
#####
```

```
class HitCounter: # extra o(N) space
```

```
    def __init__(self):
        """
```

```
        Initialize your data structure here.
        """
```

```
        self.counter = Counter()
```

```
def hit(self, timestamp: int) -> None:
    """
    Record a hit.
    @param timestamp - The current timestamp
    (in seconds granularity).
    """
    self.counter[timestamp] += 1

def getHits(self, timestamp: int) -> int:
    """
```

Return the number of hits in the past 5 minutes.

```
    @param timestamp - The current timestamp
    (in seconds granularity).
    """
```

```
    return sum([v for t, v in
self.counter.items() if t + 300 > timestamp])
```

Your HitCounter object will be instantiated and called as such:

```
# obj = HitCounter()
# obj.hit(timestamp)
# param_2 = obj.getHits(timestamp)
```

◆ Quick Review

Key Insights

- Use a sliding window of 300 seconds to count hits.
- Since timestamps are monotonically increasing, outdated hits can be removed from the front.
- A queue (or circular array) is an efficient data structure to maintain and remove expired hit records.
- Grouping consecutive hits at the same timestamp can optimize space.

Space & Time

Time Complexity: $O(1)$ per hit and getHits operation on average, since each hit is enqueued and later dequeued only once.

Space Complexity: $O(W)$ where W is the window size (300 seconds) in the worst case.

Solution

The solution uses a queue (or deque) data structure where each element stores a (timestamp, count) pair. When a hit is recorded, if the last element of the queue shares the same timestamp, then its count is incremented; otherwise, a new element is added. When getHits(timestamp) is called, the algorithm removes elements from the front of the queue whose timestamp is out of the 300-second window relative to the current timestamp, and then sums the counts of the remaining elements to return the result. This ensures that only relevant hits within the past 5 minutes are counted. The design scales well when hit frequency is high because hits at the same timestamp are grouped together instead of inserting individual records.

Binary Search

#528 Random Pick with Weight

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Binary Search · Prefix Sum · Randomized

Lists: Grind 169

Problem

Description

You are given a 0-indexed array of positive integers w where $w[i]$ describes the weight of the i th index.

You need to implement the function pickIndex(), which randomly picks an index in the range $[0, w.length - 1]$ (inclusive) and returns it. The probability of picking an index i is $w[i] / \text{sum}(w)$.

- For example, if $w = [1, 3]$, the probability of picking index 0 is $1 / (1 + 3) = 0.25$ (i.e., 25%), and the probability of picking index 1 is $3 / (1 + 3) = 0.75$ (i.e., 75%).

Example 1:

```

Input
["Solution","pickIndex"]
[[[1]],[]]
Output
[null,0]

Explanation
Solution solution = new Solution([1]);

```

Example 2:

```

Input
["Solution","pickIndex","pickIndex","pickIndex","pickIndex","pickIndex"]
[[[1,3]],[],[],[],[],[]]
Output
[null,1,1,1,1,0]

Explanation
Solution solution = new Solution([1, 3]);

```

Constraints:

- $1 \leq w.length \leq 104$
- $1 \leq w[i] \leq 105$
- pickIndex will be called at most 104 times.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def __init__(self, w: list[int]):
        self.prefix = list(itertools.accumulate(w))

    def pickIndex(self) -> int:
        target = random.randint(0, self.prefix[-1] - 1)
        return
        bisect.bisect_right(range(len(self.prefix)),
                           target,
                           key=lambda m:
                           self.prefix[m])

```

• LeetCode

```

class Solution:
    def __init__(self, w: List[int]):
        self.s = [0]
        for c in w:
            self.s.append(self.s[-1] + c)

    def pickIndex(self) -> int:
        x = random.randint(1, self.s[-1])
        left, right = 1, len(self.s) - 1
        while left < right:

```



```

        mid = (left + right) >> 1
        if self.s[mid] >= x:
            right = mid
        else:
            left = mid + 1
    return left - 1

```

Your Solution object will be instantiated and called as such:

```
# obj = Solution(w)
```

```
# param_1 = obj.pickIndex()
```

• SimplyLeet

```
import random
import bisect
```

```
class Solution:
    def __init__(self, w):
        # Build the prefix sum array.
        self.prefix_sums = []
        current_sum = 0
        for weight in w:
            current_sum += weight

```

```

        self.prefix_sums.append(current_sum)
        # Store the total sum of weights.
        self.total_sum = current_sum

    def pickIndex(self):
        # Get a random target in the range [1,
        total_sum]
        target = random.randint(1, self.total_sum)
        # Use bisect_left to find the proper index.
        index =
        bisect.bisect_left(self.prefix_sums, target)
        return index

```

Example usage:

```
# solution = Solution([1, 3])
# print(solution.pickIndex())
```

• LC.ca

```

class Solution:
    def __init__(self, w: List[int]):
        self.s = [0]
        for c in w:
            self.s.append(self.s[-1] + c)

    def pickIndex(self) -> int:
        x = random.randint(1, self.s[-1])
        left, right = 1, len(self.s) - 1
        while left < right:

```

```

        mid = (left + right) >> 1
        if self.s[mid] >= x:
            right = mid
        else:
            left = mid + 1
    return left - 1

```

Your Solution object will be instantiated and called as such:

```
# obj = Solution(w)
```

```
# param_1 = obj.pickIndex()
```

◆ Quick Review

Key Insights

- Build a prefix sum array from the weights.
- Calculate the total sum of the weights.
- Generate a random number in the range [1, total sum] and use binary search to find the corresponding index in the prefix sum array.
- Using binary search ensures efficient selection in $O(\log n)$ time for each call of `pickIndex()`.

Space & Time

Time Complexity: $O(n)$ for constructing the prefix sum array, and $O(\log n)$ per `pickIndex()` call.

Space Complexity: $O(n)$ for storing the prefix sum array.

Solution

The solution leverages a prefix sum array where each entry at index i represents the cumulative sum of weights up to i . After computing the total weight, a random integer is generated between 1 and the total sum (inclusive). A binary search is then used to find the first index where the prefix sum is greater than or equal to the random number. This approach ensures that the index is selected with a probability proportional to its weight.

Binary Search

#852 Peak Index in a Mountain Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Denny Zhang

Problem

Description

You are given an integer mountain array `arr` of length `n` where the values increase to a peak element and then decrease.

Return the index of the peak element.

Your task is to solve it in $O(\log(n))$ time complexity.

Example 1:

Input: `arr = [0,1,0]`

Output: 1

Example 2:

Input: `arr = [0,2,1,0]`

Output: 1

Example 3:

Input: `arr = [0,10,5,2]`

Output: 1

Constraints:

- $3 \leq \text{arr.length} \leq 105$
- $0 \leq \text{arr}[i] \leq 106$
- `arr` is guaranteed to be a mountain array.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def peakIndexInMountainArray(self, arr:
list[int]) -> int:
        l = 0
        r = len(arr) - 1

        while l < r:
            m = (l + r) // 2
            if arr[m] >= arr[m + 1]:
                r = m
            else:

                l = m + 1

        return l
```

• LeetDoocs

```
class Solution:
    def peakIndexInMountainArray(self, arr:
List[int]) -> int:
        left, right = 1, len(arr) - 2
        while left < right:
            mid = (left + right) >> 1
            if arr[mid] > arr[mid + 1]:
                right = mid
            else:
                left = mid + 1
        return left
```

• SimplyLeet

```
def peakIndexInMountainArray(arr):
    # Initialize pointers for the binary search
    left, right = 0, len(arr) - 1
    # Continue while there is more than one element
    in the search space
    while left < right:
        mid = (left + right) // 2 # Find the mid
        index
        # If the sequence is still increasing, move
        the left pointer to mid+1
        if arr[mid] < arr[mid + 1]:
            left = mid + 1
        else:
```

```
        # Otherwise, the peak is at mid or to
        the left of mid
        right = mid
        # When left and right converge, we have the
        peak index
        return left

# Example usage:
# print(peakIndexInMountainArray([0, 1, 0])) #
Output: 1
```

• LC.ca

```
class Solution:
    def peakIndexInMountainArray(self, arr:
List[int]) -> int:
        left, right = 1, len(arr) - 2
        while left < right:
            mid = (left + right) >> 1
            if arr[mid] > arr[mid + 1]:
                right = mid
            else:
                left = mid + 1
        return left
```

◆ Quick Review

Key Insights

- The array is strictly increasing up to a peak and then strictly decreasing.

- A binary search technique is applicable because the array exhibits a single transition point (peak) between increasing and decreasing order.
- By comparing mid and $\text{mid}+1$, we can decide which half of the array to continue the search in.

Space & Time

Time Complexity: $O(\log(n))$ – The binary search approach reduces the search interval by half at each step.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We use a binary search strategy:

- Initialize two pointers, left and right , at the beginning and end of the array.
- In each iteration, calculate mid . If $\text{arr}[\text{mid}]$ is less than $\text{arr}[\text{mid} + 1]$, then the peak is in the right half; otherwise, it is in the left half.
- When left equals right , we have found the peak index. This approach leverages the mountain array properties and ensures an $O(\log(n))$ solution.

Binary Search

#981 Time Based Key-Value Store

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Binary Search · Design

Lists: Grind 169

Problem

Description

Design a time-based key-value data structure that can store multiple values for the same key at different time stamps and retrieve the key's value at a certain timestamp.

Implement the `TimeMap` class:

- `TimeMap()` Initializes the object of the data structure.
- `void set(String key, String value, int timestamp)` Stores the key `key` with the value `value` at the given time `timestamp`.
- `String get(String key, int timestamp)` Returns a value such that `set` was called previously, with `timestamp_prev ≤ timestamp`. If there are multiple such values, it returns the value associated with the largest `timestamp_prev`. If there are no values, it returns `""`.

Example 1:

Input

```
["TimeMap", "set", "get", "get", "set", "get", "get"]
[[], ["foo", "bar", 1], ["foo", 1], ["foo", 3], ["foo", "bar2", 4], ["foo", 4], ["foo", 5]]
```

Output

```
[null, null, "bar", "bar", null, "bar2", "bar2"]
```

Explanation

```
TimeMap timeMap = new TimeMap();
```

Constraints:

- $1 \leq \text{key.length}, \text{value.length} \leq 100$
- key and value consist of lowercase English letters and digits.
- $1 \leq \text{timestamp} \leq 10^7$
- All the timestamps timestamp of set are strictly increasing.
- At most $2 * 10^5$ calls will be made to set and get.

Community Solutions (Python)

- WalkCC

class TimeMap:

```
def __init__(self):
    self.values = collections.defaultdict(list)
    self.timestamps = collections.defaultdict(list)
```

```
def set(self, key: str, value: str, timestamp: int) -> None:
    self.values[key].append(value)
    self.timestamps[key].append(timestamp)
```

```
def get(self, key: str, timestamp: int) -> str:
    if key not in self.timestamps:
        return ''
    i = bisect.bisect(self.timestamps[key], timestamp)
    return self.values[key][i - 1] if i > 0 else ''
```

- LeetCode

```
class TimeMap:
    def __init__(self):
        self.ktv = defaultdict(list)

    def set(self, key: str, value: str, timestamp:
int) -> None:
        self.ktv[key].append((timestamp, value))

    def get(self, key: str, timestamp: int) -> str:
        if key not in self.ktv:
            return ''
```

```
        tv = self.ktv[key]
        i = bisect_right(tv, (timestamp,
chr(127)))
        return tv[i - 1][1] if i else ''
```

```
# Your TimeMap object will be instantiated and
called as such:
# obj = TimeMap()
# obj.set(key,value,timestamp)
# param_2 = obj.get(key,timestamp)
```

• SimplyLeet

```
# TimeMap class in Python with line-by-line
comments.
class TimeMap:
    def __init__(self):
        # Initialize the dictionary to hold key and
list of (timestamp, value) pairs.
        self.store = {}

    def set(self, key: str, value: str, timestamp:
int) -> None:
        # Append the new (timestamp, value) pair
for the given key.
        if key not in self.store:
            self.store[key] = []

        self.store[key].append((timestamp, value))

    def get(self, key: str, timestamp: int) -> str:
        # Return the value for the largest
timestamp <= the given timestamp.
        if key not in self.store:
            return ""

        values = self.store[key]
        left, right = 0, len(values) - 1
```

```
# Binary search for the rightmost index
with timestamp <= given timestamp.
while left <= right:
    mid = (left + right) // 2
    if values[mid][0] <= timestamp:
        left = mid + 1
    else:
        right = mid - 1

# If no valid timestamp exists, right will
be -1.
if right == -1:
```

```
    return ""
    return values[right][1]
```

```
# Example usage:
# timeMap = TimeMap()
# timeMap.set("foo", "bar", 1)
# print(timeMap.get("foo", 1)) # Output: "bar"
# print(timeMap.get("foo", 3)) # Output: "bar"
# timeMap.set("foo", "bar2", 4)
# print(timeMap.get("foo", 4)) # Output: "bar2"

# print(timeMap.get("foo", 5)) # Output: "bar2"
```

• LC.ca

Round 2 · High · Medium

p.711

```
class TimeMap:
    def __init__(self):
        self.ktv = defaultdict(list)

    def set(self, key: str, value: str, timestamp:
int) -> None:
        self.ktv[key].append((timestamp, value))

    def get(self, key: str, timestamp: int) -> str:
        if key not in self.ktv:
            return ''

        tv = self.ktv[key]
        i = bisect_right(tv, (timestamp,
chr(127)))
        return tv[i - 1][1] if i else ''

# Your TimeMap object will be instantiated and
called as such:
# obj = TimeMap()
# obj.set(key,value,timestamp)
# param_2 = obj.get(key,timestamp)
```

◆ Quick Review

Key Insights

- Use a dictionary (hash table) to map each key to a list of (timestamp, value) pairs.

Round 2 · High · Medium

p.712

- Timestamps are strictly increasing when setting the values, which allows efficient appending.
- For getting the value, perform a binary search on the list of timestamps to find the largest timestamp that is $\leq$ the queried timestamp.
- If no valid timestamp exists, return an empty string.

Space & Time

Time Complexity:

- set operation: $O(1)$ per call (amortized).
- get operation: $O(\log N)$ per call, where N is the number of timestamps stored for the key.

Space Complexity: $O(\text{total number of set calls})$, to store all (timestamp, value) pairs.

Solution

We use a hash table to store keys mapping to a list of their timestamps and corresponding values. Since the timestamps are inserted in strictly increasing order, the list remains sorted. For the get operation, a binary search is applied to quickly locate the appropriate value for a given timestamp. This approach minimizes the lookup time and efficiently handles the problem constraints.

Binary Search

#1011 Capacity to Ship Packages Within D Days

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Denny Zhang

Problem

Description

A conveyor belt has packages that must be shipped from one port to another within D days.

The i th package on the conveyor belt has a weight of $weights[i]$. Each day, we load the ship with packages on the conveyor belt (in the order given by $weights$). We may not load more weight than the maximum weight capacity of the ship.

Return the least weight capacity of the ship that will result in all the packages on the conveyor belt being shipped within D days.

Example 1:

Input: weights = [1,2,3,4,5,6,7,8,9,10], days = 5

Output: 15

Explanation: A ship capacity of 15 is the minimum to ship all the packages in 5 days like this:

1st day: 1, 2, 3, 4, 5

2nd day: 6, 7

3rd day: 8

4th day: 9

5th day: 10

Example 2:

Input: weights = [3,2,2,4,1,4], days = 3

Output: 6

Explanation: A ship capacity of 6 is the minimum to ship all the packages in 3 days like this:

1st day: 3, 2

2nd day: 2, 4

3rd day: 1, 4

Example 3:

Input: weights = [1,2,3,1,1], days = 4

Output: 3

Explanation:

1st day: 1

2nd day: 2

3rd day: 3

4th day: 1, 1

Constraints:

- $1 \leq \text{days} \leq \text{weights.length} \leq 5 * 10^4$
- $1 \leq \text{weights}[i] \leq 500$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def shipWithinDays(self, weights: list[int],
days: int) -> int:
        def shipDays(shipCapacity: int) -> int:
            shipDays = 1
            capacity = 0
            for weight in weights:
                if capacity + weight > shipCapacity:
                    shipDays += 1
                    capacity = weight
                else:
                    capacity += weight
            return shipDays

        l = max(weights)
        r = sum(weights)
        return bisect.bisect_left(range(l, r), True,
                                key=lambda m:
                                shipDays(m) <= days) + l
```

• LeetCode

```
class Solution:
    def shipWithinDays(self, weights: List[int],
days: int) -> int:
        def check(mx):
            ws, cnt = 0, 1
            for w in weights:
                ws += w
                if ws > mx:
                    cnt += 1
                    ws = w
            return cnt <= days
```

```
        left, right = max(weights), sum(weights) +
1
        return left + bisect_left(range(left,
right), True, key=check)
```

• SimplyLeet

Round 2 · High · Medium

p.717

```
def shipWithinDays(weights, days):
    # Helper function to determine if a given
capacity can ship packages within 'days' days.
    def canShip(capacity):
        current_load = 0 # Current total weight on
the ship for the current day.
        required_days = 1 # Start with the first
day.
        for weight in weights:
            # If adding the current weight exceeds
capacity, increment the day counter.
            if current_load + weight > capacity:
                required_days += 1
                current_load = 0
```

```
        current_load += weight
        return required_days <= days

    # Establish lower and upper bounds for binary
search.
    low, high = max(weights), sum(weights)
    while low < high:
        mid = (low + high) // 2
        # If mid capacity can ship all packages
within the allowed days, try a lower capacity.
        if canShip(mid):
            high = mid
```

Round 2 · High · Medium

p.718

```

    else:
        low = mid + 1
    return low

```

```

# Example usage:
weights = [1,2,3,4,5,6,7,8,9,10]
days = 5
print(shipWithinDays(weights, days))

```

• LC.ca

```

class Solution:
    def shipWithinDays(self, weights: List[int],
days: int) -> int:
        def check(mx):
            ws, cnt = 0, 1
            for w in weights:
                ws += w
                if ws > mx:
                    cnt += 1
                    ws = w
            return cnt <= days

```

```

        left, right = max(weights), sum(weights) +
1
        return left + bisect_left(range(left,
right), True, key=check)

```

◆ Quick Review

Key Insights

- The minimal possible capacity is the maximum value in weights; no ship can carry a package heavier than its capacity.
- The maximum possible capacity is the sum of all weights; this allows shipping all packages in one day.
- The problem is monotonic: if a capacity C works, any capacity greater than C will also work.
- Use binary search on the capacity range and a greedy approach to simulate shipping for each candidate capacity.

Space & Time

Time Complexity: $O(n * \log(\text{sum}(\text{weights})))$ where n is the number of packages.

Space Complexity: $O(1)$ additional space.

Solution

The solution employs binary search to find the minimum ship capacity. Start by setting the search bounds: the lower bound is the maximum weight (since each package must be shipped individually if needed) and the upper bound is the sum of all weights (shipping everything in one go). For each midpoint capacity in the binary search, use a greedy method to simulate the shipping process—accumulating weights until adding another package would exceed the candidate ship capacity, at which point you increment the day count. If the number of days required exceeds D , the capacity is insufficient and the lower bound is increased; otherwise, adjust the upper bound. Continue until the bounds converge.

Binary Search

#1060 Missing Element in Sorted Array

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Premium 98

Problem

Description

Given an integer array `nums` which is sorted in ascending order and all of its elements are unique and given also an integer `k`, return the `k`th missing number starting from the leftmost number of the array.

Example 1:

Input: `nums = [4,7,9,10]`, `k = 1`
 Output: 5
 Explanation: The first missing number is 5.

Example 2:

Input: `nums = [4,7,9,10]`, `k = 3`
 Output: 8
 Explanation: The missing numbers are `[5,6,8,...]`, hence the third missing number is 8.

Example 3:

Input: `nums = [1,2,4]`, `k = 3`
 Output: 6
 Explanation: The missing numbers are `[3,5,6,7,...]`, hence the third missing number is 6.

Constraints:

- $1 \leq \text{nums.length} \leq 5 * 10^4$
- $1 \leq \text{nums}[i] \leq 10^7$
- `nums` is sorted in ascending order, and all the elements are unique.
- $1 \leq k \leq 10^8$

Follow up: Can you find a logarithmic time complexity (i.e., $O(\log(n))$) solution?

Community Solutions (Python)

• LeetCode

```
class Solution:
    def missingElement(self, nums: List[int], k: int) -> int:
        def missing(i: int) -> int:
            return nums[i] - nums[0] - i

        n = len(nums)
        if k > missing(n - 1):
            return nums[n - 1] + k - missing(n - 1)
        l, r = 0, n - 1
        while l < r:
```

```

        mid = (l + r) >> 1
        if missing(mid) >= k:
            r = mid
        else:
            l = mid + 1
    return nums[l - 1] + k - missing(l - 1)

```

• SimplyLeet

```

# Python solution with comprehensive comments

def missingElement(nums, k):
    # Calculate the missing count until index i:
    missing(i) = nums[i] - nums[0] - i.
    def missing_count(index):
        return nums[index] - nums[0] - index

    n = len(nums)

    # If kth missing number is beyond the last
    element.

```

```

    if k > missing_count(n - 1):
        return nums[-1] + k - missing_count(n - 1)

    # Perform binary search to find the smallest
    index where missing_count(index) >= k.
    left, right = 0, n - 1
    while left < right:
        mid = left + (right - left) // 2
        if missing_count(mid) < k:
            left = mid + 1
        else:
            right = mid

    # After binary search, left is the smallest
    index such that missing_count(left) >= k.
    # The kth missing number is after nums[left -
    1].
    return nums[left - 1] + k - missing_count(left
    - 1)

# Example usage:
print(missingElement([4,7,9,10], 3)) # Output: 8

```

• LC.ca

```

class Solution:
    def missingElement(self, nums: List[int], k:
int) -> int:
        def missing(i: int) -> int:
            return nums[i] - nums[0] - i

        n = len(nums)
        if k > missing(n - 1):
            return nums[n - 1] + k - missing(n - 1)
        l, r = 0, n - 1
        while l < r:

            mid = (l + r) >> 1
            if missing(mid) >= k:
                r = mid
            else:
                l = mid + 1
        return nums[l - 1] + k - missing(l - 1)

```

◆ Quick Review

Key Insights

- The array is sorted and every element is unique.
- For any index i , the number of missing numbers up to that point can be computed as $\text{nums}[i] - \text{nums}[0] - i$.
- When $\text{missing}(i)$ is less than k , the k th missing element is further to the right.
- Use binary search to find the smallest index where the count of missing numbers is greater than or equal to k .

- If all missing numbers before the last array element are counted and the k th missing still hasn't been reached, calculate the result by continuing past the last element.

Space & Time

Time Complexity: $O(\log(n))$

Space Complexity: $O(1)$

Solution

We compute the number of missing elements until an index i with the formula: $\text{missing}(i) = \text{nums}[i] - \text{nums}[0] - i$. We then use binary search over the indices to find the smallest index where $\text{missing}(i) \geq k$. If the binary search index equals the length of the array, it means the k th missing number lies beyond the last element. Otherwise, the k th missing number is derived as: $\text{nums}[\text{index} - 1] + (k - \text{missing}(\text{index} - 1))$. This approach uses binary search for logarithmic time and does not require extra auxiliary data structures.

Binary Search

#1062 Longest Repeating Substring MED

LeetCode · SimplyLeet · WalkCC · LeetCode

String · Binary Search · Dynamic Programming · Rolling Hash · Suffix Array · Hash Function

Lists: Denny Zhang

Problem

Description

Given a string s , return the length of the longest repeating substrings. If no repeating substring exists, return 0.

Example 1:

Input: $s = \text{"abcd"}$
 Output: 0
 Explanation: There is no repeating substring.

Example 2:

Input: $s = \text{"abbaba"}$
 Output: 2
 Explanation: The longest repeating substrings are "ab" and "ba", each of which occurs twice.

Example 3:

Input: $s = \text{"aabcaabdaab"}$

Output: 3

Explanation: The longest repeating substring is "aab", which occurs 3 times.

Constraints:

- $1 \leq s.length \leq 2000$
- s consists of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestRepeatingSubstring(self, s: str) -> int:
        n = len(s)
        ans = 0
        # dp[i][j] := the number of repeating
        # characters of s[0..i) and s[0..j)
        dp = [[0] * (n + 1) for _ in range(n + 1)]

        for i in range(1, n + 1):
            for j in range(i + 1, n + 1):
                if s[i - 1] == s[j - 1]:
                    dp[i][j] = 1 + dp[i - 1][j - 1]
                    ans = max(ans, dp[i][j])

        return ans
```

• LeetCode


```

class Solution:
    def longestRepeatingSubstring(self, s: str) ->
int:
    n = len(s)
    f = [[0] * n for _ in range(n)]
    ans = 0
    for i in range(1, n):
        for j in range(i):
            if s[i] == s[j]:
                f[i][j] = 1 + (f[i - 1][j - 1]
if j else 0)
                ans = max(ans, f[i][j])

    return ans

```

• SimplyLeet

```

def longestRepeatingSubstring(s):
    n = len(s)
    # Convert characters to integers for rolling
hash computation.
    nums = [ord(c) - ord('a') for c in s]

    # Base value for rolling hash.
    a = 26
    # Modulus value to avoid overflow.
    mod = 2**32

```

```

    # Helper function to check if any substring of
given length L repeats.
    def search(L):
        h = 0
        for i in range(L):
            h = (h * a + nums[i]) % mod
        seen = {h}
        # Precompute a^L % mod for removing the
first digit in the rolling hash.
        aL = pow(a, L, mod)
        for start in range(1, n - L + 1):
            h = (h * a - nums[start - 1] * aL +
nums[start + L - 1]) % mod

        if h in seen:
            return True
        seen.add(h)
        return False

    left, right = 1, n
    ans = 0
    while left <= right:
        mid = left + (right - left) // 2
        if search(mid):

```

```

        ans = mid
        left = mid + 1
    else:
        right = mid - 1
return ans

```

```

# Example usage:
s = "aabcaabdaab"
print(longestRepeatingSubstring(s)) # Expected
output: 3

```

• LC.ca

```

class Solution:
    def longestRepeatingSubstring(self, s: str) ->
int:
    n = len(s)
    dp = [[0] * n for _ in range(n)]
    ans = 0
    for i in range(n):
        for j in range(i + 1, n):
            if s[i] == s[j]:
                dp[i][j] = dp[i - 1][j - 1] + 1
            else 1
                ans = max(ans, dp[i][j])

    return ans

```

◆ Quick Review

Key Insights

- Use binary search to efficiently check for the existence of a repeating substring of a given length.
- Apply a rolling hash technique to compute hash values for substrings in $O(1)$ time after initial processing.
- Use a hash set to detect duplicate hash values (indicating a repeating substring).
- Adjust binary search bounds based on whether a repeating substring of the current length exists.
- Alternative approaches include dynamic programming and suffix arrays, but the binary search with rolling hash is efficient and conceptually clean.

Space & Time

Time Complexity: $O(n \log n)$ on average, where n is the length of the string. (Each binary search step takes $O(n)$ time for computing rolling hashes.)

Space Complexity: $O(n)$ for storing hash values in the hash set.

Solution

The solution uses binary search to decide on the candidate length L of the repeating substring. For each candidate length L , we use a rolling hash to compute hash values for all substrings of length L and store them in a hash set. If any hash is repeated, it confirms the presence of a repeating substring of length L , and we try a longer length. Otherwise, we decrease the candidate length. The rolling hash is computed using a base (26, for lowercase letters) and a modulus (2^{32} to limit integer overflow) which allows efficient hash recomputation for sliding windows.

This approach efficiently narrows down the maximum repeating substring length.

Round 2 · High · Medium

p.733

Binary Search

#1533 Find the Index of the Large Integer

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Interactive

Lists: Premium 98

Problem

Description

We have an integer array `arr`, where all the integers in `arr` are equal except for one integer which is larger than the rest of the integers. You will not be given direct access to the array, instead, you will have an API `ArrayReader` which have the following functions:

- `int compareSub(int l, int r, int x, int y)`: where $0 \leq l, r, x, y < \text{ArrayReader.length}()$, $l \leq r$ and $x \leq y$. The function compares the sum of sub-array `arr[l..r]` with the sum of the sub-array `arr[x..y]` and returns: 1 if `arr[l]+arr[l+1]+...+arr[r] > arr[x]+arr[x+1]+...+arr[y]`.
- 0 if `arr[l]+arr[l+1]+...+arr[r] == arr[x]+arr[x+1]+...+arr[y]`.
- -1 if `arr[l]+arr[l+1]+...+arr[r] < arr[x]+arr[x+1]+...+arr[y]`.

You are allowed to call `compareSub()` 20 times at most. You can assume both functions work in $O(1)$ time.

Round 2 · High · Medium

p.734

Return the index of the array `arr` which has the largest integer.

Example 1:

Input: `arr = [7,7,7,7,10,7,7,7]`

Output: 4

Explanation: The following calls to the API `reader.compareSub(0, 0, 1, 1)` // returns 0 this is a query comparing the sub-array (0, 0) with the sub array (1, 1), (i.e. compares `arr[0]` with `arr[1]`). Thus we know that `arr[0]` and `arr[1]` doesn't contain the largest element.
`reader.compareSub(2, 2, 3, 3)` // returns 0, we can exclude `arr[2]` and `arr[3]`.
`reader.compareSub(4, 4, 5, 5)` // returns 1, thus for sure `arr[4]` is the largest element in the array.
 Notice that we made only 3 calls, so the answer is valid.

Example 2:

Input: `nums = [6,6,12]`

Output: 2

Constraints:

- $2 \leq \text{arr.length} \leq 5 * 10^5$
- $1 \leq \text{arr}[i] \leq 100$
- All elements of `arr` are equal except for one element which is larger than all other elements.

Follow up:

- What if there are two numbers in `arr` that are bigger than all other numbers?

- What if there is one number that is bigger than other numbers and one number that is smaller than other numbers?

Community Solutions (Python)

• WalkCC

```
# """
# This is ArrayReader's API interface.
# You should not implement it, or speculate about
# its implementation
# """
# class ArrayReader(object):
#     # Compares the sum of arr[l..r] with the sum of
#     # arr[x..y]
#     # return 1 if sum(arr[l..r]) > sum(arr[x..y])
#     # return 0 if sum(arr[l..r]) == sum(arr[x..y])
#     # return -1 if sum(arr[l..r]) < sum(arr[x..y])
#     def compareSub(self, l: int, r: int, x: int, y:
# int) -> int:
```

```
#
# # Returns the length of the array
# def length(self) -> int:
#

class Solution:
    def getIndex(self, reader: 'ArrayReader') ->
int:
    l = 0
    r = reader.length() - 1

    while l < r:
        m = (l + r) // 2
        if (r - l) % 2 == 0:
            res = reader.compareSub(l, m - 1, m + 1, r)
            if res == 0:
                return m
            if res == 1:
                r = m - 1
            else: # res == -1
```

```
        l = m + 1
    else:
        res = reader.compareSub(l, m, m + 1, r)
        # res is either 1 or -1.
        if res == 1:
            r = m
        else: # res == -1
            l = m + 1

    return l
```

• LeetCode

```
# """
# This is ArrayReader's API interface.
# You should not implement it, or speculate about
its implementation
# """
# class ArrayReader(object):
# # Compares the sum of arr[l..r] with the sum of
arr[x..y]
# # return 1 if sum(arr[l..r]) > sum(arr[x..y])
# # return 0 if sum(arr[l..r]) == sum(arr[x..y])
# # return -1 if sum(arr[l..r]) < sum(arr[x..y])
# def compareSub(self, l: int, r: int, x: int, y:
int) -> int:
```

```
#
# # Returns the length of the array
# def length(self) -> int:
#

class Solution:
    def getIndex(self, reader: 'ArrayReader') ->
int:
    left, right = 0, reader.length() - 1
    while left < right:

        t1, t2, t3 = (
            left,
            left + (right - left) // 3,
            left + ((right - left) // 3) * 2 +
1,
        )
        cmp = reader.compareSub(t1, t2, t2 + 1,
t3)

        if cmp == 0:
            left = t3 + 1
        elif cmp == 1:
            right = t2

    else:
        left, right = t2 + 1, t3
    return left
```

• SimplyLeet

Round 2 · High · Medium

p.739

```
# Python solution with line-by-line comments
class Solution:
    def getIndex(self, reader):
        # Initialize search bounds
        left, right = 0, reader.length() - 1

        # Binary search for the candidate index
        while left < right:
            length = right - left + 1
            if length % 2 == 0:

                # Even length: split into two equal
halves.

                mid = left + length // 2
                # Compare left half [left, mid-1]
with right half [mid, right]
                comp = reader.compareSub(left, mid
- 1, mid, right)
                if comp == 1:
                    # Large integer is in the left
half

                    right = mid - 1
                else:
                    # Large integer is in the right
half

                    left = mid
```

Round 2 · High · Medium

p.740

```

        else:
            # Odd length: divide into two
            halves of equal size, ignoring middle element.
            half = length // 2
            # Compare left half [left,
            left+half-1] with right half [right-half+1, right]
            comp = reader.compareSub(left,
            left + half - 1, right - half + 1, right)
            if comp == 1:
                # Large integer in left half
                right = left + half - 1
            elif comp == -1:
                # Large integer in right half

                left = right - half + 1
            else:
                # Halves are equal, so the
                middle element is the large one.
                return left + half
        return left

```

• LC.ca

```

# """
# This is ArrayReader's API interface.
# You should not implement it, or speculate about
its implementation
# """
# class ArrayReader(object):
# # Compares the sum of arr[l..r] with the sum of
arr[x..y]
# # return 1 if sum(arr[l..r]) > sum(arr[x..y])
# # return 0 if sum(arr[l..r]) == sum(arr[x..y])
# # return -1 if sum(arr[l..r]) < sum(arr[x..y])
# def compareSub(self, l: int, r: int, x: int, y:
int) -> int:

#
# # Returns the length of the array
# def length(self) -> int:
#

class Solution:
    def getIndex(self, reader: 'ArrayReader') ->
int:
        left, right = 0, reader.length() - 1
        while left < right:

```

```

        t1, t2, t3 = (
            left,
            left + (right - left) // 3,
            left + ((right - left) // 3) * 2 +
1,
        )
    cmp = reader.compareSub(t1, t2, t2 + 1,
t3)

    if cmp == 0:
        left = t3 + 1
    elif cmp == 1:
        right = t2

    else:
        left, right = t2 + 1, t3
    return left

```

◆ Quick Review

Key Insights

- Use binary search to efficiently narrow down the position of the unique large value.
- At each step, divide the candidate range into two halves (or almost two halves when the range size is odd) and use compareSub to determine which half contains the large integer.
- Comparing two subarrays of equal length will reveal the region that contains the extra value via the sum difference.

- Handle odd-length ranges carefully by isolating the middle element if needed.

Space & Time

Time Complexity: $O(\log n)$ – because binary search is performed over the range.

Space Complexity: $O(1)$ – only a few pointers are maintained, no additional space is required.

Solution

The solution applies a binary search approach over the array indices. At every iteration:

- Calculate the current segment length. – If the length is even, split the segment evenly and compare the left half versus the right half using compareSub. – If the length is odd, compare two halves of equal size while ignoring the middle element. If the sums are equal, then the middle index holds the large value. – Update the search bounds based on the result of compareSub. – Continue until the search space is reduced to one index, which is the answer.
- This method ensures that you locate the large integer with $O(\log n)$ calls to compareSub, well within the limit of 20 calls.

Matrix

Round 2 · High · Medium · 14 questions

Round 2 · High · Medium

p.745

Matrix

#36 Valid Sudoku

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Matrix

Lists: Grind 169

Problem

Description

Determine if a 9 x 9 Sudoku board is valid. Only the filled cells need to be validated according to the following rules:

1. Each row must contain the digits 1–9 without repetition.
2. Each column must contain the digits 1–9 without repetition.
3. Each of the nine 3 x 3 sub-boxes of the grid must contain the digits 1–9 without repetition.

Note:

- A Sudoku board (partially filled) could be valid but is not necessarily solvable.
- Only the filled cells need to be validated according to the mentioned rules.

Example 1:

Round 2 · High · Medium

p.746

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9

Input: board =

```
[["5","3",".",".","7",".",".",".","."],
["6",".",".","1","9","5",".",".","."],
[".","9","8",".",".",".","6","."],
["8",".",".","6",".",".","3",".","1"],
["4",".","8","3",".",".","","1"],
["7",".","2",".",".","6"],
["6",".","2","8","."],
["4","1","9","5"],
["8","7","9"]]
```

Example 2:

Round 2 · High · Medium

p.747

Input: board =

```
[["8","3",".",".","7",".",".","."],
["6",".","1","9","5",".","."],
[".","9","8",".",".","6","."],
["8",".","6",".","3",".","1"],
["4",".","8","3",".","."],
["7",".","2",".","6"],
["6","2","8","."],
["4","1","9","5"],
["8","7","9"]]
```

Constraints:

- board.length == 9
- board[i].length == 9
- board[i][j] is a digit 1-9 or '.'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isValidSudoku(self, board: list[list[str]]) -> bool:
        seen = set()

        for i in range(9):
            for j in range(9):
                c = board[i][j]
                if c == '.':
                    continue
                if (c + '@row ' + str(i) in seen or
```

Round 2 · High · Medium

p.748

```

        c + '@col ' + str(j) in seen or
        c + '@box ' + str(i // 3) + str(j
// 3) in seen):
        return False
    seen.add(c + '@row ' + str(i))
    seen.add(c + '@col ' + str(j))
    seen.add(c + '@box ' + str(i // 3) + str(j
// 3))

    return True

```

• LeetCode

```

class Solution:
    def isValidSudoku(self, board:
List[List[str]]) -> bool:
        row = [[False] * 9 for _ in range(9)]
        col = [[False] * 9 for _ in range(9)]
        sub = [[False] * 9 for _ in range(9)]
        for i in range(9):
            for j in range(9):
                c = board[i][j]
                if c == '.':
                    continue

```

```

        num = int(c) - 1
        k = i // 3 * 3 + j // 3
        if row[i][num] or col[j][num] or
sub[k][num]:
            return False
        row[i][num] = True
        col[j][num] = True
        sub[k][num] = True

    return True

```

• SimplyLeet

```

# Python solution with inline comments
def isValidSudoku(board):
    # Initialize sets for rows, columns, and boxes
    rows = [set() for _ in range(9)]
    cols = [set() for _ in range(9)]
    boxes = [set() for _ in range(9)]

    # Iterate over each cell in the board
    for r in range(9):
        for c in range(9):

```

```

        num = board[r][c]
        # Skip empty cells
        if num == '.':
            continue
        # Compute box index using integer
division
        box_index = (r // 3) * 3 + (c // 3)
        # Check if the number is already
present in the row, column, or box
        if num in rows[r] or num in cols[c] or
num in boxes[box_index]:
            return False
        # Add number to row, column, and box
sets

```

```

        rows[r].add(num)
        cols[c].add(num)
        boxes[box_index].add(num)

    return True

# Example usage:
if __name__ == "__main__":
    board = [
        ["5","3",".", ".", ".", "7",".", ".", ".", "."],

```

```

        ["6",".", ".", ".", "1","9","5",".", ".", "."],
        [".","9","8",".", ".", ".", ".", "6","."],
        ["8",".", ".", ".", "6",".", ".", ".", "3"],
        ["4",".", ".", "8",".", "3",".", ".", "1"],
        ["7",".", ".", ".", "2",".", ".", ".", "6"],
        [".","6",".", ".", ".", ".", "2","8","."],
        [".",".", ".", "4","1","9",".", ".", "5"],
        [".",".", ".", ".", "8",".", ".", "7","9"]
    ]
    print(isValidSudoku(board)) # Expected output:
    True

```

• LC.ca

```

class Solution:
    def isValidSudoku(self, board:
List[List[str]]) -> bool:
        row = [[False] * 9 for _ in range(9)]
        col = [[False] * 9 for _ in range(9)]
        sub = [[False] * 9 for _ in range(9)]
        for i in range(9):
            for j in range(9):
                c = board[i][j]
                if c == '.':
                    continue

```

```

        num = int(c) - 1
        k = i // 3 * 3 + j // 3
        if row[i][num] or col[j][num] or
sub[k][num]:
            return False
        row[i][num] = True
        col[j][num] = True
        sub[k][num] = True
    return True

```

◆ Quick Review

Key Insights

- Validate each row by checking for duplicate numbers while ignoring empty cells.
- Validate each column in the same way.
- Divide the board into nine 3x3 sub-boxes and check each for duplicate numbers.
- Use hash sets (or equivalent data structures) to keep track of seen digits for rows, columns, and boxes.
- The board size is fixed (9x9), so the overall time complexity is constant with respect to input size.

Space & Time

Time Complexity: $O(1)$ — Since the board size is fixed at 9x9.

Space Complexity: $O(1)$ — The extra space used by the sets is bounded by the fixed board size.

Solution

The solution uses three main data structures: an array of sets for rows, an array of sets for columns, and an array of sets for 3x3 sub-boxes. For each cell, if it is not empty, check if the number is already present in the corresponding row, column, or box. If found, return false immediately. Otherwise, add the number to the respective sets. Compute the index for the 3x3 sub-box using the formula: $(\text{row_index} // 3) * 3 + (\text{col_index} // 3)$. If no duplicates are found after traversing the entire board, the board is considered valid.

Matrix

#48 Rotate Image

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Math · Matrix

Lists: Grind 169 | CodeSignal

Problem


Description

You are given an $n \times n$ 2D matrix representing an image, rotate the image by 90 degrees (clockwise).

You have to rotate the image in-place, which means you have to modify the input 2D matrix directly. DO NOT allocate another 2D matrix and do the rotation.

Example 1:

1	2	3
4	5	6
7	8	9



7	4	1
8	5	2
9	6	3


Round 2 · High · Medium

p.755

Input: matrix = [[1,2,3],[4,5,6],[7,8,9]]
 Output: [[7,4,1],[8,5,2],[9,6,3]]

Example 2:

5	1	9	11
2	4	8	10
13	3	6	7
15	14	12	16



15	13	2	5
14	3	4	1
12	6	8	9
16	7	10	11

Input: matrix =
 [[5,1,9,11],[2,4,8,10],[13,3,6,7],[15,14,12,16]]
 Output:
 [[15,13,2,5],[14,3,4,1],[12,6,8,9],[16,7,10,11]]

Constraints:

- $n == \text{matrix.length} == \text{matrix}[i].\text{length}$
- $1 \leq n \leq 20$
- $-1000 \leq \text{matrix}[i][j] \leq 1000$

Community Solutions (Python)

- WalkCC

Round 2 · High · Medium

p.756

```
class Solution:
    def rotate(self, matrix: list[list[int]]) ->
None:
        matrix.reverse()
        for i, j in
itertools.combinations(range(len(matrix)), 2):
            matrix[i][j], matrix[j][i] = matrix[j][i],
matrix[i][j]
```

```
class Solution:
    def rotate(self, matrix: list[list[int]]) ->
None:
        for mn in range(len(matrix) // 2):
            mx = len(matrix) - mn - 1
            for i in range(mn, mx):
                offset = i - mn
                top = matrix[mn][i]
                matrix[mn][i] = matrix[mx - offset][mn]
                matrix[mx - offset][mn] = matrix[mx][mx -
offset]
                matrix[mx][mx - offset] = matrix[i][mx]

            matrix[i][mx] = top
```

• LeetCode

```
class Solution:
    def rotate(self, matrix: List[List[int]]) ->
None:
        n = len(matrix)
        for i in range(n >> 1):
            for j in range(n):
                matrix[i][j], matrix[n - i - 1][j]
= matrix[n - i - 1][j], matrix[i][j]
            for i in range(n):
                for j in range(i):
                    matrix[i][j], matrix[j][i] =
matrix[j][i], matrix[i][j]
```

• SimplyLeet

```
def rotate(matrix):
    n = len(matrix)
    # Transpose the matrix by swapping elements
across the diagonal
    for i in range(n):
        for j in range(i, n):
            matrix[i][j], matrix[j][i] =
matrix[j][i], matrix[i][j]
    # Reverse each row to complete the clockwise
rotation
    for i in range(n):
        matrix[i].reverse()
```

```
# Example usage:
matrix = [[1,2,3],[4,5,6],[7,8,9]]
rotate(matrix)
print(matrix) # Outputs: [[7,4,1],[8,5,2],[9,6,3]]
```

• LC.ca

```
class Solution(object):
    def rotate(self, matrix):
        """
        :type matrix: List[List[int]]
        :rtype: void Do not return anything, modify
        matrix in-place instead.
        """
        if len(matrix) == 0:
            return
        h = len(matrix)
        w = len(matrix[0])

        for i in range(0, h): # mirror
            for j in range(0, w / 2):
                matrix[i][j], matrix[i][w - j - 1] =
                matrix[i][w - j - 1], matrix[i][j]

        for i in range(0, h): # transpos
            for j in range(0, w - 1 - i):
                matrix[i][j], matrix[w - 1 - j][h - 1 - i]
                = matrix[w - 1 - j][h - 1 - i], matrix[i][j]

        #####
```

Round 2 · High · Medium

p.759

```
class Solution:
    def rotate(self, matrix: List[List[int]]) ->
    None:
        n = len(matrix)
        for i in range(n >> 1):
            for j in range(n):
                matrix[i][j], matrix[n - i - 1][j]
                = matrix[n - i - 1][j], matrix[i][j]
            for i in range(n):
                for j in range(i):
                    matrix[i][j], matrix[j][i] =
                    matrix[j][i], matrix[i][j]
```

◆ Quick Review

Key Insights

- The rotation can be achieved in two steps: first transpose the matrix and then reverse each row.
- Transposing the matrix swaps $\text{matrix}[i][j]$ with $\text{matrix}[j][i]$, converting rows into columns.
- Reversing each row post-transposition yields the desired clockwise rotation.
- This approach works in-place with $O(1)$ extra space.

Space & Time

Time Complexity: $O(n^2)$ since every element is processed during the transpose and row reversal.

Round 2 · High · Medium

p.760

Space Complexity: $O(1)$ as the operations are performed in-place without extra data structures.

Solution

The solution leverages a two-step in-place transformation:

- **Transpose the matrix:** Iterate through the matrix and swap each element (i, j) with element (j, i) for $j \geq i$.
- **Reverse each row:** After the transpose, each row is reversed to rotate the matrix 90 degrees clockwise. The key trick is to realize that a transpose followed by a reverse of each row achieves the rotation without needing extra space.

Round 2 · High · Medium

p.761

Matrix

#54 Spiral Matrix

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix · Simulation

Lists: Grind 169 | CodeSignal

Problem


Description

Given an $m \times n$ matrix, return all elements of the matrix in spiral order.

Example 1:


Round 2 · High · Medium

p.762



Input: matrix = [[1,2,3],[4,5,6],[7,8,9]]
 Output: [1,2,3,6,9,8,7,4,5]

Example 2:



Input: matrix = [[1,2,3,4],[5,6,7,8],[9,10,11,12]]
 Output: [1,2,3,4,8,12,11,10,9,5,6,7]

Constraints:

- `m == matrix.length`
- `n == matrix[i].length`
- `1 <= m, n <= 10`
- `-100 <= matrix[i][j] <= 100`

Community Solutions (Python)

- WalkCC

```

class Solution:
    def spiralOrder(self, matrix: list[list[int]])
    -> list[int]:
        if not matrix:
            return []

        m = len(matrix)
        n = len(matrix[0])
        ans = []
        r1 = 0
        c1 = 0

        r2 = m - 1
        c2 = n - 1

        # Repeatedly add matrix[r1..r2][c1..c2] to
        `ans`.
        while len(ans) < m * n:
            j = c1
            while j <= c2 and len(ans) < m * n:
                ans.append(matrix[r1][j])
                j += 1
            i = r1 + 1

```

```

        while i <= r2 - 1 and len(ans) < m * n:
            ans.append(matrix[i][c2])
            i += 1
        j = c2
        while j >= c1 and len(ans) < m * n:
            ans.append(matrix[r2][j])
            j -= 1
        i = r2 - 1
        while i >= r1 + 1 and len(ans) < m * n:
            ans.append(matrix[i][c1])

            i -= 1
            r1 += 1
            c1 += 1
            r2 -= 1
            c2 -= 1

        return ans

```

- LeetCode

```
class Solution:
    def spiralOrder(self, matrix: List[List[int]])
-> List[int]:
    m, n = len(matrix), len(matrix[0])
    dirs = (0, 1, 0, -1, 0)
    vis = [[False] * n for _ in range(m)]
    i = j = k = 0
    ans = []
    for _ in range(m * n):
        ans.append(matrix[i][j])
        vis[i][j] = True

        x, y = i + dirs[k], j + dirs[k + 1]
        if x < 0 or x >= m or y < 0 or y >= n
or vis[x][y]:
            k = (k + 1) % 4
            i += dirs[k]
            j += dirs[k + 1]
    return ans
```

• SimplyLeet

Round 2 · High · Medium

p.767

```
def spiralOrder(matrix):
    # Initialize the boundaries of the matrix
    top, bottom = 0, len(matrix) - 1
    left, right = 0, len(matrix[0]) - 1
    result = [] # This will store the spiral order

    # Loop until the pointers cross each other
    while top <= bottom and left <= right:
        # Traverse from left to right on the
current top row
        for col in range(left, right + 1):
            result.append(matrix[top][col])
            top += 1 # Move the top boundary down

        # Traverse downwards on the current right
column
        for row in range(top, bottom + 1):
            result.append(matrix[row][right])
            right -= 1 # Move the right boundary to the
left

        # Traverse from right to left on the
current bottom row, if valid
        if top <= bottom:
```

Round 2 · High · Medium

p.768

```

        for col in range(right, left - 1, -1):
            result.append(matrix[bottom][col])
        bottom -= 1 # Move the bottom boundary
up
        # Traverse upwards on the current left
column, if valid
        if left <= right:
            for row in range(bottom, top - 1, -1):
                result.append(matrix[row][left])
            left += 1 # Move the left boundary
right
    return result

# Example usage:
# print(spiralOrder([[1,2,3],[4,5,6],[7,8,9]]))

```

• LC.ca

```

class Solution:
    def spiralOrder(self, matrix: List[List[int]])
-> List[int]:
        m, n = len(matrix), len(matrix[0])
        dirs = (0, 1, 0, -1, 0)
        i = j = k = 0
        ans = []
        vis = set()
        for _ in range(m * n):
            ans.append(matrix[i][j])
            vis.add((i, j))

```

```

        x, y = i + dirs[k], j + dirs[k + 1]
        if not 0 <= x < m or not 0 <= y < n or
(x, y) in vis:
            k = (k + 1) % 4
            i = i + dirs[k]
            j = j + dirs[k + 1]
        return ans

```

♦ Quick Review

Key Insights

- Use four boundaries: top, bottom, left, and right to track the current layer.
- Traverse right across the top row, then down the right column, then left on the bottom row, and finally up the left column.
- After each direction, adjust the corresponding boundary to move inward.
- Continue the traversal until the boundaries cross.

Space & Time

Time Complexity: $O(m * n)$ since each element is visited exactly once.

Space Complexity: $O(1)$ extra space (excluding the space required for the output list).

Solution

The solution uses a simulation approach with boundary pointers (top, bottom, left, right). At each step, the boundaries indicate the current submatrix to traverse. You

iterate as follows:

- Traverse from left to right along the top boundary.
- Traverse downward along the right boundary.
- If the top boundary has not passed the bottom, traverse from right to left along the bottom boundary.
- If the left boundary has not passed the right, traverse upward along the left boundary.

After each traversal, the boundaries are adjusted inward. This continues until all elements have been visited.

Round 2 · High · Medium

p.771

Matrix

#59 Spiral Matrix II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix · Simulation

Lists: CodeSignal

Problem


Description

Given a positive integer n , generate an $n \times n$ matrix filled with elements from 1 to n^2 in spiral order.

Example 1:

Round 2 · High · Medium

p.772



Input: n = 3
Output: [[1,2,3],[8,9,4],[7,6,5]]

Example 2:

Input: n = 1
Output: [[1]]

Constraints:

- 1 ≤ n ≤ 20

Community Solutions (Python)

• WalkCC

```
class Solution:
    def generateMatrix(self, n: int) -> list[list[int]]:
        ans = [[0] * n for _ in range(n)]
        count = 1

        for mn in range(n // 2):
            mx = n - mn - 1
            for i in range(mn, mx):
                ans[mn][i] = count
                count += 1

            for i in range(mn, mx):
                ans[i][mx] = count
                count += 1
            for i in range(mx, mn, -1):
                ans[mx][i] = count
                count += 1
            for i in range(mx, mn, -1):
                ans[i][mn] = count
                count += 1

        if n % 2 == 1:
            ans[n // 2][n // 2] = count

        return ans
```

• LeetCode

```

class Solution:
    def generateMatrix(self, n: int) ->
List[List[int]]:
    ans = [[0] * n for _ in range(n)]
    dirs = (0, 1, 0, -1, 0)
    i = j = k = 0
    for v in range(1, n * n + 1):
        ans[i][j] = v
        x, y = i + dirs[k], j + dirs[k + 1]
        if x < 0 or x >= n or y < 0 or y >= n
or ans[x][y]:
            k = (k + 1) % 4

        i, j = i + dirs[k], j + dirs[k + 1]
    return ans

```

• SimplyLeet

```

# Function to generate spiral matrix of size n x n
def generateMatrix(n):
    # Initialize an n x n matrix with zeros
    matrix = [[0] * n for _ in range(n)]
    # Initialize the starting number to fill
    num = 1
    # Set the initial boundaries for rows and
columns
    top, bottom = 0, n - 1
    left, right = 0, n - 1

```

```

        # Continue the loop until the boundaries are
valid
        while top <= bottom and left <= right:
            # Traverse from left to right along the top
row
            for j in range(left, right + 1):
                matrix[top][j] = num
                num += 1
            top += 1 # Move the top boundary down

            # Traverse from top to bottom along the
right column
            for i in range(top, bottom + 1):

                matrix[i][right] = num
                num += 1
            right -= 1 # Move the right boundary left

            # Traverse from right to left along the
bottom row if still in valid bounds
            if top <= bottom:
                for j in range(right, left - 1, -1):
                    matrix[bottom][j] = num
                    num += 1
                bottom -= 1 # Move the bottom boundary
up

```



```

        # Traverse from bottom to top along the
        left column if still in valid bounds
        if left <= right:
            for i in range(bottom, top - 1, -1):
                matrix[i][left] = num
                num += 1
            left += 1 # Move the left boundary
    right

    return matrix

```

```

# Example usage:
print(generateMatrix(3))

```

• LC.ca

```

class Solution:
    def generateMatrix(self, n: int) ->
    List[List[int]]:
        ans = [[0] * n for _ in range(n)]
        dirs = ((0, 1), (1, 0), (0, -1), (-1, 0))
        i = j = k = 0
        for v in range(1, n * n + 1):
            ans[i][j] = v
            x, y = i + dirs[k][0], j + dirs[k][1]
            if x < 0 or y < 0 or x >= n or y >= n
            or ans[x][y]:
                k = (k + 1) % 4

```

```

                x, y = i + dirs[k][0], j +
                dirs[k][1]
                i, j = x, y
            return ans

```

♦ Quick Review

Key Insights

- Use four pointers (top, bottom, left, right) to keep track of the matrix boundaries.
- Traverse the matrix in a spiral: move right along the top row, then down the right column, left along the bottom row, and finally up the left column.
- Update the boundary pointers after each complete traversal of one side.
- Continue until all numbers from 1 to n^2 are filled into the matrix.

Space & Time

Time Complexity: $O(n^2)$

Space Complexity: $O(n^2)$ (required for the output matrix)

Solution

We simulate the spiral movement by initializing four pointers: top, bottom, left, and right, which denote the current boundaries of the matrix that have not yet been filled. For each iteration, we:

- Traverse from left to right along the top boundary, then increment the top pointer.
- Traverse from top to bottom

along the right boundary, then decrement the right pointer.
– If there are rows remaining, traverse from right to left along the bottom boundary, then decrement the bottom pointer. – If there are columns remaining, traverse from bottom to top along the left boundary, then increment the left pointer. This algorithm ensures that the matrix is filled in a spiral order until all numbers are placed.

Round 2 · High · Medium

p.779

Matrix

#64 Minimum Path Sum

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Matrix

Lists: Denny Zhang

Problem

Description

Given a $m \times n$ grid filled with non-negative numbers, find a path from top left to bottom right, which minimizes the sum of all numbers along its path.

Note: You can only move either down or right at any point in time.

Example 1:

Round 2 · High · Medium

p.780

1	3	1
1	5	1
4	2	1

Input: grid = [[1,3,1],[1,5,1],[4,2,1]]

Output: 7

Explanation: Because the path 1 → 3 → 1 → 1 → 1 minimizes the sum.

Example 2:

Input: grid = [[1,2,3],[4,5,6]]

Output: 12

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$

Round 2 · High · Medium

p.781

- $1 \leq m, n \leq 200$
- $0 \leq \text{grid}[i][j] \leq 200$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def minPathSum(self, grid: list[list[int]]) -> int:
        m = len(grid)
        n = len(grid[0])

        for i in range(m):
            for j in range(n):
                if i > 0 and j > 0:
                    grid[i][j] += min(grid[i - 1][j],
grid[i][j - 1])
                elif i > 0:
                    grid[i][0] += grid[i - 1][0]
                elif j > 0:
                    grid[0][j] += grid[0][j - 1]

        return grid[m - 1][n - 1]
```

• LeetCode

Round 2 · High · Medium

p.782

```

class Solution:
    def minPathSum(self, grid: List[List[int]]) ->
int:
    m, n = len(grid), len(grid[0])
    f = [[0] * n for _ in range(m)]
    f[0][0] = grid[0][0]
    for i in range(1, m):
        f[i][0] = f[i - 1][0] + grid[i][0]
    for j in range(1, n):
        f[0][j] = f[0][j - 1] + grid[0][j]
    for i in range(1, m):
        for j in range(1, n):
            f[i][j] = min(f[i - 1][j], f[i][j -
1]) + grid[i][j]
    return f[-1][-1]

```

• SimplyLeet

```

# Python solution using dynamic programming
def minPathSum(grid):
    # Number of rows and columns in the grid
    rows = len(grid)
    cols = len(grid[0])

    # Initialize the DP table directly in grid to
save space
    # Fill the first row by accumulating the sum
from the left
    for j in range(1, cols):
        grid[0][j] += grid[0][j-1]

```

```

    # Fill the first column by accumulating the sum
above
    for i in range(1, rows):
        grid[i][0] += grid[i-1][0]

    # Compute the minimum path sum for the rest of
the cells
    for i in range(1, rows):
        for j in range(1, cols):
            # The cell value is added to the
minimum of the top or left cell
            grid[i][j] += min(grid[i-1][j],
grid[i][j-1])

    # The bottom-right cell contains the answer
    return grid[rows-1][cols-1]

```

```

# Example usage
example_grid = [[1,3,1],[1,5,1],[4,2,1]]
print(minPathSum(example_grid)) # Output should be
7

```

• LC.ca

```

class Solution:
    def minPathSum(self, grid: List[List[int]]) -> int:
        m, n = len(grid), len(grid[0])
        f = [[0] * n for _ in range(m)]
        f[0][0] = grid[0][0]
        for i in range(1, m):
            f[i][0] = f[i - 1][0] + grid[i][0]
        for j in range(1, n):
            f[0][j] = f[0][j - 1] + grid[0][j]
        for i in range(1, m):
            for j in range(1, n):
                f[i][j] = min(f[i - 1][j], f[i][j - 1]) + grid[i][j]
        return f[-1][-1]

```

◆ Quick Review

Key Insights

- The problem can be solved using Dynamic Programming.
- The optimal solution for each cell depends on the minimum path sum of the cell above it and the cell to its left.
- Since you can only move right or down, only two adjacent cells affect the current cell's result.
- Edge cases include the first row and first column, which have only one direction of movement.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the number of rows and columns.

Space Complexity: $O(m * n)$ for the DP table (this can be optimized to $O(n)$ with in-place computations).

Solution

We use a Dynamic Programming (DP) approach where we create a DP table (or use the grid itself) that stores the minimum path sum to each cell. For each cell (i, j) :

- If $i = 0$ (first row), the only possible path is from the left.
- If $j = 0$ (first column), the only possible path is from above.
- For other cells, the minimum path sum is the cell value plus the minimum of the two possible previous cells (above and left). This ensures that by the time we reach the bottom-right cell, we have accumulated the minimum sum required to reach that cell.

Matrix

#73 Set Matrix Zeroes

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Matrix

Lists: Grind 169

Problem


Description

Given an $m \times n$ integer matrix `matrix`, if an element is 0, set its entire row and column to 0's.

You must do it in place.

Example 1:

1	2	3
4	5	6
7	8	9



7	4	1
8	5	2
9	6	3

Input: `matrix = [[1,1,1],[1,0,1],[1,1,1]]`
 Output: `[[1,0,1],[0,0,0],[1,0,1]]`

Example 2:

Round 2 · High · Medium

p.787

5	1	9	11
2	4	8	10
13	3	6	7
15	14	12	16



15	13	2	5
14	3	4	1
12	6	8	9
16	7	10	11

Input: `matrix = [[0,1,2,0],[3,4,5,2],[1,3,1,5]]`
 Output: `[[0,0,0,0],[0,4,5,0],[0,3,1,0]]`

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[0].\text{length}$
- $1 \leq m, n \leq 200$
- $-231 \leq \text{matrix}[i][j] \leq 231 - 1$

Follow up:

- A straightforward solution using $O(mn)$ space is probably a bad idea.
- A simple improvement uses $O(m + n)$ space, but still not the best solution.
- Could you devise a constant space solution?

Community Solutions (Python)

- WalkCC

Round 2 · High · Medium

p.788

```
class Solution:
    def setZeroes(self, matrix: list[list[int]]) ->
None:
    m = len(matrix)
    n = len(matrix[0])
    shouldFillFirstRow = 0 in matrix[0]
    shouldFillFirstCol = 0 in
list(zip(*matrix))[0]
```

Store the information in the first row and the first column.

```
for i in range(1, m):
    for j in range(1, n):
```

```
    if matrix[i][j] == 0:
        matrix[i][0] = 0
        matrix[0][j] = 0
```

Fill 0s for the matrix except the first row and the first column.

```
for i in range(1, m):
    for j in range(1, n):
        if matrix[i][0] == 0 or matrix[0][j] == 0:
            matrix[i][j] = 0
```

```
# Fill 0s for the first row if needed.
if shouldFillFirstRow:
    matrix[0] = [0] * n
```

```
# Fill 0s for the first column if needed.
if shouldFillFirstCol:
    for row in matrix:
        row[0] = 0
```

• LeetCode

```
class Solution:
    def setZeroes(self, matrix: List[List[int]])
-> None:
        m, n = len(matrix), len(matrix[0])
        row = [False] * m
        col = [False] * n
        for i in range(m):
            for j in range(n):
                if matrix[i][j] == 0:
                    row[i] = col[j] = True
        for i in range(m):
```

```
            for j in range(n):
                if row[i] or col[j]:
                    matrix[i][j] = 0
```

• SimplyLeet

Python solution for the Set Matrix Zeroes problem

```
def setZeroes(matrix):
    # Step 1: Determine the dimensions of matrix
    m = len(matrix)
    n = len(matrix[0])

    # Determine if first row or first column should
    be zeroed
    first_row_zero = any(matrix[0][j] == 0 for j in
range(n))
    first_col_zero = any(matrix[i][0] == 0 for i in
range(m))
```

```
    # Step 2: Use first row and column as markers
    for i in range(1, m):
        for j in range(1, n):
            if matrix[i][j] == 0:
                matrix[i][0] = 0 # mark the first
element of the row
                matrix[0][j] = 0 # mark the first
element of the column
```

```
    # Step 3: Zero out cells based on the markers
in first row and column
    for i in range(1, m):
```

```
        for j in range(1, n):
            if matrix[i][0] == 0 or matrix[0][j] ==
0:
                matrix[i][j] = 0

    # Step 4: Zero out the first row if needed
    if first_row_zero:
        for j in range(n):
            matrix[0][j] = 0

    # Step 5: Zero out the first column if needed

    if first_col_zero:
        for i in range(m):
            matrix[i][0] = 0

# Example usage:
# matrix = [[1,1,1],[1,0,1],[1,1,1]]
# setZeroes(matrix)
# print(matrix)
```

• [LC.ca](#)


```

class Solution:
    def setZeroes(self, matrix: List[List[int]])
-> None:
        m, n = len(matrix), len(matrix[0])
        i0 = any(v == 0 for v in matrix[0])
        j0 = any(matrix[i][0] == 0 for i in
range(m))
        for i in range(1, m):
            for j in range(1, n):
                if matrix[i][j] == 0:
                    matrix[i][0] = matrix[0][j] =
0
        for i in range(1, m):
            for j in range(1, n):
                if matrix[i][0] == 0 or
matrix[0][j] == 0:
                    matrix[i][j] = 0
        if i0:
            for j in range(n):
                matrix[0][j] = 0
        if j0:
            for i in range(m):
                matrix[i][0] = 0

```

◆ Quick Review

Key Insights

- The challenge is to mark rows and columns that need to be zeroed without using extra space.

- A common trick is to use the first row and first column as markers.
- Special care is needed for the first row and first column, as they are also used as markers.
- The algorithm involves two passes: one for marking and one for updating the matrix based on the markers.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(1)$ (in-place, using only constant extra space)

Solution

We can solve this problem by using the first row and first column of the matrix as markers. First, determine if the first row or first column needs to be zeroed, then traverse the rest of the matrix and mark the corresponding first row and column positions if a zero is encountered. Finally, update the matrix backward using the markers, and finally update the first row and column if needed.

The algorithm works in three main steps:

- Check if the first row or first column should be zeroed (store this in two variables).
- Iterate through the rest of the matrix; for any element that is zero, mark its row and column by setting the corresponding element in the first row and first column to zero.
- Iterate through the matrix again (excluding the first row and column) and set elements to zero if their row or column marker is zero.
- Finally, zero out the first row and/or first column if needed.

This approach avoids extra space usage apart from the two boolean flags while ensuring that the matrix is updated correctly.

Round 2 · High · Medium

p.795

Matrix

#74 Search a 2D Matrix

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Matrix

Lists: Grind 169

Problem

Description

You are given an $m \times n$ integer matrix `matrix` with the following two properties:

- Each row is sorted in non-decreasing order.
- The first integer of each row is greater than the last integer of the previous row.

Given an integer `target`, return `true` if `target` is in `matrix` or `false` otherwise.

You must write a solution in $O(\log(m * n))$ time complexity.

Example 1:

Round 2 · High · Medium

p.796

1	3	5	7
10	11	16	20
23	30	34	60


Input: matrix =
[[1,3,5,7],[10,11,16,20],[23,30,34,60]], target = 3
Output: true

Example 2:

Round 2 · High · Medium

p.797

5	1	9	11
2	4	8	10
13	3	6	7
15	14	12	16



15	13	2	5
14	3	4	1
12	6	8	9
16	7	10	11

Input: matrix =
[[1,3,5,7],[10,11,16,20],[23,30,34,60]], target = 13
Output: false

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[i].\text{length}$
- $1 \leq m, n \leq 100$
- $-104 \leq \text{matrix}[i][j], \text{target} \leq 104$

Community Solutions (Python)

- WalkCC

Round 2 · High · Medium

p.798

```
class Solution:
    def searchMatrix(self, matrix: list[list[int]],
target: int) -> bool:
    if not matrix:
        return False

    m = len(matrix)
    n = len(matrix[0])
    l = 0
    r = m * n

    while l < r:
        mid = (l + r) // 2
        i = mid // n
        j = mid % n
        if matrix[i][j] == target:
            return True
        if matrix[i][j] < target:
            l = mid + 1
        else:
            r = mid

    return False
```

• LeetCode

Round 2 · High · Medium

p.799

```
class Solution:
    def searchMatrix(self, matrix:
List[List[int]], target: int) -> bool:
        m, n = len(matrix), len(matrix[0])
        left, right = 0, m * n - 1
        while left < right:
            mid = (left + right) >> 1
            x, y = divmod(mid, n)
            if matrix[x][y] >= target:
                right = mid
            else:
                left = mid + 1
        return matrix[left // n][left % n] ==
target
```

• SimplyLeet

```
def searchMatrix(matrix, target):
    # Get number of rows and columns
    if not matrix or not matrix[0]:
        return False

    rows, cols = len(matrix), len(matrix[0])
    left, right = 0, rows * cols - 1

    # Perform binary search over the virtual
    flattened array
    while left <= right:
```

Round 2 · High · Medium

p.800

```
# Find the middle index
mid = (left + right) // 2
# Map the middle index to row and column in
the matrix
row = mid // cols
col = mid % cols
mid_value = matrix[row][col]

# Check if the mid_value equals the target
if mid_value == target:
    return True

elif mid_value < target:
    left = mid + 1 # Move to the right
subarray
else:
    right = mid - 1 # Move to the left
subarray

# Target not found
return False

# Example usage:
# matrix = [[1,3,5,7],[10,11,16,20],[23,30,34,60]]

# target = 3
# print(searchMatrix(matrix, target))
```

• LC.ca

Round 2 · High · Medium

p.801

```
class Solution:
    def searchMatrix(self, matrix:
List[List[int]], target: int) -> bool:
        m, n = len(matrix), len(matrix[0])
        left, right = 0, m * n - 1
        while left <= right: # must be <=, not <,
for matrax=[[1]],target=1
            mid = (left + right) >> 1
            x, y = divmod(mid, n) # note: divide
column count
            if matrix[x][y] == target:
                return True
            elif matrix[x][y] > target:

                right = mid - 1
            else:
                left = mid + 1
        return False

#####

class Solution:
    def searchMatrix(self, matrix:
List[List[int]], target: int) -> bool:
        m, n = len(matrix), len(matrix[0])
```

Round 2 · High · Medium

p.802

```

left, right = 0, m * n - 1
while left < right:
    mid = (left + right) >> 1
    x, y = divmod(mid, n)
    if matrix[x][y] >= target:
        right = mid
    else:
        left = mid + 1
return matrix[left // n][left % n] ==
target

```

◆ Quick Review

Key Insights

- The matrix is essentially a flattened sorted array due to its properties.
- Use binary search to efficiently determine if the target exists.
- Map the 1D binary search index to the corresponding 2D matrix coordinates using division and modulus operations.
- Achieve $O(\log(m*n))$ time complexity by navigating the matrix as if it were a single sorted list.

Space & Time

Time Complexity: $O(\log(m * n))$

Space Complexity: $O(1)$

Solution

We can treat the matrix as a flattened sorted list without physically copying the elements. Given the total number of

elements = $m * n$, perform binary search on indices ranging from 0 to $(m*n - 1)$. For any index mid , calculate its corresponding row as $mid // n$ and column as $mid \% n$. Compare the element at $matrix[row][col]$ with the target. If it equals target, return true; if it is less than target, search the right half; otherwise, search the left half.

Matrix

#221 Maximal Square

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Matrix

Lists: Grind 169

Problem

Description


Given an $m \times n$ binary matrix filled with 0's and 1's, find the largest square containing only 1's and return its area.

Example 1:

1	0	1	0	0
1	0	1	1	1
1	1	1	1	1
1	0	0	1	0

Input: matrix = `[["1","0","1","0","0"],["1","0","1","1","1"],["1","1","1","1","1"],["1","0","0","1","0"]]`
 Output: 4

Example 2:



Input: matrix = [["0","1"],["1","0"]]
Output: 1

Example 3:

Input: matrix = [["0"]]
Output: 0

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[i].\text{length}$
- $1 \leq m, n \leq 300$
- $\text{matrix}[i][j]$ is '0' or '1'.

Round 2 · High · Medium

p.807

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maximalSquare(self, matrix: list[list[str]])
    -> int:
        m = len(matrix)
        n = len(matrix[0])
        dp = [[0] * n for _ in range(m)]
        maxLength = 0

        for i in range(m):
            for j in range(n):
                if i == 0 or j == 0 or matrix[i][j] == '0':
                    dp[i][j] = 1 if matrix[i][j] == '1' else 0
                else:
                    dp[i][j] = min(dp[i - 1][j - 1], dp[i - 1][j], dp[i][j - 1]) + 1
                    maxLength = max(maxLength, dp[i][j])

        return maxLength * maxLength
```

• LeetCode

Round 2 · High · Medium

p.808


```
class Solution:
    def maximalSquare(self, matrix:
List[List[str]]) -> int:
        m, n = len(matrix), len(matrix[0])
        dp = [[0] * (n + 1) for _ in range(m + 1)]
        mx = 0
        for i in range(m):
            for j in range(n):
                if matrix[i][j] == '1':
                    dp[i + 1][j + 1] = min(dp[i][j]
+ 1], dp[i + 1][j], dp[i][j + 1]) + 1
                    mx = max(mx, dp[i + 1][j + 1])

        return mx * mx
```

• SimplyLeet

Python implementation of the Maximal Square problem

```
def maximalSquare(matrix):
    # Check if matrix is empty
    if not matrix or not matrix[0]:
        return 0
    rows = len(matrix)
    cols = len(matrix[0])
    # DP table with extra row and col for ease of
    boundary conditions
    dp = [[0] * (cols + 1) for _ in range(rows +
1)]
```

```
max_side = 0
# Traverse through each cell in the matrix
for i in range(1, rows + 1):
    for j in range(1, cols + 1):
        # If the current cell is '1'
        if matrix[i - 1][j - 1] == '1':
            # Compute the minimum value among
            the top, left, and top-left neighbors and add 1
            dp[i][j] = min(dp[i - 1][j],
dp[i][j - 1], dp[i - 1][j - 1]) + 1
            max_side = max(max_side, dp[i][j])
        # The area of the square is the square of the
        side length
```

```
return max_side * max_side
```

Example usage:

```
matrix = [
    ["1", "0", "1", "0", "0"],
    ["1", "0", "1", "1", "1"],
    ["1", "1", "1", "1", "1"],
    ["1", "0", "0", "1", "0"]
]
print(maximalSquare(matrix))
```

• LC.ca

```

...
eg. a 10*10 square with full of 1s
    this square move 1 line down
    this square move 1 line right

    so there needs an extra single 1 at bottom
    right, to make it a larger full square
...

class Solution:
    def maximalSquare(self, matrix:
List[List[str]]) -> int:
        m, n = len(matrix), len(matrix[0])

        dp = [[0] * (n + 1) for _ in range(m + 1)]
        mx = 0
        for i in range(m):
            for j in range(n):
                if matrix[i][j] == '1':
                    dp[i + 1][j + 1] = 1 +
min(dp[i][j + 1], dp[i + 1][j], dp[i][j])
                    mx = max(mx, dp[i + 1][j + 1])
        return mx * mx

```

◆ Quick Review

Key Insights

- Use dynamic programming to build a solution where each cell in the DP table represents the side length of the largest square ending at that cell.

- If the current cell is '1', its DP value is the minimum of the top, left, and top-left neighbors plus one.
- Update a global maximum side length during the process to determine the area.
- The square's area is the square of its maximum side length.

Space & Time

Time Complexity: $O(m * n)$, where m is the number of rows and n is the number of columns.

Space Complexity: $O(m * n)$ for the DP table (this can be optimized to $O(n)$ with space reduction techniques).

Solution

We use a two-dimensional dynamic programming approach. We define a DP matrix with one extra row and column to simplify boundary conditions. For each cell (i, j) in the input matrix, if the value is '1', then we calculate the DP value as the minimum of its top, left, and top-left neighboring DP values plus one. This value represents the side length of the square ending at (i, j) . We also keep track of the maximum side length found and finally compute the area of the maximal square by squaring the maximum side length.

Matrix

#311 Sparse Matrix Multiplication

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Matrix

Lists: Premium 98

Problem

Description

Given two sparse matrices `mat1` of size `m x k` and `mat2` of size `k x n`, return the result of `mat1 x mat2`. You may assume that multiplication is always possible.

Example 1:

1	0	0
-1	0	3

×

7	0	0
0	0	0
0	0	1

=

7	0	0
-7	0	3

Input: `mat1 = [[1,0,0],[-1,0,3]]`, `mat2 = [[7,0,0],[0,0,0],[0,0,1]]`
 Output: `[[7,0,0],[-7,0,3]]`

Example 2:

Input: `mat1 = [[0]]`, `mat2 = [[0]]`
 Output: `[[0]]`

Constraints:

- `m == mat1.length`
- `k == mat1[i].length == mat2.length`
- `n == mat2[i].length`
- `1 <= m, n, k <= 100`
- `-100 <= mat1[i][j], mat2[i][j] <= 100`

Community Solutions (Python)

• WalkCC

```
class Solution:
    def multiply(self, mat1: list[list[int]],
                mat2: list[list[int]]) ->
        list[list[int]]:
        m = len(mat1)
        n = len(mat2)
        l = len(mat2[0])
        ans = [[0] * l for _ in range(m)]

        for i in range(m):
            for j in range(l):

                for k in range(n):
                    ans[i][j] += mat1[i][k] * mat2[k][j]

        return ans
```

• LeetCode

```
class Solution:
    def multiply(self, mat1: List[List[int]],
mat2: List[List[int]]) -> List[List[int]]:
        m, n = len(mat1), len(mat2[0])
        ans = [[0] * n for _ in range(m)]
        for i in range(m):
            for j in range(n):
                for k in range(len(mat2)):
                    ans[i][j] += mat1[i][k] *
mat2[k][j]
        return ans
```

• SimplyLeet

```
def sparseMatrixMultiplication(mat1, mat2):
    # Get dimensions of mat1 and mat2
    m, k = len(mat1), len(mat1[0])
    n = len(mat2[0])
    # Initialize result matrix with zeros
    result = [[0] * n for _ in range(m)]

    # Iterate through each row in mat1
    for i in range(m):
        # For each element in row i of mat1
```

```
        for p in range(k):
            if mat1[i][p] != 0: # Only consider
non-zero element in mat1
                # Multiply with corresponding
elements in the row of mat2
                for j in range(n):
                    if mat2[p][j] != 0: # Only
consider non-zero element in mat2
                        result[i][j] += mat1[i][p]
* mat2[p][j]
        return result

# Example usage:
print(sparseMatrixMultiplication([[1,0,0],[-1,0,3]
], [[7,0,0],[0,0,0],[0,0,1]]))
```

• LC.ca

```
class Solution:
    def multiply(self, mat1: List[List[int]],
mat2: List[List[int]]) -> List[List[int]]:
        def f(mat: List[List[int]]) ->
List[List[int]]:
            g = [[] for _ in range(len(mat))]
            for i, row in enumerate(mat):
                for j, x in enumerate(row):
                    if x:
                        g[i].append((j, x))
        return g
```

```

g1 = f(mat1)
g2 = f(mat2)
m, n = len(mat1), len(mat2[0])
ans = [[0] * n for _ in range(m)]
for i in range(m):
    for k, x in g1[i]:
        for j, y in g2[k]:
            ans[i][j] += x * y
return ans

```

◆ Quick Review

Key Insights

- Exploit sparsity by iterating only over non-zero elements.
- For each non-zero element in mat1, multiply it with corresponding non-zero elements in mat2.
- Accumulate the product into the result matrix.
- This approach reduces unnecessary multiplication operations when many elements are zero.

Space & Time

Time Complexity: $O(m * k * n)$ in the worst-case, but significantly lower when many elements are zero.

Space Complexity: $O(m * n)$ for the resulting matrix.

Solution

We loop through each row of mat1 and for every non-zero element, we traverse the corresponding row in mat2. By checking if an element in both matrices is non-zero before multiplying, we avoid redundant calculations. This

optimization leverages the sparsity of the matrices. We use standard nested loops along with conditional checks to implement this efficient multiplication.

Matrix

#498 Diagonal Traverse

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Matrix · Simulation


Lists: CodeSignal

Problem

Description

Given an $m \times n$ matrix `mat`, return an array of all the elements of the array in a diagonal order.

Example 1:



Input: `mat = [[1,2,3],[4,5,6],[7,8,9]]`

Output: `[1,2,4,7,5,3,6,8,9]`

Example 2:

Input: `mat = [[1,2],[3,4]]`

Output: `[1,2,3,4]`

Constraints:

- `m == mat.length`
- `n == mat[i].length`
- `1 <= m, n <= 104`
- `1 <= m * n <= 104`

- $-105 \leq \text{mat}[i][j] \leq 105$

Community Solutions (Python)

• LeetCode

```
class Solution:
    def findDiagonalOrder(self, mat:
List[List[int]]) -> List[int]:
    m, n = len(mat), len(mat[0])
    ans = []
    for k in range(m + n - 1):
        t = []
        i = 0 if k < n else k - n + 1
        j = k if k < n else n - 1
        while i < m and j >= 0:
            t.append(mat[i][j])

            i += 1
            j -= 1
        if k % 2 == 0:
            t = t[::-1]
        ans.extend(t)
    return ans
```

• SimplyLeet

```
def findDiagonalOrder(mat):
    if not mat or not mat[0]:
        return []

    m, n = len(mat), len(mat[0])
    result = []
    row, col = 0, 0
    direction = 1 # 1 means upward, -1 means
downward
```

```
    while len(result) < m * n:

        result.append(mat[row][col])
        if direction == 1: # moving upward
            if col == n - 1: # Hit the right
boundary
                row += 1
                direction = -1
            elif row == 0: # Hit the top boundary
                col += 1
                direction = -1
        else:
            row -= 1
```

```

        col += 1
    else: # moving downward
        if row == m - 1: # Hit the bottom
            boundary
            col += 1
            direction = 1
        elif col == 0: # Hit the left boundary
            row += 1
            direction = 1
        else:
            row += 1

        col -= 1
    return result

# Example usage:
print(findDiagonalOrder([[1,2,3],[4,5,6],[7,8,9]]))

```

• LC.ca

Round 2 · High · Medium

p.823

```

class Solution:
    def findDiagonalOrder(self, mat:
List[List[int]]) -> List[int]:
        m, n = len(mat), len(mat[0])
        ans = []
        for k in range(m + n - 1):
            t = []
            i = 0 if k < n else k - n + 1
            j = k if k < n else n - 1
            while i < m and j >= 0:
                t.append(mat[i][j])

                i += 1
                j -= 1
            if k % 2 == 0:
                t = t[::-1]
            ans.extend(t)
        return ans

```

◆ Quick Review

Key Insights

- Traverse the matrix diagonally with alternating directions: upward (top-right) and downward (bottom-left).
- When the traversal reaches a boundary (top, bottom, left, or right), adjust the position and change the direction.
- Use simulation to process each element exactly once while managing the indices based on the current direction.

Space & Time

Round 2 · High · Medium

p.824

Time Complexity: $O(m * n)$ where m is the number of rows and n is the number of columns, as every element is processed once.

Space Complexity: $O(m * n)$ for the output array (excluding the input matrix), while additional space usage remains constant.

Solution

The solution simulates the diagonal traversal using a direction flag (1 for upward and -1 for downward).

- Begin at the matrix's top-left corner.
- Append the current element to the result.
- Depending on the current direction: If moving upward, decrease the row and increase the column. If moving downward, increase the row and decrease the column.
- Adjust for boundary hits: When moving upward: If at the last column, increment the row and switch direction. If at the first row, increment the column and switch direction. When moving downward: If at the last row, increment the column and switch direction. If at the first column, increment the row and switch direction.

This process continues until all elements are added to the result array.

Matrix

#531 Lonely Pixel I

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Matrix

Lists: Premium 98

Problem

Description

Given an $m \times n$ picture consisting of black 'B' and white 'W' pixels, return the number of black lonely pixels.

A black lonely pixel is a character 'B' that located at a specific position where the same row and same column don't have any other black pixels.

Example 1:

W	W	B
W	B	W
B	W	W

Input: picture =
 [["W","W","B"],["W","B","W"],["B","W","W"]]
 Output: 3
 Explanation: All the three 'B's are black lonely pixels.

Example 2:

B	B	B
B	B	W
B	B	B

Input: picture =
 [["B","B","B"],["B","B","W"],["B","B","B"]]
 Output: 0

Constraints:

- $m == \text{picture.length}$
- $n == \text{picture}[i].\text{length}$
- $1 \leq m, n \leq 500$
- $\text{picture}[i][j]$ is 'W' or 'B'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findLonelyPixel(self, picture:
list[list[str]]) -> int:
    m = len(picture)
    n = len(picture[0])
    ans = 0
    rows = [0] * m # rows[i] := the number of B's
in rows i
    cols = [0] * n # cols[i] := the number of B's
in cols i

    for i in range(m):
        for j in range(n):

            if picture[i][j] == 'B':
                rows[i] += 1
                cols[j] += 1

    for i in range(m):
        if rows[i] == 1: # Only have to examine the
rows if rows[i] == 1.
            for j in range(n):
                # After meeting a 'B' in this rows, break
and search the next row.
                if picture[i][j] == 'B':
                    if cols[j] == 1:
```

```
        ans += 1
        break

    return ans
```

• LeetCode

```
class Solution:
    def findLonelyPixel(self, picture:
List[List[str]]) -> int:
        rows = [0] * len(picture)
        cols = [0] * len(picture[0])
        for i, row in enumerate(picture):
            for j, x in enumerate(row):
                if x == "B":
                    rows[i] += 1
                    cols[j] += 1

        ans = 0

        for i, row in enumerate(picture):
            for j, x in enumerate(row):
                if x == "B" and rows[i] == 1 and
cols[j] == 1:
                    ans += 1

        return ans
```

• SimplyLeet

```
# Function to find the number of lonely black
pixels in a picture.
def findLonelyPixel(picture):
    # Get dimensions of the picture
    rows = len(picture)
    cols = len(picture[0]) if rows > 0 else 0

    # Initialize row and column counts with zeros
    rowCount = [0] * rows
    colCount = [0] * cols

    # First pass: Count the number of 'B's in each
    row and column.
    for i in range(rows):
        for j in range(cols):
            if picture[i][j] == 'B':
                rowCount[i] += 1
                colCount[j] += 1

    # Second pass: Check for lonely black pixels.
    lonelyPixels = 0
    for i in range(rows):
```

```
        for j in range(cols):
            if picture[i][j] == 'B' and
rowCount[i] == 1 and colCount[j] == 1:
                lonelyPixels += 1

    return lonelyPixels

# Example test
picture = [["W","W","B"], ["W","B","W"],
["B","W","W"]]
print(findLonelyPixel(picture)) # Output: 3
```

• LC.ca

```
class Solution:
    def findLonelyPixel(self, picture:
List[List[str]]) -> int:
        m, n = len(picture), len(picture[0])
        rows, cols = [0] * m, [0] * n
        for i in range(m):
            for j in range(n):
                if picture[i][j] == 'B':
                    rows[i] += 1
                    cols[j] += 1

        res = 0
```

```

        for i in range(m):
            if rows[i] == 1:
                for j in range(n):
                    if picture[i][j] == 'B' and
cols[j] == 1:
                        res += 1
                        break
        return res

```

◆ Quick Review

Key Insights

- Count the number of 'B' pixels in each row.
- Count the number of 'B' pixels in each column.
- A pixel is considered lonely if the count in its corresponding row and column is exactly 1.
- Two passes through the matrix can efficiently solve the problem: one for counting and one for checking conditions.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the dimensions of the picture.

Space Complexity: $O(m + n)$, used to store the row and column counts.

Solution

The solution works by first scanning the entire matrix to compute the number of black pixels in each row and each column using two arrays/lists. Once these counts are available, a second pass through the matrix determines for

each black pixel if it is lonely by checking if its corresponding row and column count equals 1. The approach uses simple array data structures for tallying counts and ensures that each element is processed only a constant number of times.

Matrix

#723 Candy Crush

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Matrix · Simulation

Lists: Premium 98

Problem

Description

This question is about implementing a basic elimination algorithm for Candy Crush.

Given an $m \times n$ integer array board representing the grid of candy where $\text{board}[i][j]$ represents the type of candy. A value of $\text{board}[i][j] == 0$ represents that the cell is empty.

The given board represents the state of the game following the player's move. Now, you need to restore the board to a stable state by crushing candies according to the following rules:


- If three or more candies of the same type are adjacent vertically or horizontally, crush them all at the same time – these positions become empty.
- After crushing all candies simultaneously, if an empty space on the board has candies on top of itself, then these candies will drop until they hit a candy or bottom at the same time. No new candies will drop outside the

top boundary.

- After the above steps, there may exist more candies that can be crushed. If so, you need to repeat the above steps.
- If there does not exist more candies that can be crushed (i.e., the board is stable), then return the current board.

You need to perform the above rules until the board becomes stable, then return the stable board.

Example 1:



```
Input: board = [[110,5,112,113,114],[210,211,5,213,214]
,[310,311,3,313,314],[410,411,412,5,414],[5,1,512,3,3],
[610,4,1,613,614],[710,1,2,713,714],[810,1,2,1,1],[1,1,
2,2,2],[4,1,4,4,1014]]
Output: [[0,0,0,0,0],[0,0,0,0,0],[0,0,0,0,0],[110,0,0,0
,114],[210,0,0,0,214],[310,0,0,113,314],[410,0,0,213,41
4],[610,211,112,313,614],[710,311,412,613,714],[810,411
,512,713,1014]]
```

Example 2:

```
Input: board = [[1,3,5,5,2],[3,4,3,3,1],[3,2,4,5,2],[2,
4,4,5,5],[1,4,4,1,1]]
Output: [[1,3,0,0,0],[3,4,0,5,2],[3,2,0,3,1],[2,4,0,5,2
],[1,4,3,1,1]]
```

Constraints:

- $m == \text{board.length}$
- $n == \text{board}[i].\text{length}$
- $3 \leq m, n \leq 50$
- $1 \leq \text{board}[i][j] \leq 2000$

Community Solutions (Python)

- [LeetDoocs](#)

```
class Solution:
    def candyCrush(self, board: List[List[int]])
-> List[List[int]]:
        m, n = len(board), len(board[0])
        run = True
        while run:
            run = False
            for i in range(m):
                for j in range(2, n):
                    if board[i][j] and
abs(board[i][j]) == abs(board[i][j - 1]) == abs(
board[i][j - 2]
):
                        run = True
                        board[i][j] = board[i][j -
1] = board[i][j - 2] = -abs(
board[i][j]
)
                for j in range(n):
                    for i in range(2, m):
                        if board[i][j] and
abs(board[i][j]) == abs(board[i - 1][j]) == abs(
board[i - 2][j]
):
```

```

        run = True
        board[i][j] = board[i -
1][j] = board[i - 2][j] = -abs(
            board[i][j]
        )
    if run:
        for j in range(n):
            k = m - 1
            for i in range(m - 1, -1, -1):
                if board[i][j] > 0:
                    board[k][j] =
board[i][j]
            k -= 1
        while k >= 0:
            board[k][j] = 0
            k -= 1
    return board

```

• SimplyLeet

```

def candyCrush(board):
    # Retrieve board dimensions
    m, n = len(board), len(board[0])
    # Loop until no crushable candies remain
    while True:
        # Set to hold all positions that need to be
        crushed
        crush_set = set()

        # Check horizontally for crushable candies
        for i in range(m):

            for j in range(n - 2):
                # Check if candy exists and three
                consecutive candies are the same
                if board[i][j] != 0 and
                abs(board[i][j]) == abs(board[i][j+1]) ==
                abs(board[i][j+2]):
                    crush_set |= {(i, j), (i,
                    j+1), (i, j+2)}

                # Check vertically for crushable candies
                for i in range(m - 2):
                    for j in range(n):
                        if board[i][j] != 0 and
                        abs(board[i][j]) == abs(board[i+1][j]) ==
                        abs(board[i+2][j]):
                            crush_set |= {(i, j), (i+1,
                            j), (i+2, j)}

```



```

    # If no candies need to be crushed, the
    board is stable; break out of the loop.
    if not crush_set:
        break

    # Set all marked positions to 0 (simulate
    crushing)
    for i, j in crush_set:
        board[i][j] = 0

    # Apply gravity: For each column, move
    uncrushed candies downwards

```

```

    for j in range(n):
        # Pointer for placing the next candy
        from bottom-up
        write_index = m - 1
        # Traverse each column from bottom to
        top
        for i in range(m - 1, -1, -1):
            if board[i][j] != 0:
                board[write_index][j] =
                board[i][j]
                write_index -= 1
        # Fill remaining positions with 0
        for i in range(write_index, -1, -1):

```

```

        board[i][j] = 0

    return board

# Example usage:
board = [[110,5,112,113,114],[210,211,5,213,214],[
310,311,3,313,314],
        [410,411,412,5,414],[5,1,512,3,3],[610,4,
1,613,614],
        [710,1,2,713,714],[810,1,2,1,1],[1,1,2,2,
2],[4,1,4,4,1014]]
print(candyCrush(board))

```

• LC.ca

```

class Solution:
    def candyCrush(self, board: List[List[int]])
    -> List[List[int]]:
        m, n = len(board), len(board[0])
        run = True
        while run:
            run = False
            for i in range(m):
                for j in range(n - 2):
                    if (
                        board[i][j] != 0

```

```

        and abs(board[i][j]) ==
abs(board[i][j + 1])
        and abs(board[i][j]) ==
abs(board[i][j + 2])
    ):
        run = True
        board[i][j] = board[i][j +
1] = board[i][j + 2] = -abs(
            board[i][j]
        )
    for j in range(n):
        for i in range(m - 2):
            if (

```

```

        board[i][j] != 0
        and abs(board[i][j]) ==
abs(board[i + 1][j])
        and abs(board[i][j]) ==
abs(board[i + 2][j])
    ):
        run = True
        board[i][j] = board[i +
1][j] = board[i + 2][j] = -abs(
            board[i][j]
        )
    if run:
        for j in range(n):

```

```

        curr = m - 1
        for i in range(m - 1, -1, -1):
            if board[i][j] > 0:
                board[curr][j] =
board[i][j]
                curr -= 1
            while curr > -1:
                board[curr][j] = 0
                curr -= 1
    return board

```

◆ Quick Review

Key Insights

- Use a simulation approach: repeatedly detect crushable candies and simulate their removal.
- Traverse horizontally and vertically to mark candies that are in groups of three or more.
- Instead of removing candies immediately, mark them for removal and update the board once for the entire pass.
- After crushing, simulate gravity by shifting non-zero candies downward in each column.
- Continue the process until no further candy crushes can be detected.

Space & Time

Time Complexity: $O((m \cdot n)^2)$ in the worst case since each iteration includes scanning the $m \times n$ grid and then applying gravity.

Space Complexity: $O(mn)$ for maintaining additional markers or flags (or using the board itself with in-place modifications).

Solution

The approach involves iteratively scanning the board to identify candy groups that should be crushed using two passes:

- **Horizontal Check** – For each row, check segments of 3 consecutive candies.
- **Vertical Check** – For each column, check segments of 3 consecutive candies. If any candy qualifies, mark its position for crushing. After marking, update the board by turning all marked positions into 0.

Then, for each column, drop down non-zero candies such that all empty cells (zeros) are at the top. Repeat this process until no group of 3 or more adjacent candies is found. Key data structures include the board array and a temporary marker (can be implicit by using signs or a separate boolean matrix) for tracking candies to crush. The algorithm ensures a stable board before returning it.

Matrix

#835 Image Overlap

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Matrix

Lists: CodeSignal

Problem

Description

You are given two images, `img1` and `img2`, represented as binary, square matrices of size $n \times n$. A binary matrix has only 0s and 1s as values.

We translate one image however we choose by sliding all the 1 bits left, right, up, and/or down any number of units. We then place it on top of the other image. We can then calculate the overlap by counting the number of positions that have a 1 in both images.

Note also that a translation does not include any kind of rotation. Any 1 bits that are translated outside of the matrix borders are erased.

Return the largest possible overlap.

Example 1:

1	1	0
0	1	0
0	1	0

img1

0	0	0
0	1	1
0	0	1

img2

Input: `img1 = [[1,1,0],[0,1,0],[0,1,0]]`, `img2 = [[0,0,0],[0,1,1],[0,0,1]]`

Output: 3

Explanation: We translate `img1` to right by 1 unit and down by 1 unit.

The number of positions that have a 1 in both images is 3 (shown in red).

Example 2:

Input: `img1 = [[1]]`, `img2 = [[1]]`

Output: 1

Example 3:

Input: `img1 = [[0]]`, `img2 = [[0]]`

Output: 0

Constraints:

- `n == img1.length == img1[i].length`

- `n == img2.length == img2[i].length`
- `1 <= n <= 30`
- `img1[i][j]` is either 0 or 1.
- `img2[i][j]` is either 0 or 1.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def largestOverlap(self, img1: list[list[int]],
                      img2: list[list[int]]) -> int:
        MAGIC = 100
        ones1 = [(i, j)
                  for i, row in enumerate(img1)
                  for j, num in enumerate(row)
                  if num == 1]
        ones2 = [(i, j)
                  for i, row in enumerate(img2)
                  for j, num in enumerate(row)
                  if num == 1]

        offsetCount = collections.Counter()

        for ax, ay in ones1:
            for bx, by in ones2:
                offsetCount[(ax - bx) * MAGIC + (ay - by)]
                += 1

        return max(offsetCount.values()) if
        offsetCount else 0
```

• LeetCode

```
class Solution:
    def largestOverlap(self, img1:
List[List[int]], img2: List[List[int]]) -> int:
        n = len(img1)
        cnt = Counter()
        for i in range(n):
            for j in range(n):
                if img1[i][j]:
                    for h in range(n):
                        for k in range(n):
                            if img2[h][k]:
                                cnt[(i - h, j -
k)] += 1
        return max(cnt.values()) if cnt else 0
```

• SimplyLeet

```
def largestOverlap(img1, img2):
    # Get the dimension of the square images
    n = len(img1)

    # Gather positions (coordinates) where there is
    a 1 in img1 and img2
    ones_img1 = []
    ones_img2 = []
    for i in range(n):
        for j in range(n):
            if img1[i][j] == 1:
```

```
        ones_img1.append((i, j))
        if img2[i][j] == 1:
            ones_img2.append((i, j))

    # Dictionary to count translation shifts that
    align 1's of img1 with img2
    translation_count = {}
    max_overlap = 0

    # Iterate over every pair of 1's from both
    images
    for (x1, y1) in ones_img1:
        for (x2, y2) in ones_img2:
            # Calculate the translation vector
            needed to align img1's 1 with img2's 1
            delta = (x2 - x1, y2 - y1)
            # Update the count for this translation
            vector
            translation_count[delta] =
translation_count.get(delta, 0) + 1
            # Update max_overlap if we find a
            higher count
            max_overlap = max(max_overlap,
translation_count[delta])

    return max_overlap
```

Example usage:

```
img1 = [[1,1,0],[0,1,0],[0,1,0]]
img2 = [[0,0,0],[0,1,1],[0,0,1]]
print(largestOverlap(img1, img2)) # Expected
output: 3
```

• LC.ca

```
class Solution:
    def largestOverlap(self, img1:
List[List[int]], img2: List[List[int]]) -> int:
        n = len(img1)
        cnt = Counter()
        for i in range(n):
            for j in range(n):
                if img1[i][j]:
                    for h in range(n):
                        for k in range(n):
                            if img2[h][k]:
                                cnt[(i - h, j -
k)] += 1
        return max(cnt.values()) if cnt else 0
```

◆ Quick Review

Key Insights

- You only need to consider the positions of 1's in each matrix.
- The overlap count for a given translation can be computed by comparing the relative differences between

the positions of 1's in the two images.

- Using a hash map (or dictionary) to count how many times each translation vector appears lets you compute the maximum overlap efficiently.
- The translation operation is equivalent to shifting the positions, so the best possible translation is the one that aligns the most pairs of 1's.

Space & Time

Time Complexity: $O(n^4)$ in the worst-case when using a nested iteration over positions for both matrices, though the average situation is typically better since not all cells are 1.

Space Complexity: $O(n^2)$ for storing the positions of 1's and the counts for translation vectors.

Solution

The solution involves the following steps:

- Traverse both images to record the coordinates of all 1's.
- For every pair of 1's (one from the first image, one from the second), compute the translation vector (offset) that aligns them.
- Use a hash map to count how many times each offset appears.
- The maximum count recorded in the hash map is the largest overlap achievable by any translation. This approach avoids simulating every possible slide over the matrix and focuses only on relative alignments of 1's.

Matrix

#1198 Find Smallest Common Element in All Rows

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Binary Search · Matrix · Counting Lists: Premium 98

Problem

Description

Given an $m \times n$ matrix `mat` where every row is sorted in strictly increasing order, return the smallest common element in all rows.

If there is no common element, return -1 .

Example 1:

Input: `mat =`
`[[1,2,3,4,5],[2,4,5,8,10],[3,5,7,9,11],[1,3,5,7,9]]`
 Output: 5

Example 2:

Input: `mat = [[1,2,3],[2,3,4],[2,3,5]]`
 Output: 2

Constraints:

- $m == \text{mat.length}$
- $n == \text{mat}[i].\text{length}$
- $1 \leq m, n \leq 500$

- $1 \leq \text{mat}[i][j] \leq 104$
- `mat[i]` is sorted in strictly increasing order.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def smallestCommonElement(self, mat:
list[list[int]]) -> int:
        MAX = 10000
        count = [0] * (MAX + 1)

        for row in mat:
            for a in row:
                count[a] += 1
                if count[a] == len(mat):
                    return a

        return -1
```

• LeetDoocs

```
class Solution:
    def smallestCommonElement(self, mat:
List[List[int]]) -> int:
        cnt = Counter()
        for row in mat:
            for x in row:
                cnt[x] += 1
            if cnt[x] == len(mat):
                return x
        return -1
```

• SimplyLeet

```
def smallestCommonElement(mat):
    # Get number of rows and columns
    m = len(mat)
    n = len(mat[0])

    # Helper function: binary search for target in
a row
    def binary_search(row, target):
        left, right = 0, n - 1
        while left <= right:
            mid = (left + right) // 2
```

```
        if row[mid] == target:
            return True
        elif row[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return False
```

```
# Iterate over each candidate in the first row
for candidate in mat[0]:
```

```
        is_common = True
        # Check each subsequent row for the
candidate using binary search
        for i in range(1, m):
            if not binary_search(mat[i],
candidate):
                is_common = False
                break
        if is_common:
            return candidate
    return -1
```

```
# Example usage:
print(smallestCommonElement([[1,2,3,4,5],[2,4,5,8,
10],[3,5,7,9,11],[1,3,5,7,9]]))
```

• LC.ca


```

class Solution:
    def smallestCommonElement(self, mat:
List[List[int]]) -> int:
        cnt = Counter()
        for row in mat:
            for x in row:
                cnt[x] += 1
                if cnt[x] == len(mat):
                    return x
        return -1

```

◆ Quick Review

Key Insights

- Since every row is strictly increasing, binary search can be used efficiently to check for the presence of an element in a row.
- Iterating over the first row's elements and verifying their presence in all other rows is an effective approach.
- Early termination is possible if an element is not found in any one row.
- Time complexity is manageable given the constraints, even using binary search for each candidate across rows.

Space & Time

Time Complexity: $O(m * n * \log(n))$ where m is the number of rows and n is the number of columns (binary search in each row for each candidate from the first row).

Space Complexity: $O(1)$ (only a few auxiliary variables are used).

Solution

The solution iterates over each element (candidate) in the first row of the matrix. For each candidate, it checks if the candidate is present in every other row using binary search, taking advantage of the sorted order. If a candidate is found in every row, return it immediately as it will be the smallest such common element because the first row is processed in order. If no candidate is common across all rows, return -1 .

Trees & BST

Round 2 · High · Medium · 24 questions

Round 2 · High · Medium

p.859

Trees & BST

#98 Validate Binary Search Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BST

Lists: Grind 169

Problem

Description

Given the root of a binary tree, determine if it is a valid binary search tree (BST).


A valid BST is defined as follows:

- The left subtree of a node contains only nodes with keys strictly less than the node's key.
- The right subtree of a node contains only nodes with keys strictly greater than the node's key.
- Both the left and right subtrees must also be binary search trees.

Example 1:


Round 2 · High · Medium

p.860



Input: root = [2,1,3]
Output: true

Example 2:



Input: root = [5,1,4,null,null,3,6]
Output: false
Explanation: The root node's value is 5 but its right child's value is 4.

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- $-231 \leq \text{Node.val} \leq 231 - 1$

Community Solutions (Python)

- WalkCC

```

class Solution:
    def isValidBST(self, root: TreeNode | None) ->
    bool:
        def isValidBST(root: TreeNode | None,
                        minNode: TreeNode | None,
                        maxNode: TreeNode | None) -> bool:
            if not root:
                return True
            if minNode and root.val <= minNode.val:
                return False
            if maxNode and root.val >= maxNode.val:
                return False

            return (isValidBST(root.left, minNode, root)
and
                    isValidBST(root.right, root,
maxNode))

        return isValidBST(root, None, None)

```

```

class Solution:
    def isValidBST(self, root: TreeNode | None) ->
    bool:
        stack = []
        pred = None

        while root or stack:
            while root:
                stack.append(root)
                root = root.left
            root = stack.pop()

            if pred and pred.val >= root.val:
                return False
            pred = root
            root = root.right

        return True

```

- LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def isValidBST(self, root: Optional[TreeNode])
-> bool:
        def dfs(root: Optional[TreeNode]) -> bool:
            if root is None:
                return True
            if not dfs(root.left):
                return False
            nonlocal prev
            if prev >= root.val:
                return False
            prev = root.val
            return dfs(root.right)

        prev = -inf

        return dfs(root)
```

• SimplyLeet

Round 2 · High · Medium

p.865

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # Node's value
        self.left = left # Left child
        self.right = right # Right child

class Solution:
    def isValidBST(self, root: TreeNode) -> bool:
        # Helper function to recursively validate
        nodes

        def validate(node, low, high):
            if not node:
                return True # An empty
node/subtree is valid
            # If current node's value is out of the
            allowed range, return False
            if not (low < node.val < high):
                return False
            # Recursively check the left subtree
            and right subtree
            return validate(node.left, low,
node.val) and validate(node.right, node.val, high)

        return validate(root, float("-inf"),
float("inf"))
```

• LC.ca

Round 2 · High · Medium

p.866

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val

# self.left = left
# self.right = right

class Solution: # recursive
    def isValidBST(self, root: Optional[TreeNode])
-> bool:
        def dfs(root, mi, ma):
            if not root:
                return True
            return mi < root.val < ma and
dfs(root.left, mi, root.val) and dfs(root.right,
root.val, ma)

```

```

        return dfs(root, -math.inf, math.inf)

class Solution: # iterative
    def isValidBST(self, root: TreeNode) -> bool:
        # store (node, lower, upper) tuples in a
single queue
        queue = [(root, None, None)]

        # another option:
        # queue = [(root, -math.inf, math.inf)]

while queue:
    node, lower, upper = queue.pop(0)

    if node is None:
        continue

    val = node.val
    if lower is not None and val <= lower:
        return False

```

```

        if upper is not None and val >= upper:
            return False

```

```

        queue.append((node.right, val, upper))
        queue.append((node.left, lower, val))

```

```

    return True

```

```

#####

```

```

class Solution:
    def isValidBST(self, root: Optional[TreeNode])
    -> bool:
        def isValidBST(root: Optional[TreeNode],
                        minNode: Optional[TreeNode],
maxNode: Optional[TreeNode]) -> bool:
            if not root:
                return True
            if minNode and root.val <= minNode.val:
                return False
            if maxNode and root.val >= maxNode.val:
                return False

```

```

        return isValidBST(root.left, minNode, root)
        and isValidBST(root.right, root, maxNode)

```

```

        return isValidBST(root, None, None)

```

```

#####

```

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):

```

```

# self.val = val
# self.left = left
# self.right = right
class Solution:
    def isValidBST(self, root: TreeNode) -> bool:
        def dfs(root):
            nonlocal prev
            if root is None:
                return True
            if not dfs(root.left):

```

```

        return False
    if prev >= root.val:
        return False
    prev = root.val
    if not dfs(root.right):
        return False
    return True

prev = -inf
return dfs(root)

```

```
#####
```

```

# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

```

```

class Solution(object):
    def isValidBST(self, root):
        """
        :type root: TreeNode
        :rtype: bool
        """
        prev = -float("inf")
        stack = [(1, root)]
        while stack:
            p = stack.pop()

            if not p[1]:
                continue
            if p[0] == 0:
                if p[1].val <= prev:
                    return False
                prev = p[1].val
            else:
                stack.append((1, p[1].right))
                stack.append((0, p[1]))
                stack.append((1, p[1].left))

        return True

```

◆ Quick Review

Key Insights

- Every node must adhere to a range: all nodes in the left subtree should be less than the current node, and all nodes in the right subtree should be greater.

- A recursive depth-first search approach works well by passing down the valid range (lower and upper bounds) for each node.
- Using the limits (negative and positive infinity) simplifies handling edge cases at the tree root.
- An in-order traversal would alternatively yield a sorted order for node values, but updating ranges directly is more intuitive for recursive implementation.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes, as each node is visited once.

Space Complexity: $O(H)$, where H is the height of the tree, due to the recursion stack.

Solution

We use a recursive approach to validate the BST property. We maintain two variables, low and high, representing the valid range for the current node's value. For the left child, the valid range becomes (low, node.val) and for the right child, it becomes (node.val, high). By checking if each node's value falls within its respective range, we can ensure that the tree satisfies the BST properties.

Trees & BST

#102 Binary Tree Level Order Traversal

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · BFS


Lists: Grind 169

Problem

Description

Given the root of a binary tree, return the level order traversal of its nodes' values. (i.e., from left to right, level by level).

Example 1:



Input: root = [3,9,20,null,null,15,7]
 Output: [[3],[9,20],[15,7]]

Example 2:

Input: root = [1]
 Output: [[1]]

Example 3:

Input: root = []
 Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 2000].
- $-1000 \leq \text{Node.val} \leq 1000$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def levelOrder(self, root: TreeNode | None) -> list[list[int]]:
        if not root:
            return []

        ans = []
        q = collections.deque([root])

        while q:
            currLevel = []

            for _ in range(len(q)):
                node = q.popleft()
                currLevel.append(node.val)
                if node.left:
                    q.append(node.left)
                if node.right:
                    q.append(node.right)
            ans.append(currLevel)

        return ans
  
```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def levelOrder(self, root: Optional[TreeNode])
-> List[List[int]]:
        ans = []
        if root is None:

            return ans
        q = deque([root])
        while q:
            t = []
            for _ in range(len(q)):
                node = q.popleft()
                t.append(node.val)
                if node.left:
                    q.append(node.left)
                if node.right:

                    q.append(node.right)
            ans.append(t)
        return ans

```

• SimplyLeet

Round 2 · High · Medium

p.877

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

from collections import deque

class Solution:

    def levelOrder(self, root: TreeNode) ->
List[List[int]]:
        # Return an empty list if the tree is
empty.

        if not root:
            return []

        result = [] # Final result list containing
values of each level.
        queue = deque([root]) # Initialize the
queue with the root node.

        # Process nodes level by level.
        while queue:

```

Round 2 · High · Medium

p.878

```

        level_size = len(queue) # Number of
nodes at the current level.
        level_nodes = [] # List to hold node
values at the current level.

        for _ in range(level_size):
            node = queue.popleft() # Dequeue
the next node.
            level_nodes.append(node.val)
            # Enqueue left child if it exists.
            if node.left:
                queue.append(node.left)
            # Enqueue right child if it exists.

            if node.right:
                queue.append(node.right)
        # Append the current level's values to
the result.
        result.append(level_nodes)
    return result

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def levelOrder(self, root: Optional[TreeNode])
-> List[List[int]]:
        ans = []
        if root is None:

            return ans
        q = deque([root])
        while q:
            t = []
            for _ in range(len(q)):
                node = q.popleft()
                t.append(node.val)
                if node.left:
                    q.append(node.left)
                if node.right:

                    q.append(node.right)
            ans.append(t)
        return ans

```

◆ Quick Review

Key Insights

- Use Breadth-First Search (BFS) to traverse the tree.
- A queue data structure is ideal for maintaining the order of nodes at each level.
- For every level, process all nodes currently in the queue and add their children for the next level.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is processed once.

Space Complexity: $O(n)$, which accounts for the space used by the queue in the worst-case scenario.

Solution

This solution leverages a BFS approach using a queue. The algorithm starts by enqueueing the root node. Then, it enters a loop that continues until the queue is empty. For each iteration—representing one level of the tree—it computes the number of nodes at that level (`levelSize`), processes these nodes by dequeuing them, and collects their values. Their non-null children are then enqueued. The list of values for each level is added to the final result list. Key data structures include the queue (for managing nodes) and a list (to store level values).

Trees & BST

#103 Binary Tree Zigzag Level Order Traversal

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · BFS


Lists: Grind 169

Problem

Description

Given the root of a binary tree, return the zigzag level order traversal of its nodes' values. (i.e., from left to right, then right to left for the next level and alternate between).

Example 1:



Input: root = [3,9,20,null,null,15,7]
 Output: [[3],[20,9],[15,7]]

Example 2:

Input: root = [1]
 Output: [[1]]

Example 3:

Input: root = []
 Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 2000].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def zigzagLevelOrder(self, root: TreeNode | None) -> list[list[int]]:
        if not root:
            return []

        ans = []
        dq = collections.deque([root])
        isLeftToRight = True

        while dq:
            currLevel = []
            for _ in range(len(dq)):
                if isLeftToRight:
                    node = dq.popleft()
                    currLevel.append(node.val)
                    if node.left:
                        dq.append(node.left)
                    if node.right:
                        dq.append(node.right)
                else:

```

```

        node = dq.pop()
        currLevel.append(node.val)
        if node.right:
            dq.appendleft(node.right)
        if node.left:
            dq.appendleft(node.left)
        ans.append(currLevel)
        isLeftToRight = not isLeftToRight

    return ans

```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def zigzagLevelOrder(self, root: Optional[TreeNode]) -> List[List[int]]:
        ans = []
        if root is None:

```

```

        return ans
        q = deque([root])
        ans = []
        left = 1
        while q:
            t = []
            for _ in range(len(q)):
                node = q.popleft()
                t.append(node.val)
                if node.left:

```

```

                    q.append(node.left)
                    if node.right:
                        q.append(node.right)
            ans.append(t if left else t[::-1])
            left ^= 1
        return ans

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right

def zigzagLevelOrder(root):
    # If the tree is empty, return an empty list
    if not root:

```

```

        return []

    # Initialize the result list and the queue with
    the root node.
    result = []
    queue = [root]
    left_to_right = True # Flag to indicate the
    current direction of traversal

    # Process the tree level by level.
    while queue:
        level_size = len(queue)

        current_level = []

        # Process all nodes for the current level.
        for i in range(level_size):
            node = queue.pop(0) # Dequeue the node
            from the front of the queue
            current_level.append(node.val) #
            Append the node's value

            # Enqueue the left child if it exists.
            if node.left:
                queue.append(node.left)

```

```

        # Enqueue the right child if it exists.
        if node.right:
            queue.append(node.right)

        # If the current level's order is
        right-to-left, reverse the list.
        if not left_to_right:
            current_level.reverse()

        # Append the current level's values to the
        result.
        result.append(current_level)

        # Toggle the direction for the next level.
        left_to_right = not left_to_right

    return result

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
...
can also use list.insert()

```



```
>>> my_list = [2, 3, 4]
>>> my_list.insert(0, 1) # Insert 1 at the head of
the list
>>> print(my_list) # Output: [1, 2, 3, 4]
```

```
>>> a = deque([])
>>> a
deque([])
>>> a.append(1)
>>> a.append(2)
```

```
>>> a.append(3)
>>> a
deque([1, 2, 3])
>>>
>>> a.append(0, 555)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: deque.append() takes exactly one
argument (2 given)
>>> a.insert(0, 555)
>>> a
```

```
deque([555, 1, 2, 3])
'''
```

```
from collections import deque
```

```
class Solution:
    def zigzagLevelOrder(self, root:
Optional[TreeNode]) -> List[List[int]]:
        ans = []
        if root is None:
            return ans
```

```
        q = deque([root])
        ans = []
        left = True
        while q:
            t = []
            for _ in range(len(q)):
                node = q.popleft()
                t.append(node.val)
                if node.left:
                    q.append(node.left)
```

```

        if node.right:
            q.append(node.right)
        ans.append(t if left else t[::-1])
        left = (not left)
    return ans

```

```
#####
```

```

# Definition for a binary tree node.
# class TreeNode(object):

```

```

# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None
from collections import deque

```

```

class Solution(object):
    def zigzagLevelOrder(self, root):
        """

```

```

:type root: TreeNode
:rtype: List[List[int]]
"""
stack = deque([root])
ans = []
odd = True
while stack:
    level = []
    for k in range(0, len(stack)):
        top = stack.popleft()

```

```

        if top is None:
            continue
        level.append(top.val)
        stack.append(top.left)
        stack.append(top.right)
    if level:
        if odd:
            ans.append(level)
        else:
            ans.append(level[::-1])

```

```

    odd = not odd
    return ans

```

```

#####

```

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left

```

```

# self.right = right
class Solution:
    def zigzagLevelOrder(self, root:
Optional[TreeNode]) -> List[List[int]]:
    ans = []
    if root is None:
        return ans
    q = deque([root])
    ans = []
    left = 1
    while q:

```

```

    t = []
    for _ in range(len(q)):
        node = q.popleft()
        t.append(node.val)
        if node.left:
            q.append(node.left)
        if node.right:
            q.append(node.right)
    ans.append(t if left else t[::-1])
    left ^= 1

```

```

    return ans

```

◆ Quick Review

Key Insights

- Use a breadth-first search (BFS) approach to traverse the tree level by level.
- Alternate the direction of traversal on each level to achieve the zigzag pattern.
- A flag or counter can be used to determine the order of nodes for the current level.
- Using a double-ended data structure or simply reversing the list at every alternate level is an effective approach.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is visited exactly once.

Space Complexity: $O(n)$, as in the worst-case scenario (e.g., a complete binary tree) the queue used for BFS may hold

$O(n)$ nodes at once.

Solution

We can solve this problem using a breadth-first search (BFS). Start by initializing a queue to hold nodes for each level. For each level:

- Determine the number of nodes at the current level.
- Iterate over these nodes, appending their children for the next level.
- Depending on a flag (or level count), either append node values in the natural order (left to right) or reverse the order (right to left) before adding to the result.

Key data structures:

- A queue for BFS traversal.
- A list (or similar container) to store current level values. The main trick is toggling the order of insertion or reversing the list at alternate levels to achieve the "zigzag" pattern.

Trees & BST

#105 Construct Binary Tree from Preorder and Inorder Traversal

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Tree · DFS


Lists: Grind 169

Problem

Description

Given two integer arrays `preorder` and `inorder` where `preorder` is the preorder traversal of a binary tree and `inorder` is the inorder traversal of the same tree, construct and return the binary tree.

Example 1:



Input: preorder = [3,9,20,15,7], inorder =
[9,3,15,20,7]
Output: [3,9,20,null,null,15,7]

Example 2:

Input: preorder = [-1], inorder = [-1]
Output: [-1]

Constraints:

- $1 \leq \text{preorder.length} \leq 3000$
- $\text{inorder.length} == \text{preorder.length}$
- $-3000 \leq \text{preorder}[i], \text{inorder}[i] \leq 3000$

Round 2 · High · Medium

p.897

- preorder and inorder consist of unique values.
- Each value of inorder also appears in preorder.
- preorder is guaranteed to be the preorder traversal of the tree.
- inorder is guaranteed to be the inorder traversal of the tree.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def buildTree(
        self,
        preorder: list[int],
        inorder: list[int],
    ) -> TreeNode | None:
        inToIndex = {num: i for i, num in
            enumerate(inorder)}

        def build(
            preStart: int,

```

Round 2 · High · Medium

p.898

```

    preEnd: int,
    inStart: int,
    inEnd: int,
) -> TreeNode | None:
    if preStart > preEnd:
        return None

```

```

rootVal = preorder[preStart]
rootInIndex = inToIndex[rootVal]
leftSize = rootInIndex - inStart

```

```

root = TreeNode(rootVal)
root.left = build(preStart + 1, preStart +
leftSize,
                    inStart, rootInIndex - 1)
root.right = build(preStart + leftSize + 1,
                    preEnd, rootInIndex + 1,
inEnd)
return root

return build(0, len(preorder) - 1, 0,
len(inorder) - 1)

```

• LeetCode

Round 2 · High · Medium

p.899

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def buildTree(self, preorder: List[int],
inorder: List[int]) -> Optional[TreeNode]:
        def dfs(i: int, j: int, n: int) ->
Optional[TreeNode]:
            if n <= 0:

```

```

                return None
                v = preorder[i]
                k = d[v]
                l = dfs(i + 1, j, k - j)
                r = dfs(i + 1 + k - j, k + 1, n - k + j
- 1)

                return TreeNode(v, l, r)

        d = {v: i for i, v in enumerate(inorder)}
        return dfs(0, 0, len(preorder))

```

Round 2 · High · Medium

p.900

```

class Solution:
    def getBinaryTrees(self, preOrder: List[int],
inOrder: List[int]) -> List[TreeNode]:
        def dfs(i: int, j: int, n: int) ->
List[TreeNode]:
            if n <= 0:
                return [None]
            v = preOrder[i]
            ans = []
            for k in d[v]:
                if j <= k < j + n:
                    for l in dfs(i + 1, j, k - j):
                        for r in dfs(i + 1 + k - j,
k + 1, n - 1 - (k - j)):
                            ans.append(TreeNode(v,
l, r))
            return ans

            d = defaultdict(list)
            for i, x in enumerate(inOrder):
                d[x].append(i)
            return dfs(0, 0, len(preOrder))

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def buildTree(self, preorder: list[int],
inorder: list[int]) -> TreeNode:
        # Create a hashmap to store value->index
        for inorder traversal

            inorder_index = {value: idx for idx, value
in enumerate(inorder)}

            # Recursive helper function to build the
tree
            def helper(pre_start, in_left, in_right):
                # Base case: if there are no elements
to construct subtree
                if in_left > in_right:
                    return None

                # Pick up pre_start element as a root
                root_val = preorder[pre_start]

```

```

    root = TreeNode(root_val)

    # Get the inorder index of the root
    in_index = inorder_index[root_val]

    # Number of nodes in left subtree
    left_subtree_size = in_index - in_left

    # Recursively build the left subtree
    root.left = helper(pre_start + 1,
in_left, in_index - 1)

    # Recursively build the right subtree
    root.right = helper(pre_start +
left_subtree_size + 1, in_index + 1, in_right)

    return root

    # Initiate recursion from the start of
preorder and full inorder range
    return helper(0, 0, len(inorder) - 1)

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right

# best
class Solution:
    def buildTree(self, preorder: List[int],
inorder: List[int]) -> Optional[TreeNode]:

        if not preorder: # or, not inorder
            return None
        v = preorder[0]
        i = inorder.index(v)
        root = TreeNode(v)
        root.left = self.buildTree(preorder[1:
i+1], inorder[:i])
        root.right =
self.buildTree(preorder[i+1:], inorder[i+1:])
        return root

```



```
# inorder.index(preorder_val)
class Solution:
    def buildTree(self, preorder: List[int],
inorder: List[int]) -> Optional[TreeNode]:
        def dfs(pleft, pright, ileft,  iright):
            if pleft > pright or ileft > iright:
                return None
            k = inorder.index(preorder[pleft])
            root = TreeNode(preorder[pleft])
            root.left = dfs(pleft + 1, pleft + (k -
ileft), ileft, k - 1)
            root.right = dfs(pleft + 1 + (k -
ileft), pright, k + 1, iright)
```

```
        return root
```

```
        n = len(preorder)
        return dfs(0, n - 1, 0, n - 1)
```

```
# build dict for val => index
```

```
class Solution:
    def buildTree(self, preorder: List[int],
inorder: List[int]) -> Optional[TreeNode]:
        def dfs(pleft, pright, ileft,  iright):
```

```
            if pleft > pright or ileft > iright:
                return None
            v = preorder[pleft]
            k = d[v]
            root = TreeNode(v)
            root.left = dfs(pleft + 1, pleft + (k -
ileft), ileft, k - 1)
            root.right = dfs(pleft + 1 + (k -
ileft), pright, k + 1, iright)
            return root
```

```
        d = {v: i for i, v in enumerate(inorder)}
```

```
        n = len(preorder)
        return dfs(0, n - 1, 0, n - 1)
```

```
#####
```

```
...
```

```
search for index, also use .index(val)
```

```
>>> a = [1,2,3,4,5]
```

```
>>> a.index(3)
```

```
2
```

```

'''
class Solution(object):
    def buildTree(self, preorder, inorder):
        """
        :type preorder: List[int]
        :type inorder: List[int]
        :rtype: TreeNode
        """

        self.preindex = 0

        ind = {v: i for i, v in enumerate(inorder)}
        head = self.dc(0, len(preorder) - 1, preorder,
inorder, ind)
        return head

    def dc(self, start, end, preorder, inorder, ind):
        if start <= end:
            mid = ind[preorder[self.preindex]]
            self.preindex += 1
            root = TreeNode(inorder[mid])
            root.left = self.dc(start, mid - 1, preorder,
inorder, ind)

```

Round 2 · High · Medium

p.907

```

        root.right = self.dc(mid + 1, end, preorder,
inorder, ind)
        return root

#####

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left

# self.right = right
class Solution:
    def buildTree(self, preorder: List[int],
inorder: List[int]) -> Optional[TreeNode]:
        def dfs(i: int, j: int, n: int):
            if n <= 0:
                return None
            v = preorder[i]
            k = d[v]
            l = dfs(i + 1, j, k - j)
            r = dfs(i + 1 + k - j, k + 1, n - k + j
- 1)

            return TreeNode(v, l, r)

        d = {v: i for i, v in enumerate(inorder)}
        return dfs(0, 0, len(preorder))

```

Round 2 · High · Medium

p.908

◆ Quick Review

Key Insights

- Preorder traversal always starts with the root node.
- Inorder traversal splits the tree into left subtree and right subtree based on the root.
- A hash table (map/dictionary) can be used to quickly locate the index of the root in the inorder array.
- Use recursion to construct left and right subtrees.

Space & Time

Time Complexity: $O(n)$ since we process every node exactly once.

Space Complexity: $O(n)$ due to the recursion stack and the hash map used for inorder indices.

Solution

We solve the problem using a recursive divide-and-conquer approach. First, create a hash map to store each value's index in the inorder array for quick lookups. The preorder array always provides the current root. Use the inorder array to divide the tree into left and right subtrees. Recursively construct the subtrees by adjusting the preorder index and the boundaries of the inorder array. The algorithm leverages recursion and a hash map to achieve $O(n)$ time complexity.

Trees & BST

#199 Binary Tree Right Side View

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BFS

Lists: Grind 169

Problem

Description


Given the root of a binary tree, imagine yourself standing on the right side of it, return the values of the nodes you can see ordered from top to bottom.

Example 1:

Input: root = [1,2,3,null,5,null,4]

Output: [1,3,4]

Explanation:




Example 2:

Input: root = [1,2,3,4,null,null,null,5]

Output: [1,3,4,5]

Explanation:



Round 2 · High · Medium

p.911

Example 3:

Input: root = [1,null,3]

Output: [1,3]

Example 4:

Input: root = []

Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 100].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def rightSideView(self, root: TreeNode | None) -> list[int]:
        if not root:
            return []

        ans = []
        q = collections.deque([root])

        while q:
            size = len(q)
```

Round 2 · High · Medium

p.912

```

for i in range(size):
    root = q.popleft()
    if i == size - 1:
        ans.append(root.val)
    if root.left:
        q.append(root.left)
    if root.right:
        q.append(root.right)

return ans

```

```

class Solution:
    def rightSideView(self, root: TreeNode | None) -> list[int]:
        ans = []

        def dfs(root: TreeNode | None, depth: int) -> None:
            if not root:
                return

            if depth == len(ans):
                ans.append(root.val)

            dfs(root.right, depth + 1)
            dfs(root.left, depth + 1)

        dfs(root, 0)
        return ans

```

• LeetCode

Round 2 · High · Medium

p.913

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def rightSideView(self, root: Optional[TreeNode]) -> List[int]:
        ans = []
        if root is None:

```

```

            return ans
        q = deque([root])
        while q:
            ans.append(q[0].val)
            for _ in range(len(q)):
                node = q.popleft()
                if node.right:
                    q.append(node.right)
                if node.left:
                    q.append(node.left)

```

```

        return ans

```

• SimplyLeet

Round 2 · High · Medium

p.914

```
# Python solution using BFS
from collections import deque
```

```
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # Node value
        self.left = left # Left child
        self.right = right # Right child
```

```
def rightSideView(root):
```

```
    # If the tree is empty, return an empty list
    if not root:
        return []
```

```
    right_view = [] # List to store right side view
    queue = deque([root]) # Initialize queue with
the root node
```

```
    while queue:
        level_size = len(queue) # Number of nodes
in current level
        for i in range(level_size):
```

```
            node = queue.popleft() # Get the front
node from the queue
            # if it's the last node in the current
level, add it.
            if i == level_size - 1:
                right_view.append(node.val)
            # add left child if exists
            if node.left:
                queue.append(node.left)
            # add right child if exists
            if node.right:
                queue.append(node.right)
```

```
    return right_view
```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def rightSideView(self, root:
Optional[TreeNode]) -> List[int]:
        ans = []
        if root is None:
```

```

        return ans
    q = deque([root])
    while q:
        ans.append(q[-1].val) # add last node
of previous level traversal results
        for _ in range(len(q)):
            node = q.popleft()
            if node.left:
                q.append(node.left)
            if node.right:
                q.append(node.right)

```

```

    return ans

```

```

#####

```

```

# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

```

```

from collections import deque

```

```

class Solution(object):
    def rightSideView(self, root):
        """
        :type root: TreeNode
        :rtype: List[int]
        """

```

```

    def dfs(root, h):
        if root:
            if h == len(ans):
                ans.append(root.val)
                # pre-order, all the way to the right
                dfs(root.right, h + 1)
                dfs(root.left, h + 1)

```

```

    ans = []
    dfs(root, 0)

```

```

    return ans

```

◆ Quick Review

Key Insights

- The problem is essentially about capturing the last node at each level (or the rightmost node) in a binary tree.
- A Breadth-First Search (BFS) approach naturally works by processing nodes level by level.

- Alternatively, Depth-First Search (DFS) can be adapted by prioritizing the right subtree first, ensuring that the first node encountered at each depth is the visible one.
- The number of nodes is capped at 100, so both approaches are efficient.

Space & Time

Time Complexity: $O(n)$ – we visit every node in the tree.

Space Complexity: $O(n)$ – in the worst case, the queue (or call stack for DFS) may hold all nodes of the tree.

Solution

We can solve this problem using BFS (level-order traversal). The idea is to traverse the tree level by level using a queue. For each level, the last element added to the level represents the rightmost view from that level. We then append that element's value to our result list.

Alternatively, a DFS approach can be used where we visit the right subtree first. By keeping track of the current depth and checking if it's the first time we've encountered this depth, we can add the first node visited at that depth (which will be the rightmost one) to the result.

In both approaches, appropriate data structures (a queue for BFS or recursion with a list for DFS) are used to store nodes and track the traversal state.

Trees & BST

#230 Kth Smallest Element in a BST MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BST


Lists: Grind 169

Problem

Description

Given the root of a binary search tree, and an integer k , return the k th smallest value (1-indexed) of all the values of the nodes in the tree.

Example 1:




Input: root = [3,1,4,null,2], k = 1
Output: 1

Example 2:

Round 2 · High · Medium

p.921



Input: root = [5,3,6,2,4,null,null,1], k = 3
Output: 3

Constraints:

- The number of nodes in the tree is n.
- $1 \leq k \leq n \leq 104$
- $0 \leq \text{Node.val} \leq 104$

Follow up: If the BST is modified often (i.e., we can do insert and delete operations) and you need to find the kth smallest frequently, how would you optimize?

Round 2 · High · Medium

p.922

Community Solutions (Python)

• WalkCC

```
class Solution:
    def kthSmallest(self, root: TreeNode | None, k:
int) -> int:
        def countNodes(root: TreeNode | None) -> int:
            if not root:
                return 0
            return 1 + countNodes(root.left) +
countNodes(root.right)

        leftCount = countNodes(root.left)

        if leftCount == k - 1:

            return root.val
        if leftCount >= k:
            return self.kthSmallest(root.left, k)
        return self.kthSmallest(root.right, k - 1 -
leftCount) # leftCount < k
```

```
class Solution:
    def kthSmallest(self, root: TreeNode | None, k:
int) -> int:
        rank = 0
        ans = 0

        def traverse(root: TreeNode | None) -> None:
            nonlocal rank
            nonlocal ans
            if not root:
                return

            traverse(root.left)
            rank += 1
            if rank == k:
                ans = root.val
                return
            traverse(root.right)

        traverse(root)
        return ans
```

```

class Solution:
    def kthSmallest(self, root: TreeNode | None, k:
int) -> int:
        stack = []

        while root:
            stack.append(root)
            root = root.left

        for _ in range(k - 1):
            root = stack.pop()

            root = root.right
            while root:
                stack.append(root)
                root = root.left

        return stack[-1].val

```

• LeetCode

Round 2 · High · Medium

p.925

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def kthSmallest(self, root:
Optional[TreeNode], k: int) -> int:
        stk = []
        while root or stk:

            if root:
                stk.append(root)
                root = root.left
            else:
                root = stk.pop()
                k -= 1
                if k == 0:
                    return root.val
                root = root.right

```

• SimplyLeet

Round 2 · High · Medium

p.926

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def kthSmallest(self, root: TreeNode, k: int)
-> int:
        # Initialize counter and result
```

```
        self.k = k
        self.ans = None

        # In-order traversal: left -> node -> right
        def inorder(node):
            if not node:
                return
            # Traverse left subtree
            inorder(node.left)
            # Process current node
```

```
                self.k -= 1
                if self.k == 0:
                    self.ans = node.val
                    return
            # Traverse right subtree if kth element
            not found
            inorder(node.right)

        inorder(root)
        return self.ans
```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
#     self.val = val
#     self.left = left
#     self.right = right
class Solution:
    def kthSmallest(self, root:
Optional[TreeNode], k: int) -> int:
        stk = []
        while root or stk:
```

```

    if root:
        stk.append(root)
        root = root.left
    else:
        root = stk.pop()
        k -= 1
        if k == 0:
            return root.val
        root = root.right

```

```
#####
```

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution_followUp:
    def kthSmallest(self, root: TreeNode, k: int)
    -> int:

```

```

node = self.build(root)
return self.dfs(node, k)

```

```

class MyTreeNode:
    def __init__(self, x):
        self.val = x
        self.count = 1 # default 1
        self.left = None
        self.right = None

```

```

def build(self, root: TreeNode) -> MyTreeNode:
    if not root:
        return None
    node = self.MyTreeNode(root.val) # default
count already is 1
    node.left = self.build(root.left)
    node.right = self.build(root.right)
    if node.left:
        node.count += node.left.count
    if node.right:
        node.count += node.right.count

```

```

        return node

    def dfs(self, node: MyTreeNode, k: int) -> int:
        if node.left:
            cnt = node.left.count
            if k < cnt + 1:
                return self.dfs(node.left, k)
            elif k > cnt + 1:
                return self.dfs(node.right, k - 1
- cnt)
            else: # k == cnt+1
                return node.val
        else:
            if k == 1: # cannot move to beginning
of dfs()
                return node.val
            return self.dfs(node.right, k - 1)

```

◆ Quick Review

Key Insights

- An in-order traversal of a BST visits nodes in increasing order.
- Use a counter to keep track of the number of nodes visited.
- As soon as the counter reaches k, the current node's value is the kth smallest.

- For frequent modifications (insertions/deletions) and queries, consider augmenting the BST with subtree counts.

Space & Time

Time Complexity: $O(n)$ in the worst-case scenario (when $k \approx n$), but on average the traversal may stop early.

Space Complexity: $O(h)$ where h is the height of the tree ($O(n)$ in the worst-case for a skewed tree).

Solution

We perform an in-order traversal of the BST which visits nodes in ascending order. During the traversal, we maintain a counter (or decrement k) to check when we have encountered the k th smallest element. When the counter reaches 0 (or k becomes 0), we capture the current node's value as our answer and return it. This approach is both simple and efficient for the given constraints.

Trees & BST

#235 Lowest Common Ancestor of a Binary Search Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BST

Lists: Grind 169


Problem

Description

Given a binary search tree (BST), find the lowest common ancestor (LCA) node of two given nodes in the BST.

According to the definition of LCA on Wikipedia: “The lowest common ancestor is defined between two nodes p and q as the lowest node in T that has both p and q as descendants (where we allow a node to be a descendant of itself).”

Example 1:




Input: root = [6,2,8,0,4,7,9,null,null,3,5], p = 2, q = 8

Output: 6

Explanation: The LCA of nodes 2 and 8 is 6.

Example 2:



Input: root = [6,2,8,0,4,7,9,null,null,3,5], p = 2, q = 4

Output: 2

Explanation: The LCA of nodes 2 and 4 is 2, since a node can be a descendant of itself according to the LCA definition.

Example 3:

Input: root = [2,1], p = 2, q = 1

Output: 2

Constraints:

- The number of nodes in the tree is in the range [2, 105].
- $-109 \leq \text{Node.val} \leq 109$
- All Node.val are unique.
- $p \neq q$
- p and q will exist in the BST.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def lowestCommonAncestor(
        self,
        root: 'TreeNode',
        p: 'TreeNode',
        q: 'TreeNode',
    ) -> 'TreeNode':
        if root.val > max(p.val, q.val):
            return self.lowestCommonAncestor(root.left,
p, q)
        if root.val < min(p.val, q.val):
            return self.lowestCommonAncestor(root.right,
p, q)
        return root
  
```

• LeetCode


```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def lowestCommonAncestor(
        self, root: 'TreeNode', p: 'TreeNode', q:
        'TreeNode'
    ) -> 'TreeNode':
        while 1:
            if root.val < min(p.val, q.val):
                root = root.right
            elif root.val > max(p.val, q.val):
                root = root.left
            else:
                return root

```

• SimplyLeet

Round 2 · High · Medium

p.937

```

def lowestCommonAncestor(root, p, q):
    # Iterate until we find the lowest common
    ancestor
    while root:
        # If both p and q are less than the current
        node, move to left subtree
        if p.val < root.val and q.val < root.val:
            root = root.left
        # If both are greater, move to right
        subtree
        elif p.val > root.val and q.val > root.val:
            root = root.right
        else:
            # Found the split point; return current
            node as LCA
            return root

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def lowestCommonAncestor(
        self, root: 'TreeNode', p: 'TreeNode', q:
        'TreeNode'
    ) -> 'TreeNode':

```

Round 2 · High · Medium

p.938

```

    ) -> 'TreeNode':
        while 1:
            if root.val < min(p.val, q.val): # no
            =, so root is p or q is in else block
                root = root.right
            elif root.val > max(p.val, q.val):
                root = root.left
            else:
                return root

#####

```

```

class Solution:
    def lowestCommonAncestor(self, root, p, q):
        if root is None:
            return root

        if root.val > max(p.val, q.val):
            return
        self.lowestCommonAncestor(root.left, p, q)
        elif root.val < min(p.val, q.val):
            return
        self.lowestCommonAncestor(root.right, p, q)

```

```

        else:
            return root

#####

class Solution(object):
    def lowestCommonAncestor(self, root, p, q):
        """
        :type root: TreeNode
        :type p: TreeNode

        :type q: TreeNode
        :rtype: TreeNode
        """
        a, b = sorted([p.val, q.val])
        while not a <= root.val <= b:
            if a > root.val:
                root = root.right
            else:
                root = root.left
        return root

```

◆ Quick Review

Key Insights

- Utilize the BST property: for any node, values in the left subtree are less and values in the right subtree are greater.
- If both p and q have values less than the current node, the LCA must lie in the left subtree.

- If both p and q have values greater than the current node, the LCA must lie in the right subtree.
- When one node value is less than and the other is greater than (or equal to) the current node, the current node is the lowest common ancestor.
- An iterative approach is efficient and uses constant extra space.

Space & Time

Time Complexity: $O(h)$, where h is the height of the tree ($O(\log n)$ for balanced trees and $O(n)$ in the worst-case skewed tree).

Space Complexity: $O(1)$ with an iterative approach ($O(h)$ if using recursion due to call stack).

Solution

The solution leverages the BST properties. Algorithm:

– Start at the root node. – While the current node exists: If both p and q have values less than the current node's value, move to the left child. If both p and q have values greater than the current node's value, move to the right child. Otherwise, the current node is the split point where p and q diverge, so it is the lowest common ancestor. – Return the current node.

This approach efficiently finds the LCA using a single traversal from the root.

Trees & BST

#236 Lowest Common Ancestor of a Binary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Grind 169


Problem

Description

Given a binary tree, find the lowest common ancestor (LCA) of two given nodes in the tree.


According to the definition of LCA on Wikipedia: “The lowest common ancestor is defined between two nodes p and q as the lowest node in T that has both p and q as descendants (where we allow a node to be a descendant of itself).”

Example 1:



Input: root = [3,5,1,6,2,0,8,null,null,7,4], p = 5, q = 1
 Output: 3
 Explanation: The LCA of nodes 5 and 1 is 3.

Example 2:



Input: root = [3,5,1,6,2,0,8,null,null,7,4], p = 5, q = 4
 Output: 5
 Explanation: The LCA of nodes 5 and 4 is 5, since a node can be a descendant of itself according to the LCA definition.

Example 3:

Input: root = [1,2], p = 1, q = 2
 Output: 1

Constraints:

- The number of nodes in the tree is in the range [2, 105].
- $-109 \leq \text{Node.val} \leq 109$
- All Node.val are unique.
- $p \neq q$
- p and q will exist in the tree.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def lowestCommonAncestor(
        self,
        root: 'TreeNode',
        p: 'TreeNode',
        q: 'TreeNode',
    ) -> 'TreeNode':
        if not root or root == p or root == q:
            return root
        left = self.lowestCommonAncestor(root.left, p, q)
        right = self.lowestCommonAncestor(root.right, p, q)
        if left and right:
            return root
        return left or right
```

```
class Solution:
    def lowestCommonAncestor(
        self,
        root: TreeNode | None,
        p: TreeNode | None,
        q: TreeNode | None,
    ) -> TreeNode | None:
        q_ = collections.deque([root])
        parent = {root: None}
        ancestors = set() # p's ancestors
```

```
# Iterate until we find both p and q.
while p not in parent or q not in parent:
    root = q_.popleft()
    if root.left:
        parent[root.left] = root
        q_.append(root.left)
    if root.right:
        parent[root.right] = root
        q_.append(root.right)
```

```

# Insert all the p's ancestors.
while p:
    ancestors.add(p)
    p = parent[p] # `p` becomes None in the end.

# Go up from q until we meet any of p's
ancestors.
while q not in ancestors:
    q = parent[q]

return q

```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

class Solution:
    def lowestCommonAncestor(

```

```

self, root: "TreeNode", p: "TreeNode", q:
"TreeNode"
) -> "TreeNode":
    if root in (None, p, q):
        return root
    left =
self.lowestCommonAncestor(root.left, p, q)
    right =
self.lowestCommonAncestor(root.right, p, q)
    return root if left and right else (left or
right)

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

class Solution:
    def lowestCommonAncestor(self, root:
'TreeNode', p: 'TreeNode', q: 'TreeNode') ->
'TreeNode':
        # Base case: if the current node is None,
return None.

```

```

        if root is None:
            return None

        # If the current node is either p or q,
        return the current node.
        if root == p or root == q:
            return root

        # Search for p and q in the left and right
        subtrees.
        left =
        self.lowestCommonAncestor(root.left, p, q)
        right =
        self.lowestCommonAncestor(root.right, p, q)

        # If both left and right are non-null, the
        current node is the LCA.
        if left and right:
            return root

        # Otherwise, return the non-null node.
        return left if left is not None else right

```

• LC.ca

Round 2 · High · Medium

p.949

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def lowestCommonAncestor(
        self, root: 'TreeNode', p: 'TreeNode', q:
        'TreeNode'

        ) -> 'TreeNode':
        if root is None or root == p or root == q:
            return root
        left =
        self.lowestCommonAncestor(root.left, p, q)
        right =
        self.lowestCommonAncestor(root.right, p, q)
        return root if left and right else (left or
        right)

#####

# Definition for a binary tree node.

```

Round 2 · High · Medium

p.950

```
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

class Solution(object):
    def lowestCommonAncestor(self, root, p, q):
        """
        :type root: TreeNode
        :type p: TreeNode
        :type q: TreeNode
        :rtype: TreeNode
        """

        if not root:
            return root

        left = self.lowestCommonAncestor(root.left, p, q)
        right = self.lowestCommonAncestor(root.right, p, q)
```

```
        if left and right:
            return root

        if root == p or root == q:
            return root

        if left:
            return left
        if right:
            return right
        return None
```

◆ Quick Review

Key Insights

- Use a depth-first search (DFS) traversal to explore the tree.
- If the current node is either p or q, it could be the LCA.
- Recursively find p and q in the left and right subtrees.
- If both left and right recursive calls return a non-null value, then the current node is the LCA.
- If only one side returns a non-null value, propagate that value up the recursion.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree due to the recursion stack ($O(n)$ in the worst case of a skewed tree).

Solution

We solve the problem using recursion with depth-first search. The algorithm checks if the current node matches p or q . It then recursively searches the left and right subtrees. If both subtrees contain one of p or q , the current node is the LCA. Otherwise, the algorithm forwards the non-null result up the recursion. This approach naturally handles cases where one node is an ancestor of the other.

Trees & BST

#250 Count Univalue Subtrees

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98


Problem

Description

Given the root of a binary tree, return the number of uni-value subtrees.

A uni-value subtree means all nodes of the subtree have the same value.

Example 1:



Input: root = [5,1,5,5,5,null,5]
Output: 4

Example 2:

Input: root = []
Output: 0

Example 3:

Input: root = [5,5,5,5,5,null,5]
Output: 6

Constraints:

- The number of the node in the tree will be in the range [0, 1000].
- $-1000 \leq \text{Node.val} \leq 1000$

Community Solutions (Python)

Round 2 · High · Medium

p.955

• WalkCC

```

class Solution:
    def countUnivalSubtrees(self, root: TreeNode | None) -> int:
        ans = 0

        def isUnival(root: TreeNode | None, val: int) -> bool:
            nonlocal ans
            if not root:
                return True

            if isUnival(root.left, root.val) &
                isUnival(root.right, root.val):

                ans += 1
                return root.val == val

            return False

        return isUnival(root, math.inf)
    
```

• LeetCode

Round 2 · High · Medium

p.956

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
```

```
# self.val = val
# self.left = left
# self.right = right
```

```
class Solution:
```

```
    def countUnivalSubtrees(self, root:
Optional[TreeNode]) -> int:
```

```
        def dfs(root):
            if root is None:
```

```
                return True
```

```
            l, r = dfs(root.left), dfs(root.right)
```

```
            if not l or not r:
```

```
                return False
```

```
            a = root.val if root.left is None else
```

```
root.left.val
```

```
            b = root.val if root.right is None else
```

```
root.right.val
```

```
            if a == b == root.val:
```

```
                nonlocal ans
```

```
                ans += 1
```

```
                return True
```

```
        return False
```

```
    ans = 0
```

```
    dfs(root)
```

```
    return ans
```

• SimplyLeet

```
# Definition for a binary tree node.
```

```
class TreeNode:
```

```
    def __init__(self, val=0, left=None,
right=None):
```

```
        self.val = val
```

```
        self.left = left
```

```
        self.right = right
```

```
class Solution:
```

```
    def countUnivalSubtrees(self, root: TreeNode)
```

```
-> int:
```

```
        # Initialize count
```

```

        self.count = 0

        # Helper function performs postorder DFS
        def isUnival(node):
            # Base case, null nodes are considered
            unival = True
            if node is None:
                return True

            # Recursive calls for left and right
            subtrees
            leftUnival = isUnival(node.left)

            rightUnival = isUnival(node.right)

            # Check if left child exists and its
            value does not match current node's value
            if node.left and node.left.val !=
            node.val:
                leftUnival = False
            # Check if right child exists and its
            value does not match current node's value
            if node.right and node.right.val !=
            node.val:
                rightUnival = False

            # If either left or right subtree is
            not unival, then the current subtree is not
            unival

```

```

        if leftUnival and rightUnival:
            # Increment the count if the
            subtree rooted at the current node is unival
            self.count += 1
            return True
        else:
            return False

        # Start recursion from root
        isUnival(root)
        return self.count

# Example usage:
# Construct the binary tree: [5,1,5,5,5,null,5]
# root = TreeNode(5)
# root.left = TreeNode(1, TreeNode(5),
#                       TreeNode(5))
# root.right = TreeNode(5, None, TreeNode(5))
# print(Solution().countUnivalSubtrees(root)) #
Output: 4

```

- LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
```

```
class Solution:
    def countUnivalSubtrees(self, root: TreeNode)
-> int:
        count = 0
```

```
def is_unival(node):
    nonlocal count
    if not node:
        return True

    left_unival = is_unival(node.left)
    right_unival = is_unival(node.right)

    if left_unival and right_unival:
```

```
        if node.left and node.val !=
node.left.val:
            return False
        if node.right and node.val !=
node.right.val:
            return False
        count += 1
        return True

    return False

    is_unival(root)
```

```
    return count
```

```
class Solution: # return 2 values
    def countUnivalSubtrees(self, root: TreeNode)
-> int:
        def is_unival(node):
            if not node:
                return True, 0

            left_unival, left_count =
is_unival(node.left)
```

```

        right_unival, right_count =
is_unival(node.right)

        if left_unival and right_unival:
            if node.left and node.val !=
node.left.val:
                return False, 0
            if node.right and node.val !=
node.right.val:
                return False, 0
            return True, left_count +
right_count + 1

        return False, 0

_, count = is_unival(root)
return count

```

◆ Quick Review

Key Insights

- Perform a postorder traversal (left, right, root) to determine whether each subtree is univalue.
- Treat null nodes as univalue to ease the recursion.
- A node's subtree is univalue if both left and right subtrees are univalue and, if they exist, their values equal the current node's value.

- Use a global or external counter that increments every time a univalue subtree is confirmed.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack.

Solution

We use a recursive DFS approach (postorder traversal) to solve the problem. For each node, we recursively check its left and right subtrees to decide if they are univalue. By default, if a child is null, it is considered univalue. Then, we verify the current node's value against its non-null children. If both conditions hold (subtrees are univalue and matching values), we count the current node's subtree as univalue. This DFS checks in a bottom-up manner, ensuring that the properties required for a subtree to be univalue are maintained. The key trick is to return a boolean indicating if the subtree at each node is univalue and, while backtracking, increment the counter appropriately.

Trees & BST

#285 Inorder Successor in BST

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · BST

Lists: Grind 169


Problem

Description

Given the root of a binary search tree and a node p in it, return the in-order successor of that node in the BST. If the given node has no in-order successor in the tree, return null.

The successor of a node p is the node with the smallest key greater than $p.val$.

Example 1:




Input: root = [2,1,3], p = 1

Output: 2

Explanation: 1's in-order successor node is 2. Note that both p and the return value is of TreeNode type.

Example 2:



Input: root = [5,3,6,2,4,null,null,1], p = 6

Output: null

Explanation: There is no in-order successor of the current node, so the answer is null.

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- $-105 \leq \text{Node.val} \leq 105$
- All Nodes will have unique values.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def inorderSuccessor(
        self,
        root: TreeNode | None,
        p: TreeNode | None,
    ) -> TreeNode | None:
        if not root:
            return None
        if root.val <= p.val:
            return self.inorderSuccessor(root.right, p)

        return self.inorderSuccessor(root.left, p) or root
  
```

• LeetDoocs

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def inorderSuccessor(self, root: TreeNode, p:
        TreeNode) -> Optional[TreeNode]:
  
```



```

ans = None
while root:
    if root.val > p.val:
        ans = root
        root = root.left
    else:
        root = root.right
return ans

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def inorderSuccessor(self, root: TreeNode, p:
TreeNode) -> TreeNode:
        # If p has a right subtree, the successor
        is the left-most node in that subtree.

```

```

if p.right:
    successor = p.right
    while successor.left:
        successor = successor.left
    return successor

```

Otherwise, start from the root and keep track of a potential successor.

```

successor = None
current = root
while current:

```

```

        if p.val < current.val:
            # current is a candidate, so record
            it and move left to find a smaller candidate
            successor = current
            current = current.left
        else:
            # If p.val is greater or equal, go
            right.
            current = current.right
    return successor

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def inorderSuccessor(self, root: TreeNode, p:
TreeNode) -> Optional[TreeNode]:

        ans = None
        while root:
            if root.val > p.val:
                ans = root
                root = root.left
            else:
                root = root.right
        return ans

```

◆ Quick Review

Key Insights

- In a BST, an in-order traversal yields nodes in ascending order.
- If the node *p* has a right subtree, the in-order successor is the left-most node in that right subtree.
- If *p* has no right subtree, the successor is one of its ancestors; traverse from the root towards *p* while keeping

track of the last node that is greater than *p*.

- An iterative approach is efficient and avoids recursion stack overhead.

Space & Time

Time Complexity: $O(h)$, where *h* is the height of the tree. In the worst-case (skewed tree), it can be $O(n)$.

Space Complexity: $O(1)$ as only a constant amount of extra space is used.

Solution

We solve the problem using two cases:

- If node *p* has a right subtree, we simply find the left-most node in that right subtree.
- If node *p* does not have a right subtree, initiate a search from the BST root. As we traverse: If the current node's value is greater than *p.val*, the current node is a candidate for the successor; move to the left subtree to look for a smaller candidate. Otherwise, move to the right subtree. By the end of the traversal, the candidate (if exists) is the in-order successor. This approach benefits from the BST property, limiting traversal to $O(h)$ time.

Trees & BST

#298 Binary Tree Longest Consecutive Sequence

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98

Problem


Description

Given the root of a binary tree, return the length of the longest consecutive sequence path.

A consecutive sequence path is a path where the values increase by one along the path.

Note that the path can start at any node in the tree, and you cannot go from a node to its parent in the path.

Example 1:




Input: root = [1,null,3,2,4,null,null,null,5]

Output: 3

Explanation: Longest consecutive sequence path is 3-4-5, so return 3.

Example 2:



Input: root = [2,null,3,2,null,1]

Output: 2

Explanation: Longest consecutive sequence path is 2-3, not 3-2-1, so return 2.

Constraints:

- The number of nodes in the tree is in the range [1, 3 * 10⁴].
- -3 * 10⁴ ≤ Node.val ≤ 3 * 10⁴

Community Solutions (Python)

Round 2 · High · Medium

p.975

• WalkCC

```

class Solution:
    def longestConsecutive(self, root: TreeNode | None) -> int:
        if not root:
            return 0

        def dfs(root: TreeNode | None, target: int, length: int, maxLength: int) -> int:
            if not root:
                return maxLength
            if root.val == target:
                length += 1

            maxLength = max(maxLength, length)
            else:
                length = 1
            return max(dfs(root.left, root.val + 1, length, maxLength),
                      dfs(root.right, root.val + 1, length, maxLength))

        return dfs(root, root.val, 0, 0)
  
```

• LeetCode

Round 2 · High · Medium

p.976

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def longestConsecutive(self, root:
Optional[TreeNode]) -> int:
        def dfs(root: Optional[TreeNode]) -> int:
            if root is None:
```

```
                return 0
            l = dfs(root.left) + 1
            r = dfs(root.right) + 1
            if root.left and root.left.val -
root.val != 1:
                l = 1
            if root.right and root.right.val -
root.val != 1:
                r = 1
            t = max(l, r)
            nonlocal ans
            ans = max(ans, t)
```

```
        return t
```

```
    ans = 0
    dfs(root)
    return ans
```

• SimplyLeet

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def longestConsecutive(self, root: TreeNode)
-> int:
        # Initialize a variable to keep track of
the maximum sequence length.
```

```

self.max_length = 0

# Helper function to perform DFS traversal.
def dfs(node, parent_val, current_length):
    if not node:
        return
    # Check if the current node's value is
consecutive.
    if node.val == parent_val + 1:
        current_length += 1
    else:

        current_length = 1
    # Update the maximum sequence length
found so far.
    self.max_length = max(self.max_length,
current_length)
    # Recursively traverse the left and
right children.
    dfs(node.left, node.val,
current_length)
    dfs(node.right, node.val,
current_length)

    # Start DFS with an initial condition that
forces the first node to count.
    dfs(root, root.val - 1, 0)
    return self.max_length

```

• LC.ca

Round 2 · High · Medium

p.979

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right

# dfs
class Solution:
    def longestConsecutive(self, root: TreeNode)
-> int:

        maxlen = 0

        # curLen: consecutive length until parent
node
        def dfs(root: TreeNode, lastVal: int,
curLen: int) -> None:
            nonlocal maxlen
            if not root:
                return

            curLen = (curLen + 1) if (not lastVal)
and root.val == lastVal + 1 else 1
            maxlen = max(maxlen, curLen)

```

Round 2 · High · Medium

p.980

```

        dfs(root.left, root.val, curLen)
        dfs(root.right, root.val, curLen)

```

```

dfs(root, None, 0)
return maxLen

```

```
#####
```

```
# Definition for a binary tree node.
```

```

# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

```

```
from collections import deque
```

```

# iterative, bfs
class Solution:

```

```

    def longestConsecutive(self, root: TreeNode)
-> int:
        if not root:
            return 0

        max_length = 0
        # Queue elements are tuples: (current node,
        length of consecutive sequence)
        queue = deque([(root, 1)])

        while queue:
            current, length = queue.popleft()

            max_length = max(max_length, length)

            for neighbor in [current.left,
                             current.right]:
                if neighbor:
                    # Check if neighbor is
                    consecutive
                    if neighbor.val == current.val
                    + 1:
                        queue.append((neighbor,
                                      length + 1))
                    else:
                        queue.append((neighbor,
                                      1))

```

```

        return max_length

#####

class Solution:
    def longestConsecutive(self, root:
Optional[TreeNode]) -> int:
        def dfs(root: Optional[TreeNode]) -> int:
            if root is None:
                return 0
            l = dfs(root.left) + 1

            r = dfs(root.right) + 1
            if root.left and root.left.val -
root.val != 1:
                l = 1
            if root.right and root.right.val -
root.val != 1:
                r = 1
            t = max(l, r)
            nonlocal ans
            ans = max(ans, t)
            return t

        ans = 0
        dfs(root)
        return ans

```

◆ Quick Review

Round 2 · High · Medium

p.983

Key Insights

- Use a Depth-First Search (DFS) to traverse the binary tree.
- Pass the current node's value and the length of the current consecutive sequence in the recursive call.
- When a node's value is exactly one greater than its parent's value, increment the sequence length; otherwise, reset it to 1.
- Keep track of the maximum sequence length found during the traversal.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree, since each node is visited once.

Space Complexity: $O(H)$, where H is the height of the tree. In the worst-case of a skewed tree, the recursion stack can be $O(N)$.

Solution

The solution involves using DFS to explore every node in the binary tree while keeping track of the length of the consecutive sequence. For each node, if the node's value is one greater than its parent's value, we increment the sequence length; otherwise, we reset the sequence length to 1. During the traversal, a global variable is updated if the current sequence length exceeds the previously recorded maximum. This approach ensures that every possible consecutive path is considered.

Round 2 · High · Medium

p.984

Trees & BST

#314 Binary Tree Vertical Order Traversal

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Tree · BFS · Sorting

Lists: Premium 98


Problem

Description

Given the root of a binary tree, return the vertical order traversal of its nodes' values. (i.e., from top to bottom, column by column).

If two nodes are in the same row and column, the order should be from left to right.


Example 1:



Input: root = [3,9,20,null,null,15,7]


Output: [[9],[3,15],[20],[7]]

Example 2:



Input: root = [3,9,8,4,0,1,7]
Output: [[4],[9],[3,0,1],[8],[7]]

Example 3:



Input: root = [1,2,3,4,10,9,11,null,5,null,null,null,null,6]
Output: [[4],[2,5],[1,10,9,6],[3],[11]]

Constraints:

- The number of nodes in the tree is in the range [0, 100].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

- WalkCC

```

class Solution:
    def verticalOrder(self, root: TreeNode | None) ->
list[list[int]]:
    if not root:
        return []

    range_ = [0] * 2

    def getRange(root: TreeNode | None, x: int) ->
None:
    if not root:
        return

```

```

        range_[0] = min(range_[0], x)
        range_[1] = max(range_[1], x)

        getRange(root.left, x - 1)
        getRange(root.right, x + 1)

        getRange(root, 0) # Get the leftmost and the
rightmost x index.

        ans = [[] for _ in range(range_[1] - range_[0]
+ 1)]

```

```

        q = collections.deque([(root, -range_[0])]) #
(TreeNode, x)

        while q:
            node, x = q.popleft()
            ans[x].append(node.val)
            if node.left:
                q.append((node.left, x - 1))
            if node.right:
                q.append((node.right, x + 1))

        return ans

```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def verticalOrder(self, root:
Optional[TreeNode]) -> List[List[int]]:
        def dfs(root, depth, offset):
            if root is None:

```

```

        return
    d[offset].append((depth, root.val))
    dfs(root.left, depth + 1, offset - 1)
    dfs(root.right, depth + 1, offset + 1)

    d = defaultdict(list)
    dfs(root, 0, 0)
    ans = []
    for _, v in sorted(d.items()):
        v.sort(key=lambda x: x[0])

    ans.append([x[1] for x in v])
    return ans

```

• SimplyLeet

```

from collections import deque, defaultdict

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, x):
        self.val = x
        self.left = None
        self.right = None

class Solution:

```

```

def verticalOrder(self, root: TreeNode):
    # Return empty list if the tree is empty
    if not root:
        return []

    # Dictionary to hold nodes for each column
    index
    col_table = defaultdict(list)
    # Use deque for efficient FIFO queue for
    BFS; store tuple (node, column)
    queue = deque([(root, 0)])

    while queue:
        node, col = queue.popleft()
        # Append current node value to the list
        # of its column
        col_table[col].append(node.val)
        # If there is a left child, assign
        column col-1
        if node.left:
            queue.append((node.left, col - 1))
        # If there is a right child, assign
        column col+1
        if node.right:
            queue.append((node.right, col +
1))

```

```

    # Extract the columns in sorted order to
    build the final result
    sorted_cols = sorted(col_table.keys())
    result = [col_table[col] for col in
sorted_cols]
    return result

```

- LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

'''
>>> d = {1:'a', -1:'b', 10:'c', -100:'d'}
>>>

```

```

>>> d.items()
{1: 'a', 10: 'c', -100: 'd', -1: 'b'}

>>> ds = sorted(d.items())
>>> ds
[(-100, 'd'), (-1, 'b'), (1, 'a'), (10, 'c')]

>>> sorted(d.items(), key=lambda l: l[1])
[(1, 'a'), (-1, 'b'), (10, 'c'), (-100, 'd')]
>>> dict(sorted(d.items(), key=lambda l: l[1]))

```

```

{1: 'a', 10: 'c', -100: 'd', -1: 'b'}

>>> [v for i,v in sorted(d.items(), key=lambda l:
l[0])]
['d', 'b', 'a', 'c']
'''

```

```

class Solution:
    def verticalOrder(self, root:
Optional[TreeNode]) -> List[List[int]]:
        if root is None:
            return []

```

```

q = deque([(root, 0)])
d = defaultdict(list)
while q:
    for _ in range(len(q)):
        root, offset = q.popleft()
        d[offset].append(root.val)
        if root.left:
            q.append((root.left, offset -
1))
        if root.right:
            q.append((root.right, offset +
1))

```

```

return [v for _, v in sorted(d.items())]

```

```

#####

```

```

from sortedcontainers import SortedDict
from typing import List
from collections import deque

```

```

# similar to above, but use existing data structure
class Solution:

```

```

def verticalOrder(self, root: TreeNode) ->
List[List[int]]:
    if not root:
        return []

    # Use SortedDict to automatically sort by
column index
    sorted_dict = SortedDict()

    # Queue for breadth-first search; stores
pairs of (node, column_index)
    queue = deque([(0, root)])

```

```

while queue:
    node, column = queue.popleft()

    if column not in sorted_dict:
        sorted_dict[column] = [node.val]
    else:
        sorted_dict[column].append(node.val)

    # Add left and right children to the
queue
    if node.left:

```

```

        queue.append((node.left, column -
1))
        if node.right:
            queue.append((node.right, column +
1))

        # Return the values from sorted_dict as a
list of lists
        return list(sorted_dict.values())

#####

# another SortedDict, but different handling on
None node, before enqueue vs after enqueue

```

```

class Solution:
    def verticalOrder(self, root: TreeNode):
        if not root:
            return []

        # Initialize a SortedDict to maintain
columns in sorted order.
        column_table = SortedDict()
        queue = deque([(root, 0)]) # (node,
column_index)

        while queue:

```

```

        node, column = queue.popleft()

        if node is not None:
            if column not in column_table:
                column_table[column] = []
            column_table[column].append(node.v
al)

            # Add left and right children to
the queue.
            queue.append((node.left, column -
1))
            queue.append((node.right, column +
1))

```

```

        # Extract the values from the sorted
dictionary and return them.
        return list(column_table.values())

#####

```

```

from collections import defaultdict

```

```

class Solution(object):
    def verticalOrder(self, root):
        """
        :type root: TreeNode
        :rtype: List[List[int]]
        """

        def dfs(p, i, j, res): # i -> depth, j ->
            vertical-shift
            if p:
                res[j].append((p.val, i))

        self.leftMost = min(j, self.leftMost)
        dfs(p.left, i + 1, j - 1, res)
        dfs(p.right, i + 1, j + 1, res)

        self.leftMost = float("inf")
        ans = []
        res = defaultdict(list)
        dfs(root, 0, 0, res)
        i = self.leftMost
        while True:
            if not res[i]:
                break
            ans.append([item[0] for item in
sorted(res[i], key=lambda a: a[1])])
            i += 1
        return ans

```

◆ Quick Review

Key Insights

- Use Breadth-First Search (BFS) to traverse the tree level by level.
- Track the column index for each node where root is column 0, left child is column-1, and right child is column+1.
- Group nodes by their column indices using a hash table (or dictionary).
- BFS traversal naturally handles left-to-right ordering within the same level.
- Finally, sort the keys (column indices) to generate results from leftmost column to rightmost column.

Space & Time

Time Complexity: $O(n \log n)$ in the worst-case due to sorting the column indices (n nodes in the worst-case if each node is on a unique column).

Space Complexity: $O(n)$ for storing the node information and the resultant structure.

Solution

The solution uses BFS with a queue that stores nodes along with their corresponding column index. For each visited node, add its value to a dictionary keyed by its column index. When processing children, update the column index appropriately (left child: col-1, right child: col+1). After the traversal, sort the keys of the dictionary and compile the node values in order from smallest to largest column index.

This approach guarantees that nodes at the same level will be processed in left-to-right order, matching the required output.

Round 2 · High · Medium

p.1001

Trees & BST

#333 Largest BST Subtree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Dynamic Programming · Tree · DFS · BST

Lists: Premium 98

Problem

Description

Given the root of a binary tree, find the largest subtree, which is also a Binary Search Tree (BST), where the largest means subtree has the largest number of nodes.

A Binary Search Tree (BST) is a tree in which all the nodes follow the below-mentioned properties:


- The left subtree values are less than the value of their parent (root) node's value.
- The right subtree values are greater than the value of their parent (root) node's value.

Note: A subtree must include all of its descendants.

Example 1:

Round 2 · High · Medium

p.1002



Input: root = [10,5,15,1,8,null,7]

Output: 3

Explanation: The Largest BST Subtree in this case is the highlighted one. The return value is the subtree's size, which is 3.

Example 2:

Input: root =

[4,2,7,2,3,5,null,2,null,null,null,null,null,1]

Output: 2

Constraints:

- The number of nodes in the tree is in the range [0, 104].
- $-104 \leq \text{Node.val} \leq 104$

Follow up: Can you figure out ways to solve it with $O(n)$ time complexity?

Community Solutions (Python)

• WalkCC

```
from dataclasses import dataclass
```

```
@dataclass(frozen=True)
```

```
class T:
```

```
    mn: int # the minimum value in the subtree
```

```
    mx: int # the maximum value in the subtree
```

```
    size: int # the size of the subtree
```

```
class Solution:
```

```
    def largestBSTSubtree(self, root: TreeNode | None) -> int:
```

```
        def dfs(root: TreeNode | None) -> T:
```

```
            if not root:
```

```
                return T(math.inf, -math.inf, 0)
```

```
            l = dfs(root.left)
```

```
            r = dfs(root.right)
```

```
            if l.mx < root.val < r.mn:
```

```

        return T(min(l.mn, root.val), max(r.mx,
root.val), 1 + l.size + r.size)

    # Mark one as invalid, but still record the
    size of children.
    # Return (-inf, inf) because no node will be
    > inf or < -inf.
    return T(-math.inf, math.inf, max(l.size,
r.size))

return dfs(root).size

```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def largestBSTSubtree(self, root:
Optional[TreeNode]) -> int:
        def dfs(root):
            if root is None:

```

```

        return inf, -inf, 0
        lmi, lmx, ln = dfs(root.left)
        rmi, rmx, rn = dfs(root.right)
        nonlocal ans
        if lmx < root.val < rmi:
            ans = max(ans, ln + rn + 1)
            return min(lmi, root.val),
max(rmx, root.val), ln + rn + 1
        return -inf, inf, 0

ans = 0

dfs(root)
return ans

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right

class Solution:
    def largestBSTSubtree(self, root: TreeNode) ->
int:
        # Global variable to store the size of the
        largest BST subtree found.

```

```

self.max_bst_size = 0

def helper(node):
    # Returns a tuple of: (isBST, size,
    min_val, max_val)
    if not node:
        # An empty node is considered a BST
        of size 0 with extreme min and max values.
        return (True, 0, float('inf'),
        float('-inf'))

    left_is_bst, left_size, left_min,
    left_max = helper(node.left)
    right_is_bst, right_size, right_min,
    right_max = helper(node.right)

```

```

        # Check if the current node forms a BST
        with its left and right subtrees.
        if left_is_bst and right_is_bst and
        node.val > left_max and node.val < right_min:
            current_size = left_size +
            right_size + 1
            self.max_bst_size =
            max(self.max_bst_size, current_size)
            current_min = min(left_min,
            node.val)
            current_max = max(right_max,
            node.val)
            return (True, current_size,
            current_min, current_max)
        else:
            # Not a BST: Return false and the
            size of the largest BST found so far in its
            subtrees.

            return (False, max(left_size,
            right_size), 0, 0)

        helper(root)
        return self.max_bst_size

```

- LC.ca

```
'''
>>> float("inf")
inf
>>> float("inf") + 1
inf
>>> import math
>>> float("inf") == math.inf
True
'''
# Definition for a binary tree node.
```

```
# class TreeNode:
# def __init__(self, val=0, left=None,
# right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def largestBSTSubtree(self, root:
Optional[TreeNode]) -> int:
        def dfs(root):
            if root is None:
                return inf, -inf, 0
```

```
lmi, lmx, ln = dfs(root.left)
rmi, rmx, rn = dfs(root.right)
nonlocal ans
# this if can check if left/right is a
bst
# eg. if left is not bst, the returned
lmi=-inf and lmx=inf, so 'lmx < root.val < rmi' is
false
if lmx < root.val < rmi:
    ans = max(ans, ln + rn + 1)
    # lmi is guranteed < root.val,
except when left(right) is None
    return min(lmi, root.val),
max(rmx, root.val), ln + rn + 1
    return -inf, inf, 0 # if this one not
bst, then all parents will not be bst
```

```
ans = 0
dfs(root)
return ans
```

```
#####
```

```
# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
```

```

# self.val = x
# self.left = None
# self.right = None

class Solution(object):
    def largestBSTSubtree(self, root):
        """
        :type root: TreeNode
        :rtype: int
        """

    def helper(root):
        if not root:
            return (0, 0, float("inf"), float("-inf"))
        # numBST, maxNumBST, min, max
        lnumBST, lmaxNumBST, lmin, lmax =
        helper(root.left)
        rnumBST, rmaxNumBST, rmin, rmax =
        helper(root.right)
        numBST = -1
        if lmax < root.val < rmin and lnumBST != -1
        and rnumBST != -1:
            numBST = 1 + lnumBST + rnumBST
            maxNumBST = max(1, lmaxNumBST, rmaxNumBST,
            numBST)

```

```

        return numBST, maxNumBST, min(lmin, rmin,
        root.val), max(lmax, rmax, root.val)

    return helper(root)[1]

```

◆ Quick Review

Key Insights

- Use a bottom-up DFS (post-order traversal) to process each node.
- For each node, determine whether its left and right subtrees are BSTs.
- Maintain additional data per subtree: size (number of nodes), minimum value, and maximum value.
- A subtree rooted at the current node is a BST if:
 - Its left subtree is a BST.
 - Its right subtree is a BST.
 - The maximum value in the left subtree is less than the current node's value.
 - The current node's value is less than the minimum value in the right subtree.
- Track the size of the largest BST identified during the traversal.

Space & Time

Time Complexity: $O(n)$ (each node is visited once)

Space Complexity: $O(h)$, where h is the height of the tree (due to recursion stack)

Solution

Perform a post-order traversal of the binary tree. For each node, return a tuple (or object in some languages) that provides:

- Whether the subtree rooted at this node is a BST.
- The size of the subtree if it is a BST (or the size of the largest BST found within it).
- The minimum value in the subtree.
- The maximum value in the subtree.

At each node, if both left and right children form valid BSTs and the current node's value fits the BST condition (greater than left's max and less than right's min), then the subtree rooted at the node is also a BST. Calculate its size and update the current global maximum if necessary. If the condition fails, return the maximum BST size found in its subtrees.

Trees & BST

#366 Find Leaves of Binary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · Sorting

Lists: Premium 98


Problem

Description

Given the root of a binary tree, collect a tree's nodes as if you were doing this:

- Collect all the leaf nodes.
- Remove all the leaf nodes.
- Repeat until the tree is empty.

Example 1:



Input: root = [1,2,3,4,5]

Output: [[4,5,3],[2],[1]]

Explanation:

[[3,5,4],[2],[1]] and [[3,4,5],[2],[1]] are also considered correct answers since per each level it does not matter the order on which elements are returned.

Example 2:

Input: root = [1]

Output: [[1]]

Constraints:

- The number of nodes in the tree is in the range [1, 100].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def findLeaves(self, root: TreeNode | None) -> list[list[int]]:
        ans = []

        def depth(root: TreeNode | None) -> int:
            """Returns the depth of the root
            (0-indexed)."""
            if not root:
                return -1

            l = depth(root.left)

            r = depth(root.right)
            h = 1 + max(l, r)

            if len(ans) == h: # Meet a leaf
                ans.append([])

            ans[h].append(root.val)
            return h

        depth(root)

        return ans
```

- LeetCode


```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def findLeaves(self, root: Optional[TreeNode])
-> List[List[int]]:
        def dfs(root: Optional[TreeNode]) -> int:
            if root is None:
                return 0
            l, r = dfs(root.left), dfs(root.right)
            h = max(l, r)
            if len(ans) == h:
                ans.append([])
            ans[h].append(root.val)
            return h + 1

        ans = []
        dfs(root)

        return ans

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # The value of the node.
        self.left = left # Left child node.
        self.right = right # Right child node.

class Solution:
    def findLeaves(self, root: TreeNode):
        result = [] # This will store lists of
nodes at each removal step.

        def getHeight(node):
            # Base case: when the node is None,
return 0.
            if not node:
                return 0
            # Recursively find the height of left
and right subtrees.
            left_height = getHeight(node.left)
            right_height = getHeight(node.right)
            # Current node's height is max of
left/right heights plus one.
            current_height = max(left_height,
right_height) + 1

```

```

        # Ensure the result list has a sublist
        for the current height.
        if len(result) < current_height:
            result.append([])
        # Append the current node value to the
        corresponding height bucket.
        result[current_height -
1].append(node.val)
        return current_height

    getHeight(root) # Start recursion from the
    root.
    return result

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def findLeaves(self, root: TreeNode) ->
List[List[int]]:
        def dfs(root, prev, t):
            if root is None:

```

```

        return
        if root.left is None and root.right is
None:
            t.append(root.val)
            if prev.left == root:
                prev.left = None
            else:
                prev.right = None
            dfs(root.left, root, t)
            dfs(root.right, root, t)

    res = []
    prev = TreeNode(left=root)
    while prev.left:
        t = []
        dfs(prev.left, prev, t)
        res.append(t)
    return res

```

◆ Quick Review

Key Insights

- Use a postorder traversal (bottom-up approach) to process the tree.
- Determine the "height" of each node (distance from the deepest leaf) where leaves have a height of 1.
- Group nodes by their computed height to simulate the layer-wise removal of leaves.

- The same node may become a leaf after its children are removed.

Space & Time

Time Complexity: $O(n)$ — Each node is visited exactly once.

Space Complexity: $O(n)$ — Space used for the recursion stack and the result list.

Solution

The solution uses a recursive helper function that performs a postorder traversal of the tree. For each node, it computes its height as $\max(\text{height of left, height of right}) + 1$. Nodes with the same height are collected together in a list. This idea mirrors the process of removing leaves level by level, where leaves (height 1) are removed first, then their parents become new leaves, and so on. Data Structures used:

– An array or list to maintain the groups of node values by their height. Algorithmic Approach: – Use Depth-First Search (DFS) via recursion. – Postorder traversal ensures that leaf nodes (base cases) are processed before their parents. – Append the node's value to the correct list based on the computed height.

Trees & BST

#437 Path Sum III

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS · Prefix Sum

Lists: Grind 169


Problem

Description

Given the root of a binary tree and an integer `targetSum`, return the number of paths where the sum of the values along the path equals `targetSum`.

The path does not need to start or end at the root or a leaf, but it must go downwards (i.e., traveling only from parent nodes to child nodes).

Example 1:



Input: root = [10,5,-3,3,2,null,11,3,-2,null,1],
targetSum = 8
Output: 3
Explanation: The paths that sum to 8 are shown.

Example 2:

Input: root = [5,4,8,11,null,13,4,7,2,null,null,5,1],
targetSum = 22
Output: 3

Constraints:

- The number of nodes in the tree is in the range [0, 1000].
- $-109 \leq \text{Node.val} \leq 109$
- $-1000 \leq \text{targetSum} \leq 1000$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def pathSum(self, root: TreeNode | None, summ: int) -> int:
        if not root:
            return 0

        def dfs(root: TreeNode, summ: int) -> int:
            if not root:
                return 0
            return (int(summ == root.val) +
                    dfs(root.left, summ - root.val) +
                    dfs(root.right, summ - root.val))

        return (dfs(root, summ) +
                self.pathSum(root.left, summ) +
                self.pathSum(root.right, summ))
    
```

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def pathSum(self, root: Optional[TreeNode],
targetSum: int) -> int:
        def dfs(node, s):
            if node is None:
                return 0
            s += node.val
            ans = cnt[s - targetSum]
            cnt[s] += 1
            ans += dfs(node.left, s)
            ans += dfs(node.right, s)
            cnt[s] -= 1
            return ans

        cnt = Counter({0: 1})

        return dfs(root, 0)

```

• SimplyLeet

```

# Python solution for Path Sum III
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # value of the node
        self.left = left # left child
        self.right = right # right child

class Solution:
    def pathSum(self, root: TreeNode, targetSum:
int) -> int:
        # Helper function to perform DFS traversal

        def dfs(node, current_sum, prefix_sums):
            if not node:
                return 0

            # Update current sum by including the
current node's value
            current_sum += node.val
            # Check the number of times
(current_sum - targetSum) has occurred
            num_paths =
prefix_sums.get(current_sum - targetSum, 0)

            # Update the prefix sum hash map with
the current sum

```

```

        prefix_sums[current_sum] =
prefix_sums.get(current_sum, 0) + 1

        # Recursively search in the left and
right subtree
        num_paths += dfs(node.left,
current_sum, prefix_sums)
        num_paths += dfs(node.right,
current_sum, prefix_sums)

        # Backtrack: remove the current sum
count before returning to the parent node
        prefix_sums[current_sum] -= 1
        return num_paths

    # Initial prefix sum hashmap with 0 sum
seen once (empty path)
    prefix_sums = {0: 1}
    return dfs(root, 0, prefix_sums)

```

• LC.ca

Round 2 · High · Medium

p.1027

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def pathSum(self, root: Optional[TreeNode],
targetSum: int) -> int:
        def dfs(node, s):
            if node is None:
                return 0
            s += node.val
            ans = cnt[s - targetSum]
            cnt[s] += 1
            ans += dfs(node.left, s)
            ans += dfs(node.right, s)
            cnt[s] -= 1
            return ans

        cnt = Counter({0: 1})

        return dfs(root, 0)

```

◆ Quick Review

Key Insights

Round 2 · High · Medium

p.1028

- A path can start and end at any nodes, as long as it goes downwards.
- Using DFS helps traverse the tree and explore all possible paths.
- A prefix sum technique with a hash map allows quick calculation of the sum of any subpath.
- Backtracking is used to remove a path's prefix sum when moving back up the recursion tree.
- Although a brute-force approach exists, it is inefficient compared to the prefix sum method.

Space & Time

Time Complexity: $O(n)$ on average (where n is the number of nodes), though worst-case can be $O(n^2)$ in a skewed tree.

Space Complexity: $O(n)$ due to recursion stack and the hash map that stores prefix sums.

Solution

The solution employs a DFS traversal of the tree along with a hash map that records the frequency of prefix sums encountered so far. At each node, the current prefix sum is updated by adding the node's value. We then check the hash map for how many times (current prefix sum - targetSum) has occurred, since this value indicates the existence of a subpath summing to targetSum. We update the hash map with the current prefix sum, continue the DFS to traverse child nodes, and then backtrack by decrementing the current prefix sum's count before retreating to the parent node. This ensures that sums from

other branches are not incorrectly combined.

Trees & BST

#545 Boundary of Binary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98

Problem

Description

The boundary of a binary tree is the concatenation of the root, the left boundary, the leaves ordered from left-to-right, and the reverse order of the right boundary.

The left boundary is the set of nodes defined by the following:

- The root node's left child is in the left boundary. If the root does not have a left child, then the left boundary is empty.
- If a node is in the left boundary and has a left child, then the left child is in the left boundary.
- If a node is in the left boundary, has no left child, but has a right child, then the right child is in the left boundary.
- The leftmost leaf is not in the left boundary.


The right boundary is similar to the left boundary, except it is the right side of the root's right subtree.

Again, the leaf is not part of the right boundary, and the right boundary is empty if the root does not have a right child.

The leaves are nodes that do not have any children. For this problem, the root is not a leaf.

Given the root of a binary tree, return the values of its boundary.

Example 1:




Input: root = [1,null,2,3,4]

Output: [1,3,4,2]

Explanation:

- The left boundary is empty because the root does not have a left child.
 - The right boundary follows the path starting from the root's right child 2 -> 4. 4 is a leaf, so the right boundary is [2].
 - The leaves from left to right are [3,4].
- Concatenating everything results in [1] + [] + [3,4] + [2] = [1,3,4,2].

Example 2:



Round 2 · High · Medium

p.1033

Input: root = [1,2,3,4,5,6,null,null,null,7,8,9,10]

Output: [1,2,4,7,8,9,10,6,3]

Explanation:

- The left boundary follows the path starting from the root's left child 2 -> 4. 4 is a leaf, so the left boundary is [2].
- The right boundary follows the path starting from the root's right child 3 -> 6 -> 10. 10 is a leaf, so the right boundary is [3,6], and in reverse order is [6,3].
- The leaves from left to right are [4,7,8,9,10].

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- -1000 <= Node.val <= 1000

Community Solutions (Python)

- WalkCC

Round 2 · High · Medium

p.1034

```
class Solution:
    def boundaryOfBinaryTree(self, root: TreeNode |
None) -> list[int]:
        if not root:
            return []

        ans = [root.val]

        def dfs(root: TreeNode | None, lb: bool, rb:
bool):
            """
            1. root.left is left boundary if root is left
            boundary.
```

```
            root.right if left boundary if root.left
            is None.
            2. Same applies for right boundary.
            3. If root is left boundary, add it before 2
            children - preorder.
            If root is right boundary, add it after 2
            children - postorder.
            4. A leaf that is neither left/right
            boundary belongs to the bottom.
            """
            if not root:
                return
            if lb:
                ans.append(root.val)
```

```
            if not lb and not rb and not root.left and
            not root.right:
                ans.append(root.val)

            dfs(root.left, lb, rb and not root.right)
            dfs(root.right, lb and not root.left, rb)
            if rb:
                ans.append(root.val)

            dfs(root.left, True, False)
            dfs(root.right, False, True)

        return ans
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def boundaryOfBinaryTree(self, root:
Optional[TreeNode]) -> List[int]:
        def dfs(nums: List[int], root:
Optional[TreeNode], i: int):
            if root is None:
```

```

        return
    if i == 0:
        if root.left != root.right:
            nums.append(root.val)
            if root.left:
                dfs(nums, root.left, i)
            else:
                dfs(nums, root.right, i)
    elif i == 1:
        if root.left == root.right:

```

```

            nums.append(root.val)
        else:
            dfs(nums, root.left, i)
            dfs(nums, root.right, i)
    else:
        if root.left != root.right:
            nums.append(root.val)
            if root.right:
                dfs(nums, root.right, i)
            else:

```

```

        dfs(nums, root.left, i)

```

```

        ans = [root.val]
        if root.left == root.right:
            return ans
        left, leaves, right = [], [], []
        dfs(left, root.left, 0)
        dfs(leaves, root, 1)
        dfs(right, root.right, 2)
        ans += left + leaves + right[::-1]

```

```

        return ans

```

• SimplyLeet

Definition for a binary tree node.

```

class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

```

```

class Solution:

```

```

    def boundaryOfBinaryTree(self, root: TreeNode)
-> list:
        if not root:

```

```

        return []

    # Check if a node is a leaf.
    def isLeaf(node):
        return node and not node.left and not
node.right

    # Add left boundary nodes (excluding
leaves).
    def addLeftBoundary(node, boundary):
        while node:
            if not isLeaf(node):
                boundary.append(node.val)
            # Prefer left child, if not present
then right child.
            if node.left:
                node = node.left
            else:
                node = node.right

    # Add leaf nodes using DFS.
    def addLeaves(node, leaves):
        if not node:

```

```

        return
    if isLeaf(node):
        leaves.append(node.val)
    else:
        addLeaves(node.left, leaves)
        addLeaves(node.right, leaves)

    # Add right boundary nodes in reverse order
(excluding leaves).
    def addRightBoundary(node, boundary):
        temp = []

        while node:
            if not isLeaf(node):
                temp.append(node.val)
            if node.right:
                node = node.right
            else:
                node = node.left
        # Append in reverse order.
        while temp:
            boundary.append(temp.pop())

```

```

boundary = []
# Add root value if it is not a leaf.
if not isLeaf(root):
    boundary.append(root.val)

addLeftBoundary(root.left, boundary)
addLeaves(root, boundary)
addRightBoundary(root.right, boundary)

```

```

return boundary

```

```

# Example usage:
# Construct a binary tree and call
boundaryOfBinaryTree(root)

```

- LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def boundaryOfBinaryTree(self, root: TreeNode)
-> List[int]:
        self.res = []
        if not root:

```

```

        return self.res
# root
if not self.is_leaf(root):
    self.res.append(root.val)

# left boundary
t = root.left
while t:
    if not self.is_leaf(t):
        self.res.append(t.val)

```

```

t = t.left if t.left else t.right

```

```

# leaves
self.add_leaves(root)

# right boundary(reverse order)
s = []
t = root.right
while t:
    if not self.is_leaf(t):

```

```

        s.append(t.val)
        t = t.right if t.right else t.left
    while s:
        self.res.append(s.pop())

    # output
    return self.res

def add_leaves(self, root):
    if self.is_leaf(root):

```

```

        self.res.append(root.val)
        return
    if root.left:
        self.add_leaves(root.left)
    if root.right:
        self.add_leaves(root.right)

def is_leaf(self, node) -> bool:
    return node and node.left is None and
    node.right is None

```

◆ Quick Review

Key Insights

- Break down the problem into three parts:
- Left boundary (excluding leaves).
- All leaves (do a DFS to find leaves).

- Right boundary (excluding leaves, then reverse the collected list).
- Special handling is needed when the tree doesn't have a left or right subtree.
- The root is always included and is never considered as a leaf, even if it has no children.
- Avoid duplication of leaf nodes in the left/right boundaries.

Space & Time

Time Complexity: $O(n)$ – Every node is visited at most once.
 Space Complexity: $O(n)$ – $O(n)$ space is used for the output list and recursion stack (in worst case tree is skewed).

Solution

The solution uses a Depth-First Search (DFS) strategy to traverse the tree and collect the boundary nodes in three steps:

– For the left boundary, start at `root.left` and traverse down, adding nodes (skipping leaves). – For the leaves, perform a DFS from the root and add any node that has no children. – For the right boundary, start at `root.right` and traverse down, adding nodes (skipping leaves). Since the right boundary must be added in reverse order, we reverse the collected list after traversal. – Finally, concatenate these lists in the order: root, left boundary, leaves, reversed right boundary.

The approach makes use of helper functions to clearly separate the logic for gathering left boundary, leaves, and right boundary. This modular approach simplifies

reasoning and avoids edge case mistakes such as including duplicate leaf nodes.

Round 2 · High · Medium

p.1045

Trees & BST

#549 Binary Tree Longest Consecutive Sequence II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98

Problem

Description

Given the root of a binary tree, return the length of the longest consecutive path in the tree.

A consecutive path is a path where the values of the consecutive nodes in the path differ by one. This path can be either increasing or decreasing.


- For example, [1,2,3,4] and [4,3,2,1] are both considered valid, but the path [1,2,4,3] is not valid.

On the other hand, the path can be in the child-Parent-child order, where not necessarily be parent-child order.

Example 1:

Round 2 · High · Medium

p.1046




Input: root = [1,2,3]
 Output: 2
 Explanation: The longest consecutive path is [1, 2] or [2, 1].

Example 2:

Round 2 · High · Medium

p.1047



Input: root = [2,1,3]
 Output: 3
 Explanation: The longest consecutive path is [1, 2, 3] or [3, 2, 1].

Constraints:

- The number of nodes in the tree is in the range [1, 3 * 10⁴].
- -3 * 10⁴ ≤ Node.val ≤ 3 * 10⁴

Community Solutions (Python)

Round 2 · High · Medium

p.1048

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def longestConsecutive(self, root: TreeNode)
    -> int:
        def dfs(root):
            if root is None:
```

```
                return [0, 0]
            nonlocal ans
            incr = decr = 1
            i1, d1 = dfs(root.left)
            i2, d2 = dfs(root.right)
            if root.left:
                if root.left.val + 1 == root.val:
                    incr = i1 + 1
                if root.left.val - 1 == root.val:
                    decr = d1 + 1
```

```
        if root.right:
            if root.right.val + 1 == root.val:
                incr = max(incr, i2 + 1)
            if root.right.val - 1 == root.val:
                decr = max(decr, d2 + 1)
        ans = max(ans, incr + decr - 1)
        return [incr, decr]
```

```
ans = 0
dfs(root)
```

```
return ans
```

• SimplyLeet

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None, right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def longestConsecutive(self, root: TreeNode)
    -> int:
        # Initialize global maximum path length
        variable
```

```

        self.max_length = 0

    def dfs(node):
        if not node:
            return (0, 0) # (increasing,
decreasing)

        inc = dec = 1 # Each node counts as
length 1

        # Process left child
        if node.left:

            left_inc, left_dec =
dfs(node.left)

            # If left child's value is
consecutive increasing with respect to node
            if node.left.val == node.val + 1:
                inc = max(inc, left_inc + 1)
            # If left child's value is
consecutive decreasing with respect to node
            elif node.left.val == node.val -
1:

                dec = max(dec, left_dec + 1)

        # Process right child
        if node.right:

```

```

                right_inc, right_dec =
dfs(node.right)

                # If right child's value is
consecutive increasing with respect to node
                if node.right.val == node.val + 1:
                    inc = max(inc, right_inc + 1)
                # If right child's value is
consecutive decreasing with respect to node
                elif node.right.val == node.val -
1:

                    dec = max(dec, right_dec + 1)

                # Update the global maximum path by
combining increasing and decreasing paths
                self.max_length = max(self.max_length,
inc + dec - 1)

            return (inc, dec)

        dfs(root)
        return self.max_length

```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

...
>>> float('-inf')+1
-inf
```

```
>>> float('-inf')-1
-inf
...
class Solution(object):
    def longestConsecutive(self, root):
        """
        :type root: TreeNode
        :rtype: int
        """
```

```
def dfs(root):
    if not root:
        # val, increasing length,
        decreasing length, max length
        return float('-inf'), 0, 0, 0
    # inc/dec starting from root
    inc = dec = 1
    left, leftInc, leftDec, leftMax =
dfs(root.left)
    right, rightInc, rightDec, rightMax =
dfs(root.right)
    if root.val + 1 == left:
        inc = max(leftInc + 1, inc)

    if root.val - 1 == left:
        dec = max(leftDec + 1, dec)
    if root.val + 1 == right:
        inc = max(rightInc + 1, inc)
    if root.val - 1 == right:
        dec = max(rightDec + 1, dec)
    # note: max may occur in children
    subtree, not current root's tree
    return root.val, inc, dec, max(inc +
dec - 1, leftMax, rightMax, inc, dec)
    # -1 because root is counted twice,
    even left/right both null
```

```
return dfs(root)[3]
```

```
#####
```

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
# right=None):
# self.val = val
# self.left = left
# self.right = right
```

```
class Solution:
    def longestConsecutive(self, root: TreeNode)
    -> int:
        def dfs(root):
            if root is None:
                return [0, 0]
            nonlocal ans
            incr = decr = 1
            i1, d1 = dfs(root.left)
            i2, d2 = dfs(root.right)
            if root.left:
```

```
if root.left.val + 1 == root.val:
    incr = i1 + 1
if root.left.val - 1 == root.val:
    decr = d1 + 1
```

```
if root.right:
    if root.right.val + 1 == root.val:
        incr = max(incr, i2 + 1) #
```

incr1 and incr2 must be in different tree branches,
eg. left child cannot be both increasing and
decreasing. Same as decr1 and decr2

```
if root.right.val - 1 == root.val:
    decr = max(decr, d2 + 1)
```

ans = max(ans, incr + decr - 1) # -1
because root is counted twice, even left/right both
null

```
return [incr, decr]
```

```
ans = 0
dfs(root)
return ans
```

```
#####
```

```
class Solution: # issue with this solution,
    duplicated dfs() search
    def longestConsecutive(self, root: TreeNode)
    -> int:
```

```

def helper(node: TreeNode, diff: int) ->
int:
    if not node:
        return 0

    left = right = 0
    if node.left and node.val -
node.left.val == diff:
        left = 1 + helper(node.left, diff)
    if node.right and node.val -
node.right.val == diff:
        right = 1 + helper(node.right,
diff)

    return max(left, right)
if not root:
    return 0

res = helper(root, 1) + helper(root, -1) +
1 # +1 is increasing, -1 is decreasing
return max(res,
max(longestConsecutive(root.left),
longestConsecutive(root.right)))

```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to visit every node.
- For each node, compute:

- The longest increasing consecutive path starting from that node.
- The longest decreasing consecutive path starting from that node.
- Combine the increasing and decreasing paths (subtracting one to avoid double counting the current node) to potentially form a longer consecutive sequence through the node.
- Update a global maximum to record the longest path seen.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes since every node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack.

Solution

The solution employs a DFS traversal. At every node, we calculate two values:

– inc – the length of the longest increasing consecutive sequence starting from that node. – dec – the length of the longest decreasing consecutive sequence starting from that node.

For each child of the current node:

– If the child's value is $\text{node.val} + 1$, update the increasing sequence. – If the child's value is $\text{node.val} - 1$, update the decreasing sequence.

Then, the longest path through the current node is $\text{inc} + \text{dec} - 1$ (subtracting one to avoid counting the current node

twice). A global variable maintains the maximum length encountered. This process works regardless of the path direction because it combines both increasing and decreasing parts.

Round 2 · High · Medium

p.1059

Trees & BST

#582 Kill Process

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Tree · DFS · BFS

Lists: Premium 98

Problem

Description

You have n processes forming a rooted tree structure. You are given two integer arrays `pid` and `ppid`, where `pid[i]` is the ID of the i th process and `ppid[i]` is the ID of the i th process's parent process.

Each process has only one parent process but may have multiple children processes. Only one process has `ppid[i] = 0`, which means this process has no parent process (the root of the tree).


When a process is killed, all of its children processes will also be killed.

Given an integer `kill` representing the ID of a process you want to kill, return a list of the IDs of the processes that will be killed. You may return the answer in any order.

Example 1:

Round 2 · High · Medium

p.1060



Input: pid = [1,3,10,5], ppid = [3,0,5,3], kill = 5
 Output: [5,10]
 Explanation: The processes colored in red are the processes that should be killed.

Example 2:

Input: pid = [1], ppid = [0], kill = 1
 Output: [1]

Constraints:

- $n == \text{pid.length}$
- $n == \text{ppid.length}$

- $1 \leq n \leq 5 * 10^4$
- $1 \leq \text{pid}[i] \leq 5 * 10^4$
- $0 \leq \text{ppid}[i] \leq 5 * 10^4$
- Only one process has no parent.
- All the values of pid are unique.
- kill is guaranteed to be in pid.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def killProcess(
        self,
        pid: list[int],
        ppid: list[int],
        kill: int,
    ) -> list[int]:
        ans = []
        tree = collections.defaultdict(list)

        for v, u in zip(pid, ppid):
            if u == 0:
                continue
            tree[u].append(v)

        def dfs(u: int) -> None:
            ans.append(u)
            for v in tree.get(u, []):
                dfs(v)
  
```

```
dfs(kill)
return ans
```

- LeetCode

```
class Solution:
    def killProcess(self, pid: List[int], ppid:
List[int], kill: int) -> List[int]:
        def dfs(i: int):
            ans.append(i)
            for j in g[i]:
                dfs(j)

        g = defaultdict(list)
        for i, p in zip(pid, ppid):
            g[p].append(i)
```

```
ans = []
dfs(kill)
return ans
```

- SimplyLeet

Python solution using BFS

```
from collections import deque
```

```
def killProcess(pid, ppid, kill):
    # Build a mapping from parent process id to
    list of child process ids
    tree = {}
    for child, parent in zip(pid, ppid):
        if parent not in tree:
            tree[parent] = []
```

```
tree[parent].append(child)
```

```
# Use a deque for BFS traversal
result = []
queue = deque([kill])
```

```
while queue:
    current = queue.popleft() # Get process
from the left of the queue
    result.append(current) # Add the current
process to the result
    if current in tree:
```



```

        # Enqueue all children of the current
process
        for child in tree[current]:
            queue.append(child)

    return result

```

```

# Example usage:
# pid = [1, 3, 10, 5]
# ppid = [3, 0, 5, 3]
# kill = 5

```

```

# Expected output: [5, 10]

```

• LC.ca

```

class Solution:
    def killProcess(self, pid: List[int], ppid:
List[int], kill: int) -> List[int]:
        def dfs(i: int):
            ans.append(i)
            for j in g[i]:
                dfs(j)

        g = defaultdict(list)
        for i, p in zip(pid, ppid):
            g[p].append(i)

```

```

ans = []
dfs(kill)
return ans

```

Round 2 · High · Medium

p.1065

◆ Quick Review

Key Insights

- Build a mapping (dictionary/hash table) from a parent process ID to its list of child process IDs.
- Use Depth-First Search (DFS) or Breadth-First Search (BFS) starting from the kill process to traverse the tree and collect all descendant processes.
- The problem requires handling a tree structure that may have multiple children per node.

Space & Time

Time Complexity: $O(n)$, where n is the number of processes (both building the mapping and the tree traversal take linear time).

Space Complexity: $O(n)$, for storing the mapping and the traversal queue/stack.

Solution

We first construct a mapping from each process's parent ID to its list of child IDs. This allows us to quickly access all children of any process. Once the mapping is built, we perform a tree traversal (either BFS or DFS) starting from the process to kill. During the traversal, we record every process encountered. This collected list represents all processes that will be terminated. The data structures used are a dictionary (for the mapping) and either a stack (for DFS) or a queue (for BFS), making the solution efficient and straightforward.

Round 2 · High · Medium

p.1066

Trees & BST

#662 Maximum Width of Binary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · BFS

Lists: Grind 169

Problem

Description


Given the root of a binary tree, return the maximum width of the given tree.

The maximum width of a tree is the maximum width among all levels.

The width of one level is defined as the length between the end-nodes (the leftmost and rightmost non-null nodes), where the null nodes between the end-nodes that would be present in a complete binary tree extending down to that level are also counted into the length calculation.

It is guaranteed that the answer will in the range of a 32-bit signed integer.

Example 1:




Input: root = [1,3,2,5,3,null,9]

Output: 4

Explanation: The maximum width exists in the third level with length 4 (5,3,null,9).

Example 2:



Input: root = [1,3,2,5,null,null,9,6,null,7]


Output: 7

Explanation: The maximum width exists in the fourth level with length 7 (6,null,null,null,null,null,7).

Example 3:

Round 2 · High · Medium

p.1069



Input: root = [1,3,2,5]

Output: 2

Explanation: The maximum width exists in the second level with length 2 (3,2).

Constraints:

- The number of nodes in the tree is in the range [1, 3000].
- $-100 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

Round 2 · High · Medium

p.1070

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def widthOfBinaryTree(self, root: Optional[TreeNode]) -> int:
        ans = 0
        q = deque([(root, 1)])

        while q:
            ans = max(ans, q[-1][1] - q[0][1] + 1)
            for _ in range(len(q)):
                root, i = q.popleft()
                if root.left:
                    q.append((root.left, i << 1))
                if root.right:
                    q.append((root.right, i << 1 | 1))
            return ans
```

• SimplyLeet

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
```

```
from collections import deque
```

```
class Solution:
```

```
    def widthOfBinaryTree(self, root: TreeNode) -> int:
        if not root:
            return 0

        max_width = 0
        # Initialize queue with a tuple of (node, index)
        queue = deque([(root, 1)])

        while queue:
            level_length = len(queue)
```

```

        # Get the index of first node at
        current level
        _, first_index = queue[0]
        # Process nodes at current level
        for i in range(level_length):
            node, index = queue.popleft()
            # Check and add left child with
            index 2 * index
            if node.left:
                queue.append((node.left, 2 *
            index))
            # Check and add right child with
            index 2 * index + 1
            if node.right:
                queue.append((node.right, 2 *
            index + 1))
        # For the last node, calculate
        width of the level
        if i == level_length - 1:
            current_width = index -
            first_index + 1
            max_width = max(max_width,
            current_width)
        return max_width

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
#         right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def widthOfBinaryTree(self, root:
        Optional[TreeNode]) -> int:
        ans = 0
        q = deque([(root, 1)])

        while q:
            ans = max(ans, q[-1][1] - q[0][1] + 1)
            for _ in range(len(q)):
                root, i = q.popleft()
                if root.left:
                    q.append((root.left, i << 1))
                if root.right:
                    q.append((root.right, i << 1 |
            1))
        return ans

```

◆ Quick Review

Key Insights

- Perform a level order traversal (BFS) of the tree.
- Assign an index to each node as if the tree were complete (e.g., root is index 1, left child is 2

- i, right child is 2
- i+1).
- For each level, compute the width as (last_index - first_index + 1).
- Using indices allows capturing gaps between nodes due to nulls in between.
- Edge cases include handling very skewed trees.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree.

Space Complexity: $O(N)$, due to the queue used for BFS which in the worst case holds nodes of the last level.

Solution

We use a Breadth-First Search (BFS) approach with a queue to traverse the tree level by level. At each level, we pair each node with an index representing its position assuming a complete binary tree. This allows us to compute the width of each level using the formula: $\text{width} = \text{last_index} - \text{first_index} + 1$. We then track the maximum width seen throughout the traversal. The use of indices handles cases where there are null gaps between nodes.

Trees & BST

#863 All Nodes Distance K in Binary Tree MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Tree · DFS · BFS

Lists: Grind 169


Problem

Description

Given the root of a binary tree, the value of a target node target, and an integer k, return an array of the values of all nodes that have a distance k from the target node.

You can return the answer in any order.

Example 1:



Input: root = [3,5,1,6,2,0,8,null,null,7,4], target = 5, k = 2
 Output: [7,4,1]
 Explanation: The nodes that are a distance 2 from the target node (with value 5) have values 7, 4, and 1.

Example 2:

Input: root = [1], target = 1, k = 3
 Output: []

Constraints:

- The number of nodes in the tree is in the range [1, 500].
- $0 \leq \text{Node.val} \leq 500$
- All the values Node.val are unique.
- target is the value of one of the nodes in the tree.
- $0 \leq k \leq 1000$

Community Solutions (Python)

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, x):
#         self.val = x
#         self.left = None
#         self.right = None

class Solution:
    def distanceK(self, root: TreeNode, target:
    TreeNode, k: int) -> List[int]:
  
```

```
def dfs(root, fa):
    if root is None:
        return
    g[root] = fa
    dfs(root.left, root)
    dfs(root.right, root)
```

```
def dfs2(root, fa, k):
    if root is None:
        return
```

```
        if k == 0:
            ans.append(root.val)
            return
        for nxt in (root.left, root.right,
g[root]):
            if nxt != fa:
                dfs2(nxt, root, k - 1)

g = {}
dfs(root, None)
ans = []
```

```
dfs2(target, None, k)
return ans
```

• SimplyLeet

Round 2 · High · Medium

p.1079

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right
```

```
def distanceK(root: TreeNode, target: int, k: int):
    from collections import defaultdict, deque
```

```
# Build an adjacency list to represent the
graph
graph = defaultdict(list)
```

```
# Helper function to build the graph using DFS
def buildGraph(node, parent):
    if node is None:
        return
    if parent is not None:
        # Add bidirectional connection between
node and parent
        graph[node.val].append(parent.val)
```

Round 2 · High · Medium

p.1080


```
graph[parent.val].append(node.val)
buildGraph(node.left, node)
buildGraph(node.right, node)
```

```
buildGraph(root, None)
```

```
# BFS from target node
visited = set([target])
queue = deque([target])
current_distance = 0
```

```
while queue:
    if current_distance == k:
        # Return all nodes at distance k
        return list(queue)
    size = len(queue)
    for _ in range(size):
        current = queue.popleft()
        for neighbor in graph[current]:
            if neighbor not in visited:
```

```
visited.add(neighbor)
queue.append(neighbor)
current_distance += 1
```

```
return [] # If no nodes at distance k
```

```
# Example usage:
# Construct tree or use a tree builder before
calling distanceK function.
```

• LC.ca

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None
```

```
class Solution:
    def distanceK(self, root: TreeNode, target:
TreeNode, k: int) -> List[int]:
```

```
def parents(root, prev):
    nonlocal p
    if root is None:
        return
    p[root] = prev
    parents(root.left, root)
    parents(root.right, root)
```

```
def dfs(root, k):
    nonlocal ans, vis
```

```
    if root is None or root.val in vis:
        return
    vis.add(root.val)
    if k == 0:
        ans.append(root.val)
        return
    dfs(root.left, k - 1)
    dfs(root.right, k - 1)
    dfs(p[root], k - 1)
```

```
p = {}
parents(root, None)
ans = []
vis = set()
dfs(target, k)
return ans
```

◆ Quick Review

Round 2 · High · Medium

p.1083

Key Insights

- Convert the binary tree into an undirected graph so that we can easily traverse both down to the children and up to the parent.
- Use a DFS or simple recursive traversal to create an adjacency list representing the graph.
- Perform a BFS starting from the target node to explore nodes level-by-level until reaching distance k.
- Use a visited set (or equivalent) to avoid revisiting nodes and to prevent cycles.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree. Each node is visited once when constructing the graph and once in the BFS.

Space Complexity: $O(N)$, for storing the graph and the additional data structures used in BFS and DFS.

Solution

The solution involves two main steps:

- Build an undirected graph representation of the binary tree using a DFS. For each node, add the node's value as a key in an adjacency list and connect it with its children (if any). This creates bidirectional links between parent and child.
- Use a BFS starting from the target node and explore nodes level-by-level. Keep a counter for the current distance, and once the counter reaches k , return all node values found at that level. To avoid processing nodes more than once, maintain a set of visited nodes.

Round 2 · High · Medium

p.1084

Trees & BST

#1120 Maximum Average Subtree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98

Problem


Description

Given the root of a binary tree, return the maximum average value of a subtree of that tree. Answers within 10⁻⁵ of the actual answer will be accepted.

A subtree of a tree is any node of that tree plus all its descendants.

The average value of a tree is the sum of its values, divided by the number of nodes.

Example 1:



Input: root = [5,6,1]

Output: 6.00000

Explanation:

For the node with value = 5 we have an average of $(5 + 6 + 1) / 3 = 4$.

For the node with value = 6 we have an average of $6 / 1 = 6$.

For the node with value = 1 we have an average of $1 / 1 = 1$.

So the answer is 6 which is the maximum.

Example 2:

Input: root = [0,null,1]

Output: 1.00000

Constraints:

- The number of nodes in the tree is in the range [1, 104].
- $0 \leq \text{Node.val} \leq 105$

Community Solutions (Python)

• WalkCC

```
from dataclasses import dataclass

@dataclass(frozen=True)
class T:
    summ: int
    count: int
    maxAverage: int

class Solution:
    def maximumAverageSubtree(self, root: TreeNode | None) -> float:
        def maximumAverage(root: TreeNode | None) -> T:
            if not root:
                return T(0, 0, 0)

            left = maximumAverage(root.left)
            right = maximumAverage(root.right)

            summ = root.val + left.summ + right.summ
```

```
count = 1 + left.count + right.count
maxAverage = max(summ / count,
left.maxAverage, right.maxAverage)
return T(summ, count, maxAverage)

return maximumAverage(root).maxAverage
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def maximumAverageSubtree(self, root:
Optional[TreeNode]) -> float:
        def dfs(root):
            if root is None:
```

```

        return 0, 0
    ls, ln = dfs(root.left)
    rs, rn = dfs(root.right)
    s = root.val + ls + rs
    n = 1 + ln + rn
    nonlocal ans
    ans = max(ans, s / n)
    return s, n

```

```
ans = 0
```

```
dfs(root)
return ans
```

• SimplyLeet

Definition for a binary tree node.

```
class TreeNode:
```

```
    def __init__(self, val=0, left=None,
right=None):
```

```
        self.val = val # value of the node
        self.left = left # left child
        self.right = right # right child

```

```
class Solution:
```

```
    def maximumAverageSubtree(self, root):
        # Initialize the maximum average to
        negative infinity

```

```
self.maxAvg = float('-inf')
```

```
# Recursive DFS function that returns a
tuple (subtree sum, node count)
```

```
def dfs(node):
```

```
    if not node:
```

```
        return (0, 0) # return sum 0 and
count 0 if node is null

```

```
        leftSum, leftCount = dfs(node.left) #
traverse left subtree

```

```
        rightSum, rightCount = dfs(node.right)
# traverse right subtree

```

```
        currentSum = node.val + leftSum +
rightSum # total sum at current node

```

```
        currentCount = 1 + leftCount +
rightCount # total node count

```

```
        currentAvg = currentSum / currentCount
# average value for current subtree

```

```
        self.maxAvg = max(self.maxAvg,
currentAvg) # update maximum average if necessary
        return (currentSum, currentCount)

```

```
dfs(root)
return self.maxAvg
```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def maximumAverageSubtree(self, root:
Optional[TreeNode]) -> float:
        def dfs(root):
            if root is None:

                return 0, 0
            ls, ln = dfs(root.left)
            rs, rn = dfs(root.right)
            s = root.val + ls + rs
            n = 1 + ln + rn
            nonlocal ans
            ans = max(ans, s / n)
            return s, n

        ans = 0

        dfs(root)
        return ans

```

◆ Quick Review

Key Insights

- Use a depth-first search (DFS) to traverse every subtree.
- At each node, compute the sum and count of nodes in its subtree.
- Calculate the average for the current subtree and update a global maximum if needed.
- Processing every node exactly once yields an efficient solution.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes (each node is visited once).

Space Complexity: $O(h)$, where h is the height of the tree (space for recursion stack).

Solution

We perform a post-order DFS traversal, calculating at each node both the cumulative sum and the count of nodes in its subtree. This enables us to compute the average value in constant time for any subtree. We maintain a global variable to track the maximum average encountered during the traversal. The primary technique is recursive DFS combined with tuple (or pair) aggregation. The simplicity of the recursion and in-place aggregation ensures an efficient and clear solution.

Trees & BST

#1490 Clone N-ary Tree

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Hash Table · Tree · DFS · BFS

Lists: Premium 98

Problem

Description


Given a root of an N-ary tree, return a deep copy (clone) of the tree.

Each node in the n-ary tree contains a val (int) and a list (List[Node]) of its children.

```
class Node {
    public int val;
    public List<Node> children;
}
```

N-ary-Tree input serialization is represented in their level order traversal, each group of children is separated by the null value (See examples).


Example 1:



Input: root = [1,null,3,2,4,null,5,6]

Output: [1,null,3,2,4,null,5,6]

Example 2:



Input: root = [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14]
 Output: [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14]

Constraints:

- The depth of the n-ary tree is less than or equal to 1000.
- The total number of nodes is between [0, 104].

Follow up: Can your solution work for the graph problem?

Community Solutions (Python)

• LeetCode

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
        None else []
"""

class Solution:

    def cloneTree(self, root: 'Node') -> 'Node':
        if root is None:
            return None
        children = [self.cloneTree(child) for
        child in root.children]
        return Node(root.val, children)
  
```

• SimplyLeet


```

# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
None else []

def cloneTree(root):
    # Base case: if root is None, return None.
    if root is None:
        return None

    # Create a clone for the current node.
    cloned_node = Node(root.val)

    # Recursively clone each child and add it to
the cloned_node's children list.
    for child in root.children:
        cloned_child = cloneTree(child)
        cloned_node.children.append(cloned_child)

    # Return the cloned node.

    return cloned_node

# Example usage:
# original_tree = Node(1, [Node(3, [Node(5),
Node(6)]), Node(2), Node(4)])
# cloned_tree = cloneTree(original_tree)

```

• LC.ca

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
None else []
"""

class Solution:

    def cloneTree(self, root: 'Node') -> 'Node':
        if root is None:
            return None
        children = [self.cloneTree(child) for
child in root.children]
        return Node(root.val, children)

```

♦ Quick Review

Key Insights

- The tree is N-ary, meaning every node can have multiple children.
- A deep copy requires creating new nodes for each node in the tree.
- Both DFS (recursive) or BFS (iterative) traversals can be used to clone the tree.

- The overall time complexity is $O(N)$ since each node is visited once.
- We can follow similar cloning techniques used in cloning graphs, using either recursion or an auxiliary data structure like a queue.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree, since each node is visited exactly once.

Space Complexity: $O(N)$, accounting for the recursive call stack (or an explicit queue in BFS) and the space needed to store the clone.

Solution

We perform a DFS traversal starting from the root node. For each node we encounter, we:

- Create a new node with the same value.
- Recursively clone all its children and attach the cloned children to the new node. This approach ensures that every original node is copied into a new node and that the parent-child relationships are preserved. Edge cases such as an empty tree (null root) are also handled by returning null immediately.

Trees & BST

#1522 Diameter of N-Ary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Tree · DFS

Lists: Premium 98

Problem


Description

Given a root of an N -ary tree, you need to compute the length of the diameter of the tree.

The diameter of an N -ary tree is the length of the longest path between any two nodes in the tree. This path may or may not pass through the root.


(Nary-Tree input serialization is represented in their level order traversal, each group of children is separated by the null value.)

Example 1:




Input: root = [1,null,3,2,4,null,5,6]
 Output: 3
 Explanation: Diameter is shown in red color.

Example 2:



Input: root = [1,null,2,null,3,4,null,5,null,6]
 Output: 4

Example 3:



Input: root = [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14]
Output: 7

Constraints:

- The depth of the n-ary tree is less than or equal to 1000.
- The total number of nodes is between [1, 104].

Community Solutions (Python)

- WalkCC

Round 2 · High · Medium

p.1103

```

class Solution:
    def diameter(self, root: 'Node') -> int:
        ans = 0

        def maxDepth(root: 'Node') -> int:
            """Returns the maximum depth of the subtree
            rooted at `root`."""
            nonlocal ans
            maxSubDepth1 = 0
            maxSubDepth2 = 0
            for child in root.children:
                maxSubDepth = maxDepth(child)
                if maxSubDepth > maxSubDepth1:
                    maxSubDepth2 = maxSubDepth1
                    maxSubDepth1 = maxSubDepth
                elif maxSubDepth > maxSubDepth2:
                    maxSubDepth2 = maxSubDepth
            ans = max(ans, maxSubDepth1 + maxSubDepth2)
            return 1 + maxSubDepth1

        maxDepth(root)

        return ans
  
```

- LeetCode

Round 2 · High · Medium

p.1104

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
None else []
"""

```

```

class Solution:

```

```

    def diameter(self, root: 'Node') -> int:
        """
        :type root: 'Node'
        :rtype: int
        """

        def dfs(root):
            if root is None:
                return 0
            nonlocal ans

```

```

                m1 = m2 = 0
                for child in root.children:
                    t = dfs(child)
                    if t > m1:
                        m2, m1 = m1, t
                    elif t > m2:
                        m2 = t
                ans = max(ans, m1 + m2)
            return 1 + m1

```

```

        ans = 0
        dfs(root)
        return ans

```

• SimplyLeet

```

# Definition for a Node.
# class Node:
#     def __init__(self, val=None, children=None):
#         self.val = val
#         self.children = children if children is not None
#         else []

```

```

class Solution:
    def diameter(self, root: 'Node') -> int:
        # Global variable to keep track of the
        maximum diameter found so far.
        self.max_diameter = 0

```

```

def dfs(node):
    if not node:
        return 0 # Base case: if the node
is None, return depth 0

    # Variables to hold the top two maximum
depths from the children.
    max_depth1, max_depth2 = 0, 0

    # Iterate over each child to compute
their maximum depths.
    for child in node.children:

        child_depth = dfs(child) # Recurse
on the child

        # Update the two largest depths
        if child_depth > max_depth1:
            max_depth2 = max_depth1
            max_depth1 = child_depth
        elif child_depth > max_depth2:
            max_depth2 = child_depth

    # Update the maximum diameter. The
candidate is the sum of the two deepest paths.
    self.max_diameter =
max(self.max_diameter, max_depth1 + max_depth2)

```

```

        # Return the maximum depth from this
node.

        return max_depth1 + 1

    dfs(root)
    return self.max_diameter

```

• LC.ca

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
None else []
"""

class Solution:

```

```
def diameter(self, root: 'Node') -> int:
    """
    :type root: 'Node'
    :rtype: int
    """
```

```
def dfs(root):
    if root is None:
        return 0
    nonlocal ans
```

```
    m1 = m2 = 0
    for child in root.children:
        t = dfs(child)
        if t > m1:
            m2, m1 = m1, t
        elif t > m2:
            m2 = t
    ans = max(ans, m1 + m2)
    return 1 + m1
```

```
ans = 0
dfs(root)
return ans
```

◆ Quick Review

Key Insights

- The diameter of the tree can be found by considering the two largest depths from any node among its children.

- Use a depth-first search (DFS) to compute the maximum depth from each node.
- For each node, the potential diameter is the sum of the two longest depths among its children.
- Update a global variable tracking the maximum diameter as you recurse through the tree.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack (worst-case $O(n)$ for a skewed tree).

Solution

Use a DFS approach to traverse the tree. At each node, calculate the longest and second longest depth among its children. The sum of these two depths gives the candidate diameter through that node. Maintain a global variable to track the maximum diameter found so far. Return the maximum depth from the node (i.e., one plus the largest depth among its children) to be used in parent's calculations.

Trees & BST

#1650 Lowest Common Ancestor of a Binary Tree III MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Tree · Two Pointers

Lists: Premium 98

Problem

Description


Given two nodes of a binary tree p and q , return their lowest common ancestor (LCA).

Each node will have a reference to its parent node. The definition for Node is below:

```
class Node {
    public int val;
    public Node left;
    public Node right;
    public Node parent;
}
```

According to the definition of LCA on Wikipedia: "The lowest common ancestor of two nodes p and q in a tree T is the lowest node that has both p and q as descendants (where we allow a node to be a descendant of itself)."

Example 1:




Input: root = [3,5,1,6,2,0,8,null,null,7,4], p = 5, q = 1

Output: 3

Explanation: The LCA of nodes 5 and 1 is 3.

Example 2:



Input: root = [3,5,1,6,2,0,8,null,null,7,4], p = 5, q = 4

Output: 5

Explanation: The LCA of nodes 5 and 4 is 5 since a node can be a descendant of itself according to the LCA definition.

Example 3:

Input: root = [1,2], p = 1, q = 2

Output: 1

Constraints:

- The number of nodes in the tree is in the range [2, 105].
- $-109 \leq \text{Node.val} \leq 109$
- All Node.val are unique.
- $p \neq q$
- p and q exist in the tree.

Community Solutions (Python)

• WalkCC

```

class Solution:
    # Same as 160. Intersection of Two Linked Lists
    def lowestCommonAncestor(self, p: 'Node', q: 'Node') -> 'Node':
        a = p
        b = q

        while a != b:
            a = a.parent if a else q
            b = b.parent if b else p

        return a
  
```

• LeetCode

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val):
        self.val = val
        self.left = None
        self.right = None
        self.parent = None
"""

class Solution:
    def lowestCommonAncestor(self, p: "Node", q:
"Node") -> "Node":
        vis = set()
        node = p
        while node:
            vis.add(node)
            node = node.parent
        node = q
        while node not in vis:
            node = node.parent
        return node

```

• SimplyLeet

Round 2 · High · Medium

p.1115

```

# Definition for a Node.
# class Node:
# def __init__(self, val=0, left=None, right=None,
parent=None):
# self.val = val
# self.left = left
# self.right = right
# self.parent = parent

def lowestCommonAncestor(p, q):
    # Compute the depth for node p

    def getDepth(node):
        depth = 0
        while node:
            depth += 1
            node = node.parent
        return depth

    depthP = getDepth(p)
    depthQ = getDepth(q)

```

Round 2 · High · Medium

p.1116

```
# Bring both nodes to the same depth level
while depthP > depthQ:
    p = p.parent
    depthP -= 1
while depthQ > depthP:
    q = q.parent
    depthQ -= 1

# Traverse up until we find the common ancestor
while p != q:
```

```
    p = p.parent
    q = q.parent
```

```
return p
```

• LC.ca

```
"""
# Definition for a Node.
class Node:
    def __init__(self, val):
        self.val = val
        self.left = None
        self.right = None
        self.parent = None
"""
```

Round 2 · High · Medium

p.1117

```
class Solution:
    def lowestCommonAncestor(self, p: 'Node', q:
'Node') -> 'Node':
        a, b = p, q
        while a != b:
            a = a.parent if a.parent else q
            b = b.parent if b.parent else p
        return a
```

```
#####
```

```
# ref: https://iameageek.medium.com/leetcode-1650-lowest-common-ancestor-of-a-binary-tree-iii-6d008b93376c
```

```
def lowestCommonAncestor(self, a: 'Node', b:
'Node') -> 'Node':
    ancestors = set()

    while a is not None:
        ancestors.add(a)
        a = a.parent
```

Round 2 · High · Medium

p.1118

```

while b is not None:
    if b in ancestors:
        return b
    b = b.parent

# or another one
def lowestCommonAncestor(self, a: 'Node', b:
'Node') -> 'Node':
    pointerA, pointerB = a, b

    while pointerA != pointerB:

        pointerA = pointerA.parent if pointerA
else b
        pointerB = pointerB.parent if pointerA
else a

    return pointerA

```

◆ Quick Review

Key Insights

- Each node has a parent pointer, so you can traverse from any node to the root.
- By determining the depth of each node, you can “align” the two nodes to the same depth.
- When both nodes are aligned, traverse upward simultaneously until they meet.

- This approach guarantees finding the LCA with $O(h)$ time complexity, where h is the height of the tree.

Space & Time

Time Complexity: $O(h)$, where h is the height of the tree.

Space Complexity: $O(1)$

Solution

We first compute the depth of both nodes by moving up their parent pointers. If one node is deeper, move it upward until both nodes are at the same level. Then, move both nodes upward one step at a time until they meet; the meeting point is the lowest common ancestor. This approach uses constant extra space and leverages the parent pointer for an efficient solution.

DFS

Round 2 · High · Medium · 16 questions

Round 2 · High · Medium

p.1121

DFS

#130 Surrounded Regions

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Union Find · Matrix

Lists: Denny Zhang

Problem

Description

You are given an $m \times n$ matrix board containing letters 'X' and 'O', capture regions that are surrounded:

- **Connect:** A cell is connected to adjacent cells horizontally or vertically.
- **Region:** To form a region connect every 'O' cell.
- **Surround:** A region is surrounded if none of the 'O' cells in that region are on the edge of the board. Such regions are completely enclosed by 'X' cells.

To capture a surrounded region, replace all 'O's with 'X's in-place within the original board. You do not need to return anything.


Example 1:

Input: board = `[["X","X","X","X"],["X","O","O","X"],["X","X","O","X"],["X","O","X","X"]]`

Output: `[["X","X","X","X"],["X","X","X","X"],["X","X","X","X"],["X","O","X","X"]]`

Round 2 · High · Medium

p.1122

Explanation:

In the above diagram, the bottom region is not captured because it is on the edge of the board and cannot be surrounded.

Example 2:

Input: board = [["X"]]

Output: [["X"]]

Constraints:

- $m == \text{board.length}$
- $n == \text{board}[i].\text{length}$
- $1 \leq m, n \leq 200$
- $\text{board}[i][j]$ is 'X' or 'O'.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def solve(self, board: list[list[str]]) -> None:
        if not board:
            return

        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(board)
        n = len(board[0])
        q = collections.deque()

        for i in range(m):
            for j in range(n):
                if i * j == 0 or i == m - 1 or j == n - 1:
                    if board[i][j] == 'O':
                        q.append((i, j))
                        board[i][j] = '*'

        # Mark the grids that stretch from the four
        # sides with '*'.
        while q:
            i, j = q.popleft()
```

```

for dx, dy in DIRS:
    x = i + dx
    y = j + dy
    if x < 0 or x == m or y < 0 or y == n:
        continue
    if board[x][y] != '0':
        continue
    q.append((x, y))
    board[x][y] = '*'

```

```

for row in board:
    for i, c in enumerate(row):
        row[i] = '0' if c == '*' else 'X'

```

• LeetCode

```

class Solution:
    def solve(self, board: List[List[str]]) ->
None:
    def dfs(i: int, j: int):
        if not (0 <= i < m and 0 <= j < n and
board[i][j] == "0"):
            return
        board[i][j] = "."
        for a, b in pairwise((-1, 0, 1, 0,
-1)):
            dfs(i + a, j + b)

    m, n = len(board), len(board[0])

```

Round 2 · High · Medium

p.1125

```

for i in range(m):
    dfs(i, 0)
    dfs(i, n - 1)
for j in range(n):
    dfs(0, j)
    dfs(m - 1, j)
for i in range(m):
    for j in range(n):
        if board[i][j] == ".":
            board[i][j] = "0"

```

```

elif board[i][j] == "0":
    board[i][j] = "X"

```

```

class Solution:
    def solve(self, board: List[List[str]]) ->
None:
    def find(x: int) -> int:
        if p[x] != x:
            p[x] = find(p[x])
        return p[x]

    m, n = len(board), len(board[0])
    p = list(range(m * n + 1))
    for i in range(m):

```

Round 2 · High · Medium

p.1126

```

        for j in range(n):
            if board[i][j] == "0":
                if i in (0, m - 1) or j in (0,
n - 1):
                    p[find(i * n + j)] = find(m
* n)
                else:
                    for a, b in pairwise((-1,
0, 1, 0, -1)):
                        x, y = i + a, j + b
                        if board[x][y] == "0":
                            p[find(x * n + y)]
= find(i * n + j)
                        for i in range(m):
                            for j in range(n):
                                if board[i][j] == "0" and find(i *
n + j) != find(m * n):
                                    board[i][j] = "X"

```

• SimplyLeet

Round 2 · High · Medium

p.1127

```

def solve(board):
    if not board or not board[0]:
        return

    rows, cols = len(board), len(board[0])

    def dfs(r, c):
        # Return if out of bounds or current cell
is not '0'
        if r < 0 or c < 0 or r >= rows or c >= cols
or board[r][c] != '0':
            return

        # Mark the cell as safe
        board[r][c] = '*'
        # Explore four possible directions
        dfs(r + 1, c)
        dfs(r - 1, c)
        dfs(r, c + 1)
        dfs(r, c - 1)

    # Traverse the board borders
    for i in range(rows):

```

Round 2 · High · Medium

p.1128


```

    if board[i][0] == '0':
        dfs(i, 0)
    if board[i][cols - 1] == '0':
        dfs(i, cols - 1)
for j in range(cols):
    if board[0][j] == '0':
        dfs(0, j)
    if board[rows - 1][j] == '0':
        dfs(rows - 1, j)

```

```

# Flip surrounded '0's and restore safe cells
for i in range(rows):
    for j in range(cols):
        if board[i][j] == '0':
            board[i][j] = 'X'
        elif board[i][j] == '*':
            board[i][j] = '0'

```

• LC.ca

Round 2 · High · Medium

p.1129

```

class Solution:
    def solve(self, board: List[List[str]]) ->
None:
    """
    Do not return anything, modify board
    in-place instead.
    """

    def dfs(i, j):
        board[i][j] = '.'
        for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:
            x, y = i + a, j + b

            if 0 <= x < m and 0 <= y < n and
board[x][y] == '0':
                dfs(x, y)

        m, n = len(board), len(board[0])
        for i in range(m):
            for j in range(n):
                if board[i][j] == '0' and (
                    i == 0 or i == m - 1 or j == 0
or j == n - 1
                ):
                    dfs(i, j)

```

Round 2 · High · Medium

p.1130

```

for i in range(m):
    for j in range(n):
        if board[i][j] == 'O':
            board[i][j] = 'X'
        elif board[i][j] == '.':
            board[i][j] = 'O'

```

◆ Quick Review

Key Insights

- Regions connected to the border are not captured.
- Use DFS/BFS from border 'O's to mark safe cells.
- After marking, flip all remaining 'O's to 'X' and then revert the marked cells back to 'O'.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(m * n)$ in the worst case due to recursion stack or queue storage for DFS/BFS

Solution

The solution involves two main steps. First, traverse the board's borders and run a DFS or BFS from any encountered 'O', marking it (e.g., using ' ') to indicate that it should not be flipped. Next, go through each cell in the board: flip any remaining 'O' (i.e., not marked safe) to 'X' and change the temporary marker ' ' back to 'O'. This approach ensures only the regions completely surrounded by 'X' are captured.

DFS

#133 Clone Graph

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · DFS · BFS · Graph

Lists: Grind 169

Problem

Description

Given a reference of a node in a connected undirected graph.

Return a deep copy (clone) of the graph.

Each node in the graph contains a value (int) and a list (List[Node]) of its neighbors.

```

class Node {
    public int val;
    public List<Node> neighbors;
}

```


Test case format:

For simplicity, each node's value is the same as the node's index (1-indexed). For example, the first node with `val == 1`, the second node with `val == 2`, and so on. The graph is represented in the test case using an adjacency list.

An adjacency list is a collection of unordered lists used to represent a finite graph. Each list describes the set of neighbors of a node in the graph.

The given node will always be the first node with val = 1. You must return the copy of the given node as a reference to the cloned graph.

Example 1:



Input: adjList = [[2,4],[1,3],[2,4],[1,3]]

Output: [[2,4],[1,3],[2,4],[1,3]]

Explanation: There are 4 nodes in the graph.

1st node (val = 1)'s neighbors are 2nd node (val = 2) and 4th node (val = 4).

2nd node (val = 2)'s neighbors are 1st node (val = 1) and 3rd node (val = 3).

3rd node (val = 3)'s neighbors are 2nd node (val = 2) and 4th node (val = 4).

4th node (val = 4)'s neighbors are 1st node (val = 1) and 3rd node (val = 3).

Example 2:



Input: adjList = [[]]
 Output: [[]]
 Explanation: Note that the input contains one empty list. The graph consists of only one node with val = 1 and it does not have any neighbors.

Example 3:

Input: adjList = []
 Output: []
 Explanation: This an empty graph, it does not have any nodes.

Constraints:

- The number of nodes in the graph is in the range [0, 100].
- $1 \leq \text{Node.val} \leq 100$
- Node.val is unique for each node.
- There are no repeated edges and no self-loops in the graph.
- The Graph is connected and all nodes can be visited starting from the given node.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def cloneGraph(self, node: 'Node') -> 'Node':
        if not node:
            return None

        q = collections.deque([node])
        map = {node: Node(node.val)}

        while q:
            u = q.popleft()
  
```

```

    for v in u.neighbors:
        if v not in map:
            map[v] = Node(v.val)
            q.append(v)
            map[u].neighbors.append(map[v])

    return map[node]

```

```

class Solution:
    def cloneGraph(self, node: 'Node') -> 'Node':
        if not node:
            return None
        if node in self.map:
            return self.map[node]

        newNode = Node(node.val, [])
        self.map[node] = newNode

```

```

        for neighbor in node.neighbors:
            self.map[node].neighbors.append(self.cloneGraph(neighbor))

        return newNode

    map = {}

```

• LeetCode

Round 2 · High · Medium

p.1137

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val = 0, neighbors = None):
        self.val = val
        self.neighbors = neighbors if neighbors is
not None else []
"""

```

```

from typing import Optional

```

```

class Solution:
    def cloneGraph(self, node: Optional["Node"])
-> Optional["Node"]:
        def dfs(node):
            if node is None:
                return None
            if node in g:
                return g[node]
            cloned = Node(node.val)
            g[node] = cloned

```

```

            for nxt in node.neighbors:
                cloned.neighbors.append(dfs(nxt))
            return cloned

```

```

        g = defaultdict()
        return dfs(node)

```

• SimplyLeet

Round 2 · High · Medium

p.1138

```
# Definition for a Node.
class Node:
    def __init__(self, val, neighbors=None):
        self.val = val
        # Initialize neighbors with empty list if
        None is provided
        self.neighbors = neighbors if neighbors is
        not None else []

def cloneGraph(node: 'Node') -> 'Node':
    # Hash map to track original nodes and their
    corresponding clones
    old_to_new = {}
```

```
def dfs(current):
    # Base case: if the current node is None,
    return None
    if current is None:
        return None
    # If the node has already been cloned,
    return the clone to avoid duplication and cycle
    issues
    if current in old_to_new:
        return old_to_new[current]
    # Create a clone for the current node
    copy = Node(current.val)
```

```
old_to_new[current] = copy
# Recursively clone all neighbors
for neighbor in current.neighbors:
    copy.neighbors.append(dfs(neighbor))
return copy

return dfs(node)
```

- LC.ca

```
"""
# Definition for a Node.
class Node:
    def __init__(self, val = 0, neighbors = None):
        self.val = val
        self.neighbors = neighbors if neighbors is
        not None else []
"""

class Solution:
```

```
def cloneGraph(self, node: 'Node') -> 'Node':
    visited = defaultdict()

    def clone(node):
        if node is None:
            return None
        if node in visited:
            return visited[node]
        c = Node(node.val)
        visited[node] = c

        for e in node.neighbors:
            c.neighbors.append(clone(e))
        return c

    return clone(node)
```

◆ Quick Review

Key Insights

- The graph is undirected and connected.
- Each node must be cloned exactly once to prevent cycles and repetitions.
- Use a mapping/dictionary to keep track of cloned nodes.
- A depth-first search (DFS) or breadth-first search (BFS) can be used to traverse and clone the graph.
- Handle edge cases where the graph might be empty.

Space & Time

Time Complexity: $O(N + E)$, where N is the number of nodes and E is the number of edges, because each node and each edge are traversed once.

Space Complexity: $O(N)$, to store the mapping from original nodes to their clones, and for the recursion/queue used in traversal.

Solution

We employ a depth-first search (DFS) approach to clone the graph. The main idea is to create a mapping (hash table) that links each original node to its corresponding cloned node. During the DFS traversal, if a node is encountered for the first time, a new node is created and stored in the mapping; then, the algorithm recursively clones all of its neighbors. For an already visited node, the cloned node is directly returned from the mapping. This ensures that every node is cloned only once and correctly handles cycles in the graph.

DFS

#200 Number of Islands

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Union Find · Matrix

Lists: Grind 169

Problem

Description

Given an $m \times n$ 2D binary grid `grid` which represents a map of '1's (land) and '0's (water), return the number of islands.

An island is surrounded by water and is formed by connecting adjacent lands horizontally or vertically. You may assume all four edges of the grid are all surrounded by water.

Example 1:

```
Input: grid = [
  ["1","1","1","1","0"],
  ["1","1","0","1","0"],
  ["1","1","0","0","0"],
  ["0","0","0","0","0"]
]
Output: 1
```

Example 2:

```
Input: grid = [
  ["1","1","0","0","0"],
  ["1","1","0","0","0"],
  ["0","0","1","0","0"],
  ["0","0","0","1","1"]
]
Output: 3
```

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$
- $1 \leq m, n \leq 300$
- `grid[i][j]` is '0' or '1'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numIslands(self, grid: list[list[str]]) -> int:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(grid)
        n = len(grid[0])

        def bfs(r, c):
            q = collections.deque([(r, c)])
            grid[r][c] = '2' # Mark '2' as visited.
            while q:
```



```

i, j = q.popleft()
for dx, dy in DIRS:
    x = i + dx
    y = j + dy
    if x < 0 or x == m or y < 0 or y == n:
        continue
    if grid[x][y] != '1':
        continue
    q.append((x, y))
    grid[x][y] = '2' # Mark '2' as visited.

```

```
ans = 0
```

```

for i in range(m):
    for j in range(n):
        if grid[i][j] == '1':
            bfs(i, j)
            ans += 1

```

```
return ans
```

Round 2 · High · Medium

p.1145

```

class Solution:
    def numIslands(self, grid: list[list[str]]) ->
    int:
        m = len(grid)
        n = len(grid[0])

        def dfs(i: int, j: int) -> None:
            if i < 0 or i == m or j < 0 or j == n:
                return
            if grid[i][j] != '1':
                return

```

```

        grid[i][j] = '2' # Mark '2' as visited.
        dfs(i + 1, j)
        dfs(i - 1, j)
        dfs(i, j + 1)
        dfs(i, j - 1)

```

```
ans = 0
```

```
for i in range(m):
```

```

    for j in range(n):
        if grid[i][j] == '1':
            dfs(i, j)
            ans += 1

```

```
return ans
```

• LeetCode

Round 2 · High · Medium

p.1146

```

class Solution:
    def numIslands(self, grid: List[List[str]]) ->
int:
    def dfs(i, j):
        grid[i][j] = '0'
        for a, b in pairwise(dirs):
            x, y = i + a, j + b
            if 0 <= x < m and 0 <= y < n and
grid[x][y] == '1':
                dfs(x, y)

    ans = 0

    dirs = (-1, 0, 1, 0, -1)
    m, n = len(grid), len(grid[0])
    for i in range(m):
        for j in range(n):
            if grid[i][j] == '1':
                dfs(i, j)
                ans += 1

    return ans

```

Round 2 · High · Medium

p.1147

```

class Solution:
    def numIslands(self, grid: List[List[str]]) ->
int:
    def bfs(i, j):
        grid[i][j] = '0'
        q = deque([(i, j)])
        while q:
            i, j = q.popleft()
            for a, b in pairwise(dirs):
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n
and grid[x][y] == '1':
                    q.append((x, y))
                    grid[x][y] = 0

    ans = 0
    dirs = (-1, 0, 1, 0, -1)
    m, n = len(grid), len(grid[0])
    for i in range(m):
        for j in range(n):
            if grid[i][j] == '1':
                bfs(i, j)

    ans += 1

    return ans

```

Round 2 · High · Medium

p.1148

```
class Solution:
    def numIslands(self, grid: List[List[str]]) ->
    int:
        def find(x):
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        dirs = (0, 1, 0)
        m, n = len(grid), len(grid[0])
        p = list(range(m * n))

        for i in range(m):
            for j in range(n):
                if grid[i][j] == '1':
                    for a, b in pairwise(dirs):
                        x, y = i + a, j + b
                        if x < m and y < n and
grid[x][y] == '1':
                            p[find(i * n + j)] =
find(x * n + y)
            return sum(
                grid[i][j] == '1' and i * n + j ==
find(i * n + j)
                for i in range(m)
                for j in range(n)
            )
```

• SimplyLeet

Round 2 · High · Medium

p.1149

```
# Function to count the number of islands in the
grid
def numIslands(grid):
    # If the grid is empty, return 0
    if not grid:
        return 0
    rows, cols = len(grid), len(grid[0])
    islands = 0

    # Depth-First Search helper to mark the entire
    island as visited
    def dfs(r, c):

        # Check if the current cell is out of
        bounds or water
        if r < 0 or r >= rows or c < 0 or c >= cols
        or grid[r][c] == "0":
            return

        # Mark the cell as visited by setting it to
        water
        grid[r][c] = "0"

        # Explore all four adjacent cells (up,
        down, left, right)
        dfs(r + 1, c)
        dfs(r - 1, c)
        dfs(r, c + 1)
        dfs(r, c - 1)
```

Round 2 · High · Medium

p.1150

```

# Iterate over every cell in the grid
for r in range(rows):
    for c in range(cols):
        # If a land cell is found, conduct DFS
        # to visit the entire island
        if grid[r][c] == "1":
            islands += 1
            dfs(r, c)

return islands

```

• LC.ca

```

# dfs
class Solution:
    def numIslands(self, grid: List[List[str]]) ->
    int:
        def dfs(i, j):
            if not (0 <= i < m and 0 <= j < n and
            grid[i][j] == '1'):
                return
            grid[i][j] = '0'
            for a, b in pairwise(dirs):
                x, y = i + a, j + b
                dfs(x, y)

```

```

ans = 0
dirs = (-1, 0, 1, 0, -1)
m, n = len(grid), len(grid[0])
for i in range(m):
    for j in range(n):
        if grid[i][j] == '1':
            dfs(i, j)
            ans += 1

return ans

```

#####

```

# bfs
from collections import deque

class Solution:
    def numIslands(self, grid: List[List[str]]) ->
    int:
        if not grid:
            return 0

```

```

m, n = len(grid), len(grid[0])
directions = [(0, 1), (0, -1), (1, 0), (-1,
0)]

```

```
islands = 0
```

```
for i in range(m):
```

```
    for j in range(n):
```

```
        if grid[i][j] == "1":
```

```
            islands += 1
```

```
            q = deque([(i, j)]) # no need
to reset q, q already drained from previous bfs

```

```
        while q:
```

```
            x, y = q.popleft()
```

```
            for dx, dy in directions:
```

```
                nx, ny = x + dx, y + dy
```

```
                if 0 <= nx < m and 0 <=
```

```
ny < n and grid[nx][ny] == "1":
```

```
                    q.append((nx, ny))
```

```
                    grid[nx][ny] = "0"
```

```
# mark as visited
```

```
    return islands
```

```
#####
```

```
# union find
```

```
# similar to UF in https://leetcode.ca/2016-09-30-305-Number-of-Islands-II/
```

```
class Solution:
```

```
    def numIslands(self, grid: List[List[str]]) ->
int:
```

```
        def check(i, j):
```

```
            return 0 <= i < m and 0 <= j < n and
```

```
grid[i][j] == "1" # Check for "1" instead of 1
```

```
        def find(x):
```

```
            if p[x] != x:
```

```
                p[x] = find(p[x])
```

```
            return p[x]
```

```
m, n = len(grid), len(grid[0])
```

```
p = list(range(m * n))
```

```
cur = 0 # Initialize cur as 0 instead of n
```

```
for i in range(m):
```

```
    for j in range(n):
```

```
        if grid[i][j] == "1":
```

```
            cur += 1 # Increment cur when
```

```
encountering "1"
```

```

        for x, y in [(-1, 0), (1, 0),
(0, -1), (0, 1)]:
            if check(i + x, j + y) and
find(i * n + j) != find((i + x) * n + j + y):
                p[find(i * n + j)] =
find((i + x) * n + j + y)
                cur -= 1

    return cur

```

```
#####
```

```

...
>>> a = set()
>>>
>>> a |= {(1, 1)}
>>> a
{(1, 1)}
>>>
>>> a |= {(2, 2)}
>>> a
{(1, 1), (2, 2)}

```

Round 2 · High · Medium

p.1155

```

>>>
>>> a.add((3, 3))
>>> a
{(1, 1), (3, 3), (2, 2)}
>>>
>>> (1,1) in a
True
>>> (10,10) in a
False
...

```

```

class Solution(object):
    def numIslands(self, grid):
        """
        :type grid: List[List[str]]
        :rtype: int
        """
        visited = set()
        ans = 0

        def dfs(grid, i, j, visited):

```

Round 2 · High · Medium

p.1156

```

        if i < 0 or i >= len(grid) or j < 0 or j >=
len(grid[0]) or grid[i][j] == "0" or (i, j) in
visited:
            return False
        visited |= {(i, j)}
        for di, dj in [(-1, 0), (1, 0), (0, 1), (0,
-1)]:
            newi, newj = i + di, j + dj
            dfs(grid, newi, newj, visited)
        return True

    for i in range(0, len(grid)):
        for j in range(0, len(grid[0])):

            if dfs(grid, i, j, visited):
                ans += 1
    return ans

```

◆ Quick Review

Key Insights

- Use graph traversal (DFS/BFS) to explore connected components of land.
- Iterate over the matrix and start a traversal (e.g., DFS) when encountering an unvisited "1".
- During traversal, mark visited cells to avoid duplicate counting.
- Each DFS/BFS call completely explores one island.

- Alternative approach can use Union Find to group connected lands.

Space & Time

Time Complexity: $O(m * n)$, as each cell is visited at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario due to recursion stack (DFS) or auxiliary data structures.

Solution

We solve the problem using a DFS-based approach. The algorithm iterates over each cell in the grid; when a "1" (land) is encountered, the DFS is initiated to mark the entire island (all connected "1"s) as visited (by setting them to "0"). This prevents recounting the same island. The island count is incremented each time a new island is discovered and fully explored.

DFS

#207 Course Schedule

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Graph · Topological Sort

Lists: Grind 169

Problem

Description

There are a total of `numCourses` courses you have to take, labeled from 0 to `numCourses - 1`. You are given an array `prerequisites` where `prerequisites[i] = [ai, bi]` indicates that you must take course `bi` first if you want to take course `ai`.

- For example, the pair `[0, 1]`, indicates that to take course 0 you have to first take course 1.

Return true if you can finish all courses. Otherwise, return false.

Example 1:

```
Input: numCourses = 2, prerequisites = [[1,0]]
Output: true
Explanation: There are a total of 2 courses to take.
To take course 1 you should have finished course 0. So
it is possible.
```

Example 2:

```
Input: numCourses = 2, prerequisites = [[1,0],[0,1]]
Output: false
Explanation: There are a total of 2 courses to take.
To take course 1 you should have finished course 0, and
to take course 0 you should also have finished course
1. So it is impossible.
```

Constraints:

- $1 \leq \text{numCourses} \leq 2000$
- $0 \leq \text{prerequisites.length} \leq 5000$
- $\text{prerequisites}[i].\text{length} == 2$
- $0 \leq a_i, b_i < \text{numCourses}$
- All the pairs `prerequisites[i]` are unique.

Community Solutions (Python)

- WalkCC

```
from enum import Enum

class State(Enum):
    INIT = 0
    VISITING = 1
    VISITED = 2

class Solution:
```



```
def canFinish(self, numCourses: int,
prerequisites: list[list[int]]) -> bool:
    graph = [[] for _ in range(numCourses)]
    states = [State.INIT] * numCourses
```

```
    for v, u in prerequisites:
        graph[u].append(v)
```

```
    def hasCycle(u: int) -> bool:
        if states[u] == State.VISITING:
            return True
```

```
        if states[u] == State.VISITED:
            return False
        states[u] = State.VISITING
        if any(hasCycle(v) for v in graph[u]):
            return True
        states[u] = State.VISITED
        return False
```

```
    return not any(hasCycle(i) for i in
range(numCourses))
```

• LeetCode

Round 2 · High · Medium

p.1161

```
class Solution:
    def canFinish(self, numCourses: int,
prerequisites: List[List[int]]) -> bool:
        g = [[] for _ in range(numCourses)]
        indeg = [0] * numCourses
        for a, b in prerequisites:
            g[b].append(a)
            indeg[a] += 1
        q = [i for i, x in enumerate(indeg) if x ==
0]
        for i in q:
            numCourses -= 1

            for j in g[i]:
                indeg[j] -= 1
                if indeg[j] == 0:
                    q.append(j)
        return numCourses == 0
```

• SimplyLeet

Round 2 · High · Medium

p.1162

```

from collections import deque

def canFinish(numCourses, prerequisites):
    # Create a list to hold the number of
    # prerequisites (indegree) for each course.
    indegree = [0] * numCourses
    # Create an adjacency list for each course.
    graph = {i: [] for i in range(numCourses)}

    # Build the graph and update indegree for each
    # course.
    for course, prereq in prerequisites:

        graph[prereq].append(course)
        indegree[course] += 1

    # Initialize a queue with courses that have no
    # prerequisites.
    queue = deque([i for i in range(numCourses) if
        indegree[i] == 0])

    # Counter for courses that can be completed.
    visited = 0

    while queue:

```

```

        course = queue.popleft() # Remove course
        with no remaining prerequisites.
        visited += 1 # Increment count of completed
        courses.
        # For each neighbor course, reduce its
        # indegree.
        for neighbor in graph[course]:
            indegree[neighbor] -= 1 # Remove
            dependency.
            # If no more prerequisites, add to
            # queue.
            if indegree[neighbor] == 0:
                queue.append(neighbor)

        # If we processed all courses, it's possible to
        # finish.

    return visited == numCourses

# Example usage:
print(canFinish(2, [[1,0]])) # Expected output:
True
print(canFinish(2, [[1,0],[0,1]])) # Expected
output: False

```

- LC.ca

```
...
>>> from collections import defaultdict
```

```
>>> a = defaultdict(int)
>>> a
defaultdict(<type 'int'>, {})
>>>
>>> a['hehehe']
0
>>> a
```

```
defaultdict(<type 'int'>, {'hehehe': 0})
```

```
>>> a = defaultdict(list)
>>> a
defaultdict(<type 'list'>, {})
>>> a['hehehe']
[]
>>> a
defaultdict(<type 'list'>, {'hehehe': []})
```

```
...
class Solution:
    def canFinish(self, numCourses: int,
                  prerequisites: List[List[int]]) -> bool:
        g = defaultdict(list)
        indeg = [0] * numCourses
        for a, b in prerequisites:
            g[b].append(a)
            indeg[a] += 1
        cnt = 0
        q = deque([i for i, v in enumerate(indeg)
                  if v == 0])
```

```
        while q:
            i = q.popleft()
            cnt += 1
            for j in g[i]:
                indeg[j] -= 1
                if indeg[j] == 0:
                    q.append(j)
        return cnt == numCourses
```

◆ Quick Review

Key Insights

- This problem can be modeled as a directed graph where courses are nodes and prerequisites are edges.
- If the graph has a cycle, then it is not possible to complete all courses.

- We can detect cycles using either Topological Sort (BFS/Kahn's algorithm) or DFS-based cycle detection.
- The time complexity is $O(V + E)$, where V is the number of courses and E is the number of prerequisite pairs.

Space & Time

Time Complexity: $O(V + E)$

Space Complexity: $O(V + E)$

Solution

We solve the problem using the Topological Sort algorithm (BFS/Kahn's algorithm). First, we build a graph (using an adjacency list) and track the indegree (number of prerequisites) for each course. We then find all courses with an indegree of zero (courses with no prerequisites) and add them to a processing queue. While the queue is not empty, we remove a course from the queue, reduce the indegree for its neighboring courses, and if any neighbor's indegree becomes zero, we add it to the queue as well. If we successfully process all courses, then it means that the graph was acyclic and a valid schedule exists, returning true. If not, a cycle exists, and we return false.

DFS

#210 Course Schedule II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Graph · Topological Sort

Lists: Grind 169

Problem

Description

There are a total of `numCourses` courses you have to take, labeled from 0 to `numCourses - 1`. You are given an array `prerequisites` where `prerequisites[i] = [ai, bi]` indicates that you must take course `bi` first if you want to take course `ai`.

- For example, the pair `[0, 1]`, indicates that to take course 0 you have to first take course 1.

Return the ordering of courses you should take to finish all courses. If there are many valid answers, return any of them. If it is impossible to finish all courses, return an empty array.

Example 1:

Input: `numCourses = 2, prerequisites = [[1,0]]`

Output: `[0,1]`

Explanation: There are a total of 2 courses to take. To take course 1 you should have finished course 0. So the correct course order is `[0,1]`.

Example 2:

Input: numCourses = 4, prerequisites =
[[1,0],[2,0],[3,1],[3,2]]

Output: [0,2,1,3]

Explanation: There are a total of 4 courses to take. To take course 3 you should have finished both courses 1 and 2. Both courses 1 and 2 should be taken after you finished course 0.

So one correct course order is [0,1,2,3]. Another correct ordering is [0,2,1,3].

Example 3:

Input: numCourses = 1, prerequisites = []

Output: [0]

Constraints:

- $1 \leq \text{numCourses} \leq 2000$
- $0 \leq \text{prerequisites.length} \leq \text{numCourses} * (\text{numCourses} - 1)$
- $\text{prerequisites}[i].\text{length} == 2$
- $0 \leq a_i, b_i < \text{numCourses}$
- $a_i \neq b_i$
- All the pairs $[a_i, b_i]$ are distinct.

Community Solutions (Python)

- WalkCC

```
from enum import Enum
```

```
class State(Enum):
```

```
    INIT = 0
```

```
    VISITING = 1
```

```
    VISITED = 2
```

```
class Solution:
```

```
    def findOrder(
```

```
        self,
```

```
        numCourses: int,
```

```
        prerequisites: list[list[int]],
```

```
    ) -> list[int]:
```

```
        ans = []
```

```
        graph = [[] for _ in range(numCourses)]
```

```
        states = [State.INIT] * numCourses
```

```
        for v, u in prerequisites:
```

```

graph[u].append(v)

def hasCycle(u: int) -> bool:
    if states[u] == State.VISITING:
        return True
    if states[u] == State.VISITED:
        return False
    states[u] = State.VISITING
    if any(hasCycle(v) for v in graph[u]):
        return True

```

```

states[u] = State.VISITED
ans.append(u)
return False

if any(hasCycle(i) for i in
range(numCourses)):
    return []

return ans[::-1]

```

```

class Solution:
    def findOrder(
        self,
        numCourses: int,
        prerequisites: list[list[int]],
    ) -> list[int]:
        ans = []
        graph = [[] for _ in range(numCourses)]
        inDegrees = [0] * numCourses
        q = collections.deque()

```

```

        # Build the graph.
        for v, u in prerequisites:
            graph[u].append(v)
            inDegrees[v] += 1

        # Perform topological sorting.
        q = collections.deque([i for i, d in
enumerate(inDegrees) if d == 0])

        while q:

```

```

u = q.popleft()
ans.append(u)
for v in graph[u]:
    inDegrees[v] -= 1
    if inDegrees[v] == 0:
        q.append(v)

return ans if len(ans) == numCourses else []

```

• LeetCode

```

class Solution:
    def findOrder(self, numCourses: int,
prerequisites: List[List[int]]) -> List[int]:
    g = defaultdict(list)
    indeg = [0] * numCourses
    for a, b in prerequisites:
        g[b].append(a)
        indeg[a] += 1
    ans = []
    q = deque(i for i, x in enumerate(indeg) if
x == 0)
    while q:

```

```

i = q.popleft()
ans.append(i)
for j in g[i]:
    indeg[j] -= 1
    if indeg[j] == 0:
        q.append(j)
return ans if len(ans) == numCourses else
[]

```

• SimplyLeet

```

# Python solution using BFS (Kahn's Algorithm)
from collections import deque

def findOrder(numCourses, prerequisites):
    # Build the graph and the in-degree array
    in_degree = [0] * numCourses
    graph = {i: [] for i in range(numCourses)}
    for course, pre in prerequisites:
        graph[pre].append(course)
        in_degree[course] += 1

```

```
# Initialize the queue with all courses having
zero in-degree
queue = deque([i for i in range(numCourses) if
in_degree[i] == 0])
order = []
```

```
# Process courses
while queue:
    course = queue.popleft()
    order.append(course)
    for neighbor in graph[course]:
```

```
        in_degree[neighbor] -= 1
        if in_degree[neighbor] == 0:
            queue.append(neighbor)
```

```
# If the ordering includes all courses, return
it; otherwise, return an empty list
return order if len(order) == numCourses else
[]
```

```
# Example usage:
# print(findOrder(4, [[1,0],[2,0],[3,1],[3,2]]))
```

• LC.ca

Round 2 · High · Medium

p.1175

```
class Solution:
    def findOrder(self, numCourses: int,
prerequisites: List[List[int]]) -> List[int]:
        g = defaultdict(list)
        indeg = [0] * numCourses
        for a, b in prerequisites:
            g[b].append(a)
            indeg[a] += 1
        q = deque([i for i, v in enumerate(indeg)
if v == 0])
        ans = [] # added from previous question
        while q:
```

```
            i = q.popleft()
            ans.append(i)
            # assumption is only one path
            # as in question 'You may assume that
there are no duplicate edges in the input
prerequisites.'
            for j in g[i]:
                indeg[j] -= 1
                if indeg[j] == 0:
                    q.append(j)
            return ans if len(ans) == numCourses else
[]
```

◆ Quick Review

Key Insights

- This is a topological sort problem on a directed graph.

Round 2 · High · Medium

p.1176

- The courses are nodes, and prerequisites form directed edges.
- We can solve the problem by either:
 - Using DFS with cycle detection and a stack for topological ordering.
 - Using BFS (Kahn's Algorithm) to process nodes with zero in-degrees.
- If at any point no node with a zero in-degree is available, a cycle exists and no valid order can be produced.

Space & Time

Time Complexity: $O(V + E)$ where V is the number of courses and E is the number of prerequisite pairs.

Space Complexity: $O(V + E)$ for storing the graph and in-degree array.

Solution

We treat the problem as a topological sort on a directed acyclic graph (DAG). For a BFS approach using Kahn's algorithm:

- Build an adjacency list for the graph and an in-degree array to record the number of prerequisites (incoming edges) for each course.
 - Initialize a queue with all courses that have an in-degree of zero (i.e., courses that can be taken without prerequisites).
 - Dequeue nodes from the queue, add them to the ordering, and reduce the in-degree of their neighbors.
 - If a neighbor's in-degree becomes zero, add it to the queue.
 - If the final ordering includes all courses, return the ordering; otherwise, return an empty array indicating a cycle.
- The chosen data structures are:

- A list (or array) for the in-degree counts.
- A dictionary (or array of lists) for the adjacency list.
- A queue to process nodes with zero in-degree.

DFS

#261 Graph Valid Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Union Find · Graph

Lists: Grind 169 | Premium 98


Problem

Description

You have a graph of n nodes labeled from 0 to $n - 1$. You are given an integer n and a list of edges where $\text{edges}[i] = [a_i, b_i]$ indicates that there is an undirected edge between nodes a_i and b_i in the graph.


Return true if the edges of the given graph make up a valid tree, and false otherwise.

Example 1:



Input: $n = 5$, $\text{edges} = [[0,1],[0,2],[0,3],[1,4]]$
Output: true

Example 2:



Input: $n = 5$, $edges = [[0,1],[1,2],[2,3],[1,3],[1,4]]$

Output: false

Constraints:

- $1 \leq n \leq 2000$
- $0 \leq edges.length \leq 5000$
- $edges[i].length == 2$
- $0 \leq ai, bi < n$
- $ai \neq bi$
- There are no self-loops or repeated edges.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def validTree(self, n: int, edges:
list[list[int]]) -> bool:
        if n == 0 or len(edges) != n - 1:
            return False

```

```

        graph = [[] for _ in range(n)]
        q = collections.deque([0])
        seen = {0}

```

```

        for u, v in edges:

```

```

            graph[u].append(v)
            graph[v].append(u)

```

```

        while q:
            u = q.popleft()
            for v in graph[u]:
                if v not in seen:
                    q.append(v)
                    seen.add(v)

```

```

        return len(seen) == n

```

```

class UnionFind:
    def __init__(self, n: int):
        self.count = n
        self.id = list(range(n))
        self.rank = [0] * n

    def unionByRank(self, u: int, v: int) -> None:
        i = self._find(u)
        j = self._find(v)
        if i == j:
            return
        if self.rank[i] < self.rank[j]:
            self.id[i] = j
        elif self.rank[i] > self.rank[j]:
            self.id[j] = i
        else:
            self.id[i] = j
            self.rank[j] += 1
        self.count -= 1

```

```

def _find(self, u: int) -> int:
    if self.id[u] != u:
        self.id[u] = self._find(self.id[u])
    return self.id[u]

```

```

class Solution:
    def validTree(self, n: int, edges:
list[list[int]]) -> bool:
        if n == 0 or len(edges) != n - 1:
            return False

```

```

        uf = UnionFind(n)

        for u, v in edges:
            uf.unionByRank(u, v)

        return uf.count == 1

```

• LeetCode

```
class Solution:
    def validTree(self, n: int, edges:
List[List[int]]) -> bool:
        def find(x: int) -> int:
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        p = list(range(n))
        for a, b in edges:
            pa, pb = find(a), find(b)
```

```
            if pa == pb:
                return False
            p[pa] = pb
            n -= 1
        return n == 1
```

```
class Solution:
    def validTree(self, n: int, edges:
List[List[int]]) -> bool:
        def dfs(i: int):
            vis.add(i)
            for j in g[i]:
                if j not in vis:
                    dfs(j)

        if len(edges) != n - 1:
            return False
```

Round 2 · High · Medium

p.1185

```
g = [[] for _ in range(n)]
for a, b in edges:
    g[a].append(b)
    g[b].append(a)
vis = set()
dfs(0)
return len(vis) == n
```

• SimplyLeet

Python solution with DFS

```
def validTree(n, edges):
    # If the number of edges is not exactly n - 1,
    it cannot be a tree
    if len(edges) != n - 1:
        return False
```

```
    # Build the graph as an adjacency list
    graph = {i: [] for i in range(n)}
    for a, b in edges:
```

Round 2 · High · Medium

p.1186

```

graph[a].append(b)
graph[b].append(a)

visited = set()

# DFS helper function that accepts current node
and its parent node
def dfs(node, parent):
    # Add the current node to visited
    visited.add(node)
    # Traverse all the neighboring nodes

    for neighbor in graph[node]:
        # Skip the neighbor if it is the parent
        to avoid trivial cycle detection
        if neighbor == parent:
            continue
        # If neighbor is already visited, a
        cycle is detected
        if neighbor in visited:
            return False
        # Recurse deeper; if cycle is found in
        recursion, propagate False upward
        if not dfs(neighbor, node):
            return False

```

Round 2 · High · Medium

p.1187

```

return True

# Start DFS from node 0, with parent set to -1
(non-existent)
if not dfs(0, -1):
    return False

# Check if all nodes were visited (i.e., the
graph is connected)
return len(visited) == n

# Example usage:

print(validTree(5, [[0, 1], [0, 2], [0, 3], [1,
4]])) # Expected output: True
print(validTree(5, [[0, 1], [1, 2], [2, 3], [1, 3],
[1, 4]])) # Expected output: False

```

• LC.ca

```

class Solution:
    def validTree(self, n: int, edges:
List[List[int]]) -> bool:
        def find(x):
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        p = list(range(n))
        for a, b in edges:
            if find(a) == find(b):

```

Round 2 · High · Medium

p.1188

```

        return False
    p[find(a)] = find(b)
    n -= 1
    return n == 1

```

◆ Quick Review

Key Insights

- A tree must be acyclic and connected.
- For n nodes, a valid tree must have exactly $n - 1$ edges.
- A DFS/BFS traversal can check connectivity and simultaneously detect cycles.
- The Union-Find (disjoint set) approach provides an alternative for cycle detection and connectivity verification.

Space & Time

Time Complexity: $O(n + e)$, where e is the number of edges

Space Complexity: $O(n + e)$ for storing the graph and traversal/union-find structures

Solution

We can solve the problem using either a DFS/BFS approach or a Union-Find algorithm. For the DFS/BFS method, we first check if the number of edges is exactly $n - 1$. If not, it cannot be a tree. Then, we build an adjacency list for the graph and perform a DFS starting from node 0 while keeping track of visited nodes and parent information to avoid false-positive cycle detections. If during DFS we encounter a node that has been visited (and is not the parent), then a cycle exists. Finally, if all nodes are reached

from node 0 and no cycle is found, the graph is a valid tree. Using Union-Find, we iterate over each edge and try to union the two nodes. If they already share the same parent (i.e., are in the same set), a cycle exists. After processing all edges, we confirm that exactly one connected component exists.

DFS

#310 Minimum Height Trees

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Graph · Topological Sort

Lists: Grind 169

Problem

Description


A tree is an undirected graph in which any two vertices are connected by exactly one path. In other words, any connected graph without simple cycles is a tree.

Given a tree of n nodes labelled from 0 to $n - 1$, and an array of $n - 1$ edges where $\text{edges}[i] = [a_i, b_i]$ indicates that there is an undirected edge between the two nodes a_i and b_i in the tree, you can choose any node of the tree as the root. When you select a node x as the root, the result tree has height h . Among all possible rooted trees, those with minimum height (i.e. $\min(h)$) are called minimum height trees (MHTs).

Return a list of all MHTs' root labels. You can return the answer in any order.

The height of a rooted tree is the number of edges on the longest downward path between the root and a leaf.

Example 1:




Input: $n = 4$, $\text{edges} = [[1,0],[1,2],[1,3]]$

Output: [1]

Explanation: As shown, the height of the tree is 1 when the root is the node with label 1 which is the only MHT.

Example 2:



Input: $n = 6$, $\text{edges} = [[3,0],[3,1],[3,2],[3,4],[5,4]]$

Output: [3,4]

Constraints:

- $1 \leq n \leq 2 * 10^4$
- $\text{edges.length} == n - 1$
- $0 \leq a_i, b_i < n$

- $a_i \neq b_i$
- All the pairs (a_i, b_i) are distinct.
- The given input is guaranteed to be a tree and there will be no repeated edges.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findMinHeightTrees(self, n: int, edges:
list[list[int]]) -> list[int]:
        if n == 1 or not edges:
            return [0]

        ans = []
        graph = collections.defaultdict(set)

        for u, v in edges:
            graph[u].add(v)

            graph[v].add(u)

        for label, children in graph.items():
            if len(children) == 1:
                ans.append(label)

        while n > 2:
            n -= len(ans)
            nextLeaves = []
            for leaf in ans:
```

```
                u = next(iter(graph[leaf]))
                graph[u].remove(leaf)
                if len(graph[u]) == 1:
                    nextLeaves.append(u)
            ans = nextLeaves
```

```
        return ans
```

• LeetCode

```
class Solution:
    def findMinHeightTrees(self, n: int, edges:
List[List[int]]) -> List[int]:
        if n == 1:
            return [0]
        g = [[] for _ in range(n)]
        degree = [0] * n
        for a, b in edges:
            g[a].append(b)
            g[b].append(a)
            degree[a] += 1
            degree[b] += 1
```

```

        degree[b] += 1
    q = deque(i for i in range(n) if degree[i]
== 1)
    ans = []
    while q:
        ans.clear()
        for _ in range(len(q)):
            a = q.popleft()
            ans.append(a)
            for b in g[a]:
                degree[b] -= 1

        if degree[b] == 1:
            q.append(b)

    return ans

```

• SimplyLeet

```

# Python solution with line-by-line comments
class Solution:
    def findMinHeightTrees(self, n: int, edges:
List[List[int]]) -> List[int]:
        # Special case: if there's only one node,
        return it as the only MHT root.
        if n == 1:
            return [0]

        # Build the graph using an adjacency list
        with sets for neighbors.
        graph = [set() for _ in range(n)]
        for u, v in edges:

```

```

        graph[u].add(v)
        graph[v].add(u)

        # Initialize the first layer of leaves
        (nodes with one neighbor).
        leaves = [i for i in range(n) if
len(graph[i]) == 1]

        # Remove leaves layer by layer until at
        most 2 nodes remain.
        remaining_nodes = n
        while remaining_nodes > 2:
            remaining_nodes -= len(leaves)

        new_leaves = []
        # Remove current leaves.
        for leaf in leaves:
            # Since it's a leaf, there is only
            one neighbor.

            neighbor = graph[leaf].pop()
            # Remove the connection from the
            neighbor to the leaf.
            graph[neighbor].remove(leaf)
            # If the neighbor becomes a leaf,
            add it to the new leaves list.
            if len(graph[neighbor]) == 1:
                new_leaves.append(neighbor)

```

```

        leaves = new_leaves

    # The remaining nodes are the roots of the
    MHTs.
    return leaves

```

• LC.ca

```

from typing import List, Set
from collections import defaultdict

class Solution:
    def findMinHeightTrees(self, n: int, edges:
List[List[int]]) -> List[int]:
        if n == 1:
            return [0]

        graph = defaultdict(set)

```

```

        for a, b in edges:
            graph[a].add(b)
            graph[b].add(a)

        # just check neighbours[] size, not
        indgree[] for each node
        leaves = [i for i in range(n) if
len(graph[i]) == 1]

        while n > 2:
            n -= len(leaves)
            new_leaves = []

            for leaf in leaves:
                # should be only one in hashset
                pop(), because it's a leaf node
                neighbor = graph[leaf].pop()
                graph[neighbor].remove(leaf)
                if len(graph[neighbor]) == 1:
                    new_leaves.append(neighbor)
            leaves = new_leaves

        return leaves

```

```
#####
```

```
class Solution:
    def findMinHeightTrees(self, n: int, edges:
List[List[int]]) -> List[int]:
    if n == 1:
        return [0]
    g = defaultdict(list)
    degree = [0] * n
    for a, b in edges:
        g[a].append(b)
```

```
        g[b].append(a)
        degree[a] += 1 # not needed, just
len(g[a]) is good enough, will update it in next
round updating
        degree[b] += 1
    q = deque()
    for i in range(n):
        if degree[i] == 1:
            q.append(i)
    ans = []
    while q:
        n = len(q)
```

```
ans.clear()
for _ in range(n):
    a = q.popleft()
    ans.append(a)
    for b in g[a]:
        degree[b] -= 1
        if degree[b] == 1: # final
round only 2 left (a,b), then degree[b] here will
be 0, and no more node enqueue
            q.append(b)

return ans
```

◆ Quick Review

Key Insights

- A tree can have at most two centers that will serve as the roots for its minimum height trees.
- Instead of trying all possible roots, we can remove leaves layer by layer until we are left with the central node(s).
- The process is similar to a topological sort on the tree.
- When there are 2 or fewer nodes left, these nodes are the centers (roots) of MHTs.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since we process each node and edge once.

Space Complexity: $O(n)$ for storing the graph and additional data structures like the leaves list.

Solution

The solution builds an undirected graph from the input, then identifies and removes node-leaves iteratively. A leaf is defined as a node connected to only one other node. By repeatedly removing the leaves and updating the graph (decreasing neighbor degrees), the algorithm peels away the outer layers of the tree. When only one or two nodes remain, these nodes are the centers of the tree and form the roots of the minimum height trees.

Round 2 · High · Medium

p.1201

DFS

#323 Number of Connected Components in an Undirected Graph

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Union Find · Graph

Lists: Grind 169 | Premium 98

Problem


Description

You have a graph of n nodes. You are given an integer n and an array `edges` where `edges[i] = [ai, bi]` indicates that there is an edge between a_i and b_i in the graph. Return the number of connected components in the graph.

Example 1:

Round 2 · High · Medium


p.1202



Input: $n = 5$, $edges = [[0,1],[1,2],[3,4]]$

Output: 2

Example 2:



Input: $n = 5$, $edges = [[0,1],[1,2],[2,3],[3,4]]$

Output: 1

Constraints:

- $1 \leq n \leq 2000$
- $1 \leq edges.length \leq 5000$
- $edges[i].length == 2$
- $0 \leq a_i \leq b_i < n$
- $a_i \neq b_i$
- There are no repeated edges.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def countComponents(self, n: int, edges:
list[list[int]]) -> int:
    ans = 0
    graph = [[] for _ in range(n)]
    seen = set()

    for u, v in edges:
        graph[u].append(v)
        graph[v].append(u)
```

```
def bfs(node: int, seen: set[int]) -> None:
    q = collections.deque([node])
    seen.add(node)

    while q:
        u = q.pop()
        for v in graph[u]:
            if v not in seen:
                q.append(v)
                seen.add(v)
```

```
for i in range(n):
    if i not in seen:
        bfs(i, seen)
        ans += 1

return ans
```

```
class UnionFind:
    def __init__(self, n: int):
        self.count = n
        self.id = list(range(n))
        self.rank = [0] * n

    def unionByRank(self, u: int, v: int) -> None:
        i = self._find(u)
        j = self._find(v)
        if i == j:

            return
        if self.rank[i] < self.rank[j]:
            self.id[i] = j
        elif self.rank[i] > self.rank[j]:
            self.id[j] = i
        else:
            self.id[i] = j
            self.rank[j] += 1
        self.count -= 1
```

```
def _find(self, u: int) -> int:
    if self.id[u] != u:
        self.id[u] = self._find(self.id[u])
    return self.id[u]
```

```
class Solution:
    def countComponents(self, n: int, edges:
list[list[int]]) -> int:
        uf = UnionFind(n)
```

```
        for u, v in edges:
            uf.unionByRank(u, v)
```

```
        return uf.count
```

• LeetCode

```
class Solution:
    def countComponents(self, n: int, edges:
List[List[int]]) -> int:
        def dfs(i: int) -> int:
            if i in vis:
                return 0
            vis.add(i)
            for j in g[i]:
                dfs(j)
            return 1
```

```
g = [[] for _ in range(n)]
for a, b in edges:
    g[a].append(b)
    g[b].append(a)
vis = set()
return sum(dfs(i) for i in range(n))
```

```
class UnionFind:
    def __init__(self, n):
        self.p = list(range(n))
        self.size = [1] * n

    def find(self, x):
        if self.p[x] != x:
            self.p[x] = self.find(self.p[x])
        return self.p[x]
```

```
    def union(self, a, b):
        pa, pb = self.find(a), self.find(b)
        if pa == pb:
            return False
        if self.size[pa] > self.size[pb]:
            self.p[pb] = pa
            self.size[pa] += self.size[pb]
        else:
            self.p[pa] = pb
            self.size[pb] += self.size[pa]
```



```
return True
```

```
class Solution:
    def countComponents(self, n: int, edges:
List[List[int]]) -> int:
        uf = UnionFind(n)
        for a, b in edges:
            n -= uf.union(a, b)
        return n
```

```
class Solution:
    def countComponents(self, n: int, edges:
List[List[int]]) -> int:
        g = [[] for _ in range(n)]
        for a, b in edges:
            g[a].append(b)
            g[b].append(a)
        vis = set()
        ans = 0
        for i in range(n):
            if i in vis:
```

```
        continue
        vis.add(i)
        q = deque([i])
        while q:
            a = q.popleft()
            for b in g[a]:
                if b not in vis:
                    vis.add(b)
                    q.append(b)

        ans += 1
```

```
return ans
```

• SimplyLeet

```
# Function to count the number of connected
components in an undirected graph.
def count_components(n, edges):
    # Build the adjacency list for the graph.
    graph = {i: [] for i in range(n)}
    for a, b in edges:
        graph[a].append(b)
        graph[b].append(a)

    # Set to keep track of visited nodes.
    visited = set()
```

```
# Helper function to perform DFS traversal
starting from a node.
```

```
def dfs(node):
    stack = [node]
    while stack:
        current = stack.pop()
        if current not in visited:
            visited.add(current)
            for neighbor in graph[current]:
                if neighbor not in visited:
```

```
                    stack.append(neighbor)
```

```
# Count connected components by initiating DFS
on unvisited nodes.
```

```
count = 0
for i in range(n):
    if i not in visited:
        dfs(i)
        count += 1
return count
```

```
# Example usage:
```

```
print(count_components(5, [[0,1],[1,2],[3,4]])) #
Expected output: 2
```

• LC.ca

Round 2 · High · Medium

p.1211

```
...
```

```
example-2: [[0, 1], [1, 2], [2, 3], [3, 4]]
```

```
for loop
```

```
[0, 1]
    0 parent set to 1
    1 parent set to 1
```

```
[1, 2]
```

```
    1 parent set to 2
    2 parent set to 2
```

```
[2, 3]
    2 parent set to 3
    3 parent set to 3
```

```
final:
    0 parent set to 1
    1 parent set to 2
```

Round 2 · High · Medium

p.1212

```

2 parent set to 3
3 parent set to 3

==> so filter for 'i == find(i)'
...

class Solution:
    def countComponents(self, n: int, edges:
List[List[int]]) -> int:
        def find(x):
            if p[x] != x:
                p[x] = find(p[x])

            return p[x]

        p = list(range(n))
        for a, b in edges:
            p[find(a)] = find(b)
        return sum(i == find(i) for i in range(n))
# or len(set(parents))

#####

class Solution: # more literal

```

Round 2 · High · Medium

p.1213

```

def countComponents(self, n: int, edges:
List[List[int]]) -> int:
    # Initialize parent array for union-find
    parent = [i for i in range(n)]

    # Perform union-find on each edge
    for u, v in edges:
        self.union(u, v, parent)

    # Count the number of distinct parents
    (connected components)
    count = len(set(self.find(x, parent) for x
in range(n)))

    return count

    def union(self, x: int, y: int, parent:
List[int]) -> None:
        rootX = self.find(x, parent)
        rootY = self.find(y, parent)
        if rootX != rootY:
            parent[rootX] = rootY

    def find(self, x: int, parent: List[int]) ->
int:
        if parent[x] != x:

            parent[x] = self.find(parent[x],
parent)
        return parent[x]

```

Round 2 · High · Medium

p.1214

◆ Quick Review

Key Insights

- Use an adjacency list to represent the graph.
- A DFS (or BFS) traversal from an unvisited node will traverse an entire connected component.
- Each new DFS/BFS call from an unvisited node signals a separate connected component.
- Alternatively, the Union-Find data structure can be used to union connected nodes and count distinct sets.

Space & Time

Time Complexity: $O(n + e)$, where n is the number of nodes and e is the number of edges.

Space Complexity: $O(n + e)$ due to the storage of the graph and tracking of visited nodes.

Solution

We solve the problem by iterating over each node and performing a DFS for any node that hasn't been visited. The DFS uses an iterative approach with a stack to traverse all nodes in a connected component. Each DFS initiation corresponds to one connected component. This method efficiently explores the graph using an adjacency list and marks nodes as visited to prevent redundant traversals.

Round 2 · High · Medium

p.1215

DFS

#417 Pacific Atlantic Water Flow

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Matrix

Lists: Grind 169

Problem

Description

There is an $m \times n$ rectangular island that borders both the Pacific Ocean and Atlantic Ocean. The Pacific Ocean touches the island's left and top edges, and the Atlantic Ocean touches the island's right and bottom edges.

The island is partitioned into a grid of square cells. You are given an $m \times n$ integer matrix `heights` where `heights[r][c]` represents the height above sea level of the cell at coordinate (r, c) .


The island receives a lot of rain, and the rain water can flow to neighboring cells directly north, south, east, and west if the neighboring cell's height is less than or equal to the current cell's height. Water can flow from any cell adjacent to an ocean into the ocean.

Return a 2D list of grid coordinates `result` where `result[i] = [ri, ci]` denotes that rain water can flow from cell (ri, ci) to both the Pacific and Atlantic oceans.

Example 1:

Round 2 · High · Medium

p.1216



Input: heights = [[1,2,2,3,5],[3,2,3,4,4],[2,4,5,3,1],[6,7,1,4,5],[5,1,1,2,4]]

Output: [[0,4],[1,3],[1,4],[2,2],[3,0],[3,1],[4,0]]

Explanation: The following cells can flow to the Pacific and Atlantic oceans, as shown below:

[0,4]: [0,4] -> Pacific Ocean

[0,4] -> Atlantic Ocean

[1,3]: [1,3] -> [0,3] -> Pacific Ocean

[1,3] -> [1,4] -> Atlantic Ocean

[1,4]: [1,4] -> [1,3] -> [0,3] -> Pacific Ocean

Example 2:

Input: heights = [[1]]

Output: [[0,0]]

Explanation: The water can flow from the only cell to the Pacific and Atlantic oceans.

Constraints:

- $m == \text{heights.length}$
- $n == \text{heights}[r].\text{length}$
- $1 \leq m, n \leq 200$
- $0 \leq \text{heights}[r][c] \leq 105$

Community Solutions (Python)

• WalkCC

class Solution:

def pacificAtlantic(self, heights: list[list[int]]) -> list[list[int]]:

DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))

m = len(heights)

n = len(heights[0])

qP = collections.deque()

qA = collections.deque()

seenP = [[False] * n for _ in range(m)]

seenA = [[False] * n for _ in range(m)]

```

for i in range(m):
    qP.append((i, 0))
    qA.append((i, n - 1))
    seenP[i][0] = True
    seenA[i][n - 1] = True

```

```

for j in range(n):
    qP.append((0, j))
    qA.append((m - 1, j))
    seenP[0][j] = True

```

```

seenA[m - 1][j] = True

```

```

def bfs(q: deque, seen: list[list[bool]]):
    while q:
        i, j = q.popleft()
        h = heights[i][j]
        for dx, dy in DIRS:
            x = i + dx
            y = j + dy
            if x < 0 or x == m or y < 0 or y == n:

```

```

        continue
        if seen[x][y] or heights[x][y] < h:
            continue
        q.append((x, y))
        seen[x][y] = True

```

```

bfs(qP, seenP)
bfs(qA, seenA)

```

```

return [[i, j] for i in range(m) for j in
range(n) if seenP[i][j] and seenA[i][j]]

```

```

class Solution:

```

```

    def pacificAtlantic(self, heights:
list[list[int]]) -> list[list[int]]:
        m = len(heights)
        n = len(heights[0])
        seenP = [[False] * n for _ in range(m)]
        seenA = [[False] * n for _ in range(m)]

```

```

        def dfs(i: int, j: int, h: int, seen:
list[list[bool]]) -> None:
            if i < 0 or i == m or j < 0 or j == n:
                return

```

```

if seen[i][j] or heights[i][j] < h:
    return

```

```

seen[i][j] = True
dfs(i + 1, j, heights[i][j], seen)
dfs(i - 1, j, heights[i][j], seen)
dfs(i, j + 1, heights[i][j], seen)
dfs(i, j - 1, heights[i][j], seen)

```

```

for i in range(m):

```

```

    dfs(i, 0, 0, seenP)
    dfs(i, n - 1, 0, seenA)

```

```

    for j in range(n):
        dfs(0, j, 0, seenP)
        dfs(m - 1, j, 0, seenA)

```

```

    return [[i, j]
            for i in range(m)
            for j in range(n)

```

```

            if seenP[i][j] and seenA[i][j]]

```

• LeetCode

```

class Solution:
    def pacificAtlantic(self, heights:
List[List[int]]) -> List[List[int]]:
        def bfs(q: Deque[Tuple[int, int]], vis:
List[List[bool]]) -> None:
            while q:
                x, y = q.popleft()
                for dx, dy in pairwise(dirs):
                    nx, ny = x + dx, y + dy
                    if (
                        <= nx < m
                            and 0 <= ny < n
                                and not vis[nx][ny]
                                and heights[nx][ny] >=
heights[x][y]
                                    ):
                                        vis[nx][ny] = True
                                        q.append((nx, ny))

m, n = len(heights), len(heights[0])
vis1 = [[False] * n for _ in range(m)]
vis2 = [[False] * n for _ in range(m)]
q1: Deque[Tuple[int, int]] = deque()

```

```
q2: Deque[Tuple[int, int]] = deque()
dirs = (-1, 0, 1, 0, -1)
```

```
for i in range(m):
    q1.append((i, 0))
    vis1[i][0] = True
    q2.append((i, n - 1))
    vis2[i][n - 1] = True
```

```
for j in range(n):
```

```
    q1.append((0, j))
    vis1[0][j] = True
    q2.append((m - 1, j))
    vis2[m - 1][j] = True
```

```
bfs(q1, vis1)
bfs(q2, vis2)
```

```
return [(i, j) for i in range(m) for j in
range(n) if vis1[i][j] and vis2[i][j]]
```

• SimplyLeet

Python solution for Pacific Atlantic Water Flow

class Solution:

```
    def pacificAtlantic(self, heights:
List[List[int]]) -> List[List[int]]:
        if not heights or not heights[0]:
            return []
```

```
        rows, cols = len(heights), len(heights[0])
        pacific_reachable = [[False for _ in
range(cols)] for _ in range(rows)]
        atlantic_reachable = [[False for _ in
range(cols)] for _ in range(rows)]
```

```
        # Directions: up, down, left, right
        directions = [(-1, 0), (1, 0), (0, -1), (0,
1)]
```

```
        # DFS function to mark reachable cells
        def dfs(r, c, reachable):
            reachable[r][c] = True
            for dr, dc in directions:
                new_r, new_c = r + dr, c + dc
                # Check bounds and conditions to
                flow water
```



```

        if (0 <= new_r < rows and 0 <=
new_c < cols and
            not reachable[new_r][new_c]
and
            heights[new_r][new_c] >=
heights[r][c]):
            dfs(new_r, new_c, reachable)

# DFS from Pacific ocean borders (top row
and left column)
for i in range(rows):
    dfs(i, 0, pacific_reachable)
for j in range(cols):
    dfs(0, j, pacific_reachable)

# DFS from Atlantic ocean borders (bottom
row and right column)
for i in range(rows):
    dfs(i, cols - 1, atlantic_reachable)
for j in range(cols):
    dfs(rows - 1, j, atlantic_reachable)

# Collect cells that can reach both oceans
result = []
for r in range(rows):

```

Round 2 · High · Medium

p.1225

```

        for c in range(cols):
            if pacific_reachable[r][c] and
atlantic_reachable[r][c]:
                result.append([r, c])
        return result

```

• LC.ca

```

class Solution:
    def pacificAtlantic(self, heights:
List[List[int]]) -> List[List[int]]:
        def bfs(q, vis):
            while q:
                for _ in range(len(q)):
                    i, j = q.popleft()
                    for a, b in [[0, -1], [0, 1],
[1, 0], [-1, 0]]:
                        x, y = i + a, j + b
                        if (
                            0 <= x < m

```

Round 2 · High · Medium

p.1226

```

                                and 0 <= y < n
                                and (x, y) not in vis
                                and heights[x][y] >=
heights[i][j]
                                ):
                                vis.add((x, y))
                                q.append((x, y))

m, n = len(heights), len(heights[0])
vis1, vis2 = set(), set()
q1 = deque()

q2 = deque()
for i in range(m):
    for j in range(n):
        if i == 0 or j == 0:
            vis1.add((i, j))
            q1.append((i, j))
        if i == m - 1 or j == n - 1:
            vis2.add((i, j))
            q2.append((i, j))
bfs(q1, vis1)

bfs(q2, vis2)
return [
    (i, j)
    for i in range(m)
    for j in range(n)
    if (i, j) in vis1 and (i, j) in vis2
]
```

Round 2 · High · Medium

p.1227

◆ Quick Review

Key Insights

- Reverse the typical DFS/BFS perspective: start from the oceans and work inward.
- Maintain two matrices (or sets) to keep track of cells reachable from the Pacific and Atlantic edges.
- A cell can be added to the result if it is reached in both traversals.
- The height constraint ensures water flows only in non-increasing order.

Space & Time

Time Complexity: $O(m * n)$ where m is the number of rows and n is the number of columns. Each cell is processed up to twice (once from each ocean).

Space Complexity: $O(m * n)$ for the visited matrices and the recursion/queue used in DFS/BFS.

Solution

The solution uses depth-first search (DFS) from the ocean boundaries. We initialize two boolean matrices, one for each ocean. For the Pacific, we start DFS from the leftmost column and top row; for the Atlantic, we start from the rightmost column and bottom row. In each DFS, if the next cell has a height equal to or greater than the current cell and hasn't been visited, we continue the search. Finally, we scan through the grid to collect cells that are reachable from both oceans. This reverse search simplifies the path

Round 2 · High · Medium

p.1228

validations and leverages the allowed water flow conditions.

Round 2 · High · Medium

p.1229

DFS

#490 The Maze

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Matrix

Lists: Premium 98

Problem

Description

There is a ball in a maze with empty spaces (represented as 0) and walls (represented as 1). The ball can go through the empty spaces by rolling up, down, left or right, but it won't stop rolling until hitting a wall. When the ball stops, it could choose the next direction.


Given the $m \times n$ maze, the ball's start position and the destination, where $\text{start} = [\text{startrow}, \text{startcol}]$ and $\text{destination} = [\text{destinationrow}, \text{destinationcol}]$, return true if the ball can stop at the destination, otherwise return false.

You may assume that the borders of the maze are all walls (see examples).

Example 1:

Round 2 · High · Medium

p.1230




Input: maze = `[[0,0,1,0,0],[0,0,0,0,0],[0,0,0,1,0],[1,1,0,1,1],[0,0,0,0,0]]`, start = `[0,4]`, destination = `[4,4]`

Output: true

Explanation: One possible way is : left -> down -> left -> down -> right -> down -> right.

Example 2:



Input: maze = `[[0,0,1,0,0],[0,0,0,0,0],[0,0,0,1,0],[1,1,0,1,1],[0,0,0,0,0]]`, start = `[0,4]`, destination = `[3,2]`

Output: false

Explanation: There is no way for the ball to stop at the destination. Notice that you can pass through the destination but you cannot stop there.

Example 3:

Input: maze = [[0,0,0,0,0],[1,1,0,0,1],[0,0,0,0,0],[0,1,0,0,1],[0,1,0,0,0]], start = [4,3], destination = [0,1]
 Output: false

Constraints:

- $m == \text{maze.length}$
- $n == \text{maze[i].length}$
- $1 \leq m, n \leq 100$
- maze[i][j] is 0 or 1.
- $\text{start.length} == 2$
- $\text{destination.length} == 2$
- $0 \leq \text{startrow}, \text{destinationrow} < m$
- $0 \leq \text{startcol}, \text{destinationcol} < n$
- Both the ball and the destination exist in an empty space, and they will not be in the same position initially.
- The maze contains at least 2 empty spaces.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def hasPath(
        self,
        maze: list[list[int]],
        start: list[int],
        destination: list[int],
    ) -> bool:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(maze)
        n = len(maze[0])
```

```
        q = collections.deque([(start[0], start[1])])
        seen = {(start[0], start[1])}

        def isValid(x: int, y: int) -> bool:
            return 0 <= x < m and 0 <= y < n and
                maze[x][y] == 0

        while q:
            i, j = q.popleft()
            for dx, dy in DIRS:
                x = i
```

```

y = j
while isValid(x + dx, y + dy):
    x += dx
    y += dy
    if [x, y] == destination:
        return True
    if (x, y) in seen:
        continue
    q.append((x, y))
    seen.add((x, y))

```

```

return False

```

• LeetCode

```

class Solution:
    def hasPath(
        self, maze: List[List[int]], start:
List[int], destination: List[int]
    ) -> bool:
        def dfs(i, j):
            if vis[i][j]:
                return
            vis[i][j] = True
            if [i, j] == destination:
                return

```

Round 2 · High · Medium

p.1235

```

        for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:
            x, y = i, j
            while 0 <= x + a < m and 0 <= y + b
< n and maze[x + a][y + b] == 0:
                x, y = x + a, y + b
                dfs(x, y)

        m, n = len(maze), len(maze[0])
        vis = [[False] * n for _ in range(m)]
        dfs(start[0], start[1])
        return vis[destination[0]][destination[1]]

```

• SimplyLeet

```

# Function to determine if there is a valid path in
the maze from start to destination
def hasPath(maze, start, destination):
    from collections import deque
    rows, cols = len(maze), len(maze[0])
    # Create a visited 2D array to keep track of
processed cells
    visited = [[False] * cols for _ in range(rows)]
    queue = deque([tuple(start)])
    visited[start[0]][start[1]] = True

    # Directions: up, down, left, right

```

Round 2 · High · Medium

p.1236

```

directions = [(-1, 0), (1, 0), (0, -1), (0, 1)]

# Process nodes in BFS fashion
while queue:
    current = queue.popleft()
    # Check if the current position is the
destination
    if [current[0], current[1]] ==
destination:
        return True
    # Explore all four possible directions
    for dx, dy in directions:

```

```

        x, y = current
        # Roll the ball until it hits a wall or
boundary
        while 0 <= x + dx < rows and 0 <= y +
dy < cols and maze[x + dx][y + dy] == 0:
            x += dx
            y += dy
        # If this new stop has not been
visited, add it to the queue
        if not visited[x][y]:
            visited[x][y] = True
            queue.append((x, y))
    return False

```

```

# Example usage:
maze = [[0,0,1,0,0],
        [0,0,0,0,0],
        [0,0,0,1,0],
        [1,1,0,1,1],
        [0,0,0,0,0]]
start = [0,4]
destination = [4,4]
print(hasPath(maze, start, destination)) #
Expected output: True

```

• LC.ca

```

class Solution:
    def hasPath(
        self, maze: List[List[int]], start:
List[int], destination: List[int]
    ) -> bool:
        m, n = len(maze), len(maze[0])
        q = deque([start])
        rs, cs = start
        vis = {(rs, cs)}
        while q:
            i, j = q.popleft()

```

```

        for a, b in [[0, -1], [0, 1], [-1, 0],
[1, 0]]:
            x, y = i, j
            while 0 <= x + a < m and 0 <= y + b
< n and maze[x + a][y + b] == 0:
                x, y = x + a, y + b
                if [x, y] == destination:
                    return True
                if (x, y) not in vis:
                    vis.add((x, y))
                    q.append((x, y))
        return False

```

◆ Quick Review

Key Insights

- The ball continues moving in one direction until it hits a wall (or the boundary, which is considered a wall).
- You must simulate the ball's rolling rather than taking a single step.
- Even if the ball passes through the destination, it is only valid if it can stop exactly there.
- A breadth-first search (BFS) or depth-first search (DFS) strategy is ideal because we need to explore all stopping positions.
- Use a visited structure to avoid revisiting positions and thus prevent infinite loops.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the maze.

Space Complexity: $O(m * n)$ due to the storage used by the visited matrix and the BFS/DFS queue or recursion stack.

Solution

We simulate the ball's movement using a breadth-first search (BFS) approach. Each time we pick a stop position, we try to roll in all four directions until the ball hits a wall. At that point, if the new stopping position has not been visited, we mark it as visited and add it to our queue. If we ever reach the destination exactly, we return true. The key is simulating the "roll" by iteratively moving in a direction until no further move is possible.

We store the coordinates in a queue (or stack for DFS) along with a 2D visited array to monitor which positions have been processed. The directions are maintained in a list/array, and for each direction, we keep moving until the next cell would be out of bounds or a wall.

DFS

#694 Number of Distinct Islands

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · DFS · BFS · Union Find · Hash Function

Lists: Premium 98

Problem

Description

You are given an $m \times n$ binary matrix grid. An island is a group of 1's (representing land) connected 4-directionally (horizontal or vertical.) You may assume all four edges of the grid are surrounded by water.

An island is considered to be the same as another if and only if one island can be translated (and not rotated or reflected) to equal the other.

Return the number of distinct islands.

Example 1:

1	1	0	0	0
1	1	0	0	0
0	0	0	1	1
0	0	0	1	1

Input: grid =

[[1,1,0,0,0],[1,1,0,0,0],[0,0,0,1,1],[0,0,0,1,1]]

Output: 1

Example 2:

1	1	0	1	1
1	0	0	0	0
0	0	0	0	1
1	1	0	1	1

Input: grid =
[[1,1,0,1,1],[1,0,0,0,0],[0,0,0,0,1],[1,1,0,1,1]]
Output: 3

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$
- $1 \leq m, n \leq 50$
- $\text{grid}[i][j]$ is either 0 or 1.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numDistinctIslands(self, grid:
list[list[int]]) -> int:
        seen = set()

        def dfs(i: int, j: int, i0: int, j0: int):
            if i < 0 or i == len(grid) or j < 0 or j ==
len(grid[0]):
                return
            if grid[i][j] == 0 or (i, j) in seen:
                return

            seen.add((i, j))
            island.append((i - i0, j - j0))
            dfs(i + 1, j, i0, j0)
            dfs(i - 1, j, i0, j0)
            dfs(i, j + 1, i0, j0)
            dfs(i, j - 1, i0, j0)

        islands = set() # all the different islands

        for i in range(len(grid)):
```

```

for j in range(len(grid[0])):
    island = []
    dfs(i, j, i, j)
    if island:
        islands.add(frozenset(island))

return len(islands)

```

• LeetCode

```

class Solution:
    def numDistinctIslands(self, grid:
List[List[int]]) -> int:
        def dfs(i: int, j: int, k: int):
            grid[i][j] = 0
            path.append(str(k))
            dirs = (-1, 0, 1, 0, -1)
            for h in range(1, 5):
                x, y = i + dirs[h - 1], j + dirs[h]
                if 0 <= x < m and 0 <= y < n and
grid[x][y]:
                    dfs(x, y, h)

```

```

path.append(str(-k))

```

```

paths = set()
path = []
m, n = len(grid), len(grid[0])
for i, row in enumerate(grid):
    for j, x in enumerate(row):
        if x:
            dfs(i, j, 0)
            paths.add("".join(path))

```

```

path.clear()
return len(paths)

```

• SimplyLeet

```

def numDistinctIslands(grid):
    # Get the number of rows and columns
    rows, cols = len(grid), len(grid[0])
    # Set to store unique island shapes
    distinct_shapes = set()

    # DFS function to explore the island and record
    its relative shape
    def dfs(r, c, origin_r, origin_c, shape):
        # Check boundaries and if the cell is
        unvisited land
        if r < 0 or c < 0 or r >= rows or c >= cols
        or grid[r][c] == 0:

```

```

        return
    # Mark the cell as visited by setting it to
    0
    grid[r][c] = 0
    # Record the relative position from the
    origin of this island
    shape.append((r - origin_r, c - origin_c))
    # Explore all 4-directionally adjacent
    cells
    dfs(r + 1, c, origin_r, origin_c, shape)
    dfs(r - 1, c, origin_r, origin_c, shape)
    dfs(r, c + 1, origin_r, origin_c, shape)
    dfs(r, c - 1, origin_r, origin_c, shape)

```

```

# Iterate over every cell in the grid
for i in range(rows):
    for j in range(cols):
        if grid[i][j] == 1:
            shape = []
            dfs(i, j, i, j, shape)
            # Convert the shape list to a tuple
            of sorted coordinates for canonical form
            distinct_shapes.add(tuple(sorted(s
            hape)))

```

```

# The number of distinct islands is the size of
the set
return len(distinct_shapes)

# Example usage:
grid = [[1,1,0,0,0],
        [1,1,0,0,0],
        [0,0,0,1,1],
        [0,0,0,1,1]]
print(numDistinctIslands(grid)) # Output: 1

```

• LC.ca

```

class Solution:
    def numDistinctIslands(self, grid:
    List[List[int]]) -> int:
        def dfs(i: int, j: int, k: int):
            grid[i][j] = 0
            path.append(str(k))
            dirs = (-1, 0, 1, 0, -1)
            for h in range(1, 5):
                x, y = i + dirs[h - 1], j + dirs[h]
                if 0 <= x < m and 0 <= y < n and
            grid[x][y]:
                dfs(x, y, h)

```

```

        path.append(str(-k))

    paths = set()
    path = []
    m, n = len(grid), len(grid[0])
    for i, row in enumerate(grid):
        for j, x in enumerate(row):
            if x:
                dfs(i, j, 0)
                paths.add("".join(path))

    path.clear()
    return len(paths)

```

◆ Quick Review

Key Insights

- Use DFS/BFS to traverse and mark each island.
- For each island, record its shape relative to a starting coordinate to allow translation invariance.
- Represent the island's shape as a unique signature (e.g., tuple or string of relative positions).
- Use a set (or hash table) to store unique shapes.

Space & Time

Time Complexity: $O(m * n)$ as each cell is visited once.

Space Complexity: $O(m * n)$ in the worst-case due to recursion stack and storing island shapes.

Solution

Use depth-first search (DFS) to explore the grid. When an unvisited land cell (1) is encountered, initiate a DFS to traverse the entire island. Record each land cell's coordinates relative to the starting position of the island. This relative representation is invariant under translation, making it possible to compare island shapes. Store the canonical shape signature (for example, as a tuple or string) in a set to ensure uniqueness. Finally, return the size of the set as the number of distinct islands.

DFS

#695 Max Area of Island

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Union Find · Matrix

Lists: Denny Zhang

Problem

Description

You are given an $m \times n$ binary matrix grid. An island is a group of 1's (representing land) connected 4-directionally (horizontal or vertical.) You may assume all four edges of the grid are surrounded by water.

The area of an island is the number of cells with a value 1 in the island.

Return the maximum area of an island in grid. If there is no island, return 0.

Example 1:

0	0	1	0	0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	1	1	1	0	0	0
0	1	1	0	1	0	0	0	0	0	0	0	0
0	1	0	0	1	1	0	0	1	0	1	0	0
0	1	0	0	1	1	0	0	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0

Input: grid = [[0,0,1,0,0,0,0,1,0,0,0,0,0],[0,0,0,0,0,0,0,0,1,1,1,0,0],[0,1,1,0,1,0,0,0,0,0,0,0,0],[0,1,0,0,1,1,0,0,1,0,1,0,0],[0,1,0,0,1,1,0,0,1,1,1,0,0],[0,0,0,0,0,0,0,0,0,0,0,1,0],[0,0,0,0,0,0,0,0,1,1,1,0,0],[0,0,0,0,0,0,0,0,1,1,0,0,0]]

Output: 6

Explanation: The answer is not 11, because the island must be connected 4-directionally.

Example 2:

Input: grid = [[0,0,0,0,0,0,0,0,0]]

Output: 0

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$

- $1 \leq m, n \leq 50$
- `grid[i][j]` is either 0 or 1.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxAreaOfIsland(self, grid: list[list[int]])
    -> int:
        def dfs(i: int, j: int) -> int:
            if i < 0 or i == len(grid) or j < 0 or j ==
            len(grid[0]):
                return 0
            if grid[i][j] != 1:
                return 0

            grid[i][j] = 2

            return 1 + dfs(i + 1, j) + dfs(i - 1, j) +
            dfs(i, j + 1) + dfs(i, j - 1)

        return max(dfs(i, j) for i in range(len(grid))
        for j in range(len(grid[0])))
```

• LeetCode

```
class Solution:
    def maxAreaOfIsland(self, grid:
    List[List[int]]) -> int:
        def dfs(i: int, j: int) -> int:
            if grid[i][j] == 0:
                return 0
            ans = 1
            grid[i][j] = 0
            dirs = (-1, 0, 1, 0, -1)
            for a, b in pairwise(dirs):
                x, y = i + a, j + b

                if 0 <= x < m and 0 <= y < n:
                    ans += dfs(x, y)
            return ans

        m, n = len(grid), len(grid[0])
        return max(dfs(i, j) for i in range(m) for
        j in range(n))
```

• SimplyLeet

```

# Function to compute the maximum area of an island
in a grid using DFS
def maxAreaOfIsland(grid):
    # Determine the dimensions of the grid
    rows = len(grid)
    cols = len(grid[0]) if rows else 0

    # Helper function for performing DFS on the
    grid
    def dfs(r, c):
        # Check if the current cell is
        out-of-bounds or is water
        if r < 0 or c < 0 or r >= rows or c >= cols
        or grid[r][c] == 0:

            return 0
        # Mark the current cell as visited by
        setting it to 0 (water)
        grid[r][c] = 0
        area = 1 # Count the current cell as part
        of the island
        # Recursively explore all four directions
        area += dfs(r + 1, c)
        area += dfs(r - 1, c)
        area += dfs(r, c + 1)
        area += dfs(r, c - 1)
        return area

```

```

max_area = 0
# Iterate through each cell in the grid
for r in range(rows):
    for c in range(cols):
        if grid[r][c] == 1:
            # Update max_area with the area
            obtained from the DFS
            max_area = max(max_area, dfs(r,
            c))
    return max_area

```

• LC.ca

```

class Solution:
    def maxAreaOfIsland(self, grid:
    List[List[int]]) -> int:
        def dfs(i: int, j: int) -> int:
            if grid[i][j] == 0:
                return 0
            ans = 1
            grid[i][j] = 0
            dirs = (-1, 0, 1, 0, -1)
            for a, b in pairwise(dirs):
                x, y = i + a, j + b

```



```

        if 0 <= x < m and 0 <= y < n:
            ans += dfs(x, y)
        return ans

    m, n = len(grid), len(grid[0])
    return max(dfs(i, j) for i in range(m) for
j in range(n))

```

◆ Quick Review

Key Insights

- Traverse each cell in the grid.
- Use Depth-First Search (DFS) or Breadth-First Search (BFS) to explore connected 1's.
- Mark visited cells to prevent double counting (e.g., by setting them to 0).
- Maintain and update a maximum area count during traversal.
- The grid's boundaries are naturally water, so no extra edge processing is needed.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the grid since every cell is visited at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario for the recursion stack (or auxiliary data structures) when the grid is entirely land.

Solution

The solution leverages a DFS approach to iterate over each cell in the given grid. When a cell with a value of 1 is encountered, a DFS is initiated from that cell to calculate the area of the island connected to it. Cells are marked as visited (by changing them from 1 to 0) during the DFS to avoid revisiting. The DFS explores all four directions (up, down, left, right) and accumulates the count of connected cells. The maximum area found among all islands is returned at the end. This approach efficiently computes the largest island area with a simple traversal mechanism.

DFS

#721 Accounts Merge

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · DFS · BFS · Union Find

Lists: Grind 169

Problem

Description

Given a list of accounts where each element `accounts[i]` is a list of strings, where the first element `accounts[i][0]` is a name, and the rest of the elements are emails representing emails of the account.

Now, we would like to merge these accounts. Two accounts definitely belong to the same person if there is some common email to both accounts. Note that even if two accounts have the same name, they may belong to different people as people could have the same name. A person can have any number of accounts initially, but all of their accounts definitely have the same name.

After merging the accounts, return the accounts in the following format: the first element of each account is the name, and the rest of the elements are emails in sorted order. The accounts themselves can be returned in any order.

Example 1:

Input: `accounts = [{"John","johnsmith@mail.com","john_newyork@mail.com"}, {"John","johnsmith@mail.com","john00@mail.com"}, {"Mary","mary@mail.com"}, {"John","johnnybravo@mail.com"}]`

Output: `[{"John","john00@mail.com","john_newyork@mail.com","johnsmith@mail.com"}, {"Mary","mary@mail.com"}, {"John","johnnybravo@mail.com"}]`

Explanation:

The first and second John's are the same person as they have the common email "johnsmith@mail.com".

The third John and Mary are different people as none of their email addresses are used by other accounts.

We could return these lists in any order, for example the answer `[['Mary', 'mary@mail.com'], ['John', 'johnnybravo@mail.com'], ['John', 'john00@mail.com', 'john_newyork@mail.com', 'johnsmith@mail.com']]` would still be accepted.

Example 2:

Input: `accounts = [{"Gabe","Gabe0@m.co","Gabe3@m.co","Gabe1@m.co"}, {"Kevin","Kevin3@m.co","Kevin5@m.co","Kevin0@m.co"}, {"Ethan","Ethan5@m.co","Ethan4@m.co","Ethan0@m.co"}, {"Hanzo","Hanzo3@m.co","Hanzo1@m.co","Hanzo0@m.co"}, {"Fern","Fern5@m.co","Fern1@m.co","Fern0@m.co"}]`

Output: `[{"Ethan","Ethan0@m.co","Ethan4@m.co","Ethan5@m.co"}, {"Gabe","Gabe0@m.co","Gabe1@m.co","Gabe3@m.co"}, {"Hanzo","Hanzo0@m.co","Hanzo1@m.co","Hanzo3@m.co"}, {"Kevin","Kevin0@m.co","Kevin3@m.co","Kevin5@m.co"}, {"Fern","Fern0@m.co","Fern1@m.co","Fern5@m.co"}]`

Constraints:

- `1 <= accounts.length <= 1000`

Round 2 · High · Medium

p.1259

Round 2 · High · Medium

p.1260

- `2 <= accounts[i].length <= 10`
- `1 <= accounts[i][j].length <= 30`
- `accounts[i][0]` consists of English letters.
- `accounts[i][j]` (for `j > 0`) is a valid email.

Community Solutions (Python)

• LeetCode

```
class UnionFind:
    def __init__(self, n):
        self.p = list(range(n))
        self.size = [1] * n

    def find(self, x):
        if self.p[x] != x:
            self.p[x] = self.find(self.p[x])
        return self.p[x]

    def union(self, a, b):
        pa, pb = self.find(a), self.find(b)
        if pa == pb:
            return False
        if self.size[pa] > self.size[pb]:
            self.p[pb] = pa
            self.size[pa] += self.size[pb]
        else:
            self.p[pa] = pb
            self.size[pb] += self.size[pa]
```

return True

```
class Solution:
    def accountsMerge(self, accounts:
List[List[str]]) -> List[List[str]]:
        uf = UnionFind(len(accounts))
        d = {}
        for i, (_, *emails) in
enumerate(accounts):
            for email in emails:
                if email in d:
                    uf.union(i, d[email])
                else:
                    d[email] = i
        g = defaultdict(set)
        for i, (_, *emails) in
enumerate(accounts):
            root = uf.find(i)
            g[root].update(emails)
        return [[accounts[root][0]] +
sorted(emails) for root, emails in g.items()]
```

• SimplyLeet

Python solution for Accounts Merge

Define the UnionFind class for disjoint set operations

class UnionFind:

```
def __init__(self):
    self.parent = {}
```

Find with path compression

```
def find(self, x):
    if self.parent.setdefault(x, x) != x:
```

```
        self.parent[x] =
self.find(self.parent[x])
    return self.parent[x]
```

Union two nodes

```
def union(self, x, y):
    self.parent[self.find(x)] = self.find(y)
```

class Solution:

```
def accountsMerge(self, accounts):
    uf = UnionFind()
```

```
email_to_name = {}
```

Iterate over each account and union

emails

```
for account in accounts:
    name = account[0]
    first_email = account[1]
    for email in account[1:]:
        uf.union(first_email, email)
        email_to_name[email] = name
```

Group emails by their root parent

```
groups = {}
for email in email_to_name:
    parent = uf.find(email)
    groups.setdefault(parent,
[[]]).append(email)
```

Return the merged account list with sorted emails

```
result = []
for parent, emails in groups.items():
    result.append([email_to_name[parent]]
+ sorted(emails))
```

```
return result
```

```
# Example usage:
# accounts = [["John","johnsmith@mail.com","john_n
ewyork@mail.com"],
# ["John","johnsmith@mail.com","john00@mail.com"],
# ["Mary","mary@mail.com"],
# ["John","johnnybravo@mail.com"]]
# solution = Solution()
# print(solution.accountsMerge(accounts))
```

• LC.ca

```
class Solution:
    def accountsMerge(self, accounts:
List[List[str]]) -> List[List[str]]:
        def find(x):
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        n = len(accounts)
        p = list(range(n))
        email_id = {}
```

```
for i, account in enumerate(accounts):
    name = account[0]
    for email in account[1:]:
        if email in email_id:
            p[find(i)] =
find(email_id[email])
        else:
            email_id[email] = i
mp = defaultdict(set)
for i, account in enumerate(accounts):
    for email in account[1:]:
```

```
mp[find(i)].add(email)
```

```
ans = []
for i, emails in mp.items():
    t = [accounts[i][0]]
    t.extend(sorted(emails))
    ans.append(t)
return ans
```

◆ Quick Review

Key Insights

- Two accounts are merged if they share at least one common email.
- Treat each email as a node in a graph; accounts represent connected components.
- Both DFS/BFS and Union-Find can be used to find connected components.

- Keep a mapping from email to name because all emails in a connected component belong to the same person.
- After grouping, sort the emails and prepend the name to form the final account list.

Space & Time

Time Complexity: $O(N * M * \log(M))$ where N is the number of accounts and M is the number of emails per account (owing to sorting and union operations).

Space Complexity: $O(N * M)$, used for storing email mappings and union-find structures.

Solution

The solution involves using a Union-Find (Disjoint Set Union) data structure to merge emails that are connected by being on the same account. For each account, we union the first email with every other email in that account. We also maintain a mapping from email to owner name. After processing all accounts, we group the emails by their root parent derived from the union-find structure. For each group, we sort the emails and prepend the owner's name. Finally, we return the merged list of accounts.

DFS

#785 Is Graph Bipartite?

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Union Find · Graph

Lists: Denny Zhang

Problem

Description


There is an undirected graph with n nodes, where each node is numbered between 0 and $n - 1$. You are given a 2D array `graph`, where `graph[u]` is an array of nodes that node u is adjacent to. More formally, for each v in `graph[u]`, there is an undirected edge between node u and node v . The graph has the following properties:

- There are no self-edges (`graph[u]` does not contain u).
- There are no parallel edges (`graph[u]` does not contain duplicate values).
- If v is in `graph[u]`, then u is in `graph[v]` (the graph is undirected).
- The graph may not be connected, meaning there may be two nodes u and v such that there is no path between them.

A graph is bipartite if the nodes can be partitioned into two independent sets A and B such that every edge in the graph connects a node in set A and a node in set B .


Return true if and only if it is bipartite.

Example 1:



Input: graph = [[1,2,3],[0,2],[0,1,3],[0,2]]
 Output: false
 Explanation: There is no way to partition the nodes into two independent sets such that every edge connects a node in one and a node in the other.

Example 2:



Input: graph = [[1,3],[0,2],[1,3],[0,2]]
 Output: true
 Explanation: We can partition the nodes into two sets: {0, 2} and {1, 3}.

Constraints:

- graph.length == n
- 1 <= n <= 100
- 0 <= graph[u].length < n
- 0 <= graph[u][i] <= n - 1
- graph[u] does not contain u.

- All the values of `graph[u]` are unique.
- If `graph[u]` contains `v`, then `graph[v]` contains `u`.

Community Solutions (Python)

• WalkCC

```
from enum import Enum

class Color(Enum):
    WHITE = 0
    RED = 1
    GREEN = 2

class Solution:

    def isBipartite(self, graph: list[list[int]]) -> bool:
        colors = [Color.WHITE] * len(graph)

        for i in range(len(graph)):
            # This node has been colored, so do nothing.
            if colors[i] != Color.WHITE:
                continue
            # Always paint red for a white node.
            colors[i] = Color.RED
            # BFS.
```

Round 2 · High · Medium

p.1271

```
q = collections.deque([i])
while q:
    for _ in range(len(q)):
        u = q.popleft()
        for v in graph[u]:
            if colors[v] == colors[u]:
                return False
            if colors[v] == Color.WHITE:
                colors[v] = Color.RED if colors[u]
                == Color.GREEN else Color.GREEN
                q.append(v)

return True
```

• LeetCode

```
class Solution:
    def isBipartite(self, graph: List[List[int]])
    -> bool:
        def dfs(a: int, c: int) -> bool:
            color[a] = c
            for b in graph[a]:
                if color[b] == c or (color[b] == 0
                and not dfs(b, -c)):
                    return False
            return True

        n = len(graph)
```

Round 2 · High · Medium

p.1272


```
color = [0] * n
for i in range(n):
    if color[i] == 0 and not dfs(i, 1):
        return False
return True
```

```
class Solution:
    def isBipartite(self, graph: List[List[int]])
    -> bool:
        def find(x: int) -> int:
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        p = list(range(len(graph)))
        for a, bs in enumerate(graph):
            for b in bs:
```

```
                pa, pb = find(a), find(b)
                if pa == pb:
                    return False
                p[pb] = find(bs[0])
        return True
```

• SimplyLeet

Round 2 · High · Medium

p.1273

```
# Using Breadth-First Search (BFS)
from collections import deque

def isBipartite(graph):
    n = len(graph)
    # color array: -1 means uncolored, 0 and 1 are
    the two colors.
    colors = [-1] * n

    # Iterate through each node to handle
    disconnected components
    for i in range(n):

        if colors[i] == -1:
            # Start BFS from node i with color 0
            queue = deque([i])
            colors[i] = 0
            while queue:
                node = queue.popleft()
                # Check each neighbor of the
                current node
                for neighbor in graph[node]:
                    if colors[neighbor] == -1:
                        # Assign the opposite
                        color to the neighbor
```

Round 2 · High · Medium

p.1274

```

        colors[neighbor] = 1 -
colors[node]
        queue.append(neighbor)
    elif colors[neighbor] ==
colors[node]:
        # Found a neighbor with
the same color: not bipartite
        return False
    return True

```

Example usage:

```

if __name__ == '__main__':
    graph_example = [[1,3],[0,2],[1,3],[0,2]]

```

```

    print(isBipartite(graph_example)) # Expected
output: True

```

• LC.ca

```

class Solution:
    def isBipartite(self, graph: List[List[int]])
-> bool:
        def dfs(u, c):
            color[u] = c
            for v in graph[u]:
                if not color[v]:
                    if not dfs(v, 3 - c):
                        return False
                elif color[v] == c:
                    return False

```

Round 2 · High · Medium

p.1275

```

        return True

    n = len(graph)
    color = [0] * n
    for i in range(n):
        if not color[i] and not dfs(i, 1):
            return False
    return True

```

◆ Quick Review

Key Insights

- A graph is bipartite if you can assign two colors to each node such that no two adjacent nodes share the same color.
- The challenge may require exploring disconnected components.
- BFS or DFS can be used to attempt the two-coloring, and a conflict during coloring indicates the graph is not bipartite.
- Union Find is an alternative approach, but BFS/DFS is more intuitive for graph coloring problems.

Space & Time

Time Complexity: $O(V + E)$ where V is the number of nodes and E is the number of edges, since every node and edge is traversed at most once.

Space Complexity: $O(V)$, for the color array and the queue (or recursion stack) used in BFS (or DFS).

Solution

Round 2 · High · Medium

p.1276

We use a BFS/DFS approach to assign colors (0 and 1) to nodes. A color array, initially set to -1 for all nodes, is maintained. For each unvisited node (color -1), we start a BFS (or DFS) and assign it an initial color (e.g., 0). For every visited neighbor, if it has not been colored, assign it the opposite color; if it is already colored with the same color as the current node, a conflict has occurred, and the graph is not bipartite. This method gracefully handles disconnected components by iterating over all nodes.

Round 2 · High · Medium

p.1277

DFS

#1236 Web Crawler

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

String · DFS · BFS · Interactive

Lists: Premium 98

Problem

Description

Given a url `startUrl` and an interface `HtmlParser`, implement a web crawler to crawl all links that are under the same hostname as `startUrl`.

Return all urls obtained by your web crawler in any order.

Your crawler should:

- Start from the page: `startUrl`
- Call `HtmlParser.getUrls(url)` to get all urls from a webpage of given url.
- Do not crawl the same link twice.
- Explore only the links that are under the same hostname as `startUrl`.

Round 2 · High · Medium

p.1278

http://example.org:8888/foo/bar#bang

hostname

host

As shown in the example url above, the hostname is example.org. For simplicity sake, you may assume all urls use http protocol without any port specified. For example, the urls `http://leetcode.com/problems` and `http://leetcode.com/contest` are under the same hostname, while urls `http://example.org/test` and `http://example.com/abc` are not under the same hostname.

The HtmlParser interface is defined as such:


```
interface HtmlParser {
    // Return a list of all urls from a webpage of given url.
    public List<String> getUrls(String url);
}
```

Below are two examples explaining the functionality of the problem, for custom testing purposes you'll have three variables `urls`, `edges` and `startUrl`. Notice that you will only have access to `startUrl` in your code, while `urls` and `edges` are not directly accessible to you in code.

Note: Consider the same URL with the trailing slash "/" as a different URL. For example,


"`http://news.yahoo.com`", and "`http://news.yahoo.com/`" are different urls.

Example 1:



```
Input:
urls = [
    "http://news.yahoo.com",
    "http://news.yahoo.com/news",
    "http://news.yahoo.com/news/topics/",
    "http://news.google.com",
    "http://news.yahoo.com/us"
]
```

Example 2:



```

Input:
urls = [
    "http://news.yahoo.com",
    "http://news.yahoo.com/news",
    "http://news.yahoo.com/news/topics/",
    "http://news.google.com"
]
edges = [[0,2],[2,1],[3,2],[3,1],[3,0]]

```

Constraints:

- $1 \leq \text{urls.length} \leq 1000$
- $1 \leq \text{urls}[i].\text{length} \leq 300$
- startUrl is one of the urls.
- Hostname label must be from 1 to 63 characters long, including the dots, may contain only the ASCII letters

from 'a' to 'z', digits from '0' to '9' and the hyphen-minus character ('-').

- The hostname may not start or end with the hyphen-minus character ('-').
- See: https://en.wikipedia.org/wiki/Hostname#Restrictions_on_valid_hostnames
- You may assume there're no duplicates in url library.

Community Solutions (Python)

• WalkCC

```

# """
# This is HtmlParser's API interface.
# You should not implement it, or speculate about
# its implementation
# """
# class HtmlParser(object):
#     def getUrls(self, url: str) -> list[str]:

class Solution:
    def crawl(self, startUrl: str, htmlParser:
'HtmlParser') -> list[str]:
        q = collections.deque([startUrl])

```

```

seen = {startUrl}
hostname = startUrl.split('/')[2]

while q:
    currUrl = q.popleft()
    for url in htmlParser.getUrls(currUrl):
        if url in seen:
            continue
        if hostname in url:
            q.append(url)

    seen.add(url)

return seen

```

• LeetCode

```

# """
# This is HtmlParser's API interface.
# You should not implement it, or speculate about
# its implementation
# """
# class HtmlParser(object):
# def getUrls(self, url):
# """
# :type url: str
# :rtype List[str]
# """

```

```

class Solution:
    def crawl(self, startUrl: str, htmlParser:
'HtmlParser') -> List[str]:
        def host(url):
            url = url[7:]
            return url.split('/')[0]

        def dfs(url):
            if url in ans:
                return
            ans.add(url)
            for next in htmlParser.getUrls(url):
                if host(url) == host(next):
                    dfs(next)

        ans = set()
        dfs(startUrl)
        return list(ans)

```

• SimplyLeet

Python Code Solution

```

class Solution:
    def crawl(self, startUrl: str, htmlParser:
'HtmlParser') -> list[str]:
        from collections import deque

        # Helper function to extract hostname from
a url string.
        def get_hostname(url: str) -> str:
            # Remove protocol prefix "http://"
            hostname = url.split("://")[1]

            # The hostname is the part before the
first "/" (if exists)
            return hostname.split("/")[0]

        # Extract the hostname from the starting
URL.
        start_hostname = get_hostname(startUrl)

        # Initialize a set to keep track of visited
urls and a queue for BFS.
        visited = set([startUrl])
        queue = deque([startUrl])

```

```

        # Perform BFS traversal of the web links.
        while queue:
            current_url = queue.popleft()
            # Retrieve all URLs found at the
current URL.
            for url in
htmlParser.getUrls(current_url):
                # Check if URL is under the same
hostname and has not been visited.
                if url not in visited and
get_hostname(url) == start_hostname:
                    visited.add(url)
                    queue.append(url)

        # Return all visited URLs as a list.
        return list(visited)

```

• LC.ca

```

# """
# This is HtmlParser's API interface.
# You should not implement it, or speculate about
its implementation
# """
# class HtmlParser(object):
# def getUrls(self, url):
# """
# :type url: str
# :rtype List[str]
# """

```

```

from collections import deque

class Solution:
    def crawl(self, startUrl: str, htmlParser:
'HtmlParser') -> List[str]:
        visited = set()
        domain = self.get_domain(startUrl)
        queue = deque([startUrl])
        visited.add(startUrl)

        while queue:
            current_url = queue.popleft()
            for url in
htmlParser.getUrls(current_url):
                if url not in visited and
self.get_domain(url) == domain:
                    visited.add(url)
                    queue.append(url)

        return list(visited)

    def get_domain(self, url: str) -> str:

        return url.split('/')[2]

```

◆ Quick Review

Key Insights

- Treat the set of webpages as a graph where each URL is a node and an edge exists from one URL to another if it is returned by `HtmlParser.getUrls(url)`.
- Use either breadth-first search (BFS) or depth-first search (DFS) to traverse the webpages.
- Extract the hostname from a URL to ensure that only URLs on the same domain are crawled.
- Maintain a visited set (or equivalent) to avoid processing the same URL multiple times.

Space & Time

Time Complexity: $O(N + E)$ where N is the number of nodes (URLs) and E is the number of edges (links between URLs) encountered during the crawl.

Space Complexity: $O(N)$ for storing the visited URLs and the queue/stack used for traversal.

Solution

The solution uses a BFS (or DFS) approach to traverse the web graph:

- Extract the hostname from `startUrl` (for instance, by splitting on `"/"` and then the first `"/"` that follows).
- Initialize a visited set to keep track of URLs that have been crawled.
- Use a queue to implement the BFS. Enqueue the `startUrl`.
- While the queue is not empty, dequeue a URL and retrieve its outgoing links using `HtmlParser.getUrls(url)`.
- For each URL obtained: Extract its hostname and compare it with the hostname of `startUrl`. If it matches and it has not been visited before, add it to the visited set and enqueue it.
- Finally, return the set of visited URLs.

This approach guarantees that each URL is processed only once and that only URLs from the same hostname are crawled.

Round 2 · High · Medium

p.1289

DFS

#1242 Web Crawler Multithreaded

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Concurrency

Lists: Denny Zhang

Problem

Description

Given a URL `startUrl` and an interface `HtmlParser`, implement a Multi-threaded web crawler to crawl all links that are under the same hostname as `startUrl`. Return all URLs obtained by your web crawler in any order.

Your crawler should:

- Start from the page: `startUrl`
- Call `HtmlParser.getUrls(url)` to get all URLs from a webpage of a given URL.
- Do not crawl the same link twice.
- Explore only the links that are under the same hostname as `startUrl`.

Round 2 · High · Medium

p.1290

http://example.org:8888/foo/bar#bang

hostname
host

As shown in the example URL above, the hostname is example.org. For simplicity's sake, you may assume all URLs use HTTP protocol without any port specified. For example, the URLs `http://leetcode.com/problems` and `http://leetcode.com/contest` are under the same hostname, while URLs `http://example.org/test` and `http://example.com/abc` are not under the same hostname.


The `HtmlParser` interface is defined as such:

```
interface HtmlParser {
    // Return a list of all urls from a webpage of given url.
    // This is a blocking call, that means it will do HTTP
    // request and return when this request is finished.
    public List<String> getUrls(String url);
}
```

Note that `getUrls(String url)` simulates performing an HTTP request. You can treat it as a blocking function call that waits for an HTTP request to finish. It is guaranteed that `getUrls(String url)` will return the URLs within 15ms. Single-threaded solutions will exceed the time limit so, can your multi-threaded web crawler do better?

Below are two examples explaining the functionality of the problem. For custom testing purposes, you'll have three variables `urls`, `edges` and `startUrl`. Notice that you will only have access to `startUrl` in your code, while `urls` and `edges` are not directly accessible to you in code.


Example 1:



Input:

```
urls = [
    "http://news.yahoo.com",
    "http://news.yahoo.com/news",
    "http://news.yahoo.com/news/topics/",
    "http://news.google.com",
    "http://news.yahoo.com/us"
]
```

Example 2:



```
Input:
urls = [
    "http://news.yahoo.com",
    "http://news.yahoo.com/news",
    "http://news.yahoo.com/news/topics/",
    "http://news.google.com"
]
edges = [[0,2],[2,1],[3,2],[3,1],[3,0]]
```

Constraints:

- $1 \leq \text{urls.length} \leq 1000$
- $1 \leq \text{urls}[i].\text{length} \leq 300$
- startUrl is one of the urls.
- Hostname label must be from 1 to 63 characters long, including the dots, may contain only the ASCII letters

from 'a' to 'z', digits from '0' to '9' and the hyphen-minus character ('-').

- The hostname may not start or end with the hyphen-minus character ('-').
- See: https://en.wikipedia.org/wiki/Hostname#Restrictions_on_valid_hostnames
- You may assume there're no duplicates in the URL library.

Follow up:

1. Assume we have 10,000 nodes and 1 billion URLs to crawl. We will deploy the same software onto each node. The software can know about all the nodes. We have to minimize communication between machines and make sure each node does equal amount of work. How would your web crawler design change?
2. What if one node fails or does not work?
3. How do you know when the crawler is done?

Community Solutions (Python)

- SimplyLeet

```

import threading
from queue import Queue
from urllib.parse import urlparse

def crawl(startUrl, htmlParser):
    # Extract hostname of the startUrl
    start_host = urlparse(startUrl).netloc
    # Thread-safe set for visited URLs and a lock
    for safety
        visited = set()
        visited_lock = threading.Lock()

```

```

    # Queue to process URLs; initialized with
    startUrl
    url_queue = Queue()
    url_queue.put(startUrl)

    def worker():
        while True:
            try:
                current_url =
url_queue.get(timeout=1) # timeout to avoid
blocking indefinitely
            except:

```

```

                break
            # Retrieve URLs from the current page
            using htmlParser
            for next_url in
htmlParser.getUrls(current_url):
                # Check if next_url shares the same
                hostname as startUrl
                if urlparse(next_url).netloc ==
start_host:
                    with visited_lock:
                        if next_url not in
visited:
                            visited.add(next_url)
                            url_queue.put(next_url
)
                    url_queue.task_done()

```

```

    # Mark startUrl as visited
    with visited_lock:
        visited.add(startUrl)

    # Launch worker threads (8 threads used as
    example)
    threads = []
    for _ in range(8):
        t = threading.Thread(target=worker)
        t.start()

```

```

        threads.append(t)

    # Wait until all URLs have been processed
    url_queue.join()

    # Ensure all threads terminate
    for t in threads:
        t.join()

    return list(visited)

```

- LC.ca

```

from threading import Thread, Lock
from urllib.parse import urlparse # To parse URLs
and extract hostname for comparison
from typing import List

class Solution:
    def crawl(self, startUrl: str, htmlParser:
'HtmlParser') -> List[str]:
        hostname = urlparse(startUrl).hostname
        seen = set([startUrl]) # Keep track of seen
URLs to avoid re-crawling the same URL
        lock = Lock() # Lock for thread-safe
operations on the seen set

```

```

def dfs(url):
    for next_url in
htmlParser.getUrls(url):
        if urlparse(next_url).hostname ==
hostname and next_url not in seen:
            with lock: # Ensure
thread-safe write to the seen set
                if next_url not in seen:
                    seen.add(next_url)
                    dfs(next_url)

```

```

def worker():
    while True:

```

```

        with lock:
            if not unseen:
                return
            url = unseen.pop()
            dfs(url)

        unseen = [startUrl] # Stack of URLs to be
crawled
        threads = []
        num_threads = 10 # Or any number you deem
appropriate based on the problem constraints

```

```
# Start threads
for _ in range(num_threads):
    thread = Thread(target=worker)
    thread.start()
    threads.append(thread)
```

```
# Wait for all threads to finish
for thread in threads:
    thread.join()
```

```
return list(seen)
```

```
#####
```

```
# reference https://levelup.gitconnected.com/multi-threaded-python-web-crawler-for-https-pages-el03f0839b71
```

```
class Crawler(threading.Thread):
    def __init__(self,base_url,
links_to_crawl,have_visited,
error_links,url_lock):
```

```
        threading.Thread.__init__(self)
        print(f"Web Crawler worker
{threading.current_thread()} has Started")
        self.base_url = base_url
        self.links_to_crawl = links_to_crawl
```

```
self.have_visited = have_visited
self.error_links = error_links
self.url_lock = url_lock
```

```
def run(self):
    # we create a ssl context so that our
script can crawl
    # the https sties with ssl_handshake.

    #Create a SSLContext object with
default settings.
    my_ssl = ssl.create_default_context()
```

```
        # by default when creating a default
ssl context and making an handshake
        # we verify the hostname with the
certificate but our objective is to crawl
        # the webpage so we will not be
checking the validity of the cerfificate.
        my_ssl.check_hostname = False

        # in this case we are not verifying the
certificate and any
        # certificate is accepted in this mode.
        my_ssl.verify_mode = ssl.CERT_NONE
```

```

        # we are defining an infinite while
loop so that all the links in our
        # queue are processed.

    while True:

        # In this part of the code we
create a global lock on our queue of
        # links so that no two threads can
access the queue at same time
        self.url_lock.acquire()
        print(f"Queue Size:
{self.links_to_crawl.qsize()}")
        link = self.links_to_crawl.get()

        self.url_lock.release()

        # if the link is None the queue is
exhausted or the threads are yet
        # process the links.

        if link is None:
            break

        # if The link is already visited we
break the execution.
        if link in self.have_visited:

```

```

        print(f"The link {link} is
already visited")
        break

    try:
        # This method constructs a
full "absolute" URL by combining the
        # base url with other url. this
uses components of the base URL,
        # in particular the addressing
scheme, the network
        # location and the path, to
provide missing components
        # in the relative URL.
        # in short we repair our
relative url if it is broken.

```

```

        link =
urljoin(self.base_url,link)

        # we use the header parameter
to "spoof" the "User-Agent" header
        # value which is used by the
browser to identify itself. This is
        # because some servers will
only allow the connection if it comes
        # from a verified browser. In
this case we are using FireFox header.
        req = Request(link, headers=
{'User-Agent': 'Mozilla/5.0'})

        # we are opening the url using
a ssl handshake.
        response = urlopen(req,
context=my_ssl)

```

Round 2 · High · Medium

p.1303

```

        print(f"The URL
{response.geturl()} crawled with \
        status
{response.getcode()}")

        # this returns the html
representation of the webpage
        soup =
BeautifulSoup(response.read(),"html.parser")

        # in this case we are finding
all the links in the page.
        for a_tag in
soup.find_all('a'):
            # we are checking of the
link is already visited and (network location part)
is our

```

Round 2 · High · Medium

p.1304


```

                # base url itself.
                if (a_tag.get("href") not
in self.have_visited) and (urlparse(link).netloc
== "www.python.org"):
                    self.links_to_crawl.put(a_tag.get("href"))

                else:
                    print(f"The link
{a_tag.get('href')} is already visited or is not
part \
                    of the website")

                    print(f"Adding {link} to the
crawled list")
                    self.have_visited.add(link)

            except URLError as e:
                print(f"URL {link} threw this
error {e.reason} while trying to parse")

                self.error_links.append(link)

            finally:
                self.links_to_crawl.task_done(
)

```

Round 2 · High · Medium

p.1305

```

if __name__ == '__main__':
    print("The Crawler is started")
    base_url = input("Please Enter Website to Crawl
> ")
    number_of_threads = input("Please Enter number
of Threads > ")

    links_to_crawl = queue.Queue()
    url_lock = threading.Lock()
    links_to_crawl.put(base_url)

    have_visited = set()

    crawler_threads = []
    error_links = []
    #base_url, links_to_crawl,have_visited,
error_links,url_lock
    for i in range(int(number_of_threads)):
        crawler = Crawler(base_url = base_url,
                        links_to_crawl=
links_to_crawl,
                        have_visited=
have_visited,
                        error_links= error_links,
                        url_lock=url_lock)

```

Round 2 · High · Medium

p.1306

```

        crawler.start()
        crawler_threads.append(crawler)

    for crawler in crawler_threads:
        crawler.join()

    print(f"Total Number of pages visited are
    {len(have_visited)}")
    print(f"Total Number of Errornous links:
    {len(error_links)}")

#####

import multiprocessing
from bs4 import BeautifulSoup
from queue import Queue, Empty
from concurrent.futures import ThreadPoolExecutor
from urllib.parse import urljoin, urlparse
import requests

# ThreadPoolExecutor
class MultiThreadedCrawler:

```

```

    def __init__(self, seed_url):
        self.seed_url = seed_url
        self.root_url =
        '{}://{}'.format(urlparse(self.seed_url).scheme,
        urlparse(self.seed_url).netloc)
        self.pool = ThreadPoolExecutor(max_workers=5)
        self.scraped_pages = set([])
        self.crawl_queue = Queue()
        self.crawl_queue.put(self.seed_url)

    def parse_links(self, html):

        soup = BeautifulSoup(html, 'html.parser')
        Anchor_Tags = soup.find_all('a', href=True)
        for link in Anchor_Tags:
            url = link['href']
            if url.startswith('/') or
            url.startswith(self.root_url):
                url = urljoin(self.root_url, url)
            if url not in self.scraped_pages:
                self.crawl_queue.put(url)

        def scrape_info(self, html):

```

```
soup = BeautifulSoup(html, "html5lib")
web_page_paragraph_contents = soup('p')
text = ''
for para in web_page_paragraph_contents:
    if not ('https:' in str(para.text)):
        text = text + str(para.text).strip()
    print(f'\n <---Text Present in The WebPage is
    --->\n', text, '\n')
return
```

```
def post_scrape_callback(self, res):
```

```
    result = res.result()
    if result and result.status_code == 200:
        self.parse_links(result.text)
        self.scrape_info(result.text)
```

```
def scrape_page(self, url):
    try:
        res = requests.get(url, timeout=(3, 30))
        return res
    except requests.RequestException:
```

```
return
```

```
def run_web_crawler(self):
    while True:
        try:
            print("\n Name of the current executing process: ",
            multiprocessing.current_process().name, '\n')
            target_url = self.crawl_queue.get(timeout=60)
            if target_url not in self.scraped_pages:
                print("Scraping URL: {}".format(target_url))
```

```
        self.current_scraping_url =
        "{}".format(target_url)
        self.scraped_pages.add(target_url)
        job = self.pool.submit(self.scrape_page,
        target_url)
        job.add_done_callback(self.post_scrape_callback)
```

```
    except Empty:
        return
    except Exception as e:
        print(e)
        continue
```

```
def info(self):
    print('\n Seed URL is: ', self.seed_url, '\n')
    print('Scraped pages are: ', self.scraped_pages,
          '\n')

if __name__ == '__main__':
    cc =
    MultiThreadedCrawler("https://www.leetcode.ca")
    cc.run_web_crawler()
    cc.info()
```

◆ Quick Review

Key Insights

- Use multiple threads to perform concurrent crawling, which speeds up the process over a single-threaded approach.
- Maintain a thread-safe set (visited) to ensure no URL is processed twice.
- Utilize a work queue to manage URLs pending processing.
- Filter URLs based on the hostname extracted from startUrl.
- Apply synchronization (locks or concurrent data structures) to avoid race conditions when multiple threads access shared data.

Space & Time

Time Complexity: $O(N)$ where N is the number of URLs processed, with additional constant factors due to threading overhead.

Space Complexity: $O(N)$ for storing the visited URLs and the work queue.

Solution

The solution initializes a shared work queue with the startUrl and a thread-safe visited set. Each worker thread dequeues a URL, retrieves its linked URLs via the blocking `HtmlParser.getUrls()` call, and filters the resulting URLs by comparing their hostnames with the start URL's hostname. New URLs that have not yet been visited are added to both the visited set and the work queue. The use of a thread pool or multiple worker threads along with synchronization ensures that the crawler processes URLs concurrently while avoiding race conditions. The approach follows a BFS-like pattern spread across multiple threads.

Graphs

Round 2 · High · Medium · 3 questions

Round 2 · High · Medium

p.1313

Graphs

#128 Longest Consecutive Sequence

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Union Find

Lists: Grind 169 | CodeSignal

Problem

Description

Given an unsorted array of integers `nums`, return the length of the longest consecutive elements sequence.

You must write an algorithm that runs in $O(n)$ time.

Example 1:

Input: `nums = [100,4,200,1,3,2]`

Output: 4

Explanation: The longest consecutive elements sequence is `[1, 2, 3, 4]`. Therefore its length is 4.

Example 2:

Input: `nums = [0,3,7,2,5,8,4,6,0,1]`

Output: 9

Example 3:

Input: `nums = [1,0,1,2]`

Output: 3

Round 2 · High · Medium

p.1314

Constraints:

- $0 \leq \text{nums.length} \leq 105$
- $-109 \leq \text{nums}[i] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestConsecutive(self, nums: list[int]) -> int:
        ans = 0
        seen = set(nums)

        for num in seen:
            # `num` is the start of a sequence.
            if num - 1 in seen:
                continue
            length = 0

            while num in seen:
                num += 1
                length += 1
            ans = max(ans, length)

        return ans
```

• LeetCode

```
class Solution:
    def longestConsecutive(self, nums: List[int]) -> int:
        s = set(nums)
        ans = 0
        d = defaultdict(int)
        for x in nums:
            y = x
            while y in s:
                s.remove(y)
                y += 1

            d[x] = d[y] + y - x
            ans = max(ans, d[x])
        return ans
```

• SimplyLeet

```
# Python solution for Longest Consecutive Sequence

def longestConsecutive(nums):
    # Create a set from the list for O(1) look-up times
    num_set = set(nums)
    longest_streak = 0

    # Iterate through each number in the set
    for num in num_set:
        # Only start counting if 'num-1' is not present in the set
```

```

    if num - 1 not in num_set:
        current_num = num
        current_streak = 1

    # Count consecutive numbers starting
    from current_num
    while current_num + 1 in num_set:
        current_num += 1
        current_streak += 1

    # Update the longest streak if current
    one is longer

    longest_streak = max(longest_streak,
        current_streak)

    return longest_streak

# Example usage:
print(longestConsecutive([100, 4, 200, 1, 3, 2])) #
Output: 4

```

- LC.ca

```

class Solution:
    def longestConsecutive(self, nums: List[int])
    -> int:

        s = set(nums)
        ans = 0
        for x in nums:
            if x - 1 not in s:
                y = x + 1
                while y in s:
                    y += 1
                ans = max(ans, y - x)

        return ans

```

◆ Quick Review

Key Insights

- Utilize a set (or hash table) for constant time look-ups.
- Only start counting a sequence from a number that is the beginning (its previous number is not in the set).
- Each number is visited only once, ensuring an $O(n)$ time complexity.
- Update the maximum sequence length as you discover longer consecutive sequences.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

We first insert all the numbers into a hash set to enable $O(1)$ membership checks. For each number in the set, we check if it is the start of a sequence by ensuring that the number minus one is not in the set. If it is the beginning, we then count all consecutive numbers following it, updating the maximum sequence length when necessary. This method guarantees that we only count each number once, satisfying the $O(n)$ time requirement.

Round 2 · High · Medium

p.1319

Graphs

#277 Find the Celebrity

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · Graph · Interactive

Lists: Premium 98

Problem

Description

Suppose you are at a party with n people labeled from 0 to $n - 1$ and among them, there may exist one celebrity. The definition of a celebrity is that all the other $n - 1$ people know the celebrity, but the celebrity does not know any of them.

Now you want to find out who the celebrity is or verify that there is not one. You are only allowed to ask questions like: "Hi, A. Do you know B?" to get information about whether A knows B. You need to find out the celebrity (or verify there is not one) by asking as few questions as possible (in the asymptotic sense).

You are given an integer n and a helper function `bool knows(a, b)` that tells you whether a knows b. Implement a function `int findCelebrity(n)`. There will be exactly one celebrity if they are at the party.


Return the celebrity's label if there is a celebrity at the party. If there is no celebrity, return -1 .

Round 2 · High · Medium

p.1320

Note that the $n \times n$ 2D array `graph` given as input is not directly available to you, and instead only accessible through the helper function `knows`. `graph[i][j] == 1` represents person i knows person j , whereas `graph[i][j] == 0` represents person j does not know person i .

Example 1:




Input: `graph = [[1,1,0],[0,1,0],[1,1,1]]`
 Output: 1
 Explanation: There are three persons labeled with 0, 1 and 2. `graph[i][j] = 1` means person i knows person j , otherwise `graph[i][j] = 0` means person i does not know person j . The celebrity is the person labeled as 1 because both 0 and 2 know him but 1 does not know anybody.

Round 2 · High · Medium

p.1321

Example 2:



Input: `graph = [[1,0,1],[1,1,0],[0,1,1]]`
 Output: -1
 Explanation: There is no celebrity.

Constraints:

- $n == \text{graph.length} == \text{graph}[i].\text{length}$
- $2 \leq n \leq 100$
- `graph[i][j]` is 0 or 1.
- `graph[i][i] == 1`

Follow up: If the maximum number of allowed calls to the API knows is $3 * n$, could you find a solution without exceeding the maximum number of calls?

Round 2 · High · Medium

p.1322

Community Solutions (Python)

• WalkCC

```
# The knows API is already defined for you.
# Returns a bool, whether a knows b
# Def knows(a: int, b: int) -> bool:

class Solution:
    def findCelebrity(self, n: int) -> int:
        candidate = 0

        # Everyone knows the celebrity.

        for i in range(1, n):
            if knows(candidate, i):
                candidate = i

        # The candidate knows nobody and everyone knows
        the celebrity.
        for i in range(n):
            if i < candidate and knows(candidate, i) or
            not knows(i, candidate):
                return -1
            if i > candidate and not knows(i, candidate):
                return -1

        return candidate
```

• LeetCode

Round 2 · High · Medium

p.1323

```
# The knows API is already defined for you.
# return a bool, whether a knows b
# def knows(a: int, b: int) -> bool:
```

```
class Solution:
    def findCelebrity(self, n: int) -> int:
        ans = 0
        for i in range(1, n):
            if knows(ans, i):

                ans = i
        for i in range(n):
            if ans != i:
                if knows(ans, i) or not knows(i,
ans):

                    return -1
        return ans
```

• SimplyLeet

Round 2 · High · Medium

p.1324

```
# Function definition for findCelebrity using the
helper knows(a, b) API.
def findCelebrity(n):
    # Step 1: Candidate Selection.
    candidate = 0
    for i in range(1, n):
        # If candidate knows i, then candidate
        cannot be the celebrity.
        if knows(candidate, i):
            candidate = i # Select new candidate.

    # Step 2: Verification.

    # Check if candidate does not know any other
    person.
    for i in range(n):
        if i != candidate:
            # Either candidate knows someone or
            someone doesn't know candidate.
            if knows(candidate, i) or not knows(i,
            candidate):
                return -1
    return candidate

# Example helper function (for testing, this is
normally provided in the problem context).
def knows(a, b):
```

```
# This function should return True if person a
knows person b, otherwise False.
pass
```

• LC.ca

```
# The knows API is already defined for you.
# return a bool, whether a knows b
# def knows(a: int, b: int) -> bool:
```

```
class Solution:
    def findCelebrity(self, n: int) -> int:
        ans = 0
        for i in range(1, n):
            if knows(ans, i):
```

```
                ans = i
        for i in range(n):
            if ans != i:
                if knows(ans, i) or not knows(i,
            ans):
                    return -1
        return ans
```

◆ Quick Review

Key Insights

- Use pairwise comparisons to eliminate non-celebrities.
- If a candidate knows another person, the candidate cannot be the celebrity.

- Verify the candidate by checking that they do not know anyone and that everyone knows them.
- Achieve the solution in two passes: one for candidate selection and the other for validation.

Space & Time

Time Complexity: $O(n)$ – One pass for candidate selection and one for validation.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

The approach consists of two main steps:

– Candidate Selection: Iterate through the people and maintain a candidate. For every person i , if the current candidate knows i , then set $\text{candidate} = i$. This works because if candidate knows i , candidate cannot be the celebrity. – Verification: Check that the candidate does not know any other person and that every other person knows the candidate. If both conditions hold, candidate is the celebrity; otherwise, return -1 .

The main data structures used are simple integer variables. The algorithm leverages a two-pass technique to reduce the number of knows API calls while ensuring correctness.

Graphs

#1136 Parallel Courses

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Graph · Topological Sort

Lists: Premium 98


Problem

Description

You are given an integer n , which indicates that there are n courses labeled from 1 to n . You are also given an array `relations` where `relations[i] = [prevCoursei, nextCoursei]`, representing a prerequisite relationship between course `prevCoursei` and course `nextCoursei`: course `prevCoursei` has to be taken before course `nextCoursei`.

In one semester, you can take any number of courses as long as you have taken all the prerequisites in the previous semester for the courses you are taking. Return the minimum number of semesters needed to take all courses. If there is no way to take all the courses, return -1 .

Example 1:



Input: $n = 3$, relations = $[[1,3],[2,3]]$


Output: 2

Explanation: The figure above represents the given graph.

In the first semester, you can take courses 1 and 2.

In the second semester, you can take course 3.

Example 2:



Input: $n = 3$, relations = $[[1,2],[2,3],[3,1]]$

Output: -1

Explanation: No course can be studied because they are prerequisites of each other.

Constraints:

- $1 \leq n \leq 5000$
- $1 \leq \text{relations.length} \leq 5000$
- $\text{relations}[i].\text{length} == 2$
- $1 \leq \text{prevCoursei}, \text{nextCoursei} \leq n$
- $\text{prevCoursei} \neq \text{nextCoursei}$

- All the pairs [prevCoursei, nextCoursei] are unique.

Community Solutions (Python)

• WalkCC

```
from enum import Enum

class State(Enum):
    INIT = 0
    VISITING = 1
    VISITED = 2

class Solution:

    def minimumSemesters(self, n: int, relations:
list[list[int]]) -> int:
        graph = [[] for _ in range(n)]
        states = [State.INIT] * n
        depth = [1] * n

        for u, v in relations:
            graph[u - 1].append(v - 1)

        def hasCycle(u: int) -> bool:
            if states[u] == State.VISITING:
```

```
        return True
        if states[u] == State.VISITED:
            return False
        states[u] = State.VISITING
        for v in graph[u]:
            if hasCycle(v):
                return True
            depth[u] = max(depth[u], 1 + depth[v])
        states[u] = State.VISITED
        return False
```

```
        if any(hasCycle(i) for i in range(n)):
            return -1
        return max(depth)
```

```
class Solution:
    def minimumSemesters(self, n: int, relations:
list[list[int]]) -> int:
        graph = [[] for _ in range(n)]
        inDegrees = [0] * n

        # Build the graph.
        for u, v in relations:
            graph[u - 1].append(v - 1)
            inDegrees[v - 1] += 1
```

```
# Perform topological sorting.
q = collections.deque([i for i, d in
enumerate(inDegrees) if d == 0])
```

```
step = 0
while q:
    for _ in range(len(q)):
        u = q.popleft()
        n -= 1
        for v in graph[u]:
            inDegrees[v] -= 1
```

```
        if inDegrees[v] == 0:
            q.append(v)
        step += 1

    return step if n == 0 else -1
```

• LeetCode

Round 2 · High · Medium

p.1333

```
class Solution:
    def minimumSemesters(self, n: int, relations:
List[List[int]]) -> int:
        g = defaultdict(list)
        indeg = [0] * n
        for prev, nxt in relations:
            prev, nxt = prev - 1, nxt - 1
            g[prev].append(nxt)
            indeg[nxt] += 1
        q = deque(i for i, v in enumerate(indeg) if
v == 0)
        ans = 0
```

```
        while q:
            ans += 1
            for _ in range(len(q)):
                i = q.popleft()
                n -= 1
                for j in g[i]:
                    indeg[j] -= 1
                    if indeg[j] == 0:
                        q.append(j)
        return -1 if n else ans
```

• SimplyLeet

Round 2 · High · Medium

p.1334

Python Implementation

```

from collections import deque

def minimumSemesters(numCourses, relations):
    # Build graph and in-degree dictionary.
    graph = {i: [] for i in range(1, numCourses + 1)}
    in_degree = {i: 0 for i in range(1, numCourses + 1)}

    # Construct graph from relations.

    for pre, course in relations:
        graph[pre].append(course)
        in_degree[course] += 1

    # Initialize queue with all courses having no prerequisites.
    queue = deque()
    for course in range(1, numCourses + 1):
        if in_degree[course] == 0:
            queue.append(course)

```

```

semesters = 0
takenCourses = 0

# Process each level (semester) using BFS.
while queue:
    semesters += 1
    current_level_size = len(queue)
    for _ in range(current_level_size):
        course = queue.popleft()
        takenCourses += 1

```

```

        # Decrement in-degree for neighboring courses.
        for neighbor in graph[course]:
            in_degree[neighbor] -= 1
            if in_degree[neighbor] == 0:
                queue.append(neighbor)

    # If not all courses have been taken, a cycle exists.
    return semesters if takenCourses == numCourses else -1

```

Example usage:

```

print(minimumSemesters(3, [[1, 3], [2, 3]])) #
Expected output: 2

```

• [LC.ca](#)


```

class Solution:
    def minimumSemesters(self, n: int, relations:
List[List[int]]) -> int:
        g = defaultdict(list)
        indeg = [0] * n
        for prev, nxt in relations:
            prev, nxt = prev - 1, nxt - 1
            g[prev].append(nxt)
            indeg[nxt] += 1
        q = deque(i for i, v in enumerate(indeg) if
v == 0)
        ans = 0

        while q:
            ans += 1
            for _ in range(len(q)):
                i = q.popleft()
                n -= 1
                for j in g[i]:
                    indeg[j] -= 1
                    if indeg[j] == 0:
                        q.append(j)
            return -1 if n else ans

```

◆ Quick Review

Key Insights

- Model the problem as a directed graph where nodes are courses and edges represent prerequisites.

- Use topological sorting (Kahn's algorithm) to process courses with zero prerequisites.
- Each level (or layer) of processing represents one semester.
- Detect cycles by ensuring all courses are eventually processed; if not, a cycle exists.
- The number of BFS levels gives the minimum number of semesters necessary.

Space & Time

Time Complexity: $O(n + m)$ where n is the number of courses and m is the length of relations (prerequisites).

Space Complexity: $O(n + m)$ for storing the graph, in-degree counts, and the processing queue.

Solution

We solve the problem using a topological sort with Kahn's algorithm. First, construct a graph representation using an adjacency list and compute the in-degree for each course. Initialize a queue with courses that have no prerequisites (in-degree == 0). Then, process courses level by level. Each level processed represents one semester. For every course taken in that semester, decrease the in-degree of all its neighbors. Add any neighbor with an in-degree that becomes zero to the queue for the next semester. If, after processing, the total courses taken is less than n , then a cycle exists and return -1. Otherwise, the number of semesters processed is the answer.

BFS

Round 2 · High · Medium · 6 questions

Round 2 · High · Medium

p.1339

BFS

#286 Walls and Gates

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Premium 98

Problem

Description

You are given an $m \times n$ grid rooms initialized with these three possible values.



- -1 A wall or an obstacle.
- 0 A gate.
- INF Infinity means an empty room. We use the value $2^{31} - 1 = 2147483647$ to represent INF as you may assume that the distance to a gate is less than 2147483647.

Fill each empty room with the distance to its nearest gate. If it is impossible to reach a gate, it should be filled with INF.

Example 1:

Round 2 · High · Medium

p.1340



Input: rooms = [[2147483647,-1,0,2147483647],[2147483647,2147483647,2147483647,-1],[2147483647,-1,2147483647,-1],[0,-1,2147483647,2147483647]]
Output: [[3,-1,0,1],[2,2,1,-1],[1,-1,2,-1],[0,-1,3,4]]

Example 2:

Input: rooms = [[-1]]
Output: [[-1]]

Constraints:

- $m == \text{rooms.length}$
- $n == \text{rooms}[i].\text{length}$
- $1 \leq m, n \leq 250$
- $\text{rooms}[i][j]$ is -1, 0, or 231 - 1.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def wallsAndGates(self, rooms: list[list[int]])
    -> None:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        INF = 2**31 - 1
        m = len(rooms)
        n = len(rooms[0])
        q = collections.deque((i, j)
                               for i in range(m)
                               for j in range(n)
                               if rooms[i][j] == 0)
```

```
while q:
    i, j = q.popleft()
    for dx, dy in DIRS:
        x = i + dx
        y = j + dy
        if x < 0 or x == m or y < 0 or y == n:
            continue
        if rooms[x][y] != INF:
            continue
```

```
rooms[x][y] = rooms[i][j] + 1
q.append((x, y))
```

- LeetDoocs

```

class Solution:
    def wallsAndGates(self, rooms:
List[List[int]]) -> None:
        """
        Do not return anything, modify rooms
in-place instead.
        """
        m, n = len(rooms), len(rooms[0])
        inf = 2**31 - 1
        q = deque([(i, j) for i in range(m) for j
in range(n) if rooms[i][j] == 0])
        d = 0
        while q:
            d += 1
            for _ in range(len(q)):
                i, j = q.popleft()
                for a, b in [[0, 1], [0, -1], [1,
0], [-1, 0]]:
                    x, y = i + a, j + b
                    if 0 <= x < m and 0 <= y < n
and rooms[x][y] == inf:
                        rooms[x][y] = d
                        q.append((x, y))

```

• SimplyLeet

```

from collections import deque

def wallsAndGates(rooms):
    if not rooms or not rooms[0]:
        return

    rows, cols = len(rooms), len(rooms[0])
    # Directions: up, down, left, right.
    directions = [(-1, 0), (1, 0), (0, -1), (0, 1)]
    queue = deque()

    # Enqueue all gates (cells with 0).
    for r in range(rows):
        for c in range(cols):
            if rooms[r][c] == 0:
                queue.append((r, c))

    # Perform multi-source BFS.
    while queue:
        r, c = queue.popleft()

```

```

        for dr, dc in directions:
            nr, nc = r + dr, c + dc
            # Process neighbor if it is within
            bounds and is an empty room.
            if 0 <= nr < rows and 0 <= nc < cols
            and rooms[nr][nc] == 2147483647:
                rooms[nr][nc] = rooms[r][c] + 1 #
            Update distance.
                queue.append((nr, nc))

```

• LC.ca

```

class Solution:
    def wallsAndGates(self, rooms:
    List[List[int]]) -> None:
        """
        Do not return anything, modify rooms
        in-place instead.
        """
        m, n = len(rooms), len(rooms[0])
        inf = 2**31 - 1
        q = deque([(i, j) for i in range(m) for j
        in range(n) if rooms[i][j] == 0])
        d = 0
        while q:

```

```

            d += 1
            for _ in range(len(q)):
                i, j = q.popleft()
                for a, b in [[0, 1], [0, -1], [1,
                0], [-1, 0]]:
                    x, y = i + a, j + b
                    # bfs so first reaching inf is
                    always the shortest
                    # set its value from inf to d,
                    same as marking as visited
                    if 0 <= x < m and 0 <= y < n
                    and rooms[x][y] == inf:
                        rooms[x][y] = d
                        q.append((x, y))

```

#####

```

class Solution: # dfs
    def wallsAndGates(self, rooms:
    List[List[int]]) -> None:
        if not rooms or not rooms[0]:
            return

        m, n = len(rooms), len(rooms[0])
        isVisited = [[False] * n for _ in range(m)]
        # isVisited map shared among all gates

```

```

def fill(i: int, j: int, distance: int) ->
None:
    nonlocal m, n

    # rooms[i][j] <= 0 meaning == -1 or ==
0, both stop
    if i < 0 or i >= m or j < 0 or j >= n
or rooms[i][j] <= 0 or isVisited[i][j]:
        return

    rooms[i][j] = min(rooms[i][j],
distance)
    isVisited[i][j] = True

```

```

    fill(i - 1, j, distance + 1)
    fill(i, j + 1, distance + 1)
    fill(i + 1, j, distance + 1)
    fill(i, j - 1, distance + 1)

    isVisited[i][j] = False

    for i in range(m):
        for j in range(n):

```

```

        if rooms[i][j] == 0: # start from
every gate
            fill(i, j, 0)

#####

from collections import deque

INF = 2147483647

class Solution(object):
    def wallsAndGates(self, rooms):
        """
        :type rooms: List[List[int]]
        :rtype: void Do not return anything, modify
rooms in-place instead.
        """
        queue = deque([])
        directions = [(1, 0), (-1, 0), (0, 1), (0, -1)]
        for i in range(0, len(rooms)):
            for j in range(0, len(rooms[0])):

```

```

    if rooms[i][j] == 0:
        queue.append((i, j))

    while queue:
        i, j = queue.popleft()
        for di, dj in directions:
            p, q = i + di, j + dj
            if 0 <= p < len(rooms) and 0 <= q <
len(rooms[0]) and rooms[p][q] == INF:
                rooms[p][q] = rooms[i][j] + 1
                queue.append((p, q))

```

◆ Quick Review

Key Insights

- Use a multi-source Breadth-First Search (BFS) starting from all the gates, as this allows spreading the shortest distance to each empty room.
- Instead of starting from every empty room separately, starting from all gates simultaneously helps update all distances in one pass.
- Only update a room if the new calculated distance is smaller than its current value to prevent unnecessary processing.
- The problem is essentially a multi-source shortest path problem on a grid.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the grid, as each cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario used by the BFS queue.

Solution

The solution uses a multi-source BFS algorithm by initially enqueueing all the gate positions (cells with value 0). Then, for each cell obtained from the queue, check its four possible neighbors (up, down, left, right). If a neighbor is an empty room (INF), update its distance to be the current cell's distance plus one, and then enqueue the neighbor to process further. This ensures that each empty room gets the minimum distance from any gate.

BFS

#322 Coin Change

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · BFS

Lists: Grind 169

Problem

Description

You are given an integer array `coins` representing coins of different denominations and an integer `amount` representing a total amount of money.

Return the fewest number of coins that you need to make up that amount. If that amount of money cannot be made up by any combination of the coins, return `-1`.

You may assume that you have an infinite number of each kind of coin.

Example 1:

Input: `coins = [1,2,5]`, `amount = 11`Output: `3`Explanation: `11 = 5 + 5 + 1`

Example 2:

Input: `coins = [2]`, `amount = 3`Output: `-1`

Example 3:

Input: `coins = [1]`, `amount = 0`Output: `0`

Constraints:

- `1 <= coins.length <= 12`
- `1 <= coins[i] <= 231 - 1`
- `0 <= amount <= 104`

Community Solutions (Python)

• WalkCC

```
class Solution:
    def coinChange(self, coins: list[int], amount:
int) -> int:
        # dp[i] := the minimum number of coins to make
up i
        dp = [0] + [amount + 1] * amount

        for coin in coins:
            for i in range(coin, amount + 1):
                dp[i] = min(dp[i], dp[i - coin] + 1)

        return -1 if dp[amount] == amount + 1 else
dp[amount]
```

• LeetDoocs


```

class Solution:
    def coinChange(self, coins: List[int], amount:
int) -> int:
        m, n = len(coins), amount
        f = [[inf] * (n + 1) for _ in range(m + 1)]
        f[0][0] = 0
        for i, x in enumerate(coins, 1):
            for j in range(n + 1):
                f[i][j] = f[i - 1][j]
                if j >= x:
                    f[i][j] = min(f[i][j], f[i][j
- x] + 1)

        return -1 if f[m][n] >= inf else f[m][n]

```

```

class Solution:
    def coinChange(self, coins: List[int], amount:
int) -> int:
        n = amount
        f = [0] + [inf] * n
        for x in coins:
            for j in range(x, n + 1):
                f[j] = min(f[j], f[j - x] + 1)
        return -1 if f[n] >= inf else f[n]

```

• SimplyLeet

```

def coinChange(coins, amount):
    # Create a dp array with amount+1 entries set
to a value greater than any possible minimum number
of coins.
    dp = [amount + 1] * (amount + 1)
    # Base case: 0 coins are needed to form amount
0.
    dp[0] = 0

    # Build up the dp table.
    for current_amount in range(1, amount + 1):
        for coin in coins:
            if coin <= current_amount:

                dp[current_amount] =
min(dp[current_amount], dp[current_amount - coin]
+ 1)

        # If dp[amount] remains greater than amount, no
solution was found.
        return dp[amount] if dp[amount] != amount + 1
        else -1

```

• LC.ca

```
'''
>>> float("inf")
inf
>>> float("inf") + 1
inf
'''

class Solution(object):
    def coinChange(self, coins, amount):
        '''
        :type coins: List[int]
        :type amount: int
        :rtype: int
        '''

        dp = [float("inf")] * (amount + 1)
        dp[0] = 0
        for i in range(1, amount + 1):
            for coin in coins:
                if i - coin >= 0:
```

```
                dp[i] = min(dp[i], dp[i - coin] + 1)
        return dp[-1] if dp[-1] != float("inf") else -1

#####

class Solution:
    def coinChange(self, coins: List[int], amount:
int) -> int:
        dp = [amount + 1] * (amount + 1)
        dp[0] = 0
        for coin in coins:

            for j in range(coin, amount + 1):
                dp[j] = min(dp[j], dp[j - coin] +
1)

        return -1 if dp[-1] > amount else dp[-1]
```

◆ Quick Review

Key Insights

- Use a dynamic programming approach to build up a solution from smaller subproblems.
- Define $dp[x]$ as the minimum number of coins required for amount x .
- The recurrence relation is: $dp[x] = \min(dp[x], dp[x - \text{coin}] + 1)$ for each coin.
- Initialize $dp[0] = 0$ and set other values to a number larger than the maximum possible coins (amount+1 works as an infinity substitute).

- An alternative approach is through BFS, treating each amount as a node and each coin as an edge, but the DP method is more intuitive here.

Space & Time

Time Complexity: $O(\text{amount} * n)$ where n is the number of coins.

Space Complexity: $O(\text{amount})$ for the dp array.

Solution

We use a bottom-up dynamic programming approach. The idea is to use a one-dimensional dp array where each element at index x represents the minimum number of coins required to form the amount x . Starting from $dp[0]$ which is 0, we iterate through all amounts up to the target amount and update $dp[x]$ based on each coin denomination. If after filling the dp array the value $dp[\text{amount}]$ remains unchanged (i.e., greater than amount), it means the amount cannot be formed and we return -1.

BFS

#542 01 Matrix

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Grind 169 | Denny Zhang

Problem

Description

Given an $m \times n$ binary matrix `mat`, return the distance of the nearest 0 for each cell.

The distance between two cells sharing a common edge is 1.

Example 1:

0	0	0
0	1	0
0	0	0

Input: mat = [[0,0,0],[0,1,0],[0,0,0]]
Output: [[0,0,0],[0,1,0],[0,0,0]]

Example 2:

Round 2 · High · Medium

p.1359

0	0	0
0	1	0
1	1	1

Input: mat = [[0,0,0],[0,1,0],[1,1,1]]
Output: [[0,0,0],[0,1,0],[1,2,1]]

Constraints:

- $m == \text{mat.length}$
- $n == \text{mat}[i].\text{length}$
- $1 \leq m, n \leq 104$
- $1 \leq m * n \leq 104$
- $\text{mat}[i][j]$ is either 0 or 1.
- There is at least one 0 in mat.

Round 2 · High · Medium

p.1360

Note: This question is the same as 1765:

<https://leetcode.com/problems/map-of-highest-peak/>

Community Solutions (Python)

• WalkCC

```
class Solution:
    def updateMatrix(self, mat: list[list[int]]) -> list[list[int]]:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(mat)
        n = len(mat[0])
        q = collections.deque()
        seen = [[False] * n for _ in range(m)]

        for i in range(m):
            for j in range(n):
                if mat[i][j] == 0:
                    q.append((i, j))
                    seen[i][j] = True

        while q:
            i, j = q.popleft()
            for dx, dy in DIRS:
                x = i + dx
                y = j + dy
                if x < 0 or x == m or y < 0 or y == n:
```

```
        continue
    if seen[x][y]:
        continue
    mat[x][y] = mat[i][j] + 1
    q.append((x, y))
    seen[x][y] = True
```

```
    return mat
```

• LeetCode

```
class Solution:
    def updateMatrix(self, mat: List[List[int]]) -> List[List[int]]:
        m, n = len(mat), len(mat[0])
        ans = [[-1] * n for _ in range(m)]
        q = deque()
        for i, row in enumerate(mat):
            for j, x in enumerate(row):
                if x == 0:
                    ans[i][j] = 0
                    q.append((i, j))
```

```

    dirs = (-1, 0, 1, 0, -1)
    while q:
        i, j = q.popleft()
        for a, b in pairwise(dirs):
            x, y = i + a, j + b
            if 0 <= x < m and 0 <= y < n and
ans[x][y] == -1:
                ans[x][y] = ans[i][j] + 1
                q.append((x, y))
    return ans

```

• SimplyLeet

```

from collections import deque

def updateMatrix(mat):
    # Get the dimensions of the matrix
    rows, cols = len(mat), len(mat[0])
    # Initialize a queue for multi-source BFS
    queue = deque()

    # Initialize the matrix: use -1 to indicate
    # unvisited cells
    for r in range(rows):

```

```

        for c in range(cols):
            if mat[r][c] == 0:
                # Enqueue all 0 cells as starting
points
                queue.append((r, c))
            else:
                # Mark 1 cells as unvisited (or
infinite distance)
                mat[r][c] = -1

        # Directions for neighbors (up, down, left,
right)
        directions = [(1, 0), (-1, 0), (0, 1), (0, -1)]

    # BFS traversal to update distances
    while queue:
        r, c = queue.popleft()
        for dr, dc in directions:
            nr, nc = r + dr, c + dc
            # Check boundaries and if neighbor has
not been visited
            if 0 <= nr < rows and 0 <= nc < cols
and mat[nr][nc] == -1:
                # Update neighbor's distance
                mat[nr][nc] = mat[r][c] + 1

```

```

        # Enqueue the neighbor for further
BFS processing
        queue.append((nr, nc))

    return mat

# Example usage:
# print(updateMatrix([[0,0,0],[0,1,0],[1,1,1]]))

```

• LC.ca

```

class Solution:
    def updateMatrix(self, mat: List[List[int]])
-> List[List[int]]:
        m, n = len(mat), len(mat[0])
        ans = [[-1] * n for _ in range(m)]
        q = deque()
        for i, row in enumerate(mat):
            for j, x in enumerate(row):
                if x == 0:
                    ans[i][j] = 0
                    q.append((i, j))

```

```

        dirs = (-1, 0, 1, 0, -1)
        while q:
            i, j = q.popleft()
            for a, b in pairwise(dirs):
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n and
ans[x][y] == -1:
                    ans[x][y] = ans[i][j] + 1
                    q.append((x, y))

        return ans

```

◆ Quick Review

Key Insights

- Treat all 0s as starting points for a multi-source BFS.
- Update each neighboring cell's distance if a shorter distance is found.
- Utilize a queue to process cells in order of increasing distance.
- Ensure that each cell is processed only once by marking visited cells.

Space & Time

Time Complexity: $O(m$

$n)$ since each cell is processed once.

Space Complexity: $O(m$

$n)$ for the queue and the answer matrix.

Solution

We use a multi-source Breadth-First Search (BFS) where every 0 in the matrix is initially added to the BFS queue. Each neighboring cell that is a 1 will have its distance updated to be one plus the distance of the current cell if it hasn't been visited yet (or if a smaller distance is found). This ensures that the first time a 1 is reached, it is reached by the shortest path. Data structures used include a queue (or deque) for BFS and the matrix itself for storing distances.

BFS

#994 Rotting Oranges

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Grind 169

Problem

Description


You are given an $m \times n$ grid where each cell can have one of three values:

- 0 representing an empty cell,
- 1 representing a fresh orange, or
- 2 representing a rotten orange.

Every minute, any fresh orange that is 4-directionally adjacent to a rotten orange becomes rotten.

Return the minimum number of minutes that must elapse until no cell has a fresh orange. If this is impossible, return -1.

Example 1:



Round 2 · High · Medium

p.1367

Round 2 · High · Medium

p.1368

Input: grid = [[2,1,1],[1,1,0],[0,1,1]]

Output: 4

Example 2:

Input: grid = [[2,1,1],[0,1,1],[1,0,1]]

Output: -1

Explanation: The orange in the bottom left corner (row 2, column 0) is never rotten, because rotting only happens 4-directionally.

Example 3:

Input: grid = [[0,2]]

Output: 0

Explanation: Since there are already no fresh oranges at minute 0, the answer is just 0.

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$
- $1 \leq m, n \leq 10$
- $\text{grid}[i][j]$ is 0, 1, or 2.

Community Solutions (Python)

- LeetDoocs

class Solution:

def orangesRotting(self, grid:

List[List[int]]) -> int:

m, n = len(grid), len(grid[0])

cnt = 0

q = deque()

for i, row in enumerate(grid):

for j, x in enumerate(row):

if x == 2:

q.append((i, j))

elif x == 1:

cnt += 1

ans = 0

dirs = (-1, 0, 1, 0, -1)

while q and cnt:

ans += 1

for _ in range(len(q)):

i, j = q.popleft()

for a, b in pairwise(dirs):

x, y = i + a, j + b

if 0 <= x < m and 0 <= y < n

and grid[x][y] == 1:

grid[x][y] = 2

q.append((x, y))

cnt -= 1

if cnt == 0:

return ans

return -1 if cnt else 0

• SimplyLeet

```

from collections import deque

def orangesRotting(grid):
    # Dimensions of the grid
    rows, cols = len(grid), len(grid[0])
    # Queue to perform BFS, containing tuples of
    (row, col)
    rotten_queue = deque()
    fresh_count = 0

    # Step 1: Initialize the queue with all
    initially rotten oranges and count fresh oranges.

    for r in range(rows):
        for c in range(cols):
            if grid[r][c] == 2:
                rotten_queue.append((r, c))
            elif grid[r][c] == 1:
                fresh_count += 1

    # If there are no fresh oranges, return 0
    immediately.
    if fresh_count == 0:
        return 0

```

```

    minutes_passed = 0
    # Directions for adjacent cells: up, down,
    left, right.
    directions = [(-1,0), (1,0), (0,-1), (0,1)]

    # Step 2: BFS to rot adjacent fresh oranges.
    while rotten_queue:
        # Process oranges that are rotten in the
        current minute.
        for _ in range(len(rotten_queue)):
            x, y = rotten_queue.popleft()

            for dx, dy in directions:
                nx, ny = x + dx, y + dy
                # Check if the neighbor is inside
                the grid and is a fresh orange.
                if 0 <= nx < rows and 0 <= ny <
                cols and grid[nx][ny] == 1:
                    # Mark the fresh orange as
                    rotten.

                    grid[nx][ny] = 2
                    fresh_count -= 1
                    # Add the new rotten orange
                    into the queue.

                    rotten_queue.append((nx, ny))

        # Increment minutes after processing one
        level.

```

```

        minutes_passed += 1

    # If there are still fresh oranges left, return
    -1.
    return minutes_passed - 1 if fresh_count == 0
    else -1

# Example usage:
# grid = [[2,1,1],[1,1,0],[0,1,1]]
# print( orangesRotting(grid))

```

• LC.ca

```

class Solution:
    def orangesRotting(self, grid:
List[List[int]]) -> int:
        m, n = len(grid), len(grid[0])
        q = deque()
        cnt = 0
        for i in range(m):
            for j in range(n):
                if grid[i][j] == 2:
                    q.append((i, j))
                elif grid[i][j] == 1:

```

```

                                cnt += 1
        ans = 0
        while q and cnt:
            ans += 1
            for _ in range(len(q)):
                i, j = q.popleft()
                for a, b in [[0, 1], [0, -1], [1,
0], [-1, 0]]:
                    x, y = i + a, j + b
                    if 0 <= x < m and 0 <= y < n
and grid[x][y] == 1:
                        cnt -= 1

                                grid[x][y] = 2
                                q.append((x, y))
        return ans if cnt == 0 else -1

```

◆ Quick Review

Key Insights

- Use Breadth-First Search (BFS) starting from all initially rotten oranges.
- Process the grid level by level, where each level represents one minute.
- Keep a count of fresh oranges and decrement it as they become rotten.
- If after the BFS there are still fresh oranges, then it is impossible to rot every orange.

Space & Time

Time Complexity: $O(m * n)$ as every cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst case for the queue and auxiliary storage.

Solution

We begin by scanning the grid to locate all rotten oranges and count the fresh ones. These rotten oranges are added to a queue for a BFS. For each minute (BFS level), we process all the current rotten oranges and try to infect their adjacent fresh oranges. If an adjacent cell contains a fresh orange, we mark it rotten, add it to the queue, and decrement our fresh count. We continue this process level by level (minute by minute) until there are no fresh oranges left. If, after processing, there are still fresh oranges remaining, we return -1 , indicating that not all oranges can become rotten.

BFS

#1197 Minimum Knight Moves

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode


BFS

Lists: Grind 169 | Premium 98

Problem

Description

In an infinite chess board with coordinates from $-\infty$ to $+\infty$, you have a knight at square $[0, 0]$. A knight has 8 possible moves it can make, as illustrated below. Each move is two squares in a cardinal direction, then one square in an orthogonal direction.



Return the minimum number of steps needed to move the knight to the square $[x, y]$. It is guaranteed the answer exists.

Example 1:

Input: $x = 2, y = 1$
 Output: 1
 Explanation: $[0, 0] \rightarrow [2, 1]$

Example 2:

Input: $x = 5, y = 5$
 Output: 4
 Explanation: $[0, 0] \rightarrow [2, 1] \rightarrow [4, 2] \rightarrow [3, 4] \rightarrow [5, 5]$

Constraints:

- $-300 \leq x, y \leq 300$
- $0 \leq |x| + |y| \leq 300$

Community Solutions (Python)

• LeetCode

```
class Solution:
    def minKnightMoves(self, x: int, y: int) ->
    int:
        q = deque([(0, 0)])
        ans = 0
        vis = {(0, 0)}
        dirs = ((-2, 1), (-1, 2), (1, 2), (2, 1),
                (2, -1), (1, -2), (-1, -2), (-2, -1))
        while q:
            for _ in range(len(q)):
                i, j = q.popleft()
                if (i, j) == (x, y):
```

```

        return ans
    for a, b in dirs:
        c, d = i + a, j + b
        if (c, d) not in vis:
            vis.add((c, d))
            q.append((c, d))

    ans += 1
    return -1

```

```

class Solution:
    def minKnightMoves(self, x: int, y: int) ->
int:
    def extend(m1, m2, q):
        for _ in range(len(q)):
            i, j = q.popleft()
            step = m1[(i, j)]
            for a, b in (
                (-2, 1),
                (-1, 2),
                (1, 2),

```

```

                (2, 1),
                (2, -1),
                (1, -2),
                (-1, -2),
                (-2, -1),
            ):
                x, y = i + a, j + b
                if (x, y) in m1:
                    continue
                if (x, y) in m2:

```

```

                    return step + 1 + m2[(x,
y)]

                q.append((x, y))
                m1[(x, y)] = step + 1

            return -1

        if (x, y) == (0, 0):
            return 0
        q1, q2 = deque([(0, 0)]), deque([(x, y)])
        m1, m2 = {(0, 0): 0}, {(x, y): 0}
        while q1 and q2:

```

```

            t = extend(m1, m2, q1) if len(q1) <=
len(q2) else extend(m2, m1, q2)
            if t != -1:
                return t
        return -1

```

• [SimplyLeet](#)

```
# Python solution using BFS:
from collections import deque

def minKnightMoves(x, y):
    # Normalize target coordinates using symmetry
    x, y = abs(x), abs(y)

    # Possible knight moves (dx, dy)
    moves = [(2, 1), (1, 2), (-1, 2), (-2, 1), (-2, -1), (-1, -2), (1, -2), (2, -1)]
```

```
    # BFS queue stores tuples of the form
    (current_x, current_y, move_count)
    queue = deque([(0, 0, 0)])

    # Set to track visited positions to prevent
    cycles
    visited = {(0, 0)}

    # Process BFS
    while queue:
        curr_x, curr_y, move_count =
        queue.popleft()
```

```
        # Return the current move count if target
        is reached
        if curr_x == x and curr_y == y:
            return move_count

        # Explore all knight moves
        for dx, dy in moves:
            next_x, next_y = curr_x + dx, curr_y +
            dy

            # Prune search space: only process
            positions with valid bounds
            if (next_x, next_y) not in visited and
            next_x >= -1 and next_y >= -1:

                visited.add((next_x, next_y))
                queue.append((next_x, next_y,
                move_count + 1))

        return -1 # Unreachable because a solution
        always exists

    # Example test
    #print(minKnightMoves(5, 5))
```

• LC.ca

```

class Solution:
    def minKnightMoves(self, x: int, y: int) ->
int:
    q = deque([(0, 0)])
    ans = 0
    vis = {(0, 0)}
    dirs = ((-2, 1), (-1, 2), (1, 2), (2, 1),
(2, -1), (1, -2), (-1, -2), (-2, -1))
    while q:
        for _ in range(len(q)):
            i, j = q.popleft()
            if (i, j) == (x, y):
                return ans
            for a, b in dirs:
                c, d = i + a, j + b
                if (c, d) not in vis:
                    vis.add((c, d))
                    q.append((c, d))
            ans += 1
    return -1

```

◆ Quick Review

Key Insights

- Leverage the symmetry of the chessboard by converting target coordinates to their absolute values.
- Use a Breadth-First Search (BFS) because all moves have equal weight, ensuring the first encounter with the target is the minimal path.

Round 2 · High · Medium

p.1383

- Prune the search space by only considering positions that are likely beneficial (e.g., positions having coordinates not far less than -1).
- Track visited states to avoid redundant processing.
- Optimize by bounding the search region relative to the target distance.

Space & Time

Time Complexity: $O(n)$ where n is the number of positions processed during the BFS

Space Complexity: $O(n)$ due to the storage requirements for the BFS queue and the visited set

Solution

We solve the problem using a BFS starting from the origin. First, normalize the target by taking the absolute value of x and y . BFS is then applied by exploring all eight possible knight moves at each step. Each move generates a new position that is enqueued if it hasn't been visited and falls within a reasonable bound (e.g., coordinates ≥ -1). The search terminates as soon as the target position is reached.

Round 2 · High · Medium

p.1384

BFS

#1730 Shortest Path to Get Food

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Grind 169

Problem

Description

You are starving and you want to eat food as quickly as possible. You want to find the shortest path to arrive at any food cell.

You are given an $m \times n$ character matrix, `grid`, of these different types of cells:

- '*' is your location. There is exactly one '*' cell.
- '#' is a food cell. There may be multiple food cells.
- 'O' is free space, and you can travel through these cells.
- 'X' is an obstacle, and you cannot travel through these cells.


You can travel to any adjacent cell north, east, south, or west of your current location if there is not an obstacle.

Return the length of the shortest path for you to reach any food cell. If there is no path for you to reach food, return -1.

Round 2 · High · Medium

p.1385

Example 1:



Input: `grid = [ ["X","X","X","X","X","X"], ["X","*","O","O","O","X"], ["X","O","O","#","O","X"], ["X","X","X","X","X","X"] ]`


Output: 3

Explanation: It takes 3 steps to reach the food.

Example 2:


Round 2 · High · Medium

p.1386



Input: grid = [["X","X","X","X","X"],["X","*","X","0","X"],["X","0","X","#","X"],["X","X","X","X","X"]]
 Output: -1
 Explanation: It is not possible to reach the food.

Example 3:



Input: grid = [["X","X","X","X","X","X","X","X"],["X","*","0","X","0","#","0","X"],["X","0","0","X","0","0","X","X"],["X","0","0","0","0","0","#","0","X"],["X","X","X","X","X","X","X","X"]]
 Output: 6
 Explanation: There can be multiple food cells. It only takes 6 steps to reach the bottom food.

Example 4:

Input: grid = [["X","X","X","X","X","X","X","X"],["X","*","0","X","0","#","0","X"],["X","0","0","X","0","0","X","X"],["X","0","0","0","0","0","#","0","X"],["0","0","0","0","0","0","0","0"]]
 Output: 5

Constraints:

- `m == grid.length`
- `n == grid[i].length`
- `1 <= m, n <= 200`
- `grid[row][col]` is `'*'`, `'X'`, `'O'`, or `'#'`.
- The grid contains exactly one `'*'`.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def getFood(self, grid: list[list[str]]) -> int:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(grid)
        n = len(grid[0])
        ans = 0
        q =
        collections.deque([self._getStartLocation(grid)])

        while q:
            for _ in range(len(q)):
```

```
        i, j = q.popleft()
        for dx, dy in DIRS:
            x = i + dx
            y = j + dy
            if x < 0 or x == m or y < 0 or y == n:
                continue
            if grid[x][y] == 'X':
                continue
            if grid[x][y] == '#':
                return ans + 1
```

```
        q.append((x, y))
        grid[x][y] = 'X' # Mark as visited.
        ans += 1
```

```
        return -1
```

```
    def _getStartLocation(self, grid:
list[list[str]]) -> tuple[int, int]:
        for i, row in enumerate(grid):
            for j, cell in enumerate(row):
                if cell == '*':
```

```
        return (i, j)
```

• LeetCode

```

class Solution:
    def getFood(self, grid: List[List[str]]) ->
    int:
        m, n = len(grid), len(grid[0])
        i, j = next((i, j) for i in range(m) for j
in range(n) if grid[i][j] == '*')
        q = deque([(i, j)])
        dirs = (-1, 0, 1, 0, -1)
        ans = 0
        while q:
            ans += 1
            for _ in range(len(q)):
                i, j = q.popleft()
                for a, b in pairwise(dirs):
                    x, y = i + a, j + b
                    if 0 <= x < m and 0 <= y < n:
                        if grid[x][y] == '#':
                            return ans
                        if grid[x][y] == 'O':
                            grid[x][y] = 'X'
                            q.append((x, y))
        return -1

```

• SimplyLeet

```

from collections import deque

def getFood(grid):
    # number of rows and columns
    rows = len(grid)
    cols = len(grid[0])
    # finding the start position
    start = None
    for r in range(rows):
        for c in range(cols):
            if grid[r][c] == '*':
                start = (r, c)
                break
            if start is not None:
                break

    # initialize queue for BFS, with (row, col,
steps)
    queue = deque([(start[0], start[1], 0)])
    # mark visited cells to avoid reprocessing
    visited = [[False] * cols for _ in range(rows)]

```

```

visited[start[0]][start[1]] = True

# possible moves: up, right, down, left
directions = [(-1, 0), (0, 1), (1, 0), (0, -1)]

# begin BFS
while queue:
    row, col, steps = queue.popleft()
    # if food is found, return current number
of steps
    if grid[row][col] == '#':

        return steps
    # check all adjacent cells
    for dr, dc in directions:
        newRow, newCol = row + dr, col + dc
        # check boundaries, obstacles, and
visited status
        if 0 <= newRow < rows and 0 <= newCol <
cols and not visited[newRow][newCol] and
grid[newRow][newCol] != 'X':
            visited[newRow][newCol] = True
            queue.append((newRow, newCol,
steps + 1))

# if no food is reachable return -1

```

```

return -1

# Example usage:
grid = [
    ["X","X","X","X","X","X"],
    ["X","*","0","0","0","X"],
    ["X","0","0","#","0","X"],
    ["X","X","X","X","X","X"]
]
print(getFood(grid)) # Expected output: 3

```

• LC.ca

```

class Solution:
    def getFood(self, grid: List[List[str]]) ->
int:
        m, n = len(grid), len(grid[0])
        i, j = next((i, j) for i in range(m) for j
in range(n) if grid[i][j] == '*')
        q = deque([(i, j)])
        dirs = (-1, 0, 1, 0, -1)
        ans = 0
        while q:
            ans += 1
            for _ in range(len(q)):

```

```

        i, j = q.popleft()
        for a, b in pairwise(dirs):
            x, y = i + a, j + b
            if 0 <= x < m and 0 <= y < n:
                if grid[x][y] == '#':
                    return ans
                if grid[x][y] == '0':
                    grid[x][y] = 'X'
                    q.append((x, y))

    return -1

```

◆ Quick Review

Key Insights

- Use Breadth-First Search (BFS) to explore the grid level by level.
- The BFS guarantees that when a food cell ('#') is reached, the current distance is the shortest.
- Maintain a visited structure to avoid re-processing the same cell.
- Consider boundary and obstacle ('X') conditions while traversing.

Space & Time

Time Complexity: $O(m * n)$, where m is the number of rows and n is the number of columns since each cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst case due to the queue and visited matrix.

Solution

We use Breadth-First Search (BFS) starting from the cell marked with '*'. The BFS uses a queue to explore neighbors in the four allowed directions (up, down, left, right). Each level of BFS represents one step in the path. As soon as a cell containing food ('#') is reached, the current step count is returned. We also maintain a visited matrix to ensure that cells are not revisited, which could otherwise result in cycles or redundant processing. This approach ensures that the first time we encounter food, we have found the shortest possible path.

Round 3 · High · Hard

Round 3

High Priority · Hard

23 questions · 8 patterns

Arrays & Hashing ☐ #41

Sliding Window ☐ #76

Binary Search ☐ #☐ #31 ☐ #41 ☐ #66 ☐ #71 ☐ #77 ☐ #1235

Matrix ☐ #1074

Trees & BST ☐ #12 ☐ #297

DFS ☐ #26 ☐ #32 ☐ #68 ☐ #103 ☐ #1192

Graph ☐ #30 ☐ #1153

Bit ☐ #12 ☐ #31 ☐ #77 ☐ #815

Round 2 · High · Medium

p.1397

Arrays & Hashing

Round 3 · High · Hard · 1 question

Round 3 · High · Hard

p.1398

Arrays & Hashing

#41 First Missing Positive

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table

Lists: Grind 169

Problem

Description

Given an unsorted integer array `nums`. Return the smallest positive integer that is not present in `nums`. You must implement an algorithm that runs in $O(n)$ time and uses $O(1)$ auxiliary space.

Example 1:

Input: `nums = [1,2,0]`
 Output: 3
 Explanation: The numbers in the range `[1,2]` are all in the array.

Example 2:

Input: `nums = [3,4,-1,1]`
 Output: 2
 Explanation: 1 is in the array but 2 is missing.

Example 3:

Input: `nums = [7,8,9,11,12]`
 Output: 1
 Explanation: The smallest positive integer 1 is missing.

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $-231 \leq \text{nums}[i] \leq 231 - 1$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def firstMissingPositive(self, nums: list[int])
    -> int:
        n = len(nums)

        # Correct slot:
        # nums[i] = i + 1
        # nums[i] - 1 = i
        # nums[nums[i] - 1] = nums[i]
        for i in range(n):
            while nums[i] > 0 and nums[i] <= n and
            nums[nums[i] - 1] != nums[i]:
```



```

    nums[nums[i] - 1], nums[i] = nums[i],
    nums[nums[i] - 1]

```

```

    for i, num in enumerate(nums):
        if num != i + 1:
            return i + 1

```

```

    return n + 1

```

- LeetCode

```

class Solution:
    def firstMissingPositive(self, nums:
List[int]) -> int:
        n = len(nums)
        for i in range(n):
            while 1 <= nums[i] <= n and nums[i] !=
nums[nums[i] - 1]:
                j = nums[i] - 1
                nums[i], nums[j] = nums[j],
nums[i]
            for i in range(n):
                if nums[i] != i + 1:
                    return i + 1

```

```

    return n + 1

```

- SimplyLeet

```

# Python implementation of First Missing Positive

def firstMissingPositive(nums):
    n = len(nums)
    # Rearrange numbers to their correct positions
    for i in range(n):
        while 1 <= nums[i] <= n and nums[nums[i] -
1] != nums[i]:
            correct_index = nums[i] - 1
            nums[i], nums[correct_index] =
nums[correct_index], nums[i]

```

```

# Identify the first missing positive integer
for i in range(n):
    if nums[i] != i + 1:
        return i + 1
return n + 1

```

```

# Example usage:
print(firstMissingPositive([3, 4, -1, 1])) #
Output should be 2

```

- LC.ca

```

class Solution:
    def firstMissingPositive(self, nums:
List[int]) -> int:
        def swap(i, j):
            nums[i], nums[j] = nums[j], nums[i]

        n = len(nums)
        for i in range(n):
            while 1 <= nums[i] <= n and nums[i] !=
nums[nums[i] - 1]:
                swap(i, nums[i] - 1)
        for i in range(n):
            if i + 1 != nums[i]:
                return i + 1
        return n + 1

```

◆ Quick Review

Key Insights

- The missing positive number must be in the range $[1, n+1]$ where n is the length of the array.
- Use the array indices to place numbers in their correct positions (i.e., number i should be at index $i-1$) with in-place swapping.
- After rearrangement, scan the array: the first index where the number is not equal to $\text{index}+1$ indicates the missing positive integer.

Space & Time

Time Complexity: $O(n)$ as each number is swapped at most once.

Space Complexity: $O(1)$ since the rearrangement occurs in-place.

Solution

The approach involves reordering the array so that each positive integer in the range $[1, n]$ is placed at the index corresponding to its value minus one. This means for any valid number i , $\text{nums}[i-1]$ should be i . We iterate through the array and perform in-place swaps to achieve this order. After this process, another pass through the array reveals the first position where the number does not match the expected value ($i+1$), which is the smallest missing positive. If every position is correct, then the missing value is $n+1$.

Sliding Window

Round 3 · High · Hard · 1 question

Round 3 · High · Hard

p.1405

Sliding Window

#76 Minimum Window Substring

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Sliding Window

Lists: Grind 169

Problem

Description

Given two strings s and t of lengths m and n respectively, return the minimum window substring of s such that every character in t (including duplicates) is included in the window. If there is no such substring, return the empty string `""`.

The testcases will be generated such that the answer is unique.

Example 1:

Input: $s = \text{"ADOBECODEBANC"}, t = \text{"ABC"}$

Output: `"BANC"`

Explanation: The minimum window substring `"BANC"` includes `'A'`, `'B'`, and `'C'` from string t .

Example 2:

Input: $s = \text{"a"}, t = \text{"a"}$

Output: `"a"`

Explanation: The entire string s is the minimum window.

Round 3 · High · Hard

p.1406

Example 3:

Input: s = "a", t = "aa"

Output: ""

Explanation: Both 'a's from t must be included in the window.

Since the largest window of s only has one 'a', return empty string.

Constraints:

- m == s.length
- n == t.length
- 1 <= m, n <= 105
- s and t consist of uppercase and lowercase English letters.

Follow up: Could you find an algorithm that runs in $O(m + n)$ time?

Community Solutions (Python)

- WalkCC

```
class Solution:
    def minWindow(self, s: str, t: str) -> str:
        count = collections.Counter(t)
        required = len(t)
        bestLeft = -1
        minLength = len(s) + 1
```

```
    l = 0
    for r, c in enumerate(s):
        count[c] -= 1
```

```
        if count[c] >= 0:
            required -= 1
        while required == 0:
            if r - l + 1 < minLength:
                bestLeft = l
                minLength = r - l + 1
            count[s[l]] += 1
            if count[s[l]] > 0:
                required += 1
            l += 1
```

```
    return '' if bestLeft == -1 else s[bestLeft:
        bestLeft + minLength]
```

- LeetDoocs

```

class Solution:
    def minWindow(self, s: str, t: str) -> str:
        need = Counter(t)
        window = Counter()
        cnt = l = 0
        k, mi = -1, inf
        for r, c in enumerate(s):
            window[c] += 1
            if need[c] >= window[c]:
                cnt += 1

        while cnt == len(t):
            if r - l + 1 < mi:
                mi = r - l + 1
                k = l
            if need[s[l]] >= window[s[l]]:
                cnt -= 1
            window[s[l]] -= 1
            l += 1
        return "" if k < 0 else s[k : k + mi]

```

• SimplyLeet

Round 3 · High · Hard

p.1409

Python implementation of Minimum Window Substring solution

```

def minWindow(s: str, t: str) -> str:
    # If t is longer than s, no solution exists.
    if len(t) > len(s):
        return ""

    # Dictionary to store the frequency of each
    # character in t.
    t_freq = {}
    for char in t:
        t_freq[char] = t_freq.get(char, 0) + 1

    # Dictionary to keep track of the characters in
    # the current window.
    window_freq = {}

    required = len(t_freq) # Unique characters in t
    # to be matched
    formed = 0 # Number of unique characters in the
    # current window that meet required frequency

    # (window_length, left, right)
    answer = (float('inf'), None, None)

```

Round 3 · High · Hard

p.1410

```

l = 0 # Left pointer
# Start sliding the window with the right
pointer.
for r, char in enumerate(s):
    # Add the current character to the window
    counter.
    window_freq[char] = window_freq.get(char,
0) + 1

    # Check if the current character's
    frequency in the window matches that in t.
    if char in t_freq and window_freq[char] ==
t_freq[char]:
        formed += 1

```

```

# Try and contract the window if all
characters are matched.
while l <= r and formed == required:
    character = s[l]

    # Update answer if the current window
    is smaller than the previous best.
    if r - l + 1 < answer[0]:
        answer = (r - l + 1, l, r)

    # The character at the left is going to
    be removed from the window.

```

```

        window_freq[character] -= 1
        if character in t_freq and
window_freq[character] < t_freq[character]:
            formed -= 1

        # Move the left pointer ahead.
        l += 1

    # Return the smallest window, if found.
    return "" if answer[0] == float('inf') else
s[answer[1]:answer[2]+1]

# Example usage:
if __name__ == "__main__":
    s = "ADOBECODEBANC"
    t = "ABC"
    print(minWindow(s, t)) # Expected output:
"BANC"

```

• LC.ca

```

...
>>> deq = collections.deque([])
>>> deq.append(11)
>>> deq.append(22)
>>> deq.append(33)
>>>
>>> deq[0]
11
...

```

```

import collections

class Solution:
    def minWindow(self, s: str, t: str) -> str:
        cnt = 0
        need = collections.Counter(t)
        start, end = len(s), 3 * len(s) # Arbitrary
        3, just make sure end-start is larger than input s
        length, for later min-check
        d = {}
        # Using queue to store indexes, good for a
        large amount of API calls
        deq = collections.deque([])

        for i, c in enumerate(s):
            if c in need:
                deq.append(i)
                d[c] = d.get(c, 0) + 1
                # '=' also +1, because it's
                increased already one line above :)
                if d[c] <= need[c]:
                    cnt += 1
                while deq and d[s[deq[0]]] >
                need[s[deq[0]]]:
                    d[s[deq.popleft()]] -= 1
                    if cnt == len(t) and deq[-1] -
                    deq[0] < end - start:

```

Round 3 · High · Hard

p.1413

```

        start, end = deq[0], deq[-1]
        return s[start:end + 1]

#####

from collections import Counter

class Solution:
    def minWindow(self, s: str, t: str) -> str:
        ans = ''

        m, n = len(s), len(t)
        if m < n:
            return ans
        need = Counter(t)
        window = Counter()
        i, cnt, mi = 0, 0, inf
        for j, c in enumerate(s):
            window[c] += 1
            if need[c] >= window[c]: # >= , because
            1 line above already +=1
                cnt += 1

```

Round 3 · High · Hard

p.1414

```

while cnt == n:
    if j - i + 1 < mi: # in while
        mi = j - i + 1
        ans = s[i : j + 1]
    c = s[j]
    if need[c] >= window[c]: # char in
window but not in need, need[c]=0, window[c]=1..2..
        cnt -= 1
        window[c] -= 1
    j += 1
return ans

```

◆ Quick Review

Key Insights

- Use a sliding window approach to dynamically adjust the window boundaries.
- Utilize a hash table (or dictionary) to keep track of the frequency of characters required (from t) and the frequency within the current window.
- Expand the window until it satisfies the requirements, then contract to minimize the window.
- Efficiently update counts and check valid window with two pointers (left and right).

Space & Time

Time Complexity: $O(m + n)$, where m is the length of s and n is the length of t .

Space Complexity: $O(m + n)$ in the worst case due to the storage for the hash maps.

Solution

We use a sliding window technique with two pointers representing the window's boundaries. First, create a frequency map of characters from t . Then, expand the right pointer to include characters in the window until the window contains all required characters (meeting or exceeding the counts in t). Once the requirement is met, try contracting the window by moving the left pointer to reduce its size while still maintaining a valid window. Throughout the process, if a valid window is found that is smaller than the previously recorded minimum window, update the minimum. The solution leans heavily on hash tables to track counts and compares them against the required counts from t .

Binary Search

Round 3 · High · Hard · 7 questions

Round 3 · High · Hard

p.1417

Binary Search

#4 Median of Two Sorted Arrays

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Divide and Conquer

Lists: Grind 169

Problem

Description

Given two sorted arrays `nums1` and `nums2` of size `m` and `n` respectively, return the median of the two sorted arrays.

The overall run time complexity should be $O(\log(m+n))$.

Example 1:

Input: `nums1 = [1,3]`, `nums2 = [2]`

Output: `2.00000`

Explanation: merged array = `[1,2,3]` and median is 2.

Example 2:

Input: `nums1 = [1,2]`, `nums2 = [3,4]`

Output: `2.50000`

Explanation: merged array = `[1,2,3,4]` and median is $(2 + 3) / 2 = 2.5$.

Constraints:

- `nums1.length == m`
- `nums2.length == n`

Round 3 · High · Hard

p.1418

- $0 \leq m \leq 1000$
- $0 \leq n \leq 1000$
- $1 \leq m + n \leq 2000$
- $-106 \leq \text{nums1}[i], \text{nums2}[i] \leq 106$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findMedianSortedArrays(self, nums1:
list[int], nums2: list[int]) -> float:
        n1 = len(nums1)
        n2 = len(nums2)
        if n1 > n2:
            return self.findMedianSortedArrays(nums2,
nums1)

        l = 0
        r = n1
```

```
while l <= r:
    partition1 = (l + r) // 2
    partition2 = (n1 + n2 + 1) // 2 - partition1
    maxLeft1 = -2**31 if partition1 == 0 else
nums1[partition1 - 1]
    maxLeft2 = -2**31 if partition2 == 0 else
nums2[partition2 - 1]
    minRight1 = 2**31 - 1 if partition1 == n1
else nums1[partition1]
    minRight2 = 2**31 - 1 if partition2 == n2
else nums2[partition2]
    if maxLeft1 <= minRight2 and maxLeft2 <=
minRight1:
        return (max(maxLeft1, maxLeft2) +
min(minRight1, minRight2)) * 0.5 if (n1 + n2) % 2
== 0 else max(maxLeft1, maxLeft2)
    elif maxLeft1 > minRight2:

        r = partition1 - 1
    else:
        l = partition1 + 1
```

• LeetCode

```

class Solution:
    def findMedianSortedArrays(self, nums1:
List[int], nums2: List[int]) -> float:
        def f(i: int, j: int, k: int) -> int:
            if i >= m:
                return nums2[j + k - 1]
            if j >= n:
                return nums1[i + k - 1]
            if k == 1:
                return min(nums1[i], nums2[j])
            p = k // 2
            x = nums1[i + p - 1] if i + p - 1 < m
        else inf
            y = nums2[j + p - 1] if j + p - 1 < n
        else inf
            return f(i + p, j, k - p) if x < y else
f(i, j + p, k - p)

        m, n = len(nums1), len(nums2)
        a = f(0, 0, (m + n + 1) // 2)
        b = f(0, 0, (m + n + 2) // 2)
        return (a + b) / 2

```

• SimplyLeet

Round 3 · High · Hard

p.1421

```

# Function to find the median of two sorted arrays
def findMedianSortedArrays(nums1, nums2):
    # Ensure nums1 is the smaller array for faster
    binary search
    if len(nums1) > len(nums2):
        nums1, nums2 = nums2, nums1

    m, n = len(nums1), len(nums2)
    low, high = 0, m

    while low <= high:

        # Partition indexes for nums1 and nums2
        partition_nums1 = (low + high) // 2
        partition_nums2 = (m + n + 1) // 2 -
partition_nums1

        # If partition is 0 then there are no
        elements on left side. Use -infinity for
        comparison.
        maxLeftNums1 = float('-inf') if
partition_nums1 == 0 else nums1[partition_nums1 -
1]

        # If partition is length then there are no
        elements on right side. Use infinity for
        comparison.
        minRightNums1 = float('inf') if
partition_nums1 == m else nums1[partition_nums1]

        # Similarly for nums2

```

Round 3 · High · Hard

p.1422

```

        maxLeftNums2 = float('-inf') if
partition_nums2 == 0 else nums2[partition_nums2 -
1]
        minRightNums2 = float('inf') if
partition_nums2 == n else nums2[partition_nums2]

        # Check if we have found the correct
partition
        if maxLeftNums1 <= minRightNums2 and
maxLeftNums2 <= minRightNums1:
            # If total number of elements is even
            if (m + n) % 2 == 0:
                return (max(maxLeftNums1,
maxLeftNums2) + min(minRightNums1, minRightNums2))
/ 2.0
            else:
                return float(max(maxLeftNums1,
maxLeftNums2))

```

```

# If not, adjust binary search range
elif maxLeftNums1 > minRightNums2:
    # Move partition in nums1 to left
    high = partition_nums1 - 1
else:
    # Move partition in nums1 to right
    low = partition_nums1 + 1

```

```

# Example usage:
print(findMedianSortedArrays([1,3], [2])) #
Output: 2.0

```

Round 3 · High · Hard

p.1423

```

print(findMedianSortedArrays([1,2], [3,4])) #
Output: 2.5

```

• LC.ca

```

class Solution:
    def findMedianSortedArrays(self, nums1:
List[int], nums2: List[int]) -> float:

        def findKth(i: int, j: int, k: int) ->
float:
            if i >= m:
                return nums2[j + k - 1]
            if j >= n:
                return nums1[i + k - 1]
            if k == 1:
                return min(nums1[i], nums2[j])

            # Division a // b : floordiv(a, b)
            midVal1 = nums1[i + k // 2 - 1] if i +
k // 2 - 1 < m else math.inf
            midVal2 = nums2[j + k // 2 - 1] if j +
k // 2 - 1 < n else math.inf

            if midVal1 < midVal2:
                return findKth(i + k // 2, j, k - k
// 2) # '+' or '-' k//2
            else:
                return findKth(i, j + k // 2, k - k
// 2)

```

Round 3 · High · Hard

p.1424

```

    m = len(nums1)
    n = len(nums2)
    # Division a // b : floordiv(a, b)
    left = (m + n + 1) // 2
    right = (m + n + 2) // 2
    return (findKth(0, 0, left) + findKth(0, 0,
right)) / 2.0

```

```
#####
```

```
# iteration version
```

```

class Solution(object):
    def findMedianSortedArrays(self, nums1, nums2):
        a, b = sorted((nums1, nums2), key=len)
        m, n = len(a), len(b)
        after = (m + n - 1) / 2
        lo, hi = 0, m
        while lo < hi:
            i = (lo + hi) / 2
            if after - i - 1 < 0 or a[i] >= b[after - i -
1]:
                hi = i

```

```

        else:
            lo = i + 1
            i = lo
            nextfew = sorted(a[i:i + 2] + b[after - i:after
- i + 2])
            return (nextfew[0] + nextfew[1 - (m + n) % 2])
/ 2.0

```

♦ Quick Review

Key Insights

- The median divides the sorted array into two equal halves.
- We can perform binary search on the smaller array to partition both arrays correctly.
- We need to ensure that every element in the left partition is less than or equal to every element in the right partition.
- Edge cases must be handled when partitions are at the boundaries of either array.

Space & Time

Time Complexity: $O(\log(\min(m, n)))$

Space Complexity: $O(1)$

Solution

The solution uses a binary search technique on the smaller array. We define a partition index i for `nums1` and compute the corresponding index j for `nums2` such that $i + j$ equals half of the total number of elements. The algorithm checks if the largest element on the left side of `nums1` (if any) is

less than or equal to the smallest element on the right side of `nums2`, and vice versa for `nums2`. If the condition holds, we have found the correct partition and compute the median. If not, adjust the binary search range until the correct partition is found. This partitioning strategy allows us to find the median in logarithmic time.

Round 3 · High · Hard

p.1427

Binary Search

#315 Count of Smaller Numbers After Self

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Divide and Conquer · Binary Indexed Tree · Segment Tree · Merge Sort · Ordered Set

Lists: Denny Zhang

Problem

Description

Given an integer array `nums`, return an integer array `counts` where `counts[i]` is the number of smaller elements to the right of `nums[i]`.

Example 1:

Input: `nums = [5,2,6,1]`

Output: `[2,1,1,0]`

Explanation:

To the right of 5 there are 2 smaller elements (2 and 1).

To the right of 2 there is only 1 smaller element (1).

To the right of 6 there is 1 smaller element (1).

To the right of 1 there is 0 smaller element.

Example 2:

Input: `nums = [-1]`

Output: `[0]`

Round 3 · High · Hard

p.1428

Example 3:**Input:** nums = [-1,-1]**Output:** [0,0]**Constraints:**

- $1 \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$

Community Solutions (Python)

• WalkCC

```
class FenwickTree:
    def __init__(self, n: int):
        self.sums = [0] * (n + 1)

    def add(self, i: int, delta: int) -> None:
        while i < len(self.sums):
            self.sums[i] += delta
            i += FenwickTree.lowbit(i)

    def get(self, i: int) -> int:
```

```
        summ = 0
        while i > 0:
            summ += self.sums[i]
            i -= FenwickTree.lowbit(i)
        return summ
```

```
@staticmethod
def lowbit(i: int) -> int:
    return i & -i
```

```
class Solution:
    def countSmaller(self, nums: list[int]) -> list[int]:
        ans = []
        ranks = self._getRanks(nums)
        tree = FenwickTree(len(ranks))

        for num in reversed(nums):
            ans.append(tree.get(ranks[num] - 1))
            tree.add(ranks[num], 1)
```

```

    return ans[::-1]

    def _getRanks(self, nums: list[int]) ->
    dict[int, int]:
        ranks = collections.Counter()
        rank = 0
        for num in sorted(set(nums)):
            rank += 1
            ranks[num] = rank
        return ranks

```

```

from dataclasses import dataclass

```

```

@dataclass
class Item:
    num: int = 0
    index: int = 0

```

```

class Solution:

```

```

    def countSmaller(self, nums: list[int]) ->
    list[int]:
        n = len(nums)
        ans = [0] * n
        items = [Item(num, i) for i, num in
        enumerate(nums)]

```

```

        self._mergeSort(items, 0, n - 1, ans)
        return ans

```

```

    def _mergeSort(
        self,

```

```

        items: list[Item],
        l: int,
        r: int,
        ans: list[int],
    ) -> None:
        if l >= r:
            return

```

```

        m = (l + r) // 2
        self._mergeSort(items, l, m, ans)

```



```
self._mergeSort(items, m + 1, r, ans)
self._merge(items, l, m, r, ans)
```

```
def _merge(
    self,
    items: list[Item],
    l: int,
    m: int,
    r: int,
    ans: list[int],
```

```
) -> None:
    sorted = [Item()] * (r - l + 1)
    k = 0 # sorted's index
    i = l # left's index
    j = m + 1 # right's index
    rightCount = 0 # the number of numbers <
items[i].num

    while i <= m and j <= r:
        if items[i].num > items[j].num:
            rightCount += 1
```

```
        sorted[k] = items[j]
        k += 1
        j += 1
    else:
        ans[items[i].index] += rightCount
        sorted[k] = items[i]
        k += 1
        i += 1
```

Put the possible remaining left part into the sorted array.

```
while i <= m:
    ans[items[i].index] += rightCount
    sorted[k] = items[i]
    k += 1
    i += 1
```

Put the possible remaining right part into the sorted array.

```
while j <= r:
    sorted[k] = items[j]
    k += 1
```

```
j += 1
```

```
items[l:l + len(sorted)] = sorted
```

• LeetDoocs

```

class BinaryIndexedTree:
    def __init__(self, n):
        self.n = n
        self.c = [0] * (n + 1)

    @staticmethod
    def lowbit(x):
        return x & -x

    def update(self, x, delta):
        while x <= self.n:
            self.c[x] += delta
            x += BinaryIndexedTree.lowbit(x)

    def query(self, x):
        s = 0
        while x > 0:
            s += self.c[x]
            x -= BinaryIndexedTree.lowbit(x)
        return s

```

Round 3 · High · Hard

p.1435

```

class Solution:
    def countSmaller(self, nums: List[int]) -> List[int]:
        alls = sorted(set(nums))
        m = {v: i for i, v in enumerate(alls, 1)}
        tree = BinaryIndexedTree(len(m))
        ans = []
        for v in nums[::-1]:
            x = m[v]
            tree.update(x, 1)
            ans.append(tree.query(x - 1))
        return ans[::-1]

class Node:
    def __init__(self):
        self.l = 0
        self.r = 0
        self.v = 0

class SegmentTree:
    def __init__(self, n):
        self.tr = [Node() for _ in range(n << 2)]

```

Round 3 · High · Hard

p.1436

```
self.build(1, 1, n)
```

```
def build(self, u, l, r):
    self.tr[u].l = l
    self.tr[u].r = r
    if l == r:
        return
    mid = (l + r) >> 1
    self.build(u << 1, l, mid)
    self.build(u << 1 | 1, mid + 1, r)
```

```
def modify(self, u, x, v):
    if self.tr[u].l == x and self.tr[u].r == x:
        self.tr[u].v += v
        return
    mid = (self.tr[u].l + self.tr[u].r) >> 1
    if x <= mid:
        self.modify(u << 1, x, v)
    else:
        self.modify(u << 1 | 1, x, v)
```

Round 3 · High · Hard

p.1437

```
self.pushup(u)
```

```
def query(self, u, l, r):
    if self.tr[u].l >= l and self.tr[u].r <= r:
        return self.tr[u].v
    mid = (self.tr[u].l + self.tr[u].r) >> 1
    v = 0
    if l <= mid:
        v += self.query(u << 1, l, r)
    if r > mid:
```

```
        v += self.query(u << 1 | 1, l, r)
    return v
```

```
def pushup(self, u):
    self.tr[u].v = self.tr[u << 1].v +
self.tr[u << 1 | 1].v
```

```
class Solution:
    def countSmaller(self, nums: List[int]) ->
List[int]:
        s = sorted(set(nums))
```

Round 3 · High · Hard

p.1438

```

m = {v: i for i, v in enumerate(s, 1)}
tree = SegmentTree(len(s))
ans = []
for v in nums[::-1]:
    x = m[v]
    ans.append(tree.query(1, 1, x - 1))
    tree.modify(1, x, 1)
return ans[::-1]

```

```

class Solution:
    def countSmaller(self, nums: list[int]) -> list[int]:
        self.right_smaller_counts = [0] * len(nums)

        nums_indices = [(num, idx) for idx, num in enumerate(nums)]
        self.merge_sort(nums_indices)

        return self.right_smaller_counts

    def combine_arrays(

```

```

self,
    left_nums_indices: list[tuple[int, int]],
    right_nums_indices: list[tuple[int, int]],
) -> list[tuple[int, int]]:
    merged_nums_indices: list[tuple[int, int]]
    = []
    left_idx, right_idx = 0, 0

    while left_idx < len(left_nums_indices)
    and right_idx < len(right_nums_indices):
        if left_nums_indices[left_idx][0] <=
        right_nums_indices[right_idx][0]:
            # Iterated left side element
            finalizes its right smaller count.

            left_num_idx =
            left_nums_indices[left_idx][1]
            self.right_smaller_counts[left_num_idx] += right_idx

            merged_nums_indices.append(left_nums_indices[left_idx])
            left_idx += 1
            continue

            merged_nums_indices.append(right_nums_indices[right_idx])
            right_idx += 1

```

```

        while left_idx < len(left_nums_indices):
            # Iterated left side element finalizes
            its right smaller count.
            left_num_idx =
            left_nums_indices[left_idx][1]
            self.right_smaller_counts[left_num_idx
            ] += len(right_nums_indices)

            merged_nums_indices.append(left_nums_i
            ndices[left_idx])
            left_idx += 1

        while right_idx < len(right_nums_indices):
            merged_nums_indices.append(right_nums_
            indices[right_idx])

            right_idx += 1

        return merged_nums_indices

    def merge_sort(self, nums_indices:
    list[tuple[int, int]]) -> list[tuple[int, int]]:
        if len(nums_indices) == 1:
            return nums_indices # Single element.

        split_idx = len(nums_indices) // 2

```

```

        left_nums_indices =
        self.merge_sort(nums_indices[:split_idx])
        right_nums_indices =
        self.merge_sort(nums_indices[split_idx:])

        return
        self.combine_arrays(left_nums_indices,
        right_nums_indices)

```

• SimplyLeet

```

# Python solution using merge sort approach

def countSmaller(nums):
    n = len(nums)
    # Initialize counts array for results
    counts = [0] * n
    # Pair each element with its original index
    enum = [(num, i) for i, num in enumerate(nums)]

    # Recursive function for merge sort with
    counting

```

```
def mergeSort(arr):
    if len(arr) <= 1:
        return arr
    # Split the array into two halves
    mid = len(arr) // 2
    left = mergeSort(arr[:mid])
    right = mergeSort(arr[mid:])
    merged = []
    l = r = 0
    # r_count counts number of elements taken
    from right array
```

```
    r_count = 0
    while l < len(left) and r < len(right):
        # If left element is less or equal, it
        means current right elements are smaller
        if left[l][0] <= right[r][0]:
            # update count for original index
            counts[left[l][1]] += r_count
            merged.append(left[l])
            l += 1
        else:
            # right element is smaller, so
            increment r_count and add it to merged
```

```
        merged.append(right[r])
        r += 1
        r_count += 1
    # Append remaining elements from left
    reduce count by accumulated right elements
    while l < len(left):
        counts[left[l][1]] += r_count
        merged.append(left[l])
        l += 1
    # Append any remaining elements from right
    side
    while r < len(right):
```

```
        merged.append(right[r])
        r += 1
    return merged

# Start the merge sort process
mergeSort(enum)
return counts

# Example test
print(countSmaller([5, 2, 6, 1])) # Expected
output: [2, 1, 1, 0]
```

• LC.ca

```
'''
```

bisect: maintaining a list in sorted order without having to sort the list after each insertion.
<https://docs.python.org/3/library/bisect.html>

bisect.bisect_left()

Locate the insertion point for *x* in *a* to maintain sorted order.

bisect.bisect_right() or **bisect.bisect()**

Similar to **bisect_left()**, but returns an insertion point which comes after (to the right of) any existing entries of *x* in *a*.

bisect.insort_left(a, x, lo=0, hi=len(a), *, key=None)

Insert *x* in *a* in sorted order.

Keep in mind that the $O(\log n)$ search is dominated by the slow $O(n)$ insertion step.

bisect.insort_right(a, x, lo=0, hi=len(a), *, key=None)

bisect.insort(a, x, lo=0, hi=len(a), *, key=None)

Similar to **insort_left()**, but inserting *x* in *a* after any existing entries of *x*.

```
>>> import bisect
>>> bisect.bisect_left([1,2,3], 2)
1
>>> bisect.bisect_right([1,2,3], 2)
2
```

```
>>> a = [1, 1, 1, 2, 3]
>>> bisect.insort_left(a, 1.0)
>>> a
```

```
[1.0, 1, 1, 1, 2, 3]
```

```
>>> a = [1, 1, 1, 2, 3]
>>> bisect.insort_right(a, 1.0)
>>> a
[1, 1, 1, 1.0, 2, 3]
```

```
>>> a = [1, 1, 1, 2, 3]
>>> bisect.insort(a, 1.0)
>>> a
```

```
[1, 1, 1, 1.0, 2, 3]
'''

import bisect

class Solution(object):
    def countSmaller(self, nums):
        """
        :type nums: List[int]
        :rtype: List[int]

        """
        ans = []
        bst = []
        for num in reversed(nums):
            idx = bisect.bisect_left(bst, num)
            ans.append(idx)
            bisect.insort(bst, num)
        return ans[::-1]

#####
```

Round 3 · High · Hard

p.1447

```
class BinaryIndexedTree:
    def __init__(self, n):
        self.n = n
        self.c = [0] * (n + 1)

    @staticmethod
    def lowbit(x):
        return x & -x

    def update(self, x, delta):
        while x <= self.n:
            self.c[x] += delta
            x += BinaryIndexedTree.lowbit(x)

    def query(self, x):
        s = 0
        while x > 0:
            s += self.c[x]
            x -= BinaryIndexedTree.lowbit(x)
```

Round 3 · High · Hard

p.1448


```
return s
```

```
class Solution:
    def countSmaller(self, nums: List[int]) ->
    List[int]:
        alls = sorted(set(nums))
        m = {v: i for i, v in enumerate(alls, 1)}
        tree = BinaryIndexedTree(len(m))
        ans = []
        for v in nums[::-1]:
            x = m[v]
            tree.update(x, 1)
            ans.append(tree.query(x - 1))
        return ans[::-1]
```

◆ Quick Review

Key Insights

- The brute force solution takes $O(n^2)$ time, which is too slow for large input arrays.
- An efficient approach leverages the concept of Merge Sort to both sort the array and count the number of smaller elements by tracking the number of elements that "jump over" during the merge step.
- Other efficient methods include Binary Indexed Trees (Fenwick Tree) or Balanced Binary Search Trees.

Space & Time

Round 3 · High · Hard

p.1449

Time Complexity: $O(n \log n)$ on average due to the merge sort based approach.

Space Complexity: $O(n)$ for auxiliary arrays used during sorting and counting.

Solution

We use a modified Merge Sort to sort the array while counting the number of smaller elements to the right. We keep track of the original indices by pairing each element with its index. During the merge process, when an element from the left subarray is placed into the temporary array, we add the count of elements from the right subarray that have already been placed (since these represent smaller elements that were to its right). This efficient divide and conquer strategy guarantees an overall $O(n \log n)$ time complexity. Care must be taken to update the counts based on indices and during the merge step.

Round 3 · High · Hard

p.1450

Binary Search

#410 Split Array Largest Sum

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Dynamic Programming · Greedy · Prefix Sum

Lists: Denny Zhang

Problem

Description

Given an integer array `nums` and an integer `k`, split `nums` into `k` non-empty subarrays such that the largest sum of any subarray is minimized.

Return the minimized largest sum of the split.

A subarray is a contiguous part of the array.

Example 1:

Input: `nums = [7,2,5,10,8]`, `k = 2`

Output: 18

Explanation: There are four ways to split `nums` into two subarrays.

The best way is to split it into `[7,2,5]` and `[10,8]`, where the largest sum among the two subarrays is only 18.

Example 2:

Input: `nums = [1,2,3,4,5]`, `k = 2`

Output: 9

Explanation: There are four ways to split `nums` into two subarrays.

The best way is to split it into `[1,2,3]` and `[4,5]`, where the largest sum among the two subarrays is only 9.

Constraints:

- `1 <= nums.length <= 1000`
- `0 <= nums[i] <= 106`
- `1 <= k <= min(50, nums.length)`

Community Solutions (Python)

• WalkCC

```
class Solution:
    def splitArray(self, nums: list[int], k: int) -> int:
        prefix = list(itertools.accumulate(nums,
            initial=0))

        @functools.lru_cache(None)
        def dp(i: int, k: int) -> int:
            """
            Returns the minimum of the maximum sum to
            split the first i numbers into
            k groups.
            """
```

```

    if k == 1:
        return prefix[i]
    return min(max(dp[j, k - 1], prefix[i] -
prefix[j])
               for j in range(k - 1, i))

    return dp(len(nums), k)

```

```

class Solution:
    def splitArray(self, nums: list[int], k: int) ->
int:
    n = len(nums)
    # dp[i][k] := the minimum of the maximum sum to
split the first i numbers
    # into k groups
    dp = [[math.inf] * (k + 1) for _ in range(n +
1)]
    prefix = list(itertools.accumulate(nums,
initial=0))

    for i in range(1, n + 1):
        dp[i][1] = prefix[i]

```

Round 3 · High · Hard

p.1453

```

    for l in range(2, k + 1):
        for i in range(l, n + 1):
            for j in range(l - 1, i):
                dp[i][l] = min(dp[i][l], max(dp[j][l -
1], prefix[i] - prefix[j]))

    return dp[n][k]

```

```

class Solution:
    def splitArray(self, nums: list[int], k: int) ->
int:
    l = max(nums)
    r = sum(nums) + 1

    def numGroups(maxSumInGroup: int) -> int:
        groupCount = 1
        sumInGroup = 0

        for num in nums:

```

Round 3 · High · Hard

p.1454

```

    if sumInGroup + num <= maxSumInGroup:
        sumInGroup += num
    else:
        groupCount += 1
        sumInGroup = num

    return groupCount

```

```

while l < r:
    m = (l + r) // 2

```

```

    if numGroups(m) > k:
        l = m + 1
    else:
        r = m

    return l

```

• LeetCode

```

class Solution:
    def splitArray(self, nums: List[int], k: int)
    -> int:
        def check(mx):
            s, cnt = inf, 0
            for x in nums:
                s += x
                if s > mx:
                    s = x
                    cnt += 1
            return cnt <= k

```

Round 3 · High · Hard

p.1455

```

    left, right = max(nums), sum(nums)
    return left + bisect_left(range(left,
right + 1), True, key=check)

```

• SimplyLeet

Python solution for Split Array Largest Sum

```

def splitArray(nums, k):
    # Helper function to determine if we can split
    array into <= k subarrays with max subarray sum <=
    max_sum
    def canSplit(max_sum):
        subarray_count = 1 # start with one
        subarray
        current_sum = 0
        for num in nums:
            if current_sum + num > max_sum:
                subarray_count += 1 # start a new
                subarray

```

Round 3 · High · Hard

p.1456

```

        current_sum = num # reset the sum
with the current number
    else:
        current_sum += num # add num to
current subarray
        return subarray_count <= k # valid if we
used k or fewer subarrays

    low, high = max(nums), sum(nums) # boundaries
for binary search
    while low < high:
        mid = (low + high) // 2 # candidate maximum
subarray sum
        if canSplit(mid):
            high = mid # candidate is valid, try to
find a smaller one

```

```

    else:
        low = mid + 1 # candidate too small,
increase it
    return low

```

```

# Example usage:
nums = [7,2,5,10,8]
k = 2
print(splitArray(nums, k)) # Expected output: 18

```

• LC.ca

Round 3 · High · Hard

p.1457

```

class Solution:
    def splitArray(self, nums: List[int], k: int)
-> int:
        def check(mx):
            s, cnt = inf, 0
            for x in nums:
                s += x
                if s > mx:
                    s = x
                    cnt += 1
            return cnt <= k

        left, right = max(nums), sum(nums)
        return left + bisect_left(range(left,
right + 1), True, key=check)

```

◆ Quick Review

Key Insights

- Use binary search over the answer: the candidate value is the maximum subarray sum.
- The lower bound of the search is $\max(\text{nums})$ and the upper bound is $\text{sum}(\text{nums})$.
- A greedy approach is used to check if a candidate maximum sum allows splitting nums into at most k subarrays.
- Adjust the binary search bounds based on the feasibility check.

Round 3 · High · Hard

p.1458

Space & Time

Time Complexity: $O(n * \log(\text{sum}(\text{nums})))$ – n elements are processed for each binary search iteration.

Space Complexity: $O(1)$ – Constant extra space is used.

Solution

The solution employs binary search to determine the smallest maximum subarray sum possible. We set the initial lower bound to the maximum element and the upper bound to the total sum of the array. For each candidate value (mid), we check if the array can be split into k or fewer subarrays without any subarray sum exceeding mid using a greedy approach. If possible, we try to lower the candidate; if not, we increase the candidate. This continues until the optimal solution is found.

Binary Search

#668 Kth Smallest Number in Multiplication Table

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Binary Search

Lists: Denny Zhang

Problem

Description

Nearly everyone has used the Multiplication Table. The multiplication table of size $m \times n$ is an integer matrix mat where $\text{mat}[i][j] == i * j$ (1-indexed).

Given three integers m , n , and k , return the k th smallest element in the $m \times n$ multiplication table.

Example 1:

1	2	3
2	4	6
3	6	9

1	2	2	3	3	4	6	6	9
---	---	---	---	---	---	---	---	---

Input: $m = 3, n = 3, k = 5$
 Output: 3
 Explanation: The 5th smallest number is 3.

Example 2:

1	2	3
2	4	6

1	2	2	3	4	6
---	---	---	---	---	---

Round 3 · High · Hard

p.1461

Input: $m = 2, n = 3, k = 6$
 Output: 6
 Explanation: The 6th smallest number is 6.

Constraints:

- $1 \leq m, n \leq 3 \cdot 10^4$
- $1 \leq k \leq m \cdot n$

Community Solutions (Python)

• LeetCode

```
class Solution:
    def findKthNumber(self, m: int, n: int, k: int)
    -> int:
        left, right = 1, m * n
        while left < right:
            mid = (left + right) >> 1
            cnt = 0
            for i in range(1, m + 1):
                cnt += min(mid // i, n)
            if cnt >= k:
                right = mid
        else:
            left = mid + 1
        return left
```

• SimplyLeet

Round 3 · High · Hard

p.1462

Python solution using binary search

```
class Solution:
    def findKthNumber(self, m: int, n: int, k: int)
    -> int:
        def count_less_equal(x: int) -> int:
            # Count numbers in the m x n
            multiplication table <= x
            count = 0
            for i in range(1, m + 1):
                # For each row, count of numbers <=
                x is minimum of n or x // i
                count += min(x // i, n)
```

return count

```
# Binary search range from 1 to m*n
low, high = 1, m * n
while low < high:
    mid = (low + high) // 2
    if count_less_equal(mid) >= k:
        high = mid # mid might be the
        answer
    else:
        low = mid + 1 # answer must be
        greater than mid
```

return low

• LC.ca

Round 3 · High · Hard

p.1463

```
class Solution:
    def findKthNumber(self, m: int, n: int, k: int)
    -> int:
        left, right = 1, m * n
        while left < right:
            mid = (left + right) >> 1
            cnt = 0
            for i in range(1, m + 1):
                cnt += min(mid // i, n)
            if cnt >= k:
                right = mid
        else:
            left = mid + 1
        return left
```

◆ Quick Review

Key Insights

- The table elements are not fully sorted row-wise or column-wise globally, but there is a pattern that can be exploited.
- For any candidate number x , we can count how many numbers in the table are less than or equal to x by summing over rows: for row i , there are $\min(x // i, n)$ numbers.
- Binary search can be used over the range $[1, m \cdot n]$ to find the smallest x such that the count is at least k .

Round 3 · High · Hard

p.1464

- The counting function is the key to connecting the candidate value with the order statistic.

Space & Time

Time Complexity: $O(m * \log(m*n))$

Space Complexity: $O(1)$

Solution

We use binary search over the possible values in the multiplication table, which range from 1 to $m*n$. For each midpoint x in the binary search, we compute the number of elements in the table that are less than or equal to x by iterating over each row and summing up $\min(x // i, n)$. This count tells us whether the k th smallest element is to the left or right of x . We then adjust the binary search boundaries accordingly. The approach leverages the multiplicative structure of the table and the monotonicity of the counting function.

Binary Search

#710 Random Pick with Blacklist

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Math · Binary Search · Sorting · Randomized

Lists: Denny Zhang

Problem

Description

You are given an integer n and an array of unique integers `blacklist`. Design an algorithm to pick a random integer in the range $[0, n - 1]$ that is not in `blacklist`. Any integer that is in the mentioned range and not in `blacklist` should be equally likely to be returned.

Optimize your algorithm such that it minimizes the number of calls to the built-in random function of your language.

Implement the `Solution` class:

- `Solution(int n, int[] blacklist)` Initializes the object with the integer n and the blacklisted integers `blacklist`.
- `int pick()` Returns a random integer in the range $[0, n - 1]$ and not in `blacklist`.

Example 1:

```

Input
["Solution", "pick", "pick", "pick", "pick", "pick",
"pick", "pick"]
[[7, [2, 3, 5]], [], [], [], [], [], [], []]
Output
[null, 0, 4, 1, 6, 1, 0, 4]

```

Explanation

```
Solution solution = new Solution(7, [2, 3, 5]);
```

Constraints:

- $1 \leq n \leq 109$
- $0 \leq \text{blacklist.length} \leq \min(105, n - 1)$
- $0 \leq \text{blacklist}[i] < n$
- All the values of blacklist are unique.
- At most $2 * 10^4$ calls will be made to pick.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def __init__(self, n: int, blacklist:
list[int]):
        self.validRange = n - len(blacklist)
        self.dict = {}

        maxAvailable = n - 1

        for b in blacklist:
            self.dict[b] = -1

```

```

for b in blacklist:
    if b < self.validRange:
        # Find the slot that haven't been used.
        while maxAvailable in self.dict:
            maxAvailable -= 1
        self.dict[b] = maxAvailable
        maxAvailable -= 1

def pick(self) -> int:
    value = random.randint(0, self.validRange - 1)

    if value in self.dict:
        return self.dict[value]

    return value

```

• LeetCode

```

class Solution:
    def __init__(self, n: int, blacklist:
List[int]):
        self.k = n - len(blacklist)
        self.d = {}
        i = self.k
        black = set(blacklist)
        for b in blacklist:
            if b < self.k:
                while i in black:
                    i += 1

```

```

        self.d[b] = i
        i += 1

```

```

def pick(self) -> int:
    x = randrange(self.k)
    return self.d.get(x, x)

```

Your Solution object will be instantiated and called as such:

```
# obj = Solution(n, blacklist)
```

```
# param_1 = obj.pick()
```

• SimplyLeet

```
import random
```

```
class Solution:
```

```
    def __init__(self, n: int, blacklist: list):
```

```
        # Total numbers valid for picking
```

```
        self.W = n - len(blacklist)
```

```
        # Mapping for blacklisted numbers in the
        valid range
```

```
        self.mapping = {}
```

```
        # Set to hold blacklisted numbers that are
```

```
>= W
```

```
        blacklist_set = set(blacklist)
```

```

        # Pointer to find available numbers in [W,
        n-1]

```

```
        available = self.W
```

```
        for b in blacklist:
```

```
            if b < self.W:
```

```
                # Find next available
```

```
non-blacklisted number in the tail region.
```

```
                while available in blacklist_set:
```

```
                    available += 1
```

```
                self.mapping[b] = available
```

```
                available += 1
```

```
    def pick(self) -> int:
```

```
        # Generate a random number in the range [0,
        self.W-1]
```

```
        idx = random.randint(0, self.W - 1)
```

```
        # If the number is mapped, return the
        mapping; else it is valid.
```

```
        return self.mapping.get(idx, idx)
```

• LC.ca

```

class Solution:
    def __init__(self, n: int, blacklist:
List[int]):
        self.k = n - len(blacklist)
        self.d = {}
        i = self.k
        black = set(blacklist)
        for b in blacklist:
            if b < self.k:
                while i in black:
                    i += 1

                self.d[b] = i
                i += 1

    def pick(self) -> int:
        x = randrange(self.k)
        return self.d.get(x, x)

```

Your Solution object will be instantiated and called as such:

```
# obj = Solution(n, blacklist)
```

```
# param_1 = obj.pick()
```

◆ Quick Review

Key Insights

- The valid output range has a total count of available numbers $W = n - \text{len}(\text{blacklist})$.
- Instead of repeatedly generating random numbers and checking against a set, we can transform the range into a contiguous block $[0, W - 1]$ and map any blacklist number within that block to a valid number in the upper range.
- Mapping is precomputed: for any blacklisted number x in $[0, W-1]$, assign it a valid number from $[W, n-1]$ that is not blacklisted.
- This approach ensures $O(1)$ pick time after an $O(b)$ precomputation where b = length of blacklist.

Space & Time

Time Complexity: constructor – $O(b)$ for initializing data structures, pick() – $O(1)$ per call.

Space Complexity: $O(b)$ for storing the blacklist set and mapping.

Solution

We compute the count of valid numbers, $W = n - \text{len}(\text{blacklist})$. For numbers in the blacklist that are less than W , we need to map them to a valid number from the range $[W, n-1]$. To do this, we first store all blacklist numbers in the higher range $[W, n-1]$ in a set for quick lookup. Then for each blacklisted value in $[0, W-1]$, we iterate from W upwards to find a number that is not blacklisted and assign it in a mapping. During the pick() function, we generate a random integer r in $[0, W-1]$. If r is in our mapping, we return its corresponding mapped value; otherwise, we return r as it is already a valid number.

Binary Search

#774 Minimize Max Distance to Gas Station HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search

Lists: Denny Zhang

Problem

Description

You are given an integer array `stations` that represents the positions of the gas stations on the x-axis. You are also given an integer `k`.

You should add `k` new gas stations. You can add the stations anywhere on the x-axis, and not necessarily on an integer position.

Let `penalty()` be the maximum distance between adjacent gas stations after adding the `k` new stations.

Return the smallest possible value of `penalty()`. Answers within 10^{-6} of the actual answer will be accepted.

Example 1:

Input: `stations = [1,2,3,4,5,6,7,8,9,10]`, `k = 9`
 Output: `0.50000`

Example 2:

Input: `stations = [23,24,36,39,46,56,57,65,84,98]`, `k = 1`

Output: `14.00000`

Constraints:

- $10 \leq \text{stations.length} \leq 2000$
- $0 \leq \text{stations}[i] \leq 108$
- `stations` is sorted in a strictly increasing order.
- $1 \leq k \leq 106$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def minmaxGasDist(self, stations: list[int], k:
int) -> float:
        ERR = 1e-6
        l = 0
        r = stations[-1] - stations[0]

        def possible(k: int, m: float) -> bool:
            """
            Returns True if can use <= k gas stations to
            ensure that each adjacent
            distance between gas stations <= m.
```

```

"""
for a, b in zip(stations, stations[1:]):
    diff = b - a
    if diff > m:
        k -= math.ceil(diff / m) - 1
        if k < 0:
            return False
    return True

while r - l > ERR:

    m = (l + r) / 2
    if possible(k, m):
        r = m
    else:
        l = m

return l

```

• LeetCode

Round 3 · High · Hard

p.1475

```

class Solution:
    def minmaxGasDist(self, stations: List[int],
k: int) -> float:
        def check(x):
            return sum(int((b - a) / x) for a, b in
pairwise(stations)) <= k

        left, right = 0, 1e8
        while right - left > 1e-6:
            mid = (left + right) / 2
            if check(mid):
                right = mid

        else:
            left = mid
        return left

```

• SimplyLeet

```

import math

def minmaxGasDist(stations, k):
    # Initialize the search range between 0 and the
    maximum gap between existing stations
    low = 0.0
    high = max(stations[i+1] - stations[i] for i in
range(len(stations)-1))

    # Use binary search with a tolerance of 1e-6
    while high - low > 1e-6:
        mid = (low + high) / 2.0

```

Round 3 · High · Hard

p.1476

```

        used = 0
        # Calculate additional stations needed for
each gap
        for i in range(1, len(stations)):
            gap = stations[i] - stations[i-1]
            used += math.ceil(gap / mid) - 1
        # If the used stations are less than or
equal to k, try a smaller distance
        if used <= k:
            high = mid
        else:
            low = mid

```

```

return high

```

```

# Example usage:
if __name__ == '__main__':
    print("{:.5f}".format(minmaxGasDist([1,2,3,4,5
,6,7,8,9,10], 9)))
    print("{:.5f}".format(minmaxGasDist([23,24,36,
39,46,56,57,65,84,98], 1)))

```

• LC.ca

Round 3 · High · Hard

p.1477

```

class Solution:
    def minmaxGasDist(self, stations: List[int],
k: int) -> float:
        def check(x):
            return sum(int((b - a) / x) for a, b in
pairwise(stations)) <= k

        left, right = 0, 1e8
        while right - left > 1e-6:
            mid = (left + right) / 2
            if check(mid):
                right = mid

        else:
            left = mid
        return left

```

◆ Quick Review

Key Insights

- The optimal "penalty" (i.e., the maximum distance between adjacent stations after placement) can be found using binary search over a continuous distance range.
- For each candidate distance (mid), simulate the number of gas stations required by iterating over each adjacent station gap and calculating:
 - $\text{stations_required} = \text{ceil}(\text{gap} / \text{mid}) - 1$
- If the total number of required additional stations is less than or equal to k, then mid is a feasible candidate;

Round 3 · High · Hard

p.1478

otherwise, it is too small.

- Continue the binary search until the gap between low and high is within the desired tolerance ($1e-6$).

Space & Time

Time Complexity: $O(n * \log(\text{max_gap}/\text{precision}))$, where n is the number of stations.

Space Complexity: $O(1)$ (ignoring input storage)

Solution

We use binary search over the possible maximum distances. The algorithm sets $\text{low} = 0$ and $\text{high} =$ the maximum gap between adjacent stations. For a candidate distance mid , we count how many new stations are required for all gaps by computing $\text{ceil}(\text{gap} / \text{mid}) - 1$. If the total required is within k , we try a smaller candidate; otherwise, we increase the candidate distance. This approach efficiently zooms into the minimum possible maximum distance (penalty).

Binary Search

#1235 Maximum Profit in Job Scheduling

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Dynamic Programming · Sorting
Lists: Grind 169

Problem


Description

We have n jobs, where every job is scheduled to be done from $\text{startTime}[i]$ to $\text{endTime}[i]$, obtaining a profit of $\text{profit}[i]$.

You're given the startTime , endTime and profit arrays, return the maximum profit you can take such that there are no two jobs in the subset with overlapping time range.


If you choose a job that ends at time X you will be able to start another job that starts at time X .

Example 1:




Input: `startTime = [1,2,3,3], endTime = [3,4,5,6], profit = [50,10,40,70]`
 Output: 120
 Explanation: The subset chosen is the first and fourth job.
 Time range [1-3]+[3-6] , we get profit of 120 = 50 + 70.

Example 2:



Input: `startTime = [1,2,3,4,6], endTime = [3,5,10,6,9], profit = [20,20,100,70,60]`
 Output: 150
 Explanation: The subset chosen is the first, fourth and fifth job.
 Profit obtained 150 = 20 + 70 + 60.

Example 3:



Input: `startTime = [1,1,1], endTime = [2,3,4], profit = [5,6,4]`
 Output: 6

Constraints:

- $1 \leq \text{startTime.length} == \text{endTime.length} == \text{profit.length} \leq 5 \times 10^4$
- $1 \leq \text{startTime}[i] < \text{endTime}[i] \leq 10^9$
- $1 \leq \text{profit}[i] \leq 10^4$

Community Solutions (Python)

- [WalkCC](#)

```
class Solution:
    def jobScheduling(
        self,
        startTime: list[int],
        endTime: list[int],
        profit: list[int],
    ) -> int:
        jobs = sorted([(s, e, p) for s, e, p in
            zip(startTime, endTime, profit)])

        # Will use binary search to find the first
        # available startTime

        for i in range(len(startTime)):
            startTime[i] = jobs[i][0]

        @functools.lru_cache(None)
        def dp(i: int) -> int:
            """Returns the maximum profit to schedule
            jobs[i..n)."""
            if i == len(startTime):
                return 0
            j = bisect.bisect_left(startTime,
                jobs[i][1])
            return max(jobs[i][2] + dp(j), dp(i + 1))

        return dp(0)
```

• LeetCode

Round 3 · High · Hard

p.1483

```
class Solution:
    def jobScheduling(
        self, startTime: List[int], endTime:
        List[int], profit: List[int]
    ) -> int:
        @cache
        def dfs(i):
            if i >= n:
                return 0
            _, e, p = jobs[i]
            j = bisect_left(jobs, e, lo=i + 1,
                key=lambda x: x[0])

            return max(dfs(i + 1), p + dfs(j))

        jobs = sorted(zip(startTime, endTime,
            profit))
        n = len(profit)
        return dfs(0)
```

• SimplyLeet

Round 3 · High · Hard

p.1484

```
import bisect

def jobScheduling(startTime, endTime, profit):
    # Combine the jobs and sort them by end time.
    jobs = sorted(zip(startTime, endTime, profit),
key=lambda x: x[1])
```

```
    # dp_end_times stores the end times and
dp_profits stores the corresponding maximum
profit.
```

```
    dp_end_times = [0]
    dp_profits = [0]
```

```
    # Process each job in order.
    for s, e, p in jobs:
        # Find the index of the last job that
finishes before or at the start of the current job
using binary search.
```

```
        index = bisect.bisect_right(dp_end_times,
s)
        current_profit = dp_profits[index - 1] + p
    # Profit if we take current job.
```

```
    # If taking current job increases the
overall profit, update the dp arrays.
    if current_profit > dp_profits[-1]:
        dp_end_times.append(e)
        dp_profits.append(current_profit)
```

```
    return dp_profits[-1]
```

```
# Example usage:
```

```
print(jobScheduling([1,2,3,3], [3,4,5,6],
[50,10,40,70]))
```

• LC.ca

```
class Solution:
```

```
    def jobScheduling(
        self, startTime: List[int], endTime:
List[int], profit: List[int]
    ) -> int:
        @cache
        def dfs(i):
            if i >= n:
                return 0
            _, e, p = jobs[i]
            j = bisect_left(jobs, e, lo=i + 1,
key=lambda x: x[0])
```

```
        return max(dfs(i + 1), p + dfs(j))
```

```
        jobs = sorted(zip(startTime, endTime,
profit))
        n = len(profit)
        return dfs(0)
```

◆ Quick Review

Key Insights

- Sort jobs based on their end times to simplify finding non-conflicting jobs.
- Use dynamic programming where the state represents the maximum profit until a given point.
- For each job, perform a binary search to find the last job that ends before the current job's start time.
- The recurrence relation is: $dp[i] = \max(dp[i-1], profit[i] + dp[index])$ where $dp[index]$ is from the latest non-conflicting job.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting and binary search for each job.

Space Complexity: $O(n)$ for the dp arrays.

Solution

The solution begins by sorting the jobs by their end times. For each job, a binary search is used to quickly locate the last job that does not conflict with the current job (i.e., whose end time is less than or equal to the current job's start time). Then, using dynamic programming, the maximum profit is updated by choosing whether to include the current job or skip it. Two parallel arrays (or lists) are maintained: one for the end times of the jobs considered so far, and another to store the corresponding maximum profit at those times. This approach efficiently computes the answer while keeping time complexity within acceptable bounds.

Round 3 · High · Hard

p.1487

Matrix

Round 3 · High · Hard · 1 question

Round 3 · High · Hard

p.1488

Matrix

#1074 Number of Submatrices That Sum to Target HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Matrix · Prefix Sum

Lists: Denny Zhang

Problem

Description

Given a matrix and a target, return the number of non-empty submatrices that sum to target.

A submatrix $x1, y1, x2, y2$ is the set of all cells $matrix[x][y]$ with $x1 \leq x \leq x2$ and $y1 \leq y \leq y2$.

Two submatrices $(x1, y1, x2, y2)$ and $(x1', y1', x2', y2')$ are different if they have some coordinate that is different: for example, if $x1 \neq x1'$.

Example 1:

0	1	0
1	1	1
0	1	0

Input: matrix = $[[0,1,0],[1,1,1],[0,1,0]]$, target = 0

Output: 4

Explanation: The four 1x1 submatrices that only contain 0.

Example 2:

Input: matrix = $[[1,-1],[-1,1]]$, target = 0

Output: 5

Explanation: The two 1x2 submatrices, plus the two 2x1 submatrices, plus the 2x2 submatrix.

Example 3:

Input: matrix = [[904]], target = 0
 Output: 0

Constraints:

- $1 \leq \text{matrix.length} \leq 100$
- $1 \leq \text{matrix}[0].\text{length} \leq 100$
- $-1000 \leq \text{matrix}[i][j] \leq 1000$
- $-10^8 \leq \text{target} \leq 10^8$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numSubmatrixSumTarget(self, matrix:
list[list[int]], target: int) -> int:
        m = len(matrix)
        n = len(matrix[0])
        ans = 0

        # Transfer each row in the matrix to the prefix
        sum.
        for row in matrix:
            for i in range(1, n):
                row[i] += row[i - 1]
```

```
for baseCol in range(n):
    for j in range(baseCol, n):
        prefixCount = collections.Counter({0: 1})
        summ = 0
        for i in range(m):
            if baseCol > 0:
                summ -= matrix[i][baseCol - 1]
            summ += matrix[i][j]
            ans += prefixCount[summ - target]

        prefixCount[summ] += 1

return ans
```

• LeetDoocs

```
class Solution:
    def numSubmatrixSumTarget(self, matrix:
List[List[int]], target: int) -> int:
        def f(nums: List[int]) -> int:
            d = defaultdict(int)
            d[0] = 1
            cnt = s = 0
            for x in nums:
                s += x
                cnt += d[s - target]
                d[s] += 1
```

```

        return cnt

    m, n = len(matrix), len(matrix[0])
    ans = 0
    for i in range(m):
        col = [0] * n
        for j in range(i, m):
            for k in range(n):
                col[k] += matrix[j][k]
            ans += f(col)

    return ans

```

• SimplyLeet

Python solution for counting number of submatrices that sum to target

```

def numSubmatrixSumTarget(matrix, target):
    # Dimensions of the matrix
    n = len(matrix)
    m = len(matrix[0])
    result = 0

    # Loop over all possible starting rows
    for top in range(n):

```

```

        # Initialize a list to store the column
        sums between row 'top' and 'bottom'
        col_sums = [0] * m

        # Extend the submatrix from row 'top'
        downwards to 'bottom'
        for bottom in range(top, n):
            # Update column sums for the current
            bottom row
            for col in range(m):
                col_sums[col] +=
                matrix[bottom][col]

            # Use a dictionary to count cumulative
            sums in the 1D array (col_sums)

```

```

        sum_count = {0: 1}
        curr_sum = 0
        # Iterate over each column in col_sums
        to find the number of subarrays with sum equal to
        target
        for value in col_sums:
            curr_sum += value
            if (curr_sum - target) in
sum_count:
                result += sum_count[curr_sum -
target]
            # Update the dictionary with the
current sum count
            sum_count[curr_sum] =
sum_count.get(curr_sum, 0) + 1

```

```
return result
```

```

# Example usage:
if __name__ == "__main__":
    matrix = [[0, 1, 0], [1, 1, 1], [0, 1, 0]]
    target = 0
    print(numSubmatrixSumTarget(matrix, target)) #
Expected Output: 4

```

- LC.ca

```

class Solution:
    def numSubmatrixSumTarget(self, matrix:
List[List[int]], target: int) -> int:
        def f(nums: List[int]) -> int:
            d = defaultdict(int)
            d[0] = 1
            cnt = s = 0
            for x in nums:
                s += x
                cnt += d[s - target]
                d[s] += 1

        return cnt

        m, n = len(matrix), len(matrix[0])
        ans = 0
        for i in range(m):
            col = [0] * n
            for j in range(i, m):
                for k in range(n):
                    col[k] += matrix[j][k]
                    ans += f(col)

        return ans

```

◆ Quick Review

Key Insights

- Transform the 2D submatrix sum problem into multiple 1D subarray sum problems.

- Precompute cumulative sums by fixing two row boundaries and then summing columns between these rows.
- Use a hash table (or dictionary) to count the number of subarrays with a given sum, reducing the problem to the classic "subarray sum equals k" question.
- Iterate over all possible row pairs and for each, treat the column sums as an array where the target sum is sought.

Space & Time

Time Complexity: $O(n^2 * m)$, where n is the number of rows and m is the number of columns.

Space Complexity: $O(m)$, used by the hash table for storing cumulative sums along the columns.

Solution

The idea is to iterate over all pairs of rows (top and bottom) and, for each pair, compute the cumulative column sums between these rows. This converts the problem into finding the number of subarrays in a one-dimensional array (the column sums) that add up to the target. For each such 1D problem, use a hash table to track the frequency of cumulative sums seen so far and count how many times the difference (current sum – target) has been seen. This count gives the number of subarrays ending at the current index which sum to target. Combining counts from all possible row pairs yields the final result.

Trees & BST

Round 3 · High · Hard · 2 questions

Trees & BST

#124 Binary Tree Maximum Path Sum

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Dynamic Programming · Tree · DFS

Lists: Grind 169

Problem


Description

A path in a binary tree is a sequence of nodes where each pair of adjacent nodes in the sequence has an edge connecting them. A node can only appear in the sequence at most once. Note that the path does not need to pass through the root.

The path sum of a path is the sum of the node's values in the path.

Given the root of a binary tree, return the maximum path sum of any non-empty path.

Example 1:




Input: root = [1,2,3]

Output: 6

Explanation: The optimal path is 2 -> 1 -> 3 with a path sum of $2 + 1 + 3 = 6$.

Example 2:



Input: root = [-10,9,20,null,null,15,7]

Output: 42

Explanation: The optimal path is 15 → 20 → 7 with a path sum of 15 + 20 + 7 = 42.

Constraints:

- The number of nodes in the tree is in the range [1, 3 * 10⁴].
- -1000 ≤ Node.val ≤ 1000

Community Solutions (Python)

- WalkCC

```

class Solution:
    def maxPathSum(self, root: TreeNode | None) -> int:
        ans = -math.inf

        def maxPathSumDownFrom(root: TreeNode | None) -> int:
            """
            Returns the maximum path sum starting from
            the current root, where
            root.val is always included.
            """
            nonlocal ans

            if not root:
                return 0

            l = max(0, maxPathSumDownFrom(root.left))
            r = max(0, maxPathSumDownFrom(root.right))
            ans = max(ans, root.val + l + r)
            return root.val + max(l, r)

        maxPathSumDownFrom(root)
        return ans
  
```

- LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def maxPathSum(self, root: Optional[TreeNode])
-> int:
        def dfs(root: Optional[TreeNode]) -> int:
            if root is None:
                return 0
            left = max(0, dfs(root.left))
            right = max(0, dfs(root.right))
            nonlocal ans
            ans = max(ans, root.val + left + right)
            return root.val + max(left, right)

        ans = -inf
        dfs(root)
        return ans
```

• SimplyLeet

```
# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

class Solution:
    def maxPathSum(self, root: TreeNode) -> int:
        # Initialize global maximum path sum with a
        very small number

        self.global_max = float('-inf')

        def dfs(node: TreeNode) -> int:
            if not node:
                return 0

            # Recursively calculate max gain from
            left and right children, discard negative gains
            left_gain = max(dfs(node.left), 0)
            right_gain = max(dfs(node.right), 0)
```

```

        # Current maximum including both left
and right gains
        current_max = node.val + left_gain +
right_gain

        # Update global maximum path sum if
current_sum is higher
        self.global_max = max(self.global_max,
current_max)

        # Return the node's value plus the max
of one child gain (only one side can be chosen for
parent path)
        return node.val + max(left_gain,
right_gain)

    def dfs(self, root: Optional[TreeNode]) -> int:
        if not root:
            return 0
        left = self.dfs(root.left)
        right = self.dfs(root.right)
        current_max = root.val + max(left, right)
        self.global_max = max(self.global_max, current_max)
        return current_max

```

• LC.ca

Round 3 · High · Hard

p.1505

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def maxPathSum(self, root: Optional[TreeNode])
-> int:
        def dfs(root: Optional[TreeNode]) -> int:
            if not root:
                return 0
            left = max(0, dfs(root.left))
            right = max(0, dfs(root.right))
            nonlocal ans
            ans = max(ans, root.val + left + right)
            return root.val + max(left, right)

        ans = -inf
        dfs(root)
        return ans

```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to explore the tree in a post-order fashion.

Round 3 · High · Hard

p.1506

- For each node, calculate the maximum gain including that node and optionally one of its subtrees.
- A node's best contribution to its parent can be either just its own value or its value plus the maximum gain from either the left or right child.
- Update a global maximum path sum that considers the possibility of combining both left and right gains with the current node.
- Handle negative values by comparing gains against 0 to avoid decreasing the path sum.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree, since each node is visited exactly once.

Space Complexity: $O(n)$ in the worst-case scenario due to the recursion stack ($O(H)$ where H is the height of the tree, which can be $O(n)$ for skewed trees).

Solution

We use a recursive DFS approach (post-order traversal) to compute the maximum gain from each subtree. For every node:

- Recursively compute the maximum gain from the left and right child, ensuring that negative gains are treated as 0.
- The maximum gain that can be provided to the node's parent is the node's value plus the larger gain from its left or right child.
- Simultaneously, the potential maximum path sum including the current node as the highest (or root) node is the node's value plus both left and right gains.
- Update a global variable that holds the maximum path

sum found so far. – Return the maximum gain (node value plus one best child) to be used by the node's parent.

This strategy ensures that every possible path sum is considered while avoiding double-counting nodes.

Trees & BST

#297 Serialize and Deserialize Binary Tree

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Tree · DFS · BFS · Design

Lists: Grind 169

Problem

Description

Serialization is the process of converting a data structure or object into a sequence of bits so that it can be stored in a file or memory buffer, or transmitted across a network connection link to be reconstructed later in the same or another computer environment.


Design an algorithm to serialize and deserialize a binary tree. There is no restriction on how your serialization/deserialization algorithm should work. You just need to ensure that a binary tree can be serialized to a string and this string can be deserialized to the original tree structure.

Clarification: The input/output format is the same as how LeetCode serializes a binary tree. You do not necessarily need to follow this format, so please be creative and come up with different approaches yourself.

Round 3 · High · Hard

p.1509

Example 1:



Input: root = [1,2,3,null,null,4,5]

Output: [1,2,3,null,null,4,5]

Example 2:

Input: root = []

Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 104].
- $-1000 \leq \text{Node.val} \leq 1000$

Round 3 · High · Hard

p.1510

Community Solutions (Python)

• WalkCC

```
class Codec:
    def serialize(self, root: 'TreeNode') -> str:
        """Encodes a tree to a single string."""
        if not root:
            return ''

        s = ''
        q = collections.deque([root])

        while q:

            node = q.popleft()
            if node:
                s += str(node.val) + ' '
                q.append(node.left)
                q.append(node.right)
            else:
                s += 'n '

        return s
```

Round 3 · High · Hard

p.1511

```
def deserialize(self, data: str) -> 'TreeNode':
    """Decodes your encoded data to tree."""
    if not data:
        return None

    vals = data.split()
    root = TreeNode(vals[0])
    q = collections.deque([root])

    for i in range(1, len(vals), 2):

        node = q.popleft()
        if vals[i] != 'n':
            node.left = TreeNode(vals[i])
            q.append(node.left)
        if vals[i + 1] != 'n':
            node.right = TreeNode(vals[i + 1])
            q.append(node.right)

    return root
```

• LeetCodecs

Round 3 · High · Hard

p.1512


```
# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None
```

```
class Codec:
    def serialize(self, root):
```

```
        """Encodes a tree to a single string.
```

```
        :type root: TreeNode
        :rtype: str
        """
```

```
        if root is None:
            return ""
        q = deque([root])
        ans = []
        while q:
```

```
            node = q.popleft()
            if node:
                ans.append(str(node.val))
                q.append(node.left)
                q.append(node.right)
            else:
                ans.append("#")
        return ",".join(ans)
```

```
    def deserialize(self, data):
```

```
        """Decodes your encoded data to tree.
```

```
        :type data: str
        :rtype: TreeNode
        """
```

```
        if not data:
            return None
        vals = data.split(",")
        root = TreeNode(int(vals[0]))
        q = deque([root])
```

```

i = 1
while q:
    node = q.popleft()
    if vals[i] != "#":
        node.left = TreeNode(int(vals[i]))
        q.append(node.left)
    i += 1
    if vals[i] != "#":
        node.right =
TreeNode(int(vals[i]))
        q.append(node.right)

```

```

i += 1
return root

```

Your Codec object will be instantiated and called as such:

```

# codec = Codec()
# codec.deserialize(codec.serialize(root))

```

- SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val # value of the node
        self.left = left # left child
        self.right = right # right child

class Codec:
    # Encodes a tree to a single string using BFS.
    def serialize(self, root):

```

```

        if not root:
            return "" # Return empty string for an
empty tree
        result = []
        queue = [root] # Initialize the queue with
the root node
        while queue:
            node = queue.pop(0)
            if node:
                result.append(str(node.val)) #
Append the value of the node
                queue.append(node.left) # Enqueue
left child
                queue.append(node.right) # Enqueue
right child

```

```

        else:
            result.append("null") #
Placeholder for null node
            return ",".join(result) # Join list with
commas to form the string

# Decodes the encoded data to reconstruct the
tree using BFS.
def deserialize(self, data):
    if not data:
        return None # Return None for empty
string
        nodes = data.split(",")
        root = TreeNode(int(nodes[0])) # Create
root node from first element

```

```

queue = [root]
i = 1 # Index for traversing nodes list
while queue:
    current = queue.pop(0)
    if i < len(nodes) and nodes[i] !=
>null":
        current.left =
TreeNode(int(nodes[i]))
        queue.append(current.left)
        i += 1
        if i < len(nodes) and nodes[i] !=
>null":
            current.right =
TreeNode(int(nodes[i]))

```

Round 3 · High · Hard

p.1517

```

queue.append(current.right)
i += 1
return root

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode(object):
# def __init__(self, x):
# self.val = x
# self.left = None
# self.right = None

```

```

class Codec:
    def serialize(self, root):

```

```

        """Encodes a tree to a single string.

```

```

        :type root: TreeNode
        :rtype: str
        """

```

```

        if root is None:
            return ''
        res = []

```

Round 3 · High · Hard

p.1518

```
def preorder(root):
    if root is None:
        res.append("#,")
        return
    res.append(str(root.val) + ",")
    preorder(root.left)
    preorder(root.right)

preorder(root)
return ''.join(res)
```

```
def deserialize(self, data):
    """Decodes your encoded data to tree.

    :type data: str
    :rtype: TreeNode
    """
    if not data:
        return None
    vals = data.split(',')

```

```
def inner():
    first = vals.pop(0)
    if first == '#':
        return None
    return TreeNode(int(first), inner(),
inner())
    # seems using a constructor
    __init__(val, left, right)

    return inner()
```

```
# Your Codec object will be instantiated and called
as such:
# ser = Codec()
# deser = Codec()
# ans = deser.deserialize(ser.serialize(root))

#####

from collections import deque
```

```
# bfs, each level based
class Codec:
    def serialize(self, root):
        """Encodes a tree to a single string.

        :type root: TreeNode
        :rtype: str
        """
        ret = []
        queue = deque([root])

        while queue:
            top = queue.popleft()
            if not top:
                ret.append("None")
                continue
            else:
                ret.append(str(top.val))
                queue.append(top.left)
                queue.append(top.right)
        return ",".join(ret)
```

Round 3 · High · Hard

p.1521

```
def deserialize(self, data):
    """Decodes your encoded data to tree.

    :type data: str
    :rtype: TreeNode
    """
    data = data.split(",")
    if data[0] == "None":
        return None

    root = TreeNode(int(data[0]))
    queue = deque([root])
    i = 0
    while queue and i < len(data):
        top = queue.popleft()
        i += 1
        left = right = None
        if i < len(data) and data[i] != "None":
            left = TreeNode(int(data[i]))
            queue.append(left)
```

Round 3 · High · Hard

p.1522

```

        i += 1
        if i < len(data) and data[i] != "None":
            right = TreeNode(int(data[i]))
            queue.append(right)

        top.left = left
        top.right = right

    return root

```

```

# Your Codec object will be instantiated and called
as such:
# codec = Codec()
# codec.deserialize(codec.serialize(root))

```

◆ Quick Review

Key Insights

- The tree can be traversed using either depth-first search (DFS) or breadth-first search (BFS).
- Using BFS (level-order traversal) simplifies reconstruction since nodes are processed level-by-level.
- Placeholder markers (e.g., "null") are used to indicate missing children.
- Both serialization and deserialization processes should handle empty trees and edge cases gracefully.
- Ensuring a one-to-one mapping between the tree structure and the serialized string is critical.

Space & Time

Time Complexity: $O(n)$ for both serialization and deserialization, where n is the number of nodes in the tree.

Space Complexity: $O(n)$ due to the storage used for the serialized string and additional data structures (queue) during processing.

Solution

This solution uses level-order traversal (BFS) for both serialization and deserialization. During serialization, each node is processed sequentially with missing children recorded as "null". For deserialization, the string is split by commas, and a queue is used to rebuild the tree level by level. The root is created from the first token, then subsequent tokens are assigned as left and right children accordingly, ensuring an accurate reconstruction of the original tree.

DFS

Round 3 · High · Hard · 5 questions

Round 3 · High · Hard

p.1525

DFS

#269 Alien Dictionary

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · DFS · BFS · Graph · Topological Sort

Lists: Grind 169 | Premium 98

Problem

Description

There is a new alien language that uses the English alphabet. However, the order of the letters is unknown to you.

You are given a list of strings words from the alien language's dictionary. Now it is claimed that the strings in words are sorted lexicographically by the rules of this new language.

If this claim is incorrect, and the given arrangement of string in words cannot correspond to any order of letters, return "".

Otherwise, return a string of the unique letters in the new alien language sorted in lexicographically increasing order by the new language's rules. If there are multiple solutions, return any of them.

Example 1:

Round 3 · High · Hard

p.1526

Input: words = ["wrt","wrf","er","ett","rftt"]
 Output: "wertf"

Example 2:

Input: words = ["z","x"]
 Output: "zx"

Example 3:

Input: words = ["z","x","z"]
 Output: ""
 Explanation: The order is invalid, so return "".

Constraints:

- $1 \leq \text{words.length} \leq 100$
- $1 \leq \text{words}[i].\text{length} \leq 100$
- words[i] consists of only lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def alienOrder(self, words: list[str]) -> str:
        graph = {}
        inDegrees = [0] * 26

        self._buildGraph(graph, words, inDegrees)
        return self._topology(graph, inDegrees)

    def _buildGraph(
        self,
```

```
graph: dict[str, set[str]],
words: list[str],
inDegrees: list[int],
) -> None:
    # Create a node for each character in each
    word.
    for word in words:
        for c in word:
            if c not in graph:
                graph[c] = set()

    for first, second in zip(words, words[1:]):
        length = min(len(first), len(second))
        for j in range(length):
            u = first[j]
            v = second[j]
            if u != v:
                if v not in graph[u]:
                    graph[u].add(v)
                    inDegrees[ord(v) - ord('a')] += 1
                break # The order of characters after
    this are meaningless.
```



```

        # First = 'ab', second = 'a' . invalid
        if j == length - 1 and len(first) >
len(second):
            graph.clear()
            return

    def _topology(self, graph: dict[str, set[str]],
inDegrees: list[int]) -> str:
        s = ''
        q = collections.deque()

        for c in graph:

            if inDegrees[ord(c) - ord('a')] == 0:
                q.append(c)

        while q:
            u = q.pop()
            s += u
            for v in graph[u]:
                inDegrees[ord(v) - ord('a')] -= 1
                if inDegrees[ord(v) - ord('a')] == 0:
                    q.append(v)

    # Words = ['z', 'x', 'y', 'x']
    return s if len(s) == len(graph) else ''

```

• LeetCode

Round 3 · High · Hard

p.1529

```

class Solution:
    def alienOrder(self, words: List[str]) -> str:
        g = [[False] * 26 for _ in range(26)]
        s = [False] * 26
        cnt = 0
        n = len(words)
        for i in range(n - 1):
            for c in words[i]:
                if cnt == 26:
                    break

                o = ord(c) - ord('a')
                if not s[o]:
                    cnt += 1
                    s[o] = True

            m = len(words[i])
            for j in range(m):
                if j >= len(words[i + 1]):
                    return ''

                c1, c2 = words[i][j], words[i +
1][j]

                if c1 == c2:

```

Round 3 · High · Hard

p.1530

```

        continue
    o1, o2 = ord(c1) - ord('a'),
ord(c2) - ord('a')
    if g[o2][o1]:
        return ''
    g[o1][o2] = True
    break
for c in words[n - 1]:
    if cnt == 26:
        break
    o = ord(c) - ord('a')

    if not s[o]:
        cnt += 1
        s[o] = True

indegree = [0] * 26
for i in range(26):
    for j in range(26):
        if i != j and s[i] and s[j] and
g[i][j]:
            indegree[j] += 1
q = deque()

```

Round 3 · High · Hard

p.1531

```

ans = []
for i in range(26):
    if s[i] and indegree[i] == 0:
        q.append(i)
while q:
    t = q.popleft()
    ans.append(chr(t + ord('a')))
    for i in range(26):
        if s[i] and i != t and g[t][i]:
            indegree[i] -= 1

    if indegree[i] == 0:
        q.append(i)
return '' if len(ans) < cnt else
''.join(ans)

```

• SimplyLeet

```

# Python solution for Alien Dictionary
from collections import defaultdict, deque

def alienOrder(words):
    # Create data structures: graph and in_degree
    # for each unique character.
    graph = defaultdict(set)
    in_degree = {char: 0 for word in words for char
in word}

    # Build the graph.
    for i in range(len(words) - 1):

```

Round 3 · High · Hard

p.1532

```

        word1, word2 = words[i], words[i+1]
        # Check for invalid case where prefix of
        next word is the same but word1 is longer.
        if len(word1) > len(word2) and
word1.startswith(word2):
            return ""
        # Compare the words character by character.
        for c1, c2 in zip(word1, word2):
            if c1 != c2:
                # If we haven't already recorded
this edge, add it.
                if c2 not in graph[c1]:
                    graph[c1].add(c2)

```

```

            in_degree[c2] += 1
            break

```

```

# Prepare for topological sort by adding all
nodes with in_degree == 0.
queue = deque([char for char in in_degree if
in_degree[char] == 0])
order = []

```

```

# Kahn's algorithm for topological sort.
while queue:
    current = queue.popleft()

```

```

        order.append(current)
        for neighbor in graph[current]:
            in_degree[neighbor] -= 1
            if in_degree[neighbor] == 0:
                queue.append(neighbor)

# Check if topological sort is possible.
if len(order) == len(in_degree):
    return "".join(order)
else:

```

```

    return ""

```

```

# Example usage:
print(alienOrder(["wrt","wrf","er","ett","rftt"]))

```

• LC.ca

```

'''
>>> a = ['a','b','c']
>>> zip(a, a[1:])
[('a', 'b'), ('b', 'c')]

# same as pairwise()
>>> from itertools import pairwise
>>> pairwise(a)
<itertools.pairwise object at 0x108458730>

```

```
>>> list(pairwise(a))
[('a', 'b'), ('b', 'c')]
...

...
>>> inDegree = {'a':11, 'b':22, 'c':33}

>>> [print(c) for c in inDegree.keys()]
a
b
```

```
c
[None, None, None]

# same as for c in inDegree.keys()
>>> [print(c) for c in inDegree]
a
b
c
[None, None, None]
```

```
>>> [print(c) for c in inDegree.values()]
11
22
33
[None, None, None]

>>> [print(c) for c in inDegree.items()]
('a', 11)
('b', 22)
('c', 33)
```

```
[None, None, None]
...

from collections import defaultdict, deque
from typing import Dict, List, Set

class Solution:
    def alienOrder(self, words: List[str]) -> str:
        graph = defaultdict(set)
        inDegree = defaultdict(int)
```

```

        self._buildGraph(graph, words, inDegree)
        return self._topology(graph, inDegree)

    def _buildGraph(self, graph: Dict[str,
Set[str]], words: List[str], inDegree: Dict[str,
int]) -> None:
        # Create a node for each character in each
word
        for word in words:
            for c in word:
                inDegree[c] = 0 # necessary for
final char counting

        for first, second in zip(words,
words[1:]): # or pairwise(words)

            length = min(len(first), len(second))
            for j in range(length):
                u = first[j]
                v = second[j]
                if u != v:
                    if v not in graph[u]:
                        graph[u].add(v)
                        inDegree[v] += 1
                    break # Later characters'
order is meaningless
                if j == length - 1 and len(first) >
len(second):

```

Round 3 · High · Hard

p.1537

```

        # If 'ab' comes before 'a',
it's an invalid order
        graph.clear()
        return

    def _topology(self, graph: Dict[str,
Set[str]], inDegree: Dict[str, int]) -> str:
        result = ''
        q = deque([c for c in inDegree if
inDegree[c] == 0])

        while q:
            u = q.popleft()

            result += u
            for v in graph[u]:
                inDegree[v] -= 1
                if inDegree[v] == 0:
                    q.append(v)

        # If there are remaining characters in
inDegree, it means there's a cycle
        if any(inDegree.values()):
            return ''

```

Round 3 · High · Hard

p.1538

```

        # Words = ['z', 'x', 'y', 'x']
        return result if len(result) ==
len(indegree) else ''

```

```
#####
```

```
import collections
```

```
class Node(object):
    def __init__(self, val):

```

```

        self.val = val
        self.neighbors = []

```

```

    def connect(self, node):
        self.neighbors.append(node)

```

```

    def getNbrs(self):
        return self.neighbors

```

```

class Solution(object):
    def alienOrder(self, words):
        """
        :type words: List[str]
        :rtype: str
        """

        def dfs(root, graph, visited):
            visited[root] = 1
            for nbr in graph[root].getNbrs():

```

```

                if visited[nbr.val] == 0:
                    if not dfs(nbr.val, graph, visited):
                        return False
                elif visited[nbr.val] == 1:
                    return False

```

```

        visited[root] = 2
        self.ans += root
        return True

```

```

self.ans = ""
graph = {}
visited = collections.defaultdict(int)
self.topNum = 0
for i in range(0, len(words) - 1):
    a = words[i]
    b = words[i + 1]
    i = 0
    while i < len(a) and i < len(b):
        if a[i] != b[i]:

```

```

        nodeA = nodeB = None
        if a[i] not in graph:
            nodeA = Node(a[i])
            graph[a[i]] = nodeA
        else:
            nodeA = graph[a[i]]
        if b[i] not in graph:
            nodeB = Node(b[i])
            graph[b[i]] = nodeB
        else:

```

```

            nodeB = graph[b[i]]
            nodeA.connect(nodeB)
            break
        i += 1
    if i < len(a) and i >= len(b):
        return ""

    for c in graph:
        if visited[c] == 0:
            if not dfs(c, graph, visited):

    return ""

    unusedSet = set()
    for word in words:
        for c in word:
            unusedSet.add(c)

    for c in unusedSet:
        if c not in graph:
            self.ans += c

    return self.ans[::-1]

```

◆ Quick Review

Key Insights

- Build a graph where each unique character is a node.
- Compare adjacent words to determine the first different character, which gives the ordering (directed edge).

- Ensure all characters, even those without edges, are represented.
- Use a topological sort (using DFS or BFS) to determine a valid character ordering.
- Handle edge cases, such as words where a longer word appears before its prefix which is invalid.

Space & Time

Time Complexity: $O(N * L + V + E)$ where N is the number of words, L is the average length of each word, V is the number of unique characters, and E is the number of edges derived from adjacent word comparisons.

Space Complexity: $O(V + E)$ for storing the graph and additional data structures for topological sorting.

Solution

We use a graph-based approach to solve the problem. First, we initialize a graph (adjacency list) and in-degree dictionary for all unique characters. We then iterate through each pair of consecutive words, comparing them character by character to find the first pair that differs. This difference implies an order: the first character should come before the second character. We add a directed edge accordingly and update the in-degree for the target character.

To get the final ordering, we perform a topological sort using Kahn's algorithm:

- Enqueue characters with in-degree zero.
- Remove characters from the queue and append them to our result.
- Decrement the in-degree of neighboring nodes.
- If a

neighbor's in-degree becomes zero, enqueue it. If the result's length equals the number of unique nodes, we have a valid ordering; otherwise, a cycle exists, and we return an empty string.

DFS

#329 Longest Increasing Path in a Matrix

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Graph · Topological Sort · Memoization · Matrix

Lists: Grind 169

Problem

Description

Given an $m \times n$ integers matrix, return the length of the longest increasing path in matrix.

From each cell, you can either move in four directions: left, right, up, or down. You may not move diagonally or move outside the boundary (i.e., wrap-around is not allowed).

Example 1:

9	9	4
6	6	8
2	1	1

Input: matrix = [[9,9,4],[6,6,8],[2,1,1]]

Output: 4

Explanation: The longest increasing path is [1, 2, 6, 9].

Example 2:

3 →	4 →	5
3	2	↓ 6
2	2	1

Input: matrix = [[3,4,5],[3,2,6],[2,2,1]]

Output: 4

Explanation: The longest increasing path is [3, 4, 5, 6]. Moving diagonally is not allowed.

Example 3:

Input: matrix = [[1]]

Output: 1

Constraints:

- $m == \text{matrix.length}$
- $n == \text{matrix}[i].\text{length}$

Round 3 · High · Hard

p.1547

- $1 \leq m, n \leq 200$
- $0 \leq \text{matrix}[i][j] \leq 231 - 1$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestIncreasingPath(self, matrix:
list[list[int]]) -> int:
        m = len(matrix)
        n = len(matrix[0])

        @functools.lru_cache(None)
        def dfs(i: int, j: int, prev: int) -> int:
            if i < 0 or i == m or j < 0 or j == n:
                return 0
            if matrix[i][j] <= prev:
                return 0

            curr = matrix[i][j]
            return 1 + max(dfs(i + 1, j, curr),
                        dfs(i - 1, j, curr),
                        dfs(i, j + 1, curr),
                        dfs(i, j - 1, curr))

            return max(dfs(i, j, -math.inf) for i in
range(m) for j in range(n))
```

• LeetCode

Round 3 · High · Hard

p.1548

```

class Solution:
    def longestIncreasingPath(self, matrix:
List[List[int]]) -> int:
        @cache
        def dfs(i: int, j: int) -> int:
            ans = 0
            for a, b in pairwise((-1, 0, 1, 0,
-1)):
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n and
matrix[x][y] > matrix[i][j]:
                    ans = max(ans, dfs(x, y))
            return ans + 1

        m, n = len(matrix), len(matrix[0])
        return max(dfs(i, j) for i in range(m) for
j in range(n))

```

• SimplyLeet

Round 3 · High · Hard

p.1549

```

# Python implementation for Longest Increasing Path
in a Matrix

class Solution:
    def longestIncreasingPath(self, matrix:
List[List[int]]) -> int:
        # Return 0 if matrix is empty
        if not matrix or not matrix[0]:
            return 0

        # Dimensions of the matrix
        m, n = len(matrix), len(matrix[0])

        # Create a memoization table (same
dimensions as the matrix) initialized to 0
        memo = [[0] * n for _ in range(m)]

        # Define the DFS function to explore the
matrix
        def dfs(i, j):
            # If result is already computed, return
it

            if memo[i][j]:
                return memo[i][j]

            # Start with a path length of 1 (the
cell itself)

```

Round 3 · High · Hard

p.1550

```

        max_len = 1
        # Define directions: up, down, left,
right      directions = [(-1, 0), (1, 0), (0, -1),
(0, 1)]

        # Explore each direction
        for dx, dy in directions:
            ni, nj = i + dx, j + dy
            # If neighbor is within bounds and
its value is greater
            if 0 <= ni < m and 0 <= nj < n and
matrix[ni][nj] > matrix[i][j]:
                length = 1 + dfs(ni, nj)

        max_len = max(max_len, length)

        # Store computed longest length in
memoization table
        memo[i][j] = max_len
        return max_len

    # Compute the longest increasing path
starting from each cell
    longest_path = 0
    for i in range(m):
        for j in range(n):

```

Round 3 · High · Hard

p.1551

```

        longest_path = max(longest_path,
dfs(i, j))

    return longest_path

```

• LC.ca

```

class Solution:
    def longestIncreasingPath(self, matrix:
List[List[int]]) -> int:
        @cache
        def dfs(i: int, j: int) -> int:
            ans = 0
            for a, b in pairwise((-1, 0, 1, 0,
-1)):
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n and
matrix[x][y] > matrix[i][j]:
                    ans = max(ans, dfs(x, y))
            return ans + 1

    m, n = len(matrix), len(matrix[0])
    return max(dfs(i, j) for i in range(m) for
j in range(n))

```

◆ Quick Review

Key Insights

- The problem can be modeled as a graph where each cell is a node connected to its neighbors if the neighbor's value

Round 3 · High · Hard

p.1552

is greater.

- A Depth First Search (DFS) with memoization is effective to avoid recalculations.
- Dynamic Programming (DP) is used to cache the longest path found from each cell.
- Since each cell can be a start point, compute DFS for every cell and update the global maximum.

Space & Time

Time Complexity: $O(m * n)$ – each cell is computed once due to memoization.

Space Complexity: $O(m * n)$ – space is used for the memoization cache and recursion stack.

Solution

We solve the problem using a DFS approach combined with memoization to keep track of already computed longest paths from each cell. For each cell in the matrix, a DFS is initiated to explore all four valid directions (up, down, left, right). If a neighboring cell has a larger value, the DFS recurses from that neighbor. The result for each cell is cached in a 2D memoization table, ensuring that each cell is processed only once. After processing all cells, the maximum value in the memo table represents the longest increasing path in the entire matrix.

DFS

#685 Redundant Connection II

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · BFS · Union Find · Graph

Lists: Denny Zhang

Problem

Description

In this problem, a rooted tree is a directed graph such that, there is exactly one node (the root) for which all other nodes are descendants of this node, plus every node has exactly one parent, except for the root node which has no parents.


The given input is a directed graph that started as a rooted tree with n nodes (with distinct values from 1 to n), with one additional directed edge added. The added edge has two different vertices chosen from 1 to n , and was not an edge that already existed.

The resulting graph is given as a 2D-array of edges. Each element of edges is a pair $[ui, vi]$ that represents a directed edge connecting nodes ui and vi , where ui is a parent of child vi .

Return an edge that can be removed so that the resulting graph is a rooted tree of n nodes. If there are multiple answers, return the answer that occurs last in


the given 2D-array.

Example 1:



Input: edges = [[1,2],[1,3],[2,3]]
Output: [2,3]

Example 2:



Input: edges = [[1,2],[2,3],[3,4],[4,1],[1,5]]
Output: [4,1]

Constraints:

- $n == \text{edges.length}$
- $3 \leq n \leq 1000$
- $\text{edges}[i].\text{length} == 2$
- $1 \leq u_i, v_i \leq n$
- $u_i \neq v_i$

Community Solutions (Python)

• WalkCC

```

class UnionFind:
    def __init__(self, n: int):
        self.id = list(range(n))
        self.rank = [0] * n

    def unionByRank(self, u: int, v: int) -> bool:
        i = self._find(u)
        j = self._find(v)
        if i == j:
            return False

        if self.rank[i] < self.rank[j]:
            self.id[i] = j
        elif self.rank[i] > self.rank[j]:
            self.id[j] = i
        else:
            self.id[i] = j
            self.rank[j] += 1
        return True

    def _find(self, u: int) -> int:

```

```

    if self.id[u] != u:
        self.id[u] = self._find(self.id[u])
    return self.id[u]

```

```

class Solution:
    def findRedundantDirectedConnection(
        self, edges: list[list[int]],
    ) -> list[int]:
        ids = [0] * (len(edges) + 1)

```

```

        nodeWithTwoParents = 0

```

```

        for _, v in edges:
            ids[v] += 1
            if ids[v] == 2:
                nodeWithTwoParents = v

```

```

        def
        findRedundantDirectedConnection(skippedEdgeIndex:
int) -> list[int]:
            uf = UnionFind(len(edges) + 1)

```

```

for i, edge in enumerate(edges):
    if i == skippedEdgeIndex:
        continue
    if not uf.unionByRank(edge[0], edge[1]):
        return edge

return []

# If there is no edge with two ids, don't skip
any edge.
if nodeWithTwoParents == 0:

    return findRedundantDirectedConnection(-1)

for i in reversed(range(len(edges))):
    _, v = edges[i]
    if v == nodeWithTwoParents:
        # Try to delete the edges[i].
        if not findRedundantDirectedConnection(i):
            return edges[i]

```

• LeetCode

Round 3 · High · Hard

p.1559

```

class Solution:
    def findRedundantDirectedConnection(self,
edges: List[List[int]]) -> List[int]:
        def find(x: int) -> int:
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

        n = len(edges)
        ind = [0] * n
        for _, v in edges:
            ind[v - 1] += 1
            dup = [i for i, (_, v) in enumerate(edges)
if ind[v - 1] == 2]
            p = list(range(n))
            if dup:
                for i, (u, v) in enumerate(edges):
                    if i == dup[1]:
                        continue
                    pu, pv = find(u - 1), find(v - 1)
                    if pu == pv:
                        return edges[dup[0]]

```

Round 3 · High · Hard

p.1560


```

        p[pu] = pv
        return edges[dup[1]]
    for i, (u, v) in enumerate(edges):
        pu, pv = find(u - 1), find(v - 1)
        if pu == pv:
            return edges[i]
        p[pu] = pv

```

```

class UnionFind:
    __slots__ = "p", "size"

    def __init__(self, n: int):
        self.p: List[int] = list(range(n))
        self.size: List[int] = [1] * n

    def find(self, x: int) -> int:
        if self.p[x] != x:
            self.p[x] = self.find(self.p[x])

```

```

        return self.p[x]

    def union(self, a: int, b: int) -> bool:
        pa, pb = self.find(a), self.find(b)
        if pa == pb:
            return False
        if self.size[pa] > self.size[pb]:
            self.p[pb] = pa
            self.size[pa] += self.size[pb]
        else:

```

```

        self.p[pa] = pb
        self.size[pb] += self.size[pa]
        return True

```

```

class Solution:
    def findRedundantDirectedConnection(self,
edges: List[List[int]]) -> List[int]:
        n = len(edges)
        ind = [0] * n
        for _, v in edges:

```

```

            ind[v - 1] += 1
            dup = [i for i, (_, v) in enumerate(edges)
if ind[v - 1] == 2]
            uf = UnionFind(n)
            if dup:
                for i, (u, v) in enumerate(edges):
                    if i == dup[1]:
                        continue
                    if not uf.union(u - 1, v - 1):
                        return edges[dup[0]]
            return edges[dup[1]]

```

```

        for i, (u, v) in enumerate(edges):
            if not uf.union(u - 1, v - 1):
                return edges[i]

```

• SimplyLeet

```

# Python solution with line-by-line comments
class Solution:
    def findRedundantDirectedConnection(self,
edges):
        n = len(edges)
        candidate1 = None
        candidate2 = None
        # Dictionary to record the parent for each
child node.
        childParent = {}
        # First pass: detect if a node has two
parents.
        for u, v in edges:
            if v in childParent:
                candidate1 = [childParent[v], v] #
earlier edge providing parent
                candidate2 = [u, v] # later edge
providing parent
            else:
                childParent[v] = u

        # Initialize Union-Find structure: each
node is its own parent.
        root = [i for i in range(n+1)]

        # Find function with path compression.

```

```

def find(x):
    while root[x] != x:
        root[x] = root[root[x]]
        x = root[x]
    return x

cycle_edge = None
# Second pass: union-find to detect cycle,
skip candidate2 if conflict exists.
for u, v in edges:
    # If candidate2 exists and current edge
is candidate2, skip it.

    if candidate2 and [u, v] == candidate2:
        continue
    parent_u = find(u)
    parent_v = find(v)
    # If they have the same root, a cycle
is detected.
    if parent_u == parent_v:
        cycle_edge = [u, v]
    else:
        root[parent_v] = parent_u

```

```
# Decide which edge to remove based on
conflict and cycle detection.
if candidate1:
    return candidate1 if cycle_edge else
candidate2
    return cycle_edge
```

• LC.ca

```
class UnionFind:
    def __init__(self, n):
        self.p = list(range(n))
        self.n = n

    def union(self, a, b):
        if self.find(a) == self.find(b):
            return False
        self.p[self.find(a)] = self.find(b)
        self.n -= 1

    def find(self, x):
        if self.p[x] != x:
            self.p[x] = self.find(self.p[x])
        return self.p[x]

    def findRedundantDirectedConnection(self,
edges: List[List[int]]) -> List[int]:
    return True
```

Round 3 · High · Hard

p.1565

```
n = len(edges)
p = list(range(n + 1))
uf = UnionFind(n + 1)
conflict = cycle = None
for i, (u, v) in enumerate(edges):
    if p[v] != v:
        conflict = i
    else:
        p[v] = u
        if not uf.union(u, v):
```

```
        cycle = i
    if conflict is None:
        return edges[cycle]
    v = edges[conflict][1]
    if cycle is not None:
        return [p[v], v]
    return edges[conflict]
```

◆ Quick Review

Key Insights

- The extra edge introduces one of two issues:
- A node has two parents.
- A cycle exists in the graph.
- First, scan through edges to detect if any node gets two parents. If found, there are two candidate edges.
- Use the Union-Find data structure (Disjoint Set Union) to detect cycles.

Round 3 · High · Hard

p.1566

- Depending on whether a conflict (two parents) exists and/or a cycle is detected, decide which edge to remove:
- If a node has two parents, test by ignoring the second candidate edge during union-find. If no cycle forms, the second candidate is removable; otherwise, remove the first candidate.
- If no conflict exists, simply remove the edge that causes a cycle.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution consists of two main steps:

– Traverse the list of edges to detect whether any node has two parents. If such a node is found, mark the two candidate edges (candidate1: the earlier edge, candidate2: the later edge). – Use the Union-Find algorithm to iterate over the edges: If there is a conflict (i.e., a node with two parents), skip candidate2 initially. While processing edges with Union-Find, check if adding an edge would create a cycle. Finally, based on the presence of a cycle and the two-parent conflict, decide which candidate edge to return.

Key Data Structures:

– An array (or hash map) to track the parent of each node. – The Union-Find (Disjoint Set Union) structure to union sets and detect cycles.

DFS

#1036 Escape a Large Maze

HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · DFS · BFS

Lists: Denny Zhang

Problem

Description

There is a 1 million by 1 million grid on an XY-plane, and the coordinates of each grid square are (x, y) .

We start at the source = $[sx, sy]$ square and want to reach the target = $[tx, ty]$ square. There is also an array of blocked squares, where each $blocked[i] = [xi, yi]$ represents a blocked square with coordinates (xi, yi) .

Each move, we can walk one square north, east, south, or west if the square is not in the array of blocked squares. We are also not allowed to walk outside of the grid.

Return true if and only if it is possible to reach the target square from the source square through a sequence of valid moves.

Example 1:

Input: blocked = [[0,1],[1,0]], source = [0,0], target = [0,2]

Output: false

Explanation: The target square is inaccessible starting from the source square because we cannot move.

We cannot move north or east because those squares are blocked.

We cannot move south or west because we cannot go outside of the grid.

Example 2:

Input: blocked = [], source = [0,0], target = [999999,999999]

Output: true

Explanation: Because there are no blocked cells, it is possible to reach the target square.

Constraints:

- $0 \leq \text{blocked.length} \leq 200$
- $\text{blocked}[i].\text{length} == 2$
- $0 \leq x_i, y_i < 106$
- $\text{source.length} == \text{target.length} == 2$
- $0 \leq s_x, s_y, t_x, t_y < 106$
- $\text{source} \neq \text{target}$
- It is guaranteed that source and target are not blocked.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def isEscapePossible(
        self,
        blocked: list[list[int]],
        source: list[int],
        target: list[int]
    ) -> bool:
        def dfs(i: int, j: int, target: list[int],
            seen: set) -> bool:
            if i < 0 or i >= 10**6 or j < 0 or j >=
            10**6:
                return False
```

```
        if (i, j) in blocked or (i, j) in seen:
            return False
        seen.add((i, j))
        return (len(seen) > (1 + 199) * 199 // 2 or
            [i, j] == target or
                dfs(i + 1, j, target, seen) or
                dfs(i - 1, j, target, seen) or
                dfs(i, j + 1, target, seen) or
                dfs(i, j - 1, target, seen))
```

```
        blocked = set(tuple(b) for b in blocked)
```

```
        return (dfs(source[0], source[1], target,
            set()) and
                dfs(target[0], target[1], source,
            set()))
```

- LeetCode

```
class Solution:
    def isEscapePossible(
        self, blocked: List[List[int]], source:
        List[int], target: List[int]
    ) -> bool:
        def dfs(source: List[int], target:
        List[int], vis: set) -> bool:
            vis.add(tuple(source))
            if len(vis) > m:
                return True
            for a, b in pairwise(dirs):
                x, y = source[0] + a, source[1] + b

                if 0 <= x < n and 0 <= y < n and
                (x, y) not in s and (x, y) not in vis:
                    if [x, y] == target or dfs([x,
                    y], target, vis):
                        return True
            return False

        s = {(x, y) for x, y in blocked}
        dirs = (-1, 0, 1, 0, -1)
        n = 10**6
        m = len(blocked) ** 2 // 2
        return dfs(source, target, set()) and
        dfs(target, source, set())
```

• SimplyLeet

Round 3 · High · Hard

p.1571

```
from collections import deque

def isEscapePossible(blocked, source, target):
    blocked_set = {(x, y) for x, y in blocked}
    n = len(blocked)
    if n < 2:
        return True
    max_steps = n * (n - 1) // 2

    def bfs(start, finish):
        queue = deque([start])
        seen = {tuple(start)}
        directions = [(0, 1), (1, 0), (0, -1), (-1,
        0)]

        steps = 0

        while queue:
            x, y = queue.popleft()
            if [x, y] == finish:
                return True
            for dx, dy in directions:
```

Round 3 · High · Hard

p.1572

```

        nx, ny = x + dx, y + dy
        if 0 <= nx < 10**6 and 0 <= ny <
10**6 and (nx, ny) not in blocked_set and (nx, ny)
not in seen:
            queue.append((nx, ny))
            seen.add((nx, ny))
            steps += 1
            if steps > max_steps:
                return True
    return False

    return bfs(source, target) and bfs(target,
source)

```

```

# Example usage:
print(isEscapePossible([[0,1],[1,0]], [0,0],
[0,2])) # Expected output: false
print(isEscapePossible([], [0,0],
[999999,999999])) # Expected output: true

```

• LC.ca

Round 3 · High · Hard

p.1573

```

class Solution:
    def isEscapePossible(
        self, blocked: List[List[int]], source:
List[int], target: List[int]
    ) -> bool:
        def dfs(source, target, seen):
            x, y = source
            if (
                not (0 <= x < 10**6 and 0 <= y <
10**6)
                or (x, y) in blocked
                or (x, y) in seen
            ):
                return False
            seen.add((x, y))
            if len(seen) > 20000 or source ==
target:
                return True
            for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:
                next = [x + a, y + b]
                if dfs(next, target, seen):
                    return True
            return False

        blocked = set((x, y) for x, y in blocked)
        return dfs(source, target, set()) and
dfs(target, source, set())

```

Round 3 · High · Hard

p.1574

◆ Quick Review

Key Insights

- The maximum number of blocked cells is 200, which limits the size of a potential completely enclosed region.
- Even if the target is unreachable by direct search, if the BFS (or DFS) explores more than a calculated threshold (based on the blocked cells count), then the region is not fully enclosed.
- Perform two BFS explorations: one starting from the source and one from the target. If neither search is trapped within a small region, then an escape/path exists.
- Early exit in BFS if the target is reached or if the number of visited nodes exceeds the maximum possible enclosed area.

Space & Time

Time Complexity: $O(B^2)$ in the worst case, where B is the number of blocked cells (with $B \leq 200$).

Space Complexity: $O(B^2)$ due to the storage of visited cells in the worst-case scenario.

Solution

The solution uses the Breadth-First Search (BFS) algorithm. A key observation here is that the blocked cells can only enclose a limited area. We compute a threshold (maxSteps) based on the number of blocked cells ($n*(n-1)/2$) to determine if the search is confined. The BFS from the source (and similarly from the target) stops in two cases: if

the target is reached or if the visited area exceeds the threshold, indicating that the region is not fully enclosed by blocked cells. Running BFS from both the source and target ensures that neither is trapped inside a small banned region.

DFS

#1192 Critical Connections in a Network

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

DFS · Graph · Biconnected Component

Lists: Denny Zhang

Problem


Description

There are n servers numbered from 0 to $n - 1$ connected by undirected server-to-server connections forming a network where $\text{connections}[i] = [a_i, b_i]$ represents a connection between servers a_i and b_i . Any server can reach other servers directly or indirectly through the network.

A critical connection is a connection that, if removed, will make some servers unable to reach some other server.

Return all critical connections in the network in any order.

Example 1:



Input: $n = 4$, $\text{connections} = [[0,1],[1,2],[2,0],[1,3]]$
 Output: $[[1,3]]$
 Explanation: $[[3,1]]$ is also accepted.

Example 2:

Input: $n = 2$, $\text{connections} = [[0,1]]$
 Output: $[[0,1]]$

Constraints:

- $2 \leq n \leq 105$
- $n - 1 \leq \text{connections.length} \leq 105$
- $0 \leq a_i, b_i \leq n - 1$

Round 3 · High · Hard

p.1577

Round 3 · High · Hard

p.1578

- $a_i \neq b_i$
- There are no repeated connections.

Community Solutions (Python)

• LeetCode

```
class Solution:
    def criticalConnections(
        self, n: int, connections: List[List[int]]
    ) -> List[List[int]]:
        def tarjan(a: int, fa: int):
            nonlocal now
            now += 1
            dfn[a] = low[a] = now
            for b in g[a]:
                if b == fa:
                    continue
                if not dfn[b]:
                    tarjan(b, a)
                    low[a] = min(low[a], low[b])
                    if low[b] > dfn[a]:
                        ans.append([a, b])
                else:
                    low[a] = min(low[a], dfn[b])

            g = [[] for _ in range(n)]
```

```
for a, b in connections:
    g[a].append(b)
    g[b].append(a)
```

```
dfn = [0] * n
low = [0] * n
now = 0
ans = []
tarjan(0, -1)
return ans
```

• SimplyLeet

```
def criticalConnections(n, connections):
    # Build the adjacency list for the graph
    graph = [[] for _ in range(n)]
    for u, v in connections:
        graph[u].append(v)
        graph[v].append(u)

    # Initialize discovery time and low arrays
    disc = [-1] * n # Discovery times, -1 indicates
    unvisited
    low = [0] * n # Low values for each node
```

```
time = 0 # Global timer
bridges = [] # List to store critical
connections (bridges)
```

```
def dfs(u, parent):
    nonlocal time
    disc[u] = low[u] = time
    time += 1
```

```
# Traverse all adjacent nodes
for v in graph[u]:
```

```
    if disc[v] == -1: # If v is not visited
        dfs(v, u)
        low[u] = min(low[u], low[v])
        # If the smallest reachable vertex
from v is after u's discovery, edge (u,v) is a
bridge
```

```
        if low[v] > disc[u]:
            bridges.append([u, v])
        elif v != parent: # v is visited and is
not the parent (back edge)
            low[u] = min(low[u], disc[v])
```

```
# Start DFS from node 0
```

```
dfs(0, -1)
return bridges
```

```
# Example usage:
```

```
n = 4
connections = [[0,1],[1,2],[2,0],[1,3]]
print(criticalConnections(n, connections)) #
Output: [[1, 3]]
```

• LC.ca

```
class Solution:
    def criticalConnections(
        self, n: int, connections: List[List[int]]
    ) -> List[List[int]]:
        def tarjan(a: int, fa: int):
            nonlocal now
            now += 1
            dfn[a] = low[a] = now
            for b in g[a]:
                if b == fa:
```

```

        continue
    if not dfn[b]:
        tarjan(b, a)
        low[a] = min(low[a], low[b])
        if low[b] > dfn[a]:
            ans.append([a, b])
    else:
        low[a] = min(low[a], dfn[b])

g = [[] for _ in range(n)]

for a, b in connections:
    g[a].append(b)
    g[b].append(a)

dfn = [0] * n
low = [0] * n
now = 0
ans = []
tarjan(0, -1)
return ans

```

◆ Quick Review

Key Insights

- The problem is about identifying bridges in an undirected graph.
- Use Depth-First Search (DFS) to explore the graph.
- During DFS, track discovery times and low-link values for each node.

- For an edge (u, v) , if $\text{low}[v] > \text{disc}[u]$ then (u, v) is a critical connection.
- Tarjan's algorithm is well-suited for solving this problem in $O(n + m)$ time.

Space & Time

Time Complexity: $O(n + m)$ where n is the number of nodes and m is the number of connections.

Space Complexity: $O(n + m)$ for the graph representation and DFS recursion stack.

Solution

We solve this problem using a DFS-based approach (Tarjan's algorithm). First, we transform the list of connections into an adjacency list for efficient graph traversal. We then start a DFS from an arbitrary node (node 0, since the graph is connected) while maintaining two arrays: one for discovery times (disc) and one for low-link values (low). The discovery time is the time when the node is first visited, and the low-link value is the smallest discovery time that can be reached from that node (including via back edges). For each edge (u, v) , after a DFS call on v , if $\text{low}[v] > \text{disc}[u]$ then the edge (u, v) is a bridge because v (and its descendants) cannot reach u or any of its ancestors without this edge. This edge is then recorded as a critical connection.

Graphs

Round 3 · High · Hard · 2 questions

Round 3 · High · Hard

p.1585

Graphs

#305 Number of Islands II

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Union Find · Matrix

Lists: Premium 98

Problem

Description

You are given an empty 2D binary grid `grid` of size $m \times n$. The grid represents a map where 0's represent water and 1's represent land. Initially, all the cells of grid are water cells (i.e., all the cells are 0's).

We may perform an add land operation which turns the water at position into a land. You are given an array `positions` where `positions[i] = [ri, ci]` is the position (ri, ci) at which we should operate the i th operation.

Return an array of integers `answer` where `answer[i]` is the number of islands after turning the cell (ri, ci) into a land.

An island is surrounded by water and is formed by connecting adjacent lands horizontally or vertically. You may assume all four edges of the grid are all surrounded by water.

Example 1:

Round 3 · High · Hard

p.1586

3 →	4 →	5
3	2	↓ 6
2	2	1

Input: m = 3, n = 3, positions =
[[0,0],[0,1],[1,2],[2,1]]

Output: [1,1,2,3]

Explanation:

Initially, the 2d grid is filled with water.

- Operation #1: addLand(0, 0) turns the water at grid[0][0] into a land. We have 1 island.
- Operation #2: addLand(0, 1) turns the water at grid[0][1] into a land. We still have 1 island.
- Operation #3: addLand(1, 2) turns the water at grid[1][2] into a land. We have 2 islands.
- Operation #4: addLand(2, 1) turns the water at grid[2][1] into a land. We have 3 islands.

Example 2:

Input: m = 1, n = 1, positions = [[0,0]]

Output: [1]

Constraints:

- 1 ≤ m, n, positions.length ≤ 104
- 1 ≤ m * n ≤ 104
- positions[i].length == 2
- 0 ≤ ri < m
- 0 ≤ ci < n

Follow up: Could you solve it in time complexity O(k log(mn)), where k == positions.length?

Community Solutions (Python)

- WalkCC

```

class UnionFind:
    def __init__(self, n: int):
        self.id = [-1] * n
        self.rank = [0] * n

    def unionByRank(self, u: int, v: int) -> None:
        i = self.find(u)
        j = self.find(v)
        if i == j:
            return

        if self.rank[i] < self.rank[j]:
            self.id[i] = j
        elif self.rank[i] > self.rank[j]:
            self.id[j] = i
        else:
            self.id[i] = j
            self.rank[j] += 1

    def find(self, u: int) -> int:
        if self.id[u] != u:

```

```

        self.id[u] = self.find(self.id[u])
        return self.id[u]

```

```

class Solution:
    def numIslands2(
        self,
        m: int,
        n: int,
        positions: list[list[int]],

```

```

    ) -> list[int]:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        ans = []
        seen = [[False] * n for _ in range(m)]
        uf = UnionFind(m * n)
        count = 0

        def getId(i: int, j: int, n: int) -> int:
            return i * n + j

```

```

for i, j in positions:
    if seen[i][j]:
        ans.append(count)
        continue
    seen[i][j] = True
    id = getId(i, j, n)
    uf.id[id] = id
    count += 1
    for dx, dy in DIRS:
        x = i + dx

        y = j + dy
        if x < 0 or x == m or y < 0 or y == n:
            continue
        neighborId = getId(x, y, n)
        if uf.id[neighborId] == -1: # water
            continue
        currentParent = uf.find(id)
        neighborParent = uf.find(neighborId)
        if currentParent != neighborParent:
            uf.unionByRank(currentParent,
neighborParent)

        count -= 1
    ans.append(count)

return ans

```

• LeetCode

Round 3 · High · Hard

p.1591

```

class UnionFind:
    def __init__(self, n: int):
        self.p = list(range(n))
        self.size = [1] * n

    def find(self, x: int):
        if self.p[x] != x:
            self.p[x] = self.find(self.p[x])
        return self.p[x]

    def union(self, a: int, b: int) -> bool:
        pa, pb = self.find(a - 1), self.find(b - 1)
        if pa == pb:
            return False
        if self.size[pa] > self.size[pb]:
            self.p[pb] = pa
            self.size[pa] += self.size[pb]
        else:
            self.p[pa] = pb
            self.size[pb] += self.size[pa]

```

Round 3 · High · Hard

p.1592


```
return True
```

```
class Solution:
```

```
    def numIslands2(self, m: int, n: int,
positions: List[List[int]]) -> List[int]:
        uf = UnionFind(m * n)
        grid = [[0] * n for _ in range(m)]
        ans = []
        dirs = (-1, 0, 1, 0, -1)
        cnt = 0
```

```
        for i, j in positions:
            if grid[i][j]:
                ans.append(cnt)
                continue
            grid[i][j] = 1
            cnt += 1
            for a, b in pairwise(dirs):
                x, y = i + a, j + b
                if (
                    0 <= x < m
```

Round 3 · High · Hard

p.1593

```
                and 0 <= y < n
                and grid[x][y]
                and uf.union(i * n + j, x * n +
y)
            ):
                cnt -= 1
            ans.append(cnt)
        return ans
```

• SimplyLeet

Python implementation using Union Find with path compression and union by rank.

```
class UnionFind:
```

```
    def __init__(self, size):
        # parent: stores parent of each node. -1
        # indicates water.
        self.parent = [-1] * size
        self.rank = [0] * size
        self.count = 0 # number of islands
```

```
    def find(self, x):
        # Find with path compression.
```

Round 3 · High · Hard

p.1594

```

        if self.parent[x] != x:
            self.parent[x] =
self.find(self.parent[x])
            return self.parent[x]

    def union(self, x, y):
        # Union the sets; if union is successful,
        reduce island count.
        rootX = self.find(x)
        rootY = self.find(y)
        if rootX == rootY:
            return False

        # union by rank.
        if self.rank[rootX] < self.rank[rootY]:
            self.parent[rootX] = rootY
        elif self.rank[rootX] > self.rank[rootY]:
            self.parent[rootY] = rootX
        else:
            self.parent[rootY] = rootX
            self.rank[rootX] += 1
        self.count -= 1
        return True

```

Round 3 · High · Hard

p.1595

```

class Solution:
    def numIslands2(self, m, n, positions):
        uf = UnionFind(m * n)
        result = []
        directions = [(0, 1), (1, 0), (0, -1), (-1,
0)]

        for r, c in positions:
            idx = r * n + c
            # If the position is already land,
            skip.

            if uf.parent[idx] != -1:
                result.append(uf.count)
                continue

            # Mark this cell as land forming a new
            island.

            uf.parent[idx] = idx
            uf.count += 1

            # Union with all 4 neighbors if they
            are already land.
            for dr, dc in directions:

```

Round 3 · High · Hard

p.1596

```

        nr, nc = r + dr, c + dc
        nidx = nr * n + nc
        if 0 <= nr < m and 0 <= nc < n and
uf.parent[nidx] != -1:
            uf.union(idx, nidx)
        result.append(uf.count)
    return result

# Example run:
# s = Solution()
# print(s.numIslands2(3, 3,
[[0,0],[0,1],[1,2],[2,1]]))

```

• LC.ca

```

class Solution:
    def numIslands2(self, m: int, n: int,
positions: List[List[int]]) -> List[int]:
        def check(i, j):
            return 0 <= i < m and 0 <= j < n and
grid[i][j] == 1

        def find(x):
            if p[x] != x:
                p[x] = find(p[x])
            return p[x]

```

```

        p = list(range(m * n))
        grid = [[0] * n for _ in range(m)]
        res = []
        cur = 0 # current island count
        for i, j in positions:
            if grid[i][j] == 1: # already counted,
same as previous island count
                res.append(cur)
                continue
            grid[i][j] = 1
            cur += 1

            for x, y in [(-1, 0), (1, 0), (0, -1),
(0, 1)]:
                if check(i + x, j + y) and find(i *
n + j) != find((i + x) * n + j + y):
                    p[find(i * n + j)] = find((i +
x) * n + j + y)

                    cur -= 1
                    # 1 0 1
                    # 0 1 0
                    # for above, if setting
1st-row 2nd-col 0-to-1
                    # cur will -1 for 3 times
                    # but cur += 1 after grid[i][j]
= 1 so cur final result is 1
                    res.append(cur)

```

```

        return res

#####

class UnionFind(object):
    def __init__(self, m, n):
        self.dad = [i for i in range(m * n)]
        self.rank = [0 for _ in range(m * n)]

    def find(self, x):

        if self.dad[x] != x:
            self.dad[x] = self.find(self.dad[x])
        return self.dad[x]

    def union(self, xy): # xy is a tuple (x,y), with
                        # x and y value inside
        x, y = map(self.find, xy)
        if x == y:
            return False
        if self.rank[x] > self.rank[y]:
            self.dad[y] = x

```

```

        elif self.rank[x] < self.rank[y]:
            self.dad[x] = y
        else:
            self.dad[y] = x # search to the left, to find
                             parent
            self.rank[x] += 1 # now x a parent, so +1 its
                             rank
        return True

class Solution(object):
    def numIslands2(self, m, n, positions):

        """
        :type m: int
        :type n: int
        :type positions: List[List[int]]
        :rtype: List[int]
        """

        uf = UnionFind(m, n)
        ans = 0
        ret = []
        dirs = [(0, -1), (0, 1), (1, 0), (-1, 0)]

```

```

grid = set()
for i, j in positions:
    ans += 1
    grid |= {(i, j)}
    for di, dj in dirs:
        ni, nj = i + di, j + dj
        if 0 <= ni < m and 0 <= nj < n and (ni, nj)
in grid:
    if uf.union((ni * n + nj, i * n + j)):
        ans -= 1
    ret.append(ans)

return ret

#####

class UnionFind:
    def __init__(self, n: int):
        self.p = list(range(n))
        self.size = [1] * n

```

Round 3 · High · Hard

p.1601

```

def find(self, x: int):
    if self.p[x] != x:
        self.p[x] = self.find(self.p[x])
    return self.p[x]

def union(self, a: int, b: int) -> bool:
    pa, pb = self.find(a - 1), self.find(b - 1)
    if pa == pb:
        return False

    if self.size[pa] > self.size[pb]:
        self.p[pb] = pa
        self.size[pa] += self.size[pb]
    else:
        self.p[pa] = pb
        self.size[pb] += self.size[pa]
    return True

class Solution:

```

Round 3 · High · Hard

p.1602

```

def numIslands2(self, m: int, n: int,
positions: List[List[int]]) -> List[int]:
    uf = UnionFind(m * n)
    grid = [[0] * n for _ in range(m)]
    ans = []
    dirs = (-1, 0, 1, 0, -1)
    cnt = 0
    for i, j in positions:
        if grid[i][j]:
            ans.append(cnt)
            continue

        grid[i][j] = 1
        cnt += 1
        for a, b in pairwise(dirs):
            x, y = i + a, j + b
            if (
                0 <= x < m
                and 0 <= y < n
                and grid[x][y]
                and uf.union(i * n + j, x * n +
y)
            ):
                cnt -= 1
        ans.append(cnt)
    return ans

```

◆ Quick Review

Round 3 · High · Hard

p.1603

Key Insights

- Use Union Find (Disjoint Set Union) to efficiently manage connected components (islands).
- Map 2D grid coordinates to a 1D index.
- When adding a new land, check its 4 neighbors (up, down, left, right) and perform unions if they are also lands.
- Avoid duplicate processing if the same cell is added multiple times.
- Maintain a counter for islands that is updated during each union operation.

Space & Time

Time Complexity: $O(K * \alpha(m$

$n))$ where K is the number of operations and α is the inverse Ackermann function (almost constant time per union/find).

Space Complexity: $O(m$

$n)$ for storing the union-find structure.

Solution

The solution relies on the Union Find data structure to dynamically merge islands. When a new land cell is added:

- If the cell is already land, record the current island count.
- Otherwise, mark the cell as land and increment the island count.
- Check the four neighboring cells; if any neighbor is land, attempt to union the current cell with that neighbor.
- If a union actually happens (i.e., the two cells were in separate islands), decrement the island counter.
- Output the island count after each add-land operation. The key trick is converting 2D indices to a 1D representation ($r * n$

Round 3 · High · Hard

p.1604

+ c) and using path compression along with union by rank in the Union Find implementation for optimal performance.

Round 3 · High · Hard

p.1605

Graphs

#1153 String Transforms Into Another String

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Union Find

Lists: Denny Zhang

Problem

Description

Given two strings `str1` and `str2` of the same length, determine whether you can transform `str1` into `str2` by doing zero or more conversions.

In one conversion you can convert all occurrences of one character in `str1` to any other lowercase English character.

Return true if and only if you can transform `str1` into `str2`.

Example 1:

Input: `str1 = "aabcc", str2 = "ccdee"`

Output: true

Explanation: Convert 'c' to 'e' then 'b' to 'd' then 'a' to 'c'. Note that the order of conversions matter.

Example 2:

Round 3 · High · Hard

p.1606

Input: str1 = "leetcode", str2 = "codeleet"

Output: false

Explanation: There is no way to transform str1 to str2.

Constraints:

- $1 \leq \text{str1.length} == \text{str2.length} \leq 104$
- str1 and str2 contain only lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def canConvert(self, str1: str, str2: str) ->
    bool:
        if str1 == str2:
            return True

        mappings = {}

        for a, b in zip(str1, str2):
            if mappings.get(a, b) != b:
                return False
```

```
        mappings[a] = b

        # No letter in the str1 maps to > 1 letter in
        the str2 and there is at
        # lest one temporary letter can break any
        loops.
        return len(set(str2)) < 26
```

• LeetDoocs

```
class Solution:
    def canConvert(self, str1: str, str2: str) ->
    bool:
        if str1 == str2:
            return True
        if len(set(str2)) == 26:
            return False
        d = {}
        for a, b in zip(str1, str2):
            if a not in d:
                d[a] = b

            elif d[a] != b:
                return False
        return True
```

• SimplyLeet


```
# Python solution with line-by-line comments
def canConvert(str1: str, str2: str) -> bool:
    # If both strings are equal, no conversion is
    # needed.
    if str1 == str2:
        return True

    # Dictionary to hold the mapping from
    # characters in str1 to characters in str2
    mapping = {}

    # Iterate through the characters of both
    # strings simultaneously.
```

```
    for c1, c2 in zip(str1, str2):
        # If c1 is already in the mapping, check
        # for consistency.
        if c1 in mapping:
            if mapping[c1] != c2:
                # Found an inconsistent mapping.
                return False
        else:
            mapping[c1] = c2

    # Count distinct characters in str2.
```

```
    # If str2 uses all 26 letters, then there is no
    # extra character available to break cycles.
    if len(set(str2)) == 26:
        return False

    return True

# Example Usage:
#print(canConvert("aabcc", "ccdee")) # Expected
# True
#print(canConvert("leetcode", "codeleet")) #
# Expected False
```

• LC.ca

```
class Solution:
    def canConvert(self, str1: str, str2: str) ->
    bool:
        if str1 == str2:
            return True
        if len(set(str2)) == 26:
            return False
        d = {}
        for a, b in zip(str1, str2):
            if a not in d:
                d[a] = b

            elif d[a] != b:
                return False
        return True
```

◆ Quick Review

Key Insights

- Each character mapping from str1 to str2 must be consistent; one character in str1 always maps to the same character in str2.
- A bijective mapping is not necessary; multiple characters in str1 can map to the same character in str2.
- The order of conversions matters. When there is a cycle (e.g., mapping 'a'→'b' and 'b'→'a'), an intermediate unused character might be required to break the cycle.
- If str2 uses all 26 letters, no temporary character is available to break mapping cycles. In this case, transformation is possible only if str1 is already equal to str2.

Space & Time

Time Complexity: $O(n)$, where n is the length of the strings, because we traverse the strings to build and check the mapping.

Space Complexity: $O(1)$ or $O(26)$, which is constant space, to store the mapping for the 26 lowercase English letters.

Solution

The approach is as follows:

- If str1 equals str2, return true immediately.
- Build a mapping from each character in str1 to the corresponding character in str2. While creating the mapping, check that each character maps consistently.
- Count the number of distinct characters in str2. If this number is 26 and str1 is

not already equal to str2, it is impossible to perform the transformation because there is no spare character available as a temporary placeholder. – If the mapping is consistent and there exists at least one character not used in str2 (i.e., distinct characters in str2 is less than 26), the transformation is possible. The important trick is identifying the need for a free character when the target mapping is “full” (i.e., covers all 26 letters). Without a free character, you cannot avoid cycles where a temporary intermediate is necessary.

BFS

Round 3 · High · Hard · 4 questions

Round 3 · High · Hard

p.1613

BFS

#127 Word Ladder

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · BFS

Lists: Grind 169

Problem

Description

A transformation sequence from word `beginWord` to word `endWord` using a dictionary `wordList` is a sequence of words `beginWord -> s1 -> s2 -> ... -> sk` such that:

- Every adjacent pair of words differs by a single letter.
- Every s_i for $1 \leq i \leq k$ is in `wordList`. Note that `beginWord` does not need to be in `wordList`.
- $s_k == endWord$

Given two words, `beginWord` and `endWord`, and a dictionary `wordList`, return the number of words in the shortest transformation sequence from `beginWord` to `endWord`, or 0 if no such sequence exists.

Example 1:

Round 3 · High · Hard

p.1614

Input: beginWord = "hit", endWord = "cog", wordList = ["hot","dot","dog","lot","log","cog"]

Output: 5

Explanation: One shortest transformation sequence is "hit" -> "hot" -> "dot" -> "dog" -> "cog", which is 5 words long.

Example 2:

Input: beginWord = "hit", endWord = "cog", wordList = ["hot","dot","dog","lot","log"]

Output: 0

Explanation: The endWord "cog" is not in wordList, therefore there is no valid transformation sequence.

Constraints:

- $1 \leq \text{beginWord.length} \leq 10$
- $\text{endWord.length} == \text{beginWord.length}$
- $1 \leq \text{wordList.length} \leq 5000$
- $\text{wordList}[i].\text{length} == \text{beginWord.length}$
- beginWord, endWord, and wordList[i] consist of lowercase English letters.
- $\text{beginWord} \neq \text{endWord}$
- All the words in wordList are unique.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def ladderLength(
        self,
        beginWord: str,
        endWord: str,
        wordList: list[str],
    ) -> int:
        wordSet = set(wordList)
        if endWord not in wordSet:
            return 0
```

```
q = collections.deque([beginWord])

step = 1
while q:
    for _ in range(len(q)):
        wordList = list(q.popleft())
        for i, cache in enumerate(wordList):
            for c in string.ascii_lowercase:
                wordList[i] = c
```

```

        word = ''.join(wordList)
        if word == endWord:
            return step + 1
        if word in wordSet:
            q.append(word)
            wordSet.remove(word)
        wordList[i] = cache
        step += 1

    return 0

```

• LeetCode

```

class Solution:
    def ladderLength(self, beginWord: str,
endWord: str, wordList: List[str]) -> int:
        words = set(wordList)
        q = deque([beginWord])
        ans = 1
        while q:
            ans += 1
            for _ in range(len(q)):
                s = q.popleft()
                s = list(s)

```

```

        for i in range(len(s)):
            ch = s[i]
            for j in range(26):
                s[i] = chr(ord('a') + j)
                t = ''.join(s)
                if t not in words:
                    continue
                if t == endWord:
                    return ans
            q.append(t)

        words.remove(t)
        s[i] = ch

    return 0

```

• SimplyLeet

```

# Python solution using BFS

from collections import defaultdict, deque

def ladderLength(beginWord, endWord, wordList):
    # Create a set for fast lookup
    wordSet = set(wordList)
    if endWord not in wordSet:
        return 0

```

```

# Preprocessing: build a dictionary of
intermediate states to words
L = len(beginWord)
all_combo_dict = defaultdict(list)
for word in wordSet:
    for i in range(L):
        # Key with a wildcard at the i-th
        position
        intermediate = word[:i] + '*' +
        word[i+1:]
        all_combo_dict[intermediate].append(word)

# Initialize BFS

```

```

queue = deque([beginWord, 1])
visited = set([beginWord])

while queue:
    current_word, level = queue.popleft()
    for i in range(L):
        intermediate = current_word[:i] + '*'
        + current_word[i+1:]
        # Check all words sharing the same
        intermediate state
        for word in
        all_combo_dict[intermediate]:
            if word == endWord:

```

```

        return level + 1
        if word not in visited:
            visited.add(word)
            queue.append((word, level +
1))

        # Once processed, clear the
        intermediate state list to save processing time
        all_combo_dict[intermediate] = []
        return 0

# Example usage:
# print(ladderLength("hit", "cog",
["hot", "dot", "dog", "lot", "log", "cog"]))

```

• LC.ca

```

# native BFS
class Solution:
    def ladderLength(self, beginWord: str,
endWord: str, wordList: List[str]) -> int:
        words = set(wordList)
        q = deque([beginWord])
        ans = 1
        while q:
            ans += 1
            for _ in range(len(q)):
                s = q.popleft()

```

```

        s = list(s) # must convert to list,
cannot directly update s[i]='a'
        for i in range(len(s)):
            ch = s[i]
            for j in range(26):
                # will re-generate s

itself
                # but s not in words-set, s
removed after added to words-set
                s[i] = chr(ord('a') + j)
                t = ''.join(s)
                if t not in words:
                    continue

            if t == endWord:
                return ans
            q.append(t)
            words.remove(t) #
equivalent to set t as visited
            s[i] = ch # restore

        return 0

#####

```

```

'''
BFS BFS
1. q1, q2 " " -> " " -> " "
2. m1, m2
3.
4.
'''

class Solution:
    def ladderLength(self, beginWord: str,
endWord: str, wordList: List[str]) -> int:

        def extend(m1, m2, q):
            for _ in range(len(q)):
                s = q.popleft() # so, every time,
starting from start-word... and later if t in m1
will skip repeated part...
                step = m1[s]
                s = list(s) # s = "abc", list(s)
=> ['a', 'b', 'c']
                for i in range(len(s)):
                    ch = s[i]
                    for j in range(26):
                        s[i] = chr(ord('a') + j)
                        t = ''.join(s)

```

```

        if t in m1 or t not in
words:
        continue
        if t in m2:
            return step + 1 +
m2[t]
            m1[t] = step + 1
            q.append(t)
            s[i] = ch
        return -1

words = set(wordList)

```

```

    if endWord not in words:
        return 0
    q1, q2 = deque([beginWord]),
deque([endWord])
    m1, m2 = {beginWord: 0}, {endWord: 0}
    while q1 and q2:
        t = extend(m1, m2, q1) if len(q1) <=
len(q2) else extend(m2, m1, q2)
        if t != -1:
            return t + 1
    return 0

```

Round 3 · High · Hard

p.1623

```

#####

import string
from collections import deque

class Solution(object):
    def ladderLength(self, beginWord, endWord,
wordList):
        """
        :type beginWord: str

        :type endWord: str
        :type wordList: Set[str]
        :rtype: int
        """

    def getNbrs(src, dest, wordList):
        res = []
        for c in string.ascii_lowercase:
            for i in range(0, len(src)):
                newWord = src[:i] + c + src[i + 1:]

```

Round 3 · High · Hard

p.1624


```

        if newWord == src:
            continue
        if newWord in wordList or newWord ==
dest:
            # about yield https://stackoverflow.co
m/questions/231767/what-does-the-yield-keyword-do-
in-python
            # yield is a keyword that is used like
return, except the function will return a
generator.
            yield newWord

queue = deque([beginWord])
length = 0
while queue:

    length += 1
    for k in range(0, len(queue)):
        top = queue.popleft()
        for nbr in getNbrs(top, endWord,
wordList):
            wordList.remove(nbr)
            if nbr == endWord:
                return length + 1
            queue.append(nbr)
    return 0

```

Round 3 · High · Hard

p.1625

```

#####

while q1 and q2:
    if len(q1) <= len(q2):
        # Prioritize the queue with fewer elements
        for expansion
            extend(m1, m2, q1)
        else:
            extend(m2, m1, q2)

def extend(m1, m2, q):
    # New round of expansion
    for _ in range(len(q)):
        p = q.popleft()
        step = m1[p]
        for t in next(p):
            if t in m1:
                # Already visited before
                continue
            if t in m2:

                # The other direction has been
                searched, indicating that a shortest path has been
                found

                return step + 1 + m2[t]
            q.append(t)
            m1[t] = step + 1

```

◆ Quick Review

Round 3 · High · Hard

p.1626

Key Insights

- Use Breadth-First Search (BFS) to explore transformations level by level to guarantee the shortest path.
- Preprocess words by generating intermediate representations (e.g., replacing each letter with a wildcard) to quickly find neighbors.
- Use a visited set to avoid cycles.
- Each transformation counts as a level in the BFS, so the answer is the number of levels traversed.

Space & Time

Time Complexity: $O(N * M^2)$ where N is the number of words and M is the length of each word.

Space Complexity: $O(N * M)$ for the storage of intermediate mappings and the BFS queue.

Solution

We solve the problem using a BFS approach. First, we preprocess the wordList by creating a mapping of all possible intermediate states. For example, for the word "dog" and wildcard position, we generate " og", "d g", and "do*". This helps us find all adjacent words (neighbors) that are only one letter different in constant time. We then initialize a queue with the beginWord and initiate the BFS. For each word dequeued, we explore all valid transformations using the precomputed intermediate mapping. If the endWord is reached during this process, we return the current level (transformation sequence length). Otherwise, we continue the BFS until all possibilities are exhausted. If no valid transformation sequence is found, we

return 0.

BFS

#317 Shortest Distance from All Buildings

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Premium 98

Problem

Description

You are given an $m \times n$ grid of values 0, 1, or 2, where:

- each 0 marks an empty land that you can pass by freely,
- each 1 marks a building that you cannot pass through, and
- each 2 marks an obstacle that you cannot pass through.

You want to build a house on an empty land that reaches all buildings in the shortest total travel distance. You can only move up, down, left, and right.

Return the shortest travel distance for such a house. If it is not possible to build such a house according to the above rules, return -1.

The total travel distance is the sum of the distances between the houses of the friends and the meeting

point.

Example 1:

1	0	2	0	1
0	0	0	0	0
0	0	1	0	0

Input: grid = [[1,0,2,0,1],[0,0,0,0,0],[0,0,1,0,0]]

Output: 7

Explanation: Given three buildings at (0,0), (0,4), (2,2), and an obstacle at (0,2).

The point (1,2) is an ideal empty land to build a house, as the total travel distance of 3+3+1=7 is minimal.

So return 7.

Example 2:

Input: grid = [[1,0]]

Output: 1

Example 3:

Input: grid = [[1]]

Output: -1

Constraints:

- $m == \text{grid.length}$
- $n == \text{grid}[i].\text{length}$
- $1 \leq m, n \leq 50$
- $\text{grid}[i][j]$ is either 0, 1, or 2.
- There will be at least one building in the grid.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def shortestDistance(self, grid:
list[list[int]]) -> int:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))
        m = len(grid)
        n = len(grid[0])
        nBuildings = sum(a == 1 for row in grid for a
in row)
        ans = math.inf
        # dist[i][j] := the total distance of
grid[i][j] (0) to reach all the
# buildings (1)
        dist = [[0] * n for _ in range(m)]
```

```
# reachCount[i][j] := the number of buildings
(1) grid[i][j] (0) can reach
reachCount = [[0] * n for _ in range(m)]
```

```
def bfs(row: int, col: int) -> bool:
    q = collections.deque([(row, col)])
    seen = {(row, col)}
    seenBuildings = 1
```

```
    step = 1
    while q:
```

```
        for _ in range(len(q)):
            i, j = q.popleft()
            for dx, dy in DIRS:
                x = i + dx
                y = j + dy
                if x < 0 or x == m or y < 0 or y == n:
                    continue
                if (x, y) in seen:
                    continue
                seen.add((x, y))
```

```

        if not grid[x][y]:
            dist[x][y] += step
            reachCount[x][y] += 1
            q.append((x, y))
        elif grid[x][y] == 1:
            seenBuildings += 1
        step += 1

```

```

# True if all the buildings (1) are connected
return seenBuildings == nBuildings

```

```

for i in range(m):
    for j in range(n):
        if grid[i][j] == 1: # BFS from this
            building.
            if not bfs(i, j):
                return -1

```

```

for i in range(m):
    for j in range(n):
        if reachCount[i][j] == nBuildings:

```

```

        ans = min(ans, dist[i][j])

```

```

return -1 if ans == math.inf else ans

```

• LeetCode

```

class Solution:
    def shortestDistance(self, grid:
        List[List[int]]) -> int:
        m, n = len(grid), len(grid[0])
        q = deque()
        total = 0
        cnt = [[0] * n for _ in range(m)]
        dist = [[0] * n for _ in range(m)]
        for i in range(m):
            for j in range(n):
                if grid[i][j] == 1:

```

```

                    total += 1
                    q.append((i, j))
                    d = 0
                    vis = set()
                    while q:
                        d += 1
                        for _ in range(len(q)):
                            r, c = q.popleft()
                            for a, b in [[0, 1],
                                [0, -1], [1, 0], [-1, 0]]:
                                x, y = r + a, c + b

```

```

                                if (
<= x < m                                and 0 <= y < n
                                and grid[x][y]
== 0                                and (x, y) not
in vis                                ):
                                cnt[x][y] += 1
                                dist[x][y] +=
d                                q.append((x,
y))                                vis.add((x,
y))

    ans = inf
    for i in range(m):
        for j in range(n):
            if grid[i][j] == 0 and cnt[i][j] ==
total:
                ans = min(ans, dist[i][j])
    return -1 if ans == inf else ans

```

• SimplyLeet

Round 3 · High · Hard

p.1635

```

from collections import deque

def shortestDistance(grid):
    if not grid or not grid[0]:
        return -1

    m, n = len(grid), len(grid[0])
    total_buildings = 0
    total_distance = [[0] * n for _ in range(m)]
    reachable_count = [[0] * n for _ in range(m)]

    def bfs(start_i, start_j):
        visited = [[False] * n for _ in range(m)]
        q = deque([(start_i, start_j, 0)])
        visited[start_i][start_j] = True

        while q:
            i, j, dist = q.popleft()
            for x, y in [(i-1, j), (i+1, j), (i,
j-1), (i, j+1)]:
                if 0 <= x < m and 0 <= y < n and
not visited[x][y] and grid[x][y] == 0:

```

Round 3 · High · Hard

p.1636

```

        visited[x][y] = True
        total_distance[x][y] += dist +
1
        reachable_count[x][y] += 1
        q.append((x, y, dist+1))

    for i in range(m):
        for j in range(n):
            if grid[i][j] == 1:
                total_buildings += 1
                bfs(i, j)

    min_distance = float('inf')
    for i in range(m):
        for j in range(n):
            if grid[i][j] == 0 and
            reachable_count[i][j] == total_buildings:
                min_distance = min(min_distance,
                total_distance[i][j])

    return min_distance if min_distance !=
    float('inf') else -1

# Example test:

print(shortestDistance([[1,0,2,0,1],[0,0,0,0,0],[0
,0,1,0,0]]))

```

• LC.ca

Round 3 · High · Hard

p.1637

```

class Solution:
    def shortestDistance(self, grid:
    List[List[int]]) -> int:
        m, n = len(grid), len(grid[0])
        q = deque()
        total = 0 # total building count
        # cnt: how many buildings can reach (x,y)
        # if there is a column all blockers, then
        (x,y) cannot reach every building
        cnt = [[0] * n for _ in range(m)]
        dist = [[0] * n for _ in range(m)]
        for i in range(m):
            for j in range(n):
                if grid[i][j] == 1:
                    total += 1
                    q.append((i, j))
                    d = 0
                    vis = set() # reset for every
                    free-land

                    while q:
                        d += 1
                        for _ in range(len(q)):
                            r, c = q.popleft()

```

Round 3 · High · Hard

p.1638

```

        for a, b in [[0, 1],
[0, -1], [1, 0], [-1, 0]]:
            x, y = r + a, c + b
            if (
                0 <= x < m
                and 0 <= y < n
                and grid[x][y]
== 0
                and (x, y) not
in vis
            ):
                cnt[x][y] += 1
                dist[x][y] +=
d
            q.append((x,
y))
            vis.add((x,
y))
    ans = inf
    for i in range(m):
        for j in range(n):
            if grid[i][j] == 0 and cnt[i][j] ==
total:
                ans = min(ans, dist[i][j])
    return -1 if ans == inf else ans

```

◆ Quick Review

Key Insights

- We need to compute the Manhattan distance from each building (cell value 1) to all empty lands (cell value 0) by performing a BFS.
- Maintain two matrices: one for cumulative total distances and another for the count of buildings that reached a particular empty cell.
- Ensure that only those empty cells reachable by every building are considered.
- The ideal cell will have a reach count equal to the total number of buildings.
- If multiple buildings cannot all reach any specific empty cell, then return -1.

Space & Time

Time Complexity: $O(k * m * n)$ where k is the number of buildings and m, n are the grid dimensions.

Space Complexity: $O(m * n)$ due to the auxiliary matrices and visited state tracking used in BFS.

Solution

We first iterate through the grid to locate all buildings while counting them. For each building, we perform a BFS to compute the shortest distance to each reachable empty cell. During the BFS, we update the cumulative distance for that cell and record that it was reached by one additional building. After processing all the buildings, we scan the grid to find the empty cell that is reachable from all buildings and has the minimum cumulative distance. If no such cell exists, we return -1.

BFS

#773 Sliding Puzzle

HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · BFS · Matrix

Lists: Denny Zhang

Problem

Description

On an 2×3 board, there are five tiles labeled from 1 to 5, and an empty square represented by 0. A move consists of choosing 0 and a 4-directionally adjacent number and swapping it.

The state of the board is solved if and only if the board is $[[1,2,3],[4,5,0]]$.

Given the puzzle board `board`, return the least number of moves required so that the state of the board is solved. If it is impossible for the state of the board to be solved, return -1 .

Example 1:

1	2	3
4	0	5

Input: `board = [[1,2,3],[4,0,5]]`

Output: 1

Explanation: Swap the 0 and the 5 in one move.

Example 2:

1	2	3
5	4	

Input: board = [[1,2,3],[5,4,0]]

Output: -1

Explanation: No number of moves will make the board solved.

Example 3:

4	1	2
5		3

Input: board = [[4,1,2],[5,0,3]]

Output: 5

Explanation: 5 is the smallest number of moves that solves the board.

An example path:

After move 0: [[4,1,2],[5,0,3]]

After move 1: [[4,1,2],[0,5,3]]

After move 2: [[0,1,2],[4,5,3]]

After move 3: [[1,0,2],[4,5,3]]

Constraints:

- board.length == 2
- board[i].length == 3
- 0 <= board[i][j] <= 5
- Each value board[i][j] is unique.

Community Solutions (Python)

• LeetCode

```
class Solution:
    def slidingPuzzle(self, board:
List[List[int]]) -> int:
        t = [None] * 6

        def gets():
            for i in range(2):
                for j in range(3):
                    t[i * 3 + j] = str(board[i][j])
            return ''.join(t)

        def setb(s):
            for i in range(2):
                for j in range(3):
                    board[i][j] = int(s[i * 3 + j])

        def f():
            res = []
            i, j = next((i, j) for i in range(2)
for j in range(3) if board[i][j] == 0)
            for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:
                x, y = i + a, j + b
```

```
        if 0 <= x < 2 and 0 <= y < 3:
            board[i][j], board[x][y] =
board[x][y], board[i][j]
            res.append(gets())
            board[i][j], board[x][y] =
board[x][y], board[i][j]
            return res

        start = gets()
        end = "123450"
        if start == end:
            return 0

        vis = {start}
        q = deque([start])
        ans = 0
        while q:
            ans += 1
            for _ in range(len(q)):
                x = q.popleft()
                setb(x)
                for y in f():
                    if y == end:

                        return ans
                    if y not in vis:
                        vis.add(y)
                        q.append(y)

        return -1
```

```

class Solution:
    def slidingPuzzle(self, board:
List[List[int]]) -> int:
        m, n = 2, 3
        seq = []
        start, end = '', '123450'
        for i in range(m):
            for j in range(n):
                if board[i][j] != 0:
                    seq.append(board[i][j])
                    start += str(board[i][j])

```

```

        def check(seq):
            n = len(seq)
            cnt = sum(seq[i] > seq[j] for i in
range(n) for j in range(i, n))
            return cnt % 2 == 0

```

```

        def f(s):
            ans = 0
            for i in range(m * n):
                if s[i] != '0':

```

```

                    num = ord(s[i]) - ord('1')
                    ans += abs(i // n - num // n) +
abs(i % n - num % n)
                    return ans

```

```

            if not check(seq):
                return -1
            q = [(f(start), start)]
            dist = {start: 0}
            while q:
                _, state = heappop(q)

```

```

            if state == end:
                return dist[state]
            p1 = state.index('0')
            i, j = p1 // n, p1 % n
            s = list(state)
            for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n:
                    p2 = x * n + y
                    s[p1], s[p2] = s[p2], s[p1]

```

```

        next = ''.join(s)
        s[p1], s[p2] = s[p2], s[p1]
        if next not in dist or
dist[next] > dist[state] + 1:
            dist[next] = dist[state] +
1
            heappush(q, (dist[next] +
f(next), next))
        return -1

```

• SimplyLeet

```

import collections

def slidingPuzzle(board):
    # Flatten board into a string representation.
    start = ''.join(str(num) for row in board for
num in row)
    goal = "123450"
    # Mapping of indices to their neighbors in a
2x3 board.
    neighbors = {0: [1, 3], 1: [0,2,4], 2: [1,5],
3: [0,4], 4: [1,3,5], 5: [2,4]}

```

```

# Initialize BFS with the starting state and
move count 0.
q = collections.deque([(start, 0)])
visited = set([start])

```

```

# Process states in the queue.
while q:
    state, moves = q.popleft()
    if state == goal:
        return moves
    zero_index = state.index('0')

```

```

# Try all valid moves by swapping 0 with
its neighbours.
for neighbor in neighbors[zero_index]:
    new_state = list(state)
    new_state[zero_index],
new_state[neighbor] = new_state[neighbor],
new_state[zero_index]
    new_state_str = ''.join(new_state)
    if new_state_str not in visited:
        visited.add(new_state_str)
        q.append((new_state_str, moves +
1))
    return -1

```

```

# Example usage:
print(slidingPuzzle([[1,2,3],[4,0,5]]))

```

• LC.ca

```

class Solution:
    def slidingPuzzle(self, board:
List[List[int]]) -> int:
        m, n = 2, 3
        seq = []
        start, end = '', '123450'
        for i in range(m):
            for j in range(n):
                if board[i][j] != 0:
                    seq.append(board[i][j])
                    start += str(board[i][j])

```

```

        def check(seq):
            n = len(seq)
            cnt = sum(seq[i] > seq[j] for i in
range(n) for j in range(i, n))
            return cnt % 2 == 0

```

```

        def f(s):
            ans = 0
            for i in range(m * n):
                if s[i] != '0':

```

```

                    num = ord(s[i]) - ord('1')
                    ans += abs(i // n - num // n) +
abs(i % n - num % n)
                    return ans

```

```

            if not check(seq):
                return -1
            q = [(f(start), start)]
            dist = {start: 0}
            while q:
                _, state = heappop(q)

```

```

            if state == end:
                return dist[state]
            p1 = state.index('0')
            i, j = p1 // n, p1 % n
            s = list(state)
            for a, b in [[0, -1], [0, 1], [1, 0],
[-1, 0]]:

```

```

                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n:
                    p2 = x * n + y
                    s[p1], s[p2] = s[p2], s[p1]

```

```

        next = ''.join(s)
        s[p1], s[p2] = s[p2], s[p1]
        if next not in dist or
dist[next] > dist[state] + 1:
            dist[next] = dist[state] +
1
            heappush(q, (dist[next] +
f(next), next))
        return -1

```

◆ Quick Review

Key Insights

- Represent the board as a string (e.g. "123450") to simplify state comparison and manipulation.
- Use Breadth First Search (BFS) to explore all possible moves since the puzzle has a small fixed state space ($6! = 720$ possible states).
- Precompute the valid neighbor indices for each board position to efficiently generate new states.
- Track visited states to avoid loops and redundant work.

Space & Time

Time Complexity: $O(1)$ in practice, since there are at most 720 states ($6!$ permutations).

Space Complexity: $O(1)$ for the same reason, as the state space is fixed and limited.

Solution

We solve the problem using a Breadth First Search (BFS) approach. The board configuration is flattened into a string to represent states. A mapping from each index to its possible neighbors (i.e., positions to which 0 can move) is precomputed. BFS is then used to explore every valid move from the initial state while marking states as visited to avoid cycles. When the goal state ("123450") is reached, we return the number of moves taken. If BFS completes without reaching the goal, the puzzle is unsolvable and we return -1 .

BFS

#815 Bus Routes

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · BFS

Lists: Grind 169

Problem

Description

You are given an array `routes` representing bus routes where `routes[i]` is a bus route that the *i*th bus repeats forever.

- For example, if `routes[0] = [1, 5, 7]`, this means that the 0th bus travels in the sequence `1 -> 5 -> 7 -> 1 -> 5 -> 7 -> 1 -> ...` forever.

You will start at the bus stop `source` (You are not on any bus initially), and you want to go to the bus stop `target`. You can travel between bus stops by buses only.

Return the least number of buses you must take to travel from `source` to `target`. Return `-1` if it is not possible.

Example 1:

Input: `routes = [[1,2,7],[3,6,7]]`, `source = 1`, `target = 6`

Output: 2

Explanation: The best strategy is take the first bus to the bus stop 7, then take the second bus to the bus stop 6.

Example 2:

Input: `routes = [[7,12],[4,5,15],[6],[15,19],[9,12,13]]`, `source = 15`, `target = 12`

Output: -1

Constraints:

- `1 <= routes.length <= 500`.
- `1 <= routes[i].length <= 105`
- All the values of `routes[i]` are unique.
- `sum(routes[i].length) <= 105`
- `0 <= routes[i][j] < 106`
- `0 <= source, target < 106`

Community Solutions (Python)

- WalkCC


```

class Solution:
    def numBusesToDestination(
        self,
        routes: list[list[int]],
        source: int,
        target: int,
    ) -> int:
        if source == target:
            return 0

        graph = collections.defaultdict(list)
        usedBuses = set()

        for i in range(len(routes)):
            for route in routes[i]:
                graph[route].append(i)

        q = collections.deque([source])

        step = 1

```

Round 3 · High · Hard

p.1657

```

while q:
    for _ in range(len(q)):
        for bus in graph[q.popleft()]:
            if bus in usedBuses:
                continue
            usedBuses.add(bus)
            for nextRoute in routes[bus]:
                if nextRoute == target:
                    return step
                q.append(nextRoute)

    step += 1

return -1

```

• LeetCode

```

class Solution:
    def numBusesToDestination(
        self, routes: List[List[int]], source:
        int, target: int
    ) -> int:
        if source == target:
            return 0

        g = defaultdict(list)
        for i, route in enumerate(routes):
            for stop in route:
                g[stop].append(i)

```

Round 3 · High · Hard

p.1658

```

if source not in g or target not in g:
    return -1
q = [(source, 0)]
vis_bus = set()
vis_stop = {source}
for stop, bus_count in q:
    if stop == target:
        return bus_count
    for bus in g[stop]:
        if bus not in vis_bus:

```

```

                vis_bus.add(bus)
                for next_stop in routes[bus]:
                    if next_stop not in
vis_stop:
                    vis_stop.add(next_stop
)
                    q.append((next_stop,
bus_count + 1))
            return -1

```

• SimplyLeet

Round 3 · High · Hard

p.1659

```

from collections import defaultdict, deque

def numBusesToDestination(routes, source, target):
    # If source equals target, no buses are needed.
    if source == target:
        return 0

    # Build mapping: bus stop -> list of bus
    indices
    stop_to_buses = defaultdict(list)
    for bus_index, stops in enumerate(routes):

        for stop in stops:
            stop_to_buses[stop].append(bus_index)

    visited_buses = set()
    visited_stops = set([source])
    queue = deque()

    # Initialize queue with all buses that include
    the source stop.
    for bus in stop_to_buses[source]:
        visited_buses.add(bus)

```

Round 3 · High · Hard

p.1660

```

        queue.append((bus, 1))

# Perform BFS over bus routes.
while queue:
    current_bus, buses_taken = queue.popleft()
    for stop in routes[current_bus]:
        if stop == target:
            return buses_taken
        if stop not in visited_stops:
            visited_stops.add(stop)

            for neighbor_bus in
stop_to_buses[stop]:
                if neighbor_bus not in
visited_buses:
                    visited_buses.add(neighbor
_bus)
                    queue.append((neighbor_bus
, buses_taken + 1))
            return -1

# Example usage:
# routes = [[1,2,7],[3,6,7]]
# print(numBusesToDestination(routes, 1, 6))

```

• LC.ca

Round 3 · High · Hard

p.1661

```

class Solution:
    def numBusesToDestination(
        self, routes: List[List[int]], source:
int, target: int
    ) -> int:
        if source == target:
            return 0

        s = [set(r) for r in routes]
        d = defaultdict(list)
        for i, r in enumerate(routes):

            for v in r:
                d[v].append(i)

        g = defaultdict(list)
        for ids in d.values():
            m = len(ids)
            for i in range(m):
                for j in range(i + 1, m):
                    a, b = ids[i], ids[j]
                    g[a].append(b)

```

Round 3 · High · Hard

p.1662

```

                g[b].append(a)
q = deque(d[source])
ans = 1
vis = set(d[source])
while q:
    for _ in range(len(q)):
        i = q.popleft()
        if target in s[i]:
            return ans
        for j in g[i]:
            if j not in vis:
                vis.add(j)
                q.append(j)
    ans += 1
return -1

```

◆ Quick Review

Key Insights

- Model the problem as a graph where nodes are bus routes.
- Two bus routes are connected if they share at least one common bus stop.
- Use a hash map to quickly look up bus routes available at any bus stop.
- Employ Breadth-First Search (BFS) on the bus routes graph to find the minimum number of bus transfers.

- Mark visited bus routes and bus stops to avoid redundant work.

Space & Time

Time Complexity: $O(N * L)$ in the worst-case scenario, where N is the number of bus routes and L is the average number of stops per route.

Space Complexity: $O(N * L)$ due to storage of the stop-to-buses mapping and BFS auxiliary data structures.

Solution

We begin by constructing a mapping from each bus stop to the list of buses passing through it. This allows quick retrieval of buses available at a given stop. Starting at the source stop, initialize the BFS queue with all buses that can be boarded at that stop. Then, for each bus route in the queue, inspect each stop along that route. If a stop is the target, return the number of buses taken so far. Otherwise, for each unvisited stop, add all adjacent bus routes (buses that pass through this stop) to the BFS queue with an incremented bus count. Continue until the queue is empty, which means the target cannot be reached.

Round 4 · Mid · Easy

Round 4

Mid Priority · Easy

12 questions · 4 patterns

Linked List #21 #141 #206 #234 #706 #876 #1474

Stack #20 #232 #844

Trie #14

Greedy #409

Linked List

Round 4 · Mid · Easy · 7 questions

Round 3 · High · Hard

p.1665

Round 4 · Mid · Easy

p.1666

Linked List

#21 Merge Two Sorted Lists

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Recursion

Lists: Grind 169

Problem


Description

You are given the heads of two sorted linked lists list1 and list2.

Merge the two lists into one sorted list. The list should be made by splicing together the nodes of the first two lists.

Return the head of the merged linked list.

Example 1:



Round 4 · Mid · Easy

p.1667

Input: list1 = [1,2,4], list2 = [1,3,4]
 Output: [1,1,2,3,4,4]

Example 2:

Input: list1 = [], list2 = []
 Output: []

Example 3:

Input: list1 = [], list2 = [0]
 Output: [0]

Constraints:

- The number of nodes in both lists is in the range [0, 50].
- $-100 \leq \text{Node.val} \leq 100$
- Both list1 and list2 are sorted in non-decreasing order.

Community Solutions (Python)

- WalkCC

Round 4 · Mid · Easy

p.1668

```
class Solution:
    def mergeTwoLists(
        self,
        list1: ListNode | None,
        list2: ListNode | None,
    ) -> ListNode | None:
        if not list1 or not list2:
            return list1 if list1 else list2
        if list1.val > list2.val:
            list1, list2 = list2, list1
```

```
        list1.next = self.mergeTwoLists(list1.next,
list2)
        return list1
```

• LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def mergeTwoLists(
        self, list1: Optional[ListNode], list2:
Optional[ListNode]
    ) -> Optional[ListNode]:
        if list1 is None or list2 is None:
```

```
            return list1 or list2
            if list1.val <= list2.val:
                list1.next =
self.mergeTwoLists(list1.next, list2)
                return list1
            else:
                list2.next = self.mergeTwoLists(list1,
list2.next)
                return list2
```

• SimplyLeet

```
# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

def mergeTwoLists(list1: ListNode, list2:
ListNode) -> ListNode:
    # Create a dummy node to serve as the start of
the merged list
    dummy = ListNode(-1)
    current = dummy
```

```

# Iterate while both lists have nodes
while list1 and list2:
    # Compare the values of the current nodes
    of both lists
    if list1.val <= list2.val:
        current.next = list1 # Link the smaller
        node to the merged list
        list1 = list1.next # Move the pointer
        in list1
    else:
        current.next = list2
        list2 = list2.next

```

```

    current = current.next # Move the pointer
    in the merged list

```

```

# Once one list is exhausted, attach the
remaining nodes from the other list

```

```

if list1:
    current.next = list1
else:
    current.next = list2

```

```

# The merged list head is at dummy.next
return dummy.next

```

• LC.ca

Round 4 · Mid · Easy

p.1671

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
# without ops after while loop
class Solution:
    def mergeTwoLists(self, list1:
Optional[ListNode], list2: Optional[ListNode]) ->
Optional[ListNode]:
        dummy = ListNode()
        current = dummy

```

```

        while list1 or list2:
            v1 = list1.val if list1 else
float('inf')
            v2 = list2.val if list2 else
float('inf')

```

```

            if v1 < v2:
                current.next = list1
                list1 = list1.next
            else:
                current.next = list2
                list2 = list2.next

```

Round 4 · Mid · Easy

p.1672


```
        current = current.next
```

```
    return dummy.next
```

```
#####
```

```
# Definition for singly-linked list.
```

```
# class ListNode:
```

```
# def __init__(self, val=0, next=None):
```

```
# self.val = val
```

```
# self.next = next
```

```
class Solution:
```

```
    def mergeTwoLists(
```

```
        self, list1: Optional[ListNode], list2:
```

```
Optional[ListNode]
```

```
    ) -> Optional[ListNode]:
```

```
        dummy = ListNode()
```

```
        curr = dummy
```

```
        while list1 and list2:
```

```
            if list1.val <= list2.val:
```

```
                curr.next = list1
```

```
        list1 = list1.next
```

```
    else:
```

```
        curr.next = list2
```

```
        list2 = list2.next
```

```
        curr = curr.next
```

```
    curr.next = list1 or list2
```

```
    return dummy.next
```

```
"""
```

```
    curr.next = list1 or list2
```

```
better than:
```

```
    if list1:
```

```
        current.next = list1
```

```
    if list2:
```

```
        current.next = list2
```

```
"""
```

```
#####
```

```
# recursion
```

```
class ListNode:
```

```
    def __init__(self, val=0, next=None):
```

```
        self.val = val
```

```
        self.next = next
```

```
class Solution:
```

```
    def mergeTwoLists(self, l1: ListNode, l2:
ListNode) -> ListNode:
```

```
        if not l1: # If l1 is empty, return l2
```

```
            return l2
```

```
        if not l2: # If l2 is empty, return l1
```

```
            return l1
```

```
        if l1.val < l2.val:
```

```
            l1.next = self.mergeTwoLists(l1.next,
```

```
l2)
```

```
            return l1
```

```
        else:
```

```
            l2.next = self.mergeTwoLists(l1,
```

```
l2.next)
```

```
            return l2
```

◆ Quick Review

Key Insights

- Leverage the fact that both lists are sorted.
- Use two pointers (one for each list) to compare nodes sequentially.
- Create a dummy node to simplify edge cases and build the resulting list.
- Iterate until one list is exhausted, then append the remaining nodes from the other list.
- There is also a recursive approach that can simplify the logic but may use more stack space.

Space & Time

Time Complexity: $O(n + m)$, where n and m are the lengths of the two lists.

Space Complexity: $O(1)$ for the iterative solution ($O(n + m)$ stack space for the recursive solution).

Solution

We use an iterative approach with a dummy node to merge two sorted linked lists. Two pointers traverse list1 and list2; at each step, the node with the smaller value is appended to our merged list. This process repeats until one of the lists is exhausted, after which we append the remaining nodes from the other list. This approach works in one pass over both lists, ensuring linear time complexity and constant extra space usage.

Linked List

#141 Linked List Cycle

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Linked List · Two Pointers

Lists: Grind 169

Problem


Description

Given head, the head of a linked list, determine if the linked list has a cycle in it.

There is a cycle in a linked list if there is some node in the list that can be reached again by continuously following the next pointer. Internally, pos is used to denote the index of the node that tail's next pointer is connected to. Note that pos is not passed as a parameter.

Return true if there is a cycle in the linked list. Otherwise, return false.

Example 1:




Input: head = [3,2,0,-4], pos = 1

Output: true

Explanation: There is a cycle in the linked list, where the tail connects to the 1st node (0-indexed).


Example 2:



Input: head = [1,2], pos = 0

Output: true

Explanation: There is a cycle in the linked list, where the tail connects to the 0th node.

Example 3:

Input: head = [1], pos = -1

Output: false

Explanation: There is no cycle in the linked list.

Constraints:

- The number of the nodes in the list is in the range [0, 104].
- $-105 \leq \text{Node.val} \leq 105$
- pos is -1 or a valid index in the linked-list.

Follow up: Can you solve it using $O(1)$ (i.e. constant) memory?

Community Solutions (Python)

- WalkCC

```
class Solution:
    def hasCycle(self, head: ListNode) -> bool:
        slow = head
        fast = head
```

```
        while fast and fast.next:
            slow = slow.next
            fast = fast.next.next
            if slow == fast:
                return True
```

```
        return False
```

• LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, x):
#         self.val = x
#         self.next = None
```

```
class Solution:
    def hasCycle(self, head: Optional[ListNode])
    -> bool:
        s = set()
```

```

while head:
    if head in s:
        return True
    s.add(head)
    head = head.next
return False

```

• SimplyLeet

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next

class Solution:
    def hasCycle(self, head: Optional[ListNode])
-> bool:
    # Initialize two pointers, slow and fast,
    starting from the head.
    slow = head

```

```

fast = head

# Traverse the list until fast pointer
reaches the end.
while fast and fast.next:
    slow = slow.next # Move slow pointer
    one step.
    fast = fast.next.next # Move fast
    pointer two steps.

    # If both pointers meet, there is a
    cycle.
    if slow == fast:
        return True

# If loop terminates, there is no cycle.
return False

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, x):
# self.val = x
# self.next = None

class Solution:
    def hasCycle(self, head: ListNode) -> bool:
        slow = fast = head
        while fast and fast.next:

```

```

slow, fast = slow.next, fast.next.next
if slow == fast:
    return True
return False

```

◆ Quick Review

Key Insights

- Use two pointers (slow and fast) to traverse the list.
- The fast pointer moves two steps at a time and the slow pointer moves one step.
- If the fast pointer ever meets the slow pointer, a cycle exists.
- If the fast pointer reaches the end of the list (null), the list has no cycle.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes in the linked list.

Space Complexity: $O(1)$ as only two pointers are used.

Solution

We use the Floyd's Cycle Detection algorithm (Tortoise and Hare approach). Two pointers, slow and fast, are initialized at the head of the list. The slow pointer moves one step at a time while the fast pointer moves two steps. If the linked list has a cycle, the fast pointer will eventually meet the slow pointer. If the fast pointer reaches the end (i.e., a null pointer), the linked list does not have a cycle.

Linked List

#206 Reverse Linked List

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Recursion


Lists: Grind 169

Problem

Description

Given the head of a singly linked list, reverse the list, and return the reversed list.


Example 1:



Input: head = [1,2,3,4,5]

Output: [5,4,3,2,1]

Example 2:



Input: head = [1,2]
Output: [2,1]

Example 3:

Input: head = []
Output: []

Constraints:

- The number of nodes in the list is the range [0, 5000].
- $-5000 \leq \text{Node.val} \leq 5000$

Follow up: A linked list can be reversed either iteratively or recursively. Could you implement both?

Community Solutions (Python)

• WalkCC

```

class Solution:
    def reverseList(self, head: ListNode | None) ->
    ListNode | None:
        if not head or not head.next:
            return head

        newHead = self.reverseList(head.next)
        head.next.next = head
        head.next = None
        return newHead
  
```

```

class Solution:
    def reverseList(self, head: ListNode | None) ->
    ListNode | None:
        prev = None
        curr = head

        while curr:
            next = curr.next
            curr.next = prev
            prev = curr
            curr = next
  
```

return prev

• LeetDoocs

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def reverseList(self, head: ListNode) ->
    ListNode:
        dummy = ListNode()
        curr = head
        while curr:
            next = curr.next
            curr.next = dummy.next
            dummy.next = curr
            curr = next
        return dummy.next
```

• SimplyLeet

```
# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val # Store the node value
        self.next = next # Pointer to the next node

# Iterative approach to reverse the linked list.
def reverseList(head):
    prev = None # Initially, there is no previous
    node
    current = head # Start with the head of the
    list

    while current:
        next_node = current.next # Temporarily
        store the next node
        current.next = prev # Reverse the current
        node's pointer
        prev = current # Move previous pointer one
        step forward
        current = next_node # Move to the next node
        in the list
    return prev # The new head of the reversed list

# Recursive approach to reverse the linked list.
def reverseListRecursive(head):
    # Base case: empty list or single node
```



```

if not head or not head.next:
    return head
new_head = reverseListRecursive(head.next) #
Reverse the rest of the list
head.next.next = head # Reverse the pointer for
the current node
head.next = None # Set the next pointer of the
current node to None
return new_head # Return the new head of the
reversed list

```

- LC.ca

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def reverseList(self, head:
Optional[ListNode]) -> Optional[ListNode]:
        pre, p = None, head
        while p:
            q = p.next

```

```

p.next = pre
pre = p
p = q
return pre

```

```

class Solution(object): # recursive
    def reverseList(self, root):
        if not root or not root.next:
            return root

```

```

ret = self.reverseList(root.next)
root.next.next = root # root.next is end of the
newly reversed list ret
root.next = None
return ret

```

```

class Solution:
    def reverseList(self, head: ListNode) ->
ListNode:
        dummy = ListNode()

```

```
curr = head
while curr:
    next = curr.next
    curr.next = dummy.next
    dummy.next = curr
    curr = next
return dummy.next
```

◆ Quick Review

Key Insights

- The problem can be solved using either an iterative or recursive approach.
- Iterative method uses constant space by updating pointers in-place.
- Recursive method leverages the call stack to reverse the list, which could use $O(n)$ space in the worst case.
- Both approaches traverse each node exactly once, yielding linear time complexity.

Space & Time

Time Complexity: $O(n)$ for both iterative and recursive approaches, where n is the number of nodes.

Space Complexity: $O(1)$ for the iterative approach and $O(n)$ for the recursive approach because of the recursion call stack.

Solution

We can solve the problem by iterating through the list while reversing the next pointers of each node. We maintain two

pointers (previous and current) and update them as we traverse the list, so that by the end, the previous pointer is at the new head of the reversed list.

Alternatively, a recursive approach reverses the rest of the list first and then adjusts the pointers so that the original head becomes the tail.

Linked List

#234 Palindrome Linked List

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers · Recursion


Lists: Grind 169

Problem

Description


Given the head of a singly linked list, return true if it is a palindrome or false otherwise.

Example 1:



Input: head = [1,2,2,1]
Output: true

Example 2:



Input: head = [1,2]
Output: false

Constraints:

- The number of nodes in the list is in the range [1, 105].
- $0 \leq \text{Node.val} \leq 9$

Follow up: Could you do it in $O(n)$ time and $O(1)$ space?

Community Solutions (Python)

- WalkCC

```

class Solution:
    def isPalindrome(self, head: ListNode) -> bool:
        def reverseList(head: ListNode) -> ListNode:
            prev = None
            curr = head
            while curr:
                next = curr.next
                curr.next = prev
                prev = curr
                curr = next

```

```

        return prev

```

```

slow = head
fast = head

```

```

while fast and fast.next:
    slow = slow.next
    fast = fast.next.next

```

```

if fast:

```

```

        slow = slow.next
        slow = reverseList(slow)

```

```

while slow:
    if slow.val != head.val:
        return False
    slow = slow.next
    head = head.next

```

```

return True

```

• LeetCode

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def isPalindrome(self, head:
Optional[ListNode]) -> bool:
        slow, fast = head, head.next
        while fast and fast.next:
            slow, fast = slow.next, fast.next.next

```

```

pre, cur = None, slow.next
while cur:
    t = cur.next
    cur.next = pre
    pre, cur = cur, t
while pre:
    if pre.val != head.val:
        return False
    pre, head = pre.next, head.next
return True

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def isPalindrome(self, head: ListNode) ->
bool:
    # Find the midpoint using slow and fast
    pointers
    slow, fast = head, head

```

```

while fast and fast.next:
    slow = slow.next # Move slow pointer by
1
    fast = fast.next.next # Move fast
pointer by 2

# Reverse the second half of the list
prev = None
while slow:
    next_node = slow.next # Store the next
node
    slow.next = prev # Reverse the current
node's pointer
    prev = slow # Move prev to current node

    slow = next_node # Move to next node

# Compare the first half and the reversed
second half node by node
left, right = head, prev
while right: # Only need to compare till
the end of second half
    if left.val != right.val:
        return False # Mismatch found, not
a palindrome
    left = left.next # Move left pointer
    right = right.next # Move right pointer

```

```
    return True # All nodes matched, it's a
palindrome
```

• LC.ca

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def isPalindrome(self, head:
Optional[ListNode]) -> bool:
        slow, fast = head, head.next
        while fast and fast.next:
            slow, fast = slow.next, fast.next.next

        pre, cur = None, slow.next
        while cur:
            t = cur.next
            cur.next = pre
            pre, cur = cur, t
        while pre:
            if pre.val != head.val:
                return False
            pre, head = pre.next, head.next
        return True
```

◆ Quick Review

Key Insights

- Use two pointers (slow and fast) to reach the middle of the list.
- Reverse the second half of the list.
- Compare the nodes from the start of the list to the reversed half.
- This approach meets the $O(n)$ time and $O(1)$ space follow-up requirement.

Space & Time

Time Complexity: $O(n)$ – each node is visited a constant number of times.

Space Complexity: $O(1)$ – only constant extra space is used (in-place reversal).

Solution

We first use the two-pointer technique to locate the midpoint of the linked list. The slow pointer moves one step at a time while the fast pointer moves two steps. Once the middle is found, we reverse the second half of the list in-place. Finally, we traverse both halves concurrently to compare the node values. If all corresponding values match, the linked list is a palindrome. A potential gotcha is handling the odd number of nodes; in that case, we ignore the middle element when comparing.

Linked List

#706 Design HashMap

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Linked List · Design · Hash Function

Lists: Denny Zhang

Problem

Description

Design a HashMap without using any built-in hash table libraries.

Implement the MyHashMap class:

- MyHashMap() initializes the object with an empty map.
- void put(int key, int value) inserts a (key, value) pair into the HashMap. If the key already exists in the map, update the corresponding value.
- int get(int key) returns the value to which the specified key is mapped, or -1 if this map contains no mapping for the key.
- void remove(key) removes the key and its corresponding value if the map contains the mapping for the key.

Example 1:

Input

```
["MyHashMap", "put", "put", "get", "get", "put", "get", "remove", "get"]
```

```
[[], [1, 1], [2, 2], [1], [3], [2, 1], [2], [2], [2]]
```

Output

```
[null, null, null, 1, -1, null, 1, null, -1]
```

Explanation

```
MyHashMap myHashMap = new MyHashMap();
```

Constraints:

- $0 \leq \text{key}, \text{value} \leq 106$
- At most 104 calls will be made to put, get, and remove.

Community Solutions (Python)

• LeetDoocs

```
class MyHashMap:
    def __init__(self):
        self.data = [-1] * 1000001

    def put(self, key: int, value: int) -> None:
        self.data[key] = value

    def get(self, key: int) -> int:
        return self.data[key]
```

```
def remove(self, key: int) -> None:
    self.data[key] = -1
```

Your MyHashMap object will be instantiated and called as such:

```
# obj = MyHashMap()
# obj.put(key,value)
# param_2 = obj.get(key)
# obj.remove(key)
```

• SimplyLeet

```
# ListNode class for handling collisions in each bucket
class ListNode:
    def __init__(self, key, value, next=None):
        self.key = key # Store key
        self.value = value # Store value
        self.next = next # Pointer to the next node

class MyHashMap:
    def __init__(self):
        self.size = 1000 # Define the bucket size
```

```
self.buckets = [None] * self.size #
Initialize buckets as an array of None
```

```
def hash(self, key):
    # Simple modulo based hash function to map
    keys to bucket indices
    return key % self.size
```

```
def put(self, key: int, value: int) -> None:
    index = self.hash(key)
    if self.buckets[index] is None:
        # If the bucket is empty, create a new
        node and place it here
```

```
self.buckets[index] = ListNode(key,
value)
    else:
        current = self.buckets[index]
        while True:
            if current.key == key:
                # If key exists, update its
                value

                current.value = value
                return
            if current.next is None:
                break
```



```

        current = current.next
        # Append a new node at the end of the
list if key not found
        current.next = ListNode(key, value)

```

```

def get(self, key: int) -> int:
    index = self.hash(key)
    current = self.buckets[index]
    while current:
        if current.key == key:
            # Key found; return its value

```

```

        return current.value
        current = current.next
    # Key not found; return -1
    return -1

```

```

def remove(self, key: int) -> None:
    index = self.hash(key)
    current = self.buckets[index]
    if current is None:
        return

```

```

        if current.key == key:
            # Remove head node if it contains the
key
            self.buckets[index] = current.next
        else:
            prev = current
            current = current.next
            while current:
                if current.key == key:
                    # Bypass the node to remove it
from the list
                    prev.next = current.next

```

```

            return
            prev, current = current,
current.next

```

• LC.ca

```

class MyHashMap:
    def __init__(self):
        self.data = [-1] * 1000001

    def put(self, key: int, value: int) -> None:
        self.data[key] = value

    def get(self, key: int) -> int:
        return self.data[key]

```

```
def remove(self, key: int) -> None:
    self.data[key] = -1
```

Your MyHashMap object will be instantiated and called as such:

```
# obj = MyHashMap()
# obj.put(key,value)
# param_2 = obj.get(key)
# obj.remove(key)
```

#####

```
class MyHashMap: # hash to bucket
```

```
    def __init__(self):
```

```
        """
```

```
        Initialize your data structure here.
```

```
        """
```

```
        self.size = 1000
```

```
        self.buckets = [[] for _ in
range(self.size)]
```

```
def _hash(self, key: int) -> int:
```

```
    """
```

```
    Generate a hash for a given key.
```

```
    """
```

```
    return key % self.size
```

```
def put(self, key: int, value: int) -> None:
```

```
    """
```

```
    Value will always be non-negative.
```

```
    """
```

```
        hash_key = self._hash(key)
```

```
        for i, (k, v) in
```

```
enumerate(self.buckets[hash_key]):
```

```
            if k == key:
```

```
                self.buckets[hash_key][i] = (key,
value)
```

```
                return
```

```
                self.buckets[hash_key].append((key,
value))
```

```
def get(self, key: int) -> int:
```

```
    """
```

```
    Returns the value to which the specified
key is mapped, or -1 if this map contains no
mapping for the key.
```

```

        """
        hash_key = self._hash(key)
        for k, v in self.buckets[hash_key]:
            if k == key:
                return v
        return -1

    def remove(self, key: int) -> None:
        """
        Removes the mapping of the specified value
        key if this map contains a mapping for the key.

```

```

        """
        hash_key = self._hash(key)
        for i, (k, v) in
enumerate(self.buckets[hash_key]):
            if k == key:
                del self.buckets[hash_key][i]
                return

#####

class MyHashMap: # open addressing with linear
probing

```

```

    def __init__(self):
        self.capacity = 10000
        self.data = [None] * self.capacity
        self.DELETED = (None, None) # Marker for
deleted entries

    def _hash(self, key):
        return key % self.capacity

    def put(self, key: int, value: int) -> None:
        idx = self._hash(key)

        while self.data[idx] not in (None,
self.DELETED, key):
            idx = (idx + 1) % self.capacity
            self.data[idx] = (key, value)

    def get(self, key: int) -> int:
        idx, found_key = self._find(key)
        if found_key:
            return self.data[idx][1]
        return -1

```

```

def remove(self, key: int) -> None:
    idx, found_key = self._find(key)
    if found_key:
        self.data[idx] = self.DELETED

def _find(self, key):
    original_idx = idx = self._hash(key)
    while self.data[idx] is not None:
        if self.data[idx] != self.DELETED and
self.data[idx][0] == key:
            return idx, True

    idx = (idx + 1) % self.capacity
    if idx == original_idx: # Came full
circle
        break
    return original_idx, False

```

◆ Quick Review

Key Insights

- Use a fixed-size array of buckets where each bucket uses a linked list to handle collisions.
- Choose a simple hash function (e.g., modulo operation) to map a key to its corresponding bucket.
- For each bucket, traverse the linked list to update an existing key or append a new node if the key does not exist.
- Removal requires updating pointers in the linked list to bypass the removed node.

Space & Time

Time Complexity: Average $O(1)$ for put, get, and remove; worst-case $O(n)$ when many keys hash to the same bucket.
Space Complexity: $O(n)$, where n is the number of key-value pairs stored.

Solution

We implement the HashMap using an array of fixed size (e.g., 1000). Each array element is the head of a linked list that manages collisions using chaining. The key is hashed using a modulo operation. For the put operation, if a key already exists in the chain, its value is updated; otherwise, a new node is appended. The get operation traverses the chain at the computed bucket index to find and return the corresponding value, or returns -1 if not found. The remove operation also traverses the list, carefully reconnecting the links to remove the target node while handling edge cases such as removal at the head.

Linked List

#876 Middle of the Linked List

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers

Lists: Grind 169

Problem

Description

Given the head of a singly linked list, return the middle node of the linked list.

If there are two middle nodes, return the second middle node.

Example 1:




Input: head = [1,2,3,4,5]

Output: [3,4,5]

Explanation: The middle node of the list is node 3.

Example 2:



Input: head = [1,2,3,4,5,6]

Output: [4,5,6]

Explanation: Since the list has two middle nodes with values 3 and 4, we return the second one.

Constraints:

- The number of nodes in the list is in the range [1, 100].
- $1 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def middleNode(self, head: ListNode) ->
    ListNode:
        slow = head
        fast = head

        while fast and fast.next:
            slow = slow.next
            fast = fast.next.next

        return slow
  
```

• LeetDoocs

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def middleNode(self, head: ListNode) ->
    ListNode:
        slow = fast = head
        while fast and fast.next:
            slow, fast = slow.next, fast.next.next

        return slow
```

• SimplyLeet

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next

class Solution:
    def middleNode(self, head: ListNode) ->
    ListNode:
        # Initialize slow and fast pointers at the
        head of the list
        slow = head
```

```
fast = head
# Loop until fast reaches the end of the
list
while fast and fast.next:
    # Move slow pointer one step
    slow = slow.next
    # Move fast pointer two steps
    fast = fast.next.next
# When loop ends, slow pointer points to
the middle node
return slow
```

• LC.ca

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def middleNode(self, head: ListNode) ->
    ListNode:
        slow = fast = head
        while fast and fast.next:
            slow, fast = slow.next, fast.next.next

        return slow
```

◆ Quick Review

Key Insights

- Use two pointers (slow and fast) to traverse the list.
- The fast pointer moves two steps at a time while the slow pointer moves one step.
- When the fast pointer reaches the end, the slow pointer will be at the middle node.
- This approach naturally handles both odd and even lengths of the list.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list.

Space Complexity: $O(1)$, constant extra space.

Solution

The solution uses the two-pointer technique:

- Initialize two pointers, slow and fast, at the head.
- Traverse the list, moving slow by one step and fast by two steps during each iteration.
- When fast reaches the end or has no next node, slow will be at the middle.
- Return the slow pointer (the middle node).

This strategy ensures linear time complexity and constant space, efficiently finding the middle node even for lists with an even number of elements.

Linked List

#1474 Delete N Nodes After M Nodes of Linked List

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List

Lists: Premium 98

Problem

Description

You are given the head of a linked list and two integers m and n .

Traverse the linked list and remove some nodes in the following way:

- Start with the head as the current node.
- Keep the first m nodes starting with the current node.
- Remove the next n nodes
- Keep repeating steps 2 and 3 until you reach the end of the list.

Return the head of the modified list after removing the mentioned nodes.

Example 1:



Input: head = [1,2,3,4,5,6,7,8,9,10,11,12,13], m = 2, n = 3

Output: [1,2,6,7,11,12]


Explanation: Keep the first (m = 2) nodes starting from the head of the linked List (1 ->2) show in black nodes.

Delete the next (n = 3) nodes (3 -> 4 -> 5) show in read nodes.

Continue with the same procedure until reaching the tail of the Linked List.

Head of the linked list after removing nodes is returned.

Example 2:



Input: head = [1,2,3,4,5,6,7,8,9,10,11], m = 1, n = 3

Output: [1,5,9]

Explanation: Head of linked list after removing nodes is returned.

Constraints:

- The number of nodes in the list is in the range [1, 104].
- $1 \leq \text{Node.val} \leq 106$
- $1 \leq m, n \leq 1000$

Follow up: Could you solve this problem by modifying the list in-place?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def deleteNodes(
        self,
        head: ListNode | None,
        m: int,
        n: int,
    ) -> ListNode | None:
        curr = head
        prev = None # prev.next == curr

        while curr:
            # Set the m-th node as `prev`.
            for _ in range(m):
                if not curr:
                    break
                prev = curr
                curr = curr.next
            # Set the (m + n + 1)-th node as `curr`.
            for _ in range(n):
                if not curr:
```



```

        break
    curr = curr.next
    # Delete the nodes [m + 1..n - 1].
    prev.next = curr

return head

```

• LeetCode

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def deleteNodes(self, head: ListNode, m: int,
n: int) -> ListNode:
    pre = head
    while pre:
        for _ in range(m - 1):

```

```

        if pre:
            pre = pre.next
        if pre is None:
            return head
        cur = pre
        for _ in range(n):
            if cur:
                cur = cur.next
            pre.next = None if cur is None else
cur.next
            pre = pre.next

return head

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def deleteNodes(self, head: ListNode, m: int,
n: int) -> ListNode:
    # Initialize the current pointer at the
head of the list.
    current = head

```

```

while current:
    # Keep m nodes: traverse m-1 next
    pointers
    for i in range(1, m):
        if current is None: # if list ends,
            break
        current = current.next
    # If we have reached the end of the
    list, break the loop.
    if current is None:
        break
    # Start deletion from the next node
    after the m-th node.

```

```

temp = current.next
# Delete next n nodes.
for j in range(n):
    if temp is None:
        break
    temp = temp.next
# Connect the m-th node to the node
after the n deleted nodes.
current.next = temp
# Move current pointer to continue from
the next kept node.
current = temp

```

```
return head
```

• LC.ca

Round 4 · Mid · Easy

p.1723

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def deleteNodes(self, head: ListNode, m: int,
n: int) -> ListNode:
    pre = head
    while pre:
        for _ in range(m - 1):
            if pre:
                pre = pre.next
            if pre is None:
                return head
            cur = pre
            for _ in range(n):
                if cur:
                    cur = cur.next
                pre.next = None if cur is None else
cur.next
            pre = pre.next
        return head

```

◆ Quick Review

Key Insights

- Traverse the linked list while keeping count.

Round 4 · Mid · Easy

p.1724

- Skip m nodes (i.e., keep them) and then proceed to delete the next n nodes.
- Update pointers to bypass the removed nodes.
- This is an in-place modification problem, so no additional list structure is needed.
- Be cautious about edge cases when the list has fewer than m or n nodes remaining.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the list.

Space Complexity: $O(1)$ since the modification is done in-place.

Solution

The solution involves iterating through the linked list with a pointer. For each cycle:

– Traverse m nodes and keep them. – Then, from the $(m+1)$ th node, skip and remove the next n nodes by adjusting the 'next' pointer of the m -th node to point to the node after these n nodes. – Continue until you reach the end of the list. This in-place approach avoids using extra space and maintains $O(N)$ time complexity.

Stack

Round 4 · Mid · Easy · 3 questions

Stack

#20 Valid Parentheses

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Stack

Lists: Grind 169

Problem

Description

Given a string *s* containing just the characters '(', ')', '{', '}', '[' and ']', determine if the input string is valid.

An input string is valid if:

1. Open brackets must be closed by the same type of brackets.
2. Open brackets must be closed in the correct order.
3. Every close bracket has a corresponding open bracket of the same type.

Example 1:

Input: *s* = "()"

Output: true

Example 2:

Input: *s* = "()[]{}"

Output: true

Example 3:

Input: *s* = "([)"

Output: false

Example 4:

Input: *s* = "([)]"

Output: true

Example 5:

Input: *s* = "([)]"

Output: false

Constraints:

- $1 \leq s.length \leq 104$
- *s* consists of parentheses only '()[]{}'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isValid(self, s: str) -> bool:
        stack = []

        for c in s:
            if c == '(':
                stack.append(')')
            elif c == '{':
                stack.append('}')
            elif c == '[':
```

```

        stack.append(']')
    elif not stack or stack.pop() != c:
        return False

    return not stack

```

- LeetCode

```

class Solution:
    def isValid(self, s: str) -> bool:
        stk = []
        d = {'(': ')', '[': ']', '{': '}'
        for c in s:
            if c in '({[':
                stk.append(c)
            elif not stk or stk.pop() + c not in d:
                return False
        return not stk

```

- SimplyLeet

```

def isValid(s: str) -> bool:
    # Mapping from closing brackets to their
    # corresponding opening brackets.
    paren_map = {')': '(', '}': '{', ']': '['}
    # Initialize an empty list to use as a stack.
    stack = []
    # Process each character in the string.
    for char in s:
        # If the character is a closing bracket.
        if char in paren_map:
            # Pop the top element from the stack if
            # available, else use a placeholder.

```

```

            top_element = stack.pop() if stack
            else '#'
            # Check if the mapping does not match
            # the top element.
            if paren_map[char] != top_element:
                return False
            else:
                # For opening brackets, push onto the
                # stack.
                stack.append(char)
            # Return True if all brackets are matched
            # (i.e., the stack is empty).
            return not stack

```

```

# Example usage
print(isValid("()[]{}")) # Expected output: True

```

- LC.ca

```

class Solution:
    def isValid(self, s: str) -> bool:
        stk = []
        d = {'(': ')', '[': ']', '{': '}'
        for c in s:
            if c in '({[':
                stk.append(c)
            elif not stk or stk.pop() + c not in d:
                return False
        return not stk

```

◆ Quick Review

Key Insights

- The problem can be solved using a stack data structure.
- When encountering an opening bracket, push it onto the stack.
- When encountering a closing bracket, check if the top of the stack is its matching opening bracket.
- If the stack is empty before finding a match or if there is a mismatch, the string is invalid.
- At the end, the string is valid only if the stack is empty, meaning all brackets were properly matched.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, since we process each character once.

Space Complexity: $O(n)$ in the worst case for the stack when all characters are opening brackets.

Solution

We use a stack to check for valid parentheses. The algorithm iterates through the string, pushing opening brackets onto the stack and popping them when a corresponding closing bracket is encountered. A mapping of closing to opening brackets helps in validating the matching pairs. By ensuring that the stack is empty at the end, we confirm that every opening bracket was properly closed.

Stack

#232 Implement Queue using Stacks

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Stack · Design · Queue

Lists: Grind 169

Problem

Description

Implement a first in first out (FIFO) queue using only two stacks. The implemented queue should support all the functions of a normal queue (push, peek, pop, and empty).

Implement the MyQueue class:

- `void push(int x)` Pushes element `x` to the back of the queue.
- `int pop()` Removes the element from the front of the queue and returns it.
- `int peek()` Returns the element at the front of the queue.
- `boolean empty()` Returns true if the queue is empty, false otherwise.

Notes:

- You must use only standard operations of a stack, which means only push to top, peek/pop from top, size,

and is empty operations are valid.

- Depending on your language, the stack may not be supported natively. You may simulate a stack using a list or deque (double-ended queue) as long as you use only a stack's standard operations.

Example 1:

```
Input
["MyQueue", "push", "push", "peek", "pop", "empty"]
[[], [1], [2], [], [], []]
Output
[null, null, null, 1, 1, false]

Explanation
MyQueue myQueue = new MyQueue();
```

Constraints:

- $1 \leq x \leq 9$
- At most 100 calls will be made to push, pop, peek, and empty.
- All the calls to pop and peek are valid.

Follow-up: Can you implement the queue such that each operation is amortized $O(1)$ time complexity? In other words, performing n operations will take overall $O(n)$ time even if one of those operations may take longer.

Community Solutions (Python)

- WalkCC

```

class MyQueue:
    def __init__(self):
        self.input = []
        self.output = []

    def push(self, x: int) -> None:
        self.input.append(x)

    def pop(self) -> int:
        self.peak()

        return self.output.pop()

    def peek(self) -> int:
        if not self.output:
            while self.input:
                self.output.append(self.input.pop())
        return self.output[-1]

    def empty(self) -> bool:
        return not self.input and not self.output

```

• LeetCode

```

class MyQueue:
    def __init__(self):
        self.stk1 = []
        self.stk2 = []

    def push(self, x: int) -> None:
        self.stk1.append(x)

    def pop(self) -> int:
        self.move()

        return self.stk2.pop()

    def peek(self) -> int:
        self.move()
        return self.stk2[-1]

    def empty(self) -> bool:
        return not self.stk1 and not self.stk2

    def move(self):

```



```

    if not self.stk2:
        while self.stk1:
            self.stk2.append(self.stk1.pop())

```

Your MyQueue object will be instantiated and called as such:

```

# obj = MyQueue()
# obj.push(x)
# param_2 = obj.pop()
# param_3 = obj.peek()

# param_4 = obj.empty()

```

• SimplyLeet

```

class MyQueue:
    def __init__(self):
        # Stack for enqueue operations.
        self.input_stack = []
        # Stack for dequeue operations.
        self.output_stack = []

    def push(self, x: int) -> None:
        # Always push new elements onto the input
        # stack.
        self.input_stack.append(x)

```

```

    def pop(self) -> int:
        # Ensure the output stack has the current
        # queue order.
        if not self.output_stack:
            while self.input_stack:
                # Move all elements from
                # input_stack to output_stack.
                self.output_stack.append(self.inpu
                t_stack.pop())
            # The top element of output_stack is the
            # front of the queue.
            return self.output_stack.pop()

```

```

    def peek(self) -> int:
        # Ensure the output stack has the current
        # queue order.
        if not self.output_stack:
            while self.input_stack:
                # Move all elements from
                # input_stack to output_stack.
                self.output_stack.append(self.inpu
                t_stack.pop())
            # Peek the top of output_stack without
            # popping.
            return self.output_stack[-1]

```

```

    def empty(self) -> bool:

```

```

        # The queue is empty if both stacks are
        empty.
        return not self.input_stack and not
        self.output_stack

```

• LC.ca

```

class MyQueue:
    def __init__(self):
        self.stk1 = []
        self.stk2 = [] # reversed order

    def push(self, x: int) -> None:
        self.stk1.append(x)

    def pop(self) -> int:
        self.move()

        return self.stk2.pop()

    def peek(self) -> int:
        self.move()
        return self.stk2[-1]

    def empty(self) -> bool:
        return not self.stk1 and not self.stk2

    def move(self):

```

```

        if not self.stk2: # only when stk2 is empty
            while self.stk1:
                self.stk2.append(self.stk1.pop())

# Your MyQueue object will be instantiated and
called as such:
# obj = MyQueue()
# obj.push(x)
# param_2 = obj.pop()
# param_3 = obj.peek()

# param_4 = obj.empty()

#####

class MyQueue:

    def __init__(self):
        self.sk = []
        self.rsk = [] # reversed

```

```
def push(self, x: int) -> None:
    self.sk.append(x);

def pop(self) -> int:
    self.peek()
    return self.rsk.pop()

def peek(self) -> int:
    if self.rsk:

        return self.rsk[-1]
    else:
        while self.sk:
            self.rsk.append(self.sk.pop())
        return self.rsk[-1]

def empty(self) -> bool:

    return not (self.sk or self.rsk)
```

```
# Your MyQueue object will be instantiated and
called as such:
# obj = MyQueue()
# obj.push(x)
# param_2 = obj.pop()
# param_3 = obj.peek()
# param_4 = obj.empty()
```

◆ Quick Review

Key Insights

- Use two stacks:
- One (input stack) for enqueue operations.
- The other (output stack) for dequeue operations.
- When performing pop or peek, if the output stack is empty, transfer all elements from the input stack to the output stack. This reversal makes the oldest element accessible.
- Each element is moved at most once between the stacks, ensuring an amortized $O(1)$ cost per operation.
- The space complexity is $O(n)$ where n is the number of elements in the queue.

Space & Time

Time Complexity: Amortized $O(1)$ per operation

Space Complexity: $O(n)$

Solution

We maintain two stacks: an input stack for pushing new elements and an output stack for popping or peeking the front element of the queue. When we need to access the front of the queue (either during a pop or peek), we check if the output stack is empty. If it is, we transfer all elements from the input stack to the output stack. This reversal of order allows us to simulate the FIFO nature of the queue. When performing a push, the element is simply added to the input stack.

Stack

#844 Backspace String Compare

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · String · Stack · Simulation

Lists: Grind 169

Problem

Description

Given two strings *s* and *t*, return true if they are equal when both are typed into empty text editors. '#' means a backspace character.

Note that after backspacing an empty text, the text will continue empty.

Example 1:

Input: *s* = "ab#c", *t* = "ad#c"
 Output: true
 Explanation: Both *s* and *t* become "ac".

Example 2:

Input: *s* = "ab##", *t* = "c#d#"
 Output: true
 Explanation: Both *s* and *t* become "".

Example 3:

Input: s = "a#c", t = "b"

Output: false

Explanation: s becomes "c" while t becomes "b".

Constraints:

- $1 \leq s.length, t.length \leq 200$
- s and t only contain lowercase letters and '#' characters.

Follow up: Can you solve it in $O(n)$ time and $O(1)$ space?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def backspaceCompare(self, s: str, t: str) -> bool:
        def backspace(s: str) -> str:
            stack = []
            for c in s:
                if c != '#':
                    stack.append(c)
                elif stack:
                    stack.pop()
            return stack

        return backspace(s) == backspace(t)
```

```
class Solution:
    def backspaceCompare(self, s: str, t: str) -> bool:
        i = len(s) - 1 # s' index
        j = len(t) - 1 # t's index

        while True:
            # Delete characters of s if needed.
            backspace = 0
            while i >= 0 and (s[i] == '#' or backspace > 0):
                backspace += 1 if s[i] == '#' else -1
                i -= 1

            # Delete characters of t if needed.
            backspace = 0
            while j >= 0 and (t[j] == '#' or backspace > 0):
                backspace += 1 if t[j] == '#' else -1
                j -= 1

            if i >= 0 and j >= 0 and s[i] == t[j]:
                i -= 1
                j -= 1
            else:
                break

        return i == -1 and j == -1
```

• LeetCode

```

class Solution:
    def backspaceCompare(self, s: str, t: str) ->
bool:
    i, j, skip1, skip2 = len(s) - 1, len(t) -
1, 0, 0
    while i >= 0 or j >= 0:
        while i >= 0:
            if s[i] == '#':
                skip1 += 1
                i -= 1
            elif skip1:
                skip1 -= 1
                i -= 1
            else:
                break
        while j >= 0:
            if t[j] == '#':
                skip2 += 1
                j -= 1
            elif skip2:
                skip2 -= 1
                j -= 1
            else:
                break
        if i >= 0 and j >= 0:
            if s[i] != t[j]:
                return False
        elif i >= 0 or j >= 0:
            return False
        i, j = i - 1, j - 1
    return True

```

```

else:
    break
    if i >= 0 and j >= 0:
        if s[i] != t[j]:
            return False
        elif i >= 0 or j >= 0:
            return False
        i, j = i - 1, j - 1
    return True

```

• SimplyLeet

```

# Function to compare two strings with backspace
processing
def backspaceCompare(s: str, t: str) -> bool:
    # Pointers for both strings starting at the end
    i, j = len(s) - 1, len(t) - 1

    # Process both strings until all characters are
    compared
    while i >= 0 or j >= 0:
        # Process string s: skip backspaced
        characters
        skip = 0
        while i >= 0:

```

```

    if s[i] == '#':
        skip += 1
        i -= 1
    elif skip > 0:
        skip -= 1
        i -= 1
    else:
        break

```

Process string t: skip backspaced characters

```

skip = 0
while j >= 0:
    if t[j] == '#':
        skip += 1
        j -= 1
    elif skip > 0:
        skip -= 1
        j -= 1
    else:
        break

```

Compare the current valid characters from s and t

```

    if i >= 0 and j >= 0:
        if s[i] != t[j]:
            return False
    elif i >= 0 or j >= 0:
        return False

```

```

    i -= 1
    j -= 1

```

return True

Example test cases

```

print(backspaceCompare("ab#c", "ad#c")) # Expected
output: True
print(backspaceCompare("ab##", "c#d#")) # Expected
output: True
print(backspaceCompare("a#c", "b")) # Expected
output: False

```

• LC.ca

```

class Solution:
    def backspaceCompare(self, s: str, t: str) ->
bool:
    i, j, skip1, skip2 = len(s) - 1, len(t) -
1, 0, 0
    while i >= 0 or j >= 0:
        while i >= 0:
            if s[i] == '#':
                skip1 += 1
                i -= 1
            elif skip1:
                skip1 -= 1
                i -= 1
            else:
                break
        while j >= 0:
            if t[j] == '#':
                skip2 += 1
                j -= 1
            elif skip2:
                skip2 -= 1
                j -= 1
            else:
                break
        if i >= 0 and j >= 0:
            if s[i] != t[j]:
                return False
        elif i >= 0 or j >= 0:
            return False
        i, j = i - 1, j - 1
    return True

```

```

else:
    break
    if i >= 0 and j >= 0:
        if s[i] != t[j]:
            return False
        elif i >= 0 or j >= 0:
            return False
        i, j = i - 1, j - 1
    return True

```

◆ Quick Review

Key Insights

- '#' represents a backspace operation that deletes the previous character.
- A direct simulation using a stack can reconstruct the final string.
- A two-pointer approach allows comparison from the end of the strings while skipping backspaced characters.
- The two-pointer technique can achieve $O(n)$ time and $O(1)$ space by processing the strings in reverse.

Space & Time

Time Complexity: $O(n)$ where n is the length of the longer string.

Space Complexity: $O(1)$ using the two-pointer strategy ($O(n)$ if using stacks).

Solution

The solution uses two pointers starting from the end of each string. For each string, maintain a skip counter to account for backspace characters. As you iterate backwards:

- If a '#' is encountered, increment the skip counter.
- If a normal character is encountered while the skip counter is positive, skip that character by decrementing the counter.
- Once a valid character is found, compare the characters from both strings. If all corresponding characters match and both pointers finish processing at the same time, the strings are considered equal.

Trie

Round 4 · Mid · Easy · 1 question

Round 4 · Mid · Easy

p.1753

Round 4 · Mid · Easy

p.1754

Trie

#14 Longest Common Prefix

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Trie

Lists: Grind 169

Problem

Description

Write a function to find the longest common prefix string amongst an array of strings.

If there is no common prefix, return an empty string "".

Example 1:

Input: strs = ["flower","flow","flight"]
Output: "fl"

Example 2:

Input: strs = ["dog","racecar","car"]
Output: ""
Explanation: There is no common prefix among the input strings.

Constraints:

- $1 \leq \text{strs.length} \leq 200$
- $0 \leq \text{strs}[i].\text{length} \leq 200$
- $\text{strs}[i]$ consists of only lowercase English letters if it is non-empty.

Round 4 · Mid · Easy

p.1755

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestCommonPrefix(self, strs: list[str])
-> str:
    if not strs:
        return ''

    for i in range(len(strs[0])):
        for j in range(1, len(strs)):
            if i == len(strs[j]) or strs[j][i] !=
strs[0][i]:
                return strs[0][:i]

    return strs[0]
```

• LeetDoocs

```
class Solution:
    def longestCommonPrefix(self, strs: List[str])
-> str:
        for i in range(len(strs[0])):
            for s in strs[1:]:
                if len(s) <= i or s[i] !=
strs[0][i]:
                    return s[:i]
        return strs[0]
```

• SimplyLeet

Round 4 · Mid · Easy

p.1756

```
def longestCommonPrefix(strs):
    # If no strings exist, return empty
    if not strs:
        return ""
    # Iterate over each index of the first string
    for i in range(len(strs[0])):
        current_char = strs[0][i]
        # Compare the current character with the
        corresponding one in the rest of the strings
        for s in strs[1:]:
            # If index exceeds length or characters
            differ, return the prefix found so far

            if i >= len(s) or s[i] != current_char:
                return strs[0][:i]
        # If no mismatch is found, return the complete
        first string
        return strs[0]

# Example usage:
print(longestCommonPrefix(["flower", "flow",
"flight"]))
```

• LC.ca

Round 4 · Mid · Easy

p.1757

```
class Solution:
    def longestCommonPrefix(self, strs: List[str])
-> str:
        for i in range(len(strs[0])):
            for s in strs[1:]:
                if len(s) <= i or s[i] !=
strs[0][i]:
                    return s[:i]
        return strs[0]
```

◆ Quick Review

Key Insights

- Compare characters of the strings one column at a time (vertical scanning).
- The longest possible prefix is limited by the length of the shortest string in the array.
- As soon as a mismatch is found, the prefix up to that point is the answer.
- Handle edge cases such as an empty input array or strings with no common prefix.

Space & Time

Time Complexity: $O(N \cdot M)$ where N is the number of strings and M is the length of the shortest string.

Space Complexity: $O(1)$ additional space (ignoring output space).

Solution

We use the vertical scanning approach:

Round 4 · Mid · Easy

p.1758

– If the input array is empty, return an empty string. – Iterate character by character over the first string. – For each character, compare it with the corresponding character in each of the other strings. – If a mismatch or end of any string is encountered, return the substring from the start up to the current index. – If no mismatch is found, the entire first string is the common prefix.

This solution ensures that we only traverse the necessary part of the strings and stops early if a discrepancy is discovered.

Greedy

Round 4 · Mid · Easy · 1 question

Round 4 · Mid · Easy

p.1759

Round 4 · Mid · Easy

p.1760

Greedy

#409 Longest Palindrome

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Greedy

Lists: Grind 169

Problem

Description

Given a string *s* which consists of lowercase or uppercase letters, return the length of the longest palindrome that can be built with those letters.

Letters are case sensitive, for example, "Aa" is not considered a palindrome.

Example 1:

Input: *s* = "abcccd" (Note: original image has typo 'dd')

Output: 7

Explanation: One longest palindrome that can be built is "dccaccd", whose length is 7.

Example 2:

Input: *s* = "a"

Output: 1

Explanation: The longest palindrome that can be built is "a", whose length is 1.

Constraints:

- $1 \leq s.length \leq 2000$
- *s* consists of lowercase and/or uppercase English letters only.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestPalindrome(self, s: str) -> int:
        ans = 0
        count = collections.Counter(s)

        for c in count.values():
            ans += c if c % 2 == 0 else c - 1

        hasOddCount = any(c % 2 == 1 for c in count.values())
        return ans + hasOddCount
```

• LeetDoocs

```
class Solution:
    def longestPalindrome(self, s: str) -> int:
        cnt = Counter(s)
        ans = sum(v // 2 * 2 for v in cnt.values())
        ans += int(ans < len(s))
        return ans
```

```
class Solution:
    def longestPalindrome(self, s: str) -> int:
        odd = defaultdict(int)
        cnt = 0
        for c in s:
            odd[c] ^= 1
            cnt += 1 if odd[c] else -1
        return len(s) - cnt + 1 if cnt else len(s)
```

• SimplyLeet

```
# Python solution for the Longest Palindrome
problem

def longestPalindrome(s: str) -> int:
    # Create a dictionary to store the frequency of
    each character
    char_counts = {}
    for char in s:
        # Increment the count for the character
        char_counts[char] = char_counts.get(char,
0) + 1

    # Variable to store the length of the
    palindrome
```

```
palindrome_length = 0
# Flag to indicate whether we can add one more
character in the middle for an odd count
center_added = False

# Iterate through the character counts
for count in char_counts.values():
    # Add the largest even number from count to
    palindrome_length
    palindrome_length += (count // 2) * 2
    # If count is odd and we haven't added the
    center character yet
    if count % 2 == 1 and not center_added:

        palindrome_length += 1
        center_added = True

    return palindrome_length

# Example usage:
print(longestPalindrome("abcccd"))
```

• LC.ca

```
class Solution:
    def longestPalindrome(self, s: str) -> int:
        cnt = Counter(s)
        ans = 0
        for v in cnt.values():
            ans += v - (v & 1)
            ans += (ans & 1 ^ 1) and (v & 1)
        return ans
```

◆ Quick Review

Key Insights

- Count the frequency of each character.
- For each character, you can use pairs (even count or odd count minus one) to form the palindrome.
- If any character has an odd frequency, you can add one extra character as the center of the palindrome.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(1)$ since the frequency counter will have at most 52 keys (considering both uppercase and lowercase letters).

Solution

The solution uses a hash table to count the frequency of each character. For every character, the maximum number of characters that can be used in a palindrome is determined by taking the largest even number that is less than or equal to its frequency (which is frequency -

(frequency % 2)). If there exists any character with an odd frequency, one of these can be used as the center of the palindrome, adding an extra character to the total length. This greedy approach efficiently computes the length of the largest possible palindrome from the available characters.

Round 5 · Mid · Medium

Round 5

Mid Priority · Medium

56 questions · 6 patterns

Linked List #2 #19 #24 #61 #146 #148 #328 #369 #430 #622
#707 #708 #1171

Stack #143 #150 #155 #173 #227 #255 #394 #439 #484
#735 #739 #962 #1214 #1762

Heap #215 #253 #347 #505 #621 #658 #787 #973 #1057
#1167 #1424

Trie #139 #208 #211 #616 #692 #720 #1166

Backtracking #17 #22 #39 #46 #79 #113

Greedy #134 #179 #280 #954 #2131

Linked List

Round 5 · Mid · Medium · 13 questions

Round 4 · Mid · Easy

p.1767

Round 5 · Mid · Medium

p.1768

Linked List

#2 Add Two Numbers

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Math · Recursion

Lists: Grind 169




Problem

Description

You are given two non-empty linked lists representing two non-negative integers. The digits are stored in reverse order, and each of their nodes contains a single digit. Add the two numbers and return the sum as a linked list.

You may assume the two numbers do not contain any leading zero, except the number 0 itself.

Example 1:



Input: l1 = [2,4,3], l2 = [5,6,4]

Output: [7,0,8]

Explanation: 342 + 465 = 807.

Example 2:

Input: l1 = [0], l2 = [0]

Output: [0]

Example 3:

Input: l1 = [9,9,9,9,9,9,9], l2 = [9,9,9,9]

Output: [8,9,9,9,0,0,0,1]

Constraints:

- The number of nodes in each linked list is in the range [1, 100].
- $0 \leq \text{Node.val} \leq 9$
- It is guaranteed that the list represents a number that does not have leading zeros.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def addTwoNumbers(self, l1: ListNode, l2:
ListNode) -> ListNode:
        dummy = ListNode(0)
        curr = dummy
        carry = 0

        while carry or l1 or l2:
            if l1:
                carry += l1.val
                l1 = l1.next

            if l2:
                carry += l2.val
                l2 = l2.next
            curr.next = ListNode(carry % 10)
            carry //= 10
            curr = curr.next

        return dummy.next
```

• LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def addTwoNumbers(
        self, l1: Optional[ListNode], l2:
Optional[ListNode]
    ) -> Optional[ListNode]:
        dummy = ListNode()

        carry, curr = 0, dummy
        while l1 or l2 or carry:
            s = (l1.val if l1 else 0) + (l2.val if
l2 else 0) + carry
            carry, val = divmod(s, 10)
            curr.next = ListNode(val)
            curr = curr.next
            l1 = l1.next if l1 else None
            l2 = l2.next if l2 else None
        return dummy.next
```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

def addTwoNumbers(l1, l2):
    # Initialize a dummy head node to simplify the
    # result list construction.
    dummy = ListNode()
    current = dummy

    carry = 0

    # Traverse through both linked lists until all
    # nodes and any carry are processed.
    while l1 or l2 or carry:
        # Get the current values; if a node is
        # missing, use 0.
        val1 = l1.val if l1 else 0
        val2 = l2.val if l2 else 0

        # Calculate the sum of current digits and
        # carry.
        total = val1 + val2 + carry

```

```

        carry = total // 10 # Update carry for next
        iteration.
        digit = total % 10 # Compute current digit.

        # Append the digit as a new node in the
        # result list.
        current.next = ListNode(digit)
        current = current.next

        # Move to the next nodes if available.
        if l1:
            l1 = l1.next

        if l2:
            l2 = l2.next

        # Return the head of the newly formed linked
        # list.
        return dummy.next

```

- LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def addTwoNumbers(
        self, l1: Optional[ListNode], l2:
Optional[ListNode]
    ) -> Optional[ListNode]:
        dummy = ListNode()

        carry, curr = 0, dummy
        while l1 or l2 or carry:
            s = (l1.val if l1 else 0) + (l2.val if
l2 else 0) + carry
            carry, val = divmod(s, 10)
            curr.next = ListNode(val)
            curr = curr.next
            l1 = l1.next if l1 else None
            l2 = l2.next if l2 else None
        return dummy.next

```

◆ Quick Review

Key Insights

- The digits are stored in reverse order, meaning the first node is the least significant digit.
- A carry may need to propagate through subsequent nodes.

- The two linked lists might be of different lengths.
- Continue processing until both lists are fully traversed and any remaining carry is handled.

Space & Time

Time Complexity: $O(\max(n, m))$, where n and m are the lengths of the two linked lists.

Space Complexity: $O(\max(n, m) + 1)$ for the result list in the worst case.

Solution

The solution iterates through both linked lists node by node. For each pair of corresponding nodes, their values and any carry from the previous operation are summed. The result is decomposed into a current digit (modulo 10) and an updated carry (total divided by 10). A new node is created for the current digit and appended to the resulting linked list. This process continues until all nodes are processed and any leftover carry is added as a final node.

Linked List

#19 Remove Nth Node From End of List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers


Lists: Grind 169

Problem

Description

Given the head of a linked list, remove the n th node from the end of the list and return its head.

Example 1:



Input: head = [1,2,3,4,5], n = 2

Output: [1,2,3,5]

Example 2:

Round 5 · Mid · Medium

p.1777

Input: head = [1], n = 1

Output: []

Example 3:

Input: head = [1,2], n = 1

Output: [1]

Constraints:

- The number of nodes in the list is sz .
- $1 \leq sz \leq 30$
- $0 \leq \text{Node.val} \leq 100$
- $1 \leq n \leq sz$

Follow up: Could you do this in one pass?

Community Solutions (Python)

• WalkCC

```

class Solution:
    def removeNthFromEnd(self, head: ListNode, n:
int) -> ListNode:
        slow = head
        fast = head

        for _ in range(n):
            fast = fast.next
        if not fast:
            return head.next
  
```

Round 5 · Mid · Medium

p.1778

```

while fast.next:
    slow = slow.next
    fast = fast.next
    slow.next = slow.next.next

return head

```

• LeetCode

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def removeNthFromEnd(self, head:
Optional[ListNode], n: int) -> Optional[ListNode]:
        dummy = ListNode(next=head)
        fast = slow = dummy
        for _ in range(n):

            fast = fast.next
        while fast.next:
            slow, fast = slow.next, fast.next
        slow.next = slow.next.next
        return dummy.next

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val # Holds the value of the
node
        self.next = next # Points to the next node
in the list

class Solution:
    def removeNthFromEnd(self, head: ListNode, n:
int) -> ListNode:
        # Create a dummy node to simplify edge
cases (like removing the head)
        dummy = ListNode(0)

        dummy.next = head

        # Initialize two pointers starting at the
dummy node
        first = dummy
        second = dummy

        # Advance first pointer n+1 steps ahead to
maintain a gap of n nodes
        for _ in range(n + 1):
            first = first.next

```

```

        # Move both pointers until first reaches
the end of the list
        while first:
            first = first.next
            second = second.next

        # Skip the node to remove it from the list
        second.next = second.next.next

        # Return the adjusted list starting from
dummy.next
        return dummy.next

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def removeNthFromEnd(self, head:
Optional[ListNode], n: int) -> Optional[ListNode]:
        dummy = ListNode(next=head)
        fast = slow = dummy
        for _ in range(n):

```

```

        fast = fast.next
    while fast.next:
        slow, fast = slow.next, fast.next
    slow.next = slow.next.next
    return dummy.next

```

◆ Quick Review

Key Insights

- Use a dummy node at the start to simplify edge cases, especially when the head is removed.
- Utilize two pointers (fast and slow) and maintain a gap of $n+1$ nodes between them.
- By moving both pointers simultaneously, the slow pointer will point to the node immediately before the target node when the fast pointer reaches the end.

Space & Time

Time Complexity: $O(L)$, where L is the length of the list

Space Complexity: $O(1)$

Solution

We use a two-pointer approach to achieve a one-pass solution. A dummy node is created and linked to the head to handle edge cases gracefully (such as needing to remove the head). The fast pointer is advanced $n+1$ steps ahead of the slow pointer to maintain the required gap. Then, both pointers are moved simultaneously until the fast pointer reaches the end of the list. At that point, the slow pointer is immediately before the node to be deleted. We then adjust

the pointers to remove the target node from the list.

Round 5 · Mid · Medium

p.1783

Linked List

#24 Swap Nodes in Pairs

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Recursion

Lists: Grind 169

Problem

Description

Given a linked list, swap every two adjacent nodes and return its head. You must solve the problem without modifying the values in the list's nodes (i.e., only nodes themselves may be changed.)

Example 1:


Input: head = [1,2,3,4]

Output: [2,1,4,3]

Explanation:

Round 5 · Mid · Medium

p.1784



Example 2:

Input: head = []

Output: []

Example 3:

Input: head = [1]

Output: [1]

Example 4:

Input: head = [1,2,3]

Output: [2,1,3]

Constraints:

- The number of nodes in the list is in the range [0, 100].
- $0 \leq \text{Node.val} \leq 100$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def swapPairs(self, head: ListNode) -> ListNode:
        def getLength(head: ListNode) -> int:
            length = 0
            while head:
                length += 1
                head = head.next
            return length

        length = getLength(head)

        dummy = ListNode(0, head)
        prev = dummy
        curr = head

        for _ in range(length // 2):
            next = curr.next
            curr.next = next.next
            next.next = prev.next
            prev.next = next
            prev = curr

        curr = curr.next

        return dummy.next
  
```

• LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def swapPairs(self, head: Optional[ListNode])
-> Optional[ListNode]:
    if head is None or head.next is None:
        return head
    t = self.swapPairs(head.next.next)

    p = head.next
    p.next = head
    head.next = t
    return p
```

• SimplyLeet

```
# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def swapPairs(self, head: ListNode) ->
ListNode:
    # Create a dummy node and point it to the
    head of the list.
    dummy = ListNode(0)
```

```
dummy.next = head
# Initialize the pointer 'prev' to dummy.
prev = dummy

# Traverse the list while at least two
nodes remain.
while head and head.next:
    # Identify the two nodes to swap.
    first_node = head
    second_node = head.next

    # Swapping the nodes by reassigning
    pointers.
    prev.next = second_node # Link
previous node to second node.
    first_node.next = second_node.next #
Link first node to node after second.
    second_node.next = first_node # Link
second node to first to complete swap.

    # Update the pointers for next pair
processing.
    prev = first_node # Move 'prev' pointer
to the end of the swapped pair.
    head = first_node.next # Move 'head'
pointer to the next pair.

    # Return the new head of the list.

return dummy.next
```

- LC.ca

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def swapPairs(self, head: Optional[ListNode])
-> Optional[ListNode]:
        dummy = ListNode(next=head)
        pre, cur = dummy, head
        while cur and cur.next:

            t = cur.next
            cur.next = t.next
            t.next = cur
            pre.next = t
            pre, cur = cur, cur.next
        return dummy.next
```

- ◆ Quick Review

Key Insights

- Use a dummy node to simplify edge cases, such as when the list is empty or has only one node.
- Process nodes in pairs and update the pointers accordingly.
- Ensure that the swapping does not lose the connection between the swapped pair and the rest of the list.

- The solution can be implemented iteratively or recursively.

Space & Time

Time Complexity: $O(n)$ – Each node is processed once.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We use an iterative approach with a dummy node. The dummy node points to the head of the list. We maintain a pointer that traverses the list in pairs. For each pair, we adjust the next pointers to swap the two nodes. By the end, the dummy's next pointer will be the new head of the modified list. The main idea is to keep track of the previous node (prev) before the pair, and then adjust pointers accordingly using temporary variables to hold references to the nodes being swapped.

Linked List

#61 Rotate List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers




Lists: Grind 169

Problem

Description

Given the head of a linked list, rotate the list to the right by k places.

Example 1:







Input: head = [1,2,3,4,5], $k = 2$

Output: [4,5,1,2,3]

Example 2:

Round 5 · Mid · Medium

p.1791



Input: head = [0,1,2], $k = 4$

Output: [2,0,1]

Constraints:

- The number of nodes in the list is in the range [0, 500].
- $-100 \leq \text{Node.val} \leq 100$
- $0 \leq k \leq 2 * 10^9$

Community Solutions (Python)

- [WalkCC](#)

Round 5 · Mid · Medium

p.1792

```
class Solution:
    def rotateRight(self, head: ListNode, k: int) ->
    ListNode:
        if not head or not head.next or k == 0:
            return head

        tail = head
        length = 1
        while tail.next:
            tail = tail.next
            length += 1

        tail.next = head # Circle the list.

        t = length - k % length
        for _ in range(t):
            tail = tail.next
        newHead = tail.next
        tail.next = None

        return newHead
```

• LeetCode

Round 5 · Mid · Medium

p.1793

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def rotateRight(self, head:
    Optional[ListNode], k: int) -> Optional[ListNode]:
        if head is None or head.next is None:
            return head
        cur, n = head, 0

        while cur:
            n += 1
            cur = cur.next
        k %= n
        if k == 0:
            return head
        fast = slow = head
        for _ in range(k):
            fast = fast.next
        while fast.next:
            fast, slow = fast.next, slow.next

        ans = slow.next
        slow.next = None
        fast.next = head
        return ans
```

• SimplyLeet

Round 5 · Mid · Medium

p.1794

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, x):
#         self.val = x
#         self.next = None

class Solution:
    def rotateRight(self, head: ListNode, k: int)
    -> ListNode:
        if not head or not head.next or k == 0: #
        Check if rotation is necessary
            return head
```

```
        # Compute the length of the list and locate
        the tail node
        length = 1
        tail = head
        while tail.next:
            tail = tail.next
            length += 1

        # Calculate the effective number of
        rotations
        k = k % length
```

```
        if k == 0:
            return head

        # Find the new tail: (length - k - 1) steps
        from the head
        stepsToNewTail = length - k - 1
        newTail = head
        for _ in range(stepsToNewTail):
            newTail = newTail.next

        # Set the new head and break the list

        newHead = newTail.next
        newTail.next = None

        # Connect the old tail to the original head
        tail.next = head

        return newHead
```

- LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def rotateRight(self, head:
Optional[ListNode], k: int) -> Optional[ListNode]:
        if head is None or head.next is None:
            return head
        cur, n = head, 0

        while cur:
            n += 1
            cur = cur.next
        k %= n
        if k == 0:
            return head
        fast = slow = head
        for _ in range(k):
            fast = fast.next
        while fast.next:
            fast, slow = fast.next, slow.next

        ans = slow.next
        slow.next = None
        fast.next = head
        return ans

```

◆ Quick Review

Key Insights

- Calculate the length of the list and locate the tail.
- Use k modulo length to handle rotations greater than the list size.
- Find the new tail of the list at position $(\text{length} - k - 1)$ and update pointers accordingly.
- Link the original tail to the original head to form a circular list, then break the circle at the new tail.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list, since we traverse the list a few times.

Space Complexity: $O(1)$, as the rotation is done in place without extra data structures.

Solution

We first traverse the list to determine its length and identify the tail node. The effective rotation is calculated by taking k modulo the length of the list. If the result is zero, the list remains unchanged. Otherwise, we locate the new tail node by moving $(\text{length} - k - 1)$ steps from the head. The node after the new tail becomes the new head. Finally, we break the link after the new tail and connect the old tail to the original head to finalize the rotated list.

Linked List

#146 LRU Cache

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Linked List · Design

Lists: Grind 169

Problem

Description

Design a data structure that follows the constraints of a Least Recently Used (LRU) cache.

Implement the LRUCache class:

- `LRUCache(int capacity)` Initialize the LRU cache with positive size capacity.
- `int get(int key)` Return the value of the key if the key exists, otherwise return -1.
- `void put(int key, int value)` Update the value of the key if the key exists. Otherwise, add the key-value pair to the cache. If the number of keys exceeds the capacity from this operation, evict the least recently used key.

The functions `get` and `put` must each run in $O(1)$ average time complexity.

Example 1:

Input

```
["LRUCache", "put", "put", "get", "put", "get", "put",
 "get", "get", "get"]
[[2], [1, 1], [2, 2], [1], [3, 3], [2], [4, 4], [1],
 [3], [4]]
```

Output

```
[null, null, null, 1, null, -1, null, -1, 3, 4]
```

Explanation

```
LRUCache lruCache = new LRUCache(2);
```

Constraints:

- $1 \leq \text{capacity} \leq 3000$
- $0 \leq \text{key} \leq 104$
- $0 \leq \text{value} \leq 105$
- At most $2 * 10^5$ calls will be made to `get` and `put`.

Community Solutions (Python)

• WalkCC

```
class Node:
    def __init__(self, key: int, value: int):
        self.key = key
        self.value = value
        self.prev = None
        self.next = None

class LRUCache:
    def __init__(self, capacity: int):
```



```

self.capacity = capacity
self.keyToNode = {}
self.head = Node(-1, -1)
self.tail = Node(-1, -1)
self.join(self.head, self.tail)

```

```

def get(self, key: int) -> int:
    if key not in self.keyToNode:
        return -1

```

```

node = self.keyToNode[key]
self.remove(node)
self.moveToHead(node)
return node.value

```

```

def put(self, key: int, value: int) -> None:
    if key in self.keyToNode:
        node = self.keyToNode[key]
        node.value = value
        self.remove(node)

```

```

self.moveToHead(node)
return

```

```

if len(self.keyToNode) == self.capacity:
    lastNode = self.tail.prev
    del self.keyToNode[lastNode.key]
    self.remove(lastNode)

```

```

self.moveToHead(Node(key, value))
self.keyToNode[key] = self.head.next

```

```

def join(self, node1: Node, node2: Node):
    node1.next = node2
    node2.prev = node1

```

```

def moveToHead(self, node: Node):
    self.join(node, self.head.next)
    self.join(self.head, node)

```

```

def remove(self, node: Node):

```

```

    self.join(node.prev, node.next)

```

• LeetCode

```
class Node:
    def __init__(self, key: int = 0, val: int = 0):
        self.key = key
        self.val = val
        self.prev = None
        self.next = None
```

```
class LRUCache:
```

```
    def __init__(self, capacity: int):
        self.size = 0
        self.capacity = capacity
        self.cache = {}
        self.head = Node()
        self.tail = Node()
        self.head.next = self.tail
        self.tail.prev = self.head
```

```
    def get(self, key: int) -> int:
```

```
        if key not in self.cache:
            return -1
        node = self.cache[key]
        self.remove_node(node)
        self.add_to_head(node)
        return node.val
```

```
    def put(self, key: int, value: int) -> None:
        if key in self.cache:
            node = self.cache[key]
```

```
            self.remove_node(node)
            node.val = value
            self.add_to_head(node)
        else:
            node = Node(key, value)
            self.cache[key] = node
            self.add_to_head(node)
            self.size += 1
            if self.size > self.capacity:
                node = self.tail.prev
```

```

        self.cache.pop(node.key)
        self.remove_node(node)
        self.size -= 1

```

```

def remove_node(self, node):
    node.prev.next = node.next
    node.next.prev = node.prev

```

```

def add_to_head(self, node):
    node.next = self.head.next

```

```

    node.prev = self.head
    self.head.next = node
    node.next.prev = node

```

Your LRUCache object will be instantiated and called as such:

```

# obj = LRUCache(capacity)
# param_1 = obj.get(key)
# obj.put(key,value)

```

• SimplyLeet

Define a node for the doubly linked list.

```

class Node:

```

```

    def __init__(self, key, value):
        self.key = key # key of the node
        self.value = value # value of the node
        self.prev = None # pointer to previous node
        self.next = None # pointer to next node

```

```

class LRUCache:

```

```

    def __init__(self, capacity):

```

```

        self.capacity = capacity # maximum
        capacity of the cache
        self.cache = {} # dictionary mapping key ->
        node

```

Create dummy head and tail nodes.

```

self.head = Node(0, 0)
self.tail = Node(0, 0)
self.head.next = self.tail
self.tail.prev = self.head

```

Helper function to remove a node from the linked list.

```

def _remove(self, node):

```

```

        prev_node = node.prev
        next_node = node.next
        prev_node.next = next_node
        next_node.prev = prev_node

    # Helper function to add a node right after the
    head.
    def _add(self, node):
        node.prev = self.head
        node.next = self.head.next
        self.head.next.prev = node

        self.head.next = node

    # Returns the value of the key if it exists;
    otherwise, returns -1.
    def get(self, key):
        if key in self.cache:
            node = self.cache[key]
            # Move the accessed node to the head
            (marking it as recently used)
            self._remove(node)
            self._add(node)
            return node.value

```

```

        return -1

    # Updates the value of the key if it exists;
    otherwise, adds the key-value pair.
    # Removes the least recently used key if the
    cache exceeds its capacity.
    def put(self, key, value):
        if key in self.cache:
            # Remove the old node.
            self._remove(self.cache[key])
            node = Node(key, value)
            self._add(node)

        self.cache[key] = node
        if len(self.cache) > self.capacity:
            # Remove the least recently used
            element.

            lru = self.tail.prev
            self._remove(lru)
            del self.cache[lru.key]

```

- LC.ca

```

class Node:
    def __init__(self, key=0, val=0):
        self.key = key
        self.val = val
        self.prev = None
        self.next = None

class LRUCache:
    def __init__(self, capacity: int):

        self.cache = {} # key ==> Node(val)
        self.head = Node() # dummy node
        self.tail = Node() # dummy node
        self.capacity = capacity
        self.size = 0
        self.head.next = self.tail # note: key
        self.tail.prev = self.head

        def get(self, key: int) -> int:
            if key not in self.cache:

```

```

        return -1
        node = self.cache[key]
        self.move_to_head(node)
        return node.val

    def put(self, key: int, value: int) -> None:
        if key in self.cache:
            node = self.cache[key]
            node.val = value
            self.move_to_head(node)
        else:
            node = Node(key, value)
            self.cache[key] = node
            self.add_to_head(node)
            self.size += 1
            if self.size > self.capacity:
                tail = self.remove_tail()
                self.cache.pop(tail.key)
                self.size -= 1

```

```
def move_to_head(self, node):
    self.remove_node(node)
    self.add_to_head(node)

def remove_node(self, node):
    node.prev.next = node.next
    node.next.prev = node.prev

def add_to_head(self, node):
    node.next = self.head.next
```

```
    node.prev = self.head
    self.head.next = node
    node.next.prev = node

def remove_tail(self):
    node = self.tail.prev
    self.remove_node(node)
    return node
```

```
# Your LRUCache object will be instantiated and
called as such:
# obj = LRUCache(capacity)
# param_1 = obj.get(key)
# obj.put(key,value)
```

```
#####
```

```
'''
```

```
example:
```

```
>>> od = collections.OrderedDict()
>>>
>>> od[1]=1
>>> od[2]=2
>>> od[3]=3
>>>
>>> od
OrderedDict([(1, 1), (2, 2), (3, 3)])
>>> od.move_to_end(1)
>>> od
```

```
OrderedDict([(2, 2), (3, 3), (1, 1)])
>>>
>>> od.get(1)
1
>>> od.popitem()
(1, 1)
>>> od
OrderedDict([(2, 2), (3, 3)])
>>>
>>> od[1]=1
```

```
>>> od
OrderedDict([(2, 2), (3, 3), (1, 1)])
>>>
>>> od.popitem(last=False)
(2, 2)
>>> od
OrderedDict([(3, 3), (1, 1)])
...

```

```
import collections

class LRUCache:
    def __init__(self, capacity: 'int'):
        self.cache = collections.OrderedDict()
        self.remain = capacity

    def get(self, key: 'int') -> 'int':
        if key not in self.cache:
            return -1

        self.cache.move_to_end(key) # meaning end
        is the most recently used
        return self.cache.get(key)

    def put(self, key: 'int', value: 'int') ->
    'None':
        if key not in self.cache:
            if self.remain > 0:
                self.remain -= 1
            else:
                self.cache.popitem(last=False) #
                pop start position
        else:
```

```

        self.cache.pop(key)
        self.cache[key] = value # add to end of
dict, meaning most recently used

```

```
#####
```

```

### below solution with no ordered-dict
class DLinkedList:

```

```

    def __init__(self, key=0, value=0):

```

```

        self.key = key
        self.val = value
        self.next = None
        self.prev = None

```

```

class DLinkedList:

```

```

    def __init__(self):
        self.head = DLinkedList() # dummy head, its
next is real head

```

```

        self.tail = DLinkedList() # dummy tail, its
prev is real tail
        self.head.next, self.tail.prev =
self.tail, self.head

```

```

    def add_node(self, node): # add to head
        node.prev, node.next = self.head,
self.head.next
        node.next.prev, node.prev.next = node,
node

```

```

    def remove_node(self, node):
        node.prev.next, node.next.prev =
node.next, node.prev
        return node.key

```

```

    def remove_tail(self):
        return self.remove_node(self.tail.prev) #
dummy tail's prev is real tail

```

```

    def move_to_head(self, node):
        self.remove_node(node)
        self.add_node(node)

```

```

# def __repr__(self):
# ans = []

```



```
# h = self.head
# while h:
# ans.append(str(h.val))
# h = h.next
# return '<DLinkedList:
{}>'.format('->'.join(ans))
```

```
class LRUCache:
```

```
def __init__(self, capacity: int):
    self.capacity = capacity
```

```
self.nodes_map = {}
self.cache_list = DLinkedList()
```

```
def get(self, key: int) -> int:
    node = self.nodes_map.get(key, None)
    if node:
        self.cache_list.move_to_head(node)
        return node.val
    else:
        return -1
```

```
def put(self, key: int, value: int) -> None:
    node = self.nodes_map.get(key, None)
    if not node:
        if self.capacity > 0:
            self.capacity -= 1
        else:
            rm_key =
self.cache_list.remove_tail()
            self.nodes_map.pop(rm_key) # note
api
        new_node = DLinkedList(key, value)

    self.nodes_map[key] = new_node
    self.cache_list.add_node(new_node)
    else:
        self.cache_list.move_to_head(node)
        node.val = value
```

◆ Quick Review

Key Insights

- Use a Hash Table (dictionary) to allow $O(1)$ access to nodes by key.
- Use a Doubly Linked List to record the usage order. The head represents the most recently used element while the tail represents the least recently used.
- When a key is accessed or updated, move the corresponding node to the front (head) of the list.

- When the cache exceeds capacity, remove the node at the tail end which is the least recently used entry.

Space & Time

Time Complexity: $O(1)$ per get and put operation.

Space Complexity: $O(\text{capacity})$, storing key-node pairs and linked list nodes.

Solution

The solution combines a hash table (or dictionary) and a doubly linked list:

- The hash table stores keys and corresponding node addresses.
- The doubly linked list maintains the order of usage with the most recently used items near the head.
- For `get(key)`, check the dictionary. If found, move the node to the head and return its value; if not, return `-1`.
- For `put(key, value)`, if the key exists, update its value and move it to the head. If not, create a new node and add it right after the head. If the capacity is reached, remove the tail node (least recently used) and update the dictionary accordingly.
- Dummy head and tail nodes are used to simplify node addition and removal procedures without worrying about edge cases.

Linked List

#148 Sort List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers · Divide and Conquer · Sorting · Merge Sort


Lists: Grind 169

Problem

Description

Given the head of a linked list, return the list after sorting it in ascending order.


Example 1:



Input: head = [4,2,1,3]

Output: [1,2,3,4]

Example 2:



Input: head = [-1,5,3,4,0]

Output: [-1,0,3,4,5]

Example 3:

Input: head = []

Output: []

Constraints:

- The number of nodes in the list is in the range [0, 5 * 10⁴].
- -105 ≤ Node.val ≤ 105

Follow up: Can you sort the linked list in $O(n \log n)$ time and $O(1)$ memory (i.e. constant space)?

Community Solutions (Python)

- WalkCC

class Solution:

def sortList(self, head: ListNode) -> ListNode:

def split(head: ListNode, k: int) -> ListNode:

while k > 1 and head:

head = head.next

k -= 1

rest = head.next if head else None

if head:

head.next = None

return rest

def merge(l1: ListNode, l2: ListNode) -> tuple:

dummy = ListNode(0)

tail = dummy

while l1 and l2:

if l1.val > l2.val:

l1, l2 = l2, l1

tail.next = l1

l1 = l1.next

```

    tail = tail.next
    tail.next = l1 if l1 else l2
    while tail.next:
        tail = tail.next

    return dummy.next, tail

length = 0
curr = head
while curr:

```

```

    length += 1
    curr = curr.next

dummy = ListNode(0, head)

k = 1
while k < length:
    curr = dummy.next
    tail = dummy
    while curr:

```

```

    l = curr
    r = split(l, k)
    curr = split(r, k)
    mergedHead, mergedTail = merge(l, r)
    tail.next = mergedHead
    tail = mergedTail
    k *= 2

    return dummy.next

```

• LeetDoocs

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def sortList(self, head: Optional[ListNode])
-> Optional[ListNode]:
        if head is None or head.next is None:
            return head
        slow, fast = head, head.next

```

```

while fast and fast.next:
    slow = slow.next
    fast = fast.next.next
l1, l2 = head, slow.next
slow.next = None
l1, l2 = self.sortList(l1),
self.sortList(l2)
dummy = ListNode()
tail = dummy
while l1 and l2:
    if l1.val <= l2.val:

```

```

        tail.next = l1
        l1 = l1.next
    else:
        tail.next = l2
        l2 = l2.next
    tail = tail.next
tail.next = l1 or l2
return dummy.next

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def sortList(self, head: ListNode) ->
ListNode:
    # Base case: if list is empty or contains a
    single element.
    if not head or not head.next:

```

```

        return head

```

```

    # Function to split the list into two
    halves
    def getMid(head):
        slow, fast = head, head.next
        # Move fast twice as fast as slow to
        locate middle
        while fast and fast.next:
            slow = slow.next
            fast = fast.next.next
        return slow

```

```

# Split the linked list into two halves
mid = getMid(head)
right_head = mid.next
mid.next = None # cut the list into two
parts

# Recursive calls to sort each half
left = self.sortList(head)
right = self.sortList(right_head)

```

```

# Merge two sorted lists
return merge(left, right)

def merge(l1, l2):
    dummy = ListNode(0)
    tail = dummy
    # Merge the lists by comparing nodes
    while l1 and l2:
        if l1.val < l2.val:
            tail.next = l1

```

```

        l1 = l1.next
    else:
        tail.next = l2
        l2 = l2.next
    tail = tail.next
    # Append the remaining nodes
    tail.next = l1 if l1 else l2
    return dummy.next

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def sortList(self, head: ListNode) ->
    ListNode:
        if head is None or head.next is None:
            return head
        slow, fast = head, head.next

```

```

while fast and fast.next:
    slow, fast = slow.next, fast.next.next
    t = slow.next
    slow.next = None
    l1, l2 = self.sortList(head),
self.sortList(t)
    dummy = ListNode()
    cur = dummy
    while l1 and l2:
        if l1.val <= l2.val:
            cur.next = l1
            l1 = l1.next
        else:
            cur.next = l2
            l2 = l2.next
        cur = cur.next
    cur.next = l1 or l2 # add the rest
    return dummy.next

```

◆ Quick Review

Key Insights

- Merge sort is well-suited for linked lists; it naturally relies on splitting the list and merging sorted sublists.
- The slow and fast pointer technique is used to find the middle point of the list.
- A recursive merge sort implementation may use $O(\log n)$ space for recursion stack. To achieve $O(1)$ extra space, an iterative bottom-up merge sort approach would be

preferred.

- Merging two sorted lists is efficient with linked lists because of their sequential access.

Space & Time

Time Complexity: $O(n \log n)$

Space Complexity: $O(1)$ extra space for an iterative bottom-up approach ($O(\log n)$ for recursive calls if using recursion)

Solution

The solution employs the merge sort algorithm tailored for linked lists. The overall steps are:

- Use the slow and fast pointers to find the middle of the list and split it into two halves.
- Recursively sort the two halves.
- Merge the two sorted halves by comparing node values.
- Ensure that during merging, pointers are adjusted properly to build the sorted list.

For an $O(1)$ space solution, an iterative bottom-up merge sort can be implemented. However, the recursive merge sort is easier to understand and implement, though it uses $O(\log n)$ space on the call stack.

Linked List

#328 Odd Even Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List

Lists: Grind 169

Problem

Description


Given the head of a singly linked list, group all the nodes with odd indices together followed by the nodes with even indices, and return the reordered list.

The first node is considered odd, and the second node is even, and so on.

Note that the relative order inside both the even and odd groups should remain as it was in the input.

You must solve the problem in $O(1)$ extra space complexity and $O(n)$ time complexity.


Example 1:



Input: head = [1,2,3,4,5]

Output: [1,3,5,2,4]

Example 2:



Input: head = [2,1,3,5,6,4,7]

Output: [2,3,6,7,1,5,4]

Constraints:

- The number of nodes in the linked list is in the range [0, 104].
- $-106 \leq \text{Node.val} \leq 106$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def oddEvenList(self, head: ListNode) ->
    ListNode:
        oddHead = ListNode(0)
        evenHead = ListNode(0)
        odd = oddHead
        even = evenHead
        isOdd = True

        while head:
            if isOdd:
                odd.next = head
                odd = head
            else:
                even.next = head
                even = head
                head = head.next
                isOdd = not isOdd

        even.next = None
        odd.next = evenHead.next

        return oddHead.next
```

- LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def oddEvenList(self, head:
Optional[ListNode]) -> Optional[ListNode]:
        if head is None:
            return None
        a = head

        b = c = head.next
        while b and b.next:
            a.next = b.next
            a = a.next
            b.next = a.next
            b = b.next
        a.next = c
        return head
```

- SimplyLeet

```
# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def oddEvenList(self, head: ListNode) ->
    ListNode:
        # Check for empty list or list with one
        node.
        if not head or not head.next:

            return head

        odd = head # Point to the first (odd
        indexed) node.
        even = head.next # Point to the second
        (even indexed) node.
        even_head = even # Save the head of the
        even list.

        # Reorder the list by rearranging next
        pointers.
        while even and even.next:
            odd.next = even.next # Connect current
            odd to next odd.
            odd = odd.next # Move odd pointer.
```

Round 5 · Mid · Medium

p.1835

```
        even.next = odd.next # Connect current
        even to next even.
        even = even.next # Move even pointer.

        # Append the even list after the odd list.
        odd.next = even_head
        return head
```

• LC.ca

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def oddEvenList(self, head:
    Optional[ListNode]) -> Optional[ListNode]:
        if head is None:
            return None
        a = head

        b = c = head.next
        while b and b.next:
            a.next = b.next
            a = a.next
            b.next = a.next
            b = b.next
        a.next = c
        return head
```

Round 5 · Mid · Medium

p.1836

◆ Quick Review

Key Insights

- Use two pointers: one for odd-indexed nodes and one for even-indexed nodes.
- The first node is considered odd, and the second node is even.
- Traverse the list once, reassigning the next pointers to segregate odd and even nodes.
- Finally, attach the even list at the end of the odd list.
- The solution must use $O(1)$ extra space and $O(n)$ time complexity.

Space & Time

Time Complexity: $O(n)$ since the list is traversed once.

Space Complexity: $O(1)$ as only a few pointers are used.

Solution

Maintain two pointers, odd and even, starting from the head and head.next, respectively. As you iterate, update odd.next to point to the next odd element (skipping an even node) and even.next to point to the next even element. Once the end is reached, attach the even list to the end of the odd list. This in-place rearrangement preserves the initial ordering within odd and even groups while meeting the space and time constraints.

Linked List

#369 Plus One Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Math · Recursion

Lists: Premium 98

Problem

Description

Given a non-negative integer represented as a linked list of digits, plus one to the integer.

The digits are stored such that the most significant digit is at the head of the list.

Example 1:

Input: head = [1,2,3]
Output: [1,2,4]

Example 2:

Input: head = [0]
Output: [1]

Constraints:

- The number of nodes in the linked list is in the range [1, 100].
- $0 \leq \text{Node.val} \leq 9$
- The number represented by the linked list does not contain leading zeros except for the zero itself.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def plusOne(self, head: ListNode) -> ListNode:
        if not head:
            return ListNode(1)
        if self._addOne(head) == 1:
            return ListNode(1, head)
        return head

    def _addOne(self, node: ListNode) -> int:
        carry = self._addOne(node.next) if node.next
        else 1
```

```
summ = node.val + carry
node.val = summ % 10
return summ // 10
```

```
class Solution:
    def plusOne(self, head: ListNode) -> ListNode:
        dummy = ListNode(0)
        curr = dummy
        dummy.next = head

        while head:
            if head.val != 9:
                curr = head
                head = head.next
```

Round 5 · Mid · Medium

p.1839

```
# `curr` now points to the rightmost non-9
node.
```

```
curr.val += 1
while curr.next:
    curr.next.val = 0
    curr = curr.next
```

```
return dummy.next if dummy.val == 0 else dummy
```

• LeetDoocs

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def plusOne(self, head: Optional[ListNode]) ->
Optional[ListNode]:
        dummy = ListNode(0, head)
        target = dummy
        while head:
```

Round 5 · Mid · Medium

p.1840

```

        if head.val != 9:
            target = head
            head = head.next
        target.val += 1
        target = target.next
    while target:
        target.val = 0
        target = target.next
    return dummy if dummy.val else dummy.next

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

# Helper function to reverse the linked list
def reverseList(head):
    prev = None
    curr = head

```

```

while curr:
    nxt = curr.next # store next node
    curr.next = prev # reverse pointer
    prev = curr # move prev forward
    curr = nxt # move to next node
return prev

```

```

def plusOne(head: ListNode) -> ListNode:
    # Reverse the linked list to start addition
    # from the least significant digit
    head = reverseList(head)

```

```

    current = head
    carry = 1 # initialize carry as the plus one
    while current and carry:
        current.val += carry # add carry to the
        current node
        carry = current.val // 10 # calculate new
        carry
        current.val %= 10 # update the node's value
        # If we're at the end and there's still a
        carry, create a new node.
        if not current.next and carry:
            current.next = ListNode(0)

```

```
current = current.next
```

```
# Reverse the list back to its original order
before returning.
return reverseList(head)
```

• LC.ca

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def plusOne(self, head: ListNode) -> ListNode:
        dummy = ListNode(0, head)
        target = dummy
        while head:
```

```
        if head.val != 9:
            target = head
            head = head.next
        target.val += 1
        target = target.next
    while target:
        target.val = 0
        target = target.next
    return dummy if dummy.val else dummy.next
```

◆ Quick Review

Round 5 · Mid · Medium

p.1843

Key Insights

- The digits are stored such that the most significant digit comes first, but addition is easier starting from the least significant digit.
- Reversing the list allows simpler management of the carry during addition.
- After processing the addition, reversing the list back restores the original order.
- Consider special cases where all digits are 9, requiring a new head node to accommodate an additional digit.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list

Space Complexity: $O(1)$ extra space when using an iterative reversal method. ($O(n)$ if recursion is used)

Solution

We solve the problem by first reversing the linked list, which converts it to a format where the least significant digit is processed first. We then iterate over the reversed list, adding one and managing any carry. If we encounter a carry at the end of the list, a new node is appended to handle the overflow. Finally, we reverse the list again to restore the original order. This method efficiently handles the addition with constant extra space during processing.

Round 5 · Mid · Medium

p.1844

Linked List

#430 Flatten a Multilevel Doubly Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · DFS · Doubly Linked List

Lists: Premium 98

Problem


Description

You are given a doubly linked list, which contains nodes that have a next pointer, a previous pointer, and an additional child pointer. This child pointer may or may not point to a separate doubly linked list, also containing these special nodes. These child lists may have one or more children of their own, and so on, to produce a multilevel data structure as shown in the example below.

Given the head of the first level of the list, flatten the list so that all the nodes appear in a single-level, doubly linked list. Let curr be a node with a child list. The nodes in the child list should appear after curr and before curr.next in the flattened list.

Return the head of the flattened list. The nodes in the list must have all of their child pointers set to null.

Example 1:



Input: head =

[1,2,3,4,5,6,null,null,null,7,8,9,10,null,null,11,12]

Output: [1,2,3,7,8,11,12,9,10,4,5,6]

Explanation: The multilevel linked list in the input is shown.

After flattening the multilevel linked list it becomes:

Example 2:

Input: head = [1,2,null,3]

Output: [1,3,2]

Explanation: The multilevel linked list in the input is shown.

After flattening the multilevel linked list it becomes:

Example 3:

Input: head = []

Output: []

Explanation: There could be empty list in the input.

Constraints:

- The number of Nodes will not exceed 1000.
- $1 \leq \text{Node.val} \leq 105$

How the multilevel linked list is represented in test cases:

We use the multilevel linked list from Example 1 above:

```
1---2---3---4---5---6--NULL
|
7---8---9---10--NULL
|
11--12--NULL
```

The serialization of each level is as follows:

```
[1,2,3,4,5,6,null]
[7,8,9,10,null]
[11,12,null]
```

To serialize all levels together, we will add nulls in each level to signify no node connects to the upper node of the previous level. The serialization becomes:

```
[1, 2, 3, 4, 5, 6, null]
|
[null, null, 7, 8, 9, 10, null]
|
[ null, 11, 12, null]
```

Merging the serialization of each level and removing trailing nulls we obtain:

```
[1,2,3,4,5,6,null,null,null,7,8,9,10,null,null,11,12]
```

Community Solutions (Python)

• WalkCC

```
class Solution:
    def flatten(self, head: 'Node') -> 'Node':
        def flatten(head: 'Node', rest: 'Node') -> 'Node':
            if not head:
                return rest

            head.next = flatten(head.child,
                                flatten(head.next, rest))
            if head.next:
                head.next.prev = head
            head.child = None

        return head

    return flatten(head, None)
```

```
class Solution:
    def flatten(self, head: 'Node') -> 'Node':
        curr = head

        while curr:
            if curr.child:
                cachedNext = curr.next
                curr.next = curr.child
                curr.child.prev = curr
                curr.child = None
```



```

    tail = curr.next
    while tail.next:
        tail = tail.next
    tail.next = cachedNext
    if cachedNext:
        cachedNext.prev = tail
    curr = curr.next

```

```

return head

```

• LeetCode

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val, prev, next, child):
        self.val = val
        self.prev = prev
        self.next = next
        self.child = child
"""

```

```

class Solution:
    def flatten(self, head: 'Node') -> 'Node':
        def preorder(pre, cur):
            if cur is None:
                return pre
            cur.prev = pre
            pre.next = cur

            t = cur.next

```

```

            tail = preorder(cur, cur.child)
            cur.child = None
            return preorder(tail, t)

```

```

        if head is None:
            return None
        dummy = Node(0, None, head, None)
        preorder(dummy, head)
        dummy.next.prev = None
        return dummy.next

```

• SimplyLeet

```

# Definition for a Node.
class Node:
    def __init__(self, val, prev=None, next=None,
child=None):
        self.val = val
        self.prev = prev
        self.next = next
        self.child = child

def flatten(head):
    # If the list is empty, simply return None.

    if not head:
        return head

    # Helper function to perform DFS and return the
    tail of the flattened list.
    def dfs(node):
        current = node
        last = None
        while current:
            next_node = current.next # Store
            current.next before modifications
            # If current node has a child, we
            flatten the child list recursively.

```

```

        if current.child:
            # Recursively flatten the child
            list.

            child_tail = dfs(current.child)
            # Connect current node with the
            start of the child list.
            current.next = current.child
            current.child.prev = current
            # If there was a next node, attach
            it after the child list.
            if next_node:
                child_tail.next = next_node
                next_node.prev = child_tail

        # Set child pointer to None as
        required.

        current.child = None
        # Update the tail pointer.
        last = child_tail
    else:
        last = current
    # Move to the next node (which might
    have been updated after flattening a child).
    current = next_node
    return last

```

```
dfs(head) # Begin DFS traversal from the head.
return head
```

```
# Example to illustrate usage:
# head = Node(1)
# head.next = Node(2)
# head.child = Node(3)
# flatten(head)
```

• LC.ca

```
"""
# Definition for a Node.
class Node:
    def __init__(self, val, prev, next, child):
        self.val = val
        self.prev = prev
        self.next = next
        self.child = child
"""
```

```
class Solution:
    def flatten(self, head: 'Node') -> 'Node':
        def preorder(pre, cur):
            if cur is None:
                return pre
            cur.prev = pre
            pre.next = cur

            t = cur.next
```

```
tail = preorder(cur, cur.child)
cur.child = None
return preorder(tail, t)
```

```
if head is None:
    return None
dummy = Node(0, None, head, None)
preorder(dummy, head)
dummy.next.prev = None
return dummy.next
```

◆ Quick Review

Key Insights

- Use a depth-first traversal to process child lists before proceeding to the next pointer.
- When a node with a child is found, recursively flatten the child list.
- Splice the flattened child list in between the current node and the next node.
- Maintain previous and next pointers correctly to preserve the doubly linked list properties.
- Ensure that all child pointers are set to null after flattening.

Space & Time

Time Complexity: $O(n)$, where n is the total number of nodes, since each node is processed exactly once.

Space Complexity: $O(n)$ in the worst-case due to recursion stack in case of a highly nested list.

Solution

We can solve this problem using a recursive depth-first search (DFS) approach. The main idea is to traverse the list and, whenever a node with a child pointer is encountered, recursively flatten its child list. Once the child list is flattened, it can be spliced into the main list between the current node and the node originally following it. We then continue processing the next pointer. This ensures that the entire list is flattened in depth-first order. The recursive helper function returns the tail of the flattened segment, which is used to connect the subsequent parts of the list correctly.

Linked List

#622 Design Circular Queue

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Linked List · Design · Queue

Lists: Denny Zhang

Problem

Description

Design your implementation of the circular queue. The circular queue is a linear data structure in which the operations are performed based on FIFO (First In First Out) principle, and the last position is connected back to the first position to make a circle. It is also called "Ring Buffer".

One of the benefits of the circular queue is that we can make use of the spaces in front of the queue. In a normal queue, once the queue becomes full, we cannot insert the next element even if there is a space in front of the queue. But using the circular queue, we can use the space to store new values.

Implement the MyCircularQueue class:

- `MyCircularQueue(k)` Initializes the object with the size of the queue to be `k`.
- `int Front()` Gets the front item from the queue. If the queue is empty, return `-1`.

- `int Rear()` Gets the last item from the queue. If the queue is empty, return -1.
- `boolean enqueue(int value)` Inserts an element into the circular queue. Return true if the operation is successful.
- `boolean dequeue()` Deletes an element from the circular queue. Return true if the operation is successful.
- `boolean isEmpty()` Checks whether the circular queue is empty or not.
- `boolean isFull()` Checks whether the circular queue is full or not.

You must solve the problem without using the built-in queue data structure in your programming language.

Example 1:

```
Input
["MyCircularQueue", "enqueue", "enqueue", "enqueue",
"enqueue", "Rear", "isFull", "dequeue", "enqueue",
"Rear"]
[[3], [1], [2], [3], [4], [], [], [], [4], []]
Output
[null, true, true, true, false, 3, true, true, true, 4]
```

Explanation

```
MyCircularQueue myCircularQueue = new
MyCircularQueue(3);
```

Constraints:

- $1 \leq k \leq 1000$
- $0 \leq \text{value} \leq 1000$

- At most 3000 calls will be made to `enqueue`, `dequeue`, `Front`, `Rear`, `isEmpty`, and `isFull`.

Community Solutions (Python)

• LeetDoocs

```
class MyCircularQueue:

    def __init__(self, k: int):
        self.q = [0] * k
        self.size = 0
        self.capacity = k
        self.front = 0

    def enqueue(self, value: int) -> bool:
        if self.isFull():

            return False
        self.q[(self.front + self.size) %
self.capacity] = value
        self.size += 1
        return True

    def dequeue(self) -> bool:
        if self.isEmpty():
            return False
        self.front = (self.front + 1) %
self.capacity
        self.size -= 1
```

```

        return True

    def Front(self) -> int:
        return -1 if self.isEmpty() else
self.q[self.front]

    def Rear(self) -> int:
        if self.isEmpty():
            return -1
        return self.q[(self.front + self.size - 1)
% self.capacity]

    def isEmpty(self) -> bool:
        return self.size == 0

    def isFull(self) -> bool:
        return self.size == self.capacity

# Your MyCircularQueue object will be instantiated
and called as such:
# obj = MyCircularQueue(k)
# param_1 = obj.enqueue(value)

# param_2 = obj.dequeue()
# param_3 = obj.front()
# param_4 = obj.rear()
# param_5 = obj.isEmpty()
# param_6 = obj.isFull()

```

• SimplyLeet

Round 5 · Mid · Medium

p.1859

```

# Python implementation of MyCircularQueue

class MyCircularQueue:
    def __init__(self, k):
        # Initialize the queue with capacity k
        self.capacity = k
        self.queue = [0] * k # fixed-size array
        self.front = 0 # pointer to the front
        element
        self.rear = 0 # pointer to the next
        insertion position
        self.count = 0 # current number of elements
        in the queue

    def enqueue(self, value):
        # Check if the queue is full
        if self.isFull():
            return False
        # Insert element at rear pointer and update
        pointer using modulo for circularity
        self.queue[self.rear] = value
        self.rear = (self.rear + 1) % self.capacity
        self.count += 1
        return True

```

Round 5 · Mid · Medium

p.1860

```
def deQueue(self):
    # Check if the queue is empty
    if self.isEmpty():
        return False
    # Move front pointer to the next element
    using modulo for circularity
    self.front = (self.front + 1) %
self.capacity
    self.count -= 1
    return True
```

```
def Front(self):
    # Return -1 if the queue is empty, else
    return the front element
    if self.isEmpty():
        return -1
    return self.queue[self.front]

def Rear(self):
    # Return -1 if the queue is empty
    if self.isEmpty():
        return -1
```

```
# Since rear points to the next insertion
position, get the last element index
return self.queue[(self.rear - 1 +
self.capacity) % self.capacity]
```

```
def isEmpty(self):
    # Check if count is 0
    return self.count == 0

def isFull(self):
    # Check if the queue has reached its
    capacity
    return self.count == self.capacity
```

• LC.ca

```
class MyCircularQueue:
    def __init__(self, k: int):
        self.q = [0] * k
        self.front = 0
        self.size = 0
        self.capacity = k

    def enqueue(self, value: int) -> bool:
        if self.isFull():
            return False
```

```

        idx = (self.front + self.size) %
self.capacity
        self.q[idx] = value
        self.size += 1
        return True

    def deQueue(self) -> bool:
        if self.isEmpty():
            return False
        self.front = (self.front + 1) %
self.capacity
        self.size -= 1

        return True

    def Front(self) -> int:
        return -1 if self.isEmpty() else
self.q[self.front]

    def Rear(self) -> int:
        if self.isEmpty():
            return -1
        idx = (self.front + self.size - 1) %
self.capacity
        return self.q[idx]

```

```

    def isEmpty(self) -> bool:
        return self.size == 0

    def isFull(self) -> bool:
        return self.size == self.capacity

# Your MyCircularQueue object will be instantiated
and called as such:
# obj = MyCircularQueue(k)

# param_1 = obj.enqueue(value)
# param_2 = obj.dequeue()
# param_3 = obj.front()
# param_4 = obj.rear()
# param_5 = obj.isEmpty()
# param_6 = obj.isFull()

#####

class MyCircularQueue(object):

```



```
def __init__(self, k):
    """
    Initialize your data structure here. Set
    the size of the queue to be k.
    :type k: int
    """
    self.queue = []
    self.size = k
    self.front = 0
    self.rear = 0
```

```
def enQueue(self, value):
    """
    Insert an element into the circular queue.
    Return true if the operation is successful.
    :type value: int
    :rtype: bool
    """
    if self.rear - self.front < self.size:
        self.queue.append(value)
        self.rear += 1
```

```
        return True
    else:
        return False

def deQueue(self):
    """
    Delete an element from the circular queue.
    Return true if the operation is successful.
    :rtype: bool
    """
    if self.rear - self.front > 0:
```

```
        self.front += 1
        return True
    else:
        return False

def Front(self):
    """
    Get the front item from the queue.
    :rtype: int
    """
```

```

    if self.isEmpty():
        return -1
    else:
        return self.queue[self.front]

```

```

def Rear(self):
    """
    Get the last item from the queue.
    :rtype: int
    """

```

```

    if self.isEmpty():
        return -1
    else:
        return self.queue[self.rear - 1]

```

```

def isEmpty(self):
    """
    Checks whether the circular queue is empty
    or not.
    :rtype: bool
    """

```

```

    return self.front == self.rear

```

```

def isFull(self):
    """
    Checks whether the circular queue is full
    or not.
    :rtype: bool
    """
    return self.rear - self.front == self.size

```

```

# Your MyCircularQueue object will be instantiated
and called as such:
# obj = MyCircularQueue(k)
# param_1 = obj.enqueue(value)
# param_2 = obj.dequeue()
# param_3 = obj.front()
# param_4 = obj.rear()
# param_5 = obj.isEmpty()
# param_6 = obj.isFull()

```

◆ Quick Review

Key Insights

- Leverage a fixed-size array to store the queue elements.
- Use two pointers: one for the front and one for the rear.
- Maintain a count variable (or use pointer arithmetic) to differentiate between an empty and a full queue.

- Use modulo arithmetic to wrap pointers when they reach the end of the array.
- Ensure all operations (enqueue, dequeue, front, rear, isEmpty, isFull) have $O(1)$ time complexity.

Space & Time

Time Complexity: $O(1)$ per operation for enqueue, dequeue, front, rear, isEmpty, and isFull.

Space Complexity: $O(k)$, where k is the capacity of the circular queue.

Solution

We implement the circular queue using a fixed-size array alongside two pointers, front and rear, and a counter to keep track of the number of elements in the queue. The front pointer indicates the index of the first element, and the rear pointer indicates the index where the next element will be enqueued. When enqueueing an element, we insert the element at the rear index and then increment the rear pointer modulo the size of the array. Similarly, when dequeuing, we remove the element at the front index and increment the front pointer modulo the size. The counter helps in quickly checking if the queue is empty or full (count equals 0 or count equals capacity, respectively). This design ensures all operations remain $O(1)$ and correctly wraps around, forming a circular structure.

Linked List

#707 Design Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Design

Lists: Denny Zhang

Problem

Description

Design your implementation of the linked list. You can choose to use a singly or doubly linked list. A node in a singly linked list should have two attributes: val and next. val is the value of the current node, and next is a pointer/reference to the next node. If you want to use the doubly linked list, you will need one more attribute prev to indicate the previous node in the linked list. Assume all nodes in the linked list are 0-indexed.

Implement the MyLinkedList class:

- MyLinkedList() Initializes the MyLinkedList object.
- int get(int index) Get the value of the indexth node in the linked list. If the index is invalid, return -1.
- void addAtHead(int val) Add a node of value val before the first element of the linked list. After the insertion, the new node will be the first node of the linked list.

- **void addAtTail(int val)** Append a node of value val as the last element of the linked list.
- **void addAtIndex(int index, int val)** Add a node of value val before the indexth node in the linked list. If index equals the length of the linked list, the node will be appended to the end of the linked list. If index is greater than the length, the node will not be inserted.
- **void deleteAtIndex(int index)** Delete the indexth node in the linked list, if the index is valid.

Example 1:

Input
 ["MyLinkedList", "addAtHead", "addAtTail",
 "addAtIndex", "get", "deleteAtIndex", "get"]
 [[], [1], [3], [1, 2], [1], [1], [1]]

Output
 [null, null, null, null, 2, null, 3]

Explanation
 MyLinkedList myLinkedList = new MyLinkedList();

Constraints:

- $0 \leq \text{index}, \text{val} \leq 1000$
- Please do not use the built-in LinkedList library.
- At most 2000 calls will be made to get, addAtHead, addAtTail, addAtIndex and deleteAtIndex.

Community Solutions (Python)

- WalkCC

```
from dataclasses import dataclass
```

```
@dataclass
class ListNode:
    val: int
    next: ListNode | None = None
```

```
class MyLinkedList:
```

```
    def __init__(self):
        self.length = 0
        self.dummy = ListNode(0)

    def get(self, index: int) -> int:
        if index < 0 or index >= self.length:
            return -1
        curr = self.dummy.next
        for _ in range(index):
            curr = curr.next
```

```
return curr.val
```

```
def addAtHead(self, val: int) -> None:
    curr = self.dummy.next
    self.dummy.next = ListNode(val)
    self.dummy.next.next = curr
    self.length += 1
```

```
def addAtTail(self, val: int) -> None:
    curr = self.dummy
```

```
while curr.next:
    curr = curr.next
curr.next = ListNode(val)
self.length += 1
```

```
def addAtIndex(self, index: int, val: int) ->
None:
    if index > self.length:
        return
    curr = self.dummy
    for _ in range(index):
```

```
curr = curr.next
temp = curr.next
curr.next = ListNode(val)
curr.next.next = temp
self.length += 1
```

```
def deleteAtIndex(self, index: int) -> None:
    if index < 0 or index >= self.length:
        return
    curr = self.dummy
```

```
for _ in range(index):
    curr = curr.next
temp = curr.next
curr.next = temp.next
self.length -= 1
```

• LeetCode

```
class MyLinkedList:
    def __init__(self):
        self.dummy = ListNode()
        self.cnt = 0

    def get(self, index: int) -> int:
        if index < 0 or index >= self.cnt:
            return -1
        cur = self.dummy.next
        for _ in range(index):
```

```

        cur = cur.next
    return cur.val

def addAtHead(self, val: int) -> None:
    self.addAtIndex(0, val)

def addAtTail(self, val: int) -> None:
    self.addAtIndex(self.cnt, val)

def addAtIndex(self, index: int, val: int) ->
None:

    if index > self.cnt:
        return
    pre = self.dummy
    for _ in range(index):
        pre = pre.next
    pre.next = ListNode(val, pre.next)
    self.cnt += 1

def deleteAtIndex(self, index: int) -> None:
    if index >= self.cnt:

```

```

        return
    pre = self.dummy
    for _ in range(index):
        pre = pre.next
    t = pre.next
    pre.next = t.next
    t.next = None
    self.cnt -= 1

# Your MyLinkedList object will be instantiated and
# called as such:
# obj = MyLinkedList()
# param_1 = obj.get(index)
# obj.addAtHead(val)
# obj.addAtTail(val)
# obj.addAtIndex(index, val)
# obj.deleteAtIndex(index)

class MyLinkedList:
    def __init__(self):
        self.e = [0] * 1010
        self.ne = [0] * 1010
        self.idx = 0
        self.head = -1
        self.cnt = 0

    def get(self, index: int) -> int:
        if index < 0 or index >= self.cnt:

```

```

        return -1
    i = self.head
    for _ in range(index):
        i = self.ne[i]
    return self.e[i]

def addAtHead(self, val: int) -> None:
    self.e[self.idx] = val
    self.ne[self.idx] = self.head
    self.head = self.idx

```

```

    self.idx += 1
    self.cnt += 1

def addAtTail(self, val: int) -> None:
    self.addAtIndex(self.cnt, val)

def addAtIndex(self, index: int, val: int) ->
None:
    if index > self.cnt:
        return
    if index <= 0:

```

```

        self.addAtHead(val)
        return
    i = self.head
    for _ in range(index - 1):
        i = self.ne[i]
    self.e[self.idx] = val
    self.ne[self.idx] = self.ne[i]
    self.ne[i] = self.idx
    self.idx += 1
    self.cnt += 1

```

```

def deleteAtIndex(self, index: int) -> None:
    if index < 0 or index >= self.cnt:
        return -1
    self.cnt -= 1
    if index == 0:
        self.head = self.ne[self.head]
        return
    i = self.head
    for _ in range(index - 1):

```

```

        i = self.ne[i]
        self.ne[i] = self.ne[self.ne[i]]

```

Your MyLinkedList object will be instantiated and called as such:

```

# obj = MyLinkedList()
# param_1 = obj.get(index)
# obj.addAtHead(val)
# obj.addAtTail(val)
# obj.addAtIndex(index,val)

```

```

# obj.deleteAtIndex(index)

```

• SimplyLeet

Define a node for the singly linked list

```

class ListNode:

```

```

    def __init__(self, val=0, next=None):
        self.val = val # The value of the node
        self.next = next # Pointer to the next node

```

```

class MyLinkedList:

```

```

    def __init__(self):
        # Initialize size and a dummy head node to
        # simplify operations
        self.size = 0

```

```

        self.head = ListNode(0) # dummy head

```

```

    def get(self, index: int) -> int:
        # If index is out of bounds, return -1
        if index < 0 or index >= self.size:
            return -1
        current = self.head.next # start from the
        actual head
        for _ in range(index):
            current = current.next
        return current.val

```

```

    def addAtHead(self, val: int) -> None:
        self.addAtIndex(0, val) # add at index 0

```

```

    def addAtTail(self, val: int) -> None:
        self.addAtIndex(self.size, val) # add at
        the end

```

```

    def addAtIndex(self, index: int, val: int) ->
    None:
        # Normalize negative index to 0 (insert at
        head)
        if index < 0:

```



```

        index = 0
    # Only insert if index is valid (<= size)
    if index > self.size:
        return
    self.size += 1
    # Start from the dummy head node
    pred = self.head
    for _ in range(index):
        pred = pred.next
    # Insert new node

```

```

new_node = ListNode(val)
new_node.next = pred.next
pred.next = new_node

```

```

def deleteAtIndex(self, index: int) -> None:
    # If index is out of bounds, do nothing
    if index < 0 or index >= self.size:
        return
    self.size -= 1
    pred = self.head

```

```

    for _ in range(index):
        pred = pred.next
    # Remove the node by skipping it
    pred.next = pred.next.next

```

• LC.ca

Round 5 · Mid · Medium

p.1881

```

# doubly linked node
class ListNode:
    def __init__(self, value=0, prev=None,
next=None):
        self.value = value
        self.prev = prev
        self.next = next

class MyLinkedList:
    def __init__(self):
        self.head = ListNode(0) # Sentinel node as
pseudo-head

        self.tail = ListNode(0) # Sentinel node as
pseudo-tail
        self.head.next = self.tail
        self.tail.prev = self.head
        self.size = 0

    def get(self, index: int) -> int:
        if index < 0 or index >= self.size:
            return -1
        if index + 1 < self.size - index: # decide
start from head or tail
            curr = self.head

```

Round 5 · Mid · Medium

p.1882

```

        for _ in range(index + 1):
            curr = curr.next
    else:
        curr = self.tail
        for _ in range(self.size - index):
            curr = curr.prev
    return curr.value

def addAtHead(self, val: int) -> None:
    pred, succ = self.head, self.head.next

```

```

    self.size += 1
    to_add = ListNode(val, pred, succ)
    pred.next = to_add
    succ.prev = to_add

def addAtTail(self, val: int) -> None:
    succ, pred = self.tail, self.tail.prev
    self.size += 1
    to_add = ListNode(val, pred, succ)
    pred.next = to_add

```

```

    succ.prev = to_add

def addAtIndex(self, index: int, val: int) ->
None:
    if index > self.size:
        return
    if index < 0:
        index = 0
    if index < self.size - index: # decide
start from head or tail
        pred = self.head
        for _ in range(index):

```

```

            pred = pred.next
            succ = pred.next
        else:
            succ = self.tail
            for _ in range(self.size - index):
                succ = succ.prev
            pred = succ.prev
        self.size += 1
        to_add = ListNode(val, pred, succ)
        pred.next = to_add

```

```

succ.prev = to_add

def deleteAtIndex(self, index: int) -> None:
    if index < 0 or index >= self.size:
        return
    if index < self.size - index:
        pred = self.head
        for _ in range(index):
            pred = pred.next
        succ = pred.next.next
    else:
        succ = self.tail
        for _ in range(self.size - index - 1):
            succ = succ.prev
        pred = succ.prev.prev
    self.size -= 1
    pred.next = succ
    succ.prev = pred

# Your MyLinkedList object will be instantiated and
called as such:

```

```

# obj = MyLinkedList()
# param_1 = obj.get(index)
# obj.addAtHead(val)
# obj.addAtTail(val)
# obj.addAtIndex(index, val)
# obj.deleteAtIndex(index)

#####

# singly linked node

class MyLinkedList:
    def __init__(self):
        self.dummy = ListNode()
        self.cnt = 0

    def get(self, index: int) -> int:
        if index < 0 or index >= self.cnt:
            return -1
        cur = self.dummy.next
        for _ in range(index):

```

```

        cur = cur.next
    return cur.val

def addAtHead(self, val: int) -> None:
    self.addAtIndex(0, val)

def addAtTail(self, val: int) -> None:
    self.addAtIndex(self.cnt, val)

def addAtIndex(self, index: int, val: int) ->
None:
    if index > self.cnt:
        return
    pre = self.dummy
    for _ in range(index):
        pre = pre.next
    pre.next = ListNode(val, pre.next) #
insert in-between
    self.cnt += 1

def deleteAtIndex(self, index: int) -> None:
    if index >= self.cnt:

```

```

        return
    pre = self.dummy
    for _ in range(index):
        pre = pre.next
    t = pre.next
    pre.next = t.next
    t.next = None # don't forget to cutoff
    self.cnt -= 1

```

```

# Your MyLinkedList object will be instantiated and
called as such:
# obj = MyLinkedList()
# param_1 = obj.get(index)
# obj.addAtHead(val)
# obj.addAtTail(val)
# obj.addAtIndex(index, val)
# obj.deleteAtIndex(index)

```

◆ Quick Review

Key Insights

- Use a dummy head node to simplify insertion and deletion at the head.
- Maintain a size counter to quickly validate index-based operations.
- For every operation, traverse the list up to the required point since random access is not available.

- Ensure robustness by handling edge cases like negative indices and indices larger than the current length.

Space & Time

Time Complexity:

- get, addAtIndex, and deleteAtIndex operations require $O(n)$ time due to the need to traverse the list.
- addAtHead and addAtTail (if tail pointer is not maintained) are $O(n)$ in worst case; however, these can be optimized further if a tail pointer is added.

Space Complexity:

- $O(n)$ space for storing n nodes.

Solution

The solution employs a singly linked list with a dummy head node to ease edge-case handling. A size counter is maintained for quick index validations. Each node holds a value and a pointer to the next node. Operations operate by traversing the list to the correct insertion or deletion point. This design makes it simple to add the functionality of getting values, inserting at specified indices, and removing nodes.

Linked List

#708 Insert into a Sorted Circular Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List

Lists: Premium 98

Problem


Description

Given a Circular Linked List node, which is sorted in non-descending order, write a function to insert a value insertVal into the list such that it remains a sorted circular list. The given node can be a reference to any single node in the list and may not necessarily be the smallest value in the circular list.

If there are multiple suitable places for insertion, you may choose any place to insert the new value. After the insertion, the circular list should remain sorted.

If the list is empty (i.e., the given node is null), you should create a new single circular list and return the reference to that single node. Otherwise, you should return the originally given node.

Example 1:



Input: head = [3,4,1], insertVal = 2

Output: [3,4,1,2]

Explanation: In the figure above, there is a sorted circular list of three elements. You are given a reference to the node with value 3, and we need to insert 2 into the list. The new node should be inserted between node 1 and node 3. After the insertion, the list should look like this, and we should still return node 3.

Example 2:

Input: head = [], insertVal = 1

Output: [1]

Explanation: The list is empty (given head is null). We create a new single circular list and return the reference to that single node.

Example 3:

Input: head = [1], insertVal = 0

Output: [1,0]

Constraints:

- The number of nodes in the list is in the range [0, 5 * 10⁴].
- -106 ≤ Node.val, insertVal ≤ 106

Community Solutions (Python)

• LeetCode

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, next=None):
        self.val = val
        self.next = next
"""

```

class Solution:

```

def insert(self, head: 'Optional[Node]',
insertVal: int) -> 'Node':
    node = Node(insertVal)
    if head is None:
        node.next = node
        return node
    prev, curr = head, head.next
    while curr != head:
        if prev.val <= insertVal <= curr.val or
(
            prev.val > curr.val and (insertVal
>= prev.val or insertVal <= curr.val)
        ):
            break
        prev, curr = curr, curr.next
    prev.next = node
    node.next = curr
    return head

```

• SimplyLeet

```

# Definition for a Node.
class Node:
    def __init__(self, val, next=None):
        self.val = val
        self.next = next

class Solution:
    def insert(self, head: 'Node', insertVal: int)
-> 'Node':
        # If the list is empty, create a new node
        pointing to itself and return it.
        if not head:
            newNode = Node(insertVal)
            newNode.next = newNode
            return newNode

        curr = head
        inserted = False

        while True:
            # Case 1: insert between two nodes
            where insertVal is between curr.val and
            curr.next.val.
            if curr.val <= insertVal <=
            curr.next.val:

```

```

        inserted = True
        # Case 2: curr.val > curr.next.val,
        i.e., we are at the turning point.
        # In this case, insert if insertVal is
        greater than curr.val (largest)
        # or insertVal is less than
        curr.next.val (smallest).
        elif curr.val > curr.next.val:
            if insertVal >= curr.val or
insertVal <= curr.next.val:
                inserted = True

        if inserted:
            newNode = Node(insertVal,
curr.next)

```

```

        curr.next = newNode
        return head

        curr = curr.next
        # If we have traversed the entire
        circle and came back to head, break.
        if curr == head:
            break

        # Case 3: All values are the same or no
        valid index was found, insert arbitrarily.
        newNode = Node(insertVal, curr.next)

```

```

curr.next = newNode
return head

```

• LC.ca

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, next=None):
        self.val = val
        self.next = next
"""

class Solution:

    def insert(self, head: 'Optional[Node]',
insertVal: int) -> 'Node':
        node = Node(insertVal)
        if head is None:
            node.next = node
            return node
        prev, curr = head, head.next
        while curr != head:
            if prev.val <= insertVal <= curr.val or
(
                prev.val > curr.val and (insertVal
>= prev.val or insertVal <= curr.val)
            ):

```



```

        break
    prev, curr = curr, curr.next
    prev.next = node
    node.next = curr
    return head

```

◆ Quick Review

Key Insights

- Handle the empty list edge-case by creating a new node that points to itself.
- Traverse the list until the proper insertion point is found.
- Consider the "turning point" where the maximum value is followed by the minimum value.
- If no proper place is found in one full rotation, insert the new node arbitrarily (e.g., after the head).

Space & Time

Time Complexity: $O(n)$ in the worst-case when a full loop is required.

Space Complexity: $O(1)$ since we only use a few extra pointers.

Solution

We traverse the circular linked list starting from the arbitrary given node. For each pair of consecutive nodes (curr and next), we check if the new value can be inserted between them by confirming that:

– The new value lies between curr and next ($\text{curr.val} \leq \text{insertVal} \leq \text{next.val}$). – Or if we are at the pivot point where the maximum value meets the minimum value (i.e., $\text{curr.val} > \text{next.val}$) and then either the new value is larger than or equal to curr.val (should be appended after the maximum) or less than or equal to next.val (should be inserted before the minimum).

If no proper spot is found after one complete loop, we insert anywhere (all values in the list are the same). A pointer reference is updated accordingly to maintain the circular structure.

Linked List

#1171 Remove Zero Sum Consecutive Nodes from Linked List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Linked List

Lists: Denny Zhang

Problem

Description

Given the head of a linked list, we repeatedly delete consecutive sequences of nodes that sum to 0 until there are no such sequences.

After doing so, return the head of the final linked list. You may return any such answer.

(Note that in the examples below, all sequences are serializations of ListNode objects.)

Example 1:

Input: head = [1,2,-3,3,1]

Output: [3,1]

Note: The answer [1,2,1] would also be accepted.

Example 2:

Input: head = [1,2,3,-3,4]

Output: [1,2,4]

Example 3:

Input: head = [1,2,3,-3,-2]

Output: [1]

Constraints:

- The given linked list will contain between 1 and 1000 nodes.
- Each node in the linked list has $-1000 \leq \text{node.val} \leq 1000$.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def removeZeroSumSublists(self, head: ListNode)
-> ListNode:
    dummy = ListNode(0, head)
    prefix = 0
    prefixToNode = {0: dummy}

    while head:
        prefix += head.val
        prefixToNode[prefix] = head
        head = head.next
```

```

prefix = 0
head = dummy

while head:
    prefix += head.val
    head.next = prefixToNode[prefix].next
    head = head.next

return dummy.next

```

- LeetCode

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def removeZeroSumSublists(self, head:
Optional[ListNode]) -> Optional[ListNode]:
        dummy = ListNode(next=head)
        last = {}
        s, cur = 0, dummy

```

```

while cur:
    s += cur.val
    last[s] = cur
    cur = cur.next
s, cur = 0, dummy
while cur:
    s += cur.val
    cur.next = last[s].next
    cur = cur.next
return dummy.next

```

- SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def removeZeroSumSublists(self, head:
ListNode) -> ListNode:
        # Create a dummy node to simplify cases
        # where removal starts at the head.
        dummy = ListNode(0)

```

```

dummy.next = head

# Dictionary to store prefix sum to node
mapping.
sum_to_node = {}
prefix_sum = 0
current = dummy

# First pass: record the last occurrence of
each prefix sum.
while current:
    prefix_sum += current.val

    sum_to_node[prefix_sum] = current
    current = current.next

# Second pass: remove nodes that are part
of zero-sum sequences.
prefix_sum = 0
current = dummy
while current:
    prefix_sum += current.val
    # Link current node's next pointer to
    skip the zero-sum subsequence.
    current.next =
sum_to_node[prefix_sum].next

    current = current.next

return dummy.next

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def removeZeroSumSublists(self, head:
Optional[ListNode]) -> Optional[ListNode]:
        dummy = ListNode(next=head)
        last = {}
        s, cur = 0, dummy

        while cur:
            s += cur.val
            last[s] = cur
            cur = cur.next
        s, cur = 0, dummy
        while cur:
            s += cur.val
            cur.next = last[s].next
            cur = cur.next
        return dummy.next

```

◆ Quick Review

Key Insights

- Use a dummy node to handle edge cases such as the removal of nodes at the head.
- Compute a running (prefix) sum while traversing the list.

- Use a hash table (or hashmap) to map each prefix sum to the most recent node where that sum occurred.
- If the same prefix sum is encountered again, the nodes in between sum to 0 and can be removed by adjusting pointers.
- Perform two passes: the first to record prefix sums and their corresponding nodes, and the second to remove zero-sum sequences.

Space & Time

Time Complexity: $O(n)$ – Two passes over the list.

Space Complexity: $O(n)$ – Hash table storing at most n prefix sums.

Solution

The solution uses a two-pass algorithm with a hash table. In the first pass, compute the prefix sum for each node and store the mapping from prefix sum to the node (overwriting with the latest occurrence). This ensures that for any repeated prefix sum, the nodes in between sum to 0. In the second pass, traverse the list again and use the hash table to skip the zero-sum subsequences by redirecting the next pointer of a node to the next pointer of the stored node for that prefix sum. Key data structures are the hash table and a dummy node to simplify pointer adjustments.

Stack

Round 5 · Mid · Medium · 14 questions

Stack

#143 Reorder List

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Two Pointers · Stack · Recursion

Lists: Grind 169

Problem

Description

You are given the head of a singly linked-list. The list can be represented as:



$$L_0 \rightarrow L_1 \rightarrow \dots \rightarrow L_{n-1} \rightarrow L_n$$

Reorder the list to be on the following form:

$$L_0 \rightarrow L_n \rightarrow L_1 \rightarrow L_{n-1} \rightarrow L_2 \rightarrow L_{n-2} \rightarrow \dots$$

You may not modify the values in the list's nodes. Only nodes themselves may be changed.



Example 1:



Input: head = [1,2,3,4]

Output: [1,4,2,3]

Example 2:



Input: head = [1,2,3,4,5]

Output: [1,5,2,4,3]

Constraints:

- The number of nodes in the list is in the range [1, 5 * 10⁴].
- 1 ≤ Node.val ≤ 1000

Community Solutions (Python)

• WalkCC

```
class Solution:
    def reorderList(self, head: ListNode) -> None:
        def findMid(head: ListNode):
            prev = None
            slow = head
            fast = head

            while fast and fast.next:
                prev = slow
                slow = slow.next
                fast = fast.next.next

            prev.next = None

            return slow

        def reverse(head: ListNode) -> ListNode:
            prev = None
            curr = head

            while curr:
```

```
        next = curr.next
        curr.next = prev
        prev = curr
        curr = next
```

```
    return prev
```

```
def merge(l1: ListNode, l2: ListNode) -> None:
    while l2:
        next = l1.next
```

```
        l1.next = l2
        l1 = l2
        l2 = next
```

```
    if not head or not head.next:
        return
```

```
    mid = findMid(head)
    reversed = reverse(mid)
    merge(head, reversed)
```

• LeetCode

```
# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def reorderList(self, head:
Optional[ListNode]) -> None:
        fast = slow = head
        while fast.next and fast.next.next:
            slow = slow.next
```

```
            fast = fast.next.next
```

```
        cur = slow.next
        slow.next = None
```

```
        pre = None
        while cur:
            t = cur.next
            cur.next = pre
            pre, cur = cur, t
```

```
        cur = head
```

```
        while pre:
            t = pre.next
            pre.next = cur.next
            cur.next = pre
            cur, pre = pre.next, t
```

• SimplyLeet

```
# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def reorderList(self, head: ListNode) -> None:
        if not head or not head.next:
            return
```

```
        # Step 1: Find the middle of the linked
        list using slow and fast pointers.
```

```
        slow, fast = head, head
        while fast and fast.next:
            slow = slow.next
            fast = fast.next.next
```

```
        # Step 2: Reverse the second half of the
        list.
```

```
        prev, curr = None, slow.next
        slow.next = None # Split the list into two
        halves.
```



```

        while curr:
            next_temp = curr.next # Save next node.
            curr.next = prev # Reverse current
pointer.
            prev = curr # Move prev to current.
            curr = next_temp # Move to the next
node.

        # Step 3: Merge the two halves, alternating
nodes.
        first, second = head, prev
        while second:
            temp1 = first.next # Save next node
from first half.

            temp2 = second.next # Save next node
from reversed second half.
            first.next = second # Link first node
to second list node.
            second.next = temp1 # Link second node
to the next node in first half.
            first = temp1 # Move pointer in first
half.
            second = temp2 # Move pointer in second
half.

```

• LC.ca

Round 5 · Mid · Medium

p.1913

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def reorderList(self, head:
Optional[ListNode]) -> None:
        fast = slow = head
        while fast.next and fast.next.next:
            slow = slow.next

            fast = fast.next.next

            cur = slow.next
            slow.next = None

            pre = None
            while cur:
                t = cur.next
                cur.next = pre
                pre, cur = cur, t

            cur = head

            while pre:
                t = pre.next
                pre.next = cur.next
                cur.next = pre
                cur, pre = pre.next, t

```

Round 5 · Mid · Medium

p.1914

◆ Quick Review

Key Insights

- Find the middle of the linked list using the fast and slow pointer technique.
- Reverse the second half of the list.
- Merge the two halves, alternating nodes from the first half and the reversed second half.
- The challenge is done in-place, ensuring a linear time solution with constant extra space.

Space & Time

Time Complexity: $O(n)$ – Each of the three main steps (finding the middle, reversing, merging) takes $O(n)$ time.

Space Complexity: $O(1)$ – The algorithm is done in-place with only a few pointers.

Solution

The solution involves three key steps:

– Use two pointers (slow and fast) to identify the middle of the list. – Reverse the second half of the list. This can be done iteratively by adjusting the next pointers. – Merge the two lists: One starting from the head and the other from the head of the reversed second half. Interleave the nodes by alternating between the two lists until complete. This approach uses in-place pointer manipulation, thus ensuring no extra space is used aside from a few temporary variables.

Stack

#150 Evaluate Reverse Polish Notation

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Math · Stack

Lists: Grind 169

Problem

Description

You are given an array of strings tokens that represents an arithmetic expression in a Reverse Polish Notation.

Evaluate the expression. Return an integer that represents the value of the expression.

Note that:

- The valid operators are '+', '-', '*', and '/'.
- Each operand may be an integer or another expression.
- The division between two integers always truncates toward zero.
- There will not be any division by zero.
- The input represents a valid arithmetic expression in a reverse polish notation.
- The answer and all the intermediate calculations can be represented in a 32-bit integer.

Example 1:

Input: tokens = ["2","1","+","3","*"]

Output: 9

Explanation: $((2 + 1) * 3) = 9$

Example 2:

Input: tokens = ["4","13","5","/","+"]

Output: 6

Explanation: $(4 + (13 / 5)) = 6$

Example 3:

Input: tokens = ["10","6","9","3","+","-11","*","/","*","17","+","5","+"]

Output: 22

Explanation: $((10 * (6 / ((9 + 3) * -11))) + 17) + 5$
 $= ((10 * (6 / (12 * -11))) + 17) + 5$
 $= ((10 * (6 / -132)) + 17) + 5$
 $= ((10 * 0) + 17) + 5$
 $= (0 + 17) + 5$
 $= 17 + 5$

Constraints:

- $1 \leq \text{tokens.length} \leq 104$
- $\text{tokens}[i]$ is either an operator: "+", "-", "*", or "/", or an integer in the range [-200, 200].

Community Solutions (Python)

- [WalkCC](#)

class Solution:

def evalRPN(self, tokens: list[str]) -> int:

stack = []

op = {

 '+': lambda a, b: a + b,

 '-': lambda a, b: a - b,

 '*': lambda a, b: a * b,

 '/': lambda a, b: int(a / b),

}

for token in tokens:

 if token in op:

 b = stack.pop()

 a = stack.pop()

 stack.append(op[token](a, b))

 else:

 stack.append(int(token))

return stack.pop()

- [LeetDoocs](#)

```
import operator

class Solution:
    def evalRPN(self, tokens: List[str]) -> int:
        opt = {
            "+": operator.add,
            "-": operator.sub,
            "*": operator.mul,
            "/": operator.truediv,
        }
        s = []
        for token in tokens:
            if token in opt:
                s.append(int(opt[token](s.pop(-2),
s.pop(-1))))
            else:
                s.append(int(token))
        return s[0]
```

```
class Solution:
    def evalRPN(self, tokens: List[str]) -> int:
        nums = []
        for t in tokens:
            if len(t) > 1 or t.isdigit():
                nums.append(int(t))
            else:
                if t == "+":
                    nums[-2] += nums[-1]
                elif t == "-":
                    nums[-2] -= nums[-1]
                elif t == "*":
                    nums[-2] *= nums[-1]
                else:
                    nums[-2] = int(nums[-2] /
nums[-1])
                nums.pop()
        return nums[0]
```

- SimplyLeet

```
def evalRPN(tokens):
    # Initialize a stack to hold numbers
    stack = []
    # Loop through each token in the input list
    for token in tokens:
        if token in {"+", "-", "*", "/"}:
            # Pop the two topmost values; note the
            # order is important
            b = stack.pop()
            a = stack.pop()
            # Apply the operation based on the
            token
```

```
        if token == "+":
            result = a + b
        elif token == "-":
            result = a - b
        elif token == "*":
            result = a * b
        else:
            # Use int division to truncate
            # towards zero
            result = int(a / b)
            # Push the result back onto the stack
```

```
            stack.append(result)
        else:
            # Token is a number; push its integer
            value onto the stack
            stack.append(int(token))
            # The final result is the only element left in
            the stack
            return stack.pop()

# Example usage:
# tokens = ["2", "1", "+", "3", "*"]
# print(evalRPN(tokens)) # Output: 9
```

• LC.ca

```
class Solution:
    def evalRPN(self, tokens: List[str]) -> int:
        nums = []
        for t in tokens:
            if len(t) > 1 or t.isdigit():
                nums.append(int(t))
            else:
                if t == "+":
                    nums[-2] += nums[-1]
                elif t == "-":
```

```

        nums[-2] -= nums[-1]
    elif t == "*":
        nums[-2] *= nums[-1]
    else:
        nums[-2] = int(nums[-2] /
nums[-1])
        nums.pop()
    return nums[0]

```

◆ Quick Review

Key Insights

- Use a stack to process the tokens.
- Push numbers onto the stack.
- When an operator is encountered, pop the top two numbers, perform the operation, and push the result back.
- Handle division carefully to ensure it truncates toward zero.
- The final result remains as the only value in the stack.

Space & Time

Time Complexity: $O(n)$, where n is the number of tokens (each token is processed once).

Space Complexity: $O(n)$, in the worst-case scenario where all tokens are numbers and are stored in the stack.

Solution

We use a stack data structure to evaluate the RPN expression. Iterate through each token:

– If the token is a number, convert and push it onto the stack. – If the token is an operator, pop the top two numbers (be careful about order), perform the corresponding arithmetic operation, and push the result back onto the stack. – At the end of processing, the stack contains the result of the expression. Special attention is needed for division to ensure it truncates toward zero (using language-specific functions when necessary).

Stack

#155 Min Stack

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Stack · Design

Lists: Grind 169

Problem

Description

Design a stack that supports push, pop, top, and retrieving the minimum element in constant time.

Implement the MinStack class:

- `MinStack()` initializes the stack object.
- `void push(int val)` pushes the element `val` onto the stack.
- `void pop()` removes the element on the top of the stack.
- `int top()` gets the top element of the stack.
- `int getMin()` retrieves the minimum element in the stack.

You must implement a solution with $O(1)$ time complexity for each function.

Example 1:

Input

```
["MinStack","push","push","push","getMin","pop","top","getMin"]
[[],[-2],[0],[-3],[],[],[],[[]]]
```

Output

```
[null,null,null,null,-3,null,0,-2]
```

Explanation

Constraints:

- $-231 \leq \text{val} \leq 231 - 1$
- Methods `pop`, `top` and `getMin` operations will always be called on non-empty stacks.
- At most $3 * 10^4$ calls will be made to `push`, `pop`, `top`, and `getMin`.

Community Solutions (Python)

- WalkCC

```
class MinStack:
    def __init__(self):
        self.stack = []

    def push(self, x: int) -> None:
        mn = x if not self.stack else
min(self.stack[-1][1], x)
        self.stack.append([x, mn])

    def pop(self) -> None:
        self.stack.pop()
```

```
def top(self) -> int:
    return self.stack[-1][0]

def getMin(self) -> int:
    return self.stack[-1][1]
```

• LeetCode

```
class MinStack:
    def __init__(self):
        self.stk1 = []
        self.stk2 = [inf]

    def push(self, val: int) -> None:
        self.stk1.append(val)
        self.stk2.append(min(val, self.stk2[-1]))

    def pop(self) -> None:
```

Round 5 · Mid · Medium

p.1927

```
self.stk1.pop()
self.stk2.pop()
```

```
def top(self) -> int:
    return self.stk1[-1]
```

```
def getMin(self) -> int:
    return self.stk2[-1]
```

Your MinStack object will be instantiated and called as such:

```
# obj = MinStack()
# obj.push(val)
# obj.pop()
# param_3 = obj.top()
# param_4 = obj.getMin()
```

• SimplyLeet

Python implementation of MinStack with detailed comments

```
class MinStack:
    def __init__(self):
        # Main stack to store all values
        self.stack = []
        # Auxiliary stack to store minimum values
        self.min_stack = []

    def push(self, val: int) -> None:
```

Round 5 · Mid · Medium

p.1928


```

        # Push value onto main stack
        self.stack.append(val)
        # If min_stack is empty or the new value is
        # less than or equal to current minimum, push it onto
        # min_stack
        if not self.min_stack or val <=
self.min_stack[-1]:
            self.min_stack.append(val)

    def pop(self) -> None:
        # Pop the top value from main stack
        popped_value = self.stack.pop()
        # If the popped value is the current
        # minimum, pop it from min_stack as well

        if popped_value == self.min_stack[-1]:
            self.min_stack.pop()

    def top(self) -> int:
        # Return the top value from main stack
        # without removing it
        return self.stack[-1]

    def getMin(self) -> int:
        # Return the minimum value from min_stack
        return self.min_stack[-1]

```

```

# Example usage:
# minStack = MinStack()
# minStack.push(-2)
# minStack.push(0)
# minStack.push(-3)
# print(minStack.getMin()) # -> -3
# minStack.pop()
# print(minStack.top()) # -> 0
# print(minStack.getMin()) # -> -2

```

• LC.ca

```

class MinStack:
    def __init__(self):
        self.sk = []
        self.minsk = [inf] # trick, avoid later
        # empty-check for min-stack

    def push(self, x: int) -> None:
        self.sk.append(x)
        self.minsk.append(min(x, self.minsk[-1]))

    def pop(self) -> None:

```

```

if not self.sk:
    return

    self.sk.pop()
    self.minsk.pop()

def top(self) -> int:
    return self.sk[-1]

def getMin(self) -> int:
    return self.minsk[-1]

```

```

# Your MinStack object will be instantiated and
# called as such:
# obj = MinStack()
# obj.push(x)
# obj.pop()
# param_3 = obj.top()
# param_4 = obj.getMin()

```

```
#####
```

```

class MinStack:

def __init__(self):
    self.sk = []
    self.minsk = []

def push(self, val: int) -> None:
    self.sk.append(val)

    # empty check
    self.minsk.append(val if not self.minsk else
min(val, self.minsk[-1]))

def pop(self) -> None:
    if not self.sk:
        return

    self.sk.pop()
    self.minsk.pop()

```

```
def top(self) -> int:
    return self.sk[-1]

def getMin(self) -> int:
    return self.minsk[-1]

# optimize above, using only 1 stack, with stack
# element as tuple: (val, its associated min)
class MinStack:

    def __init__(self):

        self._stack = []

    def push(self, x: int) -> None:
        cur_min = self.getMin()
        if x < cur_min:
            cur_min = x
        self._stack.append((x, cur_min))

    def pop(self) -> None:
        self._stack.pop()
```

```
def top(self) -> int:
    if not self._stack:
        return None
    else:
        return self._stack[-1][0]

def getMin(self) -> int:
    if not self._stack:
        return float('inf')
    else:
        return self._stack[-1][1]
```

◆ Quick Review

Key Insights

- Use an auxiliary data structure to track the current minimum element.
- A common approach is to use two stacks: one regular stack for all elements and a second stack to keep track of the minimum values.
- When pushing a new element, push it onto the main stack and also update the min stack if the element is less than or equal to the current minimum.
- When popping, if the popped element is equal to the top of the min stack, pop from the min stack as well.
- This design guarantees that all operations (push, pop, top, getMin) run in constant time, $O(1)$.

Space & Time

Time Complexity: $O(1)$ per operation (push, pop, top, getMin)

Space Complexity: $O(n)$, where n is the number of elements in the stack, since an additional stack is used to keep track of minima.

Solution

The solution involves using two stacks:

– Main Stack: Stores all the elements. – Min Stack: Stores the minimum elements encountered so far.

When an element is pushed, we compare it with the top of the min stack. If it is smaller than or equal to the current minimum, we also push it onto the min stack. For the pop operation, if the element being popped is the same as the top of the min stack, we pop from the min stack as well. This approach ensures that the getMin operation always returns the top element of the min stack, which is the current minimum in constant time.

Stack

#173 Binary Search Tree Iterator

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Stack · Tree · Design · Binary Search Tree · Iterator

Lists: Denny Zhang

Problem

Description


Implement the BSTIterator class that represents an iterator over the in-order traversal of a binary search tree (BST):

- **BSTIterator(TreeNode root)** Initializes an object of the BSTIterator class. The root of the BST is given as part of the constructor. The pointer should be initialized to a non-existent number smaller than any element in the BST.
- **boolean hasNext()** Returns true if there exists a number in the traversal to the right of the pointer, otherwise returns false.
- **int next()** Moves the pointer to the right, then returns the number at the pointer.

Notice that by initializing the pointer to a non-existent smallest number, the first call to next() will return the smallest element in the BST.

You may assume that `next()` calls will always be valid. That is, there will be at least a next number in the in-order traversal when `next()` is called.

Example 1:



Input

```
["BSTIterator", "next", "next", "hasNext", "next",
"hasNext", "next", "hasNext", "next", "hasNext"]
[[[7, 3, 15, null, null, 9, 20]], [], [], [], [], [],
[], [], [], []]
```

Output

```
[null, 3, 7, true, 9, true, 15, true, 20, false]
```

Explanation

```
BSTIterator bSTIterator = new BSTIterator([7, 3, 15,
null, null, 9, 20]);
```

Constraints:

- The number of nodes in the tree is in the range [1, 105].
- $0 \leq \text{Node.val} \leq 106$
- At most 105 calls will be made to `hasNext`, and `next`.

Follow up:

- Could you implement `next()` and `hasNext()` to run in average $O(1)$ time and use $O(h)$ memory, where h is the height of the tree?

Community Solutions (Python)

- WalkCC

```
class BSTIterator:
    def __init__(self, root: TreeNode | None):
        self.i = 0
        self.vals = []
        self._inorder(root)

    def next(self) -> int:
        self.i += 1
        return self.vals[self.i - 1]
```

```
    def hasNext(self) -> bool:
        return self.i < len(self.vals)

    def _inorder(self, root: TreeNode | None) -> None:
        if not root:
            return
        self._inorder(root.left)
        self.vals.append(root.val)
        self._inorder(root.right)
```

```
class BSTIterator:
    def __init__(self, root: TreeNode | None):
        self.stack = []
        self._pushLeftsUntilNull(root)

    def next(self) -> int:
        root = self.stack.pop()
        self._pushLeftsUntilNull(root.right)
        return root.val
```

```
    def hasNext(self) -> bool:
        return self.stack

    def _pushLeftsUntilNull(self, root: TreeNode | None) -> None:
        while root:
            self.stack.append(root)
            root = root.left
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class BSTIterator:
    def __init__(self, root: TreeNode):
        def inorder(root):
            if root:
```

```

        inorder(root.left)
        self.vals.append(root.val)
        inorder(root.right)

    self.cur = 0
    self.vals = []
    inorder(root)

    def next(self) -> int:
        res = self.vals[self.cur]

```

```

        self.cur += 1
        return res

    def hasNext(self) -> bool:
        return self.cur < len(self.vals)

```

Your BSTIterator object will be instantiated and called as such:

```

# obj = BSTIterator(root)
# param_1 = obj.next()

```

```

# param_2 = obj.hasNext()

```

• SimplyLeet

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None, right=None):
#         self.val = val
#         self.left = left
#         self.right = right

class BSTIterator:
    def __init__(self, root):
        # Initialize an empty stack to simulate
        # in-order traversal

```

```

        self.stack = []
        # Push all left nodes starting from the
        # root
        self._push_left_branch(root)

    def _push_left_branch(self, node):
        # Push node and all its left children to
        # the stack
        while node:
            self.stack.append(node)
            node = node.left

```

```

def next(self):
    # Pop the top node which represents the
    next smallest element
    node = self.stack.pop()
    # If the node has a right child, push its
    left branch onto the stack
    if node.right:
        self._push_left_branch(node.right)
    # Return the node's value
    return node.val

```

```

def hasNext(self):

```

```

    # The iterator has a next element if the
    stack is not empty
    return len(self.stack) > 0

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class BSTIterator:
    def __init__(self, root: TreeNode):
        self.stack = []
        self._leftmost_inorder(root)

```

```

def _leftmost_inorder(self, node):
    while node:
        self.stack.append(node)
        node = node.left

```

```

def next(self) -> int:
    cur = self.stack.pop()
    node = cur.right
    self._leftmost_inorder(node)

```

```

    return cur.val

```

```

def hasNext(self) -> bool:
    return len(self.stack) > 0

```

Your BSTIterator object will be instantiated and called as such:

```

# obj = BSTIterator(root)
# param_1 = obj.next()
# param_2 = obj.hasNext()

```



```
#####
```

```
# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
# right=None):
# self.val = val
# self.left = left
# self.right = right
class BSTIterator:
    def __init__(self, root: TreeNode):
```

```
        self.stack = []
        while root:
            self.stack.append(root)
            root = root.left
```

```
    def next(self) -> int:
        cur = self.stack.pop()
        node = cur.right
        while node: # duplicated while block
            self.stack.append(node)
```

```
        node = node.left
        return cur.val
```

```
    def hasNext(self) -> bool:
        return len(self.stack) > 0
```

Your BSTIterator object will be instantiated and called as such:

```
# obj = BSTIterator(root)
# param_1 = obj.next()
```

```
# param_2 = obj.hasNext()
```

◆ Quick Review

Key Insights

- Use an iterative in-order traversal strategy with a stack.
- In-order traversal for BST yields sorted order.
- Pre-load the stack with all left descendants from the root.
- After returning a node, process its right subtree by pushing its left branch.
- This method maintains an average time complexity of $O(1)$ per operation and uses $O(h)$ memory, where h is the tree height.

Space & Time

Time Complexity: $O(1)$ on average per operation (amortized)

Space Complexity: $O(h)$, where h is the height of the tree

Solution

We use a stack to simulate an in-order traversal. The process involves initially pushing all left-child nodes from the root into the stack. For each call to `next()`, we pop the top element (the current smallest node) and then push all the left descendents of the node's right child. This ensures that the next smallest element is always at the top of the stack. The `hasNext()` function then simply checks if the stack is non-empty.

Stack

#227 Basic Calculator II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · String · Stack

Lists: Grind 169

Problem

Description

Given a string s which represents an expression, evaluate this expression and return its value.

The integer division should truncate toward zero.

You may assume that the given expression is always valid. All intermediate results will be in the range of $[-2^{31}, 2^{31} - 1]$.

Note: You are not allowed to use any built-in function which evaluates strings as mathematical expressions, such as `eval()`.

Example 1:

Input: $s = "3+2*2"$

Output: 7

Example 2:

Input: $s = " 3/2 "$

Output: 1

Example 3:

Input: s = " 3+5 / 2 "

Output: 5

Constraints:

- $1 \leq s.length \leq 3 \times 10^5$
- s consists of integers and operators ('+', '-', '*', '/') separated by some number of spaces.
- s represents a valid expression.
- All the integers in the expression are non-negative integers in the range $[0, 2^{31} - 1]$.
- The answer is guaranteed to fit in a 32-bit integer.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def calculate(self, s: str) -> int:
        ans = 0
        prevNum = 0
        currNum = 0
        op = '+'

        for i, c in enumerate(s):
            if c.isdigit():
                currNum = currNum * 10 + int(c)
```

```
        if not c.isdigit() and c != ' ' or i == len(s) - 1:
            if op == '+' or op == '-':
                ans += prevNum
                prevNum = currNum if op == '+' else -currNum
            elif op == '*':
                prevNum = prevNum * currNum
            elif op == '/':
                if prevNum < 0:
                    prevNum = math.ceil(prevNum / currNum)
                else:
```

```
                prevNum = prevNum // currNum
            op = c
            currNum = 0
```

```
        return ans + prevNum
```

• LeetCode

```
class Solution:
    def calculate(self, s: str) -> int:
        v, n = 0, len(s)
        sign = '+'
        stk = []
        for i, c in enumerate(s):
            if c.isdigit():
                v = v * 10 + int(c)
            if i == n - 1 or c in '+-* /':
                match sign:
```

```

        case '+':
            stk.append(v)
        case '-':
            stk.append(-v)
        case '*':
            stk.append(stk.pop() * v)
        case '/':
            stk.append(int(stk.pop() /
v))

        sign = c
        v = 0

    return sum(stk)

```

• SimplyLeet

Python solution for Basic Calculator II

```

def calculate(s: str) -> int:
    # Initialize the stack, current number, and the
    last seen operator.
    stack = []
    current_number = 0
    previous_operator = '+'

    # Loop over each character in the string along
    with one extra iteration to handle the end.
    for i, char in enumerate(s):

```

```

        if char.isdigit():
            # Build the current number by shifting
            previous digits.
            current_number = current_number * 10 +
int(char)
            # If the character is an operator (or we're
            at the end), process the previous operator.
            if char in "+-*/" or i == len(s) - 1:
                if previous_operator == '+':
                    stack.append(current_number)
                elif previous_operator == '-':
                    stack.append(-current_number)
                elif previous_operator == '*':

```

```

                    last_number = stack.pop()
                    stack.append(last_number *
current_number)
                elif previous_operator == '/':
                    last_number = stack.pop()
                    # Python division truncates toward
                    negative infinity by default, so we use int() to
                    truncate toward zero.
                    stack.append(int(last_number /
current_number))
                # Reset for the next number and update
                the operator.
                previous_operator = char
                current_number = 0
            # Sum up all numbers in the stack which gives
            the final result.

```

```
return sum(stack)
```

Example usage:

```
print(calculate("3+2*2")) # Output: 7
print(calculate(" 3/2 ")) # Output: 1
print(calculate(" 3+5 / 2 ")) # Output: 5
```

• LC.ca

```
class Solution:
```

```
    def calculate(self, s: str) -> int:
        v, n = 0, len(s)
        sign = '+'
        stk = []
        for i, c in enumerate(s):
            if c.isdigit():
                v = v * 10 + int(c)
            if i == n - 1 or c in '+-*/*':
                if sign == '+':
```

```
                    stk.append(v)
                    # for "10-2*5": when '-'
                    # encountered, var 'sign' is still '+'
                    # so '10' will be pushed to stk
                    # before setting sign to '-'
                    elif sign == '-':
                        stk.append(-v)
                    elif sign == '*':
                        stk.append(stk.pop() * v)
                    elif sign == '/':
                        stk.append(int(stk.pop() / v))
                    else:
```

```
                                print("operator not
                                supported")
```

```
                                sign = c # reset inside 'if'
                                v = 0
                                return sum(stk)
```

◆ Quick Review

Key Insights

- The expression can be processed in one pass while using a stack to handle operator precedence.
- When encountering '+' or '-' operators, the current number is directly added or subtracted (as a positive or negative value).
- For '*' and '/' operators, compute immediately with the last number from the stack and update it.

- Handle spaces and multi-digit numbers by building the number while iterating over the string.
- Division truncation toward zero must be carefully applied, which aligns with typical integer division in most languages.

Space & Time

Time Complexity: $O(n)$ where n is the length of the expression, since the algorithm processes each character once.

Space Complexity: $O(n)$ in the worst-case scenario when all numbers are stored in the stack.

Solution

We use a stack-based approach. Iterate over each character of the string while maintaining a variable for the current number and a previous operator (initialized to '+'). When an operator or the end of the string is reached, perform an action based on the previous operator:

– For '+' add the current number to the stack. – For '-' add the negative of the current number. – For '*' and '/' pop the last number from the stack, compute the result by performing multiplication or division, and then push the result back to the stack. After processing the entire string, sum the elements in the stack to get the final result. This approach ensures that multiplication and division are handled before addition and subtraction.

Stack

#255 Verify Preorder Sequence in Binary Search Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Tree · BST · Stack · Monotonic Stack


Lists: Premium 98

Problem

Description

Given an array of unique integers preorder, return true if it is the correct preorder traversal sequence of a binary search tree.

Example 1:



Input: preorder = [5,2,1,3,6]
Output: true

Example 2:

Input: preorder = [5,2,6,1,3]
Output: false

Constraints:

- $1 \leq \text{preorder.length} \leq 104$
- $1 \leq \text{preorder}[i] \leq 104$
- All the elements of preorder are unique.

Follow up: Could you do it using only constant space complexity?

Community Solutions (Python)

• WalkCC

```

class Solution:
    def verifyPreorder(self, preorder: list[int]) -> bool:
        i = 0

        def dfs(min: int, max: int) -> None:
            nonlocal i
            if i == len(preorder):
                return
            if preorder[i] < min or preorder[i] > max:
                return

            val = preorder[i]
            i += 1
            dfs(min, val)
            dfs(val, max)

        dfs(-math.inf, math.inf)
        return i == len(preorder)
  
```

```
class Solution:
    def verifyPreorder(self, preorder: list[int]) ->
list[int]:
    low = -math.inf
    stack = []
```

```
    for p in preorder:
        if p < low:
            return False
        while stack and stack[-1] < p:
            low = stack.pop()
```

```
        stack.append(p)
```

```
    return True
```

```
class Solution:
    def verifyPreorder(self, preorder: list[int]) ->
bool:
    low = -math.inf
    i = -1
```

```
    for p in preorder:
        if p < low:
            return False
        while i >= 0 and preorder[i] < p:
            low = preorder[i]
```

```
        i -= 1
        i += 1
        preorder[i] = p

    return True
```

• LeetCode

```
class Solution:
    def verifyPreorder(self, preorder: List[int])
-> bool:
        stk = []
        last = -inf
        for x in preorder:
            if x < last:
                return False
            while stk and stk[-1] < x:
                last = stk.pop()
            stk.append(x)
```

```
    return True
```

• SimplyLeet


```
# Python solution using a monotonic stack with
detailed comments
def verifyPreorder(preorder):
    lower_bound = float('-inf') # Set initial lower
    bound to negative infinity
    stack = [] # Use a stack to track the BST
    construction
    for value in preorder:
        # If the current value is less than
        lower_bound, the BST property is violated
        if value < lower_bound:
            return False
        # If current value is greater than the top
        of stack, pop elements and update lower_bound
        while stack and value > stack[-1]:

            lower_bound = stack.pop() # Update
            lower_bound as we move to the right subtree
            # Push the current value onto the stack
            stack.append(value)
    return True

# Example usage:
# print(verifyPreorder([5,2,1,3,6])) # Expected
output: True
```

• LC.ca

Round 5 · Mid · Medium

p.1961

```
class Solution:
    def verifyPreorder(self, preorder: List[int])
    -> bool:
        stk = []
        last = -inf
        for x in preorder:
            if x < last:
                return False
            while stk and stk[-1] < x:
                last = stk.pop()
            stk.append(x)

        return True

class Solution: # o(N) but modifying input list
    def verifyPreorder(self, preorder: List[int])
    -> bool:
        low = -math.inf
        i = -1

        for p in preorder:
            if p < low:
```

Round 5 · Mid · Medium

p.1962

```

        return False
    while i >= 0 and preorder[i] < p:
        low = preorder[i]
        i -= 1
    i += 1
    preorder[i] = p

    return True

```

```
#####
```

```
class Solution(object):
```

```

    # O(n) sapce complexity, **stack**
    def verifyPreorder(self, preorder):
        """
        :type preorder: List[int]
        :rtype: bool
        """
        if len(preorder) <= 1:

```

```

        return True
    stack, lastElem = [preorder[0]], None
    for i in range(1, len(preorder)):
        if lastElem > preorder[i]:
            return False
        while len(stack) > 0 and preorder[i] >
stack[-1]:
            lastElem = stack.pop()
        stack.append(preorder[i])

    return True

```

◆ Quick Review

Key Insights

- A BST has the property that every node's left subtree contains values less than the node, and every right subtree contains values greater.
- Preorder traversal processes nodes as: root, left subtree, then right subtree.
- Using a monotonic stack, we can simulate constructing the BST while tracking a lower bound for allowable node values.
- When moving to a right child, update the lower bound to ensure subsequent nodes are valid.

Space & Time

Time Complexity: $O(n)$ – Each element is processed once.

Space Complexity: $O(n)$ – In the worst-case an explicit stack is used; can be optimized to $O(1)$ by reusing the input array.

Solution

We use a monotonic stack to simulate the BST construction from the preorder sequence. The process is:

- Initialize a variable `lower_bound` to $-\infty$.
- Iterate through each value in the preorder array: If a value is less than `lower_bound`, it violates BST properties and the sequence is invalid. While the stack is not empty and the current value is greater than the element at the top of the stack, pop from the stack and update `lower_bound` to the popped value. This simulates moving to the right subtree. Push the current value onto the stack.
- If the sequence is processed without violations, return true.

This approach efficiently ensures that each value adheres to BST requirements during preorder traversal.

Stack

#394 Decode String

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Stack · Recursion

Lists: Grind 169

Problem

Description

Given an encoded string, return its decoded string.

The encoding rule is: `k[encoded_string]`, where the `encoded_string` inside the square brackets is being repeated exactly `k` times. Note that `k` is guaranteed to be a positive integer.

You may assume that the input string is always valid; there are no extra white spaces, square brackets are well-formed, etc. Furthermore, you may assume that the original data does not contain any digits and that digits are only for those repeat numbers, `k`. For example, there will not be input like `3a` or `2[4]`.

The test cases are generated so that the length of the output will never exceed 105.

Example 1:

Input: `s = "3[a]2[bc]"`
Output: `"aaabcbc"`

Example 2:

Input: `s = "3[a2[c]]"`
 Output: `"accaccacc"`

Example 3:

Input: `s = "2[abc]3[cd]ef"`
 Output: `"abcbccdcdef"`

Constraints:

- $1 \leq s.length \leq 30$
- `s` consists of lowercase English letters, digits, and square brackets '['].
- `s` is guaranteed to be a valid input.
- All the integers in `s` are in the range [1, 300].

Community Solutions (Python)

• WalkCC

```
class Solution:
    def decodeString(self, s: str) -> str:
        stack = [] # (prevStr, repeatCount)
        currStr = ''
        currNum = 0

        for c in s:
            if c.isdigit():
                currNum = currNum * 10 + int(c)
            else:
```

```
        if c == '[':
            stack.append((currStr, currNum))
            currStr = ''
            currNum = 0
        elif c == ']':
            prevStr, num = stack.pop()
            currStr = prevStr + num * currStr
        else:
            currStr += c
```

```
        return currStr
```

```
class Solution:
    def decodeString(self, s: str) -> str:
        ans = ''

        while self.i < len(s) and s[self.i] != ']':
            if s[self.i].isdigit():
                k = 0
                while self.i < len(s) and
s[self.i].isdigit():
                    k = k * 10 + int(s[self.i])
                self.i += 1
```

```

        self.i += 1 # '['
        decodedString = self.decodeString(s)
        self.i += 1 # ']'
        ans += k * decodedString
    else:
        ans += s[self.i]
        self.i += 1

    return ans

```

```
i = 0
```

• LeetCode

```

class Solution:
    def decodeString(self, s: str) -> str:
        s1, s2 = [], []
        num, res = 0, ''
        for c in s:
            if c.isdigit():
                num = num * 10 + int(c)
            elif c == '[':
                s1.append(num)
                s2.append(res)
                num, res = 0, ''
            elif c == ']':
                res = s2.pop() + res * s1.pop()
            else:
                res += c
        return res

```

Round 5 · Mid · Medium

p.1969

• SimplyLeet

```

def decodeString(s: str) -> str:
    # index pointer as a list so it can be mutable
    # within recursion
    def helper(i: int) -> (str, int):
        decoded = ""
        while i < len(s):
            if s[i].isdigit():
                # build the number (it might be
                # more than one digit)
                k = 0
                while i < len(s) and
                s[i].isdigit():
                    k = k * 10 + int(s[i])

                i += 1
                # skip the '[' character
                i += 1 # skip '['
                # recursively solve the substring
                # inside the brackets
                sub, i = helper(i)
                # append the substring repeated k
                # times
                decoded += sub * k
            elif s[i] == ']':
                # returning decoded string and new
                # index after ']'
                return decoded, i + 1

```

Round 5 · Mid · Medium

p.1970

```

        else:
            # accumulate alphabetic characters
            directly
            decoded += s[i]
            i += 1
        return decoded, i

    result, _ = helper(0)
    return result

```

Example usage

```

#print(decodeString("3[a2[c]]")) # Output:
"accaccacc"

```

• LC.ca

```

class Solution:
    def decodeString(self, s: str) -> str:
        num_stk, str_stk = [], []
        num, res = 0, ''
        for c in s:
            if c.isdigit():
                num = num * 10 + int(c)
            elif c == '[':
                num_stk.append(num)
                str_stk.append(res) # initial, res
                # being pushed is empty str ''
            else:

```

```

                num, res = 0, ''
            elif c == ']':
                res = str_stk.pop() + res *
                num_stk.pop()
            else:
                res += c
        return res

```

◆ Quick Review

Key Insights

- Use a stack or recursion to manage nested encoded patterns.
- Process the string character by character, accumulating digits for the repeat count.
- When encountering a '[', start a new context (recursively or via a new stack frame) for the encoded substring.
- When encountering a ']', finish the current context and repeat the accumulated substring as needed.
- Careful handling of multi-digit numbers is essential.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(n)$, for the recursion stack or the auxiliary stack used.

Solution

We can solve this problem using a recursive approach. When we see a digit, we determine the full repeat count.

Upon encountering a '[', we recursively read the encoded substring until the corresponding ']'. Each recursive call returns the decoded substring until it hits a ']', which is then repeated based on the previously computed count. Alternatively, an iterative solution using a stack is also feasible. As we parse the string, push the current string and number onto the stack when encountering '[', then pop and combine when ']' is reached. This approach correctly handles nested patterns.

Stack

#439 Ternary Expression Parser

MED

[LeetDoocs](#) · [SimplyLeet](#) · [WalkCC](#) · [LeetCode](#)[String](#) · [Stack](#) · [Recursion](#)

Lists: Premium 98

Problem

Description

Given a string expression representing arbitrarily nested ternary expressions, evaluate the expression, and return the result of it.

You can always assume that the given expression is valid and only contains digits, '?', ':', 'T', and 'F' where 'T' is true and 'F' is false. All the numbers in the expression are one-digit numbers (i.e., in the range [0, 9]).

The conditional expressions group right-to-left (as usual in most languages), and the result of the expression will always evaluate to either a digit, 'T' or 'F'.

Example 1:

Input: expression = "T?2:3"**Output:** "2"**Explanation:** If true, then result is 2; otherwise result is 3.

Example 2:

Input: expression = "F?1:T?4:5"

Output: "4"

Explanation: The conditional expressions group right-to-left. Using parenthesis, it is read/evaluated as:

"(F ? 1 : (T ? 4 : 5))" --> "(F ? 1 : 4)" --> "4"

or "(F ? 1 : (T ? 4 : 5))" --> "(T ? 4 : 5)" --> "4"

Example 3:

Input: expression = "T?T?F:5:3"

Output: "F"

Explanation: The conditional expressions group right-to-left. Using parenthesis, it is read/evaluated as:

"(T ? (T ? F : 5) : 3)" --> "(T ? F : 3)" --> "F"

"(T ? (T ? F : 5) : 3)" --> "(T ? F : 5)" --> "F"

Constraints:

- $5 \leq \text{expression.length} \leq 104$
- expression consists of digits, 'T', 'F', '?', and ':'.
- It is guaranteed that expression is a valid ternary expression and that each number is a one-digit number.

Community Solutions (Python)

- WalkCC

```
class Solution:
```

```
    def parseTernary(self, expression: str) -> str:
        c = expression[self.i]
```

```
        if self.i + 1 == len(expression) or
           expression[self.i + 1] == ':':
            self.i += 2
            return str(c)
```

```
        self.i += 2
        first = self.parseTernary(expression)
```

```
        second = self.parseTernary(expression)
```

```
        return first if c == 'T' else second
```

```
        i = 0
```

• LeetCode

```
class Solution:
```

```
    def parseTernary(self, expression: str) ->
str:
```

```
        stk = []
        cond = False
        for c in expression[::-1]:
            if c == ':':
                continue
            if c == '?':
                cond = True
            else:
```



```

        if cond:
            if c == 'T':
                x = stk.pop()
                stk.pop()
                stk.append(x)
            else:
                stk.pop()
            cond = False
        else:
            stk.append(c)

```

```

return stk[0]

```

• SimplyLeet

```

def parseTernary(expression: str) -> str:
    # Initialize stack to hold characters/results
    # during evaluation
    stack = []
    # Iterate backwards over the expression
    i = len(expression) - 1
    while i >= 0:
        if expression[i] == '?':
            # Pop the true and false parts from the
            # stack
            true_expr = stack.pop()
            false_expr = stack.pop()

```

```

        i -= 1 # Move to the condition
        character
        # Choose based on the condition
        if expression[i] == 'T':
            stack.append(true_expr)
        else:
            stack.append(false_expr)
        elif expression[i] == ':':
            # Skip the colon as it is only a
            # separator
            pass
        else:

```

```

        # Push digits, 'T', and 'F' onto stack
        as is
        stack.append(expression[i])
        i -= 1 # Move to the next character
        return stack[-1]

# Example usage:
print(parseTernary("T?2:3")) # Expected output "2"

```

• LC.ca

```

class Solution:
    def parseTernary(self, expression: str) -> str:
        stk = []
        cond = False
        for c in expression[::-1]:
            if c == ':':
                continue
            if c == '?':
                cond = True
            else:
                if cond:
                    if c == 'T':
                        x = stk.pop()
                        stk.pop()
                        stk.append(x)
                    else:
                        stk.pop()
                        cond = False
                else:
                    stk.append(c)

        return stk[0]

```

◆ Quick Review

Key Insights

- The ternary operator groups right-to-left. This property can be exploited by processing the string from the end to

the beginning.

- Using a stack allows us to "collapse" nested ternary expressions as soon as their three components (condition, true branch, false branch) are available.
- An iterative approach is efficient and avoids building an explicit syntax tree for the expression.
- The condition is not stored on the stack; instead, it is encountered immediately preceding the '?' operator when traversing backwards.

Space & Time

Time Complexity: $O(n)$, where n is the length of the expression, as we process each character exactly once.

Space Complexity: $O(n)$, due to the use of a stack for storing characters during evaluation.

Solution

The solution uses an iterative approach with a stack. We traverse the expression string from right to left. For every character:

- If it is a digit, 'T', or 'F' (and not a ':' or part of a pending operator), we push it onto the stack.
- When we encounter a '?' character, we know that the next available items on the stack represent the true and false results of the ternary expression. We then skip the ':' separator that comes between them.
- We then use the condition (the character immediately preceding the '?') to choose which result to push back on the stack. This method “collapses” each ternary expression into a single result, and by the time we finish traversing the string, the stack contains the final

evaluated result.

Round 5 · Mid · Medium

p.1981

Stack

#484 Find Permutation

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · Stack · Greedy

Lists: Premium 98

Problem

Description

A permutation perm of n integers of all the integers in the range $[1, n]$ can be represented as a string s of length $n - 1$ where:

- $s[i] == 'I'$ if $\text{perm}[i] < \text{perm}[i + 1]$, and
- $s[i] == 'D'$ if $\text{perm}[i] > \text{perm}[i + 1]$.

Given a string s , reconstruct the lexicographically smallest permutation perm and return it.

Example 1:

Input: $s = "I"$

Output: $[1,2]$

Explanation: $[1,2]$ is the only legal permutation that can be represented by s , where the number 1 and 2 construct an increasing relationship.

Example 2:

Round 5 · Mid · Medium

p.1982

Input: s = "DI"

Output: [2,1,3]

Explanation: Both [2,1,3] and [3,1,2] can be represented as "DI", but since we want to find the smallest lexicographical permutation, you should return [2,1,3]

Constraints:

- $1 \leq s.length \leq 105$
- s[i] is either 'I' or 'D'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findPermutation(self, s: str) -> list[int]:
        ans = []
        stack = []

        for i, c in enumerate(s):
            stack.append(i + 1)
            if c == 'I':
                while stack: # Consume all decreasings
                    ans.append(stack.pop())
```

```
stack.append(len(s) + 1)
```

```
while stack:
    ans.append(stack.pop())
```

```
return ans
```

```
class Solution:
    def findPermutation(self, s: str) -> list[int]:
        ans = [i for i in range(1, len(s) + 2)]

        # For each D* group (s[i..j]), reverse ans[i..j
        + 1].
        i = -1
        j = -1

        def getNextIndex(c: str, start: int) -> int:
            for i in range(start, len(s)):

                if s[i] == c:
                    return i
            return len(s)

        while True:
            i = getNextIndex('D', j + 1)
            if i == len(s):
                break
            j = getNextIndex('I', i + 1)
            ans[i:j + 1] = ans[i:j + 1][::-1]
```

```
return ans
```

- LeetCode

```
class Solution:
    def findPermutation(self, s: str) ->
    List[int]:
        n = len(s)
        ans = list(range(1, n + 2))
        i = 0
        while i < n:
            j = i
            while j < n and s[j] == 'D':
                j += 1
            ans[i : j + 1] = ans[i : j + 1][::-1]

            i = max(i + 1, j)
        return ans
```

- SimplyLeet

```
# Function to find the lexicographically smallest
permutation
def findPermutation(s):
    stack = [] # Use a stack to hold numbers during
a sequence of 'D's
    result = [] # Result list for the final
permutation
    # Iterate over 0 to len(s) (inclusive) to
process all characters plus one final number.
    for i in range(len(s) + 1):
        # Push i+1 since our numbers range from 1
to n
        stack.append(i + 1)
        # If we are at the end or the current
character is 'I', flush the stack.
        if i == len(s) or s[i] == 'I':

            while stack:
                result.append(stack.pop())

    return result

# Example usage:
print(findPermutation("DI"))
```

- LC.ca

```

class Solution:
    def findPermutation(self, s: str) ->
    List[int]:
        n = len(s)
        ans = list(range(1, n + 2))
        i = 0
        while i < n:
            j = i
            while j < n and s[j] == 'D':
                j += 1
            ans[i : j + 1] = ans[i : j + 1][::-1]

            i = max(i + 1, j)
        return ans

```

◆ Quick Review

Key Insights

- The answer is a permutation of numbers from 1 to n, where $n = \text{len}(s) + 1$.
- A common strategy is to use a stack to reverse the segments when a decrease ('D') is encountered.
- When an 'I' is encountered (or at the end), pop all numbers from the stack to output them in reverse order, ensuring the smallest lexicographical order.
- This greedy approach guarantees that as soon as a rising pattern is detected, we “flush” the pending decreases correctly.

Space & Time

Time Complexity: $O(n)$ – Each element is pushed and popped exactly once.

Space Complexity: $O(n)$ – For the stack and the output array.

Solution

We use a greedy approach with a stack. Iterate from 0 to n (where $n = \text{len}(s) + 1$):

- Push the current number ($i+1$) onto the stack.
- If the current character is 'I' or we have reached the end of the string, pop all elements from the stack and append them to the result. By doing so, segments corresponding to series of 'D's are reversed, while 'I's trigger the flush, ensuring that the final permutation is the lexicographically smallest.

Stack

#735 Asteroid Collision

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Stack · Simulation

Lists: Grind 169

Problem

Description

We are given an array `asteroids` of integers representing asteroids in a row. The indices of the asteroid in the array represent their relative position in space.

For each asteroid, the absolute value represents its size, and the sign represents its direction (positive meaning right, negative meaning left). Each asteroid moves at the same speed.

Find out the state of the asteroids after all collisions. If two asteroids meet, the smaller one will explode. If both are the same size, both will explode. Two asteroids moving in the same direction will never meet.

Example 1:

Input: `asteroids = [5,10,-5]`
 Output: `[5,10]`
 Explanation: The 10 and -5 collide resulting in 10. The 5 and 10 never collide.

Example 2:

Input: `asteroids = [8,-8]`
 Output: `[]`
 Explanation: The 8 and -8 collide exploding each other.

Example 3:

Input: `asteroids = [10,2,-5]`
 Output: `[10]`
 Explanation: The 2 and -5 collide resulting in -5. The 10 and -5 collide resulting in 10.

Example 4:

Input: `asteroids = [3,5,-6,2,-1,4]`
 Output: `[-6,2,4]`
 Explanation: The asteroid -6 makes the asteroid 3 and 5 explode, and then continues going left. On the other side, the asteroid 2 makes the asteroid -1 explode and then continues going right, without reaching asteroid 4.

Constraints:

- $2 \leq \text{asteroids.length} \leq 104$
- $-1000 \leq \text{asteroids}[i] \leq 1000$
- $\text{asteroids}[i] \neq 0$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def asteroidCollision(self, asteroids:
list[int]) -> list[int]:
    stack = []

    for a in asteroids:
        if a > 0:
            stack.append(a)
        else: # a < 0
            # Destroy the previous positive one(s).
            while stack and stack[-1] > 0 and stack[-1]
< -a:
```

```
        stack.pop()
        if not stack or stack[-1] < 0:
            stack.append(a)
        elif stack[-1] == -a:
            stack.pop() # Both asteroids explode.
        else: # stack[-1] > the current asteroid.
            pass # Destroy the current asteroid, so
do nothing.

    return stack
```

• LeetCode

```
class Solution:
    def asteroidCollision(self, asteroids:
List[int]) -> List[int]:
    stk = []
    for x in asteroids:
        if x > 0:
            stk.append(x)
        else:
            while stk and stk[-1] > 0 and
stk[-1] < -x:
                stk.pop()
            if stk and stk[-1] == -x:

                stk.pop()
            elif not stk or stk[-1] < 0:
                stk.append(x)

    return stk
```

• SimplyLeet


```
# Function to simulate asteroid collisions
def asteroidCollision(asteroids):
    # Initialize an empty list as a stack to hold
    # surviving asteroids
    stack = []
    # Iterate through each asteroid in the input
    # list
    for ast in asteroids:
        # Flag to check if the current asteroid has
        # exploded
        exploded = False
        # Check for potential collision while there
        # is an asteroid in stack and a collision condition
        # holds (right-moving on stack and left-moving
        # current)
        while stack and ast < 0 and stack[-1] > 0:
```

```
        # If the incoming left-moving asteroid
        # is larger than the right-moving one on stack, pop
        # the top and continue checking
        if abs(ast) > abs(stack[-1]):
            stack.pop()
            continue # Continue the loop to
            # check next potential collision
        # If they are equal, both explode: pop
        # the stack and mark explosion, break the loop
        elif abs(ast) == abs(stack[-1]):
            stack.pop()
            # Mark the incoming asteroid as
            # exploded and break out of the loop
            exploded = True
            break
```

```
        # If the asteroid did not explode during
        # collisions, push it onto the stack
        if not exploded:
            stack.append(ast)
        # Return the surviving asteroids
        return stack
```

```
# Example usage:
# print(asteroidCollision([5,10,-5])) # Output:
# [5,10]
```

• LC.ca

```

class Solution:
    def asteroidCollision(self, asteroids:
List[int]) -> List[int]:
        stk = []
        for x in asteroids:
            if x > 0:
                stk.append(x)
            else:
                while stk and stk[-1] > 0 and
stk[-1] < -x:
                    stk.pop()
                    if stk and stk[-1] == -x:
                        stk.pop()
                elif not stk or stk[-1] < 0:
                    stk.append(x)
        return stk

```

◆ Quick Review

Key Insights

- Use a stack to keep track of the surviving asteroids.
- Only a collision happens when a right-moving asteroid (positive) is on the stack and a left-moving asteroid (negative) is encountered.
- Compare the absolute values to decide if one or both asteroids explode.
- Continue processing collisions until no further collisions are possible with the current asteroid.

Space & Time

Time Complexity: $O(n)$ where n is the number of asteroids, since each asteroid is pushed and popped at most once.

Space Complexity: $O(n)$ in the worst case, when no collisions occur and all asteroids remain on the stack.

Solution

The solution uses a stack to simulate the asteroid collision process. For each asteroid in the input:

- If it is moving to the right or the stack is empty, simply add it to the stack.
- If it is moving to the left, check the asteroid on top of the stack: If the top asteroid is moving to the right, they might collide. Compare sizes: If the left-moving asteroid is larger (in absolute terms), pop the stack and continue checking. If both are equal, pop the stack and do not add the current asteroid. If the right-moving asteroid is larger, the current asteroid explodes.
- If no collision occurs for the current asteroid, push it onto the stack. Finally, the stack represents the surviving asteroids order.

Stack

#739 Daily Temperatures

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Stack · Monotonic Stack

Lists: Grind 169

Problem

Description

Given an array of integers `temperatures` represents the daily temperatures, return an array `answer` such that `answer[i]` is the number of days you have to wait after the `i`th day to get a warmer temperature. If there is no future day for which this is possible, keep `answer[i] == 0` instead.

Example 1:

Input: `temperatures = [73,74,75,71,69,72,76,73]`
 Output: `[1,1,4,2,1,1,0,0]`

Example 2:

Input: `temperatures = [30,40,50,60]`
 Output: `[1,1,1,0]`

Example 3:

Input: `temperatures = [30,60,90]`
 Output: `[1,1,0]`

Round 5 · Mid · Medium

p.1997

Constraints:

- `1 <= temperatures.length <= 105`
- `30 <= temperatures[i] <= 100`

Community Solutions (Python)

• WalkCC

```
class Solution:
    def dailyTemperatures(self, temperatures:
list[int]) -> list[int]:
        ans = [0] * len(temperatures)
        stack = [] # a decreasing stack

        for i, temperature in enumerate(temperatures):
            while stack and temperature >
temperatures[stack[-1]]:
                index = stack.pop()
                ans[index] = i - index
            stack.append(i)

        return ans
```

• LeetDoocs

Round 5 · Mid · Medium

p.1998

```

class Solution:
    def dailyTemperatures(self, temperatures:
List[int]) -> List[int]:
        stk = []
        n = len(temperatures)
        ans = [0] * n
        for i in range(n - 1, -1, -1):
            while stk and temperatures[stk[-1]] <=
temperatures[i]:
                stk.pop()
            if stk:
                ans[i] = stk[-1] - i

        stk.append(i)
        return ans

```

• SimplyLeet

```

def dailyTemperatures(temperatures):
    n = len(temperatures)
    # Initialize answer list with zeros
    answer = [0] * n
    # Stack to keep indices of temperatures with
    unresolved next warmer days
    stack = []
    # Iterate through the temperature list
    for i, temp in enumerate(temperatures):
        # Check if the current temperature resolves
        any previous indices
        while stack and temperatures[stack[-1]] <
temp:

```

```

        prev_index = stack.pop() # Get the
index of the previous colder day
        answer[prev_index] = i - prev_index #
Update the number of days waited
        # Push the current index onto the stack
        stack.append(i)
        return answer

# Example usage:
if __name__ == "__main__":
    temps = [73,74,75,71,69,72,76,73]
    print(dailyTemperatures(temps))

```

• LC.ca

```

class Solution:
    def dailyTemperatures(self, temperatures:
List[int]) -> List[int]:
        ans = [0] * len(temperatures)
        stk = []
        for i, t in enumerate(temperatures):
            while stk and temperatures[stk[-1]] <
t:
                j = stk.pop()
                ans[j] = i - j
            stk.append(i)
        return ans

```

◆ Quick Review

Key Insights

- Use a monotonic stack to keep track of indices of days with unresolved warmer temperatures.
- As you iterate through the temperatures, resolve previous days that have found a warmer temperature.
- The index differences yield the number of days waited until a warmer temperature.

Space & Time

Time Complexity: $O(n)$, where n is the number of temperatures, because each index is pushed and popped from the stack at most once.

Space Complexity: $O(n)$ for storing the stack and the answer array.

Solution

The solution uses a monotonic stack to track the indices whose next warmer temperature has not been found yet. As you iterate through the temperature list, compare the current temperature with the temperature at the index stored at the top of the stack. If the current temperature is higher, it means you've found a warmer day for the day represented by the index from the stack. Pop that index, calculate the difference between the current index and the popped index, and store the result. Push the current index onto the stack for future comparisons. This approach efficiently finds the answer without nested iterations.

Stack

#962 Maximum Width Ramp

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Stack · Monotonic Stack

Lists: Denny Zhang

Problem

Description

A ramp in an integer array `nums` is a pair (i, j) for which $i < j$ and `nums[i] ≤ nums[j]`. The width of such a ramp is $j - i$.

Given an integer array `nums`, return the maximum width of a ramp in `nums`. If there is no ramp in `nums`, return 0.

Example 1:

Input: `nums = [6,0,8,2,1,5]`

Output: 4

Explanation: The maximum width ramp is achieved at $(i, j) = (1, 5)$: `nums[1] = 0` and `nums[5] = 5`.

Example 2:

Input: `nums = [9,8,1,0,1,9,4,0,4,1]`

Output: 7

Explanation: The maximum width ramp is achieved at $(i, j) = (2, 9)$: `nums[2] = 1` and `nums[9] = 1`.

Constraints:

- $2 \leq \text{nums.length} \leq 5 * 10^4$
- $0 \leq \text{nums}[i] \leq 5 * 10^4$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxWidthRamp(self, nums: list[int]) -> int:
        ans = 0
        stack = []

        for i, num in enumerate(nums):
            if stack == [] or num <= nums[stack[-1]]:
                stack.append(i)

        for i, num in reversed(list(enumerate(nums))):
            while stack and num >= nums[stack[-1]]:
                ans = max(ans, i - stack.pop())

        return ans
```

• LeetCode

```
class Solution:
    def maxWidthRamp(self, nums: List[int]) -> int:
        stk = []
        for i, v in enumerate(nums):
            if not stk or nums[stk[-1]] > v:
                stk.append(i)

        ans = 0
        for i in range(len(nums) - 1, -1, -1):
            while stk and nums[stk[-1]] <= nums[i]:
                ans = max(ans, i - stk.pop())

        if not stk:
            break

        return ans
```

• SimplyLeet

```
# Python solution for Maximum Width Ramp

def maxWidthRamp(nums):
    n = len(nums)
    stack = []
    # Build a stack of indices where array values
    are in decreasing order
    for i in range(n):
        if not stack or nums[i] < nums[stack[-1]]:
            stack.append(i)
    max_width = 0
```

```

# Iterate backwards to find the maximum width
ramp
for j in range(n - 1, -1, -1):
    while stack and nums[j] >=
nums[stack[-1]]:
        max_width = max(max_width, j -
stack.pop())
    return max_width

# Example usage:
nums = [6, 0, 8, 2, 1, 5]
print(maxWidthRamp(nums)) # Output: 4

```

• LC.ca

```

class Solution:
    def maxWidthRamp(self, nums: List[int]) ->
int:
        stk = []
        for i, v in enumerate(nums):
            if not stk or nums[stk[-1]] > v:
                stk.append(i)
        ans = 0
        for i in range(len(nums) - 1, -1, -1):
            while stk and nums[stk[-1]] <=
nums[i]:
                ans = max(ans, i - stk.pop())

        if not stk:
            break
        return ans

```

◆ Quick Review

Key Insights

- We need to identify pairs (i, j) where $i < j$ and $\text{nums}[i] \leq \text{nums}[j]$.
- A monotonic decreasing stack can be constructed to maintain candidate starting indices.
- By iterating from left to right, we push indices onto the stack only when they have a smaller value than the last.
- Then, a backward pass from the end of the array allows us to efficiently match each element with the smallest possible candidate index from the stack.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution is built in two phases:

- Build a decreasing stack of indices by iterating through the array (only push an index if its value is less than the value at the current top of the stack).
- Iterate through the array backwards. For each index j, check and pop elements from the stack while $\text{nums}[j]$ is greater than or equal to the nums value at the popped index. For each valid ramp found, compute the width ($j - \text{popped index}$) and update the maximum width accordingly. This approach guarantees we examine each element a constant number of times ensuring an efficient solution.

Stack

#1214 Two Sum BSTs

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · Binary Search · Stack · Tree · DFS · BST · BFS
Lists: Premium 98

Problem

Description

Given the roots of two binary search trees, `root1` and `root2`, return true if and only if there is a node in the first tree and a node in the second tree whose values sum up to a given integer target.

Example 1:

9	9	8	1
5	6	2	6
8	2	6	4
6	2	2	2


9	9
8	6

Input: `root1 = [2,1,4]`, `root2 = [1,0,3]`, `target = 5`

Output: true

Explanation: 2 and 3 sum up to 5.

Example 2:



Input: root1 = [0,-10,10], root2 = [5,1,7,0,2], target = 18
 Output: false

Constraints:

- The number of nodes in each tree is in the range [1, 5000].
- $-109 \leq \text{Node.val}, \text{target} \leq 109$

Community Solutions (Python)

- WalkCC

```

class BSTIterator:
    def __init__(self, root: TreeNode | None,
                 leftToRight: bool):
        self.stack = []
        self.leftToRight = leftToRight
        self._pushUntilNone(root)

    def hasNext(self) -> bool:
        return len(self.stack) > 0

    def next(self) -> int:

```

```

        node = self.stack.pop()
        if self.leftToRight:
            self._pushUntilNone(node.right)
        else:
            self._pushUntilNone(node.left)
        return node.val

```

```

    def _pushUntilNone(self, root: TreeNode | None):
        while root:
            self.stack.append(root)

```

```

    root = root.left if self.leftToRight else
    root.right

```

```

class Solution:
    def twoSumBSTs(
        self,
        root1: TreeNode | None,
        root2: TreeNode | None,
        target: int,
    ) -> bool:

```

```

    bst1 = BSTIterator(root1, True)
    bst2 = BSTIterator(root2, False)

```

```

    l = bst1.next()
    r = bst2.next()
    while True:
        summ = l + r
        if summ == target:
            return True
        if summ < target:

```

```

            if not bst1.hasNext():
                return False
            l = bst1.next()
        else:
            if not bst2.hasNext():
                return False
            r = bst2.next()

```

Round 5 · Mid · Medium

p.2011

• LeetCode

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def twoSumBSTs(
        self, root1: Optional[TreeNode], root2:
Optional[TreeNode], target: int
    ) -> bool:

```

```

    def dfs(root: Optional[TreeNode], i: int):
        if root is None:
            return
        dfs(root.left, i)
        nums[i].append(root.val)
        dfs(root.right, i)

```

```

    nums = [[], []]
    dfs(root1, 0)
    dfs(root2, 1)

```

Round 5 · Mid · Medium

p.2012

```

i, j = 0, len(nums[1]) - 1
while i < len(nums[0]) and ~j:
    x = nums[0][i] + nums[1][j]
    if x == target:
        return True
    if x < target:
        i += 1
    else:
        j -= 1
return False

```

• SimplyLeet

```

# Definition for a binary tree node.
class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

def twoSumBSTs(root1: TreeNode, root2: TreeNode,
target: int) -> bool:
    # Create a set to store values from the first
    BST.
    values_set = set()

```

```

# Helper function to perform DFS and add node
values to the set.
def dfs_collect(node):
    if not node:
        return
    values_set.add(node.val) # Add current
node value
    dfs_collect(node.left) # Traverse left
subtree
    dfs_collect(node.right) # Traverse right
subtree

```

```

# Collect all values from first tree.

```

```

dfs_collect(root1)

# Helper function to traverse second tree and
check for complement.
def dfs_check(node):
    if not node:
        return False
    # If the complement exists in the first
tree values, return True.
    if target - node.val in values_set:
        return True
    # Recursively check left and right
children.

```

```

        return dfs_check(node.left) or
dfs_check(node.right)

    # Check the second tree.
    return dfs_check(root2)

# Example usage:
# Construct trees and call twoSumBSTs as needed.

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
# def __init__(self, val=0, left=None,
right=None):
# self.val = val
# self.left = left
# self.right = right
class Solution:
    def twoSumBSTs(
        self, root1: Optional[TreeNode], root2:
Optional[TreeNode], target: int
    ) -> bool:

```

```

def dfs(root: Optional[TreeNode], i: int):
    if root is None:
        return
    dfs(root.left, i)
    nums[i].append(root.val)
    dfs(root.right, i)

```

```

nums = [[], []]
dfs(root1, 0)
dfs(root2, 1)

```

```

i, j = 0, len(nums[1]) - 1
while i < len(nums[0]) and ~j:
    x = nums[0][i] + nums[1][j]
    if x == target:
        return True
    if x < target:
        i += 1
    else:
        j -= 1
return False

```

◆ Quick Review

Key Insights

- Traverse the first tree to collect all node values.
- Traverse the second tree to check for each node if the complement (target – current node value) exists in the first tree's values.

- Using a hash set for the first tree values allows efficient $O(1)$ lookups.
- The BST property is not essential for the hash set solution, but can be exploited if using iterator techniques for a two-pointer approach.

Space & Time

Time Complexity: $O(n + m)$ where n and m are the number of nodes in the first and second trees respectively.

Space Complexity: $O(n)$ for storing the node values from the first tree (in worst-case scenario).

Solution

We perform a DFS (or any tree traversal) on the first BST to collect all its values in a hash set. Then, we traverse the second BST and for each node, we calculate the complement (target – node value) and check whether this complement exists in the hash set. If we find such a pair, we return true. If no such pair exists after traversing the second tree completely, we return false.

This approach is straightforward and leverages extra space (the set) for constant time lookups, yielding an overall linear time complexity relative to the number of nodes.

Stack

#1762 Buildings With an Ocean View

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Stack · Monotonic Stack

Lists: Premium 98

Problem

Description

There are n buildings in a line. You are given an integer array `heights` of size n that represents the heights of the buildings in the line.

The ocean is to the right of the buildings. A building has an ocean view if the building can see the ocean without obstructions. Formally, a building has an ocean view if all the buildings to its right have a smaller height.

Return a list of indices (0-indexed) of buildings that have an ocean view, sorted in increasing order.

Example 1:

Input: `heights = [4,2,3,1]`

Output: `[0,2,3]`

Explanation: Building 1 (0-indexed) does not have an ocean view because building 2 is taller.

Example 2:

Input: heights = [4,3,2,1]
 Output: [0,1,2,3]
 Explanation: All the buildings have an ocean view.

Example 3:

Input: heights = [1,3,2,4]
 Output: [3]
 Explanation: Only building 3 has an ocean view.

Constraints:

- $1 \leq \text{heights.length} \leq 105$
- $1 \leq \text{heights}[i] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findBuildings(self, heights: list[int]) -> list[int]:
        stack = []

        for i, height in enumerate(heights):
            while stack and heights[stack[-1]] <= height:
                stack.pop()
            stack.append(i)

        return stack
```

• LeetCode

```
class Solution:
    def findBuildings(self, heights: List[int]) -> List[int]:
        ans = []
        mx = 0
        for i in range(len(heights) - 1, -1, -1):
            if heights[i] > mx:
                ans.append(i)
                mx = heights[i]
        return ans[::-1]
```

• SimplyLeet

```
def findBuildingsWithOceanView(heights):
    n = len(heights)
    result = []
    max_height = 0
    # Iterate from the rightmost building to the leftmost building.
    for i in range(n - 1, -1, -1):
        if heights[i] > max_height:
            result.append(i)
            max_height = heights[i]
    # Reverse the result to return indices in increasing order.
```

```
result.reverse()
return result
```

```
# Example usage:
# heights = [4,2,3,1]
# print(findBuildingsWithOceanView(heights)) #
Output: [0,2,3]
```

• LC.ca

```
class Solution:
    def findBuildings(self, heights: List[int]) ->
List[int]:
    ans = []
    mx = 0
    for i in range(len(heights) - 1, -1, -1):
        if heights[i] > mx:
            ans.append(i)
            mx = heights[i]
    return ans[::-1]
```

◆ Quick Review

Key Insights

- Processing the buildings from right to left simplifies the problem because the rightmost building always has an ocean view.
- Keep track of the highest building encountered while iterating from right to left.

- A building has an ocean view if its height is greater than the maximum height seen so far.
- Append indices meeting the criteria, then reverse the order of indices for the solution since they were collected in reverse order.

Space & Time

Time Complexity: $O(n)$ where n is the number of buildings.

Space Complexity: $O(n)$ for storing the result in the worst-case scenario.

Solution

Iterate over the "heights" array from right to left while maintaining a variable to store the maximum height encountered so far. For each building, if its height is greater than this maximum height, add its index to the results list and update the maximum height. Since indices are collected in reverse order, reverse the list before returning it. This approach ensures that each building is visited only once.

Heap

Round 5 · Mid · Medium · 11 questions

Round 5 · Mid · Medium

p.2023

Heap

#215 Kth Largest Element in an Array

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Divide and Conquer · Sorting · Heap · Quickselect

Lists: Grind 169

Problem

Description

Given an integer array `nums` and an integer `k`, return the `k`th largest element in the array.

Note that it is the `k`th largest element in the sorted order, not the `k`th distinct element.

Can you solve it without sorting?

Example 1:

Input: `nums = [3,2,1,5,6,4]`, `k = 2`
Output: 5

Example 2:

Input: `nums = [3,2,3,1,2,4,5,5,6]`, `k = 4`
Output: 4

Constraints:

- $1 \leq k \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$

Round 5 · Mid · Medium

p.2024

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findKthLargest(self, nums: list[int], k:
int) -> int:
        minHeap = []

        for num in nums:
            heapq.heappush(minHeap, num)
            if len(minHeap) > k:
                heapq.heappop(minHeap)

        return minHeap[0]
```

```
class Solution:
    def findKthLargest(self, nums: list[int], k:
int) -> int:
        def quickSelect(l: int, r: int, k: int) -> int:
            pivot = nums[r]

            nextSwapped = l
            for i in range(l, r):
                if nums[i] >= pivot:
                    nums[nextSwapped], nums[i] = nums[i],
nums[nextSwapped]
                    nextSwapped += 1
```

```
        nums[nextSwapped], nums[r] = nums[r],
nums[nextSwapped]

        count = nextSwapped - l + 1 # Number of nums
>= pivot
        if count == k:
            return nums[nextSwapped]
        if count > k:
            return quickSelect(l, nextSwapped - 1, k)
        return quickSelect(nextSwapped + 1, r, k -
count)

        return quickSelect(0, len(nums) - 1, k)
```

• LeetCode

```
class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        def quick_sort(l: int, r: int) -> int:
            if l == r:
                return nums[l]
            i, j = l - 1, r + 1
            x = nums[(l + r) >> 1]
            while i < j:
                while 1:
                    i += 1
```

```

        if nums[i] >= x:
            break
    while 1:
        j -= 1
        if nums[j] <= x:
            break
    if i < j:
        nums[i], nums[j] = nums[j],
nums[i]
    if j < k:
        return quick_sort(j + 1, r)

    return quick_sort(l, j)

n = len(nums)
k = n - k
return quick_sort(0, n - 1)

```

```

class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        return nlargest(k, nums)[-1]

```

```

class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        cnt = Counter(nums)
        for i in count(max(cnt), -1):
            k -= cnt[i]
            if k <= 0:
                return i

```

• SimplyLeet

```

import heapq

def findKthLargest(nums, k):
    # Create a min-heap with the first k elements
    min_heap = nums[:k]
    heapq.heapify(min_heap) # O(k) time to build
the initial heap

    # Process the remaining elements
    for num in nums[k:]:
        # If the current element is larger than the
smallest in the heap, replace it

```

```

        if num > min_heap[0]:
            heapq.heapreplace(min_heap, num) #
O(log k)
        # The smallest element in the heap is the kth
largest overall
        return min_heap[0]

# Example usage:
nums = [3, 2, 1, 5, 6, 4]
k = 2
print(findKthLargest(nums, k)) # Expected output:
5

```

• LC.ca

```

class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        lo, hi = 0, len(nums) - 1

        while lo <= hi:
            pos = self.partition(nums, lo, hi)
            if pos == k - 1: # pos starting from 0,
so -1
                return nums[pos] # partially
sorted, desc
            elif pos > k - 1:
                hi = pos - 1

```

```

        else: # pos < k - 1
            lo = pos + 1

        return -1 # or raise exception

    def partition(self, nums: List[int], lo: int,
hi: int) -> int:
        pivot, l, r = nums[lo], lo + 1, hi
        while l <= r:
            if nums[l] < pivot and nums[r] > pivot:
                # larger num at left of pivot,
easier to count for k-th lagest
                nums[l], nums[r] = nums[r],
nums[l]
                l += 1
                r -= 1
            if nums[l] >= pivot:
                l += 1
            if nums[r] <= pivot:
                r -= 1
        # use nums[l] will lead to infinate looping
nums[lo], nums[r] = nums[r], nums[lo]
        # possible there is duplicated num, but
will be covered here

```

```

        return r

#####

class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        def quick_sort(left, right, k):
            if left == right:
                return nums[left]
            i, j = left - 1, right + 1

x = nums[(left + right) >> 1]
while i < j:
    while 1:
        i += 1
        if nums[i] >= x:
            break
    while 1:
        j -= 1
        if nums[j] <= x:
            break

```

```

        if i < j:
            nums[i], nums[j] = nums[j],
nums[i]
        if j < k:
            return quick_sort(j + 1, right, k)
        return quick_sort(left, j, k)

n = len(nums)
return quick_sort(0, n - 1, n - k)

#####

class Solution:
    def findKthLargest(self, nums: List[int], k:
int) -> int:
        def quick_sort(left, right, k):
            if left == right:
                return nums[left]
            i, j = left - 1, right + 1
            x = nums[(left + right) >> 1]
            while i < j:

```

```

        while 1:
            i += 1
            if nums[i] >= x:
                break
        while 1:
            j -= 1
            if nums[j] <= x:
                break
        if i < j:
            nums[i], nums[j] = nums[j],
nums[i]

        if j < k:
            return quick_sort(j + 1, right, k)
        return quick_sort(left, j, k)

n = len(nums)
return quick_sort(0, n - 1, n - k)

```

◆ Quick Review

Key Insights

- kth largest element can be converted to finding the element at index $n-k$ if the array were sorted.
- A min-heap of size k is a practical method to maintain the k largest elements seen so far.
- The Quickselect algorithm can find the kth largest element in average $O(n)$ time with in-place partitioning.

- Choose an appropriate method based on performance needs and worst-case scenarios.

Space & Time

Time Complexity:

- Min-Heap: $O(n \log k)$ by maintaining a heap of size k .
- Quickselect: Average $O(n)$ but worst-case $O(n^2)$ with poor pivot selection.

Space Complexity:

- Min-Heap: $O(k)$
- Quickselect: $O(1)$ additional space with in-place partitioning.

Solution

We can solve the problem using either a min-heap or the Quickselect algorithm.

- Min-Heap Approach: Initialize a min-heap with the first k elements. Iterate through the remaining elements; if an element is greater than the heap's root (the smallest in the heap), replace the root. After processing all elements, the root of the min-heap is the kth largest element. –

Quickselect Approach: Choose a pivot and partition the array so that elements larger than the pivot come before it. Determine the position of the pivot; if it matches the index corresponding to the kth largest element, return it. Recurse into the appropriate part of the array.

Key points to consider include maintaining heap size for the first approach and proper pivot selection for the Quickselect approach.

Heap

#253 Meeting Rooms II

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Greedy · Sorting · Heap

Lists: Grind 169 | Premium 98

Problem

Description

Given an array of meeting time intervals `intervals` where `intervals[i] = [starti, endi]`, return the minimum number of conference rooms required.

Example 1:

Input: `intervals = [[0,30],[5,10],[15,20]]`
 Output: 2

Example 2:

Input: `intervals = [[7,10],[2,4]]`
 Output: 1

Constraints:

- $1 \leq \text{intervals.length} \leq 104$
- $0 \leq \text{starti} < \text{endi} \leq 106$

Community Solutions (Python)

- WalkCC

Round 5 · Mid · Medium

p.2035

```
class Solution:
    def minMeetingRooms(self, intervals:
list[list[int]]) -> int:
        minHeap = [] # Store the end times of each
        room.

        for start, end in sorted(intervals):
            # There's no overlap, so we can reuse the
            same room.
            if minHeap and start >= minHeap[0]:
                heapq.heappop(minHeap)
                heapq.heappush(minHeap, end)

        return len(minHeap)
```

```
class Solution:
    def minMeetingRooms(self, intervals:
list[list[int]]) -> int:
        n = len(intervals)
        ans = 0
        starts = []
        ends = []

        for start, end in intervals:
            starts.append(start)
            ends.append(end)
```

Round 5 · Mid · Medium

p.2036

```

starts.sort()
ends.sort()

j = 0
for i in range(n):
    if starts[i] < ends[j]:
        ans += 1
    else:
        j += 1

```

```

return ans

```

• LeetCode

```

class Solution:
    def minMeetingRooms(self, intervals:
List[List[int]]) -> int:
        m = max(e[1] for e in intervals)
        d = [0] * (m + 1)
        for l, r in intervals:
            d[l] += 1
            d[r] -= 1
        ans = s = 0
        for v in d:
            s += v

```

```

        ans = max(ans, s)
        return ans

```

```

class Solution:
    def minMeetingRooms(self, intervals:
List[List[int]]) -> int:
        d = defaultdict(int)
        for l, r in intervals:
            d[l] += 1
            d[r] -= 1
        ans = s = 0
        for _, v in sorted(d.items()):
            s += v
            ans = max(ans, s)

```

```

return ans

```

• SimplyLeet

```

import heapq

def minMeetingRooms(intervals):
    # Sort intervals based on start time
    intervals.sort(key=lambda x: x[0])

    # Initialize a min-heap to store end times of
    meetings that are currently using rooms
    min_heap = []

    # Process each meeting interval

```

```

for interval in intervals:
    start, end = interval

    # If the heap is not empty and the current
    meeting starts after the earliest ended meeting
    if min_heap and start >= min_heap[0]:
        heapq.heappop(min_heap) # Reuse the
    room

    # Allocate the current meeting end time
    into the heap (new or reused room)
    heapq.heappush(min_heap, end)

# The number of rooms is the number of
simultaneous meetings
return len(min_heap)

# Example usage:
intervals = [[0,30],[5,10],[15,20]]
print(minMeetingRooms(intervals)) # Output: 2

```

• LC.ca

```

'''
data = [3, 4, 6, 2, 1, 9, 0, 7, 5, 8]

list(accumulate(data, operator.mul)) # running
product
[3, 12, 72, 144, 144, 1296, 0, 0, 0, 0]

list(accumulate(data, max)) # running maximum
[3, 4, 6, 6, 6, 9, 9, 9, 9, 9]

>>> from itertools import accumulate
>>> data = [3, 4, 6, 2, 1, 9, 0, 7, 5, 8]
>>> accumulate(data)
<itertools.accumulate object at 0x10fd78440>
>>> list(accumulate(data))
[3, 7, 13, 15, 16, 25, 25, 32, 37, 45]
>>> max(accumulate(data))
45

https://docs.python.org/3/library/itertools.html#i
tertools.accumulate

```



```
# Amortize a 5% loan of 1000 with 4 annual payments
of 90
cashflows = [1000, -90, -90, -90, -90]
list(accumulate(cashflows, lambda bal, pmt:
bal*1.05 + pmt))
[1000, 960.0, 918.0, 873.9000000000001,
827.5950000000001]
'''
```

```
class Solution:
    def minMeetingRooms(self, intervals:
List[List[int]]) -> int:
        delta = [0] * 1000010
```

```
        for start, end in intervals:
            delta[start] += 1
            delta[end] -= 1
        return max(accumulate(delta))
# why not delta.sort()?
# because accumulate() will go by order
from index 0 to index final
# just like, from sortedcontainers import
SortedDict
```

```
#####
```

```
class Solution(object):
    def minMeetingRooms(self, intervals):
        """
        :type intervals: List[Interval]
        :rtype: int
        """
        meetings = []
        for i in intervals:
            meetings.append((i.start, 1))
            meetings.append((i.end, 0))
```

```
        meetings.sort()
        ans = 0
        count = 0
        for meeting in meetings:
            if meeting[1] == 1:
                count += 1
            else:
                count -= 1
            ans = max(ans, count)
        return ans
```

◆ Quick Review

Key Insights

- Sort the intervals based on start times to process meetings in order.
- Use a min-heap (priority queue) to keep track of meeting end times.

- When processing a new meeting, check if its start time is greater than or equal to the smallest end time from the heap.
- If so, it means a meeting room has been freed (the meeting with the smallest end time has ended) and can be reused.
- If not, allocate a new room by pushing the current meeting's end time into the heap.
- The size of the heap at any moment represents the number of rooms currently in use.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting and each heap operation (push/pop) taking $O(\log n)$.

Space Complexity: $O(n)$ in the worst-case scenario when all meetings overlap, requiring all end times in the heap.

Solution

The solution involves first sorting the meeting intervals by their start times. Then, we utilize a min-heap (priority queue) to keep track of the end times of ongoing meetings. For each meeting, we compare its start time with the smallest end time in the heap:

– If the meeting can start after the meeting at the top of the heap has ended (i.e., its start time is greater than or equal to the smallest end), we remove that end time from the heap (freeing up a room). – Then, we add the current meeting's end time to the heap. By the end of processing all meetings, the size of the heap represents the minimum number of conference rooms required.

Heap

#347 Top K Frequent Elements

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Divide and Conquer · Sorting · Heap · Bucket Sort · Counting · Quickselect

Lists: Denny Zhang

Problem

Description

Given an integer array `nums` and an integer `k`, return the `k` most frequent elements. You may return the answer in any order.

Example 1:

Input: `nums = [1,1,1,2,2,3]`, `k = 2`

Output: `[1,2]`

Example 2:

Input: `nums = [1]`, `k = 1`

Output: `[1]`

Example 3:

Input: `nums = [1,2,1,2,1,2,3,1,3,2]`, `k = 2`

Output: `[1,2]`

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$

- k is in the range $[1, \text{the number of unique elements in the array}]$.
- It is guaranteed that the answer is unique.

Follow up: Your algorithm's time complexity must be better than $O(n \log n)$, where n is the array's size.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def topKFrequent(self, nums: list[int], k: int)
-> list[int]:
    ans = []
    bucket = [[] for _ in range(len(nums) + 1)]

    for num, freq in
collections.Counter(nums).items():
        bucket[freq].append(num)

    for b in reversed(bucket):
        ans += b

    if len(ans) == k:
        return ans
```

• LeetCode

```
class Solution:
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:
        cnt = Counter(nums)
        return [x for x, _ in cnt.most_common(k)]
```

• SimplyLeet

Python solution using bucket sort

```
def topKFrequent(nums, k):
    # Count the frequency of each element in nums
    frequency = {}
    for num in nums:
        frequency[num] = frequency.get(num, 0) + 1

    # Create buckets where index represents
frequency
    n = len(nums)
```

```
    buckets = [[] for _ in range(n + 1)]
    for num, count in frequency.items():
        buckets[count].append(num)

    # Gather the k most frequent elements
    result = []
    # Iterate over buckets in reverse order
(highest frequency first)
    for freq in range(n, 0, -1):
        for num in buckets[freq]:
            result.append(num)
```

```

        if len(result) == k:
            return result

```

```

# Example usage:
# print(topKFrequent([1,1,1,2,2,3], 2)) -> Output:
[1, 2]

```

• LC.ca

```

'''
>>> nums = ["a", "a", "b", "c", "c"]

>>> cnt = Counter(nums)
>>> cnt
Counter({'a': 2, 'c': 2, 'b': 1})

# default to keys
>>> sorted_freqs = sorted(cnt, key=lambda x:
(-cnt[x], x))
>>> sorted_freqs

```

```

['a', 'c', 'b']
'''

from collections import Counter

class Solution:
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:
        # Count the frequency of each element in
the list
        cnt = Counter(nums)

        # Sort the elements by frequency in
decreasing order

        # diff from below solution: sorted(cnt),
not sorted(cnt.items())
        # more in https://leetcode.ca/2017-10-22-6
92-Top-K-Frequent-Words/
        sorted_freqs = sorted(cnt, key=lambda x:
(-cnt[x], x))

        return sorted_freqs[:k]

#####

class Solution:
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:

```

```

        # Count the frequency of each element in
        the list
        freqs = Counter(nums)

        # Sort the elements by frequency in
        decreasing order
        sorted_freqs = sorted(freqs.items(),
                               key=lambda x: x[1], reverse=True)
        # also passing 0J: sorted_freqs =
        sorted(freqs.items(), key=lambda x: -x[1])

        # Take the top k elements
        top_k = [num for num, freq in
                  sorted_freqs[:k]]

```

```

        return top_k

#####

'''
>>> from heapq import heappush
>>> h = []
>>> heappush(h, (3,1))
>>> heappush(h, (1,1))
>>> heappush(h, (2,1))

```

```

>>> h
[(1, 1), (3, 1), (2, 1)]

>>> from heapq import heappop
>>> heappop(h)
(1, 1)
>>> heappop(h)
(2, 1)
>>> heappop(h)
(3, 1)

```

```

'''

from collections import Counter

class Solution: # heap
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:
        cnt = Counter(nums)
        hp = []
        for num, freq in cnt.items():
            heappush(hp, (freq, num)) # freq
first, default sort by 1st element

```

```

        if len(hp) > k:
            heappop(hp)
        return [v[1] for v in hp]

```

```
#####
```

```

from typing import List
import random
# O(n log(n)) in the average case
# O(n^2) worst case

```

```

class Solution:
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:
        freq_map = {}
        for num in nums:
            freq_map[num] = freq_map.get(num, 0) +
1

        freq_list = list(freq_map.items())
        self.quick_select(freq_list, 0,
len(freq_list) - 1, k)

        return [num for num, freq in freq_list[:k]]

```

```

def quick_select(self, freq_list, start, end,
k):

```

```

    if start == end:
        return

```

```

    pivot_idx = random.randint(start, end)
    pivot_freq = freq_list[pivot_idx][1]

```

```

    left = start
    right = end

```

```

        while left <= right:
            while freq_list[left][1] > pivot_freq:
                left += 1
            while freq_list[right][1] <
pivot_freq:
                right -= 1

            if left <= right:
                freq_list[left], freq_list[right]
= freq_list[right], freq_list[left]
                left += 1

```

```

        right -= 1

        if k <= right:
            self.quick_select(freq_list, start,
right, k)
        elif k >= left:
            self.quick_select(freq_list, left,
end, k)

```

```
#####
```

```

...
>>> x = [1, 2, 3]
>>> x.append([4, 5])
>>> print(x)
[1, 2, 3, [4, 5]]

>>> x = [1, 2, 3]
>>> x.extend([4, 5])

```

```

>>> print(x)
[1, 2, 3, 4, 5]
...

class Solution:
    def topKFrequent(self, nums: List[int], k:
int) -> List[int]:
        freq = Counter(nums)

        # not always filled, max possible frequency
is len(nums) [1,1,1,1,...]
        buckets = [[] for _ in range(len(nums) +
1)]

        for num, freq in freq.items():

            buckets[freq].append(num)

        res = []
        for bucket in reversed(buckets):
            if bucket:
                res.extend(bucket)
                if len(res) >= k:
                    break

        return res[:k]

```

```
#####
```

```
class Solution(object):
    def topKFrequent(self, nums, k):
        """
        :type nums: List[int]
        :type k: int
        :rtype: List[int]
        """
```

```
d = {}
res = []
ans = []
buckets = [[] for _ in range(len(nums) + 1)]

for num in nums:
    d[num] = d.get(num, 0) + 1

for key in d:
    res.append((d[key], key))
```

```
for t in res:
    freq, key = t
    buckets[freq].append(key)
```

```
buckets.reverse()
```

```
for item in buckets:
    if item and k > 0:
        while item and k > 0:
```

```
            ans.append(item.pop())
            k -= 1
        if k == 0:
            return ans
```

```
return ans
```

◆ Quick Review

Key Insights

- Count the frequency of each element using a hash table.
- Use bucket sort to achieve linear time complexity by grouping numbers by their frequency.
- Traverse the frequency buckets from highest to lowest to collect the k most frequent elements.
- Alternative approaches include using heaps or Quickselect, but bucket sort meets the follow-up requirement of better than $O(n \log n)$ time complexity.

Space & Time

Time Complexity: $O(n)$ where n is the length of the array (bucket sort and frequency counting are linear).

Space Complexity: $O(n)$ for storing the counts and buckets.

Solution

The solution starts by counting the frequency of each element with a hash table (or dictionary). Then, we create an array of buckets where the index represents the frequency. Each bucket holds the numbers that appear with that frequency. Finally, we iterate over the buckets in reverse order (from high frequency to low) and collect elements until we have k of them. This bucket sort approach avoids the higher complexity of sorting a list of frequencies and guarantees linear time performance.

Heap

#505 The Maze II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Graph · Matrix · Heap · Shortest Path

Lists: Premium 98

Problem

Description


There is a ball in a maze with empty spaces (represented as 0) and walls (represented as 1). The ball can go through the empty spaces by rolling up, down, left or right, but it won't stop rolling until hitting a wall. When the ball stops, it could choose the next direction.

Given the $m \times n$ maze, the ball's start position and the destination, where $start = [startrow, startcol]$ and $destination = [destinationrow, destinationcol]$, return the shortest distance for the ball to stop at the destination. If the ball cannot stop at destination, return -1 .

The distance is the number of empty spaces traveled by the ball from the start position (excluded) to the destination (included).

You may assume that the borders of the maze are all walls (see examples).

Example 1:




Input: maze = `[[0,0,1,0,0],[0,0,0,0,0],[0,0,0,1,0],[1,1,0,1,1],[0,0,0,0,0]]`, start = `[0,4]`, destination = `[4,4]`

Output: 12

Explanation: One possible way is : left -> down -> left -> down -> right -> down -> right.

The length of the path is $1 + 1 + 3 + 1 + 2 + 2 + 2 = 12$.

Example 2:



Input: maze = `[[0,0,1,0,0],[0,0,0,0,0],[0,0,0,1,0],[1,1,0,1,1],[0,0,0,0,0]]`, start = `[0,4]`, destination = `[3,2]`

Output: -1

Explanation: There is no way for the ball to stop at the destination. Notice that you can pass through the destination but you cannot stop there.

Example 3:

```
Input: maze = [[0,0,0,0,0],[1,1,0,0,1],[0,0,0,0,0],[0,1,0,0,1],[0,1,0,0,0]], start = [4,3], destination = [0,1]
Output: -1
```

Constraints:

- `m == maze.length`
- `n == maze[i].length`
- `1 <= m, n <= 100`
- `maze[i][j]` is 0 or 1.
- `start.length == 2`
- `destination.length == 2`
- `0 <= startrow, destinationrow < m`
- `0 <= startcol, destinationcol < n`
- Both the ball and the destination exist in an empty space, and they will not be in the same position initially.
- The maze contains at least 2 empty spaces.

Community Solutions (Python)

- [LeetDoocs](#)

```
class Solution:
    def shortestDistance(
        self, maze: List[List[int]], start:
        List[int], destination: List[int]
    ) -> int:
        m, n = len(maze), len(maze[0])
        dirs = (-1, 0, 1, 0, -1)
        si, sj = start
        di, dj = destination
        q = deque([(si, sj)])
        dist = [[inf] * n for _ in range(m)]

        dist[si][sj] = 0
        while q:
            i, j = q.popleft()
            for a, b in pairwise(dirs):
                x, y, k = i, j, dist[i][j]
                while 0 <= x + a < m and 0 <= y + b
                < n and maze[x + a][y + b] == 0:
                    x, y, k = x + a, y + b, k + 1
                    if k < dist[x][y]:
                        dist[x][y] = k
                        q.append((x, y))

        return -1 if dist[di][dj] == inf else
        dist[di][dj]
```

- [SimplyLeet](#)

```

import heapq

def shortestDistance(maze, start, destination):
    m, n = len(maze), len(maze[0])
    # Initialize distance matrix with infinity
    dist = [[float('inf')] * n for _ in range(m)]
    dist[start[0]][start[1]] = 0
    # Priority queue (min-heap) storing tuples
    (distance, row, col)
    heap = [(0, start[0], start[1])]
    # Directions: up, down, left, right

    directions = [(-1, 0), (1, 0), (0, -1), (0, 1)]

    while heap:
        d, row, col = heapq.heappop(heap)
        # If we have reached the destination,
        return the corresponding distance
        if [row, col] == destination:
            return d
        # Skip if we already found a better path to
        (row, col)
        if d > dist[row][col]:
            continue

```

```

        for dr, dc in directions:
            next_row, next_col, count = row, col, 0
            # Roll the ball until it hits a wall
            while 0 <= next_row + dr < m and 0 <=
next_col + dc < n and maze[next_row + dr][next_col
+ dc] == 0:
                next_row += dr
                next_col += dc
                count += 1
            # Check if the new distance is shorter
            if d + count <
dist[next_row][next_col]:
                dist[next_row][next_col] = d +
count

            heapq.heappush(heap,
(dist[next_row][next_col], next_row, next_col))
            return -1

```

• LC.ca

```

class Solution:
    def shortestDistance(
        self, maze: List[List[int]], start:
List[int], destination: List[int]
    ) -> int:
        m, n = len(maze), len(maze[0])
        dirs = (-1, 0, 1, 0, -1)
        si, sj = start
        di, dj = destination
        q = deque([(si, sj)])
        dist = [[inf] * n for _ in range(m)]

        dist[si][sj] = 0
        while q:
            i, j = q.popleft()
            for a, b in pairwise(dirs):
                x, y, k = i, j, dist[i][j]
                while 0 <= x + a < m and 0 <= y + b
< n and maze[x + a][y + b] == 0:
                    x, y, k = x + a, y + b, k + 1
                    if k < dist[x][y]:
                        dist[x][y] = k
                        q.append((x, y))

        return -1 if dist[di][dj] == inf else
dist[di][dj]

```

◆ Quick Review

Key Insights

- The ball rolls until it hits a wall, so each move is not one step but a continuous travel in a direction.
- The problem can be modeled as finding the shortest path in a weighted graph where nodes are positions in the maze and edge weights are the distance traveled in that roll.
- A priority queue (min-heap) based algorithm, such as Dijkstra's algorithm, is ideal since each roll covers variable distances.
- Always check if the new computed distance for a cell is less than a previously recorded one before pushing it to the heap.

Space & Time

Time Complexity: $O(m * n * \log(m * n))$, where m and n are the maze dimensions (each cell may get processed and queued with a logarithmic cost).

Space Complexity: $O(m * n)$ for maintaining the distance matrix and the priority queue.

Solution

We solve the problem using a modified version of Dijkstra's algorithm. The maze grid is treated as a graph where each cell is a node and an edge exists between the current cell and the cell where the ball stops after rolling in one direction. For every cell popped from the priority queue, we roll the ball in all four directions until it hits a wall, counting the number of steps taken. If the distance calculated from this roll is shorter than a previously stored distance for the stopping cell, we update the distance and add it to the priority queue. This process continues until we

either reach the destination (and return the distance) or exhaust all possibilities (returning -1).

Round 5 · Mid · Medium

p.2067

Heap

#621 Task Scheduler

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Greedy · Heap · Counting

Lists: Grind 169

Problem

Description

You are given an array of CPU tasks, each labeled with a letter from A to Z, and a number n . Each CPU interval can be idle or allow the completion of one task. Tasks can be completed in any order, but there's a constraint: there has to be a gap of at least n intervals between two tasks with the same label.

Return the minimum number of CPU intervals required to complete all tasks.

Example 1:

Input: tasks = ["A","A","A","B","B","B"], $n = 2$

Output: 8

Explanation: A possible sequence is: A -> B -> idle -> A -> B -> idle -> A -> B.

After completing task A, you must wait two intervals before doing A again. The same applies to task B. In the 3rd interval, neither A nor B can be done, so you idle. By

Round 5 · Mid · Medium

p.2068

the 4th interval, you can do A again as 2 intervals have passed.

Example 2:

Input: tasks = ["A","C","A","B","D","B"], n = 1

Output: 6

Explanation: A possible sequence is: A -> B -> C -> D -> A -> B.

With a cooling interval of 1, you can repeat a task after just one other task.

Example 3:

Input: tasks = ["A","A","A", "B","B","B"], n = 3

Output: 10

Explanation: A possible sequence is: A -> B -> idle -> idle -> A -> B -> idle -> idle -> A -> B.

There are only two types of tasks, A and B, which need to be separated by 3 intervals. This leads to idling twice between repetitions of these tasks.

Constraints:

- $1 \leq \text{tasks.length} \leq 104$
- tasks[i] is an uppercase English letter.
- $0 \leq n \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def leastInterval(self, tasks: list[str], n:
int) -> int:
        count = collections.Counter(tasks)
        maxFreq = max(count.values())
        # Put the most frequent task in the slot first.
        maxFreqTaskOccupy = (maxFreq - 1) * (n + 1)
        # Get the number of tasks with same frequency
as maxFreq, we'll append them after the
        # `maxFreqTaskOccupy`.
        nMaxFreq = sum(value == maxFreq for value in
count.values())
        # max(

        # the most frequent task is frequent enough to
force some idle slots,
        # the most frequent task is not frequent enough
to force idle slots
        # )
        return max(maxFreqTaskOccupy + nMaxFreq,
len(tasks))
```

- LeetCode

```
class Solution:
    def leastInterval(self, tasks: List[str], n:
int) -> int:
        cnt = Counter(tasks)
        x = max(cnt.values())
        s = sum(v == x for v in cnt.values())
        return max(len(tasks), (x - 1) * (n + 1) +
s)
```

• SimplyLeet

```
# Python solution for Task Scheduler
from collections import Counter

def leastInterval(tasks, n):
    # Count the frequency of each task
    task_counts = Counter(tasks)
    # Find the maximum frequency among tasks
    max_freq = max(task_counts.values())
    # Count how many tasks have the maximum
frequency
    max_count = sum(1 for count in
task_counts.values() if count == max_freq)
```

```
# Calculate the minimum intervals required
based on the most frequent tasks
part_count = max_freq - 1 # Number of parts
between the most frequent task
part_length = n - (max_count - 1) # Effective
empty slots per part after placing max frequency
tasks
empty_slots = part_count * part_length # Total
empty slots to fill with idle or other tasks
available_tasks = len(tasks) - max_freq *
max_count # Remaining tasks to fill the empty slots
idles = max(0, empty_slots - available_tasks) #
Remaining idle slots if available tasks are not
enough

# The total intervals is the sum of all tasks
and idle slots inserted
return len(tasks) + idles
```

```
# Example usage:
if __name__ == "__main__":
    tasks1 = ["A","A","A","B","B","B"]
    n1 = 2
    print(leastInterval(tasks1, n1)) # Output: 8
```

• LC.ca


```
class Solution:
    def leastInterval(self, tasks: List[str], n:
int) -> int:
        cnt = Counter(tasks)
        x = max(cnt.values())
        s = sum(v == x for v in cnt.values())
        return max(len(tasks), (x - 1) * (n + 1) +
s)
```

◆ Quick Review

Key Insights

- Count the frequency of each task.
- The task with the highest frequency determines the base structure.
- Identify how many tasks share the maximum frequency.
- Calculate the total time based on the formula: $\max(\text{total tasks}, (\text{max frequency} - 1) * (n + 1) + \text{count of tasks with maximum frequency})$.
- This approach avoids simulating the entire scheduling by using a greedy strategy.

Space & Time

Time Complexity: $O(N)$ where N is the number of tasks (counting frequency, then iterating over constant 26 letters).

Space Complexity: $O(1)$ since the frequency count uses an array or hash map with fixed size (26 letters).

Solution

We first count the frequency of each task. The task(s) with the maximum frequency establish the minimum structure required. Compute the ideal idle slots based on the count of these maximum tasks and the cooling period n . Finally, compare this computed length with the total number of tasks to get the final answer. This greedy strategy avoids explicit simulation and leverages mathematical formulation.

Heap

#658 Find K Closest Elements

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Binary Search · Sorting · Heap · Two Pointers

Lists: Grind 169

Problem

Description

Given a sorted integer array `arr`, two integers `k` and `x`, return the `k` closest integers to `x` in the array. The result should also be sorted in ascending order.

An integer `a` is closer to `x` than an integer `b` if:

- $|a - x| < |b - x|$, or
- $|a - x| == |b - x|$ and $a < b$

Example 1:

Input: `arr = [1,2,3,4,5]`, `k = 4`, `x = 3`

Output: `[1,2,3,4]`

Example 2:

Input: `arr = [1,1,2,3,4,5]`, `k = 4`, `x = -1`

Output: `[1,1,2,3]`

Constraints:

- $1 \leq k \leq \text{arr.length}$
- $1 \leq \text{arr.length} \leq 104$
- `arr` is sorted in ascending order.

Round 5 · Mid · Medium

p.2075

- $-104 \leq \text{arr}[i], x \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findClosestElements(self, arr: list[int], k:
int, x: int) -> list[int]:
        l = 0
        r = len(arr) - k

        while l < r:
            m = (l + r) // 2
            if x - arr[m] <= arr[m + k] - x:
                r = m
            else:
                l = m + 1

        return arr[l:l + k]
```

• LeetCode

```
class Solution:
    def findClosestElements(self, arr: List[int],
k: int, x: int) -> List[int]:
        arr.sort(key=lambda v: abs(v - x))
        return sorted(arr[:k])
```

• SimplyLeet

Round 5 · Mid · Medium

p.2076

```
# Function to find k closest elements to x in a
sorted array
def findClosestElements(arr, k, x):
    n = len(arr)
    # Initialize the binary search bounds for the
starting index of the window
    left, right = 0, n - k
    # Binary search to find the best starting index
    while left < right:
        mid = left + (right - left) // 2
        # Compare the distance from x to arr[mid]
and arr[mid+k]
        # If arr[mid] is farther from x than
arr[mid+k], move window to the right
```

```
        if x - arr[mid] > arr[mid + k] - x:
            left = mid + 1
        else:
            # Otherwise, continue searching
towards the left
            right = mid
        # Return the k elements starting from index
`left`
        return arr[left:left + k]
```

```
# Example usage:
# arr = [1,2,3,4,5], k = 4, x = 3
```

```
# print(findClosestElements(arr, 4, 3)) # Output:
[1,2,3,4]
```

• LC.ca

```
class Solution:
    def findClosestElements(self, arr: List[int],
k: int, x: int) -> List[int]:
        left, right = 0, len(arr) - k
        while left < right:
            mid = (left + right) >> 1
            if x - arr[mid] <= arr[mid + k] - x:
                right = mid
            else:
                left = mid + 1
        return arr[left : left + k]
```

◆ Quick Review

Key Insights

- The input array is already sorted.
- A binary search can be used to efficiently narrow down the potential starting position.
- A two-pointers or sliding window approach is effective for selecting k consecutive elements.
- The comparison of absolute differences (with a tie-breaker on values) is crucial.
- Maintaining ascending order simplifies the final output.

Space & Time

Time Complexity: $O(\log n + k)$, where $O(\log n)$ for binary search and $O(k)$ for collecting results.

Space Complexity: $O(k)$, for storing the output.

Solution

We can solve this problem by first using binary search to identify the left-most window where the k closest elements might start. The binary search operates over the possible starting indices for windows of size k (from 0 to $n-k$). For each middle index mid , we compare the distances between x and $arr[mid]$ and $arr[mid + k]$ to decide whether to move the window to the right or left. Once the correct starting position is found, return the k elements starting from that index. This approach leverages the sorted property of the array to efficiently converge on the optimal window.

Round 5 · Mid · Medium

p.2079

Heap

#787 Cheapest Flights Within K Stops

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Dynamic Programming · DFS · BFS · Graph · Heap

Lists: Grind 169

Problem

Description


There are n cities connected by some number of flights. You are given an array `flights` where `flights[i] = [fromi, toi, pricei]` indicates that there is a flight from city `fromi` to city `toi` with cost `pricei`.

You are also given three integers `src`, `dst`, and `k`, return the cheapest price from `src` to `dst` with at most k stops. If there is no such route, return `-1`.

Example 1:

Round 5 · Mid · Medium

p.2080




Input: $n = 4$, flights =
 $[[0,1,100],[1,2,100],[2,0,100],[1,3,600],[2,3,200]]$,
 $src = 0$, $dst = 3$, $k = 1$
 Output: 700
 Explanation:
 The graph is shown above.
 The optimal path with at most 1 stop from city 0 to 3
 is marked in red and has cost $100 + 600 = 700$.
 Note that the path through cities $[0,1,2,3]$ is cheaper
 but is invalid because it uses 2 stops.

Example 2:

Round 5 · Mid · Medium

p.2081




Input: $n = 3$, flights =
 $[[0,1,100],[1,2,100],[0,2,500]]$, $src = 0$, $dst = 2$, $k = 1$
 Output: 200
 Explanation:
 The graph is shown above.
 The optimal path with at most 1 stop from city 0 to 2
 is marked in red and has cost $100 + 100 = 200$.

Example 3:

Round 5 · Mid · Medium

p.2082



Input: $n = 3$, flights =
 $[[0,1,100],[1,2,100],[0,2,500]]$, src = 0, dst = 2, k = 0
 Output: 500
 Explanation:
 The graph is shown above.
 The optimal path with no stops from city 0 to 2 is marked in red and has cost 500.

Constraints:

- $2 \leq n \leq 100$
- $0 \leq \text{flights.length} \leq (n * (n - 1) / 2)$
- $\text{flights}[i].\text{length} == 3$
- $0 \leq \text{from}_i, \text{to}_i < n$

- $\text{from}_i \neq \text{to}_i$
- $1 \leq \text{price}_i \leq 104$
- There will not be any multiple flights between two cities.
- $0 \leq \text{src}, \text{dst}, k < n$
- $\text{src} \neq \text{dst}$

Community Solutions (Python)

• WalkCC

```

class Solution:
    def findCheapestPrice(
        self,
        n: int,
        flights: list[list[int]],
        src: int,
        dst: int,
        k: int,
    ) -> int:
        graph = [[] for _ in range(n)]
  
```

```

for u, v, w in flights:
    graph[u].append((v, w))

return self._dijkstra(graph, src, dst, k)

def _dijkstra(
    self,
    graph: list[list[tuple[int, int]]],
    src: int,

    dst: int,
    k: int,
) -> int:
    dist = [[math.inf] * (k + 2) for _ in
range(len(graph))]

    dist[src][k + 1] = 0
    minHeap = [(dist[src][k + 1], src, k + 1)] #
(d, u, stops)

    while minHeap:
        d, u, stops = heapq.heappop(minHeap)

```

```

if u == dst:
    return d
if stops == 0 or d > dist[u][stops]:
    continue
for v, w in graph[u]:
    if d + w < dist[v][stops - 1]:
        dist[v][stops - 1] = d + w
        heapq.heappush(minHeap, (dist[v][stops -
1], v, stops - 1))

return -1

```

• LeetCode

```

class Solution:
    def findCheapestPrice(
        self, n: int, flights: List[List[int]],
src: int, dst: int, k: int
    ) -> int:
        INF = 0x3F3F3F3F
        dist = [INF] * n
        dist[src] = 0
        for _ in range(k + 1):
            backup = dist.copy()
            for f, t, p in flights:

                dist[t] = min(dist[t], backup[f] +
p)

        return -1 if dist[dst] == INF else
dist[dst]

```

- SimplyLeet

```
import heapq

def findCheapestPrice(n, flights, src, dst, k):
    # Build graph: each city maps to a list of
    # (neighboring city, flight price)
    graph = {i: [] for i in range(n)}
    for from_city, to_city, price in flights:
        graph[from_city].append((to_city, price))

    # Priority queue: (total_cost, current_city,
    # stops so far)
    heap = [(0, src, 0)]

    # Dictionary to store the best cost to reach a
    # city with a given stop count
    best = {}

    while heap:
        cost, city, stops = heapq.heappop(heap)
        # Return cost if destination is reached
        if city == dst:
            return cost
```

```
        # If the current number of stops exceeds k,
        # skip further processing
        if stops > k:
            continue

        # Expand the neighboring flights
        for neighbor, price in graph[city]:
            new_cost = cost + price
            # Only proceed if this route is cheaper
            # or not yet explored for this many stops
            if (neighbor, stops + 1) not in best or
            new_cost < best[(neighbor, stops + 1)]:
                best[(neighbor, stops + 1)] =
                new_cost

            heapq.heappush(heap, (new_cost,
            neighbor, stops + 1))
        return -1

# Example usage:
print(findCheapestPrice(4, [[0,1,100],[1,2,100],[2
,0,100],[1,3,600],[2,3,200]], 0, 3, 1))
```

- LC.ca


```

class Solution:
    def findCheapestPrice(
        self, n: int, flights: List[List[int]],
        src: int, dst: int, k: int
    ) -> int:
        INF = 0x3F3F3F3F
        dist = [INF] * n
        dist[src] = 0
        for _ in range(k + 1):
            backup = dist.copy()
            for f, t, p in flights:
                dist[t] = min(dist[t], backup[f] +
p)
            return -1 if dist[dst] == INF else
dist[dst]

```

◆ Quick Review

Key Insights

- Build the graph as an adjacency list of flights and prices.
- Use a min-heap so the cheapest partial route is always processed first.
- Track states as `(city, flightsUsed)` instead of only `city`, because the stop count matters.
- For each popped state, try all outgoing flights and push the updated cost and flight count.
- Do not expand a state once it has already used more than `k + 1` flights.

- Keep the best known cost for each `(city, flightsUsed)` state so worse duplicate states can be skipped.

Space & Time

Time Complexity: $O((V + E) * \log(V * (k + 1)))$ for the heap-based search over `(city, stopsUsed)` states.

Space Complexity: $O(E + V * (k + 1))$ for the graph plus the best-cost tracking table.

Solution

We solve the problem using a Dijkstra-style min-heap search, but the state must include both the city and how many flights have been used so far. We start from `(0, src, 0)` and always pop the cheapest state from the heap. For each state, if we reach `dst`, that cost is the answer. Otherwise, if we still have flights available within the `k + 1` limit, we explore all outgoing edges and push the next states into the heap. To avoid useless work, we store the best cost seen for each `(city, flightsUsed)` pair and only continue when the new state improves that record. This guarantees the cheapest valid route under the stop constraint.

Heap

#973 K Closest Points to Origin

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Math · Divide and Conquer · Sorting · Heap

Lists: Grind 169

Problem


Description

Given an array of points where $\text{points}[i] = [x_i, y_i]$ represents a point on the X-Y plane and an integer k , return the k closest points to the origin $(0, 0)$.

The distance between two points on the X-Y plane is the Euclidean distance (i.e., $\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$).

You may return the answer in any order. The answer is guaranteed to be unique (except for the order that it is in).

Example 1:



Input: $\text{points} = [[1,3], [-2,2]]$, $k = 1$

Output: $[[-2,2]]$

Explanation:

The distance between $(1, 3)$ and the origin is $\sqrt{10}$. The distance between $(-2, 2)$ and the origin is $\sqrt{8}$. Since $\sqrt{8} < \sqrt{10}$, $(-2, 2)$ is closer to the origin.

We only want the closest $k = 1$ points from the origin, so the answer is just $[[-2,2]]$.

Example 2:

Input: points = [[3,3],[5,-1],[-2,4]], k = 2
 Output: [[3,3],[-2,4]]
 Explanation: The answer [[-2,4],[3,3]] would also be accepted.

Constraints:

- $1 \leq k \leq \text{points.length} \leq 104$
- $-104 \leq x_i, y_i \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def kClosest(self, points: list[list[int]], k:
int) -> list[list[int]]:
        maxHeap = []

        for x, y in points:
            heapq.heappush(maxHeap, (- x * x - y * y, [x,
y]))
            if len(maxHeap) > k:
                heapq.heappop(maxHeap)

        return [pair[1] for pair in maxHeap]
```

```
class Solution:
    def kClosest(self, points: list[list[int]], k:
int) -> list[list[int]]:
        def squareDist(p: list[int]) -> int:
            return p[0] * p[0] + p[1] * p[1]

        def quickSelect(l: int, r: int, k: int) ->
None:
            pivot = points[r]

            nextSwapped = l
            for i in range(l, r):
```

```
                if squareDist(points[i]) <=
squareDist(pivot):
                    points[nextSwapped], points[i] =
points[i], points[nextSwapped]
                    nextSwapped += 1
                points[nextSwapped], points[r] = points[r],
points[nextSwapped]

            count = nextSwapped - l + 1 the number of
points <= pivot
            if count == k:
                return
            if count > k:
                quickSelect(l, nextSwapped - 1, k)
```

```

else:
    quickSelect(nextSwapped + 1, r, k - count)

quickSelect(0, len(points) - 1, k)
return points[0:k]

```

• LeetCode

```

class Solution:
    def kClosest(self, points: List[List[int]], k:
int) -> List[List[int]]:
        points.sort(key=lambda p: hypot(p[0],
p[1]))
        return points[:k]

```

• SimplyLeet

```

# Python Solution
def kClosest(points, k):
    # Sort the points based on their squared
    Euclidean distance from the origin.
    points.sort(key=lambda point: point[0]**2 +
point[1]**2)
    # Return the first k points from the sorted
    list.
    return points[:k]

# Example usage:
# points = [[1,3], [-2,2]]
# k = 1

```

```
# print(kClosest(points, k))
```

• LC.ca

```

class Solution:
    def kClosest(self, points: List[List[int]], k:
int) -> List[List[int]]:
        points.sort(key=lambda p: p[0] * p[0] +
p[1] * p[1])
        return points[:k]

```

◆ Quick Review

Key Insights

- Calculate the squared Euclidean distance ($x^2 + y^2$) to avoid the overhead of square root computation.
- Sorting the points based on their distance gives an $O(n \log n)$ solution.
- Alternatively, using a max-heap to keep track of the k closest points results in $O(n \log k)$ time, which is effective if k is much smaller than n .
- Using Quickselect can achieve an average time complexity of $O(n)$, though the worst-case is $O(n^2)$.
- The problem does not require the output to be sorted in any specific order.

Space & Time

Time Complexity: $O(n \log n)$ with a sorting approach, $O(n \log k)$ with a heap approach, or average $O(n)$ with a Quickselect approach.

Space Complexity: $O(n)$ if extra space is used for sorting, or $O(k)$ for the heap approach.

Solution

We can solve this problem by first computing the squared Euclidean distance for each point to avoid unnecessary square root operations. Then, we can sort the points based on their computed distance and select the first k points. Alternatively, a max-heap or Quickselect algorithm can be used:

- **Max-Heap:** Maintain a heap of size k while iterating through the points, ensuring that only the k points with the smallest distances are kept.
- **Quickselect:** Partition the array like in QuickSort to position the k th closest point in its correct position. Once partitioned, all points to the left of that position will be the k closest. The sorting solution is straightforward and efficient enough given the constraints.

Heap

#1057 Campus Bikes

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Greedy · Sorting · Heap

Lists: Premium 98

Problem

Description

On a campus represented on the X-Y plane, there are n workers and m bikes, with $n \leq m$.

You are given an array `workers` of length n where `workers[i] = [xi, yi]` is the position of the i th worker. You are also given an array `bikes` of length m where `bikes[j] = [xj, yj]` is the position of the j th bike. All the given positions are unique.


Assign a bike to each worker. Among the available bikes and workers, we choose the $(worker_i, bike_j)$ pair with the shortest Manhattan distance between each other and assign the bike to that worker.

If there are multiple $(worker_i, bike_j)$ pairs with the same shortest Manhattan distance, we choose the pair with the smallest worker index. If there are multiple ways to do that, we choose the pair with the smallest bike index. Repeat this process until there are no available workers.

Return an array `answer` of length `n`, where `answer[i]` is the index (0-indexed) of the bike that the `i`th worker is assigned to.

The Manhattan distance between two points `p1` and `p2` is $\text{Manhattan}(p1, p2) = |p1.x - p2.x| + |p1.y - p2.y|$.

Example 1:




Input: `workers = [[0,0],[2,1]]`, `bikes = [[1,2],[3,3]]`

Output: `[1,0]`

Explanation: Worker 1 grabs Bike 0 as they are closest (without ties), and Worker 0 is assigned Bike 1. So the output is `[1, 0]`.

Example 2:



Input: workers = [[0,0],[1,1],[2,0]], bikes = [[1,0],[2,2],[2,1]]

Output: [0,2,1]

Explanation: Worker 0 grabs Bike 0 at first. Worker 1 and Worker 2 share the same distance to Bike 2, thus Worker 1 is assigned to Bike 2, and Worker 2 will take Bike 1. So the output is [0,2,1].

Constraints:

- $n == \text{workers.length}$
- $m == \text{bikes.length}$
- $1 \leq n \leq m \leq 1000$
- $\text{workers}[i].\text{length} == \text{bikes}[j].\text{length} == 2$
- $0 \leq x_i, y_i < 1000$
- $0 \leq x_j, y_j < 1000$
- All worker and bike locations are unique.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def assignBikes(
        self,
        workers: list[list[int]],
        bikes: list[list[int]],
    ) -> list[int]:
        ans = [-1] * len(workers)
        usedBikes = [False] * len(bikes)
        # buckets[k] := (i, j), where k =
        dist(workers[i], bikes[j])
        buckets = [[] for _ in range(2001)]
```

```
def dist(p1: list[int], p2: list[int]) -> int:
    return abs(p1[0] - p2[0]) + abs(p1[1] -
p2[1])
```

```
for i, worker in enumerate(workers):
    for j, bike in enumerate(bikes):
        buckets[dist(worker, bike)].append((i, j))

for k in range(2001):
    for i, j in buckets[k]:
```

```
        if ans[i] == -1 and not usedBikes[j]:
            ans[i] = j
            usedBikes[j] = True
```

```
return ans
```

- LeetCode

```

class Solution:
    def assignBikes(
        self, workers: List[List[int]], bikes:
        List[List[int]]
    ) -> List[int]:
        n, m = len(workers), len(bikes)
        arr = []
        for i, j in product(range(n), range(m)):
            dist = abs(workers[i][0] -
bikes[j][0]) + abs(workers[i][1] - bikes[j][1])
            arr.append((dist, i, j))
        arr.sort()

        vis1 = [False] * n
        vis2 = [False] * m
        ans = [0] * n
        for _, i, j in arr:
            if not vis1[i] and not vis2[j]:
                vis1[i] = vis2[j] = True
                ans[i] = j
        return ans

```

• SimplyLeet

```

# Python solution

def assign_bikes(workers, bikes):
    # List to hold all possible (distance, worker,
    bike) tuples
    pairs = []
    for worker_index, (wx, wy) in
    enumerate(workers):
        for bike_index, (bx, by) in
    enumerate(bikes):
            # Calculate Manhattan distance between
            worker and bike
            distance = abs(wx - bx) + abs(wy - by)
            pairs.append((distance, worker_index,
            bike_index))

```



```

# Sort pairs based on distance, then worker
index, then bike index
pairs.sort()

n = len(workers) # Number of workers
assigned_workers = [False] * n # Track if
worker is already assigned
assigned_bikes = [False] * len(bikes) # Track
if bike is already assigned
result = [-1] * n # Result list to store bike
assignment for each worker

# Iterate over the sorted pairs and assign
bikes

```

```

    for distance, worker_index, bike_index in
pairs:
        if not assigned_workers[worker_index] and
not assigned_bikes[bike_index]:
            result[worker_index] = bike_index #
Assign bike to worker
            assigned_workers[worker_index] = True
# Mark worker as assigned
            assigned_bikes[bike_index] = True #
Mark bike as assigned
        return result

# Example usage:
workers = [[0,0],[2,1]]
bikes = [[1,2],[3,3]]

print(assign_bikes(workers, bikes)) # Output: [1,
0]

```

- LC.ca

```

class Solution:
    def assignBikes(
        self, workers: List[List[int]], bikes:
List[List[int]]
    ) -> List[int]:
        n, m = len(workers), len(bikes)
        arr = []
        for i, j in product(range(n), range(m)):
            dist = abs(workers[i][0] -
bikes[j][0]) + abs(workers[i][1] - bikes[j][1])
            arr.append((dist, i, j))
        arr.sort()

        vis1 = [False] * n
        vis2 = [False] * m
        ans = [0] * n
        for _, i, j in arr:
            if not vis1[i] and not vis2[j]:
                vis1[i] = vis2[j] = True
                ans[i] = j
        return ans

```

◆ Quick Review

Key Insights

- Compute the Manhattan distance between every worker and bike pair.
- To achieve the tie-breaking rules, sort all pairs by distance, then by worker index and bike index.

- Once a worker is assigned a bike, skip any other pairs involving that worker or the already assigned bike.
- The problem can also be approached using bucket sort because the Manhattan distance range is limited, but sorting is intuitive given the constraints.

Space & Time

Time Complexity: $O(n * m * \log(n * m))$ where n is the number of workers and m is the number of bikes.

Space Complexity: $O(n * m)$ for storing all worker-bike pairs.

Solution

Create all possible (worker, bike) pairs along with their Manhattan distances. Sort this list of pairs by:

– Manhattan distance, – Worker index (if distances are equal), – Bike index (if both distance and worker index are equal).

Then, iterate over the sorted list and assign bikes to workers if neither the worker nor the bike has already been assigned. This guarantees that at each step, the optimal available pairing (according to the problem's requirements) is made. This method naturally enforces the tie-breaking rules specified in the problem.

Heap

#1167 Minimum Cost to Connect Sticks

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Greedy · Heap

Lists: Premium 98

Problem

Description

You have some number of sticks with positive integer lengths. These lengths are given as an array `sticks`, where `sticks[i]` is the length of the *i*th stick.

You can connect any two sticks of lengths *x* and *y* into one stick by paying a cost of *x* + *y*. You must connect all the sticks until there is only one stick remaining.

Return the minimum cost of connecting all the given sticks into one stick in this way.

Example 1:

Input: `sticks = [2,4,3]`

Output: 14

Explanation: You start with `sticks = [2,4,3]`.

1. Combine sticks 2 and 3 for a cost of $2 + 3 = 5$. Now you have `sticks = [5,4]`.

2. Combine sticks 5 and 4 for a cost of $5 + 4 = 9$. Now you have `sticks = [9]`.

There is only one stick left, so you are done. The total cost is $5 + 9 = 14$.

Example 2:

Input: `sticks = [1,8,3,5]`

Output: 30

Explanation: You start with `sticks = [1,8,3,5]`.

1. Combine sticks 1 and 3 for a cost of $1 + 3 = 4$. Now you have `sticks = [4,8,5]`.

2. Combine sticks 4 and 5 for a cost of $4 + 5 = 9$. Now you have `sticks = [9,8]`.

3. Combine sticks 9 and 8 for a cost of $9 + 8 = 17$. Now you have `sticks = [17]`.

There is only one stick left, so you are done. The total cost is $4 + 9 + 17 = 30$.

Example 3:

Input: `sticks = [5]`

Output: 0

Explanation: There is only one stick, so you don't need to do anything. The total cost is 0.

Constraints:

- $1 \leq \text{sticks.length} \leq 104$
- $1 \leq \text{sticks}[i] \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def connectSticks(self, sticks: list[int]) ->
    int:
        ans = 0
        heapq.heapify(sticks)

        while len(sticks) > 1:
            x = heapq.heappop(sticks)
            y = heapq.heappop(sticks)
            ans += x + y
            heapq.heappush(sticks, x + y)

        return ans
```

• LeetDoocs

```
class Solution:
    def connectSticks(self, sticks: List[int]) ->
    int:
        heapify(sticks)
        ans = 0
        while len(sticks) > 1:
            z = heappop(sticks) + heappop(sticks)
            ans += z
            heappush(sticks, z)
        return ans
```

Round 5 · Mid · Medium

p.2111

• SimplyLeet

```
import heapq

def connectSticks(sticks):
    # Initialize the heap with the provided stick
    lengths
    heapq.heapify(sticks)

    # Initialize total cost to 0
    total_cost = 0

    # Continue connecting sticks until there is
    only one left

    while len(sticks) > 1:
        # Extract the two smallest sticks
        stick1 = heapq.heappop(sticks)
        stick2 = heapq.heappop(sticks)
        # Calculate the cost to connect these two
        sticks
        cost = stick1 + stick2
        # Add the cost to the total cost
        total_cost += cost
        # Push the combined stick back into the
        heap
        heapq.heappush(sticks, cost)
```

Round 5 · Mid · Medium

p.2112

```
return total_cost
```

```
# Example usage:
```

```
# print(connectSticks([2, 4, 3])) # Output: 14
```

- LC.ca

```
class Solution:
```

```
    def connectSticks(self, sticks: List[int]) ->
```

```
    int:
```

```
        heapify(sticks)
```

```
        ans = 0
```

```
        while len(sticks) > 1:
```

```
            z = heappop(sticks) + heappop(sticks)
```

```
            ans += z
```

```
            heappush(sticks, z)
```

```
        return ans
```

◆ Quick Review

Key Insights

- Always connect the two smallest sticks available to minimize the incremental cost.
- A min-heap (or priority queue) is an ideal data structure to efficiently retrieve the smallest sticks.
- The process continues until only one stick remains.

Space & Time

Time Complexity: $O(n \log n)$, where n is the number of sticks (due to heap insertion and removal operations).

Space Complexity: $O(n)$, for storing the sticks in the heap.

Solution

The solution uses a greedy algorithm with a min-heap:

- Insert all stick lengths into a min-heap.
- While there are at least two sticks in the heap: Extract the two smallest sticks. Compute the cost of connecting them (sum the two lengths). Add this cost to the total result. Push the combined stick (with length equal to their sum) back into the heap.
- Once the heap is reduced to one element, return the total cost.

This approach ensures that at every connection, the smallest possible sticks are merged, which leads to the minimum overall cost.

Heap

#1424 Diagonal Traverse II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Sorting · Heap


Lists: CodeSignal

Problem

Description

Given a 2D integer array `nums`, return all elements of `nums` in diagonal order as shown in the below images.


Example 1:



Input: `nums = [[1,2,3],[4,5,6],[7,8,9]]`

Output: `[1,4,2,7,5,3,8,6,9]`

Example 2:



Input: nums =
[[1,2,3,4,5],[6,7],[8],[9,10,11],[12,13,14,15,16]]
Output: [1,6,2,8,7,3,9,4,12,10,5,13,11,14,15,16]

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $1 \leq \text{nums}[i].\text{length} \leq 105$
- $1 \leq \sum(\text{nums}[i].\text{length}) \leq 105$
- $1 \leq \text{nums}[i][j] \leq 105$

Community Solutions (Python)

• LeetCode

```
class Solution:
    def findDiagonalOrder(self, nums:
List[List[int]]) -> List[int]:
        arr = []
        for i, row in enumerate(nums):
            for j, v in enumerate(row):
                arr.append((i + j, j, v))
        arr.sort()
        return [v[2] for v in arr]
```

• SimplyLeet

```
def findDiagonalOrder(nums):
    # Use a dictionary to collect elements by their
    diagonal key (i + j)
    from collections import defaultdict
    diagonals = defaultdict(list)

    # Iterate over each row and its elements
    for i, row in enumerate(nums):
        for j, value in enumerate(row):
            # Insert at the beginning to achieve
            the desired order for the diagonal
            diagonals[i + j].insert(0, value)
```

```

result = []
# Process the diagonals in order of increasing
key
for key in sorted(diagonals.keys()):
    result.extend(diagonals[key])
return result

# Example usage:
# nums = [[1,2,3],[4,5,6],[7,8,9]]
# print(findDiagonalOrder(nums)) # Output:
[1,4,2,7,5,3,8,6,9]

```

• LC.ca

```

class Solution:
    def findDiagonalOrder(self, nums:
List[List[int]]) -> List[int]:
        arr = []
        for i, row in enumerate(nums):
            for j, v in enumerate(row):
                arr.append((i + j, j, v))
        arr.sort()
        return [v[2] for v in arr]

```

◆ Quick Review

Key Insights

- The diagonal key for each element can be defined as (row_index + column_index).

- Elements in the same diagonal (with equal diagonal key) must be output in an order such that the element from the lower row comes before the one from a higher row.
- Because the input grid may be jagged (rows with varying lengths), you must carefully iterate only over valid indices.
- A hash table (or map) can be used to collect elements grouped by their diagonal key.
- Inserting each element at the beginning of the list for that diagonal builds the list in the desired order without needing an extra reverse step later.
- Finally, traverse the keys in increasing order (from 0 to maximum key) to produce the final result.

Space & Time

Time Complexity: $O(N)$, where N is the total number of elements (each element is processed once).

Space Complexity: $O(N)$, for storing the elements in the map and the output array.

Solution

The solution iterates over every element in `nums` while computing its diagonal key as $(i + j)$. For each element, the value is inserted at the beginning of the list corresponding to that diagonal key. This ensures that when the diagonal groups are later concatenated in order of increasing key, the internal ordering within each diagonal is correct (i.e., elements appear from the bottom of the diagonal to the top). The final result is built by iterating over the keys in sorted order and appending each list of diagonal elements.

Trie

Round 5 · Mid · Medium · 7 questions

Round 5 · Mid · Medium

p.2121

Trie

#139 Word Break

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Dynamic Programming · Trie · Memoization

Lists: Grind 169

Problem

Description

Given a string s and a dictionary of strings $wordDict$, return true if s can be segmented into a space-separated sequence of one or more dictionary words.

Note that the same word in the dictionary may be reused multiple times in the segmentation.

Example 1:

Input: $s = \text{"leetcode"}$, $wordDict = [\text{"leet"}, \text{"code"}]$

Output: true

Explanation: Return true because "leetcode" can be segmented as "leet code".

Example 2:

Round 5 · Mid · Medium

p.2122

Input: `s = "applepenapple", wordDict = ["apple","pen"]`
 Output: `true`
 Explanation: Return true because "applepenapple" can be segmented as "apple pen apple".
 Note that you are allowed to reuse a dictionary word.

Example 3:

Input: `s = "catsanddog", wordDict = ["cats","dog","sand","and","cat"]`
 Output: `false`

Constraints:

- $1 \leq s.length \leq 300$
- $1 \leq wordDict.length \leq 1000$
- $1 \leq wordDict[i].length \leq 20$
- `s` and `wordDict[i]` consist of only lowercase English letters.
- All the strings of `wordDict` are unique.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def wordBreak(self, s: str, wordDict: list[str])
    -> bool:
        wordSet = set(wordDict)

        @functools.lru_cache(None)
        def wordBreak(s: str) -> bool:
            """Returns True if s can be segmented."""
            if s in wordSet:
                return True
            return any(s[:i] in wordSet and
                      wordBreak(s[i:]) for i in range(len(s)))
```

```
        return wordBreak(s)
```

```
class Solution:
    def wordBreak(self, s: str, wordDict: list[str])
    -> bool:
        n = len(s)
        wordSet = set(wordDict)
        # dp[i] := True if s[0..i) can be segmented
        dp = [True] + [False] * n

        for i in range(1, n + 1):
            for j in range(i):
                # s[0..j) can be segmented and s[j..i) is
                in `wordSet`, so s[0..i) can
```

```

    # be segmented.
    if dp[j] and s[j:i] in wordSet:
        dp[i] = True
        break

    return dp[n]

```

```

class Solution:
    def wordBreak(self, s: str, wordDict: list[str])
    -> bool:
        n = len(s)
        maxLength = len(max(wordDict, key=len))
        wordSet = set(wordDict)
        # dp[i] := True if s[0..i) can be segmented
        dp = [True] + [False] * n

        for i in range(1, n + 1):
            for j in range(i - 1, -1, -1):

```

```

                if i - j > maxLength:
                    break
                # s[0..j) can be segmented and s[j..i) is
in the wordSet, so s[0..i)
                # can be segmented.
                if dp[j] and s[j:i] in wordSet:
                    dp[i] = True
                    break

            return dp[n]

```

• LeetCode

Round 5 · Mid · Medium

p.2125

```

class Solution:
    def wordBreak(self, s: str, wordDict:
List[str]) -> bool:
        words = set(wordDict)
        n = len(s)
        f = [True] + [False] * n
        for i in range(1, n + 1):
            f[i] = any(f[j] and s[j:i] in words for
j in range(i))
        return f[n]

```

```

class Trie:
    def __init__(self):
        self.children: List[Trie | None] = [None] *
26

        self.isEnd = False

    def insert(self, w: str):
        node = self
        for c in w:
            idx = ord(c) - ord('a')
            if not node.children[idx]:

```

Round 5 · Mid · Medium

p.2126

```

        node.children[idx] = Trie()
        node = node.children[idx]
    node.isEnd = True

```

```
class Solution:
```

```

    def wordBreak(self, s: str, wordDict:
List[str]) -> bool:
        trie = Trie()
        for w in wordDict:
            trie.insert(w)

```

```

        n = len(s)
        f = [False] * (n + 1)
        f[n] = True
        for i in range(n - 1, -1, -1):
            node = trie
            for j in range(i, n):
                idx = ord(s[j]) - ord('a')
                if not node.children[idx]:
                    break
                node = node.children[idx]

```

```

            if node.isEnd and f[j + 1]:
                f[i] = True
                break

        return f[0]

```

• SimplyLeet

Round 5 · Mid · Medium

p.2127

Python solution for the Word Break problem

```

def wordBreak(s, wordDict):
    # Convert the list to a set for O(1) look-ups
    wordSet = set(wordDict)

    # dp[i] is True if s[0:i] can be segmented into
    words from the wordDict
    n = len(s)
    dp = [False] * (n + 1)

```

```

    # Base case: empty string is always segmentable
    dp[0] = True

    # Iterate through the string positions
    for i in range(1, n + 1):
        # Check every possible partition s[0:j] and
        s[j:i]
        for j in range(i):
            # If s[0:j] can be segmented and s[j:i]
            is in the dictionary
            if dp[j] and s[j:i] in wordSet:
                dp[i] = True # s[0:i] can be
                segmented

```

Round 5 · Mid · Medium

p.2128

```

                break # No need to check further
partitions
    return dp[n]

# Example usage:
s = "leetcode"
wordDict = ["leet", "code"]
print(wordBreak(s, wordDict)) # Output: True

```

• LC.ca

```

class Solution:
    def wordBreak(self, s: str, wordDict:
List[str]) -> bool:
        words = set(wordDict)
        n = len(s)

        # dp[j] meaning s from 0 to index=i-1~ is
breakable
        dp = [True] + [False] * n # so actually
size is n+1
        for i in range(1, n + 1):
            dp[i] = any( (dp[j] and s[j:i] in
words) for j in range(i) )
        return dp[n]

```

```

#####

```

```

class Solution:
    def wordBreak(self, s: str, wordDict:
Set[str]) -> bool:
        if not s or not wordDict:
            return False

        length = len(s)

```

```

        # construct dp, dp[i] meaning position i-1
can from dict or not
        dp[-1] meaning at last-index plus 1,
checking up to last-index
        dp = [False] * (length + 1)
        dp[0] = True # initiate

        for i in range(length + 1):
            if not dp[i]:
                continue
            for word in wordDict:
                if i + len(word) > length:

```

```

        continue
    if s[i:i + len(word)] == word:
        dp[i + len(word)] = True

    return dp[length]

#####

class Solution:
    def wordBreak(self, s: str, wordDict:
List[str]) -> bool:

        words = set(wordDict)
        n = len(s)
        dp = [False] * (n + 1)
        dp[0] = True # i'th char in string, is good
        for j in range(1, n + 1):
            for i in range(j): # starting at 0,
including dp[0]
                if dp[i] and s[i:j] in words:
                    dp[j] = True # j is exclusive,
meaning True until index i-1
                    break
        return dp[-1]

```

```

#####

class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.is_end = False

    def insert(self, w):

        node = self
        for c in w:
            idx = ord(c) - ord('a')
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]
        node.is_end = True

    def search(self, w):
        node = self

```

```

    for c in w:
        idx = ord(c) - ord('a')
        if node.children[idx] is None:
            return False
        node = node.children[idx]
    return node.is_end

class Solution:
    def wordBreak(self, s: str, wordDict:
List[str]) -> bool:
        # https://docs.python.org/3/library/functo
ols.html#functools.cache

        # creating a thin wrapper around a
dictionary lookup for the function arguments
        @cache
        def dfs(s):
            return not s or any(trie.search(s[:i])
and dfs(s[i:]) for i in range(1, len(s) + 1))

        trie = Trie()
        for w in wordDict:
            trie.insert(w)
        return dfs(s)

```

```

#####

class Solution(object):
    def wordBreak(self, s, wordDict):
        """
        :type s: str
        :type wordDict: Set[str]
        :rtype: bool
        """
        queue = [0]

        ls = len(s)
        lenList = [l for l in set(map(len, wordDict))]
        visited = [0 for _ in range(0, ls + 1)]
        while queue:
            start = queue.pop(0)
            for l in lenList:
                if s[start:start + l] in wordDict:
                    if start + l == ls:
                        return True
                    if visited[start + l] == 0:
                        queue.append(start + l)
                        visited[start + l] = 1
        return False

```

◆ Quick Review

Key Insights

- Use Dynamic Programming (DP) to break the problem into subproblems.
- Create a DP array where $dp[i]$ is true if the substring $s[0:i]$ can be segmented using words from the dictionary.
- Leverage a hash set for fast look-up of dictionary words.
- For every index i , check all possible previous break points j such that $dp[j]$ is true and $s[j:i]$ is in the wordDict.
- The problem can also be approached with recursion and memoization, but DP is more straightforward for iterative implementation.

Space & Time

Time Complexity: $O(n^2)$ in the worst-case scenario, where n is the length of s (each substring check is $O(1)$ if using a hash set, and there could be $O(n^2)$ checks overall).

Space Complexity: $O(n)$ for the dp array and $O(k)$ for the hash set where k is the number of words in the dictionary.

Solution

The solution uses a Dynamic Programming (DP) approach:

- Convert wordDict to a hash set for constant time lookup.
- Create a boolean dp array of size $n+1$ (n is the length of s), where $dp[i]$ indicates if $s[0:i]$ can be segmented.
- Set $dp[0]$ to true because an empty string is always segmentable.
- For each index i from 1 to n , iterate through all indices j from 0 to i : If $dp[j]$ is true and the substring $s[j:i]$ is in the dictionary, then set $dp[i]$ to true. Break early from the inner loop when a valid break is found.
- Return $dp[n]$ as the result.

This method efficiently builds up the solution by utilizing previous results to determine if the entire string can be segmented.

Trie

#208 Implement Trie (Prefix Tree)

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Design · Trie

Lists: Grind 169

Problem

Description

A trie (pronounced as "try") or prefix tree is a tree data structure used to efficiently store and retrieve keys in a dataset of strings. There are various applications of this data structure, such as autocomplete and spellchecker.

Implement the Trie class:

- `Trie()` Initializes the trie object.
- `void insert(String word)` Inserts the string word into the trie.
- `boolean search(String word)` Returns true if the string word is in the trie (i.e., was inserted before), and false otherwise.
- `boolean startsWith(String prefix)` Returns true if there is a previously inserted string word that has the prefix prefix, and false otherwise.

Example 1:

Input

```
["Trie", "insert", "search", "search", "startsWith",
"insert", "search"]
[[], ["apple"], ["apple"], ["app"], ["app"], ["app"],
["app"]]
```

Output

```
[null, null, true, false, true, null, true]
```

Explanation

```
Trie trie = new Trie();
```

Constraints:

- $1 \leq \text{word.length}, \text{prefix.length} \leq 2000$
- word and prefix consist only of lowercase English letters.
- At most $3 * 10^4$ calls in total will be made to insert, search, and startsWith.

Community Solutions (Python)

• WalkCC

```
class TrieNode:
    def __init__(self):
        self.children: dict[str, TrieNode] = {}
        self.isWord = False

class Trie:
    def __init__(self):
        self.root = TrieNode()
```

```
def insert(self, word: str) -> None:
    node: TrieNode = self.root
    for c in word:
        node = node.children.setdefault(c,
TrieNode())
        node.isWord = True

def search(self, word: str) -> bool:
    node = self._find(word)
    return node is not None and node.isWord
```

```
def startsWith(self, prefix: str) -> bool:
    return self._find(prefix) is not None

def _find(self, prefix: str) -> TrieNode | None:
    node: TrieNode = self.root
    for c in prefix:
        if c not in node.children:
            return None
        node = node.children[c]
    return node
```

• LeetDoocs

Round 5 · Mid · Medium

p.2139

```
class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.is_end = False

    def insert(self, word: str) -> None:
        node = self
        for c in word:
            idx = ord(c) - ord('a')
            if node.children[idx] is None:

                node.children[idx] = Trie()
                node = node.children[idx]
            node.is_end = True

    def search(self, word: str) -> bool:
        node = self._search_prefix(word)
        return node is not None and node.is_end

    def startsWith(self, prefix: str) -> bool:
        node = self._search_prefix(prefix)
```

Round 5 · Mid · Medium

p.2140

```

    return node is not None

def _search_prefix(self, prefix: str):
    node = self
    for c in prefix:
        idx = ord(c) - ord('a')
        if node.children[idx] is None:
            return None
        node = node.children[idx]
    return node

```

```

# Your Trie object will be instantiated and called
as such:
# obj = Trie()
# obj.insert(word)
# param_2 = obj.search(word)
# param_3 = obj.startsWith(prefix)

```

- SimplyLeet

```

class TrieNode:
    def __init__(self):
        # Dictionary to store child nodes and a
        # flag for end of word.
        self.children = {}
        self.is_end_of_word = False

class Trie:
    def __init__(self):
        # Initialize the Trie with the root node.
        self.root = TrieNode()

    def insert(self, word: str) -> None:
        # Start from the root.
        node = self.root
        for char in word:
            # If character doesn't exist, add a new
            # TrieNode.
            if char not in node.children:
                node.children[char] = TrieNode()
            node = node.children[char]
        # Mark the end of word.

```

```

node.is_end_of_word = True

def search(self, word: str) -> bool:
    # Traverse the Trie for each character.
    node = self.root
    for char in word:
        if char not in node.children:
            return False
        node = node.children[char]
    # Returns True if word ends exactly here.

    return node.is_end_of_word

def startsWith(self, prefix: str) -> bool:
    # Traverse checking for the prefix.
    node = self.root
    for char in prefix:
        if char not in node.children:
            return False
        node = node.children[char]
    return True

```

• LC.ca

```

import collections

class TrieNode:

    def __init__(self):
        ...

        Usually, a Python dictionary throws a KeyError
        if you try to get an item with a key that is not
        currently in the dictionary.
        The defaultdict in contrast will simply create
        any items that you try to access

        https://stackoverflow.com/questions/5900578/how-does-collections-defaultdict-work

        ...

        self.child = collections.defaultdict(TrieNode)
        self.is_word = False

class Trie:

    def __init__(self):
        self.root = TrieNode()

    def insert(self, word: str) -> None:

```

```

cur = self.root
for letter in word:
    cur = cur.child[letter]
cur.is_word = True

def search(self, word: str) -> bool:
    cur = self.root
    for letter in word:
        # cur = cur.child[letter] ==> will not work, it
        # will creat a default node for this letter
        cur = cur.child.get(letter)

    if not cur:
        return False
    return cur.is_word

def startsWith(self, prefix: str) -> bool:
    cur = self.root
    for letter in prefix:
        # cur = cur.child[letter] ==> will not work, it
        # will creat a default node for this letter
        cur = cur.child.get(letter)
    if not cur:

```

```

return False
return True

# Your Trie object will be instantiated and called
# as such:
# obj = Trie()
# obj.insert(word)
# param_2 = obj.search(word)
# param_3 = obj.startsWith(prefix)

#####

class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.is_end = False

    def insert(self, word: str) -> None:
        node = self
        for c in word:

```

```

        idx = ord(c) - ord('a')
        if node.children[idx] is None:
            node.children[idx] = Trie()
        node = node.children[idx]
    node.is_end = True

def search(self, word: str) -> bool:
    node = self._search_prefix(word)
    return node is not None and node.is_end

```

```

def startsWith(self, prefix: str) -> bool:
    node = self._search_prefix(prefix)
    return node is not None

```

```

def _search_prefix(self, prefix: str):
    node = self
    for c in prefix:
        idx = ord(c) - ord('a')
        if node.children[idx] is None:
            return None

```

```

        node = node.children[idx]
    return node

```

Your Trie object will be instantiated and called as such:

```

# obj = Trie()
# obj.insert(word)
# param_2 = obj.search(word)
# param_3 = obj.startsWith(prefix)

```

below is using reduce()

```

class Trie(object):

```

```

    def __init__(self):
        T = lambda: collections.defaultdict(T)
        self.root = T()

```

```

    def insert(self, word):

```

```

        reduce(dict.__getitem__, word,
self.root)['#'] = True

    def search(self, word):
        return '#' in reduce(lambda cur, c:
cur.get(c, {}), word, self.root)

    def startsWith(self, prefix):
        return bool(reduce(lambda cur, c:
cur.get(c, {}), prefix, self.root))

```

◆ Quick Review

Key Insights

- Use a tree-like structure where each node represents a character.
- Each node stores its children in a dictionary (or hash map) for $O(1)$ average time access.
- Mark the end of a word with a boolean flag.
- Insertion is done character by character, creating new nodes as needed.
- Searching and checking prefixes follow traversals down the tree.

Space & Time

Time Complexity:

- Insert: $O(m)$ where m is the length of the word.
- Search: $O(m)$
- startsWith: $O(m)$

Space Complexity:

- $O(N * M)$ in the worst-case scenario, where N is the number of words and M is the average length of the words, due to storing each character in the Trie.

Solution

We implement the Trie using a node-based approach. Each TrieNode contains a hash map (or dictionary) of its children and a boolean indicating if the node marks the end of a word. The Trie class uses these nodes to build out the words character by character. The insert method traverses and builds the necessary nodes; the search method verifies that all characters exist in order and that the final node marks the end of a word; the startsWith method checks if the prefix exists in the Trie. No tricky edge cases exist beyond handling null or empty inputs, as constraints guarantee valid lowercase strings.

Trie

#211 Design Add and Search Words Data Structure MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · DFS · Design · Trie

Lists: Grind 169

Problem

Description

Design a data structure that supports adding new words and finding if a string matches any previously added string.

Implement the WordDictionary class:

- WordDictionary() Initializes the object.
- void addWord(word) Adds word to the data structure, it can be matched later.
- bool search(word) Returns true if there is any string in the data structure that matches word or false otherwise. word may contain dots '.' where dots can be matched with any letter.

Example:

Input

```
["WordDictionary","addWord","addWord","addWord","search",
,"search","search","search"]
[[],["bad"],["dad"],["mad"],["pad"],["bad"],[".ad"],["b.."]]
```

Output

```
[null,null,null,null,false,true,true,true]
```

Explanation

```
WordDictionary wordDictionary = new WordDictionary();
```

Constraints:

- $1 \leq \text{word.length} \leq 25$
- word in addWord consists of lowercase English letters.
- word in search consist of '.' or lowercase English letters.
- There will be at most 2 dots in word for search queries.
- At most 104 calls will be made to addWord and search.

Community Solutions (Python)

- WalkCC


```

class TrieNode:
    def __init__(self):
        self.children: dict[str, TrieNode] = {}
        self.isWord = False

class WordDictionary:
    def __init__(self):
        self.root = TrieNode()

    def addWord(self, word: str) -> None:
        node: TrieNode = self.root
        for c in word:
            node = node.children.setdefault(c,
TrieNode())
            node.isWord = True

    def search(self, word: str) -> bool:
        return self._dfs(word, 0, self.root)

    def _dfs(self, word: str, s: int, node: TrieNode)
-> bool:

```

```

        if s == len(word):
            return node.isWord
        if word[s] != '.':
            child: TrieNode = node.children.get(word[s],
None)
            return self._dfs(word, s + 1, child) if child
        else False
            return any(self._dfs(word, s + 1, child) for
child in node.children.values())

```

• LeetCode

```

class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.is_end = False

class WordDictionary:
    def __init__(self):
        self.trie = Trie()

```

```
def addWord(self, word: str) -> None:
    node = self.trie
    for c in word:
        idx = ord(c) - ord('a')
        if node.children[idx] is None:
            node.children[idx] = Trie()
        node = node.children[idx]
    node.is_end = True
```

```
def search(self, word: str) -> bool:
```

```
    def search(word, node):
        for i in range(len(word)):
            c = word[i]
            idx = ord(c) - ord('a')
            if c != '.' and node.children[idx]
is None:
                return False
            if c == '.':
                for child in node.children:
                    if child is not None and
search(word[i + 1 :], child):
                        return True
```

```
        return False
        node = node.children[idx]
        return node.is_end
```

```
        return search(word, self.trie)
```

Your WordDictionary object will be instantiated and called as such:

```
# obj = WordDictionary()
# obj.addWord(word)
```

```
# param_2 = obj.search(word)
```

• SimplyLeet

```
class TrieNode:
    def __init__(self):
        # Dictionary to store child nodes
        self.children = {}
        # Flag to indicate if a word ends here
        self.is_end = False

class WordDictionary:
    def __init__(self):
        # Initialize the Trie with an empty root
        node
```

```

self.root = TrieNode()

def addWord(self, word: str) -> None:
    # Start at the root of the Trie
    node = self.root
    for char in word:
        # Create a new node if the character is
not present
        if char not in node.children:
            node.children[char] = TrieNode()
            node = node.children[char]

    # Mark the end of the word
    node.is_end = True

def search(self, word: str) -> bool:
    # DFS helper function
    def dfs(index, node):
        if index == len(word):
            return node.is_end
        char = word[index]
        if char == '.':

```

```

        # If wildcard, search all children
        nodes
        for child in
node.children.values():
            if dfs(index + 1, child):
                return True
            return False
        else:
            # If character not found, word does
not exist
            if char not in node.children:
                return False
            return dfs(index + 1,
node.children[char])

    return dfs(0, self.root)

# Example usage:
wd = WordDictionary()
wd.addWord("bad")
wd.addWord("dad")
wd.addWord("mad")
print(wd.search("pad")) # Output: False
print(wd.search("bad")) # Output: True
print(wd.search(".ad")) # Output: True

print(wd.search("b..")) # Output: True

```

- LC.ca

```
class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.is_end = False

class WordDictionary:
    def __init__(self):
        self.trie = Trie()

    def addWord(self, word: str) -> None:
        node = self.trie
        for c in word:
            idx = ord(c) - ord('a')
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]
        node.is_end = True

    def search(self, word: str) -> bool:
```

```
        def search(word, node):
            for i in range(len(word)):
                c = word[i]
                idx = ord(c) - ord('a')
                if c != '.' and node.children[idx]
is None:
                    return False
                if c == '.':
                    for child in node.children:
                        if child is not None and
search(word[i + 1 :], child):
                            return True

            return False
            node = node.children[idx]
            return node.is_end

        return search(word, self.trie)

# Your WordDictionary object will be instantiated
and called as such:
# obj = WordDictionary()
# obj.addWord(word)

# param_2 = obj.search(word)
```

◆ Quick Review

Key Insights

- Utilize a Trie (prefix tree) to efficiently store words.
- Each node in the Trie contains a dictionary for its children and a flag indicating the end of a word.
- The search operation is implemented with a depth-first search (DFS) approach.
- When the search query contains a dot, explore all possible child nodes.

Space & Time

Time Complexity:

- addWord: $O(w)$, where w is the length of the word.
- search: Worst-case $O(26^w)$ for words with multiple wildcards, but practical performance is better given constraints.

Space Complexity: $O(n * w)$, where n is the number of words and w is the average word length.

Solution

The WordDictionary is implemented using a Trie. For adding, we insert each character of the word into the Trie. For searching, we use DFS. When encountering a wildcard '.', we recursively search all children nodes. This structure combined with DFS efficiently addresses the wildcard matching and ensures good performance for multiple operations.

Trie

#616 Add Bold Tag in String

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Trie

Lists: Premium 98

Problem

Description

You are given a string s and an array of strings $words$. You should add a closed pair of bold tag $$ and $$ to wrap the substrings in s that exist in $words$.

- If two such substrings overlap, you should wrap them together with only one pair of closed bold-tag.
- If two substrings wrapped by bold tags are consecutive, you should combine them.

Return s after adding the bold tags.

Example 1:

Input: $s = \text{"abcxyz123"}$, $words = [\text{"abc"}, \text{"123"}]$
 Output: "abcxyz123"
 Explanation: The two strings of words are substrings of s as following: "abcxyz123".
 We add $$ before each substring and $$ after each substring.

Example 2:

Input: `s = "aaabbb", words = ["aa","b"]`

Output: `"<b>aaabbb</b>"`

Explanation:

"aa" appears as a substring two times: "aaabbb" and "aaabbb".

"b" appears as a substring three times: "aaabbb", "aaabbb", and "aaabbb".

We add `<b>` before each substring and `</b>` after each substring: `"<b>a<b>a</b>a</b><b>b</b><b>b</b><b>b</b>"`.

Since the first two `<b>`'s overlap, we merge them:

`"<b>aaa</b><b>b</b><b>b</b><b>b</b>"`.

Since now the four `<b>`'s are consecutive, we merge them: `"<b>aaabbb</b>"`.

Constraints:

- `1 <= s.length <= 1000`
- `0 <= words.length <= 100`
- `1 <= words[i].length <= 1000`
- `s` and `words[i]` consist of English letters and digits.
- All the values of `words` are unique.

Note: This question is the same as 758. Bold Words in String.

Community Solutions (Python)

- WalkCC

```
class Solution:
```

```
    def addBoldTag(self, s: str, words: list[str]) -> str:
```

```
        n = len(s)
```

```
        ans = []
```

```
        # bold[i] := True if s[i] should be bolded
```

```
        bold = [0] * n
```

```
        boldEnd = -1 # s[i:boldEnd] should be bolded
```

```
        for i in range(n):
```

```
            for word in words:
```

```
                if s[i:].startswith(word):
```

```
                    boldEnd = max(boldEnd, i + len(word))
```

```
            bold[i] = boldEnd > i
```

```
        # Construct the with bold tags
```

```
        i = 0
```

```
        while i < n:
```

```
            if bold[i]:
```

```
                j = i
```

```
                while j < n and bold[j]:
```

```

        j += 1
        # `s[i..j)` should be bolded.
        ans.append('<b>' + s[i:j] + '</b>')
        i = j
    else:
        ans.append(s[i])
        i += 1

    return ''.join(ans)

```

```

class TrieNode:
    def __init__(self):
        self.children: dict[str, TrieNode] = {}
        self.isWord = False

class Solution:
    def addBoldTag(self, s: str, words: list[str]) -> str:
        n = len(s)
        ans = []

```

```

        # bold[i] := True if s[i] should be bolded
        bold = [0] * n
        root = TrieNode()

    def insert(word: str) -> None:
        node = root
        for c in word:
            node = node.children.setdefault(c,
TrieNode())
            node.isWord = True

    def find(s: str, i: int) -> int:
        node = root
        ans = -1
        for j in range(i, len(s)):
            if s[j] not in node.children:
                node.children[s[j]] = TrieNode()
            node = node.children[s[j]]
            if node.isWord:
                ans = j
        return ans

```

```

for word in words:
    insert(word)

boldEnd = -1 # `s[i..boldEnd]` should be
bolded.
for i in range(n):
    boldEnd = max(boldEnd, find(s, i))
    bold[i] = boldEnd >= i

# Construct the with bold tags

i = 0
while i < n:
    if bold[i]:
        j = i
        while j < n and bold[j]:
            j += 1
        # `s[i..j]` should be bolded.
        ans.append('<b>' + s[i:j] + '</b>')
        i = j
    else:
        ans.append(s[i])
        i += 1

return ''.join(ans)

```

• LeetCode

Round 5 · Mid · Medium

p.2167

```

class Trie:
    def __init__(self):
        self.children = [None] * 128
        self.is_end = False

    def insert(self, word):
        node = self
        for c in word:
            idx = ord(c)
            if node.children[idx] is None:
                node.children[idx] = Trie()
                node = node.children[idx]
            node.is_end = True

class Solution:
    def addBoldTag(self, s: str, words: List[str])
-> str:
        trie = Trie()
        for w in words:
            trie.insert(w)

```

Round 5 · Mid · Medium

p.2168


```

n = len(s)
pairs = []
for i in range(n):
    node = trie
    for j in range(i, n):
        idx = ord(s[j])
        if node.children[idx] is None:
            break
        node = node.children[idx]
        if node.is_end:

```

```

            pairs.append([i, j])
    if not pairs:
        return s
    st, ed = pairs[0]
    t = []
    for a, b in pairs[1:]:
        if ed + 1 < a:
            t.append([st, ed])
            st, ed = a, b
        else:

```

```

        ed = max(ed, b)
    t.append([st, ed])

ans = []
i = j = 0
while i < n:
    if j == len(t):
        ans.append(s[i:])
        break
    st, ed = t[j]

```

```

    if i < st:
        ans.append(s[i:st])
    ans.append('<b>')
    ans.append(s[st : ed + 1])
    ans.append('</b>')
    j += 1
    i = ed + 1

return ''.join(ans)

```

- SimplyLeet

```
# Python Code Solution
def addBoldTag(s, words):
    n = len(s)
    # Mark positions to be bolded
    bold = [False] * n
    # Iterate each word and mark its occurrences
    for word in words:
        start = s.find(word)
        while start != -1:
            for i in range(start, start +
len(word)):
```

```
        bold[i] = True
        start = s.find(word, start + 1)
```

```
# Build the final result with bold tags
result = []
i = 0
while i < n:
    if not bold[i]:
        result.append(s[i])
        i += 1
```

```
        else:
            result.append("<b>")
            while i < n and bold[i]:
                result.append(s[i])
                i += 1
            result.append("</b>")
    return "".join(result)
```

```
# Example usage:
# s = "abcxyz123"
```

```
# words = ["abc", "123"]
# print(addBoldTag(s, words)) #
"<b>abc</b>xyz<b>123</b>"
```

```
# Line-by-line:
# 1. Initialize an array "bold" of length n to
track which characters need bold formatting.
# 2. For each word, use string.find to locate every
occurrence and mark the corresponding indices.
# 3. Traverse the string s. If the current
character is marked bold and it is the start of a
bold segment,
# insert the opening tag before appending
characters until the end of that bold segment.
# 4. Append the closing tag after the bold segment
and continue the process.
```

• LC.ca

```

class Trie:
    def __init__(self):
        self.children = [None] * 128
        self.is_end = False

    def insert(self, word):
        node = self
        for c in word:
            idx = ord(c)
            if node.children[idx] is None:

                node.children[idx] = Trie()
                node = node.children[idx]
            node.is_end = True

class Solution:
    def addBoldTag(self, s: str, words: List[str])
-> str:
    trie = Trie()
    for w in words:
        trie.insert(w)

```

Round 5 · Mid · Medium

p.2173

```

n = len(s)
pairs = []
for i in range(n):
    node = trie
    for j in range(i, n):
        idx = ord(s[j])
        if node.children[idx] is None:
            break
        node = node.children[idx]
        if node.is_end:
            pairs.append([i, j])
if not pairs:
    return s
st, ed = pairs[0]
t = []
for a, b in pairs[1:]:
    if ed + 1 < a:
        t.append([st, ed])
        st, ed = a, b
    else:

```

Round 5 · Mid · Medium

p.2174

```

        ed = max(ed, b)
    t.append([st, ed])

```

```

ans = []
i = j = 0
while i < n:
    if j == len(t):
        ans.append(s[i:])
        break
    st, ed = t[j]

```

```

    if i < st:
        ans.append(s[i:st])
    ans.append('<b>')
    ans.append(s[st : ed + 1])
    ans.append('</b>')
    j += 1
    i = ed + 1

return ''.join(ans)

```

◆ Quick Review

Key Insights

- We need to identify every occurrence of each word in s.
- Overlapping or consecutive intervals should be merged to form one continuous bold region.
- Use an auxiliary data structure (e.g., a boolean array) to mark positions that should be bolded.

- Once the positions are marked, build the final string by inserting bold tags at the appropriate boundaries.

Space & Time

Time Complexity: $O(n * m * k)$ where n is the length of s , m is the number of words, and k is the average length of the words (due to substring matching).

Space Complexity: $O(n)$ for the auxiliary boolean array used to mark bold positions.

Solution

The solution involves two main steps:

- Identify and mark the indices in the string s that are part of any matching word. We do this by iterating over each word and marking every occurrence's start and end in a boolean array.
- Construct the final string by traversing the boolean array. When encountering a segment of consecutive true values, wrap that segment with and tags. The key idea is to merge intervals by marking every index that should be bold and then using a single pass to insert tags only when transitioning from non-bold to bold and vice versa.

Trie

#692 Top K Frequent Words

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Trie · Sorting · Heap

Lists: Grind 169

Problem

Description

Given an array of strings `words` and an integer `k`, return the `k` most frequent strings.

Return the answer sorted by the frequency from highest to lowest. Sort the words with the same frequency by their lexicographical order.

Example 1:

```

Input: words =
["i","love","leetcode","i","love","coding"], k = 2
Output: ["i","love"]
Explanation: "i" and "love" are the two most frequent words.
Note that "i" comes before "love" due to a lower alphabetical order.

```

Example 2:

```

Input: words = ["the","day","is","sunny","the","the","t
he","sunny","is","is"], k = 4
Output: ["the","is","sunny","day"]
Explanation: "the", "is", "sunny" and "day" are the
four most frequent words, with the number of occurrence
being 4, 3, 2 and 1 respectively.

```

Constraints:

- $1 \leq \text{words.length} \leq 500$
- $1 \leq \text{words}[i].\text{length} \leq 10$
- `words[i]` consists of lowercase English letters.
- `k` is in the range `[1, The number of unique words[i]]`

Follow-up: Could you solve it in $O(n \log(k))$ time and $O(n)$ extra space?

Community Solutions (Python)

• WalkCC

```

class Solution:
    def topKFrequent(self, words: list[str], k: int)
-> list[str]:
        ans = []
        bucket = [[] for _ in range(len(words) + 1)]

        for word, freq in
collections.Counter(words).items():
            bucket[freq].append(word)

        for b in reversed(bucket):
            for word in sorted(b):

```

```
ans.append(word)
if len(ans) == k:
    return ans
```

```
from dataclasses import dataclass
```

```
@dataclass(frozen=True)
```

```
class T:
    word: str
    freq: int
```

```
def __lt__(self, other):
    if self.freq == other.freq:
```

```
        return self.word > other.word
    return self.freq < other.freq
```

```
class Solution:
    def topKFrequent(self, words: list[str], k: int)
    -> list[str]:
        ans = []
        heap = []

        for word, freq in
collections.Counter(words).items():
```

```
heapq.heappush(heap, T(word, freq))
if len(heap) > k:
    heapq.heappop(heap)
```

```
while heap:
    ans.append(heapq.heappop(heap).word)
```

```
return ans[::-1]
```

• LeetCode

```
class Solution:
    def topKFrequent(self, words: List[str], k:
int) -> List[str]:
        cnt = Counter(words)
        return sorted(cnt, key=lambda x: (-cnt[x],
x))[:k]
```

• SimplyLeet

```
import heapq
from collections import Counter

def topKFrequent(words, k):
    # Count the frequency of each word
    count = Counter(words)

    # Initialize a min-heap; Python's heapq is a
    min-heap.
    # The comparator is (frequency, negative word)
    but use word in reverse order:
    # Since we want words with lower alphabetical
    order to be prioritized in the result,
```

```
    # and in the min heap if frequencies equal, we
    want the one with lexicographically larger word to
    be popped.
    heap = []

    for word, freq in count.items():
        # Push each pair into the heap. The tuple
        ordering:
        # First element is frequency and second
        element is the word (with reverse lex order trick)
        heapq.heappush(heap, (freq, -ord(word[0]),
        word))
        # However, using -ord(word[0]) is not
        robust for multi-letter words.
        # Instead, we can use tuple (freq, word)
        with a custom comparator by ensuring we invert the
        word order by using word in reverse.
        # For simplicity, we push tuple (freq,
        word) and then pop if heap size > k with a custom
        condition.
```

```
# Reset heap and use a custom comparator
implemented by the tuple (-freq, word) to simulate
max heap behavior.
heap = []
for word, freq in count.items():
    # We push with frequency as negative so
    that larger frequencies come first;
    # use word as is to satisfy lexicographical
    order when frequencies are equal.
    heapq.heappush(heap, (-freq, word))

# Extract k elements from the heap and sort
them accordingly.
result = []
```

```
for _ in range(k):
    freq, word = heapq.heappop(heap)
    result.append(word)
return result

# Example usage:
print(topKFrequent(["i","love","leetcode","i","love",
"coding"], 2)) # Output: ['i', 'love']
```

• LC.ca

Round 5 · Mid · Medium

p.2183

```
...
>>> sorted(student_objects,
key=attrgetter('grade', 'age'))
[('john', 'A', 15), ('dave', 'B', 10), ('jane',
'B', 12)]

ref: https://docs.python.org/3/howto/sorting.html#operator-module-functions
'''
class Solution:
    def topKFrequent(self, words: List[str], k:
int) -> List[str]:
        cnt = Counter(words)
        # multiple comparators "lambda x:
(-cnt[x], x)"

        return sorted(cnt, key=lambda x: (-cnt[x],
x))[:k]
        # why the returned sorted() not a tuple
        (word, count), but only word?
        # it's the property of Counter() class
        # so, also pass 0J: return
        sorted(cnt.keys(), key=lambda x: (-cnt[x], x))[:k]

...
>>> words =
["i","love","leetcode","i","love","coding"]
>>> k = 2
>>> cnt = Counter(words)
>>> cnt
```

Round 5 · Mid · Medium

p.2184


```
Counter({'i': 2, 'love': 2, 'leetcode': 1,
'coding': 1})
```

```
>>> sorted(cnt, key=lambda x: (-cnt[x], x))
['i', 'love', 'coding', 'leetcode']
```

```
# default, only for keys()
```

```
>>> sorted(cnt)
['coding', 'i', 'leetcode', 'love']
>>> sorted(cnt.keys())
['coding', 'i', 'leetcode', 'love']
```

```
>>> sorted(cnt.values())
[1, 1, 2, 2]
>>> sorted(cnt.items())
[('coding', 1), ('i', 2), ('leetcode', 1), ('love',
2)]
'''
```

```
#####
```

```
'''
```

```
>>> words = ["the","day","is","sunny","the","the",
"the","sunny","is","is"]
>>> count = collections.Counter(words)
>>> count
Counter({'the': 4, 'is': 3, 'sunny': 2, 'day': 1})
>>>
>>> count.items()
dict_items([('the', 4), ('day', 1), ('is', 3),
('sunny', 2)])
>>>
'''
```

```
'''
cmp doesn't exist in Python 3. If you really want
it, you could define it yourself:
```

```
def cmp(a, b):
    return (a > b) - (a < b)
```

```
'''
```

```
class Solution(object): #python-2
    def topKFrequent(self, words, k):
        """
```

```

:type words: List[str]
:type k: int
:rtype: List[str]
"""
count = collections.Counter(words)

# https://docs.python.org/3/howto/sorting.html#comparison-functions
def compare(x, y):
    if x[1] == y[1]:
        return cmp(x[0], y[0])

    else:
        return -cmp(x[1], y[1]) # -cmp, so
        reversed order, bigger ones at front
    return [x[0] for x in
sorted(count.items(), cmp = compare)[:k]]

```

◆ Quick Review

Key Insights

- Count the frequency of each unique word using a hash table or dictionary.
- Use a min-heap (priority queue) to keep track of the top k frequent words. In the heap, the ordering is based on frequency; if frequencies are equal, order by lexicographical order (inverted condition for min-heap).
- Alternatively, you can sort the unique words based on frequency and lexicographical order and then pick the top k.

- For an efficient solution, use a min-heap that maintains size k achieving $O(n \log k)$ time complexity.

Space & Time

Time Complexity: $O(n \log k)$ when using a heap, where n is the number of unique words.

Space Complexity: $O(n)$ for storing the frequency counts and the heap.

Solution

We first calculate the frequency of each word using a hash table (or dictionary). Then, we use a min-heap that stores pairs of (frequency, word). The comparator for the heap is defined such that the root of the heap is the word with the smallest frequency; if two words have the same frequency, the word with the lexicographically larger value is considered “smaller” in our ordering so that it gets popped when the heap size exceeds k. After processing all words, the heap contains the top k frequent words. Finally, the heap elements are extracted and reversed (if needed) to produce the final list sorted from highest to lowest frequency with lexicographic tie-breaking.

Trie

#720 Longest Word in Dictionary

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Trie · Sorting

Lists: Denny Zhang

Problem

Description

Given an array of strings words representing an English Dictionary, return the longest word in words that can be built one character at a time by other words in words.

If there is more than one possible answer, return the longest word with the smallest lexicographical order. If there is no answer, return the empty string.

Note that the word should be built from left to right with each additional character being added to the end of a previous word.

Example 1:

Input: words = ["w","wo","wor","worl","world"]

Output: "world"

Explanation: The word "world" can be built one character at a time by "w", "wo", "wor", and "worl".

Example 2:

Input: words =

["a","banana","app","appl","ap","apply","apple"]

Output: "apple"

Explanation: Both "apply" and "apple" can be built from other words in the dictionary. However, "apple" is lexicographically smaller than "apply".

Constraints:

- $1 \leq \text{words.length} \leq 1000$
- $1 \leq \text{words}[i].\text{length} \leq 30$
- words[i] consists of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestWord(self, words: list[str]) -> str:
        root = {}

        for word in words:
            node = root
            for c in word:
                if c not in node:
                    node[c] = {}
                node = node[c]
```

```

node['word'] = word

def dfs(node: dict) -> str:
    ans = node['word'] if 'word' in node else ''

    for child in node:
        if 'word' in node[child] and
len(node[child]['word']) > 0:
            childWord = dfs(node[child])
            if len(childWord) > len(ans) or (
                len(childWord) == len(ans) and
childWord < ans):

                ans = childWord

    return ans

return dfs(root)

```

• LeetCode

```

class Trie:
    def __init__(self):
        self.children: List[Optional[Trie]] =
[None] * 26
        self.is_end = False

    def insert(self, w: str):
        node = self
        for c in w:
            idx = ord(c) - ord("a")
            if node.children[idx] is None:

                node.children[idx] = Trie()
                node = node.children[idx]
            node.is_end = True

    def search(self, w: str) -> bool:
        node = self
        for c in w:
            idx = ord(c) - ord("a")
            if node.children[idx] is None:
                return False

```

```

        node = node.children[idx]
        if not node.is_end:
            return False
    return True

```

```

class Solution:
    def longestWord(self, words: List[str]) ->
    str:
        trie = Trie()
        for w in words:

```

```

            trie.insert(w)
            ans = ""
            for w in words:
                if trie.search(w) and (
                    len(ans) < len(w) or (len(ans) ==
len(w) and ans > w)
                ):
                    ans = w
            return ans

```

• SimplyLeet

Round 5 · Mid · Medium

p.2193

Python solution for Longest Word in Dictionary

```

def longestWord(words):
    # Sort words by length and lexicographical
    order

```

```

        words.sort(key=lambda x: (len(x), x))

```

```

        valid_words = set()
        result = ""

```

Process each word

```

        for word in words:
            # For single-letter words, they are valid
            if they exist in the list
            if len(word) == 1 or word[:-1] in
valid_words:
                valid_words.add(word)
                # Update result if the current word is
                longer or if equal length but lexicographically
                smaller
                if len(word) > len(result) or
(len(word) == len(result) and word < result):
                    result = word

```

```

        return result

```

```

# Example usage:
# print(longestWord(["a", "banana", "app", "appl",
"ap", "apply", "apple"]))

```

Round 5 · Mid · Medium

p.2194

• LC.ca

```

class Solution:
    def longestWord(self, words: List[str]) ->
    str:
        cnt, ans = 0, ''
        s = set(words)
        for w in s:
            n = len(w)
            if all(w[:i] in s for i in range(1,
n)):
                if cnt < n:
                    cnt, ans = n, w
                elif cnt == n and w < ans:
                    ans = w
        return ans

```

◆ Quick Review

Key Insights

- The word must be built character by character; hence every prefix (removing one character from the end at a time) must exist in the dictionary.
- Sorting the input list can help in ensuring that smaller words (prefixes) are processed first.
- Using a hash set to store valid words enables rapid prefix checks.
- Lexicographical order plays a role when words have the same length.

Space & Time

Time Complexity: $O(n * L + n \log n)$ where n is the number of words and L is the maximum length of a word ($n * L$ for prefix checking, and $n \log n$ for sorting).

Space Complexity: $O(n * L)$ for storing words in the set.

Solution

The approach begins by sorting the list of words. Sorting is done primarily by word length and secondarily by lexicographical order so that when we iterate through the words, all possible prefixes for a word have already been considered. Initialize a set to keep track of valid words that can be built one character at a time. For each word in the sorted list, check if all its prefixes (word without the last character progressively) exist in the set. If a word qualifies, add it to the set and check if it should be the new result based on length and lexicographical order. The set lookup ensures that checking each prefix is efficient.

Trie

#1166 Design File System

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Design · Trie

Lists: Denny Zhang

Problem

Description

You are asked to design a file system that allows you to create new paths and associate them with different values.

The format of a path is one or more concatenated strings of the form: / followed by one or more lowercase English letters. For example, "/leetcode" and "/leetcode/problems" are valid paths while an empty string "" and "/" are not.

Implement the FileSystem class:

- `bool createPath(string path, int value)` Creates a new path and associates a value to it if possible and returns true. Returns false if the path already exists or its parent path doesn't exist.
- `int get(string path)` Returns the value associated with path or returns -1 if the path doesn't exist.

Example 1:

```
Input:
["FileSystem","createPath","get"]
[[],["/a",1],["/a"]]
Output:
[null,true,1]
Explanation:
FileSystem fileSystem = new FileSystem();
```

Example 2:

```
Input:
["FileSystem","createPath","createPath","get","createPa
th","get"]
[[],["/leet",1],["/leet/code",2],["/leet/code"],["/c/d"
,1],["/c"]]
Output:
[null,true,true,2,false,-1]
Explanation:
FileSystem fileSystem = new FileSystem();
```

Constraints:

- $2 \leq \text{path.length} \leq 100$
- $1 \leq \text{value} \leq 109$
- Each path is valid and consists of lowercase English letters and '/'.
 - At most 104 calls in total will be made to createPath and get.

Community Solutions (Python)

- WalkCC

```

class TrieNode:
    def __init__(self, value: int = 0):
        self.children: dict[str, TrieNode] = {}
        self.value = value

class FileSystem:
    def __init__(self):
        self.root = TrieNode()

    def createPath(self, path: str, value: int) -> bool:
        node: TrieNode = self.root
        subpaths = path.split('/')

        for i in range(1, len(subpaths) - 1):
            if subpaths[i] not in node.children:
                return False
            node = node.children[subpaths[i]]

        if subpaths[-1] in node.children:

```

```

        return False
        node.children[subpaths[-1]] = TrieNode(value)
        return True

    def get(self, path: str) -> int:
        node: TrieNode = self.root
        for subpath in path.split('/')[1:]:
            if subpath not in node.children:
                return -1
            node = node.children[subpath]

        return node.value

```

• LeetCode

```

class Trie:
    def __init__(self, v: int = -1):
        self.children = {}
        self.v = v

    def insert(self, w: str, v: int) -> bool:
        node = self
        ps = w.split("/")
        for p in ps[1:-1]:
            if p not in node.children:

```



```

        return False
        node = node.children[p]
    if ps[-1] in node.children:
        return False
    node.children[ps[-1]] = Trie(v)
    return True

```

```

def search(self, w: str) -> int:
    node = self
    for p in w.split("/")[1:]:

```

```

        if p not in node.children:
            return -1
        node = node.children[p]
    return node.v

```

```

class FileSystem:
    def __init__(self):
        self.trie = Trie()

```

```

    def createPath(self, path: str, value: int) ->
bool:
        return self.trie.insert(path, value)

    def get(self, path: str) -> int:
        return self.trie.search(path)

```

Your FileSystem object will be instantiated and called as such:

```

# obj = FileSystem()
# param_1 = obj.createPath(path,value)

```

```

# param_2 = obj.get(path)

```

• SimplyLeet

```

class FileSystem:
    def __init__(self):
        # Initialize hash map to store paths and
        their values.
        self.paths = {}

    def createPath(self, path: str, value: int) ->
bool:
        # Check if the path already exists.
        if path in self.paths:
            return False

```

```

    # Find the index of the last '/' to get the
    parent path.
    last_slash = path.rfind('/')
    # If last_slash is 0, then parent is root
    level, which is invalid because "" is not a valid
    path.

```

```

    if last_slash == 0:
        parent = ""
    else:
        parent = path[:last_slash]

```

```

    # For non-root paths, the parent must
    exist.
    if parent and parent not in self.paths:

```

```

        return False

```

```

    # Store the new path with its value.
    self.paths[path] = value
    return True

```

```

    def get(self, path: str) -> int:
        # Return the value associated with the
        path, or -1 if it doesn't exist.
        return self.paths.get(path, -1)

```

```

# Example usage:
# fs = FileSystem()
# print(fs.createPath("/a", 1)) # True
# print(fs.get("/a")) # 1

```

• LC.ca

```

'''
>>> "/a/b/c".split('/')
['', 'a', 'b', 'c']
'''

```

```

class Trie:
    def __init__(self, v: int = -1):
        self.children = {}
        self.v = v

```

```

    def insert(self, w: str, v: int) -> bool:
        node = self
        ps = w.split("/")
        for p in ps[1:-1]:
            if p not in node.children:
                return False
            node = node.children[p]
        if ps[-1] in node.children:
            return False
        node.children[ps[-1]] = Trie(v)

```

```
return True
```

```
def search(self, w: str) -> int:
    node = self
    for p in w.split("/")[1:]:
        if p not in node.children:
            return -1
        node = node.children[p]
    return node.v
```

```
class FileSystem:
    def __init__(self):
        self.trie = Trie()

    def createPath(self, path: str, value: int) -> bool:
        return self.trie.insert(path, value)

    def get(self, path: str) -> int:
        return self.trie.search(path)
```

```
# Your FileSystem object will be instantiated and
called as such:
# obj = FileSystem()
# param_1 = obj.createPath(path,value)
# param_2 = obj.get(path)
```

◆ Quick Review

Key Insights

- Use a data structure (like a hash map) to store paths and associated values.
- To create a new path, check that the path does not already exist.
- Ensure the immediate parent path exists before creation.
- When retrieving a value, simply look it up in the storage structure.
- Splitting the path by "/" allows easy identification of the parent.

Space & Time

Time Complexity:

- createPath: $O(n)$ per call where n is the number of segments in the path (to process and validate parent path).
- get: $O(1)$ per call if using a hash map.

Space Complexity: $O(m * L)$ where m is the number of paths stored and L is the average length of a path.

Solution

The solution uses a hash map (dictionary) to store file paths along with their associated values. For the createPath method, the algorithm:

- Checks if the path already exists.
 - Extracts the parent path by removing the last segment.
 - Validates the existence of the parent path.
 - Inserts the new path with its corresponding value if validation passes.
- For the get method, the algorithm simply returns the stored value if

the path exists; otherwise, it returns `-1`. This approach ensures efficient lookup and validation using string manipulation and hash map operations.

Backtracking

Round 5 · Mid · Medium · 6 questions

Round 5 · Mid · Medium

p.2207

Round 5 · Mid · Medium

p.2208

Backtracking

#17 Letter Combinations of a Phone Number MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Backtracking

Lists: Grind 169

Problem

Description

Given a string containing digits from 2–9 inclusive, return all possible letter combinations that the number could represent. Return the answer in any order.

A mapping of digits to letters (just like on the telephone buttons) is given below. Note that 1 does not map to any letters.



Example 1:

Input: digits = "23"

Output: ["ad","ae","af","bd","be","bf","cd","ce","cf"]

Example 2:

Input: digits = "2"

Output: ["a","b","c"]

Constraints:

- $1 \leq \text{digits.length} \leq 4$
- $\text{digits}[i]$ is a digit in the range ['2', '9'].

Community Solutions (Python)

• WalkCC

```
class Solution:
    def letterCombinations(self, digits: str) -> List[str]:
        if not digits:
            return []

        digitToLetters = ['', '', 'abc', 'def', 'ghi',
                          'jkl', 'mno', 'pqrs', 'tuv',
                          'wxyz']
        ans = []

        def dfs(i: int, path: List[str]) -> None:

            if i == len(digits):
                ans.append(''.join(path))
                return

            for letter in
digitToLetters[int(digits[i])]:
                path.append(letter)
                dfs(i + 1, path)
                path.pop()

        dfs(0, [])

        return ans
```

• LeetCode

```
class Solution:
    def letterCombinations(self, digits: str) -> List[str]:
        if not digits:
            return []
        d = ["abc", "def", "ghi", "jkl", "mno",
            "pqrs", "tuv", "wxyz"]
        ans = [""]
        for i in digits:
            s = d[int(i) - 2]
            ans = [a + b for a in ans for b in s]
        return ans
```

```
class Solution:
    def letterCombinations(self, digits: str) -> List[str]:
        def dfs(i: int):
            if i >= len(digits):
                ans.append(''.join(t))
                return
            for c in d[int(digits[i]) - 2]:
                t.append(c)
                dfs(i + 1)
                t.pop()
```

```

    if not digits:
        return []
    d = ["abc", "def", "ghi", "jkl", "mno",
"pqrs", "tuv", "wxyz"]
    ans = []
    t = []
    dfs(0)
    return ans

```

• SimplyLeet

```

# Python solution for Letter Combinations of a
Phone Number

def letterCombinations(digits):
    # Mapping of digits to corresponding letters
    digit_to_letters = {
        "2": "abc",
        "3": "def",
        "4": "ghi",
        "5": "jkl",
        "6": "mno",

```

```

        "7": "pqrs",
        "8": "tuv",
        "9": "wxyz"
    }

```

```

    # If input is empty, return an empty list
    if not digits:
        return []

    result = []

```

```

    # Backtracking function to generate
combinations
    def backtrack(index, current_combination):
        # If the current combination length equals
the digits length,
        # add the combination to the result list.
        if index == len(digits):
            result.append(current_combination)
            return

        # Get the corresponding letters for the
current digit

```

```

        letters = digit_to_letters[digits[index]]
        # Explore all possible letter routes for
the current digit.
        for letter in letters:
            backtrack(index + 1,
current_combination + letter)

        # Initiate backtracking starting from index 0
and empty combination.
        backtrack(0, "")
        return result

# Example usage:

#print(letterCombinations("23"))

```

• LC.ca

```

class Solution:
    def letterCombinations(self, digits: str) ->
List[str]:
        if not digits:
            return []
        d = ["abc", "def", "ghi", "jkl", "mno",
"pqrs", "tuv", "wxyz"]
        ans = [""]
        for i in digits:
            s = d[int(i) - 2]
            ans = [a + b for a in ans for b in s]
        return ans

```

◆ Quick Review

Key Insights

- Use a mapping from digits to the respective letters as defined by a telephone keypad.
- Backtracking is an ideal approach to generate all possible combinations.
- Each digit adds a level in the recursive tree where you append one letter at a time.
- Handling edge cases such as an empty input string is crucial.

Space & Time

Time Complexity: $O(4^n * n)$ where n is the length of the input string (each digit may map to up to 4 letters and adding a combination takes $O(n)$ time).

Space Complexity: $O(n)$ for the recursion call stack, not counting the space required for storing the output.

Solution

We can solve the problem using a backtracking algorithm. First, we define a dictionary mapping each digit to its corresponding letters. Then we initiate a recursive function called `backtrack` that takes the current combination and the index of the next digit to process. At each call, if the combination is complete (length equal to the input string), we add it to the result list. Otherwise, we iterate through the letters corresponding to the current digit, append a letter to the combination, and recursively call `backtrack` with the next index. This approach ensures that all possible letter combinations are generated.

Backtracking

#22 Generate Parentheses

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming · Backtracking

Lists: Grind 169

Problem

Description

Given n pairs of parentheses, write a function to generate all combinations of well-formed parentheses.

Example 1:

Input: $n = 3$

Output: ["((()))", "(()())", "(())()", "()(())", "()()()"]

Example 2:

Input: $n = 1$

Output: ["()"]

Constraints:

- $1 \leq n \leq 8$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def generateParenthesis(self, n):
        ans = []

        def dfs(l: int, r: int, s: list[str]) -> None:
            if l == 0 and r == 0:
                ans.append(''.join(s))
            if l > 0:
                s.append('(')
                dfs(l - 1, r, s)

            s.pop()
            if l < r:
                s.append(')')
                dfs(l, r - 1, s)
                s.pop()

        dfs(n, n, [])
        return ans
```

- LeetDoocs

```
class Solution:
    def generateParenthesis(self, n: int) ->
List[str]:
    def dfs(l: int, r: int, t: str):
        if l > n or r > n or l < r:
            return
        if l == n and r == n:
            ans.append(t)
            return
        dfs(l + 1, r, t + "(")
        dfs(l, r + 1, t + ")")
```

```
ans = []
dfs(0, 0, "")
return ans
```

• SimplyLeet

```
def generateParenthesis(n):
    # List to store valid combinations
    result = []

    # Helper function for backtracking
    def backtrack(current, open_count,
close_count):
        # If current string reaches the maximum
length, it's a valid combination
        if len(current) == 2 * n:
            result.append(current)
            return
```

```
        # If we can add an opening parenthesis, do
it
        if open_count < n:
            backtrack(current + '(', open_count +
1, close_count)
        # If adding a closing parenthesis maintains
the validity, do it
        if close_count < open_count:
            backtrack(current + ')', open_count,
close_count + 1)

        # Start the recursion with an empty string and
zero counts for both
        backtrack("", 0, 0)
        return result
```

```
# Example usage:
print(generateParenthesis(3))
```

• LC.ca

```
class Solution:
    def generateParenthesis(self, n: int) ->
List[str]:
    def dfs(l, r, t):
        if l > n or r > n or l < r:
            return
        if l == n and r == n:
            ans.append(t)
            return
        dfs(l + 1, r, t + '(')
        dfs(l, r + 1, t + ')')
```

```
ans = []
dfs(0, 0, '')
return ans
```

```
#####
```

```
class Solution(object):
    def generateParenthesis(self, n):
        """"
```

```
:type n: int
:rtype: List[str]
"""

def dfs(left, path, res, n):
    if len(path) == 2 * n:
        if left == 0:
            res.append("".join(path))
        return
```

```
if left < n:
    path.append("(")
    dfs(left + 1, path, res, n)
    path.pop()
if left > 0:
    path.append(")")
    dfs(left - 1, path, res, n)
    path.pop()
```

```
res = []
```

```
dfs(0, [], res, n)
return res
```

◆ Quick Review

Key Insights

- Use a backtracking approach to build the valid combinations step by step.

- At any recursion step, you can add an opening parenthesis if you still have one available.
- You can add a closing parenthesis only if it would not result in an invalid sequence (i.e., the number of closing parentheses is less than the number of opening ones).
- The solution leverages recursion to explore all potential valid sequences while pruning invalid paths early.

Space & Time

Time Complexity: $O(4^n / (n^{3/2}))$

Space Complexity: $O(n)$ for the recursion stack, plus space for storing all valid combinations.

Solution

The solution uses a recursive backtracking technique to generate the combinations. The algorithm maintains counters for the number of '(' and ')' added to the current string. It recursively adds a '(' if the count is less than n and adds a ')' if doing so does not make the string invalid (i.e., if the count of ')' is less than the count of '('). When a valid combination reaches a length of $2*n$, it is added to the result list.

Backtracking

#39 Combination Sum

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Backtracking

Lists: Grind 169

Problem

Description

Given an array of distinct integers candidates and a target integer target, return a list of all unique combinations of candidates where the chosen numbers sum to target. You may return the combinations in any order.

The same number may be chosen from candidates an unlimited number of times. Two combinations are unique if the frequency of at least one of the chosen numbers is different.

The test cases are generated such that the number of unique combinations that sum up to target is less than 150 combinations for the given input.

Example 1:

Input: candidates = [2,3,6,7], target = 7

Output: [[2,2,3],[7]]

Explanation:

2 and 3 are candidates, and $2 + 2 + 3 = 7$. Note that 2 can be used multiple times.

7 is a candidate, and $7 = 7$.

These are the only two combinations.

Example 2:

Input: candidates = [2,3,5], target = 8

Output: [[2,2,2,2],[2,3,3],[3,5]]

Example 3:

Input: candidates = [2], target = 1

Output: []

Constraints:

- $1 \leq \text{candidates.length} \leq 30$
- $2 \leq \text{candidates}[i] \leq 40$
- All elements of candidates are distinct.
- $1 \leq \text{target} \leq 40$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def combinationSum(self, candidates: list[int],
                      target: int) ->
list[list[int]]:
    ans = []

    def dfs(s: int, target: int, path: list[int])
-> None:
        if target < 0:
            return
        if target == 0:
            ans.append(path.clone())

        return

        for i in range(s, len(candidates)):
            path.append(candidates[i])
            dfs(i, target - candidates[i], path)
            path.pop()

    candidates.sort()
    dfs(0, target, [])
    return ans
```

- LeetCode

```

class Solution:
    def combinationSum(self, candidates:
List[int], target: int) -> List[List[int]]:
        def dfs(i: int, s: int):
            if s == 0:
                ans.append(t[:])
                return
            if s < candidates[i]:
                return
            for j in range(i, len(candidates)):
                t.append(candidates[j])

                dfs(j, s - candidates[j])
                t.pop()

        candidates.sort()
        t = []
        ans = []
        dfs(0, target)
        return ans

```

• SimplyLeet

Round 5 · Mid · Medium

p.2227

```

# Function to find all unique combinations that sum
to the target
def combinationSum(candidates, target):
    # List to store all valid combinations
    result = []
    # Sort candidates to facilitate pruning of the
search
    candidates.sort()

    # Recursive helper function for backtracking
    def backtrack(remaining, start, path):
        # If remaining is 0, path is a valid
combination so add a copy of it to result

        if remaining == 0:
            result.append(list(path))
            return
        # Loop through candidates starting from the
current index
        for i in range(start, len(candidates)):
            # If candidate is greater than
remaining, break out for efficiency
            if candidates[i] > remaining:
                break
            # Include candidate in the current
combination path
            path.append(candidates[i])

```

Round 5 · Mid · Medium

p.2228

```

        # Recursively call backtrack with
        updated remaining and same index (allowing reuse)
        backtrack(remaining - candidates[i],
i, path)
        # Backtrack by removing the candidate
        added in this iteration
        path.pop()

        # Start backtracking from index 0 with an empty
        path
        backtrack(target, 0, [])
        return result

```

Example usage

```
print(combinationSum([2,3,6,7], 7))
```

• LC.ca

```

class Solution:
    def combinationSum(self, candidates:
List[int], target: int) -> List[List[int]]:
        # i: start index for this recursion
        # s: sum
        def dfs(i, s):
            if s == target:
                ans.append(t.copy())
                return
            if s > target:
                return

```

```

        for j in range(i, len(candidates)):
            c = candidates[j]
            t.append(c)
            dfs(j, s + c)
            t.pop()

        ans = []
        t = []
        # candidates.sort() # diff from
        combinationSum-II, no need sorting it
        dfs(0, 0)

```

return ans

#####

dp version

```

class Solution_dp:
    def combinationSum(self, candidates:
List[int], target: int) -> List[List[int]]:
        # for each-target (from 1 to target), its
        dp[i][j]
        # => so 3-D array dp[][][]
        dp = []

```

```
candidates.sort()
```

```
for i in range(1, target+1):
```

```
    cur = []
```

```
    for j in range(len(candidates)):
```

```
        if candidates[j] > i:
```

```
            break
```

```
        if candidates[j] == i:
```

```
            one = [candidates[j]]
```

```
            cur.append(one)
```

```
        break
```

```
    for a in dp[i - candidates[j] - 1]:
```

```
        if candidates[j] > a[0]:
```

```
            continue
```

```
        deepCopied = a.copy()
```

```
        deepCopied.insert(0,
```

```
candidates[j])
```

```
        cur.append(deepCopied)
```

```
    dp.append(cur)
```

```
return dp[-1]
```

```
#####
```

```
class Solution(object):
```

```
    def combinationSum(self, candidates, target):
```

```
        """
```

```
        :type candidates: List[int]
```

```
        :type target: int
```

```
        :rtype: List[List[int]]
```

```
        """
```

```
    def dfs(candidates, start, target, path, res):
```

```
        if target == 0:
```

```
            return res.append(path + [])
```

```
        for i in range(start, len(candidates)):
```

```
            if target - candidates[i] >= 0:
```

```
                path.append(candidates[i])
```

```
                dfs(candidates, i, target -  
candidates[i], path, res)
```



```

        path.pop()

    res = []
    dfs(candidates, 0, target, [], res)
    return res

#####

class Solution:
    def combinationSum(self, candidates:
List[int], target: int) -> List[List[int]]:

        def dfs(i: int, s: int):
            if s == 0:
                ans.append(t[:])
                return
            if i >= len(candidates) or s <
candidates[i]:
                return
            dfs(i + 1, s)
            t.append(candidates[i])
            dfs(i, s - candidates[i])
            t.pop()

        candidates.sort()
        t = []
        ans = []
        dfs(0, target)
        return ans

```

Round 5 · Mid · Medium

p.2233

◆ Quick Review

Key Insights

- Backtracking is a natural fit to explore all potential combinations.
- Sorting candidates simplifies the process by allowing early termination if the current candidate exceeds the remaining target.
- Recursion is used to build combinations and backtrack once the current path exceeds or meets the target.
- The same candidate can be reused in the combination, so the recursive call does not advance the index.

Space & Time

Time Complexity: $O(2^{(\text{target}/\min(\text{candidates}))})$ in the worst case where many combinations are possible.

Space Complexity: $O(\text{target}/\min(\text{candidates}))$ for the recursion stack and temporary space used to store combinations.

Solution

The solution uses a backtracking approach. First, sort the candidates to enable pruning of recursive calls once the candidate exceeds the remaining target. A recursive helper function builds up a combination, subtracting the candidate value from the target until it reaches zero (a valid combination) or becomes negative (an invalid path). The recursion allows the same candidate to be chosen again. Each valid combination is then added to the result list.

Round 5 · Mid · Medium

p.2234

Backtracking

#46 Permutations

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Backtracking

Lists: Grind 169

Problem

Description

Given an array `nums` of distinct integers, return all the possible permutations. You can return the answer in any order.

Example 1:

Input: `nums = [1,2,3]`
 Output:
`[[1,2,3],[1,3,2],[2,1,3],[2,3,1],[3,1,2],[3,2,1]]`

Example 2:

Input: `nums = [0,1]`
 Output: `[[0,1],[1,0]]`

Example 3:

Input: `nums = [1]`
 Output: `[[1]]`

Constraints:

- $1 \leq \text{nums.length} \leq 6$

- $-10 \leq \text{nums}[i] \leq 10$
- All the integers of `nums` are unique.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def permute(self, nums: list[int]) -> list[list[int]]:
        ans = []
        used = [False] * len(nums)

        def dfs(path: list[int]) -> None:
            if len(path) == len(nums):
                ans.append(path.copy())
            return

        for i, num in enumerate(nums):
            if used[i]:
                continue
            used[i] = True
            path.append(num)
            dfs(path)
            path.pop()
            used[i] = False

        dfs([])

        return ans
```

• LeetDoocs

```

class Solution:
    def permute(self, nums: List[int]) ->
List[List[int]]:
    def dfs(i: int):
        if i >= n:
            ans.append(t[:])
            return
        for j, x in enumerate(nums):
            if not vis[j]:
                vis[j] = True
                t[i] = x

                dfs(i + 1)
                vis[j] = False

    n = len(nums)
    vis = [False] * n
    t = [0] * n
    ans = []
    dfs(0)
    return ans

```

• SimplyLeet

Round 5 · Mid · Medium

p.2237

```

def permute(nums):
    result = []

    def backtrack(current, remaining):
        # If no remaining elements, current
        permutation is complete.
        if not remaining:
            result.append(current.copy())
            return
        # Explore each candidate number.
        for i in range(len(remaining)):

            # Add the chosen number to the current
            permutation.
            current.append(remaining[i])
            # Recursively generate permutations
            with the remaining numbers.
            backtrack(current, remaining[:i] +
            remaining[i+1:])
            # Backtrack by removing the chosen
            number.
            current.pop()

    backtrack([], nums)
    return result

# Example usage:
print(permute([1,2,3]))

```

• LC.ca

Round 5 · Mid · Medium

p.2238

```
'''
remove an element by its value from a set
```

```
>>> my_set = {1, 2, 3, 4, 5}
>>> my_set.remove(3)
>>> print(my_set)
{1, 2, 4, 5}
```

```
-----
```

cannot remove an element by index from a set in Python3
need to convert the set to a list

```
>>> my_set = {1, 2, 3, 4, 5}
>>> my_list = list(my_set)
>>> del my_list[2] # remove the element at index 2
>>> my_set = set(my_list)
>>> print(my_set)
{1, 2, 4, 5}
'''
```

```
class Solution: # iterative
    def permute(self, nums: List[int]) -> List[List[int]]:
        res = [[]]

        if nums is None or len(nums) == 0:
            return ans

        for num in nums:
            new_res = []

            for perm in res:
                for i in range(len(perm) + 1):
                    new_perm = perm[:i] + [num] + perm[i:]
                    new_res.append(new_perm)

            res = new_res

        return res

#####
```

```
class Solution: # iterative,
single_perm.insert(index, num)
    def permute(self, nums: List[int]) ->
List[List[int]]:
    ans = [[]]

    if nums is None or len(nums) == 0:
        return ans

    for num in nums:
        tmp_list = []
```

```
        for single_perm in ans:
            for index in
range(len(single_perm) + 1):
                single_perm.insert(index, num)
                tmp_list.append(single_perm.co
py())

                single_perm.pop(index)

        ans = tmp_list

    return ans
```

```
#####
```

```
'''
```

In Python 3, both the `discard()` and `remove()` methods of a set object are used to remove an element from the set, but there is one key difference:

- * `discard()` removes the specified element from the set if it is present, but does nothing if the element is not present.

- * `remove()` removes the specified element from the set if it is present, but raises a `KeyError` exception if the element is not present.

```
>>> a
{33, 66, 11, 44, 22, 55}
>>> a.discard(22)
>>> a
{33, 66, 11, 44, 55}
```

```
>>> a.discard(200)
>>>
>>> a.remove(200)
Traceback (most recent call last):
```

```
File "<stdin>", line 1, in <module>
```

```
KeyError: 200
```

```
...
```

```
...
```

```
>>> v1 = set([])
```

```
>>> v2 = set()
```

```
>>>
```

```
>>> v1
```

```
set()
```

```
>>> v2
```

```
set()
```

```
...
```

```
class Solution:
```

```
    def permute(self, nums: List[int]) ->
List[List[int]]:
```

```
        res = []
```

```
        visited = set([])
```

```
        def dfs(nums, path, res, visited):
```

```
            if len(path) == len(nums):
```

```
        res.append(path + [])
```

```
        return
```

```
        for i in range(0, len(nums)):
```

```
            if i not in visited:
```

```
                visited.add(i)
```

```
                path.append(nums[i])
```

```
                dfs(nums, path, res, visited)
```

```
                path.pop()
```

```
                visited.discard(i) # remove(i)
```

```
will throw exception if i not existing
```

```
        dfs(nums, [], res, visited)
```

```
        return res
```

```
#####
```

```
class Solution:
```

```
    def permute(self, nums: List[int]) ->
List[List[int]]:
```

```
        def dfs(i):
```

```
            if i == n:
```

```

        ans.append(t[:])
        return
    for j in range(n):
        if not vis[j]:
            vis[j] = True
            t[i] = nums[j]
            dfs(i + 1)
            vis[j] = False

n = len(nums)

vis = [False] * n
t = [0] * n
ans = []
dfs(0)
return ans

```

◆ Quick Review

Key Insights

- Use backtracking to generate all possible permutations.
- Build each permutation by selecting numbers one at a time.
- When the current permutation reaches the length of nums, add it to the result.
- Backtracking allows you to undo choices and explore different orderings.

Space & Time

Time Complexity: $O(n * n!)$ – There are $n!$ permutations, and constructing each permutation takes $O(n)$ time.

Space Complexity: $O(n)$ – Recursion depth is limited to n , plus additional space for the current permutation.

Solution

The solution employs a backtracking algorithm. A helper function is used to build permutations recursively. At each step, the algorithm iterates through the available numbers, adds one to the current permutation if it has not been used yet, and recurses to handle the remaining numbers. Once the current permutation matches the length of the input array, it is added to the result list. Then, by "backtracking" (removing the last number added), the algorithm explores alternative possibilities. This systematic exploration ensures that all possible permutations are generated.

Backtracking

#79 Word Search

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · String · Backtracking · Matrix

Lists: Grind 169

Problem

Description

Given an $m \times n$ grid of characters board and a string word, return true if word exists in the grid.

The word can be constructed from letters of sequentially adjacent cells, where adjacent cells are horizontally or vertically neighboring. The same letter cell may not be used more than once.

Example 1:

A	B	C	E
S	F	C	S
A	D	E	E

Input: board = [["A","B","C","E"], ["S","F","C","S"], ["A","D","E","E"]], word = "ABCCED"

Output: true

Example 2:

A	B	C	E
S	F	C	S
A	D	E	E

Input: board = `[["A","B","C","E"],["S","F","C","S"],["A","D","E","E"]]`, word = "SEE"

Output: true

Example 3:

Round 5 · Mid · Medium

p.2249

A	B	C	E
S	F	C	S
A	D	E	E

Input: board = `[["A","B","C","E"],["S","F","C","S"],["A","D","E","E"]]`, word = "ABCB"

Output: false

Constraints:

- `m == board.length`
- `n = board[i].length`
- `1 <= m, n <= 6`
- `1 <= word.length <= 15`
- board and word consists of only lowercase and uppercase English letters.

Round 5 · Mid · Medium

p.2250

Follow up: Could you use search pruning to make your solution faster with a larger board?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def exist(self, board: list[list[str]], word: str) -> bool:
        m = len(board)
        n = len(board[0])

        def dfs(i: int, j: int, s: int) -> bool:
            if i < 0 or i == m or j < 0 or j == n:
                return False
            if board[i][j] != word[s] or board[i][j] == '*':
                return False

            if s == len(word) - 1:
                return True

            cache = board[i][j]
            board[i][j] = '*'
            isExist = (dfs(i + 1, j, s + 1) or
                      dfs(i - 1, j, s + 1) or
                      dfs(i, j + 1, s + 1) or
                      dfs(i, j - 1, s + 1))
            board[i][j] = cache
```

```
return isExist
```

```
return any(dfs(i, j, 0)
           for i in range(m)
           for j in range(n))
```

• LeetDoocs

```
class Solution:
    def exist(self, board: List[List[str]], word: str) -> bool:
        def dfs(i: int, j: int, k: int) -> bool:
            if k == len(word) - 1:
                return board[i][j] == word[k]
            if board[i][j] != word[k]:
                return False
            c = board[i][j]
            board[i][j] = "0"
            for a, b in pairwise((-1, 0, 1, 0, -1)):
                if dfs(i + a, j + b, k + 1):
                    return True
            board[i][j] = c
            return False
```

```

        x, y = i + a, j + b
        ok = 0 <= x < m and 0 <= y < n and
board[x][y] != "0"
        if ok and dfs(x, y, k + 1):
            return True
        board[i][j] = c
        return False

    m, n = len(board), len(board[0])
    return any(dfs(i, j, 0) for i in range(m)
for j in range(n))

```

• SimplyLeet

```

# Define the function to solve the Word Search
problem.
def exist(board, word):
    # Get dimensions of the board.
    rows, cols = len(board), len(board[0])

    # Helper function for DFS search.
    def dfs(r, c, index):
        # If the entire word is found, return True.
        if index == len(word):
            return True

```

```

        # Check boundaries and if current cell
already used or not matching.
        if r < 0 or c < 0 or r >= rows or c >= cols
or board[r][c] != word[index]:
            return False

```

```

        # Temporarily mark the cell as visited.
temp = board[r][c]
board[r][c] = "#"

```

```

        # Explore all four directions.

```

```

        found = (dfs(r + 1, c, index + 1) or
dfs(r - 1, c, index + 1) or
dfs(r, c + 1, index + 1) or
dfs(r, c - 1, index + 1))

        # Backtrack: restore the original cell
value.
        board[r][c] = temp
        return found

    # Initiate DFS from each cell in the grid.

```

```

for i in range(rows):
    for j in range(cols):
        if dfs(i, j, 0):
            return True
    return False

# Example test case
if __name__ == "__main__":
    board = [
        ["A", "B", "C", "E"],
        ["S", "F", "C", "S"],
        ["A", "D", "E", "E"]
    ]
    word = "ABCCED"
    print(exist(board, word)) # Expected output:
    True

```

• LC.ca

Round 5 · Mid · Medium

p.2255

```

class Solution:
    def exist(self, board: List[List[str]], word:
str) -> bool:
        def dfs(i, j, cur): # cur: current char
count
            if cur == len(word):
                return True
            if (
                i < 0
                or i >= m
                or j < 0
                or j >= n
                or board[i][j] == '0'
                or word[cur] != board[i][j]
            ):
                return False
            t = board[i][j]
            board[i][j] = '0' # mark as visited
            for a, b in [[0, 1], [0, -1], [-1, 0],
[1, 0]]:
                x, y = i + a, j + b
                if dfs(x, y, cur + 1):
                    return True

```

Round 5 · Mid · Medium

p.2256

```

        board[i][j] = t
        return False

    m, n = len(board), len(board[0])
    return any(dfs(i, j, 0) for i in range(m)
              for j in range(n))

```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to explore each cell in the board as a potential starting point.
- Employ backtracking to mark cells as visited when exploring a path and revert them back when backtracking.
- Check boundary conditions and ensure the current cell's character matches the corresponding character in the word.
- Given the small grid and word sizes, a recursive solution is both feasible and straightforward.
- Search pruning can be applied by stopping a DFS branch as soon as the current path does not match the prefix of the target word.

Space & Time

Time Complexity: $O(m * n * 3^L)$, where L is the length of the word (each DFS may branch to at most 3 further cells, excluding the coming cell).

Space Complexity: $O(L)$ due to recursion call stack (where L is the word length).

Solution

The solution uses DFS with backtracking to traverse the grid. For each cell, we:

- Check if the cell matches the first character of the word.
- Recursively explore the adjacent cells (up, down, left, right) while marking the current cell as visited.
- If the path does not lead to a complete match, backtrack by resetting the visited cell.
- The DFS stops if all characters in the word have been found in sequence.

The algorithm leverages in-place modifications (or an auxiliary visited array) to avoid revisiting cells while exploring a candidate path.

Backtracking

#113 Path Sum II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Backtracking · Tree · DFS

Lists: Grind 169


Problem

Description

Given the root of a binary tree and an integer `targetSum`, return all root-to-leaf paths where the sum of the node values in the path equals `targetSum`. Each path should be returned as a list of the node values, not node references.

A root-to-leaf path is a path starting from the root and ending at any leaf node. A leaf is a node with no children.

Example 1:



Input: `root = [5,4,8,11,null,13,4,7,2,null,null,5,1]`,
`targetSum = 22`


Output: `[[5,4,11,2],[5,8,4,5]]`

Explanation: There are two paths whose sum equals `targetSum`:

$5 + 4 + 11 + 2 = 22$

$5 + 8 + 4 + 5 = 22$

Example 2:



Input: root = [1,2,3], targetSum = 5
Output: []

Example 3:

Input: root = [1,2], targetSum = 0
Output: []

Constraints:

- The number of nodes in the tree is in the range [0, 5000].
- $-1000 \leq \text{Node.val} \leq 1000$
- $-1000 \leq \text{targetSum} \leq 1000$

Round 5 · Mid · Medium

p.2261

Community Solutions (Python)

• WalkCC

```

class Solution:
    def pathSum(self, root: TreeNode, summ: int) -> list[list[int]]:
        ans = []

        def dfs(root: TreeNode, summ: int, path: list[int]) -> None:
            if not root:
                return
            if root.val == summ and not root.left and not root.right:
                ans.append(path + [root.val])
                return

            dfs(root.left, summ - root.val, path + [root.val])
            dfs(root.right, summ - root.val, path + [root.val])

        dfs(root, summ, [])
        return ans
  
```

• LeetCode

Round 5 · Mid · Medium

p.2262

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def pathSum(self, root: Optional[TreeNode],
targetSum: int) -> List[List[int]]:
        def dfs(root, s):
            if root is None:
                return
            s += root.val
            t.append(root.val)
            if root.left is None and root.right is
None and s == targetSum:
                ans.append(t[:])
                dfs(root.left, s)
                dfs(root.right, s)
                t.pop()

        ans = []

        t = []
        dfs(root, 0)
        return ans

```

• SimplyLeet

Round 5 · Mid · Medium

p.2263

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
# right=None):
#         self.val = val
#         self.left = left
#         self.right = right

def pathSum(root, targetSum):
    result = [] # This will store all valid paths

    def dfs(node, current_path, remaining_sum):
        if not node:
            return # Base case: reached the end of
a path

        # Add the current node's value to the path
        current_path.append(node.val)

        # Check if it's a leaf and if the path sum
equals targetSum
        if not node.left and not node.right and
node.val == remaining_sum:
            result.append(list(current_path)) #
Add a copy of the current path

```

Round 5 · Mid · Medium

p.2264


```

        else:
            # Continue the search in the left and
            right subtrees
            dfs(node.left, current_path,
                remaining_sum - node.val)
            dfs(node.right, current_path,
                remaining_sum - node.val)

            # Backtrack: remove the last element before
            returning to the caller
            current_path.pop()

    dfs(root, [], targetSum)
    return result

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
#         right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def pathSum(self, root: Optional[TreeNode],
        targetSum: int) -> List[List[int]]:
        def dfs(root, s):
            if root is None:

```

```

            return
            s += root.val
            t.append(root.val)
            if root.left is None and root.right is
            None and s == targetSum:
                ans.append(t[:])
            dfs(root.left, s)
            dfs(root.right, s)
            t.pop()

    ans = []

    t = []
    dfs(root, 0)
    return ans

```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to explore all paths from the root to the leaves.
- Backtracking is essential: add the current node value to the path before the recursive calls and remove it (backtrack) after exploring both subtrees.
- At each leaf, check if the accumulated sum equals targetSum.
- Maintain a running sum or subtract the current node's value from the target sum as you traverse.

Space & Time

Time Complexity: $O(N^2)$ in the worst-case scenario (where N is the number of nodes) when most nodes are leaf nodes, since each potential path copy operation can be $O(N)$.

Space Complexity: $O(N)$ for the recursion stack and the path storage, where N is the height of the tree.

Solution

We solve the problem using a recursive DFS approach combined with backtracking. The idea is to traverse the tree while maintaining the current path and the remaining sum needed to reach targetSum. When a leaf node is reached (node with no children), we check if the remaining sum equals the leaf's value – if yes, we record a copy of the current path.

Data structures and algorithmic approaches used:

- Recursion for DFS: to traverse to each leaf.
- List (or Array) for storing the current path.
- Backtracking: to remove the last added element when stepping back up the recursion stack.
- Copying the path when a valid leaf is reached to store it in the final result.

A common pitfall is to forget to backtrack, which would result in incorrect paths being recorded.

Greedy

Round 5 · Mid · Medium · 5 questions

Round 5 · Mid · Medium

p.2267

Round 5 · Mid · Medium

p.2268

Greedy

#134 Gas Station

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Greedy

Lists: Grind 169

Problem

Description

There are n gas stations along a circular route, where the amount of gas at the i th station is $gas[i]$.

You have a car with an unlimited gas tank and it costs $cost[i]$ of gas to travel from the i th station to its next ($i + 1$)th station. You begin the journey with an empty tank at one of the gas stations.

Given two integer arrays gas and $cost$, return the starting gas station's index if you can travel around the circuit once in the clockwise direction, otherwise return -1 . If there exists a solution, it is guaranteed to be unique.

Example 1:

Input: $gas = [1,2,3,4,5]$, $cost = [3,4,5,1,2]$

Output: 3

Explanation:

Start at station 3 (index 3) and fill up with 4 unit of gas. Your tank = $0 + 4 = 4$

Travel to station 4. Your tank = $4 - 1 + 5 = 8$

Travel to station 0. Your tank = $8 - 2 + 1 = 7$

Travel to station 1. Your tank = $7 - 3 + 2 = 6$

Travel to station 2. Your tank = $6 - 4 + 3 = 5$

Example 2:

Input: $gas = [2,3,4]$, $cost = [3,4,3]$

Output: -1

Explanation:

You can't start at station 0 or 1, as there is not enough gas to travel to the next station.

Let's start at station 2 and fill up with 4 unit of gas. Your tank = $0 + 4 = 4$

Travel to station 0. Your tank = $4 - 3 + 2 = 3$

Travel to station 1. Your tank = $3 - 3 + 3 = 3$

You cannot travel back to station 2, as it requires 4 unit of gas but you only have 3.

Constraints:

- $n == gas.length == cost.length$
- $1 \leq n \leq 105$
- $0 \leq gas[i], cost[i] \leq 104$
- The input is generated such that the answer is unique.

Community Solutions (Python)

• WalkCC

```

class Solution:
    def canCompleteCircuit(self, gas: list[int],
cost: list[int]) -> int:
    ans = 0
    net = 0
    summ = 0

    # Try to start from each index.
    for i in range(len(gas)):
        net += gas[i] - cost[i]
        summ += gas[i] - cost[i]

    if summ < 0:
        summ = 0
        ans = i + 1 # Start from the next index.

    return -1 if net < 0 else ans

```

• LeetCode

```

class Solution:
    def canCompleteCircuit(self, gas: List[int],
cost: List[int]) -> int:
    n = len(gas)
    i = j = n - 1
    cnt = s = 0
    while cnt < n:
        s += gas[j] - cost[j]
        cnt += 1
        j = (j + 1) % n
        while s < 0 and cnt < n:
            i -= 1
            s += gas[i] - cost[i]
            cnt += 1
    return -1 if s < 0 else i

```

• SimplyLeet

```
# Define a function to find starting gas station
index
def canCompleteCircuit(gas, cost):
    total_surplus = 0 # Total gas surplus after
    completing the route
    current_surplus = 0 # Gas surplus while
    iterating from a candidate start station
    start_index = 0 # Candidate starting station
    index

    # Loop through each station
    for i in range(len(gas)):
        # Calculate net gas after leaving station i
        surplus = gas[i] - cost[i]

        total_surplus += surplus
        current_surplus += surplus

        # If current surplus becomes negative,
        station i is not a valid start.
        if current_surplus < 0:
            # Reset starting station to the next
            station
            start_index = i + 1
            current_surplus = 0 # Reset current
            surplus

        # Check if overall surplus is non-negative
        return start_index if total_surplus >= 0 else
        -1
```

```
# Example usage:
gas = [1,2,3,4,5]
cost = [3,4,5,1,2]
print(canCompleteCircuit(gas, cost)) # Output: 3
```

• LC.ca

```
class Solution:
    def canCompleteCircuit(self, gas: List[int],
    cost: List[int]) -> int:
        n = len(gas)
        i = j = n - 1
        cnt = s = 0
        while cnt < n:
            s += gas[j] - cost[j]
            cnt += 1
            j = (j + 1) % n
            while s < 0 and cnt < n:

                i -= 1
                s += gas[i] - cost[i]
                cnt += 1
        return -1 if s < 0 else i
```

♦ Quick Review

Key Insights

- If the total amount of gas available is less than the total travel cost, then completing the circuit is impossible.

- A greedy approach can be used by keeping track of the current surplus of gas. If the surplus becomes negative at any station, the journey cannot start from any station before that point.
- Restart the journey from the next station whenever the surplus goes negative.
- The problem guarantees that if a solution exists, it is unique.

Space & Time

Time Complexity: $O(n)$, where n is the number of gas stations, as we traverse the list only once.

Space Complexity: $O(1)$, as no extra space is used other than a few variables.

Solution

The solution involves iterating over the list of stations while maintaining two variables: one for the current gas surplus and one for the total gas surplus. If at any station the current surplus drops below zero, the current starting station is not viable and we set the starting station to the next index and reset the current surplus. After one pass, if the total gas surplus is non-negative, then the identified station is the valid starting point.

Greedy

#179 Largest Number

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · Greedy · Sorting

Lists: Grind 169

Problem

Description

Given a list of non-negative integers `nums`, arrange them such that they form the largest number and return it.

Since the result may be very large, so you need to return a string instead of an integer.

Example 1:

Input: `nums = [10,2]`
Output: `"210"`

Example 2:

Input: `nums = [3,30,34,5,9]`
Output: `"9534330"`

Constraints:

- $1 \leq \text{nums.length} \leq 100$
- $0 \leq \text{nums}[i] \leq 10^9$

Community Solutions (Python)

• WalkCC

```
class LargerStrKey(str):
    def __lt__(x: str, y: str) -> bool:
        return x + y > y + x

class Solution:
    def largestNumber(self, nums: list[int]) -> str:
        return ''.join(sorted(map(str, nums),
                                key=LargerStrKey)).rstrip('0') or '0'
```

• LeetCode

```
class Solution:
    def largestNumber(self, nums: List[int]) -> str:
        nums = [str(v) for v in nums]
        nums.sort(key=cmp_to_key(lambda a, b: 1 if
                                a + b < b + a else -1))
        return "0" if nums[0] == "0" else
        "".join(nums)
```

• SimplyLeet

```
# Import cmp_to_key to convert a comparator
function to a key function for sorting
from functools import cmp_to_key
```

```
# Define the comparator function
def compare(x, y):
    # Compare two possible orders: x+y and y+x
    if x + y > y + x:
        return -1 # x should come before y
    elif x + y < y + x:
        return 1 # y should come before x
```

```
else:
    return 0 # both orders are equal
```

```
def largestNumber(nums):
    # Convert all integers to strings
    nums_str = list(map(str, nums))
    # Sort the strings using the custom comparator
    nums_str.sort(key=cmp_to_key(compare))
    # If the largest number is '0', the entire
    number is 0 so return '0'
    if nums_str[0] == "0":
```

```

        return "0"
    # Concatenate the sorted numbers and return the
    result
    return "".join(nums_str)

# Example usage:
print(largestNumber([3, 30, 34, 5, 9])) # Output:
"9534330"

```

• LC.ca

```

class Solution:
    def largestNumber(self, nums: List[int]) ->
    str:
        nums = [str(v) for v in nums]
        nums.sort(key=cmp_to_key(lambda a, b: 1 if
a + b < b + a else -1))
        return "0" if nums[0] == "0" else
"".join(nums)

```

◆ Quick Review

Key Insights

- The order in which numbers appear drastically affects the concatenated result.
- A custom sorting strategy is required where two numbers are compared by checking which concatenation (first-second vs second-first) yields a larger number.
- Edge case: When the numbers are all zeros, the result should be "0" instead of several leading zeros.

Space & Time

Time Complexity: $O(n \log n * k)$, where n is the number of elements in the array and k is the maximum number of digits in a number (due to string comparisons in the custom sort).

Space Complexity: $O(n * k)$ for storing and processing the strings.

Solution

We solve the problem by converting all integers to strings so that we can use string concatenation to compare the results. A custom comparator is defined that, given two number strings, compares the concatenated results in both orders. By sorting the array using this comparator in descending order, we ensure that placing the highest contributing number first yields the largest possible number. After sorting, a check is performed to return "0" if the highest number is "0". Finally, the sorted strings are concatenated to form the answer.

Greedy

#280 Wiggle Sort

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Greedy · Sorting

Lists: Premium 98

Problem

Description

Given an integer array `nums`, reorder it such that `nums[0] <= nums[1] >= nums[2] <= nums[3]...`

You may assume the input array always has a valid answer.

Example 1:

Input: `nums = [3,5,2,1,6,4]`

Output: `[3,5,1,6,2,4]`

Explanation: `[1,6,2,5,3,4]` is also accepted.

Example 2:

Input: `nums = [6,6,5,6,3,8]`

Output: `[6,6,5,6,3,8]`

Constraints:

- $1 \leq \text{nums.length} \leq 5 * 10^4$
- $0 \leq \text{nums}[i] \leq 10^4$
- It is guaranteed that there will be an answer for the given input `nums`.

Round 5 · Mid · Medium

p.2281

Follow up: Could you solve the problem in $O(n)$ time complexity?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def wiggleSort(self, nums: list[int]) -> None:
        # 1. If i is even, then nums[i] <= nums[i - 1].
        # 2. If i is odd, then nums[i] >= nums[i - 1].
        for i in range(1, len(nums)):
            if (i % 2 == 0 and nums[i] > nums[i - 1] or
                i % 2 == 1 and nums[i] < nums[i - 1]):
                nums[i], nums[i - 1] = nums[i - 1], nums[i]
```

• LeetCode

Round 5 · Mid · Medium

p.2282

```

class Solution:
    def wiggleSort(self, nums: List[int]) -> None:
        """
        Do not return anything, modify nums
        in-place instead.
        """
        for i in range(1, len(nums)):
            if (i % 2 == 1 and nums[i] < nums[i -
1]) or (
                i % 2 == 0 and nums[i] > nums[i -
1]
            ):
                nums[i], nums[i - 1] = nums[i - 1],
nums[i]

```

• SimplyLeet

```

# Python solution for Wiggle Sort
def wiggleSort(nums):
    # Iterate over the list starting from index 1
    for i in range(1, len(nums)):
        # Check the condition based on the parity
        of the index
        if i % 2 == 1:
            # For odd index, if previous element is
            greater, swap them
            if nums[i] < nums[i-1]:
                nums[i], nums[i-1] = nums[i-1],
nums[i]
            else:

```

```

        # For even index, if previous element
        is smaller, swap them
        if nums[i] > nums[i-1]:
            nums[i], nums[i-1] = nums[i-1],
nums[i]
        return nums

# Example usage:
# nums = [3,5,2,1,6,4]
# print(wiggleSort(nums))

```

• LC.ca

```

'''
math prove:

example [3, 5, 2, 1]
when reaching 2, all good
then next, is 1,
so here comes the question, will the number after 2
violating the wiggle rule?
answer is no.

in example, when reaching 2:

```

* if [3, 5, 2, 10], so 10 is bigger than 2, then no swap needed according to the if () conditions

* if [3, 5, 2, 1], so previous round, guaranteed that 2 is smaller than 5.

And, now a swap needed meaning the number is smaller than 2, i.e. must be smaller than 5, so wiggle rule is still good

...

class Solution:

```
def wiggleSort(self, nums: List[int]) -> None:
    """
```

Do not return anything, modify nums in-place instead.

.....

```
    for i in range(1, len(nums)):
        if (i % 2 == 1 and nums[i] < nums[i - 1]) or (
            i % 2 == 0 and nums[i] > nums[i - 1]
        ):
            nums[i], nums[i - 1] = nums[i - 1], nums[i]
```

#####

```
class Solution: # sort list first
```

```
def wiggleSort(self, nums: List[int]) -> None:
    nums.sort() # Sort the list in ascending
```

order

```
n = len(nums)
```

```
    for i in range(1, n - 1, 2):
        nums[i], nums[i + 1] = nums[i + 1],
        nums[i] # Swap adjacent elements
```

.....

If the list has an odd length, the last element will be compared with the second-to-last element and swapped if necessary to form a wiggle sequence.

.....

```
    if n % 2 == 0 or nums[-1] < nums[-2]:
        nums[-1], nums[-2] = nums[-2],
        nums[-1]
```

```
#####
import random
```

class Solution:

```
def wiggleSort(self, nums: List[int]) -> None:
    """
```

Do not return anything, modify nums in-place instead.

```

"""
    if len(nums) <= 1:
        return

    # Find the median using quickselect
    n = len(nums)
    mid = self.quickselect(nums, 0, n-1,
(n-1)//2)

    # Create an empty array with the same
length as nums

    ans = [None] * n

    # Assign the median to every other element
    i = 0
    j = n-1 if n % 2 == 1 else n-2
    for k in range(n):
        if nums[k] < mid:
            ans[j] = nums[k]
            j -= 2
        elif nums[k] > mid:

```

```

        ans[i] = nums[k]
        i += 2
    else:
        ans[j+1] = mid

    # Copy the result back to nums
    nums[:] = ans

def quickselect(self, nums, left, right, k):
    """
        Returns the k-th smallest element in
nums[left:right+1].
        """
        if left == right:
            return nums[left]

        # Partition the subarray around a pivot
element
        pivot_idx = self.partition(nums, left,
right)

        # Recursively select on the left or right
subarray
        if k == pivot_idx:

```

```

        return nums[k]
    elif k < pivot_idx:
        return self.quickselect(nums, left,
            pivot_idx-1, k)
    else:
        return self.quickselect(nums,
            pivot_idx+1, right, k)

    def partition(self, nums, left, right):
        """
        Partitions the subarray nums[left:right+1]
        using the first element as the pivot.
        Returns the final index of the pivot
        element.

        """
        pivot = nums[left]
        i, j = left+1, right
        while i <= j:
            if nums[i] <= pivot:
                i += 1
            elif nums[j] > pivot:
                j -= 1
            else:
                nums[i], nums[j] = nums[j],
nums[i]

        nums[left], nums[j] = nums[j], nums[left]
        return j

```

◆ Quick Review

Key Insights

- The alternating "wiggle" requirement can be satisfied by ensuring two conditions:
- For every odd index i , $\text{nums}[i]$ should be greater than or equal to $\text{nums}[i-1]$.
- For every even index i (where $i > 0$), $\text{nums}[i]$ should be less than or equal to $\text{nums}[i-1]$.
- A single pass through the array can ensure these conditions by swapping adjacent elements when the condition is violated.
- This greedy approach works because the local adjustments guarantee the overall wiggle pattern without needing to sort the entire array.

Space & Time

Time Complexity: $O(n)$ — Only one pass through the array is needed.

Space Complexity: $O(1)$ — The modifications are done in-place without extra data structures.

Solution

We can achieve a wiggle sort in one pass by iterating through the array and, at each index i (starting from 1), checking if the pair $(\text{nums}[i-1], \text{nums}[i])$ satisfies the wiggle condition. If i is odd, $\text{nums}[i]$ should be at least $\text{nums}[i-1]$; if not, we swap them. Conversely, if i is even and $\text{nums}[i]$ is greater than $\text{nums}[i-1]$, then we swap them. These local swaps adjust the order incrementally while

ensuring that the overall wiggle pattern is maintained. This greedy approach works because only two adjacent numbers are involved in any violation and fixing them locally does not disturb the rest of the order.

Greedy

#954 Array of Doubled Pairs

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Greedy · Sorting

Lists: CodeSignal

Problem

Description

Given an integer array of even length `arr`, return `true` if it is possible to reorder `arr` such that $arr[2 * i + 1] = 2 * arr[2 * i]$ for every $0 \leq i < \text{len}(arr) / 2$, or `false` otherwise.

Example 1:

Input: `arr = [3,1,3,6]`
Output: `false`

Example 2:

Input: `arr = [2,1,2,6]`
Output: `false`

Example 3:

Input: `arr = [4,-2,2,-4]`
Output: `true`
Explanation: We can take two groups, `[-2,-4]` and `[2,4]` to form `[-2,-4,2,4]` or `[2,4,-2,-4]`.

Constraints:

- $2 \leq \text{arr.length} \leq 3 * 10^4$
- arr.length is even.
- $-105 \leq \text{arr}[i] \leq 105$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def canReorderDoubled(self, arr: list[int]) ->
    bool:
        count = collections.Counter(arr)

        for key in sorted(count, key=abs):
            if count[key] > count[2 * key]:
                return False
            count[2 * key] -= count[key]

        return True
```

• LeetCode

```
class Solution:
    def canReorderDoubled(self, arr: List[int]) ->
    bool:
        freq = Counter(arr)
        if freq[0] & 1:
            return False
        for x in sorted(freq, key=abs):
            if freq[x <= 1] < freq[x]:
                return False
            freq[x <= 1] -= freq[x]
        return True
```

• SimplyLeet

```
from collections import Counter

def canReorderDoubled(arr):
    # Count frequency for each element in arr
    count = Counter(arr)
    # Sort the array based on the absolute value of
    elements
    for x in sorted(arr, key=abs):
        # If no x is left to be paired, skip to the
        next element
        if count[x] == 0:
            continue
```

```

    # Check if double the current number exists
    in sufficient frequency
    if count[2 * x] < count[x]:
        return False
    # Decrement the frequency for the double of
    x
    count[2 * x] -= count[x]
    return True

# Example usage:
print(canReorderDoubled([4,-2,2,-4])) # Expected
Output: True

```

- LC.ca

```

class Solution:
    def canReorderDoubled(self, arr: List[int]) ->
    bool:
        freq = Counter(arr)
        if freq[0] & 1:
            return False
        for x in sorted(freq, key=abs):
            if freq[x << 1] < freq[x]:
                return False
            freq[x << 1] -= freq[x]
        return True

```

◆ Quick Review

Key Insights

- Sort the array based on the absolute value of the elements to handle negative numbers correctly.
- Use a hash table (or frequency counter) to track the occurrence of each element.
- For each number in the sorted list, if the number is still available in the frequency map, try to pair it with its double.
- Decrement the appropriate counts and validate that it is possible to find sufficient doubles for all numbers.
- Special case: when dealing with zeros, pairing works within the same value.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting.

Space Complexity: $O(n)$ for storing the frequency map.

Solution

The solution starts by counting the frequency of each element in the array. Sorting the array by absolute value ensures that when we process an element, any potential pair (i.e., its double) comes later in the sequence. For each number, if there are available occurrences, we try to match it with its double. If there are not enough doubles available, the answer is false. If all elements can be paired appropriately by decrementing their counts in the frequency map, the function returns true.

Greedy

#2131 Longest Palindrome by Concatenating Two Letter Words

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Greedy · Counting

Lists: CodeSignal

Problem

Description

You are given an array of strings `words`. Each element of `words` consists of two lowercase English letters.

Create the longest possible palindrome by selecting some elements from `words` and concatenating them in any order. Each element can be selected at most once. Return the length of the longest palindrome that you can create. If it is impossible to create any palindrome, return 0.

A palindrome is a string that reads the same forward and backward.

Example 1:

Input: `words = ["lc","cl","gg"]`

Output: 6

Explanation: One longest palindrome is "lc" + "gg" + "cl" = "lcggcl", of length 6.

Note that "clggcl" is another longest palindrome that can be created.

Example 2:

Input: `words = ["ab","ty","yt","lc","cl","ab"]`

Output: 8

Explanation: One longest palindrome is "ty" + "lc" + "cl" + "yt" = "tylcclyt", of length 8.

Note that "lcyttycl" is another longest palindrome that can be created.

Example 3:

Input: `words = ["cc","ll","xx"]`

Output: 2

Explanation: One longest palindrome is "cc", of length 2.

Note that "ll" is another longest palindrome that can be created, and so is "xx".

Constraints:

- $1 \leq \text{words.length} \leq 105$
- `words[i].length == 2`
- `words[i]` consists of lowercase English letters.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def longestPalindrome(self, words: list[str]) ->
    int:
        ans = 0
        count = [[0] * 26 for _ in range(26)]

        for a, b in words:
            i = ord(a) - ord('a')
            j = ord(b) - ord('a')
            if count[j][i]:
                ans += 4

            count[j][i] -= 1
            else:
                count[i][j] += 1

        for i in range(26):
            if count[i][i]:
                return ans + 2

        return ans
```

• LeetCode

Round 5 · Mid · Medium

p.2299

```
class Solution:
    def longestPalindrome(self, words: List[str])
    -> int:
        cnt = Counter(words)
        ans = x = 0
        for k, v in cnt.items():
            if k[0] == k[1]:
                x += v & 1
                ans += v // 2 * 2 * 2
            else:
                ans += min(v, cnt[k[::-1]]) * 2

        ans += 2 if x else 0
        return ans
```

• SimplyLeet

```
# Python solution
from collections import Counter

class Solution:
    def longestPalindrome(self, words):
        # Count the frequency of each word
        word_count = Counter(words)
        length = 0 # Total words used count (each
        adds 2 letters)
        center_used = False
```

Round 5 · Mid · Medium

p.2300

```

        for word, count in
list(word_count.items()):
            reverse = word[::-1] # Get the reverse
of the current word

            if word == reverse:
                # For words that are palindromic
themselves (e.g. "gg")
                pairs = count // 2
                length += pairs * 2 # Each pair
adds 2 words (4 letters)
                # If an extra word exists and
center is not used, add one word in the center
                if count % 2 == 1 and not
center_used:
                    length += 1

```

```

                center_used = True
            else:
                # For words that are not
self-palindromic.
                if reverse in word_count:
                    # We only count each pair once,
so use minimum and then set the reverse count to 0.
                    pair_count = min(count,
word_count[reverse])
                    length += pair_count * 2 # each
pair contributes two words
                    # Set reverse count to 0 to
avoid double counting later.
                    word_count[reverse] = 0
                return length * 2 # Each word is of length
2

```

```

# Example usage:
# sol = Solution()
# print(sol.longestPalindrome(["lc", "cl", "gg"]))
# Output: 6

```

- LC.ca

```

class Solution:
    def longestPalindrome(self, words: List[str])
-> int:
    cnt = Counter(words)
    ans = x = 0
    for k, v in cnt.items():
        if k[0] == k[1]:
            x += v & 1
            ans += v // 2 * 2 * 2
        else:
            ans += min(v, cnt[k[::-1]]) * 2

    ans += 2 if x else 0
    return ans

```

◆ Quick Review

Key Insights

- Pair words with their reverse (e.g., "ab" with "ba") to form a palindrome.
- Handle words that are palindromic by themselves (like "gg" or "aa") differently; they can be used in pairs and one extra can be placed in the center.
- Use a hash map (or dictionary) to count frequencies of each word.
- For non-palindromic word pairs, consider each pair only once to avoid double counting.
- For self-palindromic words, count how many pairs can be used and whether an extra word can be placed in the

middle to maximize length.

Space & Time

Time Complexity: $O(n)$, where n is the number of words, since we iterate through the list and perform constant time operations per word.

Space Complexity: $O(n)$, to store counts of words in the hash map.

Solution

The approach is to count the occurrences of each word using a hash map. For each word:

- If the word is not self-palindromic, check if its reverse exists and form pairs by considering the minimum count between the word and its reverse.
- If the word is self-palindromic, form as many pairs as possible (using $\text{count} // 2$) and note if there is any word left; you can use one extra self-palindromic word as the center of the palindrome. Finally, compute the total length by multiplying the number of words used by 2, since each word contributes 2 characters.

Round 6 · Mid · Hard

Round 6

Mid Priority · Hard

28 questions · 5 patterns

Linked List #25 #432 #460

Stack #32 #42 #84 #224 #239 #716 #772 #895

Heap #23 #218 #272 #295 #358 #499 #632 #759 #1425

Trie #212 #336 #588 #642

Backtracking #37 #51 #126 #996

Round 5 · Mid · Medium

p.2305

Linked List

Round 6 · Mid · Hard · 3 questions

Round 6 · Mid · Hard

p.2306

Linked List

#25 Reverse Nodes in k-Group

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Recursion

Lists: Grind 169

Problem


Description

Given the head of a linked list, reverse the nodes of the list k at a time, and return the modified list.

k is a positive integer and is less than or equal to the length of the linked list. If the number of nodes is not a multiple of k then left-out nodes, in the end, should remain as it is.

You may not alter the values in the list's nodes, only nodes themselves may be changed.


Example 1:



Input: head = [1,2,3,4,5], k = 2

Output: [2,1,4,3,5]

Example 2:



Input: head = [1,2,3,4,5], k = 3

Output: [3,2,1,4,5]

Constraints:

- The number of nodes in the list is n .
- $1 \leq k \leq n \leq 5000$
- $0 \leq \text{Node.val} \leq 1000$

Follow-up: Can you solve the problem in $O(1)$ extra memory space?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def reverseKGroup(self, head: ListNode | None, k:
int) -> ListNode | None:
        if not head:
            return None

        tail = head

        for _ in range(k):
            # There are less than k nodes in the list, do
            nothing.
            if not tail:

                return head
                tail = tail.next

        newHead = self._reverse(head, tail)
        head.next = self.reverseKGroup(tail, k)
        return newHead

    def _reverse(
        self,
        head: ListNode | None,
```

```
        tail: ListNode | None,
    ) -> ListNode | None:
        """Reverses [head, tail)."""
        prev = None
        curr = head
        while curr != tail:
            next = curr.next
            curr.next = prev
            prev = curr
            curr = next
```

```
        return prev
```

```
class Solution:
    def reverseKGroup(self, head: ListNode, k: int)
-> ListNode:
        if not head or k == 1:
            return head

        def getLength(head: ListNode) -> int:
            length = 0
            while head:
                length += 1
                head = head.next
```

```

    return length

length = getLength(head)
dummy = ListNode(0, head)
prev = dummy
curr = head

for _ in range(length // k):
    for _ in range(k - 1):
        next = curr.next

        curr.next = next.next
        next.next = prev.next
        prev.next = next
        prev = curr
        curr = curr.next

    return dummy.next

```

• LeetCode

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def reverseKGroup(self, head: Optional[ListNode], k: int) -> Optional[ListNode]:
        def reverse(head: Optional[ListNode]) -> Optional[ListNode]:
            dummy = ListNode()
            cur = head

            while cur:
                nxt = cur.next
                cur.next = dummy.next
                dummy.next = cur
                cur = nxt
            return dummy.next

        dummy = pre = ListNode(next=head)
        while pre:
            cur = pre

```



```

    for _ in range(k):
        cur = cur.next
        if cur is None:
            return dummy.next
    node = pre.next
    nxt = cur.next
    cur.next = None
    pre.next = reverse(node)
    node.next = nxt
    pre = node

```

```

return dummy.next

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

class Solution:
    def reverseKGroup(self, head: ListNode, k:
int) -> ListNode:
    # Create a dummy node to simplify edge
cases
    dummy = ListNode(0)

```

```

dummy.next = head
group_prev = dummy

```

```

while True:
    # Find the kth node from group_prev
    kth = self.getKthNode(group_prev, k)
    if not kth:
        break
    group_next = kth.next

```

```

    # Reverse the group starting at
group_prev.next up to kth
    prev = kth.next
    curr = group_prev.next
    while curr != group_next:
        temp = curr.next # Store next node
        curr.next = prev # Reverse pointer
        prev = curr # Move prev forward
        curr = temp # Move curr forward

    # Reconnect the reversed group with the
previous part

```

```

        temp = group_prev.next # temp is the
start of the original group, becomes the new tail
        group_prev.next = kth # Connect
previous part to the kth node (new head of reversed
group)
        group_prev = temp # Move group_prev to
the end of the reversed group

    return dummy.next

# Helper function to get the kth node from the
current node
def getKthNode(self, curr, k):
    while curr and k > 0:
        curr = curr.next

    k -= 1
    return curr

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next
class Solution:
    def reverseKGroup(self, head: ListNode, k:
int) -> ListNode:
        ...

        for reverse 1->2->3->4->5, process is
like

        1->None, 2->3->4->5

        2->1->None, 3->4->5
        3->2->1->None, 4->5
        4->3->2->1->None, 5
        5->4->3->2->1->None, None
        ...

    def reverseList(head):
        pre, p = None, head
        while p:
            pnext = p.next
            p.next = pre

```

```

        pre = p
        p = pnext
    return pre

```

```

dummy = ListNode(next=head)
pre = cur = dummy
while cur.next:
    for _ in range(k):
        cur = cur.next
        if cur is None:

```

```

            return dummy.next
        t = cur.next
        cur.next = None # cut from next
        k-group, so to reverseList() for current k-group
        start = pre.next
        pre.next = reverseList(start)
        start.next = t # so now 'start' is the
        last node of k-group
        pre = cur = start # same as the reset
        dummy before while loop
    return dummy.next

```

```

#####

```

```

# Definition for singly-linked list.
# class ListNode(object):
#     def __init__(self, x):
#         self.val = x
#         self.next = None

```

```

class Solution(object):
    def reverseKGroup(self, head, k):
        """

```

```

        :type head: ListNode
        :type k: int
        :rtype: ListNode
        """

```

```

    def reverseList(head, k):
        pre = None
        cur = head
        while cur and k > 0:
            tmp = cur.next

```

```

    cur.next = pre
    pre = cur
    cur = tmp
    k -= 1
    head.next = cur
    return cur, pre

```

```

length = 0
p = head
while p:

```

```

    length += 1
    p = p.next
    if length < k:
        return head
    step = length / k
    ret = None
    pre = None
    p = head
    while p and step:
        next, newHead = reverseList(p, k)

```

```

    if ret is None:
        ret = newHead
    if pre:
        pre.next = newHead
    pre = p
    p = next
    step -= 1
    return ret

```

◆ Quick Review

Key Insights

- Use a dummy node to simplify edge cases.
- Identify groups of k nodes using a helper function.
- Reverse the nodes in each group using pointer manipulation.
- If a remaining group has fewer than k nodes, leave it unchanged.
- The reversal is done in-place with O(1) extra space.

Space & Time

Time Complexity: O(n) where n is the number of nodes (each node is processed a constant number of times)

Space Complexity: O(1) (only constant extra space is used)

Solution

The solution involves iterating through the linked list by identifying groups of k consecutive nodes. For each group:

- Find the kth node from the current starting point. If it does not exist, the remainder of the list is not reversed.
- Reverse the nodes in the identified group in-place by adjusting the next pointers.
- Reconnect the reversed group with the previous part of the list and update pointers to move to the next group.

The approach leverages a dummy node for easier handling of head modifications and uses a helper function to fetch the kth node. The in-place reversal ensures that the extra memory usage is constant.

Linked List

#432 All O'one Data Structure

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Design · Doubly-Linked List · Linked List

Lists: Denny Zhang

Problem

Description

Design a data structure to store the strings' count with the ability to return the strings with minimum and maximum counts.

Implement the AllOne class:

- AllOne() Initializes the object of the data structure.
- inc(String key) Increments the count of the string key by 1. If key does not exist in the data structure, insert it with count 1.
- dec(String key) Decrements the count of the string key by 1. If the count of key is 0 after the decrement, remove it from the data structure. It is guaranteed that key exists in the data structure before the decrement.
- getMaxKey() Returns one of the keys with the maximal count. If no element exists, return an empty string "".
- getMinKey() Returns one of the keys with the minimum count. If no element exists, return an empty

string "".

Note that each function must run in O(1) average time complexity.

Example 1:

Input

```
["AllOne", "inc", "inc", "getMaxKey", "getMinKey",
"inc", "getMaxKey", "getMinKey"]
[[], ["hello"], ["hello"], [], [], ["leet"], [], []]
```

Output

```
[null, null, null, "hello", "hello", null, "hello",
"leet"]
```

Explanation

```
AllOne allOne = new AllOne();
```

Constraints:

- $1 \leq \text{key.length} \leq 10$
- key consists of lowercase English letters.
- It is guaranteed that for each call to dec, key is existing in the data structure.
- At most $5 * 10^4$ calls will be made to inc, dec, getMaxKey, and getMinKey.

Community Solutions (Python)

- WalkCC

```

from dataclasses import dataclass

@dataclass
class Node:
    def __init__(self, count: int, key: str | None = None):
        self.count = count
        self.keys: set[str] = {key} if key else set()
        self.prev: Node | None = None
        self.next: Node | None = None

    def __eq__(self, other) -> bool:
        if not isinstance(other, Node):
            return NotImplemented
        return self.count == other.count and self.keys == other.keys

class AllOne:
    def __init__(self):
        self.keyToNode: dict[str, Node] = {}

```

```

        self.head = Node(0)
        self.tail = Node(0)
        self.head.next = self.tail
        self.tail.prev = self.head

    def inc(self, key: str) -> None:
        if key in self.keyToNode:
            self._incrementExistingKey(key)
        else:
            self._addNewKey(key)

```

```

    def dec(self, key: str) -> None:
        # It is guaranteed that key exists in the data
        # structure before the
        # decrement.
        self._decrementExistingKey(key)

    def getMaxKey(self) -> str:
        return '' if self.tail.prev == self.head \
            else next(iter(self.tail.prev.keys))

```

```

def getMinKey(self) -> str:
    return '' if self.head.next == self.tail \
        else next(iter(self.head.next.keys))

def _addNewKey(self, key: str) -> None:
    """Adds a new node with frequency 1."""
    if self.head.next.count == 1:
        self.head.next.keys.add(key)
    else:
        self._insertAfter(self.head, Node(1, key))

self.keyToNode[key] = self.head.next

def _incrementExistingKey(self, key: str) ->
None:
    """Increments the frequency of the key by 1."""
    node = self.keyToNode[key]
    node.keys.remove(key)
    if node.next == self.tail or node.next.count >
node.count + 1:
        self._insertAfter(node, Node(node.count +
1))
    node.next.keys.add(key)
    self.keyToNode[key] = node.next

```

```

if not node.keys:
    self._remove(node)

def _decrementExistingKey(self, key: str) ->
None:
    """Decrements the count of the key by 1."""
    node = self.keyToNode[key]
    node.keys.remove(key)
    if node.count > 1:
        if node.prev == self.head or node.prev.count
!= node.count - 1:
            self._insertAfter(node.prev,
Node(node.count - 1))

    node.prev.keys.add(key)
    self.keyToNode[key] = node.prev
    else:
        del self.keyToNode[key]
    if not node.keys:
        self._remove(node)

def _insertAfter(self, node: Node, newNode:
Node) -> None:
    newNode.prev = node
    newNode.next = node.next

```

```
node.next.prev = newNode
node.next = newNode
```

```
def _remove(self, node: Node) -> None:
    node.prev.next = node.next
    node.next.prev = node.prev
```

• LeetCode

```
class Node:
    def __init__(self, key='', cnt=0):
        self.prev = None
        self.next = None
        self.cnt = cnt
        self.keys = {key}
```

```
def insert(self, node):
    node.prev = self
    node.next = self.next
```

```
node.prev.next = node
node.next.prev = node
return node
```

```
def remove(self):
    self.prev.next = self.next
    self.next.prev = self.prev
```

```
class AllOne:
```

```
def __init__(self):
    self.root = Node()
    self.root.next = self.root
    self.root.prev = self.root
    self.nodes = {}
```

```
def inc(self, key: str) -> None:
    root, nodes = self.root, self.nodes
    if key not in nodes:
        if root.next == root or root.next.cnt >
```

```
1:
```

```
nodes[key] = root.insert(Node(key,
1))
```

```
else:
    root.next.keys.add(key)
    nodes[key] = root.next
```

```
else:
    curr = nodes[key]
    next = curr.next
    if next == root or next.cnt > curr.cnt
```

```
+ 1:
```

```
nodes[key] = curr.insert(Node(key,
curr.cnt + 1))
else:
```



```

        next.keys.add(key)
        nodes[key] = next
        curr.keys.discard(key)
        if not curr.keys:
            curr.remove()

def dec(self, key: str) -> None:
    root, nodes = self.root, self.nodes
    curr = nodes[key]
    if curr.cnt == 1:

```

```

        nodes.pop(key)
    else:
        prev = curr.prev
        if prev == root or prev.cnt < curr.cnt
- 1:
            nodes[key] = prev.insert(Node(key,
curr.cnt - 1))
        else:
            prev.keys.add(key)
            nodes[key] = prev
        curr.keys.discard(key)
        if not curr.keys:

```

```
curr.remove()
```

```
def getMaxKey(self) -> str:
    return next(iter(self.root.prev.keys))
```

```
def getMinKey(self) -> str:
    return next(iter(self.root.next.keys))
```

Your AllOne object will be instantiated and called as such:

```

# obj = AllOne()
# obj.inc(key)
# obj.dec(key)
# param_3 = obj.getMaxKey()
# param_4 = obj.getMinKey()

```

• SimplyLeet

```
# Node class to represent a frequency bucket in the
doubly linked list
class Node:
    def __init__(self, count):
        self.count = count # frequency count for
keys in this node
        self.keys = set() # set of keys with this
frequency
        self.prev = None # previous node in the
linked list
        self.next = None # next node in the linked
list
```

```
class AllOne:
    def __init__(self):
```

```
        # Use sentinel nodes for easier handling of
edge cases.
        self.head = Node(float('-inf'))
        self.tail = Node(float('inf'))
        self.head.next = self.tail
        self.tail.prev = self.head
        # Map each key to its corresponding node.
        self.keyToNode = {}
```

```
        # Helper: insert newNode after prevNode in the
linked list.
```

```
        def _insertNodeAfter(self, prevNode, newNode):
```

```
            newNode.prev = prevNode
            newNode.next = prevNode.next
            prevNode.next.prev = newNode
            prevNode.next = newNode
```

```
        # Helper: remove a node if it has no keys.
```

```
        def _removeNode(self, node):
            node.prev.next = node.next
            node.next.prev = node.prev
```

```
        def inc(self, key: str) -> None:
            # If key is new, treat current node as
head.
            if key not in self.keyToNode:
                curr = self.head
            else:
                curr = self.keyToNode[key]
                curr.keys.remove(key)
            # Check if the next node has count =
current count + 1.
            targetCount = (0 if curr == self.head else
curr.count) + 1
            if curr.next == self.tail or
curr.next.count != targetCount:
```

```

        # Create new node with targetCount.
        newNode = Node(targetCount)
        self._insertNodeAfter(curr, newNode)
    else:
        newNode = curr.next
        newNode.keys.add(key)
        self.keyToNode[key] = newNode
    # Remove the current node if it is not a
    sentinel and has no keys.
    if curr != self.head and len(curr.keys) ==
0:
        self._removeNode(curr)

```

```

def dec(self, key: str) -> None:
    curr = self.keyToNode[key]
    curr.keys.remove(key)
    # If decrementing causes the count to be 0,
    remove key completely.
    if curr.count == 1:
        del self.keyToNode[key]
    else:
        targetCount = curr.count - 1
        # Check if the previous node has the
        target count.

```

```

        if curr.prev == self.head or
curr.prev.count != targetCount:
            newNode = Node(targetCount)
            # Insert newNode before current.
            self._insertNodeAfter(curr.prev,
newNode)
        else:
            newNode = curr.prev
            newNode.keys.add(key)
            self.keyToNode[key] = newNode
    # Clean up current node if empty.
    if len(curr.keys) == 0:

```

```

        self._removeNode(curr)

    def getMaxKey(self) -> str:
        # The node before tail has the maximal
        count.
        return next(iter(self.tail.prev.keys)) if
self.tail.prev != self.head else ""

    def getMinKey(self) -> str:
        # The node after head has the minimal
        count.
        return next(iter(self.head.next.keys)) if
self.head.next != self.tail else ""

```

```
# Example usage:
# obj = AllOne()
# obj.inc("hello")
# obj.inc("hello")
# print(obj.getMaxKey()) # "hello"
# print(obj.getMinKey()) # "hello"
# obj.inc("leet")
# print(obj.getMaxKey()) # "hello"
# print(obj.getMinKey()) # "leet"
```

• LC.ca

```
class Node:
    def __init__(self, key='', cnt=0):
        self.prev = None
        self.next = None
        self.cnt = cnt
        self.keys = {key}

    def insert(self, node):
        node.prev = self
        node.next = self.next
```

```
node.prev.next = node
node.next.prev = node
return node
```

```
def remove(self):
    self.prev.next = self.next
    self.next.prev = self.prev
```

```
class AllOne:
```

```
    def __init__(self):
        self.root = Node()
        self.root.next = self.root
        self.root.prev = self.root
        self.nodes = {}

    def inc(self, key: str) -> None:
        root, nodes = self.root, self.nodes
        if key not in nodes:
            if root.next == root or root.next.cnt >
1:
```

```

        nodes[key] = root.insert(Node(key,
1))
    else:
        root.next.keys.add(key)
        nodes[key] = root.next
    else:
        curr = nodes[key]
        next = curr.next
        if next == root or next.cnt > curr.cnt
+ 1:
        nodes[key] = curr.insert(Node(key,
curr.cnt + 1))
    else:

```

```

        next.keys.add(key)
        nodes[key] = next
        curr.keys.discard(key)
        if not curr.keys:
            curr.remove()

def dec(self, key: str) -> None:
    root, nodes = self.root, self.nodes
    curr = nodes[key]
    if curr.cnt == 1:

```

Round 6 · Mid · Hard

p.2337

```

        nodes.pop(key)
    else:
        prev = curr.prev
        if prev == root or prev.cnt < curr.cnt
- 1:
        nodes[key] = prev.insert(Node(key,
curr.cnt - 1))
    else:
        prev.keys.add(key)
        nodes[key] = prev
        curr.keys.discard(key)
        if not curr.keys:

```

```

        curr.remove()

def getMaxKey(self) -> str:
    return next(iter(self.root.prev.keys))

def getMinKey(self) -> str:
    return next(iter(self.root.next.keys))

```

Your AllOne object will be instantiated and called as such:

```

# obj = AllOne()
# obj.inc(key)
# obj.dec(key)
# param_3 = obj.getMaxKey()
# param_4 = obj.getMinKey()

```

Round 6 · Mid · Hard

p.2338

◆ Quick Review

Key Insights

- Use a Doubly-Linked List (DLL) where each node holds a unique count value and a set of keys with that count.
- Maintain a Hash Map (keyToNode) to quickly locate the node corresponding to each key.
- Insert new nodes in the DLL when a key's count is incremented or decremented and the adjacent node does not have the target count.
- Remove nodes from the DLL when no keys remain in that count.
- The head of the DLL holds the minimum count and the tail holds the maximum count.

Space & Time

Time Complexity: $O(1)$ average for inc, dec, getMaxKey, and getMinKey.

Space Complexity: $O(n)$ where n is the number of unique keys in the data structure.

Solution

The solution uses a combination of a doubly-linked list and hash maps. Each node in the DLL represents a frequency and maintains a set of keys that have that frequency. A hash map maps keys to their corresponding node. When inc(key) or dec(key) is called, we update the key's frequency and move the key to the appropriate node in the DLL. If the adjacent node with the required frequency doesn't exist, we

create a new node and insert it in the correct position.

Removal of a key from a node and deletion of empty nodes ensures that getMinKey() and getMaxKey() can be returned directly from the head and tail of the DLL.

Linked List

#460 LFU Cache

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Linked List · Design · Doubly-Linked List

Lists: Denny Zhang

Problem

Description

Design and implement a data structure for a Least Frequently Used (LFU) cache.

Implement the LFUCache class:

- **LFUCache(int capacity)** Initializes the object with the capacity of the data structure.
- **int get(int key)** Gets the value of the key if the key exists in the cache. Otherwise, returns -1.
- **void put(int key, int value)** Update the value of the key if present, or inserts the key if not already present. When the cache reaches its capacity, it should invalidate and remove the least frequently used key before inserting a new item. For this problem, when there is a tie (i.e., two or more keys with the same frequency), the least recently used key would be invalidated.

To determine the least frequently used key, a use counter is maintained for each key in the cache. The key with the smallest use counter is the least frequently

used key.

When a key is first inserted into the cache, its use counter is set to 1 (due to the put operation). The use counter for a key in the cache is incremented either a get or put operation is called on it.

The functions get and put must each run in O(1) average time complexity.

Example 1:

```
Input
["LFUCache", "put", "put", "get", "put", "get", "get",
"put", "get", "get", "get"]
[[2], [1, 1], [2, 2], [1], [3, 3], [2], [3], [4, 4],
[1], [3], [4]]

Output
[null, null, null, 1, null, -1, 3, null, -1, 3, 4]

Explanation
// cnt(x) = the use counter for key x
```

Constraints:

- $1 \leq \text{capacity} \leq 104$
- $0 \leq \text{key} \leq 105$
- $0 \leq \text{value} \leq 109$
- At most $2 * 105$ calls will be made to get and put.

Community Solutions (Python)

- WalkCC

```
class Node:
    def __init__(self, key: int, value: int, freq:
int, it):
        self.key = key
        self.value = value
        self.freq = freq
        self.it = it
```

```
class LFUCache:
    def __init__(self, capacity: int):
```

```
        self.capacity = capacity
        self.minFreq = 0
        self.keyToNode = {}
        self.freqToList = {}
```

```
    def get(self, key: int) -> int:
        if key not in self.keyToNode:
            return -1
```

```
        node = self.keyToNode[key]
```

```
        self._touch(node)
        return node.value
```

```
    def put(self, key: int, value: int) -> None:
        if self.capacity == 0:
            return
```

```
        if key in self.keyToNode:
            node = self.keyToNode[key]
            node.value = value
```

```
        self._touch(node)
        return
```

```
        if len(self.keyToNode) == self.capacity:
            # Evict an LRU key from `minFreq` list
            keyToEvict =
self.freqToList[self.minFreq][-1]
            self.freqToList[self.minFreq].pop()
            del self.keyToNode[keyToEvict]
```

```
        self.minFreq = 1
```



```

if 1 not in self.freqToList:
    self.freqToList[1] = []
    self.freqToList[1].insert(0, key)
    self.keyToNode[key] = Node(key, value, 1, 0) #
Using index 0 as iterator

```

```

def _touch(self, node: Node) -> None:
    # Update the node's frequency
    prevFreq = node.freq
    node.freq += 1
    newFreq = node.freq

```

```

# Remove the key from `prevFreq`'s list
self.freqToList[prevFreq].remove(node.key)
if not self.freqToList[prevFreq]:
    del self.freqToList[prevFreq]
    # Update `minFreq` if needed
    if prevFreq == self.minFreq:
        self.minFreq += 1

# Insert the key to the front of `newFreq`'s
list

```

```

if newFreq not in self.freqToList:
    self.freqToList[newFreq] = []
    self.freqToList[newFreq].insert(0, node.key)
    node.it = 0 # Using index 0 as iterator

```

• LeetCode

Round 6 · Mid · Hard

p.2345

```

class Node:
    def __init__(self, key: int, value: int) ->
None:
        self.key = key
        self.value = value
        self.freq = 1
        self.prev = None
        self.next = None

```

```

class DoublyLinkedList:

```

```

def __init__(self) -> None:
    self.head = Node(-1, -1)
    self.tail = Node(-1, -1)
    self.head.next = self.tail
    self.tail.prev = self.head

```

```

def add_first(self, node: Node) -> None:
    node.prev = self.head
    node.next = self.head.next
    self.head.next.prev = node

```

Round 6 · Mid · Hard

p.2346

```
self.head.next = node
```

```
def remove(self, node: Node) -> Node:
    node.next.prev = node.prev
    node.prev.next = node.next
    node.next, node.prev = None, None
    return node
```

```
def remove_last(self) -> Node:
    return self.remove(self.tail.prev)
```

```
def is_empty(self) -> bool:
    return self.head.next == self.tail
```

```
class LFUCache:
    def __init__(self, capacity: int):
        self.capacity = capacity
        self.min_freq = 0
        self.map = defaultdict(Node)
```

```
self.freq_map =
defaultdict(DoublyLinkedList)
```

```
def get(self, key: int) -> int:
    if self.capacity == 0 or key not in
self.map:
        return -1
    node = self.map[key]
    self.incr_freq(node)
    return node.value
```

```
def put(self, key: int, value: int) -> None:
```

```
    if self.capacity == 0:
        return
    if key in self.map:
        node = self.map[key]
        node.value = value
        self.incr_freq(node)
        return
    if len(self.map) == self.capacity:
        ls = self.freq_map[self.min_freq]
        node = ls.remove_last()
```

```

        self.map.pop(node.key)
        node = Node(key, value)
        self.add_node(node)
        self.map[key] = node
        self.min_freq = 1

def incr_freq(self, node: Node) -> None:
    freq = node.freq
    ls = self.freq_map[freq]
    ls.remove(node)

```

```

    if ls.is_empty():
        self.freq_map.pop(freq)
        if freq == self.min_freq:
            self.min_freq += 1
    node.freq += 1
    self.add_node(node)

def add_node(self, node: Node) -> None:
    freq = node.freq
    ls = self.freq_map[freq]

```

```

ls.add_first(node)
self.freq_map[freq] = ls

```

```

# Your LFUCache object will be instantiated and
# called as such:
# obj = LFUCache(capacity)
# param_1 = obj.get(key)
# obj.put(key,value)

```

• SimplyLeet

```

from collections import OrderedDict

class LFUCache:
    def __init__(self, capacity: int):
        # Initialize cache capacity and helper
        # variables.
        self.capacity = capacity
        self.min_freq = 0 # Track the minimum
        # frequency in cache.
        # Map key -> (value, frequency)
        self.key_to_val_freq = {}
        # Map frequency -> OrderedDict of keys for
        # LRU ordering.

```

```

        self.freq_to_keys = {}

    def update_freq(self, key):
        # Increase frequency of the key since it
        # was accessed.
        value, freq = self.key_to_val_freq[key]
        # Remove key from its current frequency
        # list.
        self.freq_to_keys[freq].pop(key)
        # If current frequency list is empty,
        # update min_freq if necessary.
        if not self.freq_to_keys[freq]:
            del self.freq_to_keys[freq]

        if self.min_freq == freq:
            self.min_freq += 1
        new_freq = freq + 1
        # Insert key into the list corresponding to
        # the new frequency.
        if new_freq not in self.freq_to_keys:
            self.freq_to_keys[new_freq] =
OrderedDict()
        self.freq_to_keys[new_freq][key] = None
        # Update the frequency entry for the key.
        self.key_to_val_freq[key] = (value,
new_freq)

```

```

    def get(self, key: int) -> int:
        # Return -1 if the key is not present.
        if key not in self.key_to_val_freq:
            return -1
        # Update frequency due to access.
        self.update_freq(key)
        return self.key_to_val_freq[key][0]

    def put(self, key: int, value: int) -> None:
        # If capacity is zero, nothing can be
        # added.

        if self.capacity <= 0:
            return
        if key in self.key_to_val_freq:
            # Update value and frequency if key
            # exists.
            self.key_to_val_freq[key] = (value,
self.key_to_val_freq[key][1])
            self.update_freq(key)
        else:
            # If cache is full, evict the least
            # frequently used (and least recently used) key.
            if len(self.key_to_val_freq) >=
self.capacity:
                # Remove the first key from the
                # lowest frequency list.

```

```

        key_to_remove, _ = self.freq_to_keys[
self.min_freq].popitem(last=False)
        del
self.key_to_val_freq[key_to_remove]
        if not
self.freq_to_keys[self.min_freq]:
            del
self.freq_to_keys[self.min_freq]
        # Insert the new key with frequency 1.
self.key_to_val_freq[key] = (value, 1)
        if 1 not in self.freq_to_keys:
            self.freq_to_keys[1] =
OrderedDict()
            self.freq_to_keys[1][key] = None
            self.min_freq = 1 # Reset min_freq.

```

• LC.ca

```

from collections import defaultdict

class LFUCache:
    def __init__(self, capacity: int):
        self.capacity = capacity
        self.key_to_value = {}
        self.key_to_freq = defaultdict(int)
        self.freq_to_keys =
defaultdict(OrderedDict)
        self.min_freq = 0

```

```

def get(self, key: int) -> int:
    if key not in self.key_to_value:
        return -1

    # Update the frequency
    self._update_frequency(key)
    return self.key_to_value[key]

def put(self, key: int, value: int) -> None:
    if self.capacity == 0:

        return

    if key in self.key_to_value:
        # Update the value and frequency
        self.key_to_value[key] = value
        self._update_frequency(key)
    else:
        if len(self.key_to_value) >=
self.capacity:
            # Evict the least frequently used
key
            self._evict()

```

```
# Add the new key-value pair
self.key_to_value[key] = value
self.key_to_freq[key] = 1
self.freq_to_keys[1][key] = None
self.min_freq = 1
```

```
def _update_frequency(self, key: int) -> None:
    freq = self.key_to_freq[key]
    del self.freq_to_keys[freq][key]
```

```
if not self.freq_to_keys[freq]:
    # If there are no keys with the
previous frequency, update min_freq
    if self.min_freq == freq:
        self.min_freq += 1
    del self.freq_to_keys[freq]

freq += 1
self.key_to_freq[key] = freq
self.freq_to_keys[freq][key] = None
```

```
def _evict(self) -> None:
    keys = self.freq_to_keys[self.min_freq]
    evict_key, _ = keys.popitem(last=False)
    del self.key_to_value[evict_key]
    del self.key_to_freq[evict_key]
```

```
# Your LFUCache object will be instantiated and
called as such:
# obj = LFUCache(capacity)
# param_1 = obj.get(key)
```

```
# obj.put(key,value)
```

```
#####
```

```
class Node:
    def __init__(self, key: int, value: int) ->
None:
        self.key = key
        self.value = value
        self.freq = 1
```

```
self.prev = None
self.next = None
```

```
class DoublyLinkedList:
    def __init__(self) -> None:
        self.head = Node(-1, -1)
        self.tail = Node(-1, -1)
        self.head.next = self.tail
        self.tail.prev = self.head
```

```
def add_first(self, node: Node) -> None:
    node.prev = self.head
    node.next = self.head.next
    self.head.next.prev = node
    self.head.next = node
```

```
def remove(self, node: Node) -> Node:
    node.next.prev = node.prev
    node.prev.next = node.next
```

```
node.next, node.prev = None, None
return node
```

```
def remove_last(self) -> Node:
    return self.remove(self.tail.prev)
```

```
def is_empty(self) -> bool:
    return self.head.next == self.tail
```

```
class LFUCache:
    def __init__(self, capacity: int):
        self.capacity = capacity
        self.min_freq = 0
        self.map = defaultdict(Node)
        self.freq_map =
        defaultdict(DoublyLinkedList)

    def get(self, key: int) -> int:
        if self.capacity == 0 or key not in
        self.map:
            return -1
```

```

node = self.map[key]
self.incr_freq(node)
return node.value

def put(self, key: int, value: int) -> None:
    if self.capacity == 0:
        return
    if key in self.map:
        node = self.map[key]
        node.value = value

```

```

        self.incr_freq(node)
        return
    if len(self.map) == self.capacity:
        ls = self.freq_map[self.min_freq]
        node = ls.remove_last()
        self.map.pop(node.key)
    node = Node(key, value)
    self.add_node(node)
    self.map[key] = node
    self.min_freq = 1

```

```

def incr_freq(self, node: Node) -> None:
    freq = node.freq
    ls = self.freq_map[freq]
    ls.remove(node)
    if ls.is_empty():
        self.freq_map.pop(freq)
        if freq == self.min_freq:
            self.min_freq += 1
    node.freq += 1

```

```
self.add_node(node)
```

```

def add_node(self, node: Node) -> None:
    freq = node.freq
    ls = self.freq_map[freq]
    ls.add_first(node)
    self.freq_map[freq] = ls

```

Your LFUCache object will be instantiated and called as such:

```

# obj = LFUCache(capacity)
# param_1 = obj.get(key)
# obj.put(key,value)

```

◆ Quick Review

Key Insights

- Maintain a mapping from key to its value and frequency.
- Use an additional mapping from frequency to keys (maintaining the order of insertion or recent use) for quick eviction.
- Keep a global minFreq variable to track the lowest frequency present in the cache.
- When a key is accessed or updated, increase its frequency and relocate it in the frequency mapping.
- Evictions occur by removing the oldest key from the lowest frequency bucket.

Space & Time

Time Complexity: $O(1)$ average for both get and put.

Space Complexity: $O(\text{capacity})$, where capacity is the maximum number of keys stored.

Solution

The solution utilizes two main data structures:

– A hash map (keyToValFreq) that maps keys to a pair (value, frequency). – A second hash map (freqToKeys) mapping frequencies to an ordered list/set of keys. This data structure preserves the order in which keys were added to allow removal of the least recently used key when multiple keys share the same frequency.

A global minFreq variable is maintained to quickly identify the bucket with the least frequency. When a key is accessed or updated, the key is removed from its current frequency list and placed in the next higher frequency list. If the frequency list for the minimal frequency becomes empty, minFreq is updated.

This design guarantees that both get and put operations execute in constant time on average.

Stack

Round 6 · Mid · Hard · 8 questions

Round 6 · Mid · Hard

p.2363

Stack

#32 Longest Valid Parentheses

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming · Stack

Lists: Grind 169

Problem

Description

Given a string containing just the characters '(' and ')', return the length of the longest valid (well-formed) parentheses substring.

Example 1:

Input: s = "()"

Output: 2

Explanation: The longest valid parentheses substring is "()".

Example 2:

Input: s = ")()())"

Output: 4

Explanation: The longest valid parentheses substring is "()()".

Example 3:

Input: s = ""

Output: 0

Round 6 · Mid · Hard

p.2364

Constraints:

- $0 \leq s.length \leq 3 * 10^4$
- $s[i]$ is '(', or ')'.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestValidParentheses(self, s: str) ->
    int:
        s2 = ')' + s
        # dp[i] := the length of the longest valid
        # parentheses in the substring
        # s2[1..i]
        dp = [0] * len(s2)

        for i in range(1, len(s2)):
            if s2[i] == ')' and s2[i - dp[i - 1] - 1] ==
            '(':
                dp[i] = dp[i - 1] + dp[i - dp[i - 1] - 2] +
                2

        return max(dp)
```

• LeetCode

```
class Solution:
    def longestValidParentheses(self, s: str) ->
    int:
        n = len(s)
        f = [0] * (n + 1)
        for i, c in enumerate(s, 1):
            if c == ")":
                if i > 1 and s[i - 2] == "(":
                    f[i] = f[i - 2] + 2
                else:
                    j = i - f[i - 1] - 1

                    if j and s[j - 1] == "(":
                        f[i] = f[i - 1] + 2 + f[j -
                        1]

        return max(f)

class Solution:
    def longestValidParentheses(self, s: str) ->
    int:
        stack = [-1]
        ans = 0
        for i in range(len(s)):
            if s[i] == '(':
                stack.append(i)
            else:
                stack.pop()
                if not stack:
```

```

        stack.append(i)
    else:
        ans = max(ans, i - stack[-1])
return ans

```

• SimplyLeet

```

# Python solution using a stack
def longestValidParentheses(s: str) -> int:
    # Initialize the stack with the base index -1
    stack = [-1]
    max_length = 0

    # Iterate over the string with index
    for i, char in enumerate(s):
        if char == '(':
            # Push the index of '(' onto the stack

```

```

            stack.append(i)
        else:
            # Pop the last index for ')'
            stack.pop()
            if not stack:
                # If stack is empty, push current
                # index as the new base
                stack.append(i)
            else:
                # Calculate the length of the
                # current valid substring
                current_length = i - stack[-1]

```

```

        max_length = max(max_length,
        current_length)

    return max_length

```

Example usage:

```

print(longestValidParentheses("()()()))") # Output:
4

```

• LC.ca

```

class Solution:
    def longestValidParentheses(self, s: str) ->
    int:
        left = right = 0
        res = 0
        for c in s: # from left to right, '(()' =>
            # will never hit left==right
            if c == '(':
                left += 1
            else:
                right += 1
                if left == right:

```

```

        res = max(res, 2 * left)
    if left < right:
        left = right = 0

    left = right = 0 # dont forget to reset
    for c in reversed(s): # from right to left,
        '())' => will never hit left==right
        if c == '(':
            left += 1
        else:
            right += 1

    if left == right:
        res = max(res, 2 * left)
    if left > right: # reverse '<' to '>'
        left = right = 0

    return res

#####

class Solution:

```

```

def longestValidParentheses(self, s: str) ->
int:
    n = len(s)
    if n < 2:
        return 0
    dp = [0] * n
    for i in range(1, n):
        if s[i] == ')':
            if s[i - 1] == '(':
                dp[i] = 2 + (dp[i - 2] if i > 1
            else 0)
            else:
                j = i - dp[i - 1] - 1
                if j >= 0 and s[j] == '(':
                    dp[i] = 2 + dp[i - 1] +
dp[j - 1]
    return max(dp)

#####

...

>>> s="abcdefg"
>>> [print(i,",",c) for i, c in enumerate(s, 1)]

```

```

1 , a
2 , b
3 , c
4 , d
5 , e
6 , f
7 , g
[None, None, None, None, None, None, None]
>>>
...

class Solution:
    def longestValidParentheses(self, s: str) ->
    int:
        n = len(s)
        f = [0] * (n + 1)
        for i, c in enumerate(s, 1): # starting i
        from 1, not 0
            if c == ")":
                if i > 1 and s[i - 2] == "(":
                    f[i] = f[i - 2] + 2
                else:
                    j = i - f[i - 1] - 1
                    if j and s[j - 1] == "(":
                        f[i] = f[i - 1] + 2 + f[j -
1]
        return max(f)

```

◆ Quick Review

Round 6 · Mid · Hard

p.2371

Key Insights

- Use a stack to keep track of indices where potential valid substrings begin.
- Start by pushing a dummy index (-1) to handle edge cases.
- For every '(', push its index onto the stack.
- For every ')', pop from the stack. If the stack becomes empty, push the current index as a new base. Otherwise, compute the length of the valid substring using the current index minus the top index of the stack.
- An alternative is the dynamic programming approach where $dp[i]$ represents the length of the longest valid substring ending at index i .

Space & Time

Time Complexity: $O(n)$ where n is the length of the string.

Space Complexity: $O(n)$ for the stack data structure.

Solution

The stack-based solution works by maintaining indices for unmatched parentheses. Initially, a dummy index (-1) is pushed onto the stack to serve as the base for calculating subsequent valid substring lengths. As we iterate through the string:

- If the character is '(', we push its index onto the stack.
- If the character is ')', we pop from the stack. If the stack is empty afterwards, that means there is no valid starting position, so we push the current index to reset the base. Otherwise, we calculate the length of the current valid substring by subtracting the current index with the new top

Round 6 · Mid · Hard

p.2372

index of the stack and update the maximum length accordingly.

Stack

#42 Trapping Rain Water

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Dynamic Programming · Stack · Monotonic Stack

Lists: Grind 169

Problem

Description

Given n non-negative integers representing an elevation map where the width of each bar is 1, compute how much water it can trap after raining.

Example 1:



Round 6 · Mid · Hard

p.2373

Round 6 · Mid · Hard

p.2374

Input: height = [0,1,0,2,1,0,1,3,2,1,2,1]

Output: 6

Explanation: The above elevation map (black section) is represented by array [0,1,0,2,1,0,1,3,2,1,2,1]. In this case, 6 units of rain water (blue section) are being trapped.

Example 2:

Input: height = [4,2,0,3,2,5]

Output: 9

Constraints:

- $n == \text{height.length}$
- $1 \leq n \leq 2 * 10^4$
- $0 \leq \text{height}[i] \leq 10^5$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def trap(self, height: list[int]) -> int:
        n = len(height)
        l = [0] * n # l[i] := max(height[0..i])
        r = [0] * n # r[i] := max(height[i..n))

        for i, h in enumerate(height):
            l[i] = h if i == 0 else max(h, l[i - 1])

        for i, h in reversed(list(enumerate(height))):
```

```
        r[i] = h if i == n - 1 else max(h, r[i + 1])
```

```
        return sum(min(l[i], r[i]) - h
                    for i, h in enumerate(height))
```

```
class Solution:
    def trap(self, height: list[int]) -> int:
        if not height:
            return 0
```

```
        ans = 0
        l = 0
        r = len(height) - 1
        maxL = height[l]
        maxR = height[r]
```

```
        while l < r:
            if maxL < maxR:
                ans += maxL - height[l]
                l += 1
                maxL = max(maxL, height[l])
            else:
                ans += maxR - height[r]
                r -= 1
                maxR = max(maxR, height[r])
```

```
        return ans
```

• LeetCode


```
class Solution:
    def trap(self, height: List[int]) -> int:
        n = len(height)
        left = [height[0]] * n
        right = [height[-1]] * n
        for i in range(1, n):
            left[i] = max(left[i - 1], height[i])
            right[n - i - 1] = max(right[n - i],
                                     height[n - i - 1])
        return sum(min(l, r) - h for l, r, h in
                    zip(left, right, height))
```

• SimplyLeet

```
# Python solution using the two pointers approach
def trap(height):
    # Initialize left and right pointers and water
    accumulator
    left, right = 0, len(height) - 1
    left_max, right_max = 0, 0
    water = 0

    # Traverse the elevation map from both ends
    while left < right:
        # Update left_max and right_max as we
        encounter new bars
```

```
        if height[left] < height[right]:
            # left side is the determining side
            if height[left] >= left_max:
                left_max = height[left] # update
                left max since current height is higher
            else:
                water += left_max - height[left] #
                water trapped at left position
                left += 1 # move left pointer to the
                right
        else:
            # right side is the determining side
            if height[right] >= right_max:
```

```
                right_max = height[right] # update
                right max since current height is higher
            else:
                water += right_max - height[right]
            # water trapped at right position
            right -= 1 # move right pointer to the
            left
        return water

# Example usage:
print(trap([0,1,0,2,1,0,1,3,2,1,2,1])) # Expected
output: 6
```

• LC.ca

```

class Solution:
    def trap(self, height: List[int]) -> int:
        n = len(height)
        if n < 3:
            return 0

        lmx = [height[0]] * n
        rmx = [height[n - 1]] * n
        for i in range(1, n):
            lmx[i] = max(lmx[i - 1], height[i]) # i
self compare

            rmx[n - 1 - i] = max(rmx[n - i],
height[n - 1 - i])

        # no negative, worst is, min(left,right) is
        itself
        return sum( min(lmx[i], rmx[i]) -
height[i] for i in range(n) )

#####

class Solution: # zip
    def trap(self, height: List[int]) -> int:
        n = len(height)

```

```

        left = [height[0]] * n
        right = [height[-1]] * n
        for i in range(1, n):
            left[i] = max(left[i - 1], height[i])
            right[n - i - 1] = max(right[n - i],
height[n - i - 1])
        return sum(min(l, r) - h for l, r, h in
zip(left, right, height))

#####

class Solution: # find the highest bar first

    def trap(self, height: List[int]) -> int:
        if not height:
            return 0

        total_water = 0
        max_index = 0
        max_height = float('-inf')

        # Find the index of the maximum height
        for i in range(len(height)):

```

```

        if height[i] > max_height:
            max_height = height[i]
            max_index = i

    # Calculate water trapped on the left side
of the maximum height
    left_max = height[0]
    for i in range(1, max_index):
        if left_max < height[i]:
            left_max = height[i]
        total_water += left_max - height[i]

    # Calculate water trapped on the right side
of the maximum height
    right_max = height[-1]
    for i in range(len(height) - 2, max_index,
-1):
        if right_max < height[i]:
            right_max = height[i]
        total_water += right_max - height[i]

    return total_water

```

◆ Quick Review

Key Insights

- The water trapped at any index is determined by the minimum of the maximum height on its left and right minus the height at that index.

- A two-pointers approach can efficiently solve the problem in one pass by maintaining left and right boundaries.
- Alternative approaches include dynamic programming with pre-computed left max and right max arrays, or using a monotonic stack to determine boundaries.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$ when using the two-pointer approach

Solution

We use a two-pointers strategy by initializing pointers at the beginning and the end of the array. Keep track of the maximum height seen so far from the left (left_max) and the right (right_max). At each step, move the pointer with the lesser height because the water trapped is determined by the shorter boundary. Compute the water at that position as the difference between the boundary (left_max or right_max) and the current height, accumulating it into the total. This approach ensures linear time scanning with constant extra space.

Stack

#84 Largest Rectangle in Histogram HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Stack · Monotonic Stack


Lists: Grind 169

Problem

Description

Given an array of integers heights representing the histogram's bar height where the width of each bar is 1, return the area of the largest rectangle in the histogram.

Example 1:



Round 6 · Mid · Hard

p.2383


Input: heights = [2,1,5,6,2,3]

Output: 10

Explanation: The above is a histogram where width of each bar is 1.

The largest rectangle is shown in the red area, which has an area = 10 units.

Example 2:



Input: heights = [2,4]

Output: 4

Constraints:

Round 6 · Mid · Hard

p.2384

- $1 \leq \text{heights.length} \leq 105$
- $0 \leq \text{heights}[i] \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def largestRectangleArea(self, heights:
list[int]) -> int:
    ans = 0
    stack = []

    for i in range(len(heights) + 1):
        while stack and (i == len(heights) or
heights[stack[-1]] > heights[i]):
            h = heights[stack.pop()]
            w = i - stack[-1] - 1 if stack else i
            ans = max(ans, h * w)

    stack.append(i)

    return ans
```

• LeetCode

```
class Solution:
    def largestRectangleArea(self, heights:
List[int]) -> int:
        n = len(heights)
        stk = []
        left = [-1] * n
        right = [n] * n
        for i, h in enumerate(heights):
            while stk and heights[stk[-1]] >= h:
                right[stk[-1]] = i
                stk.pop()

            if stk:
                left[i] = stk[-1]
            stk.append(i)
        return max(h * (right[i] - left[i] - 1) for
i, h in enumerate(heights))
```

• SimplyLeet

```
# Python solution

def largestRectangleArea(heights):
    # Append a zero to flush out the remaining bars
    in the stack
    heights.append(0)
    stack = [] # Stack to store indices of bars
    max_area = 0

    # Iterate over each index and bar height
    for i, h in enumerate(heights):
```

```

        # Process bars that are taller than the
        current bar
        while stack and heights[stack[-1]] > h:
            height = heights[stack.pop()] # Height
            of the bar at the index popped from the stack
            # Calculate width: if stack is empty,
            width is the current index; otherwise, it's the
            difference
            width = i if not stack else i -
            stack[-1] - 1
            max_area = max(max_area, height *
            width)
            stack.append(i)

        return max_area

```

```

# Example usage:
if __name__ == "__main__":
    histogram = [2, 1, 5, 6, 2, 3]
    print(largestRectangleArea(histogram))

```

• LC.ca

```

'''
    0000000000 h 0000000000000000000000 00(<) h 000 left_i,
    right_i

    0000000000 h * (right_i - left_i - 1)0000000000
    '''

class Solution:
    def largestRectangleArea(self, heights:
    List[int]) -> int:
        n = len(heights)
        stk = []
        left = [-1] * n

        right = [n] * n
        for i, h in enumerate(heights):
            while stk and heights[stk[-1]] >= h:
                right[stk[-1]] = i
                stk.pop()
            if stk:
                left[i] = stk[-1] # same as below:
                stk[-1] in 'i - stack[-1] - 1'
                stk.append(i)

        # valid for one element input [3]

```

```
    return max(h * (right[i] - left[i] - 1) for
i, h in enumerate(heights))
```

```
#####
```

```
class Solution:
```

```
    def largestRectangleArea(self, heights:
List[int]) -> int:
        if not heights:
            return 0
```

```
        heights.append(-1) # make for loop running
        stack = []
        ans = 0
        for i in range(0, len(heights)):
            while stack and heights[i] <
heights[stack[-1]]:
                currentLowestBarIndex =
stack.pop()
                h = heights[currentLowestBarIndex]
                # stack[-1] is after pop(). why not
i-currentLowestBarIndex?
                # because not working for tie case
in the stack,
                # eg. for input=[4,2,0,3,2,5],
expected=6 but output=5
```

```
        # 000000000000 00(<) h 000 left_i,
right_i
        # 0000000000 00 h 0000000000
right-left+1
        # 00000 00 h00000right-left+1 00
0000 00000index000000+1 -2 0 -1
        w = i - stack[-1] - 1 if stack else
i
        ans = max(ans, h * w)
        stack.append(i)

        heights.pop() # restore, pop -1
        return ans
```

◆ Quick Review

Key Insights

- Each bar can potentially extend left and right as long as the neighboring bars are not shorter.
- A monotonic (increasing) stack can be used to efficiently track indices of bars.
- When encountering a bar lower than the one on the top of the stack, pop the stack and calculate the area of the rectangle formed with the popped bar height.
- A dummy bar (of height 0) is appended at the end to ensure all bars are processed.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution employs a monotonic stack to traverse the histogram once. For each bar, while the current bar's height is less than the height at the top index of the stack, it pops from the stack and calculates the area using the popped height as the smallest bar. The width for the rectangle is determined by the difference between the current index and the index from the stack (or the current index if the stack is empty). This effectively finds the maximum rectangular area possible. A sentinel value (0) is added at the end of the histogram to flush any remaining bars from the stack.

Round 6 · Mid · Hard

p.2391

Stack

#224 Basic Calculator

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · String · Stack · Recursion

Lists: Grind 169

Problem

Description

Given a string s representing a valid expression, implement a basic calculator to evaluate it, and return the result of the evaluation.

Note: You are not allowed to use any built-in function which evaluates strings as mathematical expressions, such as `eval()`.

Example 1:

Input: $s = "1 + 1"$
Output: 2

Example 2:

Input: $s = "2-1 + 2"$
Output: 3

Example 3:

Input: $s = "(1+(4+5+2)-3)+(6+8)"$
Output: 23

Round 6 · Mid · Hard

p.2392

Constraints:

- $1 \leq s.length \leq 3 \times 10^5$
- s consists of digits, '+', '-', '(', ')', and ' '.
- s represents a valid expression.
- '+' is not used as a unary operation (i.e., "+1" and "(+2 + 3)" is invalid).
- '-' could be used as a unary operation (i.e., "-1" and "-(2 + 3)" is valid).
- There will be no two consecutive operators in the input.
- Every number and running calculation will fit in a signed 32-bit integer.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def calculate(self, s: str) -> int:
        ans = 0
        num = 0
        sign = 1
        stack = [sign] # stack[-1]: the current
        environment's sign

        for c in s:
            if c.isdigit():
                num = num * 10 + int(c)
```

```
elif c == '(':
    stack.append(sign)
elif c == ')':
    stack.pop()
elif c == '+' or c == '-':
    ans += sign * num
    sign = (1 if c == '+' else -1) * stack[-1]
    num = 0

return ans + sign * num
```

• LeetCode

```
class Solution:
    def calculate(self, s: str) -> int:
        stk = []
        ans, sign = 0, 1
        i, n = 0, len(s)
        while i < n:
            if s[i].isdigit():
                x = 0
                j = i
                while j < n and s[j].isdigit():
```

```

        x = x * 10 + int(s[j])
        j += 1
        ans += sign * x
        i = j - 1
    elif s[i] == "+":
        sign = 1
    elif s[i] == "-":
        sign = -1
    elif s[i] == "(":
        stk.append(ans)

```

```

        stk.append(sign)
        ans, sign = 0, 1
    elif s[i] == ")":
        ans = stk.pop() * ans + stk.pop()
        i += 1
    return ans

```

• SimplyLeet

Python solution for the Basic Calculator problem

```

def calculate(s: str) -> int:
    # Stack to store pairs of (current result,
    # current sign) before a '('
    stack = []
    result = 0 # Running total
    number = 0 # Current number being processed
    sign = 1 # Current sign (1 or -1)

```

```

    # Iterate through each character in the string

```

```

    for char in s:
        if char.isdigit():
            # Build number digit by digit
            number = number * 10 + int(char)
        elif char in "+-":
            # Add the complete number with its sign
            to result
            result += sign * number
            # Reset number for the next potential
            number

            number = 0
            # Update the sign based on the operator

```

```

            sign = 1 if char == '+' else -1
        elif char == '(':
            # Push the current result and sign onto
            the stack
            stack.append(result)
            stack.append(sign)
            # Reset result and sign for the new
            sub-expression
            result = 0
            sign = 1
        elif char == ')':
            # Complete the number before the
            parentheses close

```

```

        result += sign * number
        number = 0
        # The top of the stack is the sign
before the parenthesis
        prev_sign = stack.pop()
        # The next value on the stack is the
result before the parenthesis
        prev_result = stack.pop()
        # Combine the result of the previous
expression with the sub-expression evaluated
        result = prev_result + prev_sign *
result
        # Skip spaces implicitly when no if
condition is met

```

```

# Add any number left after the loop ends
result += sign * number
return result

```

```

# Example usage:
print(calculate("1 + 1")) # Expected output: 2
print(calculate(" 2-1 + 2 ")) # Expected output: 3
print(calculate("(1+(4+5+2)-3)+(6+8)")) # Expected
output: 23

```

• LC.ca

```

# only + and - , no * , no /
class Solution:
    def calculate(self, s: str) -> int:
        stk = []
        ans, sign = 0, 1
        i, n = 0, len(s)
        while i < n:
            if s[i].isdigit():
                x = 0
                j = i

```

```

                # with this while, no need to do
final calculation like below solution
                while j < n and s[j].isdigit():
                    x = x * 10 + int(s[j])
                    j += 1
                ans += sign * x
                i = j - 1
            elif s[i] == "+":
                sign = 1
            elif s[i] == "-":
                sign = -1

```

```

        elif s[i] == "(":
            stk.append(ans)
            stk.append(sign)
            ans, sign = 0, 1
        elif s[i] == ")":
            ans = stk.pop() * ans + stk.pop()
            i += 1
    return ans

```

#####

```

class Solution:
    def calculate(self, s: str) -> int:
        ans = 0
        num = 0
        sign = 1
        stack = [sign] # stack[-1]: current env's sign

        for c in s:
            if c.isdigit():

```

```

        num = num * 10 + int(c)
        # num = num * 10 + (ord(c) - ord('0')) =>
also works
        elif c == '(':
            stack.append(sign)
        elif c == ')':
            stack.pop() # pop pairing sign for this ()
pair
        elif c == '+' or c == '-':
            ans += sign * num
            sign = (1 if c == '+' else -1) * stack[-1]
# after all + or -, the sign for current num
        num = 0

```

```

        return ans + sign * num # final calculation

```

◆ Quick Review

Key Insights

- Use a stack to save previous results and signs when encountering parentheses.
- Process the string left-to-right by accumulating multi-digit numbers and applying the current sign.
- When encountering a '(', push the current result and sign onto the stack and reset them.
- When encountering a ')', complete the current sub-expression then pop the previous state to incorporate the computed value.

- Handle whitespaces by simply skipping them.

Space & Time

Time Complexity: $O(n)$ where n is the length of the input string, since we process each character once.

Space Complexity: $O(n)$ in the worst-case scenario due to the use of a stack for nested parentheses.

Solution

The solution uses an iterative approach with a stack:

- Initialize variables: result to keep the running total, sign to handle addition/subtraction, and index to iterate over the string.
- As you iterate, build multi-digit numbers by checking for consecutive numerical characters.
- Adjust the total by adding the product of the current number and its sign.
- When encountering a '(', push the current result and sign onto the stack, then reset results and sign for the new sub-expression.
- When encountering a ')', complete the sub-expression by popping the previous result and sign, then combine.
- Skip over spaces.
- The final result is the computed value after the end of the string.

This approach effectively simulates recursion and correctly handles nested expressions and negation.

Stack

#239 Sliding Window Maximum

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Queue · Sliding Window · Monotonic Queue

Lists: Grind 169

Problem

Description

You are given an array of integers `nums`, there is a sliding window of size `k` which is moving from the very left of the array to the very right. You can only see the `k` numbers in the window. Each time the sliding window moves right by one position.

Return the max sliding window.

Example 1:

Input: `nums = [1,3,-1,-3,5,3,6,7]`, `k = 3`

Output: `[3,3,5,5,6,7]`

Explanation:

Window position Max

[1 3 -1] -3 5 3 6 7 3

1 [3 -1 -3] 5 3 6 7 3

1 3 [-1 -3 5] 3 6 7 5

Example 2:

Input: nums = [1], k = 1
Output: [1]

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$
- $1 \leq k \leq \text{nums.length}$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxSlidingWindow(self, nums: list[int], k:
int) -> list[int]:
        ans = []
        maxQ = collections.deque()

        for i, num in enumerate(nums):
            while maxQ and maxQ[-1] < num:
                maxQ.pop()
            maxQ.append(num)
            if i >= k and nums[i - k] == maxQ[0]: #
out-of-bounds

                maxQ.popleft()
            if i >= k - 1:
                ans.append(maxQ[0])

        return ans
```

• LeetCode

```
class Solution:
    def maxSlidingWindow(self, nums: List[int], k:
int) -> List[int]:
        q = [(-v, i) for i, v in enumerate(nums[: k
- 1])]
        heapify(q)
        ans = []
        for i in range(k - 1, len(nums)):
            heappush(q, (-nums[i], i))
            while q[0][1] <= i - k:
                heappop(q)
            ans.append(-q[0][0])
```

return ans

```
class Solution:
    def maxSlidingWindow(self, nums: List[int], k:
int) -> List[int]:
        q = deque()
        ans = []
        for i, x in enumerate(nums):
            if q and i - q[0] >= k:
                q.popleft()
            while q and nums[q[-1]] <= x:
                q.pop()
            q.append(i)
```

if i >= k - 1:
 ans.append(nums[q[0]])
return ans

- SimplyLeet

```
from collections import deque

def maxSlidingWindow(nums, k):
    # Create a deque to store indices of elements
    # in the current window
    deq = deque()
    result = []

    for i, num in enumerate(nums):
        # Remove indices from the front if they are
        # out of the current window
        if deq and deq[0] == i - k:

            deq.popleft()

        # Remove elements from the back while the
        # current number is larger
        # than the number at the indices stored in
        # the deque
        while deq and nums[deq[-1]] < num:
            deq.pop()

        # Append the current index to the deque
        deq.append(i)
```

```
        # If we've hit the window size, append the
        # current max to result
        if i >= k - 1:
            result.append(nums[deq[0]])

    return result

# Example usage:
# print(maxSlidingWindow([1,3,-1,-3,5,3,6,7], 3))
```

- LC.ca

```
from collections import deque

class Solution:
    def maxSlidingWindow(self, nums: List[int], k:
int) -> List[int]:
        q = deque()
        ans = []
        for i, v in enumerate(nums):
            if q and i - k + 1 > q[0]:
                q.popleft() # remove index if out
of window left
            while q and nums[q[-1]] <= v: # `<=`,
not `<`, to ensure the bigger index stored in deque
```

```

        q.pop()
        q.append(i)
        if i >= k - 1:
            ans.append(nums[q[0]])
    return ans

```

```
#####
```

```

class Solution(object):
    def maxSlidingWindow(self, nums, k):

```

```

        """

```

```

        :type nums: List[int]

```

```

        :type k: int

```

```

        :rtype: List[int]

```

```

        """

```

```

        if k == 0:

```

```

            return []

```

```

        ans = [0 for _ in range(len(nums) - k + 1)]

```

```

        stack = collections.deque([])

```

```

        for i in range(0, k):

```

```

            while stack and nums[stack[-1]] < nums[i]:
                stack.pop()
            stack.append(i)
            ans[0] = nums[stack[0]]
            idx = 0
            for i in range(k, len(nums)):
                idx += 1
                if stack and stack[0] == i - k:
                    stack.popleft()
                while stack and nums[stack[-1]] < nums[i]:

```

```

                    stack.pop()
                    stack.append(i)
                    ans[idx] = nums[stack[0]]

```

```

        return ans

```

◆ Quick Review

Key Insights

- Use a deque (double-ended queue) to maintain the indices of potential maximum elements in a decreasing order.
- Ensure that the deque always has the current window's indices by removing elements that fall outside the current window.
- Only add elements that are greater than the ones at the back of the deque, thus preserving a monotonic decreasing order.

- The maximum for the current window is at the front of the deque.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in `nums`, since each element is processed at most twice.

Space Complexity: $O(k)$, as the deque stores indices for at most k elements at any time.

Solution

The solution uses a monotonic deque to maintain indices of potential maximum elements in the current window. As the window slides, we:

- Remove indices from the front if they are out of the window.
- Remove indices from the back if the current element is larger, because they cannot be the maximum if the current element is in the window.
- Append the current index.
- After the first k elements, the element at the front of the deque is the maximum for that window. This approach ensures that each element is pushed and popped from the deque only once, achieving an overall $O(n)$ time complexity.

Stack

#716 Max Stack

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Stack · Design · Doubly-Linked List · Ordered Set

Lists: Denny Zhang

Problem

Description

Design a max stack data structure that supports the stack operations and supports finding the stack's maximum element.

Implement the `MaxStack` class:

- `MaxStack()` Initializes the stack object.
- `void push(int x)` Pushes element x onto the stack.
- `int pop()` Removes the element on top of the stack and returns it.
- `int top()` Gets the element on the top of the stack without removing it.
- `int peekMax()` Retrieves the maximum element in the stack without removing it.
- `int popMax()` Retrieves the maximum element in the stack and removes it. If there is more than one maximum element, only remove the top-most one.

You must come up with a solution that supports $O(1)$ for each `top` call and $O(\log n)$ for each other call.

Example 1:

Input
 ["MaxStack", "push", "push", "push", "top", "popMax",
 "top", "peekMax", "pop", "top"]
 [[], [5], [1], [5], [], [], [], [], [], []]

Output
 [null, null, null, null, 5, 5, 1, 5, 1, 5]

Explanation
 MaxStack stk = new MaxStack();

Constraints:

- $-107 \leq x \leq 107$
- At most 105 calls will be made to push, pop, top, peekMax, and popMax.
- There will be at least one element in the stack when pop, top, peekMax, or popMax is called.

Community Solutions (Python)

- LeetDoocs

```
class Node:
    def __init__(self, val=0):
        self.val = val
        self.prev: Union[Node, None] = None
        self.next: Union[Node, None] = None
```

```
class DoubleLinkedList:
    def __init__(self):
        self.head = Node()
```

```
        self.tail = Node()
        self.head.next = self.tail
        self.tail.prev = self.head
```

```
    def append(self, val) -> Node:
        node = Node(val)
        node.next = self.tail
        node.prev = self.tail.prev
        self.tail.prev = node
        node.prev.next = node
```

```
return node
```

```
@staticmethod
```

```
def remove(node) -> Node:
```

```
    node.prev.next = node.next
```

```
    node.next.prev = node.prev
```

```
    return node
```

```
def pop(self) -> Node:
```

```
    return self.remove(self.tail.prev)
```

```
def peek(self):
```

```
    return self.tail.prev.val
```

```
class MaxStack:
```

```
    def __init__(self):
```

```
        self.stk = DoubleLinkedList()
```

```
        self.sl = SortedList(key=lambda x: x.val)
```

```
def push(self, x: int) -> None:
```

```
    node = self.stk.append(x)
```

```
    self.sl.add(node)
```

```
def pop(self) -> int:
```

```
    node = self.stk.pop()
```

```
    self.sl.remove(node)
```

```
    return node.val
```

```
def top(self) -> int:
```

```
    return self.stk.peek()
```

```
def peekMax(self) -> int:
```

```
    return self.sl[-1].val
```

```
def popMax(self) -> int:
```

```
    node = self.sl.pop()
```

```
    DoubleLinkedList.remove(node)
```

```
    return node.val
```

```
# Your MaxStack object will be instantiated and  
called as such:
```

```
# obj = MaxStack()
```

```
# obj.push(x)
```

```
# param_2 = obj.pop()
```

```
# param_3 = obj.top()
```

```
# param_4 = obj.peekMax()
```

```
# param_5 = obj.popMax()
```

• SimplyLeet

```
# Python solution implementing the MaxStack data
structure
# We use collections.OrderedDict for
demonstration.
# However, in Python we do not have a built in
balanced BST;
# one can use "sortedcontainers" or manually
maintain a heap and dictionary.
# For clarity, we simulate an ordered map with a
TreeMap-like structure using sorted list of keys.
# In production, one may use "SortedDict" from
sortedcontainers.
```

```
from collections import defaultdict
```

```
class Node:
```

```
    def __init__(self, val):
        self.val = val
        self.prev = None
        self.next = None
```

```
class DoubleLinkedList:
```

```
    def __init__(self):
        # Create dummy head and tail nodes.
        self.head = Node(0)
        self.tail = Node(0)
```

```
        self.head.next = self.tail
        self.tail.prev = self.head
```

```
    def add_last(self, node):
        # Add node right before the tail.
        node.prev = self.tail.prev
        node.next = self.tail
        self.tail.prev.next = node
        self.tail.prev = node
```

```
    def pop(self, node=None):
        # If no node is provided, pop from the end.
        if not node:
            node = self.tail.prev
        node.prev.next = node.next
        node.next.prev = node.prev
        return node
```

```
    def top(self):
        # Return the last node's value.
```

```

        return self.tail.prev.val

class MaxStack:
    def __init__(self):
        self.dll = DoubleLinkedList()
        # Dictionary mapping value to list of nodes
        self.valToNodes = defaultdict(list)
        # Sorted list to simulate ordered set of
        keys.
        self.sortedKeys = []

    def push(self, x: int) -> None:
        node = Node(x)
        self.dll.add_last(node)
        self.valToNodes[x].append(node)
        # Insert x into sortedKeys if new,
        otherwise maintain count.
        if len(self.valToNodes[x]) == 1:
            # binary insertion into sortedKeys
            lo, hi = 0, len(self.sortedKeys)
            while lo < hi:
                mid = (lo + hi) // 2

```

```

        if self.sortedKeys[mid] < x:
            lo = mid + 1
        else:
            hi = mid
        self.sortedKeys.insert(lo, x)

    def pop(self) -> int:
        node = self.dll.pop()
        x = node.val
        self.valToNodes[x].pop()

    if not self.valToNodes[x]:
        # Remove x from sortedKeys
        self.sortedKeys.remove(x)
        return x

    def top(self) -> int:
        return self.dll.top()

    def peekMax(self) -> int:
        # The maximum is the last element in
        sortedKeys.

```

```

        return self.sortedKeys[-1]

    def popMax(self) -> int:
        max_val = self.peekMax()
        # Get the most recent node with max_val.
        node = self.valToNodes[max_val].pop()
        if not self.valToNodes[max_val]:
            self.sortedKeys.pop() # removes the
last element
        # Remove the node from the doubly-linked
list.
        self.dll.pop(node)

```

```

        return max_val

```

```

# Example usage:
# stk = MaxStack()
# stk.push(5)
# stk.push(1)
# stk.push(5)
# print(stk.top()) # returns 5
# print(stk.popMax()) # returns 5
# print(stk.top()) # returns 1

```

```

# print(stk.peekMax()) # returns 5
# print(stk.pop()) # returns 1
# print(stk.top()) # returns 5

```

• LC.ca

Round 6 · Mid · Hard

p.2419

```

class MaxStack(object):

    def __init__(self):
        """
        initialize your data structure here.
        """
        self.stack = []
        self.max_stack = []
        # or the same trick in min-stack
        # self.max_stack = [-math.inf]

```

```

    def push(self, x):
        """
        :type x: int
        :rtype: void
        """
        self.stack.append(x)
        self.max_stack.append(max(x,
self.max_stack[-1]) if self.max_stack else x)

    def pop(self):

```

Round 6 · Mid · Hard

p.2420

```

    """
    :rtype: int
    """
    if len(self.stack) != 0:
        self.max_stack.pop(-1)
        return self.stack.pop(-1)

```

```

def top(self):
    """
    :rtype: int

```

```

    """
    return self.stack[-1] if self.stack else
None

```

```

def peekMax(self):
    """
    :rtype: int
    """
    if len(self.max_stack) != 0:
        return self.max_stack[-1]

```

```

def popMax(self):

```

```

    """
    :rtype: int
    """
    val = self.peekMax()
    buff = []
    while self.top() != val:
        # re-use pop(), ops on both max_stack
        and stack
        buff.append(self.pop())
    self.pop()

    while len(buff) != 0:
        # re-use push(), no need to save buffer
        for max-stack
        self.push(buff.pop(-1))
    return val

```

```

#####

```

```

from sortedcontainers import SortedList

```

```

class Node:

```

```
def __init__(self, val=0):
    self.val = val
    self.prev: Union[Node, None] = None
    self.next: Union[Node, None] = None
```

```
class DoubleLinkedList:
    def __init__(self):
        self.head = Node()
        self.tail = Node()
```

```
self.head.next = self.tail
self.tail.prev = self.head
```

```
def append(self, val) -> Node:
    node = Node(val)
    node.next = self.tail
    node.prev = self.tail.prev
    self.tail.prev = node
    node.prev.next = node
    return node
```

```
@staticmethod
def remove(node) -> Node:
    node.prev.next = node.next
    node.next.prev = node.prev
    return node

def pop(self) -> Node:
    return self.remove(self.tail.prev)
```

```
def peek(self):
    return self.tail.prev.val
```

```
class MaxStack:
    def __init__(self):
        self.stk = DoubleLinkedList()
        self.sl = SortedList(key=lambda x: x.val)

    def push(self, x: int) -> None:
```



```

        node = self.stk.append(x)
        self.sl.add(node)

    def pop(self) -> int:
        node = self.stk.pop()
        self.sl.remove(node)
        return node.val

    def top(self) -> int:
        return self.stk.peek()

    def peekMax(self) -> int:
        return self.sl[-1].val

    def popMax(self) -> int:
        node = self.sl.pop()
        DoubleLinkedList.remove(node)
        return node.val

```

```

# Your MaxStack object will be instantiated and
called as such:
# obj = MaxStack()
# obj.push(x)
# param_2 = obj.pop()
# param_3 = obj.top()
# param_4 = obj.peekMax()
# param_5 = obj.popMax()

```

◆ Quick Review

Key Insights

- Use a doubly-linked list to support fast insertion and removal from the stack.
- Maintain a balanced tree (or ordered dictionary/multiset) to quickly locate the maximum element in $O(\log n)$ time.
- Map each value to its corresponding node(s) in the doubly-linked list, so that when popMax is called, you can immediately remove the corresponding node from the list.
- The double-linking allows removal of arbitrary nodes in $O(1)$, once you have the reference.

Space & Time

Time Complexity:

- push: $O(\log n)$ due to insertion in the ordered structure.
- pop: $O(1)$ for removal from the doubly-linked list, plus $O(\log n)$ for updating the tree.
- top: $O(1)$
- peekMax: $O(1)$ or $O(\log n)$ depending on the tree implementation.
- popMax: $O(\log n)$ for looking up and removal from the tree, $O(1)$ for doubly-linked list deletion.

Space Complexity:

- $O(n)$ for storing all nodes in the doubly-linked list and the tree structure mapping values to nodes.

Solution

We use two data structures:

– A doubly-linked list to store the elements in stack order. This supports fast push and pop operations and allows $O(1)$ removal when given a reference. – An ordered structure (like a balanced BST, TreeMap in Java, or SortedList in Python) that maps each value to a list of nodes that hold that value in the linked list. This allows us to quickly determine the maximum element. When there are duplicates, we remove the one closest to the top by storing nodes in order. The trick is to keep these two structures synchronized. When an element is pushed, add its node to both the doubly-linked list and the ordered structure. When an element is removed (either via pop or popMax), remove its node from both structures.

Stack

#772 Basic Calculator III

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · String · Stack · Recursion

Lists: Premium 98

Problem

Description

Implement a basic calculator to evaluate a simple expression string.

The expression string contains only non-negative integers, '+', '-', '*', '/' operators, and open '(' and closing parentheses ')'. The integer division should truncate toward zero.

You may assume that the given expression is always valid. All intermediate results will be in the range of $[-2^{31}, 2^{31} - 1]$.

Note: You are not allowed to use any built-in function which evaluates strings as mathematical expressions, such as eval().

Example 1:

Input: s = "1+1"
Output: 2

Example 2:

Input: s = "6-4/2"

Output: 4

Example 3:

Input: s = "2*(5+5*2)/3+(6/2+8)"

Output: 21

Constraints:

- $1 \leq s \leq 104$
- s consists of digits, '+', '-', '*', '/', '(', and ')'.
Note: s will not contain any spaces.
- s is a valid expression.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def calculate(self, s: str) -> int:
        nums = []
        ops = []

        def calc():
            b = nums.pop()
            a = nums.pop()
            op = ops.pop()
            if op == '+':
```

```
        nums.append(a + b)
    elif op == '-':
        nums.append(a - b)
    elif op == '*':
        nums.append(a * b)
    else: # op == '/'
        nums.append(int(a / b))
```

```
def precedes(prev: str, curr: str) -> bool:
    """
```

Returns True if the previous character is a operator and the priority of the previous operator >= the priority of the current character (operator).

```
    """
    if prev == '(':
        return False
    return prev in '*/' or curr in '+-'
```

```
i = 0
hasPrevNum = False
```

```

while i < len(s):
    c = s[i]
    if c.isdigit():
        num = int(c)
        while i + 1 < len(s) and s[i +
1].isdigit():
            num = num * 10 + int(s[i + 1])
            i += 1
        nums.append(num)
        hasPrevNum = True
    elif c == '(':
        ops.append('(')
        hasPrevNum = False
    elif c == ')':
        while ops[-1] != '(':
            calc()
        ops.pop() # Pop '('
    elif c in '+-*/':
        if not hasPrevNum: # Handle input like
"-1-(-1)"
            num.append(0)
        while ops and precedes(ops[-1], c):

```

```

        calc()
        ops.append(c)
        i += 1

while ops:
    calc()

return nums.pop()

```

• LeetCode

```

class Solution:
    def calculate(self, s: str) -> int:
        def dfs(q):
            num, sign, stk = 0, "+", []
            while q:
                c = q.popleft()
                if c.isdigit():
                    num = num * 10 + int(c)
                if c == "(":
                    num = dfs(q)

```

```

        if c in "+-*/" or not q:
            match sign:
                case "+":
                    stk.append(num)
                case "-":
                    stk.append(-num)
                case "*":
                    stk.append(stk.pop() *
num)
                case "/":
                    stk.append(int(stk.pop
() / num))

        num, sign = 0, c
        if c == ")":
            break
        return sum(stk)

    return dfs(deque(s))

```

• SimplyLeet

```

def calculate(s: str) -> int:
    # Recursive helper function that processes the
    expression starting from index 'i'
    def helper(i):
        stack = []
        num = 0
        sign = '+'
        while i < len(s):
            char = s[i]
            # Build the current number if the
            character is a digit
            if char.isdigit():
                num = num * 10 + int(char)
            # If '(' is encountered, evaluate the
            subexpression recursively
            if char == '(':
                num, i = helper(i + 1)
            # When encountering an operator or
            closing parenthesis, process the sign and number
            if (not char.isdigit() and char != ' ')
or i == len(s) - 1:
                if sign == '+':
                    stack.append(num)
                elif sign == '-':
                    stack.append(-num)

```

```

        elif sign == '*':
            prev = stack.pop()
            stack.append(prev * num)
        elif sign == '/':
            prev = stack.pop()
            stack.append(int(prev / num))
# Truncate toward zero
        sign = char
        num = 0
        # If a closing parenthesis is reached,
        end this level of recursion
        if char == ')':

            return sum(stack), i
            i += 1
            return sum(stack), i

        result, _ = helper(0)
        return result

# Example test
print(calculate("2*(5+5*2)/3+(6/2+8)"))

```

• LC.ca

Round 6 · Mid · Hard

p.2435

```

# good one, based on solution of Basic Calculator
II
class Solution: # dfs recursion
    def calculate(self, s: str) -> int:
        stack = []
        num = 0
        op = "+"
        i = 0
        while i < len(s):
            c = s[i]
            if c.isdigit():
                num = num * 10 + int(c)
            elif c == "(":
                j = self.findClosing(s, i)
                num = self.calculate(s[i+1:j]) # j
            elif c == ')':
                i = j # out if, i will +1 below
            if c in "+-*/" or i == len(s) - 1:
                if op == "+":
                    stack.append(num)
                elif op == "-":
                    stack.append(-num)
                num = 0
                op = c
            i += 1
        return sum(stack)

```

Round 6 · Mid · Hard

p.2436

```

        elif op == "*":
            stack[-1] *= num
        else:
            stack[-1] = int(stack[-1] /
num)

        op = c
        num = 0
        i += 1
    return sum(stack)

def findClosing(self, s: str, i: int) -> int:

    level = 0
    while i < len(s):
        if s[i] == "(":
            level += 1
        elif s[i] == ")":
            level -= 1
            if level == 0:
                return i
        i += 1
    return i

```

```
#####
```

```

class Solution: # iterative
    def calculate(self, s: str) -> int:
        def evaluate_expression(stack):
            res = stack.pop() if stack else 0
            # Evaluate the expression till we get
            '(' or empty stack

            while stack and stack[-1] != '(':
                sign = stack.pop()
                if sign == '+':
                    res += stack.pop()
                elif sign == '-':
                    res -= stack.pop()
                elif sign == '*':
                    res *= stack.pop()
                elif sign == '/':
                    res //= stack.pop()

```

```

        return res

    stack = [] # mix of num and operator
    n = len(s)
    i = 0
    while i < n:
        if s[i].isdigit():
            # Get the complete number and push
            it to the stack
            j = i
            while j < n and s[j].isdigit():

                j += 1
                stack.append(int(s[i:j]))
                i = j - 1
            elif s[i] in '+-*/(':
                stack.append(s[i])
            elif s[i] == ')':
                # Evaluate the expression within
                the parentheses
                res = evaluate_expression(stack)
                stack.pop() # Remove the opening
                '('
                stack.append(res)

```

```

        i += 1

    return evaluate_expression(stack)

#####

class Solution:
    def calculate(self, s: str) -> int:
        if not s:

            return 0

        # remove leading and trailing spaces and
        white spaces.
        s = s.strip().replace(" ", "")

        if not s:
            return 0

        op_stack = []
        num_stack = []

```



```

i = 0
while i < len(s):
    if s[i].isdigit():
        num = 0
        while i < len(s) and
s[i].isdigit():
            num = num * 10 + int(s[i])
            i += 1
        num_stack.append(num)
    else:

```

```

    op = s[i]
    if not op_stack:
        op_stack.append(op)
        i += 1
    elif op == '+' or op == '-':
        top = op_stack[-1]
        if top == '(':
            op_stack.append(op)
            i += 1
        else:

```

```

        self.calculate_helper(num_
stack, op_stack)
        elif op == '*' or op == '/':
            top = op_stack[-1]
            if top == '(' or top == '+' or
top == '-':
                op_stack.append(op)
                i += 1
            else:
                self.calculate_helper(num_
stack, op_stack)
        elif op == '(':
            op_stack.append(op)

```

```

            i += 1
        elif op == ')':
            while op_stack[-1] != '(':
                self.calculate_helper(num_
stack, op_stack)
            op_stack.pop()
            i += 1

        while op_stack:
            self.calculate_helper(num_stack,
op_stack)

```

```

        return num_stack[-1]

    def calculate_helper(self, num_stack,
op_stack):
        num2 = num_stack.pop()
        num1 = num_stack.pop()

        op = op_stack.pop()

        if op == '+':
            ans = num1 + num2

        elif op == '-':
            ans = num1 - num2
        elif op == '*':
            ans = num1 * num2
        else:
            ans = int(num1 / num2)

        num_stack.append(ans)

```

◆ Quick Review

Key Insights

- Use recursion to handle nested expressions defined by parentheses.
- A stack can be used to correctly evaluate the expression and respect operator precedence.
- Multiplication and division have higher precedence than addition and subtraction.

- When encountering an open parenthesis, recursively compute its value until a closing parenthesis is found.
- Use immediate calculation for multiplication and division by updating the top of the stack before moving to the next operator.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, as every character is processed.

Space Complexity: $O(n)$ due to the recursion stack or auxiliary stack used for intermediate results.

Solution

We use a recursive approach combined with a stack to solve the problem. As we iterate through the string:

- Convert digits into numbers.
- When encountering an operator or parenthesis, process the previously accumulated number based on the preceding operator.
- For multiplication or division, pop the last number from the stack, apply the operation, and push the result back.
- When a '(' is encountered, recursively evaluate the substring until the corresponding ')'. – After processing all tokens, sum the values in the stack to get the final result. This method correctly handles both operator precedence and nested expressions.

Stack

#895 Maximum Frequency Stack

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Stack · Design

Lists: Grind 169

Problem

Description

Design a stack-like data structure to push elements to the stack and pop the most frequent element from the stack.

Implement the FreqStack class:

- FreqStack() constructs an empty frequency stack.
- void push(int val) pushes an integer val onto the top of the stack.
- int pop() removes and returns the most frequent element in the stack. If there is a tie for the most frequent element, the element closest to the stack's top is removed and returned.

Example 1:

Input

```
["FreqStack", "push", "push", "push", "push", "push",
"push", "pop", "pop", "pop", "pop"]
```

```
[[], [5], [7], [5], [7], [4], [5], [], [], [], []]
```

Output

```
[null, null, null, null, null, null, null, null, 5, 7, 5, 4]
```

Explanation

```
FreqStack freqStack = new FreqStack();
```

Constraints:

- $0 \leq \text{val} \leq 109$
- At most $2 * 10^4$ calls will be made to push and pop.
- It is guaranteed that there will be at least one element in the stack before calling pop.

Community Solutions (Python)

- WalkCC

```

class FreqStack:
    def __init__(self):
        self.maxFreq = 0
        self.count = collections.Counter()
        self.countToStack =
collections.defaultdict(list)

    def push(self, val: int) -> None:
        self.count[val] += 1
        self.countToStack[self.count[val]].append(val)
        self.maxFreq = max(self.maxFreq,
self.count[val])

```

```

    def pop(self) -> int:
        val = self.countToStack[self.maxFreq].pop()
        self.count[val] -= 1
        if not self.countToStack[self.maxFreq]:
            self.maxFreq -= 1
        return val

```

• LeetCode

Round 6 · Mid · Hard

p.2447

```

class FreqStack:
    def __init__(self):
        self.cnt = defaultdict(int)
        self.q = []
        self.ts = 0

    def push(self, val: int) -> None:
        self.ts += 1
        self.cnt[val] += 1
        heappush(self.q, (-self.cnt[val],
-self.ts, val))

```

```

    def pop(self) -> int:
        val = heappop(self.q)[2]
        self.cnt[val] -= 1
        return val

```

Your FreqStack object will be instantiated and called as such:

```

# obj = FreqStack()
# obj.push(val)

```

```

# param_2 = obj.pop()

```

Round 6 · Mid · Hard

p.2448

```
class FreqStack:
    def __init__(self):
        self.cnt = defaultdict(int)
        self.d = defaultdict(list)
        self.mx = 0

    def push(self, val: int) -> None:
        self.cnt[val] += 1
        self.d[self.cnt[val]].append(val)
        self.mx = max(self.mx, self.cnt[val])
```

```
    def pop(self) -> int:
        val = self.d[self.mx].pop()
        self.cnt[val] -= 1
        if not self.d[self.mx]:
            self.mx -= 1
        return val
```

Your FreqStack object will be instantiated and called as such:

```
# obj = FreqStack()
# obj.push(val)
# param_2 = obj.pop()
```

• SimplyLeet

Round 6 · Mid · Hard

p.2449

Python implementation of FreqStack

```
class FreqStack:
    def __init__(self):
        # Dictionary mapping value to its frequency
        self.freq = {}
        # Dictionary mapping frequency to a list
        # (stack) of values
        self.group = {}
        # Variable tracking the current maximum
        # frequency in the stack
        self.maxfreq = 0
```

```
    def push(self, val: int) -> None:
        # Update frequency count for the value
        f = self.freq.get(val, 0) + 1
        self.freq[val] = f
        # Update the maximum frequency if necessary
        if f > self.maxfreq:
            self.maxfreq = f
        # Append the value to the appropriate
        # frequency group
        if f not in self.group:
```

Round 6 · Mid · Hard

p.2450

```

        self.group[f] = []
        self.group[f].append(val)

    def pop(self) -> int:
        # Remove the value from the stack of the
        # current maximum frequency
        val = self.group[self.maxfreq].pop()
        # Decrement the frequency count for the
        # value
        self.freq[val] -= 1
        # If no more elements have the current max
        # frequency, decrease maxfreq
        if not self.group[self.maxfreq]:
            self.maxfreq -= 1
        return val

```

• LC.ca

```

class FreqStack:
    def __init__(self):
        self.cnt = defaultdict(int)
        self.d = defaultdict(list)
        self.mx = 0

    def push(self, val: int) -> None:
        self.cnt[val] += 1
        self.d[self.cnt[val]].append(val)
        self.mx = max(self.mx, self.cnt[val])

```

```

def pop(self) -> int:
    val = self.d[self.mx].pop()
    self.cnt[val] -= 1
    if not self.d[self.mx]:
        self.mx -= 1
    return val

```

Your FreqStack object will be instantiated and called as such:

```

# obj = FreqStack()
# obj.push(val)
# param_2 = obj.pop()

```

◆ Quick Review

Key Insights

- Use a hash map to count the frequency of each element.
- Maintain another hash map that maps frequencies to stacks (lists) of elements.
- Track the current maximum frequency to enable O(1) retrieval of the most frequent element.
- When pushing an element, update its frequency and add it to the corresponding frequency stack.
- When popping an element, remove it from the stack corresponding to the maximum frequency and update the frequency count; if that stack becomes empty, decrease the

current maximum frequency.

Space & Time

Time Complexity: $O(1)$ per push and pop.

Space Complexity: $O(n)$, where n is the number of elements pushed onto the stack.

Solution

The solution involves using two main data structures:

- A hash map (freq) to store the frequency count for each value.
- A hash map (group) that maps a frequency to a stack (list) of values that have that frequency.

During a push:

- Increment the frequency count of the value.
- Push the value onto the stack corresponding to its updated frequency.
- Update the maximum frequency if necessary.

During a pop:

- Pop the top element from the stack corresponding to the current maximum frequency.
- Decrement its frequency count in the freq map.
- If the frequency stack becomes empty after popping, decrement the current maximum frequency.

This design enables the retrieval of the most frequent element in constant time while dealing with frequency ties by returning the element closest to the top.

Heap

Round 6 · Mid · Hard · 9 questions

Round 6 · Mid · Hard

p.2453

Round 6 · Mid · Hard

p.2454

Heap

#23 Merge k Sorted Lists

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Linked List · Divide and Conquer · Heap · Merge Sort

Lists: Grind 169

Problem

Description

You are given an array of k linked-lists `lists`, each linked-list is sorted in ascending order.

Merge all the linked-lists into one sorted linked-list and return it.

Example 1:

```
Input: lists = [[1,4,5],[1,3,4],[2,6]]
Output: [1,1,2,3,4,4,5,6]
Explanation: The linked-lists are:
[
  1->4->5,
  1->3->4,
  2->6
]
```

Example 2:

```
Input: lists = []
Output: []
```

Example 3:

```
Input: lists = [[]]
Output: []
```

Constraints:

- $k == \text{lists.length}$
- $0 \leq k \leq 104$
- $0 \leq \text{lists}[i].\text{length} \leq 500$
- $-104 \leq \text{lists}[i][j] \leq 104$
- `lists[i]` is sorted in ascending order.
- The sum of `lists[i].length` will not exceed 104.

Community Solutions (Python)

• WalkCC

```
from queue import PriorityQueue

class Solution:
    def mergeKLists(self, lists: list[ListNode]) ->
    ListNode:
        dummy = ListNode(0)
        curr = dummy
        pq = PriorityQueue()

        for i, lst in enumerate(lists):
```



```

    if lst:
        pq.put((lst.val, i, lst))

    while not pq.empty():
        _, i, minNode = pq.get()
        if minNode.next:
            pq.put((minNode.next.val, i,
minNode.next))
        curr.next = minNode
        curr = curr.next

```

```

return dummy.next

```

• LeetCode

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def mergeKLists(self, lists:
List[Optional[ListNode]]) -> Optional[ListNode]:
        setattr(ListNode, "__lt__", lambda a, b:
a.val < b.val)
        pq = [head for head in lists if head]
        heapify(pq)

```

```

dummy = cur = ListNode()
while pq:
    node = heappop(pq)
    if node.next:
        heappush(pq, node.next)
    cur.next = node
    cur = cur.next
return dummy.next

```

• SimplyLeet

```

# Definition for singly-linked list.
class ListNode:
    def __init__(self, val=0, next=None):
        self.val = val
        self.next = next

import heapq

class Solution:
    def mergeKLists(self, lists):

```

```

# Create a min-heap
min_heap = []
# Heap entry: (node value, unique id, node)
to handle duplicate values
for i, node in enumerate(lists):
    if node:
        heapq.heappush(min_heap,
(node.val, i, node))

# Dummy node to form the start of merged
list
dummy = ListNode(0)
current = dummy

```

```

# While there are elements in the heap
while min_heap:
    # Pop the smallest node from the heap
    val, idx, node =
heapq.heappop(min_heap)
    # Add the smallest node to the merged
list
    current.next = node
    current = current.next
    # If there is a next node in the list,
push it into the heap
    if node.next:

```

```

        heapq.heappush(min_heap,
(node.next.val, idx, node.next))

return dummy.next

```

• LC.ca

```

# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
class Solution:
    def mergeKLists(self, lists:
List[Optional[ListNode]]) -> Optional[ListNode]:
        setattr(ListNode, "__lt__", lambda a, b:
a.val < b.val)
        pq = [head for head in lists if head]
        heapify(pq)

```

```

dummy = cur = ListNode()
while pq:
    node = heappop(pq)
    if node.next:
        heappush(pq, node.next)
    cur.next = node
    cur = cur.next
return dummy.next

```

```
# Definition for singly-linked list.
# class ListNode:
# def __init__(self, val=0, next=None):
# self.val = val
# self.next = next
```

```
#####
```

```
import heapq
```

```
'''
__lt__(self, other) for <
__le__(self, other) for <=
__gt__(self, other) for >
__ge__(self, other) for >=
'''

'''
```

```
# or create a wrapper class, without modifying
existing node class
```

```
class NodeWrapper:
    def __init__(self, node):
        self.node = node

    def __lt__(self, other):
        return self.node.val < other.node.val
```

```
'''
```

```
# overwrite the comparison function, so the node
can be comparable
# or else, error, TypeError: '<' not supported
between instances of 'ListNode' and 'ListNode'
```

```
# define 'gt' will also work, so either lt or gt
will do the compare job for ListNode
# ListNode.__gt__ = lambda x, y: (x.val > y.val)
```

```
ListNode.__lt__ = lambda x, y: (x.val < y.val)
```

```
class Solution:
```

```

def mergeKLists(self, lists:
List[Optional[ListNode]]) -> Optional[ListNode]:
    dummy = current = ListNode()
    heap = []
    for i, node in enumerate(lists):
        if node:
            # need to override
            # or else error: TypeError: '<' not
supported between instances of 'ListNode' and
'ListNode'
            heapq.heappush(heap, node)

    while heap:

        anode = heapq.heappop(heap)

        current.next = anode
        current = current.next

        if anode.next:
            heapq.heappush(heap, anode.next)

    return dummy.next

```

```

'''
tried to use node.val to order heap, but still got
error TypeError: '<' not supported between
instances of 'ListNode' and 'ListNode':

heapq.heappush(heap, (node.val, node))

'''

#####

# based on merge 2 lists

# note for empty input
class Solution:
    def mergeKLists(self, lists:
List[Optional[ListNode]]) -> Optional[ListNode]:
        if lists is None or len(lists) == 0:
            return None
        return self.mergeKListsHelper(lists, 0,
len(lists) - 1)

    def mergeKListsHelper(self, lists:
List[Optional[ListNode]], start: int, end: int) ->
Optional[ListNode]:

        if start == end:

```

```

        return lists[start]

        mid = int((start + end) / 2)

        left = self.mergeKListsHelper(lists,
start, mid)
        # print(left)

        right = self.mergeKListsHelper(lists, mid
+ 1, end)
        # print(right)

```

```

        return self.mergeTwoLists(left, right)

        def mergeTwoLists(self, list1:
Optional[ListNode], list2: Optional[ListNode]) ->
Optional[ListNode]:
            dummy = ListNode()
            current = dummy
            # print("merging: ", list1)
            # print("merging: ", list2)
            while list1 or list2:
                v1 = list1.val if list1 else
float('inf')
                v2 = list2.val if list2 else
float('inf')

```

```

        if v1 < v2:
            current.next = list1
            list1 = list1.next
        else:
            current.next = list2
            list2 = list2.next
            current = current.next

        # print("merging result: ", dummy.next)

```

```

        return dummy.next

#####

# Definition for singly-linked list.
# class ListNode:
#     def __init__(self, val=0, next=None):
#         self.val = val
#         self.next = next

```

```

class Solution:
    def mergeKLists(self, lists: List[ListNode])
-> ListNode:
    n = len(lists)
    if n == 0:
        return None
    for i in range(n - 1): # this is o(N) times
of merge, while if use divide-mid-merge then it's
o(logN) times of merge
        lists[i + 1] =
self.mergeTwoLists(lists[i], lists[i + 1])
    return lists[-1]

```

```

    def mergeTwoLists(self, l1: ListNode, l2:
ListNode) -> ListNode:

```

```

    dummy = ListNode()
    cur = dummy
    while l1 and l2:
        if l1.val <= l2.val:
            cur.next = l1
            l1 = l1.next
        else:
            cur.next = l2
            l2 = l2.next
        cur = cur.next

```

```

    cur.next = l1 or l2
    return dummy.next

```

```

if __name__ == '__main__':

```

```

    l1 = ListNode(1)
    l1.next = ListNode(4)
    l1.next.next = ListNode(5)

```

```

    l2 = ListNode(1)
    l2.next = ListNode(3)
    l2.next.next = ListNode(4)

```

```

    l3 = ListNode(2)
    l3.next = ListNode(6)

```

```

    print(Solution().mergeKLists([l1, l2, l3]))

```

```

    print(Solution().mergeKLists([]))
    print(Solution().mergeKLists([[]]))

```

◆ Quick Review

Key Insights

- Use a min-heap (priority queue) to efficiently extract the smallest value among the current heads of the lists.

- By pushing the head of each non-empty list into the heap, we can pull the smallest node, then push that node's next element into the heap.
- Alternatively, a divide and conquer approach can be used by recursively merging pairs of lists.
- Careful management of pointers (or references) is key to constructing the new linked list.

Space & Time

Time Complexity: $O(N \log k)$ where N is the total number of nodes across all lists and k is the number of lists.

Space Complexity: $O(k)$ due to the heap containing at most one node from each list.

Solution

We use a min-heap to always retrieve the smallest node among the current nodes in the linked lists. First, initialize the heap with the head of each list. Then, continuously extract the smallest node from the heap, attach it to the merged list, and if the extracted node has a next node, insert that into the heap. The process repeats until the heap is empty.

Heap

#218 The Skyline Problem

HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Divide and Conquer · Binary Indexed Tree · Segment Tree · Line Sweep · Heap · Ordered Set

Lists: Denny Zhang

Problem

Description

A city's skyline is the outer contour of the silhouette formed by all the buildings in that city when viewed from a distance. Given the locations and heights of all the buildings, return the skyline formed by these buildings collectively.

The geometric information of each building is given in the array `buildings` where `buildings[i] = [lefti, righti, heighti]`:


- `lefti` is the x coordinate of the left edge of the *i*th building.
- `righti` is the x coordinate of the right edge of the *i*th building.
- `heighti` is the height of the *i*th building.

You may assume all buildings are perfect rectangles grounded on an absolutely flat surface at height 0.

The skyline should be represented as a list of "key points" sorted by their x-coordinate in the form $[[x_1, y_1], [x_2, y_2], \dots]$. Each key point is the left endpoint of some horizontal segment in the skyline except the last point in the list, which always has a y-coordinate 0 and is used to mark the skyline's termination where the rightmost building ends. Any ground between the leftmost and rightmost buildings should be part of the skyline's contour.

Note: There must be no consecutive horizontal lines of equal height in the output skyline. For instance, $[\dots, [2, 3], [4, 5], [7, 5], [11, 5], [12, 7], \dots]$ is not acceptable; the three lines of height 5 should be merged into one in the final output as such: $[\dots, [2, 3], [4, 5], [12, 7], \dots]$

Example 1:



Input: buildings = $[[2, 9, 10], [3, 7, 15], [5, 12, 12], [15, 20, 10], [19, 24, 8]]$

Output: $[[2, 10], [3, 15], [7, 12], [12, 0], [15, 10], [20, 8], [24, 0]]$

Explanation:

Figure A shows the buildings of the input.

Figure B shows the skyline formed by those buildings.

The red points in figure B represent the key points in the output list.

Example 2:

Input: buildings = $[[0, 2, 3], [2, 5, 3]]$

Output: $[[0, 3], [5, 0]]$

Constraints:

- $1 \leq \text{buildings.length} \leq 104$
- $0 \leq \text{left}_i < \text{right}_i \leq 231 - 1$
- $1 \leq \text{height}_i \leq 231 - 1$
- buildings is sorted by left_i in non-decreasing order.

Community Solutions (Python)

- [WalkCC](#)


```

class Solution:
    def getSkyline(self, buildings: list[list[int]])
-> list[list[int]]:
    n = len(buildings)
    if n == 0:
        return []
    if n == 1:
        left, right, height = buildings[0]
        return [[left, height], [right, 0]]

    left = self.getSkyline(buildings[:n // 2])

    right = self.getSkyline(buildings[n // 2:])
    return self._merge(left, right)

    def _merge(self, left: list[list[int]],
               right: list[list[int]]) ->
list[list[int]]:
    ans = []
    i = 0 # left's index
    j = 0 # right's index
    leftY = 0
    rightY = 0

```

```

        while i < len(left) and j < len(right):
            # Choose the powith smaller x
            if left[i][0] < right[j][0]:
                leftY = left[i][1] # Update the ongoing
`leftY`.
                self._addPoint(ans, left[i][0],
max(left[i][1], rightY))
                i += 1
            else:
                rightY = right[j][1] # Update the ongoing
`rightY`.
                self._addPoint(ans, right[j][0],
max(right[j][1], leftY))

        j += 1

        while i < len(left):
            self._addPoint(ans, left[i][0], left[i][1])
            i += 1

        while j < len(right):
            self._addPoint(ans, right[j][0],
right[j][1])
            j += 1

```

```

return ans

def _addPoint(self, ans: list[list[int]], x:
int, y: int) -> None:
    if ans and ans[-1][0] == x:
        ans[-1][1] = y
        return
    if ans and ans[-1][1] == y:
        return
    ans.append([x, y])

```

- LeetCode

```

from queue import PriorityQueue

class Solution:
    def getSkyline(self, buildings:
List[List[int]]) -> List[List[int]]:
        skys, lines, pq = [], [], PriorityQueue()
        for build in buildings:
            lines.extend([build[0], build[1]])
        lines.sort()
        city, n = 0, len(buildings)

```

```

        for line in lines:
            while city < n and buildings[city][0]
<= line:
                pq.put([-buildings[city][2],
buildings[city][0], buildings[city][1]])
                city += 1
            while not pq.empty() and
pq.queue[0][2] <= line:
                pq.get()
            high = 0
            if not pq.empty():
                high = -pq.queue[0][0]
            if len(skys) > 0 and skys[-1][1] ==
high:
                continue
            skys.append([line, high])
        return skys

```

- SimplyLeet

```
import heapq

def getSkyline(buildings):
    # List to store all events: (x, height, right)
    events = []
    # Create events for all buildings: left edge
    (start) and right edge (end)
    for left, right, height in buildings:
        # For the start point, height is negative
        for proper sorting (start events come before end
        events)
        events.append((left, -height, right))
        # For the end point, height is 0 to
        indicate removal of building from active heap
```

```
        events.append((right, 0, 0))

    # Sort events by x coordinate.
    events.sort()

    # Max heap to store active buildings (using
    negative heights)
    result = []
    # Initialize the heap with a dummy building of
    height 0 that extends infinitely
    heap = [(0, float('inf'))]
```

```
        for x, neg_height, right in events:
            # Remove buildings from the heap that have
            ended by current x.
            while heap and heap[0][1] <= x:
                heapq.heappop(heap)

            # If event is a start event (neg_height !=
            0), push it into the heap.
            if neg_height != 0:
                heapq.heappush(heap, (neg_height,
                right))

            # Current highest building height is at the
            top of the heap
```

```
        current_height = -heap[0][0]
        # If the result is empty or current height
        differs from last recorded height, add a new key
        point.
        if not result or current_height !=
        result[-1][1]:
            result.append([x, current_height])

        return result

# Example usage:
buildings =
[[2,9,10],[3,7,15],[5,12,12],[15,20,10],[19,24,8]]
print(getSkyline(buildings)) # Expected output: [[
2,10],[3,15],[7,12],[12,0],[15,10],[20,8],[24,0]]
```

• LC.ca

```
...
>>> a = []
>>> a.append((3, 5))
>>> a.append((3, -5))
>>> a.append((2, -10))
>>> a.append((2, 10))
>>> a
[(3, 5), (3, -5), (2, -10), (2, 10)]

>>> a.sort()
```

```
>>> a
[(2, -10), (2, 10), (3, -5), (3, 5)]
```

```
>>> from queue import PriorityQueue
>>> pq = PriorityQueue()
>>>
>>> pq.put([1,2,3])
>>> pq.put([-10,20,30])
```

```
>>> pq.put([11,22,33])
>>>
>>> pq
<Queue.PriorityQueue instance at 0x10aec51b8>
>>> pq.queue[0][0]
-10
>>> pq.get()
[-10, 20, 30]
>>> pq.queue[0][0]
1
```

```
...

from queue import PriorityQueue
class Solution:
    def getSkyline(self, buildings:
List[List[int]]) -> List[List[int]]:
        ans, lines, pq = [], [], PriorityQueue()
        for build in buildings:
            lines.extend([build[0], build[1]])
        lines.sort()
        curbuilding, n = 0, len(buildings)
```

```

    for line in lines: #
        a = -1
        while curbuilding < n and
            buildings[curbuilding][0] <= line:
                pq.put([-buildings[curbuilding][2]
, buildings[curbuilding][0],
buildings[curbuilding][1]]) # a = -
                curbuilding += 1
                while not pq.empty() and
                    pq.queue[0][2] <= line: # higher at heap top after
                        negated
                            pq.get() # i.e. pop(), remove
                                no-overlapping building
                                    high = 0 # a = -
                                    lines[i] <= line <= line[i] <= line[i]
                                    if not pq.empty():
                                        high = -pq.queue[0][0]
                                        if len(ans) > 0 and ans[-1][1] == high:
                                            # c = line[i] <= line[i] <= line[i]

```

```

        continue
        ans.append([line, high])
    return ans

```

#####

...

The solution uses a list called points to store the critical points and heights of the buildings. Each point is represented as a tuple (x, h), where x is the x-coordinate and h is the height. The points are sorted in ascending order based on the x-coordinate.

The solution also uses a heap to store the heights in descending order. The heap is initialized with a height of 0. For each point, if the height is negative, it means it is the start of a building, so the negative height is added to the heap. If the height is positive, it means it is the end of a building, so the corresponding negative height is removed from the heap.

After processing each point, the maximum height is obtained from the heap, and if it is different from the previous maximum height, the current point is added to the skyline.

Finally, the skyline is returned, excluding the initial point (0, 0) that was added as a starting point.

Note: The solution assumes that the input buildings is a list of tuples (left, right, height), where left and right represent the x-coordinates of the building's left and right edges, and height represents the height of the building.

```
'''
import heapq
```

```
class Solution:
    def getSkyline(self, buildings):
        # Create a list to store the critical
        points and heights
        points = []
        for left, right, height in buildings:
            points.append((left, -height)) # Start
            points.append((right, height)) # End
        of building, negative height
        of building, positive height

        # Sort the points in ascending order based
        on x-coordinate
        # If two points have the same x-coordinate,
        the one with larger height comes first

        points.sort()

        # Create a heap to store the heights in
        descending order
        heights = [0] # Initialize the heap with a
        height of 0
        skyline = [(0, 0)] # Initialize the skyline
        with a point (0, 0)

        for x, h in points:
            if h < 0:
                heapq.heappush(heights, h) # Add
            the negative height to the heap
            else:
```

```

        heights.remove(-h) # Remove the
corresponding negative height from the heap
        heapq.heapify(heights) #
Reorganize the heap

        # The current maximum height is the
first element in the heap
        max_height = -heights[0]

        # If the maximum height has changed,
add the current point to the skyline
        if max_height != skyline[-1][1]:
            skyline.append((x, max_height))

    return skyline[1:] # Exclude the initial
point (0, 0)

```

◆ Quick Review

Key Insights

- Use a line sweep approach to process building start and end events in sorted order.
- Differentiate between start events (which add to the skyline) and end events (which remove from the skyline).
- Employ a max heap (or a priority queue) to efficiently track the current highest building at each sweep line position.
- When the current maximum height changes, record the corresponding x-coordinate and new height as a key point.

- Take care to remove expired building heights from the heap as the sweep line moves past the building's end.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting events and heap operations for each building event.

Space Complexity: $O(n)$ for the event list and the heap used for processing active buildings.

Solution

The solution uses the line sweep algorithm with a max heap (or priority queue) to maintain the heights of active buildings. Each building contributes two events: one when the building starts (adding its height) and one when it ends (removing its height). These events are sorted primarily by their x-coordinate. For start events, we use negative height values to ensure they are processed before end events at the same x-coordinate. During the sweep, the heap is updated by adding new building heights and removing buildings that have ended. Whenever there is a change in the top of the heap (i.e., the current maximum height), a key point is added to the result. This key point signifies a vertical change in the skyline.

Heap

#272 Closest Binary Search Tree Value II

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · Stack · Tree · DFS · BST · Heap

Lists: Premium 98


Problem

Description

Given the root of a binary search tree, a target value, and an integer k , return the k values in the BST that are closest to the target. You may return the answer in any order.

You are guaranteed to have only one unique set of k values in the BST that are closest to the target.

Example 1:



Input: root = [4,2,5,1,3], target = 3.714286, k = 2
Output: [4,3]

Example 2:

Input: root = [1], target = 0.000000, k = 1
Output: [1]

Constraints:

- The number of nodes in the tree is n .
- $1 \leq k \leq n \leq 104$.
- $0 \leq \text{Node.val} \leq 109$
- $-109 \leq \text{target} \leq 109$

Round 6 · Mid · Hard

p.2487

Round 6 · Mid · Hard

p.2488

Follow up: Assume that the BST is balanced. Could you solve it in less than $O(n)$ runtime (where n = total nodes)?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def closestKValues(
        self,
        root: TreeNode | None,
        target: float,
        k: int,
    ) -> list[int]:
        dq = collections.deque()

        def inorder(root: TreeNode | None) -> None:

            if not root:
                return

            inorder(root.left)
            dq.append(root.val)
            inorder(root.right)

        inorder(root)

        while len(dq) > k:
```

```
            if abs(dq[0] - target) > abs(dq[-1] -
target):
                dq.popleft()
            else:
                dq.pop()

        return list(dq)
```

• LeetCode

```
# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def closestKValues(self, root: TreeNode,
target: float, k: int) -> List[int]:
        def dfs(root):
            if root is None:
```

```

        return
    dfs(root.left)
    if len(q) < k:
        q.append(root.val)
    else:
        if abs(root.val - target) >=
abs(q[0] - target):
            return
        q.popleft()
        q.append(root.val)
    dfs(root.right)

```

```

q = deque()
dfs(root)
return list(q)

```

• SimplyLeet

Python solution for Closest Binary Search Tree Value II

Definition for a binary tree node.

```

class TreeNode:
    def __init__(self, val=0, left=None,
right=None):
        self.val = val
        self.left = left
        self.right = right

```

```

class Solution:

```

```

def closestKValues(self, root: TreeNode,
target: float, k: int):
    # Helper function to initialize
predecessors stack
    def initPredecessor(node):
        while node:
            # If node value is less or equal,
push to preds and go right
            if node.val <= target:
                preds.append(node)
                node = node.right
            else:
                node = node.left

```

```

    # Helper function to initialize successors
stack
    def initSuccessor(node):
        while node:
            # If node value is greater than
target, push to succs and go left
            if node.val > target:
                succs.append(node)
                node = node.left
            else:
                node = node.right

```

```

# Helper function to get next predecessor
value
def getPredecessor():
    if not preds: return None
    node = preds.pop()
    val = node.val
    # Traverse the left subtree: push nodes
    from left child then rightmost branch
    node = node.left
    while node:
        preds.append(node)

```

```

    node = node.right
    return val

# Helper function to get next successor
value
def getSuccessor():
    if not succs: return None
    node = succs.pop()
    val = node.val
    # Traverse the right subtree: push
    nodes from right child then leftmost branch
    node = node.right

```

```

while node:
    succs.append(node)
    node = node.left
    return val

preds, succs = [], []
# Initialize the stacks with correct
boundary values
initPredecessor(root)
initSuccessor(root)

```

```

result = []
# If the closest value from predecessor and
successor are same, skip one to avoid duplicates.
if preds and succs and preds[-1].val ==
succs[-1].val:
    getPredecessor()

# Merge k closest values
for _ in range(k):
    if not preds:
        # Only successors left
        result.append(getSuccessor())

```

```

        elif not succs:
            # Only predecessors left
            result.append(getPredecessor())
        else:
            # Compare which candidate is
            closer to the target
            if abs(preds[-1].val - target) <=
abs(succs[-1].val - target):
                result.append(getPredecessor()
)
            else:
                result.append(getSuccessor())
        return result

```

• LC.ca

```

# Definition for a binary tree node.
# class TreeNode:
#     def __init__(self, val=0, left=None,
right=None):
#         self.val = val
#         self.left = left
#         self.right = right
class Solution:
    def closestKValues(self, root: TreeNode,
target: float, k: int) -> List[int]:
        def dfs(root):
            if root is None:

```

```

            return
            dfs(root.left)
            if len(q) < k:
                q.append(root.val)
            else:
                if abs(root.val - target) >=
abs(q[0] - target):
                    return
                q.popleft()
                q.append(root.val)
            dfs(root.right)

```

```

q = deque()
dfs(root)
return list(q)

```

```
#####
```

```

from collections import deque

# iteration

```

```

class Solution:
    def closestKValues(self, root: TreeNode,
target: float, k: int) -> List[int]:
        stack = []
        q = deque()

        # Iterative in-order traversal
        while stack or root:
            while root:
                stack.append(root)
                root = root.left

```

```

        root = stack.pop()

        # Process current node
        if len(q) < k:
            q.append(root.val)
        else:
            if abs(root.val - target) <
abs(q[0] - target):
                q.popleft()
                q.append(root.val)

```

```

        else:
            break
            # Early stop if the current
            value is farther than the first value in queue

            root = root.right

        return list(q)

```

◆ Quick Review

Key Insights

- The BST property allows efficient narrowing down of candidate nodes.
- Using two stacks to track predecessors (values $\leq$ target) and successors (values $>$ target) can efficiently yield the closest values.
- A modified in-order traversal helps in retrieving the next closest candidate from either side.
- Merging two sorted lists of candidate values based on absolute difference with the target is the key step.

Space & Time

Time Complexity: $O(k + \log n)$ – $O(\log n)$ to initialize the two stacks and $O(k)$ to collect k elements.

Space Complexity: $O(\log n)$ – due to the stack space in a balanced BST.

Solution

We can solve the problem by splitting it into two parts:

– Build two stacks: The predecessor stack for nodes with values less than or equal to target. The successor stack for nodes with values greater than target. To build these stacks, traverse the BST starting from the root and push nodes along the path: if the node's value is less than or equal to target, push it to the predecessor stack and go right; if greater, push it to the successor stack and go left. – Merge the two stacks: At each step, compare the top element of the predecessor stack (if available) and the successor stack (if available) based on their distance from the target. Pop from the stack with the closer candidate and add that value to the result list. For the popped node, update the given stack by moving to the next appropriate node in the in-order traversal (for predecessor, go to left child then rightmost descendant; for successor, go to right child then leftmost descendant). Repeat until k values have been collected.

Key Data Structures and Techniques:

– Two stacks to simulate reverse and forward in-order traversals. – Merge technique to choose the closest candidate at each iteration.

Heap

#295 Find Median from Data Stream HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Two Pointers · Design · Sorting · Heap

Lists: Grind 169

Problem

Description

The median is the middle value in an ordered integer list. If the size of the list is even, there is no middle value, and the median is the mean of the two middle values.

- For example, for arr = [2,3,4], the median is 3.
- For example, for arr = [2,3], the median is $(2 + 3) / 2 = 2.5$.

Implement the MedianFinder class:

- MedianFinder() initializes the MedianFinder object.
- void addNum(int num) adds the integer num from the data stream to the data structure.
- double findMedian() returns the median of all elements so far. Answers within 10⁻⁵ of the actual answer will be accepted.

Example 1:

Input

```
["MedianFinder", "addNum", "addNum", "findMedian",
"addNum", "findMedian"]
[[], [1], [2], [], [3], []]
```

Output

```
[null, null, null, 1.5, null, 2.0]
```

Explanation

```
MedianFinder medianFinder = new MedianFinder();
```

Constraints:

- $-105 \leq \text{num} \leq 105$
- There will be at least one element in the data structure before calling findMedian.
- At most $5 * 10^4$ calls will be made to addNum and findMedian.

Follow up:

- If all integer numbers from the stream are in the range $[0, 100]$, how would you optimize your solution?
- If 99% of all integer numbers from the stream are in the range $[0, 100]$, how would you optimize your solution?

Community Solutions (Python)

- WalkCC

```
class MedianFinder:
    def __init__(self):
        self.maxHeap = []
        self.minHeap = []

    def addNum(self, num: int) -> None:
        if not self.maxHeap or num <= -self.maxHeap[0]:
            heapq.heappush(self.maxHeap, -num)
        else:
            heapq.heappush(self.minHeap, num)

        # Balance the two heaps s.t.
        # |maxHeap| >= |minHeap| and |maxHeap| -
        # |minHeap| <= 1.
        if len(self.maxHeap) < len(self.minHeap):
            heapq.heappush(self.maxHeap,
                -heapq.heappop(self.minHeap))
        elif len(self.maxHeap) - len(self.minHeap) >
        1:
            heapq.heappush(self.minHeap,
                -heapq.heappop(self.maxHeap))

    def findMedian(self) -> float:
        if len(self.maxHeap) == len(self.minHeap):

            return (-self.maxHeap[0] + self.minHeap[0])
            / 2.0
        return -self.maxHeap[0]
```

- LeetCode

```

class MedianFinder:

    def __init__(self):
        self.minq = []
        self.maxq = []

    def addNum(self, num: int) -> None:
        heappush(self.minq,
        -heappushpop(self.maxq, -num))
        if len(self.minq) - len(self.maxq) > 1:
            heappush(self.maxq,
            -heappop(self.minq))

    def findMedian(self) -> float:
        if len(self.minq) == len(self.maxq):
            return (self.minq[0] - self.maxq[0]) /
2
        return self.minq[0]

# Your MedianFinder object will be instantiated and
called as such:
# obj = MedianFinder()
# obj.addNum(num)

# param_2 = obj.findMedian()

```

• SimplyLeet

```

import heapq

class MedianFinder:
    def __init__(self):
        # max-heap for the lower half (store
negatives to simulate max-heap)
        self.lowerHalf = []
        # min-heap for the upper half
        self.upperHalf = []

    def addNum(self, num: int) -> None:

        # Push to max-heap (simulate by pushing
negative)
        heapq.heappush(self.lowerHalf, -num)
        # Ensure every element in lowerHalf is <=
every element in upperHalf
        # Move maximum of lowerHalf to upperHalf
        heapq.heappush(self.upperHalf,
        -heapq.heappop(self.lowerHalf))

        # Balance the heaps: if upperHalf exceeds
lowerHalf in size
        if len(self.upperHalf) >
len(self.lowerHalf):
            heapq.heappush(self.lowerHalf,
            -heapq.heappop(self.upperHalf))

```



```

def findMedian(self) -> float:
    # If both heaps are balanced, median is
    # average of both heaps' top elements
    if len(self.lowerHalf) ==
len(self.upperHalf):
        return (-self.lowerHalf[0] +
self.upperHalf[0]) / 2.0
    # Otherwise, median is the top element of
    # the larger heap
    return -self.lowerHalf[0]

# Example usage:
# medianFinder = MedianFinder()
# medianFinder.addNum(1)

# medianFinder.addNum(2)
# print(medianFinder.findMedian()) # Outputs 1.5
# medianFinder.addNum(3)
# print(medianFinder.findMedian()) # Outputs 2.0

```

- LC.ca

```

from heapq import heappush, heappop

class MedianFinder:

def __init__(self):
    # the smaller half of the list, max heap
    # (invert min-heap)
    self.smallerHalf = []
    # the larger half of the list, min heap
    self.largerHalf = []

def addNum(self, num: int) -> None:

    # trick for smaller half, use -1*val
    # heapq in python does NOT have comparator like in
    # Java
    heappush(self.smallerHalf, -num)
    heappush(self.largerHalf,
-heappop(self.smallerHalf)) # note: not
self.smallerHalf.pop()

    if len(self.smallerHalf) < len(self.largerHalf):
        heappush(self.smallerHalf,
-heappop(self.largerHalf))

```

```

def findMedian(self) -> float:
    if len(self.smallerHalf) == len(self.largerHalf):
        return (-self.smallerHalf[0] + self.largerHalf[0])
        / 2.0
    else:
        return float(-self.smallerHalf[0])

# Your MedianFinder object will be instantiated and
# called as such:
# obj = MedianFinder()
# obj.addNum(num)

# param_2 = obj.findMedian()

# follow up
class MedianFinder:
    def __init__(self):
        self.counts = [0] * 101
        self.total = 0

    def addNum(self, num: int) -> None:

```

```

        self.counts[num] += 1
        self.total += 1

    def findMedian(self) -> float:
        if self.total % 2 == 0:
            # even number of elements
            middle1 = self.total // 2
            middle2 = middle1 + 1
            count = 0
            i1 = i2 = 0

            for i in range(101):
                count += self.counts[i]
                if count >= middle1 and i1 == 0:
                    i1 = i
                if count >= middle2:
                    i2 = i
                    break
            return (i1 + i2) / 2
        else:
            # odd number of elements

```

```

middle = self.total // 2 + 1
count = 0
for i in range(101):
    count += self.counts[i]
    if count >= middle:
        return i

```

```
#####
```

```

class MedianFinder:
    def __init__(self):
        """
        initialize your data structure here.
        """
        self.h1 = []
        self.h2 = []

    def addNum(self, num: int) -> None:
        heappush(self.h1, num)

```

```

heappush(self.h2, -heappop(self.h1))
if len(self.h2) - len(self.h1) > 1:
    heappush(self.h1, -heappop(self.h2))

```

```

def findMedian(self) -> float:
    if len(self.h2) > len(self.h1):
        return -self.h2[0]
    return (self.h1[0] - self.h2[0]) / 2

```

Your MedianFinder object will be instantiated and called as such:

```

# obj = MedianFinder()
# obj.addNum(num)
# param_2 = obj.findMedian()

```

◆ Quick Review

Key Insights

- Use two heaps: a max-heap for the lower half and a min-heap for the upper half of numbers.
- Balancing the heaps ensures the median can be computed in $O(1)$ time once numbers are added.
- Rebalance the heaps after each insertion to ensure size properties hold (difference in sizes is at most 1).

Space & Time

Time Complexity:

- addNum: $O(\log n)$ per insertion (heap operations).
- findMedian: $O(1)$ (accessing the top elements).

Space Complexity:

- $O(n)$, where n is the number of elements added in the data stream, as they are stored in the two heaps.

Solution

We maintain two heaps:

- A max-heap (lowerHalf) for the smaller half of the numbers.
 - A min-heap (upperHalf) for the larger half of the numbers.
- On inserting a new number, decide the heap membership based on the current median. After each insertion, ensure that the heaps have balanced sizes (the size difference is at most one). When computing the median:
- If both heaps have the same number of elements, the median is the average of the max of lowerHalf and the min of upperHalf.
 - If one heap has more elements, the median is the top element of that heap.

This method makes it efficient to dynamically update and retrieve the median as new numbers are inserted.

Heap**#358 Rearrange String k Distance Apart**

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Greedy · Sorting · Heap · Counting Lists: Premium 98

Problem**Description**

Given a string s and an integer k , rearrange s such that the same characters are at least distance k from each other. If it is not possible to rearrange the string, return an empty string "".

Example 1:

Input: $s = \text{"aabbcc"}, k = 3$

Output: "abcabc"

Explanation: The same letters are at least a distance of 3 from each other.

Example 2:

Input: $s = \text{"aaabc"}, k = 3$

Output: ""

Explanation: It is not possible to rearrange the string.

Example 3:

Input: s = "aaadbbcc", k = 2

Output: "abacabcd"

Explanation: The same letters are at least a distance of 2 from each other.

Constraints:

- $1 \leq s.length \leq 3 * 10^5$
- s consists of only lowercase English letters.
- $0 \leq k \leq s.length$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def rearrangeString(self, s: str, k: int) -> str:
        n = len(s)
        ans = []
        count = collections.Counter(s)
        # valid[i] := the leftmost index i can appear
        valid = collections.Counter()

        def getBestLetter(index: int) -> str:
            """Returns the valid letter that has the most count."""
```

```
maxCount = -1
bestLetter = '*'
```

```
for c in string.ascii_lowercase:
    if count[c] > 0 and count[c] > maxCount and
index >= valid[c]:
        bestLetter = c
        maxCount = count[c]

return bestLetter
```

```
for i in range(n):
    c = getBestLetter(i)
    if c == '*':
        return ''
    ans.append(c)
    count[c] -= 1
    valid[c] = i + k

return ''.join(ans)
```

• LeetCode

```

class Solution:
    def rearrangeString(self, s: str, k: int) -> str:
        cnt = Counter(s)
        pq = [(-v, c) for c, v in cnt.items()]
        heapify(pq)
        q = deque()
        ans = []
        while pq:
            v, c = heappop(pq)
            ans.append(c)

            q.append((v + 1, c))
            if len(q) >= k:
                e = q.popleft()
                if e[0]:
                    heappush(pq, e)
        return "" if len(ans) < len(s) else "".join(ans)

```

• SimplyLeet

```

from collections import Counter, deque
import heapq

def rearrangeString(s: str, k: int) -> str:
    if k == 0:
        return s # No rearrangement needed if k is 0

    # Count the frequency of each character
    char_count = Counter(s)

    # Build a max-heap using negative frequencies
    max_heap = [(-freq, ch) for ch, freq in char_count.items()]
    heapq.heapify(max_heap)

    # Queue to keep track of characters waiting for k distance
    wait_queue = deque()
    result = []

    # Process the heap until it's empty
    while max_heap:

```

```

    freq, ch = heapq.heappop(max_heap)
    result.append(ch)

    # Decrease frequency (remember freq is
    negative)
    freq += 1

    # Add current char to the wait queue
    wait_queue.append((freq, ch))

    # If we've placed k characters, we can
    re-insert the oldest one if needed

    if len(wait_queue) < k:
        continue

    freq_front, ch_front =
    wait_queue.popleft()
    if freq_front < 0:
        heapq.heappush(max_heap, (freq_front,
        ch_front))

    rearranged = ''.join(result)
    return rearranged if len(rearranged) == len(s)
    else ""

```

• LC.ca

Round 6 · Mid · Hard

p.2517

```

from heapq import heapify

class Solution:
    def rearrangeString(self, s: str, k: int) ->
    str:
        h = [(-v, c) for c, v in
        Counter(s).items()]
        heapify(h)
        q = deque()
        ans = []
        while h:
            v, c = heappop(h)

            v *= -1
            ans.append(c)
            q.append((v - 1, c)) # enqueue even if
            'v-1==0'

            # not '>='
            # use example k=1, then q size is 2,
            then q can pop one out
            if len(q) > k: # @note: this is avoid
            you pick up same char in the same k-segment.
                w, c = q.popleft()
                if w: # w!=0
                    heappush(h, (-w, c))
            return "" if len(ans) != len(s) else
            "".join(ans)

```

Round 6 · Mid · Hard

p.2518

```
#####
```

```
import collections
import heapq
```

```
class Solution:
    def rearrangeString(self, s: str, k: int) ->
    str:
        if k == 0:
            return s
```

```
counts = collections.Counter(s)
```

```
pq = []
for ch, count in counts.items():
    heapq.heappush(pq, (-count, ch))
```

```
sb = []
while pq:
    tmp = []
```

```
        d = min(k, len(s))
        for i in range(d):
            if not pq:
                return ""
            count, ch = heapq.heappop(pq)
            sb.append(ch)
            if count < -1: # pushed in as
                negated, so it's: if actual-positive-count > 1
                tmp.append((count+1, ch))
            s = s[1:]

        for count, ch in tmp:
            heapq.heappush(pq, (count, ch))

    return "".join(sb)
```

◆ Quick Review

Key Insights

- The frequency of each character dictates the feasibility and placement within the rearrangement.
- A greedy approach picks the most frequent available character to maximize the chance of successful placement.
- A max-heap (priority queue) efficiently selects the highest frequency character each time.
- A waiting queue ensures that once a character is used, it can't be reused until k positions later.
- Special case: When k is 0, no rearrangement is required because all orders satisfy the condition.

Space & Time

Time Complexity: $O(n \log C)$ where n is the length of s and C is the number of unique characters.

Space Complexity: $O(n)$ due to the storage used for the heap and wait queue.

Solution

We use a greedy algorithm with a max-heap to always select the character with the highest remaining frequency that is currently available. The algorithm first counts the frequency of each character and pushes these into a max-heap (using negative frequencies for languages like Python where `heapq` is a min-heap). Each time a character is selected, it is appended to the result string and its frequency is decreased. This character is then placed in a waiting queue to ensure it isn't reused until k positions later. Once the waiting period completes, if the character still has remaining occurrences, it is reinserted into the heap. If at any point the heap is empty and the waiting queue still contains characters that haven't been reinserted while the result string is incomplete, we conclude that a valid rearrangement is impossible.

Heap

#499 The Maze III

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · DFS · BFS · Graph · Matrix · Heap · Shortest Path

Lists: Premium 98

Problem

Description


There is a ball in a maze with empty spaces (represented as 0) and walls (represented as 1). The ball can go through the empty spaces by rolling up, down, left or right, but it won't stop rolling until hitting a wall. When the ball stops, it could choose the next direction (must be different from last chosen direction). There is also a hole in this maze. The ball will drop into the hole if it rolls onto the hole.

Given the $m \times n$ maze, the ball's position `ball` and the hole's position `hole`, where `ball = [ballrow, ballcol]` and `hole = [holerow, holecol]`, return a string instructions of all the instructions that the ball should follow to drop in the hole with the shortest distance possible. If there are multiple valid instructions, return the lexicographically minimum one. If the ball can't drop in the hole, return "impossible".

If there is a way for the ball to drop in the hole, the answer instructions should contain the characters 'u' (i.e., up), 'd' (i.e., down), 'l' (i.e., left), and 'r' (i.e., right). The distance is the number of empty spaces traveled by the ball from the start position (excluded) to the destination (included).

You may assume that the borders of the maze are all walls (see examples).

Example 1:



Input: maze = [[0,0,0,0,0],[1,1,0,0,1],[0,0,0,0,0],[0,1,0,0,1],[0,1,0,0,0]], ball = [4,3], hole = [0,1]

Output: "lul"


Explanation: There are two shortest ways for the ball to drop into the hole.

The first way is left -> up -> left, represented by "lul".

The second way is up -> left, represented by 'ul'.

Both ways have shortest distance 6, but the first way is lexicographically smaller because 'l' < 'u'. So the output is "lul".

Example 2:



Input: maze = [[0,0,0,0,0],[1,1,0,0,1],[0,0,0,0,0],[0,1,0,0,1],[0,1,0,0,0]], ball = [4,3], hole = [3,0]
 Output: "impossible"
 Explanation: The ball cannot reach the hole.

Example 3:

Input: maze = [[0,0,0,0,0,0,0],[0,0,1,0,0,1,0],[0,0,0,0,1,0,0],[1,0,0],[0,0,0,0,0,0,1]], ball = [0,4], hole = [3,5]
 Output: "dldr"

Constraints:

- $m == \text{maze.length}$

- $n == \text{maze}[i].\text{length}$
- $1 \leq m, n \leq 100$
- $\text{maze}[i][j]$ is 0 or 1.
- $\text{ball.length} == 2$
- $\text{hole.length} == 2$
- $0 \leq \text{ballrow}, \text{holerow} \leq m$
- $0 \leq \text{ballcol}, \text{holecol} \leq n$
- Both the ball and the hole exist in an empty space, and they will not be in the same position initially.
- The maze contains at least 2 empty spaces.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findShortestWay(
        self,
        maze: list[list[int]],
        ball: list[int],
        hole: list[int],
    ) -> str:
        ans = 'impossible'
        minSteps = math.inf
```

```

def dfs(i: int, j: int, dx: int, dy: int,
steps: int, path: str):
    nonlocal ans
    nonlocal minSteps
    if steps >= minSteps:
        return

    if dx != 0 or dy != 0: # Both are zeros for
the initial ball position.
        while (0 <= i + dx < len(maze) and 0 <= j +
dy < len(maze[0]) and
                maze[i + dx][j + dy] != 1):
            i += dx

```

```

            j += dy
            steps += 1
            if i == hole[0] and j == hole[1] and
steps < minSteps:
                minSteps = steps
                ans = path

            if maze[i][j] == 0 or steps + 2 < maze[i][j]:
                maze[i][j] = steps + 2 # +2 because
maze[i][j] == 0 || 1.
                if dx == 0:
                    dfs(i, j, 1, 0, steps, path + 'd')

```

```

        if dy == 0:
            dfs(i, j, 0, -1, steps, path + 'l')
        if dy == 0:
            dfs(i, j, 0, 1, steps, path + 'r')
        if dx == 0:
            dfs(i, j, -1, 0, steps, path + 'u')

    dfs(ball[0], ball[1], 0, 0, 0, '')
    return ans

```

• LeetCode

```

class Solution:
    def findShortestWay(
        self, maze: List[List[int]], ball:
List[int], hole: List[int]
    ) -> str:
        m, n = len(maze), len(maze[0])
        r, c = ball
        rh, ch = hole
        q = deque([(r, c)])
        dist = [[inf] * n for _ in range(m)]
        dist[r][c] = 0

```

```

path = [[None] * n for _ in range(m)]
path[r][c] = ''
while q:
    i, j = q.popleft()
    for a, b, d in [(-1, 0, 'u'), (1, 0, 'd'), (0, -1, 'l'), (0, 1, 'r')]:
        x, y, step = i, j, dist[i][j]
        while (
            0 <= x + a < m
            and 0 <= y + b < n
            and maze[x + a][y + b] == 0
            and (x != rh or y != ch)
        ):
            x, y = x + a, y + b
            step += 1
            if dist[x][y] > step or (
                dist[x][y] == step and
                path[i][j] + d < path[x][y]
            ):
                dist[x][y] = step
                path[x][y] = path[i][j] + d
                if x != rh or y != ch:
                    q.append((x, y))
return path[rh][ch] or 'impossible'

```

• SimplyLeet

```

import heapq

def findShortestWay(maze, ball, hole):
    m, n = len(maze), len(maze[0])
    # Directions: sorted lexicographically as 'd' < 'l' < 'r' < 'u'
    # since ASCII: d(100), l(108), r(114), u(117)
    directions = [('d', 1, 0), ('l', 0, -1), ('r', 0, 1), ('u', -1, 0)]

    # Initialize distance and path record, fill with (infy, "")
    dist = [[float('inf')] * n for _ in range(m)]

    pathRecord = [[""] * n for _ in range(m)]

    # Priority queue stores tuples: (distance, path, row, col)
    heap = [(0, "", ball[0], ball[1])]
    dist[ball[0]][ball[1]] = 0

    while heap:
        curDist, curPath, x, y = heapq.heappop(heap)
        # If we pop a state for the hole then return result.
        if [x, y] == hole:

```

```

        return curPath
    # If we have a larger distance than already
    recorded, continue.
    if curDist > dist[x][y]:
        continue
    # Explore all four directions
    for d, dx, dy in directions:
        nx, ny = x, y
        steps = 0
        # Roll the ball until hitting a wall or
        passing through the hole.
        # Check if ball goes through the hole
        during rolling.

```

```

        while 0 <= nx + dx < m and 0 <= ny + dy
        < n and maze[nx + dx][ny + dy] == 0:
            nx += dx
            ny += dy
            steps += 1
            if [nx, ny] == hole:
                break
            newDist = curDist + steps
            newPath = curPath + d
            # If we found a shorter route or equal
            distance but lexicographically smaller route.
            if newDist < dist[nx][ny] or (newDist
            == dist[nx][ny] and newPath < pathRecord[nx][ny]):

```

```

        dist[nx][ny] = newDist
        pathRecord[nx][ny] = newPath
        heapq.heappush(heap, (newDist,
        newPath, nx, ny))
        return "impossible"

```

Example usage:

```

maze = [[0,0,0,0,0],
        [1,1,0,0,1],
        [0,0,0,0,0],
        [0,1,0,0,1],

```

```

        [0,1,0,0,0]]
ball = [4, 3]
hole = [0, 1]
print(findShortestWay(maze, ball, hole)) #
Expected output: "lul"

```

• LC.ca

```

class Solution:
    def findShortestWay(
        self, maze: List[List[int]], ball:
        List[int], hole: List[int]
    ) -> str:
        m, n = len(maze), len(maze[0])
        r, c = ball
        rh, ch = hole
        q = deque([(r, c)])
        dist = [[inf] * n for _ in range(m)]
        dist[r][c] = 0

```

```

path = [[None] * n for _ in range(m)]
path[r][c] = ''
while q:
    i, j = q.popleft()
    for a, b, d in [(-1, 0, 'u'), (1, 0, 'd'), (0, -1, 'l'), (0, 1, 'r')]:
        x, y, step = i, j, dist[i][j]
        while (
            0 <= x + a < m
            and 0 <= y + b < n
            and maze[x + a][y + b] == 0
            and (x != rh or y != ch)
        ):
            x, y = x + a, y + b
            step += 1
            if dist[x][y] > step or (
                dist[x][y] == step and
                path[i][j] + d < path[x][y]
            ):
                dist[x][y] = step
                path[x][y] = path[i][j] + d
                if x != rh or y != ch:
                    q.append((x, y))
return path[rh][ch] or 'impossible'

```

◆ Quick Review

Key Insights

- The ball rolls continuously until hitting a wall or falling into the hole.
- A shortest path search is needed based on rolling distance (each step counts when passing an empty cell).
- Dijkstra's algorithm (or a modified BFS with a priority queue) is appropriate since we need to consider distance and lexicographic order.
- When expanding moves, simulate the roll until the ball stops — if a hole is passed en route, use that stopping point.
- Track the best (distance, path) for each cell to avoid redundant work.

Space & Time

Time Complexity: $O(mn \log(mn))$ where m and n are the maze dimensions

Space Complexity: $O(mn)$

Solution

We use a Dijkstra-like approach with a priority queue containing tuples of (distance, path, row, col). The priority queue is ordered by distance first and then by the instruction string lexicographically. For each cell, simulate rolling in each possible direction (up, down, left, right; iterated in alphabetical order so that the lexicographical property is maintained) until the ball stops. If the ball encounters the hole during rolling, treat that as a valid stop immediately. For every new stopping cell, if a shorter distance or lexicographically smaller path is found, update the record and add the new state to the priority queue. If

you reach the hole, return the corresponding path string. If no solution is found when the queue is empty, return "impossible".

Round 6 · Mid · Hard

p.2535

Heap

#632 Smallest Range Covering Elements from K Lists

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Greedy · Heap · Sliding Window · Sorting

Lists: Grind 169

Problem

Description

You have k lists of sorted integers in non-decreasing order. Find the smallest range that includes at least one number from each of the k lists.

We define the range $[a, b]$ is smaller than range $[c, d]$ if $b - a < d - c$ or $a < c$ if $b - a == d - c$.

Example 1:

```
Input: nums =
[[4,10,15,24,26],[0,9,12,20],[5,18,22,30]]
Output: [20,24]
Explanation:
List 1: [4, 10, 15, 24,26], 24 is in range [20,24].
List 2: [0, 9, 12, 20], 20 is in range [20,24].
List 3: [5, 18, 22, 30], 22 is in range [20,24].
```

Example 2:

Round 6 · Mid · Hard

p.2536

Input: nums = [[1,2,3],[1,2,3],[1,2,3]]

Output: [1,1]

Constraints:

- nums.length == k
- 1 <= k <= 3500
- 1 <= nums[i].length <= 50
- -105 <= nums[i][j] <= 105
- nums[i] is sorted in non-decreasing order.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def smallestRange(self, nums: list[list[int]])
-> list[int]:
    minHeap = [(row[0], i, 0) for i, row in
enumerate(nums)]
    heapq.heapify(minHeap)

    maxRange = max(row[0] for row in nums)
    minRange = heapq.nsmallest(1, minHeap)[0][0]
    ans = [minRange, maxRange]

    while len(minHeap) == len(nums):
```

```
num, r, c = heapq.heappop(minHeap)
if c + 1 < len(nums[r]):
    heapq.heappush(minHeap, (nums[r][c + 1],
r, c + 1))
    maxRange = max(maxRange, nums[r][c + 1])
    minRange = heapq.nsmallest(1,
minHeap)[0][0]
    if maxRange - minRange < ans[1] - ans[0]:
        ans[0], ans[1] = minRange, maxRange

return ans
```

• LeetDoocs

```
class Solution:
    def smallestRange(self, nums: List[List[int]])
-> List[int]:
    t = [(x, i) for i, v in enumerate(nums) for
x in v]
    t.sort()
    cnt = Counter()
    ans = [-inf, inf]
    j = 0
    for i, (b, v) in enumerate(t):
        cnt[v] += 1
        while len(cnt) == len(nums):
```

```

        a = t[j][0]
        x = b - a - (ans[1] - ans[0])
        if x < 0 or (x == 0 and a <
ans[0]):
            ans = [a, b]
            w = t[j][1]
            cnt[w] -= 1
            if cnt[w] == 0:
                cnt.pop(w)
            j += 1
    return ans

```

• SimplyLeet

```

import heapq

def smallestRange(nums):
    # Initialize the min-heap and current maximum
    value
    heap = []
    current_max = float('-inf')

    # Add the first element of each list to the
    heap
    for i, lst in enumerate(nums):
        value = lst[0]

```

```

        heapq.heappush(heap, (value, i, 0))
        current_max = max(current_max, value)

    best_range = [float('-inf'), float('inf')]

    # Process until one list is exhausted
    while heap:
        current_min, list_idx, elem_idx =
        heapq.heappop(heap)

        # Update the best range if the current
        range is smaller

        if current_max - current_min <
        best_range[1] - best_range[0] or \
            (current_max - current_min ==
        best_range[1] - best_range[0] and current_min <
        best_range[0]):
            best_range = [current_min,
            current_max]

        # If there is no next element in the same
        list, break out
        if elem_idx + 1 == len(nums[list_idx]):
            break
        next_value = nums[list_idx][elem_idx + 1]
        heapq.heappush(heap, (next_value,
        list_idx, elem_idx + 1))
        current_max = max(current_max, next_value)

```

```
return best_range
```

Example Usage:

```
nums = [[4,10,15,24,26],[0,9,12,20],[5,18,22,30]]
print(smallestRange(nums))
```

• LC.ca

```
class Solution:
    def smallestRange(self, nums: List[List[int]])
-> List[int]:
    t = [(x, i) for i, v in enumerate(nums) for
x in v]
    t.sort()
    cnt = Counter()
    ans = [-inf, inf]
    j = 0
    for b, v in t:
        cnt[v] += 1
        while len(cnt) == len(nums):
```

```
        a = t[j][0]
        x = b - a - (ans[1] - ans[0])
        if x < 0 or (x == 0 and a <
ans[0]):
            ans = [a, b]
            w = t[j][1]
            cnt[w] -= 1
            if cnt[w] == 0:
                cnt.pop(w)
            j += 1
    return ans
```

◆ Quick Review

Key Insights

- Use a min heap to maintain the smallest current element among the selected elements from each list.
- Track the current maximum among the chosen elements to easily compute the range.
- Iteratively replace the minimum element with the next element from the same list and update the current maximum.
- Stop when one of the lists is exhausted, as the range must include at least one element from each list.

Space & Time

Time Complexity: $O(N \log k)$, where N is the total number of elements across all lists.

Space Complexity: $O(k)$ for the heap, plus additional space for storing the range.

Solution

The solution uses a min heap (priority queue) that stores a tuple containing the current element, the index of the list it comes from, and its index in that list. Initially, the first element from each list is inserted into the heap and the current maximum among them is tracked. In each iteration, the smallest element is removed (heap root) which provides the current minimum, and the range $[current_min, current_max]$ is assessed. If the list from which the minimum element was extracted has a next element, that element is inserted into the heap and the current maximum is updated if necessary. This process continues until one of the lists is exhausted.

Heap

#759 Employee Free Time

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Sorting · Heap

Lists: Grind 169

Problem

Description

We are given a list schedule of employees, which represents the working time for each employee. Each employee has a list of non-overlapping Intervals, and these intervals are in sorted order.

Return the list of finite intervals representing common, positive-length free time for all employees, also in sorted order.

(Even though we are representing Intervals in the form $[x, y]$, the objects inside are Intervals, not lists or arrays. For example, `schedule[0][0].start = 1`, `schedule[0][0].end = 2`, and `schedule[0][0][0]` is not defined). Also, we wouldn't include intervals like $[5, 5]$ in our answer, as they have zero length.

Example 1:

Input: schedule = [[[1,2],[5,6]],[[1,3]],[[4,10]]]
 Output: [[3,4]]
 Explanation: There are a total of three employees, and all common free time intervals would be [-inf, 1], [3, 4], [10, inf].
 We discard any intervals that contain inf as they aren't finite.

Example 2:

Input: schedule =
 [[[1,3],[6,7]],[[2,4]],[[2,5],[9,12]]]
 Output: [[5,6],[7,9]]

Constraints:

- $1 \leq \text{schedule.length}, \text{schedule}[i].\text{length} \leq 50$
- $0 \leq \text{schedule}[i].\text{start} < \text{schedule}[i].\text{end} \leq 10^8$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def employeeFreeTime(self, schedule:
        '[[Interval]]') -> '[Interval]':
        ans = []
        intervals = []

        for s in schedule:
            intervals.extend(s)

        intervals.sort(key=lambda x: x.start)

        prevEnd = intervals[0].end

        for interval in intervals:
            if interval.start > prevEnd:
                ans.append(Interval(prevEnd,
                    interval.start))
                prevEnd = max(prevEnd, interval.end)

        return ans
```

- LeetCode

```

"""
# Definition for an Interval.
class Interval:
    def __init__(self, start: int = None, end: int = None):
        self.start = start
        self.end = end
"""

```

```

class Solution:

```

```

    def employeeFreeTime(self, schedule:
"[[Interval]]") -> "[Interval]":
        intervals = []
        for e in schedule:
            intervals.extend(e)
        intervals.sort(key=lambda x: (x.start,
x.end))
        merged = [intervals[0]]
        for x in intervals[1:]:
            if merged[-1].end < x.start:
                merged.append(x)
            else:

```

```

                merged[-1].end =
max(merged[-1].end, x.end)
        ans = []
        for a, b in pairwise(merged):
            ans.append(Interval(a.end, b.start))
        return ans

```

• SimplyLeet

```

# Class definition for Interval
class Interval:
    def __init__(self, start, end):
        self.start = start # Start time of the
interval
        self.end = end # End time of the interval

def employeeFreeTime(schedule):
    # Flatten the schedule: extract all intervals
from each employee
    allIntervals = []
    for employee in schedule:

```

```

    for interval in employee:
        allIntervals.append(interval)

    # Sort intervals by start time
    allIntervals.sort(key=lambda interval:
interval.start)

    # Merge overlapping intervals
    merged = []
    prev = allIntervals[0]
    for interval in allIntervals[1:]:

        if interval.start <= prev.end:
            # Expand the previous interval if there
            is an overlap
            prev.end = max(prev.end, interval.end)
        else:
            merged.append(prev)
            prev = interval
    merged.append(prev)

    # Collect free times from the gaps between
merged intervals
    freeTimes = []

```

```

    for i in range(1, len(merged)):
        # Only intervals with positive length are
        added
        if merged[i].start > merged[i-1].end:
            freeTimes.append(Interval(merged[i-1].
end, merged[i].start))

    return freeTimes

# Example usage:
# schedule = [
# [Interval(1,2), Interval(5,6)],

# [Interval(1,3)],
# [Interval(4,10)]
# ]
# freeTimes = employeeFreeTime(schedule)
# for interval in freeTimes:
# print(f"[{interval.start}, {interval.end}]")

```

- LC.ca

```

"""
# Definition for an Interval.
class Interval:
    def __init__(self, start: int = None, end: int = None):
        self.start = start
        self.end = end
"""

```

```

class Solution:

```

```

    def employeeFreeTime(self, schedule: "[[Interval]]") -> "[Interval]":
        intervals = []
        for e in schedule:
            intervals.extend(e)
        intervals.sort(key=lambda x: (x.start, x.end))
        merged = [intervals[0]]
        for x in intervals[1:]:
            if merged[-1].end < x.start:
                merged.append(x)
            else:

```

```

                merged[-1].end =
max(merged[-1].end, x.end)
        ans = []
        for a, b in pairwise(merged):
            ans.append(Interval(a.end, b.start))
        return ans

```

◆ Quick Review

Key Insights

- Flatten the schedule to combine all employees' intervals.
- Sort the flattened intervals by start time.
- Merge overlapping intervals to form a consolidated view of busy times.
- Identify gaps between merged intervals as the common free time.
- Ensure free intervals have strictly positive length.

Space & Time

Time Complexity: $O(N \log N)$, where N is the total number of intervals (due to sorting).

Space Complexity: $O(N)$, used for the flattened list and merged intervals.

Solution

The solution leverages interval merging. First, all intervals from every employee are collected and sorted by start time. By iterating through the sorted intervals, overlapping intervals are merged into one consolidated busy period.

Finally, the gaps between consecutive merged intervals are identified as free periods available to all employees. This approach efficiently handles the intervals with a simple sort and a linear scan.

Round 6 · Mid · Hard

p.2553

Heap

#1425 Constrained Subsequence Sum

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Queue · Sliding Window · Heap · Monotonic Queue

Lists: Denny Zhang

Problem

Description

Given an integer array `nums` and an integer `k`, return the maximum sum of a non-empty subsequence of that array such that for every two consecutive integers in the subsequence, `nums[i]` and `nums[j]`, where $i < j$, the condition $j - i \leq k$ is satisfied.

A subsequence of an array is obtained by deleting some number of elements (can be zero) from the array, leaving the remaining elements in their original order.

Example 1:

Input: `nums = [10,2,-10,5,20]`, `k = 2`

Output: 37

Explanation: The subsequence is `[10, 2, 5, 20]`.

Example 2:

Round 6 · Mid · Hard

p.2554

Input: nums = [-1,-2,-3], k = 1

Output: -1

Explanation: The subsequence must be non-empty, so we choose the largest number.

Example 3:

Input: nums = [10,-2,-10,-5,20], k = 2

Output: 23

Explanation: The subsequence is [10, -2, -5, 20].

Constraints:

- $1 \leq k \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def constrainedSubsetSum(self, nums: list[int], k: int) -> int:
        # dp[i] := the maximum the sum of non-empty
        # subsequences in nums[0..i]
        dp = [0] * len(nums)
        # dq stores dp[i - k], dp[i - k + 1], ..., dp[i - 1]
        # whose values are > 0
        # in decreasing order.
        dq = collections.deque()

        for i, num in enumerate(nums):
            if dq:
```

```
        dp[i] = max(dq[0], 0) + num
    else:
        dp[i] = num
        while dq and dq[-1] < dp[i]:
            dq.pop()
        dq.append(dp[i])
        if i >= k and dp[i - k] == dq[0]:
            dq.popleft()

    return max(dp)
```

• LeetCode

```
class Solution:
    def constrainedSubsetSum(self, nums: List[int], k: int) -> int:
        q = deque([0])
        n = len(nums)
        f = [0] * n
        ans = -inf
        for i, x in enumerate(nums):
            while i - q[0] > k:
                q.popleft()
            f[i] = max(0, f[q[0]]) + x

        ans = max(ans, f[i])
        while q and f[q[-1]] <= f[i]:
            q.pop()
        q.append(i)
        return ans
```

- SimplyLeet

```
# Python solution using a monotonic deque for
efficient window maximum computation.
from collections import deque

def constrained_subsequence_sum(nums, k):
    n = len(nums)
    dp = [0] * n # dp[i] stores the maximum
subsequence sum ending at index i
    dp[0] = nums[0]
    result = dp[0]

    # Deque to store indices in a decreasing order
of dp values

    deq = deque([0])

    for i in range(1, n):
        # Remove indices from the front which are
out of the window
        while deq and i - deq[0] > k:
            deq.popleft()

        # The maximum dp in the window is at the
front of the deque
        max_in_window = dp[deq[0]] if dp[deq[0]] >
0 else 0
```

```
        # Calculate dp[i] by either starting a new
subsequence or extending the best one from the
window
        dp[i] = nums[i] + max_in_window

        # Update overall answer
        result = max(result, dp[i])

        # Maintain the deque monotonically so that
dp indices are in decreasing order
        while deq and dp[i] >= dp[deq[-1]]:
            deq.pop()
        deq.append(i)

    return result

# Example usage:
#
print(constrained_subsequence_sum([10,2,-10,5,20],
2))
```

- LC.ca

```

class Solution:
    def constrainedSubsetSum(self, nums:
List[int], k: int) -> int:
        n = len(nums)
        dp = [0] * n
        ans = -inf
        q = deque()
        for i, v in enumerate(nums):
            if q and i - q[0] > k:
                q.popleft()
            dp[i] = max(0, 0 if not q else
dp[q[0]]) + v

            while q and dp[q[-1]] <= dp[i]:
                q.pop()
            q.append(i)
            ans = max(ans, dp[i])
        return ans

```

◆ Quick Review

Key Insights

- Use dynamic programming where $dp[i]$ represents the maximum sum of a valid subsequence ending at index i .
- For each index i , consider taking $nums[i]$ alone or appending it to a valid subsequence ending at an index j (where $i - j \leq k$) that maximizes the sum.
- To efficiently obtain the maximum $dp[j]$ in a sliding window of the last k indices, use a monotonic (decreasing) deque.

- The idea is to maintain the window maximum and update it as we proceed through the array.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$ ($O(k)$ extra space for the deque, with dp array using $O(n)$)

Solution

We solve the problem by defining a dp array where $dp[i] = nums[i] + \max(0, \text{maximum } dp \text{ value in the window } [i-k, i-1])$. The trick is to include $\max(0, \dots)$ to allow starting a new subsequence when previous sums are negative. To access the maximum dp in the window in constant time, a monotonic deque is maintained that stores indices with their dp values in decreasing order. At each index, we:

- Remove elements from the front of the deque if they are out of the current window ($i - \text{index} > k$).
- Use the front of the deque (if non-empty) as the maximum dp value from the previous k indices.
- Compute $dp[i]$ and update the global maximum answer.
- Remove from the back of the deque all indices with dp values less than the current $dp[i]$ since they are not useful.
- Insert the current index into the deque.

Trie

Round 6 · Mid · Hard · 4 questions

Round 6 · Mid · Hard

p.2561

Trie

#212 Word Search II

HAR

[LeetCode](#) · [SimplyLeet](#) · [WalkCC](#) · [LeetCode](#)

[Array](#) · [String](#) · [Backtracking](#) · [Trie](#) · [Matrix](#)

Lists: Grind 169

Problem

Description

Given an $m \times n$ board of characters and a list of strings words, return all words on the board.

Each word must be constructed from letters of sequentially adjacent cells, where adjacent cells are horizontally or vertically neighboring. The same letter cell may not be used more than once in a word.

Example 1:

Round 6 · Mid · Hard

p.2562

o	a	a	n
e	t	a	e
i	h	k	r
i	f	l	v

Input: board = [["o","a","a","n"],["e","t","a","e"],["i","h","k","r"],["i","f","l","v"]], words = ["oath","pea","eat","rain"]
Output: ["eat","oath"]

Example 2:

Round 6 · Mid · Hard

p.2563

a	b
c	d

Input: board = [["a","b"],["c","d"]], words = ["abcb"]
Output: []

Constraints:

- $m == \text{board.length}$
- $n == \text{board}[i].\text{length}$
- $1 \leq m, n \leq 12$
- $\text{board}[i][j]$ is a lowercase English letter.
- $1 \leq \text{words.length} \leq 3 * 10^4$
- $1 \leq \text{words}[i].\text{length} \leq 10$
- $\text{words}[i]$ consists of lowercase English letters.

Round 6 · Mid · Hard

p.2564

- All the strings of words are unique.

Community Solutions (Python)

• WalkCC

```
class TrieNode:
    def __init__(self):
        self.children: dict[str, TrieNode] = {}
        self.word: str | None = None

class Solution:
    def findWords(self, board: list[list[str]],
words: list[str]) -> list[str]:
        m = len(board)
        n = len(board[0])

        ans = []
        root = TrieNode()

        def insert(word: str) -> None:
            node = root
            for c in word:
                node = node.children.setdefault(c,
TrieNode())
            node.word = word

        for word in words:
```

```
insert(word)
```

```
def dfs(i: int, j: int, node: TrieNode) ->
None:
```

```
    if i < 0 or i == m or j < 0 or j == n:
        return
    if board[i][j] == '*':
        return
```

```
    c = board[i][j]
    if c not in node.children:
```

```
        return
```

```
    child = node.children[c]
    if child.word:
        ans.append(child.word)
        child.word = None
```

```
    board[i][j] = '*'
    dfs(i + 1, j, child)
    dfs(i - 1, j, child)
```

```

        dfs(i, j + 1, child)
        dfs(i, j - 1, child)
        board[i][j] = c

    for i in range(m):
        for j in range(n):
            dfs(i, j, root)

    return ans

```

• LeetCode

```

class Trie:
    def __init__(self):
        self.children: List[Trie | None] = [None] *
26
        self.ref: int = -1

    def insert(self, w: str, ref: int):
        node = self
        for c in w:
            idx = ord(c) - ord('a')
            if node.children[idx] is None:

```

```

        node.children[idx] = Trie()
        node = node.children[idx]
        node.ref = ref

```

```

class Solution:
    def findWords(self, board: List[List[str]],
words: List[str]) -> List[str]:
        def dfs(node: Trie, i: int, j: int):
            idx = ord(board[i][j]) - ord('a')
            if node.children[idx] is None:

```

```

                return
            node = node.children[idx]
            if node.ref >= 0:
                ans.append(words[node.ref])
                node.ref = -1
            c = board[i][j]
            board[i][j] = '#'
            for a, b in pairwise((-1, 0, 1, 0,
-1)):
                x, y = i + a, j + b
                if 0 <= x < m and 0 <= y < n and
board[x][y] != '#':

```



```

        dfs(node, x, y)
        board[i][j] = c

    tree = Trie()
    for i, w in enumerate(words):
        tree.insert(w, i)
    m, n = len(board), len(board[0])
    ans = []
    for i in range(m):
        for j in range(n):

```

```

        dfs(tree, i, j)
    return ans

```

• SimplyLeet

```

# Define a TrieNode for building the Trie.
class TrieNode:
    def __init__(self):
        self.children = {}
        self.word = None # Stores complete word at
the end node

class Solution:
    def findWords(self, board, words):
        # Build the Trie from the list of words.
        root = TrieNode()

```

```

        for word in words:
            node = root
            for char in word:
                if char not in node.children:
                    node.children[char] =
TrieNode()
                    node = node.children[char]
            node.word = word # Mark the end of a
word

```

```

        result = []
        m, n = len(board), len(board[0])

```

```

# Define the DFS helper function.
def dfs(i, j, node):
    char = board[i][j]
    if char not in node.children:
        return
    next_node = node.children[char]
    if next_node.word:
        result.append(next_node.word)
        next_node.word = None # Avoid
duplicate entries

```

```

        board[i][j] = '#' # Mark the current
cell as visited
        # Explore all four adjacent
directions.
        for dx, dy in ((0, 1), (1, 0), (0, -1),
(-1, 0)):
            ni, nj = i + dx, j + dy
            if 0 <= ni < m and 0 <= nj < n and
board[ni][nj] != '#':
                dfs(ni, nj, next_node)
            board[i][j] = char # Restore the
original character

        # Start DFS from each cell.

    for i in range(m):
        for j in range(n):
            dfs(i, j, root)
    return result

```

• LC.ca

```

class Trie:
    def __init__(self):
        self.children = [None] * 26
        self.w = ''
        # minor ajust, not a boolean is_end, but
the whole word
        # so it's easier for dfs to save the word

    def insert(self, w):
        node = self
        for c in w:

            idx = ord(c) - ord('a')
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]
        node.w = w

class Solution:
    def findWords(self, board: List[List[str]],
words: List[str]) -> List[str]:
        def dfs(node, i, j):

```

```

        idx = ord(board[i][j]) - ord('a')
        if node.children[idx] is None:
            return
        node = node.children[idx]
        if node.w:
            ans.add(node.w)
        c = board[i][j]
        board[i][j] = '0' # 0 for visited
already
        for a, b in [[0, -1], [0, 1], [1, 0],
                    [-1, 0]]:
            x, y = i + a, j + b

            if 0 <= x < m and 0 <= y < n and
board[x][y] != '0':
                dfs(node, x, y)
            board[i][y] = c

trie = Trie()
for w in words:
    trie.insert(w)
ans = set()
m, n = len(board), len(board[0])
for i in range(m):

```

```

        for j in range(n):
            dfs(trie, i, j)
        return list(ans)

#####

class Trie:
    def __init__(self):

26        self.children: List[Trie | None] = [None] * 26
        self.ref: int = -1

    def insert(self, w: str, ref: int):
        node = self
        for c in w:
            idx = ord(c) - ord('a')
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]

```

```
node.ref = ref
```

```
class Solution:
```

```
    def findWords(self, board: List[List[str]],
words: List[str]) -> List[str]:
```

```
        def dfs(node: Trie, i: int, j: int):
            idx = ord(board[i][j]) - ord('a')
            if node.children[idx] is None:
                return
            node = node.children[idx]
```

```
        if node.ref >= 0:
            ans.append(words[node.ref])
            node.ref = -1
        c = board[i][j]
        board[i][j] = '#'
        for a, b in pairwise((-1, 0, 1, 0,
-1)):
            x, y = i + a, j + b
            if 0 <= x < m and 0 <= y < n and
board[x][y] != '#':
                dfs(node, x, y)
        board[i][j] = c
```

```
tree = Trie()
for i, w in enumerate(words):
    tree.insert(w, i)
m, n = len(board), len(board[0])
ans = []
for i in range(m):
    for j in range(n):
        dfs(tree, i, j)
return ans
```

◆ Quick Review

Key Insights

- Use Depth-First Search (DFS) to explore possible words starting from each board cell.
- Build a Trie (prefix tree) to store the words for O(1) prefix checks and efficient pruning.
- Mark visited cells and restore them after DFS to avoid revisiting in the current search path.
- Prune search paths early if the current prefix is not part of any word in the Trie.

Space & Time

Time Complexity: $O(m * n * 4^L)$ in the worst case, where L is the maximum word length; however, using a Trie prunes unnecessary paths.

Space Complexity: $O(K)$ for the Trie storage, where K is the total number of letters in all words, plus $O(m * n)$ auxiliary

space for recursion and board marking.

Solution

We first insert all words into a Trie data structure to allow quick lookups of prefixes. Then, for each cell in the board, we perform DFS. During DFS, if the current path matches a Trie node that represents a complete word, we add that word to the result list and mark it as found to avoid duplicates. We continue exploring adjacent cells (up, down, left, right) while marking the cell as visited. After exploring, we restore the original letter for further DFS calls. This combination of Trie and DFS with backtracking efficiently finds valid words on the board.

Trie

#336 Palindrome Pairs

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · String · Trie

Lists: Grind 169

Problem

Description

You are given a 0-indexed array of unique strings words.

A palindrome pair is a pair of integers (i, j) such that:

- $0 \leq i, j < \text{words.length}$,
- $i \neq j$, and
- $\text{words}[i] + \text{words}[j]$ (the concatenation of the two strings) is a palindrome.

Return an array of all the palindrome pairs of words.

You must write an algorithm with $O(\sum \text{of words}[i].\text{length})$ runtime complexity.

Example 1:

Input: words = ["abcd","dcba","lls","s","sssll"]

Output: [[0,1],[1,0],[3,2],[2,4]]

Explanation: The palindromes are
["abccddcba","dcbaabcd","slls","llssssll"]

Example 2:

Input: words = ["bat","tab","cat"]
 Output: [[0,1],[1,0]]
 Explanation: The palindromes are ["battab","tabbat"]

Example 3:

Input: words = ["a",""]
 Output: [[0,1],[1,0]]
 Explanation: The palindromes are ["a","a"]

Constraints:

- $1 \leq \text{words.length} \leq 5000$
- $0 \leq \text{words}[i].\text{length} \leq 300$
- $\text{words}[i]$ consists of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def palindromePairs(self, words: list[str]) -> list[list[int]]:
        ans = []
        dict = {word[::-1]: i for i, word in enumerate(words)}

        for i, word in enumerate(words):
            if "" in dict and dict[""] != i and word == word[::-1]:
                ans.append([i, dict[""]])

            for j in range(1, len(word) + 1):
```

```
        l = word[:j]
        r = word[j:]
        if l in dict and dict[l] != i and r == r[::-1]:
            ans.append([i, dict[l]])
        if r in dict and dict[r] != i and l == l[::-1]:
            ans.append([dict[r], i])

        return ans
```

• LeetDoocs

```
class Solution:
    def palindromePairs(self, words: List[str]) -> List[List[int]]:
        d = {w: i for i, w in enumerate(words)}
        ans = []
        for i, w in enumerate(words):
            for j in range(len(w) + 1):
                a, b = w[:j], w[j:]
                ra, rb = a[::-1], b[::-1]
                if ra in d and d[ra] != i and b == rb:
                    ans.append([i, d[ra]])

                if j and rb in d and d[rb] != i and a == ra:
                    ans.append([d[rb], i])

        return ans
```

• SimplyLeet

Python implementation for Palindrome Pairs

```
def is_palindrome(s):
    # Helper function to check if a string is a
    # palindrome
    return s == s[::-1]

def palindromePairs(words):
    result = []
    # Create a dictionary mapping each word to its
    # index
    word_dict = {word: i for i, word in
    enumerate(words)}
```

```
    # Iterate over each word with its index
    for i, word in enumerate(words):
        # For each possible split position
        # (including edge empty string splits)
        for j in range(len(word) + 1):
            prefix = word[:j]
            suffix = word[j:]

            # If prefix is palindrome, check for
            # reversed suffix in dictionary
            if is_palindrome(prefix):
```

```
                reversed_suffix = suffix[::-1]
                # Ensure that the reversed_suffix
                # exists and is not the same word
                if reversed_suffix in word_dict
                and word_dict[reversed_suffix] != i:
                    result.append([word_dict[rever
                    sed_suffix], i])
```

```
            # Check the suffix; avoid duplicate
            # pairs when j equals len(word) because
            # the case when suffix is empty was
            # handled in the first condition
            if j != len(word) and
            is_palindrome(suffix):
                reversed_prefix = prefix[::-1]
                if reversed_prefix in word_dict
                and word_dict[reversed_prefix] != i:
```

```
                result.append([i,
                word_dict[reversed_prefix]])
```

```
    return result
```

```
# Example Usage:
words = ["abcd", "dcba", "lls", "s", "sssll"]
print(palindromePairs(words)) # Expected Output:
[[0,1],[1,0],[3,2],[2,4]]
```

• LC.ca

```

class Solution:
    def palindromePairs(self, words: List[str]) -> List[List[int]]:
        d = {w: i for i, w in enumerate(words)}
        ans = []
        for i, w in enumerate(words):
            for j in range(len(w) + 1):
                a, b = w[:j], w[j:]
                ra, rb = a[::-1], b[::-1]
                if ra in d and d[ra] != i and b ==
rb:
                    ans.append([i, d[ra]])

                if j and rb in d and d[rb] != i and
a == ra:
                    ans.append([d[rb], i])
        return ans

```

◆ Quick Review

Key Insights

- A palindrome is a string that reads the same forwards and backwards.
- For any given word, if you split it into two parts, one part (either prefix or suffix) might be a palindrome.
- If the other part's reverse exists elsewhere in the array, then that reverse can be concatenated to yield a palindrome.

- Carefully handle edge cases like empty strings, since an empty string is a palindrome by default.

Space & Time

Time Complexity: $O(n * k^2)$ in the worst case where n is the number of words and k is the maximum word length, but with the palindrome checking optimized this can be considered $O(\text{sum of word lengths})$: Each word is processed once and each split is checked in $O(k)$ time.

Space Complexity: $O(n * k)$ due to the hashmap storing all words and their indices.

Solution

We use a hashmap (dictionary) to store each word and its corresponding index. For every word, we iterate through every possible split position to divide the word into a prefix and a suffix. For each split:

– If the prefix is a palindrome, we look for the reverse of the suffix in the hashmap. If found (and not the same word), then combining the reverse with the current word forms a palindrome. – Similarly, if the suffix is a palindrome, we look for the reverse of the prefix.

We must be cautious with edge cases, especially when either part is empty, as the empty string is considered a palindrome. This method avoids checking every possible pair explicitly and leverages string reversal and palindrome checks to achieve efficiency.

Trie

#588 Design In-Memory File System

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Design · Trie

Lists: Grind 169 | Premium 98

Problem

Description

Design a data structure that simulates an in-memory file system.

Implement the `FileSystem` class:

- `FileSystem()` Initializes the object of the system.
- `List<String> ls(String path)` If path is a file path, returns a list that only contains this file's name.
- If path is a directory path, returns the list of file and directory names in this directory.
- If filePath does not exist, creates that file containing given content.
- If filePath already exists, appends the given content to original content.

Example 1:

Operation	Output	Explanation
<code>FileSystem fs = new FileSystem()</code>	<code>null</code>	The constructor returns nothing.
<code>fs.ls("/")</code>	<code>[]</code>	Initially, directory <code>/</code> has nothing. So return empty list.
<code>fs.mkdir("/a/b/c")</code>	<code>null</code>	Create directory <code>a</code> in directory <code>/</code> . Then create directory <code>b</code> in directory <code>a</code> . Finally, create directory <code>c</code> in directory <code>b</code> .
<code>fs.addContentToFile("/a/b/c/d", "hello")</code>	<code>null</code>	Create a file named <code>d</code> with content <code>"hello"</code> in directory <code>/a/b/c</code> .
<code>fs.ls("/")</code>	<code>["a"]</code>	Only directory <code>a</code> is in directory <code>/</code> .
<code>fs.readContentFromFile("/a/b/c/d")</code>	<code>"hello"</code>	Output the file content.

Input

["FileSystem", "ls", "mkdir", "addContentToFile", "ls", "readContentFromFile"]
[[], ["/"], ["/a/b/c"], ["/a/b/c/d", "hello"], ["/"], ["/a/b/c/d"]]

Output

[null, [], null, null, ["a"], "hello"]

Explanation

FileSystem fileSystem = new FileSystem();

Constraints:

- `1 <= path.length, filePath.length <= 100`
- `path` and `filePath` are absolute paths which begin with `'/'` and do not end with `'/'` except that the path is just `"/"`.
- You can assume that all directory names and file names only contain lowercase letters, and the same names will not exist in the same directory.

- You can assume that all operations will be passed valid parameters, and users will not attempt to retrieve file content or list a directory or file that does not exist.
- You can assume that the parent directory for the file in addContentToFile will exist.
- $1 \leq \text{content.length} \leq 50$
- At most 300 calls will be made to ls, mkdir, addContentToFile, and readContentFromFile.

Community Solutions (Python)

• LeetCode

```
class Trie:
    def __init__(self):
        self.name = None
        self.isFile = False
        self.content = []
        self.children = {}

    def insert(self, path, isFile):
        node = self
        ps = path.split('/')
        for p in ps[1:]:
            if p not in node.children:
                node.children[p] = Trie()
            node = node.children[p]
        node.name = ps[-1]
        node.isFile = isFile
```

```
for p in ps[1:]:
    if p not in node.children:
        node.children[p] = Trie()
    node = node.children[p]
node.isFile = isFile
if isFile:
    node.name = ps[-1]
return node

def search(self, path):
    node = self
    if path == '/':
        return node
    ps = path.split('/')
    for p in ps[1:]:
        if p not in node.children:
            return None
        node = node.children[p]
    return node
```

```

class FileSystem:
    def __init__(self):
        self.root = Trie()

    def ls(self, path: str) -> List[str]:
        node = self.root.search(path)
        if node is None:
            return []
        if node.isFile:

            return [node.name]
            return sorted(node.children.keys())

    def mkdir(self, path: str) -> None:
        self.root.insert(path, False)

    def addContentToFile(self, filePath: str,
content: str) -> None:
        node = self.root.insert(filePath, True)
        node.content.append(content)

```

```

        def readContentFromFile(self, filePath: str)
-> str:
            node = self.root.search(filePath)
            return ''.join(node.content)

# Your FileSystem object will be instantiated and
called as such:
# obj = FileSystem()
# param_1 = obj.ls(path)
# obj.mkdir(path)
# obj.addContentToFile(filePath,content)

# param_4 = obj.readContentFromFile(filePath)

```

• SimplyLeet

```

class Node:
    def __init__(self):
        # Indicates whether this node represents a
file or a directory
        self.is_file = False
        # For directories, mapping from name to
Node

        self.children = {}
        # For files, store content
        self.content = ""

class FileSystem:

```

```

def __init__(self):
    # Root directory node initialization
    self.root = Node()

def ls(self, path: str) -> list:
    # Traverse the path to obtain the target
node
    node, name = self.traverse(path)
    # If the node is a file, return the single
element list containing its name
    if node.is_file:

        return [name]
    # Otherwise, list all keys (names) in the
directory, sorted lexicographically
    return sorted(node.children.keys())

def mkdir(self, path: str) -> None:
    # Create directory by traversing the path
and creating nodes if necessary
    self.traverse(path)

def addContentToFile(self, filePath: str,
content: str) -> None:
    # Traverse to the file's node; if it does
not exist, create necessary nodes along the path

```

```

        node, name = self.traverse(filePath)
        # If not already a file, mark the node as a
file
        if not node.is_file:
            node.is_file = True
            node.content = ""
        # Append the new content to the file's
content
        node.content += content

    def readContentFromFile(self, filePath: str)
-> str:
        # Traverse to the file's node

        node, _ = self.traverse(filePath)
        # Return the content of the file
        return node.content

def traverse(self, path: str):
    # Start from the root node
    node = self.root
    if path == "/":
        # Special case for root path
        return node, ""

```

```

# Split the path by "/" and traverse the
nodes
parts = path.split("/")[1:]
for part in parts:
    # Create the node if it does not exist
    if part not in node.children:
        node.children[part] = Node()
    node = node.children[part]
# Return the target node and its name (last
part of the path)
return node, parts[-1]

```

```

# Example usage:
# fs = FileSystem()
# print(fs.ls("/")) # prints []
# fs.mkdir("/a/b/c")
# fs.addContentToFile("/a/b/c/d", "hello")
# print(fs.ls("/")) # prints ['a']
# print(fs.readContentFromFile("/a/b/c/d")) #
prints "hello"

```

• LC.ca

```

'''
>>> "/a/b/c".split('/')
['', 'a', 'b', 'c']
'''

class Trie:
    def __init__(self):
        self.name = None
        self.isFile = False
        self.content = []

        self.children = {}

    def insert(self, path, isFile):
        node = self
        ps = path.split('/')
        for p in ps[1:]:
            if p not in node.children:
                node.children[p] = Trie()
            node = node.children[p]
        node.isFile = isFile

```

```

    if isFile:
        node.name = ps[-1]
    return node

def search(self, path):
    node = self
    if path == '/':
        return node
    ps = path.split('/')
    for p in ps[1:]:

```

```

        if p not in node.children:
            return None
        node = node.children[p]
    return node

```

```

class FileSystem:
    def __init__(self):
        self.root = Trie()

```

```

def ls(self, path: str) -> List[str]:
    node = self.root.search(path)
    if node is None:
        return []
    if node.isFile:
        return [node.name]
    return sorted(node.children.keys())

def mkdir(self, path: str) -> None:
    self.root.insert(path, False)

```

```

def addContentToFile(self, filePath: str,
content: str) -> None:
    node = self.root.insert(filePath, True)
    node.content.append(content)

def readContentFromFile(self, filePath: str)
-> str:
    node = self.root.search(filePath)
    return ''.join(node.content)

```

```

# Your FileSystem object will be instantiated and
called as such:
# obj = FileSystem()
# param_1 = obj.ls(path)
# obj.mkdir(path)
# obj.addContentToFile(filePath, content)
# param_4 = obj.readContentFromFile(filePath)

```

◆ Quick Review

Key Insights

- Use a tree (trie) to represent directories and files.
- Each node can be either a directory or a file; directories maintain a mapping of names to child nodes.
- For ls, if the given path is a file, return only its name. If it's a directory, return the sorted list of names of its children.
- mkdir must create all intermediate directories if they do not already exist.
- File nodes store content and support content appending.

Space & Time

Time Complexity:

- ls: $O(k \log k)$ where k is the number of entries in the directory (due to sorting), or $O(n)$ if traversing a long path.
- mkdir, addContentToFile, readContentFromFile: $O(d)$ where d is the depth of the file path.

Space Complexity:

- $O(T)$ where T is the total number of characters stored in directory names and file contents.

Solution

The solution uses a tree-based approach where each node represents either a directory or a file. Each directory node stores its children in a hash table (or dictionary) mapping names to node objects. File nodes store their content as a string and carry a flag indicating they are files. For the ls

operation, check if the node is a file—if so, return its name; otherwise, list and sort the names of all children in that directory. The mkdir operation traverses (or creates) nodes for each segment of the provided path. The addContentToFile method accesses or creates the file node and appends the content. Finally, readContentFromFile retrieves the stored content from the file node.

Trie

#642 Design Search Autocomplete System

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Design · Trie · Data Stream · Heap

Lists: Premium 98

Problem

Description

Design a search autocomplete system for a search engine. Users may input a sentence (at least one word and end with a special character '#').

You are given a string array `sentences` and an integer array `times` both of length `n` where `sentences[i]` is a previously typed sentence and `times[i]` is the corresponding number of times the sentence was typed. For each input character except '#', return the top 3 historical hot sentences that have the same prefix as the part of the sentence already typed.

Here are the specific rules:

- The hot degree for a sentence is defined as the number of times a user typed the exactly same sentence before.
- The returned top 3 hot sentences should be sorted by hot degree (The first is the hottest one). If several

sentences have the same hot degree, use ASCII-code order (smaller one appears first).

- If less than 3 hot sentences exist, return as many as you can.
- When the input is a special character, it means the sentence ends, and in this case, you need to return an empty list.

Implement the `AutocompleteSystem` class:

- `AutocompleteSystem(String[] sentences, int[] times)` Initializes the object with the sentences and times arrays.
- `List<String> input(char c)` This indicates that the user typed the character `c`. Returns an empty array `[]` if `c == '#'` and stores the inputted sentence in the system.
- Returns the top 3 historical hot sentences that have the same prefix as the part of the sentence already typed. If there are fewer than 3 matches, return them all.

Example 1:

Input

```
["AutocompleteSystem", "input", "input", "input",
"input"]
[[["i love you", "island", "iroman", "i love
leetcode"], [5, 3, 2, 2]], ["i"], [" "], ["a"], ["#"]]
Output
[null, ["i love you", "island", "i love leetcode"], ["i
love you", "i love leetcode"], [], []]
```

Explanation

```
AutocompleteSystem obj = new AutocompleteSystem(["i
love you", "island", "iroman", "i love leetcode"], [5,
3, 2, 2]);
```

Constraints:

- $n == \text{sentences.length}$
- $n == \text{times.length}$
- $1 \leq n \leq 100$
- $1 \leq \text{sentences}[i].\text{length} \leq 100$
- $1 \leq \text{times}[i] \leq 50$
- c is a lowercase English letter, a hash '#', or space ' '.
- Each tested sentence will be a sequence of characters c that end with the character '#'.
- Each tested sentence will have a length in the range $[1, 200]$.
- The words in each input sentence are separated by single spaces.
- At most 5000 calls will be made to input.

Community Solutions (Python)

• WalkCC

```
class TrieNode:
    def __init__(self):
        self.children: dict[str, TrieNode] = {}
        self.s: str | None = None
        self.time = 0
        self.top3: list[TrieNode] = []

    def __lt__(self, other):
        if self.time == other.time:
            return self.s < other.s

        return self.time > other.time

    def update(self, node) -> None:
        if node not in self.top3:
            self.top3.append(node)
        self.top3.sort()
        if len(self.top3) > 3:
            self.top3.pop()
```

```

class AutocompleteSystem:
    def __init__(self, sentences: list[str], times:
list[int]):
        self.root = TrieNode()
        self.curr = self.root
        self.s: list[str] = []

        for sentence, time in zip(sentences, times):
            self._insert(sentence, time)

    def input(self, c: str) -> list[str]:

        if c == '#':
            self._insert(''.join(self.s), 1)
            self.curr = self.root
            self.s = []
            return []

        self.s.append(c)

        if self.curr:
            self.curr = self.curr.children.get(c, None)

```

```

        if not self.curr:
            return []
        return [node.s for node in self.curr.top3]

    def _insert(self, sentence: str, time: int) ->
None:
        node = self.root
        for c in sentence:
            node = node.children.setdefault(c,
TrieNode())
            node.s = sentence
            node.time += time

        leaf = node
        node: TrieNode = self.root
        for c in sentence:
            node = node.children[c]
            node.update(leaf)

```

- LeetCode

```
class Trie:
    def __init__(self):
        self.children = [None] * 27
        self.v = 0
        self.w = ''

    def insert(self, w, t):
        node = self
        for c in w:
            idx = 26 if c == ' ' else ord(c) -
ord('a')
```

```
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]
            node.v += t
            node.w = w

    def search(self, pref):
        node = self
        for c in pref:
            idx = 26 if c == ' ' else ord(c) -
ord('a')
```

```
            if node.children[idx] is None:
                return None
            node = node.children[idx]
        return node
```

```
class AutocompleteSystem:
    def __init__(self, sentences: List[str],
times: List[int]):
        self.trie = Trie()
        for a, b in zip(sentences, times):
```

```
            self.trie.insert(a, b)
            self.t = []

    def input(self, c: str) -> List[str]:
        def dfs(node):
            if node is None:
                return
            if node.v:
                res.append((node.v, node.w))
            for nxt in node.children:
```

```

        dfs(nxt)

    if c == '#':
        s = ''.join(self.t)
        self.trie.insert(s, 1)
        self.t = []
        return []

    res = []
    self.t.append(c)

    node = self.trie.search(''.join(self.t))
    if node is None:
        return res
    dfs(node)
    res.sort(key=lambda x: (-x[0], x[1]))
    return [v[1] for v in res[:3]]

# Your AutocompleteSystem object will be
# instantiated and called as such:
# obj = AutocompleteSystem(sentences, times)

# param_1 = obj.input(c)

```

• SimplyLeet

```

class TrieNode:
    def __init__(self):
        # Mapping from character to TrieNode
        self.children = {}
        # Dictionary mapping sentence to its
        frequency count
        self.counts = {}

class AutocompleteSystem:
    def __init__(self, sentences, times):
        self.root = TrieNode()

        self.current_node = self.root
        self.current_input = ""
        # Initialize Trie with the given sentences
        and their times
        for sentence, time in zip(sentences,
        times):
            self.addSentence(sentence, time)

        def addSentence(self, sentence, freq):
            node = self.root
            for char in sentence:
                if char not in node.children:

```

```

        node.children[char] = TrieNode()
        node = node.children[char]
        # Update sentence frequency at each
node
        node.counts[sentence] =
node.counts.get(sentence, 0) + freq

    def input(self, c):
        results = []
        if c == '#':
            # End of input: add current sentence to
Trie
            self.addSentence(self.current_input,
1)

        # Reset the current input and traversal
pointer
        self.current_input = ""
        self.current_node = self.root
        return results

        self.current_input += c
        # Move to the next Trie node if it exists
        if self.current_node:
            self.current_node =
self.current_node.children.get(c, None)
        if not self.current_node:

```

```

        return results

        # Retrieve sentences and sort them by
frequency (descending) and lexicographical order
(ascending)
        data =
list(self.current_node.counts.items())
        data.sort(key=lambda x: (-x[1], x[0]))
        for i in range(min(3, len(data))):
            results.append(data[i][0])
        return results

# Example usage:

# auto = AutocompleteSystem(["i love you",
"island", "iroman", "i love leetcode"], [5, 3, 2,
2])
# print(auto.input("i"))
# print(auto.input(" "))
# print(auto.input("a"))
# print(auto.input("#"))

```

- LC.ca

```
class Trie:
    def __init__(self):
        # 26+1, extra 1 for space ' ' node
        self.children = [None] * 27
        self.v = 0
        self.w = ''

    def insert(self, w, t):
        node = self
        for c in w:
```

```
            idx = 26 if c == ' ' else ord(c) -
ord('a')
            if node.children[idx] is None:
                node.children[idx] = Trie()
            node = node.children[idx]
            node.v += t
            node.w = w

    def search(self, pref):
        node = self
        for c in pref:
```

```
            idx = 26 if c == ' ' else ord(c) -
ord('a')
            if node.children[idx] is None:
                return None
            node = node.children[idx]
            return node
```

```
class AutocompleteSystem:
    def __init__(self, sentences: List[str],
times: List[int]):
        self.trie = Trie()
```

```
        for a, b in zip(sentences, times):
            self.trie.insert(a, b)
            self.t = [] # target

    def input(self, c: str) -> List[str]:
        def dfs(node):
            if node is None:
                return
            if node.v:
                res.append((node.v, node.w))
```

```

        for nxt in node.children:
            dfs(nxt)

    if c == '#':
        s = ''.join(self.t)
        self.trie.insert(s, 1)
        self.t = []
        return []

    res = []

    self.t.append(c)
    node = self.trie.search(''.join(self.t))
    if node is None:
        return res
    dfs(node)
    res.sort(key=lambda x: (-x[0], x[1]))
    return [v[1] for v in res[:3]]

```

Your AutocompleteSystem object will be instantiated and called as such:

```

# obj = AutocompleteSystem(sentences, times)
# param_1 = obj.input(c)

```

◆ Quick Review

Key Insights

- Use a Trie to efficiently store and retrieve sentences based on their prefixes.
- Maintain a mapping at each Trie node to record sentences (or their frequencies) that pass through that node.
- Use sorting or a min-heap to select the top three sentences based on their frequency (hot degree) and lexicographical order if frequencies tie.
- On receiving '#', update the Trie with the current sentence and reset the current search state.

Space & Time

Time Complexity: $O(L * \log k)$ per character input query, where L is the length of the current prefix and k is 3 (constant factor), thus ensuring efficient autocompletion.

Space Complexity: $O(N * L)$, where N is the number of unique sentences and L is the average sentence length, for maintaining the Trie.

Solution

The solution revolves around building a Trie where each node contains:

- A mapping to its child nodes.
- A hash map that records all sentences passing through the node along with their occurrence counts.

For each character entered (except '#'), the system:

- Traverses the Trie to find the current prefix node.
- Retrieves stored sentences at that node and sorts them based on the following criteria: descending by frequency and ascending lexicographically if frequencies are equal.
- Returns the top three sentences from the sorted result.

When the user inputs '#', the current sentence (tracked as user input so far) is added or its frequency updated in the Trie, and the input state resets.

Backtracking

Round 6 · Mid · Hard · 4 questions

Round 6 · Mid · Hard

p.2615

Round 6 · Mid · Hard

p.2616

Backtracking

#37 Sudoku Solver

HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Backtracking · Matrix

Lists: Grind 169

Problem

Description

Write a program to solve a Sudoku puzzle by filling the empty cells.

A sudoku solution must satisfy all of the following rules:

1. Each of the digits 1–9 must occur exactly once in each row.
2. Each of the digits 1–9 must occur exactly once in each column.
3. Each of the digits 1–9 must occur exactly once in each of the 9 3x3 sub-boxes of the grid.

The '.' character indicates empty cells.

Example 1:

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9

```
Input: board = [
["5","3",".",".","7",".",".",".","."],
["6",".",".","1","9","5",".",".","."],
[".","9","8",".",".",".",".","."],
[".",".","6","."],
["8",".",".","6",".",".","3"],
["4",".",".","8",".","3",".","1"],
["7",".",".","."],
["2",".",".","6"],
["6",".",".","2","8"],
[".","4","1","9",".","5"],
[".","8",".","7","9"]]
```

```
Output: [
["5","3","4","6","7","8","9","1","2"],
["6","7","2","1","9","5","3","4","8"],
["1","9","8","3","4","2","5","6","7"],
["8","5","9","7","6","1","4","2","3"],
["4","2","6","8","5","3","7","9","1"],
["7","1","3","9","2","4","8","5","6"],
["9","6","1","5","3","7","2","8","4"],
["2","8","7","4","1","9","6","3","5"],
["3","4","5","2","8","6","1","7","9"]]
```

Explanation: The input board is shown above and the only valid solution is shown below:

Constraints:

- `board.length == 9`
- `board[i].length == 9`
- `board[i][j]` is a digit or '.'.
- It is guaranteed that the input board has only one solution.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def solveSudoku(self, board: list[list[str]]) -> None:
        def isValid(row: int, col: int, c: str) -> bool:
            for i in range(9):
                if (board[i][col] == c or
                    board[row][i] == c or
                    board[3 * (row // 3) + i // 3][3 *
                        (col // 3) + i % 3] == c):
                    return False
            return True
```

```
def solve(s: int) -> bool:
    if s == 81:
        return True

    i = s // 9
    j = s % 9

    if board[i][j] != '.':
        return solve(s + 1)
```

```

for c in string.digits[1:]:
    if isValid(i, j, c):
        board[i][j] = c
        if solve(s + 1):
            return True
        board[i][j] = '.'

return False

solve(0)

```

• LeetCode

```

class Solution:
    def solveSudoku(self, board: List[List[str]])
    -> None:
        def dfs(k):
            nonlocal ok
            if k == len(t):
                ok = True
                return
            i, j = t[k]
            for v in range(9):
                if row[i][v] == col[j][v] ==
                block[i // 3][j // 3][v] == False:

```

```

        row[i][v] = col[j][v] =
        block[i // 3][j // 3][v] = True
        board[i][j] = str(v + 1)
        dfs(k + 1)
        row[i][v] = col[j][v] =
        block[i // 3][j // 3][v] = False
        if ok:
            return

        row = [[False] * 9 for _ in range(9)]
        col = [[False] * 9 for _ in range(9)]
        block = [[[False] * 9 for _ in range(3)]
        for _ in range(3)]

```

```

        t = []
        ok = False
        for i in range(9):
            for j in range(9):
                if board[i][j] == '.':
                    t.append((i, j))
                else:
                    v = int(board[i][j]) - 1
                    row[i][v] = col[j][v] =
                    block[i // 3][j // 3][v] = True
        dfs(0)

```

• SimplyLeet

```
def solveSudoku(board):
    # Check if placing the digit 'char' at
    board[row][col] is valid
    def is_valid(row, col, char):
        for i in range(9):
            # Check the row
            if board[row][i] == char:
                return False
            # Check the column
            if board[i][col] == char:
                return False

        # Check the 3x3 sub-box
        sub_row = 3 * (row // 3) + i // 3
        sub_col = 3 * (col // 3) + i % 3
        if board[sub_row][sub_col] == char:
            return False
        return True

    # Backtracking function to fill the board
    def backtrack():
        for row in range(9):
```

```
        for col in range(9):
            if board[row][col] == '.':
                for char in map(str, range(1,
10)):
                    if is_valid(row, col,
char):
                        board[row][col] = char
# Place the digit
                        if backtrack():
                            return True #
Recurse; if solution found, bubble up True
                        board[row][col] = '.'
# Backtrack if not valid further down
                            return False # No valid digit
found; trigger backtracking
                            return True # Board completely filled
```

```
backtrack()
```

```
# Example usage:
board = [
    ["5","3",".", ".", "7",".", ".", ".", "."],
    ["6",".", ".", "1","9","5",".", ".", "."],
    [".","9","8",".", ".", ".", ".", "6","."],
    ["8",".", ".", ".", "6",".", ".", ".", "3"],
    ["4",".", ".", "8",".", "3",".", ".", "1"],
    ["7",".", ".", ".", "2",".", ".", ".", "6"],
```

```

        [".","6",".", ".", ".", ".", "2", "8", "."],
        [".", ".", ".", "4", "1", "9", ".", ".", "5"],
        [".", ".", ".", ".", "8", ".", ".", "7", "9"]
    solveSudoku(board)
    print(board)

```

• LC.ca

```

class Solution:
    def solveSudoku(self, board: List[List[str]])
-> None:
    def dfs(k):
        nonlocal ok
        if k == len(t):
            ok = True
            return
        i, j = t[k]
        for v in range(9):
            if row[i][v] == col[j][v] ==
block[i // 3][j // 3][v] == False:

```

```

        row[i][v] = col[j][v] =
block[i // 3][j // 3][v] = True
        board[i][j] = str(v + 1)
        dfs(k + 1)
        row[i][v] = col[j][v] =
block[i // 3][j // 3][v] = False
        if ok:
            return

        row = [[False] * 9 for _ in range(9)]
        col = [[False] * 9 for _ in range(9)]
        block = [[[False] * 9 for _ in range(3)]
for _ in range(3)]

```

```

        t = []
        ok = False
        for i in range(9):
            for j in range(9):
                if board[i][j] == '.':
                    t.append((i, j))
                else:
                    v = int(board[i][j]) - 1
                    row[i][v] = col[j][v] =
block[i // 3][j // 3][v] = True
        dfs(0)

```

◆ Quick Review

Key Insights

- Use backtracking to explore digit placements for each empty cell.
- Validate placements by ensuring that the chosen digit is not already present in the corresponding row, column, or 3x3 sub-box.
- Employ recursion to fill the board incrementally and backtrack upon encountering invalid configurations.
- Although the worst-case time complexity is exponential, the fixed board size (9x9) keeps the problem computationally manageable.

Space & Time

Time Complexity: In the worst-case scenario, the algorithm can explore up to $O(9^n)$ possibilities where n is the number of empty cells. However, since n is at most 81 and many branches are pruned early, the practical runtime is far less.

Space Complexity: $O(n)$ due to the recursion stack, where n is the number of empty cells; given the board's fixed size, this space is constant in practice.

Solution

The solution uses a backtracking approach. For each empty cell on the board, we attempt to place a digit from '1' to '9'. Before placing a digit, we check if the placement is valid by ensuring that:

- The digit is not already present in the same row.
- The digit is not already present in the same column.
- The digit is not already present in the corresponding 3x3 sub-box.

If a valid digit is found, we place it on the board and recursively attempt to solve the rest of the board. If no valid digit can be placed, the algorithm backtracks by reverting the last move and trying a different digit. This recursive search continues until the board is completely filled with a valid Sudoku configuration.

Backtracking

#51 N-Queens

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Backtracking

Lists: Grind 169

Problem


Description

The n-queens puzzle is the problem of placing n queens on an n x n chessboard such that no two queens attack each other.

Given an integer n, return all distinct solutions to the n-queens puzzle. You may return the answer in any order.

Each solution contains a distinct board configuration of the n-queens' placement, where 'Q' and '.' both indicate a queen and an empty space, respectively.

Example 1:



Input: n = 4

Output: [[".Q..", "...Q", "Q...", "...Q."], [".Q..", "Q...", "...Q", ".Q.."]]

Explanation: There exist two distinct solutions to the 4-queens puzzle as shown above

Example 2:

Input: n = 1

Output: [["Q"]]

Constraints:

- $1 \leq n \leq 9$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def solveNQueens(self, n: int) ->
list[list[str]]:
    ans = []
    cols = [False] * n
    diag1 = [False] * (2 * n - 1)
    diag2 = [False] * (2 * n - 1)

    def dfs(i: int, board: list[int]) -> None:
        if i == n:
            ans.append(board)

        return

        for j in range(n):
            if cols[j] or diag1[i + j] or diag2[j - i +
n - 1]:
                continue
            cols[j] = diag1[i + j] = diag2[j - i + n -
1] = True
            dfs(i + 1, board + ['. ' * j + 'Q' + ' ' *
(n - j - 1)])
            cols[j] = diag1[i + j] = diag2[j - i + n -
1] = False

        dfs(0, [])

    return ans
```

• LeetCode

Round 6 · Mid · Hard

p.2631

```
class Solution:
    def solveNQueens(self, n: int) ->
List[List[str]]:
    def dfs(i: int):
        if i == n:
            ans.append(["".join(row) for row
in g])

            return
        for j in range(n):
            if col[j] + dg[i + j] + udg[n - i +
j] == 0:

                g[i][j] = "Q"
                col[j] = dg[i + j] = udg[n - i
+ j] = 1

                dfs(i + 1)
                col[j] = dg[i + j] = udg[n - i
+ j] = 0

                g[i][j] = "."

        ans = []
        g = [["."] * n for _ in range(n)]
        col = [0] * n
        dg = [0] * (n << 1)
        udg = [0] * (n << 1)
        dfs(0)

    return ans
```

• SimplyLeet

Round 6 · Mid · Hard

p.2632


```
def solveNQueens(n):
    result = []

    # Helper function to backtrack
    def backtrack(row, columns, diag1, diag2,
board):
        # Base case: if all rows are filled, add a
copy of board to result
        if row == n:
            result.append(list(board))
            return

        # Try to place a queen in each column in
the current row
        for col in range(n):
            # Check if the column or diagonals are
not under attack
            if col in columns or (row - col) in
diag1 or (row + col) in diag2:
                continue

            # Place the queen
            board[row] = '.' * col + 'Q' + '.' * (n
- col - 1)
            # Mark the column and diagonals as
occupied
            columns.add(col)
```

```
            diag1.add(row - col)
            diag2.add(row + col)

            # Move to the next row
            backtrack(row + 1, columns, diag1,
diag2, board)

            # Backtrack: remove the queen and
unmark the sets
            columns.remove(col)
            diag1.remove(row - col)
            diag2.remove(row + col)

        # Initialize board with empty rows
        board = ["." * n for _ in range(n)]
        backtrack(0, set(), set(), set(), board)
        return result

# Example usage:
print(solveNQueens(4))
```

- LC.ca

```

class Solution:
    def solveNQueens(self, n: int) ->
List[List[str]]:
    def dfs(i: int):
        if i == n:
            ans.append(["".join(row) for row
in g])
            return
        for j in range(n):
            if col[j] + dg[i + j] + udg[n - i +
j] == 0:
                g[i][j] = "Q"
                col[j] = dg[i + j] = udg[n - i
+ j] = 1

                dfs(i + 1)
                col[j] = dg[i + j] = udg[n - i
+ j] = 0

                g[i][j] = "."

        ans = []
        g = [["."] * n for _ in range(n)]
        col = [0] * n
        dg = [0] * (n << 1)
        udg = [0] * (n << 1)
        dfs(0)

    return ans

```

◆ Quick Review

Round 6 · Mid · Hard

p.2635

Key Insights

- Use backtracking to explore potential positions row by row.
- Use sets to track which columns and diagonals are under attack.
- Diagonals can be tracked by (row - col) and (row + col), ensuring that queens do not conflict along any diagonal.
- Recursively build each valid board configuration and backtrack when a conflict arises.

Space & Time

Time Complexity: $O(n!)$ in the worst-case scenario, as the solution explores permutations for queen placements.

Space Complexity: $O(n^2)$ for storing board configurations and $O(n)$ for recursion stack and conflict-tracking sets.

Solution

We use a backtracking algorithm that places queens row by row. For each row, we try placing a queen in each column. Before placing, we check if the column, main diagonal (row - col), and anti-diagonal (row + col) are free using dedicated sets. If a queen can be safely placed, we mark the column and diagonals as occupied and recurse to the next row. Once all rows are processed, the board configuration is a valid solution and is added to the result. After that, we backtrack by unmarking the sets and removing the queen from the board. This systematic approach ensures all possible configurations are considered.

Round 6 · Mid · Hard

p.2636

Backtracking

#126 Word Ladder II

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · String · Backtracking · BFS

Lists: Denny Zhang

Problem

Description

A transformation sequence from word `beginWord` to word `endWord` using a dictionary `wordList` is a sequence of words `beginWord` -> `s1` -> `s2` -> ... -> `sk` such that:

- Every adjacent pair of words differs by a single letter.
- Every `si` for $1 \leq i \leq k$ is in `wordList`. Note that `beginWord` does not need to be in `wordList`.
- `sk == endWord`

Given two words, `beginWord` and `endWord`, and a dictionary `wordList`, return all the shortest transformation sequences from `beginWord` to `endWord`, or an empty list if no such sequence exists. Each sequence should be returned as a list of the words [`beginWord`, `s1`, `s2`, ..., `sk`].

Example 1:

Input: `beginWord` = "hit", `endWord` = "cog", `wordList` = ["hot","dot","dog","lot","log","cog"]
 Output: [["hit","hot","dot","dog","cog"],["hit","hot","lot","log","cog"]]
 Explanation: There are 2 shortest transformation sequences:
 "hit" -> "hot" -> "dot" -> "dog" -> "cog"
 "hit" -> "hot" -> "lot" -> "log" -> "cog"

Example 2:

Input: `beginWord` = "hit", `endWord` = "cog", `wordList` = ["hot","dot","dog","lot","log"]
 Output: []
 Explanation: The `endWord` "cog" is not in `wordList`, therefore there is no valid transformation sequence.

Constraints:

- $1 \leq \text{beginWord.length} \leq 5$
- `endWord.length == beginWord.length`
- $1 \leq \text{wordList.length} \leq 500$
- `wordList[i].length == beginWord.length`
- `beginWord`, `endWord`, and `wordList[i]` consist of lowercase English letters.
- `beginWord != endWord`
- All the words in `wordList` are unique.
- The sum of all shortest transformation sequences does not exceed 105.

Community Solutions (Python)

- WalkCC

```
class Solution:
    def findLadders(self, beginWord: str, endWord:
str, wordList: list[str]) -> list[list[str]]:
        wordSet = set(wordList)
        if endWord not in wordList:
            return []

        # {"hit": ["hot"], "hot": ["dot", "lot"], ...}
        graph: dict[str, list[str]] =
collections.defaultdict(list)
```

```
        # Build the graph from the beginWord to the
endWord.
```

```
        if not self._bfs(beginWord, endWord, wordSet,
graph):
            return []
```

```
        ans = []
        self._dfs(graph, beginWord, endWord,
[beginWord], ans)
        return ans
```

```
    def _bfs(
        self,
        beginWord: str,
```

```
        endWord: str,
        wordSet: set[str],
        graph: dict[str, list[str]],
    ) -> bool:
        currentLevelWords = {beginWord}
```

```
        while currentLevelWords:
            for word in currentLevelWords:
                wordSet.discard(word)
                nextLevelWords = set()
```

```
            reachEndWord = False
            for parent in currentLevelWords:
                for child in self._getChildren(parent,
wordSet):
                    if child in wordSet:
                        nextLevelWords.add(child)
                        graph[parent].append(child)
                    if child == endWord:
                        reachEndWord = True
            if reachEndWord:
                return True
```

```

    currentLevelWords = nextLevelWords

    return False

    def _getChildren(self, parent: str, wordSet:
set[str]) -> list[str]:
        children = []
        s = list(parent)

        for i, cache in enumerate(s):
            for c in string.ascii_lowercase:

                if c == cache:
                    continue
                s[i] = c
                child = ''.join(s)
                if child in wordSet:
                    children.append(child)
                s[i] = cache

        return children

```

Round 6 · Mid · Hard

p.2641

```

def _dfs(
    self,
    graph: dict[str, list[str]],
    word: str,
    endWord: str,
    path: list[str],
    ans: list[list[str]],
) -> None:
    if word == endWord:
        ans.append(path.copy())

    return

    for child in graph.get(word, []):
        path.append(child)
        self._dfs(graph, child, endWord, path, ans)
        path.pop()

```

• LeetCode

```

class Solution:
    def findLadders(
        self, beginWord: str, endWord: str,
        wordList: List[str]
    ) -> List[List[str]]:
        def dfs(path, cur):
            if cur == endWord:
                ans.append(path[:])
                return
            for precursor in prev[cur]:
                path.append(precursor)

```

Round 6 · Mid · Hard

p.2642

```

        dfs(path, precursor)
        path.pop()

```

```

ans = []
words = set(wordList)
if endWord not in words:
    return ans
words.discard(beginWord)
dist = {beginWord: 0}
prev = defaultdict(set)

```

```

q = deque([beginWord])
found = False
step = 0
while q and not found:
    step += 1
    for i in range(len(q), 0, -1):
        p = q.popleft()
        s = list(p)
        for i in range(len(s)):
            ch = s[i]

```

```

        for j in range(26):
            s[i] = chr(ord('a') + j)
            t = ''.join(s)
            if dist.get(t, 0) == step:
                prev[t].add(p)
            if t not in words:
                continue
            prev[t].add(p)
            words.discard(t)
            q.append(t)

```

```

                dist[t] = step
                if endWord == t:
                    found = True
                    s[i] = ch
            if found:
                path = [endWord]
                dfs(path, endWord)
            return ans

```

• SimplyLeet

```
# Python solution for Word Ladder II
```

```
from collections import defaultdict, deque
```

```
def findLadders(beginWord, endWord, wordList):
    # Convert list to set for fast lookup
    wordSet = set(wordList)
    if endWord not in wordSet:
        return []
```

```
    # Dictionary to hold parent pointers for
    # backtracking: child -> list of parents
    parents = defaultdict(list)
    # Set for words visited at current BFS level
    level = {beginWord}
    # Flag to stop BFS when endWord is reached
    found = False
```

```
    while level and not found:
        next_level = set()
        # Remove words of current level from
        wordSet to prevent cycles
```

```
        wordSet -= level
        for word in level:
            # Try all possible one-letter
            transformations
            for i in range(len(word)):
                for char in
                'abcdefghijklmnopqrstuvwxyz':
                    candidate = word[:i] + char +
                    word[i+1:]

                    if candidate in wordSet:
                        # If candidate is the
                        endWord, note that we have found a path
                        if candidate == endWord:
                            found = True
```

```
                    # Record the current word
                    as a parent of candidate word
                    parents[candidate].append(
                    word)

                    next_level.add(candidate)
                    # Move to the next level of the BFS
                    level = next_level
```

```
    # The result list to hold all valid
    transformation sequences
    res = []
```

```
    # Helper function to backtrack from endWord to
    beginWord using the parents dictionary
```

```

def backtrack(word, path):
    if word == beginWord:
        # Append reversed path to have the
        sequence from beginWord to endWord
        res.append(path[::-1])
        return
    # Recurse over all parents of the current
word
    for parent in parents[word]:
        backtrack(parent, path + [parent])

if found:

    backtrack(endWord, [endWord])
    return res

# Example usage:
# print(findLadders("hit", "cog",
["hot","dot","dog","lot","log","cog"]))

```

- LC.ca

```

'''
remove() same as discard()

>>> a = set([11,22,33])
>>> a.remove(22)
>>> a
{33, 11}

>>> a = set([11,22,33])
>>> a.discard(22)

>>> a
{33, 11}

>>> a=set([1,2,3])
>>> a.discard(555)
>>> a.remove(555)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
KeyError: 555

```



```

'''

class Solution:
    def findLadders(
        self, beginWord: str, endWord: str,
        wordList: List[str]
    ) -> List[List[str]]:
        # endWord to beginWord
        # better than begin to end, too many paths
        # not in final result
        def dfs(path, cur):
            if cur == beginWord:
                ans.append(path[::-1])
                return
            for precursor in prev[cur]:
                path.append(precursor)
                dfs(path, precursor)
                path.pop()

        ans = []
        words = set(wordList)
        if endWord not in words:

```

Round 6 · Mid · Hard

p.2649

```

        return ans
        # no exception if beginWord not in set
        words.discard(beginWord)
        dist = {beginWord: 0}
        prev = defaultdict(set)
        q = deque([beginWord])
        found = False
        step = 0
        while q and not found:
            step += 1

```

```

        for _ in range(len(q), 0, -1):
            p = q.popleft()
            s = list(p)
            for i in range(len(s)):
                ch = s[i]
                for j in range(26):
                    s[i] = chr(ord('a') + j)
                    t = ''.join(s)
                    if dist.get(t, 0) == step:
                        prev[t].add(p) #
repeated 3 lines below

```

Round 6 · Mid · Hard

p.2650

```

        if t not in words:
            # if above '== step' met,
            then t must be removed from words[] from previous
            iterations

```

```

        continue
        prev[t].add(p) # repeated

3 lines above
        words.discard(t)
        q.append(t)
        dist[t] = step
        if endWord == t:
            found = True
        s[i] = ch

```

```

    if found:
        path = [endWord]
        dfs(path, endWord)
    return ans

```

```
#####
```

```

from collections import deque

class Solution(object):

```

```

def findLadders(self, beginWord, endWord,
wordlist):

```

```

    """
    :type beginWord: str
    :type endWord: str
    :type wordlist: Set[str]
    :rtype: List[List[int]]
    """

```

```

def getNbrs(src, dest, wordList):
    res = []

```

```

        for c in string.ascii_lowercase:
            for i in range(0, len(src)):
                newWord = src[:i] + c + src[i + 1:]
                if newWord == src:
                    continue
                if newWord in wordList or newWord ==
dest:
                    yield newWord

```

```

def bfs(beginWord, endWord, wordList):
    distance = {beginWord: 0}

```

```

queue = deque([beginWord])
length = 0
while queue:
    length += 1
    for k in range(0, len(queue)):
        top = queue.popleft()
        for nbr in getNbrs(top, endWord,
wordList):
            if nbr not in distance:
                distance[nbr] = distance[top] + 1
                queue.append(nbr)

return distance

def dfs(beginWord, endWord, wordList, path,
res, distance):
    if beginWord == endWord:
        res.append(path + [])
        return

    for nbr in getNbrs(beginWord, endWord,
wordList):
        if distance.get(nbr, -2) + 1 ==
distance[beginWord]:
            path.append(nbr)

```

```

        dfs(nbr, endWord, wordList, path, res,
distance)
        path.pop()

    res = []
    distance = bfs(endWord, beginWord, wordlist)
    dfs(beginWord, endWord, wordlist, [beginWord],
res, distance)
    return res

```

◆ Quick Review

Key Insights

- Use Breadth-First Search (BFS) to traverse level-by-level and capture only the shortest paths.
- Build a parent mapping during BFS to record which words lead to each newly discovered word.
- Once the shortest path is found via BFS, use backtracking (or DFS) from endWord to reconstruct all transformation sequences.
- Remove words as they are visited within a level to avoid unnecessary revisits and cycles.
- The solution must handle multiple shortest paths, so maintaining a list of predecessors for each word is crucial.

Space & Time

Time Complexity: $O(N * L * 26)$ where N is the number of words in the wordList and L is the length of each word. This accounts for checking all possible one-letter

transformations.

Space Complexity: $O(N * L)$ for storing the parent mapping and maintaining the BFS queue and other auxiliary data structures.

Solution

We solve the problem in two major phases:

- Use BFS to explore from beginWord to endWord. During BFS, record all parents for each word encountered. This ensures we capture all possible predecessors that yield a shortest transformation. – Once the BFS completes (once endWord is reached), use backtracking (DFS) beginning from the endWord and moving backwards via the parent mapping to reconstruct all transformation sequences in reverse. Finally, reverse each sequence to present it from beginWord to endWord.

This two-phase approach guarantees that only the shortest paths are reconstructed, as BFS ensures level-by-level exploration.

Backtracking

#996 Number of Squareful Arrays

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Math · Dynamic Programming · Backtracking

Lists: Denny Zhang

Problem

Description

An array is squareful if the sum of every pair of adjacent elements is a perfect square.

Given an integer array `nums`, return the number of permutations of `nums` that are squareful.

Two permutations `perm1` and `perm2` are different if there is some index `i` such that `perm1[i] != perm2[i]`.

Example 1:

Input: `nums = [1,17,8]`

Output: 2

Explanation: `[1,8,17]` and `[17,8,1]` are the valid permutations.

Example 2:

Input: `nums = [2,2,2]`

Output: 1

Constraints:

- $1 \leq \text{nums.length} \leq 12$

- $0 \leq \text{nums}[i] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numSquarefulPerms(self, nums: list[int]) -> int:
        ans = 0
        used = [False] * len(nums)

        def isSquare(num: int) -> bool:
            root = math.isqrt(num)
            return root * root == num

        def dfs(path: list[int]) -> None:
            nonlocal ans
            if len(path) > 1 and not isSquare(path[-1] + path[-2]):
                return
            if len(path) == len(nums):
                ans += 1
                return

            for i, a in enumerate(nums):
                if used[i]:
                    continue
```

```
        if i > 0 and nums[i] == nums[i - 1] and not used[i - 1]:
            continue
        used[i] = True
        dfs(path + [a])
        used[i] = False

        nums.sort()
        dfs([])
        return ans
```

• LeetCode

```
class Solution:
    def numSquarefulPerms(self, nums: List[int]) -> int:
        n = len(nums)
        f = [[0] * n for _ in range(1 << n)]
        for j in range(n):
            f[1 << j][j] = 1
        for i in range(1 << n):
            for j in range(n):
                if i >> j & 1:
                    for k in range(n):
```

```

                if (i >> k & 1) and k != j:
                    s = nums[j] + nums[k]
                    t = int(sqrt(s))
                    if t * t == s:
                        f[i][j] += f[i ^
(1 << j)][k]

            ans = sum(f[(1 << n) - 1][j] for j in
range(n))
            for v in Counter(nums).values():
                ans //= factorial(v)
            return ans

```

• SimplyLeet

```

from math import isqrt

# Function to check if a number is a perfect square
def is_square(num):
    root = isqrt(num)
    return root * root == num

def numSquarefulPerms(nums):
    n = len(nums)
    nums.sort() # Sort the array to handle
duplicates

```

```

visited = [False] * n
result = 0

def dfs(last_index, count):
    nonlocal result
    # If all elements have been used, we found
a valid permutation
    if count == n:
        result += 1
        return
    for i in range(n):
        if visited[i]:
            continue
        # Skip duplicates: if the current
element is the same as the previous one and the
previous one hasn't been used
        if i > 0 and nums[i] == nums[i - 1] and
not visited[i - 1]:
            continue
        # If it's the first element or the sum
with the last element is a perfect square
        if last_index == -1 or
is_square(nums[last_index] + nums[i]):
            visited[i] = True
            dfs(i, count + 1)
            visited[i] = False

```

```
dfs(-1, 0)
return result
```

Example usage:

```
if __name__ == "__main__":
    print(numSquarefulPerms([1, 17, 8])) #
Expected output: 2
    print(numSquarefulPerms([2, 2, 2])) # Expected
output: 1
```

• LC.ca

class Solution:

```
    def numSquarefulPerms(self, nums: List[int])
-> int:
    n = len(nums)
    f = [[0] * n for _ in range(1 << n)]
    for j in range(n):
        f[1 << j][j] = 1
    for i in range(1 << n):
        for j in range(n):
            if i >> j & 1:
                for k in range(n):
```

```
                if (i >> k & 1) and k != j:
                    s = nums[j] + nums[k]
                    t = int(sqrt(s))
                    if t * t == s:
                        f[i][j] += f[i ^
(1 << j)][k]

        ans = sum(f[(1 << n) - 1][j] for j in
range(n))
        for v in Counter(nums).values():
            ans //= factorial(v)
        return ans
```

◆ Quick Review

Key Insights

- The problem requires ensuring adjacent pairs sum up to a perfect square.
- Sorting the array helps in efficiently handling duplicates; when iterating in the backtracking, skip duplicate elements to avoid overcounting.
- For each valid candidate, check that the sum with the previous number (if any) is a perfect square.
- Use backtracking (DFS) to generate all valid permutations while maintaining a visited flag for each index.
- Since the array may contain duplicates, extra care is taken to only use each occurrence in a controlled order (i.e., if two identical numbers exist, only the first unvisited instance can be chosen at a given recursive call).

Space & Time

Time Complexity: $O(n^2 * 2^n)$ in the worst-case scenario due to exploring all bitmask states with n elements and checking pairwise sums.

Space Complexity: $O(n)$ for recursion and $O(n)$ for the visited flag, plus potentially $O(2^n * n)$ for memoization if implemented.

Solution

The solution uses a backtracking algorithm to explore all permutations. The array is first sorted to handle duplicates. A helper function checks if a number is a perfect square and is used to validate that the sum of every adjacent pair is a perfect square. During DFS, a visited array is maintained to track elements included so far. Skipping duplicate elements (unless the earlier duplicate was used) is essential to avoid overcounting. This approach ensures that every valid permutation is considered exactly once.

Round 7 · Low · Easy

Round 7

Low Priority · Easy**14 questions · 3 patterns**

Dynamic Programming #70 #121

Bit Manipulation #67 #136 #190 #191 #268 #338

Math #9 #13 #504 #1056 #1228 #1427

Round 6 · Mid · Hard

p.2663

Round 6 · Mid · Hard

p.2664

Dynamic Programming

Round 7 · Low · Easy · 2 questions

Round 7 · Low · Easy

p.2665

Dynamic Programming

#70 Climbing Stairs

EAS

[LeetDoocs](#) · [SimplyLeet](#) · [WalkCC](#) · [LeetCode](#)

[Math](#) · [Dynamic Programming](#) · [Memoization](#)

[Lists: Grind 169](#)

Problem

Description

You are climbing a staircase. It takes n steps to reach the top.

Each time you can either climb 1 or 2 steps. In how many distinct ways can you climb to the top?

Example 1:

Input: $n = 2$

Output: 2

Explanation: There are two ways to climb to the top.

1. 1 step + 1 step

2. 2 steps

Example 2:

Round 7 · Low · Easy

p.2666

Input: n = 3

Output: 3

Explanation: There are three ways to climb to the top.

1. 1 step + 1 step + 1 step

2. 1 step + 2 steps

3. 2 steps + 1 step

Constraints:

- $1 \leq n \leq 45$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def climbStairs(self, n: int) -> int:
        # dp[i] := the number of ways to climb to the
        # i-th stair
        dp = [1, 1] + [0] * (n - 1)

        for i in range(2, n + 1):
            dp[i] = dp[i - 1] + dp[i - 2]

        return dp[n]
```

class Solution:

def climbStairs(self, n: int) -> int:

prev1 = 1 # dp[i - 1]

prev2 = 1 # dp[i - 2]

for _ in range(2, n + 1):

dp = prev1 + prev2

prev2 = prev1

prev1 = dp

return prev1

• LeetDoocs

class Solution:

def climbStairs(self, n: int) -> int:

a, b = 0, 1

for _ in range(n):

a, b = b, a + b

return b

```
import numpy as np

class Solution:
    def climbStairs(self, n: int) -> int:
        res = np.asmatrix([(1, 1)], np.dtype("0"))
        factor = np.asmatrix([(1, 1), (1, 0)],
                               np.dtype("0"))
        n -= 1
        while n:
            if n & 1:

                res *= factor
                factor *= factor
                n >= 1
        return res[0, 0]
```

• SimplyLeet

```
# Define a function to compute the number of ways
to climb stairs.
def climbStairs(n):
    # Base case: if there is only one step, return
    1.
    if n == 1:
        return 1
    # Initialize the dp array with n+1 elements.
    dp = [0] * (n + 1)
    # There is 1 way to climb 1 step.
    dp[1] = 1
    # There are 2 ways to climb 2 steps.
```

```
dp[2] = 2
# Compute the number of ways for each step from
3 to n.
for i in range(3, n + 1):
    dp[i] = dp[i - 1] + dp[i - 2] # Sum of ways
    from the previous two steps.
return dp[n]

# Example usage:
print(climbStairs(3)) # Output: 3
```

• LC.ca

```
class Solution:
    def climbStairs(self, n: int) -> int:
        a, b = 0, 1
        for _ in range(n):
            a, b = b, a + b
        return b
```

◆ Quick Review

Key Insights

- The problem can be solved using dynamic programming because the solution for a given step depends on the solutions of the previous two steps.
- The recurrence relation is: $\text{ways}[i] = \text{ways}[i-1] + \text{ways}[i-2]$.
- Base cases: when there is 1 step, there is 1 way; when there are 2 steps, there are 2 ways.

- This is analogous to the Fibonacci sequence, shifting indices accordingly.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$ (can be optimized to $O(1)$ using constant space)

Solution

We use a dynamic programming approach to build a solution gradually. The idea is that to reach step i , you could have come from step $i-1$ (taking one step) or from step $i-2$ (taking two steps). Therefore, the total number of ways to reach step i is the sum of ways to reach step $i-1$ and $i-2$. We initialize $dp[1] = 1$ and $dp[2] = 2$, and iterate from 3 to n to fill $dp[]$. Each code solution below uses this approach with comments to explain the logic.

Dynamic Programming

#121 Best Time to Buy and Sell Stock

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Grind 169

Problem

Description

You are given an array `prices` where `prices[i]` is the price of a given stock on the i th day.

You want to maximize your profit by choosing a single day to buy one stock and choosing a different day in the future to sell that stock.

Return the maximum profit you can achieve from this transaction. If you cannot achieve any profit, return 0.

Example 1:

Input: `prices = [7,1,5,3,6,4]`

Output: 5

Explanation: Buy on day 2 (price = 1) and sell on day 5 (price = 6), profit = $6 - 1 = 5$.

Note that buying on day 2 and selling on day 1 is not allowed because you must buy before you sell.

Example 2:

Input: prices = [7,6,4,3,1]

Output: 0

Explanation: In this case, no transactions are done and the max profit = 0.

Constraints:

- $1 \leq \text{prices.length} \leq 105$
- $0 \leq \text{prices}[i] \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxProfit(self, prices: list[int]) -> int:
        sellOne = 0
        holdOne = -math.inf

        for price in prices:
            sellOne = max(sellOne, holdOne + price)
            holdOne = max(holdOne, -price)

        return sellOne
```

• LeetDoocs

class Solution:

```
def maxProfit(self, prices: List[int]) -> int:
    ans, mi = 0, inf
    for v in prices:
        ans = max(ans, v - mi)
        mi = min(mi, v)
    return ans
```

• SimplyLeet

```
def maxProfit(prices):
    # Initialize min_price as infinity and
    max_profit as 0
    min_price = float('inf')
    max_profit = 0
    # Iterate through each price in the list
    for price in prices:
        # Update the minimum price if current price
        is lower
        if price < min_price:
            min_price = price
        else:
```

```

        # Calculate current profit and update
max_profit if it is higher
        max_profit = max(max_profit, price -
min_price)
        return max_profit

```

```

# Example usage:
# result = maxProfit([7,1,5,3,6,4])
# print(result) # Output should be 5

```

• LC.ca

```

...
>>> import math
>>> a = math.inf
>>> a
inf

>>> x = float('-inf')
>>> y = float('inf')
>>> print(x < y) # Output: True
True

```

```

>>>
>>> z = -math.inf
>>> x==z
True

...

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        ans, mi = 0, inf
        for v in prices:

```

```

        mi = min(mi, v)
        ans = max(ans, v - mi)
        return ans

```

```
#####
```

```

class Solution(object):
    def maxProfit(self, prices):
        """
        :type prices: List[int]

```

```

:rtype: int
"""
if not prices:
    return 0
ans = 0
pre = prices[0]
for i in range(1, len(prices)):
    pre = min(pre, prices[i])
    ans = max(prices[i] - pre, ans)
return ans

```

◆ Quick Review

Key Insights

- Traverse the list of prices once while keeping track of the minimum price encountered so far.
- At each day, calculate the profit if you sold the stock on that day.
- Update the maximum profit when the current profit exceeds the previous maximum.
- The optimal solution makes a single pass through the array with constant space usage.

Space & Time

Time Complexity: $O(n)$ – requires one pass through the prices array.

Space Complexity: $O(1)$ – only a few extra variables are used.

Solution

We solve the problem using a simple linear scan. We maintain two key variables: one for the minimum price seen so far and another for the maximum profit calculated so far. As we iterate, we continuously update the minimum price if a lower price is found, and then compute the potential profit at the current price. If this profit exceeds the current maximum profit, we update the maximum profit. The final maximum profit is returned after iterating through all prices.

Bit Manipulation

Round 7 · Low · Easy · 6 questions

Round 7 · Low · Easy

p.2679

Bit Manipulation

#67 Add Binary

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · String · Bit Manipulation · Simulation

Lists: Grind 169

Problem

Description

Given two binary strings *a* and *b*, return their sum as a binary string.

Example 1:

Input: *a* = "11", *b* = "1"
Output: "100"

Example 2:

Input: *a* = "1010", *b* = "1011"
Output: "10101"

Constraints:

- $1 \leq a.length, b.length \leq 104$
- *a* and *b* consist only of '0' or '1' characters.
- Each string does not contain leading zeros except for the zero itself.

Community Solutions (Python)

Round 7 · Low · Easy

p.2680

• WalkCC

```
class Solution:
    def addBinary(self, a: str, b: str) -> str:
        ans = []
        carry = 0
        i = len(a) - 1
        j = len(b) - 1

        while i >= 0 or j >= 0 or carry:
            if i >= 0:
                carry += int(a[i])

            i -= 1
            if j >= 0:
                carry += int(b[j])
                j -= 1
            ans.append(str(carry % 2))
            carry //= 2

        return ''.join(reversed(ans))
```

• LeetCode

class Solution:

```
    def addBinary(self, a: str, b: str) -> str:
        ans = []
        i, j, carry = len(a) - 1, len(b) - 1, 0
        while i >= 0 or j >= 0 or carry:
            carry += (0 if i < 0 else int(a[i])) +
            (0 if j < 0 else int(b[j]))
            carry, v = divmod(carry, 2)
            ans.append(str(v))
            i, j = i - 1, j - 1
        return ''.join(ans[::-1])
```

• SimplyLeet

```
# Function to add two binary strings
def addBinary(a: str, b: str) -> str:
    i = len(a) - 1 # Pointer for string a
    j = len(b) - 1 # Pointer for string b
    carry = 0 # Initialize carry to 0
    result = [] # List to store binary digits

    # Loop until both strings are processed and no
    carry remains
    while i >= 0 or j >= 0 or carry:
        total_sum = carry # Start with the carry
        from previous step
```

```

        if i >= 0: # If there is a digit left in a
            total_sum += int(a[i]) # Convert
character to integer and add
            i -= 1 # Move to the next digit
        if j >= 0: # If there is a digit left in b
            total_sum += int(b[j]) # Convert
character to integer and add
            j -= 1 # Move to the next digit

        # Append the remainder of total_sum divided
by 2 (0 or 1)
        result.append(str(total_sum % 2))

```

```

        carry = total_sum // 2 # Update the carry

        # Reverse the result since we built it
backwards and join to form a string
        return ''.join(result[::-1])

# Sample test cases
print(addBinary("11", "1")) # Expected output:
"100"
print(addBinary("1010", "1011")) # Expected
output: "10101"

```

- [LC.ca](#)

```

class Solution:
    def addBinary(self, a: str, b: str) -> str:
        ans = []
        i, j, carry = len(a) - 1, len(b) - 1, 0
        while i >= 0 or j >= 0 or carry:
            carry += (0 if i < 0 else int(a[i])) +
(0 if j < 0 else int(b[j]))
            carry, v = divmod(carry, 2)
            ans.append(str(v))
            i, j = i - 1, j - 1
        return "".join(ans[::-1])

```

◆ Quick Review

Key Insights

- Simulate binary addition just like you add numbers digit by digit.
- Process the strings from right to left, handling the carry at each step.
- Use modulo and division operations to extract the current digit and update the carry.
- Append any remaining carry at the end.

Space & Time

Time Complexity: $O(\max(n, m))$, where n and m are the lengths of the two binary strings.

Space Complexity: $O(\max(n, m))$ for storing the resulting binary string.

Solution

The strategy is to iterate over the input strings from the end to the beginning. At each step, calculate the sum of corresponding bits and the existing carry. Use the modulo operation ($\% 2$) to determine the current digit and integer division ($// 2$) to update the carry. Once all digits have been processed, reverse the result to obtain the final binary string.

Round 7 · Low · Easy

p.2685

Bit Manipulation

#136 Single Number

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Bit Manipulation

Lists: Grind 169

Problem

Description

Given a non-empty array of integers `nums`, every element appears twice except for one. Find that single one.

You must implement a solution with a linear runtime complexity and use only constant extra space.

Example 1:

Input: `nums = [2,2,1]`

Output: 1

Example 2:

Input: `nums = [4,1,2,1,2]`

Output: 4

Example 3:

Input: `nums = [1]`

Output: 1

Constraints:

- $1 \leq \text{nums.length} \leq 3 \times 10^4$

Round 7 · Low · Easy

p.2686

- $-3 * 104 \leq \text{nums}[i] \leq 3 * 104$
- Each element in the array appears twice except for one element which appears only once.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def singleNumber(self, nums: list[int]) -> int:
        return functools.reduce(operator.xor, nums, 0)
```

• LeetDoocs

```
class Solution:
    def singleNumber(self, nums: List[int]) -> int:
        return reduce(xor, nums)
```

• SimplyLeet

```
# Function to find the single number in the list
def singleNumber(nums):
    # Initialize result to 0; any number XOR 0 is
    the number itself
    result = 0
    # Iterate through every number in the list
    for num in nums:
        # XOR the current result with the current
        number
        result ^= num # Using XOR to cancel
        duplicates
    # Return the unique number that did not cancel
    out
    return result
```

```
# Example usage
nums = [4, 1, 2, 1, 2]
print(singleNumber(nums)) # Output should be 4
```

• LC.ca

```

...
>>> from functools import reduce
>>> reduce(lambda x, y: x ^ y, [3, 5, 3])
5
...

from functools import reduce

class Solution:
    def singleNumber(self, nums: List[int]) ->
    int:

        return reduce(lambda x, y: x ^ y, nums)

#####

class Solution(object):
    def singleNumber(self, nums):
        """
        :type nums: List[int]
        :rtype: int
        """

        for i in range(1, len(nums)):
            nums[0] ^= nums[i]
        return nums[0]

```

◆ Quick Review

Key Insights

- Using bitwise XOR: XOR-ing a number with itself cancels out to 0.
- XOR is both commutative and associative, so the order of operations does not matter.
- A single traversal (linear time) and constant extra space is sufficient for solving the problem.

Space & Time

Time Complexity: $O(n)$ – We only traverse the array once.

Space Complexity: $O(1)$ – Only a single extra variable is used regardless of input size.

Solution

We can solve this problem by initializing a variable to 0 and then traversing the array while XOR-ing each element with the variable. Due to the properties of XOR, pairs of identical numbers will cancel each other out ($x \text{ XOR } x = 0$) and the unique number will be left as the final result. This approach uses constant extra space and performs in linear time.

Bit Manipulation

#190 Reverse Bits

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Divide and Conquer · Bit Manipulation

Lists: Grind 169

Problem

Description

Reverse bits of a given 32 bits signed integer.

Example 1:

Input: n = 43261596

Output: 964176192

Explanation:

Example 2:

Input: n = 2147483644

Output: 1073741822

Explanation:

Constraints:

- $0 \leq n \leq 2^{31} - 1$
- n is even.

Follow up: If this function is called many times, how would you optimize it?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def reverseBits(self, n: int) -> int:
        ans = 0

        for i in range(32):
            if n >> i & 1:
                ans |= 1 << 31 - i

        return ans
```

• LeetDoocs

```
class Solution:
    def reverseBits(self, n: int) -> int:
        ans = 0
        for i in range(32):
            ans |= (n & 1) << (31 - i)
            n >>= 1
        return ans
```

• SimplyLeet

```
# Function to reverse bits of a 32-bit binary string.
def reverse_bits(n: str) -> int:
    # Reverse the binary string using slicing.
    reversed_binary = n[::-1]
    # Convert the reversed binary string to an integer.
    result = int(reversed_binary, 2)
    return result

# Example usage:
binary_input = "00000010100101000001111010011100"

print(reverse_bits(binary_input)) # Expected output: 964176192
```

- LC.ca

```
class Solution:
    def reverseBits(self, n: int) -> int:
        res = 0
        for i in range(32):
            res |= (n & 1) << (31 - i)
            n >>= 1
        return res
```

◆ Quick Review

Key Insights

- The input is a binary string of fixed length (32 bits).
- Reversing the bits is equivalent to reversing the string.

- The reversed binary string can be converted back into its integer representation.
- Since the number of bits is constant (32), the solution can be achieved in constant time and constant space.
- For performance optimization when this function is called many times, caching may be applied for common intermediate results.

Space & Time

Time Complexity: $O(1)$ – We always process 32 bits regardless of the input.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We solve the problem by reversing the 32-character binary string representing the unsigned integer. Once reversed, the new binary string is converted into its integer value. An alternative solution involves using bitwise operations where you iterate over each bit, shift the result to the left, and add the current bit. This approach does the reversal one bit at a time. The chosen approach leverages string reversal which is straightforward to implement and understand.

Bit Manipulation

#191 Number of 1 Bits

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Divide and Conquer · Bit Manipulation

Lists: Grind 169

Problem

Description

Given a positive integer n , write a function that returns the number of set bits in its binary representation (also known as the Hamming weight).

Example 1:

Input: $n = 11$

Output: 3

Explanation:

The input binary string 1011 has a total of three set bits.

Example 2:

Input: $n = 128$

Output: 1

Explanation:

The input binary string 10000000 has a total of one set bit.

Example 3:

Input: $n = 2147483645$

Output: 30

Explanation:

The input binary string

1111111111111111111111111111101 has a total of thirty set bits.

Constraints:

- $1 \leq n \leq 2^{31} - 1$

Follow up: If this function is called many times, how would you optimize it?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def hammingWeight(self, n: int) -> int:
        ans = 0

        for i in range(32):
            if (n >> i) & 1:
                ans += 1

        return ans
```

```
class Solution:
    def hammingWeight(self, n: int) -> int:
        return n.bit_count()
```

• LeetDoocs


```
class Solution:
    def hammingWeight(self, n: int) -> int:
        ans = 0
        while n:
            n &= n - 1
            ans += 1
        return ans
```

• SimplyLeet

```
# Function to count the number of set bits in a
given integer n
def hammingWeight(n: int) -> int:
    count = 0 # Initialize count of set bits
    while n:
        # Remove the lowest set bit from n
        n &= n - 1
        count += 1 # Increment count for each set
        bit removed
    return count

# Example usage:
print(hammingWeight(11)) # Output: 3
```

• LC.ca

```
class Solution:
    def hammingWeight(self, n: int) -> int:
        ans = 0
        while n:
            n &= n - 1
            ans += 1
        return ans
```

♦ Quick Review

Key Insights

- The problem is essentially counting the number of 1s in the binary representation of a number.
- A common efficient approach uses bit manipulation: repeatedly clear the lowest set bit until n becomes zero.
- The expression $n \& (n - 1)$ clears the lowest set bit of n .
- This method only iterates as many times as there are set bits, making it efficient.

Space & Time

Time Complexity: $O(k)$, where k is the number of set bits, worst-case $O(\log n)$ (or $O(1)$ considering 32-bit integers).

Space Complexity: $O(1)$.

Solution

To solve the problem, we use the bit manipulation trick $n \& (n - 1)$ which removes the lowest set bit from n . The algorithm is as follows:

– Initialize a counter to 0. – While n is not zero, update n by $n = n \& (n - 1)$, and increment the counter. – The counter will hold the number of set bits by the end of the loop. This approach leverages efficient bit-level operations and works well even when the function is called frequently.

Round 7 · Low · Easy

p.2699

Bit Manipulation

#268 Missing Number

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Math · Binary Search · Bit Manipulation

Lists: Grind 169

Problem

Description

Given an array `nums` containing n distinct numbers in the range $[0, n]$, return the only number in the range that is missing from the array.

Example 1:

Input: `nums = [3,0,1]`

Output: 2

Explanation:

$n = 3$ since there are 3 numbers, so all numbers are in the range $[0,3]$. 2 is the missing number in the range since it does not appear in `nums`.

Example 2:

Input: `nums = [0,1]`

Output: 2

Explanation:

$n = 2$ since there are 2 numbers, so all numbers are in the range $[0,2]$. 2 is the missing number in the range

Round 7 · Low · Easy

p.2700

since it does not appear in nums.

Example 3:

Input: nums = [9,6,4,2,3,5,7,0,1]

Output: 8

Explanation:

n = 9 since there are 9 numbers, so all numbers are in the range [0,9]. 8 is the missing number in the range since it does not appear in nums.

Constraints:

- n == nums.length
- 1 <= n <= 104
- 0 <= nums[i] <= n
- All the numbers of nums are unique.

Follow up: Could you implement a solution using only O(1) extra space complexity and O(n) runtime complexity?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def missingNumber(self, nums: list[int]) -> int:
        ans = len(nums)

        for i, num in enumerate(nums):
            ans ^= i ^ num

        return ans
```

Round 7 · Low · Easy

p.2701

• LeetCode

```
class Solution:
    def missingNumber(self, nums: List[int]) ->
int:
        return reduce(xor, (i ^ v for i, v in
enumerate(nums, 1)))
```

```
class Solution:
    def missingNumber(self, nums: List[int]) ->
int:
        n = len(nums)
        return (1 + n) * n // 2 - sum(nums)
```

• SimplyLeet

```
# Function to find the missing number in the array
def missingNumber(nums):
    # Calculate n, where n is the length of the
nums array
    n = len(nums)
    # Compute the expected total sum using the
arithmetic sum formula for numbers 0 to n
    expected_sum = n * (n + 1) // 2
    # Compute the actual sum of the numbers in the
array
    actual_sum = sum(nums)
    # The missing number is the difference between
expected and actual sums
    return expected_sum - actual_sum
```

Round 7 · Low · Easy

p.2702

```
# Example Usage
print(missingNumber([3, 0, 1])) # Output: 2
```

- LC.ca

```
class Solution:
    def missingNumber(self, nums: List[int]) ->
    int:
        return reduce(xor, (i ^ v for i, v in
        enumerate(nums, 1)))
```

◆ Quick Review

Key Insights

- The range of numbers is continuous from 0 to n, but one number is missing.
- The sum of numbers from 0 to n can be computed using the arithmetic sum formula.
- Subtracting the sum of the given array from the expected sum gives the missing number.
- There is a follow-up challenge to solve the problem in O(n) time and O(1) extra space.

Space & Time

Time Complexity: O(n)

Space Complexity: O(1)

Solution

We compute the expected sum of all numbers from 0 to n using the formula $\text{sum} = n*(n+1)/2$. Then, we compute the actual sum of the numbers in the array. The missing number is obtained by subtracting the actual sum from the expected sum. This approach uses arithmetic operations (O(1) extra space) and loops over the array once (O(n) time).

Bit Manipulation

#338 Counting Bits

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Dynamic Programming · Bit Manipulation

Lists: Grind 169

Problem

Description

Given an integer n , return an array `ans` of length $n + 1$ such that for each i ($0 \leq i \leq n$), `ans[i]` is the number of 1's in the binary representation of i .

Example 1:

```
Input: n = 2
Output: [0,1,1]
Explanation:
0 --> 0
1 --> 1
2 --> 10
```

Example 2:

```
Input: n = 5
Output: [0,1,1,2,1,2]
Explanation:
0 --> 0
1 --> 1
2 --> 10
3 --> 11
4 --> 100
```

Constraints:

- $0 \leq n \leq 105$

Follow up:

- It is very easy to come up with a solution with a runtime of $O(n \log n)$. Can you do it in linear time $O(n)$ and possibly in a single pass?
- Can you do it without using any built-in function (i.e., like `__builtin_popcount` in C++)?

Community Solutions (Python)

- WalkCC

```
class Solution:
    def countBits(self, n: int) -> list[int]:
        # f(i) := i's number of 1s in bitmask
        # f(i) = f(i / 2) + i % 2
        ans = [0] * (n + 1)

        for i in range(1, n + 1):
            ans[i] = ans[i // 2] + (i & 1)

        return ans
```

- LeetCode

```
class Solution:
    def countBits(self, n: int) -> List[int]:
        return [i.bit_count() for i in range(n + 1)]
```

```
class Solution:
    def countBits(self, n: int) -> List[int]:
        ans = [0] * (n + 1)
        for i in range(1, n + 1):
            ans[i] = ans[i & (i - 1)] + 1
        return ans
```

- SimplyLeet

```
def countBits(n):
    # Create a list to store the number of 1's for
    # each number from 0 to n
    dp = [0] * (n + 1)
    # Loop over each number starting from 1, since
    # dp[0] is already 0
    for i in range(1, n + 1):
        # dp[i] is computed as the count for i//2
        # plus 1 if i is odd (i & 1)
        dp[i] = dp[i >> 1] + (i & 1)
    return dp
```

Example usage:

```
# print(countBits(5)) -> [0, 1, 1, 2, 1, 2]
```

- LC.ca

```
class Solution:
    def countBits(self, n: int) -> List[int]:
        ans = [0] * (n + 1)
        for i in range(1, n + 1):
            ans[i] = ans[i & (i - 1)] + 1
        return ans
```

- ◆ Quick Review

Key Insights

- We can use dynamic programming to build the solution by utilizing the count of bits for smaller numbers.

- The relationship $dp[i] = dp[i \gg 1] + (i \& 1)$ holds, where $i \gg 1$ gives the count for i divided by 2 (dropping the least significant bit) and $(i \& 1)$ checks if the least significant bit is 1.
- This approach computes the result for each number in constant time, resulting in a linear time solution.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

We use a dynamic programming approach where an array (or list) dp is initialized with size $n+1$. The key idea is to iterate from 1 to n and for each number, use the formula $dp[i] = dp[i \gg 1] + (i \& 1)$.

– The expression $(i \gg 1)$ effectively divides i by 2, giving the count of 1's for that half number. – The $(i \& 1)$ operation checks if i is odd, adding 1 if true. This method ensures each $dp[i]$ is computed in constant time, leading to an overall $O(n)$ time complexity.

Math

Round 7 · Low · Easy · 6 questions

Round 7 · Low · Easy

p.2709

Round 7 · Low · Easy

p.2710

Math

#9 Palindrome Number

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math

Lists: Grind 169

Problem

Description

Given an integer x , return true if x is a palindrome, and false otherwise.

Example 1:

Input: $x = 121$
 Output: true
 Explanation: 121 reads as 121 from left to right and from right to left.

Example 2:

Input: $x = -121$
 Output: false
 Explanation: From left to right, it reads -121. From right to left, it becomes 121-. Therefore it is not a palindrome.

Example 3:

Input: $x = 10$

Output: false

Explanation: Reads 01 from right to left. Therefore it is not a palindrome.

Constraints:

- $-231 \leq x \leq 231 - 1$

Follow up: Could you solve it without converting the integer to a string?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def isPalindrome(self, x: int) -> bool:
        if x < 0:
            return False

        rev = 0
        y = x

        while y:
            rev = rev * 10 + y % 10

            y //= 10

        return rev == x
```

• LeetDoocs


```
class Solution:
    def isPalindrome(self, x: int) -> bool:
        if x < 0 or (x and x % 10 == 0):
            return False
        y = 0
        while y < x:
            y = y * 10 + x % 10
            x //= 10
        return x in (y, y // 10)
```

• SimplyLeet

```
# Function to check if a number is a palindrome
without converting to a string.
def isPalindrome(x):
    # Negative numbers and numbers ending with 0
    (but not 0 itself) are not palindromic.
    if x < 0 or (x % 10 == 0 and x != 0):
        return False
    reversed_half = 0
    # Process: Reverse half of the number.
    while x > reversed_half:
        # Add the last digit of x to reversed_half.
        reversed_half = reversed_half * 10 + x % 10
```

```
# Remove the last digit from x.
x //= 10
# For numbers with odd digits, reversed_half //
10 removes the middle digit.
return x == reversed_half or x == reversed_half
// 10

# Example usage:
print(isPalindrome(121)) # Output: True
print(isPalindrome(-121)) # Output: False
print(isPalindrome(10)) # Output: False
```

• LC.ca

```
class Solution:
    def isPalindrome(self, x: int) -> bool:
        if x < 0 or (x and x % 10 == 0):
            return False
        y = 0
        while y < x:
            y = y * 10 + x % 10
            x //= 10
        return x in (y, y // 10)
```

◆ Quick Review

Key Insights

- Negative numbers are not palindromic.
- A number ending in 0 (except 0 itself) cannot be a palindrome.

- Instead of reversing the entire number (risking overflow), reverse only half of the number and compare with the other half.
- For numbers with an odd number of digits, the middle digit can be disregarded.

Space & Time

Time Complexity: $O(\log_{10}(n))$ – Each iteration reduces the number by a factor of 10.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

The solution checks for obvious non-palindrome cases first (negative numbers and numbers ending with 0 that are not 0). Then, it reverses half of the digits of the number while the original number is larger than the reversed half. Finally, it compares the remaining part of the original number with the reversed half. In the case of an odd number of digits, the middle digit is removed by dividing the reversed half by 10 before the final comparison.

Math

#13 Roman to Integer

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Math · String

Lists: Grind 169

Problem

Description

Roman numerals are represented by seven different symbols: I, V, X, L, C, D and M.

Symbol Value

I	1
V	5
X	10
L	50
C	100
D	500
M	1000

For example, 2 is written as II in Roman numeral, just two ones added together. 12 is written as XII, which is simply X + II. The number 27 is written as XXVII, which is XX + V + II.

Roman numerals are usually written largest to smallest from left to right. However, the numeral for four is not IIII. Instead, the number four is written as IV. Because

the one is before the five we subtract it making four. The same principle applies to the number nine, which is written as IX. There are six instances where subtraction is used:

- I can be placed before V (5) and X (10) to make 4 and 9.
- X can be placed before L (50) and C (100) to make 40 and 90.
- C can be placed before D (500) and M (1000) to make 400 and 900.

Given a roman numeral, convert it to an integer.

Example 1:

Input: s = "III"
Output: 3
Explanation: III = 3.

Example 2:

Input: s = "LVIII"
Output: 58
Explanation: L = 50, V= 5, III = 3.

Example 3:

Input: s = "MCMXCIV"
Output: 1994
Explanation: M = 1000, CM = 900, XC = 90 and IV = 4.

Constraints:

- $1 \leq s.length \leq 15$
- s contains only the characters ('I', 'V', 'X', 'L', 'C', 'D', 'M').

- It is guaranteed that s is a valid roman numeral in the range [1, 3999].

Community Solutions (Python)

• WalkCC

```
class Solution:
    def romanToInt(self, s: str) -> int:
        ans = 0
        roman = {'I': 1, 'V': 5, 'X': 10, 'L': 50,
                  'C': 100, 'D': 500, 'M': 1000}

        for a, b in zip(s, s[1:]):
            if roman[a] < roman[b]:
                ans -= roman[a]
            else:
                ans += roman[a]

        return ans + roman[s[-1]]
```

• LeetDoocs

```
class Solution:
    def romanToInt(self, s: str) -> int:
        d = {'I': 1, 'V': 5, 'X': 10, 'L': 50, 'C': 100, 'D': 500, 'M': 1000}
        return sum((-1 if d[a] < d[b] else 1) * d[a] for a, b in pairwise(s)) + d[s[-1]]
```

• SimplyLeet

```
# Define a function to convert Roman numeral to
integer
def romanToInt(s):
    # Map each Roman numeral to its integer value
    roman_values = {'I': 1, 'V': 5, 'X': 10, 'L':
50, 'C': 100, 'D': 500, 'M': 1000}
    total = 0 # Initialize the total integer value

    # Loop over the string using index
    for i in range(len(s)):
        # If this is not the last character and the
current value is less than the next value,
        # subtract the current value; otherwise,
add the current value.

        if i + 1 < len(s) and roman_values[s[i]] <
roman_values[s[i + 1]]:
            total -= roman_values[s[i]]
        else:
            total += roman_values[s[i]]
    return total

# Example usage:
print(romanToInt("MCMXCIV")) # Output: 1994
```

- LC.ca

```
class Solution:
    def romanToInt(self, s: str) -> int:
        d = {'I': 1, 'V': 5, 'X': 10, 'L': 50, 'C':
100, 'D': 500, 'M': 1000}
        return sum((-1 if d[a] < d[b] else 1) *
d[a] for a, b in pairwise(s)) + d[s[-1]]
```

◆ Quick Review

Key Insights

- Map each Roman numeral character to its integer value using a hash table or dictionary.
- Loop through the string and check if the current numeral is less than the next numeral; if so, subtract its value.
- Otherwise, add its value to the total.
- This approach takes advantage of the subtractive notation used in Roman numerals.

Space & Time

Time Complexity: $O(n)$, where n is the length of the input string.

Space Complexity: $O(1)$, since the mapping and additional variables use constant space.

Solution

We use a dictionary (or hash map) to store the values of each Roman numeral. As we iterate over the string, we compare the current numeral's value with the next numeral's value. When a numeral of smaller value appears before a numeral of larger value, we subtract the current

numeral's value; otherwise, we add it. This technique efficiently handles both the standard and subtractive cases in one pass through the string.

Round 7 · Low · Easy

p.2721

Math

#504 Base 7

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math

Lists: Denny Zhang

Problem

Description

Given an integer num, return a string of its base 7 representation.

Example 1:

Input: num = 100
Output: "202"

Example 2:

Input: num = -7
Output: "-10"

Constraints:

- $-107 \leq \text{num} \leq 107$

Community Solutions (Python)

- LeetDoocs

Round 7 · Low · Easy

p.2722

```
class Solution:
    def convertToBase7(self, num: int) -> str:
        if num == 0:
            return '0'
        if num < 0:
            return '-' + self.convertToBase7(-num)
        ans = []
        while num:
            ans.append(str(num % 7))
            num //= 7

        return ''.join(ans[::-1])
```

• SimplyLeet

```
# Python solution for converting an integer to its
base 7 string representation.
def convertToBase7(num):
    # Special case: If num is 0, return "0"
    if num == 0:
        return "0"

    # Flag to check if the number is negative
    negative = num < 0
    # Work with the absolute value of the number
    num = abs(num)
```

```
# List to store base 7 digits
digits = []

# Process the number until it becomes 0
while num > 0:
    # Get the remainder (digit in base 7)
    remainder = num % 7
    # Append the digit as string to the list
    digits.append(str(remainder))
```

```
# Update num by integer division by 7
num //= 7
```

```
# Reverse the list to get digits in correct
order
base7 = ''.join(reversed(digits))
# If original number was negative, attach the
minus sign
return "-" + base7 if negative else base7
```

```
# Example usage:
print(convertToBase7(100)) # Expected output:
"202"
```

```
print(convertToBase7(-7)) # Expected output: "-10"
```

• LC.ca

```

class Solution:
    def convertToBase7(self, num: int) -> str:
        if num == 0:
            return '0'
        if num < 0:
            return '-' + self.convertToBase7(-num)
        ans = []
        while num:
            ans.append(str(num % 7))
            num //= 7

        return ''.join(ans[::-1])

```

◆ Quick Review

Key Insights

- Handle the edge case where num is 0.
- Use division and modulo operations to extract base 7 digits.
- Manage negative numbers by converting to positive and reattaching the minus sign.
- Build the base 7 digits in reverse order during iteration.

Space & Time

Time Complexity: $O(\log_7(n))$ where n is the absolute value of the number.

Space Complexity: $O(\log_7(n))$ to store the digits of the number.

Solution

Round 7 · Low · Easy

p.2725

The algorithm follows these steps:

– Check if num equals 0 and return "0". – Determine if the number is negative. If so, work with its absolute value but remember the negative flag. – Iteratively compute the remainder when divided by 7 (using modulo) and update the number by integer division by 7. – Append each remainder (converted to character) to a list representing the digits. – The digits are computed in reverse order; reverse the list for the correct order. – If the original number was negative, append the "-" sign to the front. – Join the digits to form the resulting base 7 string.

This approach uses basic arithmetic (division and modulo) and a list (or equivalent structure) to accumulate the computed digits.

Round 7 · Low · Easy

p.2726

Math

#1056 Confusing Number

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math

Lists: Premium 98

Problem

Description

A confusing number is a number that when rotated 180 degrees becomes a different number with each digit valid.

We can rotate digits of a number by 180 degrees to form new digits.

- When 0, 1, 6, 8, and 9 are rotated 180 degrees, they become 0, 1, 9, 8, and 6 respectively.
- When 2, 3, 4, 5, and 7 are rotated 180 degrees, they become invalid.

Note that after rotating a number, we can ignore leading zeros.

- For example, after rotating 8000, we have 0008 which is considered as just 8.

Given an integer n , return true if it is a confusing number, or false otherwise.

Example 1:

A large black digit '6' is on the left. An arrow points to the right with the word 'rotate' written above it. On the right is a large black digit '9'.

Input: $n = 6$

Output: true

Explanation: We get 9 after rotating 6, 9 is a valid number, and $9 \neq 6$.

Example 2:

A large black number '89' is on the left. An arrow points to the right with the word 'rotate' written above it. On the right is a large black number '68'.

Input: $n = 89$

Output: true

Explanation: We get 68 after rotating 89, 68 is a valid number and $68 \neq 89$.


Round 7 · Low · Easy

p.2727

Round 7 · Low · Easy

p.2728

Example 3:



Input: $n = 11$

Output: false

Explanation: We get 11 after rotating 11, 11 is a valid number but the value remains the same, thus 11 is not a confusing number

Constraints:

- $0 \leq n \leq 109$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def confusingNumber(self, n: int) -> bool:
        s = str(n)
        rotated = {'0': '0', '1': '1', '6': '9', '8': '8', '9': '6'}
        rotatedNum = []

        for c in s[::-1]:
            if c not in rotated:
                return False
            rotatedNum.append(rotated[c])

        return ''.join(rotatedNum) != s
```

- LeetCode

```
class Solution:
    def confusingNumber(self, n: int) -> bool:
        x, y = n, 0
        d = [0, 1, -1, -1, -1, -1, 9, -1, 8, 6]
        while x:
            x, v = divmod(x, 10)
            if d[v] < 0:
                return False
            y = y * 10 + d[v]
        return y != n
```

- SimplyLeet

```

# Define the mapping for valid rotations
def is_confusing_number(n: int) -> bool:
    rotation_map = {'0': '0', '1': '1', '6': '9',
                    '8': '8', '9': '6'}

    # Convert the number to a string to process
    # each digit
    original_str = str(n)
    rotated_str = ""

    # Process each digit in reverse order to
    # simulate 180 degree rotation
    for digit in reversed(original_str):

        # If digit is not valid for rotation,
        # return false immediately
        if digit not in rotation_map:
            return False

        # Append the rotated digit to the rotated
        # string
        rotated_str += rotation_map[digit]

    # Convert rotated string to an integer (leading
    # zeros are ignored)
    rotated_num = int(rotated_str)

    # The number is confusing if the rotated number
    # is different from the original number

```

```

return rotated_num != n

```

```

# Example usage:
print(is_confusing_number(6)) # Expected output:
True
print(is_confusing_number(89)) # Expected output:
True
print(is_confusing_number(11)) # Expected output:
False

```

• LC.ca

```

class Solution:
    def confusingNumber(self, n: int) -> bool:
        x, y = n, 0
        d = [0, 1, -1, -1, -1, -1, 9, -1, 8, 6]
        while x:
            x, v = divmod(x, 10)
            if d[v] < 0:
                return False
            y = y * 10 + d[v]
        return y != n

```

◆ Quick Review

Key Insights

- Use a mapping dictionary for valid rotations: {0:0, 1:1, 6:9, 8:8, 9:6}.
- Process the digits of n from right to left to form the rotated number.

- If any digit is not present in the mapping, the number is instantly not confusing.
- Compare the rotated number with the original number to check if they are different.

Space & Time

Time Complexity: $O(d)$, where d is the number of digits in n (approximately $O(\log n)$).

Space Complexity: $O(d)$, needed to store the rotated number string.

Solution

The approach is to convert the number to a string and iterate through it in reverse. For each digit, check if it is valid using the predefined mapping. If it is, append its rotated counterpart to a new string. After processing, convert the rotated string to an integer (ignoring any leading zeros). Finally, compare the rotated number to the original number. If they are different, the number is confusing.

Math

#1228 Missing Number in Arithmetic Progression

EAS

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Math

Lists: Premium 98

Problem

Description

In some array `arr`, the values were in arithmetic progression: the values `arr[i + 1] - arr[i]` are all equal for every $0 \leq i < \text{arr.length} - 1$.

A value from `arr` was removed that was not the first or last value in the array.

Given `arr`, return the removed value.

Example 1:

Input: `arr = [5,7,11,13]`

Output: 9

Explanation: The previous array was `[5,7,9,11,13]`.

Example 2:

Input: `arr = [15,13,12]`

Output: 14

Explanation: The previous array was `[15,14,13,12]`.

Constraints:

- $3 \leq \text{arr.length} \leq 1000$
- $0 \leq \text{arr}[i] \leq 105$
- The given array is guaranteed to be a valid array.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def missingNumber(self, arr: list[int]) -> int:
        n = len(arr)
        delta = (arr[-1] - arr[0]) // n
        l = 0
        r = n - 1

        while l < r:
            m = (l + r) // 2
            if arr[m] == arr[0] + m * delta:

                l = m + 1
            else:
                r = m

        return arr[0] + l * delta
```

• LeetCode

```
class Solution:
    def missingNumber(self, arr: List[int]) ->
int:
        return (arr[0] + arr[-1]) * (len(arr) + 1)
// 2 - sum(arr)
```

```
class Solution:
    def missingNumber(self, arr: List[int]) ->
int:
        n = len(arr)
        d = (arr[-1] - arr[0]) // n
        for i in range(1, n):
            if arr[i] != arr[i - 1] + d:
                return arr[i - 1] + d
        return arr[0]
```

• SimplyLeet

```
# Function to find the missing number in the
arithmetic progression
def missingNumber(arr):
    # Calculate the expected difference using the
    first and last elements
    n = len(arr)
    diff = (arr[-1] - arr[0]) // n # integer
division because values are integers
    # Iterate over the array to find the place
    where the difference deviates from diff
    for i in range(len(arr) - 1):
        # If the difference between consecutive
        numbers is not equal to the expected diff
        if arr[i+1] - arr[i] != diff:
            # Return the missing number which
            should be current number + diff
```

```

        return arr[i] + diff
    # Return -1 if no missing number is found
    (should not happen given the problem constraints)
    return -1

# Example usage:
# print(missingNumber([5,7,11,13])) # Output
# should be 9

```

- LC.ca

```

class Solution:
    def missingNumber(self, arr: List[int]) ->
    int:
        return (arr[0] + arr[-1]) * (len(arr) + 1)
        // 2 - sum(arr)

```

◆ Quick Review

Key Insights

- The complete arithmetic progression has equal differences between consecutive numbers.
- The missing element can be found by calculating the expected common difference using the first and last elements.
- Since exactly one number is missing, iterate through the provided array to detect where the difference deviates from the expected common difference.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We first compute the expected common difference (diff) for the original arithmetic progression. Since the actual progression had one more element than the current array, the formula used is: $\text{diff} = (\text{last element} - \text{first element}) / (n)$ where n is the length of the current array. Next, iterate over the array and check the difference between each pair of consecutive elements. When the actual difference deviates from diff, it indicates that the missing number is the previous element plus diff. This approach uses a simple loop and constant extra space.

Math

#1427 Perform String Shifts

EAS

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Math · String

Lists: Premium 98

Problem

Description

You are given a string *s* containing lowercase English letters, and a matrix *shift*, where *shift[i] = [directioni, amounti]*:

- *directioni* can be 0 (for left shift) or 1 (for right shift).
- *amounti* is the amount by which string *s* is to be shifted.
- A left shift by 1 means remove the first character of *s* and append it to the end.
- Similarly, a right shift by 1 means remove the last character of *s* and add it to the beginning.

Return the final string after all operations.

Example 1:

Input: *s* = "abc", *shift* = [[0,1],[1,2]]
 Output: "cab"
 Explanation:
 [0,1] means shift to left by 1. "abc" -> "bca"
 [1,2] means shift to right by 2. "bca" -> "cab"

Round 7 · Low · Easy

p.2739

Example 2:

Input: *s* = "abcdefg", *shift* = [[1,1],[1,1],[0,2],[1,3]]
 Output: "efgabcd"
 Explanation:
 [1,1] means shift to right by 1. "abcdefg" -> "gabcdef"
 [1,1] means shift to right by 1. "gabcdef" -> "fgabcde"
 [0,2] means shift to left by 2. "fgabcde" -> "abcdefg"
 [1,3] means shift to right by 3. "abcdefg" -> "efgabcd"

Constraints:

- $1 \leq s.length \leq 100$
- *s* only contains lower case English letters.
- $1 \leq shift.length \leq 100$
- $shift[i].length == 2$
- *directioni* is either 0 or 1.
- $0 \leq amounti \leq 100$

Community Solutions (Python)

- WalkCC

Round 7 · Low · Easy

p.2740

```
class Solution:
    def stringShift(self, s: str, shift:
list[list[int]]) -> str:
    move = 0

    for direction, amount in shift:
        if direction == 0:
            move -= amount
        else:
            move += amount

    move %= len(s)
    return s[-move:] + s[:-move]
```

- LeetCode

```
class Solution:
    def stringShift(self, s: str, shift:
List[List[int]]) -> str:
    x = sum((b if a else -b) for a, b in shift)
    x %= len(s)
    return s[-x:] + s[:-x]
```

- SimplyLeet

Python solution

```
def stringShift(s, shift):
    # Initialize the net shift counter
    net_shift = 0
    for direction, amount in shift:
        if direction == 0:
            net_shift -= amount # Left shift
            reduces index
        else:
            net_shift += amount # Right shift
            increases index
```

```
# Normalize shift to be within bounds of the
string length
net_shift %= len(s)

# A right shift can be simulated by taking the
last net_shift characters
# and concatenating them with the beginning of
the string.
return s[-net_shift:] + s[:-net_shift]
```

```
# Example usage
if __name__ == "__main__":
```

```
s = "abcdefg"
shift_operations = [[1,1],[1,1],[0,2],[1,3]]
print(stringShift(s, shift_operations)) #
Output: "efgabcd"
```

- LC.ca

```
class Solution:
    def stringShift(self, s: str, shift:
List[List[int]]) -> str:
        x = sum((b if a else -b) for a, b in shift)
        x %= len(s)
        return s[-x:] + s[:-x]
```

◆ Quick Review

Key Insights

- Consolidate all shifts by calculating a net shift: subtract for left shifts and add for right shifts.
- Use modulo arithmetic with the string length to avoid unnecessary full rotations.
- Slice the string based on the effective shift to construct the final string.

Space & Time

Time Complexity: $O(n + m)$, where n is the length of the string and m is the number of shift operations.

Space Complexity: $O(n)$ for the output string.

Solution

The solution involves computing the net shift value by iterating through the shift operations. Left shifts subtract from the net value, and right shifts add to the net value. After obtaining the cumulative shift, use modulo with the string length to normalize the shift, ensuring it lies within bounds. The final string is constructed by slicing the string into two parts and concatenating them appropriately. This approach leverages string slicing as the primary operation for shifting.

Round 8 · Low · Medium

Round 8

Low Priority · Medium

32 questions · 4 patterns

**Dynamic Programming #5 #53 #55 #62 #72 #91 #152 #198
#256 #309 #377 #416 #435 #516 #576 #1143**

Bit Manipulation #78 #90 #287 #1506

Math #7 #50 #380 #1017 #3160

JavaScript #2627 #2628 #2630 #2633 #2636 #2675 #2700

Round 7 · Low · Easy

p.2745

Dynamic Programming

Round 8 · Low · Medium · 16 questions

Round 8 · Low · Medium

p.2746

Dynamic Programming

#5 Longest Palindromic Substring

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Grind 169

Problem

Description

Given a string *s*, return the longest palindromic substring in *s*.

Example 1:

Input: *s* = "babad"
 Output: "bab"
 Explanation: "aba" is also a valid answer.

Example 2:

Input: *s* = "cbabd"
 Output: "bb"

Constraints:

- $1 \leq s.length \leq 1000$
- *s* consist of only digits and English letters.

Community Solutions (Python)

- WalkCC

Round 8 · Low · Medium

p.2747

```
class Solution:
    def longestPalindrome(self, s: str) -> str:
        if not s:
            return ''

        # (start, end) indices of the longest
        # palindrome in s
        indices = [0, 0]

        def extend(s: str, i: int, j: int) ->
            tuple[int, int]:
                """
                Returns the (start, end) indices of the
                longest palindrome extended from
                the substring s[i..j].
                """
                while i >= 0 and j < len(s):
                    if s[i] != s[j]:
                        break
                    i -= 1
                    j += 1
                return i + 1, j - 1
```

Round 8 · Low · Medium

p.2748

```

for i in range(len(s)):
    l1, r1 = extend(s, i, i)
    if r1 - l1 > indices[1] - indices[0]:
        indices = l1, r1
    if i + 1 < len(s) and s[i] == s[i + 1]:
        l2, r2 = extend(s, i, i + 1)
        if r2 - l2 > indices[1] - indices[0]:
            indices = l2, r2

return s[indices[0]:indices[1] + 1]

```

```

class Solution:
    def longestPalindrome(self, s: str) -> str:
        t = '#'.join('@' + s + '$')
        p = self._manacher(t)
        maxPalindromeLength, bestCenter = max((extend,
i)
                                for i,
extend in enumerate(p))
        l = (bestCenter - maxPalindromeLength) // 2
        r = (bestCenter + maxPalindromeLength) // 2
        return s[l:r]

```

```

def _manacher(self, t: str) -> list[int]:
    """
    Returns an array `p` s.t. `p[i]` is the length
    of the longest palindrome
    centered at `t[i]`, where `t` is a string with
    delimiters and sentinels.
    """
    p = [0] * len(t)
    center = 0
    for i in range(1, len(t) - 1):
        rightBoundary = center + p[center]
        mirrorIndex = center - (i - center)

        if rightBoundary > i:
            p[i] = min(rightBoundary - i,
p[mirrorIndex])
            # Try to expand the palindrome centered at i.
            while t[i + 1 + p[i]] == t[i - 1 - p[i]]:
                p[i] += 1
            # If a palindrome centered at i expands past
            `rightBoundary`, adjust
            # the center based on the expanded
            palindrome.
            if i + p[i] > rightBoundary:
                center = i
    return p

```

• LeetCode

```

class Solution:
    def longestPalindrome(self, s: str) -> str:
        n = len(s)
        f = [[True] * n for _ in range(n)]
        k, mx = 0, 1
        for i in range(n - 2, -1, -1):
            for j in range(i + 1, n):
                f[i][j] = False
                if s[i] == s[j]:
                    f[i][j] = f[i + 1][j - 1]

                    if f[i][j] and mx < j - i + 1:
                        k, mx = i, j - i + 1
        return s[k : k + mx]

```

```

class Solution:
    def longestPalindrome(self, s: str) -> str:
        def f(l, r):
            while l >= 0 and r < n and s[l] == s[r]:
                l, r = l - 1, r + 1
            return r - l - 1

        n = len(s)
        start, mx = 0, 1
        for i in range(n):

```

```

        a = f(i, i)
        b = f(i, i + 1)
        t = max(a, b)
        if mx < t:
            mx = t
            start = i - ((t - 1) >> 1)
        return s[start : start + mx]

```

• SimplyLeet

```

# Function to expand from the center and return the
# palindrome substring
def expand_from_center(s, left, right):
    # Expand while the characters on both sides are
    # the same
    while left >= 0 and right < len(s) and s[left]
    == s[right]:
        left -= 1
        right += 1
    # Return the palindromic substring found after
    # expansion
    return s[left + 1:right]

def longest_palindromic_substring(s):

```

```

if not s:
    return ""
longest = ""
# Iterate over each index in string as
potential center
for i in range(len(s)):
    # Check for odd-length palindrome
    odd_palindrome = expand_from_center(s, i,
i)
    if len(odd_palindrome) > len(longest):
        longest = odd_palindrome
    # Check for even-length palindrome
    even_palindrome = expand_from_center(s, i,
i + 1)
    if len(even_palindrome) > len(longest):
        longest = even_palindrome
return longest

# Example usage
print(longest_palindromic_substring("babad"))

```

• LC.ca

Round 8 · Low · Medium

p.2753

```

class Solution:
    def longestPalindrome(self, s: str) -> str:
        mlen = 0
        start = end = 0
        n = len(s)
        dp = [ [False] * n for i in range(n) ]

        for j in range(n):
            for i in range(j + 1):
                dp[i][j] = (i == j) or (s[i] ==
s[j] and j - i == 1) or (s[i] == s[j] and dp[i +
1][j - 1])

                if dp[i][j] is True and j - i + 1 >
mlen:
                    mlen = j - i + 1
                    start = i
                    end = j

        return s[start: end + 1]

#####

class Solution:

```

Round 8 · Low · Medium

p.2754

```
def longestPalindrome(self, s: str) -> str:
    n = len(s)
    f = [[True] * n for _ in range(n)]
    k, mx = 0, 1
    for i in range(n - 2, -1, -1):
        for j in range(i + 1, n):
            f[i][j] = False
            if s[i] == s[j]:
                f[i][j] = f[i + 1][j - 1]
                if f[i][j] and mx < j - i + 1:
                    k, mx = i, j - i + 1
    return s[k : k + mx]
```

◆ Quick Review

Key Insights

- Every palindrome is centered around a character (for odd-length) or a pair of characters (for even-length).
- Expand from each center until the substring is no longer a palindrome.
- Update the longest palindrome when a longer one is found during the expansion.
- The solution employs a two-pointer technique without extra storage for dynamic programming tables.

Space & Time

Time Complexity: $O(n^2)$ in the worst case as we expand around each center.

Space Complexity: $O(1)$, excluding the space required for the output substring.

Solution

The approach is based on expanding around potential centers of palindromes. For each character in the string, consider it as the center for an odd-length palindrome and the gap between it and the next character for an even-length palindrome. Using two pointers, expand outward until the characters differ. Track the longest palindrome found during these expansions. This method is efficient given the input constraints, and no extra data structures are needed besides variables to store indices or substrings.

Dynamic Programming

#53 Maximum Subarray

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Divide and Conquer

Lists: Grind 169

Problem

Description

Given an integer array `nums`, find the subarray with the largest sum, and return its sum.

Example 1:

Input: `nums = [-2,1,-3,4,-1,2,1,-5,4]`

Output: 6

Explanation: The subarray `[4,-1,2,1]` has the largest sum 6.

Example 2:

Input: `nums = [1]`

Output: 1

Explanation: The subarray `[1]` has the largest sum 1.

Example 3:

Input: `nums = [5,4,-1,7,8]`

Output: 23

Explanation: The subarray `[5,4,-1,7,8]` has the largest sum 23.

Round 8 · Low · Medium

p.2757

Constraints:

- $1 \leq \text{nums.length} \leq 105$
- $-104 \leq \text{nums}[i] \leq 104$

Follow up: If you have figured out the $O(n)$ solution, try coding another solution using the divide and conquer approach, which is more subtle.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxSubArray(self, nums: list[int]) -> int:
        # dp[i] := the maximum sum subarray ending in i
        dp = [0] * len(nums)

        dp[0] = nums[0]
        for i in range(1, len(nums)):
            dp[i] = max(nums[i], dp[i - 1] + nums[i])

        return max(dp)
```

Round 8 · Low · Medium

p.2758

```
class Solution:
    def maxSubArray(self, nums: list[int]) -> int:
        ans = -math.inf
        summ = 0

        for num in nums:
            summ = max(num, summ + num)
            ans = max(ans, summ)

        return ans
```

```
from dataclasses import dataclass
```

```
@dataclass(frozen=True)
class T:
    summ: int
    # the sum of the subarray starting from the first
    # number
    maxSubarraySumLeft: int
    # the sum of the subarray ending in the last
    # number
    maxSubarraySumRight: int
```

```
maxSubarraySum: int
```

```
class Solution:
    def maxSubArray(self, nums: list[int]) -> int:
        def divideAndConquer(l: int, r: int) -> T:
            if l == r:
                return T(nums[l], nums[l], nums[l],
nums[l])
            m = (l + r) // 2
            left = divideAndConquer(l, m)

            right = divideAndConquer(m + 1, r)
            maxSubarraySumLeft =
max(left.maxSubarraySumLeft,
                                left.summ +
right.maxSubarraySumLeft)
            maxSubarraySumRight = max(
                left.maxSubarraySumRight + right.summ,
right.maxSubarraySumRight)
            maxSubarraySum =
max(left.maxSubarraySumRight +
right.maxSubarraySumLeft,
                                left.maxSubarraySum,
right.maxSubarraySum)
            summ = left.summ + right.summ
            return T(summ, maxSubarraySumLeft,
maxSubarraySumRight, maxSubarraySum)
```



```
return divideAndConquer(0, len(nums) -
1).maxSubarraySum
```

- LeetCode

```
class Solution:
    def maxSubArray(self, nums: List[int]) -> int:
        ans = f = nums[0]
        for x in nums[1:]:
            f = max(f, 0) + x
            ans = max(ans, f)
        return ans
```

```
class Solution:
    def maxSubArray(self, nums: List[int]) -> int:
        def crossMaxSub(nums, left, mid, right):
            lsum = rsum = 0
            lmx = rmx = -inf
            for i in range(mid, left - 1, -1):
                lsum += nums[i]
                lmx = max(lmx, lsum)
            for i in range(mid + 1, right + 1):
                rsum += nums[i]
```

```
rmx = max(rmx, rsum)
return lmx + rmx
```

```
def maxSub(nums, left, right):
    if left == right:
        return nums[left]
    mid = (left + right) >> 1
    lsum = maxSub(nums, left, mid)
    rsum = maxSub(nums, mid + 1, right)
    csum = crossMaxSub(nums, left, mid,
right)
```

```
return max(lsum, rsum, csum)
```

```
left, right = 0, len(nums) - 1
return maxSub(nums, left, right)
```

- SimplyLeet

```
# Python solution using Kadane's Algorithm
def maxSubArray(nums):
    # Initialize currentSum and maxSum to the first
    # element
    current_sum = max_sum = nums[0]
    # Traverse through the array starting from the
    # second element
    for num in nums[1:]:
        # If current_sum + num is less than num,
        # restart at num
        current_sum = max(num, current_sum + num)
        # Update max_sum if current_sum is greater
        max_sum = max(max_sum, current_sum)

    return max_sum

# Example usage
if __name__ == "__main__":
    sample = [-2,1,-3,4,-1,2,1,-5,4]
    print(maxSubArray(sample)) # Output: 6
```

• LC.ca

```
class Solution:
    def maxSubArray(self, nums: List[int]) -> int:
        res = cur_sum = nums[0]
        for num in nums[1:]:
            cur_sum = num + max(cur_sum, 0)
            res = max(res, cur_sum)
        return res
```

◆ Quick Review

Key Insights

- Use Kadane's Algorithm to solve the problem in $O(n)$ time.
- Keep track of a running sum (currentSum) and update it with the maximum between the current element and the sum including the current element.
- The maximum sum encountered during the iteration (maxSum) is the answer.
- A divide and conquer approach also exists, splitting the array and merging solutions, but Kadane's method is optimal for this problem.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution uses Kadane's Algorithm. Iterate over the array while maintaining two variables: currentSum, which tracks the maximum subarray sum ending at the current position, and maxSum, which stores the best sum seen so far. For each element, update currentSum to be either the current element or the sum of currentSum and the current element, whichever is higher. Update maxSum accordingly. This approach efficiently computes the maximum contiguous subarray sum in a single pass.

Dynamic Programming

#55 Jump Game

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Greedy

Lists: Grind 169

Problem

Description

You are given an integer array `nums`. You are initially positioned at the array's first index, and each element in the array represents your maximum jump length at that position.

Return `true` if you can reach the last index, or `false` otherwise.

Example 1:

Input: `nums = [2,3,1,1,4]`Output: `true`

Explanation: Jump 1 step from index 0 to 1, then 3 steps to the last index.

Example 2:

Input: `nums = [3,2,1,0,4]`Output: `false`

Explanation: You will always arrive at index 3 no matter what. Its maximum jump length is 0, which makes it impossible to reach the last index.

Constraints:

- `1 <= nums.length <= 104`
- `0 <= nums[i] <= 105`

Community Solutions (Python)

• WalkCC

```
class Solution:
    def canJump(self, nums: list[int]) -> bool:
        i = 0
        reach = 0

        while i < len(nums) and i <= reach:
            reach = max(reach, i + nums[i])
            i += 1

        return i == len(nums)
```

• LeetDoocs

```
class Solution:
    def canJump(self, nums: List[int]) -> bool:
        mx = 0
        for i, x in enumerate(nums):
            if mx < i:
                return False
            mx = max(mx, i + x)
        return True
```

• SimplyLeet

```
# Function to determine if you can reach the last
index of the array
def canJump(nums):
    # Initialize maxReach which is the max index
    reachable at start (index 0)
    maxReach = 0
    # Iterate through each element and its index
    for i, jump in enumerate(nums):
        # If current index is greater than maximum
        reachable index, return False
        if i > maxReach:
            return False
        # Update maxReach with the maximum of
        current maxReach and i + jump length at index i
```

```
        maxReach = max(maxReach, i + jump)
        # If current maxReach is already at or
        beyond the last index, return True
        if maxReach >= len(nums) - 1:
            return True
        # Return True if the end is reachable after
        processing all indices
        return True

# Example usage:
# print(canJump([2,3,1,1,4])) # Expected output:
True
# print(canJump([3,2,1,0,4])) # Expected output:
False
```

• LC.ca

```
class Solution:
    def canJump(self, nums: List[int]) -> bool:
        mx = 0
        for i, x in enumerate(nums):
            if mx < i:
                return False
            mx = max(mx, i + x)
        return True
```

◆ Quick Review

Key Insights

- The greedy approach is optimal for this problem.

- Track the furthest index reachable as you iterate through the array.
- If the current index exceeds the maximum reachable index, it is impossible to proceed further.
- The solution only requires constant extra space.

Space & Time

Time Complexity: $O(n)$ where n is the length of the array.

Space Complexity: $O(1)$

Solution

The solution uses a greedy algorithm. As you iterate through each array element, you maintain a variable (`maxReach`) that represents the furthest index you can reach so far. For each index i , if i is greater than `maxReach` then it is not possible to reach index i and hence return false immediately. Otherwise, update `maxReach` to the maximum of its current value and $i + \text{nums}[i]$. If you can iterate through the array without i exceeding `maxReach`, then you can reach the end and return true.

Round 8 · Low · Medium

p.2769

Dynamic Programming

#62 Unique Paths

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Dynamic Programming · Combinatorics

Lists: Grind 169

Problem

Description

There is a robot on an $m \times n$ grid. The robot is initially located at the top-left corner (i.e., `grid[0][0]`). The robot tries to move to the bottom-right corner (i.e., `grid[m - 1][n - 1]`). The robot can only move either down or right at any point in time.

Given the two integers m and n , return the number of possible unique paths that the robot can take to reach the bottom-right corner.

The test cases are generated so that the answer will be less than or equal to $2 * 10^9$.

Example 1:

Round 8 · Low · Medium

p.2770



Input: $m = 3, n = 7$
Output: 28

Example 2:

Input: $m = 3, n = 2$
Output: 3
Explanation: From the top-left corner, there are a total of 3 ways to reach the bottom-right corner:
1. Right -> Down -> Down
2. Down -> Down -> Right
3. Down -> Right -> Down

Constraints:

- $1 \leq m, n \leq 100$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        # dp[i][j] := the number of unique paths from
        # (0, 0) to (i, j)
        dp = [[1] * n for _ in range(m)]

        for i in range(1, m):
            for j in range(1, n):
                dp[i][j] = dp[i - 1][j] + dp[i][j - 1]

        return dp[-1][-1]
```

```
class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        dp = [1] * n

        for _ in range(1, m):
            for j in range(1, n):
                dp[j] += dp[j - 1]

        return dp[n - 1]
```

- LeetCode

```
class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        f = [[0] * n for _ in range(m)]
        f[0][0] = 1
        for i in range(m):
            for j in range(n):
                if i:
                    f[i][j] += f[i - 1][j]
                if j:
                    f[i][j] += f[i][j - 1]

        return f[-1][-1]
```

```
class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        f = [[1] * n for _ in range(m)]
        for i in range(1, m):
            for j in range(1, n):
                f[i][j] = f[i - 1][j] + f[i][j - 1]
        return f[-1][-1]
```

```
class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        f = [1] * n
        for _ in range(1, m):
            for j in range(1, n):
                f[j] += f[j - 1]
        return f[-1]
```

• SimplyLeet

Round 8 · Low · Medium

p.2773

Python solution using Dynamic Programming

```
def uniquePaths(m, n):
    # Initialize a 2D list with 1's for the first
    row and first column
    dp = [[1] * n for _ in range(m)]

    # Fill in the DP table
    for i in range(1, m):
        for j in range(1, n):
            # The number of ways to reach this cell
            is the sum of the ways from the cell above and the
            cell on the left
```

```
        dp[i][j] = dp[i-1][j] + dp[i][j-1]
    # Return the value in the bottom-right cell
    return dp[m-1][n-1]
```

```
# Example usage:
print(uniquePaths(3, 7))
```

• LC.ca

Round 8 · Low · Medium

p.2774

```

class Solution:
    def uniquePaths(self, m: int, n: int) -> int:
        # avoid setting dp[][] to 1 for i==0 or
        j==0 as initialization
        dp = [[1] * n for _ in range(m)]
        for i in range(1, m):
            for j in range(1, n):
                dp[i][j] = dp[i - 1][j] + dp[i][j -
1]
        return dp[-1][-1]

```

◆ Quick Review

Key Insights

- The robot is constrained to only move right or down.
- Any valid path consists of exactly $(m-1)$ down moves and $(n-1)$ right moves.
- The problem can be solved using Dynamic Programming where each cell is the sum of the ways to reach the cell from the top and left.
- Alternatively, a combinatorial solution is possible by calculating the number of unique sequences of moves, equivalent to choosing $(m-1)$ moves from $(m+n-2)$ moves.

Space & Time

Time Complexity: $O(mn)$

Space Complexity: $O(mn)$ – can be optimized to $O(n)$ using a 1D DP array.

Solution

We can solve the problem using a dynamic programming approach. The idea is to create a 2D DP table where each cell (i, j) represents the number of ways to reach that cell. Since the robot can only come from above or from the left, we update the DP table as follows:

– Initialize $dp[0][j] = 1$ for all j (only one way to move right in the first row). – Initialize $dp[i][0] = 1$ for all i (only one way to move down in the first column). – For each other cell, compute $dp[i][j] = dp[i-1][j] + dp[i][j-1]$. The final answer will be at $dp[m-1][n-1]$.

The combinatorial alternative computes $C(m+n-2, m-1)$ or $C(m+n-2, n-1)$ using basic factorial arithmetic, which is efficient for the given constraints.

Dynamic Programming

#72 Edit Distance

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Denny Zhang

Problem

Description

Given two strings word1 and word2, return the minimum number of operations required to convert word1 to word2.

You have the following three operations permitted on a word:

- Insert a character
- Delete a character
- Replace a character

Example 1:

```
Input: word1 = "horse", word2 = "ros"
Output: 3
Explanation:
horse -> rorse (replace 'h' with 'r')
rorse -> rose (remove 'r')
rose -> ros (remove 'e')
```

Example 2:

```
Input: word1 = "intention", word2 = "execution"
```

```
Output: 5
```

```
Explanation:
```

```
intention -> inention (remove 't')
inention -> enention (replace 'i' with 'e')
enention -> exention (replace 'n' with 'x')
exention -> exection (replace 'n' with 'c')
exection -> execution (insert 'u')
```

Constraints:

- $0 \leq \text{word1.length}, \text{word2.length} \leq 500$
- word1 and word2 consist of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def minDistance(self, word1: str, word2: str) -> int:
        m = len(word1)
        n = len(word2)
        # dp[i][j] := the minimum number of operations
        # to convert word1[0..i) to word2[0..j)
        dp = [[0] * (n + 1) for _ in range(m + 1)]

        for i in range(1, m + 1):
            dp[i][0] = i
```

```

for j in range(1, n + 1):
    dp[0][j] = j

for i in range(1, m + 1):
    for j in range(1, n + 1):
        if word1[i - 1] == word2[j - 1]:
            dp[i][j] = dp[i - 1][j - 1]
        else:
            dp[i][j] = min(dp[i - 1][j - 1], dp[i - 1][j], dp[i][j - 1]) + 1

return dp[m][n]

```

• LeetCode

```

class Solution:
    def minDistance(self, word1: str, word2: str)
    -> int:
        m, n = len(word1), len(word2)
        f = [[0] * (n + 1) for _ in range(m + 1)]
        for j in range(1, n + 1):
            f[0][j] = j
        for i, a in enumerate(word1, 1):
            f[i][0] = i
            for j, b in enumerate(word2, 1):
                if a == b:

```

```

            f[i][j] = f[i - 1][j - 1]
        else:
            f[i][j] = min(f[i - 1][j],
            f[i][j - 1], f[i - 1][j - 1]) + 1
        return f[m][n]

```

• SimplyLeet

```

# Python solution to compute the Edit Distance
def minDistance(word1, word2):
    m = len(word1)
    n = len(word2)

    # Initialize the DP table with size (m+1) x (n+1)
    dp = [[0] * (n + 1) for _ in range(m + 1)]

    # Base cases: transform word1[0:i] to empty
    # word2 requires i deletions
    for i in range(m + 1):

```

```

    dp[i][0] = i
    # Base cases: transform empty word1 to
    word2[0:j] requires j insertions
    for j in range(n + 1):
        dp[0][j] = j

    # Fill up the dp table
    for i in range(1, m + 1):
        for j in range(1, n + 1):
            # If the characters are the same, no
            new operation is needed.
            if word1[i - 1] == word2[j - 1]:

```

```

                dp[i][j] = dp[i - 1][j - 1]
            else:
                # consider insertion, deletion,
                and replacement
                dp[i][j] = 1 + min(dp[i][j - 1], #
                Insert
                                dp[i - 1][j], #
                Delete
                                dp[i - 1][j -
                1]) # Replace
            return dp[m][n]

    # Example usage:
    print(minDistance("horse", "ros"))

```

• LC.ca

Round 8 · Low · Medium

p.2781

```

class Solution:
    def minDistance(self, word1: str, word2: str)
    -> int:
        m, n = len(word1), len(word2)
        dp = [[0] * (n + 1) for _ in range(m + 1)]
        for i in range(m + 1):
            dp[i][0] = i
        for j in range(n + 1):
            dp[0][j] = j
        for i in range(1, m + 1):
            for j in range(1, n + 1):

```

```

                if word1[i - 1] == word2[j - 1]:
                    dp[i][j] = dp[i - 1][j - 1]
                else:
                    # dp[i - 1][j - 1]) meaning
                    replace.
                    # e.g. "abc" and "bf", checking
                    "ab" and "b" distance,
                    # then replace either way for "c"
                    or "f"
                    dp[i][j] = 1 + min(dp[i][j -
                    1], dp[i - 1][j], dp[i - 1][j - 1])
            return dp[-1][-1]

```

#####

Round 8 · Low · Medium

p.2782

```
class Solution:
    def minDistance(self, word1: str, word2: str)
-> int:
    m, n = len(word1), len(word2)
    f = [[0] * (n + 1) for _ in range(m + 1)]
    for j in range(1, n + 1):
        f[0][j] = j
    for i, a in enumerate(word1, 1):
        f[i][0] = i # merged for loops from
above solution, but not good for readability
        for j, b in enumerate(word2, 1):
            if a == b:
                f[i][j] = f[i - 1][j - 1]
            else:
                f[i][j] = min(f[i - 1][j],
f[i][j - 1], f[i - 1][j - 1]) + 1
    return f[m][n]
```

◆ Quick Review

Key Insights

- This is a classic dynamic programming problem.
- Use a 2D DP table where $dp[i][j]$ represents the minimum edit distance between $word1[0...i-1]$ and $word2[0...j-1]$.
- Base cases: converting an empty string to a string of length j requires j insertions, and vice versa.

- Transition: If characters match then $dp[i][j] = dp[i-1][j-1]$, otherwise take minimum among insertion, deletion, and replacement operations and add 1.

Space & Time

Time Complexity: $O(mn)$

where m and n are the lengths of $word1$ and $word2$ respectively.

Space Complexity: $O(mn)$

due to the DP table used for storing computed distances.

Solution

We use a dynamic programming approach by constructing a 2D table dp of size $(m+1) \times (n+1)$, where m and n are the lengths of $word1$ and $word2$. Each $dp[i][j]$ holds the minimum number of operations needed to convert the first i characters of $word1$ to the first j characters of $word2$. The algorithm includes the following steps:

- Initialize the base cases: $dp[0][j] = j$ (converting an empty $word1$ to j characters of $word2$ requires j insertions)
 - $dp[i][0] = i$ (converting i characters of $word1$ to an empty $word2$ requires i deletions)
 - Iterate over each character in $word1$ and $word2$: If the current characters are equal, no additional operation is needed and $dp[i][j] = dp[i-1][j-1]$. Otherwise, choose the best operation among insertion ($dp[i][j-1]$), deletion ($dp[i-1][j]$), or substitution ($dp[i-1][j-1]$) and add 1.
 - The answer is found in $dp[m][n]$.
- This method reliably computes the edit distance using a systematic DP approach, ensuring that all subproblems are solved and reused efficiently.

Dynamic Programming

#91 Decode Ways

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Grind 169

Problem

Description

You have intercepted a secret message encoded as a string of numbers. The message is decoded via the following mapping:

"1" -> 'A' "2" -> 'B' ... "25" -> 'Y' "26" -> 'Z'

However, while decoding the message, you realize that there are many different ways you can decode the message because some codes are contained in other codes ("2" and "5" vs "25").

For example, "11106" can be decoded into:

- "AAJF" with the grouping (1, 1, 10, 6)
- "KJF" with the grouping (11, 10, 6)
- The grouping (1, 11, 06) is invalid because "06" is not a valid code (only "6" is valid).

Note: there may be strings that are impossible to decode. Given a string *s* containing only digits, return the number of ways to decode it. If the entire string

cannot be decoded in any valid way, return 0.

The test cases are generated so that the answer fits in a 32-bit integer.

Example 1:

Input: *s* = "12"

Output: 2

Explanation:

"12" could be decoded as "AB" (1 2) or "L" (12).

Example 2:

Input: *s* = "226"

Output: 3

Explanation:

"226" could be decoded as "BZ" (2 26), "VF" (22 6), or "BBF" (2 2 6).

Example 3:

Input: *s* = "06"

Output: 0

Explanation:

"06" cannot be mapped to "F" because of the leading zero ("6" is different from "06"). In this case, the string is not a valid encoding, so return 0.

Constraints:

- $1 \leq s.length \leq 100$
- *s* contains only digits and may contain leading zero(s).

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numDecodings(self, s: str) -> int:
        n = len(s)
        # dp[i] := the number of ways to decode s[i..n)
        dp = [0] * n + [1]

        def isValid(a: str, b=None) -> bool:
            if b:
                return a == '1' or a == '2' and b < '7'
            return a != '0'

        if isValid(s[-1]):
            dp[n - 1] = 1

        for i in reversed(range(n - 1)):
            if isValid(s[i]):
                dp[i] += dp[i + 1]
            if isValid(s[i], s[i + 1]):
                dp[i] += dp[i + 2]

        return dp[0]
```

• LeetDoocs

```
class Solution:
    def numDecodings(self, s: str) -> int:
        n = len(s)
        f = [1] + [0] * n
        for i, c in enumerate(s, 1):
            if c != "0":
                f[i] = f[i - 1]
            if i > 1 and s[i - 2] != "0" and
int(s[i - 2 : i]) <= 26:
                f[i] += f[i - 2]
        return f[n]
```

```
class Solution:
    def numDecodings(self, s: str) -> int:
        f, g = 0, 1
        for i, c in enumerate(s, 1):
            h = g if c != "0" else 0
            if i > 1 and s[i - 2] != "0" and
int(s[i - 2 : i]) <= 26:
                h += f
            f, g = g, h
        return g
```

• SimplyLeet

```
# Python solution using dynamic programming
def num_decodings(s: str) -> int:
    # If the string starts with '0', it's not
    # decodable
    if not s or s[0] == '0':
        return 0

    # dp array where dp[i] represents the number of
    # ways to decode s[:i]
    n = len(s)
    dp = [0] * (n + 1)

    # Base case: an empty string has one decoding
    dp[0] = 1

    for i in range(1, n + 1):
        # Check for single digit decoding (s[i-1]
        # is valid if not '0')
        if s[i - 1] != '0':
            dp[i] += dp[i - 1]

        # Check for two-digit decoding (if i>=2,
        # s[i-2:i] forms a valid number between 10 and 26)
        if i >= 2 and s[i - 2] != '0':
            two_digit = int(s[i - 2:i])
```

```
        if 10 <= two_digit <= 26:
            dp[i] += dp[i - 2]

    return dp[n]

# Example usage:
print(num_decodings("12")) # Expected output: 2
```

• LC.ca

```
'''
>>> s = "abcdef"
>>> [(i,c) for i, c in enumerate(s, 1)]
[(1, 'a'), (2, 'b'), (3, 'c'), (4, 'd'), (5, 'e'),
(6, 'f')]
'''

class Solution:
    def numDecodings(self, s: str) -> int:
        # f: i-2
        # g: i-1
        f, g = 0, 1
```

```

    for i, c in enumerate(s, 1):
        h = g if c != "0" else 0
        if i > 1 and s[i - 2] != "0" and
int(s[i - 2 : i]) <= 26:
            h += f
        f, g = g, h
    return g

```

```
#####
```

```
class Solution:
```

```

    def numDecodings(self, s: str) -> int:
        if len(s) == 0:
            return 0

        length = len(s)
        dp = [0] * (length + 1) # dp[i] => at index
i, its decode ways

        dp[length] = 1 # initiator, just make
forloop flow working
        dp[length - 1] = 0 if s[length - 1] == '0'
else 1 # assumption is starting at i, so starting
as 0 is not decodable

```

```

    # start at 'len-2'
    # or else below 'dp[i+2]' will out of index
boundary
    for i in range(length - 2, -1, -1):
        if s[i] == '0':
            continue
        tem = int(s[i:i + 2])

        if tem > 26:
            dp[i] = dp[i + 1]
        else:

```

```

            dp[i] = dp[i + 1] + dp[i + 2]
        return dp[0]

```

```

# optimized from above solution, no dp array
class Solution:
    def numDecodings(self, s: str) -> int:
        n = len(s)
        # a: dp[i-2], b: dp[i-1], c: count for
current index i
        a, b, c = 0, 1, 0
        for i in range(1, n + 1):

```



```

        c = 0
        if s[i - 1] != '0':
            c += b
            if i > 1 and s[i - 2] != '0' and
(int(s[i - 2]) * 10 + int(s[i - 1])) <= 26):
                c += a
            a, b = b, c
        return c

```

#####

```

class Solution:
    def numDecodings(self, s: str) -> int:
        n = len(s)
        dp = [0] * (n + 1)
        dp[0] = 1
        for i in range(1, n + 1):
            if s[i - 1] != '0':
                dp[i] += dp[i - 1]
            if i > 1 and s[i - 2] != '0' and
(int(s[i - 2]) * 10 + int(s[i - 1])) <= 26):
                dp[i] += dp[i - 2]

        return dp[n]

```

◆ Quick Review

Key Insights

- A digit '0' cannot be mapped on its own; it must be part of "10" or "20".

- Use dynamic programming where $dp[i]$ represents the number of ways to decode the substring $s[0...i-1]$.
- For each index i , if the digit at position $i-1$ is non-zero, add $dp[i-1]$ to $dp[i]$.
- If the two-digit number formed by $s[i-2]$ and $s[i-1]$ is between 10 and 26, add $dp[i-2]$ to $dp[i]$.
- Initialize $dp[0]$ as 1 (an empty string has one way to be decoded).

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(n)$, due to the dp array (which can be optimized to $O(1)$).

Solution

We solve the problem using a dynamic programming approach. We initialize a dp array where $dp[i]$ indicates the number of ways to decode the substring $s[0...i-1]$. We set $dp[0] = 1$ since an empty string is trivially decodable. We loop through the string, and at each step:

- If $s[i-1]$ is not '0', we add $dp[i-1]$ to $dp[i]$ because $s[i-1]$ can be mapped to a valid letter.
- If the substring $s[i-2:i]$ (for $i \geq 2$) forms a valid two-digit number between 10 and 26, we add $dp[i-2]$ to $dp[i]$. This cumulative approach yields the total ways to decode the string, and if the string starts with '0' or contains invalid sequences, $dp[n]$ will eventually be 0.

Dynamic Programming

#152 Maximum Product Subarray

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Grind 169

Problem

Description

Given an integer array `nums`, find a subarray that has the largest product, and return the product.

The test cases are generated so that the answer will fit in a 32-bit integer.

Note that the product of an array with a single element is the value of that element.

Example 1:

Input: `nums = [2,3,-2,4]`

Output: 6

Explanation: [2,3] has the largest product 6.

Example 2:

Input: `nums = [-2,0,-1]`

Output: 0

Explanation: The result cannot be 2, because [-2,-1] is not a subarray.

Constraints:

- $1 \leq \text{nums.length} \leq 2 * 10^4$
- $-10 \leq \text{nums}[i] \leq 10$
- The product of any subarray of `nums` is guaranteed to fit in a 32-bit integer.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxProduct(self, nums: list[int]) -> int:
        ans = nums[0]
        dpMin = nums[0] # the minimum so far
        dpMax = nums[0] # the maximum so far

        for i in range(1, len(nums)):
            num = nums[i]
            prevMin = dpMin # dpMin[i - 1]
            prevMax = dpMax # dpMax[i - 1]

            if num < 0:
                dpMin = min(prevMax * num, num)
                dpMax = max(prevMin * num, num)
            else:
                dpMin = min(prevMin * num, num)
                dpMax = max(prevMax * num, num)

            ans = max(ans, dpMax)

        return ans
```

- LeetCode

```
class Solution:
    def maxProduct(self, nums: List[int]) -> int:
        ans = f = g = nums[0]
        for x in nums[1:]:
            ff, gg = f, g
            f = max(x, ff * x, gg * x)
            g = min(x, ff * x, gg * x)
            ans = max(ans, f)
        return ans
```

- SimplyLeet

```
def maxProduct(nums):
    # If the input array is empty, return 0 (though
    # problem guarantees at least one element)
    if not nums:
        return 0
    # Initialize maximum, minimum product ending
    # here and the result.
    max_ending_here = nums[0]
    min_ending_here = nums[0]
    global_max = nums[0]

    # Iterate over the array starting from the
    # second element.
```

```
        for num in nums[1:]:
            # If the current number is negative, swap
            # max and min.
            if num < 0:
                max_ending_here, min_ending_here =
                min_ending_here, max_ending_here
            # Update the max product at the current
            # index.
            max_ending_here = max(num, max_ending_here
            * num)
            # Update the min product at the current
            # index.
            min_ending_here = min(num, min_ending_here
            * num)
            # Update the global maximum product.
            global_max = max(global_max,
            max_ending_here)

        return global_max

# Example usage:
# print(maxProduct([2,3,-2,4])) # Output: 6
```

- LC.ca

```

class Solution:
    def maxProduct(self, nums: List[int]) -> int:
        ans = f = g = nums[0]
        for x in nums[1:]:
            ff, gg = f, g
            f = max(x, ff * x, gg * x)
            g = min(x, ff * x, gg * x)
            ans = max(ans, f)
        return ans

```

◆ Quick Review

Key Insights

- Maintain both the maximum and minimum product at the current index because a negative number can swap the max and min.
- A subarray must be contiguous.
- Use dynamic programming to update both the maximum product and the minimum product ending at each index.
- Resetting the product in case of a zero is crucial since the product becomes 0.

Space & Time

Time Complexity: $O(n)$ – The solution iterates through the array once.

Space Complexity: $O(1)$ – Constant space is used for tracking products.

Solution

The solution uses an iterative dynamic programming approach. At each index, compute the maximum product (maxEndingHere) and minimum product (minEndingHere) ending at that index. If the current element is negative, swap these two values, because multiplying a negative with a minimum product (possibly negative) might yield a larger product. After processing each element, update the global maximum. This method handles edge cases like zeros and negatives efficiently with constant additional space.

Dynamic Programming

#198 House Robber

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Grind 169

Problem

Description

You are a professional robber planning to rob houses along a street. Each house has a certain amount of money stashed, the only constraint stopping you from robbing each of them is that adjacent houses have security systems connected and it will automatically contact the police if two adjacent houses were broken into on the same night.

Given an integer array `nums` representing the amount of money of each house, return the maximum amount of money you can rob tonight without alerting the police.

Example 1:

Input: `nums = [1,2,3,1]`
 Output: 4
 Explanation: Rob house 1 (money = 1) and then rob house 3 (money = 3).
 Total amount you can rob = 1 + 3 = 4.

Example 2:

Input: `nums = [2,7,9,3,1]`

Output: 12

Explanation: Rob house 1 (money = 2), rob house 3 (money = 9) and rob house 5 (money = 1).

Total amount you can rob = 2 + 9 + 1 = 12.

Constraints:

- $1 \leq \text{nums.length} \leq 100$
- $0 \leq \text{nums}[i] \leq 400$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def rob(self, nums: list[int]) -> int:
        if not nums:
            return 0
        if len(nums) == 1:
            return nums[0]

        # dp[i] := max money of robbing nums[0..i]
        dp = [0] * len(nums)
        dp[0] = nums[0]

        dp[1] = max(nums[0], nums[1])

        for i in range(2, len(nums)):
            dp[i] = max(dp[i - 1], dp[i - 2] + nums[i])

        return dp[-1]
```

```
class Solution:
    def rob(self, nums: List[int]) -> int:
        prev1 = 0 # dp[i - 1]
        prev2 = 0 # dp[i - 2]

        for num in nums:
            dp = max(prev1, prev2 + num)
            prev2 = prev1
            prev1 = dp

        return prev1
```

• LeetCode

```
class Solution:
    def rob(self, nums: List[int]) -> int:
        @cache
        def dfs(i: int) -> int:
            if i >= len(nums):
                return 0
            return max(nums[i] + dfs(i + 2), dfs(i + 1))

        return dfs(0)
```

```
class Solution:
    def rob(self, nums: List[int]) -> int:
        n = len(nums)
        f = [0] * (n + 1)
        f[1] = nums[0]
        for i in range(2, n + 1):
            f[i] = max(f[i - 1], f[i - 2] + nums[i - 1])

        return f[n]
```

```
class Solution:
    def rob(self, nums: List[int]) -> int:
        f = g = 0
        for x in nums:
            f, g = max(f, g), f + x
        return max(f, g)
```

• SimplyLeet

```
# Definition of the Solution class
class Solution:
    def rob(self, nums: List[int]) -> int:
        # Edge case: if there are no houses, return 0

        if not nums:
            return 0

        # Initialize two variables to keep track of the max amounts
        prev1, prev2 = 0, 0
```

```
# Iterate over each house's money
for money in nums:
    # Calculate the maximum money that can
    # be robbed if we consider the current house
    current = max(prev1, prev2 + money) #
    Either skip current house or rob it
    # Update variables: prev2 becomes
    # previous best, and prev1 becomes current best
    prev2 = prev1
    prev1 = current

# Return the maximum money that can be
# robbed
return prev1
```

• LC.ca

```
# greedy
class Solution:
    def rob(self, nums: List[int]) -> int:
        not_rob, rob = 0, nums[0]
        for num in nums[1:]:
            # must max check
            # eg. first robbed 99999, then
            # following ones are just 3,3,3,3,3
            not_rob, rob = rob, max(num + not_rob,
                                     rob)
        return rob
```

```
#####

class Solution(object):
    def rob(self, nums):
        """
        :type nums: List[int]
        :rtype: int
        """
        if len(nums) == 0:
            return 0

        if len(nums) <= 2:
            return max(nums)
        dp = [0 for i in range(0, 2)]
        dp[0] = nums[0]
        dp[1] = max(nums[1], nums[0])
        for i in range(2, len(nums)):
            dp[i % 2] = max(dp[(i - 1) % 2], dp[(i - 2) %
            2] + nums[i])
        return dp[(len(nums) - 1) % 2]
```

◆ Quick Review

Key Insights

- The problem can be solved using Dynamic Programming.
- For each house, decide whether to rob it or not.
- Recurrence relation: $dp[i] = \max(dp[i-1], dp[i-2] + \text{currentHouseMoney})$

- Use iterative computation to keep track of the optimal solutions without needing extra space for the entire dp array.

Space & Time

Time Complexity: $O(n)$, where n is the number of houses.

Space Complexity: $O(1)$, using only two variables to track the previous decisions.

Solution

We use a dynamic programming approach by iterating through each house in the array. At each step, we consider two scenarios: either we rob the current house (and add its money to the total from two houses ago) or we skip it (keeping the optimal solution up to the previous house). By using two variables to store these intermediate results, we achieve an optimized space complexity of $O(1)$.

Dynamic Programming

#256 Paint House

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Premium 98

Problem

Description

There is a row of n houses, where each house can be painted one of three colors: red, blue, or green. The cost of painting each house with a certain color is different. You have to paint all the houses such that no two adjacent houses have the same color.

The cost of painting each house with a certain color is represented by an $n \times 3$ cost matrix costs.

- For example, `costs[0][0]` is the cost of painting house 0 with the color red; `costs[1][2]` is the cost of painting house 1 with color green, and so on...

Return the minimum cost to paint all houses.

Example 1:

Input: `costs = [[17,2,17],[16,16,5],[14,3,19]]`

Output: 10

Explanation: Paint house 0 into blue, paint house 1 into green, paint house 2 into blue.

Minimum cost: $2 + 5 + 3 = 10$.

Example 2:**Input:** costs = [[7,6,2]]**Output:** 2**Constraints:**

- costs.length == n
- costs[i].length == 3
- 1 <= n <= 100
- 1 <= costs[i][j] <= 20

Community Solutions (Python)

• WalkCC

```
class Solution:
    def minCost(self, costs: list[list[int]]) -> list[list[int]]:
        for i in range(1, len(costs)):
            costs[i][0] += min(costs[i - 1][1], costs[i - 1][2])
            costs[i][1] += min(costs[i - 1][0], costs[i - 1][2])
            costs[i][2] += min(costs[i - 1][0], costs[i - 1][1])

        return min(costs[-1])
```

• LeetDoocs

class Solution:

```
    def minCost(self, costs: List[List[int]]) -> int:
        a = b = c = 0
        for ca, cb, cc in costs:
            a, b, c = min(b, c) + ca, min(a, c) + cb, min(a, b) + cc
        return min(a, b, c)
```

• SimplyLeet

Python solution for the Paint House problem

```
def minCost(costs):
    # Get the number of houses
    n = len(costs)
    # Base case: If there are no houses, return 0
    cost
    if n == 0:
        return 0

    # Iterate from the second house to the last house
```

```

for i in range(1, n):
    # Update the cost for painting the current
    house red
    costs[i][0] += min(costs[i-1][1],
costs[i-1][2])
    # Update the cost for painting the current
    house blue
    costs[i][1] += min(costs[i-1][0],
costs[i-1][2])
    # Update the cost for painting the current
    house green
    costs[i][2] += min(costs[i-1][0],
costs[i-1][1])

    # The answer is the minimum cost of painting
    the last house with any of the colors
    return min(costs[-1])

```

```

# Example usage:
example_costs = [[17,2,17],[16,16,5],[14,3,19]]
print(minCost(example_costs)) # Output: 10

```

• LC.ca

Round 8 · Low · Medium

p.2811

```

class Solution:
    def minCost(self, costs: List[List[int]]) ->
    int:
        a = b = c = 0
        for ca, cb, cc in costs:
            a, b, c = min(b, c) + ca, min(a, c) +
cb, min(a, b) + cc
        return min(a, b, c)

```

#####

...

Two-dimensional dynamic array dp,
where dp[i][j] represents the minimum cost of
brushing to the i+1th house with color j,

and the state transition equation is:
 $dp[i][j] = cost[i][j] + \min(dp[i-1][(j + 1) \% 3], dp[i-1][(j + 2) \% 3])$

...

```

class Solution(object):
    def minCost(self, costs):
        """

```

Round 8 · Low · Medium

p.2812

```

:type costs: List[List[int]]
:rtype: int
"""
if not costs:
    return 0
dp = [0] * (len(costs[0]))
dp[:] = costs[0]
for i in range(1, len(costs)):
    d0 = d1 = d2 = 0
    for j in range(0, 3):

        if j == 0:
            d0 = min(dp[1], dp[2]) + costs[i][0]
        elif j == 1:
            d1 = min(dp[0], dp[2]) + costs[i][1]
        else:
            d2 = min(dp[0], dp[1]) + costs[i][2]
    dp[0], dp[1], dp[2] = d0, d1, d2
return min(dp)

```

◆ Quick Review

Key Insights

- The choice for painting each house depends on the previous house's color choices.
- Use dynamic programming to accumulate the cost while ensuring adjacent houses do not have the same color.
- The state of each house can be represented by the minimum cost when the house is painted a specific color, computed using the minimal costs of the previous house

from the other two colors.

- The problem only requires updating the cost matrix in-place for optimal space usage.

Space & Time

Time Complexity: $O(n)$, where n is the number of houses.

Space Complexity: $O(1)$ extra space if the input matrix is modified in place.

Solution

We use a dynamic programming approach with iterative in-place updating of the cost matrix. For each house starting from the second one, we update the cost for each color by adding the minimal cost of painting the previous house with either of the other two colors. This guarantees we never use the same color for adjacent houses while accumulating the minimal total cost. The final answer is the minimum value for the last house.

Dynamic Programming

#309 Best Time to Buy and Sell Stock with Cooldown

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Denny Zhang

Problem

Description

You are given an array `prices` where `prices[i]` is the price of a given stock on the `i`th day.

Find the maximum profit you can achieve. You may complete as many transactions as you like (i.e., buy one and sell one share of the stock multiple times) with the following restrictions:

- After you sell your stock, you cannot buy stock on the next day (i.e., cooldown one day).

Note: You may not engage in multiple transactions simultaneously (i.e., you must sell the stock before you buy again).

Example 1:

Input: `prices = [1,2,3,0,2]`
 Output: 3
 Explanation: `transactions = [buy, sell, cooldown, buy, sell]`

Round 8 · Low · Medium

p.2815

Example 2:

Input: `prices = [1]`
 Output: 0

Constraints:

- $1 \leq \text{prices.length} \leq 5000$
- $0 \leq \text{prices}[i] \leq 1000$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxProfit(self, prices: list[int]) -> int:
        sell = 0
        hold = -math.inf
        prev = 0

        for price in prices:
            cache = sell
            sell = max(sell, hold + price)
            hold = max(hold, prev - price)

        prev = cache

        return sell
```

• LeetDoocs

Round 8 · Low · Medium

p.2816

```

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        @cache
        def dfs(i: int, j: int) -> int:
            if i >= len(prices):
                return 0
            ans = dfs(i + 1, j)
            if j:
                ans = max(ans, prices[i] + dfs(i +
2, 0))
            else:
                ans = max(ans, -prices[i] + dfs(i +
1, 1))
            return ans
        return dfs(0, 0)

```

```

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        n = len(prices)
        f = [[0] * 2 for _ in range(n)]
        f[0][1] = -prices[0]
        for i in range(1, n):
            f[i][0] = max(f[i - 1][0], f[i - 1][1]
+ prices[i])
            f[i][1] = max(f[i - 1][1], f[i - 2][0]
- prices[i])
        return f[n - 1][0]

```

```

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        f, f0, f1 = 0, 0, -prices[0]
        for x in prices[1:]:
            f, f0, f1 = f0, max(f0, f1 + x),
max(f1, f - x)
        return f0

```

• SimplyLeet

Python Solution with line-by-line comments

```

def maxProfit(prices):
    # Handle edge case: if there are no prices,
    profit is 0
    if not prices:
        return 0

    # Initialize states:
    # s0: Maximum profit when ready to buy
    # s1: Maximum profit when holding a stock
    (bought)

```

```

# s2: Maximum profit during cooldown after a
sale
s0, s1, s2 = 0, -prices[0], 0

# Iterate over prices starting from day 1
for price in prices[1:]:
    # Store previous state values
    prev_s0, prev_s1, prev_s2 = s0, s1, s2
    # s0 can come either from staying in s0 or
    coming from a cooldown (s2)
    s0 = max(prev_s0, prev_s2)
    # s1 is updated by either keeping the stock
    or buying today with profit from s0

    s1 = max(prev_s1, prev_s0 - price)
    # s2 is achieved by selling the stock that
    we were holding (prev_s1 + price)
    s2 = prev_s1 + price
    # The final maximum profit is the maximum of
    not holding a stock (s0 or s2)
    return max(s0, s2)

# Example usage:
print(maxProfit([1,2,3,0,2]))

```

• LC.ca

Round 8 · Low · Medium

p.2819

```

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        f1, f2, f3 = -prices[0], 0, 0
        for price in prices[1:]:
            pf1, pf2, pf3 = f1, f2, f3
            f1 = max(pf1, pf3 - price)
            f2 = max(pf2, pf1 + price)
            f3 = max(pf3, pf2) # cooldown
        return f2

```

#####

```

class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        f, f0, f1 = 0, 0, -prices[0]
        for x in prices[1:]:
            f, f0, f1 = f0, max(f0, f1 + x),
            max(f1, f - x)
        return f0

```

#####

Round 8 · Low · Medium

p.2820

```

class Solution(object):
    def maxProfit(self, prices):
        """
        :type prices: List[int]
        :rtype: int
        """
        if len(prices) < 2:
            return 0
        buy = [0] * len(prices)

        sell = [0] * len(prices)
        buy[0] = -prices[0]
        buy[1] = max(-prices[1], buy[0])
        sell[0] = 0
        sell[1] = max(prices[1] - prices[0], 0)
        for i in range(2, len(prices)):
            buy[i] = max(sell[i - 2] - prices[i], buy[i - 1])
            sell[i] = max(prices[i] + buy[i - 1], sell[i - 1])
        return max(sell)

```

◆ Quick Review

Key Insights

- Use Dynamic Programming to solve the problem.
- Maintain three states to represent:
 - s0: the state where you are ready to buy (not holding any stock).

- s1: the state where you are holding a stock.
- s2: the cooldown state encountered right after selling a stock.
- Transition between states:
 - Transition into s0 can come from staying in s0 or coming out from the cooldown state s2.
 - To enter s1, you either continue holding your stock, or buy a stock using the profit available from s0.
 - The only way to reach s2 is by selling the stock from s1.
 - The final answer is determined by the maximum profit when you are not holding a stock (i.e., max(s0, s2)).

Space & Time

Time Complexity: $O(n)$, where n is the number of days.

Space Complexity: $O(1)$, since only a few variables are maintained for the states.

Solution

The solution involves updating three state variables for each day:

- s0 (ready to buy) is updated as the maximum of the previous s0 and the previous cooldown s2.
 - s1 (holding a stock) is updated as the maximum of the previous s1 or buying today with the profit from s0 minus the current price.
 - s2 (cooldown) is updated by selling the stock held in s1 (i.e., previous s1 plus the current price).
- At the end of the iteration, the maximum profit will be in state s0 (ready to buy) or s2 (cooldown), since remaining in s1 means you are holding a stock, which does not constitute a realized profit.

Dynamic Programming

#377 Combination Sum IV

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Grind 169

Problem

Description

Given an array of distinct integers `nums` and a target integer `target`, return the number of possible combinations that add up to `target`.

The test cases are generated so that the answer can fit in a 32-bit integer.

Example 1:

```
Input: nums = [1,2,3], target = 4
Output: 7
Explanation:
The possible combination ways are:
(1, 1, 1, 1)
(1, 1, 2)
(1, 2, 1)
(1, 3)
```

Example 2:

```
Input: nums = [9], target = 3
Output: 0
```

Constraints:

- $1 \leq \text{nums.length} \leq 200$
- $1 \leq \text{nums}[i] \leq 1000$
- All the elements of `nums` are unique.
- $1 \leq \text{target} \leq 1000$

Follow up: What if negative numbers are allowed in the given array? How does it change the problem? What limitation we need to add to the question to allow negative numbers?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def combinationSum4(self, nums: list[int],
target: int) -> int:
    dp = [1] + [-1] * target

    def dfs(target: int) -> int:
        if target < 0:
            return 0
        if dp[target] != -1:
            return dp[target]
```



```

        dp[target] = sum(dfs(target - num) for num in
nums)
    return dp[target]

return dfs(target)

```

- LeetCode

```

class Solution:
    def combinationSum4(self, nums: List[int],
target: int) -> int:
    f = [1] + [0] * target
    for i in range(1, target + 1):
        for x in nums:
            if i >= x:
                f[i] += f[i - x]
    return f[target]

```

- SimplyLeet

Python solution using dynamic programming

```

class Solution:
    def combinationSum4(self, nums, target):
        # Initialize dp array where dp[i] stores
the number of ways to reach sum i
        dp = [0] * (target + 1)
        # Base case: there is one way to form a sum
of 0 (use no elements)
        dp[0] = 1

        # Fill the dp table for each sum from 1 to
target

        for i in range(1, target + 1):
            for num in nums:
                if i - num >= 0:
                    dp[i] += dp[i - num]
        return dp[target]

```

- LC.ca

```

class Solution:
    def combinationSum4(self, nums: List[int],
target: int) -> int:
    f = [1] + [0] * target
    for i in range(1, target + 1):
        for x in nums:
            if i >= x:
                f[i] += f[i - x]
    return f[target]

```

◆ Quick Review

Key Insights

- Use dynamic programming to count the number of valid combinations.
- The order in which numbers are picked matters, making this akin to permutation counting.
- Initialize a DP table where $dp[i]$ represents the number of ways to form the sum i .
- $dp[0]$ is initialized to 1 since there is one way to form a sum of 0 (by choosing nothing).
- For each sum i from 1 to target, iterate over $nums$ and update $dp[i]$ based on $dp[i - num]$ when $i \geq num$.
- For negative numbers, the problem becomes unbounded due to potential infinite combinations; hence a constraint on the combination length is required to make it well-defined.

Space & Time

Time Complexity: $O(\text{target} * n)$, where n is the length of $nums$.

Space Complexity: $O(\text{target})$, for the dp array of size $\text{target} + 1$.

Solution

The problem is solved using a bottom-up dynamic programming approach. We create a dp array where each $dp[i]$ represents the number of possible combinations to reach the sum i . We iterate from 1 up to target, and for each

sum, we check every number in $nums$. If the number can contribute to the current sum ($i - num \geq 0$), we add the number of ways to form the remainder ($i - num$) to $dp[i]$. This ensures that every possible permutation is counted. When negative numbers are allowed in the input, the potential for an infinite number of combinations arises unless we impose additional constraints, such as limiting the length of the combination.

Dynamic Programming

#416 Partition Equal Subset Sum

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Grind 169

Problem

Description

Given an integer array `nums`, return `true` if you can partition the array into two subsets such that the sum of the elements in both subsets is equal or `false` otherwise.

Example 1:

Input: `nums = [1,5,11,5]`
 Output: `true`
 Explanation: The array can be partitioned as `[1, 5, 5]` and `[11]`.

Example 2:

Input: `nums = [1,2,3,5]`
 Output: `false`
 Explanation: The array cannot be partitioned into equal sum subsets.

Constraints:

- $1 \leq \text{nums.length} \leq 200$
- $1 \leq \text{nums}[i] \leq 100$

Round 8 · Low · Medium

p.2829

Community Solutions (Python)

• WalkCC

```
class Solution:
    def canPartition(self, nums: list[int]) -> bool:
        summ = sum(nums)
        if summ % 2 == 1:
            return False
        return self.knapsack_(nums, summ // 2)

    def knapsack_(self, nums: list[int], subsetSum:
int) -> bool:
        n = len(nums)
        # dp[i][j] := True if j can be formed by
        nums[0..i)

        dp = [[False] * (subsetSum + 1) for _ in
range(n + 1)]
        dp[0][0] = True

        for i in range(1, n + 1):
            num = nums[i - 1]
            for j in range(subsetSum + 1):
                if j < num:
                    dp[i][j] = dp[i - 1][j]
                else:
                    dp[i][j] = dp[i - 1][j] or dp[i - 1][j -
num]
```

Round 8 · Low · Medium

p.2830

```
return dp[n][subsetSum]
```

- LeetCode

```
class Solution:
    def canPartition(self, nums: List[int]) ->
    bool:
        m, mod = divmod(sum(nums), 2)
        if mod:
            return False
        n = len(nums)
        f = [[False] * (m + 1) for _ in range(n +
1)]
        f[0][0] = True
        for i, x in enumerate(nums, 1):
            for j in range(m + 1):
                f[i][j] = f[i - 1][j] or (j >= x
and f[i - 1][j - x])
        return f[n][m]
```

- SimplyLeet

Python solution for Partition Equal Subset Sum

```
def canPartition(nums):
    # Calculate the total sum of the array
    total_sum = sum(nums)

    # If the total sum is odd, equal partition is
    not possible
    if total_sum % 2 != 0:
        return False

    # Set the target sum to half of the total sum
    target = total_sum // 2

    # Initialize dp array where dp[i] is True if a
    subset sum i is achievable
    dp = [False] * (target + 1)
    dp[0] = True # Base case: subset sum of 0 is
    always possible

    # Iterate through each number in the array
    for num in nums:
        # Iterate backwards to avoid overwriting dp
        values needed for current iteration
```

```

    for i in range(target, num - 1, -1):
        # Update dp[i] if dp[i-num] is True
        if dp[i - num]:
            dp[i] = True

    # dp[target] is True if a subset adds up to
    target sum
    return dp[target]

# Example usage:
# print(canPartition([1, 5, 11, 5]))

```

• LC.ca

```

class Solution:
    def canPartition(self, nums: List[int]) ->
    bool:
        s = sum(nums)
        if s % 2 != 0:
            return False
        n = s >> 1
        dp = [False] * (n + 1)
        dp[0] = True
        for v in nums:
            for target in range(n, v - 1, -1): #
            including v itself

```

```

        dp[target] = dp[target] or
        dp[target - v]
        return dp[-1]

class Solution: # recursive, but over time limit in
0J
    def canPartition(self, nums: List[int]) ->
    bool:
        total_sum = sum(nums)
        if total_sum % 2 != 0:
            return False
        target_sum = total_sum // 2

        memo = {} # ==> or just use @cache for
        dfs()

        def dfs(idx, curr_sum):
            if curr_sum == target_sum:
                return True
            if curr_sum > target_sum or idx >=
            len(nums):
                return False
            if (idx, curr_sum) in memo:
                return memo[(idx, curr_sum)]

```

```

        # Explore two options: include the
        current number or exclude it
        include = dfs(idx + 1, curr_sum +
nums[idx])
        exclude = dfs(idx + 1, curr_sum)

        memo[(idx, curr_sum)] = include or
exclude
        return memo[(idx, curr_sum)]

    return dfs(0, 0)

```

#####

```

class Solution:
    def canPartition(self, nums: List[int]) ->
bool:
    m, mod = divmod(sum(nums), 2)
    if mod:
        return False
    f = [True] + [False] * m
    for x in nums:
        for j in range(m, x - 1, -1):
            f[j] = f[j] or f[j - x]

    return f[m]

```

◆ Quick Review

Round 8 · Low · Medium

p.2835

Key Insights

- The total sum of the array must be even, otherwise, partitioning into two equal subsets is impossible.
- The problem can be transformed into a "subset sum" problem where we need to check if there exists a subset of numbers that sums to half of the total sum.
- A dynamic programming approach can efficiently decide if such a subset exists.

Space & Time

Time Complexity: $O(n * \text{target})$, where target is half of the total sum.

Space Complexity: $O(\text{target})$, using a one-dimensional DP array.

Solution

First, compute the total sum of the array. If the total is odd, return false since an odd number cannot be equally partitioned. Next, set the target sum to half of the total. Use a dynamic programming array (dp) where $dp[i]$ represents whether a subset with sum i is achievable. Initialize $dp[0]$ as true because a subset sum of 0 is always possible. Then, for each number in the array, iterate through the dp array backwards from target to the number's value and update $dp[i]$ to true if $dp[i - \text{num}]$ is true. Ultimately, if $dp[\text{target}]$ is true, the array can be partitioned into two subsets with equal sum.

Round 8 · Low · Medium

p.2836

Dynamic Programming

#435 Non-overlapping Intervals

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming · Greedy · Sorting

Lists: Grind 169

Problem

Description

Given an array of intervals `intervals` where `intervals[i] = [starti, endi]`, return the minimum number of intervals you need to remove to make the rest of the intervals non-overlapping.

Note that intervals which only touch at a point are non-overlapping. For example, `[1, 2]` and `[2, 3]` are non-overlapping.

Example 1:

Input: `intervals = [[1,2],[2,3],[3,4],[1,3]]`

Output: 1

Explanation: `[1,3]` can be removed and the rest of the intervals are non-overlapping.

Example 2:

Input: `intervals = [[1,2],[1,2],[1,2]]`

Output: 2

Explanation: You need to remove two `[1,2]` to make the rest of the intervals non-overlapping.

Example 3:

Input: `intervals = [[1,2],[2,3]]`

Output: 0

Explanation: You don't need to remove any of the intervals since they're already non-overlapping.

Constraints:

- $1 \leq \text{intervals.length} \leq 105$
- $\text{intervals}[i].\text{length} == 2$
- $-5 * 10^4 \leq \text{starti} < \text{endi} \leq 5 * 10^4$

Community Solutions (Python)

- WalkCC

```

class Solution:
    def eraseOverlapIntervals(self, intervals:
list[list[int]]) -> int:
        ans = 0
        currentEnd = -math.inf

        for interval in sorted(intervals, key=lambda
x: x[1]):
            if interval[0] >= currentEnd:
                currentEnd = interval[1]
            else:
                ans += 1

        return ans

```

• LeetCode

```

class Solution:
    def eraseOverlapIntervals(self, intervals:
List[List[int]]) -> int:
        intervals.sort(key=lambda x: x[1])
        ans = len(intervals)
        pre = -inf
        for l, r in intervals:
            if pre <= l:
                ans -= 1
                pre = r
        return ans

```

• SimplyLeet

Python implementation for Non-overlapping Intervals

```

def eraseOverlapIntervals(intervals):
    # Sort intervals based on the end time
    intervals.sort(key=lambda x: x[1])

    # Initialize the count of intervals in the
    non-overlapping set
    count = 0
    # Initialize the end of the last added interval
    to the minimum possible value
    last_included_end = float('-inf')

```

```

    # Iterate over every interval in the sorted
    list
    for interval in intervals:
        # If the current interval does not overlap
        with the last added interval
        if interval[0] >= last_included_end:
            count += 1 # We include this interval
            last_included_end = interval[1] #
            Update end for subsequent comparisons
        # The minimum number of removals is total
        intervals minus non-overlapping intervals count
    return len(intervals) - count

```



```
# Example usage:
print(eraseOverlapIntervals([[1,2],[2,3],[3,4],[1,3]])) # Expected output: 1
```

- LC.ca

```
class Solution:
    def eraseOverlapIntervals(self, intervals:
List[List[int]]) -> int:
        intervals.sort(key=lambda x: x[1])
        ans, t = 0, intervals[0][1]
        for s, e in intervals[1:]:
            if s >= t:
                t = e
            else:
                ans += 1
        return ans
```

- ◆ Quick Review

Key Insights

- **Sorting is the key:** By sorting intervals by their end times, we can greedily choose the optimal set of non-overlapping intervals.
- **Greedy approach:** Always select the interval that ends the earliest, as it leaves maximum room for subsequent intervals.
- **Counting removals:** The answer is the total number of intervals minus the count of intervals selected by this greedy approach.

- **Edge observation:** Intervals that touch at the boundary are not overlapping, so the check uses “>=” instead of “>”.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(n)$ if considering the storage for the sorted list (or $O(1)$ if sorting in-place).

Solution

The solution involves first sorting the intervals by their end times. Initialize a variable to track the end of the last chosen interval. Iterate over the sorted intervals: if the current interval's start is greater than or equal to the last chosen interval's end, update the last end and include this interval in the non-overlapping set. Otherwise, count the interval as one that must be removed. The final answer is the removal count, which is computed as the total number of intervals minus the count of non-overlapping intervals selected.

Dynamic Programming

#516 Longest Palindromic Subsequence

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Denny Zhang

Problem

Description

Given a string s , find the longest palindromic subsequence's length in s .

A subsequence is a sequence that can be derived from another sequence by deleting some or no elements without changing the order of the remaining elements.

Example 1:

Input: $s = \text{"bbbab"}$

Output: 4

Explanation: One possible longest palindromic subsequence is "bbbb".

Example 2:

Input: $s = \text{"cbdd"}$

Output: 2

Explanation: One possible longest palindromic subsequence is "bb".

Round 8 · Low · Medium

p.2843

Constraints:

- $1 \leq s.length \leq 1000$
- s consists only of lowercase English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def longestPalindromeSubseq(self, s: str) -> int:
        @functools.lru_cache(None)
        def dp(i: int, j: int) -> int:
            """Returns the length of LPS(s[i..j])."""
            if i > j:
                return 0
            if i == j:
                return 1
            if s[i] == s[j]:
                return 2 + dp(i + 1, j - 1)
            return max(dp(i + 1, j), dp(i, j - 1))

        return dp(0, len(s) - 1)
```

Round 8 · Low · Medium

p.2844

```
class Solution:
    def longestPalindromeSubseq(self, s: str) ->
    int:
        n = len(s)
        # dp[i][j] := the length of LPS(s[i..j])
        dp = [[0] * n for _ in range(n)]

        for i in range(n):
            dp[i][i] = 1

        for d in range(1, n):

            for i in range(n - d):
                j = i + d
                if s[i] == s[j]:
                    dp[i][j] = 2 + dp[i + 1][j - 1]
                else:
                    dp[i][j] = max(dp[i + 1][j], dp[i][j - 1])

            return dp[0][n - 1]
```

• LeetCode

Round 8 · Low · Medium

p.2845

```
class Solution:
    def longestPalindromeSubseq(self, s: str) ->
    int:
        n = len(s)
        f = [[0] * n for _ in range(n)]
        for i in range(n):
            f[i][i] = 1
            for i in range(n - 1, -1, -1):
                for j in range(i + 1, n):
                    if s[i] == s[j]:
                        f[i][j] = f[i + 1][j - 1] + 2
                    else:
                        f[i][j] = max(f[i + 1][j],
                        f[i][j - 1])
        return f[0][n - 1]
```

• SimplyLeet

```
# Function to compute the longest palindromic
subsequence length
def longest_palindromic_subseq(s):
    n = len(s)
    # Create a 2D DP table initialized to 0
    dp = [[0] * n for _ in range(n)]

    # Base case: every single character is a
    palindrome of length 1
    for i in range(n):
        dp[i][i] = 1
```

Round 8 · Low · Medium

p.2846

```
# Build the table. 'sub_len' is the current
length of substring we are considering.
for sub_len in range(2, n+1):
    for i in range(n - sub_len + 1):
        j = i + sub_len - 1 # Compute the end
index of the substring
        if s[i] == s[j]:
            # Characters match, add 2 to the
result from the inner substring if available
            dp[i][j] = 2 + (dp[i+1][j-1] if
sub_len > 2 else 0)
        else:
            # Characters don't match, choose
the maximum from either excluding left or right
char
            dp[i][j] = max(dp[i+1][j],
dp[i][j-1])
```

```
# The answer for the full string is stored at
dp[0][n-1]
return dp[0][n-1]
```

```
# Example usage:
s = "bbbab"
print(longest_palindromic_subseq(s)) # Expected
output: 4
```

• [LC.ca](#)

Round 8 · Low · Medium

p.2847

```
class Solution:
    def longestPalindromeSubseq(self, s: str) ->
int:
        n = len(s)
        dp = [[0] * n for _ in range(n)]
        for i in range(n - 1, -1, -1):
            dp[i][i] = 1 # pre-set before inner
loop
            for j in range(i + 1, n):
                if s[i] == s[j]:
                    dp[i][j] = dp[i + 1][j - 1] + 2
                else:
```

```
                    dp[i][j] = max(dp[i + 1][j],
dp[i][j - 1])
            return dp[0][n-1]
```

#####

```
class Solution:
    def longestPalindromeSubseq(self, s: str) ->
int:
        n = len(s)
        dp = [[0] * n for _ in range(n)]
        for i in range(n):
```

Round 8 · Low · Medium

p.2848

```

        dp[i][i] = 1 # cannot move it inside
    'for j in range(1, n)'
        # because below 'for j in range'
starting at 1
        # so cannot put dp[j][j]=1 inside below
for loop
    for j in range(1, n):
        for i in range(j - 1, -1, -1):
            if s[i] == s[j]:
                dp[i][j] = dp[i + 1][j - 1] + 2
                # "aa" => i=0, j=1 => dp[1][0]
is 0, still works
            else:
                dp[i][j] = max(dp[i + 1][j],
dp[i][j - 1])

return dp[0][-1]

```

◆ Quick Review

Key Insights

- The problem can be solved efficiently using dynamic programming.
- A 2D DP table is used where $dp[i][j]$ represents the length of the longest palindromic subsequence in the substring $s[i..j]$.
- If the characters at the two ends of the substring match ($s[i] == s[j]$), they contribute to the palindrome length.

- Otherwise, the solution is the maximum sequence length by either excluding the left or the right character.
- The approach builds the solution from smaller substrings (length 1) to the entire string.

Space & Time

Time Complexity: $O(n^2)$ where n is the length of the string.

Space Complexity: $O(n^2)$ due to the usage of a 2D DP table.

Solution

We solve this problem using a bottom-up dynamic programming approach. We create a 2D table $dp[][]$ such that $dp[i][j]$ stores the longest palindromic subsequence length for the substring $s[i..j]$. For any single character, $dp[i][i]$ is 1. For substrings of length greater than 1, if the characters at positions i and j ($s[i]$ and $s[j]$) match, then $dp[i][j] = dp[i+1][j-1] + 2$. If they do not match, we take the maximum of $dp[i+1][j]$ and $dp[i][j-1]$. Finally, the $dp[0][n-1]$ holds the result for the entire string.

Dynamic Programming

#576 Out of Boundary Paths

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Dynamic Programming

Lists: Denny Zhang

Problem


Description

There is an $m \times n$ grid with a ball. The ball is initially at the position $[startRow, startColumn]$. You are allowed to move the ball to one of the four adjacent cells in the grid (possibly out of the grid crossing the grid boundary). You can apply at most $maxMove$ moves to the ball.


Given the five integers m , n , $maxMove$, $startRow$, $startColumn$, return the number of paths to move the ball out of the grid boundary. Since the answer can be very large, return it modulo $10^9 + 7$.

Example 1:

Move one time:




Move two times:




Input: $m = 2, n = 2, maxMove = 2, startRow = 0, startColumn = 0$
Output: 6

Example 2:


Move one time:



Move two times:



Move three times:



Input: $m = 1$, $n = 3$, $\text{maxMove} = 3$, $\text{startRow} = 0$,
 $\text{startColumn} = 1$
 Output: 12

Constraints:

- $1 \leq m, n \leq 50$
- $0 \leq \text{maxMove} \leq 50$
- $0 \leq \text{startRow} < m$
- $0 \leq \text{startColumn} < n$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def findPaths(self, m, n, maxMove, startRow,
startColumn):
        MOD = 1000000007

        @functools.lru_cache(None)
        def dp(k: int, i: int, j: int) -> int:
            """
            Returns the number of paths to move the ball
            at (i, j) out-of-bounds with
            k moves.
            """

            if i < 0 or i == m or j < 0 or j == n:
                return 1
            if k == 0:
                return 0
            return (dp(k - 1, i + 1, j) + dp(k - 1, i -
1, j) +
                    dp(k - 1, i, j + 1) + dp(k - 1, i, j
- 1)) % MOD

            return dp(maxMove, startRow, startColumn)
```

```

class Solution:
    def findPaths(
        self,
        m: int,
        n: int,
        maxMove: int,
        startRow: int,
        startColumn: int,
    ) -> int:
        DIRS = ((0, 1), (1, 0), (0, -1), (-1, 0))

        MOD = 1_000_000_007
        ans = 0
        # dp[i][j] := the number of paths to move the
        ball (i, j) out-of-bounds
        dp = [[0] * n for _ in range(m)]
        dp[startRow][startColumn] = 1

        for _ in range(maxMove):
            newDp = [[0] * n for _ in range(m)]
            for i in range(m):
                for j in range(n):

```

```

                    if dp[i][j] > 0:
                        for dx, dy in DIRS:
                            x = i + dx
                            y = j + dy
                            if x < 0 or x == m or y < 0 or y ==
n:
                                ans = (ans + dp[i][j]) % MOD
                            else:
                                newDp[x][y] = (newDp[x][y] +
dp[i][j]) % MOD
                        dp = newDp

        return ans

```

• LeetCode

```

class Solution:
    def findPaths(
        self, m: int, n: int, maxMove: int,
        startRow: int, startColumn: int
    ) -> int:
        @cache
        def dfs(i: int, j: int, k: int) -> int:
            if not 0 <= i < m or not 0 <= j < n:
                return int(k >= 0)
            if k <= 0:
                return 0

```



```

        ans = 0
        for a, b in pairwise(dirs):
            x, y = i + a, j + b
            ans = (ans + dfs(x, y, k - 1)) %
mod
        return ans

mod = 10**9 + 7
dirs = (-1, 0, 1, 0, -1)
return dfs(startRow, startColumn, maxMove)

```

• SimplyLeet

```

MOD = 10**9 + 7

def findPaths(m, n, maxMove, startRow,
startColumn):
    # Initialize memo dictionary to cache the
    states
    memo = {}

    def dp(i, j, movesLeft):
        # If out of grid bounds, it's a valid path
        if i < 0 or i >= m or j < 0 or j >= n:
            return 1

```

```

        # No moves left and still inside the grid
        means no valid path out
        if movesLeft == 0:
            return 0
        # Check if we already computed this state
        if (i, j, movesLeft) in memo:
            return memo[(i, j, movesLeft)]

        count = 0
        # Four possible directions: up, down, left,
        right
        count = (dp(i - 1, j, movesLeft - 1) +
                dp(i + 1, j, movesLeft - 1) +
                dp(i, j - 1, movesLeft - 1) +
                dp(i, j + 1, movesLeft - 1)) % MOD

        memo[(i, j, movesLeft)] = count
        return count

    return dp(startRow, startColumn, maxMove)

# Example usage

if __name__ == "__main__":
    print(findPaths(2, 2, 2, 0, 0)) # Expected
    output: 6

```

• LC.ca

```

class Solution:
    def findPaths(
        self, m: int, n: int, maxMove: int,
        startRow: int, startColumn: int
    ) -> int:
        @cache
        def dfs(i, j, k):
            if i < 0 or j < 0 or i >= m or j >= n:
                return 1
            if k <= 0:
                return 0

            res = 0
            for a, b in [[-1, 0], [1, 0], [0, 1],
                [0, -1]]:
                x, y = i + a, j + b
                res += dfs(x, y, k - 1)
            res %= mod
            return res

        mod = 10**9 + 7
        return dfs(startRow, startColumn, maxMove)

```

◆ Quick Review

Key Insights

- Use Dynamic Programming (DP) as the state is defined by (current row, current column, moves remaining).
- Transitions are made from a cell to its four adjacent cells.

- If the ball moves out of bounds, it contributes to one valid path.
- States can be memoized to avoid recomputation and optimize the recursive solution.

Space & Time

Time Complexity: $O(\text{maxMove} * m * n)$, since we compute each state at most once.

Space Complexity: $O(\text{maxMove} * m * n)$ for the memoization table.

Solution

We use a recursive dynamic programming (or memoization) approach. Define a function `dp(i, j, movesLeft)` that returns the number of ways to move out of bounds starting from cell `(i, j)` with `movesLeft` moves remaining. For each move from the current cell, check:

- If the cell is out of bounds, return 1 (base case).
- If no moves are left (`movesLeft == 0`) and we are still inside, return 0.
- Otherwise, use recursion to explore all four directions and sum the paths. Store results in a memoization table to avoid recomputation, then return the computed value modulo $10^9 + 7$.

Dynamic Programming

#1143 Longest Common Subsequence

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Denny Zhang

Problem

Description

Given two strings `text1` and `text2`, return the length of their longest common subsequence. If there is no common subsequence, return 0.

A subsequence of a string is a new string generated from the original string with some characters (can be none) deleted without changing the relative order of the remaining characters.

- For example, "ace" is a subsequence of "abcde".

A common subsequence of two strings is a subsequence that is common to both strings.

Example 1:

Input: `text1 = "abcde"`, `text2 = "ace"`
 Output: 3
 Explanation: The longest common subsequence is "ace" and its length is 3.

Example 2:

Input: `text1 = "abc"`, `text2 = "abc"`
 Output: 3
 Explanation: The longest common subsequence is "abc" and its length is 3.

Example 3:

Input: `text1 = "abc"`, `text2 = "def"`
 Output: 0
 Explanation: There is no such common subsequence, so the result is 0.

Constraints:

- $1 \leq \text{text1.length}, \text{text2.length} \leq 1000$
- `text1` and `text2` consist of only lowercase English characters.

Community Solutions (Python)

- WalkCC

```

class Solution:
    def longestCommonSubsequence(self, text1: str,
text2: str) -> int:
        m = len(text1)
        n = len(text2)
        # dp[i][j] := the length of LCS(text1[0..i],
text2[0..j])
        dp = [[0] * (n + 1) for _ in range(m + 1)]

        for i in range(m):
            for j in range(n):
                dp[i + 1][j + 1] = (1 + dp[i][j] if
text1[i] == text2[j]
                                else max(dp[i][j + 1],
dp[i + 1][j]))

        return dp[m][n]

```

- LeetCode

Round 8 · Low · Medium

p.2863

```

class Solution:
    def longestCommonSubsequence(self, text1: str,
text2: str) -> int:
        m, n = len(text1), len(text2)
        f = [[0] * (n + 1) for _ in range(m + 1)]
        for i in range(1, m + 1):
            for j in range(1, n + 1):
                if text1[i - 1] == text2[j - 1]:
                    f[i][j] = f[i - 1][j - 1] + 1
                else:
                    f[i][j] = max(f[i - 1][j],
f[i][j - 1])

        return f[m][n]

```

- SimplyLeet

```

# Function to find the longest common subsequence
length
def longestCommonSubsequence(text1: str, text2:
str) -> int:
    m, n = len(text1), len(text2)
    # Initialize a (m+1) x (n+1) DP table filled
with zeros
    dp = [[0] * (n + 1) for _ in range(m + 1)]

    # Build the DP table bottom-up
    for i in range(1, m + 1):
        for j in range(1, n + 1):
            # If characters match, extend the
subsequence length by 1

```

Round 8 · Low · Medium

p.2864

```

        if text1[i - 1] == text2[j - 1]:
            dp[i][j] = dp[i - 1][j - 1] + 1
        else:
            # Otherwise, choose the maximum
            # from the left or top cell
            dp[i][j] = max(dp[i - 1][j],
                           dp[i][j - 1])

    return dp[m][n]

# Example usage
if __name__ == "__main__":

    text1 = "abcde"
    text2 = "ace"
    print(longestCommonSubsequence(text1, text2))
# Output: 3

```

- LC.ca

```

class Solution:
    def longestCommonSubsequence(self, text1: str,
                                  text2: str) -> int:
        m, n = len(text1), len(text2)
        f = [[0] * (n + 1) for _ in range(m + 1)]
        for i in range(1, m + 1):
            for j in range(1, n + 1):
                if text1[i - 1] == text2[j - 1]:
                    f[i][j] = f[i - 1][j - 1] + 1
                else:
                    f[i][j] = max(f[i - 1][j],
                                   f[i][j - 1])

        return f[m][n]

```

◆ Quick Review

Key Insights

- The problem is well-suited for a dynamic programming approach.
- Use a 2D table where $dp[i][j]$ stores the length of the longest common subsequence for $text1[0...i-1]$ and $text2[0...j-1]$.
- If the characters match at the current position, extend the subsequence by 1 from $dp[i-1][j-1]$; otherwise, take the maximum from $dp[i-1][j]$ or $dp[i][j-1]$.
- Filling the table in a bottom-up manner ensures that all subproblems are solved optimally.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the lengths of `text1` and `text2`.

Space Complexity: $O(m * n)$ for maintaining the 2D DP table. Space can be optimized to $O(\min(m, n))$ if needed.

Solution

This solution uses a dynamic programming approach with a 2D array (`dp`) to record the lengths of common subsequences at each substring pair. We initialize a table of dimensions $(m+1) \times (n+1)$ with zeros, where m and n are the lengths of `text1` and `text2`. We iterate over both strings, comparing characters; if they match, we set $dp[i][j] = dp[i-1][j-1] + 1$, otherwise we choose the maximum from $dp[i-1][j]$ and $dp[i][j-1]$. The final answer is found at $dp[m][n]$.

Bit Manipulation

Round 8 · Low · Medium · 4 questions

Round 8 · Low · Medium

p.2867

Round 8 · Low · Medium

p.2868

Bit Manipulation

#78 Subsets

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Array · Backtracking · Bit Manipulation

Lists: Grind 169

Problem

Description

Given an integer array `nums` of unique elements, return all possible subsets (the power set).

The solution set must not contain duplicate subsets. Return the solution in any order.

Example 1:

Input: `nums = [1,2,3]`
 Output: `[[], [1], [2], [1,2], [3], [1,3], [2,3], [1,2,3]]`

Example 2:

Input: `nums = [0]`
 Output: `[[], [0]]`

Constraints:

- $1 \leq \text{nums.length} \leq 10$
- $-10 \leq \text{nums}[i] \leq 10$
- All the numbers of `nums` are unique.

Round 8 · Low · Medium

p.2869

Community Solutions (Python)

• WalkCC

```
class Solution:
    def subsets(self, nums: list[int]) -> list[list[int]]:
        ans = []

        def dfs(s: int, path: list[int]) -> None:
            ans.append(path)

            for i in range(s, len(nums)):
                dfs(i + 1, path + [nums[i]])

        dfs(0, [])
        return ans
```

• LeetCode

```
class Solution:
    def subsets(self, nums: List[int]) -> List[List[int]]:
        def dfs(i: int):
            if i == len(nums):
                ans.append(t[:])
                return
            dfs(i + 1)
            t.append(nums[i])
            dfs(i + 1)
            t.pop()
```

Round 8 · Low · Medium

p.2870

```
ans = []
t = []
dfs(0)
return ans
```

```
class Solution:
    def subsets(self, nums: List[int]) ->
List[List[int]]:
    ans = []
    for mask in range(1 << len(nums)):
        t = [x for i, x in enumerate(nums) if
mask >> i & 1]
        ans.append(t)
    return ans
```

- SimplyLeet

Define the function to generate all subsets using backtracking.

```
def subsets(nums):
    # This list will store all the subsets
    result = []

    # Recursive helper function to generate
    subsets.
    def backtrack(start, current):
        # Append a copy of the current subset to
        the final result.
        result.append(current.copy())
        # Iterate over the possible next elements
        starting from index 'start'
```

```
        for i in range(start, len(nums)):
            # Include nums[i] in the current
            subset.
            current.append(nums[i])
            # Recurse with the next starting index.
            backtrack(i + 1, current)
            # Backtrack: remove the last element to
            explore other possibilities.
            current.pop()
```

```
# Start the recursion with an empty path.
backtrack(0, [])
```



```
return result
```

```
# Example usage:
print(subsets([1, 2, 3]))
```

• LC.ca

```
from typing import List
```

```
class Solution: # bfs
```

```
    def subsets(self, nums: List[int]) ->
    List[List[int]]:
        if not nums:
            return [[]]
```

```
        nums.sort() # sort() not necessary if no
        duplicates
        result = [[]]
```

```
        for num in nums:
            result += [subset + [num] for subset in
            result]

        return result
```

```
class Solution: # dfs
```

```
    def subsets(self, nums: List[int]) ->
    List[List[int]]:
        def dfs(u, t):
            ans.append(t[:]) # or, t.copy()
```

```
        for i in range(u, len(nums)):
            t.append(nums[i])
            dfs(i + 1, t)
            t.pop()
```

```
        ans = []
        nums.sort() # sort() not necessary if no
        duplicates
        dfs(0, [])
        return ans
```

```
class Solution: # dfs, just pass down the final
path
```

```
    def subsets(self, nums: List[int]) ->
    List[List[int]]:
        def dfs(nums, index, path, ans):
            ans.append(path)
            [dfs(nums, i + 1, path + [nums[i]],
            ans) for i in range(index, len(nums))]
```

```
        ans = []
        dfs(nums, 0, [], ans)
        return ans
```

```

class Solution: # dfs, but running slower, since
need to reference from parent method for 'nums' and
'ans'
    def subsets(self, nums: List[int]) ->
List[List[int]]:
        def dfs(index, path):
            ans.append(path)
            [dfs(i + 1, path + [nums[i]]) for i in
range(index, len(nums))]

        ans = []
        dfs(0, [])

    return ans

```

◆ Quick Review

Key Insights

- The problem is equivalent to generating the power set of the given set.
- A recursive backtracking approach can be used to explore both including and excluding each element.
- Iterative (bit manipulation) solutions are also possible since there are 2^n subsets in total.
- The maximum array length is small (up to 10), making exponential time solutions feasible.

Space & Time

Time Complexity: $O(2^n * n)$ – There are 2^n subsets and creating each subset may take up to $O(n)$ time.

Space Complexity: $O(2^n * n)$ – Space for storing all subsets. The recursion stack has a maximum depth of $O(n)$.

Solution

We solve the problem using backtracking, where we decide for each element whether to include it in the current subset or not. We recursively build all subsets starting at an empty list and at each step appending the current subset to our result list. We can also use iterative bit manipulation but here we focus on the backtracking approach. This algorithm leverages recursion and a list to maintain the current state, leading to a clear and concise implementation.

Bit Manipulation

#90 Subsets II

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Backtracking · Bit Manipulation

Lists: Denny Zhang

Problem

Description

Given an integer array `nums` that may contain duplicates, return all possible subsets (the power set). The solution set must not contain duplicate subsets. Return the solution in any order.

Example 1:

Input: `nums = [1,2,2]`
 Output: `[[], [1], [1,2], [1,2,2], [2], [2,2]]`

Example 2:

Input: `nums = [0]`
 Output: `[[], [0]]`

Constraints:

- $1 \leq \text{nums.length} \leq 10$
- $-10 \leq \text{nums}[i] \leq 10$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def subsetsWithDup(self, nums: list[int]) -> list[list[int]]:
        ans = []

        def dfs(s: int, path: list[int]) -> None:
            ans.append(path)
            if s == len(nums):
                return

            for i in range(s, len(nums)):

                if i > s and nums[i] == nums[i - 1]:
                    continue
                dfs(i + 1, path + [nums[i]])

        nums.sort()
        dfs(0, [])
        return ans
```

• LeetDoocs

```

class Solution:
    def subsetsWithDup(self, nums: List[int]) ->
List[List[int]]:
    def dfs(i: int):
        if i == len(nums):
            ans.append(t[:])
            return
        t.append(nums[i])
        dfs(i + 1)
        x = t.pop()
        while i + 1 < len(nums) and nums[i + 1]
== x:
            i += 1
            dfs(i + 1)

        nums.sort()
        ans = []
        t = []
        dfs(0)
        return ans

```

```

class Solution:
    def subsetsWithDup(self, nums: List[int]) ->
List[List[int]]:
        nums.sort()
        n = len(nums)
        ans = []
        for mask in range(1 << n):
            ok = True
            t = []
            for i in range(n):
                if mask >> i & 1:
                    if i and (mask >> (i - 1) & 1)
== 0 and nums[i] == nums[i - 1]:
                        ok = False
                        break
                    t.append(nums[i])
            if ok:
                ans.append(t)
        return ans

```

- SimplyLeet

```
def subsetsWithDup(nums):
    # Sort the numbers to handle duplicates
    nums.sort()
    result = [] # To store all subsets
    subset = [] # Current subset being built

    def backtrack(start):
        # Append a copy of the current subset to
        the result
        result.append(list(subset))
        # Iterate over possible elements to include

        for i in range(start, len(nums)):
            # If the current element is a duplicate
            and not at the start of this loop, skip it
            if i > start and nums[i] == nums[i -
1]:
                continue
            # Include current element into the
            subset

            subset.append(nums[i])
            # Recurse with next start index
            backtrack(i + 1)
            # Backtrack and remove the last element
            added

            subset.pop()
```

```
    backtrack(0)
    return result

# Example usage:
print(subsetsWithDup([1, 2, 2]))
```

• LC.ca

```
class Solution:
    def subsetsWithDup(self, nums: List[int]) ->
    List[List[int]]:
        def dfs(u, t):
            ans.append(t[:]) # or, t.copy()
            for i in range(u, len(nums)):
                # also good for [1,5,5], second
                '5' will be skipped here if
                # but second '5' is always
                covered, because i==u when ans=[1,5] and i=2
                if i != u and nums[i] == nums[i -
1]:
                    continue
            t.append(nums[i])
```

```

        dfs(i + 1, t)
        t.pop()

```

```

ans = []
nums.sort()
dfs(0, [])
return ans

```

```

# iteration

```

```

class Solution:
    def subsetsWithDup(self, nums: List[int]) ->
    List[List[int]]:
        if not nums:
            return [[]]

        nums.sort() # Sorting is necessary to
        handle duplicates.
        result = [[]]
        start_idx = 0 # Start index for the new
        subsets
        new_subsets_size = 0

```

```

        for i in range(len(nums)):
            size = len(result)
            # If the current number is the same as
            the previous one,
            # we only want to extend the subsets
            added in the previous iteration.
            if i > 0 and nums[i] == nums[i - 1]:
                start_idx = size -
            new_subsets_size
            else:
                start_idx = 0

            new_subsets_size = 0 # Reset the size
            of newly created subsets

            for j in range(start_idx, size):
                current_subset = result[j] +
                [nums[i]]

                result.append(current_subset)
                new_subsets_size += 1

        return result

```

```

#####

```

```

class Solution:
    def subsetsWithDup(self, nums: List[int]) ->
List[List[int]]:
        nums.sort()
        n = len(nums)
        ans = []
        for mask in range(1 << n):
            ok = True
            t = []
            for i in range(n):
                if mask >> i & 1:
                    if i and (mask >> (i - 1) & 1)
== 0 and nums[i] == nums[i - 1]:
                        ok = False
                        break
                    t.append(nums[i])
            if ok:
                ans.append(t)
        return ans

```

◆ Quick Review

Key Insights

- Sort the array first so that duplicate elements are adjacent.
- Use backtracking to generate all possible subsets.
- Skip duplicate elements in the recursion to avoid forming duplicate subsets.

- The use of sorted order and conditionally skipping elements is the trick to handle duplicates.

Space & Time

Time Complexity: $O(2^n * n)$ since each subset creation may take $O(n)$ in the worst case.

Space Complexity: $O(2^n * n)$ for storing the subsets, plus $O(n)$ for recursion stack.

Solution

The solution uses a backtracking technique that builds subsets incrementally. The input array is first sorted to bring duplicates together. During recursion, if the current element is the same as the previous one and it has not been used in the current recursion path, it is skipped to prevent duplicate subsets. This ensures that each subset is unique. The algorithm stores each valid subset by copying the current subset state into the result list.

Bit Manipulation

#287 Find the Duplicate Number

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Two Pointers · Binary Search · Bit Manipulation

Lists: Grind 169

Problem

Description

Given an array of integers `nums` containing $n + 1$ integers where each integer is in the range $[1, n]$ inclusive.

There is only one repeated number in `nums`, return this repeated number.

You must solve the problem without modifying the array `nums` and using only constant extra space.

Example 1:

Input: `nums = [1,3,4,2,2]`
Output: 2

Example 2:

Input: `nums = [3,1,3,4,2]`
Output: 3

Example 3:

Input: `nums = [3,3,3,3,3]`
Output: 3

Constraints:

- $1 \leq n \leq 10^5$
- `nums.length == n + 1`
- $1 \leq \text{nums}[i] \leq n$
- All the integers in `nums` appear only once except for precisely one integer which appears two or more times.

Follow up:

- How can we prove that at least one duplicate number must exist in `nums`?
- Can you solve the problem in linear runtime complexity?

Community Solutions (Python)

• WalkCC

```
class Solution:
    def findDuplicate(self, nums: list[int]) -> int:
        slow = nums[nums[0]]
        fast = nums[nums[nums[0]]]

        while slow != fast:
            slow = nums[slow]
            fast = nums[nums[fast]]

        slow = nums[0]
```



```

while slow != fast:
    slow = nums[slow]
    fast = nums[fast]

return slow

```

- LeetCode

```

class Solution:
    def findDuplicate(self, nums: List[int]) -> int:
        def f(x: int) -> bool:
            return sum(v <= x for v in nums) > x

        return bisect_left(range(len(nums)), True, key=f)

```

- SimplyLeet

```

# Function to find the duplicate number using
Floyd's Tortoise and Hare algorithm.
def findDuplicate(nums):
    # Phase 1: Find the intersection point inside
    the cycle.
    tortoise = nums[0] # Initialize tortoise
    pointer.
    hare = nums[0] # Initialize hare pointer.
    while True:
        tortoise = nums[tortoise] # Move tortoise
        one step.
        hare = nums[nums[hare]] # Move hare two
        steps.
        if tortoise == hare: # If pointers meet,
        cycle is detected.
            break

```

```

# Phase 2: Find the entrance to the cycle.
pointer = nums[0] # Initialize a new pointer
from the start.
while pointer != tortoise:
    pointer = nums[pointer] # Move pointer one
    step.
    tortoise = nums[tortoise] # Move tortoise
    one step.
    return pointer # The duplicate number.

# Example usage:
if __name__ == "__main__":

```

```
sample = [1, 3, 4, 2, 2]
print(findDuplicate(sample)) # Expected
output: 2
```

• LC.ca

```
class Solution:
    def findDuplicate(self, nums: List[int]) ->
int:
    def f(x: int) -> bool:
        return sum(v <= x for v in nums) > x

    return bisect_left(range(len(nums)), True,
key=f)
```

◆ Quick Review

Key Insights

- Since there are $n+1$ elements with values only from 1 to n , by the pigeonhole principle at least one number must be repeated.
- The problem can be approached by mapping the array values to indices, forming a cycle structure.
- Floyd's Tortoise and Hare algorithm (cycle detection) can be applied to find the duplicate number in linear time and constant space.
- There is no need for extra data structures like hash sets or alterations to the original array, satisfying the space and array-modification constraints.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution leverages Floyd's Tortoise and Hare cycle detection algorithm. Think of each array element as a pointer to another index, forming a linked list with a cycle (the duplicate number creates the cycle). The algorithm works in two phases:

- Detect a meeting point within the cycle by having two pointers (tortoise and hare) move at different speeds.
- Find the cycle entrance by resetting one pointer to the start and moving both pointers at the same speed. The meeting point reveals the duplicate number.

This approach provides an elegant solution that satisfies both linear runtime and constant space requirements.

Bit Manipulation

#1506 Find Root of N-Ary Tree

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Hash Table · Bit Manipulation · Tree · DFS

Lists: Premium 98

Problem


Description

You are given all the nodes of an N-ary tree as an array of Node objects, where each node has a unique value.

Return the root of the N-ary tree.

Custom testing:

An N-ary tree can be serialized as represented in its level order traversal where each group of children is separated by the null value (see examples).



For example, the above tree is serialized as [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14].


The testing will be done in the following way:

1. The input data should be provided as a serialization of the tree.
2. The driver code will construct the tree from the serialized input data and put each Node object into an array in an arbitrary order.
3. The driver code will pass the array to findRoot, and your function should find and return the root Node

object in the array.

4. The driver code will take the returned Node object and serialize it. If the serialized value and the input data are the same, the test passes.

Example 1:



Input: tree = [1,null,3,2,4,null,5,6]

Output: [1,null,3,2,4,null,5,6]

Explanation: The tree from the input data is shown above.

The driver code creates the tree and gives findRoot the Node objects in an arbitrary order.

For example, the passed array could be


[Node(5),Node(4),Node(3),Node(6),Node(2),Node(1)] or

[Node(2),Node(6),Node(1),Node(3),Node(5),Node(4)].

The findRoot function should return the root Node(1), and the driver code will serialize it and compare with the input data.

The input data and serialized Node(1) are the same, so the test passes.

Example 2:



Input: tree = [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14]
 Output: [1,null,2,3,4,5,null,null,6,7,null,8,null,9,10,null,null,11,null,12,null,13,null,null,14]

Constraints:

- The total number of nodes is between [1, 5 * 10⁴].
- Each node has a unique value.

Follow up:

- Could you solve this problem in constant space complexity with a linear time algorithm?

Round 8 · Low · Medium

p.2897

Community Solutions (Python)

• WalkCC

```

class Solution:
    def findRoot(self, tree: list['Node']) -> 'Node':
        sum = 0

        for node in tree:
            sum ^= node.val
            for child in node.children:
                sum ^= child.val

        for node in tree:
            if node.val == sum:
                return node
  
```

• LeetCode

```

"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not None else []
"""

class Solution:
  
```

Round 8 · Low · Medium

p.2898

```
def findRoot(self, tree: List['Node']) ->
'Node':
    x = 0
    for node in tree:
        x ^= node.val
        for child in node.children:
            x ^= child.val
    return next(node for node in tree if
node.val == x)
```

• SimplyLeet

```
# Python implementation with detailed comments
def findRoot(tree):
    # Initialize variables to store the sum of all
    node values and children values
    total_sum = 0
    children_sum = 0

    # Iterate over each node in the tree array
    for node in tree:
        total_sum += node.val # Add current node's
value to total_sum
        # Iterate over each child of the current
node and add their values to children_sum
```

```
        for child in node.children:
            children_sum += child.val

    # The root's value is the difference between
total_sum and children_sum
    root_val = total_sum - children_sum

    # Find and return the node with the value equal
to root_val
    for node in tree:
        if node.val == root_val:
            return node

    return None # This should never happen if the
input tree is valid
```

• LC.ca

```
"""
# Definition for a Node.
class Node:
    def __init__(self, val=None, children=None):
        self.val = val
        self.children = children if children is not
None else []
"""

class Solution:
```

```
def findRoot(self, tree: List['Node']) ->
'Node':
    x = 0
    for node in tree:
        x ^= node.val
        for child in node.children:
            x ^= child.val
    return next(node for node in tree if
node.val == x)
```

◆ Quick Review

Key Insights

- Every node except the root appears as a child of another node.
- By summing all node values and then subtracting the sum of all children node values, the remaining value corresponds to the root.
- This subtraction trick allows us to identify the root in a single pass of the data structure without extra memory for tracking child nodes.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since we iterate through all nodes and their children.

Space Complexity: $O(1)$ extra space using arithmetic accumulators.

Solution

The solution calculates two sums: one for all nodes' values and one for all children nodes' values. The difference between these gives the value of the root node since every node's value (except the root's) is added once as a child. Finally, we iterate over the nodes to find the node with the computed root value and return it. This approach meets the requirement of constant space complexity and linear time.

Math

Round 8 · Low · Medium · 5 questions

Round 8 · Low · Medium

p.2903

Math

#7 Reverse Integer

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math

Lists: Grind 169

Problem

Description

Given a signed 32-bit integer x , return x with its digits reversed. If reversing x causes the value to go outside the signed 32-bit integer range $[-2^{31}, 2^{31} - 1]$, then return 0.

Assume the environment does not allow you to store 64-bit integers (signed or unsigned).

Example 1:

Input: $x = 123$
Output: 321

Example 2:

Input: $x = -123$
Output: -321

Example 3:

Input: $x = 120$
Output: 21

Round 8 · Low · Medium

p.2904

Constraints:

- $-231 \leq x \leq 231 - 1$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def reverse(self, x: int) -> int:
        ans = 0
        sign = -1 if x < 0 else 1
        x *= sign

        while x:
            ans = ans * 10 + x % 10
            x //= 10

        return 0 if ans < -2**31 or ans > 2**31 - 1
        else sign * ans
```

• LeetCode

```
class Solution:
    def reverse(self, x: int) -> int:
        ans = 0
        mi, mx = -(2**31), 2**31 - 1
        while x:
            if ans < mi // 10 + 1 or ans > mx //
10:
                return 0
            y = x % 10
            if x < 0 and y > 0:
                y -= 10

            ans = ans * 10 + y
            x = (x - y) // 10
        return ans
```

• SimplyLeet

```
# Function to reverse an integer
def reverse(x: int) -> int:
    # Define the maximum and minimum limits for a
    32-bit signed integer
    INT_MAX = 2**31 - 1
    INT_MIN = -2**31

    reversed_num = 0 # This will store the reversed
    number

    # Process each digit until x becomes 0
    while x != 0:
```

```

# Pop the last digit from x
# Use modulo 10 to get the last digit
# For negative numbers, modulo may yield
negative remainder
pop = x % 10 if x > 0 else x % -10
# Remove the last digit from x by integer
division
x = x // 10

# Check if appending the digit would cause
overflow
if reversed_num > INT_MAX // 10 or
(reversed_num == INT_MAX // 10 and pop > 7):
    return 0

if reversed_num < INT_MIN // 10 or
(reversed_num == INT_MIN // 10 and pop < -8):
    return 0

# Append the digit to the reversed number
reversed_num = reversed_num * 10 + pop

return reversed_num

# Example usage:
print(reverse(123)) # Output: 321

print(reverse(-123)) # Output: -321
print(reverse(120)) # Output: 21

```

• LC.ca

Round 8 · Low · Medium

p.2907

```

class Solution:
    def reverse(self, x: int) -> int:
        ans = 0
        mi, mx = -(2**31), 2**31 - 1
        while x:
            if ans < mi // 10 + 1 or ans > mx //
10:
                return 0
            y = x % 10
            if x < 0 and y > 0:
                y -= 10

            ans = ans * 10 + y
            x = (x - y) // 10
        return ans

```

◆ Quick Review

Key Insights

- The problem requires reversing the digits while preserving the sign.
- Watch out for integer overflow due to the reversed number exceeding the 32-bit signed integer limits.
- The approach involves repeatedly extracting the last digit from the number and appending it to a new reversed integer.
- Overflow checks are done during each iteration before actually appending the digit.

Space & Time

Round 8 · Low · Medium

p.2908

Time Complexity: $O(n)$, where n is the number of digits in the integer.

Space Complexity: $O(1)$, only constant extra space is used.

Solution

The solution uses a while loop to extract the last digit from the original number using modulo operation ($x \% 10$) and then removes this digit using integer division ($x //= 10$). A reversed number is constructed by multiplying the current reversed number by 10 and adding the extracted digit. Before appending each digit, the algorithm checks if the new reversed number will overflow a 32-bit signed integer by comparing it with the thresholds derived from `INT_MAX` and `INT_MIN`. The algorithm handles negative values by working with x directly since the modulo and division operations work properly with negatives in most languages.

Math

#50 Pow(x, n)

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Math · Recursion

Lists: Grind 169

Problem

Description

Implement `pow(x, n)`, which calculates x raised to the power n (i.e., x^n).

Example 1:

Input: $x = 2.00000$, $n = 10$
Output: `1024.00000`

Example 2:

Input: $x = 2.10000$, $n = 3$
Output: `9.26100`

Example 3:

Input: $x = 2.00000$, $n = -2$
Output: `0.25000`
Explanation: $2^{-2} = 1/2^2 = 1/4 = 0.25$

Constraints:

- $-100.0 < x < 100.0$
- $-231 \leq n \leq 231 - 1$

- n is an integer.
- Either x is not zero or $n > 0$.
- $-104 \leq x^n \leq 104$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def myPow(self, x: float, n: int) -> float:
        if n == 0:
            return 1
        if n < 0:
            return 1 / self.myPow(x, -n)
        if n % 2 == 1:
            return x * self.myPow(x, n - 1)
        return self.myPow(x * x, n // 2)
```

• LeetCode

```
class Solution:
    def myPow(self, x: float, n: int) -> float:
        def qpow(a: float, n: int) -> float:
            ans = 1
            while n:
                if n & 1:
                    ans *= a
                a *= a
                n >>= 1
            return ans
```

```
return qpow(x, n) if n >= 0 else 1 /
qpow(x, -n)
```

• SimplyLeet

```
# Python implementation of pow(x, n) using
exponentiation by squaring
def myPow(x, n):
    # if exponent is negative, invert x and make n
    positive
    if n < 0:
        x = 1 / x
        n = -n

    result = 1.0 # initialize result
    while n:
        # if the current exponent is odd, multiply
        result by the base
```

```
        if n % 2:
            result *= x
        x *= x # square the base for the next
iteration
        n //= 2 # divide the exponent by 2
    return result
```

```
# Example usage:
print(myPow(2.0, 10)) # Output: 1024.0
print(myPow(2.1, 3)) # Output: 9.261000000000001
print(myPow(2.0, -2)) # Output: 0.25
```

• LC.ca

```

class Solution:
    def myPow(self, x: float, n: int) -> float:
        def qpow(a: float, n: int) -> float:
            ans = 1
            while n:
                if n & 1:
                    ans *= a
                a *= a
                n >>= 1
            return ans

        return qpow(x, n) if n >= 0 else 1 /
qpow(x, -n)

```

♦ Quick Review

Key Insights

- Use exponentiation by squaring to reduce the number of multiplications.
- For a negative exponent, compute the power for $|n|$ and then take the reciprocal.
- Handle the edge case when n is 0 (return 1).
- The algorithm works in $O(\log n)$ time complexity.

Space & Time

Time Complexity: $O(\log |n|)$

Space Complexity: $O(1)$ for the iterative solution, or $O(\log |n|)$ if implemented recursively

Solution

We use exponentiation by squaring, an algorithm that reduces the number of multiplications by repeatedly squaring the base and halving the exponent. Here's how it works:

- If n is negative, switch x to $1/x$ and n to $-n$.
- Initialize the result as 1.
- Loop while n is greater than 0: If the current n is odd, multiply the result by x . Square the base x . Divide n by 2 (using integer division). This iterative approach efficiently computes the power in logarithmic steps. The same approach can be implemented recursively.

Math

#380 Insert Delete GetRandom O(1) MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Math · Design

Lists: Grind 169

Problem

Description

Implement the RandomizedSet class:

- RandomizedSet() Initializes the RandomizedSet object.
- bool insert(int val) Inserts an item val into the set if not present. Returns true if the item was not present, false otherwise.
- bool remove(int val) Removes an item val from the set if present. Returns true if the item was present, false otherwise.
- int getRandom() Returns a random element from the current set of elements (it's guaranteed that at least one element exists when this method is called). Each element must have the same probability of being returned.

You must implement the functions of the class such that each function works in average $O(1)$ time complexity.

Example 1:

Input

```
["RandomizedSet", "insert", "remove", "insert",
 "getRandom", "remove", "insert", "getRandom"]
[[], [1], [2], [2], [], [1], [2], []]
```

Output

```
[null, true, false, true, 2, true, false, 2]
```

Explanation

```
RandomizedSet randomizedSet = new RandomizedSet();
```

Constraints:

- $-231 \leq \text{val} \leq 231 - 1$
- At most $2 * 10^5$ calls will be made to insert, remove, and getRandom.
- There will be at least one element in the data structure when getRandom is called.

Community Solutions (Python)

- WalkCC

```
class RandomizedSet:
    def __init__(self):
        self.vals = []
        self.valToIndex = collections.defaultdict(int)
        # {val: index in vals}

    def insert(self, val: int) -> bool:
        if val in self.valToIndex:
            return False
        self.valToIndex[val] = len(self.vals)
        self.vals.append(val)

    def remove(self, val: int) -> bool:
        if val not in self.valToIndex:
            return False
        index = self.valToIndex[val]
        # The order of the following two lines is
        # important when vals.size() == 1.
        self.valToIndex[self.vals[-1]] = index
        del self.valToIndex[val]
        self.vals[index] = self.vals[-1]

    def getRandom(self) -> int:
        index = random.randint(0, len(self.vals) - 1)
        return self.vals[index]
```

Round 8 · Low · Medium

p.2917

• LeetCode

```
class RandomizedSet:
    def __init__(self):
        self.d = {}
        self.q = []

    def insert(self, val: int) -> bool:
        if val in self.d:
            return False
        self.d[val] = len(self.q)
        self.q.append(val)

    def remove(self, val: int) -> bool:
        if val not in self.d:
            return False
        i = self.d[val]
        self.d[self.q[-1]] = i
        self.q[i] = self.q[-1]
        self.q.pop()
        self.d.pop(val)
```

Round 8 · Low · Medium

p.2918

```
return True
```

```
def getRandom(self) -> int:
    return choice(self.q)
```

Your RandomizedSet object will be instantiated and called as such:

```
# obj = RandomizedSet()
# param_1 = obj.insert(val)
# param_2 = obj.remove(val)
```

```
# param_3 = obj.getRandom()
```

• SimplyLeet

```
import random
```

```
class RandomizedSet:
```

```
    def __init__(self):
```

```
        # List to store the values
```

```
        self.values = []
```

```
        # Dictionary to map a value to its index in
the list
```

```
        self.val_to_index = {}
```

```
    def insert(self, val: int) -> bool:
```

```
        # If the value is already present, do not
insert.
```

```
        if val in self.val_to_index:
```

```
            return False
```

```
        # Append the value and store its index.
```

```
        self.val_to_index[val] = len(self.values)
```

```
        self.values.append(val)
```

```
        return True
```

```
    def remove(self, val: int) -> bool:
```

```
        # If the value is not present, return
False.
```

```
        if val not in self.val_to_index:
```

```
            return False
```

```
        # Index of element to remove.
```

```
        index = self.val_to_index[val]
```

```
        # Get the last element.
```

```
        last_element = self.values[-1]
```

```
        # Place the last element at the index of
the element to remove.
```

```
        self.values[index] = last_element
```

```
        self.val_to_index[last_element] = index
```

```
        # Remove the last element from the list.
```



```

        self.values.pop()
        # Delete the mapping for the removed
        element.
        del self.val_to_index[val]
        return True

    def getRandom(self) -> int:
        # Return a random element from the list.
        return random.choice(self.values)

```

• LC.ca

```

'''
my_dict.pop('b') # vs del d[k]

my_dict = {'a': 1, 'b': 2, 'c': 3}
val = my_dict.pop('b')
print(my_dict) # {'a': 1, 'c': 3}
print(val) # 2

>>> item = my_dict.pop()
Traceback (most recent call last):

```

```

File "<stdin>", line 1, in <module>
TypeError: pop expected at least 1 argument,
got 0

```

```

>>> my_dict = {'a': 1, 'b': 2, 'c': 3}
>>> del my_dict['b']
>>> my_dict
{'a': 1, 'c': 3}
'''

```

```

from collections import defaultdict

```

```

from random import choice

```

```

class RandomizedSet(object):
    def __init__(self):
        """
        Initialize your data structure here.
        """

        # value (map) -> its index -> value (list)
        self.d = {}

```

```

    self.a = [] # introduce it purely for random

    def insert(self, val):
        """
        Inserts a value to the set. Returns true if the
        set did not already contain the specified element.
        :type val: int
        :rtype: bool
        """
        if val in self.d:
            return False

        self.a.append(val)
        self.d[val] = len(self.a) - 1
        return True

    def remove(self, val):
        """
        Removes a value from the set. Returns true if
        the set contained the specified element.
        :type val: int
        :rtype: bool
        """

```

```

        if val not in self.d:
            return False
        index = self.d[val]

        # process last index/val
        self.a[index] = self.a[-1]
        self.d[self.a[-1]] = index

        # process to be deleted index/val
        self.a.pop() # delete in list

        del self.d[val] # delete in dict
        # or, self.d.pop(val)
        return True

    def getRandom(self):
        """
        Get a random element from the set.
        :rtype: int
        """
        return self.a[random.randrange(0,
len(self.a))]

```

```
# return random.choice(self.a)
```

```
#####
```

```
# Your RandomizedSet object will be instantiated  
and called as such:
```

```
# obj = RandomizedSet()  
# param_1 = obj.insert(val)  
# param_2 = obj.remove(val)  
# param_3 = obj.getRandom()
```

```
'''  
>>> import random  
>>> random.choice([1,2,3,4,5])  
3  
>>> random.choice([1,2,3,4,5])  
4  
>>> random.choice([1,2,3,4,5])  
4  
'''  
class RandomizedSet:
```

```
def __init__(self):  
    self.m = {}  
    self.l = []
```

```
def insert(self, val: int) -> bool:  
    if val in self.m:  
        return False  
    self.m[val] = len(self.l)  
    self.l.append(val)  
    return True
```

```
def remove(self, val: int) -> bool:  
    if val not in self.m:  
        return False  
    idx = self.m[val]  
    self.l[idx] = self.l[-1]  
    self.m[self.l[-1]] = idx  
    self.l.pop()  
    self.m.pop(val)  
    return True
```

```
def getRandom(self) -> int:
    return random.choice(self.l)
```

```
# Your RandomizedSet object will be instantiated
and called as such:
# obj = RandomizedSet()
# param_1 = obj.insert(val)
# param_2 = obj.remove(val)
# param_3 = obj.getRandom()
```

◆ Quick Review

Key Insights

- Use a dynamic array (or list) to store the elements to allow $O(1)$ access by index.
- Use a hash table (or dictionary/map) to store the mapping from value to its index in the array.
- For removal, swap the element to be removed with the last element in the array, then delete the last element. This avoids expensive shifting of elements.
- `getRandom` can be implemented by generating a random index into the array.

Space & Time

Time Complexity: $O(1)$ on average for insert, remove, and `getRandom` operations.

Space Complexity: $O(n)$ where n is the number of elements stored in the data structure.

Solution

The solution involves maintaining a list to store the values and a hash map to quickly check for the presence of a value and to know its index in the list. When an element is removed, swap it with the last element of the list and update the map accordingly; then, remove the last element in $O(1)$ time. This approach ensures that insert, remove, and `getRandom` can all be achieved in constant time on average.

Math

#1017 Convert to Base -2

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

Math

Lists: Denny Zhang

Problem

Description

Given an integer n , return a binary string representing its representation in base -2 .

Note that the returned string should not have leading zeros unless the string is "0".

Example 1:

Input: $n = 2$
 Output: "110"
 Explanation: $(-2)^2 + (-2)^1 = 2$

Example 2:

Input: $n = 3$
 Output: "111"
 Explanation: $(-2)^2 + (-2)^1 + (-2)^0 = 3$

Example 3:

Input: $n = 4$
 Output: "100"
 Explanation: $(-2)^2 = 4$

Constraints:

- $0 \leq n \leq 109$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def baseNeg2(self, n: int) -> str:
        ans = []

        while n != 0:
            ans.append(str(n % 2))
            n = -(n >> 1)

        return ''.join(reversed(ans)) if ans else '0'
```

- LeetCode

```
class Solution:
    def baseNeg2(self, n: int) -> str:
        k = 1
        ans = []
        while n:
            if n % 2:
                ans.append('1')
                n -= k
            else:
                ans.append('0')

        n //= 2
        k *= -1
        return ''.join(ans[::-1]) or '0'
```

• SimplyLeet

```
# Python solution with detailed comments
def baseNeg2(n: int) -> str:
    # Special case for zero as input.
    if n == 0:
        return "0"

    digits = [] # List to store digits in base -2.

    while n != 0:
        n, remainder = divmod(n, -2) # Get quotient
        and remainder when dividing by -2.
```

```
        if remainder < 0:
            # Adjust the remainder and increment
            the quotient.
            remainder += 2
            n += 1
        # Append the remainder (digit) to the list.
        digits.append(str(remainder))

        # Reverse the digits to form the proper binary
        string and return.
        return "".join(reversed(digits))

# Example usage:
print(baseNeg2(2)) # Output: "110"
print(baseNeg2(3)) # Output: "111"
print(baseNeg2(4)) # Output: "100"
```

• LC.ca

```
class Solution:
    def baseNeg2(self, n: int) -> str:
        k = 1
        ans = []
        while n:
            if n % 2:
                ans.append('1')
                n -= k
            else:
                ans.append('0')
```

```

n //= 2
k *= -1
return ''.join(ans[::-1]) or '0'

```

◆ Quick Review

Key Insights

- Use division and remainder operations to simulate converting a number to a different base.
- When working with a negative base (-2), the remainder can become negative; adjust the remainder (by adding the base's absolute value) and update the quotient accordingly.
- For $n = 0$, simply return "0" since there is nothing to convert.
- Build the answer by collecting digits (remainders) and then reversing the order since the conversion provides digits from least significant to most significant.

Space & Time

Time Complexity: $O(\log|n|)$ – the loop runs for about the number of digits in the base -2 representation.

Space Complexity: $O(\log|n|)$ – extra space for storing the computed digits.

Solution

To solve the problem, we simulate the conversion process similar to how you would convert an integer to a positive base. However, since the base here is -2 , special handling of negative remainders is needed. In each iteration, we:

- Divide the current number by -2 .
- Calculate the remainder. If the remainder is negative, adjust it by adding 2, and increment the quotient to compensate.
- Append the computed remainder to a list (which represents the binary digit).
- Continue until the number reduces to zero. Finally, reverse the list to form the correct string representation from the most significant to the least significant digit.

Math

#3160 Find the Number of Distinct Colors Among the Balls

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Hash Table · Simulation

Lists: CodeSignal

Problem

Description

You are given an integer limit and a 2D array queries of size $n \times 2$.

There are $\text{limit} + 1$ balls with distinct labels in the range $[0, \text{limit}]$. Initially, all balls are uncolored. For every query in queries that is of the form $[x, y]$, you mark ball x with the color y . After each query, you need to find the number of colors among the balls.

Return an array result of length n , where $\text{result}[i]$ denotes the number of colors after i th query.


Note that when answering a query, lack of a color will not be considered as a color.

Example 1:

Input: $\text{limit} = 4$, $\text{queries} = [[1,4],[2,5],[1,3],[3,4]]$

Output: $[1,2,2,3]$

Explanation:




- After query 0, ball 1 has color 4.
- After query 1, ball 1 has color 4, and ball 2 has color 5.
- After query 2, ball 1 has color 3, and ball 2 has color 5.
- After query 3, ball 1 has color 3, ball 2 has color 5, and ball 3 has color 4.

Example 2:

Input: $\text{limit} = 4$, $\text{queries} = [[0,1],[1,2],[2,2],[3,4],[4,5]]$

Output: $[1,2,2,3,4]$

Explanation:



- After query 0, ball 0 has color 1.

- After query 1, ball 0 has color 1, and ball 1 has color 2.
- After query 2, ball 0 has color 1, and balls 1 and 2 have color 2.
- After query 3, ball 0 has color 1, balls 1 and 2 have color 2, and ball 3 has color 4.
- After query 4, ball 0 has color 1, balls 1 and 2 have color 2, ball 3 has color 4, and ball 4 has color 5.

Constraints:

- $1 \leq \text{limit} \leq 109$
- $1 \leq n == \text{queries.length} \leq 105$
- $\text{queries}[i].\text{length} == 2$
- $0 \leq \text{queries}[i][0] \leq \text{limit}$
- $1 \leq \text{queries}[i][1] \leq 109$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def queryResults(self, limit: int, queries:
list[list[int]]) -> list[int]:
        ans = []
        ballToColor = {}
        colorCount = collections.Counter()

        for ball, color in queries:
            if ball in ballToColor:
                prevColor = ballToColor[ball]
                colorCount[prevColor] -= 1
```

```
        if colorCount[prevColor] == 0:
            del colorCount[prevColor]
        ballToColor[ball] = color
        colorCount[color] += 1
        ans.append(len(colorCount))
```

```
    return ans
```

• LeetDoocs

```
class Solution:
    def queryResults(self, limit: int, queries:
List[List[int]]) -> List[int]:
        g = {}
        cnt = Counter()
        ans = []
        for x, y in queries:
            cnt[y] += 1
            if x in g:
                cnt[g[x]] -= 1
                if cnt[g[x]] == 0:
```

```
                    cnt.pop(g[x])
                g[x] = y
            ans.append(len(cnt))
        return ans
```

• SimplyLeet

```
# Python solution with detailed comments

def distinctColors(limit, queries):
    # Dictionary to keep track of the current color
    # of each ball
    ball_color = {}
    # Dictionary to count the frequency of each
    # color applied
    color_count = {}
    # List to store the results after each query
    results = []
```

```
    # Process each query one by one
    for ball, color in queries:
        # If the ball has been colored before,
        # update the old color frequency
        if ball in ball_color:
            old_color = ball_color[ball]
            color_count[old_color] -= 1 #
            # Decrement the count of the old color
            # If the frequency becomes zero, remove
            # the color from the dictionary
            if color_count[old_color] == 0:
                del color_count[old_color]
```

```
        # Update ball's color to the new color
        ball_color[ball] = color
        # Increase the frequency count of the new
        # color
        color_count[color] =
        color_count.get(color, 0) + 1

        # The number of distinct colors is the
        # number of keys in color_count
        results.append(len(color_count))

    return results
```

```
# Example usage:
# limit = 4
# queries = [[1,4],[2,5],[1,3],[3,4]]
# print(distinctColors(limit, queries)) # Output:
# [1,2,2,3]
```

- LC.ca

```

class Solution:
    def queryResults(self, limit: int, queries:
List[List[int]]) -> List[int]:
        g = {}
        cnt = Counter()
        ans = []
        for x, y in queries:
            cnt[y] += 1
            if x in g:
                cnt[g[x]] -= 1
                if cnt[g[x]] == 0:
                    cnt.pop(g[x])
            g[x] = y
            ans.append(len(cnt))
        return ans

```

◆ Quick Review

Key Insights

- Use a hash map to store the current color of each ball that has been updated.
- Use another hash map to store how many balls currently use each color.
- When a ball is recolored, decrement the count of its old color first.
- If an old color count becomes zero, remove it from the color-frequency map.
- Add the new color to the ball and increment its frequency.

- After each query, the number of distinct colors is the number of keys in the color-frequency map.

Space & Time

Time Complexity: $O(q)$ average for q queries, since each query does only $O(1)$ average-time hash map work.

Space Complexity: $O(\min(q, n))$ for colored balls, plus $O(q)$ in the worst case for distinct active colors.

Solution

We solve the problem using two hash maps. The first map stores the current color of each ball, and the second map stores the frequency of each active color. For every query, if the ball already has a color, we decrement the old color's count and remove that color from the frequency map if its count reaches zero. Then we update the ball with the new color and increment the new color's frequency. After processing the query, the number of distinct colors is simply the number of keys remaining in the frequency map.

JavaScript

Round 8 · Low · Medium · 7 questions

Round 8 · Low · Medium

p.2943

JavaScript

#2627 Debounce

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript

Lists: Premium 98

Problem

Description

Given a function fn and a time in milliseconds t , return a debounced version of that function.

A debounced function is a function whose execution is delayed by t milliseconds and whose execution is cancelled if it is called again within that window of time. The debounced function should also receive the passed parameters.

For example, let's say $t = 50\text{ms}$, and the function was called at 30ms , 60ms , and 100ms .

The first 2 function calls would be cancelled, and the 3rd function call would be executed at 150ms .

If instead $t = 35\text{ms}$, The 1st call would be cancelled, the 2nd would be executed at 95ms , and the 3rd would be executed at 135ms .

The above diagram shows how debounce will transform events. Each rectangle represents 100ms and

Round 8 · Low · Medium

p.2944

the debounce time is 400ms. Each color represents a different set of inputs.

Please solve it without using lodash's `_debounce()` function.

Example 1:

```
Input:
t = 50
calls = [
  {"t": 50, inputs: [1]},
  {"t": 75, inputs: [2]}
]
Output: [{"t": 125, inputs: [2]}]
Explanation:
```

Example 2:

```
Input:
t = 20
calls = [
  {"t": 50, inputs: [1]},
  {"t": 100, inputs: [2]}
]
Output: [{"t": 70, inputs: [1]}, {"t": 120, inputs: [2]}]
Explanation:
```

Example 3:

```
Input:
t = 150
calls = [
  {"t": 50, inputs: [1, 2]},
  {"t": 300, inputs: [3, 4]},
  {"t": 300, inputs: [5, 6]}
]
Output: [{"t": 200, inputs: [1,2]}, {"t": 450, inputs: [5, 6]}]
```

Constraints:

- $0 \leq t \leq 1000$
- $1 \leq \text{calls.length} \leq 10$
- $0 \leq \text{calls}[i].t \leq 1000$
- $0 \leq \text{calls}[i].\text{inputs.length} \leq 10$

Community Solutions (JavaScript / TypeScript)

• WalkCC

```
type F = (...args: number[]) => void;

function debounce(fn: F, t: number): F {
  let timeout: ReturnType<typeof setTimeout> | undefined;
  return function (...args) {
    clearTimeout(timeout);
    timeout = setTimeout(() => fn(...args), t);
  };
}
```

• LeetCode

```

type F = (...p: any[]) => any;

function debounce(fn: F, t: number): F {
  let timeout: ReturnType<typeof setTimeout> |
  undefined;

  return function (...args) {
    if (timeout !== undefined) {
      clearTimeout(timeout);
    }
    timeout = setTimeout(() => {

      fn.apply(this, args);
    }, t);
  };
}

/**
 * const log = debounce(console.log, 100);
 * log('Hello'); // cancelled
 * log('Hello'); // cancelled
 * log('Hello'); // Logged at t=100ms
 */

```

• LC.ca

```

type F = (...p: any[]) => any;

function debounce(fn: F, t: number): F {
  let timeout: ReturnType<typeof setTimeout> |
  undefined;

  return function (...args) {
    if (timeout !== undefined) {
      clearTimeout(timeout);
    }
    timeout = setTimeout(() => {

      fn.apply(this, args);
    }, t);
  };
}

/**
 * const log = debounce(console.log, 100);
 * log('Hello'); // cancelled
 * log('Hello'); // cancelled
 * log('Hello'); // Logged at t=100ms
 */

```

◆ Quick Review

Key Insights

- Each call to the debounced function resets the timer.

- Only the most recent function call (that isn't subsequently cancelled) will eventually execute.
- Passing of parameters is handled by storing the arguments of the latest call.
- The solution leverages timers (or equivalent scheduling mechanisms) and cancellation techniques.

Space & Time

Time Complexity: $O(1)$ per call, as each invocation only clears and sets a timer.

Space Complexity: $O(1)$, since only a single timer reference is maintained.

Solution

The solution creates a wrapper function that maintains an internal timer reference. On each call to the debounced function, if a timer exists, it is cancelled. A new timer is then set to execute the original function after t milliseconds with the provided arguments. This ensures that only the last invocation runs if multiple calls occur within the delay period.

JavaScript

#2628 JSON Deep Equal

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript

Lists: Premium 98

Problem

Description

Given two values $o1$ and $o2$, return a boolean value indicating whether two values, $o1$ and $o2$, are deeply equal.

For two values to be deeply equal, the following conditions must be met:

- If both values are primitive types, they are deeply equal if they pass the `===` equality check.
- If both values are arrays, they are deeply equal if they have the same elements in the same order, and each element is also deeply equal according to these conditions.
- If both values are objects, they are deeply equal if they have the same keys, and the associated values for each key are also deeply equal according to these conditions.

You may assume both values are the output of `JSON.parse`. In other words, they are valid JSON.

Please solve it without using
lodash's `_isEqual()` function

Example 1:

Input: `o1 = {"x":1,"y":2}, o2 = {"x":1,"y":2}`
 Output: `true`
 Explanation: The keys and values match exactly.

Example 2:

Input: `o1 = {"y":2,"x":1}, o2 = {"x":1,"y":2}`
 Output: `true`
 Explanation: Although the keys are in a different order, they still match exactly.

Example 3:

Input: `o1 = {"x":null,"L":[1,2,3]}, o2 = {"x":null,"L":["1","2","3"]}`
 Output: `false`
 Explanation: The array of numbers is different from the array of strings.

Example 4:

Input: `o1 = true, o2 = false`
 Output: `false`
 Explanation: `true !== false`

Constraints:

- `1 <= JSON.stringify(o1).length <= 105`
- `1 <= JSON.stringify(o2).length <= 105`
- `maxNestingDepth <= 1000`

Community Solutions (JavaScript / TypeScript)

• WalkCC

```
type JSONValue =
  | null
  | boolean
  | number
  | string
  | JSONValue[]
  | { [key: string]: JSONValue };

function areDeeplyEqual(o1: JSONValue, o2:
JSONValue): boolean {
  if (o1 === o2) {

    return true;
  }
  if (o1 === null || o1 === undefined || o2 ===
null || o2 === undefined) {
    return false;
  }
  if (typeof o1 !== 'object' || typeof o2 !==
'object') {
    return false;
  }
  if (Array.isArray(o1) !== Array.isArray(o2)) {
    return false;
  }
}
```



```

    }
    if (Object.keys(o1).length !==
Object.keys(o2).length) {
        return false;
    }
    for (const key in o1) {
        if (!areDeeplyEqual(o1[key], o2[key])) {
            return false;
        }
    }
    return true;
}

```

```

}

```

• LeetCode

```

function areDeeplyEqual(o1: any, o2: any): boolean
{
    if (o1 === null || typeof o1 !== 'object') {
        return o1 === o2;
    }
    if (typeof o1 !== typeof o2) {
        return false;
    }
    if (Array.isArray(o1) !== Array.isArray(o2)) {
        return false;
    }
}

```

```

if (Array.isArray(o1)) {
    if (o1.length !== o2.length) {
        return false;
    }
    for (let i = 0; i < o1.length; i++) {
        if (!areDeeplyEqual(o1[i], o2[i])) {
            return false;
        }
    }
    return true;
}

```

```

} else {
    const keys1 = Object.keys(o1);
    const keys2 = Object.keys(o2);
    if (keys1.length !== keys2.length) {
        return false;
    }
    for (const key of keys1) {
        if (!areDeeplyEqual(o1[key], o2[key]))
    {
        return false;
    }
}

```

```

    }
    return true;
}
}

```

◆ Quick Review

Key Insights

- Use recursion to compare nested structures.
- For primitive types, a simple equality check (`===` in JavaScript, `==` in Python, etc.) suffices.
- For arrays, verify both have the same length and compare each element recursively.
- For objects, ensure both have the same keys (order does not matter) then recursively compare the corresponding values.
- Handle cases where one operand is null or differs by type early to avoid unnecessary recursion.

Space & Time

Time Complexity: $O(N)$ in the worst-case, where N is the total number of elements/values in the JSON structure.

Space Complexity: $O(N)$ due to the recursion call stack in the worst-case scenario of deeply nested objects or arrays.

Solution

We solve this problem using a recursive approach. The algorithm first checks if both values are strictly equal, which handles primitives immediately. For arrays and objects, it verifies that the sizes (length for arrays and number of keys for objects) match. Then it recurses into each element (for arrays) or checks each key-value pair (for objects) to ensure that they are also deeply equal. For objects, the order of keys is irrelevant, so we compare key sets. This method naturally handles nested structures by delegating equality checks to the same function.

JavaScript

#2630 Memoize

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript

Lists: Premium 98

Problem

Description

Given a function `fn`, return a memoized version of that function.

A memoized function is a function that will never be called twice with the same inputs. Instead it will return a cached value.

You can assume there are 3 possible input functions: `sum`, `fib`, and `factorial`.

- `sum` accepts two integers `a` and `b` and returns `a + b`. Assume that if a value has already been cached for the arguments (`b, a`) where `a !== b`, it cannot be used for the arguments (`a, b`). For example, if the arguments are `(3, 2)` and `(2, 3)`, two separate calls should be made.
- `fib` accepts a single integer `n` and returns `1` if `n <= 1` or `fib(n - 1) + fib(n - 2)` otherwise.
- `factorial` accepts a single integer `n` and returns `1` if `n <= 1` or `factorial(n - 1) * n` otherwise.

Example 1:

Input:
 fnName = "sum"
 actions =
 ["call","call","getCallCount","call","getCallCount"]
 values = [[2,2],[2,2],[],[1,2],[]]
Output: [4,4,1,3,2]
Explanation:
 const sum = (a, b) => a + b;
 const memoizedSum = memoize(sum);

Example 2:

Input:
 fnName = "factorial"
 actions = ["call","call","call","getCallCount","call","getCallCount"]
 values = [[2],[3],[2],[],[3],[]]
Output: [2,6,2,2,6,2]
Explanation:
 const factorial = (n) => (n <= 1) ? 1 : (n * factorial(n - 1));
 const memoFactorial = memoize(factorial);

Example 3:

Input:
 fnName = "fib"
 actions = ["call","getCallCount"]
 values = [[5],[]]
Output: [8,1]
Explanation:
 fib(5) = 8 // "call"
 // "getCallCount" - total call count: 1

Constraints:

- $0 \leq a, b \leq 105$
- $1 \leq n \leq 10$
- $1 \leq \text{actions.length} \leq 105$
- $\text{actions.length} === \text{values.length}$
- $\text{actions}[i]$ is one of "call" and "getCallCount"
- fnName is one of "sum", "factorial" and "fib"

Community Solutions (JavaScript / TypeScript)

• WalkCC

```
type Fn = (...params: number[]) => number;

function memoize(fn: Fn): Fn {
  const cache: Record<string, number> = {};
  return function (...args) {
    const key = args.join(' ');
    return cache[key] === undefined ? (cache[key] =
fn(...args)) : cache[key];
  };
}
```

• LeetDoocs

```

type Fn = (...params: any) => any;

function memoize(fn: Fn): Fn {
  const cache: Record<any, any> = {};

  return function (...args) {
    if (args in cache) {
      return cache[args];
    }
    const result = fn(...args);

    cache[args] = result;
    return result;
  };
}

/**
 * let callCount = 0;
 * const memoizedFn = memoize(function (a, b) {
 *   callCount += 1;
 *   return a + b;
 * })
 * memoizedFn(2, 3) // 5
 * memoizedFn(2, 3) // 5
 * console.log(callCount) // 1
 */

```

◆ Quick Review

Key Insights

- Cache previously computed results using a data structure keyed by the function arguments.
- Ensure the key uniquely represents the argument order; simple tuple (or equivalent) suffices.
- Use closures (or object state) to maintain both the cache and a call count tracking actual function invocations.
- Expose a method (getCallCount) on the memoized function to report the number of times the original function was actually called.
- The memoized version works for functions that might take one or more parameters.

Space & Time

Time Complexity: $O(1)$ average for each call assuming efficient hashing of arguments.

Space Complexity: $O(n)$ for caching n distinct input combinations.

Solution

We use a hash-based cache (a dictionary or map) to store results for every unique set of arguments. Upon each call, we check if the key (constructed from the input arguments) exists in the cache. If not, we increment a counter, call the original function, store the result in the cache, and return it. Otherwise, we return the cached result. The memoized function is augmented with a getCallCount method to report the number of actual calls made to the original function.

JavaScript

#2633 Convert Object to JSON String

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript

Lists: Premium 98

Problem

Description

Given a value, return a valid JSON string of that value. The value can be a string, number, array, object, boolean, or null. The returned string should not include extra spaces. The order of keys should be the same as the order returned by `Object.keys()`.

Please solve it without using the built-in `JSON.stringify` method.

Example 1:

Input: object = {"y":1,"x":2}
 Output: {"y":1,"x":2}
 Explanation:
 Return the JSON representation.
 Note that the order of keys should be the same as the order returned by `Object.keys()`.

Example 2:

Round 8 · Low · Medium

p.2961

Input: object = {"a":"str","b":-12,"c":true,"d":null}
 Output: {"a":"str","b":-12,"c":true,"d":null}
 Explanation:
 The primitives of JSON are strings, numbers, booleans, and null.

Example 3:

Input: object = {"key":{"a":1,"b":[{"},null,"Hello"]}}
 Output: {"key":{"a":1,"b":[{"},null,"Hello"]}}
 Explanation:
 Objects and arrays can include other objects and arrays.

Example 4:

Input: object = true
 Output: true
 Explanation:
 Primitive types are valid inputs.

Constraints:

- value is a valid JSON value
- $1 \leq \text{JSON.stringify(object).length} \leq 105$
- $\text{maxNestingLevel} \leq 1000$
- all strings contain only alphanumeric characters

Community Solutions (JavaScript / TypeScript)

- WalkCC

Round 8 · Low · Medium

p.2962

```

type JSONValue =
  | null
  | boolean
  | number
  | string
  | JSONValue[]
  | { [key: string]: JSONValue };

function jsonStringify(object: JSONValue): string
{
  if (object === null) {

    return 'null';
  }
  if (typeof object === 'boolean' || typeof object
=== 'number') {
    return String(object);
  }
  if (typeof object === 'string') {
    return `"${object}"`;
  }
  if (Array.isArray(object)) {
    const elems = object.map((elem) =>
jsonStringify(elem));

```

```

    return `[${elems.join(',')}]`;
  }
  // typeof object === 'object'
  const pairs = Object.keys(object).map(
    (key) =>
    `"${key}":${jsonStringify(object[key])}`
  );
  return `{${pairs.join(',')}`;
}

```

• LeetCode

```

function jsonStringify(object: any): string {
  if (object === null) {
    return 'null';
  }
  if (typeof object === 'string') {
    return `"${object}"`;
  }
  if (typeof object === 'number' || typeof object
=== 'boolean') {
    return object.toString();
  }
}

```

```

    if (Array.isArray(object)) {
        return
        `[${object.map(jsonStringify).join(',')}]`;
    }
    if (typeof object === 'object') {
        return `${Object.entries(object)
            .map(([key, value]) =>
                `${jsonStringify(key)}:${jsonStringify(value)}`)
            .join(',')}`;
    }
    return '';
}

```

◆ Quick Review

Key Insights

- The solution must support all JSON primitives: strings, numbers, booleans, and null.
- Composite types such as arrays and objects require recursive traversal for proper serializing.
- For strings, wrap them in double quotes.
- For arrays, recursively process every element and join them with commas between square brackets.
- For objects, iterate keys in insertion order and combine key-value pairs (key must be a string in double quotes) with commas between curly braces.
- Use recursion carefully to handle nested structures without adding unnecessary spaces.

Space & Time

Time Complexity: $O(n)$, where n is the total number of elements/characters in the JSON structure.

Space Complexity: $O(n)$ due to the recursion stack and the constructed output string.

Solution

We use a recursive approach that inspects the type of the value and processes it accordingly:

- If the value is null, output "null".
- If the value is a boolean or number, convert it directly to its string representation.
- For a string, ensure it is enclosed in double quotes.
- For an array, iterate over each element, serialize each recursively, join them with commas, and enclose in square brackets.
- For an object, iterate over its keys in insertion order (or using `Object.keys` in JS / `LinkedHashMap` in Java), serialize both key (wrapped in double quotes) and its corresponding value, join with commas, and enclose in curly braces. This approach leverages recursion to handle arbitrarily nested objects and arrays while ensuring there are no extra spaces in the final string.

JavaScript

#2636 Promise Pool

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript · Concurrency

Lists: Premium 98

Problem

Description

Given an array of asynchronous functions `functions` and a pool limit `n`, return an asynchronous function `promisePool`. It should return a promise that resolves when all the input functions resolve.

Pool limit is defined as the maximum number promises that can be pending at once. `promisePool` should begin execution of as many functions as possible and continue executing new functions when old promises resolve. `promisePool` should execute `functions[i]` then `functions[i + 1]` then `functions[i + 2]`, etc. When the last promise resolves, `promisePool` should also resolve.

For example, if `n = 1`, `promisePool` will execute one function at a time in series. However, if `n = 2`, it first executes two functions. When either of the two functions resolve, a 3rd function should be executed (if available), and so on until there are no functions left to execute.

You can assume all functions never reject. It is acceptable for `promisePool` to return a promise that resolves any value.

Example 1:

```
Input:
functions = [
  () => new Promise(res => setTimeout(res, 300)),
  () => new Promise(res => setTimeout(res, 400)),
  () => new Promise(res => setTimeout(res, 200))
]
n = 2
Output: [[300,400,500],500]
```

Example 2:

```
Input:
functions = [
  () => new Promise(res => setTimeout(res, 300)),
  () => new Promise(res => setTimeout(res, 400)),
  () => new Promise(res => setTimeout(res, 200))
]
n = 5
Output: [[300,400,200],400]
```

Example 3:

Round 8 · Low · Medium

p.2967

Round 8 · Low · Medium

p.2968


```

Input:
functions = [
  () => new Promise(res => setTimeout(res, 300)),
  () => new Promise(res => setTimeout(res, 400)),
  () => new Promise(res => setTimeout(res, 200))
]
n = 1
Output: [[300,700,900],900]

```

Constraints:

- $0 \leq \text{functions.length} \leq 10$
- $1 \leq n \leq 10$

Community Solutions (JavaScript / TypeScript)

• WalkCC

```

type F = () => Promise<any>;

function promisePool(functions: F[], n: number):
Promise<any> {
  const next = () =>
functions[n++]?.().then(next);
  return Promise.all(functions.slice(0, n).map((f)
=> f().then(next)));
}

```

• LeetDoocs

```

type F = () => Promise<any>;

function promisePool(functions: F[], n: number):
Promise<any> {
  const wrappers = functions.map(fn => async ()
=> {
    await fn();
    const nxt = waiting.shift();
    nxt && (await nxt());
  });

  const running = wrappers.slice(0, n);

  const waiting = wrappers.slice(n);
  return Promise.all(running.map(fn => fn()));
}

```

• LC.ca

```

type F = () => Promise<any>;

function promisePool(functions: F[], n: number):
Promise<any> {
  const wrappers = functions.map(fn => async ()
=> {
    await fn();
    const nxt = waiting.shift();
    nxt && (await nxt());
  });

  const running = wrappers.slice(0, n);

```

```
const waiting = wrappers.slice(n);
return Promise.all(running.map(fn => fn()));
}
```

◆ Quick Review

Key Insights

- Use a concurrency control mechanism to limit the number of actively running promises.
- Maintain an index or pointer to track which function from the input array should be executed next.
- Upon the resolution of any promise, start a new one if available until all functions have been processed.
- The overall task is complete when there are no functions left and all running promises have resolved.

Space & Time

Time Complexity: $O(m)$, where m is the number of functions to execute.

Space Complexity: $O(n)$ for the active pool of concurrent promises.

Solution

We can solve the problem by maintaining an index that tracks which asynchronous function to execute next. We define a helper routine (or executor) that:

- Checks if there is a function left to run.
- If yes, increments the index and awaits the result of running that function.
- Once the function completes, the routine

recursively calls itself to pick up the next function. We start n such routines (to respect the pool limit) and use `Promise.all` (or the equivalent in other languages) to wait until all concurrent routines have finished executing. This structure ensures that we never run more than n promises concurrently and that a new promise starts as soon as one finishes.

JavaScript

#2675 Array of Objects to Matrix

MED

LeetCode · SimplyLeet · WalkCC · LeetCode

JavaScript · Array

Lists: Premium 98

Problem

Description

Write a function that converts an array of objects `arr` into a matrix `m`.

`arr` is an array of objects or arrays. Each item in the array can be deeply nested with child arrays and child objects. It can also contain numbers, strings, booleans, and null values.

The first row `m` should be the column names. If there is no nesting, the column names are the unique keys within the objects. If there is nesting, the column names are the respective paths in the object separated by ".".

Each of the remaining rows corresponds to an object in `arr`. Each value in the matrix corresponds to a value in an object. If a given object doesn't contain a value for a given column, the cell should contain an empty string "".

The columns in the matrix should be in lexicographically ascending order.

Example 1:

```
Input:
arr = [
  {"b": 1, "a": 2},
  {"b": 3, "a": 4}
]
Output:
[
  ["a", "b"],
```

Example 2:

```
Input:
arr = [
  {"a": 1, "b": 2},
  {"c": 3, "d": 4},
  {}
]
Output:
[
```

Example 3:

```
Input:
arr = [
  {"a": {"b": 1, "c": 2}},
  {"a": {"b": 3, "d": 4}}
]
Output:
[
  ["a.b", "a.c", "a.d"],
```

Example 4:

Input:
 arr = [
 [{"a": null}],
 [{"b": true}],
 [{"c": "x"}]
]
 Output:
 [

Example 5:

Input:
 arr = [
 {},
 {},
 {},
]
 Output:
 [

Constraints:

- arr is a valid JSON array
- $1 \leq \text{arr.length} \leq 1000$
- unique keys ≤ 1000

Community Solutions (JavaScript / TypeScript)

- WalkCC

```
function jsonToMatrix(arr: any[]): (string |
number | boolean | null)[][] {
  const isObject = (o: any) => o !== null && typeof
o === 'object';
```

// Returns the keys of a JSON-like object by recursively unwrapping the nests.

```
const getKeys = (json: any): string[] => {
  if (!isObject(json)) {
    return [''];
  }
  return Object.keys(json).reduce((acc:
string[], currKey: string) => {
    return (
```

```
      acc.push(
        ...getKeys(json[currKey]).map((nextKey:
string) =>
          nextKey === '' ? currKey :
`${currKey}.${nextKey}`
        ),
      acc
    );
  }, []);
};
```

```

const sortedKeys: string[] = [
  ...arr.reduce((acc: Set<string>, curr: any) =>
{
  getKeys(curr).forEach((key: string) =>
acc.add(key));
  return acc;
}, new Set<string>()),
].sort();

// Returns the value of `obj` keyed by
`nestedKey`.
const getValue = (
  obj: any,

```

```

  nestedKey: string
): string | number | boolean | null => {
  let value: any = obj;
  for (const key of nestedKey.split('.')) {
    if (!isObject(value) || !(key in value)) {
      return '';
    }
    value = value[key];
  }
  return isObject(value) ? '' : value;

```

```

  };

  const matrix: (string | number | boolean |
null)[][] = [sortedKeys];
  arr.forEach((obj: any) => {
    matrix.push(
      sortedKeys.map((nestedKey: string) =>
getValue(obj, nestedKey))
    );
  });

  return matrix;
}

```

• LeetCode

```

function jsonToMatrix(arr: any[]): (string |
number | boolean | null)[] {
  const dfs = (key: string, obj: any) => {
    if (
      typeof obj === 'number' ||
      typeof obj === 'string' ||
      typeof obj === 'boolean' ||
      obj === null
    ) {
      return { [key]: obj };
    }
  }

```

```

    const res: any[] = [];
    for (const [k, v] of Object.entries(obj)) {
        const newKey = key ? `${key}.${k}` :
`${k}`;
        res.push(dfs(newKey, v));
    }
    return res.flat();
};

const kv = arr.map(obj => dfs('', obj));
const keys = [

```

```

    ...new Set(
        kv
            .flat()
            .map(obj => Object.keys(obj))
            .flat(),
    ),
].sort();
const ans: any[] = [keys];
for (const row of kv) {
    const newRow: any[] = [];

```

```

        for (const key of keys) {
            const v = row.find(r =>
r.hasOwnProperty(key))?.[key];
            newRow.push(v === undefined ? '' : v);
        }
        ans.push(newRow);
    }
    return ans;
}

```

• LC.ca

```

function jsonToMatrix(arr: any[]): (string |
number | boolean | null)[] {
    const dfs = (key: string, obj: any) => {
        if (
            typeof obj === 'number' ||
            typeof obj === 'string' ||
            typeof obj === 'boolean' ||
            obj === null
        ) {
            return { [key]: obj };
        }
    }
}

```

```

    const res: any[] = [];
    for (const [k, v] of Object.entries(obj)) {
        const newKey = key ? `${key}.${k}` :
`${k}`;
        res.push(dfs(newKey, v));
    }
    return res.flat();
};

const kv = arr.map(obj => dfs('', obj));
const keys = [

```

```

    ...new Set(
        kv
            .flat()
            .map(obj => Object.keys(obj))
            .flat(),
    ),
].sort();
const ans: any[] = [keys];
for (const row of kv) {
    const newRow: any[] = [];

```

```

        for (const key of keys) {
            const v = row.find(r =>
r.hasOwnProperty(key))?.[key];
            newRow.push(v === undefined ? '' : v);
        }
        ans.push(newRow);
    }
    return ans;
}

```

◆ Quick Review

Key Insights

- Flatten each object or array into dot-separated paths such as `a.b` or `a.0.c`.
- Use DFS so nested objects and nested arrays are handled uniformly.
- For each row, store a map from flattened path to value.
- Collect every path seen across all rows into a set of column names.
- Sort the column names lexicographically to build the header row.
- For each flattened row map, output values in header order and use an empty string when a column is missing.

Space & Time

Time Complexity: $O(T + R * C \log C)$, where T is the total number of nested nodes visited, R is the number of rows, and C is the number of unique flattened columns.

Space Complexity: $O(R * C + C)$ in the worst case for the flattened row maps and the set of all column names.

Solution

We solve the problem by flattening every input row first. A DFS walks through objects and arrays, extending the current path with either the property name or the array index. When the DFS reaches a primitive value or `null`, it stores that value under the built path. While flattening each row, we also collect every path into a global set of column names. After all rows are processed, we sort the columns lexicographically, place them in the first row of the matrix, and then build each remaining row by reading values from that row's flattened map, using `""` for missing paths.

Round 8 · Low · Medium

p.2983

JavaScript

#2700 Differences Between Two Objects

MED

LeetDoocs · SimplyLeet · WalkCC · LeetCode

JavaScript

Lists: Premium 98

Problem

Description

Write a function that accepts two deeply nested objects or arrays `obj1` and `obj2` and returns a new object representing their differences.

The function should compare the properties of the two objects and identify any changes. The returned object should only contains keys where the value is different from `obj1` to `obj2`.

For each changed key, the value should be represented as an array [`obj1` value, `obj2` value]. Keys that exist in one object but not in the other should not be included in the returned object. The end result should be a deeply nested object where each leaf value is a difference array.

When comparing two arrays, the indices of the arrays are considered to be their keys.

You may assume that both objects are the output of `JSON.parse`.

Round 8 · Low · Medium

p.2984

Example 1:

```

Input:
obj1 = {}
obj2 = {
  "a": 1,
  "b": 2
}
Output: {}
Explanation: There were no modifications made to obj1.
New keys "a" and "b" appear in obj2, but keys that are
added or removed should be ignored.

```

Example 2:

```

Input:
obj1 = {
  "a": 1,
  "v": 3,
  "x": [],
  "z": {
    "a": null
  }
}

```

Example 3:

```

Input:
obj1 = {
  "a": 5,
  "v": 6,
  "z": [1, 2, 4, [2, 5, 7]]
}
obj2 = {
  "a": 5,

```

Example 4:

```

Input:
obj1 = {
  "a": {"b": 1},
}
obj2 = {
  "a": [5],
}
Output:

```

Example 5:

```

Input:
obj1 = {
  "a": [1, 2, {}],
  "b": false
}
obj2 = {
  "b": false,
  "a": [1, 2, {}]
}

```

Constraints:

- obj1 and obj2 are valid JSON objects or arrays
- $2 \leq \text{JSON.stringify(obj1).length} \leq 104$
- $2 \leq \text{JSON.stringify(obj2).length} \leq 104$

Community Solutions (JavaScript / TypeScript)

- WalkCC

```

type JSONValue =
  | null
  | boolean
  | number
  | string
  | JSONValue[]
  | { [key: string]: JSONValue };
type Obj = Record<string, JSONValue> |
Array<JSONValue>;

function objDiff(obj1: Obj, obj2: Obj): Obj {

  if (obj1 === obj2) {
    return {};
  }
  if (obj1 === null || obj2 === null) {
    return [obj1, obj2];
  }
  if (typeof obj1 !== 'object' || typeof obj2 !==
'object') {
    return [obj1, obj2];
  }
  if (Array.isArray(obj1) !== Array.isArray(obj2))
{

```

```

    return [obj1, obj2];
  }
  const ans = {};
  for (const key in obj1) {
    if (key in obj2) {
      const subDiff = objDiff(obj1[key],
obj2[key]);
      if (Object.keys(subDiff).length > 0) {
        ans[key] = subDiff;
      }
    }
  }
  return ans;
}

```

- LeetCode

```
function objDiff(obj1: any, obj2: any): any {
  if (type(obj1) !== type(obj2)) return [obj1,
obj2];
  if (!isObject(obj1)) return obj1 === obj2 ? {}
: [obj1, obj2];
  const diff: Record<string, unknown> = {};
  const sameKeys = Object.keys(obj1).filter(key
=> key in obj2);
  sameKeys.forEach(key => {
    const subDiff = objDiff(obj1[key],
obj2[key]);
    if (Object.keys(subDiff).length) diff[key]
= subDiff;
  });
  return diff;
}
```

```
}

function type(obj: unknown): string {
  return
Object.prototype.toString.call(obj).slice(8, -1);
}

function isObject(obj: unknown): obj is
Record<string, unknown> {
  return typeof obj === 'object' && obj !== null;
}
```

◆ Quick Review

Round 8 · Low · Medium

p.2989

Key Insights

- Use recursion to traverse nested objects and arrays.
- Compare only keys that exist in both objects.
- When encountering two values of different types or unequal primitives, record the difference as [value from obj1, value from obj2].
- When both values are objects or arrays, recursively compute their differences and include them only if the resulting nested diff is non-empty.
- Arrays are treated like objects with indices as keys.
- The solution requires careful type checking and handling of recursive structures.

Space & Time

Time Complexity: $O(n)$ in average, where n is the number of keys/elements in the common parts of the structures.

Space Complexity: $O(n)$ due to recursion and storage needed for the output differences.

Solution

We solve the problem using a recursive function that compares keys present in both input objects. For each common key:

- If both values are objects (or both are arrays), recursively compute their differences.
- If the recursive call returns a non-empty object, include that key in the result.
- If the two values are of different types or are primitives that differ, add the key with a difference array [value from obj1, value from obj2].
- Ignore keys that exist only in one object. This recursive approach naturally handles the

Round 8 · Low · Medium

p.2990

deeply nested structure while ensuring only common keys with different values are captured.

Round 9 · Low · Hard

Round 9

Low Priority · Hard

7 questions · 2 patterns

Dynamic Programming #10 #115 #123 #132 #312 #1000

Math #68

Round 8 · Low · Medium

p.2991

Round 8 · Low · Medium

p.2992

Dynamic Programming

Round 9 · Low · Hard · 6 questions

Round 9 · Low · Hard

p.2993

Dynamic Programming

#10 Regular Expression Matching

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming · Recursion

Lists: Denny Zhang

Problem

Description

Given an input string *s* and a pattern *p*, implement regular expression matching with support for '.' and '*' where:

- '.' Matches any single character.
- '*' Matches zero or more of the preceding element.

Return a boolean indicating whether the matching covers the entire input string (not partial).

Example 1:

Input: *s* = "aa", *p* = "a"

Output: false

Explanation: "a" does not match the entire string "aa".

Example 2:

Round 9 · Low · Hard

p.2994

Input: s = "aa", p = "a*"

Output: true

Explanation: '*' means zero or more of the preceding element, 'a'. Therefore, by repeating 'a' once, it becomes "aa".

Example 3:

Input: s = "ab", p = ".*"

Output: true

Explanation: ".*" means "zero or more (*) of any character (.)".

Constraints:

- $1 \leq s.length \leq 20$
- $1 \leq p.length \leq 20$
- s contains only lowercase English letters.
- p contains only lowercase English letters, '.', and '*'.
- It is guaranteed for each appearance of the character '*', there will be a previous valid character to match.

Community Solutions (Python)

- WalkCC

```
class Solution:
```

```
    def isMatch(self, s: str, p: str) -> bool:
```

```
        m = len(s)
```

```
        n = len(p)
```

```
        # dp[i][j] := True if s[0..i) matches p[0..j)
```

```
        dp = [[False] * (n + 1) for _ in range(m + 1)]
```

```
        dp[0][0] = True
```

```
    def isMatch(i: int, j: int) -> bool:
```

```
        return j >= 0 and p[j] == '.' or s[i] == p[j]
```

```
    for j, c in enumerate(p):
```

```
        if c == '*' and dp[0][j - 1]:
```

```
            dp[0][j + 1] = True
```

```
    for i in range(m):
```

```
        for j in range(n):
```

```
            if p[j] == '*':
```

```
                # The minimum index of '*' is 1.
```

```
                noRepeat = dp[i + 1][j - 1]
```

```
                doRepeat = isMatch(i, j - 1) and dp[i][j + 1]
```

```
                dp[i + 1][j + 1] = noRepeat or doRepeat
```

```
            elif isMatch(i, j):
```

```
                dp[i + 1][j + 1] = dp[i][j]
```

```
    return dp[m][n]
```

- LeetCode

```

class Solution:
    def isMatch(self, s: str, p: str) -> bool:
        @cache
        def dfs(i, j):
            if j >= n:
                return i == m
            if j + 1 < n and p[j + 1] == '*':
                return dfs(i, j + 2) or (
                    i < m and (s[i] == p[j] or p[j]
                        == '.') and dfs(i + 1, j)
                )
            return i < m and (s[i] == p[j] or p[j]
                == '.') and dfs(i + 1, j + 1)

        m, n = len(s), len(p)
        return dfs(0, 0)

```

```

class Solution:
    def isMatch(self, s: str, p: str) -> bool:
        m, n = len(s), len(p)
        f = [[False] * (n + 1) for _ in range(m + 1)]

        f[0][0] = True
        for i in range(m + 1):
            for j in range(1, n + 1):
                if p[j - 1] == '*':
                    f[i][j] = f[i][j - 2]
                    if i > 0 and (p[j - 2] == "."
                        or s[i - 1] == p[j - 2]):

```

```

                    f[i][j] |= f[i - 1][j]
                elif i > 0 and (p[j - 1] == "." or
                    s[i - 1] == p[j - 1]):
                    f[i][j] = f[i - 1][j - 1]
            return f[m][n]

```

• SimplyLeet

```

# Python solution using recursion with memoization
def isMatch(s: str, p: str) -> bool:
    memo = {}

    def dp(i: int, j: int) -> bool:
        # If result is already computed, return it.
        if (i, j) in memo:
            return memo[(i, j)]

        # If we have reached the end of the
        pattern, s must also be at the end.

```

```

    if j == len(p):
        ans = i == len(s)
    else:
        # Check if current characters match
        first_match = i < len(s) and (p[j] ==
s[i] or p[j] == '.')

        # If there is a '*' coming after the
current pattern character
        if j + 1 < len(p) and p[j+1] == '*':
            # Two cases:
            # 1. '*' stands for zero occurrence
            -> skip current pattern char and '*'

            # 2. first_match is true, so we try
to consume one character in s and remain at same
pattern index
            ans = dp(i, j+2) or (first_match
and dp(i+1, j))
            else:
                # No '*' means we directly check
next characters in both s and p.
                ans = first_match and dp(i+1, j+1)

        memo[(i, j)] = ans
        return ans

return dp(0, 0)

```

Round 9 · Low · Hard

p.2999

```

# Example test run
#print(isMatch("aa", "a*")) # Expected output:
True

```

• LC.ca

```

class Solution:
    def isMatch(self, s: str, p: str) -> bool:
        m, n = len(s), len(p)
        f = [[False] * (n + 1) for _ in range(m +
1)]

        f[0][0] = True
        for i in range(m + 1):
            for j in range(1, n + 1):
                if p[j - 1] == "*":
                    f[i][j] = f[i][j - 2]
                    if i > 0 and (p[j - 2] == "."
or s[i - 1] == p[j - 2]):
                        f[i][j] |= f[i - 1][j]
                    elif i > 0 and (p[j - 1] == "." or
s[i - 1] == p[j - 1]):
                        f[i][j] = f[i - 1][j - 1]
                return f[m][n]

```

◆ Quick Review

Key Insights

Round 9 · Low · Hard

p.3000

- Use recursion or dynamic programming to simulate matching as the '*' operator can affect the current and following characters.
- The '.' operator is a wildcard that matches any single character.
- When encountering a '*' in the pattern, consider both skipping the preceding element or consuming one character from the string if it matches.
- Using memoization helps to save intermediate results and efficiently process overlapping subproblems.

Space & Time

Time Complexity: $O(mn)$

where $m = \text{length}(s)$ and $n = \text{length}(p)$, due to exploring each subproblem combination.

Space Complexity: $O(mn)$

for storing the memoization table in the dynamic programming or recursive approach.

Solution

The solution uses a top-down recursive strategy with memoization to efficiently match the string against the pattern. For each character position in s and p , we check if they match directly or if the pattern character is '.' to allow any match. When encountering '*' in the pattern, two possibilities are explored: either the '*' stands for zero occurrences of the preceding element, or if there is a valid match, it accounts for one occurrence and we proceed further in the string while maintaining the same pattern index. The memoization aspect prevents repeated

computations of the same (i, j) state in the recursion, resulting in an improved overall runtime.

Dynamic Programming

#115 Distinct Subsequences

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Denny Zhang

Problem

Description

Given two strings s and t , return the number of distinct subsequences of s which equals t .

The test cases are generated so that the answer fits on a 32-bit signed integer.

Example 1:

```
Input: s = "rabbbit", t = "rabbit"
Output: 3
Explanation:
As shown below, there are 3 ways you can generate
"rabbit" from s.
rabbbit
rabbbit
rabbbit
```

Example 2:

```
Input: s = "babgbag", t = "bag"
```

```
Output: 5
```

```
Explanation:
```

As shown below, there are 5 ways you can generate "bag" from s .

```
babgbag
babgbag
babgbag
babgbag
babgbag
```

Constraints:

- $1 \leq s.length, t.length \leq 1000$
- s and t consist of English letters.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def numDistinct(self, s: str, t: str) -> int:
        m = len(s)
        n = len(t)
        dp = [[0] * (n + 1) for _ in range(m + 1)]

        for i in range(m + 1):
            dp[i][0] = 1

        for i in range(1, m + 1):
```

```

for j in range(1, n + 1):
    if s[i - 1] == t[j - 1]:
        dp[i][j] = dp[i - 1][j - 1] + dp[i - 1][j]
    else:
        dp[i][j] = dp[i - 1][j]

return dp[m][n]

```

• LeetCode

```

class Solution:
    def numDistinct(self, s: str, t: str) -> int:
        m, n = len(s), len(t)
        f = [[0] * (n + 1) for _ in range(m + 1)]
        for i in range(m + 1):
            f[i][0] = 1
        for i, a in enumerate(s, 1):
            for j, b in enumerate(t, 1):
                f[i][j] = f[i - 1][j]
                if a == b:
                    f[i][j] += f[i - 1][j - 1]
        return f[m][n]

```

```

class Solution:
    def numDistinct(self, s: str, t: str) -> int:
        n = len(t)
        f = [1] + [0] * n
        for a in s:
            for j in range(n, 0, -1):
                if a == t[j - 1]:
                    f[j] += f[j - 1]
        return f[n]

```

• SimplyLeet

```

# Python solution for Distinct Subsequences
def numDistinct(s: str, t: str) -> int:
    # Get lengths of the strings s and t
    m, n = len(s), len(t)
    # Initialize a dp table with (m+1) rows and (n+1) columns filled with 0
    dp = [[0] * (n + 1) for _ in range(m + 1)]

    # Base case: An empty string t (j = 0) is a subsequence of any prefix of s, so set dp[i][0] = 1
    for i in range(m + 1):
        dp[i][0] = 1

```

```

# Fill the dp table
for i in range(1, m + 1):
    for j in range(1, n + 1):
        # If the current characters match, add
        # the count of sequences by either using or skipping
        # s[i-1]
        if s[i - 1] == t[j - 1]:
            dp[i][j] = dp[i - 1][j - 1] + dp[i - 1][j]
        else:
            # If they do not match, skip s[i-1]
            dp[i][j] = dp[i - 1][j]

# The answer is the way to form all of t from
# all of s
return dp[m][n]

# Example usage:
print(numDistinct("rabbbit", "rabbit")) # Output:
3

```

• LC.ca

Round 9 · Low · Hard

p.3007

```

class Solution:
    def numDistinct(self, s: str, t: str) -> int:
        n = len(t)
        f = [1] + [0] * n
        for a in s:
            for j in range(n, 0, -1):
                if a == t[j - 1]:
                    f[j] += f[j - 1]
        return f[n]

```

◆ Quick Review

Key Insights

- Use dynamic programming to build the solution by comparing characters of s and t.
- Define a 2D dp table where $dp[i][j]$ represents the number of ways to form the first j characters of t using the first i characters of s.
- The recurrence relation:
 - If $s[i-1]$ equals $t[j-1]$, then include both the cases of matching this character as well as skipping it: $dp[i][j] = dp[i-1][j-1] + dp[i-1][j]$.
 - If they do not match, then simply skip the current character of s: $dp[i][j] = dp[i-1][j]$.
- Initialize $dp[i][0] = 1$ for all i since an empty t can always be formed.
- The final answer is $dp[s.length][t.length]$.

Space & Time

Round 9 · Low · Hard

p.3008

Time Complexity: $O(n * m)$, where n is the length of s and m is the length of t .

Space Complexity: $O(n * m)$, which can be optimized to $O(m)$ with careful row reuse.

Solution

We use a dynamic programming approach with a 2D table. The table has dimensions $(\text{len}(s) + 1) \times (\text{len}(t) + 1)$. Each cell $\text{dp}[i][j]$ stores the number of distinct subsequences from the first i characters of s that match the first j characters of t .

Steps:

- Initialize $\text{dp}[i][0] = 1$ for every i , because an empty t is a subsequence of any prefix of s .
- Iterate over each character of s (outer loop) and for each, iterate over t (inner loop).
- For every character in s , if it equals the current character in t , update $\text{dp}[i][j]$ as the sum of the ways by using or skipping this character. Otherwise, carry over the previous count.
- Return $\text{dp}[\text{len}(s)][\text{len}(t)]$ as the result.

Dynamic Programming

#123 Best Time to Buy and Sell Stock III

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Denny Zhang

Problem

Description

You are given an array `prices` where `prices[i]` is the price of a given stock on the i th day.

Find the maximum profit you can achieve. You may complete at most two transactions.

Note: You may not engage in multiple transactions simultaneously (i.e., you must sell the stock before you buy again).

Example 1:

Input: `prices = [3,3,5,0,0,3,1,4]`

Output: 6

Explanation: Buy on day 4 (price = 0) and sell on day 6 (price = 3), profit = $3 - 0 = 3$.

Then buy on day 7 (price = 1) and sell on day 8 (price = 4), profit = $4 - 1 = 3$.

Example 2:

Python solution for Best Time to Buy and Sell Stock III

```
def max_profit(prices):
    # If there are less than 2 days, no transaction
    # can be made.
    if not prices or len(prices) < 2:
        return 0

    n = len(prices)
    # Initialize left and right profit arrays with
    # zeros.
    left_profits = [0] * n

    right_profits = [0] * n

    # Forward pass: compute max profit until day i
    # with one transaction.
    min_price = prices[0]
    for i in range(1, n):
        # Update the min price encountered so far.
        min_price = min(min_price, prices[i])
        # Maximum profit if sold on day i.
        left_profits[i] = max(left_profits[i - 1],
                             prices[i] - min_price)
```

```
    # Backward pass: compute max profit from day i
    # to end with one transaction.
    max_price = prices[-1]
    for i in range(n - 2, -1, -1):
        # Update max price seen from the right.
        max_price = max(max_price, prices[i])
        # Maximum profit possible starting from day
        # i.
        right_profits[i] = max(right_profits[i +
        1], max_price - prices[i])

    # Combine the two profits to find the maximum
    # total profit.
    max_total_profit = 0

    for i in range(n):
        max_total_profit = max(max_total_profit,
                               left_profits[i] + right_profits[i])

    return max_total_profit

# Example usage:
prices = [3,3,5,0,0,3,1,4]
print(max_profit(prices)) # Expected output: 6
```

• LC.ca

```

'''
- `f1`  []
- `f2`  []
- `f3`  []
- `f4`  []
'''
class Solution:
    def maxProfit(self, prices: List[int]) -> int:
        # []
        f1, f2, f3, f4 = -prices[0], 0, -prices[0],
0

        for price in prices[1:]:
            f1 = max(f1, -price)
            f2 = max(f2, f1 + price)
            f3 = max(f3, f2 - price)
            f4 = max(f4, f3 + price)
        return f4

class Solution:
    def maxProfit(self, prices: List[int]) -> int:

```

```

        if not prices:
            return 0

        # 0 to i, max profit
        left_max = [0] * len(prices)

        tmp_min = prices[0]
        for i in range(len(prices)):
            tmp_min = min(tmp_min, prices[i])
            left_max[i] = max(left_max[i-1],
prices[i]-tmp_min) if i > 0 else 0

        # i to end, max profit
        right_max = [0] * len(prices)

        tmp_max = prices[-1]
        for i in range(len(prices)-1, -1, -1):
            tmp_max = max(tmp_max, prices[i])
            right_max[i] = max(right_max[i+1],
tmp_max-prices[i]) if i < len(prices)-1 else 0

        return max(left_max[i] + right_max[i] for
i in range(len(prices)))

```



```
#####
```

```
class Solution(object):
    def maxProfit(self, prices):
        """
        :type prices: List[int]
        :rtype: int
        """
        buy1 = buy2 = float("-inf")

        sell1 = sell2 = 0

        for i in range(len(prices)):
            sell1 = max(prices[i] + buy1, sell1)
            buy1 = max(buy1, -prices[i])
            sell2 = max(sell2, prices[i] + buy2)
            buy2 = max(sell1 - prices[i], buy2)
        return max(sell1, sell2)
```

◆ Quick Review

Key Insights

- The problem requires the optimal selection of at most two buy-sell transactions.
- One effective strategy is to split the problem into two parts:
 - For each day i , calculate the maximum profit obtainable from day 0 to i with one transaction.

- For each day i , calculate the maximum profit obtainable from day i to the end with one transaction.
- Finally, for each day, the sum of these two profits (left and right) gives a candidate answer.
- Alternatively, a state-based dynamic programming approach can track 4 states corresponding to the two transactions.

Space & Time

Time Complexity: $O(n)$ where n is the number of days

Space Complexity: $O(n)$ due to extra arrays used for forward and backward computations (this can be optimized to $O(1)$ with a state variable approach)

Solution

We use a two-pass dynamic programming approach. In the first pass (forward), we compute an array "leftProfits" where $\text{leftProfits}[i]$ is the maximum profit from day 0 to day i with one transaction. We do this by keeping track of the minimum price observed so far and calculating the profit if sold on day i . In the second pass (backward), we compute an array "rightProfits" where $\text{rightProfits}[i]$ is the maximum profit obtainable from day i to the last day with one transaction by tracking the maximum price seen in reverse order. Finally, for every index i , the sum $\text{leftProfits}[i] + \text{rightProfits}[i]$ represents the maximum profit achievable by completing the first transaction on or before day i and the second transaction starting from day i . The answer is the maximum sum over all i .

Dynamic Programming

#132 Palindrome Partitioning II

HAR

LeetCode · SimplyLeet · WalkCC · LeetCode

String · Dynamic Programming

Lists: Denny Zhang

Problem

Description

Given a string s , partition s such that every substring of the partition is a palindrome.

Return the minimum cuts needed for a palindrome partitioning of s .

Example 1:

Input: $s = \text{"aab"}$

Output: 1

Explanation: The palindrome partitioning $[\text{"aa"}, \text{"b"}]$ could be produced using 1 cut.

Example 2:

Input: $s = \text{"a"}$

Output: 0

Example 3:

Input: $s = \text{"ab"}$

Output: 1

Constraints:

- $1 \leq s.length \leq 2000$
- s consists of lowercase English letters only.

Community Solutions (Python)

• WalkCC

```
class Solution:
    def minCut(self, s: str) -> int:
        n = len(s)
        # isPalindrome[i][j] := True if s[i..j] is a
        # palindrome
        isPalindrome = [[True] * n for _ in range(n)]
        # dp[i] := the minimum cuts needed for a
        # palindrome partitioning of s[0..i]
        dp = [n] * n

        for l in range(2, n + 1):
            i = 0
```

```

    for j in range(l - 1, n):
        isPalindrome[i][j] = s[i] == s[j] and
isPalindrome[i + 1][j - 1]
        i += 1

```

```

for i in range(n):
    if isPalindrome[0][i]:
        dp[i] = 0
        continue

```

Try all the possible partitions.

```

for j in range(i):
    if isPalindrome[j + 1][i]:
        dp[i] = min(dp[i], dp[j] + 1)

return dp[-1]

```

• LeetCode

```

class Solution:
    def minCut(self, s: str) -> int:
        n = len(s)
        g = [[True] * n for _ in range(n)]
        for i in range(n - 1, -1, -1):
            for j in range(i + 1, n):
                g[i][j] = s[i] == s[j] and g[i +
1][j - 1]
        f = list(range(n))
        for i in range(1, n):
            for j in range(i + 1):

```

```

        if g[j][i]:
            f[i] = min(f[i], 1 + f[j - 1])
    if j else 0)
    return f[-1]

```

• SimplyLeet

```

def minCut(s: str) -> int:
    n = len(s)
    # Precompute palindrome table where
    isPal[i][j] is True if s[i:j+1] is a palindrome.
    isPal = [[False] * n for _ in range(n)]
    # dp[i] will store the minimum cuts required
    for substring s[0...i]
    dp = [0] * n
    for i in range(n):
        min_cut = i # maximum cuts (cut every
character)
        for j in range(i + 1):
            # Check if s[j...i] is a palindrome:

```

```

        if s[j] == s[i] and (i - j < 2 or
isPal[j+1][i-1]):
            isPal[j][i] = True # Mark
substring as palindrome.
            # If the palindrome starts at index
0, no cut is needed.
            min_cut = 0 if j == 0 else
min(min_cut, dp[j-1] + 1)
            dp[i] = min_cut # Record minimum cuts
needed for s[0...i]
            return dp[-1]

# Example usage:
# print(minCut("aab")) # Output: 1

```

• LC.ca

```

class Solution:
    def minCut(self, s: str) -> int:
        pal = [[False for j in range(0, len(s))]
for i in range(0, len(s))]
        dp = [len(s) for _ in range(0, len(s) + 1)]
        for i in range(0, len(s)):
            for j in range(0, i + 1):
                # also ok if, (j + 1 >= i - 1)
                if (s[i] == s[j]) and ((j + 1 > i -
1) or (pal[i - 1][j + 1])):
                    pal[i][j] = True
                    dp[i + 1] = min(dp[i + 1],
dp[j] + 1) if j != 0 else 0

```

```

        # 'if j != 0' to ensure from
start to i is a palindrome
        return dp[-1]

#####

class Solution: # clear cache
    def minCut(self, s: str) -> int:
        @cache
        def dfs(i):
            if i >= n - 1:

                return 0
            ans = inf
            for j in range(i, n):
                if g[i][j]:
                    ans = min(ans, dfs(j + 1) + (j
< n - 1))
            return ans

            n = len(s)
            g = [[True] * n for _ in range(n)]
            for i in range(n - 1, -1, -1):

```

```

        for j in range(i + 1, n):
            g[i][j] = s[i] == s[j] and g[i +
1][j - 1]
        ans = dfs(0)
        dfs.cache_clear() # nice, clear cache !!!
        return ans

#####

class Solution:
    def minCut(self, s: str) -> int:

        n = len(s)
        is_pal = [[False] * n for _ in range(n)]
        for i in range(n - 1, -1, -1):
            for j in range(i, n):
                is_pal[i][j] = s[i] == s[j] and (j
- i < 3 or is_pal[i + 1][j - 1])
            min_cut = [i for i in range(n)]
            for i in range(n):
                for j in range(0, i + 1):
                    if is_pal[j][i]:
                        min_cut[i] = min(min_cut[i],
min_cut[j - 1] + 1) if j > 0 else 0

        return min_cut[-1]

```

◆ Quick Review

Key Insights

- Precompute a 2D table (isPal) where isPal[i][j] indicates if the substring s[i...j] is a palindrome.
- Use dynamic programming: dp[i] represents the minimum cuts needed for the substring s[0...i].
- For every index i, iterate over all possible starting indices j. If s[j...i] is a palindrome, update dp[i] using dp[j-1] + 1 (or 0 if j is 0).
- Time complexity is dominated by the O(n²) palindrome checking and state transitions.

Space & Time

Time Complexity: O(n²)

Space Complexity: O(n²) (due to the 2D table for palindrome checks, plus O(n) for the dp array)

Solution

The solution relies on dynamic programming with a two-step approach. First, precompute a 2D table (isPal) where each cell indicates if a substring is a palindrome. This is done in O(n²) time by leveraging the fact that a substring s[j...i] is a palindrome if the characters at the ends match and the inner substring s[j+1...i-1] is also a palindrome (or if the length is less than 3). Then, use a dp array where dp[i] holds the minimum number of cuts needed for the substring s[0...i]. For each i, check every index j (0 ≤ j ≤ i) and if s[j...i] is a palindrome, update dp[i] to be the minimum of its current value and dp[j-1] + 1 (with a special case when j is 0, meaning no cut is needed). This two-phase approach efficiently breaks down the problem and computes the answer while avoiding

redundant palindrome checks.

Round 9 · Low · Hard

p.3027

Dynamic Programming

#312 Burst Balloons

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Denny Zhang

Problem

Description

You are given n balloons, indexed from 0 to $n - 1$. Each balloon is painted with a number on it represented by an array `nums`. You are asked to burst all the balloons.

If you burst the i th balloon, you will get $\text{nums}[i - 1] * \text{nums}[i] * \text{nums}[i + 1]$ coins. If $i - 1$ or $i + 1$ goes out of bounds of the array, then treat it as if there is a balloon with a 1 painted on it.

Return the maximum coins you can collect by bursting the balloons wisely.

Example 1:

Input: `nums = [3,1,5,8]`

Output: 167

Explanation:

`nums = [3,1,5,8] --> [3,5,8] --> [3,8] --> [8] --> []`

`coins = 3*1*5 + 3*5*8 + 1*3*8 + 1*8*1 = 167`

Example 2:

Round 9 · Low · Hard

p.3028

Input: nums = [1,5]

Output: 10

Constraints:

- $n == \text{nums.length}$
- $1 \leq n \leq 300$
- $0 \leq \text{nums}[i] \leq 100$

Community Solutions (Python)

• WalkCC

```
class Solution:
    def maxCoins(self, nums: list[int]) -> int:
        n = len(nums)
        nums = [1] + nums + [1]

        @functools.lru_cache(None)
        def dp(i: int, j: int) -> int:
            """Returns maxCoins(nums[i..j])."""
            if i > j:
                return 0

            return max(dp(i, k - 1) +
                       dp(k + 1, j) +
                       nums[i - 1] * nums[k] * nums[j + 1]
                       for k in range(i, j + 1))

        return dp(1, n)
```

```
class Solution:
    def maxCoins(self, nums: list[int]) -> int:
        n = len(nums)
        # dp[i][j] := maxCoins(nums[i..j])
        dp = [[0] * (n + 2) for _ in range(n + 2)]

        nums = [1] + nums + [1]

        for d in range(n):
            for i in range(1, n - d + 1):
                j = i + d
                for k in range(i, j + 1):
                    dp[i][j] = max(
                        dp[i][j],
                        dp[i][k - 1] +
                        dp[k + 1][j] +
                        nums[i - 1] * nums[k] * nums[j + 1])

        return dp[1][n]
```

• LeetCode

```
class Solution:
    def maxCoins(self, nums: List[int]) -> int:
        n = len(nums)
        arr = [1] + nums + [1]
        f = [[0] * (n + 2) for _ in range(n + 2)]
        for i in range(n - 1, -1, -1):
            for j in range(i + 2, n + 2):
                for k in range(i + 1, j):
                    f[i][j] = max(f[i][j], f[i][k]
+ f[k][j] + arr[i] * arr[k] * arr[j])
        return f[0][-1]
```

• SimplyLeet

```
# Define the function to compute maximum coins
def maxCoins(nums):
    # Add virtual balloons with value 1 at both
    boundaries
    n = len(nums)
    extended = [1] + nums + [1]
    # Initialize dp table with 0's; dimensions:
    (n+2) x (n+2)
    dp = [[0] * (n + 2) for _ in range(n + 2)]

    # Calculate dp values for intervals with length
    2 to n+1
    for length in range(2, n + 2): # interval
    length from 2 to n+1
```

```
        for left in range(0, n + 2 - length):
            right = left + length
            # Iterate through all possible last
            burst balloons in the interval (left, right)
            for k in range(left + 1, right):
                # Coins when bursting balloon k
                last in the interval:
                # value from bursting:
                extended[left] * extended[k] * extended[right]
                # plus coins from subintervals
                [left, k] and [k, right]
                coins = extended[left] *
                extended[k] * extended[right] + dp[left][k] +
                dp[k][right]
                dp[left][right] =
                max(dp[left][right], coins)
```

```
# The final answer is stored in dp[0][n+1]
which represents bursting all real balloons.
return dp[0][n+1]

# Example test cases
print(maxCoins([3,1,5,8])) # Expected output: 167
print(maxCoins([1,5])) # Expected output: 10
```

• LC.ca


```

class Solution:
    def maxCoins(self, nums: List[int]) -> int:
        nums = [1] + nums + [1]
        n = len(nums)
        dp = [[0] * n for _ in range(n)]
        for l in range(2, n):
            for i in range(n - l):
                j = i + l
                for k in range(i + 1, j):
                    dp[i][j] = max(
                        dp[i][j], dp[i][k] +
dp[k][j] + nums[i] * nums[k] * nums[j]
                    )
        return dp[0][-1]

```

◆ Quick Review

Key Insights

- Transform the problem by adding virtual balloons with value 1 at both ends to simplify edge handling.
- Use interval dynamic programming where $dp[i][j]$ represents the maximum coins that can be obtained by bursting all balloons between indices i and j (exclusive).
- Realize that the order of bursting balloons matters; choose the last balloon to burst in each interval to break dependencies into independent subproblems.
- The solution involves iterating over all possible intervals and determining the optimal burst order.

Space & Time

Time Complexity: $O(n^3)$ – Three nested loops: one for the interval length, one for the starting index, and one for the final balloon choice.

Space Complexity: $O(n^2)$ – A 2D dynamic programming table is used to store the maximum coins for each interval.

Solution

The problem is solved using interval dynamic programming. First, extend the input array by adding 1 at both ends to handle edge cases seamlessly. Define a dp table where $dp[i][j]$ represents the maximum coins that can be obtained by bursting all the balloons between i and j (not including i and j , as these are the virtual boundaries). For every interval $[i, j]$, iterate through the possible choices for the last balloon k to burst in that interval. The value for bursting balloon k will be determined by the coins from bursting it last (i.e., $nums[i] * nums[k] * nums[j]$) plus the maximum coins from the two resulting subintervals: (i, k) and (k, j) . By filling the dp table in increasing interval lengths, you build up to the solution for the overall interval that covers the entire (extended) array.

Dynamic Programming

#1000 Minimum Cost to Merge Stones

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · Dynamic Programming

Lists: Denny Zhang

Problem

Description

There are n piles of stones arranged in a row. The i th pile has $\text{stones}[i]$ stones.

A move consists of merging exactly k consecutive piles into one pile, and the cost of this move is equal to the total number of stones in these k piles.

Return the minimum cost to merge all piles of stones into one pile. If it is impossible, return -1 .

Example 1:

Input: $\text{stones} = [3, 2, 4, 1]$, $k = 2$

Output: 20

Explanation: We start with $[3, 2, 4, 1]$.

We merge $[3, 2]$ for a cost of 5, and we are left with $[5, 4, 1]$.

We merge $[4, 1]$ for a cost of 5, and we are left with $[5, 5]$.

We merge $[5, 5]$ for a cost of 10, and we are left with $[10]$.

The total cost was 20, and this is the minimum possible.

Example 2:

Input: $\text{stones} = [3, 2, 4, 1]$, $k = 3$

Output: -1

Explanation: After any merge operation, there are 2 piles left, and we can't merge anymore. So the task is impossible.

Example 3:

Input: $\text{stones} = [3, 5, 1, 2, 6]$, $k = 3$

Output: 25

Explanation: We start with $[3, 5, 1, 2, 6]$.

We merge $[5, 1, 2]$ for a cost of 8, and we are left with $[3, 8, 6]$.

We merge $[3, 8, 6]$ for a cost of 17, and we are left with $[17]$.

The total cost was 25, and this is the minimum possible.

Constraints:

- $n == \text{stones.length}$

- $1 \leq n \leq 30$
- $1 \leq \text{stones}[i] \leq 100$
- $2 \leq k \leq 30$

Community Solutions (Python)

• LeetCode

```
class Solution:
    def mergeStones(self, stones: List[int], K:
int) -> int:
        n = len(stones)
        if (n - 1) % (K - 1):
            return -1
        s = list(accumulate(stones, initial=0))
        f = [[[inf] * (K + 1) for _ in range(n +
1)] for _ in range(n + 1)]
        for i in range(1, n + 1):
            f[i][i][1] = 0
        for l in range(2, n + 1):
            for i in range(1, n - l + 2):
                j = i + l - 1
                for k in range(1, K + 1):
                    for h in range(i, j):
                        f[i][j][k] =
min(f[i][j][k], f[i][h][1] + f[h + 1][j][k - 1])
                        f[i][j][1] = f[i][j][K] + s[j] -
s[i - 1]
            return f[1][n][1]
```

• SimplyLeet

```
# Python solution for merging stones with detailed
comments.
class Solution:
    def mergeStones(self, stones: List[int], k:
int) -> int:
        n = len(stones)
        # Check for merge feasibility.
        if (n - 1) % (k - 1) != 0:
            return -1

        # Precompute prefix sums for fast range sum
queries.
        prefix = [0] * (n + 1)

        for i in range(n):
            prefix[i + 1] = prefix[i] + stones[i]

        # Initialize dp table, dp[i][j] represents
minimal cost to merge stones[i...j].
        dp = [[0] * n for _ in range(n)]

        # Process all possible subarrays with
lengths >= k.
        for length in range(k, n + 1):
            for i in range(n - length + 1):
                j = i + length - 1
```

```

        dp[i][j] = float('inf')
        # Try all valid partitions with
step size (k-1).
        for mid in range(i, j, k - 1):
            dp[i][j] = min(dp[i][j],
dp[i][mid] + dp[mid+1][j])
            # If this segment can be merged
into one pile, add cost.
            if (length - 1) % (k - 1) == 0:
                dp[i][j] += prefix[j+1] -
prefix[i]
        return dp[0][n-1]

```

Example usage:

```

# solution = Solution()
# print(solution.mergeStones([3, 2, 4, 1], 2)) #
Expected output: 20

```

• LC.ca

Round 9 · Low · Hard

p.3039

```

class Solution:
    def mergeStones(self, stones: List[int], K:
int) -> int:
        n = len(stones)
        if (n - 1) % (K - 1):
            return -1
        s = list(accumulate(stones, initial=0))
        f = [[[inf] * (K + 1) for _ in range(n +
1)] for _ in range(n + 1)]
        for i in range(1, n + 1):
            f[i][i][1] = 0
            for l in range(2, n + 1):
                for i in range(1, n - l + 2):
                    j = i + l - 1
                    for k in range(1, K + 1):
                        for h in range(i, j):
                            f[i][j][k] =
min(f[i][j][k], f[i][h][1] + f[h + 1][j][k - 1])
                            f[i][j][1] = f[i][j][K] + s[j] -
s[i - 1]
                return f[1][n][1]

```

◆ Quick Review

Key Insights

- You can only merge stones into one final pile if $(n - 1)$ is divisible by $(k - 1)$. If not, return -1 .
- A dynamic programming approach is used where $dp[i][j]$ represents the minimum cost to merge the subarray

Round 9 · Low · Hard

p.3040

stones[i...j] into as few piles as possible.

- A prefix sum array is precomputed to quickly calculate the sum of stones in any subarray.
- When a segment can be merged into one pile (i.e. when $(\text{length} - 1) \% (k - 1) == 0$), add the sum of stones in that segment to the dp value.

Space & Time

Time Complexity: $O(n^3)$ due to the triple nested loop over segments and splits.

Space Complexity: $O(n^2)$ from the dp table and $O(n)$ for the prefix sum array.

Solution

The problem is solved using a range DP approach. We first verify that merging all piles into one is feasible by checking if $(n - 1) \% (k - 1) == 0$. Next, a prefix sum array is computed to optimize range sum calculations. A 2D dp table $dp[i][j]$ is then used to store the minimum cost to merge stones[i...j]. For each interval, we consider valid split points, updating $dp[i][j]$ by combining sub-problems. If the current segment can form a single merged pile (determined by the condition on the interval length), the cost of merging that segment (the sum of stones) is added to $dp[i][j]$.

Math

Round 9 · Low · Hard · 1 question

Round 9 · Low · Hard

p.3041

Round 9 · Low · Hard

p.3042

Math

#68 Text Justification

HAR

LeetDoocs · SimplyLeet · WalkCC · LeetCode

Array · String · Simulation

Lists: CodeSignal

Problem

Description

Given an array of strings `words` and a width `maxWidth`, format the text such that each line has exactly `maxWidth` characters and is fully (left and right) justified.

You should pack your words in a greedy approach; that is, pack as many words as you can in each line. Pad extra spaces ' ' when necessary so that each line has exactly `maxWidth` characters.

Extra spaces between words should be distributed as evenly as possible. If the number of spaces on a line does not divide evenly between words, the empty slots on the left will be assigned more spaces than the slots on the right.

For the last line of text, it should be left-justified, and no extra space is inserted between words.

Note:

- A word is defined as a character sequence consisting of non-space characters only.

- Each word's length is guaranteed to be greater than 0 and not exceed `maxWidth`.
- The input array `words` contains at least one word.

Example 1:

```
Input: words = ["This", "is", "an", "example", "of",
               "text", "justification."], maxWidth = 16
Output:
[
  "This    is    an",
  "example  of text",
  "justification. "
]
```

Example 2:

```
Input: words =
["What", "must", "be", "acknowledgment", "shall", "be"],
maxWidth = 16
Output:
[
  "What   must   be",
  "acknowledgment ",
  "shall be      "
]
Explanation: Note that the last line is "shall be "
instead of "shall be", because the last line must be
left-justified instead of fully-justified.
```

Example 3:

Round 9 · Low · Hard

p.3043

Round 9 · Low · Hard

p.3044

Input: words = ["Science","is","what","we","understand",
 ,"well","enough","to","explain","to","a","computer.", "A
 rt","is","everything","else","we","do"], maxWidth = 20
 Output:

```
[
  "Science is what we",
  "understand      well",
  "enough to explain to",
  "a computer. Art is",
  "everything else we",
```

Constraints:

- $1 \leq \text{words.length} \leq 300$
- $1 \leq \text{words}[i].\text{length} \leq 20$
- $\text{words}[i]$ consists of only English letters and symbols.
- $1 \leq \text{maxWidth} \leq 100$
- $\text{words}[i].\text{length} \leq \text{maxWidth}$

Community Solutions (Python)

- WalkCC

```
class Solution:
    def fullJustify(self, words: list[str],
                    maxWidth: int) -> list[str]:
        ans = []
        row = []
        rowLetters = 0

        for word in words:
            # If we place the word in this row, it will
            # exceed the maximum width.
            # Therefore, we cannot put the word in this
            # row and have to pad spaces
            # for each word in this row.

            if rowLetters + len(word) + len(row) >
            maxWidth:
                for i in range(maxWidth - rowLetters):
                    row[i % (len(row) - 1 or 1)] += ' '
                ans.append(''.join(row))
                row = []
                rowLetters = 0
            row.append(word)
            rowLetters += len(word)

        return ans + [' '.join(row).ljust(maxWidth)]
```

- LeetCode

```

class Solution:
    def fullJustify(self, words: List[str],
maxWidth: int) -> List[str]:
    ans = []
    i, n = 0, len(words)
    while i < n:
        t = []
        cnt = len(words[i])
        t.append(words[i])
        i += 1
        while i < n and cnt + 1 + len(words[i])
<= maxWidth:

```

```

        cnt += 1 + len(words[i])
        t.append(words[i])
        i += 1
        if i == n or len(t) == 1:
            left = ' '.join(t)
            right = ' ' * (maxWidth -
len(left))
            ans.append(left + right)
            continue
        space_width = maxWidth - (cnt - len(t)
+ 1)
        w, m = divmod(space_width, len(t) - 1)

```

```

        row = []
        for j, s in enumerate(t[:-1]):
            row.append(s)
            row.append(' ' * (w + (1 if j < m
else 0)))
        row.append(t[-1])
        ans.append(''.join(row))
    return ans

```

• SimplyLeet

```

def fullJustify(words, maxWidth):
    result = [] # Final list of justified text
    lines.
    current_words = [] # Current line's words.
    current_len = 0 # Total length of current
    line's words (excluding spaces).

    # Process all words.
    for word in words:
        # Check if adding the next word would
        exceed the maxWidth when including minimum one
        space between words.
        if current_words and (current_len +
len(word) + len(current_words)) > maxWidth:
            total_spaces = maxWidth - current_len
            # Total spaces to distribute.

```



```

        gaps = len(current_words) - 1 # Number
of gaps between words.

        # If there is only one word, left
justify.
        if gaps == 0:
            line = current_words[0] + ' ' *
total_spaces
        else:
            # Calculate even spaces and extras
to add from the left.
            space_per_gap, extra =
divmod(total_spaces, gaps)
            line = ''
            for i in range(gaps):

```

```

        # Distribute an extra space to
the leftmost gaps.
        spaces = space_per_gap + (1 if
i < extra else 0)
        line += current_words[i] + ' '
* spaces
        line += current_words[-1] # Append
the last word.
        result.append(line)
        # Reset the current line to start with
the current word.
        current_words = []
        current_len = 0

        # Add the word to the current line.

```

```

        current_words.append(word)
        current_len += len(word)

        # Process the last line, which is
left-justified.
        line = ' '.join(current_words)
        line += ' ' * (maxWidth - len(line))
        result.append(line)
        return result

# Example usage:

```

```
words = ["This", "is", "an", "example", "of",
"text", "justification."]
maxWidth = 16
print(fullJustify(words, maxWidth))
```

- LC.ca

```
'''
```

Explanation:

- The function `fullJustify` takes a list of words and a maximum width `maxWidth` as inputs.
- It iterates over each word, deciding whether to add the current word to the current line (`cur`) or to start a new line.
- If adding the current word to the line would exceed `maxWidth`, it justifies the current line by adding extra spaces between words as needed, then starts a new line.
- Once all words are processed, it left-justifies the last line by joining the remaining words in `cur` with a single space and then using `.ljust(maxWidth)` to ensure the line is of maximum width.
- The result is a list of strings, where each string represents a justified line of text.

This solution carefully handles edge cases, such as when there's only one word in the line (avoiding division by zero) and ensuring the last line is left-justified instead of fully justified.

```

...

...

>>> cur = [1]
>>> len(cur) - 1
0
>>> ( len(cur) - 1 or 1 )
1
...

```

```

...

# api: string.ljust(width[, fillchar])

>>> "Hello".ljust(10)
'Hello '
>>> "Hello".ljust(10, '-')
'Hello-----'
>>> "Hello".ljust(2)
'Hello'

```

```

>>> "Hello World".ljust(20)
'Hello World '
>>> "Hello World".ljust(20, '-')
'Hello World-----'

```

```

# api: string.rjust(width[, fillchar])
>>> "Hello".rjust(10)
' Hello'
>>> "Hello".rjust(10, '-')

```

```

'-----Hello'
>>> "Hello".rjust(2)
'Hello'

>>> "Hello World".rjust(20)
' Hello World'
>>> "Hello World".rjust(20, '-')
'-----Hello World'

...

```

```

class Solution:
    def fullJustify(self, words: List[str],
maxWidth: int) -> List[str]:
        res, cur, num_of_letters = [], [], 0

        for w in words:
            # len(cur) is the space count, from
last round, right one less after adding 'w'
            if num_of_letters + len(cur) + len(w) >
maxWidth:
                for i in range(maxWidth -
num_of_letters):
                    # The "or 1" part is for
dealing with the edge case 'len(cur) == 1'

```

```

        cur[i % ( len(cur) - 1 or 1 )]
+= ' '
        res.append(''.join(cur))
        cur, num_of_letters = [], 0
        cur += [w]
        num_of_letters += len(w)

    return res + ['
'.join(cur).ljust(maxWidth)]

```

```
#####
```

```

...
>>> 17 % 3
2
>>> divmod(17,3)
(5, 2)
...
...
>>> ['This', 'is', 'an']

```

```

['This', 'is', 'an']
>>> t = ['This', 'is', 'an']
>>> spaces = [' ', ' ']

# note: here j starts from 0
>>> for j, word in enumerate(t[1:]):
...     print(j)
...     print(word)
...
0

```

```

is
1
an
...

```

```

class Solution:
    def fullJustify(self, words: List[str],
maxWidth: int) -> List[str]:
        def partition(n, cnt):
            base, mod = divmod(n, cnt)
            res = [f"{' ' * (base + (i < mod))}"
for i in range(cnt)]

```

```

        return res

    ans = []
    i, n = 0, len(words)
    while i < n: # one line per iteration
        t = [words[i]]
        cnt = len(words[i])
        i += 1 # move to next word
        while i < n and cnt + 1 + len(words[i])
            <= maxWidth: # greedy search for one line
                cnt += 1 + len(words[i])

        t.append(words[i])
        i += 1
        if i == n or len(t) == 1:
            # this is the last line or only one
            (super-long) word in a line
            left = ' '.join(t)
            right = ' ' * (maxWidth -
len(left))

            ans.append(f"{left}{right}")
            continue # so i != n , so only one
word, no need to add spaces.
                # e.g. a word with the
same length as maxWidth,
                # or, e.g. the 2nd word is
super long and cannot fit in current line

```

```

        # if i==n, then will not enter next
while
        words_width = cnt - (len(t) - 1) # pure
words total length, with no spaces
        #
spaces in-between these words in t[: len(t) - 1
        space_width = maxWidth - words_width
spaces = partition(space_width, len(t)
- 1)
        sb = [t[0]] # 1st word of this line, no
space before it
        for j, word in enumerate(t[1:]): # last
word has no following space
            sb.append(spaces[j]) # j starts
from 0

            sb.append(word)
            ans.append(''.join(sb))

    return ans

#####

class Solution:
    def fullJustify(self, words: List[str],
maxWidth: int) -> List[str]:
        left = 0
        result = []

        while left < len(words):

```

```

        right = self.findRight(left, words,
maxWidth)
        result.append(self.justify(left,
right, words, maxWidth))
        left = right + 1

```

```

    return result

```

```

def findRight(self, left, words, maxWidth):
    right = left
    word_length = len(words[right])

```

```

        while (right + 1) < len(words) and
(word_length + 1 + len(words[right+1])) <=
maxWidth:
            right += 1
            word_length += 1 + len(words[right])

```

```

    return right

```

```

def justify(self, left, right, words,
maxWidth):
    if right - left == 0:
        return self.padResult(words[left],
maxWidth)

```

```

        is_last_line = (right == len(words) - 1)
        num_spaces = right - left
        total_space = maxWidth -
self.wordsLength(left, right, words)

```

```

        space = " " if is_last_line else " " *
(total_space // num_spaces)
        remainder = 0 if is_last_line else
total_space % num_spaces

```

```

        result = ""
        for i in range(left, right):
            result += words[i]

```

```

            result += space
            if remainder > 0:
                result += " "
                remainder -= 1

```

```

        result += words[right]
        result += " " * (maxWidth - len(result))

```

```

    return result

```

```

def wordsLength(self, left, right, words):
    words_length = 0
    for i in range(left, right+1):
        words_length += len(words[i])
    return words_length

def padResult(self, result, maxWidth):
    return result + " " * (maxWidth -
len(result))

```

◆ Quick Review

Key Insights

- Use a greedy approach to fill each line with as many words as possible without exceeding `maxWidth`.
- Calculate total space to distribute by subtracting the combined words length from `maxWidth`.
- If the line has more than one word, distribute extra spaces evenly; if not, pad the line with spaces at the end.
- Handle the last line separately by left-justifying the text (i.e., adding a single space between words and padding the end).

Space & Time

Time Complexity: $O(n)$ where n is the number of words, since each word is processed once.

Space Complexity: $O(n)$ for storing the result and temporary buffers.

Solution

The solution iterates over the list of words and groups them into lines such that the combined length of the words and minimum required spaces does not exceed `maxWidth`. For each completed group (line):

– If it's not the last line, calculate extra spaces required. – If the line has multiple words, distribute the spaces evenly, adding any extra spaces from left to right. – If the line contains a single word, append spaces to the right. – For the last line, join words with a single space and pad the remaining width with spaces.

Data structures used include lists (or arrays) for current words accumulation and the final result. The logic is implemented using a simulation approach to mimic text formatting.

Chapter 2

Master Quick-Review Summary

All 9 rounds · key insights · complexity · solution ·
organised by priority & difficulty

Round 9 · Low · Hard

p.3063

Round 9 · Low · Hard

p.3064

R1 · High · Easy (43 q)

>> Round 1 · High Priority · Easy — Quick Review

43 questions · key insights · complexity · solution**#1 Two Sum [Easy] · Arrays & Hashing****Key Insights**

- Use a hash table to store the elements and their indices for fast lookup.
- For each element, calculate its complement: target minus the current element.
- Check if the complement exists in the hash table.
- If found, return the indices of the current element and its complement.
- This approach only requires a single pass through the array.

Space & Time**Time Complexity:** $O(n)$ – We traverse the list only once.**Space Complexity:** $O(n)$ – In the worst-case, we store all elements in the hash table.**Solution**

We use a hash table (dictionary) to map each array element to its index as we iterate. For each number, we compute the complement (target – number) and check if

it already exists in our hash table. If the complement is found, we return the indices. This one-pass solution ensures linear time complexity and avoids the need for nested loops.

#163 Missing Ranges [Easy] · Arrays & Hashing

Key Insights

- Use a pointer (or iterate index) to traverse the nums array while keeping track of the previous number considered.
- Check for missing numbers before the first element, between consecutive elements, and after the last element.
- When a gap is found, convert it to a range string that is either a single number (if the gap is exactly one number) or an interval.
- Handle edge cases when the nums array is empty.

Space & Time

Time Complexity: $O(n)$, where n is the length of the nums array, since we iterate through the array once.

Space Complexity: $O(1)$ additional space (ignoring the output list), as we only use a few extra variables.

Solution

The solution leverages a single pass through the input array and calculates the missing ranges by comparing the expected number with the current number in nums.

First, initialize a variable "prev" to one less than the lower bound so that missing numbers from the very start of the range can be handled uniformly. For each number in the array, if there is a gap between prev and the current number (i.e., current number is at least 2 more than prev), then add the missing range (prev+1 to current number-1) to the output list. Update "prev" to the current number and finally handle any gap between the last number and the upper bound. The primary data structures used are the list for the output; the approach is iterative.

#346 Moving Average from Data Stream [Easy] · Arrays & Hashing

Key Insights

- Use a queue (or circular buffer) to keep track of the current window elements.
- Maintain a running sum to quickly update the window's total when a new element comes in.
- When the window is full, subtract the value of the element that falls out as a new one is added.
- Each call to next() is executed in constant time, $O(1)$.

Space & Time

Time Complexity: $O(1)$ per call to next()

Space Complexity: $O(\text{window size})$

Solution

The solution leverages a queue data structure to store the elements of the sliding window. A running sum variable is maintained for efficient calculation of the moving average. When a new integer is added:

- If the size of the queue is equal to the maximum window size, remove the oldest element and subtract its value from the running sum.
- Add the new element to both the queue and the running sum.
- Compute and return the moving average as the running sum divided by the number of elements in the queue. This approach ensures that each update is done in constant time without having to sum all elements in the window repeatedly.

#760 Find Anagram Mappings [Easy] · Arrays & Hashing

Key Insights

- Since `nums2` is an anagram of `nums1`, every element in `nums1` has a corresponding position in `nums2`.
- A dictionary (or hash table) can be used to map each number in `nums2` to its indices.
- When iterating through `nums1`, we can assign the next available index from the dictionary for that number.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the arrays (single pass to build the dictionary and another pass to form the mapping).

Round 9 · Low · Hard

p.3069

Space Complexity: $O(n)$ for storing the dictionary that holds indices of elements in `nums2`.

Solution

The approach consists of two primary steps:

- Build a dictionary to record all indices where each number appears in `nums2`. This dictionary maps a number to a list of indices.
- Iterate over `nums1` and for each element, remove (or pop) one index from the corresponding list in the dictionary, and add that index to the final mapping array. This guarantees that each number from `nums1` is correctly mapped to one of its positions in `nums2` while efficiently handling duplicates.

#1426 Counting Elements [Easy] · Arrays & Hashing

Key Insights

- Use a hash table or set to store unique elements from `arr` for constant time lookups.
- Iterate through each element in `arr` and check if $x + 1$ exists in the set.
- Count each valid occurrence, including duplicates.
- This approach efficiently handles the problem without needing nested loops.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in `arr` (set lookups are $O(1)$ on average).

Round 9 · Low · Hard

p.3070

Space Complexity: $O(n)$, due to the use of an auxiliary set to store unique elements.

Solution

The solution leverages a set to achieve constant time membership checks. First, we convert `arr` to a set of unique numbers. Then, iterate through `arr`, and for each element, check if `(element + 1)` exists in the set. If it does, increment the count. This approach correctly handles duplicates by counting each occurrence during iteration.

#1534 Count Good Triplets [Easy] · Arrays & Hashing

Key Insights

- Brute force approach using three nested loops is acceptable since the maximum array length is 100.
- Check all possible triplets and count only those meeting the specified conditions.
- Absolute difference comparisons are the key condition validations.

Space & Time

Time Complexity: $O(n^3)$ where n is the length of the array.

Space Complexity: $O(1)$ extra space (ignoring input space).

Solution

We use a brute force algorithm with three nested loops, iterating over all combinations of triplets (i, j, k) ensuring $i < j < k$. For each triplet, we calculate the absolute differences:

– $|arr[i] - arr[j]| - |arr[j] - arr[k]| - |arr[i] - arr[k]|$ The triplet is counted as good if all absolute differences satisfy the conditions ($\leq a$, $\leq b$, and $\leq c$ respectively). This simple approach leverages the small input size, eliminating the need for more complex data structures or optimizations.

#383 Ransom Note [Easy] · String

Key Insights

- Use a hash table (or dictionary) to store the frequency of each character present in magazine.
- Iterate over every character in `ransomNote` and check if the corresponding count in the hash table is available and sufficient.
- Decrease the character count for every match; if any character count falls below zero, return false.
- If all characters in `ransomNote` can be accounted for, return true.

Space & Time

Time Complexity: $O(n + m)$, where n is the length of magazine and m is the length of `ransomNote`.

Space Complexity: $O(1)$, since the extra space required for the hash table is limited to at most 26 entries for lowercase English letters.

Solution

We construct a frequency map for the magazine string using a hash table. This frequency map records how many times each letter appears in magazine. Then we iterate through the ransomNote string, and for every letter, we check the corresponding frequency in the map. If a letter is not found or its frequency is zero, we immediately return false as it means the letter cannot be used to build the note. Otherwise, we decrement the frequency count by one. If we successfully process all characters in ransomNote, then it can be constructed from magazine, and we return true.

#734 Sentence Similarity [Easy] · String

Key Insights

- First, check if both sentences have the same length.
- A word is always similar to itself.
- Similarity is defined only by the given pairs; it is not transitive.
- Use a hash table (dictionary) to store bi-directional similar pairs for quick look-ups.

Space & Time

Time Complexity: $O(n + m)$, where n is the number of words in the sentences and m is the number of similar pairs.

Space Complexity: $O(m)$, for storing the similar pairs in a hash table.

Solution

- Check if the two sentences have the same length; if not, return false.
- Build a hash table where each word maps to a set of words it is similar to. For each pair $[x, y]$ in similarPairs, add y to the set for x and x to the set for y .
- Iterate over the words of both sentences simultaneously. For each corresponding pair, if the words are identical, continue; otherwise, check if one word appears in the similar set of the other.
- If all pairs meet the similarity condition, return true; otherwise, return false.

#1165 Single Row Keyboard [Easy] · String

Key Insights

- Create a mapping (hash table/dictionary) from each character to its index on the keyboard for $O(1)$ lookups.
- The cost to type each character is the absolute difference between the current position and the target character's position on the keyboard.
- Start from index 0 and sum up the distances moved for each successive character in the word.

Space & Time

Time Complexity: $O(n)$, where n is the length of the word.

Space Complexity: $O(1)$ (constant space for a mapping of 26 characters).

Solution

We solve the problem by first building a dictionary that maps each character in the keyboard layout to its corresponding index. We then initialize a variable to track the current finger position (starting at 0). For each character in the word, we look up its index, calculate the absolute difference from the current position, add that distance to the total time, and update the current position. This approach uses a single pass over the word with constant time lookups, ensuring an efficient solution.

#88 Merge Sorted Array [Easy] · Two Pointers

Key Insights

- Both arrays are already sorted in non-decreasing order.
- The merge should be done in-place into `nums1`.
- Start merging from the end to avoid overwriting elements in `nums1`.
- Use two pointers, one for `nums1` and one for `nums2`, and a tail pointer for placement.

Space & Time

Time Complexity: $O(m + n)$

Space Complexity: $O(1)$

Solution

We use a two-pointer approach starting from the end of both arrays. Initialize three pointers:

- `p1` pointing to the last valid element in `nums1`.
- `p2` pointing to the last element in `nums2`.
- `tail` pointing to the last index of `nums1`.

While both pointers are valid, compare the elements pointed by `p1` and `p2`, placing the larger element at the `tail` position. Decrement the corresponding pointer and the `tail` pointer. Finally, if any elements remain in `nums2` (i.e., `p2 >= 0`), copy them over into `nums1`. This solution leverages the extra space at the end of `nums1` and ensures the merge is done in-place in linear time.

#125 Valid Palindrome [Easy] · Two Pointers

Key Insights

- Use two pointers (one starting at the beginning and one at the end) to compare characters.
- Skip characters that are not alphanumeric.
- Compare characters in a case-insensitive manner.
- Continue until the pointers meet or cross.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, as each character is processed at most once.

Space Complexity: $O(1)$, using only a few extra variables for pointers and temporary comparisons.

Solution

The solution employs a two-pointer approach:

- Initialize two pointers: one at the start (left) and one at the end (right) of the string.
- Skip any non-alphanumeric characters using these pointers.
- Convert the remaining characters at each pointer to lowercase and compare them.
- If any mismatch is found, return false; otherwise, continue until the pointers cross.
- Return true if all corresponding characters match. This method efficiently checks for a palindrome without needing additional space for a filtered string.

#283 Move Zeroes [Easy] · Two Pointers

Key Insights

- Use a two-pointer technique: one pointer to iterate through the array and another pointer to track the position for the next non-zero element.
- For each non-zero element encountered, swap it with the element at the tracking pointer.
- This strategy minimizes operations by ensuring each non-zero element is moved at most once.

- After repositioning all non-zero elements, zeros will naturally occupy the remaining positions.

- The in-place requirement ensures space complexity remains $O(1)$.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in the array, since only one pass (or at most two simple passes) is required.

Space Complexity: $O(1)$ extra space, as the reordering is accomplished in-place.

Solution

The solution leverages a two-pointer approach. The left pointer (j) keeps track of the index where the next non-zero element should be placed, while the right pointer (i) iterates through the array. Whenever a non-zero element is encountered, it is swapped with the element at index j , and j is then incremented. This effectively gathers all non-zero elements at the beginning while moving zeros to the end. A key optimization is to perform the swap only when necessary (i.e., when i and j are not the same) to reduce the number of write operations.

#408 Valid Word Abbreviation [Easy] · Two Pointers

Key Insights

- Use two pointers: one for looping through the word and one for the abbreviation.
- When encountering a digit in abbr, accumulate the number and skip that many characters in the word.
- Check for invalid cases such as numbers with leading zeros or consecutive digit segments that might imply adjacent skipped substrings.
- Compare letters directly when they are not digits.

Space & Time

Time Complexity: $O(n + m)$ where n is the length of the word and m is the length of the abbreviation.

Space Complexity: $O(1)$ (constant extra space)

Solution

The approach uses two pointers to iterate over the word and abbreviation. For each character in the abbreviation:

- If it is a letter, compare it with the current character in the word.
- If it is a digit, first ensure that it does not start with '0'. Then, convert the sequence of digits into a number which represents how many characters in the word to skip. After processing, both pointers should have reached the end of their respective strings for the abbreviation to be valid. This approach is efficient due to its linear pass with constant extra space.

#977 Squares of a Sorted Array [Easy] · Two Pointers

Round 9 · Low · Hard

p.3079

Key Insights

- Squaring numbers can disrupt the original order since negative numbers become positive.
- A two-pointer approach can efficiently retrieve the largest square from either end of the array.
- By comparing the absolute values at the two pointers, the largest square is inserted from the back of the result array.
- This method achieves an $O(n)$ time complexity without needing to sort the squared values afterward.

Space & Time

Time Complexity: $O(n)$ – We traverse the array only once.

Space Complexity: $O(n)$ – We use an extra array to store the output.

Solution

We use a two-pointer technique to traverse the array from both ends. Initialize pointers (left and right) at the beginning and end of the array, respectively, and an index pointer starting at the end of a new results array. At each step, compare the absolute values of the elements at the left and right pointers. The larger absolute value, when squared, is placed in the result array at the current index, and the corresponding pointer is adjusted. This continues until the left pointer exceeds the right pointer. Filling the result array from the back ensures that the array remains sorted in non-decreasing order.

Round 9 · Low · Hard

p.3080

#169 Majority Element [Easy] · Sorting

Key Insights

- The majority element appears more than half the length of the array.
- A simple hash table approach can count occurrences, but a more efficient solution exists.
- The Boyer–Moore Voting Algorithm finds the majority in $O(n)$ time and $O(1)$ space.
- The algorithm maintains a candidate and a counter, adjusting the candidate when the count drops to zero.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We use the Boyer–Moore Voting Algorithm to identify the majority element. The algorithm iterates through the array using two variables: one for the candidate and one for the count. Initially, we set the candidate to an arbitrary value and count to 0. For each element, if the count is 0, we update the candidate to the current element. Then, if the current element is the same as the candidate, we increment the count; otherwise, we decrement it. Since the majority element appears more than $n/2$ times, it will remain as the candidate by the end of the iteration.

#217 Contains Duplicate [Easy] · Sorting

Key Insights

- A hash set can be used to keep track of seen numbers.
- Iterating once through the array provides an efficient solution.
- If a number is already in the set, a duplicate is found.
- Alternative approaches like sorting have a higher time complexity ($O(n \log n)$).

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution uses a hash set to track the numbers seen while iterating through the array. For each number in the array:

- Check if it exists in the hash set. – If it does, return true immediately as a duplicate is found. – Otherwise, add the number to the hash set. If the iteration completes without finding any duplicates, return false. This approach efficiently determines if any duplicates exist in a single pass.

#242 Valid Anagram [Easy] · Sorting

Key Insights

- If s and t have different lengths, they cannot be anagrams.
- An anagram requires exactly the same character frequency for both strings.
- Using a hash table (or frequency array when limited to 26 lowercase letters) provides an efficient solution.
- An alternative approach is to sort both strings and compare them.
- For Unicode inputs, a hash table is more adaptable than a fixed-size frequency array.

Space & Time

Time Complexity: $O(n)$, where n is the length of the strings (using a hash table for counting).

Space Complexity: $O(1)$ if we use a fixed-size array (for 26 letters), otherwise $O(n)$ for a hash table with a larger character set.

Solution

We use a hash table (dictionary) to count the frequency of each character in the first string s . Then, we iterate over string t and decrement the count for each character. If at any point a character count goes negative or a character is missing in the hash table, t cannot be an anagram of s . Finally, if all counts return to zero, the strings are anagrams. This approach is efficient and handles the problem in linear time.

#252 Meeting Rooms [Easy] · Sorting

Key Insights

- Sort the intervals based on their starting times.
- Compare each interval's end time with the next interval's start time.
- Detect overlap if the end time of one meeting exceeds the start time of the following meeting.
- If no overlaps are detected, all meetings can be attended.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(1)$ if sorting is done in-place; otherwise $O(n)$.

Solution

The solution starts by sorting the intervals in ascending order by their start times. After sorting, iterate through the intervals and for each consecutive pair, check if the previous meeting's end time exceeds the next meeting's start time. An overlap indicates that the person cannot attend all meetings. The primary data structure used is an array (or list) for holding the intervals. Sorting makes it easier to compare adjacent intervals for potential overlaps.

#1086 Largest Unique Number [Easy] · Sorting

Key Insights

- Use a hash table (or dictionary) to count the frequency of each number.
- Iterate over the hash table to find the numbers that appear exactly once.
- Identify the maximum number among the unique numbers.
- Return -1 if no unique number exists.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in `nums`.

Space Complexity: $O(n)$, for the hash table storing frequency counts.

Solution

The approach starts by initializing a hash table to keep track of the frequency of each element in the array. Once the frequency count is complete, we traverse the hash table to collect numbers that appear exactly once. Finally, we determine the maximum number among these unique elements. If there is no unique element, we return -1. This method is efficient because the frequency count and the search for the maximum unique element both run in linear time relative to the input size.

#1122 Relative Sort Array [Easy] · Sorting

Key Insights

Round 9 · Low · Hard

p.3085

- Use a hash table (or array for counting given the value constraints) to count the occurrences of each element in `arr1`.
- Iterate through `arr2` and for each element, append it to the result based on its count from the hash table.
- Remove the used elements from the hash table.
- Sort the remaining elements (those not in `arr2`) in ascending order and append them to the result.

Space & Time

Time Complexity: $O(n + m \log m)$ where n is the length of `arr1` and m is the number of leftover elements not in `arr2`. Given that element values are bounded (0 to 1000), a counting sort can make this nearly $O(n)$.

Space Complexity: $O(n)$ for storing the frequency counts and the result array.

Solution

The solution leverages a frequency counter (hash table) to count how many times each element appears in `arr1`. We then process `arr2` to add each number in its given order according to its frequency. Finally, we sort the remaining numbers that are not in `arr2` and append them to the result array. This approach ensures that the relative order specified by `arr2` is maintained and that the other elements are sorted in ascending order.

#35 Search Insert Position [Easy] · Binary Search

Round 9 · Low · Hard

p.3086

Key Insights

- The array is sorted, enabling the use of binary search.
- Binary search achieves a time complexity of $O(\log n)$ which meets the problem constraints.
- When the target is not found during the search, the final low pointer represents the correct insertion index.

Space & Time

Time Complexity: $O(\log n)$ due to binary search.

Space Complexity: $O(1)$ as no extra space is used.

Solution

We use binary search to find the target in the sorted array. Initialize two pointers, low and high. At each iteration, compute the mid index. If the value at mid matches the target, return mid. If the target is less than the mid value, adjust the high pointer; if greater, adjust the low pointer. When the iterations finish without finding the target, return the low pointer, which represents the correct insertion index.

#69 Sqrt(x) [Easy] · Binary Search

Key Insights

- The problem can be solved efficiently using binary search.
- Instead of iterating from 0 to x, binary search narrows down the candidate square roots quickly.

- The algorithm should handle edge cases like $x = 0$ and $x = 1$.
- By checking $\text{mid} * \text{mid}$ against x, we can determine if we need to search in the lower or upper half.

Space & Time

Time Complexity: $O(\log x)$

Space Complexity: $O(1)$

Solution

We use a binary search approach to find the integer square root. Initialize two pointers, low and high, where low is set to 0 and high is set to x. For each iteration, calculate the mid value. If mid squared equals x, we have found the square root. If mid squared is less than x, record mid as a potential answer and move the low pointer to mid + 1. Otherwise, move the high pointer to mid - 1. Continue the search until low exceeds high, then return the last recorded candidate.

#278 First Bad Version [Easy] · Binary Search

Key Insights

- The API returns true for bad versions and all versions after a bad version.
- The problem is essentially about finding a transition point in a sorted sequence which calls for a binary search approach.

- The aim is to minimize the number of calls to `isBadVersion` by leveraging the ordered nature of the versions.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We use binary search to efficiently locate the first bad version. Initialize two pointers, left and right, at 1 and n respectively. At each step, compute the mid point and call `isBadVersion(mid)`. If mid is bad, it indicates that the first bad version may be mid or to its left, so adjust right to mid. Otherwise, adjust left to mid + 1. Continue the process until the pointers converge. The converged pointer will be the first bad version. Key structures include simple integer variables for left, right, and mid, with binary search being the fundamental algorithm to reduce the search space logarithmically.

#374 Guess Number Higher or Lower [Easy] · Binary Search

Key Insights

- The problem can be effectively solved using binary search since the search space is sorted by nature (1 to n).

- Each API call helps to halve the search space, leading to an efficient solution.
- Be cautious with potential overflow when calculating the mid-point by using $\text{low} + (\text{high} - \text{low}) // 2$.

Space & Time

Time Complexity: $O(\log n)$ since we binary search the range.

Space Complexity: $O(1)$ as we use a constant amount of extra space.

Solution

We use a binary search algorithm to efficiently locate the target number. We initialize two pointers: low (1) and high (n). In each iteration, we compute the mid-point, then call the guess API. Based on the response:

- If `guess(mid)` returns 0, we have found the number.
- If it returns -1, it indicates that our guess is higher than the pick, so we update the high pointer to mid - 1.
- If it returns 1, it indicates that our guess is lower than the pick, so we update the low pointer to mid + 1.

This approach guarantees a logarithmic time solution.

#704 Binary Search [Easy] · Binary Search

Key Insights

- Utilize binary search to achieve $O(\log n)$ runtime on a sorted array.

- Use two pointers to maintain the current search region.
- Narrow down the search range by comparing the middle element with the target.
- Return the index when the target is found; otherwise, return -1 if not found.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We use an iterative binary search algorithm. We initialize two pointers, left and right, to represent the current segment of the array that could contain the target. At each iteration, we compute the middle index and compare the element there with the target. If they match, we return the index. Otherwise, based on whether the target is smaller or larger than the middle element, we adjust our pointers to narrow the search to either the left or right half of the array. This process continues until the target is found or until the pointers cross, indicating that the target is not present. No extra data structures are used, making the space complexity $O(1)$.

#422 Valid Word Square [Easy] · Matrix

Key Insights

- The k -th row should be equal to the k -th column.
- Not every word is the same length; checks must be performed within bounds.
- Iterating over each character and cross-verifying with the corresponding column character is sufficient.

Space & Time

Time Complexity: $O(N * L)$ where N is the number of words and L is the average length of the words.

Space Complexity: $O(1)$ additional space is used for the checking procedure.

Solution

We iterate through each word (row) and each character (column) in the word. For every character at position (i, j) , we confirm that there is a corresponding word at index j and that its length is greater than i . Then we check that the character at position (j, i) in that word is the same as the current character. If any of these checks fail, we return false. If all checks pass, the words form a valid word square.

#566 Reshape the Matrix [Easy] · Matrix

Key Insights

- Check if the reshape operation is possible by comparing the total elements ($m * n$ with $r * c$).
- Flatten the original matrix while preserving the row-traversing order.

- Use the flattened list to build the new matrix row by row.
- If the total number of elements does not match, simply return the original matrix.

Space & Time

Time Complexity: $O(m$

$n)$ as we iterate through all elements of the matrix once.

Space Complexity: $O(m$

$n)$ for the additional list used to store flattened elements (or the new matrix).

Solution

The solution involves first checking whether the reshape is possible by ensuring that the product of the rows and columns of the original matrix equals that of the desired matrix ($r * c$). If not, the original matrix is returned. Otherwise, the matrix is flattened into a one-dimensional list, preserving the row-order traversal. Finally, the flattened list is segmented by the new column size c to form the reshaped matrix. The primary data structures used include a list (or array) for flattening the matrix and then another list (or vector/array) to hold the final reshaped matrix.

#766 Toeplitz Matrix [Easy] · Matrix

Key Insights

Round 9 · Low · Hard

p.3093

- Each element (except those in the first row and first column) must be equal to its top-left neighbor.
- You can validate the matrix by iterating from the second row/column and comparing each element with $matrix[i-1][j-1]$.
- This simple check allows you to determine if every diagonal has the same elements without extra space.
- For streaming data (loading one row at a time), maintain the previous row for comparison.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(1)$

Solution

The problem is solved by iterating through the matrix starting from the element at position (1, 1). For every element at $matrix[i][j]$, compare it with the element at $matrix[i-1][j-1]$. If any two elements in this pair do not match, the matrix is not Toeplitz. This check ensures that every diagonal from the top-left to bottom-right has identical elements. In terms of auxiliary space, only constant extra space is used, which makes this approach efficient both in time and space.

#867 Transpose Matrix [Easy] · Matrix

Key Insights

Round 9 · Low · Hard

p.3094

- The value at position (i, j) in the original matrix moves to (j, i) in the transposed matrix.
- The transposed matrix dimensions are the reverse of the original: if the original is m x n, then the transposed is n x m.
- A simple double loop is sufficient to swap the indices, ensuring $O(m \cdot n)$ time complexity.
- This problem handles both square and rectangular matrices.

Space & Time

Time Complexity: $O(m \cdot n)$

Space Complexity: $O(m \cdot n)$ for the new matrix holding the transposed values

Solution

The solution involves iterating over each element of the original matrix and assigning it to the corresponding position in the new matrix where the indexes are swapped. We create the transposed matrix with dimensions based on the number of columns and rows of the original matrix, respectively. The approach leverages basic array manipulation with nested loops to fill in the new matrix.

#2373 Largest Local Values in a Matrix [Easy] · Matrix

Round 9 · Low · Hard

p.3095

Key Insights

- The size of the output matrix is $(n - 2) \times (n - 2)$.
- For each valid index (i, j) in the output, examine the 3 x 3 block in grid starting at `grid[i][j]` to `grid[i+2][j+2]`.
- Since each 3 x 3 block contains exactly 9 elements, finding the maximum in a block takes constant time.
- Use nested loops to slide over the grid and efficiently compute each local maximum.

Space & Time

Time Complexity: $O(n^2)$ – We iterate over $(n-2) \times (n-2)$ positions, and for each position, we look at 9 elements.

Space Complexity: $O(n^2)$ – The additional space is required for the output matrix.

Solution

The solution involves iterating over each possible 3 x 3 submatrix in grid. For each (i, j) starting index of the submatrix:

- Examine the nine elements from `grid[i][j]` to `grid[i+2][j+2]`.
 - Compute the maximum value from these elements.
 - Store the maximum in the corresponding position of the result matrix.
- Data structures used include the result matrix for storing the maximum values. The algorithmic approach involves two nested loops (one for rows and one for columns) to traverse each possible starting position of a 3 x 3 window. The simplicity of the solution comes from the

Round 9 · Low · Hard

p.3096

fixed 3x3 block size, which enables constant-time maximum computation for each block.

#100 Same Tree [Easy] · Trees & BST

Key Insights

- Use recursion to compare nodes in both trees.
- If both nodes being compared are null, they are the same.
- If one node is null while the other is not, the trees differ.
- Check if current nodes' values match, then recursively verify the left and right subtrees.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the trees, since every node is visited once.

Space Complexity: $O(n)$ in the worst-case (unbalanced tree), due to recursion call stack usage. In a balanced tree, the space would be $O(\log n)$.

Solution

The problem can be solved using a recursive approach (Depth-First Search) where we compare the current nodes from both trees. If the current nodes are both null, they match; if one is null, then the trees are not identical. For non-null nodes, check if the values are equal. Then, recursively apply the same logic on the left and right children of the nodes. This approach ensures

that both the structure and the node values are identical across the two trees.

#101 Symmetric Tree [Easy] · Trees & BST

Key Insights

- Use recursion (or iteration) to compare pairs of nodes.
- For two trees (or two nodes) to be mirror images:
 - Their root values must be the same.
 - The left subtree of one must match the right subtree of the other.
 - The right subtree of one must match the left subtree of the other.
- Alternatively, iterative approaches using a queue can be applied to compare the nodes in pairs.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree, since we traverse each node once.

Space Complexity: $O(h)$ for recursive call stack (worst-case $O(n)$ for a skewed tree) or $O(n)$ for the iterative approach using a queue.

Solution

We solve the problem using recursion by defining a helper function that takes two nodes and checks if they are mirrors of each other. The helper function:

- Returns true if both nodes are null.
- Returns false if one node is null or if their values differ.
- Recursively

compares the left child of the first node with the right child of the second node, and vice-versa. This method ensures each comparison verifies symmetry at every level. A similar logic can be applied iteratively by using a queue to compare pairs of nodes.

#104 Maximum Depth of Binary Tree [Easy] · Trees & BST

Key Insights

- Use a recursive strategy (DFS) to traverse the binary tree.
- At each node, compute the maximum depth of its left and right subtree.
- The maximum depth at the current node is 1 (current node) plus the maximum of the left and right subtree depths.
- Alternatively, an iterative approach using BFS (level-order traversal) can be used.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since every node is visited once.

Space Complexity: $O(h)$ for recursion stack in DFS, where h is the height of the tree (worst-case $O(n)$ for a skewed tree); $O(n)$ for BFS in the worst-case scenario if the tree is balanced.

Solution

Round 9 · Low · Hard

p.3099

The problem is approached using Depth-First Search (DFS) with recursion. The idea is to visit every node and calculate the maximum depth of the left and right subtrees, then add 1 for the current node. This method naturally handles the base case when a null node (empty tree) is reached by returning 0. Alternatively, a Breadth-First Search (BFS) approach can determine the tree level by level, with the total number of levels equaling the maximum depth. Key data structures include recursion call stack for DFS or a queue for BFS.

#108 Convert Sorted Array to Binary Search Tree [Easy] · Trees & BST

Key Insights

- The array is sorted in ascending order, which allows the use of a divide and conquer strategy.
- Choosing the middle element of the array as the root ensures that the left and right subtrees are roughly equal in size, yielding a balanced BST.
- A recursive approach is natural for building the tree: the base case is when the subarray is empty.
- The algorithm splits the array into two halves to recursively build the left and right subtrees.

Space & Time

Time Complexity: $O(n)$ – Every element is processed exactly once.

Round 9 · Low · Hard

p.3100

Space Complexity: $O(n)$ – $O(n)$ space is used to store the BST, and additional $O(\log n)$ space may be used for the recursion stack.

Solution

We use a divide and conquer approach with recursion. The main idea is to select the middle element of the current subarray as the root node so that the left subarray elements form the left subtree and the right subarray elements form the right subtree, ensuring the tree is balanced. The recursion continues until the subarray is empty, which serves as the base case. This strategy naturally maintains the BST property since all elements in the left subarray are smaller than the root, and all in the right subarray are larger.

#110 Balanced Binary Tree [Easy] · Trees & BST

Key Insights

- Use a bottom-up DFS (Depth-First Search) traversal to compute the height of each subtree.
- If any subtree is found unbalanced, propagate an error indicator (e.g., -1) upward, allowing early termination.
- An empty tree is considered balanced.
- The key check is to ensure that for each node, the absolute difference between the heights of its left and right subtrees is no more than 1.

Space & Time

Round 9 · Low · Hard

p.3101

Time Complexity: $O(n)$ where n is the number of nodes in the tree, since we visit each node once.

Space Complexity: $O(h)$ where h is the height of the tree, which corresponds to the recursion stack ($O(n)$ in the worst-case of a skewed tree).

Solution

The solution uses a recursive DFS approach. A helper function recursively computes the height of a subtree. If a subtree is found to be unbalanced (the height difference between its left and right subtrees exceeds 1), the helper returns a special value (e.g., -1) to indicate that the subtree is unbalanced. This value is then propagated up the recursion to terminate further unnecessary computations. The main function checks the result of the helper function and returns true if the tree is balanced (i.e., the helper did not return -1) and false otherwise.

#112 Path Sum [Easy] · Trees & BST

Key Insights

- Use a depth-first search (DFS) strategy to traverse the tree from the root to the leaves.
- At each node, subtract the node's value from the targetSum.
- Check if the node is a leaf; if yes and the remaining targetSum equals the node's value then a valid path is found.

Round 9 · Low · Hard

p.3102

- If the tree is empty, return false as no path exists.

Space & Time

Time Complexity: $O(N)$ where N is the number of nodes, since every node may need to be visited.

Space Complexity: $O(H)$ where H is the height of the tree, accounting for the recursion stack (worst-case $O(N)$ for a skewed tree).

Solution

The approach uses a recursive depth-first search. The algorithm subtracts the current node's value from targetSum as it moves down the tree. When it reaches a leaf, it checks if the modified targetSum is equal to the leaf's value. If yes, it returns true. The recursion continues until either a valid path is found or all paths are exhausted. The primary data structure used here is the call stack for recursion.

#226 Invert Binary Tree [Easy] · Trees & BST

Key Insights

- The problem can be solved using recursion (depth-first search) by swapping the children of every node.
- Alternatively, an iterative approach (using a queue for breadth-first search) can be used.
- The recursion terminates when a null node is encountered.

- The operation is performed in-place, modifying the original tree structure.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$ where h is the height of the tree, which is the recursion stack space in the recursive approach. In worst-case (skewed tree), this could be $O(n)$.

Solution

We solve the problem using recursion (depth-first search). The approach is to:

- Check the base case: if the node is null, simply return null.
- Recursively invert the left subtree and the right subtree.
- Swap the left and right children of the current node.
- Return the current node after the swap. The recursive approach elegantly handles all nodes and ensures the inversion process covers the complete tree.

#270 Closest Binary Search Tree Value [Easy] · Trees & BST

Key Insights

- BST property allows us to traverse the tree efficiently by comparing the target with node values.
- We can iteratively traverse the tree, moving left if the target is smaller than the current node's value and

moving right if greater.

- At each node, calculate the absolute difference between node's value and the target, and update the answer if the difference is smaller—or equal with a smaller candidate value.
- This approach minimizes the search area by leveraging the inherent ordering in BSTs.

Space & Time

Time Complexity: $O(h)$ on average, where h is the height of the tree ($O(\log n)$ for a balanced tree, and $O(n)$ in the worst-case scenario for a skewed tree).

Space Complexity: $O(1)$ for an iterative solution (or $O(h)$ if using recursion due to the call stack).

Solution

We solve the problem by performing an iterative traversal of the BST. Starting at the root, we maintain a variable to store the value closest to the target. For every node encountered, we:

- Compute the absolute difference between the node's value and the target.
- Update the stored closest value if the computed difference is less than the current smallest difference; if the difference is the same, choose the smaller node value.
- Depending on whether the target is less than or greater than the current node's value, move to the left or right child accordingly. This approach efficiently defines a search path that is consistent with the properties of a BST.

#501 Find Mode in Binary Search Tree [Easy] · Trees & BST

Key Insights

- The BST in-order traversal yields a sorted order, which helps in counting consecutive duplicates.
- A two-pass in-order traversal can be used: one pass to determine the maximum frequency (maxCount) and a second pass to collect all values that match this frequency.
- The solution can be implemented without extra space (apart from recursion stack) by leveraging constant space counters and state variables.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is processed twice (once in each traversal).

Space Complexity: $O(h)$ where h is the height of the tree due to the recursion stack. $O(1)$ extra space is used aside from this.

Solution

We perform an in-order traversal of the BST twice. The first traversal determines the maximum frequency of any value in the BST by comparing the current node's value with the previously visited node. Since the in-order traversal processes nodes in sorted order,

duplicate values are consecutive, making the counting simple. After calculating `maxCount`, we reset our state and perform a second in-order traversal to collect all node values with frequency equal to `maxCount`. The key trick is to leverage the BST property (sorted in-order traversal) to count frequencies in one pass without using additional data structures like hash maps.

#543 Diameter of Binary Tree [Easy] · Trees & BST

Key Insights

- Use depth-first search (DFS) to explore the tree.
- The diameter at a node is the sum of the depths of its left and right subtrees.
- Track the maximum diameter found during the DFS traversal.
- Recursively calculating the depth of subtrees provides an efficient solution.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree.

Space Complexity: $O(n)$, in the worst-case scenario due to the recursion stack (skewed tree).

Solution

The solution employs a DFS traversal of the binary tree. At every node in the tree, compute the depth of the left and right subtrees. The sum of these depths gives the

potential diameter (longest path through that node). Update a global maximum diameter variable if this sum exceeds the current maximum. Finally, return the maximum diameter after completing the DFS.

#572 Subtree of Another Tree [Easy] · Trees & BST

Key Insights

- Use Depth-First Search (DFS) to visit each node in the main tree.
- At each node, check if the subtree starting from that node is identical to `subRoot`.
- A helper function is used to compare the structure and node values of two trees.
- Consider recursion as the primary approach to traverse and compare trees.
- Early exit when a mismatch is found can optimize average-case performance.

Space & Time

Time Complexity: $O(m * n)$ in the worst case, where m is the number of nodes in the main tree and n is the number of nodes in the `subRoot` tree.

Space Complexity: $O(\max(\text{depth of main tree}, \text{depth of subRoot}))$ due to recursive call stack.

Solution

We solve the problem by performing a DFS traversal on the main tree. For each visited node, a helper function

checks whether the subtree starting at that node is identical to subRoot by comparing node values and recursively checking both left and right subtrees. This method leverages recursion effectively and provides early termination if mismatches occur, allowing for efficient traversal and comparison.

#463 Island Perimeter [Easy] · DFS

Key Insights

- Each land cell contributes 4 edges if isolated.
- For every adjacent pair of land cells (neighbors horizontally or vertically), the shared edge should be subtracted from the total perimeter.
- Instead of DFS/BFS, we can iterate the grid and check the four neighbors for each land cell.
- Alternatively, one may use DFS/BFS to visit all connected cells if needed, but an iterative solution is straightforward with $O(\text{rows} * \text{cols})$ time complexity.

Space & Time

Time Complexity: $O(n * m)$, where n is the number of rows and m is the number of columns.

Space Complexity: $O(1)$ for the iterative solution (excluding the input storage).

Solution

We iterate through every cell in the grid. For each land cell, we initially add 4 to the perimeter. Then we check

its four neighbors (up, down, left, right). For every neighbor that is also land, we subtract 1 from the current cell's contribution since that side is not exposed. This simple approach ensures that shared borders between adjacent land cells are not over-counted.

#733 Flood Fill [Easy] · DFS

Key Insights

- The flood fill operation can be executed using either DFS (Depth First Search) or BFS (Breadth First Search).
- A check is needed at the beginning to determine if the new color is the same as the original; if so, no changes are needed.
- The algorithm recursively or iteratively visits each pixel that is directly adjacent and matches the original color, then updates its color.
- The problem has a worst-case scenario where all pixels are the same color, leading to a full traversal of the grid.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the dimensions of the image grid, because in the worst-case all cells are visited.

Space Complexity: $O(m * n)$ in the worst-case due to the recursion stack (or BFS queue) if the entire image needs to be filled.

Solution

We use a DFS approach to perform the flood fill. The algorithm starts at the given starting pixel and uses recursion to fill connected pixels with the same original color. The following steps are key:

- Check if the starting pixel's color is already the target color; if so, return the image immediately.
- Define a recursive helper function that: Checks boundary conditions. Updates the current pixel's color. Recursively calls itself on the four adjacent directions (up, down, left, right).
- Return the modified image after the recursion completes.

Data structures used:

- Recursion call stack for the DFS implementation.
- Variables to store the image dimensions and original color.

R2 · High · Medium (116 q)

>> Round 2 · High Priority · Medium — Quick Review

116 questions · key insights · complexity · solution

#57 Insert Interval [Medium] · Arrays & Hashing

Key Insights

- The intervals array is sorted by start times.
- Three main steps are used:
- Add intervals that come completely before the new interval.
- Merge overlapping intervals with the new interval.
- Append intervals that start after the new (merged) interval.
- Iterating through the list once allows an $O(n)$ solution.

Space & Time

Time Complexity: $O(n)$ since we scan the list once.

Space Complexity: $O(n)$ to store the resulting list of intervals.

Solution

Traverse the intervals list and perform the following:

- Append all intervals that end before the new interval starts.
- For intervals that overlap with newInterval,

update newInterval to merge them (using min for start and max for end). – Append the new merged interval. – Append all remaining intervals. This ensures that all intervals remain non-overlapping and sorted with a single pass.

#238 Product of Array Except Self [Medium] · Arrays & Hashing

Key Insights

- The product for each index can be computed by multiplying the product of all numbers to the left and the product of all numbers to the right of that index.
- Division is not allowed, so we must compute the cumulative products explicitly.
- Create two passes: one forward pass for prefix products and one backward pass for suffix products.
- Use the output array to store the cumulative product values, achieving $O(1)$ extra space by not using additional arrays.

Space & Time

Time Complexity: $O(n)$ because we iterate through the array twice.

Space Complexity: $O(1)$ extra space (excluding the output array).

Solution

The approach uses two passes over the input array:

– Initialize the answer array and fill it with prefix products. For each index i , $\text{answer}[i]$ contains the product of all elements before i . – Traverse the array from the end using a running suffix product. Update the answer array by multiplying the current suffix product with the prefix product already stored. This method leverages an output array to store intermediate results and uses a temporary variable for the suffix product, ensuring constant extra space.

#281 Zigzag Iterator [Medium] · Arrays & Hashing

Key Insights

- Use a FIFO (first-in-first-out) structure to maintain the order of vectors which still have remaining elements.
- For each call to `next()`, select the next element from the current vector and then move that vector to the end of the queue if it still contains elements.
- The approach ensures proper zigzag (or cyclic) ordering among the vectors.
- When extending to k vectors, the same approach works: initialize the queue with each non-empty vector.

Space & Time

Time Complexity: $O(1)$ per `next()` call on average, with an overall $O(n)$ for n total elements.

Space Complexity: $O(k)$ for the queue structure (where k is the number of vectors), plus the storage for the input

vectors.

Solution

We can solve the Zigzag Iterator problem by using a queue (or deque) to hold pairs of the vector's current index and the vector itself. For each call to next(), we remove the front element from the queue, return the current item from its associated vector, and if the vector has more elements, reinsert the pair (with an updated index) into the back of the queue. This cyclic process naturally interleaves the elements from the provided vectors.

#307 Range Sum Query – Mutable [Medium] · Arrays & Hashing

Key Insights

- To efficiently support both point updates and range sum queries, a specialized data structure like a Binary Indexed Tree (Fenwick Tree) or a Segment Tree is ideal.
- Both data structures provide update and sum query operations in $O(\log n)$ time.
- The challenge is to maintain a mutable array while enabling fast sum computation over any specified range.

Space & Time

Time Complexity: $O(\log n)$ for both update and sumRange operations

Space Complexity: $O(n)$ to store the additional tree structure

Solution

We can use a Binary Indexed Tree (Fenwick Tree) to solve this problem. The Fenwick Tree allows for efficient prefix sum calculations and point updates, both in $O(\log n)$ time. The approach involves:

- Building the Fenwick Tree using the initial array.
- On an update, computing the difference between the new value and the current value, then propagating the change appropriately in the tree.
- For the sum range query, computing the prefix sum up to right and subtracting the prefix sum up to left-1.

#454 4Sum II [Medium] · Arrays & Hashing

Key Insights

- Use a hash table to store the frequency of all possible sums from nums1 and nums2.
- For each pair from nums3 and nums4, compute the complement (negative sum) and look it up in the hash table.
- This reduces the time complexity from $O(n^4)$ to $O(n^2)$.

Space & Time

Time Complexity: $O(n^2)$ where n is the length of each array.

Space Complexity: $O(n^2)$ due to the storage of pair sums in the hash table.

Solution

The solution employs a hash table to precompute and store the frequency of sums of all pairs from the first two arrays. Then, by iterating over all pairs from the third and fourth arrays, we calculate the required complement that would result in a total sum of zero. This complement is checked in the hash table, and the frequency (number of valid pairs) is added to the result. This method efficiently finds the number of valid quadruplets without resorting to a brute-force $O(n^4)$ approach.

#525 Contiguous Array [Medium] · Arrays & Hashing

Key Insights

- Convert 0s to -1 so that a balanced subarray (equal number of 0s and 1s) sums to zero.
- Use a prefix sum to track the cumulative sum as you iterate through the array.
- Use a hash table (or dictionary) to store the first occurrence of each prefix sum.
- When the same prefix sum is encountered again, it indicates that the subarray between these indices is balanced.

Space & Time

Time Complexity: $O(n)$ – We iterate through the array once.

Space Complexity: $O(n)$ – In the worst case, we store an entry for each prefix sum.

Solution

We use a technique where we convert each 0 in the array to -1. Then, we compute a running prefix sum as we iterate through the array. A prefix sum of 0 at any point indicates that the subarray from the beginning to the current index is balanced. For each prefix sum, if we have seen it before, the subarray between the previous index (where the prefix sum was first seen) and the current index sums to 0. We update the maximum length accordingly. We use a hash table to store the first index where each prefix sum occurs.

#560 Subarray Sum Equals K [Medium] · Arrays & Hashing

Key Insights

- Utilize the prefix sum technique to simplify sum calculation for subarrays.
- Use a hash table (or dictionary/map) to store the frequency of prefix sums encountered.
- For each element, check if (current prefix sum - k) exists in the hash table; if it does, it indicates a valid

subarray.

- Initialize the hash table with {0: 1} to manage cases where a subarray's sum is exactly k.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

We solve the problem by maintaining a running sum (prefix sum) and using a hash table to record the number of times each sum has occurred. As we iterate through the array, we update the running sum. For each new sum, we check if (current sum - k) exists in our hash table. If it does, the value associated with (current sum - k) gives the number of subarrays ending at the current index that sum to k. This method efficiently computes the result in a single pass through the array.

#1010 Pairs of Songs With Total Durations Divisible by 60 [Medium] · Arrays & Hashing

Key Insights

- Each song duration's effect is determined by its remainder when divided by 60.
- For any given module value r (where $0 \leq r < 60$), the complementary module needed to complete 60 is $(60 - r) \% 60$.

Round 9 · Low · Hard

p.3119

- Special handling is needed when r is 0 (or 30, since $30 + 30 = 60$) to avoid double-counting.
- A frequency counter for remainders can be efficiently used to count valid pairs in one pass through the list.

Space & Time

Time Complexity: $O(n)$, where n is the number of songs.

Space Complexity: $O(60) \rightarrow O(1)$, since we only need to keep track of 60 possible remainders.

Solution

We iterate over the list of song durations and calculate the remainder of each duration modulo 60. A frequency array (or hash table) is used to record how many times each remainder has been seen so far. For each song, we look up the frequency of its complementary remainder (i.e., $(60 - \text{current_remainder}) \% 60$) already encountered. This frequency indicates how many valid pairs can be formed with the current song. After counting pairs with the current song, we update the frequency of its remainder for future pairings.

#1272 Remove Interval [Medium] · Arrays & Hashing

Key Insights

- Iterate over each interval and check if it overlaps with the toBeRemoved interval.

Round 9 · Low · Hard

p.3120

- If the current interval is completely outside of toBeRemoved, add it to the result as is.
- For overlapping intervals, compute the non-overlapping portions:
 - If the left side of the interval lies before toBeRemoved, include [start, min(end, toBeRemoved.start)].
 - If the right side of the interval lies after toBeRemoved, include [max(start, toBeRemoved.end), end].

Space & Time

Time Complexity: $O(n)$, where n is the number of intervals.

Space Complexity: $O(n)$, for storing the resulting intervals.

Solution

The approach involves iterating through the list of sorted disjoint intervals and checking for overlap with the toBeRemoved interval. For each interval:

- If it is completely before or after the removal interval, add it directly to the result.
- If it overlaps with toBeRemoved, determine the remaining portion(s): The left side of the interval that does not intersect (if any). The right side of the interval that does not intersect (if any). Use simple conditional checks and list manipulation to form the final set of intervals.

#1429 First Unique Number [Medium] · Arrays & Hashing

Key Insights

- Use a queue to maintain the order of potential unique numbers.
- Use a hash table (or map) to count the occurrences of each number.
- When showing the first unique number, remove numbers from the front of the queue if their frequency is more than one.
- All operations (add and showFirstUnique) can be done in $O(1)$ amortized time.

Space & Time

Time Complexity: $O(1)$ amortized for each add and showFirstUnique call.

Space Complexity: $O(n)$, where n is the number of distinct integers encountered.

Solution

We design the FirstUnique class using two main data structures: a queue and a hash map (or dictionary). The queue preserves the order of insertion for unique candidates, while the hash map holds the frequency count for each integer.

When a new number is added:

- Increment its count in the hash map.
- If the count is one, append it to the queue.

To retrieve the first unique number:

– While the queue is not empty and the count for the element at the front is greater than one, remove it from the queue. – Return the front element of the queue if it exists, otherwise return -1 .

This approach ensures that we are always up-to-date with the current first unique number.

#3159 Find Occurrences of an Element in an Array [Medium] · Arrays & Hashing

Key Insights

- Scan the array once and record every index where `nums[i] == x``.
- Store those indices in order inside a `positions` array.
- For each query `k``, check whether the `k``-th occurrence exists.
- If it exists, return `positions[k - 1]``.
- Otherwise, return `-1``.
- This lets every query be answered in constant time after the preprocessing pass.

Space & Time

Time Complexity: $O(n + q)$, where n is the length of `nums`` and q is the number of queries.

Space Complexity: $O(m)$, where m is the number of occurrences of `x``.

Solution

We solve the problem by first precomputing all occurrences of `x``. As we iterate through `nums``, every time we see `x``, we append its index to an array called `positions``. Then, for each query value `k``, we check whether `k`` is within the bounds of `positions``. If it is, the answer is `positions[k - 1]``; otherwise, the required occurrence does not exist, so the answer is `-1``. This avoids rescanning the array for every query.

#8 String to Integer (atoi) [Medium] · String

Key Insights

- Skip leading spaces in the string.
- Detect and handle the optional sign indicator ('+' or '-').
- Read and convert digit characters until a non-digit is encountered.
- Handle cases with no digits or leading zeros.
- Clamp the result to fit within the 32-bit signed integer range.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string.

Space Complexity: $O(1)$ as only a constant amount of extra space is used.

Solution

We iterate through the string to first eliminate any leading whitespace. Next, we check for a sign character to determine if the resulting integer should be negative. We then convert each consecutive digit into an integer, all the while updating the result and checking for potential overflow. If an overflow is detected, we immediately return the appropriate maximum or minimum 32-bit signed integer boundary value. This approach ensures that we strictly follow the conversion rules and efficiently handle edge cases.

#249 Group Shifted Strings [Medium] · String

Key Insights

- All strings in the same group have the same relative differences between adjacent characters.
- Normalize each string's pattern by converting it to a key representation based on the differences between consecutive letters.
- A single character string forms its own group as there are no differences to compute.

Space & Time

Time Complexity: $O(n * m)$ where n is the number of strings and m is the average length of the strings (to compute each string's key).

Space Complexity: $O(n * m)$ to store the keys and grouping in the hash table.

Solution

The solution involves iterating over each string and generating a unique key that represents the relative differences between adjacent characters. This key is calculated by computing the difference (mod 26) between every consecutive pair in the string. Strings that share the same key can be grouped together since they can be shifted into one another. A hash table (or dictionary) is used to store and group the strings by their computed key. This approach handles edge cases like single-character strings separately.

#271 Encode and Decode Strings [Medium] · String

Key Insights

- Use length-prefix encoding for each string to avoid ambiguity.
- Append a delimiter (e.g., "#") after the length to separate it from the string content.
- This method ensures that even if the original strings contain special characters, including the delimiter, the decoding remains unambiguous.
- The approach handles edge cases, such as empty strings.

Space & Time

Time Complexity: $O(n)$, where n is the total number of characters across all strings.

Space Complexity: $O(n)$, required for the encoded string and the resulting decoded list.

Solution

We encode by iterating over each string and building a new string that contains the length of the string, a delimiter (e.g., "#"), and then the string itself. This way, during decoding, we can safely determine how many characters to extract by first reading until the delimiter to get the length, and then reading that number of characters. This length-prefix approach avoids conflicts with potential characters in the strings.

#635 Design Log Storage System [Medium] · String

Key Insights

- Store logs in a simple data structure such as a list or array.
- Convert the timestamps into strings; due to fixed-length and zero padding, lexicographical string comparisons are valid.
- Use a mapping to determine the slice length for each granularity, e.g. "Year" corresponds to the first 4 characters, "Month" corresponds to the first 7 characters, etc.
- During retrieval, compare the relevant substring of each stored timestamp against the start and end query timestamp slices.

- The inclusive nature of the query simplifies checking conditions.

Space & Time

Time Complexity: $O(N)$ per retrieval query, where N is the number of logs stored (since each log is checked). In the worst case, all log entries are examined.

Space Complexity: $O(N)$ where N is the number of logs stored.

Solution

We implement a `LogSystem` class that supports putting log entries and retrieving them using a specified granularity.

- **Data Structures:** Use a list/array to store log entries as tuples (or objects) containing the timestamp and log ID. A dictionary/map is used to store the mapping from granularity string to the substring index (e.g., "Year" => 4, "Month" => 7, etc.).
- For the `put()` method, simply append the new log entry into the list.
- For the `retrieve()` method, compute the substring length based on the provided granularity. Then iterate over all stored logs and compare the sliced timestamp against the sliced start and end query values. If it falls within the bounds, include the log ID in the result.
- Since the timestamps have fixed width with leading zeros, slicing and lexicographical comparisons are both valid and simple.

#1138 Alphabet Board Path [Medium] · String

Key Insights

- Determine the row and column for each letter (using integer division and modulus).
- Because "z" is in a special row that only contains one column, the move order is critical to avoid invalid positions.
- For target letter "z", perform horizontal moves before vertical moves; for all other letters, do vertical moves before horizontal moves.
- Build the answer step-by-step by moving from your current position to each target letter.

Space & Time

Time Complexity: $O(N)$ where N is the length of target (each move is processed in constant time)

Space Complexity: $O(N)$ for the resulting string of moves

Solution

The approach calculates the destination coordinates for each character. For each letter in the target:

- Convert the letter to its board coordinates (row and col). For example, for letter c, $\text{row} = (\text{ord}(c) - \text{ord}('a')) // 5$ and $\text{col} = (\text{ord}(c) - \text{ord}('a')) \% 5$. Note that "z" (the 26th letter) always maps to (5, 0).
- To move from the current position (curRow, curCol) to the target's (targetRow, targetCol), decide on the order of moves: If the target letter is "z", first move horizontally (left or right) then

vertically (up or down) to avoid stepping into invalid regions from the extra row. For any other letter, move vertically (up or down) first, then horizontally (left or right). – Append "!" to the moves after reaching each letter. – Update the current position to the letter's coordinates.

This method ensures you only traverse valid cells on the board and minimizes moves by taking the optimal coordinate difference each step.

#11 Container With Most Water [Medium] · Two Pointers

Key Insights

- The container's area is determined by the distance between two lines multiplied by the shorter line's height.
- A brute-force approach checking all possible pairs would result in $O(n^2)$ time complexity, which is inefficient for large inputs.
- A two-pointer approach allows us to solve the problem in $O(n)$ time by initializing pointers at both ends and gradually moving them inward.
- By always moving the pointer corresponding to the shorter line, we are more likely to find a taller line that could potentially form a container with a larger capacity.

Space & Time

Time Complexity: $O(n)$ – The algorithm makes a single pass through the array with two pointers.

Space Complexity: $O(1)$ – Only a few extra variables are used, regardless of the input size.

Solution

The optimal approach uses a two-pointer strategy:

- Start with two pointers at each end of the array.
- Calculate the area formed between the two lines at these pointers.
- Update the maximum area if the calculated value is larger.
- Move the pointer pointing to the shorter line inward, as this is the only possible way to potentially increase the height of the shorter line and thus the area.
- Continue until the two pointers meet.

This method effectively reduces the number of comparisons and ensures that every possible pair that can yield a larger area is considered.

#15 3Sum [Medium] · Two Pointers

Key Insights

- Sort the array to efficiently skip duplicates and facilitate the two pointers approach.
- Use a fixed pointer for the first element, and then use two additional pointers (left and right) to find pairs that sum to the negation of the fixed element.
- Skip duplicate elements to ensure unique triplets.

Space & Time

Round 9 · Low · Hard

p.3131

Time Complexity: $O(n^2)$ where n is the number of elements in the array.

Space Complexity: $O(1)$ extra space (ignoring the space required for the output).

Solution

We first sort the array to bring duplicates together and allow the two-pointer technique to work efficiently. Then, iterate over the array and for each element, use two pointers (one at the beginning right after the current element and one at the end of the array) to find a pair that sums with the current element to zero. Skip duplicates for the current element and within the inner loop to avoid duplicate triplets. This ensures that each valid triplet is unique and the overall runtime remains efficient.

#16 3Sum Closest [Medium] · Two Pointers

Key Insights

- Sort the array to make it easier to use the two pointers technique.
- Fix one element and then use two pointers (left and right) to find the best pair that, along with the fixed element, produces a sum closest to the target.
- Update the best sum whenever a closer sum is found.
- The sorted order helps eliminate unnecessary iterations and allows efficient movement of pointers.

Round 9 · Low · Hard

p.3132

Space & Time

Time Complexity: $O(n^2)$

Space Complexity: $O(\log n)$ to $O(n)$ depending on the sorting algorithm used (typically in-place sort, so extra space might be minimal)

Solution

The solution starts by sorting the array. For each index in the array (considered as the first element of the triple), we use two pointers: one starting just after the fixed element and the other at the end of the array. We calculate the sum of the fixed element and the two pointers. If this sum is closer to the target than the current best, we update our best sum. Depending on whether the current sum is less than or greater than the target, we adjust the left or right pointer to try and get closer to the target. This two-pointer method leverages the sorted order of the array for efficient summing.

#31 Next Permutation [Medium] · Two Pointers

Key Insights

- The problem is solved by identifying the first pair from the right where the left element is less than the right element.
- Once this pivot is found, find the rightmost element that is greater than this pivot.
- Swap these two elements.

Round 9 · Low · Hard

p.3133

- Reverse the subarray that comes after the pivot to ensure the next permutation is the next lexicographically larger sequence.
- If no such pivot exists, the array is in descending order and should be reversed to get the smallest permutation.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We start by scanning the array from the right to find the first index (i) such that $\text{nums}[i]$ is less than $\text{nums}[i+1]$. This index 'i' represents the pivot point where the next permutation can be formed. If we find such an index, we then scan from the right again to locate the first number that is greater than $\text{nums}[i]$ (let this index be j), and swap these two numbers. Finally, we reverse the subarray starting from index $i+1$ to the end of the array, which gives us the next permutation in lexicographical order. This approach leverages in-place modifications using basic operations like swapping and reversing, thus meeting the constant space requirement.

#75 Sort Colors [Medium] · Two Pointers

Key Insights

- Use the Dutch National Flag algorithm to partition the array into three sections.

Round 9 · Low · Hard

p.3134

- Maintain three pointers: one for the next position for 0 (low), one for the next position for 2 (high), and one to traverse the array (mid).
- Swap elements appropriately to ensure that all 0's are at the beginning, all 2's are at the end, and 1's remain in the middle.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution uses a three-pointer approach (Dutch National Flag algorithm). Initialize three pointers: low, mid, and high. As you iterate through the array:

- If the element is 0, swap it with the element at the low pointer and increment both low and mid.
 - If the element is 1, simply move mid forward.
 - If the element is 2, swap it with the element at the high pointer and decrement high.
- This single pass method ensures that the array is sorted in-place with constant extra space. Key details include correctly handling swaps without losing any data and ensuring that the pointers are updated in the proper order.

#161 One Edit Distance [Medium] · Two Pointers

Key Insights

- The difference in length between s and t must be at most 1.
- When lengths are equal, the only possibility is that exactly one character is different (replacement).
- When one string is longer by one character, the difference can be fixed by one insertion (or deletion in the other direction).
- Use two pointers to compare corresponding characters and carefully handle the first difference.

Space & Time

Time Complexity: $O(n)$, where n is the length of the shorter string, as we traverse both strings at most once.

Space Complexity: $O(1)$, using only constant extra space.

Solution

The solution uses a two-pointer approach. First, we check if the length difference is greater than 1, in which case they cannot be one edit apart. For the two primary cases:

- When strings are of equal length, iterate through both strings and count the number of differing characters. There should be exactly one difference.
- When lengths differ by one, use pointers for both strings. Upon encountering a difference, increment the pointer for the longer string only and ensure that no previous difference has been encountered. If a second discrepancy is found, return false. The key is to handle the edge case where the difference is at the end of the

string.

#167 Two Sum II – Input Array Is Sorted [Medium] · Two Pointers

Key Insights

- The array is sorted in non-decreasing order, which makes the two pointers technique very effective.
- Using two pointers (one at the start and one at the end) allows checking pairs and adjusting the pointers based on the sum.
- Since the problem guarantees exactly one solution and requires constant space, using two pointers satisfies both constraints.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in the array, because in the worst-case scenario each element is visited at most once.

Space Complexity: $O(1)$, as no additional data structure is used that scales with the input size.

Solution

We use the two pointers approach:

- Initialize two pointers, left at the beginning (index 0) and right at the end (last index) of the array.
- While left is less than right, compute the sum of the elements at the two pointers.
- If the sum equals the target, return the indices as [left+1, right+1] (to adjust for the

1-indexed array). – If the sum is less than the target, move the left pointer to the right to increase the sum. – If the sum is greater than the target, move the right pointer to the left to decrease the sum. This approach leverages the sorted property and uses constant extra space.

#186 Reverse Words in a String II [Medium] · Two Pointers

Key Insights

- Reverse the entire array first to obtain the reversed order of characters.
- Then, reverse each individual word to restore the correct letter order.
- Doing so leverages the two-pointer technique for in-place reversals.
- This method accomplishes the goal in $O(n)$ time while using $O(1)$ extra space.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution employs a two-step reversal process using the two-pointer technique. First, reverse the entire character array to reverse the overall order of characters. Although this reverses both the words and

the letters in each word, it sets up the scenario where each word's letters are in reverse order. The next step is to iterate through the array and locate each word (using the space character as a delimiter) and then reverse each word individually to correct its letter order. This process is done in-place without any additional data structures, conserving space.

#189 Rotate Array [Medium] · Two Pointers

Key Insights

- The rotation amount k might be greater than the length of the array, so use $k \% n$ (where n is the array length) to compute the effective rotation.
- A smart in-place solution involves reversing sections of the array:
 - Reverse the entire array.
 - Reverse the first k elements.
 - Reverse the remaining $n-k$ elements.
- Alternative approaches include using extra space or performing k single-step rotations (which is less optimal).

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the array.

Space Complexity: $O(1)$ if done in-place.

Solution

Round 9 · Low · Hard

p.3139

We solve the problem using the "reverse" technique to perform the rotation in-place. First, we adjust k by taking k modulo the length of the array to handle cases where k exceeds the array length. Then, we reverse the entire array. Next, we reverse the first k elements to restore their correct order. Finally, we reverse the remaining $n-k$ elements. This sequence yields the array rotated to the right by k positions while strictly using $O(1)$ extra space.

#532 K-diff Pairs in an Array [Medium] · Two Pointers

Key Insights

- For $k < 0$, the answer is always 0 because the absolute difference cannot be negative.
- When k is 0, you are looking for numbers that appear at least twice.
- For $k > 0$, it is efficient to use a hash set or a hash map to check if each number's complement ($\text{number} + k$) exists.
- Uniqueness is ensured by processing each number only once from the set of unique numbers.

Space & Time

Time Complexity: $O(n)$, where n is the length of the array since we iterate over the array and then over the unique numbers.

Round 9 · Low · Hard

p.3140

Space Complexity: $O(n)$, for storing the frequency count or the unique elements in a hash map or hash set.

Solution

The solution first checks if k is negative, returning 0 immediately if true. For k equals 0, we count numbers that appear more than once using a frequency map. For k greater than 0, a set of unique numbers is created and for each number, we check if $(\text{number} + k)$ is also present in the set. Each valid case contributes to the result count, ensuring that only unique pairs are considered.

#923 3Sum With Multiplicity [Medium] · Two Pointers

Key Insights

- Use a frequency counter for elements since $\text{arr}[i]$ is in the range $[0, 100]$.
- Iterate through all possible combinations of numbers x , y , and z with $x \leq y \leq z$.
- For each valid triplet (x, y, z) that sums to target, count the number of ways using combinatorial formulas:
 - All distinct: $\text{count}[x] * \text{count}[y] * \text{count}[z]$.
 - Two numbers the same: use $nC2$ for the repeated element.
 - All three the same: use $nC3$.

- Use modulo arithmetic ($\text{MOD} = 10^9 + 7$) at every step to avoid overflow.

Space & Time

Time Complexity: $O(M^2)$, where $M = 101$ (the range of possible numbers), which is efficient given the constraints.

Space Complexity: $O(M)$ for the frequency array.

Solution

The solution first counts the frequency of each number in the array. Then it iterates through all valid triplets (x, y, z) (with $x \leq y \leq z$) and checks if they sum up to the target. Depending on whether x , y , and z are distinct or have duplicates, the combination count is computed accordingly:

- If x , y , and z are all the same, compute the combination as $nC3$.
- If only two are the same, compute as $nC2$ multiplied by the count of the third number.
- If all are distinct, use the direct product of the counts. This covers all cases while applying the modulo operation to the result.

#986 Interval List Intersections [Medium] · Two Pointers

Key Insights

- Use two pointers to traverse both lists simultaneously.

- Determine the overlapping interval by taking the maximum of the start points and the minimum of the end points.
- If the overlapping condition ($\text{start} \leq \text{end}$) is met, add the intersection to the result.
- Move the pointer that has the smaller end value to explore further potential intersections efficiently.
- Achieve an overall time complexity of $O(m+n)$, where m and n are the lengths of the two interval lists.

Space & Time

Time Complexity: $O(m + n)$ where m and n are the lengths of `firstList` and `secondList` respectively.

Space Complexity: $O(1)$ extra space (excluding the space used for the output).

Solution

The solution uses a two pointers approach. Initialize two indices, one for each list, and iterate through both lists. For each pair of intervals, the potential intersection is computed by:

- Calculating the maximum of the two start values.
- Calculating the minimum of the two end values. If the maximum start is less than or equal to the minimum end, there is an intersection, and it is added to the result list. The pointer corresponding to the interval with the smaller end value is then incremented since that interval cannot intersect with any further intervals from the other list. This process repeats until one of the lists is

fully traversed.

#1055 Shortest Way to Form String [Medium] · Two Pointers

Key Insights

- Check if every character in target exists in source; if not, the answer is -1 .
- Use a greedy two-pointer strategy: iterate over target while scanning source for matching characters.
- During each pass over source, match as many characters as possible from target.
- Count the number of passes (subsequences) required until the target is completely matched.
- An optimized approach may precompute indices for binary search; however, the basic greedy simulation is efficient for the given constraints.

Space & Time

Time Complexity: $O(n * m)$ in the worst-case, where n = length of source and m = length of target.

Space Complexity: $O(1)$ for the greedy simulation ($O(n)$ if additional data structures for binary search are used).

Solution

The solution uses a greedy two-pointer approach. We simulate the process of reading the target by repeatedly scanning through the source string. Each scan through source is treated as one subsequence concatenated to

form the target. A pointer in the target moves forward whenever a matching character is found in source. If a complete pass through source does not advance the target pointer, then it is impossible to form the target and we return -1. The algorithm counts the number of passes required until the target is fully matched.

#2563 Count the Number of Fair Pairs [Medium] · Two Pointers

Key Insights

- Sorting the array simplifies searching for pairs with required sum limits.
- For each element, binary search is used to locate:
- The first index where the complement makes the sum $\geq$ lower.
- The last index where the complement makes the sum $\leq$ upper.
- This avoids the need for a nested loop through the entire array, yielding a more efficient solution.

Space & Time

Time Complexity: $O(n \log n)$, primarily due to sorting and (for each element) binary search operations.

Space Complexity: $O(n)$ if sorting creates a copy; $O(1)$ additional space if sorting in-place.

Solution

Round 9 · Low · Hard

p.3145

The solution starts by sorting the input array. For each element at index i in the sorted array, we search for indices j (where $j > i$) such that the sum of $\text{nums}[i]$ and $\text{nums}[j]$ falls between lower and upper. Two binary search operations are performed:

- To find the smallest index j where $\text{nums}[i] + \text{nums}[j] \geq \text{lower}$.
 - To find the largest index j where $\text{nums}[i] + \text{nums}[j] \leq \text{upper}$.
- Subtracting the two positions (and summing the counts for each i) gives the total number of fair pairs. Edge cases (like no valid pair) are automatically handled by the binary search approach.

#3 Longest Substring Without Repeating Characters [Medium] · Sliding Window

Key Insights

- Use the sliding window technique to maintain a window of unique characters.
- Use a hash table (or dictionary) to quickly check for duplicate characters and update window boundaries.
- The window is expanded by moving a pointer (right) and contracted by moving another pointer (left) when duplicates are encountered.
- Time complexity can be maintained at $O(n)$ by ensuring each character is processed a limited number of times.

Space & Time

Round 9 · Low · Hard

p.3146

Time Complexity: $O(n)$, where n is the length of the input string.

Space Complexity: $O(\min(n, m))$, where m is the size of the character set (constant for fixed character sets).

Solution

The solution uses a sliding window technique with two pointers, left and right, to keep track of the current substring without repeating characters. A hash table (dictionary or map) is used to store the most recent index of each character encountered. As the right pointer moves forward, if a duplicate character is found within the current window, the left pointer is advanced to the position right after where the duplicate was last seen. This ensures that the window always contains unique characters. The maximum length is updated throughout the process.

#159 Longest Substring with At Most Two Distinct Characters [Medium] · Sliding Window

Key Insights

- Use a sliding window to maintain a contiguous substring.
- Keep track of character frequencies within the current window using a hash map or dictionary.
- When the window contains more than two distinct characters, shrink the window from the left until only two remain.

- Update the maximum length each time the window satisfies the condition.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, since each character is processed at most twice.

Space Complexity: $O(1)$ because the hash map holds at most 3 entries (usually up to 2 distinct characters).

Solution

We use the sliding window technique with two pointers (left and right) to represent the current window. A hash table/dictionary records the count of each character inside the window. As we extend the right pointer, we update the count. If the hash table contains more than two distinct characters, we start moving the left pointer to reduce the number of distinct characters until we have at most two distinct keys. The length of the current valid window is then used to update our maximum length answer.

#340 Longest Substring with At Most K Distinct Characters [Medium] · Sliding Window

Key Insights

- Use the sliding window technique to efficiently scan through the string.
- Maintain a hash table (or dictionary) that keeps track of the frequency of characters in the current window.

- When the number of distinct characters exceeds k , increment the left pointer to shrink the window until the condition is met.
- Update the maximum length of valid substrings throughout the process.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, as each character is processed at most twice (once when expanding the window and once when contracting it).

Space Complexity: $O(\min(m, k))$, where m is the size of the charset (or number of unique characters in s), since at most $k+1$ keys are stored.

Solution

This problem is solved using the sliding window technique. The algorithm keeps track of a window defined by two indices (left and right pointers). A hash map is used to count the frequency of each character in the current window. As the right pointer extends the window, the count of the character is incremented. When the count of distinct keys in the map exceeds k , the left pointer is moved forward while decrementing the count of the leftmost character, until the condition is restored. The maximum window length encountered is recorded and returned as the result.

#424 Longest Repeating Character Replacement [Medium] · Sliding Window

Key Insights

- Use a Sliding Window algorithm to keep track of a contiguous substring.
- Maintain a frequency count of letters in the current window.
- Keep track of the count of the most frequent character in the window.
- If the current window size minus the maximum frequency is greater than k , shrink the window.
- The maximum window size seen during the process is the answer.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, since the window traverses the string in linear time.

Space Complexity: $O(1)$ because the frequency count for letters (A-Z) has a constant size.

Solution

The problem is solved using the sliding window technique. We initialize two pointers representing the left and right bounds of the current window. As we expand the window by moving the right pointer, we update the frequency count of the letter at the right pointer and track the letter that appears most

frequently. The key observation is that the number of characters that need to be replaced in the current window is the window's size minus the count of the most frequently occurring letter. If this number exceeds k , we increment the left pointer to shrink the window until the condition is satisfied again. The maximum size of the window during the process gives the answer.

#438 Find All Anagrams in a String [Medium] · Sliding Window

Key Insights

- Use a sliding window of length equal to p over s .
- Maintain frequency counts of characters in p and within the current window in s .
- Compare frequency counts to determine if the current window is an anagram.
- Update the window efficiently by adding one character and removing the character that goes out of the window.

Space & Time

Time Complexity: $O(n)$ where n is the length of s .

Space Complexity: $O(1)$ since we use fixed-size frequency arrays or hash maps.

Solution

We use the sliding window technique to iterate over s with a window size equal to the length of p . First, count the frequencies of characters in p . Then, iterate over s

and update a frequency count for the current window. When the window size equals the length of p , compare the two frequency counts. If they match, record the start index. To slide the window efficiently, increment the count for the incoming character and decrement the count for the outgoing character. This way, we achieve an $O(n)$ solution that is optimal for the problem constraints.

#487 Max Consecutive Ones II [Medium] · Sliding Window

Key Insights

- Utilize a sliding window to maintain the current segment.
- The window can contain at most one 0.
- Expand the window with the right pointer and contract it with the left pointer when the constraint is violated.
- The maximum window size observed during this process is the desired result.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in the array.

Space Complexity: $O(1)$.

Solution

We use a sliding window approach to efficiently find the longest contiguous subarray that contains at most one

0. Two pointers (left and right) define the window. As we iterate with the right pointer, counting zeros in the window, we adjust the left pointer when the count of zeros exceeds one. The length of the valid window gives us the number of consecutive 1's (considering one flipped 0), and we update the maximum length accordingly.

#1100 Find K-Length Substrings with No Repeated Characters [Medium] · Sliding Window

Key Insights

- Use a sliding window of fixed size k.
- Track character counts within the current window using a hash table (or frequency array).
- When a duplicate is found in the window, slide the window from the left to remove duplicates.
- Count a valid window when its size equals k and all characters are unique.
- Optimize by adjusting the window immediately after counting a valid substring.

Space & Time

Time Complexity: $O(n)$ where n is the length of s, as each character is visited at most twice.

Space Complexity: $O(1)$ since the frequency tracking is over a fixed set (26 lowercase letters).

Solution

Round 9 · Low · Hard

p.3153

We use a sliding window approach with two pointers: left and right. A hash table (or frequency array) tracks the occurrences of characters in the window. As we move the right pointer, we update the frequency and check for duplicates. If a duplicate is encountered, we increment the left pointer until the duplicate is removed. Once the window size exactly equals k, we have a valid substring, so we increment our count and then slide the window forward. This ensures that every possible k-length substring is examined once.

#1248 Count Number of Nice Subarrays [Medium] · Sliding Window

Key Insights

- Convert the problem to counting subarrays with a specific "prefix sum" value where the sum represents the count of odd numbers.
- Use a hash table (or dictionary) to store the frequency of prefix sums encountered.
- The number of nice subarrays ending at the current index equals the frequency of (current prefix sum - k) seen so far.
- This approach allows you to find the answer in one pass over the array.

Space & Time

Time Complexity: $O(n)$, where n is the number of elements in nums, because we traverse the array once.

Round 9 · Low · Hard

p.3154

Space Complexity: $O(n)$, in the worst case where each prefix sum is unique and stored in the hash table.

Solution

We solve the problem using a prefix sum approach. First, we iterate through the array and convert it into a prefix sum representing the cumulative count of odd numbers encountered. As we compute the prefix sums, we use a hash table to record how many times each prefix sum appears. For every current prefix sum value, we check how many times we have seen a prefix sum equal to (current prefix sum - k). The idea is that if $\text{prefix}[j] - \text{prefix}[i]$ equals k, then between indices $i+1$ and j , there are exactly k odd numbers. This gives us the count of valid subarrays ending at the current index.

This method works efficiently as it requires only a single pass through the array and maintains a running frequency count using the hash table.

#49 Group Anagrams [Medium] · Sorting

Key Insights

- Anagrams have the same characters when sorted alphabetically.
- Sorting each string provides a unique key for grouping.
- Use a hash table (dictionary) to map each sorted key to a list of its original strings.

- Finally, gather all the lists from the dictionary to form the result.

Space & Time

Time Complexity: $O(n * k \log k)$, where n is the number of strings and k is the maximum length of a string (due to sorting each string).

Space Complexity: $O(n * k)$ for storing the grouped strings in the hash table.

Solution

The primary idea is to iterate through each string, sort it to obtain a key that represents its character composition, and then group the original string under this key in a hash table. After processing all strings, the values of the hash table will be the groups of anagrams. This approach efficiently groups anagrams by leveraging sorting and a dictionary for constant-time lookups.

#56 Merge Intervals [Medium] · Sorting

Key Insights

- Sort the intervals based on their start values.
- Use a pointer or iteration to compare the current interval with the last merged interval.
- If the current interval overlaps with the previous one, merge them by updating the end time.

- Otherwise, add the current interval to the results as a new merged interval.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(n)$ to store the merged intervals.

Solution

First, sort the array of intervals based on the starting time. Then, iterate through the sorted list and check if the current interval overlaps with the last interval in the merged list. If an overlap is detected (i.e., the start of the current interval is less than or equal to the end of the last merged interval), update the last merged interval's end to be the maximum of both ends. If there is no overlap, simply add the current interval to the merged list. This approach efficiently merges overlapping intervals and ensures that all intervals are processed in a single pass after sorting.

#1152 Analyze User Website Visit Pattern [Medium]

· Sorting

Key Insights

- Combine the inputs into a single list of events and sort them by timestamp.
- Group the visits by each user, preserving the sorted order.

Round 9 · Low · Hard

p.3157

- Generate all unique combinations of 3-sequences for each user (order matters and non-contiguous visits are allowed).
- Use a dictionary (or hash map) to map each 3-sequence pattern to the set of users who visited that pattern.
- The problem constraints are small, so generating combinations per user (even using triple nested loops) is acceptable.
- When multiple patterns have the same frequency, choose the lexicographically smallest one.

Space & Time

Time Complexity: $O(U * L^3)$ where U is the number of users and L is the maximum length of visits per user (since we generate combinations of 3 visits per user). Sorting all events costs $O(N \log N)$ for N total events.

Space Complexity: $O(P)$ where P is the number of unique 3-sequence patterns stored in the dictionary, plus $O(N)$ for grouping the events.

Solution

– Aggregate the events by user while sorting them based on timestamp. – For each user, generate all sets of 3-website patterns (to avoid duplicate patterns per user). – Use a hash map to count the number of unique users for each pattern. – Traverse the map to find the pattern with the highest count. In case of a tie, compare patterns lexicographically. – Return the pattern that

Round 9 · Low · Hard

p.3158

meets the criteria.

#33 Search in Rotated Sorted Array [Medium] · Binary Search

Key Insights

- The array is originally sorted and then rotated at an unknown pivot.
- Use a modified binary search to accommodate the rotation.
- At every step, determine which half of the array is properly sorted.
- Narrow down the search based on the sorted half and the target's value.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

We perform a binary search while accounting for the possibility that the array is rotated. At each iteration, we check if the left-to-mid portion is sorted. If it is, and the target falls within that range, we search the left half; otherwise, we search the right half. If the left half is not sorted, then the right half must be sorted, and similarly, we check if the target is in that sorted half. This approach ensures logarithmic time complexity while handling the rotation correctly.

#34 Find First and Last Position of Element in Sorted Array [Medium] · Binary Search

Key Insights

- The array is sorted, allowing the use of binary search.
- Two binary searches can efficiently find the leftmost (first) and rightmost (last) indices of the target.
- Handling edge cases where the target is not present by checking bounds.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

The approach is to use binary search twice:

- The first binary search finds the leftmost occurrence of the target by continuing to search in the left half even after finding the target.
- The second binary search finds the rightmost occurrence by searching in the right half. The algorithm maintains pointers for the current search bounds and narrows the search space using comparisons. If the target is not found during the first binary search, the solution immediately returns $[-1, -1]$. This method leverages the sorted property of the array, ensuring $O(\log n)$ time complexity with constant extra space.

#153 Find Minimum in Rotated Sorted Array [Medium] · Binary Search

Key Insights

- The rotated sorted array consists of two sorted subarrays.
- A modified binary search can be used to locate the pivot point where the rotation occurs.
- By comparing the middle element to the rightmost element, we can decide which half of the array contains the minimum.

Space & Time

Time Complexity: $O(\log n)$

Space Complexity: $O(1)$

Solution

Use a binary search approach. Initialize two pointers, left and right, at the start and end of the array. While left is less than right, calculate the middle index. If the middle element is greater than the right element, then the minimum is in the right half, so move the left pointer to $\text{mid} + 1$; otherwise, the minimum must be in the left half including mid, so adjust the right pointer to mid. When the loop terminates, left equals right and points to the minimum element. This approach leverages the properties of the rotated sorted array to achieve logarithmic time complexity.

#300 Longest Increasing Subsequence [Medium] · Binary Search

Key Insights

- A dynamic programming solution with $O(n^2)$ time complexity exists, but we focus on an optimized approach.
- Using binary search and a dp array, we can achieve $O(n \log n)$ time.
- The dp array maintains the smallest possible tail for increasing subsequences of various lengths.
- For each number, perform binary search to decide if it can extend or replace a value in dp.

Space & Time

Time Complexity: $O(n \log n)$ where n is the length of nums

Space Complexity: $O(n)$ for storing the dp array

Solution

We use a dynamic programming approach combined with binary search. The idea is to maintain a dp array where $\text{dp}[i]$ holds the smallest tail value for any increasing subsequence with length $i+1$. For each number in the array, we perform a binary search on dp to find its correct insertion position:

- If the number is larger than all elements, we append it, effectively extending the current longest subsequence.
- Otherwise, we replace the first element in dp that is

greater than or equal to the current number to keep the tail as small as possible. This process ensures that `dp` remains sorted and its length gives the length of the longest increasing subsequence.

#362 Design Hit Counter **[Medium]** · Binary Search

Key Insights

- Use a sliding window of 300 seconds to count hits.
- Since timestamps are monotonically increasing, outdated hits can be removed from the front.
- A queue (or circular array) is an efficient data structure to maintain and remove expired hit records.
- Grouping consecutive hits at the same timestamp can optimize space.

Space & Time

Time Complexity: $O(1)$ per hit and `getHits` operation on average, since each hit is enqueued and later dequeued only once.

Space Complexity: $O(W)$ where W is the window size (300 seconds) in the worst case.

Solution

The solution uses a queue (or deque) data structure where each element stores a (timestamp, count) pair. When a hit is recorded, if the last element of the queue shares the same timestamp, then its count is incremented; otherwise, a new element is added. When

`getHits(timestamp)` is called, the algorithm removes elements from the front of the queue whose timestamp is out of the 300-second window relative to the current timestamp, and then sums the counts of the remaining elements to return the result. This ensures that only relevant hits within the past 5 minutes are counted. The design scales well when hit frequency is high because hits at the same timestamp are grouped together instead of inserting individual records.

#528 Random Pick with Weight **[Medium]** · Binary Search

Key Insights

- Build a prefix sum array from the weights.
- Calculate the total sum of the weights.
- Generate a random number in the range $[1, \text{total sum}]$ and use binary search to find the corresponding index in the prefix sum array.
- Using binary search ensures efficient selection in $O(\log n)$ time for each call of `pickIndex()`.

Space & Time

Time Complexity: $O(n)$ for constructing the prefix sum array, and $O(\log n)$ per `pickIndex()` call.

Space Complexity: $O(n)$ for storing the prefix sum array.

Solution

The solution leverages a prefix sum array where each entry at index i represents the cumulative sum of weights up to i . After computing the total weight, a random integer is generated between 1 and the total sum (inclusive). A binary search is then used to find the first index where the prefix sum is greater than or equal to the random number. This approach ensures that the index is selected with a probability proportional to its weight.

#852 Peak Index in a Mountain Array [Medium] · Binary Search

Key Insights

- The array is strictly increasing up to a peak and then strictly decreasing.
- A binary search technique is applicable because the array exhibits a single transition point (peak) between increasing and decreasing order.
- By comparing mid and $mid+1$, we can decide which half of the array to continue the search in.

Space & Time

Time Complexity: $O(\log(n))$ – The binary search approach reduces the search interval by half at each step.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We use a binary search strategy:

- Initialize two pointers, left and right, at the beginning and end of the array.
- In each iteration, calculate mid. If $arr[mid]$ is less than $arr[mid + 1]$, then the peak is in the right half; otherwise, it is in the left half.
- When left equals right, we have found the peak index. This approach leverages the mountain array properties and ensures an $O(\log(n))$ solution.

#981 Time Based Key-Value Store [Medium] · Binary Search

Key Insights

- Use a dictionary (hash table) to map each key to a list of (timestamp, value) pairs.
- Timestamps are strictly increasing when setting the values, which allows efficient appending.
- For getting the value, perform a binary search on the list of timestamps to find the largest timestamp that is $\leq$ the queried timestamp.
- If no valid timestamp exists, return an empty string.

Space & Time

Time Complexity:

- set operation: $O(1)$ per call (amortized).
- get operation: $O(\log N)$ per call, where N is the number of timestamps stored for the key.

Space Complexity: $O(\text{total number of set calls})$, to store all (timestamp, value) pairs.

Solution

We use a hash table to store keys mapping to a list of their timestamps and corresponding values. Since the timestamps are inserted in strictly increasing order, the list remains sorted. For the get operation, a binary search is applied to quickly locate the appropriate value for a given timestamp. This approach minimizes the lookup time and efficiently handles the problem constraints.

#1011 Capacity to Ship Packages Within D Days

[Medium] · Binary Search

Key Insights

- The minimal possible capacity is the maximum value in weights; no ship can carry a package heavier than its capacity.
- The maximum possible capacity is the sum of all weights; this allows shipping all packages in one day.
- The problem is monotonic: if a capacity C works, any capacity greater than C will also work.
- Use binary search on the capacity range and a greedy approach to simulate shipping for each candidate capacity.

Space & Time

Time Complexity: $O(n * \log(\text{sum}(\text{weights})))$ where n is the number of packages.

Space Complexity: $O(1)$ additional space.

Solution

The solution employs binary search to find the minimum ship capacity. Start by setting the search bounds: the lower bound is the maximum weight (since each package must be shipped individually if needed) and the upper bound is the sum of all weights (shipping everything in one go). For each midpoint capacity in the binary search, use a greedy method to simulate the shipping process—accumulating weights until adding another package would exceed the candidate ship capacity, at which point you increment the day count. If the number of days required exceeds D , the capacity is insufficient and the lower bound is increased; otherwise, adjust the upper bound. Continue until the bounds converge.

#1060 Missing Element in Sorted Array **[Medium]** · Binary Search

Key Insights

- The array is sorted and every element is unique.
- For any index i , the number of missing numbers up to that point can be computed as $\text{nums}[i] - \text{nums}[0] - i$.
- When $\text{missing}(i)$ is less than k , the k th missing element is further to the right.

- Use binary search to find the smallest index where the count of missing numbers is greater than or equal to k .
- If all missing numbers before the last array element are counted and the k th missing still hasn't been reached, calculate the result by continuing past the last element.

Space & Time

Time Complexity: $O(\log(n))$

Space Complexity: $O(1)$

Solution

We compute the number of missing elements until an index i with the formula: $\text{missing}(i) = \text{nums}[i] - \text{nums}[0] - i$. We then use binary search over the indices to find the smallest index where $\text{missing}(i) \geq k$. If the binary search index equals the length of the array, it means the k th missing number lies beyond the last element. Otherwise, the k th missing number is derived as: $\text{nums}[\text{index} - 1] + (k - \text{missing}(\text{index} - 1))$. This approach uses binary search for logarithmic time and does not require extra auxiliary data structures.

#1062 Longest Repeating Substring [Medium] · Binary Search

Key Insights

- Use binary search to efficiently check for the existence of a repeating substring of a given length.

- Apply a rolling hash technique to compute hash values for substrings in $O(1)$ time after initial processing.
- Use a hash set to detect duplicate hash values (indicating a repeating substring).
- Adjust binary search bounds based on whether a repeating substring of the current length exists.
- Alternative approaches include dynamic programming and suffix arrays, but the binary search with rolling hash is efficient and conceptually clean.

Space & Time

Time Complexity: $O(n \log n)$ on average, where n is the length of the string. (Each binary search step takes $O(n)$ time for computing rolling hashes.)

Space Complexity: $O(n)$ for storing hash values in the hash set.

Solution

The solution uses binary search to decide on the candidate length L of the repeating substring. For each candidate length L , we use a rolling hash to compute hash values for all substrings of length L and store them in a hash set. If any hash is repeated, it confirms the presence of a repeating substring of length L , and we try a longer length. Otherwise, we decrease the candidate length. The rolling hash is computed using a base (26, for lowercase letters) and a modulus (2^{32} to limit integer overflow) which allows efficient hash recomputation for sliding windows. This approach

efficiently narrows down the maximum repeating substring length.

#1533 Find the Index of the Large Integer [Medium]

· Binary Search

Key Insights

- Use binary search to efficiently narrow down the position of the unique large value.
- At each step, divide the candidate range into two halves (or almost two halves when the range size is odd) and use `compareSub` to determine which half contains the large integer.
- Comparing two subarrays of equal length will reveal the region that contains the extra value via the sum difference.
- Handle odd-length ranges carefully by isolating the middle element if needed.

Space & Time

Time Complexity: $O(\log n)$ – because binary search is performed over the range.

Space Complexity: $O(1)$ – only a few pointers are maintained, no additional space is required.

Solution

The solution applies a binary search approach over the array indices. At every iteration:

– Calculate the current segment length. – If the length is even, split the segment evenly and compare the left half versus the right half using `compareSub`. – If the length is odd, compare two halves of equal size while ignoring the middle element. If the sums are equal, then the middle index holds the large value. – Update the search bounds based on the result of `compareSub`. – Continue until the search space is reduced to one index, which is the answer.

This method ensures that you locate the large integer with $O(\log n)$ calls to `compareSub`, well within the limit of 20 calls.

#36 Valid Sudoku [Medium] · Matrix

Key Insights

- Validate each row by checking for duplicate numbers while ignoring empty cells.
- Validate each column in the same way.
- Divide the board into nine 3×3 sub-boxes and check each for duplicate numbers.
- Use hash sets (or equivalent data structures) to keep track of seen digits for rows, columns, and boxes.
- The board size is fixed (9×9), so the overall time complexity is constant with respect to input size.

Space & Time

Time Complexity: $O(1)$ — Since the board size is fixed at 9×9 .

Space Complexity: $O(1)$ — The extra space used by the sets is bounded by the fixed board size.

Solution

The solution uses three main data structures: an array of sets for rows, an array of sets for columns, and an array of sets for 3×3 sub-boxes. For each cell, if it is not empty, check if the number is already present in the corresponding row, column, or box. If found, return false immediately. Otherwise, add the number to the respective sets. Compute the index for the 3×3 sub-box using the formula: $(\text{row_index} // 3) * 3 + (\text{col_index} // 3)$. If no duplicates are found after traversing the entire board, the board is considered valid.

#48 Rotate Image [Medium] · Matrix

Key Insights

- The rotation can be achieved in two steps: first transpose the matrix and then reverse each row.
- Transposing the matrix swaps $\text{matrix}[i][j]$ with $\text{matrix}[j][i]$, converting rows into columns.
- Reversing each row post-transposition yields the desired clockwise rotation.
- This approach works in-place with $O(1)$ extra space.

Space & Time

Time Complexity: $O(n^2)$ since every element is processed during the transpose and row reversal.

Space Complexity: $O(1)$ as the operations are performed in-place without extra data structures.

Solution

The solution leverages a two-step in-place transformation:

- Transpose the matrix: Iterate through the matrix and swap each element (i, j) with element (j, i) for $j \geq i$.
- Reverse each row: After the transpose, each row is reversed to rotate the matrix 90 degrees clockwise. The key trick is to realize that a transpose followed by a reverse of each row achieves the rotation without needing extra space.

#54 Spiral Matrix [Medium] · Matrix

Key Insights

- Use four boundaries: top, bottom, left, and right to track the current layer.
- Traverse right across the top row, then down the right column, then left on the bottom row, and finally up the left column.
- After each direction, adjust the corresponding boundary to move inward.
- Continue the traversal until the boundaries cross.

Space & Time

Time Complexity: $O(m * n)$ since each element is visited exactly once.

Space Complexity: $O(1)$ extra space (excluding the space required for the output list).

Solution

The solution uses a simulation approach with boundary pointers (top, bottom, left, right). At each step, the boundaries indicate the current submatrix to traverse.

You iterate as follows:

- Traverse from left to right along the top boundary.
 - Traverse downward along the right boundary.
 - If the top boundary has not passed the bottom, traverse from right to left along the bottom boundary.
 - If the left boundary has not passed the right, traverse upward along the left boundary.
- After each traversal, the boundaries are adjusted inward. This continues until all elements have been visited.

#59 Spiral Matrix II [Medium] · Matrix

Key Insights

- Use four pointers (top, bottom, left, right) to keep track of the matrix boundaries.
- Traverse the matrix in a spiral: move right along the top row, then down the right column, left along the bottom row, and finally up the left column.
- Update the boundary pointers after each complete traversal of one side.

Round 9 · Low · Hard

p.3175

- Continue until all numbers from 1 to n^2 are filled into the matrix.

Space & Time

Time Complexity: $O(n^2)$

Space Complexity: $O(n^2)$ (required for the output matrix)

Solution

We simulate the spiral movement by initializing four pointers: top, bottom, left, and right, which denote the current boundaries of the matrix that have not yet been filled. For each iteration, we:

- Traverse from left to right along the top boundary, then increment the top pointer.
 - Traverse from top to bottom along the right boundary, then decrement the right pointer.
 - If there are rows remaining, traverse from right to left along the bottom boundary, then decrement the bottom pointer.
 - If there are columns remaining, traverse from bottom to top along the left boundary, then increment the left pointer.
- This algorithm ensures that the matrix is filled in a spiral order until all numbers are placed.

#64 Minimum Path Sum [Medium] · Matrix

Key Insights

- The problem can be solved using Dynamic Programming.

Round 9 · Low · Hard

p.3176

- The optimal solution for each cell depends on the minimum path sum of the cell above it and the cell to its left.
- Since you can only move right or down, only two adjacent cells affect the current cell's result.
- Edge cases include the first row and first column, which have only one direction of movement.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the number of rows and columns.

Space Complexity: $O(m * n)$ for the DP table (this can be optimized to $O(n)$ with in-place computations).

Solution

We use a Dynamic Programming (DP) approach where we create a DP table (or use the grid itself) that stores the minimum path sum to each cell. For each cell (i, j) :

- If $i = 0$ (first row), the only possible path is from the left.
- If $j = 0$ (first column), the only possible path is from above.
- For other cells, the minimum path sum is the cell value plus the minimum of the two possible previous cells (above and left). This ensures that by the time we reach the bottom-right cell, we have accumulated the minimum sum required to reach that cell.

#73 Set Matrix Zeroes [Medium] · Matrix

Round 9 · Low · Hard

p.3177

Key Insights

- The challenge is to mark rows and columns that need to be zeroed without using extra space.
- A common trick is to use the first row and first column as markers.
- Special care is needed for the first row and first column, as they are also used as markers.
- The algorithm involves two passes: one for marking and one for updating the matrix based on the markers.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(1)$ (in-place, using only constant extra space)

Solution

We can solve this problem by using the first row and first column of the matrix as markers. First, determine if the first row or first column needs to be zeroed, then traverse the rest of the matrix and mark the corresponding first row and column positions if a zero is encountered. Finally, update the matrix backward using the markers, and finally update the first row and column if needed.

The algorithm works in three main steps:

- Check if the first row or first column should be zeroed (store this in two variables).
- Iterate through the rest of the matrix; for any element that is zero, mark its row

Round 9 · Low · Hard

p.3178

and column by setting the corresponding element in the first row and first column to zero. – Iterate through the matrix again (excluding the first row and column) and set elements to zero if their row or column marker is zero. – Finally, zero out the first row and/or first column if needed.

This approach avoids extra space usage apart from the two boolean flags while ensuring that the matrix is updated correctly.

#74 Search a 2D Matrix [Medium] · Matrix

Key Insights

- The matrix is essentially a flattened sorted array due to its properties.
- Use binary search to efficiently determine if the target exists.
- Map the 1D binary search index to the corresponding 2D matrix coordinates using division and modulus operations.
- Achieve $O(\log(m*n))$ time complexity by navigating the matrix as if it were a single sorted list.

Space & Time

Time Complexity: $O(\log(m * n))$

Space Complexity: $O(1)$

Solution

We can treat the matrix as a flattened sorted list without physically copying the elements. Given the total number of elements = $m * n$, perform binary search on indices ranging from 0 to $(m*n - 1)$. For any index mid , calculate its corresponding row as $mid // n$ and column as $mid \% n$. Compare the element at $matrix[row][col]$ with the target. If it equals target, return true; if it is less than target, search the right half; otherwise, search the left half.

#221 Maximal Square [Medium] · Matrix

Key Insights

- Use dynamic programming to build a solution where each cell in the DP table represents the side length of the largest square ending at that cell.
- If the current cell is '1', its DP value is the minimum of the top, left, and top-left neighbors plus one.
- Update a global maximum side length during the process to determine the area.
- The square's area is the square of its maximum side length.

Space & Time

Time Complexity: $O(m * n)$, where m is the number of rows and n is the number of columns.

Space Complexity: $O(m * n)$ for the DP table (this can be optimized to $O(n)$ with space reduction techniques).

Solution

We use a two-dimensional dynamic programming approach. We define a DP matrix with one extra row and column to simplify boundary conditions. For each cell (i, j) in the input matrix, if the value is '1', then we calculate the DP value as the minimum of its top, left, and top-left neighboring DP values plus one. This value represents the side length of the square ending at (i, j) . We also keep track of the maximum side length found and finally compute the area of the maximal square by squaring the maximum side length.

#311 Sparse Matrix Multiplication [Medium] · Matrix

Key Insights

- Exploit sparsity by iterating only over non-zero elements.
- For each non-zero element in mat1, multiply it with corresponding non-zero elements in mat2.
- Accumulate the product into the result matrix.
- This approach reduces unnecessary multiplication operations when many elements are zero.

Space & Time

Time Complexity: $O(m * k * n)$ in the worst-case, but significantly lower when many elements are zero.

Space Complexity: $O(m * n)$ for the resulting matrix.

Solution

We loop through each row of mat1 and for every non-zero element, we traverse the corresponding row in mat2. By checking if an element in both matrices is non-zero before multiplying, we avoid redundant calculations. This optimization leverages the sparsity of the matrices. We use standard nested loops along with conditional checks to implement this efficient multiplication.

#498 Diagonal Traverse [Medium] · Matrix

Key Insights

- Traverse the matrix diagonally with alternating directions: upward (top-right) and downward (bottom-left).
- When the traversal reaches a boundary (top, bottom, left, or right), adjust the position and change the direction.
- Use simulation to process each element exactly once while managing the indices based on the current direction.

Space & Time

Time Complexity: $O(m * n)$ where m is the number of rows and n is the number of columns, as every element is processed once.

Space Complexity: $O(m * n)$ for the output array (excluding the input matrix), while additional space usage remains constant.

Solution

The solution simulates the diagonal traversal using a direction flag (1 for upward and -1 for downward).

- Begin at the matrix's top-left corner.
- Append the current element to the result.
- Depending on the current direction: If moving upward, decrease the row and increase the column. If moving downward, increase the row and decrease the column.
- Adjust for boundary hits: When moving upward: If at the last column, increment the row and switch direction. If at the first row, increment the column and switch direction. When moving downward: If at the last row, increment the column and switch direction. If at the first column, increment the row and switch direction. This process continues until all elements are added to the result array.

#531 Lonely Pixel I [Medium] · Matrix

Key Insights

- Count the number of 'B' pixels in each row.
- Count the number of 'B' pixels in each column.
- A pixel is considered lonely if the count in its corresponding row and column is exactly 1.
- Two passes through the matrix can efficiently solve the problem: one for counting and one for checking conditions.

Space & Time

Round 9 · Low · Hard

p.3183

Time Complexity: $O(m * n)$, where m and n are the dimensions of the picture.

Space Complexity: $O(m + n)$, used to store the row and column counts.

Solution

The solution works by first scanning the entire matrix to compute the number of black pixels in each row and each column using two arrays/lists. Once these counts are available, a second pass through the matrix determines for each black pixel if it is lonely by checking if its corresponding row and column count equals 1. The approach uses simple array data structures for tallying counts and ensures that each element is processed only a constant number of times.

#723 Candy Crush [Medium] · Matrix

Key Insights

- Use a simulation approach: repeatedly detect crushable candies and simulate their removal.
- Traverse horizontally and vertically to mark candies that are in groups of three or more.
- Instead of removing candies immediately, mark them for removal and update the board once for the entire pass.
- After crushing, simulate gravity by shifting non-zero candies downward in each column.

Round 9 · Low · Hard

p.3184

- Continue the process until no further candy crushes can be detected.

Space & Time

Time Complexity: $O((m$

$n)^2)$ in the worst case since each iteration includes scanning the $m \times n$ grid and then applying gravity.

Space Complexity: $O(m$

$n)$ for maintaining additional markers or flags (or using the board itself with in-place modifications).

Solution

The approach involves iteratively scanning the board to identify candy groups that should be crushed using two passes:

– Horizontal Check – For each row, check segments of 3 consecutive candies. – Vertical Check – For each column, check segments of 3 consecutive candies. If any candy qualifies, mark its position for crushing. After marking, update the board by turning all marked positions into 0. Then, for each column, drop down non-zero candies such that all empty cells (zeros) are at the top. Repeat this process until no group of 3 or more adjacent candies is found. Key data structures include the board array and a temporary marker (can be implicit by using signs or a separate boolean matrix) for tracking candies to crush. The algorithm ensures a stable board before returning it.

#835 Image Overlap [Medium] · Matrix

Key Insights

- You only need to consider the positions of 1's in each matrix.
- The overlap count for a given translation can be computed by comparing the relative differences between the positions of 1's in the two images.
- Using a hash map (or dictionary) to count how many times each translation vector appears lets you compute the maximum overlap efficiently.
- The translation operation is equivalent to shifting the positions, so the best possible translation is the one that aligns the most pairs of 1's.

Space & Time

Time Complexity: $O(n^4)$ in the worst-case when using a nested iteration over positions for both matrices, though the average situation is typically better since not all cells are 1.

Space Complexity: $O(n^2)$ for storing the positions of 1's and the counts for translation vectors.

Solution

The solution involves the following steps:

– Traverse both images to record the coordinates of all 1's. – For every pair of 1's (one from the first image, one from the second), compute the translation vector (offset) that aligns them. – Use a hash map to count how many

times each offset appears. – The maximum count recorded in the hash map is the largest overlap achievable by any translation. This approach avoids simulating every possible slide over the matrix and focuses only on relative alignments of 1's.

#1198 Find Smallest Common Element in All Rows [Medium] · Matrix

Key Insights

- Since every row is strictly increasing, binary search can be used efficiently to check for the presence of an element in a row.
- Iterating over the first row's elements and verifying their presence in all other rows is an effective approach.
- Early termination is possible if an element is not found in any one row.
- Time complexity is manageable given the constraints, even using binary search for each candidate across rows.

Space & Time

Time Complexity: $O(m * n * \log(n))$ where m is the number of rows and n is the number of columns (binary search in each row for each candidate from the first row).

Space Complexity: $O(1)$ (only a few auxiliary variables are used).

Solution

The solution iterates over each element (candidate) in the first row of the matrix. For each candidate, it checks if the candidate is present in every other row using binary search, taking advantage of the sorted order. If a candidate is found in every row, return it immediately as it will be the smallest such common element because the first row is processed in order. If no candidate is common across all rows, return -1.

#98 Validate Binary Search Tree [Medium] · Trees & BST

Key Insights

- Every node must adhere to a range: all nodes in the left subtree should be less than the current node, and all nodes in the right subtree should be greater.
- A recursive depth-first search approach works well by passing down the valid range (lower and upper bounds) for each node.
- Using the limits (negative and positive infinity) simplifies handling edge cases at the tree root.
- An in-order traversal would alternatively yield a sorted order for node values, but updating ranges directly is more intuitive for recursive implementation.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes, as each node is visited once.

Space Complexity: $O(H)$, where H is the height of the tree, due to the recursion stack.

Solution

We use a recursive approach to validate the BST property. We maintain two variables, low and high, representing the valid range for the current node's value. For the left child, the valid range becomes (low, node.val) and for the right child, it becomes (node.val, high). By checking if each node's value falls within its respective range, we can ensure that the tree satisfies the BST properties.

#102 Binary Tree Level Order Traversal [Medium] · Trees & BST

Key Insights

- Use Breadth-First Search (BFS) to traverse the tree.
- A queue data structure is ideal for maintaining the order of nodes at each level.
- For every level, process all nodes currently in the queue and add their children for the next level.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is processed once.

Space Complexity: $O(n)$, which accounts for the space used by the queue in the worst-case scenario.

Solution

This solution leverages a BFS approach using a queue. The algorithm starts by enqueueing the root node. Then, it enters a loop that continues until the queue is empty. For each iteration—representing one level of the tree—it computes the number of nodes at that level (levelSize), processes these nodes by dequeuing them, and collects their values. Their non-null children are then enqueued. The list of values for each level is added to the final result list. Key data structures include the queue (for managing nodes) and a list (to store level values).

#103 Binary Tree Zigzag Level Order Traversal [Medium] · Trees & BST

Key Insights

- Use a breadth-first search (BFS) approach to traverse the tree level by level.
- Alternate the direction of traversal on each level to achieve the zigzag pattern.
- A flag or counter can be used to determine the order of nodes for the current level.
- Using a double-ended data structure or simply reversing the list at every alternate level is an effective approach.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, because each node is visited exactly once.

Space Complexity: $O(n)$, as in the worst-case scenario (e.g., a complete binary tree) the queue used for BFS may hold $O(n)$ nodes at once.

Solution

We can solve this problem using a breadth-first search (BFS). Start by initializing a queue to hold nodes for each level. For each level:

- Determine the number of nodes at the current level. – Iterate over these nodes, appending their children for the next level. – Depending on a flag (or level count), either append node values in the natural order (left to right) or reverse the order (right to left) before adding to the result. Key data structures:
- A queue for BFS traversal. – A list (or similar container) to store current level values. The main trick is toggling the order of insertion or reversing the list at alternate levels to achieve the "zigzag" pattern.

#105 Construct Binary Tree from Preorder and Inorder Traversal **[Medium]** · Trees & BST

Key Insights

- Preorder traversal always starts with the root node.

Round 9 · Low · Hard

p.3191

- Inorder traversal splits the tree into left subtree and right subtree based on the root.
- A hash table (map/dictionary) can be used to quickly locate the index of the root in the inorder array.
- Use recursion to construct left and right subtrees.

Space & Time

Time Complexity: $O(n)$ since we process every node exactly once.

Space Complexity: $O(n)$ due to the recursion stack and the hash map used for inorder indices.

Solution

We solve the problem using a recursive divide-and-conquer approach. First, create a hash map to store each value's index in the inorder array for quick lookups. The preorder array always provides the current root. Use the inorder array to divide the tree into left and right subtrees. Recursively construct the subtrees by adjusting the preorder index and the boundaries of the inorder array. The algorithm leverages recursion and a hash map to achieve $O(n)$ time complexity.

#199 Binary Tree Right Side View **[Medium]** · Trees & BST

Key Insights

- The problem is essentially about capturing the last node at each level (or the rightmost node) in a binary

Round 9 · Low · Hard

p.3192

tree.

- A Breadth-First Search (BFS) approach naturally works by processing nodes level by level.
- Alternatively, Depth-First Search (DFS) can be adapted by prioritizing the right subtree first, ensuring that the first node encountered at each depth is the visible one.
- The number of nodes is capped at 100, so both approaches are efficient.

Space & Time

Time Complexity: $O(n)$ – we visit every node in the tree.

Space Complexity: $O(n)$ – in the worst case, the queue (or call stack for DFS) may hold all nodes of the tree.

Solution

We can solve this problem using BFS (level-order traversal). The idea is to traverse the tree level by level using a queue. For each level, the last element added to the level represents the rightmost view from that level. We then append that element's value to our result list. Alternatively, a DFS approach can be used where we visit the right subtree first. By keeping track of the current depth and checking if it's the first time we've encountered this depth, we can add the first node visited at that depth (which will be the rightmost one) to the result.

In both approaches, appropriate data structures (a queue for BFS or recursion with a list for DFS) are used

to store nodes and track the traversal state.

#230 Kth Smallest Element in a BST [Medium] · Trees & BST

Key Insights

- An in-order traversal of a BST visits nodes in increasing order.
- Use a counter to keep track of the number of nodes visited.
- As soon as the counter reaches k , the current node's value is the k th smallest.
- For frequent modifications (insertions/deletions) and queries, consider augmenting the BST with subtree counts.

Space & Time

Time Complexity: $O(n)$ in the worst-case scenario (when $k \approx n$), but on average the traversal may stop early.

Space Complexity: $O(h)$ where h is the height of the tree ($O(n)$ in the worst-case for a skewed tree).

Solution

We perform an in-order traversal of the BST which visits nodes in ascending order. During the traversal, we maintain a counter (or decrement k) to check when we have encountered the k th smallest element. When the counter reaches 0 (or k becomes 0), we capture the current node's value as our answer and return it. This

approach is both simple and efficient for the given constraints.

#235 Lowest Common Ancestor of a Binary Search Tree [Medium] · Trees & BST

Key Insights

- Utilize the BST property: for any node, values in the left subtree are less and values in the right subtree are greater.
- If both p and q have values less than the current node, the LCA must lie in the left subtree.
- If both p and q have values greater than the current node, the LCA must lie in the right subtree.
- When one node value is less than and the other is greater than (or equal to) the current node, the current node is the lowest common ancestor.
- An iterative approach is efficient and uses constant extra space.

Space & Time

Time Complexity: $O(h)$, where h is the height of the tree ($O(\log n)$ for balanced trees and $O(n)$ in the worst-case skewed tree).

Space Complexity: $O(1)$ with an iterative approach ($O(h)$ if using recursion due to call stack).

Solution

The solution leverages the BST properties. Algorithm:

– Start at the root node. – While the current node exists: If both p and q have values less than the current node's value, move to the left child. If both p and q have values greater than the current node's value, move to the right child. Otherwise, the current node is the split point where p and q diverge, so it is the lowest common ancestor. – Return the current node.

This approach efficiently finds the LCA using a single traversal from the root.

#236 Lowest Common Ancestor of a Binary Tree [Medium] · Trees & BST

Key Insights

- Use a depth-first search (DFS) traversal to explore the tree.
- If the current node is either p or q, it could be the LCA.
- Recursively find p and q in the left and right subtrees.
- If both left and right recursive calls return a non-null value, then the current node is the LCA.
- If only one side returns a non-null value, propagate that value up the recursion.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree due to the recursion stack ($O(n)$ in the worst case of a

skewed tree).

Solution

We solve the problem using recursion with depth-first search. The algorithm checks if the current node matches p or q. It then recursively searches the left and right subtrees. If both subtrees contain one of p or q, the current node is the LCA. Otherwise, the algorithm forwards the non-null result up the recursion. This approach naturally handles cases where one node is an ancestor of the other.

#250 Count Univalue Subtrees [Medium] · Trees & BST

Key Insights

- Perform a postorder traversal (left, right, root) to determine whether each subtree is univalue.
- Treat null nodes as univalue to ease the recursion.
- A node's subtree is univalue if both left and right subtrees are univalue and, if they exist, their values equal the current node's value.
- Use a global or external counter that increments every time a univalue subtree is confirmed.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack.

Solution

We use a recursive DFS approach (postorder traversal) to solve the problem. For each node, we recursively check its left and right subtrees to decide if they are univalue. By default, if a child is null, it is considered univalue. Then, we verify the current node's value against its non-null children. If both conditions hold (subtrees are univalue and matching values), we count the current node's subtree as univalue. This DFS checks in a bottom-up manner, ensuring that the properties required for a subtree to be univalue are maintained. The key trick is to return a boolean indicating if the subtree at each node is univalue and, while backtracking, increment the counter appropriately.

#285 Inorder Successor in BST [Medium] · Trees & BST

Key Insights

- In a BST, an in-order traversal yields nodes in ascending order.
- If the node p has a right subtree, the in-order successor is the left-most node in that right subtree.
- If p has no right subtree, the successor is one of its ancestors; traverse from the root towards p while keeping track of the last node that is greater than p.

- An iterative approach is efficient and avoids recursion stack overhead.

Space & Time

Time Complexity: $O(h)$, where h is the height of the tree. In the worst-case (skewed tree), it can be $O(n)$.

Space Complexity: $O(1)$ as only a constant amount of extra space is used.

Solution

We solve the problem using two cases:

– If node p has a right subtree, we simply find the left-most node in that right subtree. – If node p does not have a right subtree, initiate a search from the BST root. As we traverse: If the current node's value is greater than $p.val$, the current node is a candidate for the successor; move to the left subtree to look for a smaller candidate. Otherwise, move to the right subtree. By the end of the traversal, the candidate (if exists) is the in-order successor. This approach benefits from the BST property, limiting traversal to $O(h)$ time.

#298 Binary Tree Longest Consecutive Sequence

[Medium] · Trees & BST

Key Insights

- Use a Depth-First Search (DFS) to traverse the binary tree.

Round 9 · Low · Hard

p.3199

- Pass the current node's value and the length of the current consecutive sequence in the recursive call.
- When a node's value is exactly one greater than its parent's value, increment the sequence length; otherwise, reset it to 1.
- Keep track of the maximum sequence length found during the traversal.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree, since each node is visited once.

Space Complexity: $O(H)$, where H is the height of the tree. In the worst-case of a skewed tree, the recursion stack can be $O(N)$.

Solution

The solution involves using DFS to explore every node in the binary tree while keeping track of the length of the consecutive sequence. For each node, if the node's value is one greater than its parent's value, we increment the sequence length; otherwise, we reset the sequence length to 1. During the traversal, a global variable is updated if the current sequence length exceeds the previously recorded maximum. This approach ensures that every possible consecutive path is considered.

#314 Binary Tree Vertical Order Traversal [Medium]

· Trees & BST

Round 9 · Low · Hard

p.3200

Key Insights

- Use Breadth-First Search (BFS) to traverse the tree level by level.
- Track the column index for each node where root is column 0, left child is column-1, and right child is column+1.
- Group nodes by their column indices using a hash table (or dictionary).
- BFS traversal naturally handles left-to-right ordering within the same level.
- Finally, sort the keys (column indices) to generate results from leftmost column to rightmost column.

Space & Time

Time Complexity: $O(n \log n)$ in the worst-case due to sorting the column indices (n nodes in the worst-case if each node is on a unique column).

Space Complexity: $O(n)$ for storing the node information and the resultant structure.

Solution

The solution uses BFS with a queue that stores nodes along with their corresponding column index. For each visited node, add its value to a dictionary keyed by its column index. When processing children, update the column index appropriately (left child: $col-1$, right child: $col+1$). After the traversal, sort the keys of the dictionary and compile the node values in order from

smallest to largest column index. This approach guarantees that nodes at the same level will be processed in left-to-right order, matching the required output.

#333 Largest BST Subtree [Medium] · Trees & BST

Key Insights

- Use a bottom-up DFS (post-order traversal) to process each node.
- For each node, determine whether its left and right subtrees are BSTs.
- Maintain additional data per subtree: size (number of nodes), minimum value, and maximum value.
- A subtree rooted at the current node is a BST if:
 - Its left subtree is a BST.
 - Its right subtree is a BST.
 - The maximum value in the left subtree is less than the current node's value.
 - The current node's value is less than the minimum value in the right subtree.
- Track the size of the largest BST identified during the traversal.

Space & Time

Time Complexity: $O(n)$ (each node is visited once)

Space Complexity: $O(h)$, where h is the height of the tree (due to recursion stack)

Solution

Perform a post-order traversal of the binary tree. For each node, return a tuple (or object in some languages) that provides:

- Whether the subtree rooted at this node is a BST. – The size of the subtree if it is a BST (or the size of the largest BST found within it). – The minimum value in the subtree. – The maximum value in the subtree.

At each node, if both left and right children form valid BSTs and the current node's value fits the BST condition (greater than left's max and less than right's min), then the subtree rooted at the node is also a BST. Calculate its size and update the current global maximum if necessary. If the condition fails, return the maximum BST size found in its subtrees.

#366 Find Leaves of Binary Tree [Medium] · Trees & BST

Key Insights

- Use a postorder traversal (bottom-up approach) to process the tree.
- Determine the "height" of each node (distance from the deepest leaf) where leaves have a height of 1.
- Group nodes by their computed height to simulate the layer-wise removal of leaves.
- The same node may become a leaf after its children are removed.

Round 9 · Low · Hard

p.3203

Space & Time

Time Complexity: $O(n)$ — Each node is visited exactly once.

Space Complexity: $O(n)$ — Space used for the recursion stack and the result list.

Solution

The solution uses a recursive helper function that performs a postorder traversal of the tree. For each node, it computes its height as $\max(\text{height of left, height of right}) + 1$. Nodes with the same height are collected together in a list. This idea mirrors the process of removing leaves level by level, where leaves (height 1) are removed first, then their parents become new leaves, and so on. Data Structures used:

- An array or list to maintain the groups of node values by their height. Algorithmic Approach: – Use Depth-First Search (DFS) via recursion. – Postorder traversal ensures that leaf nodes (base cases) are processed before their parents. – Append the node's value to the correct list based on the computed height.

#437 Path Sum III [Medium] · Trees & BST

Key Insights

- A path can start and end at any nodes, as long as it goes downwards.
- Using DFS helps traverse the tree and explore all possible paths.

Round 9 · Low · Hard

p.3204

- A prefix sum technique with a hash map allows quick calculation of the sum of any subpath.
- Backtracking is used to remove a path's prefix sum when moving back up the recursion tree.
- Although a brute-force approach exists, it is inefficient compared to the prefix sum method.

Space & Time

Time Complexity: $O(n)$ on average (where n is the number of nodes), though worst-case can be $O(n^2)$ in a skewed tree.

Space Complexity: $O(n)$ due to recursion stack and the hash map that stores prefix sums.

Solution

The solution employs a DFS traversal of the tree along with a hash map that records the frequency of prefix sums encountered so far. At each node, the current prefix sum is updated by adding the node's value. We then check the hash map for how many times (current prefix sum – targetSum) has occurred, since this value indicates the existence of a subpath summing to targetSum. We update the hash map with the current prefix sum, continue the DFS to traverse child nodes, and then backtrack by decrementing the current prefix sum's count before retreating to the parent node. This ensures that sums from other branches are not incorrectly combined.

#545 Boundary of Binary Tree [Medium] · Trees & BST

Key Insights

- Break down the problem into three parts:
- Left boundary (excluding leaves).
- All leaves (do a DFS to find leaves).
- Right boundary (excluding leaves, then reverse the collected list).
- Special handling is needed when the tree doesn't have a left or right subtree.
- The root is always included and is never considered as a leaf, even if it has no children.
- Avoid duplication of leaf nodes in the left/right boundaries.

Space & Time

Time Complexity: $O(n)$ – Every node is visited at most once.

Space Complexity: $O(n)$ – $O(n)$ space is used for the output list and recursion stack (in worst case tree is skewed).

Solution

The solution uses a Depth-First Search (DFS) strategy to traverse the tree and collect the boundary nodes in three steps:

– For the left boundary, start at `root.left` and traverse down, adding nodes (skipping leaves). – For the leaves, perform a DFS from the root and add any node that has no children. – For the right boundary, start at `root.right` and traverse down, adding nodes (skipping leaves). Since the right boundary must be added in reverse order, we reverse the collected list after traversal. – Finally, concatenate these lists in the order: root, left boundary, leaves, reversed right boundary.

The approach makes use of helper functions to clearly separate the logic for gathering left boundary, leaves, and right boundary. This modular approach simplifies reasoning and avoids edge case mistakes such as including duplicate leaf nodes.

#549 Binary Tree Longest Consecutive Sequence II

[Medium] · Trees & BST

Key Insights

- Use Depth-First Search (DFS) to visit every node.
- For each node, compute:
 - The longest increasing consecutive path starting from that node.
 - The longest decreasing consecutive path starting from that node.
 - Combine the increasing and decreasing paths (subtracting one to avoid double counting the current node) to potentially form a longer consecutive sequence

through the node.

- Update a global maximum to record the longest path seen.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes since every node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack.

Solution

The solution employs a DFS traversal. At every node, we calculate two values:

– `inc` – the length of the longest increasing consecutive sequence starting from that node. – `dec` – the length of the longest decreasing consecutive sequence starting from that node.

For each child of the current node:

– If the child's value is `node.val + 1`, update the increasing sequence. – If the child's value is `node.val - 1`, update the decreasing sequence.

Then, the longest path through the current node is `inc + dec - 1` (subtracting one to avoid counting the current node twice). A global variable maintains the maximum length encountered. This process works regardless of the path direction because it combines both increasing and decreasing parts.

#582 Kill Process [Medium] · Trees & BST

Key Insights

- Build a mapping (dictionary/hash table) from a parent process ID to its list of child process IDs.
- Use Depth-First Search (DFS) or Breadth-First Search (BFS) starting from the kill process to traverse the tree and collect all descendant processes.
- The problem requires handling a tree structure that may have multiple children per node.

Space & Time

Time Complexity: $O(n)$, where n is the number of processes (both building the mapping and the tree traversal take linear time).

Space Complexity: $O(n)$, for storing the mapping and the traversal queue/stack.

Solution

We first construct a mapping from each process's parent ID to its list of child IDs. This allows us to quickly access all children of any process. Once the mapping is built, we perform a tree traversal (either BFS or DFS) starting from the process to kill. During the traversal, we record every process encountered. This collected list represents all processes that will be terminated. The data structures used are a dictionary (for the mapping) and either a stack (for DFS) or a queue (for BFS), making the solution efficient and straightforward.

#662 Maximum Width of Binary Tree [Medium] · Trees & BST

Key Insights

- Perform a level order traversal (BFS) of the tree.
- Assign an index to each node as if the tree were complete (e.g., root is index 1, left child is 2
- i , right child is 2
- $i+1$).
- For each level, compute the width as $(\text{last_index} - \text{first_index} + 1)$.
- Using indices allows capturing gaps between nodes due to nulls in between.
- Edge cases include handling very skewed trees.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree.

Space Complexity: $O(N)$, due to the queue used for BFS which in the worst case holds nodes of the last level.

Solution

We use a Breadth-First Search (BFS) approach with a queue to traverse the tree level by level. At each level, we pair each node with an index representing its position assuming a complete binary tree. This allows us to compute the width of each level using the formula: $\text{width} = \text{last_index} - \text{first_index} + 1$. We then track the

maximum width seen throughout the traversal. The use of indices handles cases where there are null gaps between nodes.

#863 All Nodes Distance K in Binary Tree [Medium]

· Trees & BST

Key Insights

- Convert the binary tree into an undirected graph so that we can easily traverse both down to the children and up to the parent.
- Use a DFS or simple recursive traversal to create an adjacency list representing the graph.
- Perform a BFS starting from the target node to explore nodes level-by-level until reaching distance k.
- Use a visited set (or equivalent) to avoid revisiting nodes and to prevent cycles.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree. Each node is visited once when constructing the graph and once in the BFS.

Space Complexity: $O(N)$, for storing the graph and the additional data structures used in BFS and DFS.

Solution

The solution involves two main steps:

- Build an undirected graph representation of the binary tree using a DFS. For each node, add the node's value as

a key in an adjacency list and connect it with its children (if any). This creates bidirectional links between parent and child. – Use a BFS starting from the target node and explore nodes level-by-level. Keep a counter for the current distance, and once the counter reaches k, return all node values found at that level. To avoid processing nodes more than once, maintain a set of visited nodes.

#1120 Maximum Average Subtree [Medium] · Trees & BST

Key Insights

- Use a depth-first search (DFS) to traverse every subtree.
- At each node, compute the sum and count of nodes in its subtree.
- Calculate the average for the current subtree and update a global maximum if needed.
- Processing every node exactly once yields an efficient solution.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes (each node is visited once).

Space Complexity: $O(h)$, where h is the height of the tree (space for recursion stack).

Solution

We perform a post-order DFS traversal, calculating at each node both the cumulative sum and the count of nodes in its subtree. This enables us to compute the average value in constant time for any subtree. We maintain a global variable to track the maximum average encountered during the traversal. The primary technique is recursive DFS combined with tuple (or pair) aggregation. The simplicity of the recursion and in-place aggregation ensures an efficient and clear solution.

#1490 Clone N-ary Tree [Medium] · Trees & BST

Key Insights

- The tree is N-ary, meaning every node can have multiple children.
- A deep copy requires creating new nodes for each node in the tree.
- Both DFS (recursive) or BFS (iterative) traversals can be used to clone the tree.
- The overall time complexity is $O(N)$ since each node is visited once.
- We can follow similar cloning techniques used in cloning graphs, using either recursion or an auxiliary data structure like a queue.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the tree, since each node is visited exactly once.

Space Complexity: $O(N)$, accounting for the recursive call stack (or an explicit queue in BFS) and the space needed to store the clone.

Solution

We perform a DFS traversal starting from the root node. For each node we encounter, we:

- Create a new node with the same value.
- Recursively clone all its children and attach the cloned children to the new node. This approach ensures that every original node is copied into a new node and that the parent-child relationships are preserved. Edge cases such as an empty tree (null root) are also handled by returning null immediately.

#1522 Diameter of N-Ary Tree [Medium] · Trees & BST

Key Insights

- The diameter of the tree can be found by considering the two largest depths from any node among its children.
- Use a depth-first search (DFS) to compute the maximum depth from each node.
- For each node, the potential diameter is the sum of the two longest depths among its children.
- Update a global variable tracking the maximum diameter as you recurse through the tree.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes, since each node is visited once.

Space Complexity: $O(h)$, where h is the height of the tree, due to the recursion stack (worst-case $O(n)$ for a skewed tree).

Solution

Use a DFS approach to traverse the tree. At each node, calculate the longest and second longest depth among its children. The sum of these two depths gives the candidate diameter through that node. Maintain a global variable to track the maximum diameter found so far. Return the maximum depth from the node (i.e., one plus the largest depth among its children) to be used in parent's calculations.

#1650 Lowest Common Ancestor of a Binary Tree III [Medium] · Trees & BST

Key Insights

- Each node has a parent pointer, so you can traverse from any node to the root.
- By determining the depth of each node, you can “align” the two nodes to the same depth.
- When both nodes are aligned, traverse upward simultaneously until they meet.

Round 9 · Low · Hard

p.3215

- This approach guarantees finding the LCA with $O(h)$ time complexity, where h is the height of the tree.

Space & Time

Time Complexity: $O(h)$, where h is the height of the tree.

Space Complexity: $O(1)$

Solution

We first compute the depth of both nodes by moving up their parent pointers. If one node is deeper, move it upward until both nodes are at the same level. Then, move both nodes upward one step at a time until they meet; the meeting point is the lowest common ancestor. This approach uses constant extra space and leverages the parent pointer for an efficient solution.

#130 Surrounded Regions [Medium] · DFS

Key Insights

- Regions connected to the border are not captured.
- Use DFS/BFS from border 'O's to mark safe cells.
- After marking, flip all remaining 'O's to 'X' and then revert the marked cells back to 'O'.

Space & Time

Time Complexity: $O(m * n)$

Space Complexity: $O(m * n)$ in the worst case due to recursion stack or queue storage for DFS/BFS

Solution

Round 9 · Low · Hard

p.3216

The solution involves two main steps. First, traverse the board's borders and run a DFS or BFS from any encountered 'O', marking it (e.g., using ' ') to indicate that it should not be flipped. Next, go through each cell in the board: flip any remaining 'O' (i.e., not marked safe) to 'X' and change the temporary marker ' ' back to 'O'. This approach ensures only the regions completely surrounded by 'X' are captured.

#133 Clone Graph [Medium] · DFS

Key Insights

- The graph is undirected and connected.
- Each node must be cloned exactly once to prevent cycles and repetitions.
- Use a mapping/dictionary to keep track of cloned nodes.
- A depth-first search (DFS) or breadth-first search (BFS) can be used to traverse and clone the graph.
- Handle edge cases where the graph might be empty.

Space & Time

Time Complexity: $O(N + E)$, where N is the number of nodes and E is the number of edges, because each node and each edge are traversed once.

Space Complexity: $O(N)$, to store the mapping from original nodes to their clones, and for the recursion/queue used in traversal.

Solution

We employ a depth-first search (DFS) approach to clone the graph. The main idea is to create a mapping (hash table) that links each original node to its corresponding cloned node. During the DFS traversal, if a node is encountered for the first time, a new node is created and stored in the mapping; then, the algorithm recursively clones all of its neighbors. For an already visited node, the cloned node is directly returned from the mapping. This ensures that every node is cloned only once and correctly handles cycles in the graph.

#200 Number of Islands [Medium] · DFS

Key Insights

- Use graph traversal (DFS/BFS) to explore connected components of land.
- Iterate over the matrix and start a traversal (e.g., DFS) when encountering an unvisited "1".
- During traversal, mark visited cells to avoid duplicate counting.
- Each DFS/BFS call completely explores one island.
- Alternative approach can use Union Find to group connected lands.

Space & Time

Time Complexity: $O(m * n)$, as each cell is visited at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario due to recursion stack (DFS) or auxiliary data structures.

Solution

We solve the problem using a DFS-based approach. The algorithm iterates over each cell in the grid; when a "1" (land) is encountered, the DFS is initiated to mark the entire island (all connected "1"s) as visited (by setting them to "0"). This prevents recounting the same island. The island count is incremented each time a new island is discovered and fully explored.

#207 Course Schedule [Medium] · DFS

Key Insights

- This problem can be modeled as a directed graph where courses are nodes and prerequisites are edges.
- If the graph has a cycle, then it is not possible to complete all courses.
- We can detect cycles using either Topological Sort (BFS/Kahn's algorithm) or DFS-based cycle detection.
- The time complexity is $O(V + E)$, where V is the number of courses and E is the number of prerequisite pairs.

Space & Time

Time Complexity: $O(V + E)$

Space Complexity: $O(V + E)$

Solution

We solve the problem using the Topological Sort algorithm (BFS/Kahn's algorithm). First, we build a graph (using an adjacency list) and track the indegree (number of prerequisites) for each course. We then find all courses with an indegree of zero (courses with no prerequisites) and add them to a processing queue. While the queue is not empty, we remove a course from the queue, reduce the indegree for its neighboring courses, and if any neighbor's indegree becomes zero, we add it to the queue as well. If we successfully process all courses, then it means that the graph was acyclic and a valid schedule exists, returning true. If not, a cycle exists, and we return false.

#210 Course Schedule II [Medium] · DFS

Key Insights

- This is a topological sort problem on a directed graph.
- The courses are nodes, and prerequisites form directed edges.
- We can solve the problem by either:
 - Using DFS with cycle detection and a stack for topological ordering.
 - Using BFS (Kahn's Algorithm) to process nodes with zero in-degrees.
- If at any point no node with a zero in-degree is available, a cycle exists and no valid order can be

produced.

Space & Time

Time Complexity: $O(V + E)$ where V is the number of courses and E is the number of prerequisite pairs.

Space Complexity: $O(V + E)$ for storing the graph and in-degree array.

Solution

We treat the problem as a topological sort on a directed acyclic graph (DAG). For a BFS approach using Kahn's algorithm:

- Build an adjacency list for the graph and an in-degree array to record the number of prerequisites (incoming edges) for each course.
 - Initialize a queue with all courses that have an in-degree of zero (i.e., courses that can be taken without prerequisites).
 - Dequeue nodes from the queue, add them to the ordering, and reduce the in-degree of their neighbors.
 - If a neighbor's in-degree becomes zero, add it to the queue.
 - If the final ordering includes all courses, return the ordering; otherwise, return an empty array indicating a cycle.
- The chosen data structures are:
- A list (or array) for the in-degree counts.
 - A dictionary (or array of lists) for the adjacency list.
 - A queue to process nodes with zero in-degree.

#261 Graph Valid Tree [Medium] · DFS

Round 9 · Low · Hard

p.3221

Key Insights

- A tree must be acyclic and connected.
- For n nodes, a valid tree must have exactly $n - 1$ edges.
- A DFS/BFS traversal can check connectivity and simultaneously detect cycles.
- The Union-Find (disjoint set) approach provides an alternative for cycle detection and connectivity verification.

Space & Time

Time Complexity: $O(n + e)$, where e is the number of edges

Space Complexity: $O(n + e)$ for storing the graph and traversal/union-find structures

Solution

We can solve the problem using either a DFS/BFS approach or a Union-Find algorithm. For the DFS/BFS method, we first check if the number of edges is exactly $n - 1$. If not, it cannot be a tree. Then, we build an adjacency list for the graph and perform a DFS starting from node 0 while keeping track of visited nodes and parent information to avoid false-positive cycle detections. If during DFS we encounter a node that has been visited (and is not the parent), then a cycle exists. Finally, if all nodes are reached from node 0 and no cycle is found, the graph is a valid tree.

Round 9 · Low · Hard

p.3222

Using Union-Find, we iterate over each edge and try to union the two nodes. If they already share the same parent (i.e., are in the same set), a cycle exists. After processing all edges, we confirm that exactly one connected component exists.

#310 Minimum Height Trees [Medium] · DFS

Key Insights

- A tree can have at most two centers that will serve as the roots for its minimum height trees.
- Instead of trying all possible roots, we can remove leaves layer by layer until we are left with the central node(s).
- The process is similar to a topological sort on the tree.
- When there are 2 or fewer nodes left, these nodes are the centers (roots) of MHTs.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since we process each node and edge once.

Space Complexity: $O(n)$ for storing the graph and additional data structures like the leaves list.

Solution

The solution builds an undirected graph from the input, then identifies and removes node-leaves iteratively. A leaf is defined as a node connected to only one other node. By repeatedly removing the leaves and updating

the graph (decreasing neighbor degrees), the algorithm peels away the outer layers of the tree. When only one or two nodes remain, these nodes are the centers of the tree and form the roots of the minimum height trees.

#323 Number of Connected Components in an Undirected Graph [Medium] · DFS

Key Insights

- Use an adjacency list to represent the graph.
- A DFS (or BFS) traversal from an unvisited node will traverse an entire connected component.
- Each new DFS/BFS call from an unvisited node signals a separate connected component.
- Alternatively, the Union-Find data structure can be used to union connected nodes and count distinct sets.

Space & Time

Time Complexity: $O(n + e)$, where n is the number of nodes and e is the number of edges.

Space Complexity: $O(n + e)$ due to the storage of the graph and tracking of visited nodes.

Solution

We solve the problem by iterating over each node and performing a DFS for any node that hasn't been visited. The DFS uses an iterative approach with a stack to traverse all nodes in a connected component. Each DFS initiation corresponds to one connected component.

This method efficiently explores the graph using an adjacency list and marks nodes as visited to prevent redundant traversals.

#417 Pacific Atlantic Water Flow [Medium] · DFS

Key Insights

- Reverse the typical DFS/BFS perspective: start from the oceans and work inward.
- Maintain two matrices (or sets) to keep track of cells reachable from the Pacific and Atlantic edges.
- A cell can be added to the result if it is reached in both traversals.
- The height constraint ensures water flows only in non-increasing order.

Space & Time

Time Complexity: $O(m * n)$ where m is the number of rows and n is the number of columns. Each cell is processed up to twice (once from each ocean).

Space Complexity: $O(m * n)$ for the visited matrices and the recursion/queue used in DFS/BFS.

Solution

The solution uses depth-first search (DFS) from the ocean boundaries. We initialize two boolean matrices, one for each ocean. For the Pacific, we start DFS from the leftmost column and top row; for the Atlantic, we start from the rightmost column and bottom row. In each DFS,

if the next cell has a height equal to or greater than the current cell and hasn't been visited, we continue the search. Finally, we scan through the grid to collect cells that are reachable from both oceans. This reverse search simplifies the path validations and leverages the allowed water flow conditions.

#490 The Maze [Medium] · DFS

Key Insights

- The ball continues moving in one direction until it hits a wall (or the boundary, which is considered a wall).
- You must simulate the ball's rolling rather than taking a single step.
- Even if the ball passes through the destination, it is only valid if it can stop exactly there.
- A breadth-first search (BFS) or depth-first search (DFS) strategy is ideal because we need to explore all stopping positions.
- Use a visited structure to avoid revisiting positions and thus prevent infinite loops.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the maze.

Space Complexity: $O(m * n)$ due to the storage used by the visited matrix and the BFS/DFS queue or recursion stack.

Solution

We simulate the ball's movement using a breadth-first search (BFS) approach. Each time we pick a stop position, we try to roll in all four directions until the ball hits a wall. At that point, if the new stopping position has not been visited, we mark it as visited and add it to our queue. If we ever reach the destination exactly, we return true. The key is simulating the "roll" by iteratively moving in a direction until no further move is possible. We store the coordinates in a queue (or stack for DFS) along with a 2D visited array to monitor which positions have been processed. The directions are maintained in a list/array, and for each direction, we keep moving until the next cell would be out of bounds or a wall.

#694 Number of Distinct Islands [Medium] · DFS

Key Insights

- Use DFS/BFS to traverse and mark each island.
- For each island, record its shape relative to a starting coordinate to allow translation invariance.
- Represent the island's shape as a unique signature (e.g., tuple or string of relative positions).
- Use a set (or hash table) to store unique shapes.

Space & Time

Time Complexity: $O(m * n)$ as each cell is visited once.

Space Complexity: $O(m * n)$ in the worst-case due to recursion stack and storing island shapes.

Solution

Use depth-first search (DFS) to explore the grid. When an unvisited land cell (1) is encountered, initiate a DFS to traverse the entire island. Record each land cell's coordinates relative to the starting position of the island. This relative representation is invariant under translation, making it possible to compare island shapes. Store the canonical shape signature (for example, as a tuple or string) in a set to ensure uniqueness. Finally, return the size of the set as the number of distinct islands.

#695 Max Area of Island [Medium] · DFS

Key Insights

- Traverse each cell in the grid.
- Use Depth-First Search (DFS) or Breadth-First Search (BFS) to explore connected 1's.
- Mark visited cells to prevent double counting (e.g., by setting them to 0).
- Maintain and update a maximum area count during traversal.
- The grid's boundaries are naturally water, so no extra edge processing is needed.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the grid since every cell is visited at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario for the recursion stack (or auxiliary data structures) when the grid is entirely land.

Solution

The solution leverages a DFS approach to iterate over each cell in the given grid. When a cell with a value of 1 is encountered, a DFS is initiated from that cell to calculate the area of the island connected to it. Cells are marked as visited (by changing them from 1 to 0) during the DFS to avoid revisiting. The DFS explores all four directions (up, down, left, right) and accumulates the count of connected cells. The maximum area found among all islands is returned at the end. This approach efficiently computes the largest island area with a simple traversal mechanism.

#721 Accounts Merge [Medium] · DFS

Key Insights

- Two accounts are merged if they share at least one common email.
- Treat each email as a node in a graph; accounts represent connected components.
- Both DFS/BFS and Union-Find can be used to find connected components.

Round 9 · Low · Hard

p.3229

- Keep a mapping from email to name because all emails in a connected component belong to the same person.
- After grouping, sort the emails and prepend the name to form the final account list.

Space & Time

Time Complexity: $O(N * M * \log(M))$ where N is the number of accounts and M is the number of emails per account (owing to sorting and union operations).

Space Complexity: $O(N * M)$, used for storing email mappings and union-find structures.

Solution

The solution involves using a Union-Find (Disjoint Set Union) data structure to merge emails that are connected by being on the same account. For each account, we union the first email with every other email in that account. We also maintain a mapping from email to owner name. After processing all accounts, we group the emails by their root parent derived from the union-find structure. For each group, we sort the emails and prepend the owner's name. Finally, we return the merged list of accounts.

#785 Is Graph Bipartite? [Medium] · DFS

Key Insights

- A graph is bipartite if you can assign two colors to each node such that no two adjacent nodes share the

Round 9 · Low · Hard

p.3230

same color.

- The challenge may require exploring disconnected components.
- BFS or DFS can be used to attempt the two-coloring, and a conflict during coloring indicates the graph is not bipartite.
- Union Find is an alternative approach, but BFS/DFS is more intuitive for graph coloring problems.

Space & Time

Time Complexity: $O(V + E)$ where V is the number of nodes and E is the number of edges, since every node and edge is traversed at most once.

Space Complexity: $O(V)$, for the color array and the queue (or recursion stack) used in BFS (or DFS).

Solution

We use a BFS/DFS approach to assign colors (0 and 1) to nodes. A color array, initially set to -1 for all nodes, is maintained. For each unvisited node (color -1), we start a BFS (or DFS) and assign it an initial color (e.g., 0). For every visited neighbor, if it has not been colored, assign it the opposite color; if it is already colored with the same color as the current node, a conflict has occurred, and the graph is not bipartite. This method gracefully handles disconnected components by iterating over all nodes.

#1236 Web Crawler [Medium] · DFS

Round 9 · Low · Hard

p.3231

Key Insights

- Treat the set of webpages as a graph where each URL is a node and an edge exists from one URL to another if it is returned by `HtmlParser.getUrls(url)`.
- Use either breadth-first search (BFS) or depth-first search (DFS) to traverse the webpages.
- Extract the hostname from a URL to ensure that only URLs on the same domain are crawled.
- Maintain a visited set (or equivalent) to avoid processing the same URL multiple times.

Space & Time

Time Complexity: $O(N + E)$ where N is the number of nodes (URLs) and E is the number of edges (links between URLs) encountered during the crawl.

Space Complexity: $O(N)$ for storing the visited URLs and the queue/stack used for traversal.

Solution

The solution uses a BFS (or DFS) approach to traverse the web graph:

- Extract the hostname from `startUrl` (for instance, by splitting on `"/"` and then the first `"/"` that follows).
- Initialize a visited set to keep track of URLs that have been crawled.
- Use a queue to implement the BFS. Enqueue the `startUrl`.
- While the queue is not empty, dequeue a URL and retrieve its outgoing links using `HtmlParser.getUrls(url)`.
- For each URL obtained: Extract

Round 9 · Low · Hard

p.3232

its hostname and compare it with the hostname of startUrl. If it matches and it has not been visited before, add it to the visited set and enqueue it. – Finally, return the set of visited URLs.

This approach guarantees that each URL is processed only once and that only URLs from the same hostname are crawled.

#1242 Web Crawler Multithreaded [Medium] · DFS

Key Insights

- Use multiple threads to perform concurrent crawling, which speeds up the process over a single-threaded approach.
- Maintain a thread-safe set (visited) to ensure no URL is processed twice.
- Utilize a work queue to manage URLs pending processing.
- Filter URLs based on the hostname extracted from startUrl.
- Apply synchronization (locks or concurrent data structures) to avoid race conditions when multiple threads access shared data.

Space & Time

Time Complexity: $O(N)$ where N is the number of URLs processed, with additional constant factors due to threading overhead.

Space Complexity: $O(N)$ for storing the visited URLs and the work queue.

Solution

The solution initializes a shared work queue with the startUrl and a thread-safe visited set. Each worker thread dequeues a URL, retrieves its linked URLs via the blocking HtmlParser.getUrls() call, and filters the resulting URLs by comparing their hostnames with the start URL's hostname. New URLs that have not yet been visited are added to both the visited set and the work queue. The use of a thread pool or multiple worker threads along with synchronization ensures that the crawler processes URLs concurrently while avoiding race conditions. The approach follows a BFS-like pattern spread across multiple threads.

#128 Longest Consecutive Sequence [Medium] · Graphs

Key Insights

- Utilize a set (or hash table) for constant time look-ups.
- Only start counting a sequence from a number that is the beginning (its previous number is not in the set).
- Each number is visited only once, ensuring an $O(n)$ time complexity.
- Update the maximum sequence length as you discover longer consecutive sequences.

Space & TimeTime Complexity: $O(n)$ Space Complexity: $O(n)$ **Solution**

We first insert all the numbers into a hash set to enable $O(1)$ membership checks. For each number in the set, we check if it is the start of a sequence by ensuring that the number minus one is not in the set. If it is the beginning, we then count all consecutive numbers following it, updating the maximum sequence length when necessary. This method guarantees that we only count each number once, satisfying the $O(n)$ time requirement.

#277 Find the Celebrity [Medium] · Graphs**Key Insights**

- Use pairwise comparisons to eliminate non-celebrities.
- If a candidate knows another person, the candidate cannot be the celebrity.
- Verify the candidate by checking that they do not know anyone and that everyone knows them.
- Achieve the solution in two passes: one for candidate selection and the other for validation.

Space & Time

Time Complexity: $O(n)$ – One pass for candidate selection and one for validation.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

The approach consists of two main steps:

– **Candidate Selection:** Iterate through the people and maintain a candidate. For every person i , if the current candidate knows i , then set $\text{candidate} = i$. This works because if candidate knows i , candidate cannot be the celebrity. – **Verification:** Check that the candidate does not know any other person and that every other person knows the candidate. If both conditions hold, candidate is the celebrity; otherwise, return -1 .

The main data structures used are simple integer variables. The algorithm leverages a two-pass technique to reduce the number of knows API calls while ensuring correctness.

#1136 Parallel Courses [Medium] · Graphs**Key Insights**

- Model the problem as a directed graph where nodes are courses and edges represent prerequisites.
- Use topological sorting (Kahn's algorithm) to process courses with zero prerequisites.
- Each level (or layer) of processing represents one semester.

- Detect cycles by ensuring all courses are eventually processed; if not, a cycle exists.
- The number of BFS levels gives the minimum number of semesters necessary.

Space & Time

Time Complexity: $O(n + m)$ where n is the number of courses and m is the length of relations (prerequisites).

Space Complexity: $O(n + m)$ for storing the graph, in-degree counts, and the processing queue.

Solution

We solve the problem using a topological sort with Kahn's algorithm. First, construct a graph representation using an adjacency list and compute the in-degree for each course. Initialize a queue with courses that have no prerequisites (in-degree == 0). Then, process courses level by level. Each level processed represents one semester. For every course taken in that semester, decrease the in-degree of all its neighbors. Add any neighbor with an in-degree that becomes zero to the queue for the next semester. If, after processing, the total courses taken is less than n , then a cycle exists and return -1 . Otherwise, the number of semesters processed is the answer.

#286 Walls and Gates [Medium] · BFS

Key Insights

- Use a multi-source Breadth-First Search (BFS) starting from all the gates, as this allows spreading the shortest distance to each empty room.
- Instead of starting from every empty room separately, starting from all gates simultaneously helps update all distances in one pass.
- Only update a room if the new calculated distance is smaller than its current value to prevent unnecessary processing.
- The problem is essentially a multi-source shortest path problem on a grid.

Space & Time

Time Complexity: $O(m * n)$ where m and n are the dimensions of the grid, as each cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst-case scenario used by the BFS queue.

Solution

The solution uses a multi-source BFS algorithm by initially enqueueing all the gate positions (cells with value 0). Then, for each cell obtained from the queue, check its four possible neighbors (up, down, left, right). If a neighbor is an empty room (INF), update its distance to be the current cell's distance plus one, and then enqueue the neighbor to process further. This ensures that each empty room gets the minimum distance from any gate.

#322 Coin Change [Medium] · BFS

Key Insights

- Use a dynamic programming approach to build up a solution from smaller subproblems.
- Define $dp[x]$ as the minimum number of coins required for amount x .
- The recurrence relation is: $dp[x] = \min(dp[x], dp[x - \text{coin}] + 1)$ for each coin.
- Initialize $dp[0] = 0$ and set other values to a number larger than the maximum possible coins ($\text{amount} + 1$ works as an infinity substitute).
- An alternative approach is through BFS, treating each amount as a node and each coin as an edge, but the DP method is more intuitive here.

Space & Time

Time Complexity: $O(\text{amount} * n)$ where n is the number of coins.

Space Complexity: $O(\text{amount})$ for the dp array.

Solution

We use a bottom-up dynamic programming approach. The idea is to use a one-dimensional dp array where each element at index x represents the minimum number of coins required to form the amount x . Starting from $dp[0]$ which is 0, we iterate through all amounts up to the target amount and update $dp[x]$ based on each

coin denomination. If after filling the dp array the value $dp[\text{amount}]$ remains unchanged (i.e., greater than amount), it means the amount cannot be formed and we return -1.

#542 01 Matrix [Medium] · BFS

Key Insights

- Treat all 0s as starting points for a multi-source BFS.
- Update each neighboring cell's distance if a shorter distance is found.
- Utilize a queue to process cells in order of increasing distance.
- Ensure that each cell is processed only once by marking visited cells.

Space & Time

Time Complexity: $O(m$

$n)$ since each cell is processed once.

Space Complexity: $O(m$

$n)$ for the queue and the answer matrix.

Solution

We use a multi-source Breadth-First Search (BFS) where every 0 in the matrix is initially added to the BFS queue. Each neighboring cell that is a 1 will have its distance updated to be one plus the distance of the current cell if it hasn't been visited yet (or if a smaller distance is found). This ensures that the first time a 1 is reached, it

is reached by the shortest path. Data structures used include a queue (or deque) for BFS and the matrix itself for storing distances.

#994 Rotting Oranges [Medium] · BFS

Key Insights

- Use Breadth-First Search (BFS) starting from all initially rotten oranges.
- Process the grid level by level, where each level represents one minute.
- Keep a count of fresh oranges and decrement it as they become rotten.
- If after the BFS there are still fresh oranges, then it is impossible to rot every orange.

Space & Time

Time Complexity: $O(m * n)$ as every cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst case for the queue and auxiliary storage.

Solution

We begin by scanning the grid to locate all rotten oranges and count the fresh ones. These rotten oranges are added to a queue for a BFS. For each minute (BFS level), we process all the current rotten oranges and try to infect their adjacent fresh oranges. If an adjacent cell contains a fresh orange, we mark it rotten, add it to the

queue, and decrement our fresh count. We continue this process level by level (minute by minute) until there are no fresh oranges left. If, after processing, there are still fresh oranges remaining, we return -1 , indicating that not all oranges can become rotten.

#1197 Minimum Knight Moves [Medium] · BFS

Key Insights

- Leverage the symmetry of the chessboard by converting target coordinates to their absolute values.
- Use a Breadth-First Search (BFS) because all moves have equal weight, ensuring the first encounter with the target is the minimal path.
- Prune the search space by only considering positions that are likely beneficial (e.g., positions having coordinates not far less than -1).
- Track visited states to avoid redundant processing.
- Optimize by bounding the search region relative to the target distance.

Space & Time

Time Complexity: $O(n)$ where n is the number of positions processed during the BFS

Space Complexity: $O(n)$ due to the storage requirements for the BFS queue and the visited set

Solution

We solve the problem using a BFS starting from the origin. First, normalize the target by taking the absolute value of x and y . BFS is then applied by exploring all eight possible knight moves at each step. Each move generates a new position that is enqueued if it hasn't been visited and falls within a reasonable bound (e.g., coordinates ≥ -1). The search terminates as soon as the target position is reached.

#1730 Shortest Path to Get Food [Medium] · BFS

Key Insights

- Use Breadth-First Search (BFS) to explore the grid level by level.
- The BFS guarantees that when a food cell ('#') is reached, the current distance is the shortest.
- Maintain a visited structure to avoid re-processing the same cell.
- Consider boundary and obstacle ('X') conditions while traversing.

Space & Time

Time Complexity: $O(m * n)$, where m is the number of rows and n is the number of columns since each cell is processed at most once.

Space Complexity: $O(m * n)$ in the worst case due to the queue and visited matrix.

Solution

Round 9 · Low · Hard

p.3243

We use Breadth-First Search (BFS) starting from the cell marked with '*'. The BFS uses a queue to explore neighbors in the four allowed directions (up, down, left, right). Each level of BFS represents one step in the path. As soon as a cell containing food ('#') is reached, the current step count is returned. We also maintain a visited matrix to ensure that cells are not revisited, which could otherwise result in cycles or redundant processing. This approach ensures that the first time we encounter food, we have found the shortest possible path.

Round 9 · Low · Hard

p.3244

R3 · High · Hard (23 q)

>> Round 3 · High Priority · Hard — Quick Review

23 questions · key insights · complexity · solution

#41 First Missing Positive [Hard] · Arrays & Hashing

Key Insights

- The missing positive number must be in the range $[1, n+1]$ where n is the length of the array.
- Use the array indices to place numbers in their correct positions (i.e., number i should be at index $i-1$) with in-place swapping.
- After rearrangement, scan the array: the first index where the number is not equal to $\text{index}+1$ indicates the missing positive integer.

Space & Time

Time Complexity: $O(n)$ as each number is swapped at most once.

Space Complexity: $O(1)$ since the rearrangement occurs in-place.

Solution

The approach involves reordering the array so that each positive integer in the range $[1, n]$ is placed at the index

corresponding to its value minus one. This means for any valid number i , $\text{nums}[i-1]$ should be i . We iterate through the array and perform in-place swaps to achieve this order. After this process, another pass through the array reveals the first position where the number does not match the expected value $(i+1)$, which is the smallest missing positive. If every position is correct, then the missing value is $n+1$.

#76 Minimum Window Substring [Hard] · Sliding Window

Key Insights

- Use a sliding window approach to dynamically adjust the window boundaries.
- Utilize a hash table (or dictionary) to keep track of the frequency of characters required (from t) and the frequency within the current window.
- Expand the window until it satisfies the requirements, then contract to minimize the window.
- Efficiently update counts and check valid window with two pointers (left and right).

Space & Time

Time Complexity: $O(m + n)$, where m is the length of s and n is the length of t .

Space Complexity: $O(m + n)$ in the worst case due to the storage for the hash maps.

Solution

We use a sliding window technique with two pointers representing the window's boundaries. First, create a frequency map of characters from *t*. Then, expand the right pointer to include characters in the window until the window contains all required characters (meeting or exceeding the counts in *t*). Once the requirement is met, try contracting the window by moving the left pointer to reduce its size while still maintaining a valid window. Throughout the process, if a valid window is found that is smaller than the previously recorded minimum window, update the minimum. The solution leans heavily on hash tables to track counts and compares them against the required counts from *t*.

#4 Median of Two Sorted Arrays [Hard] · Binary Search

Key Insights

- The median divides the sorted array into two equal halves.
- We can perform binary search on the smaller array to partition both arrays correctly.
- We need to ensure that every element in the left partition is less than or equal to every element in the right partition.
- Edge cases must be handled when partitions are at the boundaries of either array.

Round 9 · Low · Hard

p.3247

Space & Time

Time Complexity: $O(\log(\min(m, n)))$

Space Complexity: $O(1)$

Solution

The solution uses a binary search technique on the smaller array. We define a partition index *i* for *nums1* and compute the corresponding index *j* for *nums2* such that *i* + *j* equals half of the total number of elements. The algorithm checks if the largest element on the left side of *nums1* (if any) is less than or equal to the smallest element on the right side of *nums2*, and vice versa for *nums2*. If the condition holds, we have found the correct partition and compute the median. If not, adjust the binary search range until the correct partition is found. This partitioning strategy allows us to find the median in logarithmic time.

#315 Count of Smaller Numbers After Self [Hard] · Binary Search

Key Insights

- The brute force solution takes $O(n^2)$ time, which is too slow for large input arrays.
- An efficient approach leverages the concept of Merge Sort to both sort the array and count the number of smaller elements by tracking the number of elements that "jump over" during the merge step.

Round 9 · Low · Hard

p.3248

- Other efficient methods include Binary Indexed Trees (Fenwick Tree) or Balanced Binary Search Trees.

Space & Time

Time Complexity: $O(n \log n)$ on average due to the merge sort based approach.

Space Complexity: $O(n)$ for auxiliary arrays used during sorting and counting.

Solution

We use a modified Merge Sort to sort the array while counting the number of smaller elements to the right. We keep track of the original indices by pairing each element with its index. During the merge process, when an element from the left subarray is placed into the temporary array, we add the count of elements from the right subarray that have already been placed (since these represent smaller elements that were to its right). This efficient divide and conquer strategy guarantees an overall $O(n \log n)$ time complexity. Care must be taken to update the counts based on indices and during the merge step.

#410 Split Array Largest Sum **[Hard]** · Binary Search

Key Insights

- Use binary search over the answer: the candidate value is the maximum subarray sum.

- The lower bound of the search is $\max(\text{nums})$ and the upper bound is $\text{sum}(\text{nums})$.
- A greedy approach is used to check if a candidate maximum sum allows splitting nums into at most k subarrays.
- Adjust the binary search bounds based on the feasibility check.

Space & Time

Time Complexity: $O(n * \log(\text{sum}(\text{nums})))$ – n elements are processed for each binary search iteration.

Space Complexity: $O(1)$ – Constant extra space is used.

Solution

The solution employs binary search to determine the smallest maximum subarray sum possible. We set the initial lower bound to the maximum element and the upper bound to the total sum of the array. For each candidate value (mid), we check if the array can be split into k or fewer subarrays without any subarray sum exceeding mid using a greedy approach. If possible, we try to lower the candidate; if not, we increase the candidate. This continues until the optimal solution is found.

#668 Kth Smallest Number in Multiplication Table **[Hard]** · Binary Search

Key Insights

- The table elements are not fully sorted row-wise or column-wise globally, but there is a pattern that can be exploited.
- For any candidate number x , we can count how many numbers in the table are less than or equal to x by summing over rows: for row i , there are $\min(x // i, n)$ numbers.
- Binary search can be used over the range $[1, m \cdot n]$ to find the smallest x such that the count is at least k .
- The counting function is the key to connecting the candidate value with the order statistic.

Space & Time

Time Complexity: $O(m \cdot \log(m \cdot n))$

Space Complexity: $O(1)$

Solution

We use binary search over the possible values in the multiplication table, which range from 1 to $m \cdot n$. For each midpoint x in the binary search, we compute the number of elements in the table that are less than or equal to x by iterating over each row and summing up $\min(x // i, n)$. This count tells us whether the k th smallest element is to the left or right of x . We then adjust the binary search boundaries accordingly. The approach leverages the multiplicative structure of the table and the monotonicity of the counting function.

#710 Random Pick with Blacklist [Hard] · Binary Search

Key Insights

- The valid output range has a total count of available numbers $W = n - \text{len}(\text{blacklist})$.
- Instead of repeatedly generating random numbers and checking against a set, we can transform the range into a contiguous block $[0, W - 1]$ and map any blacklist number within that block to a valid number in the upper range.
- Mapping is precomputed: for any blacklisted number x in $[0, W - 1]$, assign it a valid number from $[W, n - 1]$ that is not blacklisted.
- This approach ensures $O(1)$ pick time after an $O(b)$ precomputation where $b = \text{length of blacklist}$.

Space & Time

Time Complexity: constructor – $O(b)$ for initializing data structures, $\text{pick}()$ – $O(1)$ per call.

Space Complexity: $O(b)$ for storing the blacklist set and mapping.

Solution

We compute the count of valid numbers, $W = n - \text{len}(\text{blacklist})$. For numbers in the blacklist that are less than W , we need to map them to a valid number from the range $[W, n - 1]$. To do this, we first store all blacklist numbers in the higher range $[W, n - 1]$ in a set for quick

lookup. Then for each blacklisted value in $[0, W-1]$, we iterate from W upwards to find a number that is not blacklisted and assign it in a mapping. During the `pick()` function, we generate a random integer r in $[0, W-1]$. If r is in our mapping, we return its corresponding mapped value; otherwise, we return r as it is already a valid number.

#774 Minimize Max Distance to Gas Station [Hard] · Binary Search

Key Insights

- The optimal "penalty" (i.e., the maximum distance between adjacent stations after placement) can be found using binary search over a continuous distance range.
- For each candidate distance (mid), simulate the number of gas stations required by iterating over each adjacent station gap and calculating:
 - $stations_required = \lceil gap / mid \rceil - 1$
 - If the total number of required additional stations is less than or equal to k , then mid is a feasible candidate; otherwise, it is too small.
 - Continue the binary search until the gap between low and $high$ is within the desired tolerance ($1e-6$).

Space & Time

Time Complexity: $O(n * \log(\max_gap / \text{precision}))$, where n is the number of stations.

Space Complexity: $O(1)$ (ignoring input storage)

Solution

We use binary search over the possible maximum distances. The algorithm sets $low = 0$ and $high =$ the maximum gap between adjacent stations. For a candidate distance mid , we count how many new stations are required for all gaps by computing $\lceil gap / mid \rceil - 1$. If the total required is within k , we try a smaller candidate; otherwise, we increase the candidate distance. This approach efficiently zooms into the minimum possible maximum distance (penalty).

#1235 Maximum Profit in Job Scheduling [Hard] · Binary Search

Key Insights

- Sort jobs based on their end times to simplify finding non-conflicting jobs.
- Use dynamic programming where the state represents the maximum profit until a given point.
- For each job, perform a binary search to find the last job that ends before the current job's start time.
- The recurrence relation is: $dp[i] = \max(dp[i-1], \text{profit}[i] + dp[\text{index}])$ where $dp[\text{index}]$ is from the latest non-conflicting job.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting and binary search for each job.

Space Complexity: $O(n)$ for the dp arrays.

Solution

The solution begins by sorting the jobs by their end times. For each job, a binary search is used to quickly locate the last job that does not conflict with the current job (i.e., whose end time is less than or equal to the current job's start time). Then, using dynamic programming, the maximum profit is updated by choosing whether to include the current job or skip it. Two parallel arrays (or lists) are maintained: one for the end times of the jobs considered so far, and another to store the corresponding maximum profit at those times. This approach efficiently computes the answer while keeping time complexity within acceptable bounds.

#1074 Number of Submatrices That Sum to Target **[Hard]** · Matrix

Key Insights

- Transform the 2D submatrix sum problem into multiple 1D subarray sum problems.
- Precompute cumulative sums by fixing two row boundaries and then summing columns between these rows.
- Use a hash table (or dictionary) to count the number of subarrays with a given sum, reducing the problem to the

classic "subarray sum equals k" question.

- Iterate over all possible row pairs and for each, treat the column sums as an array where the target sum is sought.

Space & Time

Time Complexity: $O(n^2 * m)$, where n is the number of rows and m is the number of columns.

Space Complexity: $O(m)$, used by the hash table for storing cumulative sums along the columns.

Solution

The idea is to iterate over all pairs of rows (top and bottom) and, for each pair, compute the cumulative column sums between these rows. This converts the problem into finding the number of subarrays in a one-dimensional array (the column sums) that add up to the target. For each such 1D problem, use a hash table to track the frequency of cumulative sums seen so far and count how many times the difference (current sum - target) has been seen. This count gives the number of subarrays ending at the current index which sum to target. Combining counts from all possible row pairs yields the final result.

#124 Binary Tree Maximum Path Sum **[Hard]** · Trees & BST

Key Insights

- Use Depth-First Search (DFS) to explore the tree in a post-order fashion.
- For each node, calculate the maximum gain including that node and optionally one of its subtrees.
- A node's best contribution to its parent can be either just its own value or its value plus the maximum gain from either the left or right child.
- Update a global maximum path sum that considers the possibility of combining both left and right gains with the current node.
- Handle negative values by comparing gains against 0 to avoid decreasing the path sum.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the tree, since each node is visited exactly once.

Space Complexity: $O(n)$ in the worst-case scenario due to the recursion stack ($O(H)$ where H is the height of the tree, which can be $O(n)$ for skewed trees).

Solution

We use a recursive DFS approach (post-order traversal) to compute the maximum gain from each subtree. For every node:

- Recursively compute the maximum gain from the left and right child, ensuring that negative gains are treated as 0.
- The maximum gain that can be provided to the node's parent is the node's value plus the larger gain

from its left or right child. – Simultaneously, the potential maximum path sum including the current node as the highest (or root) node is the node's value plus both left and right gains. – Update a global variable that holds the maximum path sum found so far. – Return the maximum gain (node value plus one best child) to be used by the node's parent.

This strategy ensures that every possible path sum is considered while avoiding double-counting nodes.

#297 Serialize and Deserialize Binary Tree [Hard] · Trees & BST

Key Insights

- The tree can be traversed using either depth-first search (DFS) or breadth-first search (BFS).
- Using BFS (level-order traversal) simplifies reconstruction since nodes are processed level-by-level.
- Placeholder markers (e.g., "null") are used to indicate missing children.
- Both serialization and deserialization processes should handle empty trees and edge cases gracefully.
- Ensuring a one-to-one mapping between the tree structure and the serialized string is critical.

Space & Time

Time Complexity: $O(n)$ for both serialization and deserialization, where n is the number of nodes in the tree.

Space Complexity: $O(n)$ due to the storage used for the serialized string and additional data structures (queue) during processing.

Solution

This solution uses level-order traversal (BFS) for both serialization and deserialization. During serialization, each node is processed sequentially with missing children recorded as "null". For deserialization, the string is split by commas, and a queue is used to rebuild the tree level by level. The root is created from the first token, then subsequent tokens are assigned as left and right children accordingly, ensuring an accurate reconstruction of the original tree.

#269 Alien Dictionary [Hard] · DFS

Key Insights

- Build a graph where each unique character is a node.
- Compare adjacent words to determine the first different character, which gives the ordering (directed edge).
- Ensure all characters, even those without edges, are represented.
- Use a topological sort (using DFS or BFS) to determine a valid character ordering.

- Handle edge cases, such as words where a longer word appears before its prefix which is invalid.

Space & Time

Time Complexity: $O(N * L + V + E)$ where N is the number of words, L is the average length of each word, V is the number of unique characters, and E is the number of edges derived from adjacent word comparisons.

Space Complexity: $O(V + E)$ for storing the graph and additional data structures for topological sorting.

Solution

We use a graph-based approach to solve the problem. First, we initialize a graph (adjacency list) and in-degree dictionary for all unique characters. We then iterate through each pair of consecutive words, comparing them character by character to find the first pair that differs. This difference implies an order: the first character should come before the second character. We add a directed edge accordingly and update the in-degree for the target character.

To get the final ordering, we perform a topological sort using Kahn's algorithm:

- Enqueue characters with in-degree zero.
 - Remove characters from the queue and append them to our result.
 - Decrement the in-degree of neighboring nodes.
 - If a neighbor's in-degree becomes zero, enqueue it.
- If the result's length equals the number of unique nodes,

we have a valid ordering; otherwise, a cycle exists, and we return an empty string.

#329 Longest Increasing Path in a Matrix [Hard] · DFS

Key Insights

- The problem can be modeled as a graph where each cell is a node connected to its neighbors if the neighbor's value is greater.
- A Depth First Search (DFS) with memoization is effective to avoid recalculations.
- Dynamic Programming (DP) is used to cache the longest path found from each cell.
- Since each cell can be a start point, compute DFS for every cell and update the global maximum.

Space & Time

Time Complexity: $O(m * n)$ – each cell is computed once due to memoization.

Space Complexity: $O(m * n)$ – space is used for the memoization cache and recursion stack.

Solution

We solve the problem using a DFS approach combined with memoization to keep track of already computed longest paths from each cell. For each cell in the matrix, a DFS is initiated to explore all four valid directions (up, down, left, right). If a neighboring cell has a larger value,

the DFS recurses from that neighbor. The result for each cell is cached in a 2D memoization table, ensuring that each cell is processed only once. After processing all cells, the maximum value in the memo table represents the longest increasing path in the entire matrix.

#685 Redundant Connection II [Hard] · DFS

Key Insights

- The extra edge introduces one of two issues:
- A node has two parents.
- A cycle exists in the graph.
- First, scan through edges to detect if any node gets two parents. If found, there are two candidate edges.
- Use the Union-Find data structure (Disjoint Set Union) to detect cycles.
- Depending on whether a conflict (two parents) exists and/or a cycle is detected, decide which edge to remove:
- If a node has two parents, test by ignoring the second candidate edge during union-find. If no cycle forms, the second candidate is removable; otherwise, remove the first candidate.
- If no conflict exists, simply remove the edge that causes a cycle.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution consists of two main steps:

- Traverse the list of edges to detect whether any node has two parents. If such a node is found, mark the two candidate edges (candidate1: the earlier edge, candidate2: the later edge).
- Use the Union-Find algorithm to iterate over the edges: If there is a conflict (i.e., a node with two parents), skip candidate2 initially. While processing edges with Union-Find, check if adding an edge would create a cycle. Finally, based on the presence of a cycle and the two-parent conflict, decide which candidate edge to return.

Key Data Structures:

- An array (or hash map) to track the parent of each node.
- The Union-Find (Disjoint Set Union) structure to union sets and detect cycles.

#1036 Escape a Large Maze [Hard] · DFS

Key Insights

- The maximum number of blocked cells is 200, which limits the size of a potential completely enclosed region.
- Even if the target is unreachable by direct search, if the BFS (or DFS) explores more than a calculated threshold (based on the blocked cells count), then the region is not fully enclosed.
- Perform two BFS explorations: one starting from the source and one from the target. If neither search is

trapped within a small region, then an escape/path exists.

- Early exit in BFS if the target is reached or if the number of visited nodes exceeds the maximum possible enclosed area.

Space & Time

Time Complexity: $O(B^2)$ in the worst case, where B is the number of blocked cells (with $B \leq 200$).

Space Complexity: $O(B^2)$ due to the storage of visited cells in the worst-case scenario.

Solution

The solution uses the Breadth-First Search (BFS) algorithm. A key observation here is that the blocked cells can only enclose a limited area. We compute a threshold (maxSteps) based on the number of blocked cells ($n*(n-1)/2$) to determine if the search is confined. The BFS from the source (and similarly from the target) stops in two cases: if the target is reached or if the visited area exceeds the threshold, indicating that the region is not fully enclosed by blocked cells. Running BFS from both the source and target ensures that neither is trapped inside a small banned region.

#1192 Critical Connections in a Network [Hard] · DFS

Key Insights

- The problem is about identifying bridges in an undirected graph.
- Use Depth-First Search (DFS) to explore the graph.
- During DFS, track discovery times and low-link values for each node.
- For an edge (u, v) , if $\text{low}[v] > \text{disc}[u]$ then (u, v) is a critical connection.
- Tarjan's algorithm is well-suited for solving this problem in $O(n + m)$ time.

Space & Time

Time Complexity: $O(n + m)$ where n is the number of nodes and m is the number of connections.

Space Complexity: $O(n + m)$ for the graph representation and DFS recursion stack.

Solution

We solve this problem using a DFS-based approach (Tarjan's algorithm). First, we transform the list of connections into an adjacency list for efficient graph traversal. We then start a DFS from an arbitrary node (node 0, since the graph is connected) while maintaining two arrays: one for discovery times (disc) and one for low-link values (low). The discovery time is the time when the node is first visited, and the low-link value is the smallest discovery time that can be reached from that node (including via back edges). For each edge (u, v) , after a DFS call on v , if $\text{low}[v] > \text{disc}[u]$ then the edge (u, v) is a bridge because v (and its descendants) cannot

reach u or any of its ancestors without this edge. This edge is then recorded as a critical connection.

#305 Number of Islands II [Hard] · Graphs

Key Insights

- Use Union Find (Disjoint Set Union) to efficiently manage connected components (islands).
- Map 2D grid coordinates to a 1D index.
- When adding a new land, check its 4 neighbors (up, down, left, right) and perform unions if they are also lands.
- Avoid duplicate processing if the same cell is added multiple times.
- Maintain a counter for islands that is updated during each union operation.

Space & Time

Time Complexity: $O(K * \alpha(m$

$n))$ where K is the number of operations and α is the inverse Ackermann function (almost constant time per union/find).

Space Complexity: $O(m$

$n)$ for storing the union-find structure.

Solution

The solution relies on the Union Find data structure to dynamically merge islands. When a new land cell is added:

– If the cell is already land, record the current island count. – Otherwise, mark the cell as land and increment the island count. – Check the four neighboring cells; if any neighbor is land, attempt to union the current cell with that neighbor. – If a union actually happens (i.e., the two cells were in separate islands), decrement the island counter. – Output the island count after each add-land operation. The key trick is converting 2D indices to a 1D representation ($r * n + c$) and using path compression along with union by rank in the Union Find implementation for optimal performance.

#1153 String Transforms Into Another String [Hard] · Graphs

Key Insights

- Each character mapping from str1 to str2 must be consistent; one character in str1 always maps to the same character in str2.
- A bijective mapping is not necessary; multiple characters in str1 can map to the same character in str2.
- The order of conversions matters. When there is a cycle (e.g., mapping 'a'→'b' and 'b'→'a'), an intermediate unused character might be required to break the cycle.
- If str2 uses all 26 letters, no temporary character is available to break mapping cycles. In this case, transformation is possible only if str1 is already equal to str2.

Space & Time

Time Complexity: $O(n)$, where n is the length of the strings, because we traverse the strings to build and check the mapping.

Space Complexity: $O(1)$ or $O(26)$, which is constant space, to store the mapping for the 26 lowercase English letters.

Solution

The approach is as follows:

- If str1 equals str2, return true immediately.
- Build a mapping from each character in str1 to the corresponding character in str2. While creating the mapping, check that each character maps consistently.
- Count the number of distinct characters in str2. If this number is 26 and str1 is not already equal to str2, it is impossible to perform the transformation because there is no spare character available as a temporary placeholder.
- If the mapping is consistent and there exists at least one character not used in str2 (i.e., distinct characters in str2 is less than 26), the transformation is possible. The important trick is identifying the need for a free character when the target mapping is “full” (i.e., covers all 26 letters). Without a free character, you cannot avoid cycles where a temporary intermediate is necessary.

#127 Word Ladder [Hard] · BFS

Key Insights

- Use Breadth-First Search (BFS) to explore transformations level by level to guarantee the shortest path.
- Preprocess words by generating intermediate representations (e.g., replacing each letter with a wildcard) to quickly find neighbors.
- Use a visited set to avoid cycles.
- Each transformation counts as a level in the BFS, so the answer is the number of levels traversed.

Space & Time

Time Complexity: $O(N * M^2)$ where N is the number of words and M is the length of each word.

Space Complexity: $O(N * M)$ for the storage of intermediate mappings and the BFS queue.

Solution

We solve the problem using a BFS approach. First, we preprocess the wordList by creating a mapping of all possible intermediate states. For example, for the word "dog" and wildcard position, we generate "og", "d g", and "do*". This helps us find all adjacent words (neighbors) that are only one letter different in constant time. We then initialize a queue with the beginWord and initiate the BFS. For each word dequeued, we explore all valid transformations using the precomputed intermediate mapping. If the endWord is reached during this process, we return the current level (transformation

sequence length). Otherwise, we continue the BFS until all possibilities are exhausted. If no valid transformation sequence is found, we return 0.

#317 Shortest Distance from All Buildings [Hard] · BFS

Key Insights

- We need to compute the Manhattan distance from each building (cell value 1) to all empty lands (cell value 0) by performing a BFS.
- Maintain two matrices: one for cumulative total distances and another for the count of buildings that reached a particular empty cell.
- Ensure that only those empty cells reachable by every building are considered.
- The ideal cell will have a reach count equal to the total number of buildings.
- If multiple buildings cannot all reach any specific empty cell, then return -1.

Space & Time

Time Complexity: $O(k * m * n)$ where k is the number of buildings and m, n are the grid dimensions.

Space Complexity: $O(m * n)$ due to the auxiliary matrices and visited state tracking used in BFS.

Solution

We first iterate through the grid to locate all buildings while counting them. For each building, we perform a BFS to compute the shortest distance to each reachable empty cell. During the BFS, we update the cumulative distance for that cell and record that it was reached by one additional building. After processing all the buildings, we scan the grid to find the empty cell that is reachable from all buildings and has the minimum cumulative distance. If no such cell exists, we return -1 .

#773 Sliding Puzzle **[Hard]** · BFS

Key Insights

- Represent the board as a string (e.g. "123450") to simplify state comparison and manipulation.
- Use Breadth First Search (BFS) to explore all possible moves since the puzzle has a small fixed state space ($6! = 720$ possible states).
- Precompute the valid neighbor indices for each board position to efficiently generate new states.
- Track visited states to avoid loops and redundant work.

Space & Time

Time Complexity: $O(1)$ in practice, since there are at most 720 states ($6!$ permutations).

Space Complexity: $O(1)$ for the same reason, as the state space is fixed and limited.

Solution

We solve the problem using a Breadth First Search (BFS) approach. The board configuration is flattened into a string to represent states. A mapping from each index to its possible neighbors (i.e., positions to which 0 can move) is precomputed. BFS is then used to explore every valid move from the initial state while marking states as visited to avoid cycles. When the goal state ("123450") is reached, we return the number of moves taken. If BFS completes without reaching the goal, the puzzle is unsolvable and we return -1 .

#815 Bus Routes **[Hard]** · BFS

Key Insights

- Model the problem as a graph where nodes are bus routes.
- Two bus routes are connected if they share at least one common bus stop.
- Use a hash map to quickly look up bus routes available at any bus stop.
- Employ Breadth-First Search (BFS) on the bus routes graph to find the minimum number of bus transfers.
- Mark visited bus routes and bus stops to avoid redundant work.

Space & Time

Time Complexity: $O(N * L)$ in the worst-case scenario, where N is the number of bus routes and L is the average number of stops per route.

Space Complexity: $O(N * L)$ due to storage of the stop-to-buses mapping and BFS auxiliary data structures.

Solution

We begin by constructing a mapping from each bus stop to the list of buses passing through it. This allows quick retrieval of buses available at a given stop. Starting at the source stop, initialize the BFS queue with all buses that can be boarded at that stop. Then, for each bus route in the queue, inspect each stop along that route. If a stop is the target, return the number of buses taken so far. Otherwise, for each unvisited stop, add all adjacent bus routes (buses that pass through this stop) to the BFS queue with an incremented bus count. Continue until the queue is empty, which means the target cannot be reached.

R4 · Mid · Easy (12 q)

>> Round 4 · Mid Priority · Easy — Quick Review

12 questions · key insights · complexity · solution

#21 Merge Two Sorted Lists [Easy] · Linked List

Key Insights

- Leverage the fact that both lists are sorted.
- Use two pointers (one for each list) to compare nodes sequentially.
- Create a dummy node to simplify edge cases and build the resulting list.
- Iterate until one list is exhausted, then append the remaining nodes from the other list.
- There is also a recursive approach that can simplify the logic but may use more stack space.

Space & Time

Time Complexity: $O(n + m)$, where n and m are the lengths of the two lists.

Space Complexity: $O(1)$ for the iterative solution ($O(n + m)$ stack space for the recursive solution).

Solution

We use an iterative approach with a dummy node to merge two sorted linked lists. Two pointers traverse

list1 and list2; at each step, the node with the smaller value is appended to our merged list. This process repeats until one of the lists is exhausted, after which we append the remaining nodes from the other list. This approach works in one pass over both lists, ensuring linear time complexity and constant extra space usage.

#141 Linked List Cycle [Easy] · Linked List

Key Insights

- Use two pointers (slow and fast) to traverse the list.
- The fast pointer moves two steps at a time and the slow pointer moves one step.
- If the fast pointer ever meets the slow pointer, a cycle exists.
- If the fast pointer reaches the end of the list (null), the list has no cycle.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes in the linked list.

Space Complexity: $O(1)$ as only two pointers are used.

Solution

We use the Floyd's Cycle Detection algorithm (Tortoise and Hare approach). Two pointers, slow and fast, are initialized at the head of the list. The slow pointer moves one step at a time while the fast pointer moves two steps. If the linked list has a cycle, the fast pointer

will eventually meet the slow pointer. If the fast pointer reaches the end (i.e., a null pointer), the linked list does not have a cycle.

#206 Reverse Linked List [Easy] · Linked List

Key Insights

- The problem can be solved using either an iterative or recursive approach.
- Iterative method uses constant space by updating pointers in-place.
- Recursive method leverages the call stack to reverse the list, which could use $O(n)$ space in the worst case.
- Both approaches traverse each node exactly once, yielding linear time complexity.

Space & Time

Time Complexity: $O(n)$ for both iterative and recursive approaches, where n is the number of nodes.

Space Complexity: $O(1)$ for the iterative approach and $O(n)$ for the recursive approach because of the recursion call stack.

Solution

We can solve the problem by iterating through the list while reversing the next pointers of each node. We maintain two pointers (previous and current) and update them as we traverse the list, so that by the end, the previous pointer is at the new head of the reversed list.

Alternatively, a recursive approach reverses the rest of the list first and then adjusts the pointers so that the original head becomes the tail.

#234 Palindrome Linked List [Easy] · Linked List

Key Insights

- Use two pointers (slow and fast) to reach the middle of the list.
- Reverse the second half of the list.
- Compare the nodes from the start of the list to the reversed half.
- This approach meets the $O(n)$ time and $O(1)$ space follow-up requirement.

Space & Time

Time Complexity: $O(n)$ – each node is visited a constant number of times.

Space Complexity: $O(1)$ – only constant extra space is used (in-place reversal).

Solution

We first use the two-pointer technique to locate the midpoint of the linked list. The slow pointer moves one step at a time while the fast pointer moves two steps. Once the middle is found, we reverse the second half of the list in-place. Finally, we traverse both halves concurrently to compare the node values. If all corresponding values match, the linked list is a

palindrome. A potential gotcha is handling the odd number of nodes; in that case, we ignore the middle element when comparing.

#706 Design HashMap [Easy] · Linked List

Key Insights

- Use a fixed-size array of buckets where each bucket uses a linked list to handle collisions.
- Choose a simple hash function (e.g., modulo operation) to map a key to its corresponding bucket.
- For each bucket, traverse the linked list to update an existing key or append a new node if the key does not exist.
- Removal requires updating pointers in the linked list to bypass the removed node.

Space & Time

Time Complexity: Average $O(1)$ for put, get, and remove; worst-case $O(n)$ when many keys hash to the same bucket.

Space Complexity: $O(n)$, where n is the number of key-value pairs stored.

Solution

We implement the HashMap using an array of fixed size (e.g., 1000). Each array element is the head of a linked list that manages collisions using chaining. The key is hashed using a modulo operation. For the put operation,

if a key already exists in the chain, its value is updated; otherwise, a new node is appended. The get operation traverses the chain at the computed bucket index to find and return the corresponding value, or returns -1 if not found. The remove operation also traverses the list, carefully reconnecting the links to remove the target node while handling edge cases such as removal at the head.

#876 Middle of the Linked List [Easy] · Linked List

Key Insights

- Use two pointers (slow and fast) to traverse the list.
- The fast pointer moves two steps at a time while the slow pointer moves one step.
- When the fast pointer reaches the end, the slow pointer will be at the middle node.
- This approach naturally handles both odd and even lengths of the list.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list.

Space Complexity: $O(1)$, constant extra space.

Solution

The solution uses the two-pointer technique:

- Initialize two pointers, slow and fast, at the head.
- Traverse the list, moving slow by one step and fast by

two steps during each iteration. – When fast reaches the end or has no next node, slow will be at the middle. – Return the slow pointer (the middle node).

This strategy ensures linear time complexity and constant space, efficiently finding the middle node even for lists with an even number of elements.

#1474 Delete N Nodes After M Nodes of Linked List [Easy] · Linked List

Key Insights

- Traverse the linked list while keeping count.
- Skip m nodes (i.e., keep them) and then proceed to delete the next n nodes.
- Update pointers to bypass the removed nodes.
- This is an in-place modification problem, so no additional list structure is needed.
- Be cautious about edge cases when the list has fewer than m or n nodes remaining.

Space & Time

Time Complexity: $O(N)$, where N is the number of nodes in the list.

Space Complexity: $O(1)$ since the modification is done in-place.

Solution

The solution involves iterating through the linked list with a pointer. For each cycle:

– Traverse m nodes and keep them. – Then, from the $(m+1)$ th node, skip and remove the next n nodes by adjusting the 'next' pointer of the m -th node to point to the node after these n nodes. – Continue until you reach the end of the list. This in-place approach avoids using extra space and maintains $O(N)$ time complexity.

#20 Valid Parentheses [Easy] · Stack

Key Insights

- The problem can be solved using a stack data structure.
- When encountering an opening bracket, push it onto the stack.
- When encountering a closing bracket, check if the top of the stack is its matching opening bracket.
- If the stack is empty before finding a match or if there is a mismatch, the string is invalid.
- At the end, the string is valid only if the stack is empty, meaning all brackets were properly matched.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string, since we process each character once.

Space Complexity: $O(n)$ in the worst case for the stack when all characters are opening brackets.

Solution

We use a stack to check for valid parentheses. The algorithm iterates through the string, pushing opening brackets onto the stack and popping them when a corresponding closing bracket is encountered. A mapping of closing to opening brackets helps in validating the matching pairs. By ensuring that the stack is empty at the end, we confirm that every opening bracket was properly closed.

#232 Implement Queue using Stacks [Easy] · Stack

Key Insights

- Use two stacks:
- One (input stack) for enqueue operations.
- The other (output stack) for dequeue operations.
- When performing pop or peek, if the output stack is empty, transfer all elements from the input stack to the output stack. This reversal makes the oldest element accessible.
- Each element is moved at most once between the stacks, ensuring an amortized $O(1)$ cost per operation.
- The space complexity is $O(n)$ where n is the number of elements in the queue.

Space & Time

Time Complexity: Amortized $O(1)$ per operation

Space Complexity: $O(n)$

Solution

We maintain two stacks: an input stack for pushing new elements and an output stack for popping or peeking the front element of the queue. When we need to access the front of the queue (either during a pop or peek), we check if the output stack is empty. If it is, we transfer all elements from the input stack to the output stack. This reversal of order allows us to simulate the FIFO nature of the queue. When performing a push, the element is simply added to the input stack.

#844 Backspace String Compare [Easy] · Stack

Key Insights

- '#' represents a backspace operation that deletes the previous character.
- A direct simulation using a stack can reconstruct the final string.
- A two-pointer approach allows comparison from the end of the strings while skipping backspaced characters.
- The two-pointer technique can achieve $O(n)$ time and $O(1)$ space by processing the strings in reverse.

Space & Time

Time Complexity: $O(n)$ where n is the length of the longer string.

Space Complexity: $O(1)$ using the two-pointer strategy ($O(n)$ if using stacks).

Solution

Round 9 · Low · Hard

p.3283

The solution uses two pointers starting from the end of each string. For each string, maintain a skip counter to account for backspace characters. As you iterate backwards:

- If a '#' is encountered, increment the skip counter.
- If a normal character is encountered while the skip counter is positive, skip that character by decrementing the counter.
- Once a valid character is found, compare the characters from both strings. If all corresponding characters match and both pointers finish processing at the same time, the strings are considered equal.

#14 Longest Common Prefix [Easy] · Trie

Key Insights

- Compare characters of the strings one column at a time (vertical scanning).
- The longest possible prefix is limited by the length of the shortest string in the array.
- As soon as a mismatch is found, the prefix up to that point is the answer.
- Handle edge cases such as an empty input array or strings with no common prefix.

Space & Time

Time Complexity: $O(N \cdot M)$ where N is the number of strings and M is the length of the shortest string.

Round 9 · Low · Hard

p.3284

Space Complexity: $O(1)$ additional space (ignoring output space).

Solution

We use the vertical scanning approach:

- If the input array is empty, return an empty string.
- Iterate character by character over the first string.
- For each character, compare it with the corresponding character in each of the other strings.
- If a mismatch or end of any string is encountered, return the substring from the start up to the current index.
- If no mismatch is found, the entire first string is the common prefix.

This solution ensures that we only traverse the necessary part of the strings and stops early if a discrepancy is discovered.

#409 Longest Palindrome [Easy] · Greedy

Key Insights

- Count the frequency of each character.
- For each character, you can use pairs (even count or odd count minus one) to form the palindrome.
- If any character has an odd frequency, you can add one extra character as the center of the palindrome.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(1)$ since the frequency counter will have at most 52 keys (considering both uppercase and lowercase letters).

Solution

The solution uses a hash table to count the frequency of each character. For every character, the maximum number of characters that can be used in a palindrome is determined by taking the largest even number that is less than or equal to its frequency (which is $\text{frequency} - (\text{frequency} \% 2)$). If there exists any character with an odd frequency, one of these can be used as the center of the palindrome, adding an extra character to the total length. This greedy approach efficiently computes the length of the largest possible palindrome from the available characters.

R5 · Mid · Medium (56 q)

>> Round 5 · Mid Priority · Medium — Quick Review

56 questions · key insights · complexity · solution

#2 Add Two Numbers [Medium] · Linked List

Key Insights

- The digits are stored in reverse order, meaning the first node is the least significant digit.
- A carry may need to propagate through subsequent nodes.
- The two linked lists might be of different lengths.
- Continue processing until both lists are fully traversed and any remaining carry is handled.

Space & Time

Time Complexity: $O(\max(n, m))$, where n and m are the lengths of the two linked lists.

Space Complexity: $O(\max(n, m) + 1)$ for the result list in the worst case.

Solution

The solution iterates through both linked lists node by node. For each pair of corresponding nodes, their values and any carry from the previous operation are summed. The result is decomposed into a current digit (modulo

10) and an updated carry (total divided by 10). A new node is created for the current digit and appended to the resulting linked list. This process continues until all nodes are processed and any leftover carry is added as a final node.

#19 Remove Nth Node From End of List [Medium] · Linked List

Key Insights

- Use a dummy node at the start to simplify edge cases, especially when the head is removed.
- Utilize two pointers (fast and slow) and maintain a gap of $n+1$ nodes between them.
- By moving both pointers simultaneously, the slow pointer will point to the node immediately before the target node when the fast pointer reaches the end.

Space & Time

Time Complexity: $O(L)$, where L is the length of the list

Space Complexity: $O(1)$

Solution

We use a two-pointer approach to achieve a one-pass solution. A dummy node is created and linked to the head to handle edge cases gracefully (such as needing to remove the head). The fast pointer is advanced $n+1$ steps ahead of the slow pointer to maintain the required gap. Then, both pointers are moved simultaneously until

the fast pointer reaches the end of the list. At that point, the slow pointer is immediately before the node to be deleted. We then adjust the pointers to remove the target node from the list.

#24 Swap Nodes in Pairs [Medium] · Linked List

Key Insights

- Use a dummy node to simplify edge cases, such as when the list is empty or has only one node.
- Process nodes in pairs and update the pointers accordingly.
- Ensure that the swapping does not lose the connection between the swapped pair and the rest of the list.
- The solution can be implemented iteratively or recursively.

Space & Time

Time Complexity: $O(n)$ – Each node is processed once.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We use an iterative approach with a dummy node. The dummy node points to the head of the list. We maintain a pointer that traverses the list in pairs. For each pair, we adjust the next pointers to swap the two nodes. By the end, the dummy's next pointer will be the new head of the modified list. The main idea is to keep track of the

previous node (prev) before the pair, and then adjust pointers accordingly using temporary variables to hold references to the nodes being swapped.

#61 Rotate List [Medium] · Linked List

Key Insights

- Calculate the length of the list and locate the tail.
- Use k modulo length to handle rotations greater than the list size.
- Find the new tail of the list at position $(\text{length} - k - 1)$ and update pointers accordingly.
- Link the original tail to the original head to form a circular list, then break the circle at the new tail.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list, since we traverse the list a few times.

Space Complexity: $O(1)$, as the rotation is done in place without extra data structures.

Solution

We first traverse the list to determine its length and identify the tail node. The effective rotation is calculated by taking k modulo the length of the list. If the result is zero, the list remains unchanged. Otherwise, we locate the new tail node by moving $(\text{length} - k - 1)$ steps from the head. The node after the new tail becomes the new head. Finally, we break the link after the new tail and

connect the old tail to the original head to finalize the rotated list.

#146 LRU Cache [Medium] · Linked List

Key Insights

- Use a Hash Table (dictionary) to allow $O(1)$ access to nodes by key.
- Use a Doubly Linked List to record the usage order. The head represents the most recently used element while the tail represents the least recently used.
- When a key is accessed or updated, move the corresponding node to the front (head) of the list.
- When the cache exceeds capacity, remove the node at the tail end which is the least recently used entry.

Space & Time

Time Complexity: $O(1)$ per get and put operation.

Space Complexity: $O(\text{capacity})$, storing key-node pairs and linked list nodes.

Solution

The solution combines a hash table (or dictionary) and a doubly linked list:

- The hash table stores keys and corresponding node addresses.
- The doubly linked list maintains the order of usage with the most recently used items near the head.
- For `get(key)`, check the dictionary. If found, move the node to the head and return its value; if not, return

- 1. – For `put(key, value)`, if the key exists, update its value and move it to the head. If not, create a new node and add it right after the head. If the capacity is reached, remove the tail node (least recently used) and update the dictionary accordingly.
- Dummy head and tail nodes are used to simplify node addition and removal procedures without worrying about edge cases.

#148 Sort List [Medium] · Linked List

Key Insights

- Merge sort is well-suited for linked lists; it naturally relies on splitting the list and merging sorted sublists.
- The slow and fast pointer technique is used to find the middle point of the list.
- A recursive merge sort implementation may use $O(\log n)$ space for recursion stack. To achieve $O(1)$ extra space, an iterative bottom-up merge sort approach would be preferred.
- Merging two sorted lists is efficient with linked lists because of their sequential access.

Space & Time

Time Complexity: $O(n \log n)$

Space Complexity: $O(1)$ extra space for an iterative bottom-up approach ($O(\log n)$ for recursive calls if using recursion)

Solution

The solution employs the merge sort algorithm tailored for linked lists. The overall steps are:

- Use the slow and fast pointers to find the middle of the list and split it into two halves.
 - Recursively sort the two halves.
 - Merge the two sorted halves by comparing node values.
 - Ensure that during merging, pointers are adjusted properly to build the sorted list.
- For an $O(1)$ space solution, an iterative bottom-up merge sort can be implemented. However, the recursive merge sort is easier to understand and implement, though it uses $O(\log n)$ space on the call stack.

#328 Odd Even Linked List [Medium] · Linked List

Key Insights

- Use two pointers: one for odd-indexed nodes and one for even-indexed nodes.
- The first node is considered odd, and the second node is even.
- Traverse the list once, reassigning the next pointers to segregate odd and even nodes.
- Finally, attach the even list at the end of the odd list.
- The solution must use $O(1)$ extra space and $O(n)$ time complexity.

Space & Time

Time Complexity: $O(n)$ since the list is traversed once.

Space Complexity: $O(1)$ as only a few pointers are used.

Solution

Maintain two pointers, odd and even, starting from the head and head.next, respectively. As you iterate, update odd.next to point to the next odd element (skipping an even node) and even.next to point to the next even element. Once the end is reached, attach the even list to the end of the odd list. This in-place rearrangement preserves the initial ordering within odd and even groups while meeting the space and time constraints.

#369 Plus One Linked List [Medium] · Linked List

Key Insights

- The digits are stored such that the most significant digit comes first, but addition is easier starting from the least significant digit.
- Reversing the list allows simpler management of the carry during addition.
- After processing the addition, reversing the list back restores the original order.
- Consider special cases where all digits are 9, requiring a new head node to accommodate an additional digit.

Space & Time

Time Complexity: $O(n)$, where n is the number of nodes in the list

Space Complexity: $O(1)$ extra space when using an iterative reversal method. ($O(n)$ if recursion is used)

Solution

We solve the problem by first reversing the linked list, which converts it to a format where the least significant digit is processed first. We then iterate over the reversed list, adding one and managing any carry. If we encounter a carry at the end of the list, a new node is appended to handle the overflow. Finally, we reverse the list again to restore the original order. This method efficiently handles the addition with constant extra space during processing.

#430 Flatten a Multilevel Doubly Linked List

[Medium] · Linked List

Key Insights

- Use a depth-first traversal to process child lists before proceeding to the next pointer.
- When a node with a child is found, recursively flatten the child list.
- Splice the flattened child list in between the current node and the next node.
- Maintain previous and next pointers correctly to preserve the doubly linked list properties.
- Ensure that all child pointers are set to null after flattening.

Space & Time

Time Complexity: $O(n)$, where n is the total number of nodes, since each node is processed exactly once.

Space Complexity: $O(n)$ in the worst-case due to recursion stack in case of a highly nested list.

Solution

We can solve this problem using a recursive depth-first search (DFS) approach. The main idea is to traverse the list and, whenever a node with a child pointer is encountered, recursively flatten its child list. Once the child list is flattened, it can be spliced into the main list between the current node and the node originally following it. We then continue processing the next pointer. This ensures that the entire list is flattened in depth-first order. The recursive helper function returns the tail of the flattened segment, which is used to connect the subsequent parts of the list correctly.

#622 Design Circular Queue [Medium] · Linked List

Key Insights

- Leverage a fixed-size array to store the queue elements.
- Use two pointers: one for the front and one for the rear.
- Maintain a count variable (or use pointer arithmetic) to differentiate between an empty and a full queue.
- Use modulo arithmetic to wrap pointers when they reach the end of the array.

- Ensure all operations (enqueue, dequeue, front, rear, isEmpty, isFull) have $O(1)$ time complexity.

Space & Time

Time Complexity: $O(1)$ per operation for enqueue, dequeue, front, rear, isEmpty, and isFull.

Space Complexity: $O(k)$, where k is the capacity of the circular queue.

Solution

We implement the circular queue using a fixed-size array alongside two pointers, front and rear, and a counter to keep track of the number of elements in the queue. The front pointer indicates the index of the first element, and the rear pointer indicates the index where the next element will be enqueued. When enqueueing an element, we insert the element at the rear index and then increment the rear pointer modulo the size of the array. Similarly, when dequeuing, we remove the element at the front index and increment the front pointer modulo the size. The counter helps in quickly checking if the queue is empty or full (count equals 0 or count equals capacity, respectively). This design ensures all operations remain $O(1)$ and correctly wraps around, forming a circular structure.

#707 Design Linked List [Medium] · Linked List

Key Insights

- Use a dummy head node to simplify insertion and deletion at the head.
- Maintain a size counter to quickly validate index-based operations.
- For every operation, traverse the list up to the required point since random access is not available.
- Ensure robustness by handling edge cases like negative indices and indices larger than the current length.

Space & Time

Time Complexity:

- get, addAtIndex, and deleteAtIndex operations require $O(n)$ time due to the need to traverse the list.
- addAtHead and addAtTail (if tail pointer is not maintained) are $O(n)$ in worst case; however, these can be optimized further if a tail pointer is added.

Space Complexity:

- $O(n)$ space for storing n nodes.

Solution

The solution employs a singly linked list with a dummy head node to ease edge-case handling. A size counter is maintained for quick index validations. Each node holds a value and a pointer to the next node. Operations operate by traversing the list to the correct insertion or deletion point. This design makes it simple to add the functionality of getting values, inserting at specified

indices, and removing nodes.

#708 Insert into a Sorted Circular Linked List

[Medium] · Linked List

Key Insights

- Handle the empty list edge-case by creating a new node that points to itself.
- Traverse the list until the proper insertion point is found.
- Consider the "turning point" where the maximum value is followed by the minimum value.
- If no proper place is found in one full rotation, insert the new node arbitrarily (e.g., after the head).

Space & Time

Time Complexity: $O(n)$ in the worst-case when a full loop is required.

Space Complexity: $O(1)$ since we only use a few extra pointers.

Solution

We traverse the circular linked list starting from the arbitrary given node. For each pair of consecutive nodes (curr and next), we check if the new value can be inserted between them by confirming that:

- The new value lies between curr and next ($\text{curr.val} \leq \text{insertVal} \leq \text{next.val}$).
- Or if we are at the pivot point where the maximum value meets the minimum value

(i.e., $\text{curr.val} > \text{next.val}$) and then either the new value is larger than or equal to curr.val (should be appended after the maximum) or less than or equal to next.val (should be inserted before the minimum).

If no proper spot is found after one complete loop, we insert anywhere (all values in the list are the same). A pointer reference is updated accordingly to maintain the circular structure.

#1171 Remove Zero Sum Consecutive Nodes from Linked List [Medium] · Linked List

Key Insights

- Use a dummy node to handle edge cases such as the removal of nodes at the head.
- Compute a running (prefix) sum while traversing the list.
- Use a hash table (or hashmap) to map each prefix sum to the most recent node where that sum occurred.
- If the same prefix sum is encountered again, the nodes in between sum to 0 and can be removed by adjusting pointers.
- Perform two passes: the first to record prefix sums and their corresponding nodes, and the second to remove zero-sum sequences.

Space & Time

Time Complexity: $O(n)$ – Two passes over the list.

Space Complexity: $O(n)$ – Hash table storing at most n prefix sums.

Solution

The solution uses a two-pass algorithm with a hash table. In the first pass, compute the prefix sum for each node and store the mapping from prefix sum to the node (overwriting with the latest occurrence). This ensures that for any repeated prefix sum, the nodes in between sum to 0. In the second pass, traverse the list again and use the hash table to skip the zero-sum subsequences by redirecting the next pointer of a node to the next pointer of the stored node for that prefix sum. Key data structures are the hash table and a dummy node to simplify pointer adjustments.

#143 Reorder List [Medium] · Stack

Key Insights

- Find the middle of the linked list using the fast and slow pointer technique.
- Reverse the second half of the list.
- Merge the two halves, alternating nodes from the first half and the reversed second half.
- The challenge is done in-place, ensuring a linear time solution with constant extra space.

Space & Time

Time Complexity: $O(n)$ – Each of the three main steps (finding the middle, reversing, merging) takes $O(n)$ time.

Space Complexity: $O(1)$ – The algorithm is done in-place with only a few pointers.

Solution

The solution involves three key steps:

- Use two pointers (slow and fast) to identify the middle of the list.
- Reverse the second half of the list. This can be done iteratively by adjusting the next pointers.
- Merge the two lists: One starting from the head and the other from the head of the reversed second half. Interleave the nodes by alternating between the two lists until complete. This approach uses in-place pointer manipulation, thus ensuring no extra space is used aside from a few temporary variables.

#150 Evaluate Reverse Polish Notation [Medium] · Stack

Key Insights

- Use a stack to process the tokens.
- Push numbers onto the stack.
- When an operator is encountered, pop the top two numbers, perform the operation, and push the result back.
- Handle division carefully to ensure it truncates toward zero.

- The final result remains as the only value in the stack.

Space & Time

Time Complexity: $O(n)$, where n is the number of tokens (each token is processed once).

Space Complexity: $O(n)$, in the worst-case scenario where all tokens are numbers and are stored in the stack.

Solution

We use a stack data structure to evaluate the RPN expression. Iterate through each token:

- If the token is a number, convert and push it onto the stack.
- If the token is an operator, pop the top two numbers (be careful about order), perform the corresponding arithmetic operation, and push the result back onto the stack.
- At the end of processing, the stack contains the result of the expression. Special attention is needed for division to ensure it truncates toward zero (using language-specific functions when necessary).

#155 Min Stack [Medium] · Stack

Key Insights

- Use an auxiliary data structure to track the current minimum element.
- A common approach is to use two stacks: one regular stack for all elements and a second stack to keep track

of the minimum values.

- When pushing a new element, push it onto the main stack and also update the min stack if the element is less than or equal to the current minimum.
- When popping, if the popped element is equal to the top of the min stack, pop from the min stack as well.
- This design guarantees that all operations (push, pop, top, getMin) run in constant time, $O(1)$.

Space & Time

Time Complexity: $O(1)$ per operation (push, pop, top, getMin)

Space Complexity: $O(n)$, where n is the number of elements in the stack, since an additional stack is used to keep track of minima.

Solution

The solution involves using two stacks:

- Main Stack: Stores all the elements.
- Min Stack: Stores the minimum elements encountered so far.

When an element is pushed, we compare it with the top of the min stack. If it is smaller than or equal to the current minimum, we also push it onto the min stack. For the pop operation, if the element being popped is the same as the top of the min stack, we pop from the min stack as well. This approach ensures that the getMin operation always returns the top element of the min stack, which is the current minimum in constant

time.

#173 Binary Search Tree Iterator [Medium] · Stack

Key Insights

- Use an iterative in-order traversal strategy with a stack.
- In-order traversal for BST yields sorted order.
- Pre-load the stack with all left descendants from the root.
- After returning a node, process its right subtree by pushing its left branch.
- This method maintains an average time complexity of $O(1)$ per operation and uses $O(h)$ memory, where h is the tree height.

Space & Time

Time Complexity: $O(1)$ on average per operation (amortized)

Space Complexity: $O(h)$, where h is the height of the tree

Solution

We use a stack to simulate an in-order traversal. The process involves initially pushing all left-child nodes from the root into the stack. For each call to `next()`, we pop the top element (the current smallest node) and then push all the left descendants of the node's right child. This ensures that the next smallest element is always at the top of the stack. The `hasNext()` function

then simply checks if the stack is non-empty.

#227 Basic Calculator II [Medium] · Stack

Key Insights

- The expression can be processed in one pass while using a stack to handle operator precedence.
- When encountering '+' or '-' operators, the current number is directly added or subtracted (as a positive or negative value).
- For '*' and '/' operators, compute immediately with the last number from the stack and update it.
- Handle spaces and multi-digit numbers by building the number while iterating over the string.
- Division truncation toward zero must be carefully applied, which aligns with typical integer division in most languages.

Space & Time

Time Complexity: $O(n)$ where n is the length of the expression, since the algorithm processes each character once.

Space Complexity: $O(n)$ in the worst-case scenario when all numbers are stored in the stack.

Solution

We use a stack-based approach. Iterate over each character of the string while maintaining a variable for the current number and a previous operator (initialized

to '+'). When an operator or the end of the string is reached, perform an action based on the previous operator:

– For '+' add the current number to the stack. – For '-' add the negative of the current number. – For '*' and '/' pop the last number from the stack, compute the result by performing multiplication or division, and then push the result back to the stack. After processing the entire string, sum the elements in the stack to get the final result. This approach ensures that multiplication and division are handled before addition and subtraction.

#255 Verify Preorder Sequence in Binary Search Tree [Medium] · Stack

Key Insights

- A BST has the property that every node's left subtree contains values less than the node, and every right subtree contains values greater.
- Preorder traversal processes nodes as: root, left subtree, then right subtree.
- Using a monotonic stack, we can simulate constructing the BST while tracking a lower bound for allowable node values.
- When moving to a right child, update the lower bound to ensure subsequent nodes are valid.

Space & Time

Time Complexity: $O(n)$ – Each element is processed once.

Space Complexity: $O(n)$ – In the worst-case an explicit stack is used; can be optimized to $O(1)$ by reusing the input array.

Solution

We use a monotonic stack to simulate the BST construction from the preorder sequence. The process is:

– Initialize a variable `lower_bound` to $-\infty$. – Iterate through each value in the preorder array: If a value is less than `lower_bound`, it violates BST properties and the sequence is invalid. While the stack is not empty and the current value is greater than the element at the top of the stack, pop from the stack and update `lower_bound` to the popped value. This simulates moving to the right subtree. Push the current value onto the stack. – If the sequence is processed without violations, return true. This approach efficiently ensures that each value adheres to BST requirements during preorder traversal.

#394 Decode String [Medium] · Stack

Key Insights

- Use a stack or recursion to manage nested encoded patterns.
- Process the string character by character, accumulating digits for the repeat count.

- When encountering a '[', start a new context (recursively or via a new stack frame) for the encoded substring.
- When encountering a ']', finish the current context and repeat the accumulated substring as needed.
- Careful handling of multi-digit numbers is essential.

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(n)$, for the recursion stack or the auxiliary stack used.

Solution

We can solve this problem using a recursive approach. When we see a digit, we determine the full repeat count. Upon encountering a '[', we recursively read the encoded substring until the corresponding ']'. Each recursive call returns the decoded substring until it hits a ']', which is then repeated based on the previously computed count. Alternatively, an iterative solution using a stack is also feasible. As we parse the string, push the current string and number onto the stack when encountering '[', then pop and combine when ']' is reached. This approach correctly handles nested patterns.

#439 Ternary Expression Parser [Medium] · Stack

Key Insights

- The ternary operator groups right-to-left. This property can be exploited by processing the string from the end to the beginning.
- Using a stack allows us to "collapse" nested ternary expressions as soon as their three components (condition, true branch, false branch) are available.
- An iterative approach is efficient and avoids building an explicit syntax tree for the expression.
- The condition is not stored on the stack; instead, it is encountered immediately preceding the '?' operator when traversing backwards.

Space & Time

Time Complexity: $O(n)$, where n is the length of the expression, as we process each character exactly once.

Space Complexity: $O(n)$, due to the use of a stack for storing characters during evaluation.

Solution

The solution uses an iterative approach with a stack. We traverse the expression string from right to left. For every character:

- If it is a digit, 'T', or 'F' (and not a ':' or part of a pending operator), we push it onto the stack.
- When we encounter a '?' character, we know that the next available items on the stack represent the true and false results of the ternary expression. We then skip the ':' separator that comes between them.
- We then use the condition (the character immediately preceding the '?')

to choose which result to push back on the stack. This method “collapses” each ternary expression into a single result, and by the time we finish traversing the string, the stack contains the final evaluated result.

#484 Find Permutation [Medium] · Stack

Key Insights

- The answer is a permutation of numbers from 1 to n , where $n = \text{len}(s) + 1$.
- A common strategy is to use a stack to reverse the segments when a decrease ('D') is encountered.
- When an 'I' is encountered (or at the end), pop all numbers from the stack to output them in reverse order, ensuring the smallest lexicographical order.
- This greedy approach guarantees that as soon as a rising pattern is detected, we “flush” the pending decreases correctly.

Space & Time

Time Complexity: $O(n)$ – Each element is pushed and popped exactly once.

Space Complexity: $O(n)$ – For the stack and the output array.

Solution

We use a greedy approach with a stack. Iterate from 0 to n (where $n = \text{len}(s) + 1$):

– Push the current number ($i+1$) onto the stack. – If the current character is 'I' or we have reached the end of the string, pop all elements from the stack and append them to the result. By doing so, segments corresponding to series of 'D's are reversed, while 'I's trigger the flush, ensuring that the final permutation is the lexicographically smallest.

#735 Asteroid Collision [Medium] · Stack

Key Insights

- Use a stack to keep track of the surviving asteroids.
- Only a collision happens when a right-moving asteroid (positive) is on the stack and a left-moving asteroid (negative) is encountered.
- Compare the absolute values to decide if one or both asteroids explode.
- Continue processing collisions until no further collisions are possible with the current asteroid.

Space & Time

Time Complexity: $O(n)$ where n is the number of asteroids, since each asteroid is pushed and popped at most once.

Space Complexity: $O(n)$ in the worst case, when no collisions occur and all asteroids remain on the stack.

Solution

The solution uses a stack to simulate the asteroid collision process. For each asteroid in the input:

- If it is moving to the right or the stack is empty, simply add it to the stack.
- If it is moving to the left, check the asteroid on top of the stack: If the top asteroid is moving to the right, they might collide. Compare sizes: If the left-moving asteroid is larger (in absolute terms), pop the stack and continue checking. If both are equal, pop the stack and do not add the current asteroid. If the right-moving asteroid is larger, the current asteroid explodes.
- If no collision occurs for the current asteroid, push it onto the stack. Finally, the stack represents the surviving asteroids order.

#739 Daily Temperatures [Medium] · Stack

Key Insights

- Use a monotonic stack to keep track of indices of days with unresolved warmer temperatures.
- As you iterate through the temperatures, resolve previous days that have found a warmer temperature.
- The index differences yield the number of days waited until a warmer temperature.

Space & Time

Time Complexity: $O(n)$, where n is the number of temperatures, because each index is pushed and popped from the stack at most once.

Space Complexity: $O(n)$ for storing the stack and the answer array.

Solution

The solution uses a monotonic stack to track the indices whose next warmer temperature has not been found yet. As you iterate through the temperature list, compare the current temperature with the temperature at the index stored at the top of the stack. If the current temperature is higher, it means you've found a warmer day for the day represented by the index from the stack. Pop that index, calculate the difference between the current index and the popped index, and store the result. Push the current index onto the stack for future comparisons. This approach efficiently finds the answer without nested iterations.

#962 Maximum Width Ramp [Medium] · Stack

Key Insights

- We need to identify pairs (i, j) where $i < j$ and $\text{nums}[i] \leq \text{nums}[j]$.
- A monotonic decreasing stack can be constructed to maintain candidate starting indices.
- By iterating from left to right, we push indices onto the stack only when they have a smaller value than the last.
- Then, a backward pass from the end of the array allows us to efficiently match each element with the

smallest possible candidate index from the stack.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution is built in two phases:

- Build a decreasing stack of indices by iterating through the array (only push an index if its value is less than the value at the current top of the stack).
- Iterate through the array backwards. For each index j , check and pop elements from the stack while $\text{nums}[j]$ is greater than or equal to the nums value at the popped index. For each valid ramp found, compute the width (j – popped index) and update the maximum width accordingly. This approach guarantees we examine each element a constant number of times ensuring an efficient solution.

#1214 Two Sum BSTs [Medium] · Stack

Key Insights

- Traverse the first tree to collect all node values.
- Traverse the second tree to check for each node if the complement ($\text{target} - \text{current node value}$) exists in the first tree's values.
- Using a hash set for the first tree values allows efficient $O(1)$ lookups.

Round 9 · Low · Hard

p.3315

- The BST property is not essential for the hash set solution, but can be exploited if using iterator techniques for a two-pointer approach.

Space & Time

Time Complexity: $O(n + m)$ where n and m are the number of nodes in the first and second trees respectively.

Space Complexity: $O(n)$ for storing the node values from the first tree (in worst-case scenario).

Solution

We perform a DFS (or any tree traversal) on the first BST to collect all its values in a hash set. Then, we traverse the second BST and for each node, we calculate the complement ($\text{target} - \text{node value}$) and check whether this complement exists in the hash set. If we find such a pair, we return true. If no such pair exists after traversing the second tree completely, we return false. This approach is straightforward and leverages extra space (the set) for constant time lookups, yielding an overall linear time complexity relative to the number of nodes.

#1762 Buildings With an Ocean View [Medium] · Stack

Key Insights

Round 9 · Low · Hard

p.3316

- Processing the buildings from right to left simplifies the problem because the rightmost building always has an ocean view.
- Keep track of the highest building encountered while iterating from right to left.
- A building has an ocean view if its height is greater than the maximum height seen so far.
- Append indices meeting the criteria, then reverse the order of indices for the solution since they were collected in reverse order.

Space & Time

Time Complexity: $O(n)$ where n is the number of buildings.

Space Complexity: $O(n)$ for storing the result in the worst-case scenario.

Solution

Iterate over the "heights" array from right to left while maintaining a variable to store the maximum height encountered so far. For each building, if its height is greater than this maximum height, add its index to the results list and update the maximum height. Since indices are collected in reverse order, reverse the list before returning it. This approach ensures that each building is visited only once.

#215 Kth Largest Element in an Array [Medium] · Heap

Key Insights

- kth largest element can be converted to finding the element at index $n-k$ if the array were sorted.
- A min-heap of size k is a practical method to maintain the k largest elements seen so far.
- The Quickselect algorithm can find the kth largest element in average $O(n)$ time with in-place partitioning.
- Choose an appropriate method based on performance needs and worst-case scenarios.

Space & Time

Time Complexity:

- Min-Heap: $O(n \log k)$ by maintaining a heap of size k .
- Quickselect: Average $O(n)$ but worst-case $O(n^2)$ with poor pivot selection.

Space Complexity:

- Min-Heap: $O(k)$
- Quickselect: $O(1)$ additional space with in-place partitioning.

Solution

We can solve the problem using either a min-heap or the Quickselect algorithm.

- Min-Heap Approach: Initialize a min-heap with the first k elements. Iterate through the remaining elements;

if an element is greater than the heap's root (the smallest in the heap), replace the root. After processing all elements, the root of the min-heap is the kth largest element. – Quickselect Approach: Choose a pivot and partition the array so that elements larger than the pivot come before it. Determine the position of the pivot; if it matches the index corresponding to the kth largest element, return it. Recurse into the appropriate part of the array.

Key points to consider include maintaining heap size for the first approach and proper pivot selection for the Quickselect approach.

#253 Meeting Rooms II [Medium] · Heap

Key Insights

- Sort the intervals based on start times to process meetings in order.
- Use a min-heap (priority queue) to keep track of meeting end times.
- When processing a new meeting, check if its start time is greater than or equal to the smallest end time from the heap.
- If so, it means a meeting room has been freed (the meeting with the smallest end time has ended) and can be reused.
- If not, allocate a new room by pushing the current meeting's end time into the heap.

Round 9 · Low · Hard

p.3319

- The size of the heap at any moment represents the number of rooms currently in use.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting and each heap operation (push/pop) taking $O(\log n)$.

Space Complexity: $O(n)$ in the worst-case scenario when all meetings overlap, requiring all end times in the heap.

Solution

The solution involves first sorting the meeting intervals by their start times. Then, we utilize a min-heap (priority queue) to keep track of the end times of ongoing meetings. For each meeting, we compare its start time with the smallest end time in the heap:

- If the meeting can start after the meeting at the top of the heap has ended (i.e., its start time is greater than or equal to the smallest end), we remove that end time from the heap (freeing up a room). – Then, we add the current meeting's end time to the heap. By the end of processing all meetings, the size of the heap represents the minimum number of conference rooms required.

#347 Top K Frequent Elements [Medium] · Heap

Key Insights

- Count the frequency of each element using a hash table.

Round 9 · Low · Hard

p.3320

- Use bucket sort to achieve linear time complexity by grouping numbers by their frequency.
- Traverse the frequency buckets from highest to lowest to collect the k most frequent elements.
- Alternative approaches include using heaps or Quickselect, but bucket sort meets the follow-up requirement of better than $O(n \log n)$ time complexity.

Space & Time

Time Complexity: $O(n)$ where n is the length of the array (bucket sort and frequency counting are linear).

Space Complexity: $O(n)$ for storing the counts and buckets.

Solution

The solution starts by counting the frequency of each element with a hash table (or dictionary). Then, we create an array of buckets where the index represents the frequency. Each bucket holds the numbers that appear with that frequency. Finally, we iterate over the buckets in reverse order (from high frequency to low) and collect elements until we have k of them. This bucket sort approach avoids the higher complexity of sorting a list of frequencies and guarantees linear time performance.

#505 The Maze II [Medium] · Heap

Key Insights

- The ball rolls until it hits a wall, so each move is not one step but a continuous travel in a direction.
- The problem can be modeled as finding the shortest path in a weighted graph where nodes are positions in the maze and edge weights are the distance traveled in that roll.
- A priority queue (min-heap) based algorithm, such as Dijkstra's algorithm, is ideal since each roll covers variable distances.
- Always check if the new computed distance for a cell is less than a previously recorded one before pushing it to the heap.

Space & Time

Time Complexity: $O(m * n * \log(m * n))$, where m and n are the maze dimensions (each cell may get processed and queued with a logarithmic cost).

Space Complexity: $O(m * n)$ for maintaining the distance matrix and the priority queue.

Solution

We solve the problem using a modified version of Dijkstra's algorithm. The maze grid is treated as a graph where each cell is a node and an edge exists between the current cell and the cell where the ball stops after rolling in one direction. For every cell popped from the priority queue, we roll the ball in all four directions until it hits a wall, counting the number of steps taken. If the distance calculated from this roll is shorter than a

previously stored distance for the stopping cell, we update the distance and add it to the priority queue. This process continues until we either reach the destination (and return the distance) or exhaust all possibilities (returning -1).

#621 Task Scheduler [Medium] · Heap

Key Insights

- Count the frequency of each task.
- The task with the highest frequency determines the base structure.
- Identify how many tasks share the maximum frequency.
- Calculate the total time based on the formula:
 $\text{max}(\text{total tasks}, (\text{max frequency} - 1) * (n + 1) + \text{count of tasks with maximum frequency})$.
- This approach avoids simulating the entire scheduling by using a greedy strategy.

Space & Time

Time Complexity: $O(N)$ where N is the number of tasks (counting frequency, then iterating over constant 26 letters).

Space Complexity: $O(1)$ since the frequency count uses an array or hash map with fixed size (26 letters).

Solution

Round 9 · Low · Hard

p.3323

We first count the frequency of each task. The task(s) with the maximum frequency establish the minimum structure required. Compute the ideal idle slots based on the count of these maximum tasks and the cooling period n . Finally, compare this computed length with the total number of tasks to get the final answer. This greedy strategy avoids explicit simulation and leverages mathematical formulation.

#658 Find K Closest Elements [Medium] · Heap

Key Insights

- The input array is already sorted.
- A binary search can be used to efficiently narrow down the potential starting position.
- A two-pointers or sliding window approach is effective for selecting k consecutive elements.
- The comparison of absolute differences (with a tie-breaker on values) is crucial.
- Maintaining ascending order simplifies the final output.

Space & Time

Time Complexity: $O(\log n + k)$, where $O(\log n)$ for binary search and $O(k)$ for collecting results.

Space Complexity: $O(k)$, for storing the output.

Solution

Round 9 · Low · Hard

p.3324

We can solve this problem by first using binary search to identify the left-most window where the k closest elements might start. The binary search operates over the possible starting indices for windows of size k (from 0 to $n-k$). For each middle index mid , we compare the distances between x and $arr[mid]$ and $arr[mid + k]$ to decide whether to move the window to the right or left. Once the correct starting position is found, return the k elements starting from that index. This approach leverages the sorted property of the array to efficiently converge on the optimal window.

#787 Cheapest Flights Within K Stops [Medium] · Heap

Key Insights

- Build the graph as an adjacency list of flights and prices.
- Use a min-heap so the cheapest partial route is always processed first.
- Track states as $(city, flightsUsed)$ instead of only $city$, because the stop count matters.
- For each popped state, try all outgoing flights and push the updated cost and flight count.
- Do not expand a state once it has already used more than $k + 1$ flights.
- Keep the best known cost for each $(city, flightsUsed)$ state so worse duplicate states can be

skipped.

Space & Time

Time Complexity: $O((V + E) * \log(V * (k + 1)))$ for the heap-based search over $(city, stopsUsed)$ states.

Space Complexity: $O(E + V * (k + 1))$ for the graph plus the best-cost tracking table.

Solution

We solve the problem using a Dijkstra-style min-heap search, but the state must include both the city and how many flights have been used so far. We start from $(0, src, 0)$ and always pop the cheapest state from the heap. For each state, if we reach dst , that cost is the answer. Otherwise, if we still have flights available within the $k + 1$ limit, we explore all outgoing edges and push the next states into the heap. To avoid useless work, we store the best cost seen for each $(city, flightsUsed)$ pair and only continue when the new state improves that record. This guarantees the cheapest valid route under the stop constraint.

#973 K Closest Points to Origin [Medium] · Heap

Key Insights

- Calculate the squared Euclidean distance $(x^2 + y^2)$ to avoid the overhead of square root computation.
- Sorting the points based on their distance gives an $O(n \log n)$ solution.

- Alternatively, using a max-heap to keep track of the k closest points results in $O(n \log k)$ time, which is effective if k is much smaller than n .
- Using Quickselect can achieve an average time complexity of $O(n)$, though the worst-case is $O(n^2)$.
- The problem does not require the output to be sorted in any specific order.

Space & Time

Time Complexity: $O(n \log n)$ with a sorting approach, $O(n \log k)$ with a heap approach, or average $O(n)$ with a Quickselect approach.

Space Complexity: $O(n)$ if extra space is used for sorting, or $O(k)$ for the heap approach.

Solution

We can solve this problem by first computing the squared Euclidean distance for each point to avoid unnecessary square root operations. Then, we can sort the points based on their computed distance and select the first k points. Alternatively, a max-heap or Quickselect algorithm can be used:

- Max-Heap: Maintain a heap of size k while iterating through the points, ensuring that only the k points with the smallest distances are kept.
- Quickselect: Partition the array like in QuickSort to position the k th closest point in its correct position. Once partitioned, all points to the left of that position will be the k closest. The sorting solution is straightforward and efficient enough

given the constraints.

#1057 Campus Bikes [Medium] · Heap

Key Insights

- Compute the Manhattan distance between every worker and bike pair.
- To achieve the tie-breaking rules, sort all pairs by distance, then by worker index and bike index.
- Once a worker is assigned a bike, skip any other pairs involving that worker or the already assigned bike.
- The problem can also be approached using bucket sort because the Manhattan distance range is limited, but sorting is intuitive given the constraints.

Space & Time

Time Complexity: $O(n * m * \log(n * m))$ where n is the number of workers and m is the number of bikes.

Space Complexity: $O(n * m)$ for storing all worker-bike pairs.

Solution

Create all possible (worker, bike) pairs along with their Manhattan distances. Sort this list of pairs by:

- Manhattan distance,
- Worker index (if distances are equal),
- Bike index (if both distance and worker index are equal).

Then, iterate over the sorted list and assign bikes to workers if neither the worker nor the bike has already

been assigned. This guarantees that at each step, the optimal available pairing (according to the problem's requirements) is made. This method naturally enforces the tie-breaking rules specified in the problem.

#1167 Minimum Cost to Connect Sticks [Medium] · Heap

Key Insights

- Always connect the two smallest sticks available to minimize the incremental cost.
- A min-heap (or priority queue) is an ideal data structure to efficiently retrieve the smallest sticks.
- The process continues until only one stick remains.

Space & Time

Time Complexity: $O(n \log n)$, where n is the number of sticks (due to heap insertion and removal operations).

Space Complexity: $O(n)$, for storing the sticks in the heap.

Solution

The solution uses a greedy algorithm with a min-heap:

- Insert all stick lengths into a min-heap.
- While there are at least two sticks in the heap: Extract the two smallest sticks. Compute the cost of connecting them (sum the two lengths). Add this cost to the total result. Push the combined stick (with length equal to their sum) back into the heap.
- Once the heap is reduced to one

element, return the total cost.

This approach ensures that at every connection, the smallest possible sticks are merged, which leads to the minimum overall cost.

#1424 Diagonal Traverse II [Medium] · Heap

Key Insights

- The diagonal key for each element can be defined as $(\text{row_index} + \text{column_index})$.
- Elements in the same diagonal (with equal diagonal key) must be output in an order such that the element from the lower row comes before the one from a higher row.
- Because the input grid may be jagged (rows with varying lengths), you must carefully iterate only over valid indices.
- A hash table (or map) can be used to collect elements grouped by their diagonal key.
- Inserting each element at the beginning of the list for that diagonal builds the list in the desired order without needing an extra reverse step later.
- Finally, traverse the keys in increasing order (from 0 to maximum key) to produce the final result.

Space & Time

Time Complexity: $O(N)$, where N is the total number of elements (each element is processed once).

Space Complexity: $O(N)$, for storing the elements in the map and the output array.

Solution

The solution iterates over every element in `nums` while computing its diagonal key as $(i + j)$. For each element, the value is inserted at the beginning of the list corresponding to that diagonal key. This ensures that when the diagonal groups are later concatenated in order of increasing key, the internal ordering within each diagonal is correct (i.e., elements appear from the bottom of the diagonal to the top). The final result is built by iterating over the keys in sorted order and appending each list of diagonal elements.

#139 Word Break [Medium] · Trie

Key Insights

- Use Dynamic Programming (DP) to break the problem into subproblems.
- Create a DP array where `dp[i]` is true if the substring `s[0:i]` can be segmented using words from the dictionary.
- Leverage a hash set for fast look-up of dictionary words.
- For every index `i`, check all possible previous break points `j` such that `dp[j]` is true and `s[j:i]` is in the `wordDict`.

- The problem can also be approached with recursion and memoization, but DP is more straightforward for iterative implementation.

Space & Time

Time Complexity: $O(n^2)$ in the worst-case scenario, where `n` is the length of `s` (each substring check is $O(1)$ if using a hash set, and there could be $O(n^2)$ checks overall).

Space Complexity: $O(n)$ for the dp array and $O(k)$ for the hash set where `k` is the number of words in the dictionary.

Solution

The solution uses a Dynamic Programming (DP) approach:

- Convert `wordDict` to a hash set for constant time lookup.
- Create a boolean dp array of size `n+1` (`n` is the length of `s`), where `dp[i]` indicates if `s[0:i]` can be segmented.
- Set `dp[0]` to true because an empty string is always segmentable.
- For each index `i` from 1 to `n`, iterate through all indices `j` from 0 to `i`: If `dp[j]` is true and the substring `s[j:i]` is in the dictionary, then set `dp[i]` to true. Break early from the inner loop when a valid break is found.
- Return `dp[n]` as the result.

This method efficiently builds up the solution by utilizing previous results to determine if the entire string can be segmented.

#208 Implement Trie (Prefix Tree) [Medium] · Trie

Key Insights

- Use a tree-like structure where each node represents a character.
- Each node stores its children in a dictionary (or hash map) for $O(1)$ average time access.
- Mark the end of a word with a boolean flag.
- Insertion is done character by character, creating new nodes as needed.
- Searching and checking prefixes follow traversals down the tree.

Space & Time

Time Complexity:

- Insert: $O(m)$ where m is the length of the word.
- Search: $O(m)$
- `startsWith`: $O(m)$

Space Complexity:

- $O(N * M)$ in the worst-case scenario, where N is the number of words and M is the average length of the words, due to storing each character in the Trie.

Solution

We implement the Trie using a node-based approach. Each `TrieNode` contains a hash map (or dictionary) of its children and a boolean indicating if the node marks the end of a word. The Trie class uses these nodes to build

out the words character by character. The insert method traverses and builds the necessary nodes; the search method verifies that all characters exist in order and that the final node marks the end of a word; the `startsWith` method checks if the prefix exists in the Trie. No tricky edge cases exist beyond handling null or empty inputs, as constraints guarantee valid lowercase strings.

#211 Design Add and Search Words Data Structure [Medium] · Trie

Key Insights

- Utilize a Trie (prefix tree) to efficiently store words.
- Each node in the Trie contains a dictionary for its children and a flag indicating the end of a word.
- The search operation is implemented with a depth-first search (DFS) approach.
- When the search query contains a dot, explore all possible child nodes.

Space & Time

Time Complexity:

- `addWord`: $O(w)$, where w is the length of the word.
- `search`: Worst-case $O(26^w)$ for words with multiple wildcards, but practical performance is better given constraints.

Space Complexity: $O(n * w)$, where n is the number of words and w is the average word length.

Solution

The WordDictionary is implemented using a Trie. For adding, we insert each character of the word into the Trie. For searching, we use DFS. When encountering a wildcard '.', we recursively search all children nodes. This structure combined with DFS efficiently addresses the wildcard matching and ensures good performance for multiple operations.

#616 Add Bold Tag in String [Medium] · Trie

Key Insights

- We need to identify every occurrence of each word in s .
- Overlapping or consecutive intervals should be merged to form one continuous bold region.
- Use an auxiliary data structure (e.g., a boolean array) to mark positions that should be bolded.
- Once the positions are marked, build the final string by inserting bold tags at the appropriate boundaries.

Space & Time

Time Complexity: $O(n * m * k)$ where n is the length of s , m is the number of words, and k is the average length of the words (due to substring matching).

Space Complexity: $O(n)$ for the auxiliary boolean array used to mark bold positions.

Solution

The solution involves two main steps:

– Identify and mark the indices in the string s that are part of any matching word. We do this by iterating over each word and marking every occurrence's start and end in a boolean array. – Construct the final string by traversing the boolean array. When encountering a segment of consecutive true values, wrap that segment with and tags.

The key idea is to merge intervals by marking every index that should be bold and then using a single pass to insert tags only when transitioning from non-bold to bold and vice versa.

#692 Top K Frequent Words [Medium] · Trie

Key Insights

- Count the frequency of each unique word using a hash table or dictionary.
- Use a min-heap (priority queue) to keep track of the top k frequent words. In the heap, the ordering is based on frequency; if frequencies are equal, order by lexicographical order (inverted condition for min-heap).
- Alternatively, you can sort the unique words based on frequency and lexicographical order and then pick the top k .

- For an efficient solution, use a min-heap that maintains size k achieving $O(n \log k)$ time complexity.

Space & Time

Time Complexity: $O(n \log k)$ when using a heap, where n is the number of unique words.

Space Complexity: $O(n)$ for storing the frequency counts and the heap.

Solution

We first calculate the frequency of each word using a hash table (or dictionary). Then, we use a min-heap that stores pairs of (frequency, word). The comparator for the heap is defined such that the root of the heap is the word with the smallest frequency; if two words have the same frequency, the word with the lexicographically larger value is considered “smaller” in our ordering so that it gets popped when the heap size exceeds k . After processing all words, the heap contains the top k frequent words. Finally, the heap elements are extracted and reversed (if needed) to produce the final list sorted from highest to lowest frequency with lexicographic tie-breaking.

#720 Longest Word in Dictionary [Medium] · Trie

Key Insights

- The word must be built character by character; hence every prefix (removing one character from the end at a

time) must exist in the dictionary.

- Sorting the input list can help in ensuring that smaller words (prefixes) are processed first.
- Using a hash set to store valid words enables rapid prefix checks.
- Lexicographical order plays a role when words have the same length.

Space & Time

Time Complexity: $O(n * L + n \log n)$ where n is the number of words and L is the maximum length of a word ($n * L$ for prefix checking, and $n \log n$ for sorting).

Space Complexity: $O(n * L)$ for storing words in the set.

Solution

The approach begins by sorting the list of words. Sorting is done primarily by word length and secondarily by lexicographical order so that when we iterate through the words, all possible prefixes for a word have already been considered. Initialize a set to keep track of valid words that can be built one character at a time. For each word in the sorted list, check if all its prefixes (word without the last character progressively) exist in the set. If a word qualifies, add it to the set and check if it should be the new result based on length and lexicographical order. The set lookup ensures that checking each prefix is efficient.

#1166 Design File System [Medium] · Trie

Key Insights

- Use a data structure (like a hash map) to store paths and associated values.
- To create a new path, check that the path does not already exist.
- Ensure the immediate parent path exists before creation.
- When retrieving a value, simply look it up in the storage structure.
- Splitting the path by "/" allows easy identification of the parent.

Space & Time

Time Complexity:

- createPath: $O(n)$ per call where n is the number of segments in the path (to process and validate parent path).
- get: $O(1)$ per call if using a hash map.

Space Complexity: $O(m * L)$ where m is the number of paths stored and L is the average length of a path.

Solution

The solution uses a hash map (dictionary) to store file paths along with their associated values. For the createPath method, the algorithm:

- Checks if the path already exists.
- Extracts the parent path by removing the last segment.
- Validates the existence of the parent path.
- Inserts the new path with

its corresponding value if validation passes. For the get method, the algorithm simply returns the stored value if the path exists; otherwise, it returns -1 . This approach ensures efficient lookup and validation using string manipulation and hash map operations.

#17 Letter Combinations of a Phone Number

[Medium] · Backtracking

Key Insights

- Use a mapping from digits to the respective letters as defined by a telephone keypad.
- Backtracking is an ideal approach to generate all possible combinations.
- Each digit adds a level in the recursive tree where you append one letter at a time.
- Handling edge cases such as an empty input string is crucial.

Space & Time

Time Complexity: $O(4^n * n)$ where n is the length of the input string (each digit may map to up to 4 letters and adding a combination takes $O(n)$ time).

Space Complexity: $O(n)$ for the recursion call stack, not counting the space required for storing the output.

Solution

We can solve the problem using a backtracking algorithm. First, we define a dictionary mapping each

digit to its corresponding letters. Then we initiate a recursive function called backtrack that takes the current combination and the index of the next digit to process. At each call, if the combination is complete (length equal to the input string), we add it to the result list. Otherwise, we iterate through the letters corresponding to the current digit, append a letter to the combination, and recursively call backtrack with the next index. This approach ensures that all possible letter combinations are generated.

#22 Generate Parentheses [Medium] · Backtracking

Key Insights

- Use a backtracking approach to build the valid combinations step by step.
- At any recursion step, you can add an opening parenthesis if you still have one available.
- You can add a closing parenthesis only if it would not result in an invalid sequence (i.e., the number of closing parentheses is less than the number of opening ones).
- The solution leverages recursion to explore all potential valid sequences while pruning invalid paths early.

Space & Time

Time Complexity: $O(4^n / (n^{(3/2)}))$

Space Complexity: $O(n)$ for the recursion stack, plus space for storing all valid combinations.

Solution

The solution uses a recursive backtracking technique to generate the combinations. The algorithm maintains counters for the number of '(' and ')' added to the current string. It recursively adds a '(' if the count is less than n and adds a ')' if doing so does not make the string invalid (i.e., if the count of ')' is less than the count of '('). When a valid combination reaches a length of $2*n$, it is added to the result list.

#39 Combination Sum [Medium] · Backtracking

Key Insights

- Backtracking is a natural fit to explore all potential combinations.
- Sorting candidates simplifies the process by allowing early termination if the current candidate exceeds the remaining target.
- Recursion is used to build combinations and backtrack once the current path exceeds or meets the target.
- The same candidate can be reused in the combination, so the recursive call does not advance the index.

Space & Time

Time Complexity: $O(2^{(target/\min(candidates))})$ in the worst case where many combinations are possible.

Space Complexity: $O(target/\min(candidates))$ for the recursion stack and temporary space used to store

combinations.

Solution

The solution uses a backtracking approach. First, sort the candidates to enable pruning of recursive calls once the candidate exceeds the remaining target. A recursive helper function builds up a combination, subtracting the candidate value from the target until it reaches zero (a valid combination) or becomes negative (an invalid path). The recursion allows the same candidate to be chosen again. Each valid combination is then added to the result list.

#46 Permutations [Medium] · Backtracking

Key Insights

- Use backtracking to generate all possible permutations.
- Build each permutation by selecting numbers one at a time.
- When the current permutation reaches the length of nums, add it to the result.
- Backtracking allows you to undo choices and explore different orderings.

Space & Time

Time Complexity: $O(n * n!)$ – There are $n!$ permutations, and constructing each permutation takes $O(n)$ time.

Space Complexity: $O(n)$ – Recursion depth is limited to n , plus additional space for the current permutation.

Solution

The solution employs a backtracking algorithm. A helper function is used to build permutations recursively. At each step, the algorithm iterates through the available numbers, adds one to the current permutation if it has not been used yet, and recurses to handle the remaining numbers. Once the current permutation matches the length of the input array, it is added to the result list. Then, by "backtracking" (removing the last number added), the algorithm explores alternative possibilities. This systematic exploration ensures that all possible permutations are generated.

#79 Word Search [Medium] · Backtracking

Key Insights

- Use Depth-First Search (DFS) to explore each cell in the board as a potential starting point.
- Employ backtracking to mark cells as visited when exploring a path and revert them back when backtracking.
- Check boundary conditions and ensure the current cell's character matches the corresponding character in the word.
- Given the small grid and word sizes, a recursive solution is both feasible and straightforward.

- Search pruning can be applied by stopping a DFS branch as soon as the current path does not match the prefix of the target word.

Space & Time

Time Complexity: $O(m * n * 3^L)$, where L is the length of the word (each DFS may branch to at most 3 further cells, excluding the coming cell).

Space Complexity: $O(L)$ due to recursion call stack (where L is the word length).

Solution

The solution uses DFS with backtracking to traverse the grid. For each cell, we:

- Check if the cell matches the first character of the word.
 - Recursively explore the adjacent cells (up, down, left, right) while marking the current cell as visited.
 - If the path does not lead to a complete match, backtrack by resetting the visited cell.
 - The DFS stops if all characters in the word have been found in sequence.
- The algorithm leverages in-place modifications (or an auxiliary visited array) to avoid revisiting cells while exploring a candidate path.

#113 Path Sum II [Medium] · Backtracking

Key Insights

- Use Depth-First Search (DFS) to explore all paths from the root to the leaves.

Round 9 · Low · Hard

p.3345

- Backtracking is essential: add the current node value to the path before the recursive calls and remove it (backtrack) after exploring both subtrees.
- At each leaf, check if the accumulated sum equals targetSum.
- Maintain a running sum or subtract the current node's value from the target sum as you traverse.

Space & Time

Time Complexity: $O(N^2)$ in the worst-case scenario (where N is the number of nodes) when most nodes are leaf nodes, since each potential path copy operation can be $O(N)$.

Space Complexity: $O(N)$ for the recursion stack and the path storage, where N is the height of the tree.

Solution

We solve the problem using a recursive DFS approach combined with backtracking. The idea is to traverse the tree while maintaining the current path and the remaining sum needed to reach targetSum. When a leaf node is reached (node with no children), we check if the remaining sum equals the leaf's value – if yes, we record a copy of the current path.

Data structures and algorithmic approaches used:

- Recursion for DFS: to traverse to each leaf.
- List (or Array) for storing the current path.
- Backtracking: to remove the last added element when stepping back up the recursion stack.
- Copying the path when a valid leaf

Round 9 · Low · Hard

p.3346

is reached to store it in the final result.

A common pitfall is to forget to backtrack, which would result in incorrect paths being recorded.

#134 Gas Station [Medium] · Greedy

Key Insights

- If the total amount of gas available is less than the total travel cost, then completing the circuit is impossible.
- A greedy approach can be used by keeping track of the current surplus of gas. If the surplus becomes negative at any station, the journey cannot start from any station before that point.
- Restart the journey from the next station whenever the surplus goes negative.
- The problem guarantees that if a solution exists, it is unique.

Space & Time

Time Complexity: $O(n)$, where n is the number of gas stations, as we traverse the list only once.

Space Complexity: $O(1)$, as no extra space is used other than a few variables.

Solution

The solution involves iterating over the list of stations while maintaining two variables: one for the current gas surplus and one for the total gas surplus. If at any

station the current surplus drops below zero, the current starting station is not viable and we set the starting station to the next index and reset the current surplus. After one pass, if the total gas surplus is non-negative, then the identified station is the valid starting point.

#179 Largest Number [Medium] · Greedy

Key Insights

- The order in which numbers appear drastically affects the concatenated result.
- A custom sorting strategy is required where two numbers are compared by checking which concatenation (first-second vs second-first) yields a larger number.
- Edge case: When the numbers are all zeros, the result should be "0" instead of several leading zeros.

Space & Time

Time Complexity: $O(n \log n * k)$, where n is the number of elements in the array and k is the maximum number of digits in a number (due to string comparisons in the custom sort).

Space Complexity: $O(n * k)$ for storing and processing the strings.

Solution

We solve the problem by converting all integers to strings so that we can use string concatenation to

compare the results. A custom comparator is defined that, given two number strings, compares the concatenated results in both orders. By sorting the array using this comparator in descending order, we ensure that placing the highest contributing number first yields the largest possible number. After sorting, a check is performed to return "0" if the highest number is "0". Finally, the sorted strings are concatenated to form the answer.

#280 Wiggle Sort [Medium] · Greedy

Key Insights

- The alternating "wiggle" requirement can be satisfied by ensuring two conditions:
- For every odd index i , $\text{nums}[i]$ should be greater than or equal to $\text{nums}[i-1]$.
- For every even index i (where $i > 0$), $\text{nums}[i]$ should be less than or equal to $\text{nums}[i-1]$.
- A single pass through the array can ensure these conditions by swapping adjacent elements when the condition is violated.
- This greedy approach works because the local adjustments guarantee the overall wiggle pattern without needing to sort the entire array.

Space & Time

Time Complexity: $O(n)$ — Only one pass through the array is needed.

Space Complexity: $O(1)$ — The modifications are done in-place without extra data structures.

Solution

We can achieve a wiggle sort in one pass by iterating through the array and, at each index i (starting from 1), checking if the pair $(\text{nums}[i-1], \text{nums}[i])$ satisfies the wiggle condition. If i is odd, $\text{nums}[i]$ should be at least $\text{nums}[i-1]$; if not, we swap them. Conversely, if i is even and $\text{nums}[i]$ is greater than $\text{nums}[i-1]$, then we swap them. These local swaps adjust the order incrementally while ensuring that the overall wiggle pattern is maintained. This greedy approach works because only two adjacent numbers are involved in any violation and fixing them locally does not disturb the rest of the order.

#954 Array of Doubled Pairs [Medium] · Greedy

Key Insights

- Sort the array based on the absolute value of the elements to handle negative numbers correctly.
- Use a hash table (or frequency counter) to track the occurrence of each element.
- For each number in the sorted list, if the number is still available in the frequency map, try to pair it with its double.
- Decrement the appropriate counts and validate that it is possible to find sufficient doubles for all numbers.

- Special case: when dealing with zeros, pairing works within the same value.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting.

Space Complexity: $O(n)$ for storing the frequency map.

Solution

The solution starts by counting the frequency of each element in the array. Sorting the array by absolute value ensures that when we process an element, any potential pair (i.e., its double) comes later in the sequence. For each number, if there are available occurrences, we try to match it with its double. If there are not enough doubles available, the answer is false. If all elements can be paired appropriately by decrementing their counts in the frequency map, the function returns true.

#2131 Longest Palindrome by Concatenating Two Letter Words [Medium] · Greedy

Key Insights

- Pair words with their reverse (e.g., "ab" with "ba") to form a palindrome.
- Handle words that are palindromic by themselves (like "gg" or "aa") differently; they can be used in pairs and one extra can be placed in the center.
- Use a hash map (or dictionary) to count frequencies of each word.

- For non-palindromic word pairs, consider each pair only once to avoid double counting.
- For self-palindromic words, count how many pairs can be used and whether an extra word can be placed in the middle to maximize length.

Space & Time

Time Complexity: $O(n)$, where n is the number of words, since we iterate through the list and perform constant time operations per word.

Space Complexity: $O(n)$, to store counts of words in the hash map.

Solution

The approach is to count the occurrences of each word using a hash map. For each word:

- If the word is not self-palindromic, check if its reverse exists and form pairs by considering the minimum count between the word and its reverse.
- If the word is self-palindromic, form as many pairs as possible (using $\text{count} // 2$) and note if there is any word left; you can use one extra self-palindromic word as the center of the palindrome. Finally, compute the total length by multiplying the number of words used by 2, since each word contributes 2 characters.

R6 · Mid · Hard (28 q)

>> Round 6 · Mid Priority · Hard — Quick Review

28 questions · key insights · complexity · solution

#25 Reverse Nodes in k-Group [Hard] · Linked List

Key Insights

- Use a dummy node to simplify edge cases.
- Identify groups of k nodes using a helper function.
- Reverse the nodes in each group using pointer manipulation.
- If a remaining group has fewer than k nodes, leave it unchanged.
- The reversal is done in-place with $O(1)$ extra space.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes (each node is processed a constant number of times)

Space Complexity: $O(1)$ (only constant extra space is used)

Solution

The solution involves iterating through the linked list by identifying groups of k consecutive nodes. For each group:

- Find the kth node from the current starting point. If it does not exist, the remainder of the list is not reversed.
- Reverse the nodes in the identified group in-place by adjusting the next pointers.
- Reconnect the reversed group with the previous part of the list and update pointers to move to the next group. The approach leverages a dummy node for easier handling of head modifications and uses a helper function to fetch the kth node. The in-place reversal ensures that the extra memory usage is constant.

#432 All O'one Data Structure [Hard] · Linked List

Key Insights

- Use a Doubly-Linked List (DLL) where each node holds a unique count value and a set of keys with that count.
- Maintain a Hash Map (keyToNode) to quickly locate the node corresponding to each key.
- Insert new nodes in the DLL when a key's count is incremented or decremented and the adjacent node does not have the target count.
- Remove nodes from the DLL when no keys remain in that count.
- The head of the DLL holds the minimum count and the tail holds the maximum count.

Space & Time

Time Complexity: $O(1)$ average for inc, dec, getMaxKey, and getMinKey.

Space Complexity: $O(n)$ where n is the number of unique keys in the data structure.

Solution

The solution uses a combination of a doubly-linked list and hash maps. Each node in the DLL represents a frequency and maintains a set of keys that have that frequency. A hash map maps keys to their corresponding node. When `inc(key)` or `dec(key)` is called, we update the key's frequency and move the key to the appropriate node in the DLL. If the adjacent node with the required frequency doesn't exist, we create a new node and insert it in the correct position. Removal of a key from a node and deletion of empty nodes ensures that `getMinKey()` and `getMaxKey()` can be returned directly from the head and tail of the DLL.

#460 LFU Cache [Hard] · Linked List

Key Insights

- Maintain a mapping from key to its value and frequency.
- Use an additional mapping from frequency to keys (maintaining the order of insertion or recent use) for quick eviction.
- Keep a global `minFreq` variable to track the lowest frequency present in the cache.
- When a key is accessed or updated, increase its frequency and relocate it in the frequency mapping.

Round 9 · Low · Hard

p.3355

- Evictions occur by removing the oldest key from the lowest frequency bucket.

Space & Time

Time Complexity: $O(1)$ average for both get and put.

Space Complexity: $O(\text{capacity})$, where capacity is the maximum number of keys stored.

Solution

The solution utilizes two main data structures:

- A hash map (`keyToValFreq`) that maps keys to a pair (value, frequency).
- A second hash map (`freqToKeys`) mapping frequencies to an ordered list/set of keys. This data structure preserves the order in which keys were added to allow removal of the least recently used key when multiple keys share the same frequency.

A global `minFreq` variable is maintained to quickly identify the bucket with the least frequency. When a key is accessed or updated, the key is removed from its current frequency list and placed in the next higher frequency list. If the frequency list for the minimal frequency becomes empty, `minFreq` is updated.

This design guarantees that both get and put operations execute in constant time on average.

#32 Longest Valid Parentheses [Hard] · Stack

Key Insights

Round 9 · Low · Hard

p.3356

- Use a stack to keep track of indices where potential valid substrings begin.
- Start by pushing a dummy index (-1) to handle edge cases.
- For every '(', push its index onto the stack.
- For every ')', pop from the stack. If the stack becomes empty, push the current index as a new base. Otherwise, compute the length of the valid substring using the current index minus the top index of the stack.
- An alternative is the dynamic programming approach where $dp[i]$ represents the length of the longest valid substring ending at index i .

Space & Time

Time Complexity: $O(n)$ where n is the length of the string.

Space Complexity: $O(n)$ for the stack data structure.

Solution

The stack-based solution works by maintaining indices for unmatched parentheses. Initially, a dummy index (-1) is pushed onto the stack to serve as the base for calculating subsequent valid substring lengths. As we iterate through the string:

- If the character is '(', we push its index onto the stack.
- If the character is ')', we pop from the stack. If the stack is empty afterwards, that means there is no valid starting position, so we push the current index to reset

the base. Otherwise, we calculate the length of the current valid substring by subtracting the current index with the new top index of the stack and update the maximum length accordingly.

#42 Trapping Rain Water [Hard] · Stack

Key Insights

- The water trapped at any index is determined by the minimum of the maximum height on its left and right minus the height at that index.
- A two-pointers approach can efficiently solve the problem in one pass by maintaining left and right boundaries.
- Alternative approaches include dynamic programming with pre-computed left max and right max arrays, or using a monotonic stack to determine boundaries.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$ when using the two-pointer approach

Solution

We use a two-pointers strategy by initializing pointers at the beginning and the end of the array. Keep track of the maximum height seen so far from the left ($left_max$) and the right ($right_max$). At each step, move the pointer with the lesser height because the water trapped is

determined by the shorter boundary. Compute the water at that position as the difference between the boundary (left_max or right_max) and the current height, accumulating it into the total. This approach ensures linear time scanning with constant extra space.

#84 Largest Rectangle in Histogram [Hard] · Stack

Key Insights

- Each bar can potentially extend left and right as long as the neighboring bars are not shorter.
- A monotonic (increasing) stack can be used to efficiently track indices of bars.
- When encountering a bar lower than the one on the top of the stack, pop the stack and calculate the area of the rectangle formed with the popped bar height.
- A dummy bar (of height 0) is appended at the end to ensure all bars are processed.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

The solution employs a monotonic stack to traverse the histogram once. For each bar, while the current bar's height is less than the height at the top index of the stack, it pops from the stack and calculates the area using the popped height as the smallest bar. The width

for the rectangle is determined by the difference between the current index and the index from the stack (or the current index if the stack is empty). This effectively finds the maximum rectangular area possible. A sentinel value (0) is added at the end of the histogram to flush any remaining bars from the stack.

#224 Basic Calculator [Hard] · Stack

Key Insights

- Use a stack to save previous results and signs when encountering parentheses.
- Process the string left-to-right by accumulating multi-digit numbers and applying the current sign.
- When encountering a '(', push the current result and sign onto the stack and reset them.
- When encountering a ')', complete the current sub-expression then pop the previous state to incorporate the computed value.
- Handle whitespaces by simply skipping them.

Space & Time

Time Complexity: $O(n)$ where n is the length of the input string, since we process each character once.

Space Complexity: $O(n)$ in the worst-case scenario due to the use of a stack for nested parentheses.

Solution

The solution uses an iterative approach with a stack:

– Initialize variables: result to keep the running total, sign to handle addition/subtraction, and index to iterate over the string. – As you iterate, build multi-digit numbers by checking for consecutive numerical characters. – Adjust the total by adding the product of the current number and its sign. – When encountering a '(', push the current result and sign onto the stack, then reset results and sign for the new sub-expression. – When encountering a ')', complete the sub-expression by popping the previous result and sign, then combine. – Skip over spaces. – The final result is the computed value after the end of the string.

This approach effectively simulates recursion and correctly handles nested expressions and negation.

#239 Sliding Window Maximum [Hard] · Stack

Key Insights

- Use a deque (double-ended queue) to maintain the indices of potential maximum elements in a decreasing order.
- Ensure that the deque always has the current window's indices by removing elements that fall outside the current window.
- Only add elements that are greater than the ones at the back of the deque, thus preserving a monotonic decreasing order.

Round 9 · Low · Hard

p.3361

- The maximum for the current window is at the front of the deque.

Space & Time

Time Complexity: $O(n)$ where n is the number of elements in nums, since each element is processed at most twice.

Space Complexity: $O(k)$, as the deque stores indices for at most k elements at any time.

Solution

The solution uses a monotonic deque to maintain indices of potential maximum elements in the current window. As the window slides, we:

- Remove indices from the front if they are out of the window.
 - Remove indices from the back if the current element is larger, because they cannot be the maximum if the current element is in the window.
 - Append the current index.
 - After the first k elements, the element at the front of the deque is the maximum for that window.
- This approach ensures that each element is pushed and popped from the deque only once, achieving an overall $O(n)$ time complexity.

#716 Max Stack [Hard] · Stack

Key Insights

- Use a doubly-linked list to support fast insertion and removal from the stack.

Round 9 · Low · Hard

p.3362

- Maintain a balanced tree (or ordered dictionary/multiset) to quickly locate the maximum element in $O(\log n)$ time.
- Map each value to its corresponding node(s) in the doubly-linked list, so that when `popMax` is called, you can immediately remove the corresponding node from the list.
- The double-linking allows removal of arbitrary nodes in $O(1)$, once you have the reference.

Space & Time

Time Complexity:

- push: $O(\log n)$ due to insertion in the ordered structure.
- pop: $O(1)$ for removal from the doubly-linked list, plus $O(\log n)$ for updating the tree.
- top: $O(1)$
- peekMax: $O(1)$ or $O(\log n)$ depending on the tree implementation.
- popMax: $O(\log n)$ for looking up and removal from the tree, $O(1)$ for doubly-linked list deletion.

Space Complexity:

- $O(n)$ for storing all nodes in the doubly-linked list and the tree structure mapping values to nodes.

Solution

We use two data structures:

- A doubly-linked list to store the elements in stack order. This supports fast push and pop operations and allows $O(1)$ removal when given a reference.
- An ordered structure (like a balanced BST, `TreeMap` in Java, or `SortedList` in Python) that maps each value to a list of nodes that hold that value in the linked list. This allows us to quickly determine the maximum element. When there are duplicates, we remove the one closest to the top by storing nodes in order. The trick is to keep these two structures synchronized. When an element is pushed, add its node to both the doubly-linked list and the ordered structure. When an element is removed (either via `pop` or `popMax`), remove its node from both structures.

#772 Basic Calculator III [Hard] · Stack

Key Insights

- Use recursion to handle nested expressions defined by parentheses.
- A stack can be used to correctly evaluate the expression and respect operator precedence.
- Multiplication and division have higher precedence than addition and subtraction.
- When encountering an open parenthesis, recursively compute its value until a closing parenthesis is found.
- Use immediate calculation for multiplication and division by updating the top of the stack before moving

to the next operator.

Space & Time

Time Complexity: $O(n)$ where n is the length of the string, as every character is processed.

Space Complexity: $O(n)$ due to the recursion stack or auxiliary stack used for intermediate results.

Solution

We use a recursive approach combined with a stack to solve the problem. As we iterate through the string:

- Convert digits into numbers.
- When encountering an operator or parenthesis, process the previously accumulated number based on the preceding operator.
- For multiplication or division, pop the last number from the stack, apply the operation, and push the result back.
- When a '(' is encountered, recursively evaluate the substring until the corresponding ')'. – After processing all tokens, sum the values in the stack to get the final result. This method correctly handles both operator precedence and nested expressions.

#895 Maximum Frequency Stack **[Hard]** · Stack

Key Insights

- Use a hash map to count the frequency of each element.
- Maintain another hash map that maps frequencies to stacks (lists) of elements.

- Track the current maximum frequency to enable $O(1)$ retrieval of the most frequent element.
- When pushing an element, update its frequency and add it to the corresponding frequency stack.
- When popping an element, remove it from the stack corresponding to the maximum frequency and update the frequency count; if that stack becomes empty, decrease the current maximum frequency.

Space & Time

Time Complexity: $O(1)$ per push and pop.

Space Complexity: $O(n)$, where n is the number of elements pushed onto the stack.

Solution

The solution involves using two main data structures:

- A hash map (freq) to store the frequency count for each value.
- A hash map (group) that maps a frequency to a stack (list) of values that have that frequency.

During a push:

- Increment the frequency count of the value.
- Push the value onto the stack corresponding to its updated frequency.
- Update the maximum frequency if necessary.

During a pop:

- Pop the top element from the stack corresponding to the current maximum frequency.
- Decrement its frequency count in the freq map.
- If the frequency stack

becomes empty after popping, decrement the current maximum frequency. This design enables the retrieval of the most frequent element in constant time while dealing with frequency ties by returning the element closest to the top.

#23 Merge k Sorted Lists [Hard] · Heap

Key Insights

- Use a min-heap (priority queue) to efficiently extract the smallest value among the current heads of the lists.
- By pushing the head of each non-empty list into the heap, we can pull the smallest node, then push that node's next element into the heap.
- Alternatively, a divide and conquer approach can be used by recursively merging pairs of lists.
- Careful management of pointers (or references) is key to constructing the new linked list.

Space & Time

Time Complexity: $O(N \log k)$ where N is the total number of nodes across all lists and k is the number of lists.

Space Complexity: $O(k)$ due to the heap containing at most one node from each list.

Solution

We use a min-heap to always retrieve the smallest node among the current nodes in the linked lists. First, initialize the heap with the head of each list. Then,

continuously extract the smallest node from the heap, attach it to the merged list, and if the extracted node has a next node, insert that into the heap. The process repeats until the heap is empty.

#218 The Skyline Problem [Hard] · Heap

Key Insights

- Use a line sweep approach to process building start and end events in sorted order.
- Differentiate between start events (which add to the skyline) and end events (which remove from the skyline).
- Employ a max heap (or a priority queue) to efficiently track the current highest building at each sweep line position.
- When the current maximum height changes, record the corresponding x-coordinate and new height as a key point.
- Take care to remove expired building heights from the heap as the sweep line moves past the building's end.

Space & Time

Time Complexity: $O(n \log n)$ due to sorting events and heap operations for each building event.

Space Complexity: $O(n)$ for the event list and the heap used for processing active buildings.

Solution

The solution uses the line sweep algorithm with a max heap (or priority queue) to maintain the heights of active buildings. Each building contributes two events: one when the building starts (adding its height) and one when it ends (removing its height). These events are sorted primarily by their x-coordinate. For start events, we use negative height values to ensure they are processed before end events at the same x-coordinate. During the sweep, the heap is updated by adding new building heights and removing buildings that have ended. Whenever there is a change in the top of the heap (i.e., the current maximum height), a key point is added to the result. This key point signifies a vertical change in the skyline.

#272 Closest Binary Search Tree Value II [Hard] · Heap

Key Insights

- The BST property allows efficient narrowing down of candidate nodes.
- Using two stacks to track predecessors (values $\leq$ target) and successors (values $>$ target) can efficiently yield the closest values.
- A modified in-order traversal helps in retrieving the next closest candidate from either side.
- Merging two sorted lists of candidate values based on absolute difference with the target is the key step.

Round 9 · Low · Hard

p.3369

Space & Time

Time Complexity: $O(k + \log n)$ – $O(\log n)$ to initialize the two stacks and $O(k)$ to collect k elements.

Space Complexity: $O(\log n)$ – due to the stack space in a balanced BST.

Solution

We can solve the problem by splitting it into two parts:

– **Build two stacks:** The predecessor stack for nodes with values less than or equal to target. The successor stack for nodes with values greater than target. To build these stacks, traverse the BST starting from the root and push nodes along the path: if the node's value is less than or equal to target, push it to the predecessor stack and go right; if greater, push it to the successor stack and go left. – **Merge the two stacks:** At each step, compare the top element of the predecessor stack (if available) and the successor stack (if available) based on their distance from the target. Pop from the stack with the closer candidate and add that value to the result list. For the popped node, update the given stack by moving to the next appropriate node in the in-order traversal (for predecessor, go to left child then rightmost descendant; for successor, go to right child then leftmost descendant). Repeat until k values have been collected.

Key Data Structures and Techniques:

Round 9 · Low · Hard

p.3370

- Two stacks to simulate reverse and forward in-order traversals.
- Merge technique to choose the closest candidate at each iteration.

#295 Find Median from Data Stream [Hard] · Heap

Key Insights

- Use two heaps: a max-heap for the lower half and a min-heap for the upper half of numbers.
- Balancing the heaps ensures the median can be computed in $O(1)$ time once numbers are added.
- Rebalance the heaps after each insertion to ensure size properties hold (difference in sizes is at most 1).

Space & Time

Time Complexity:

- addNum: $O(\log n)$ per insertion (heap operations).
- findMedian: $O(1)$ (accessing the top elements).

Space Complexity:

- $O(n)$, where n is the number of elements added in the data stream, as they are stored in the two heaps.

Solution

We maintain two heaps:

- A max-heap (lowerHalf) for the smaller half of the numbers.
- A min-heap (upperHalf) for the larger half of the numbers. On inserting a new number, decide the heap membership based on the current median. After each insertion, ensure that the heaps have balanced

sizes (the size difference is at most one). When computing the median: – If both heaps have the same number of elements, the median is the average of the max of lowerHalf and the min of upperHalf. – If one heap has more elements, the median is the top element of that heap.

This method makes it efficient to dynamically update and retrieve the median as new numbers are inserted.

#358 Rearrange String k Distance Apart [Hard] · Heap

Key Insights

- The frequency of each character dictates the feasibility and placement within the rearrangement.
- A greedy approach picks the most frequent available character to maximize the chance of successful placement.
- A max-heap (priority queue) efficiently selects the highest frequency character each time.
- A waiting queue ensures that once a character is used, it can't be reused until k positions later.
- Special case: When k is 0, no rearrangement is required because all orders satisfy the condition.

Space & Time

Time Complexity: $O(n \log C)$ where n is the length of s and C is the number of unique characters.

Space Complexity: $O(n)$ due to the storage used for the heap and wait queue.

Solution

We use a greedy algorithm with a max-heap to always select the character with the highest remaining frequency that is currently available. The algorithm first counts the frequency of each character and pushes these into a max-heap (using negative frequencies for languages like Python where heapq is a min-heap). Each time a character is selected, it is appended to the result string and its frequency is decreased. This character is then placed in a waiting queue to ensure it isn't reused until k positions later. Once the waiting period completes, if the character still has remaining occurrences, it is reinserted into the heap. If at any point the heap is empty and the waiting queue still contains characters that haven't been reinserted while the result string is incomplete, we conclude that a valid rearrangement is impossible.

#499 The Maze III [Hard] · Heap

Key Insights

- The ball rolls continuously until hitting a wall or falling into the hole.
- A shortest path search is needed based on rolling distance (each step counts when passing an empty cell).

- Dijkstra's algorithm (or a modified BFS with a priority queue) is appropriate since we need to consider distance and lexicographic order.
- When expanding moves, simulate the roll until the ball stops — if a hole is passed en route, use that stopping point.
- Track the best (distance, path) for each cell to avoid redundant work.

Space & Time

Time Complexity: $O(mn \log(mn))$ where m and n are the maze dimensions

Space Complexity: $O(mn)$

Solution

We use a Dijkstra-like approach with a priority queue containing tuples of (distance, path, row, col). The priority queue is ordered by distance first and then by the instruction string lexicographically. For each cell, simulate rolling in each possible direction (up, down, left, right; iterated in alphabetical order so that the lexicographical property is maintained) until the ball stops. If the ball encounters the hole during rolling, treat that as a valid stop immediately. For every new stopping cell, if a shorter distance or lexicographically smaller path is found, update the record and add the new state to the priority queue. If you reach the hole, return the corresponding path string. If no solution is found when the queue is empty, return "impossible".

#632 Smallest Range Covering Elements from K Lists [Hard] · Heap

Key Insights

- Use a min heap to maintain the smallest current element among the selected elements from each list.
- Track the current maximum among the chosen elements to easily compute the range.
- Iteratively replace the minimum element with the next element from the same list and update the current maximum.
- Stop when one of the lists is exhausted, as the range must include at least one element from each list.

Space & Time

Time Complexity: $O(N \log k)$, where N is the total number of elements across all lists.

Space Complexity: $O(k)$ for the heap, plus additional space for storing the range.

Solution

The solution uses a min heap (priority queue) that stores a tuple containing the current element, the index of the list it comes from, and its index in that list. Initially, the first element from each list is inserted into the heap and the current maximum among them is tracked. In each iteration, the smallest element is removed (heap root) which provides the current minimum, and the range $[current_min, current_max]$ is

assessed. If the list from which the minimum element was extracted has a next element, that element is inserted into the heap and the current maximum is updated if necessary. This process continues until one of the lists is exhausted.

#759 Employee Free Time [Hard] · Heap

Key Insights

- Flatten the schedule to combine all employees' intervals.
- Sort the flattened intervals by start time.
- Merge overlapping intervals to form a consolidated view of busy times.
- Identify gaps between merged intervals as the common free time.
- Ensure free intervals have strictly positive length.

Space & Time

Time Complexity: $O(N \log N)$, where N is the total number of intervals (due to sorting).

Space Complexity: $O(N)$, used for the flattened list and merged intervals.

Solution

The solution leverages interval merging. First, all intervals from every employee are collected and sorted by start time. By iterating through the sorted intervals, overlapping intervals are merged into one consolidated

busy period. Finally, the gaps between consecutive merged intervals are identified as free periods available to all employees. This approach efficiently handles the intervals with a simple sort and a linear scan.

#1425 Constrained Subsequence Sum [Hard] · Heap

Key Insights

- Use dynamic programming where $dp[i]$ represents the maximum sum of a valid subsequence ending at index i .
- For each index i , consider taking $nums[i]$ alone or appending it to a valid subsequence ending at an index j (where $i - j \leq k$) that maximizes the sum.
- To efficiently obtain the maximum $dp[j]$ in a sliding window of the last k indices, use a monotonic (decreasing) deque.
- The idea is to maintain the window maximum and update it as we proceed through the array.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$ ($O(k)$ extra space for the deque, with dp array using $O(n)$)

Solution

We solve the problem by defining a dp array where $dp[i] = nums[i] + \max(0, \text{maximum dp value in the window } [i-k, i-1])$. The trick is to include $\max(0, \dots)$ to allow starting a new subsequence when previous sums are

negative. To access the maximum dp in the window in constant time, a monotonic deque is maintained that stores indices with their dp values in decreasing order. At each index, we:

- Remove elements from the front of the deque if they are out of the current window ($i - \text{index} > k$).
- Use the front of the deque (if non-empty) as the maximum dp value from the previous k indices.
- Compute $dp[i]$ and update the global maximum answer.
- Remove from the back of the deque all indices with dp values less than the current $dp[i]$ since they are not useful.
- Insert the current index into the deque.

#212 Word Search II [Hard] · Trie

Key Insights

- Use Depth-First Search (DFS) to explore possible words starting from each board cell.
- Build a Trie (prefix tree) to store the words for $O(1)$ prefix checks and efficient pruning.
- Mark visited cells and restore them after DFS to avoid revisiting in the current search path.
- Prune search paths early if the current prefix is not part of any word in the Trie.

Space & Time

Time Complexity: $O(m * n * 4^L)$ in the worst case, where L is the maximum word length; however, using a Trie prunes unnecessary paths.

Space Complexity: $O(K)$ for the Trie storage, where K is the total number of letters in all words, plus $O(m * n)$ auxiliary space for recursion and board marking.

Solution

We first insert all words into a Trie data structure to allow quick lookups of prefixes. Then, for each cell in the board, we perform DFS. During DFS, if the current path matches a Trie node that represents a complete word, we add that word to the result list and mark it as found to avoid duplicates. We continue exploring adjacent cells (up, down, left, right) while marking the cell as visited. After exploring, we restore the original letter for further DFS calls. This combination of Trie and DFS with backtracking efficiently finds valid words on the board.

#336 Palindrome Pairs [Hard] · Trie

Key Insights

- A palindrome is a string that reads the same forwards and backwards.
- For any given word, if you split it into two parts, one part (either prefix or suffix) might be a palindrome.
- If the other part's reverse exists elsewhere in the array, then that reverse can be concatenated to yield a palindrome.
- Carefully handle edge cases like empty strings, since an empty string is a palindrome by default.

Round 9 · Low · Hard

p.3379

Space & Time

Time Complexity: $O(n * k^2)$ in the worst case where n is the number of words and k is the maximum word length, but with the palindrome checking optimized this can be considered $O(\text{sum of word lengths})$: Each word is processed once and each split is checked in $O(k)$ time.

Space Complexity: $O(n * k)$ due to the hashmap storing all words and their indices.

Solution

We use a hashmap (dictionary) to store each word and its corresponding index. For every word, we iterate through every possible split position to divide the word into a prefix and a suffix. For each split:

- If the prefix is a palindrome, we look for the reverse of the suffix in the hashmap. If found (and not the same word), then combining the reverse with the current word forms a palindrome.
- Similarly, if the suffix is a palindrome, we look for the reverse of the prefix.

We must be cautious with edge cases, especially when either part is empty, as the empty string is considered a palindrome. This method avoids checking every possible pair explicitly and leverages string reversal and palindrome checks to achieve efficiency.

#588 Design In-Memory File System [Hard] · Trie

Key Insights

- Use a tree (trie) to represent directories and files.

Round 9 · Low · Hard

p.3380

- Each node can be either a directory or a file; directories maintain a mapping of names to child nodes.
- For ls, if the given path is a file, return only its name. If it's a directory, return the sorted list of names of its children.
- mkdir must create all intermediate directories if they do not already exist.
- File nodes store content and support content appending.

Space & Time

Time Complexity:

- ls: $O(k \log k)$ where k is the number of entries in the directory (due to sorting), or $O(n)$ if traversing a long path.
- mkdir, addContentToFile, readContentFromFile: $O(d)$ where d is the depth of the file path.

Space Complexity:

- $O(T)$ where T is the total number of characters stored in directory names and file contents.

Solution

The solution uses a tree-based approach where each node represents either a directory or a file. Each directory node stores its children in a hash table (or dictionary) mapping names to node objects. File nodes store their content as a string and carry a flag indicating they are files. For the ls operation, check if the node is a

file—if so, return its name; otherwise, list and sort the names of all children in that directory. The mkdir operation traverses (or creates) nodes for each segment of the provided path. The addContentToFile method accesses or creates the file node and appends the content. Finally, readContentFromFile retrieves the stored content from the file node.

#642 Design Search Autocomplete System [Hard] · Trie

Key Insights

- Use a Trie to efficiently store and retrieve sentences based on their prefixes.
- Maintain a mapping at each Trie node to record sentences (or their frequencies) that pass through that node.
- Use sorting or a min-heap to select the top three sentences based on their frequency (hot degree) and lexicographical order if frequencies tie.
- On receiving '#', update the Trie with the current sentence and reset the current search state.

Space & Time

Time Complexity: $O(L * \log k)$ per character input query, where L is the length of the current prefix and k is 3 (constant factor), thus ensuring efficient autocomplete.

Space Complexity: $O(N * L)$, where N is the number of unique sentences and L is the average sentence length, for maintaining the Trie.

Solution

The solution revolves around building a Trie where each node contains:

- A mapping to its child nodes.
- A hash map that records all sentences passing through the node along with their occurrence counts.

For each character entered (except '#'), the system:

- Traverses the Trie to find the current prefix node.
- Retrieves stored sentences at that node and sorts them based on the following criteria: descending by frequency and ascending lexicographically if frequencies are equal.
- Returns the top three sentences from the sorted result.

When the user inputs '#', the current sentence (tracked as user input so far) is added or its frequency updated in the Trie, and the input state resets.

#37 Sudoku Solver [Hard] · Backtracking

Key Insights

- Use backtracking to explore digit placements for each empty cell.
- Validate placements by ensuring that the chosen digit is not already present in the corresponding row, column,

or 3x3 sub-box.

- Employ recursion to fill the board incrementally and backtrack upon encountering invalid configurations.
- Although the worst-case time complexity is exponential, the fixed board size (9x9) keeps the problem computationally manageable.

Space & Time

Time Complexity: In the worst-case scenario, the algorithm can explore up to $O(9^n)$ possibilities where n is the number of empty cells. However, since n is at most 81 and many branches are pruned early, the practical runtime is far less.

Space Complexity: $O(n)$ due to the recursion stack, where n is the number of empty cells; given the board's fixed size, this space is constant in practice.

Solution

The solution uses a backtracking approach. For each empty cell on the board, we attempt to place a digit from '1' to '9'. Before placing a digit, we check if the placement is valid by ensuring that:

- The digit is not already present in the same row.
- The digit is not already present in the same column.
- The digit is not already present in the corresponding 3x3 sub-box.

If a valid digit is found, we place it on the board and recursively attempt to solve the rest of the board. If no valid digit can be placed, the algorithm backtracks by

reverting the last move and trying a different digit. This recursive search continues until the board is completely filled with a valid Sudoku configuration.

#51 N-Queens [Hard] · Backtracking

Key Insights

- Use backtracking to explore potential positions row by row.
- Use sets to track which columns and diagonals are under attack.
- Diagonals can be tracked by (row - col) and (row + col), ensuring that queens do not conflict along any diagonal.
- Recursively build each valid board configuration and backtrack when a conflict arises.

Space & Time

Time Complexity: $O(n!)$ in the worst-case scenario, as the solution explores permutations for queen placements.

Space Complexity: $O(n^2)$ for storing board configurations and $O(n)$ for recursion stack and conflict-tracking sets.

Solution

We use a backtracking algorithm that places queens row by row. For each row, we try placing a queen in each column. Before placing, we check if the column, main

diagonal (row - col), and anti-diagonal (row + col) are free using dedicated sets. If a queen can be safely placed, we mark the column and diagonals as occupied and recurse to the next row. Once all rows are processed, the board configuration is a valid solution and is added to the result. After that, we backtrack by unmarking the sets and removing the queen from the board. This systematic approach ensures all possible configurations are considered.

#126 Word Ladder II [Hard] · Backtracking

Key Insights

- Use Breadth-First Search (BFS) to traverse level-by-level and capture only the shortest paths.
- Build a parent mapping during BFS to record which words lead to each newly discovered word.
- Once the shortest path is found via BFS, use backtracking (or DFS) from endWord to reconstruct all transformation sequences.
- Remove words as they are visited within a level to avoid unnecessary revisits and cycles.
- The solution must handle multiple shortest paths, so maintaining a list of predecessors for each word is crucial.

Space & Time

Time Complexity: $O(N * L * 26)$ where N is the number of words in the wordList and L is the length of each word.

This accounts for checking all possible one-letter transformations.

Space Complexity: $O(N * L)$ for storing the parent mapping and maintaining the BFS queue and other auxiliary data structures.

Solution

We solve the problem in two major phases:

– Use BFS to explore from beginWord to endWord.

During BFS, record all parents for each word encountered. This ensures we capture all possible predecessors that yield a shortest transformation. – Once the BFS completes (once endWord is reached), use backtracking (DFS) beginning from the endWord and moving backwards via the parent mapping to reconstruct all transformation sequences in reverse. Finally, reverse each sequence to present it from beginWord to endWord.

This two-phase approach guarantees that only the shortest paths are reconstructed, as BFS ensures level-by-level exploration.

#996 Number of Squareful Arrays [Hard] · Backtracking

Key Insights

- The problem requires ensuring adjacent pairs sum up to a perfect square.

- Sorting the array helps in efficiently handling duplicates; when iterating in the backtracking, skip duplicate elements to avoid overcounting.
- For each valid candidate, check that the sum with the previous number (if any) is a perfect square.
- Use backtracking (DFS) to generate all valid permutations while maintaining a visited flag for each index.
- Since the array may contain duplicates, extra care is taken to only use each occurrence in a controlled order (i.e., if two identical numbers exist, only the first unvisited instance can be chosen at a given recursive call).

Space & Time

Time Complexity: $O(n^2 * 2^n)$ in the worst-case scenario due to exploring all bitmask states with n elements and checking pairwise sums.

Space Complexity: $O(n)$ for recursion and $O(n)$ for the visited flag, plus potentially $O(2^n * n)$ for memoization if implemented.

Solution

The solution uses a backtracking algorithm to explore all permutations. The array is first sorted to handle duplicates. A helper function checks if a number is a perfect square and is used to validate that the sum of every adjacent pair is a perfect square. During DFS, a visited array is maintained to track elements included so

far. Skipping duplicate elements (unless the earlier duplicate was used) is essential to avoid overcounting. This approach ensures that every valid permutation is considered exactly once.

Round 9 · Low · Hard

p.3389

R7 · Low · Easy (14 q)

>> Round 7 · Low Priority · Easy — Quick Review

14 questions · key insights · complexity · solution

#70 Climbing Stairs [Easy] · Dynamic Programming

Key Insights

- The problem can be solved using dynamic programming because the solution for a given step depends on the solutions of the previous two steps.
- The recurrence relation is: $\text{ways}[i] = \text{ways}[i-1] + \text{ways}[i-2]$.
- Base cases: when there is 1 step, there is 1 way; when there are 2 steps, there are 2 ways.
- This is analogous to the Fibonacci sequence, shifting indices accordingly.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$ (can be optimized to $O(1)$ using constant space)

Solution

We use a dynamic programming approach to build a solution gradually. The idea is that to reach step i , you could have come from step $i-1$ (taking one step) or from

Round 9 · Low · Hard

p.3390

step $i-2$ (taking two steps). Therefore, the total number of ways to reach step i is the sum of ways to reach step $i-1$ and $i-2$. We initialize $dp[1] = 1$ and $dp[2] = 2$, and iterate from 3 to n to fill $dp[]$. Each code solution below uses this approach with comments to explain the logic.

#121 Best Time to Buy and Sell Stock [Easy] · Dynamic Programming

Key Insights

- Traverse the list of prices once while keeping track of the minimum price encountered so far.
- At each day, calculate the profit if you sold the stock on that day.
- Update the maximum profit when the current profit exceeds the previous maximum.
- The optimal solution makes a single pass through the array with constant space usage.

Space & Time

Time Complexity: $O(n)$ – requires one pass through the prices array.

Space Complexity: $O(1)$ – only a few extra variables are used.

Solution

We solve the problem using a simple linear scan. We maintain two key variables: one for the minimum price seen so far and another for the maximum profit

calculated so far. As we iterate, we continuously update the minimum price if a lower price is found, and then compute the potential profit at the current price. If this profit exceeds the current maximum profit, we update the maximum profit. The final maximum profit is returned after iterating through all prices.

#67 Add Binary [Easy] · Bit Manipulation

Key Insights

- Simulate binary addition just like you add numbers digit by digit.
- Process the strings from right to left, handling the carry at each step.
- Use modulo and division operations to extract the current digit and update the carry.
- Append any remaining carry at the end.

Space & Time

Time Complexity: $O(\max(n, m))$, where n and m are the lengths of the two binary strings.

Space Complexity: $O(\max(n, m))$ for storing the resulting binary string.

Solution

The strategy is to iterate over the input strings from the end to the beginning. At each step, calculate the sum of corresponding bits and the existing carry. Use the modulo operation ($\% 2$) to determine the current digit

and integer division ($//$ 2) to update the carry. Once all digits have been processed, reverse the result to obtain the final binary string.

#136 Single Number [Easy] · Bit Manipulation

Key Insights

- Using bitwise XOR: XOR-ing a number with itself cancels out to 0.
- XOR is both commutative and associative, so the order of operations does not matter.
- A single traversal (linear time) and constant extra space is sufficient for solving the problem.

Space & Time

Time Complexity: $O(n)$ – We only traverse the array once.

Space Complexity: $O(1)$ – Only a single extra variable is used regardless of input size.

Solution

We can solve this problem by initializing a variable to 0 and then traversing the array while XOR-ing each element with the variable. Due to the properties of XOR, pairs of identical numbers will cancel each other out ($x \text{ XOR } x = 0$) and the unique number will be left as the final result. This approach uses constant extra space and performs in linear time.

#190 Reverse Bits [Easy] · Bit Manipulation

Key Insights

- The input is a binary string of fixed length (32 bits).
- Reversing the bits is equivalent to reversing the string.
- The reversed binary string can be converted back into its integer representation.
- Since the number of bits is constant (32), the solution can be achieved in constant time and constant space.
- For performance optimization when this function is called many times, caching may be applied for common intermediate results.

Space & Time

Time Complexity: $O(1)$ – We always process 32 bits regardless of the input.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

We solve the problem by reversing the 32-character binary string representing the unsigned integer. Once reversed, the new binary string is converted into its integer value. An alternative solution involves using bitwise operations where you iterate over each bit, shift the result to the left, and add the current bit. This approach does the reversal one bit at a time. The chosen approach leverages string reversal which is straightforward to implement and understand.

#191 Number of 1 Bits [Easy] · Bit Manipulation

Key Insights

- The problem is essentially counting the number of 1s in the binary representation of a number.
- A common efficient approach uses bit manipulation: repeatedly clear the lowest set bit until n becomes zero.
- The expression $n \& (n - 1)$ clears the lowest set bit of n .
- This method only iterates as many times as there are set bits, making it efficient.

Space & Time

Time Complexity: $O(k)$, where k is the number of set bits, worst-case $O(\log n)$ (or $O(1)$ considering 32-bit integers).

Space Complexity: $O(1)$.

Solution

To solve the problem, we use the bit manipulation trick $n \& (n - 1)$ which removes the lowest set bit from n . The algorithm is as follows:

– Initialize a counter to 0. – While n is not zero, update n by $n = n \& (n - 1)$, and increment the counter. – The counter will hold the number of set bits by the end of the loop. This approach leverages efficient bit-level operations and works well even when the function is called frequently.

#268 Missing Number [Easy] · Bit Manipulation

Key Insights

- The range of numbers is continuous from 0 to n , but one number is missing.
- The sum of numbers from 0 to n can be computed using the arithmetic sum formula.
- Subtracting the sum of the given array from the expected sum gives the missing number.
- There is a follow-up challenge to solve the problem in $O(n)$ time and $O(1)$ extra space.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We compute the expected sum of all numbers from 0 to n using the formula $\text{sum} = n*(n+1)/2$. Then, we compute the actual sum of the numbers in the array. The missing number is obtained by subtracting the actual sum from the expected sum. This approach uses arithmetic operations ($O(1)$ extra space) and loops over the array once ($O(n)$ time).

#338 Counting Bits [Easy] · Bit Manipulation

Key Insights

- We can use dynamic programming to build the solution by utilizing the count of bits for smaller numbers.
- The relationship $dp[i] = dp[i \gg 1] + (i \& 1)$ holds, where $i \gg 1$ gives the count for i divided by 2 (dropping the least significant bit) and $(i \& 1)$ checks if the least significant bit is 1.
- This approach computes the result for each number in constant time, resulting in a linear time solution.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(n)$

Solution

We use a dynamic programming approach where an array (or list) dp is initialized with size $n+1$. The key idea is to iterate from 1 to n and for each number, use the formula $dp[i] = dp[i \gg 1] + (i \& 1)$.

– The expression $(i \gg 1)$ effectively divides i by 2, giving the count of 1's for that half number. – The $(i \& 1)$ operation checks if i is odd, adding 1 if true. This method ensures each $dp[i]$ is computed in constant time, leading to an overall $O(n)$ time complexity.

#9 Palindrome Number [Easy] · Math

Key Insights

- Negative numbers are not palindromic.

Round 9 · Low · Hard

p.3397

- A number ending in 0 (except 0 itself) cannot be a palindrome.
- Instead of reversing the entire number (risking overflow), reverse only half of the number and compare with the other half.
- For numbers with an odd number of digits, the middle digit can be disregarded.

Space & Time

Time Complexity: $O(\log_{10}(n))$ – Each iteration reduces the number by a factor of 10.

Space Complexity: $O(1)$ – Only constant extra space is used.

Solution

The solution checks for obvious non-palindrome cases first (negative numbers and numbers ending with 0 that are not 0). Then, it reverses half of the digits of the number while the original number is larger than the reversed half. Finally, it compares the remaining part of the original number with the reversed half. In the case of an odd number of digits, the middle digit is removed by dividing the reversed half by 10 before the final comparison.

#13 Roman to Integer [Easy] · Math

Key Insights

Round 9 · Low · Hard

p.3398

- Map each Roman numeral character to its integer value using a hash table or dictionary.
- Loop through the string and check if the current numeral is less than the next numeral; if so, subtract its value.
- Otherwise, add its value to the total.
- This approach takes advantage of the subtractive notation used in Roman numerals.

Space & Time

Time Complexity: $O(n)$, where n is the length of the input string.

Space Complexity: $O(1)$, since the mapping and additional variables use constant space.

Solution

We use a dictionary (or hash map) to store the values of each Roman numeral. As we iterate over the string, we compare the current numeral's value with the next numeral's value. When a numeral of smaller value appears before a numeral of larger value, we subtract the current numeral's value; otherwise, we add it. This technique efficiently handles both the standard and subtractive cases in one pass through the string.

#504 Base 7 [Easy] · Math

Key Insights

- Handle the edge case where num is 0.

Round 9 · Low · Hard

p.3399

- Use division and modulo operations to extract base 7 digits.
- Manage negative numbers by converting to positive and reattaching the minus sign.
- Build the base 7 digits in reverse order during iteration.

Space & Time

Time Complexity: $O(\log_7(n))$ where n is the absolute value of the number.

Space Complexity: $O(\log_7(n))$ to store the digits of the number.

Solution

The algorithm follows these steps:

– Check if num equals 0 and return "0". – Determine if the number is negative. If so, work with its absolute value but remember the negative flag. – Iteratively compute the remainder when divided by 7 (using modulo) and update the number by integer division by 7. – Append each remainder (converted to character) to a list representing the digits. – The digits are computed in reverse order; reverse the list for the correct order. – If the original number was negative, append the "-" sign to the front. – Join the digits to form the resulting base 7 string.

This approach uses basic arithmetic (division and modulo) and a list (or equivalent structure) to accumulate the computed digits.

Round 9 · Low · Hard

p.3400

#1056 Confusing Number [Easy] · Math

Key Insights

- Use a mapping dictionary for valid rotations: {0:0, 1:1, 6:9, 8:8, 9:6}.
- Process the digits of n from right to left to form the rotated number.
- If any digit is not present in the mapping, the number is instantly not confusing.
- Compare the rotated number with the original number to check if they are different.

Space & Time

Time Complexity: $O(d)$, where d is the number of digits in n (approximately $O(\log n)$).

Space Complexity: $O(d)$, needed to store the rotated number string.

Solution

The approach is to convert the number to a string and iterate through it in reverse. For each digit, check if it is valid using the predefined mapping. If it is, append its rotated counterpart to a new string. After processing, convert the rotated string to an integer (ignoring any leading zeros). Finally, compare the rotated number to the original number. If they are different, the number is confusing.

#1228 Missing Number in Arithmetic Progression [Easy] · Math

Key Insights

- The complete arithmetic progression has equal differences between consecutive numbers.
- The missing element can be found by calculating the expected common difference using the first and last elements.
- Since exactly one number is missing, iterate through the provided array to detect where the difference deviates from the expected common difference.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

We first compute the expected common difference (diff) for the original arithmetic progression. Since the actual progression had one more element than the current array, the formula used is: $\text{diff} = (\text{last element} - \text{first element}) / (n)$ where n is the length of the current array. Next, iterate over the array and check the difference between each pair of consecutive elements. When the actual difference deviates from diff, it indicates that the missing number is the previous element plus diff. This approach uses a simple loop and constant extra space.

#1427 Perform String Shifts [Easy] · Math**Key Insights**

- Consolidate all shifts by calculating a net shift: subtract for left shifts and add for right shifts.
- Use modulo arithmetic with the string length to avoid unnecessary full rotations.
- Slice the string based on the effective shift to construct the final string.

Space & Time

Time Complexity: $O(n + m)$, where n is the length of the string and m is the number of shift operations.

Space Complexity: $O(n)$ for the output string.

Solution

The solution involves computing the net shift value by iterating through the shift operations. Left shifts subtract from the net value, and right shifts add to the net value. After obtaining the cumulative shift, use modulo with the string length to normalize the shift, ensuring it lies within bounds. The final string is constructed by slicing the string into two parts and concatenating them appropriately. This approach leverages string slicing as the primary operation for shifting.

Round 9 · Low · Hard

p.3403

R8 · Low · Medium (32 q)

>> Round 8 · Low Priority · Medium — Quick Review

32 questions · key insights · complexity · solution

#5 Longest Palindromic Substring [Medium] · Dynamic Programming**Key Insights**

- Every palindrome is centered around a character (for odd-length) or a pair of characters (for even-length).
- Expand from each center until the substring is no longer a palindrome.
- Update the longest palindrome when a longer one is found during the expansion.
- The solution employs a two-pointer technique without extra storage for dynamic programming tables.

Space & Time

Time Complexity: $O(n^2)$ in the worst case as we expand around each center.

Space Complexity: $O(1)$, excluding the space required for the output substring.

Solution

The approach is based on expanding around potential centers of palindromes. For each character in the string,

Round 9 · Low · Hard

p.3404

consider it as the center for an odd-length palindrome and the gap between it and the next character for an even-length palindrome. Using two pointers, expand outward until the characters differ. Track the longest palindrome found during these expansions. This method is efficient given the input constraints, and no extra data structures are needed besides variables to store indices or substrings.

#53 Maximum Subarray [Medium] · Dynamic Programming

Key Insights

- Use Kadane's Algorithm to solve the problem in $O(n)$ time.
- Keep track of a running sum (currentSum) and update it with the maximum between the current element and the sum including the current element.
- The maximum sum encountered during the iteration (maxSum) is the answer.
- A divide and conquer approach also exists, splitting the array and merging solutions, but Kadane's method is optimal for this problem.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

Round 9 · Low · Hard

p.3405

The solution uses Kadane's Algorithm. Iterate over the array while maintaining two variables: currentSum, which tracks the maximum subarray sum ending at the current position, and maxSum, which stores the best sum seen so far. For each element, update currentSum to be either the current element or the sum of currentSum and the current element, whichever is higher. Update maxSum accordingly. This approach efficiently computes the maximum contiguous subarray sum in a single pass.

#55 Jump Game [Medium] · Dynamic Programming

Key Insights

- The greedy approach is optimal for this problem.
- Track the furthest index reachable as you iterate through the array.
- If the current index exceeds the maximum reachable index, it is impossible to proceed further.
- The solution only requires constant extra space.

Space & Time

Time Complexity: $O(n)$ where n is the length of the array.

Space Complexity: $O(1)$

Solution

The solution uses a greedy algorithm. As you iterate through each array element, you maintain a variable (maxReach) that represents the furthest index you can reach so far. For each index i , if i is greater than

Round 9 · Low · Hard

p.3406

maxReach then it is not possible to reach index i and hence return false immediately. Otherwise, update maxReach to the maximum of its current value and $i + \text{nums}[i]$. If you can iterate through the array without i exceeding maxReach, then you can reach the end and return true.

#62 Unique Paths [Medium] · Dynamic Programming

Key Insights

- The robot is constrained to only move right or down.
- Any valid path consists of exactly $(m-1)$ down moves and $(n-1)$ right moves.
- The problem can be solved using Dynamic Programming where each cell is the sum of the ways to reach the cell from the top and left.
- Alternatively, a combinatorial solution is possible by calculating the number of unique sequences of moves, equivalent to choosing $(m-1)$ moves from $(m+n-2)$ moves.

Space & Time

Time Complexity: $O(mn)$

Space Complexity: $O(mn)$ – can be optimized to $O(n)$ using a 1D DP array.

Solution

Round 9 · Low · Hard

p.3407

We can solve the problem using a dynamic programming approach. The idea is to create a 2D DP table where each cell (i, j) represents the number of ways to reach that cell. Since the robot can only come from above or from the left, we update the DP table as follows:

- Initialize $\text{dp}[0][j] = 1$ for all j (only one way to move right in the first row).
 - Initialize $\text{dp}[i][0] = 1$ for all i (only one way to move down in the first column).
 - For each other cell, compute $\text{dp}[i][j] = \text{dp}[i-1][j] + \text{dp}[i][j-1]$.
- The final answer will be at $\text{dp}[m-1][n-1]$.

The combinatorial alternative computes $C(m+n-2, m-1)$ or $C(m+n-2, n-1)$ using basic factorial arithmetic, which is efficient for the given constraints.

#72 Edit Distance [Medium] · Dynamic Programming

Key Insights

- This is a classic dynamic programming problem.
- Use a 2D DP table where $\text{dp}[i][j]$ represents the minimum edit distance between $\text{word1}[0\dots i-1]$ and $\text{word2}[0\dots j-1]$.
- Base cases: converting an empty string to a string of length j requires j insertions, and vice versa.
- Transition: If characters match then $\text{dp}[i][j] = \text{dp}[i-1][j-1]$, otherwise take minimum among insertion, deletion, and replacement operations and add 1.

Round 9 · Low · Hard

p.3408

Space & Time

Time Complexity: $O(mn)$

, where m and n are the lengths of word1 and word2 respectively.

Space Complexity: $O(mn)$

due to the DP table used for storing computed distances.

Solution

We use a dynamic programming approach by constructing a 2D table dp of size $(m+1) \times (n+1)$, where m and n are the lengths of word1 and word2. Each $dp[i][j]$ holds the minimum number of operations needed to convert the first i characters of word1 to the first j characters of word2. The algorithm includes the following steps:

- Initialize the base cases: $dp[0][j] = j$ (converting an empty word1 to j characters of word2 requires j insertions) $dp[i][0] = i$ (converting i characters of word1 to an empty word2 requires i deletions) – Iterate over each character in word1 and word2: If the current characters are equal, no additional operation is needed and $dp[i][j] = dp[i-1][j-1]$. Otherwise, choose the best operation among insertion ($dp[i][j-1]$), deletion ($dp[i-1][j]$), or substitution ($dp[i-1][j-1]$) and add 1. – The answer is found in $dp[m][n]$.

This method reliably computes the edit distance using a systematic DP approach, ensuring that all subproblems

are solved and reused efficiently.

#91 Decode Ways [Medium] · Dynamic Programming

Key Insights

- A digit '0' cannot be mapped on its own; it must be part of "10" or "20".
- Use dynamic programming where $dp[i]$ represents the number of ways to decode the substring $s[0...i-1]$.
- For each index i , if the digit at position $i-1$ is non-zero, add $dp[i-1]$ to $dp[i]$.
- If the two-digit number formed by $s[i-2]$ and $s[i-1]$ is between 10 and 26, add $dp[i-2]$ to $dp[i]$.
- Initialize $dp[0]$ as 1 (an empty string has one way to be decoded).

Space & Time

Time Complexity: $O(n)$, where n is the length of the string.

Space Complexity: $O(n)$, due to the dp array (which can be optimized to $O(1)$).

Solution

We solve the problem using a dynamic programming approach. We initialize a dp array where $dp[i]$ indicates the number of ways to decode the substring $s[0...i-1]$. We set $dp[0] = 1$ since an empty string is trivially decodable. We loop through the string, and at each step:

– If $s[i-1]$ is not '0', we add $dp[i-1]$ to $dp[i]$ because $s[i-1]$ can be mapped to a valid letter. – If the substring $s[i-2:i]$ (for $i \geq 2$) forms a valid two-digit number between 10 and 26, we add $dp[i-2]$ to $dp[i]$. This cumulative approach yields the total ways to decode the string, and if the string starts with '0' or contains invalid sequences, $dp[n]$ will eventually be 0.

#152 Maximum Product Subarray [Medium] · Dynamic Programming

Key Insights

- Maintain both the maximum and minimum product at the current index because a negative number can swap the max and min.
- A subarray must be contiguous.
- Use dynamic programming to update both the maximum product and the minimum product ending at each index.
- Resetting the product in case of a zero is crucial since the product becomes 0.

Space & Time

Time Complexity: $O(n)$ – The solution iterates through the array once.

Space Complexity: $O(1)$ – Constant space is used for tracking products.

Solution

Round 9 · Low · Hard

p.3411

The solution uses an iterative dynamic programming approach. At each index, compute the maximum product ($maxEndingHere$) and minimum product ($minEndingHere$) ending at that index. If the current element is negative, swap these two values, because multiplying a negative with a minimum product (possibly negative) might yield a larger product. After processing each element, update the global maximum. This method handles edge cases like zeros and negatives efficiently with constant additional space.

#198 House Robber [Medium] · Dynamic Programming

Key Insights

- The problem can be solved using Dynamic Programming.
- For each house, decide whether to rob it or not.
- Recurrence relation: $dp[i] = \max(dp[i-1], dp[i-2] + \text{currentHouseMoney})$
- Use iterative computation to keep track of the optimal solutions without needing extra space for the entire dp array.

Space & Time

Time Complexity: $O(n)$, where n is the number of houses.

Space Complexity: $O(1)$, using only two variables to track the previous decisions.

Round 9 · Low · Hard

p.3412

Solution

We use a dynamic programming approach by iterating through each house in the array. At each step, we consider two scenarios: either we rob the current house (and add its money to the total from two houses ago) or we skip it (keeping the optimal solution up to the previous house). By using two variables to store these intermediate results, we achieve an optimized space complexity of $O(1)$.

#256 Paint House [Medium] · Dynamic Programming

Key Insights

- The choice for painting each house depends on the previous house's color choices.
- Use dynamic programming to accumulate the cost while ensuring adjacent houses do not have the same color.
- The state of each house can be represented by the minimum cost when the house is painted a specific color, computed using the minimal costs of the previous house from the other two colors.
- The problem only requires updating the cost matrix in-place for optimal space usage.

Space & Time

Time Complexity: $O(n)$, where n is the number of houses.

Space Complexity: $O(1)$ extra space if the input matrix is modified in place.

Solution

We use a dynamic programming approach with iterative in-place updating of the cost matrix. For each house starting from the second one, we update the cost for each color by adding the minimal cost of painting the previous house with either of the other two colors. This guarantees we never use the same color for adjacent houses while accumulating the minimal total cost. The final answer is the minimum value for the last house.

#309 Best Time to Buy and Sell Stock with Cooldown [Medium] · Dynamic Programming

Key Insights

- Use Dynamic Programming to solve the problem.
- Maintain three states to represent:
 - s_0 : the state where you are ready to buy (not holding any stock).
 - s_1 : the state where you are holding a stock.
 - s_2 : the cooldown state encountered right after selling a stock.
- Transition between states:
 - Transition into s_0 can come from staying in s_0 or coming out from the cooldown state s_2 .

- To enter s_1 , you either continue holding your stock, or buy a stock using the profit available from s_0 .
- The only way to reach s_2 is by selling the stock from s_1 .
- The final answer is determined by the maximum profit when you are not holding a stock (i.e., $\max(s_0, s_2)$).

Space & Time

Time Complexity: $O(n)$, where n is the number of days.

Space Complexity: $O(1)$, since only a few variables are maintained for the states.

Solution

The solution involves updating three state variables for each day:

– s_0 (ready to buy) is updated as the maximum of the previous s_0 and the previous cooldown s_2 . – s_1 (holding a stock) is updated as the maximum of the previous s_1 or buying today with the profit from s_0 minus the current price. – s_2 (cooldown) is updated by selling the stock held in s_1 (i.e., previous s_1 plus the current price). At the end of the iteration, the maximum profit will be in state s_0 (ready to buy) or s_2 (cooldown), since remaining in s_1 means you are holding a stock, which does not constitute a realized profit.

#377 Combination Sum IV [Medium] · Dynamic Programming

Round 9 · Low · Hard

p.3415

Key Insights

- Use dynamic programming to count the number of valid combinations.
- The order in which numbers are picked matters, making this akin to permutation counting.
- Initialize a DP table where $dp[i]$ represents the number of ways to form the sum i .
- $dp[0]$ is initialized to 1 since there is one way to form a sum of 0 (by choosing nothing).
- For each sum i from 1 to target, iterate over $nums$ and update $dp[i]$ based on $dp[i - num]$ when $i \geq num$.
- For negative numbers, the problem becomes unbounded due to potential infinite combinations; hence a constraint on the combination length is required to make it well-defined.

Space & Time

Time Complexity: $O(\text{target} * n)$, where n is the length of $nums$.

Space Complexity: $O(\text{target})$, for the dp array of size $\text{target} + 1$.

Solution

The problem is solved using a bottom-up dynamic programming approach. We create a dp array where each $dp[i]$ represents the number of possible combinations to reach the sum i . We iterate from 1 up to target, and for each sum, we check every number in $nums$. If the

Round 9 · Low · Hard

p.3416

number can contribute to the current sum ($i - \text{num} \geq 0$), we add the number of ways to form the remainder ($i - \text{num}$) to $\text{dp}[i]$. This ensures that every possible permutation is counted. When negative numbers are allowed in the input, the potential for an infinite number of combinations arises unless we impose additional constraints, such as limiting the length of the combination.

#416 Partition Equal Subset Sum [Medium] · Dynamic Programming

Key Insights

- The total sum of the array must be even, otherwise, partitioning into two equal subsets is impossible.
- The problem can be transformed into a "subset sum" problem where we need to check if there exists a subset of numbers that sums to half of the total sum.
- A dynamic programming approach can efficiently decide if such a subset exists.

Space & Time

Time Complexity: $O(n * \text{target})$, where target is half of the total sum.

Space Complexity: $O(\text{target})$, using a one-dimensional DP array.

Solution

Round 9 · Low · Hard

p.3417

First, compute the total sum of the array. If the total is odd, return false since an odd number cannot be equally partitioned. Next, set the target sum to half of the total. Use a dynamic programming array (dp) where $\text{dp}[i]$ represents whether a subset with sum i is achievable. Initialize $\text{dp}[0]$ as true because a subset sum of 0 is always possible. Then, for each number in the array, iterate through the dp array backwards from target to the number's value and update $\text{dp}[i]$ to true if $\text{dp}[i - \text{num}]$ is true. Ultimately, if $\text{dp}[\text{target}]$ is true, the array can be partitioned into two subsets with equal sum.

#435 Non-overlapping Intervals [Medium] · Dynamic Programming

Key Insights

- Sorting is the key: By sorting intervals by their end times, we can greedily choose the optimal set of non-overlapping intervals.
- Greedy approach: Always select the interval that ends the earliest, as it leaves maximum room for subsequent intervals.
- Counting removals: The answer is the total number of intervals minus the count of intervals selected by this greedy approach.
- Edge observation: Intervals that touch at the boundary are not overlapping, so the check uses " $\geq$ " instead of " $>$ ".

Round 9 · Low · Hard

p.3418

Space & Time

Time Complexity: $O(n \log n)$ due to sorting the intervals.

Space Complexity: $O(n)$ if considering the storage for the sorted list (or $O(1)$ if sorting in-place).

Solution

The solution involves first sorting the intervals by their end times. Initialize a variable to track the end of the last chosen interval. Iterate over the sorted intervals: if the current interval's start is greater than or equal to the last chosen interval's end, update the last end and include this interval in the non-overlapping set. Otherwise, count the interval as one that must be removed. The final answer is the removal count, which is computed as the total number of intervals minus the count of non-overlapping intervals selected.

#516 Longest Palindromic Subsequence [Medium] · Dynamic Programming

Key Insights

- The problem can be solved efficiently using dynamic programming.
- A 2D DP table is used where $dp[i][j]$ represents the length of the longest palindromic subsequence in the substring $s[i..j]$.
- If the characters at the two ends of the substring match ($s[i] == s[j]$), they contribute to the palindrome

length.

- Otherwise, the solution is the maximum sequence length by either excluding the left or the right character.
- The approach builds the solution from smaller substrings (length 1) to the entire string.

Space & Time

Time Complexity: $O(n^2)$ where n is the length of the string.

Space Complexity: $O(n^2)$ due to the usage of a 2D DP table.

Solution

We solve this problem using a bottom-up dynamic programming approach. We create a 2D table $dp[][]$ such that $dp[i][j]$ stores the longest palindromic subsequence length for the substring $s[i..j]$. For any single character, $dp[i][i]$ is 1. For substrings of length greater than 1, if the characters at positions i and j ($s[i]$ and $s[j]$) match, then $dp[i][j] = dp[i+1][j-1] + 2$. If they do not match, we take the maximum of $dp[i+1][j]$ and $dp[i][j-1]$. Finally, the $dp[0][n-1]$ holds the result for the entire string.

#576 Out of Boundary Paths [Medium] · Dynamic Programming

Key Insights

- Use Dynamic Programming (DP) as the state is defined by (current row, current column, moves remaining).

- Transitions are made from a cell to its four adjacent cells.
- If the ball moves out of bounds, it contributes to one valid path.
- States can be memoized to avoid recomputation and optimize the recursive solution.

Space & Time

Time Complexity: $O(\text{maxMove} * m * n)$, since we compute each state at most once.

Space Complexity: $O(\text{maxMove} * m * n)$ for the memoization table.

Solution

We use a recursive dynamic programming (or memoization) approach. Define a function $\text{dp}(i, j, \text{movesLeft})$ that returns the number of ways to move out of bounds starting from cell (i, j) with movesLeft moves remaining. For each move from the current cell, check:

- If the cell is out of bounds, return 1 (base case).
- If no moves are left ($\text{movesLeft} == 0$) and we are still inside, return 0.
- Otherwise, use recursion to explore all four directions and sum the paths. Store results in a memoization table to avoid recomputation, then return the computed value modulo $10^9 + 7$.

#1143 Longest Common Subsequence [Medium] · Dynamic Programming

Round 9 · Low · Hard

p.3421

Key Insights

- The problem is well-suited for a dynamic programming approach.
- Use a 2D table where $\text{dp}[i][j]$ stores the length of the longest common subsequence for $\text{text1}[0\dots i-1]$ and $\text{text2}[0\dots j-1]$.
- If the characters match at the current position, extend the subsequence by 1 from $\text{dp}[i-1][j-1]$; otherwise, take the maximum from $\text{dp}[i-1][j]$ or $\text{dp}[i][j-1]$.
- Filling the table in a bottom-up manner ensures that all subproblems are solved optimally.

Space & Time

Time Complexity: $O(m * n)$, where m and n are the lengths of text1 and text2 .

Space Complexity: $O(m * n)$ for maintaining the 2D DP table. Space can be optimized to $O(\min(m, n))$ if needed.

Solution

This solution uses a dynamic programming approach with a 2D array (dp) to record the lengths of common subsequences at each substring pair. We initialize a table of dimensions $(m+1) \times (n+1)$ with zeros, where m and n are the lengths of text1 and text2 . We iterate over both strings, comparing characters; if they match, we set $\text{dp}[i][j] = \text{dp}[i-1][j-1] + 1$, otherwise we choose the maximum from $\text{dp}[i-1][j]$ and $\text{dp}[i][j-1]$. The final answer is found at $\text{dp}[m][n]$.

Round 9 · Low · Hard

p.3422

#78 Subsets [Medium] · Bit Manipulation

Key Insights

- The problem is equivalent to generating the power set of the given set.
- A recursive backtracking approach can be used to explore both including and excluding each element.
- Iterative (bit manipulation) solutions are also possible since there are 2^n subsets in total.
- The maximum array length is small (up to 10), making exponential time solutions feasible.

Space & Time

Time Complexity: $O(2^n * n)$ – There are 2^n subsets and creating each subset may take up to $O(n)$ time.

Space Complexity: $O(2^n * n)$ – Space for storing all subsets. The recursion stack has a maximum depth of $O(n)$.

Solution

We solve the problem using backtracking, where we decide for each element whether to include it in the current subset or not. We recursively build all subsets starting at an empty list and at each step appending the current subset to our result list. We can also use iterative bit manipulation but here we focus on the backtracking approach. This algorithm leverages recursion and a list to maintain the current state, leading to a clear and concise implementation.

#90 Subsets II [Medium] · Bit Manipulation

Key Insights

- Sort the array first so that duplicate elements are adjacent.
- Use backtracking to generate all possible subsets.
- Skip duplicate elements in the recursion to avoid forming duplicate subsets.
- The use of sorted order and conditionally skipping elements is the trick to handle duplicates.

Space & Time

Time Complexity: $O(2^n * n)$ since each subset creation may take $O(n)$ in the worst case.

Space Complexity: $O(2^n * n)$ for storing the subsets, plus $O(n)$ for recursion stack.

Solution

The solution uses a backtracking technique that builds subsets incrementally. The input array is first sorted to bring duplicates together. During recursion, if the current element is the same as the previous one and it has not been used in the current recursion path, it is skipped to prevent duplicate subsets. This ensures that each subset is unique. The algorithm stores each valid subset by copying the current subset state into the result list.

#287 Find the Duplicate Number [Medium] · Bit Manipulation

Key Insights

- Since there are $n+1$ elements with values only from 1 to n , by the pigeonhole principle at least one number must be repeated.
- The problem can be approached by mapping the array values to indices, forming a cycle structure.
- Floyd's Tortoise and Hare algorithm (cycle detection) can be applied to find the duplicate number in linear time and constant space.
- There is no need for extra data structures like hash sets or alterations to the original array, satisfying the space and array-modification constraints.

Space & Time

Time Complexity: $O(n)$

Space Complexity: $O(1)$

Solution

The solution leverages Floyd's Tortoise and Hare cycle detection algorithm. Think of each array element as a pointer to another index, forming a linked list with a cycle (the duplicate number creates the cycle). The algorithm works in two phases:

- Detect a meeting point within the cycle by having two pointers (tortoise and hare) move at different speeds. – Find the cycle entrance by resetting one pointer to the

start and moving both pointers at the same speed. The meeting point reveals the duplicate number.

This approach provides an elegant solution that satisfies both linear runtime and constant space requirements.

#1506 Find Root of N-Ary Tree [Medium] · Bit Manipulation

Key Insights

- Every node except the root appears as a child of another node.
- By summing all node values and then subtracting the sum of all children node values, the remaining value corresponds to the root.
- This subtraction trick allows us to identify the root in a single pass of the data structure without extra memory for tracking child nodes.

Space & Time

Time Complexity: $O(n)$ where n is the number of nodes, since we iterate through all nodes and their children.

Space Complexity: $O(1)$ extra space using arithmetic accumulators.

Solution

The solution calculates two sums: one for all nodes' values and one for all children nodes' values. The difference between these gives the value of the root

node since every node's value (except the root's) is added once as a child. Finally, we iterate over the nodes to find the node with the computed root value and return it. This approach meets the requirement of constant space complexity and linear time.

#7 Reverse Integer [Medium] · Math

Key Insights

- The problem requires reversing the digits while preserving the sign.
- Watch out for integer overflow due to the reversed number exceeding the 32-bit signed integer limits.
- The approach involves repeatedly extracting the last digit from the number and appending it to a new reversed integer.
- Overflow checks are done during each iteration before actually appending the digit.

Space & Time

Time Complexity: $O(n)$, where n is the number of digits in the integer.

Space Complexity: $O(1)$, only constant extra space is used.

Solution

The solution uses a while loop to extract the last digit from the original number using modulo operation ($x \% 10$) and then removes this digit using integer division (x

$/= 10$). A reversed number is constructed by multiplying the current reversed number by 10 and adding the extracted digit. Before appending each digit, the algorithm checks if the new reversed number will overflow a 32-bit signed integer by comparing it with the thresholds derived from `INT_MAX` and `INT_MIN`. The algorithm handles negative values by working with x directly since the modulo and division operations work properly with negatives in most languages.

#50 Pow(x, n) [Medium] · Math

Key Insights

- Use exponentiation by squaring to reduce the number of multiplications.
- For a negative exponent, compute the power for $|n|$ and then take the reciprocal.
- Handle the edge case when n is 0 (return 1).
- The algorithm works in $O(\log n)$ time complexity.

Space & Time

Time Complexity: $O(\log |n|)$

Space Complexity: $O(1)$ for the iterative solution, or $O(\log |n|)$ if implemented recursively

Solution

We use exponentiation by squaring, an algorithm that reduces the number of multiplications by repeatedly squaring the base and halving the exponent. Here's how

it works:

– If n is negative, switch x to $1/x$ and n to $-n$. – Initialize the result as 1. – Loop while n is greater than 0: If the current n is odd, multiply the result by x . Square the base x . Divide n by 2 (using integer division). This iterative approach efficiently computes the power in logarithmic steps. The same approach can be implemented recursively.

#380 Insert Delete GetRandom O(1) [Medium] · Math

Key Insights

- Use a dynamic array (or list) to store the elements to allow $O(1)$ access by index.
- Use a hash table (or dictionary/map) to store the mapping from value to its index in the array.
- For removal, swap the element to be removed with the last element in the array, then delete the last element. This avoids expensive shifting of elements.
- `getRandom` can be implemented by generating a random index into the array.

Space & Time

Time Complexity: $O(1)$ on average for insert, remove, and `getRandom` operations.

Space Complexity: $O(n)$ where n is the number of elements stored in the data structure.

Solution

The solution involves maintaining a list to store the values and a hash map to quickly check for the presence of a value and to know its index in the list. When an element is removed, swap it with the last element of the list and update the map accordingly; then, remove the last element in $O(1)$ time. This approach ensures that insert, remove, and `getRandom` can all be achieved in constant time on average.

#1017 Convert to Base -2 [Medium] · Math

Key Insights

- Use division and remainder operations to simulate converting a number to a different base.
- When working with a negative base (-2), the remainder can become negative; adjust the remainder (by adding the base's absolute value) and update the quotient accordingly.
- For $n = 0$, simply return "0" since there is nothing to convert.
- Build the answer by collecting digits (remainders) and then reversing the order since the conversion provides digits from least significant to most significant.

Space & Time

Time Complexity: $O(\log|n|)$ – the loop runs for about the number of digits in the base -2 representation.

Space Complexity: $O(\log|n|)$ – extra space for storing the computed digits.

Solution

To solve the problem, we simulate the conversion process similar to how you would convert an integer to a positive base. However, since the base here is -2 , special handling of negative remainders is needed. In each iteration, we:

- Divide the current number by -2 . – Calculate the remainder. If the remainder is negative, adjust it by adding 2, and increment the quotient to compensate. – Append the computed remainder to a list (which represents the binary digit). – Continue until the number reduces to zero. Finally, reverse the list to form the correct string representation from the most significant to the least significant digit.

#3160 Find the Number of Distinct Colors Among the Balls [Medium] · Math

Key Insights

- Use a hash map to store the current color of each ball that has been updated.
- Use another hash map to store how many balls currently use each color.
- When a ball is recolored, decrement the count of its old color first.

Round 9 · Low · Hard

p.3431

- If an old color count becomes zero, remove it from the color–frequency map.
- Add the new color to the ball and increment its frequency.
- After each query, the number of distinct colors is the number of keys in the color–frequency map.

Space & Time

Time Complexity: $O(q)$ average for q queries, since each query does only $O(1)$ average–time hash map work.

Space Complexity: $O(\min(q, n))$ for colored balls, plus $O(q)$ in the worst case for distinct active colors.

Solution

We solve the problem using two hash maps. The first map stores the current color of each ball, and the second map stores the frequency of each active color. For every query, if the ball already has a color, we decrement the old color's count and remove that color from the frequency map if its count reaches zero. Then we update the ball with the new color and increment the new color's frequency. After processing the query, the number of distinct colors is simply the number of keys remaining in the frequency map.

#2627 Debounce [Medium] · JavaScript

Key Insights

- Each call to the debounced function resets the timer.

Round 9 · Low · Hard

p.3432

- Only the most recent function call (that isn't subsequently cancelled) will eventually execute.
- Passing of parameters is handled by storing the arguments of the latest call.
- The solution leverages timers (or equivalent scheduling mechanisms) and cancellation techniques.

Space & Time

Time Complexity: $O(1)$ per call, as each invocation only clears and sets a timer.

Space Complexity: $O(1)$, since only a single timer reference is maintained.

Solution

The solution creates a wrapper function that maintains an internal timer reference. On each call to the debounced function, if a timer exists, it is cancelled. A new timer is then set to execute the original function after t milliseconds with the provided arguments. This ensures that only the last invocation runs if multiple calls occur within the delay period.

#2628 JSON Deep Equal [Medium] · JavaScript

Key Insights

- Use recursion to compare nested structures.
- For primitive types, a simple equality check ($===$ in JavaScript, $==$ in Python, etc.) suffices.

Round 9 · Low · Hard

p.3433

- For arrays, verify both have the same length and compare each element recursively.
- For objects, ensure both have the same keys (order does not matter) then recursively compare the corresponding values.
- Handle cases where one operand is null or differs by type early to avoid unnecessary recursion.

Space & Time

Time Complexity: $O(N)$ in the worst-case, where N is the total number of elements/values in the JSON structure.

Space Complexity: $O(N)$ due to the recursion call stack in the worst-case scenario of deeply nested objects or arrays.

Solution

We solve this problem using a recursive approach. The algorithm first checks if both values are strictly equal, which handles primitives immediately. For arrays and objects, it verifies that the sizes (length for arrays and number of keys for objects) match. Then it recurses into each element (for arrays) or checks each key-value pair (for objects) to ensure that they are also deeply equal. For objects, the order of keys is irrelevant, so we compare key sets. This method naturally handles nested structures by delegating equality checks to the same function.

#2630 Memoize [Medium] · JavaScript

Round 9 · Low · Hard

p.3434

Key Insights

- Cache previously computed results using a data structure keyed by the function arguments.
- Ensure the key uniquely represents the argument order; simple tuple (or equivalent) suffices.
- Use closures (or object state) to maintain both the cache and a call count tracking actual function invocations.
- Expose a method (getCallCount) on the memoized function to report the number of times the original function was actually called.
- The memoized version works for functions that might take one or more parameters.

Space & Time

Time Complexity: $O(1)$ average for each call assuming efficient hashing of arguments.

Space Complexity: $O(n)$ for caching n distinct input combinations.

Solution

We use a hash-based cache (a dictionary or map) to store results for every unique set of arguments. Upon each call, we check if the key (constructed from the input arguments) exists in the cache. If not, we increment a counter, call the original function, store the result in the cache, and return it. Otherwise, we return the cached result. The memoized function is augmented

with a getCallCount method to report the number of actual calls made to the original function.

#2633 Convert Object to JSON String [Medium] · JavaScript

Key Insights

- The solution must support all JSON primitives: strings, numbers, booleans, and null.
- Composite types such as arrays and objects require recursive traversal for proper serializing.
- For strings, wrap them in double quotes.
- For arrays, recursively process every element and join them with commas between square brackets.
- For objects, iterate keys in insertion order and combine key-value pairs (key must be a string in double quotes) with commas between curly braces.
- Use recursion carefully to handle nested structures without adding unnecessary spaces.

Space & Time

Time Complexity: $O(n)$, where n is the total number of elements/characters in the JSON structure.

Space Complexity: $O(n)$ due to the recursion stack and the constructed output string.

Solution

We use a recursive approach that inspects the type of the value and processes it accordingly:

– If the value is null, output "null". – If the value is a boolean or number, convert it directly to its string representation. – For a string, ensure it is enclosed in double quotes. – For an array, iterate over each element, serialize each recursively, join them with commas, and enclose in square brackets. – For an object, iterate over its keys in insertion order (or using Object.keys in JS / LinkedHashMap in Java), serialize both key (wrapped in double quotes) and its corresponding value, join with commas, and enclose in curly braces. This approach leverages recursion to handle arbitrarily nested objects and arrays while ensuring there are no extra spaces in the final string.

#2636 Promise Pool [Medium] · JavaScript

Key Insights

- Use a concurrency control mechanism to limit the number of actively running promises.
- Maintain an index or pointer to track which function from the input array should be executed next.
- Upon the resolution of any promise, start a new one if available until all functions have been processed.
- The overall task is complete when there are no functions left and all running promises have resolved.

Space & Time

Time Complexity: $O(m)$, where m is the number of functions to execute.

Space Complexity: $O(n)$ for the active pool of concurrent promises.

Solution

We can solve the problem by maintaining an index that tracks which asynchronous function to execute next. We define a helper routine (or executor) that:

- Checks if there is a function left to run. – If yes, increments the index and awaits the result of running that function. – Once the function completes, the routine recursively calls itself to pick up the next function. We start n such routines (to respect the pool limit) and use Promise.all (or the equivalent in other languages) to wait until all concurrent routines have finished executing. This structure ensures that we never run more than n promises concurrently and that a new promise starts as soon as one finishes.

#2675 Array of Objects to Matrix [Medium] · JavaScript

Key Insights

- Flatten each object or array into dot-separated paths such as `a.b` or `a.0.c`.
- Use DFS so nested objects and nested arrays are handled uniformly.
- For each row, store a map from flattened path to value.

- Collect every path seen across all rows into a set of column names.
- Sort the column names lexicographically to build the header row.
- For each flattened row map, output values in header order and use an empty string when a column is missing.

Space & Time

Time Complexity: $O(T + R * C \log C)$, where T is the total number of nested nodes visited, R is the number of rows, and C is the number of unique flattened columns.

Space Complexity: $O(R * C + C)$ in the worst case for the flattened row maps and the set of all column names.

Solution

We solve the problem by flattening every input row first. A DFS walks through objects and arrays, extending the current path with either the property name or the array index. When the DFS reaches a primitive value or `null`, it stores that value under the built path. While flattening each row, we also collect every path into a global set of column names. After all rows are processed, we sort the columns lexicographically, place them in the first row of the matrix, and then build each remaining row by reading values from that row's flattened map, using `""` for missing paths.

#2700 Differences Between Two Objects [Medium] · JavaScript

Key Insights

- Use recursion to traverse nested objects and arrays.
- Compare only keys that exist in both objects.
- When encountering two values of different types or unequal primitives, record the difference as [value from obj1, value from obj2].
- When both values are objects or arrays, recursively compute their differences and include them only if the resulting nested diff is non-empty.
- Arrays are treated like objects with indices as keys.
- The solution requires careful type checking and handling of recursive structures.

Space & Time

Time Complexity: $O(n)$ in average, where n is the number of keys/elements in the common parts of the structures.

Space Complexity: $O(n)$ due to recursion and storage needed for the output differences.

Solution

We solve the problem using a recursive function that compares keys present in both input objects. For each common key:

- If both values are objects (or both are arrays), recursively compute their differences.
- If the recursive

call returns a non-empty object, include that key in the result. – If the two values are of different types or are primitives that differ, add the key with a difference array [value from obj1, value from obj2]. – Ignore keys that exist only in one object. This recursive approach naturally handles the deeply nested structure while ensuring only common keys with different values are captured.

Round 9 · Low · Hard

p.3441

R9 · Low · Hard (7 q)

>> Round 9 · Low Priority · Hard — Quick Review

7 questions · key insights · complexity · solution

#10 Regular Expression Matching [Hard] · Dynamic Programming

Key Insights

- Use recursion or dynamic programming to simulate matching as the '*' operator can affect the current and following characters.
- The '.' operator is a wildcard that matches any single character.
- When encountering a '*' in the pattern, consider both skipping the preceding element or consuming one character from the string if it matches.
- Using memoization helps to save intermediate results and efficiently process overlapping subproblems.

Space & Time

Time Complexity: $O(mn)$

where $m = \text{length}(s)$ and $n = \text{length}(p)$, due to exploring each subproblem combination.

Space Complexity: $O(m)$

Round 9 · Low · Hard

p.3442

n) for storing the memoization table in the dynamic programming or recursive approach.

Solution

The solution uses a top-down recursive strategy with memoization to efficiently match the string against the pattern. For each character position in s and p , we check if they match directly or if the pattern character is '.' to allow any match. When encountering '.' in the pattern, two possibilities are explored: either the '.' stands for zero occurrences of the preceding element, or if there is a valid match, it accounts for one occurrence and we proceed further in the string while maintaining the same pattern index. The memoization aspect prevents repeated computations of the same (i, j) state in the recursion, resulting in an improved overall runtime.

#115 Distinct Subsequences [Hard] · Dynamic Programming

Key Insights

- Use dynamic programming to build the solution by comparing characters of s and t .
- Define a 2D dp table where $dp[i][j]$ represents the number of ways to form the first j characters of t using the first i characters of s .
- The recurrence relation:
- If $s[i-1]$ equals $t[j-1]$, then include both the cases of matching this character as well as skipping it: $dp[i][j] =$

$dp[i-1][j-1] + dp[i-1][j]$.

- If they do not match, then simply skip the current character of s : $dp[i][j] = dp[i-1][j]$.
- Initialize $dp[i][0] = 1$ for all i since an empty t can always be formed.
- The final answer is $dp[s.length][t.length]$.

Space & Time

Time Complexity: $O(n * m)$, where n is the length of s and m is the length of t .

Space Complexity: $O(n * m)$, which can be optimized to $O(m)$ with careful row reuse.

Solution

We use a dynamic programming approach with a 2D table. The table has dimensions $(len(s) + 1) \times (len(t) + 1)$. Each cell $dp[i][j]$ stores the number of distinct subsequences from the first i characters of s that match the first j characters of t .

Steps:

- Initialize $dp[i][0] = 1$ for every i , because an empty t is a subsequence of any prefix of s .
- Iterate over each character of s (outer loop) and for each, iterate over t (inner loop).
- For every character in s , if it equals the current character in t , update $dp[i][j]$ as the sum of the ways by using or skipping this character. Otherwise, carry over the previous count.
- Return $dp[len(s)][len(t)]$ as the result.

#123 Best Time to Buy and Sell Stock III [Hard] · Dynamic Programming

Key Insights

- The problem requires the optimal selection of at most two buy-sell transactions.
- One effective strategy is to split the problem into two parts:
 - For each day i , calculate the maximum profit obtainable from day 0 to i with one transaction.
 - For each day i , calculate the maximum profit obtainable from day i to the end with one transaction.
- Finally, for each day, the sum of these two profits (left and right) gives a candidate answer.
- Alternatively, a state-based dynamic programming approach can track 4 states corresponding to the two transactions.

Space & Time

Time Complexity: $O(n)$ where n is the number of days

Space Complexity: $O(n)$ due to extra arrays used for forward and backward computations (this can be optimized to $O(1)$ with a state variable approach)

Solution

We use a two-pass dynamic programming approach. In the first pass (forward), we compute an array "leftProfits" where $\text{leftProfits}[i]$ is the maximum profit from day 0 to day i with one transaction. We do this by

keeping track of the minimum price observed so far and calculating the profit if sold on day i . In the second pass (backward), we compute an array "rightProfits" where $\text{rightProfits}[i]$ is the maximum profit obtainable from day i to the last day with one transaction by tracking the maximum price seen in reverse order. Finally, for every index i , the sum $\text{leftProfits}[i] + \text{rightProfits}[i]$ represents the maximum profit achievable by completing the first transaction on or before day i and the second transaction starting from day i . The answer is the maximum sum over all i .

#132 Palindrome Partitioning II [Hard] · Dynamic Programming

Key Insights

- Precompute a 2D table (isPal) where $\text{isPal}[i][j]$ indicates if the substring $s[i...j]$ is a palindrome.
- Use dynamic programming: $\text{dp}[i]$ represents the minimum cuts needed for the substring $s[0...i]$.
- For every index i , iterate over all possible starting indices j . If $s[j...i]$ is a palindrome, update $\text{dp}[i]$ using $\text{dp}[j-1] + 1$ (or 0 if j is 0).
- Time complexity is dominated by the $O(n^2)$ palindrome checking and state transitions.

Space & Time

Time Complexity: $O(n^2)$

Space Complexity: $O(n^2)$ (due to the 2D table for palindrome checks, plus $O(n)$ for the dp array)

Solution

The solution relies on dynamic programming with a two-step approach. First, precompute a 2D table (isPal) where each cell indicates if a substring is a palindrome. This is done in $O(n^2)$ time by leveraging the fact that a substring $s[j...i]$ is a palindrome if the characters at the ends match and the inner substring $s[j+1...i-1]$ is also a palindrome (or if the length is less than 3). Then, use a dp array where $dp[i]$ holds the minimum number of cuts needed for the substring $s[0...i]$. For each i , check every index j ($0 \leq j \leq i$) and if $s[j...i]$ is a palindrome, update $dp[i]$ to be the minimum of its current value and $dp[j-1] + 1$ (with a special case when j is 0, meaning no cut is needed). This two-phase approach efficiently breaks down the problem and computes the answer while avoiding redundant palindrome checks.

#312 Burst Balloons [Hard] · Dynamic Programming

Key Insights

- Transform the problem by adding virtual balloons with value 1 at both ends to simplify edge handling.
- Use interval dynamic programming where $dp[i][j]$ represents the maximum coins that can be obtained by bursting all balloons between indices i and j (exclusive).

- Realize that the order of bursting balloons matters; choose the last balloon to burst in each interval to break dependencies into independent subproblems.
- The solution involves iterating over all possible intervals and determining the optimal burst order.

Space & Time

Time Complexity: $O(n^3)$ – Three nested loops: one for the interval length, one for the starting index, and one for the final balloon choice.

Space Complexity: $O(n^2)$ – A 2D dynamic programming table is used to store the maximum coins for each interval.

Solution

The problem is solved using interval dynamic programming. First, extend the input array by adding 1 at both ends to handle edge cases seamlessly. Define a dp table where $dp[i][j]$ represents the maximum coins that can be obtained by bursting all the balloons between i and j (not including i and j , as these are the virtual boundaries). For every interval $[i, j]$, iterate through the possible choices for the last balloon k to burst in that interval. The value for bursting balloon k will be determined by the coins from bursting it last (i.e., $nums[i] * nums[k] * nums[j]$) plus the maximum coins from the two resulting subintervals: (i, k) and (k, j) . By filling the dp table in increasing interval lengths, you build up to the solution for the overall interval that

covers the entire (extended) array.

#1000 Minimum Cost to Merge Stones [Hard] · Dynamic Programming

Key Insights

- You can only merge stones into one final pile if $(n - 1)$ is divisible by $(k - 1)$. If not, return -1 .
- A dynamic programming approach is used where $dp[i][j]$ represents the minimum cost to merge the subarray $stones[i..j]$ into as few piles as possible.
- A prefix sum array is precomputed to quickly calculate the sum of stones in any subarray.
- When a segment can be merged into one pile (i.e. when $(length - 1) \% (k - 1) == 0$), add the sum of stones in that segment to the dp value.

Space & Time

Time Complexity: $O(n^3)$ due to the triple nested loop over segments and splits.

Space Complexity: $O(n^2)$ from the dp table and $O(n)$ for the prefix sum array.

Solution

The problem is solved using a range DP approach. We first verify that merging all piles into one is feasible by checking if $(n - 1) \% (k - 1) == 0$. Next, a prefix sum array is computed to optimize range sum calculations. A 2D dp table $dp[i][j]$ is then used to store the minimum

cost to merge $stones[i..j]$. For each interval, we consider valid split points, updating $dp[i][j]$ by combining sub-problems. If the current segment can form a single merged pile (determined by the condition on the interval length), the cost of merging that segment (the sum of stones) is added to $dp[i][j]$.

#68 Text Justification [Hard] · Math

Key Insights

- Use a greedy approach to fill each line with as many words as possible without exceeding `maxWidth`.
- Calculate total space to distribute by subtracting the combined words length from `maxWidth`.
- If the line has more than one word, distribute extra spaces evenly; if not, pad the line with spaces at the end.
- Handle the last line separately by left-justifying the text (i.e., adding a single space between words and padding the end).

Space & Time

Time Complexity: $O(n)$ where n is the number of words, since each word is processed once.

Space Complexity: $O(n)$ for storing the result and temporary buffers.

Solution

The solution iterates over the list of words and groups them into lines such that the combined length of the

words and minimum required spaces does not exceed `maxWidth`. For each completed group (line):

- If it's not the last line, calculate extra spaces required.
- If the line has multiple words, distribute the spaces evenly, adding any extra spaces from left to right.
- If the line contains a single word, append spaces to the right.
- For the last line, join words with a single space and pad the remaining width with spaces.

Data structures used include lists (or arrays) for current words accumulation and the final result. The logic is implemented using a simulation approach to mimic text formatting.
